

WRITTEN BY
MICHAEL ALEXANDER
DICK KUSLEIKA

Microsoft®
Excel 365
BIBLE

THE COMPREHENSIVE TUTORIAL RESOURCE

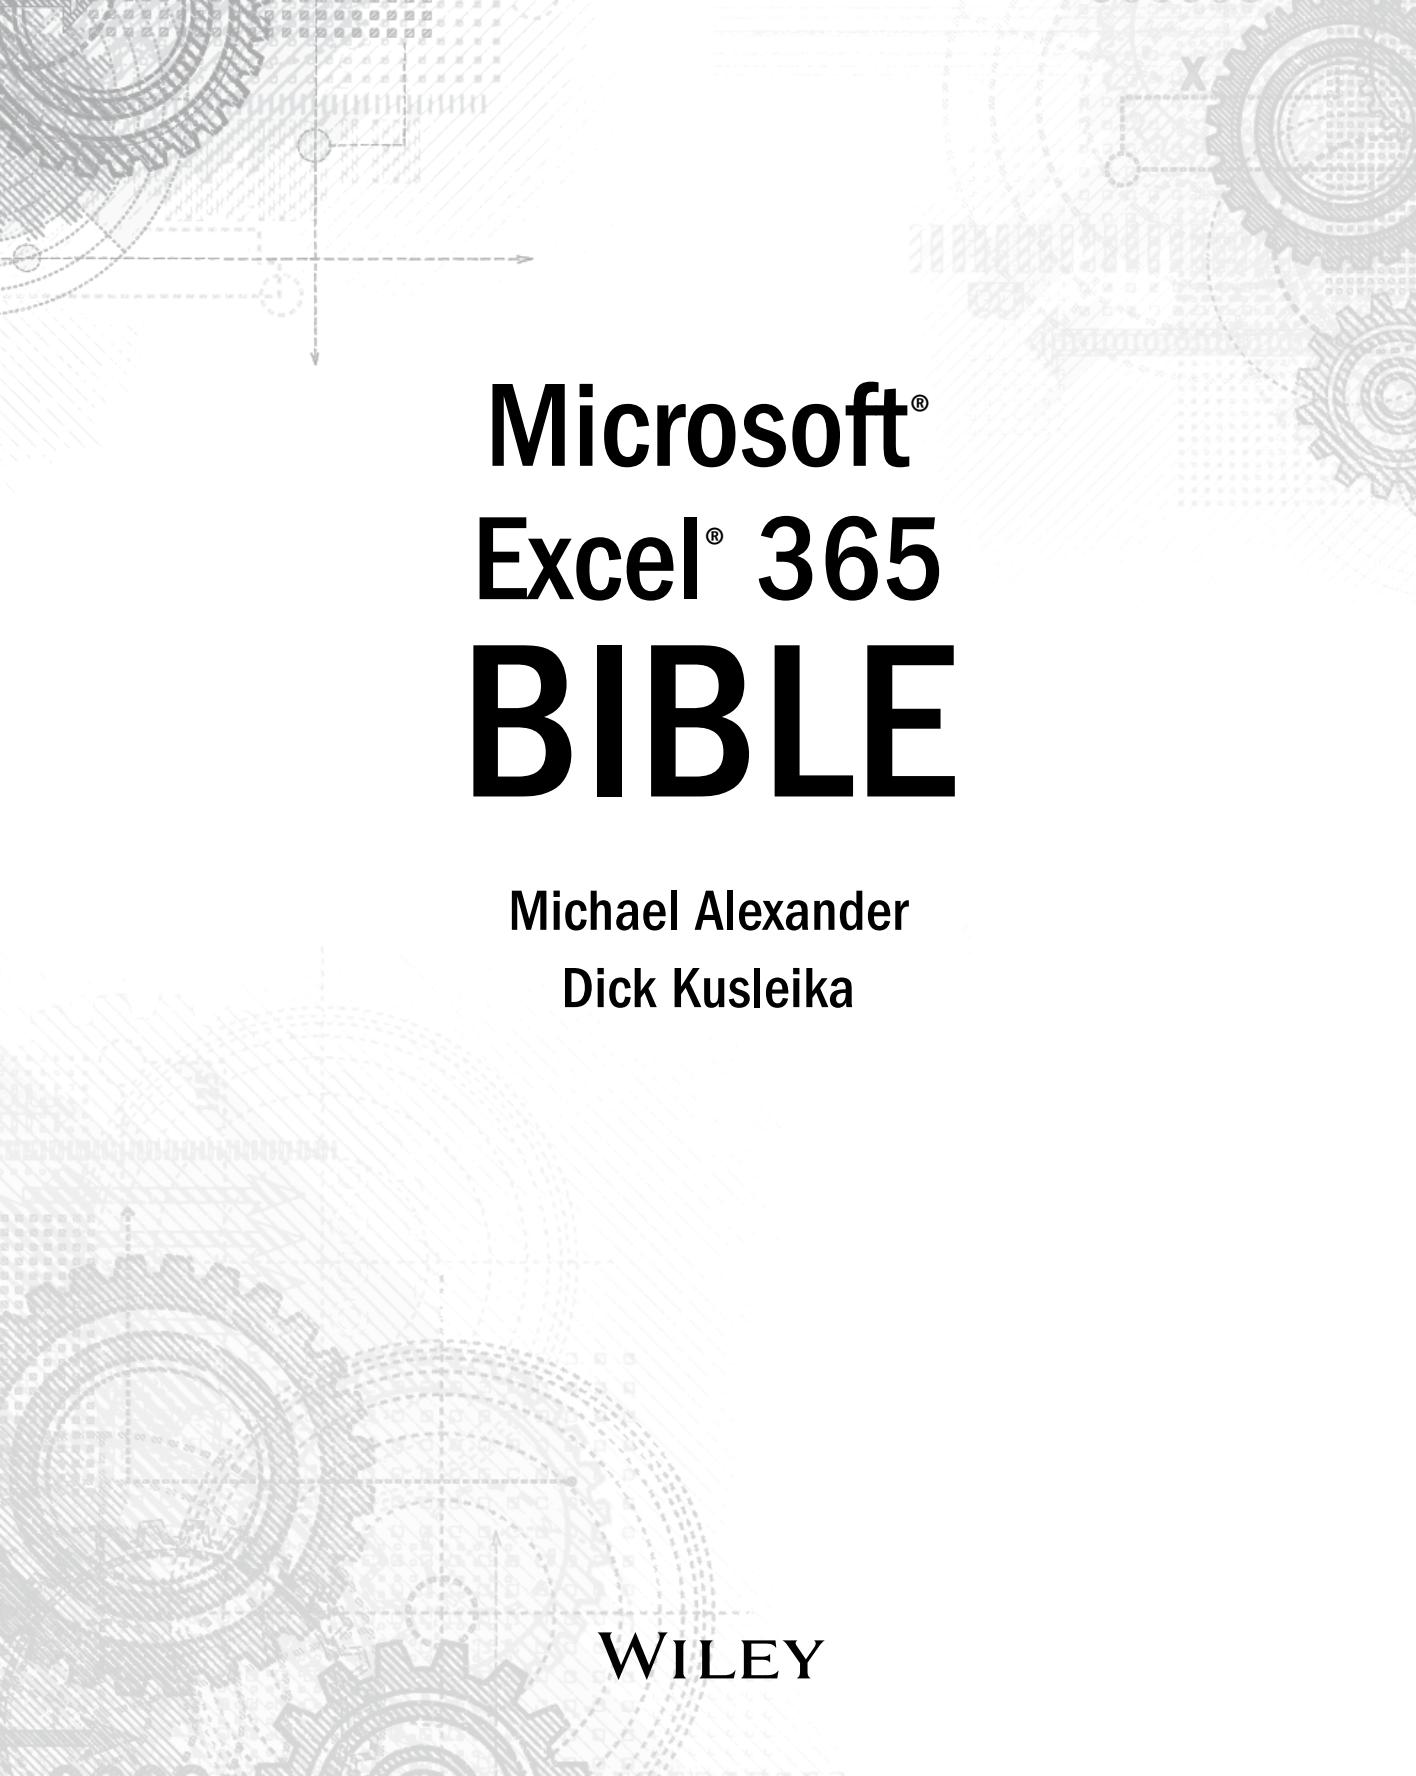
EXPLORE THE LATEST
TOOLS AND FEATURES

MASTER INSIDE TRICKS
AND BEST PRACTICES

DISCOVER WHAT YOU
NEED, WHEN YOU NEED IT

WILEY

Microsoft®
Excel® 365
Bible



Microsoft® Excel® 365 BIBLE

Michael Alexander

Dick Kusleika

WILEY

Copyright © 2022 by John Wiley & Sons, Inc. All rights reserved.

Published by John Wiley & Sons, Inc., Hoboken, New Jersey.

Published simultaneously in Canada and the United Kingdom.

978-1-119-83510-3

978-1-119-83522-6 (ebk.)

978-1-119-83523-3 (ebk.)

No part of this publication may be reproduced, stored in a retrieval system, or transmitted in any form or by any means, electronic, mechanical, photocopying, recording, scanning, or otherwise, except as permitted under Section 107 or 108 of the 1976 United States Copyright Act, without either the prior written permission of the Publisher, or authorization through payment of the appropriate per-copy fee to the Copyright Clearance Center, Inc., 222 Rosewood Drive, Danvers, MA 01923, (978) 750-8400, fax (978) 750-4470, or on the web at www.copyright.com. Requests to the Publisher for permission should be addressed to the Permissions Department, John Wiley & Sons, Inc., 111 River Street, Hoboken, NJ 07030, (201) 748-6011, fax (201) 748-6008, or online at <http://www.wiley.com/go/permission>.

Limit of Liability/Disclaimer of Warranty: The publisher and the author make no representations or warranties with respect to the accuracy or completeness of the contents of this work and specifically disclaim all warranties, including without limitation warranties of fitness for a particular purpose. No warranty may be created or extended by sales or promotional materials. The advice and strategies contained herein may not be suitable for every situation. This work is sold with the understanding that the publisher is not engaged in rendering legal, accounting, or other professional services. If professional assistance is required, the services of a competent professional person should be sought. Neither the publisher nor the author shall be liable for damages arising herefrom. The fact that an organization or Website is referred to in this work as a citation and/or a potential source of further information does not mean that the author or the publisher endorses the information the organization or Website may provide or recommendations it may make. Further, readers should be aware the Internet Websites listed in this work may have changed or disappeared between when this work was written and when it is read.

For general information on our other products and services or for technical support, please contact our Customer Care Department within the United States at (800) 762-2974, outside the United States at (317) 572-3993 or fax (317) 572-4002.

Wiley also publishes its books in a variety of electronic formats. Some content that appears in print may not be available in electronic formats. For more information about Wiley products, visit our web site at www.wiley.com.

Library of Congress Control Number: 2021951349

Trademarks: WILEY and the Wiley logo are trademarks or registered trademarks of John Wiley & Sons, Inc. and/or its affiliates, in the United States and other countries, and may not be used without written permission. Microsoft and Excel are registered trademarks of Microsoft Corporation in the United States and/or other countries. All other trademarks are the property of their respective owners. John Wiley & Sons, Inc. is not associated with any product or vendor mentioned in this book.

Cover image: © aleksandarvelasevic/GettyImages

Cover design: Wiley

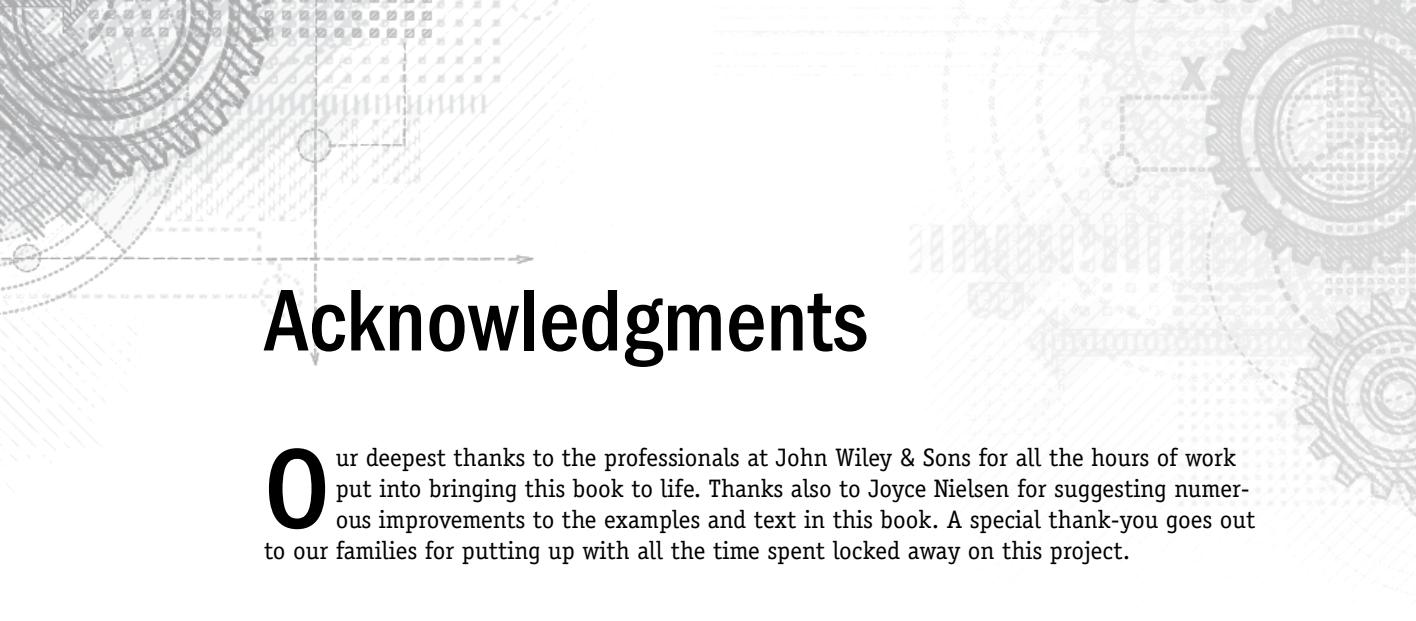
About the Authors

Michael Alexander is a senior consultant at Slalom Consulting with more than 15 years' experience in data management and reporting. He is the author of more than a dozen books on business analysis using Microsoft Excel and has been named Microsoft Excel MVP for his contributions to the Excel community.

Dick Kusleika is a 12-time Microsoft Excel MVP and has been working with Microsoft Office for more than 20 years. Dick develops Access- and Excel-based solutions for his clients and has conducted training seminars on Office products in the United States and Australia. Dick also writes a popular Excel-related blog at www.dailydoseofexcel.com.

About the Technical Editor

Joyce J. Nielsen has worked in the publishing industry for more than 30 years as a technical writer/editor, development editor, and project manager, specializing in Microsoft Office, Windows, Internet, and general technology content for leading educational, retail, and online publishers. She is the author or co-author of more than 50 computer books and 2,100 online articles. Joyce holds a bachelor of science degree in Quantitative Business Analysis from Indiana University's Kelley School of Business in Bloomington, Indiana. She currently resides in Arizona.



Acknowledgments

Our deepest thanks to the professionals at John Wiley & Sons for all the hours of work put into bringing this book to life. Thanks also to Joyce Nielsen for suggesting numerous improvements to the examples and text in this book. A special thank-you goes out to our families for putting up with all the time spent locked away on this project.

Contents at a Glance

Introduction	xxxix
Part I: Getting Started with Excel	1
Chapter 1: Introducing Excel	3
Chapter 2: Entering and Editing Worksheet Data	27
Chapter 3: Performing Basic Worksheet Operations	53
Chapter 4: Working with Excel Ranges and Tables.....	73
Chapter 5: Formatting Worksheets	121
Chapter 6: Understanding Excel Files and Templates	157
Chapter 7: Printing Your Work	177
Chapter 8: Customizing the Excel User Interface.....	199
Part II: Working with Formulas and Functions	209
Chapter 9: Introducing Formulas and Functions	211
Chapter 10: Understanding and Using Array Formulas.....	247
Chapter 11: Using Formulas for Common Mathematical Operations	271
Chapter 12: Using Formulas to Manipulate Text.....	285
Chapter 13: Using Formulas with Dates and Times	301
Chapter 14: Using Formulas for Conditional Analysis	329
Chapter 15: Using Formulas for Matching and Lookups	347
Chapter 16: Using Formulas with Tables and Conditional Formatting.....	365
Chapter 17: Making Your Formulas Error-Free.....	379
Part III: Creating Charts and Other Visualizations	399
Chapter 18: Getting Started with Excel Charts	401
Chapter 19: Using Advanced Charting Techniques.....	441
Chapter 20: Creating Sparkline Graphics.....	475
Chapter 21: Visualizing with Custom Number Formats and Shapes	489
Part IV: Managing and Analyzing Data	519
Chapter 22: Importing and Cleaning Data.....	521
Chapter 23: Using Data Validation	553
Chapter 24: Creating and Using Worksheet Outlines.....	567
Chapter 25: Linking and Consolidating Worksheets.....	577
Chapter 26: Introducing PivotTables	595

Contents at a Glance

Chapter 27: Analyzing Data with PivotTables	617
Chapter 28: Performing Spreadsheet What-If Analysis	651
Chapter 29: Analyzing Data Using Goal Seeking and Solver	675
Chapter 30: Analyzing Data with the Analysis ToolPak.....	697
Chapter 31: Protecting Your Work.....	709
Part V: Understanding Power Pivot and Power Query	721
Chapter 32: Introducing Power Pivot.....	723
Chapter 33: Working Directly with the Internal Data Model	747
Chapter 34: Adding Formulas to Power Pivot.....	757
Chapter 35: Introducing Power Query.....	777
Chapter 36: Transforming Data with Power Query	805
Chapter 37: Making Queries Work Together.....	837
Chapter 38: Enhancing Power Query Productivity	855
Part VI: Automating Excel	867
Chapter 39: Introducing Visual Basic for Applications	869
Chapter 40: Creating Custom Worksheet Functions.....	899
Chapter 41: Creating UserForms.....	913
Chapter 42: Using UserForm Controls in a Worksheet	935
Chapter 43: Working with Excel Events	949
Chapter 44: Seeing Some VBA Examples	963
Chapter 45: Creating Custom Excel Add-Ins.....	979
Index.....	989

Contents

Introduction	xxxix
<hr/>	
Part I: Getting Started with Excel	1
<hr/>	
Chapter 1: Introducing Excel	3
Understanding What Excel Is Used For.....	3
Understanding Workbooks and Worksheets.....	4
Moving around a Worksheet	4
Navigating with your keyboard	7
Navigating with your mouse	8
Using the Ribbon.....	9
Ribbon tabs.....	9
Contextual tabs.....	11
Types of commands on the Ribbon	12
Accessing the Ribbon by using your keyboard.....	12
Using Shortcut Menus	15
Customizing Your Quick Access Toolbar.....	16
Working with Dialog Boxes.....	18
Navigating dialog boxes	18
Using tabbed dialog boxes.....	19
Using Task Panes	20
Creating Your First Excel Workbook.....	20
Getting started on your worksheet.....	20
Filling in the month names	21
Entering the sales data	22
Formatting the numbers.....	23
Making your worksheet look a bit fancier.....	23
Summing the values	24
Creating a chart	25
Printing your worksheet.....	25
Saving your workbook	26
Chapter 2: Entering and Editing Worksheet Data	27
Exploring Data Types.....	27
Numeric values	27
Text entries	28

Contents

Formulas	29
Error values	29
Entering Text and Values into Your Worksheets.....	30
Entering numbers.....	30
Entering text.....	30
Using Enter mode.....	31
Entering Dates and Times into Your Worksheets.....	31
Entering date values.....	31
Entering time values.....	32
Modifying Cell Contents.....	32
Deleting the contents of a cell.....	32
Replacing the contents of a cell	33
Editing the contents of a cell.....	33
Learning some handy data-entry techniques	35
Automatically moving the selection after entering data	35
Selecting a range of input cells before entering data.....	36
Using Ctrl+Enter to place information into multiple cells simultaneously....	36
Changing modes	36
Entering decimal points automatically	36
Using AutoFill to enter a series of values.....	37
Using AutoComplete to automate data entry	37
Forcing text to appear on a new line within a cell	38
Using AutoCorrect for shorthand data entry.....	39
Entering numbers with fractions.....	40
Using a form for data entry.....	40
Entering the current date or time into a cell.....	41
Applying Number Formatting	42
Using automatic number formatting	43
Formatting numbers by using the Ribbon	43
Using shortcut keys to format numbers.....	43
Formatting numbers by using the Format Cells dialog box	45
Adding your own custom number formats	47
Using Excel on a Tablet.....	47
Exploring Excel's tablet interface.....	48
Entering formulas on a tablet	49
Introducing the Draw Ribbon.....	49
Chapter 3: Performing Basic Worksheet Operations	53
Learning the Fundamentals of Excel Worksheets	53
Working with Excel windows.....	53
Moving and resizing windows	54
Switching among windows	55
Closing windows	55
Activating a worksheet	56
Adding a new worksheet to your workbook	56

Deleting a worksheet you no longer need.....	57
Changing the name of a worksheet.....	57
Changing a sheet tab color	58
Rearranging your worksheets.....	58
Hiding and unhiding a worksheet	60
Controlling the Worksheet View.....	60
Zooming in or out for a better view.....	60
Viewing a worksheet in multiple windows	61
Comparing sheets side by side.....	62
Splitting the worksheet window into panes	63
Keeping the titles in view by freezing panes.....	63
Monitoring cells with a Watch Window.....	65
Working with Rows and Columns	66
Selecting rows and columns.....	66
Inserting rows and columns.....	66
Deleting rows and columns.....	68
Changing column widths and row heights	68
Changing column widths	69
Changing row heights.....	69
Hiding rows and columns	70
Chapter 4: Working with Excel Ranges and Tables	73
Understanding Cells and Ranges	73
Selecting ranges.....	74
Selecting complete rows and columns	75
Selecting noncontiguous ranges	75
Selecting multi-sheet ranges.....	76
Selecting special types of cells.....	79
Selecting cells by searching	81
Copying or Moving Ranges	83
Copying by using Ribbon commands	84
Copying by using shortcut menu commands	85
Copying by using shortcut keys.....	85
Copying or moving by using drag-and-drop	87
Copying to adjacent cells.....	88
Copying a range to other sheets	89
Using the Office Clipboard to paste.....	89
Pasting in special ways	91
Using the Paste Special dialog box	92
Performing mathematical operations without formulas.....	94
Skipping blanks when pasting.....	94
Transposing a range	94
Using Names to Work with Ranges	95
Creating range names in your workbooks	96
Using the Name box	96

Contents

Using the New Name dialog box.....	96
Using the Create Names from Selection dialog box	97
Managing names	99
Adding Comments to Cells.....	100
Showing comments.....	101
Replying to comments	102
Editing comments and replies	102
Deleting comments and replies	103
Resolving comment threads.....	103
Adding Notes to Cells	104
Showing notes	105
Formatting notes	106
Editing notes.....	108
Deleting notes	108
Working with Tables.....	108
Understanding a table's structure	108
The header row.....	109
The data body	109
The total row	109
The resizing handle.....	110
Creating a table	110
Adding data to a table	111
Sorting and filtering table data.....	111
Sorting a table	112
Filtering a table.....	114
Filtering a table with slicers.....	116
Changing the table's appearance	117
Chapter 5: Formatting Worksheets.....	121
Getting to Know the Formatting Tools	121
Using the formatting tools on the Home tab.....	122
Using the Mini toolbar	123
Using the Format Cells dialog box.....	124
Formatting Your Worksheet.....	124
Using fonts to format your worksheet.....	124
Changing text alignment.....	127
Choosing horizontal alignment options	127
Choosing vertical alignment options	129
Wrapping or shrinking text to fit the cell	129
Merging worksheet cells to create additional text space.....	129
Displaying text at an angle	131
Using colors and shading.....	131
Adding borders and lines.....	132
Using Conditional Formatting.....	135

Specifying conditional formatting	135
Using graphical conditional formats	135
Using data bars	135
Using color scales	137
Using icon sets.....	138
Creating formula-based rules	139
Understanding relative and absolute references.....	141
Conditional formatting formula examples.....	142
Identifying weekend days.....	142
Highlighting a row based on a value	142
Displaying alternate-row shading	143
Creating checkerboard shading	144
Shading groups of rows	144
Working with conditional formats	144
Managing rules.....	145
Copying cells that contain conditional formatting	145
Deleting conditional formatting	146
Locating cells that contain conditional formatting	146
Using Named Styles for Easier Formatting	146
Applying styles.....	147
Modifying an existing style	147
Creating new styles	149
Merging styles from other workbooks	150
Controlling styles with templates.....	150
Understanding Document Themes	150
Applying a theme.....	152
Customizing a theme	153
Chapter 6: Understanding Excel Files and Templates	157
Creating a New Workbook	157
Opening an Existing Workbook.....	158
Filtering filenames	160
Choosing your file display preferences	161
Saving a Workbook	161
Using AutoRecover.....	163
Recovering versions of the current workbook.....	164
Recovering unsaved work	164
Configuring AutoRecover.....	165
Password-Protecting a Workbook	165
Organizing Your Files	166
Other Workbook Info Options	166
Protect Workbook options	166
Check for Issues options	167
Version History	167

Contents

Manage Workbook options.....	167
Browser View options.....	168
Compatibility Mode section	168
Closing Workbooks	168
Safeguarding Your Work	168
Working with Templates	169
Exploring Excel templates	169
Viewing templates	169
Creating a workbook from a template.....	170
Modifying a template.....	172
Using default templates	172
Using the workbook template to change workbook defaults.....	173
Creating a worksheet template	174
Editing your template	174
Resetting the default workbook.....	174
Using custom workbook templates	174
Creating custom templates.....	174
Saving your custom templates.....	175
Using custom templates	176
Chapter 7: Printing Your Work	177
Doing Basic Printing	177
Changing Your Page View	179
Normal view	179
Page Layout view	180
Page Break Preview.....	181
Adjusting Common Page Setup Settings	183
Choosing your printer	184
Specifying what you want to print	184
Changing page orientation	185
Specifying paper size.....	185
Printing multiple copies of a report.....	186
Adjusting the page margins.....	186
Understanding page breaks.....	187
Inserting a page break.....	187
Removing manual page breaks.....	188
Printing row and column titles	188
Scaling printed output.....	189
Printing cell gridlines	189
Printing row and column headers	190
Using a background image.....	190
Adding a Header or a Footer to Your Reports	192
Selecting a predefined header or footer.....	192
Understanding header and footer element codes	192
Exploring other header and footer options	194

Exploring Other Print-Related Topics.....	194
Copying Page Setup settings across sheets.....	195
Preventing certain cells from being printed	195
Preventing objects from being printed	196
Creating custom views of your worksheet	197
Creating PDF files.....	198
Chapter 8: Customizing the Excel User Interface	199
Customizing the Quick Access Toolbar	199
About the Quick Access Toolbar	200
Adding new commands to the Quick Access Toolbar.....	201
Other Quick Access Toolbar actions.....	203
Customizing the Ribbon.....	204
Why you may want to customize the Ribbon	205
What can be customized	205
How to customize the Ribbon	205
Creating a new tab.....	206
Creating a new group	206
Adding commands to a new group	207
Resetting the Ribbon.....	208
Part II: Working with Formulas and Functions	209
Chapter 9: Introducing Formulas and Functions	211
Understanding Formula Basics	211
Using operators in formulas.....	213
Understanding operator precedence in formulas.....	214
Using functions in your formulas	216
Examples of formulas that use functions	216
Function arguments	217
More about functions	218
Entering Formulas into Your Worksheets	218
Entering formulas by pointing	220
Pasting range names into formulas	220
Inserting functions into formulas.....	221
Function entry tips	223
Editing Formulas.....	224
Using Cell References in Formulas	225
Using relative, absolute, and mixed references.....	225
Changing the types of your references	227
Referencing cells outside the worksheet.....	228
Referencing cells in other worksheets	228
Referencing cells in other workbooks	228
Introducing Formula Variables	229

Contents

Understanding the LET function.....	230
Formula variables in action.....	231
Using Formulas in Tables.....	232
Summarizing data in a table.....	232
Using formulas within a table	234
Referencing data in a table.....	235
Correcting Common Formula Errors.....	237
Handling circular references	238
Specifying when formulas are calculated.....	238
Using Advanced Naming Techniques.....	239
Using names for constants.....	240
Using names for formulas	240
Using range intersections	241
Applying names to existing references.....	243
Working with Formulas.....	244
Not hard-coding values	244
Using the Formula bar as a calculator	244
Making an exact copy of a formula.....	244
Converting formulas to values.....	245
Chapter 10: Understanding and Using Array Formulas	247
Understanding Legacy Array Formulas	248
Example of a legacy array formula.....	248
Editing legacy array formulas.....	249
Introducing Dynamic Arrays.....	250
Understanding spill ranges	252
Referencing spill ranges	254
Exploring Dynamic Array Functions	255
The SORT function	256
The SORTBY function	257
The UNIQUE function	258
The RANDARRAY function	259
The SEQUENCE function	260
The FILTER function	262
Using multiple conditions with the FILTER function.....	263
Filtering records that contain a search term	264
The XLOOKUP function.....	265
XLOOKUP with wildcards.....	268
Chapter 11: Using Formulas for Common Mathematical Operations	271
Calculating Percentages.....	271
Calculating percent of goal.....	271
Calculating percent variance.....	272
Calculating percent variance with negative values.....	273
Calculating a percent distribution	274

Calculating a running total.....	275
Applying a percent increase or decrease to values	276
Dealing with divide-by-zero errors.....	277
Rounding Numbers.....	278
Rounding numbers using formulas.....	279
Rounding to the nearest penny	279
Rounding to significant digits.....	280
Counting Values in a Range.....	282
Using Excel's Conversion Functions	283
Chapter 12: Using Formulas to Manipulate Text	285
Working with Text	285
Using Text Functions.....	286
Joining text strings.....	286
Setting text to sentence case.....	288
Removing spaces from a text string	289
Extracting parts of a text string	290
Finding a particular character in a text string.....	291
Finding the second instance of a character.....	292
Substituting text strings.....	293
Counting specific characters in a cell.....	294
Adding a line break within a formula.....	295
Cleaning strange characters from text fields	296
Padding numbers with zeros	297
Formatting the numbers in a text string	297
Using the DOLLAR function	299
Chapter 13: Using Formulas with Dates and Times.	301
Understanding How Excel Handles Dates and Times.....	301
Understanding date serial numbers.....	301
Entering dates	302
Understanding time serial numbers	303
Entering times	304
Formatting dates and times	305
Problems with dates	306
Excel's leap year bug	306
Pre-1900 dates	306
Inconsistent date entries.....	307
Using Excel's Date and Time Functions	307
Getting the current date and time.....	308
Calculating age	308
Calculating the number of days between two dates	309
Calculating the number of workdays between two dates.....	310
Using NETWORKDAYS.INTL	310

Contents

Generating a list of business days excluding holidays	311
Extracting parts of a date.....	313
Calculating number of years and months between dates.....	314
Converting dates to Julian date formats.....	315
Calculating the percent of year completed and remaining	316
Returning the last date of a given month.....	317
Using the EOMONTH function	318
Calculating the calendar quarter for a date.....	318
Calculating the fiscal quarter for a date	319
Returning a fiscal month from a date	320
Calculating the date of the Nth weekday of the month	321
Calculating the date of the last weekday of the month	322
Extracting parts of a time.....	323
Calculating elapsed time	324
Rounding time values	325
Converting decimal hours, minutes, or seconds to a time	326
Adding hours, minutes, or seconds to a time	326
Chapter 14: Using Formulas for Conditional Analysis	329
Understanding Conditional Analysis	329
Checking if a simple condition is met	329
Checking for multiple conditions.....	330
Validating conditional data	331
Looking up values	332
Checking if Condition1 AND Condition2 are met.....	333
Referring to logical conditions in cells	334
Checking if Condition1 OR Condition2 are met.....	335
Performing Conditional Calculations.....	336
Summing all values that meet a certain condition.....	336
Summing greater than zero.....	338
Summing all values that meet two or more conditions	339
Summing if values fall between a given date range	340
Using SUMIFS.....	341
Getting a count of values that meet a certain condition	341
Getting a count of values that meet two or more conditions.....	342
Finding nonstandard characters.....	343
Getting the average of all numbers that meet a certain condition	344
Getting the average of all numbers that meet two or more conditions	344
Chapter 15: Using Formulas for Matching and Lookups	347
Introducing Lookup Formulas.....	347
Leveraging Excel's Lookup Functions	348
Looking up an exact value based on a left lookup column.....	348
Looking up an exact value based on any lookup column	351
Looking up values horizontally	352

Hiding errors returned by lookup functions	353
Finding the closest match from a list of banded values	354
Finding the closest match with the INDEX and MATCH functions	356
Looking up values from multiple tables.....	357
Looking up a value based on a two-way matrix.....	359
Using default values for match.....	360
Finding a value based on multiple criteria.....	361
Returning text with SUMPRODUCT	362
Finding the last value in a column	362
Finding the last number using LOOKUP	363
Chapter 16: Using Formulas with Tables and Conditional Formatting	365
Highlighting Cells That Meet Certain Criteria.....	365
Highlighting cells based on the value of another cell	367
Highlighting Values That Exist in List1 but Not List2	369
Highlighting Values That Exist in List1 and List2.....	371
Highlighting Based on Dates	372
Highlighting days between two dates	374
Highlighting dates based on a due date.....	376
Chapter 17: Making Your Formulas Error-Free	379
Finding and Correcting Formula Errors.....	379
Mismatched parentheses	380
Cells are filled with hash marks	381
Blank cells are not blank	381
Extra space characters	382
Formulas returning an error.....	382
#DIV/0! errors	383
#N/A errors.....	383
#NAME? errors	384
#NULL! errors	384
#NUM! errors.....	384
#REF! errors	385
#SPILL! errors	385
#VALUE! errors	386
Operator precedence problems	386
Formulas are not calculated.....	387
Problems with decimal precision	387
“Phantom link” errors.....	388
Using Excel Auditing Tools	388
Identifying cells of a particular type.....	388
Viewing formulas	389
Tracing cell relationships	389
Identifying precedents	390
Identifying dependents	390

Contents

Tracing error values.....	391
Fixing circular reference errors	391
Using the background error-checking feature	391
Using Formula Evaluator	392
Searching and Replacing.....	393
Searching for information	393
Replacing information	394
Searching for formatting.....	395
Spell-checking your worksheets	396
Using AutoCorrect.....	396
Part III: Creating Charts and Other Visualizations	399
Chapter 18: Getting Started with Excel Charts	401
What Is a Chart?	401
How Excel handles charts	402
Embedded charts.....	403
Chart sheets	404
Parts of a chart	405
Chart limitations.....	408
Basic Steps for Creating a Chart.....	408
Creating the chart	408
Switching the row and column orientation.....	410
Changing the chart type	410
Applying a chart layout	412
Applying a chart style	413
Adding and deleting chart elements	413
Formatting chart elements	413
Modifying and Customizing Charts	414
Moving and resizing a chart.....	414
Converting an embedded chart to a chart sheet	415
Copying a chart	416
Deleting a chart.....	416
Adding chart elements	416
Moving and deleting chart elements	416
Formatting chart elements	416
Copying a chart's formatting.....	417
Renaming a chart.....	418
Printing charts.....	418
Understanding Chart Types	419
Choosing a chart type.....	419
Column charts	421
Bar charts	423
Line charts	424

Pie charts.....	426
XY (scatter) charts	427
Area charts.....	428
Radar charts	429
Surface charts.....	430
Bubble charts	430
Stock charts	431
Newer Chart Types for Excel	431
Histogram charts.....	431
Pareto charts	433
Waterfall charts	434
Box & whisker charts.....	434
Sunburst charts	436
Treemap charts	437
Funnel charts	437
Map charts	438
Chapter 19: Using Advanced Charting Techniques	441
Selecting Chart Elements	441
Selecting with the mouse.....	442
Selecting with the keyboard	443
Selecting with the Chart Elements control.....	443
Exploring the User Interface Choices for Modifying Chart Elements	444
Using the Format task pane.....	444
Using the chart customization buttons	445
Using the Ribbon	446
Using the Mini toolbar	446
Modifying the Chart Area.....	447
Modifying the Plot Area.....	448
Working with Titles in a Chart.....	449
Working with a Legend.....	450
Working with Gridlines.....	452
Modifying the Axes.....	452
Modifying the value axis.....	452
Modifying the category axis	456
Working with Data Series.....	460
Deleting or hiding a data series	461
Adding a new data series to a chart.....	462
Changing data used by a series	462
Changing the data range by dragging the range outline.....	463
Using the Edit Series dialog box	463
Editing the Series formula	464
Displaying data labels in a chart	465
Handling missing data	467

Contents

Adding error bars	468
Adding a trendline	468
Creating combination charts.....	470
Displaying a data table	472
Creating Chart Templates	473
Chapter 20: Creating Sparkline Graphics.....	475
Sparkline Types	475
Creating Sparklines	477
Customizing Sparklines.....	480
Sizing Sparkline cells.....	480
Handling hidden or missing data.....	480
Changing the Sparkline type	481
Changing Sparkline colors and line width.....	481
Highlighting certain data points	481
Adjusting Sparkline axis scaling.....	482
Faking a reference line.....	483
Specifying a Date Axis.....	484
Auto-Updating Sparklines	486
Displaying a Sparkline for a Dynamic Range.....	486
Chapter 21: Visualizing with Custom Number Formats and Shapes .. .	489
Visualizing with Number Formatting	489
Doing basic number formatting	489
Using shortcut keys to format numbers	490
Using the Format Cells dialog box to format numbers.....	491
Getting fancy with custom number formatting.....	492
Formatting numbers in thousands and millions	494
Hiding and suppressing zeros.....	495
Applying custom format colors.....	496
Formatting dates and times	497
Using symbols to enhance reporting	498
Using Shapes and Icons as Visual Elements	502
Inserting a shape	502
Inserting SVG icon graphics.....	504
Inserting 3D models	504
Formatting shapes and icons	506
Enhancing Excel reports with shapes	507
Creating visually appealing containers with shapes.....	507
Layering shapes to save space	509
Constructing your own infographic widgets with shapes.....	509
Creating dynamic labels	510
Creating linked pictures	510

Using SmartArt and WordArt	513
SmartArt basics.....	513
WordArt basics.....	514
Working with Other Graphics Types.....	515
About graphics files	515
Inserting screenshots.....	516
Displaying a worksheet background image.....	516
Using the Equation Editor.....	516
Part IV: Managing and Analyzing Data	519
Chapter 22: Importing and Cleaning Data	521
Importing Data.....	521
Importing from a file	522
Spreadsheet file formats	522
Database file formats	522
Text file formats	523
HTML files	523
XML files	524
Importing vs. opening	524
Importing a text file.....	525
Copying and pasting data.....	528
Cleaning Up Data.....	529
Removing duplicate rows.....	529
Identifying duplicate rows.....	530
Splitting text.....	531
Using Text to Columns.....	532
Using Flash Fill.....	533
Changing the case of text	536
Removing extra spaces.....	537
Removing strange characters	538
Converting values	538
Classifying values.....	538
Joining columns	540
Rearranging columns	541
Randomizing the rows	541
Extracting a filename from a URL.....	541
Matching text in a list	542
Changing vertical data to horizontal data	543
Filling gaps in an imported report.....	545
Checking spelling	547
Replacing or removing text in cells.....	547
Adding text to cells.....	548

Contents

Fixing trailing minus signs	549
Following a data cleaning checklist	549
Exporting Data	550
Exporting to a text file	550
CSV files	550
TXT files	550
PRN files.....	551
Exporting to other file formats	551
Chapter 23: Using Data Validation.	553
About Data Validation	553
Specifying Validation Criteria	554
Types of Validation Criteria You Can Apply.....	555
Creating a Drop-Down List	557
Using Formulas for Data Validation Rules	558
Understanding Cell References.....	559
Data Validation Formula Examples	560
Accepting text only.....	561
Accepting a larger value than the previous cell.....	561
Accepting nonduplicate entries only	561
Accepting text that begins with a specific character.....	561
Accepting dates by the day of the week.....	562
Accepting only values that don't exceed a total.....	563
Creating a dependent list	563
Using Data Validation without Restricting Entry	564
Showing an input message	564
Making suggested entries.....	564
Chapter 24: Creating and Using Worksheet Outlines	567
Introducing Worksheet Outlines.....	567
Creating an Outline.....	570
Preparing the data	571
Creating an outline automatically	572
Creating an outline manually.....	572
Working with Outlines.....	574
Displaying levels	574
Adding data to an outline	575
Removing an outline	575
Adjusting the outline symbols	575
Hiding the outline symbols	575
Chapter 25: Linking and Consolidating Worksheets	577
Linking Workbooks	577
Creating External Reference Formulas.....	578

Understanding link formula syntax.....	578
Creating a link formula by pointing.....	579
Pasting links.....	580
Working with External Reference Formulas	580
Creating links to unsaved workbooks	580
Opening a workbook with external reference formulas.....	581
Changing the startup prompt.....	582
Updating links.....	582
Changing the link source	583
Severing links.....	583
Avoiding Potential Problems with External Reference Formulas.....	583
Renaming or moving a source workbook	584
Using the Save As command	584
Modifying a source workbook	584
Using Intermediary links.....	585
Consolidating Worksheets	585
Consolidating worksheets by using formulas	587
Consolidating worksheets by using Paste Special	587
Consolidating worksheets by using the Consolidate dialog box.....	588
Viewing a workbook consolidation example.....	590
Refreshing a consolidation	592
Learning more about consolidation	593
Chapter 26: Introducing PivotTables	595
About PivotTables	595
A PivotTable example.....	596
Data appropriate for a PivotTable	598
Creating a PivotTable Automatically	600
Creating a PivotTable Manually.....	602
Specifying the data	602
Specifying the location for the PivotTable	603
Laying out the PivotTable	603
Formatting the PivotTable	607
Modifying the PivotTable.....	609
Seeing More PivotTable Examples.....	611
What is the daily total new deposit amount for each branch?	611
Which day of the week accounts for the most deposits?.....	611
How many accounts were opened at each branch, broken down by account type?	613
How much money was used to open the accounts?	613
What types of accounts do tellers open most often?	614
In which branch do tellers open the most checking accounts for new customers?	615
Learning More	616

Contents

Chapter 27: Analyzing Data with PivotTables	617
Working with Non-Numeric Data.....	617
Grouping PivotTable Items	619
Grouping items manually	619
Grouping items automatically	621
Grouping by date.....	621
Grouping by time.....	625
Using a PivotTable to Create a Frequency Distribution	626
Creating a Calculated Field or Calculated Item	628
Creating a calculated field.....	630
Inserting a calculated item.....	632
Filtering PivotTables with Slicers	635
Filtering PivotTables with a Timeline.....	637
Referencing Cells within a PivotTable	638
Creating PivotCharts.....	640
A PivotChart example	640
More about PivotCharts.....	643
Using the Data Model	644
Chapter 28: Performing Spreadsheet What-If Analysis	651
Looking at a What-If Example	651
Exploring Types of What-If Analyses	653
Performing manual what-if analysis.....	653
Creating data tables	653
Creating a one-input data table.....	654
Creating a two-input data table.....	657
Using Scenario Manager	661
Defining scenarios	662
Displaying scenarios	664
Modifying scenarios	666
Merging scenarios	666
Generating a scenario report	666
Analyzing Data with Artificial Intelligence	668
Using Excel's suggestions	668
Querying analyzed data	671
Chapter 29: Analyzing Data Using Goal Seeking and Solver	675
Exploring What-If Analysis, in Reverse	675
Using Single-Cell Goal Seeking	675
Looking at a goal-seeking example.....	676
Learning more about goal seeking.....	678
Introducing Solver	678
Looking at appropriate problems for Solver	679
Seeing a simple Solver example.....	679
Exploring Solver options	685

Seeing Some Solver Examples.....	686
Solving simultaneous linear equations	686
Minimizing shipping costs.....	688
Allocating resources	691
Optimizing an investment portfolio.....	693
Chapter 30: Analyzing Data with the Analysis ToolPak	697
The Analysis ToolPak: An Overview	697
Installing the Analysis ToolPak Add-In	698
Using the Analysis Tools.....	698
Introducing the Analysis ToolPak Tools.....	699
Analysis of variance	699
Correlation	700
Covariance.....	701
Descriptive statistics	701
Exponential smoothing	701
F-Test (two-sample test for variance).....	701
Fourier analysis	702
Histogram	703
Moving average.....	704
Random number generation.....	705
Rank and percentile	706
Regression.....	706
Sampling.....	707
t-Test	707
z-Test (two-sample test for means).....	708
Chapter 31: Protecting Your Work	709
Types of Protection.....	709
Protecting a Worksheet.....	710
Unlocking cells	710
Sheet protection options.....	712
Assigning user permissions.....	713
Protecting a Workbook	714
Requiring a password to open a workbook.....	714
Protecting a workbook's structure	715
Protecting a VBA Project.....	716
Related Topics	717
Saving a worksheet as a PDF file	717
Marking a workbook as final.....	717
Inspecting a workbook.....	718
Using a digital signature.....	719
Getting a digital ID.....	719
Signing a workbook	719

Part V: Understanding Power Pivot and Power Query	721
Chapter 32: Introducing Power Pivot	723
Understanding the Power Pivot Internal Data Model.....	723
The Power Pivot Ribbon	724
Linking Excel tables to Power Pivot	725
Preparing your Excel tables.....	726
Adding your Excel tables to the data model	727
Creating relationships between your PowerPivot tables.....	728
Managing existing relationships	730
Using Power Pivot data in reporting	732
Loading Data from Other Data Sources.....	733
Loading data from relational databases.....	733
Loading data from SQL Server	733
Loading data from other relational database systems	738
Loading data from flat files	738
Loading data from external Excel files.....	739
Loading data from text files.....	741
Loading data from the Clipboard.....	742
Refreshing and managing external data connections	742
Manually refreshing your Power Pivot data	743
Setting up automatic refreshing	743
Editing your data connection.....	745
Chapter 33: Working Directly with the Internal Data Model.	747
Directly Feeding the Internal Data Model.....	747
Managing Relationships in the Internal Data Model	754
Managing Queries & Connections	755
Chapter 34: Adding Formulas to Power Pivot	757
Enhancing Power Pivot Data with Calculated Columns	757
Creating your first calculated column	758
Formatting your calculated columns	759
Referencing calculated columns in other calculations	760
Hiding calculated columns from end users.....	760
Utilizing DAX to Create Calculated Columns	762
Identifying DAX functions safe for calculated columns.....	762
Building DAX-driven calculated columns	764
Month sorting in Power Pivot-driven PivotTables.....	765
Referencing fields from other tables	768
Nesting functions.....	770
Understanding Calculated Measures.....	770
Editing and deleting calculated measures	773
Using Cube Functions to Free Your Data.....	773

Chapter 35: Introducing Power Query	777
Understanding Power Query Basics	777
Understanding query steps	784
Viewing the Advanced Query Editor	785
Refreshing Power Query data	786
Managing existing queries	787
Understanding column-level actions	788
Understanding table actions	790
Getting Data from External Sources.....	792
Importing data from files	793
Getting data from Excel workbooks.....	794
Getting data from CSV and text files	795
Getting data from PDF files	795
Importing data from database systems.....	796
Importing data from relational and OLAP databases	796
Importing data from Azure databases.....	797
Importing data using ODBC connections to nonstandard databases	797
Getting Data from Other Data Systems.....	797
Managing Data Source Settings.....	798
Editing data source settings	798
Data Profiling with Power Query	800
Data profiling options.....	800
Data profiling quick actions.....	801
Chapter 36: Transforming Data with Power Query.....	805
Performing Common Transformation Tasks	805
Removing duplicate records.....	805
Filling in blank fields	808
Filling in empty strings.....	808
Concatenating columns.....	809
Changing case.....	811
Finding and replacing specific text	811
Trimming and cleaning text.....	812
Extracting the left, right, and middle values.....	814
Extracting first and last characters	815
Extracting middle characters.....	816
Splitting columns using character markers	816
Unpivoting columns	819
Unpivoting other columns	820
Pivoting columns	821
Creating Custom Columns.....	823
Concatenating with a custom column	824
Understanding data type conversions	826
Spicing up custom columns with functions.....	826
Adding conditional logic to custom columns.....	829

Contents

Grouping and Aggregating Data.....	830
Working with Custom Data Types.....	832
Chapter 37: Making Queries Work Together	837
Reusing Query Steps	837
Understanding the Append Feature.....	841
Creating the needed base queries	841
Appending the data.....	842
Understanding the Merge Feature	845
Understanding Power Query joins	845
Merging queries	846
Understanding fuzzy matching	851
Chapter 38: Enhancing Power Query Productivity	855
Implementing Some Power Query Productivity Tips	855
Getting quick information about your queries	855
Organizing queries in groups	856
Selecting columns in your queries faster.....	857
Renaming query steps	857
Quickly creating reference tables	859
Copying queries to save time	859
Viewing query dependencies.....	860
Setting a default load behavior	860
Preventing automatic data type changes.....	861
Avoiding Power Query Performance Issues.....	862
Using views instead of tables	862
Letting your back-end database servers do some crunching	863
Upgrading to 64-bit Excel.....	863
Disabling privacy settings to improve performance	864
Disabling relationship detection.....	864
Part VI: Automating Excel	867
Chapter 39: Introducing Visual Basic for Applications	869
Introducing VBA Macros	869
Displaying the Developer Tab	870
Learning about Macro Security	871
Saving Workbooks That Contain Macros	872
Looking at Two Types of VBA Macros.....	873
VBA Sub procedures	873
VBA functions.....	874
Creating VBA Macros	876
Recording VBA macros	876
Recording your actions to create VBA code: the basics	876
Recording a macro: a simple example	877
Examining the macro	878

Testing the macro	879
Editing the macro	879
Relative versus absolute recording	880
Another example	881
Running the macro	881
Examining the macro	882
Rerecording the macro	883
Testing the macro	883
More about recording VBA macros	884
Storing macros in your Personal Macro Workbook	884
Assigning a macro to a shortcut key	885
Assigning a macro to a button.....	885
Adding a macro to your Quick Access Toolbar.....	886
Writing VBA code	887
The basics: entering and editing code	887
The Excel object model.....	888
Objects and collections.....	889
Properties.....	889
Methods	891
The Range object.....	892
Variables.....	892
Controlling execution.....	893
A macro that can't be recorded.....	895
Learning More	896
Chapter 40: Creating Custom Worksheet Functions	899
Introducing VBA Functions	899
Seeing a Simple Example	900
Creating a custom function	900
Using the function in a worksheet	901
Analyzing the custom function	901
Learning about Function Procedures	902
Executing Function Procedures.....	904
Calling custom functions from a procedure	904
Using custom functions in a worksheet formula	904
Using Function Procedure Arguments	905
Creating a function with no arguments	905
Creating a function with one argument.....	906
Creating another function with one argument	906
Creating a function with two arguments	907
Creating a function with a range argument	908
Creating a simple but useful function	909
Debugging Custom Functions	910
Inserting Custom Functions.....	910
Learning More	912

Contents

Chapter 41: Creating UserForms	913
Understanding Why to Create UserForms.....	913
Exploring UserForm Alternatives.....	914
Using the InputBox function	914
Using the MsgBox function.....	915
Creating UserForms: An Overview.....	918
Working with UserForms	919
Adding controls	919
Changing the properties of a control	920
Handling events.....	921
Displaying a UserForm	922
Looking at a UserForm Example.....	923
Creating the UserForm	923
Testing the UserForm.....	924
Creating an event handler procedure	925
Looking at Another UserForm Example	926
Creating the UserForm	926
Creating event handler procedures	928
Showing the UserForm	930
Testing the UserForm.....	931
Making the macro available from a worksheet button	931
Making the macro available on your Quick Access Toolbar	932
Enhancing UserForms	932
Adding accelerator keys	932
Controlling tab order	933
Learning More	933
Chapter 42: Using UserForm Controls in a Worksheet	935
Understanding Why to Use Controls on a Worksheet.....	935
Using Controls	938
Adding a control	938
Learning about Design mode	938
Adjusting properties.....	938
Using common properties.....	939
Linking controls to cells.....	940
Creating macros for controls	941
Reviewing the Available ActiveX Controls.....	942
CheckBox	942
ComboBox	942
CommandButton	943
Image	944
Label	944
ListBox	944
OptionButton	945

ScrollBar	945
SpinBox.....	946
TextBox	946
ToggleButton.....	947
Chapter 43: Working with Excel Events.	949
Understanding Events.....	949
Entering Event-Handler VBA Code.....	950
Using Workbook-Level Events	951
Using the Open event	952
Using the SheetActivate event	953
Using the NewSheet event.....	954
Using the BeforeSave event	954
Using the BeforeClose event	954
Working with Worksheet Events.....	955
Using the Change event	956
Monitoring a specific range for changes	956
Using the SelectionChange event	958
Using the BeforeRightClick event	959
Using Special Application Events	959
Using the OnTime event	960
Using the OnKey event.....	961
Chapter 44: Seeing Some VBA Examples.	963
Working with Ranges	963
Copying a range	964
Copying a variable-size range.....	965
Selecting to the end of a row or column.....	966
Selecting a row or column.....	966
Moving a range	967
Looping through a range efficiently	967
Prompting for a cell value.....	968
Determining the type of selection	970
Identifying a multiple selection	970
Counting selected cells	971
Working with Workbooks	972
Saving all workbooks.....	972
Saving and closing all workbooks	972
Creating a workbook.....	972
Working with Charts	973
Modifying the chart type	973
Modifying chart properties	974
Applying chart formatting	974
VBA Speed Tips	975
Turning off screen updating	975

Contents

Preventing alert messages	975
Simplifying object references.....	976
Declaring variable types.....	976
Chapter 45: Creating Custom Excel Add-Ins.....	979
Understanding Add-Ins.....	979
Working with Add-Ins	980
Understanding When to Create Add-Ins	982
Creating Add-Ins.....	982
Looking at an Add-In Example	983
Learning about Module1	984
Learning about the UserForm.....	984
Testing the workbook	985
Adding descriptive information.....	985
Creating the user interface for your add-in macro.....	986
Protecting the project.....	986
Creating the add-in	987
Installing the add-in	987
Index	989

Introduction

Welcome to the world of Excel. Okay, that's a bit cheesy. But if you look around the business world, the financial world, the manufacturing world, and any other industry you can think of, you will see people using Excel. Excel is everywhere. It is by far the most popular program in the history of business applications. So, we truly are living in a world of Excel. This is probably why you've picked up this book. You need a way to accelerate your learning and get up to speed quickly.

Well, worry not, dear reader. Whether you're boning up on Excel for a new job (congratulations, by the way), for a school project, or just for home use, this book is perfect for you.

In this book, we've organized everything that one would need to know to get up and running quickly with Excel. And we've made certain that this book contains many useful examples and lots of tips and tricks that cover all the essential aspects of Excel—from the basics to more advanced topics.

Looking at What's New in Excel 365

Here's an overview of what's available to those with Microsoft 365 subscriptions as of the October 2021 update and those using the stand-alone (perpetual license) version of Office/Excel 2021.

Dynamic array inherent to Excel's calculation engine Dynamic array behavior is now fundamentally part of Excel's calculation engine. When any function uses an array that returns multiple values, the results will be output to a spill range. This even includes older functions that weren't designed to output arrays. With dynamic arrays, Excel moves into a new era where you no longer need to be a formula guru to leverage the power of arrays. The underlying changes to Excel's calculation engine make legacy array formulas obsolete.

New dynamic array functions With the introduction of dynamic arrays, Microsoft released several new functions that leverage dynamic arrays to improve the ability to carry out complex formula operations with ease. These new functions can remove duplicates, extract unique values, filter data, dynamically sort data, and perform sophisticated look-ups.

Formula variables Formula variables enable you to create a container of sorts that holds the results of a function or calculation for use later in other calculations. Armed with the new LET function, you can simplify your formulas and potentially improve performance.

Power Query Custom Data Types The Custom Data feature of Power Query allows you to store multiple columns of data in one column as metadata. You can think of a Custom Data Type as a kind of container that allows you to store the data for many columns, then use that data elsewhere in your workbook.

Power Query import from PDF Power Query now allows you to import and transform data directly from PDF files.

Power Query Data Profiler Power Query's new data profiling capabilities allow you to know your data and identify potential issues before using it. You can leverage data profiling in Power Query to get a better understanding of your data and address problem areas before they become a problem later in your reporting processes.

Import custom 3D models Excel now allows you to import your own custom 3D model graphics, including 3D Manufacturing, Filmbox, Binary GL Transmission, Polygon, and StereoLithography files.

Threaded comments Threaded comments allow for discussions between contributors directly within a workbook. With the look and feel of comments on a blog or online forum, anyone with access to the workbook can enter their own comments or reply to yours. Traditional comments are now called notes.

Action Pen With the new Action Pen feature, you can handwrite directly in the cells and Excel will convert writing into data! Primarily designed for mobile devices, Action Pen makes it easy to edit your workbooks on the go.

Is This Book for You?

This book is designed to enhance the skillset of users at all levels (beginning, intermediate, and even advanced users).

Start at the beginning if you're new to Excel. Part I covers everything you'll need to get familiar with entering data, managing workbooks, formatting worksheets, and printing. You can then move on to Part II, where you'll discover the ins and outs of Excel formulas and functions.

If you're a seasoned analyst, hoping to enhance your data visualization and analytic toolset, check out Part III and Part IV. We've included many examples and tips for analyzing data and creating visually appealing Excel dashboards.

If you've been working with an earlier version of Excel, this book is for you too! Part V covers the Power Pivot and Power Query toolsets. In the past, these features were free Microsoft add-ins that were used peripherally. Now they've become an essential part of how Excel manages data and interacts with external data sources.

If you want to learn the basics of Visual Basic for Applications (VBA) programming, you'll find what you need in Part VI. This part offers a robust set of chapters that will get you started leveraging VBA to automate and enhance your Excel solutions.

Software Versions

This book covers features released as of the October 2021 update of Microsoft 365. The functionality explored here will be available to those on Microsoft 365 subscriptions and those using the stand-alone (perpetual license) version of Office/Excel 2021 for the desktop. Please note that this book is *not* applicable to Microsoft Excel for Mac.

Excel is available in several versions, including a web version and a version for tablets and phones. Though this book was written for the desktop version of Excel, much of the information here will also apply to the web and tablet versions.

Over the last few years, Microsoft has adopted an agile release cycle, releasing updates to Microsoft 365 practically on a monthly basis. This is great news for those who love seeing new features added to Excel. It's not so great if you're trying to document the features of these tools in a book.

Our assumption is that Microsoft will continue to add new bells and whistles to Excel at a rapid pace after publication of this book. Thus, you may encounter new functionality not covered in this book. That being said, Excel has a broad feature set, much of which is stable and here to stay. So, even though changes will be made to Excel, they won't be so drastic as to turn this book into a doorstop. The core functionality covered in these chapters will remain relevant—even if the mechanics change a bit.

Conventions Used in This Book

Take a minute to scan the following sections to learn some of the typographical and organizational conventions that this book uses.

Excel commands

Excel uses a context-sensitive Ribbon system. The words along the top (such as File, Home, Insert, Page Layout, and so on) are known as *tabs*. Click a tab, and the Ribbon displays the commands for the selected tab. Each command has a name, which is (usually) displayed next to or below the icon. The commands are arranged in groups, and the group name appears at the bottom of the Ribbon.

The convention we use is to indicate the tab name, followed by the group name, followed by the command name. So, the command used to toggle word wrap within a cell is indicated as follows:

Home \leftrightarrow Alignment \leftrightarrow Wrap Text

You'll learn more about the Ribbon user interface in Chapter 1, "Introducing Excel."

Typographical conventions

Anything that you're supposed to type using the keyboard appears in a **bold monospaced font**. Lengthy input usually appears on a separate line. Here's an example:

```
= "Part Name: " &VLOOKUP(PartNumber, PartList, 2)
```

Names of the keys on your keyboard appear in normal type. When two keys should be pressed simultaneously, they're connected with a plus sign, like this: "Press Ctrl+C to copy the selected cells."

The four "arrow" keys are collectively known as the navigation keys.

Excel's built-in worksheet functions appear in monospaced font in uppercase like this: "Note the SUMPRODUCT function used in cell C20."

Mouse conventions

You'll come across some of the following mouse-related terms, which are all standard fare:

Mouse pointer This is the small graphic figure that moves onscreen when you move your mouse. The mouse pointer is usually an arrow, but it changes shape when you move to certain areas of the screen or when you're performing certain actions.

Point Move the mouse so that the mouse pointer is on a specific item; for example, "Point to the Paste button on the Home tab."

Click Press the left mouse button once and release it immediately.

Right-click Press the right mouse button once and release it immediately. The right mouse button is used in Excel to open shortcut menus that are appropriate for whatever is currently selected.

Double-click Press the left mouse button twice in rapid succession.

Drag Press the left mouse button and keep it pressed while you move the mouse. Dragging is often used to select a range of cells or to change the size of an object.

For Touch Screen Users

If you happen to be using one of these devices, you probably already know the basic touch gestures.

This book doesn't cover specific touch screen gestures, but these three guidelines should work most of the time:

- When you read "click," you should tap. Quickly touching and releasing your finger on a button is the same as clicking it with a mouse.
- When you read "double-click," tap twice. Touching twice in rapid succession is equivalent to double-clicking.
- When you read "right-click," press and hold your finger on the item until a menu appears. Tap an item on the pop-up menu to execute the command.

Make sure you enable Touch mode from the Quick Access Toolbar. Touch mode increases the spacing between the Ribbon commands, making it less likely that you'll touch the wrong command. If the Touch mode command is not in your Quick Access Toolbar, touch the rightmost control and select Touch/Mouse Mode. This command toggles between normal mode and Touch mode.

How This Book Is Organized

Notice that the book is divided into six main parts.

Part I: Getting Started with Excel This part consists of eight chapters that provide background about Excel. These chapters are considered required reading for Excel newcomers, but even experienced users will find new information on threaded comments and working with Excel on mobile devices.

Part II: Working with Formulas and Functions The chapters in Part II cover everything that you need to know to become proficient with performing calculations in Excel. Chapter 10 is required reading, even for seasoned professionals, as it covers the new dynamic array functionality inherent in Excel's calculation engine.

Part III: Creating Charts and Other Visualizations The chapters in Part III describe how to create effective charts. In addition, you'll find chapters on the conditional formatting visualization features and Sparkline graphics and a chapter with lots of tips on integrating graphics into your worksheet.

Part IV: Managing and Analyzing Data Data analysis is the focus of the chapters in Part IV. Here you'll find chapters focusing on data validation, PivotTables, conditional analyses, and more.

Part V: Understanding Power Pivot and Power Query The chapters in Part V take an in-depth look at the functionality found in Power Pivot and Power Query. Here you'll discover how to develop powerful reporting solutions with Power Pivot as well as how to leverage Power Query to automate and steps for cleaning and transforming data.

Part VI: Automating Excel Part VI is for those who want to customize Excel for their own use or who are designing workbooks or add-ins that are to be used by others. It starts with an introduction to recording macros and VBA programming, and then it provides coverage of UserForms, events, and add-ins.

How to Use This Book

This book is obviously not meant to be read cover to cover. Instead, it's a reference book that you can consult when you need help with the following situations:

- You're stuck while trying to do something.
- You need to do something that you've never done before.
- You have some time on your hands, and you're interested in learning something new about Excel.

The index is comprehensive, and each chapter typically focuses on a single broad topic. Don't be discouraged if some of the material is over your head. Most users get by just fine by using only a small subset of Excel's total capabilities. In fact, the 80/20 rule applies here: 80 percent of Excel users use only 20 percent of its features. However, knowing only 20 percent of Excel's features still gives you a lot of power at your fingertips.

What's on the Website

This book contains many examples, and you can download the workbooks for those examples from the Web. The files are arranged in directories that correspond to the chapters.

The URL is www.wiley.com/go/excel365bible.

Please note that the URL is case sensitive, so use all lowercase letters.

Part I

Getting Started with Excel

The chapters in this part are intended to provide essential background information for working with Excel. Here, you'll see how to make use of the basic features that are required for every Excel user. If you've used Excel (or even a different spreadsheet program) in the past, much of this information may seem like a review. Even so, it's likely that you'll find quite a few new tricks and techniques in these chapters.

IN THIS PART

- Chapter 1**
Introducing Excel
- Chapter 2**
Entering and Editing Worksheet Data
- Chapter 3**
Performing Basic Worksheet Operations
- Chapter 4**
Working with Excel Ranges and Tables
- Chapter 5**
Formatting Worksheets
- Chapter 6**
Understanding Excel Files and Templates
- Chapter 7**
Printing Your Work
- Chapter 8**
Customizing the Excel User Interface

Introducing Excel

IN THIS CHAPTER

Understanding what Excel is used for

Learning the parts of an Excel window

Moving around a worksheet

Introducing the Ribbon, shortcut menus, dialog boxes, and task panes

Introducing Excel with a step-by-step hands-on session

This chapter is an introductory overview of Excel 2022. If you're already familiar with a previous version of Excel, reading (or at least skimming) this chapter is still a good idea.

Understanding What Excel Is Used For

Excel is the world's most widely used spreadsheet software and is part of the Microsoft Office suite. Other spreadsheet software is available, but Excel is by far the most popular and has been the world standard for many years.

Much of the appeal of Excel is its versatility. Excel's forte, of course, is performing numerical calculations, but Excel is also useful for nonnumeric applications. Here are just a few uses for Excel:

Crunching numbers Create budgets, tabulate expenses, analyze survey results, and perform just about any type of financial analysis you can think of.

Creating charts Create a variety of highly customizable charts.

Organizing lists Use the row-and-column layout to store lists efficiently.

Manipulating text Clean up and standardize text-based data.

Accessing other data Import data from a variety of sources such as databases, text files, web pages, and many others.

Creating graphical dashboards Summarize a large amount of business information in a concise format.

Creating graphics and diagrams Use shapes and illustrations to create professional-looking diagrams.

Automating complex tasks Perform a tedious task with a single mouse click with Excel's macro capabilities.

Understanding Workbooks and Worksheets

An Excel file is called a *workbook*. You can have as many workbooks open as you need, and each one appears in its own window. By default, Excel workbooks use an .xlsx file extension.

Note

In old versions of Excel, every workbook opened in a single Excel window. Beginning with Excel 2013, each workbook opens in its own window. This change makes Excel work more like other Office applications and gives you the opportunity to put different workbooks on different monitors more easily.

The tabs in a workbook are called worksheets. Each workbook contains one or more worksheets, and each worksheet consists of individual cells. Each cell can contain a number, a formula, or text. A worksheet also has an invisible drawing layer, which holds charts, images, and diagrams. Objects on the drawing layer sit over the top of the cells, but they are not *in* the cells like a number or formula. You switch to a different worksheet by clicking its tab at the bottom of the workbook window. In addition, a workbook can store chart sheets: a chart sheet displays a single chart and is accessible by clicking a tab.

Don't be intimidated by all the different elements that appear within Excel's window. You don't need to know what all of them mean to use Excel effectively. And after you become familiar with the various parts, it all starts to make sense and you'll feel right at home.

Figure 1.1 shows you the more important bits and pieces of Excel. As you look at the figure, refer to Table 1.1 for a brief explanation of the items shown.

Moving around a Worksheet

This section describes various ways to navigate the cells in a worksheet.

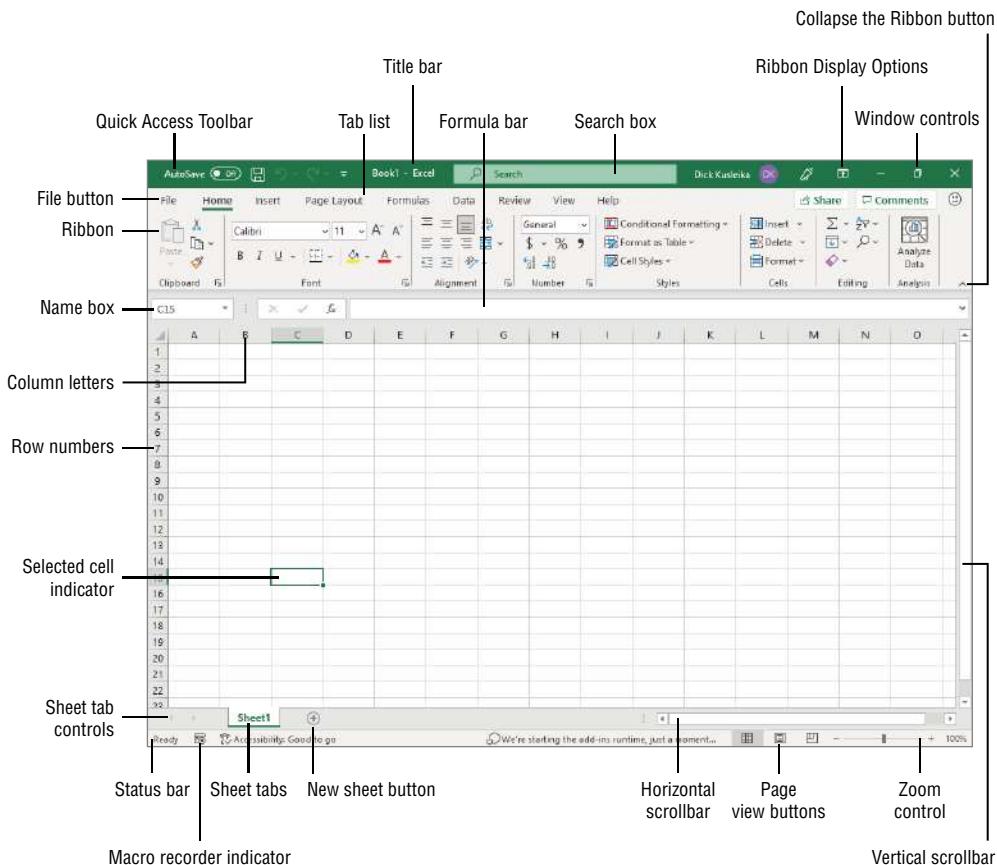
Every worksheet consists of rows (numbered 1 through 1,048,576) and columns (labeled A through XFD). Column labeling works like this: After column Z comes column AA, which is followed by AB, AC, and so on. After column AZ comes BA, BB, and so on. After column ZZ is AAA, AAB, and so on.

The intersection of a row and a column is a single cell, and each cell has a unique address made up of its column letter and row number. For example, the address of the upper-left cell is A1. The address of the cell at the lower right of a worksheet is XFD1048576.

At any given time, one cell is the active cell. The active cell is the cell that accepts keyboard input, and its contents can be edited. You can identify the active cell by its darker border, as shown in Figure 1.2. If more than one cell is selected, the dark border surrounds the entire selection, and the active cell is the light-colored cell within the border. Its address appears in the Name box. Depending on the technique you use to navigate through a workbook, you may or may not change the active cell when you navigate.

FIGURE 1.1

The Excel screen has many useful elements that you will use often.

**TABLE 1.1 Parts of the Excel Screen That You Need to Know**

Name	Description
Collapse the Ribbon button	Click this button to hide the Ribbon temporarily. Double-click any Ribbon tab to make the Ribbon remain visible. Ctrl+F1 is the shortcut key that does the same task.
Column letters	Letters range from A to XFD—one for each of the 16,384 columns in the worksheet. You can click a column heading to select an entire column or click between the column letters and drag to change the column width.
File button	Click this button to open Backstage view, which contains many options for working with your document (including printing) and setting Excel options.
Formula bar	When you enter information or formulas into a cell, it appears in this bar.

Part I: Getting Started with Excel

TABLE 1.1 (Continued)

Name	Description
Horizontal scrollbar	Use this tool to scroll the sheet horizontally.
Macro recorder indicator	Click to start recording a Visual Basic for Applications (VBA) macro. The icon changes while your actions are being recorded. Click again to stop recording.
Name box	This box displays the active cell address or the name of the selected cell, range, or object.
New sheet button	Add a new worksheet by clicking the New sheet button (which is displayed after the last sheet tab).
Page view buttons	Click these buttons to change the way the worksheet is displayed.
Quick Access Toolbar	This customizable toolbar holds commonly used commands. The Quick Access Toolbar is always visible, regardless of which tab is selected.
Ribbon	This is the main location for Excel commands. Clicking an item in the tab list changes the Ribbon that is displayed.
Ribbon Display Options	A drop-down control that offers three options related to displaying the Ribbon.
Row numbers	Numbers range from 1 to 1,048,576—one for each row in the worksheet. You can click a row number to select an entire row or click between the row numbers and drag to change the row height.
Search box	Use this control to find commands or have Excel issue a command automatically. Alt+Q is the shortcut to access the Search box.
Selected cell indicator	This dark outline indicates the currently selected cell or range of cells. (There are 17,179,869,184 cells on each worksheet.)
Sheet tabs	Each of these notebook-like tabs represents a different sheet in the workbook. A workbook can have any number of sheets, and each sheet has its name displayed in a sheet tab.
Sheet tab controls	Use these buttons to scroll the sheet tabs to display tabs that aren't visible. You can also right-click to get a list of sheets.
Status bar	This bar displays various messages as well as summary information about the range of cells selected. Right-click the status bar to change which messages are displayed.
Tab list	Use these commands to display a different Ribbon.
Title bar	This displays the name of the program and the name of the current workbook. It also holds the Quick Access Toolbar (on the left), the Search box, and some control buttons that you can use to modify the window (on the right).
Vertical scrollbar	Use this tool to scroll the sheet vertically.
Window controls	There are three controls for minimizing the current window, maximizing or restoring the current window, and closing the current window, which are common to virtually all Windows applications.
Zoom control	Use this to zoom your worksheet in and out.

FIGURE 1.2

The active cell is the one with the dark border—in this case, cell C11.

	A	B	C	D
1		This Year	Last Year	
2	January	8,097	8,371	
3	February	7,985	7,567	
4	March	8,441	7,512	
5	April	8,088	7,453	
6	May	8,204	8,664	
7	June	7,114	7,466	
8	July	7,040	7,794	
9	August	7,265	7,018	
10	September	8,459	8,032	
11	October	8,982	8,637	
12	November	7,337	7,127	
13	December	7,799	7,331	
14				
15				

The row and column headings of the active cell appear in a different color to make it easier to identify the row and column of the active cell.

NOTE

Excel is also available for devices that use a touch interface. This book assumes you have a traditional keyboard and mouse, so it doesn't cover the touch-related commands. Note that the drop-down control in the Quick Access Toolbar has a Touch/Mouse Mode command. In Touch mode, the Ribbon and Quick Access Toolbar icons are placed further apart.

Navigating with your keyboard

Not surprisingly, you can use the standard navigational keys on your keyboard to move around a worksheet. These keys work just as you'd expect: the down arrow moves the active cell down one row, the right arrow moves it one column to the right, and so on. PgUp and PgDn move the active cell up or down one full window. (The actual number of rows moved depends on the number of rows displayed in the window.)

TIP

You can use the keyboard to scroll through the worksheet without changing the active cell by turning on Scroll Lock, which is useful if you need to view another area of your worksheet and then quickly return to your original location. Just press Scroll Lock and use the navigation keys to scroll through the worksheet. When you want to return to the original position (the active cell), press Ctrl+Backspace and then press Scroll Lock again to turn it off. When Scroll Lock is turned on, Excel displays *Scroll Lock* in the status bar at the bottom of the window.

Part I: Getting Started with Excel

The Num Lock key on your keyboard controls the way the keys on the numeric keypad behave. When Num Lock is on, the keys on your numeric keypad generate numbers. Many keyboards have a separate set of navigation (arrow) keys located to the left of the numeric keypad. The state of the Num Lock key doesn't affect these keys.

Table 1.2 summarizes all the worksheet movement keys available in Excel.

TABLE 1.2 Excel Worksheet Movement Keys

Key	Action
Up arrow (\uparrow) or Shift+Enter	Moves the active cell up one row
Down arrow (\downarrow) or Enter	Moves the active cell down one row
Left arrow (\leftarrow) or Shift+Tab	Moves the active cell one column to the left
Right arrow (\rightarrow) or Tab	Moves the active cell one column to the right
PgUp	Moves the active cell up one screen
PgDn	Moves the active cell down one screen
Alt+PgDn	Moves the active cell right one screen
Alt+PgUp	Moves the active cell left one screen
Ctrl+Backspace	Scrolls the screen so that the active cell is visible
Ctrl+Home	Moves the active cell to A1
Ctrl+End	Moves the active cell to the bottom-rightmost cell on the worksheet's used range
\uparrow^*	Scrolls the screen up one row (active cell does not change)
\downarrow^*	Scrolls the screen down one row (active cell does not change)
\leftarrow^*	Scrolls the screen left one column (active cell does not change)
\rightarrow^*	Scrolls the screen right one column (active cell does not change)

* With Scroll Lock on

Navigating with your mouse

To change the active cell by using the mouse, just click another cell and it becomes the active cell. If the cell that you want to activate isn't visible in the workbook window, you can use the scrollbars to scroll the window in any direction. To scroll one cell, click either of the arrows on the scrollbar. To scroll by a complete screen, click either side of the scrollbar's scroll box. To scroll faster, drag the scroll box or right-click anywhere on the scrollbar for a menu of shortcuts.

Tip

If your mouse has a wheel, you can use it to scroll vertically. Also, if you click the wheel and move the mouse in any direction, the worksheet scrolls automatically in that direction. The more you move the mouse, the faster you scroll.

Press Ctrl while you use the mouse wheel to zoom the worksheet. If you prefer to use the mouse wheel to zoom the worksheet without pressing Ctrl, choose File \Rightarrow Options and select the Advanced section. Place a check mark next to the Zoom on Roll with IntelliMouse check box.

Using the scrollbars or scrolling with your mouse doesn't change the active cell. It simply scrolls the worksheet. To change the active cell, you must click a new cell after scrolling.

Using the Ribbon

The *Ribbon* is the primary way of interacting with Excel other than entering data into cells. The words above the icons are known as *tabs*: the Home tab, the Insert tab, and so on. The term *Ribbon* is used in two different ways: when you click a tab, you are said to be displaying a different Ribbon and the whole structure of tabs, groups, and controls is known as the Ribbon.

The Ribbon can be either hidden or visible. It's your choice. To toggle the Ribbon's visibility, press Ctrl+F1 (or double-click a tab at the top). If the Ribbon is hidden, it temporarily appears when you click a tab and hides itself when you click in the worksheet. The title bar has a control named Ribbon Display Options (next to the Minimize button). Click the control and choose one of three Ribbon options: Auto-Hide Ribbon, Show Tabs, or Show Tabs and Commands.

Ribbon tabs

The commands available on the Ribbon vary, depending upon which tab is selected. The Ribbon is arranged into groups of related commands. Here's a quick overview of Excel's tabs:

Home You'll probably spend most of your time with the Home tab selected. This tab contains the basic Clipboard commands, formatting commands, style commands, commands to insert and delete rows or columns, plus an assortment of worksheet editing commands.

Insert Select this tab when you need to insert something into a worksheet—a table, a diagram, a chart, a symbol, and so on.

Page Layout This tab contains commands that affect the overall appearance of your worksheet, including some settings that deal with printing.

Formulas Use this tab to insert a formula, name a cell or a range, access the formula auditing tools, or control the way Excel performs calculations.

Data Excel's data-related commands are on this tab, including data validation and sorting commands.

Review This tab contains tools to check spelling, translate words, add comments, or protect sheets.

View The View tab contains commands that control various aspects of how a sheet is viewed. Some commands on this tab are also available in the status bar.

Developer This tab isn't visible by default. It contains commands that are useful for programmers. To display the Developer tab, choose File \Rightarrow Options and then select Customize Ribbon. In the Customize the Ribbon section on the right, make sure that Main Tabs is selected in the drop-down control and place a check mark next to Developer.

Help This tab provides ways to get help, make suggestions, and access other aspects of Microsoft's community.

Add-Ins This tab is visible only if you loaded an older workbook or add-in that customizes the menu or toolbars. Because the old menus and toolbars were replaced by the Ribbon, these user interface customizations appear on the Add-Ins tab.

The preceding list contains the standard Ribbon tabs. Excel may display additional Ribbon tabs based on what's selected or resulting from add-ins that are installed.

NOTE

Although the File button shares space with the tabs, it's not actually a tab. Clicking the File button displays a different screen (known as Backstage view), where you perform actions with your documents. This screen has commands along the left side. To exit the Backstage view, click the back-arrow button in the upper-left corner.

The appearance of the commands on the Ribbon varies, depending on the width of the Excel window. When the Excel window is too narrow to display everything, the commands adapt; some of them might seem to be missing, but the commands are still available. Figure 1.3 shows the Home tab of the Ribbon with all controls fully visible. Figure 1.4 shows the Ribbon when Excel's window is made narrower. Notice that some of the descriptive text is gone, but the icons remain. Figure 1.5 shows the extreme case when the window is made very narrow. Most groups display a single icon; however, if you click the icon, all the group commands are available to you.

FIGURE 1.3

The Home tab of the Ribbon



FIGURE 1.4

The Home tab when Excel's window is made narrower

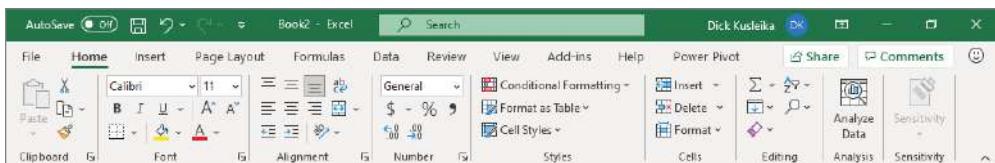
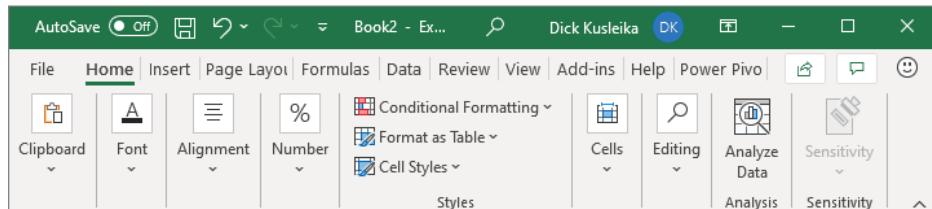


FIGURE 1.5

The Home tab when Excel's window is made very narrow



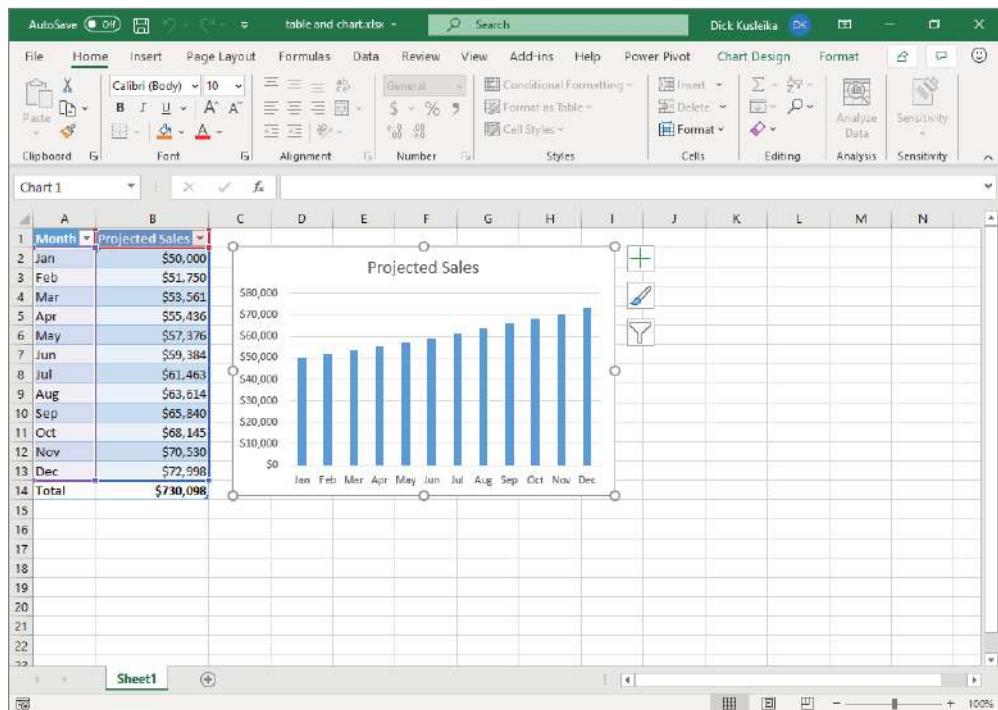
Contextual tabs

In addition to the standard tabs, Excel includes contextual tabs. Whenever an object (such as a chart, a table, or an illustration) is selected, specific tools for working with that object are made available on the Ribbon.

Figure 1.6 shows the contextual tabs that appear when a chart is selected. In this case, it has two contextual tabs: Chart Design and Format. When contextual tabs appear, you can, of course, continue to use all the other tabs.

FIGURE 1.6

When you select an object, contextual tabs contain tools for working with that object.



Types of commands on the Ribbon

When you hover your mouse pointer over a Ribbon command, you'll see a ScreenTip that contains the command's name and a brief description. For the most part, the commands on the Ribbon work just as you would expect. You'll find several different styles of commands on the Ribbon.

Simple buttons Click the button, and it does its thing. An example of a simple button is the Increase Font Size button in the Font group of the Home tab. Some buttons perform the action immediately; others display a dialog box so that you can enter additional information. Button controls may or may not be accompanied by a descriptive label.

Toggle buttons A toggle button is clickable and conveys some type of information by displaying two different colors. An example is the Bold button in the Font group of the Home tab. If the active cell isn't bold, the Bold button displays in its normal color. If the active cell is already bold, the Bold button displays a different background color. If you click the Bold button, it toggles the Bold attribute for the selection.

Simple drop-downs If the Ribbon command has a small down arrow, the command is a drop-down menu. Click it, and additional commands appear below it. An example is the Conditional Formatting command in the Styles group of the Home tab. When you click this control, you see several options related to conditional formatting. Style galleries, like the Data Types gallery on the Data tab, also have drop-down arrows to reveal more styles if they are available.

Split buttons A split button control combines a one-click button with a drop-down. If you click the button part, the command is executed. If you click the drop-down part (a down arrow), you choose from a list of related commands. An example of a split button is the Merge & Center command in the Alignment group of the Home tab (see Figure 1.7). Clicking the left part of this control merges and centers text in the selected cells. If you click the arrow part of the control (on the right), you get a list of commands related to merging cells.

Check boxes A check box control turns something on or off. An example is the Gridlines control in the Show group of the View tab. When the Gridlines check box is checked, the sheet displays gridlines. When the control isn't checked, the gridlines don't appear.

Spin buttons Excel's Ribbon has only one spin button control: the Scale to Fit group of the Page Layout tab. Click the top part of the spin button to increase the value; click the bottom part of the spin button to decrease the value.

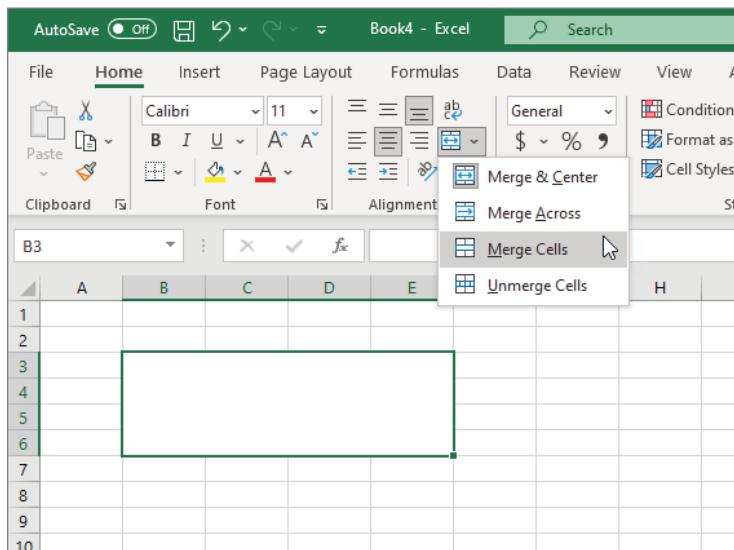
Some of the Ribbon groups contain a small icon in the bottom-right corner, known as a *dialog box launcher*. For example, if you examine the groups on the Home tab, you find dialog box launchers for the Clipboard, Font, Alignment, and Number groups—but not the Styles, Cells, and Editing groups. Click the icon, and Excel displays a dialog box or task pane. The dialog box launchers often provide options that aren't available on the Ribbon.

Accessing the Ribbon by using your keyboard

At first glance, you may think that the Ribbon is completely mouse centric. After all, the commands don't display the traditional underlined letter to indicate the Alt+keystrokes. But in fact, the Ribbon is very keyboard friendly. The trick is to press the Alt key to display the pop-up key-tips. Each Ribbon control has a letter (or series of letters) that you type to issue the command.

FIGURE 1.7

The Merge & Center command is a split button control.

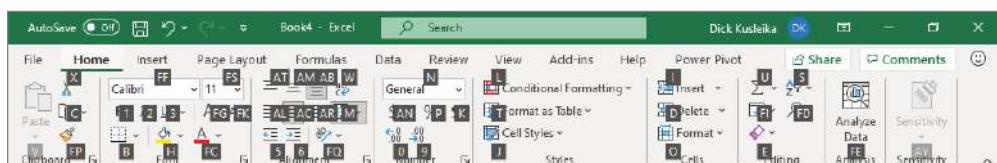
**TIP**

You don't need to hold down the Alt key while you type keytip letters.

Figure 1.8 shows how the Home tab looks after you press the Alt key to display the keytips and then the H key to display the keytips for the Home tab. If you press one of the keytips, the screen then displays more keytips. For example, to use the keyboard to align the cell contents to the left, press Alt, followed by H (for Home), and then AL (for Align Left).

FIGURE 1.8

Pressing Alt displays the keytips.



Nobody will memorize all of these keys, but if you're a keyboard fan, it takes just a few times before you memorize the keystrokes required for commands that you use frequently.

After you press Alt, you can also use the left- and right-arrow keys to scroll through the tabs. When you reach the proper tab, press the down arrow to enter the Ribbon. Then use left- and right-arrow keys to scroll through the Ribbon commands. If you hold down the Ctrl key while in

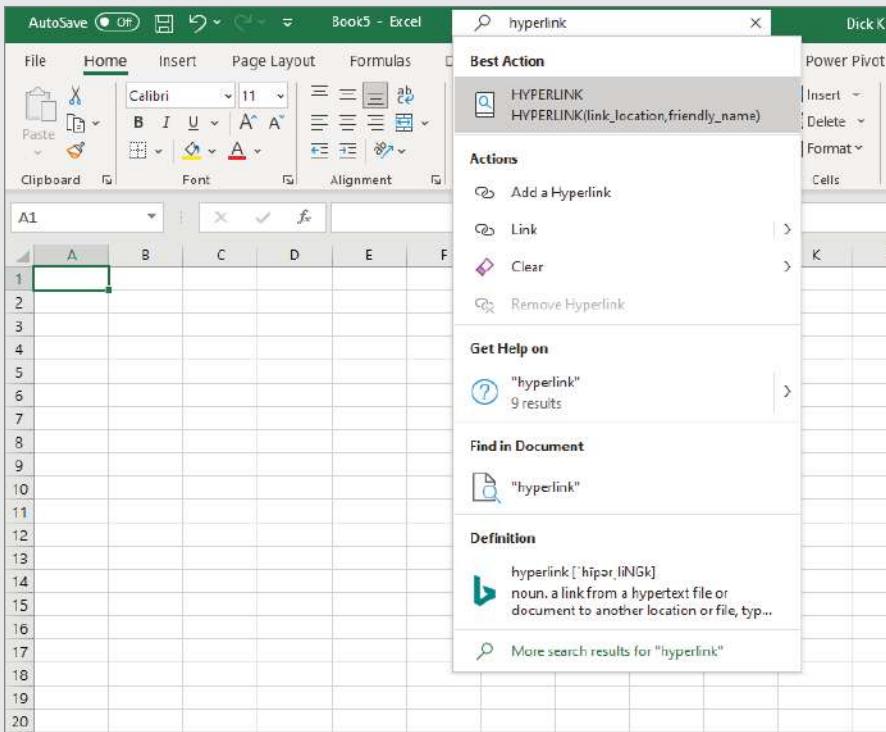
the Ribbon, the left and right arrows will jump to the first control in the previous or next group, respectively. When you reach the command you need, press Enter to execute it. This method isn't as efficient as using the keytips, but it's a quick way to take a look at the commands available.

TIP

Often, you'll want to repeat a particular command. Excel provides a way to simplify that. For example, if you apply a particular style to a cell (by choosing Home \Rightarrow Styles \Rightarrow Cell Styles), you can activate another cell and press Ctrl+Y (or F4) to repeat the command.

Searching for Commands

Excel has a Search box in the title bar for finding commands. If you're unsure of where to find a command, try typing it in the box. For example, if you want to insert a hyperlink on the current worksheet, type **hyperlink** in the Search box (Alt+Q will put your cursor in the box). Excel displays a list of potentially relevant commands and some help topics. If you see the command you want, click it (or use the arrow keys and press Enter). The command is executed. In this example, the HYPERLINK() worksheet function is the "best action." Selecting that option will launch the Function Arguments dialog box for the HYPERLINK() worksheet function just as if you had selected Formulas \Rightarrow Function Library \Rightarrow Lookup & Reference \Rightarrow HYPERLINK.



You don't have to be an Excel newcomer to appreciate the Search box. The Ribbon contains numerous commands, and even Excel experts can't be expected to remember where all of them are.

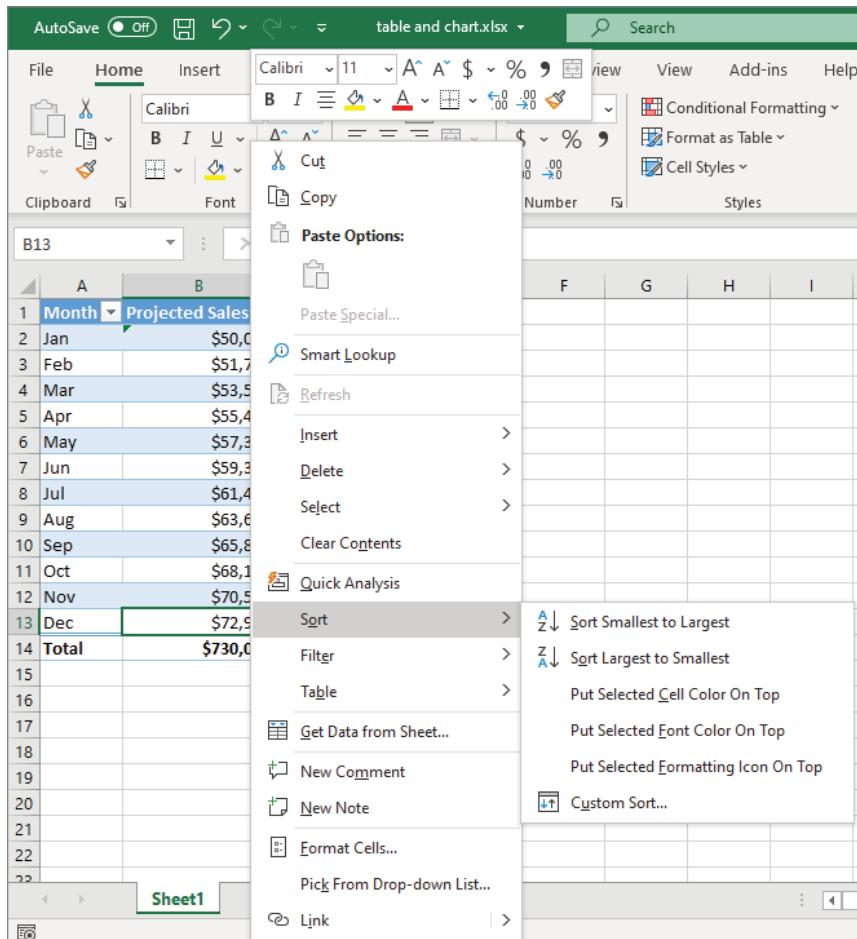
Using Shortcut Menus

In addition to the Ribbon, Excel features many shortcut menus, which you access by right-clicking just about anything within Excel. Shortcut menus don't contain every relevant command, just those that are most commonly used for whatever is selected.

As an example, Figure 1.9 shows the shortcut menu that appears when you right-click a cell in a table. The shortcut menu appears at the mouse-pointer position, which makes selecting a command fast and efficient. The shortcut menu that appears depends on what you're doing at the time. For example, if you're working with a chart, the shortcut menu contains commands that are pertinent to the selected chart element.

FIGURE 1.9

Right-click to display a shortcut menu of commands you're most likely to use.



The box above the shortcut menu—the Mini toolbar—contains commonly used tools from the Home tab. The Mini toolbar was designed to reduce the distance your mouse has to travel around the screen. Just right-click, and common formatting tools are near your mouse pointer. The Mini toolbar is particularly useful when a tab other than Home is displayed. If you use a tool on the Mini toolbar, the toolbar remains displayed in case you want to perform other formatting on the selection.

Customizing Your Quick Access Toolbar

The Ribbon is efficient, but many users prefer to have certain commands available at all times without having to click a tab. The solution is to customize your Quick Access Toolbar. Typically, the Quick Access Toolbar appears on the left side of the title bar, above the Ribbon. Alternatively, you can display the Quick Access Toolbar below the Ribbon; just right-click the Quick Access Toolbar and choose Show Quick Access Toolbar Below the Ribbon.

Displaying the Quick Access Toolbar below the Ribbon provides a bit more room for icons, but it also means that you see one less row of your worksheet.

By default, the Quick Access Toolbar contains four tools: AutoSave, Save, Undo, and Redo. You can customize the Quick Access Toolbar by adding other commands that you use often or removing the default controls. To add a command from the Ribbon to your Quick Access Toolbar, right-click the command and choose Add to Quick Access Toolbar. If you click the down arrow to the right of the Quick Access Toolbar, you will see a drop-down menu with some additional commands that you might want to place in your Quick Access Toolbar.

Excel has quite a few commands (mostly obscure ones) that aren't available on the Ribbon. In most cases, the only way to access these commands is to add them to your Quick Access Toolbar. Right-click the Quick Access Toolbar and choose Customize Quick Access Toolbar. You see the Excel Options dialog box, shown in Figure 1.10. This section of the Excel Options dialog box is your one-stop shop for Quick Access Toolbar customization.



See Chapter 8, “Customizing the Excel User Interface,” for more information about customizing your Quick Access Toolbar.

CAUTION

You can't reverse every action, however. Generally, anything that you do using the File button can't be undone. For example, if you save a file and realize that you've overwritten a good copy with a bad one, Undo can't save the day. You're just out of luck unless you have a backup of the file. Also, changes made by a macro can't be undone. In fact, executing a macro that changes the workbook clears the Undo list.

Changing Your Mind

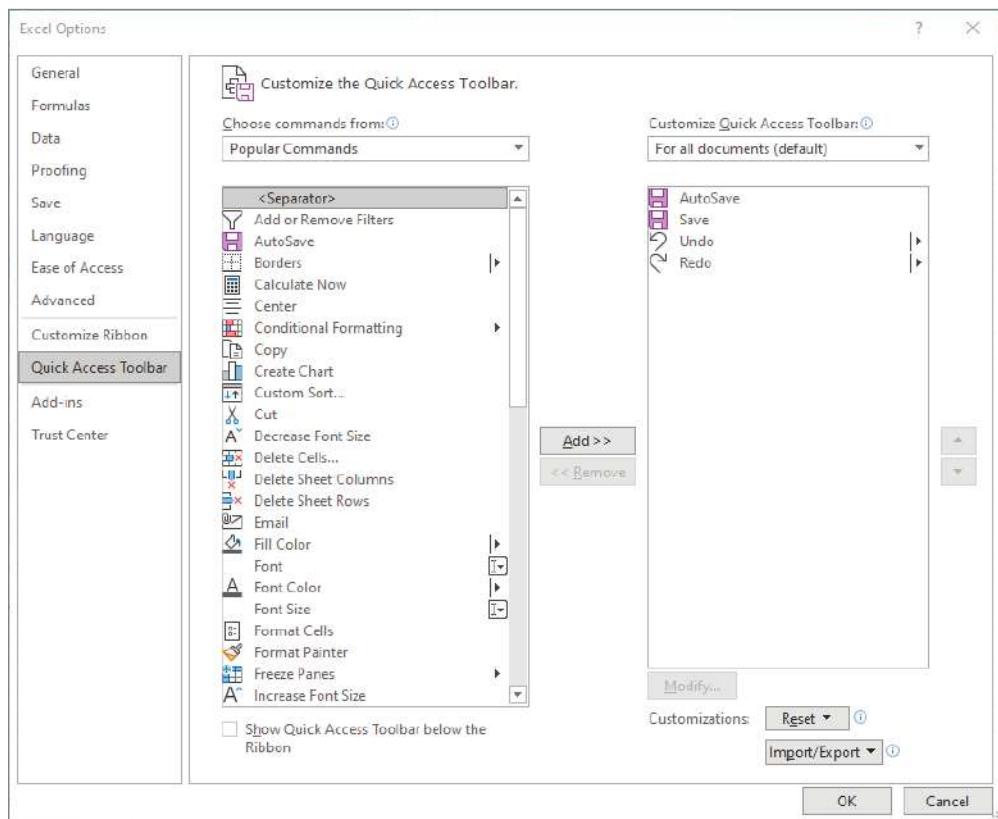
You can reverse almost every action in Excel by using the Undo command, located on the Quick Access Toolbar. Click Undo (or press Ctrl+Z) after issuing a command in error and it's as if you never issued the command. You can reverse the effects of the past 100 actions that you performed.

If you click the arrow on the right side of the Undo button, you see a list of the actions that you can reverse. Click an item in that list to undo that action and all the subsequent actions you performed.

The Redo button, also on the Quick Access Toolbar, performs the opposite of the Undo button: Redo reissues commands that have been undone. If nothing has been undone, this command is not available.

FIGURE 1.10

Add new icons to your Quick Access Toolbar by using the Quick Access Toolbar section of the Excel Options dialog box.

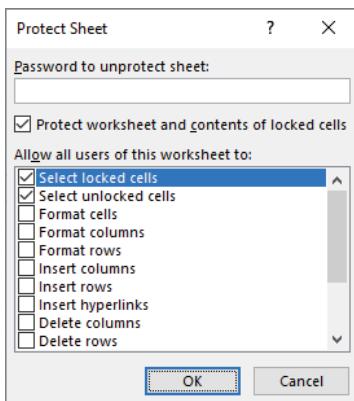


Working with Dialog Boxes

Many Excel commands display a dialog box, which is simply a way of getting more information from you. For example, if you choose Review \Rightarrow Protect \Rightarrow Protect Sheet, Excel can't carry out the command until you tell it what parts of the sheet you want to protect. Therefore, it displays the Protect Sheet dialog box, shown in Figure 1.11.

FIGURE 1.11

Excel uses a dialog box to get additional information about a command.



Excel dialog boxes vary in the way they work. You'll find two types of dialog boxes.

Typical dialog box A modal dialog box takes the focus away from the spreadsheet. When this type of dialog box is displayed, you can't do anything in the worksheet until you dismiss the dialog box. Clicking OK performs the specified actions and clicking Cancel (or pressing Esc) closes the dialog box without taking any action. Most Excel dialog boxes are this type.

Stay-on-top dialog box A modeless dialog box works in a manner similar to a toolbar. When a modeless dialog box is displayed, you can continue working in Excel, and the dialog box remains open. Changes made in a modeless dialog box take effect immediately. An example of a modeless dialog box is the Find and Replace dialog box. You can leave this dialog box open while you continue to use your worksheet. A modeless dialog box has a Close button but no OK button.

Most people find working with dialog boxes to be quite straightforward and natural. If you've used other programs, you'll feel right at home. You can manipulate the controls either with your mouse or directly from the keyboard.

Navigating dialog boxes

Navigating dialog boxes is generally easy—you simply click the control that you want to activate.

Although dialog boxes were designed with mouse users in mind, you can also use the keyboard. Every dialog box control has text associated with it, and this text always has one underlined letter (called a *hot key* or an *accelerator key*). You can access the control from the keyboard by pressing Alt and then the underlined letter. You can also press Tab to cycle through all the controls on a dialog box. Pressing Shift+Tab cycles through the controls in reverse order.

TIP

When a control that accepts text entry is selected, a cursor appears in the control. For drop-down and spin button controls, the default text is highlighted. Use Alt+down arrow to drop the list down and use the up and down arrows to change the spin button value. For all other controls, a dotted outline surrounds the control to let you know that it's selected. You can use the spacebar to activate a selected control.

Using tabbed dialog boxes

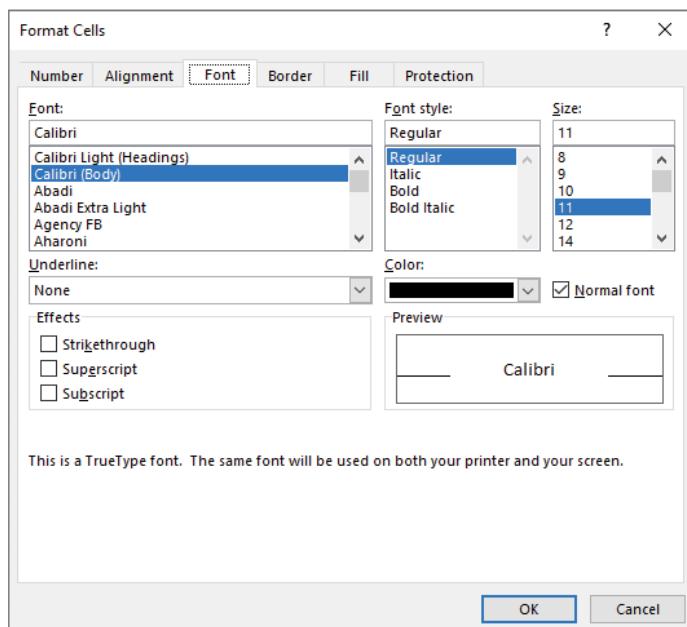
Several Excel dialog boxes are “tabbed” dialog boxes; that is, they include notebook-like tabs, each of which is associated with a different panel.

When you select a tab, the dialog box changes to display a new panel containing a new set of controls. The Format Cells dialog box, shown in Figure 1.12, is a good example. It has six tabs, which makes it functionally equivalent to six different dialog boxes.

Tabbed dialog boxes are quite convenient because you can make several changes in a single dialog box. After you make all your setting changes, click OK or press Enter.

FIGURE 1.12

Use the dialog box tabs to select different functional areas of the dialog box.



TIP

To select a tab by using the keyboard, press **Ctrl+PgUp** or **Ctrl+PgDn**, or simply press the first letter of the tab that you want to activate.

Using Task Panes

Another user interface element is the *task pane*. Task panes appear automatically in response to several commands. For example, to work with a picture that you've inserted, right-click the image and choose Format Picture. Excel responds by displaying the Format Picture task pane, shown in Figure 1.13. The task pane is similar to a dialog box except that you can keep it visible as long as you like.

Many of the task panes are complex. The Format Picture task pane has four icons along the top. Clicking an icon changes the command lists displayed next. Click an item in a command list and it expands to show the options.

There's no OK button in a task pane. When you're finished using a task pane, click the Close button (X) in the upper-right corner.

By default, a task pane is docked on the right side of the Excel window, but you can move it anywhere you like by clicking its title bar and dragging. Excel remembers the last position, so the next time you use that task pane, it will be right where you left it. To re-dock the task pane, double-click the task pane's title bar.

TIP

If you prefer to use your keyboard to work within a task pane, you may find that common dialog box keys such as Tab, spacebar, the arrow keys, and Alt key combinations don't seem to work. The trick is to press F6. After doing so, you'll find that the task pane works well using only a keyboard. For example, use the Tab key to activate a section title and then press Enter to expand the section.

Creating Your First Excel Workbook

This section presents an introductory, hands-on session with Excel. If you haven't used Excel, you may want to follow along on your computer to get a feel for how this software works.

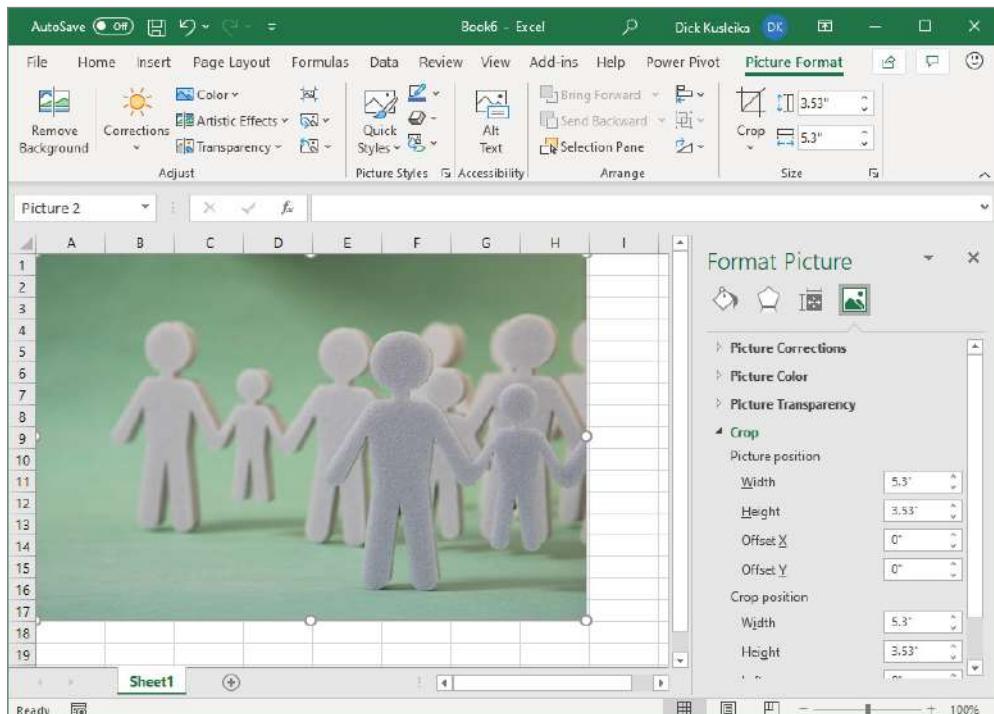
In this example, you create a simple monthly sales projection table plus a chart that depicts the data.

Getting started on your worksheet

Start Excel and make sure you have an empty workbook displayed. To create a new, blank workbook, press **Ctrl+N** (the shortcut key for File \leftrightarrow New \leftrightarrow Blank Workbook). Enter some sales projections in the new workbook.

FIGURE 1.13

The Format Picture task pane, docked on the right side of the window



The sales projections will consist of two columns of information. Column A will contain the month names, and column B will store the projected sales numbers. You start by entering some descriptive titles into the worksheet. Here's how to begin:

1. Select cell A1 (the upper-left cell in the worksheet) by using the navigation (arrow) keys, if necessary. The Name box displays the cell's address.
2. Type **Month** into cell A1 and press Enter. Depending on your setup, either Excel moves the selection to a different cell or the pointer remains in cell A1.
3. Select cell B1, type **Projected Sales**, and press Enter. The text extends beyond the cell width, but don't worry about that for now.

Filling in the month names

In this step, you enter the month names in column A.

1. **Select cell A2 and type Jan (an abbreviation for January).** At this point, you can enter the other month name abbreviations manually, or you can let Excel do some of the work by taking advantage of the AutoFill feature.

2. **Make sure that cell A2 is selected.** Notice that the active cell is displayed with a heavy outline. At the bottom-right corner of the outline, you'll see a small square known as the *fill handle*. Move your mouse pointer over the fill handle, click, and drag down until you've highlighted from cell A2 down to cell A13.
3. **Release the mouse button, and Excel automatically fills in the month names.**

Your worksheet should resemble the one shown in Figure 1.14.

FIGURE 1.14

Your worksheet after you've entered the column headings and month names

	A	B	C	D
1	Month	Projected Sales		
2	Jan			
3	Feb			
4	Mar			
5	Apr			
6	May			
7	Jun			
8	Jul			
9	Aug			
10	Sep			
11	Oct			
12	Nov			
13	Dec			
14				
15				

Entering the sales data

Next, you provide the sales projection numbers in column B. Assume that January's sales are projected to be \$50,000 and that sales will increase by 3.5 percent in each subsequent month.

1. **Select cell B2 and type 50000, the projected sales for January.** You could type a dollar sign and comma to make the number more legible, but you do the number formatting a bit later.
2. **To enter a formula to calculate the projected sales for February, move to cell B3 and type the following:**
`=B2*103.5%`
When you press Enter, the cell displays 51750. The formula returns the contents of cell B2, multiplied by 103.5%. In other words, February sales are projected to be 103.5% of the January sales—a 3.5% increase.
3. **The projected sales for subsequent months use a similar formula, but rather than retype the formula for each cell in column B, take advantage of the AutoFill feature.** Make sure that cell B3 is selected. Click the cell's fill handle, drag down to cell B13, and release the mouse button.

At this point, your worksheet should resemble the one shown in Figure 1.15. Keep in mind that, except for cell B2, the values in column B are calculated with formulas. To demonstrate, try changing the projected sales value for the initial month, January (in cell B2). You'll find that the formulas recalculate and return different values. All these formulas depend on the initial value in cell B2.

FIGURE 1.15

Your worksheet after you've created the formulas

	A	B	C	D
1	Month	Projected Sales		
2	Jan	50000		
3	Feb	51750		
4	Mar	53561.25		
5	Apr	55435.89		
6	May	57376.15		
7	Jun	59384.32		
8	Jul	61462.77		
9	Aug	63613.96		
10	Sep	65840.45		
11	Oct	68144.87		
12	Nov	70529.94		
13	Dec	72998.49		
14				
15				
16				

Formatting the numbers

The values in the worksheet are difficult to read because they aren't formatted. In this step, you apply a number format to make the numbers easier to read and more consistent in appearance.

1. **Select the numbers by clicking cell B2 and dragging down to cell B13.** Don't drag the fill handle this time, though, because you're selecting cells, not filling a range.
2. **Access the Ribbon and choose Home.** In the Number group, click the drop-down Number Format control (it initially displays General), and select Currency from the list. The numbers now display with a currency symbol and two decimal places. That's much better, but the decimal places aren't necessary for this type of projection.
3. **Make sure that the range B2 : B13 is selected, choose Home ➔ Number, and click the Decrease Decimal button.** One of the decimal places disappears. Click that button a second time and the values are displayed with no decimal places.

Making your worksheet look a bit fancier

At this point, you have a functional worksheet, but it could use some help in the appearance department. Converting this range to an "official" (and attractive) Excel table is a snap.

1. **Activate any cell within the range A1 : B13.**
2. **Choose Insert \Rightarrow Tables \Rightarrow Table.** Excel displays the Create Table dialog box to make sure that it guessed the range properly.
3. **Click OK to close the Create Table dialog box.** Excel applies its default table formatting and displays its Table Design contextual tab.

Your worksheet should look like Figure 1.16.

FIGURE 1.16

Your worksheet after you've converted the range to a table

	A	B	C
1	Month	Projected Sales	
2	Jan	\$50,000	
3	Feb	\$51,750	
4	Mar	\$53,561	
5	Apr	\$55,436	
6	May	\$57,376	
7	Jun	\$59,384	
8	Jul	\$61,463	
9	Aug	\$63,614	
10	Sep	\$65,840	
11	Oct	\$68,145	
12	Nov	\$70,530	
13	Dec	\$72,998	
14			

If you don't like the default table style, just select another one from the Table Design \Rightarrow Table Styles group. Notice that you can get a preview of different table styles by moving your mouse over the Ribbon. When you find one you like, click it, and the style will be applied to your table.



See Chapter 4, "Working with Excel Ranges and Tables," for more information on Excel tables.

Summing the values

The worksheet displays the monthly projected sales, but what about the total projected sales for the year? Because this range is a table, it's simple.

1. **Activate any cell in the table.**
2. **Choose Table Design \Rightarrow Table Style Options \Rightarrow Total Row.** Excel automatically adds a new row to the bottom of your table, including a formula that calculates the total of the Projected Sales column.
3. **If you'd prefer to see a different summary formula (for example, average), click cell B14 and choose a different summary formula from the drop-down list.**

Creating a chart

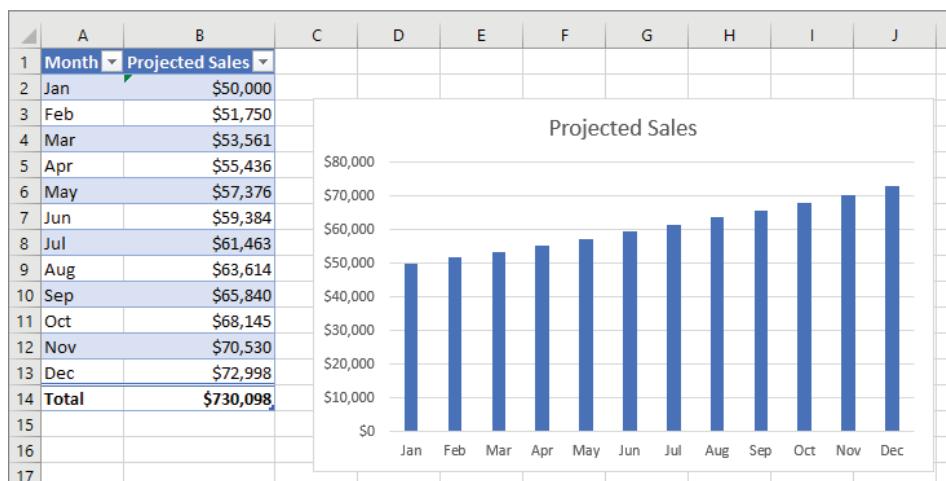
How about a chart that shows the projected sales for each month?

1. **Activate any cell in the table.**
2. **Choose Insert ▷ Charts ▷ Recommended Charts.** Excel displays some suggested chart type options.
3. **In the Insert Chart dialog box, click the second recommended chart (a column chart), and click OK.** Excel inserts the chart in the center of the window. To move the chart to another location, click its border and drag it.
4. **Click the chart and choose a style using the Chart Design ▷ Chart Styles options.**

Figure 1.17 shows the worksheet with a column chart. Your chart may look different, depending on the chart style you selected.

FIGURE 1.17

The table and chart



This workbook is available on this book's website at www.wiley.com/go/excel365bible. The filename is table and chart.xlsx.

Printing your worksheet

Printing your worksheet is easy (assuming that you have a printer attached and that it works properly).

1. **Make sure that the chart isn't selected.** If a chart is selected, the chart will print on a page by itself. To deselect the chart, just press Esc or click any cell.

2. **To make use of Excel's handy Page Layout view, click the Page Layout button on the right side of the status bar.** Excel displays the worksheet page by page so that you can easily see how your printed output will look. In Page Layout view, you can tell immediately whether the chart is too wide to fit on one page. If the chart is too wide, click and drag a corner of the chart to resize it or just move the chart below the table of numbers. Click the Normal button to return to the default view.
3. **When you're ready to print, choose File ➔ Print.** At this point, you can change some print settings. For example, you can choose to print in landscape rather than portrait orientation. Make the change, and you see the result in the preview window.
4. **When you're satisfied, click the large Print button in the upper-left corner.** The page is printed, and you're returned to your workbook.

Saving your workbook

Until now, everything that you've done has occurred in your computer's memory. If the power should fail, all may be lost—unless Excel's AutoRecover feature happened to kick in. It's time to save your work to a file on your hard drive.

1. **Click the Save button on the Quick Access Toolbar.** (This button looks like an old-fashioned floppy disk, popular in the previous century.) Because the workbook hasn't been saved yet and still has its default name, Excel responds with a Save this file dialog box that lets you choose the location for the workbook file. The Choose a Location drop-down lists some recently used locations, or you can click More options to see the Save as Backstage screen. From there, you can click Browse to navigate to any location on your computer.
2. **Click Browse.** Excel displays the Save As dialog box.
3. **In the File Name field, enter a name (such as Monthly Sales Projection).** If you like, you can specify a different location.
4. **Click Save or press Enter.** Excel saves the workbook as a file. The workbook remains open so that you can work with it some more.

NOTE

By default, Excel saves a backup copy of your work automatically every 10 minutes. To adjust the AutoRecover setting (or turn it off), choose File ➔ Options and click the Save tab of the Excel Options dialog box. However, you should never rely on Excel's AutoRecover feature. Saving your work frequently is a good idea.

If you've followed along, you probably have realized that creating this workbook was not difficult. But, of course, you've barely scratched the surface of Excel. The remainder of this book covers these tasks (and many, many more) in much greater detail.

TIP

Excel's Backstage view has a section for pinned folders at the top of the list of recent folders. If you use a particular folder often, you can pin it to the top of the list to make it more accessible. To pin a folder, find it in the list of recent folders, hover your mouse pointer over the folder, and click the push pin icon.

Entering and Editing Worksheet Data

IN THIS CHAPTER

- Understanding the types of data you can use
- Entering text and values into your worksheets
- Entering dates and times into your worksheets
- Modifying and editing information
- Using built-in and custom number formats
- Using Excel on a tablet

This chapter describes what you need to know about entering and modifying data in your worksheets. As you'll see, Excel doesn't treat all data equally. Therefore, you need to learn about the various types of data you can use in an Excel worksheet.

Exploring Data Types

An Excel workbook file can hold any number of worksheets, and each worksheet is made up of more than 17 billion cells. A cell can hold any of four basic types of data.

- A numeric value
- Text
- A formula
- An error

A worksheet can also hold charts, diagrams, pictures, buttons, and other objects. These objects aren't contained in cells. Instead, they reside on the worksheet's drawing layer, which is an invisible layer on top of each worksheet.



Numeric values

Numeric values represent a quantity of some type: sales amounts, number of employees, atomic weights, test scores, and so on. Values also can be dates (such as Feb 26, 2022) or times (such as 3:24 AM).



Excel can display values in many different formats. In the section “[Applying Number Formatting](#)” later in this chapter, you’ll see how different format options can affect the display of numeric values.

Excel’s Numeric Limitations

You may be curious about the types of values Excel can handle. In other words, how large can a number be? How accurate are large numbers?

Excel’s numbers are precise up to 15 digits. For example, if you enter a large value, such as 123,456,789,123,456,789 (18 digits), Excel actually stores it with only 15 digits of precision. This 18-digit number displays as 123,456,789,123,456,000. This precision may seem quite limiting, but in practice, it rarely causes any problems.

One situation in which the 15-digit precision can cause a problem is when entering credit card numbers. Most credit card numbers are 16 digits, but Excel can handle only 15 digits, so it substitutes a zero for the last credit card digit. Even worse, you may not even realize that Excel made the card number invalid. The solution? Enter the credit card numbers as text. The easiest way is to preformat the cell as Text. (Choose Home ➔ Number, and choose Text from the Number Format drop-down list.) Or you can precede the credit card number with an apostrophe. Either method prevents Excel from interpreting the entry as a number.

Here are some of Excel’s other numeric limits:

Largest positive number: 9.9E+307

Smallest negative number: -9.9E+307

Smallest positive number: 2.2251E-308

Largest negative number: -2.2251E-308

These numbers are expressed in scientific notation. For example, the largest positive number is “9.9 times 10 to the 307th power”—in other words, 99 followed by 306 zeros. Keep in mind, though, that this number has only 15 digits of precision.

Text entries

Most worksheets also include text in some of the cells. Text can serve as data (for example, a list of employee names), labels for values, headings for columns, or instructions about the worksheet. Text is often used to clarify what the values in a worksheet mean or where the numbers came from.

Text that begins with a number is still considered text. For example, if you type **12 Employees** into a cell, Excel considers the entry to be text rather than a numeric value. Consequently, you can’t use this cell for numeric calculations. If you need to indicate that the number 12 refers to employees, enter **12** into a cell and then type **Employees** into the cell to the right.

Formulas

Formulas are what make a spreadsheet a spreadsheet. Excel enables you to enter flexible formulas that use the values (or even text) in cells to calculate a result. When you enter a formula into a cell, the formula's result appears in the cell. If you change any of the cells used by a formula, the formula recalculates and shows the new result.

Formulas can be simple mathematical expressions, or they can use some of the powerful functions that are built into Excel. Figure 2.1 shows an Excel worksheet set up to calculate a monthly loan payment. The worksheet contains values, text, and formulas. The cells in column A contain text. Column B contains four values and two formulas. The formulas are in cells B6 and B10. Column D, for reference, shows the actual contents of the cells in column B.

FIGURE 2.1

You can use values, text, and formulas to create useful Excel worksheets.

A	B	C	D	E
1 Loan Payment Calculator				
2				
3			Column B Contents	
4 Purchase Amount:	\$475,000		475000	
5 Down Payment Pct:	20%		0.2	
6 Loan Amount:	\$380,000		=B4*(1-B5)	
7 Term (months):	360		360	
8 Interest Rate (APR):	6.25%		0.0625	
9				
10 Monthly Payment:	\$2,339.73		=PMT(B8/12,B7,-B6)	
11				
12				



This workbook, named `loan payment calculator.xlsx`, is available on this book's website at www.wiley.com/go/excel365bible.



You can find out much more about formulas in Part II, "Working with Formulas and Functions."

Error values

The fourth data type cells can hold is an error value. Error values are the results of formulas that contain an error, like the #VALUE! error that results from trying to do addition on a text entry. Error values are primarily used by Excel's calculation engine so that formulas that use the results of other formulas continue to show an error.

Entering Text and Values into Your Worksheets

If you've ever worked in a Windows application, you'll find that entering data into worksheet cells is simple and intuitive. And while there are differences in how Excel stores and displays the different data types, for the most part it just works.

Entering numbers

To enter a numeric value into a cell, select the appropriate cell, type the value, and then press Enter, Tab, or one of the arrow navigation keys. The value is displayed in the cell and appears in the Formula bar when the cell is selected. You can include decimal points and currency symbols when entering values, along with plus signs, minus signs, percent signs, and commas (to separate thousands). If you precede a value with a minus sign or enclose it in parentheses, Excel considers it to be a negative number.

Entering text

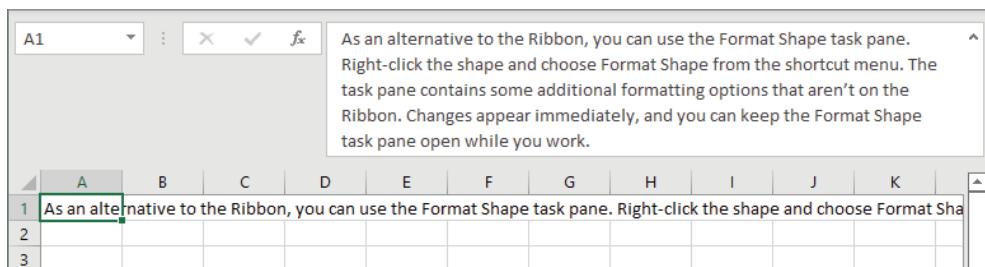
Entering text into a cell is just as easy as entering a value: activate the cell, type the text, and then press Enter or a navigation key. A cell can contain a maximum of about 32,000 characters—more than enough to store a typical chapter in this book. Even though a cell can hold a huge number of characters, you'll find that it's not actually possible to display all of these characters.

TIP

If you type an exceptionally long text entry into a cell, the Formula bar may not show all the text. To display more of the text in the Formula bar, click the bottom of the Formula bar and drag down to increase the height (see Figure 2.2). Also useful is the Ctrl+Shift+U keyboard shortcut. Pressing this key combination toggles the height of the Formula bar to show either one row or the previous size.

FIGURE 2.2

The Formula bar, expanded in height to show more information in the cell



What happens when you enter text that's longer than its column's current width? If the cells to the immediate right are blank, Excel displays the text in its entirety, appearing to spill the entry into adjacent cells. If an adjacent cell isn't blank, Excel displays as much of the text as

possible. (The full text is contained in the cell; it's just not displayed.) If you need to display a long text string in a cell that's adjacent to a nonblank cell, you have a few choices:

- Edit your text to make it shorter.
- Increase the width of the column (drag the border in the column letter display).
- Use a smaller font.
- Wrap the text within the cell so that it occupies more than one line. Choose Home \Rightarrow Alignment \Rightarrow Wrap Text to toggle wrapping on and off for the selected cell or range.

Using Enter mode

The left side of Excel's status bar normally displays "Ready," indicating that Excel is ready for you to enter or edit data in the worksheet. If you start typing numbers or text in a cell, the status bar changes to display "Enter" to indicate you're in *Enter mode*. The most common modes for Excel to be in are Ready, Enter, and Edit. See "Modifying Cell Contents" later in this chapter for more information about Edit mode.

In Enter mode, you are actively entering something into a cell. As you type, the text shows in the cell and in the Formula bar. You haven't actually changed the contents of the cell until you leave Enter mode, which commits the value to the cell. To leave Enter mode, you can press Enter, Tab, or just about any navigation key on your keyboard (like PgUp or Home). The value you typed is committed to the cell, and the status bar changes back to say "Ready."

You can also leave Enter mode by pressing the Esc key. Pressing Esc ignores your changes and returns the cell to its previous value.

Entering Dates and Times into Your Worksheets

Excel treats dates and times as special types of numeric values. Dates and times are values that are formatted so that they appear as dates or times. If you work with dates and times, you need to understand Excel's date and time system.

Entering date values

Excel handles dates by using a serial number system. The earliest date that Excel understands is January 1, 1900. This date has a serial number of 1. January 2, 1900, has a serial number of 2, and so on. This system makes it easy to deal with dates in formulas. For example, you can enter a formula to calculate the number of days between two dates.

Most of the time, you don't have to be concerned with Excel's serial number date system. You can simply enter a date in a common date format, and Excel takes care of the details behind the scenes. For example, if you need to enter June 1, 2022, you can enter the date by typing **June 1, 2022** (or use any of several different date formats). Excel interprets your entry and stores the value 44713, which is the serial number for that date.

Note

The date examples in this book use the U.S. English system. Your Windows regional settings will affect the way Excel interprets a date that you've entered. For example, depending on your regional date settings, June 1, 2022, may be interpreted as text rather than as a date. In such a case, you need to enter the date in a format that corresponds to your regional date settings—for example, 1 June 2022.



For more information about working with dates, see Chapter 13, “Using Formulas with Dates and Times.”

Entering time values

When you work with times, you extend Excel’s date serial number system to include decimals. In other words, Excel works with times by using fractional days. For example, the date serial number for June 1, 2022, is 44713. Noon on June 1, 2022 (halfway through the day), is represented internally as 44713.5 because the time fraction is added to the date serial number to get the full date/time serial number.

Again, you normally don’t have to be concerned with these serial numbers or fractional serial numbers for times. Just enter the time into a cell in a recognized format. In this case, type **June 1, 2022 12:00**.



See Chapter 13 for more information about working with time values.

Modifying Cell Contents

After you enter a value or text into a cell, you can modify it in several ways.

- Delete the cell’s contents.
- Replace the cell’s contents with something else.
- Edit the cell’s contents.

Note

You can also modify a cell by changing its formatting. However, formatting a cell affects only a cell’s appearance. Formatting doesn’t affect the cell’s contents. Later sections in this chapter cover formatting.

Deleting the contents of a cell

To delete the contents of a cell, just click the cell and press the Delete key. To delete the contents of more than one cell, select all the cells that you want to delete and then press Delete. Pressing Delete removes the cell’s contents but doesn’t remove any formatting (such as bold, italic, or the number format) that you may have applied to the cell.

For more control over what gets deleted, you can choose Home \Rightarrow Editing \Rightarrow Clear. This command's drop-down list has six choices.

- Clear All** Clears everything from the cell—its contents, formatting, and cell comment (if it has one).
- Clear Formats** Clears only the formatting and leaves the value, text, or formula.
- Clear Contents** Clears only the cell's contents and leaves the formatting. This has the same effect as pressing Delete.
- Clear Comments and Notes** Clears the comment or note (if one exists) attached to the cell.
- Clear Hyperlinks** Removes hyperlinks contained in the selected cells. The text and formatting remain, so the cell still looks like it has a hyperlink, but it no longer functions as a hyperlink.
- Remove Hyperlinks** Removes hyperlinks in the selected cells, including the cell formatting.

Note

Clearing formats doesn't clear the background colors in a range that has been designated as a table except in two circumstances: If you select the entire table or you've replaced the table style background colors manually, the table formatting will be cleared. See Chapter 4, "Working with Excel Ranges and Tables," for more about tables.

Replacing the contents of a cell

To replace the contents of a cell with something else, just activate the cell, type your new entry, and press Enter or a navigation key. Any formatting applied to the cell remains in place and is applied to the new content.

You can also replace cell contents by dragging and dropping or by copying and pasting data from another cell. In both cases, the cell formatting will be replaced by the format of the new data. To avoid pasting formatting, choose Home \Rightarrow Clipboard \Rightarrow Paste \Rightarrow Values (V), or Home \Rightarrow Clipboard \Rightarrow Paste \Rightarrow Formulas (F).

Editing the contents of a cell

If the cell contains only a few characters, replacing its contents by typing new data usually is easiest. However, if the cell contains lengthy text or a complex formula and you need to make only a slight modification, you probably want to edit the cell rather than re-enter information.

When you want to edit the contents of a cell, you can use one of the following ways to enter Edit mode:

- Double-click the cell to edit the cell contents directly in the cell.
- Select the cell and press F2 to edit the cell contents directly in the cell.
- Select the cell that you want to edit and then click inside the Formula bar to edit the cell contents in the Formula bar.

You can use whichever method you prefer. Some people find editing directly in the cell easier; others prefer to use the Formula bar to edit a cell.

NOTE

The Advanced tab of the Excel Options dialog box contains a section called Editing Options. These settings affect how editing works. (To access this dialog box, choose File ➔ Options.) If the Allow Editing Directly in Cells option isn't enabled, you can't edit a cell by double-clicking. In addition, pressing F2 allows you to edit the cell in the Formula bar (not directly in the cell).

All these methods cause Excel to go into Edit mode. (The word *Edit* appears at the left side of the status bar at the bottom of the window.) When Excel is in Edit mode, the Formula bar enables two icons: Cancel (the X) and Enter (the check mark). Figure 2.3 shows these two icons. Clicking the Cancel icon cancels editing without changing the cell's contents. (Pressing Esc has the same effect.) Clicking the Enter icon completes the editing and enters the modified contents into the cell. (Pressing Enter has the same effect, except that clicking the Enter icon doesn't change the active cell.)

FIGURE 2.3

When you're editing a cell, the Formula bar enables two new icons: Cancel (X) and Enter (check mark).

	A	B	C	D	E	F
1	Sunday	177				
2	Monday	139				
3	Tuesday	160				
4	Wednesday	144				
5	Thursday	114				
6	Friday	136				
7	Saturday	116				
8		=SUM(B1:B7)				
9						

When you begin editing a cell, the insertion point appears as a vertical bar, and you can perform the following tasks:

- Add new characters at the location of the insertion point. Move the insertion point by doing one of the following:
 - Using the navigation keys to move within the cell
 - Pressing Home to move the insertion point to the beginning of the cell
 - Pressing End to move the insertion point to the end of the cell
- Select multiple characters. Press Shift while you use the navigation keys.
- Select characters while you're editing a cell. Use the mouse. Just click and drag the mouse pointer over the characters that you want to select.
- Delete a character to the left of the insertion point. The Backspace key deletes the selected text or the character to the left of the insertion point if no characters are selected.

- Delete a character to the right of the insertion point. The Delete key also deletes the selected text. If no text is selected, it deletes the character to the right of the insertion point.

Learning some handy data-entry techniques

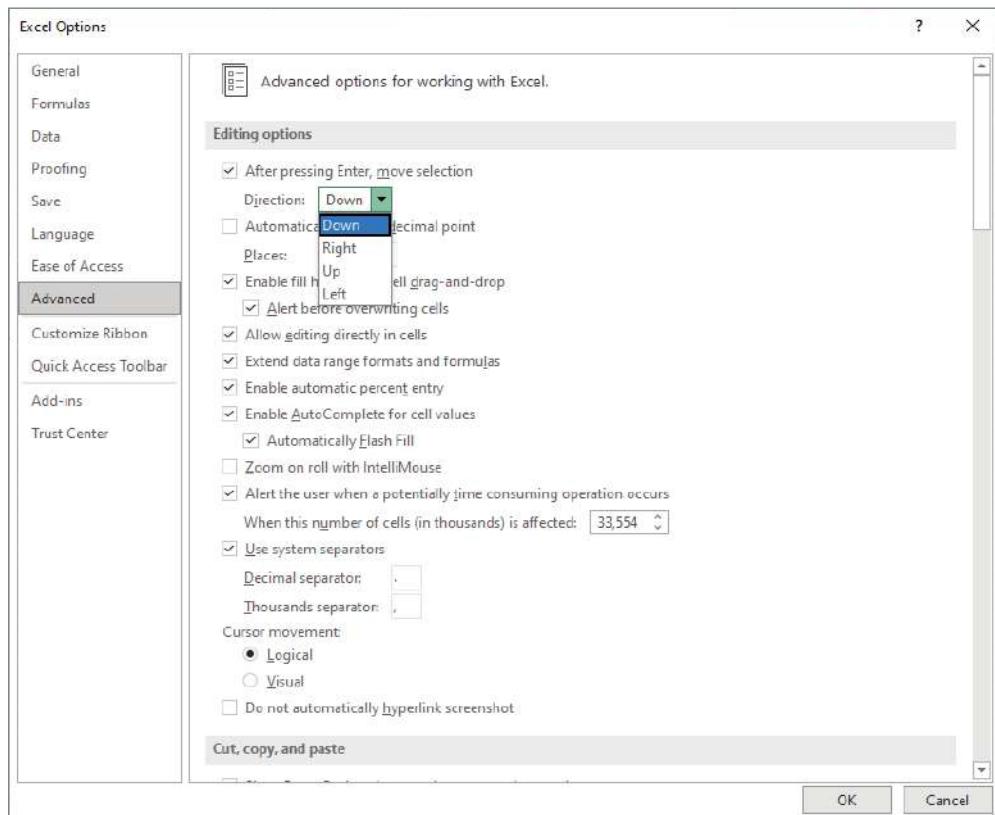
You can simplify the process of entering information into your Excel worksheets and make your work go quite a bit faster by using a number of useful tricks, which are described in the following sections.

Automatically moving the selection after entering data

By default, Excel automatically selects the next cell down when you press the Enter key after entering data into a cell. To change this setting, choose File ➔ Options and click the Advanced tab (see Figure 2.4). The check box that controls this behavior is labeled After pressing Enter, move selection. If you enable this option, you can choose the direction in which the selection moves (down, right, up, or left).

FIGURE 2.4

You can use the Advanced tab in the Excel Options dialog box to select several helpful input option settings.



Selecting a range of input cells before entering data

When a range of cells is selected, Excel automatically selects the next cell in the range when you press Enter, even if you disabled the After pressing Enter, move selection option. If the selection consists of multiple rows, Excel moves down the column; when it reaches the end of the selection in the column, it moves to the first selected cell in the next column.

To skip a cell, just press Enter without entering anything. To go backward, press Shift+Enter. If you prefer to enter the data by rows rather than by columns, press Tab rather than Enter. Excel continues to cycle through the selected range until you select a cell outside the range. Any of the navigation keys, like the arrow keys or the Home key, will change the selected range. If you want to navigate within the selected range, you must stick to Enter and Tab.

Using Ctrl+Enter to place information into multiple cells simultaneously

If you need to enter the same data into multiple cells, Excel offers a handy shortcut. Select all the cells that you want to contain the data; enter the value, text, or formula; and then press Ctrl+Enter. The same information is inserted into each cell in the selection.

Changing modes

You can press F2 to change between Enter mode and Edit mode. For example, if you're typing a long sentence in Enter mode and you realize that you spelled a word wrong, you can press F2 to change to Edit mode. In Edit mode, you can move through the sentence with your arrow keys to fix the misspelled word. You can also use the Ctrl+arrow keys to move one word at a time instead of one character at a time. You can continue to enter text in Edit mode or return to Enter mode by pressing F2 again, after which the navigation keys can be used to move to a different cell.

Entering decimal points automatically

If you need to enter lots of numbers with a fixed number of decimal places, Excel has a useful tool that works like some old adding machines. Access the Excel Options dialog box and click the Advanced tab. Select the Automatically Insert a Decimal Point check box and make sure that the Places box is set for the correct number of decimal places for the data you need to enter.

When this option is set, Excel supplies the decimal points for you automatically. For example, if you specify two decimal places, entering **12345** into a cell is interpreted as 123.45. To restore things to normal, just clear the Automatically Insert a Decimal Point check box in the Excel Options dialog box. Changing this setting doesn't affect any values that you already entered.

CAUTION

The fixed decimal places option is a global setting and applies to all workbooks (not just the active workbook). If you forget that this option is turned on, you can easily end up entering incorrect values—or cause some major confusion if someone else uses your computer.

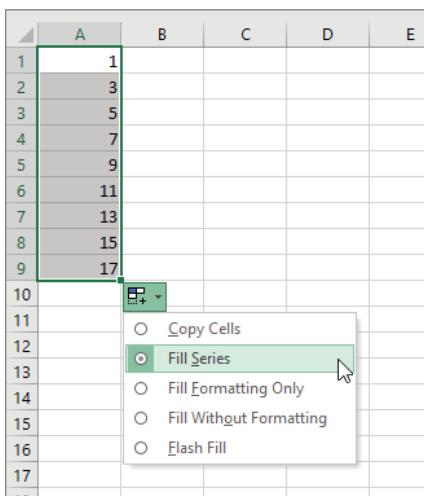
Using AutoFill to enter a series of values

The Excel AutoFill feature makes inserting a series of values or text items in a range of cells easy. It uses the *fill handle* (the small box at the lower right of the active cell or range). You can drag the fill handle to copy the cell or automatically complete a series.

Figure 2.5 shows an example. Enter **1** into cell A1, and enter **3** into cell A2. Then select both cells and drag down the fill handle to create a linear series of odd numbers. The figure also shows an icon that, when clicked, displays some additional AutoFill options. This icon appears only if the Show Paste Options button when content is pasted option is selected in the Advanced tab of the Excel Options dialog box.

FIGURE 2.5

This series was created by using AutoFill.



A screenshot of an Excel spreadsheet. Column A contains the following data: Row 1: 1; Row 2: 3; Row 3: 5; Row 4: 7; Row 5: 9; Row 6: 11; Row 7: 13; Row 8: 15; Row 9: 17. Row 10 is empty. Row 11 has a green border and contains a small square icon with a plus sign and a downward arrow. A context menu is open over this icon, listing the following options: Copy Cells, Fill Series (which is highlighted with a green background), Fill Formatting Only, Fill Without Formatting, and Flash Fill.

A	B	C	D	E
1	1			
2	3			
3	5			
4	7			
5	9			
6	11			
7	13			
8	15			
9	17			
10				
11				
12				
13				
14				
15				
16				
17				
18				

Excel uses the cells' data to guess the pattern. If you start with 1 and 2, it will guess you want each cell to go up by 1. If, as in the previous example, you start with 1 and 3, it guesses that you want the increment to be 2. Excel does a good job of guessing date patterns too. If you start with 1/31/2022 and 2/28/2022, it will fill the last day of the successive months.

TIP

If you drag the fill handle while you press and hold the right mouse button, Excel displays a shortcut menu with additional fill options. You can also use Home \Rightarrow Editing \Rightarrow Fill for even more control over automatically filling a range.

Using AutoComplete to automate data entry

The Excel AutoComplete feature makes entering the same text into multiple cells easy. With AutoComplete, you type the first few letters of a text entry into a cell, and Excel automatically

completes the entry based on other entries that you already made in the column. Besides reducing typing, this feature ensures that your entries are spelled correctly and are consistent.

Here's how it works: Suppose you're entering product information into a column. One of your products is named Widgets. The first time you enter **Widgets** into a cell, Excel remembers it. Later, when you start typing **Widgets** in that same column, Excel recognizes it by the first few letters and finishes typing it for you. Just press Enter, and you're done. To override the suggestion, just keep typing.

AutoComplete also changes the case of letters for you automatically. If you start entering **widgets** (with a lowercase *w*) in the second entry, Excel makes the *w* uppercase to be consistent with the previous entry in the column.

TIP

You also can access a mouse-oriented version of AutoComplete by right-clicking the cell and choosing Pick From Drop-down List from the shortcut menu. Excel then displays a drop-down box that has all of the text entries in the current column, and you just click the one that you want.

Keep in mind that AutoComplete works only within a contiguous column of cells. If you have a blank row, for example, AutoComplete identifies only the cell contents below the blank row.

Sometimes, Excel will use AutoComplete to try to finish a word when you don't want it to do so. If you type **canister** in a cell and then below it type the shorter word **can**, Excel will attempt to AutoComplete the entry to *canister*. When you want to type a word that starts with the same letters as an AutoComplete entry but is shorter, simply press the Delete key when you've reached the end of the word and then press Enter or a navigation key.

If you find the AutoComplete feature distracting, you can turn it off by using the Advanced tab of the Excel Options dialog box. Remove the check mark from the Enable AutoComplete for Cell Values box.

Forcing text to appear on a new line within a cell

If you have lengthy text in a cell, you can force Excel to display it in multiple lines within the cell: press Alt+Enter to start a new line in a cell.

When you add a line break, Excel automatically changes the cell's format to Wrap Text. But unlike normal text wrap, your manual line break forces Excel to break the text at a specific place within the text, which gives you more precise control over the appearance of the text than if you rely on automatic text wrapping.

TIP

To remove a manual line break, edit the cell and press Delete when the insertion point is located at the end of the line that contains the manual line break. You won't see any symbol to indicate the position of the manual line break, but the text that follows it will move up when the line break is deleted.

Using AutoCorrect for shorthand data entry

You can use the AutoCorrect feature to create shortcuts for commonly used words or phrases. For example, if you work for a company named Consolidated Data Processing Corporation, you can create an AutoCorrect entry for an abbreviation, such as *cdp*. Then, whenever you type *cdp* and take an action to trigger AutoCorrect (such as typing a space, pressing Enter, or selecting a different cell), Excel automatically changes the text to Consolidated Data Processing Corporation.

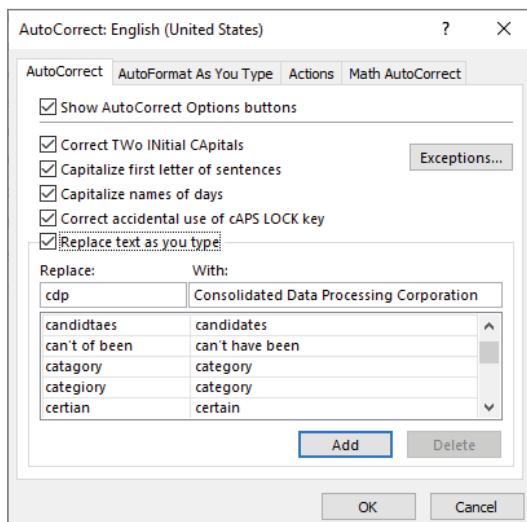
Excel includes quite a few built-in AutoCorrect terms (mostly to correct common misspellings), and you can add your own. To set up your custom AutoCorrect entries, access the Excel Options dialog box (choose File \Rightarrow Options) and click the Proofing tab. Then click the AutoCorrect Options button to display the AutoCorrect dialog box. In the dialog box, click the AutoCorrect tab, check the Replace Text as You Type option, and then enter your custom entries. (Figure 2.6 shows an example.) You can set up as many custom entries as you like. Just be careful not to use an abbreviation that might appear normally in your text.

TIP

Excel shares your AutoCorrect list with other Microsoft Office applications. For example, any AutoCorrect entries you created in Word also work in Excel.

FIGURE 2.6

AutoCorrect allows you to create shorthand abbreviations for text you enter often.



Entering numbers with fractions

Most of the time, you'll want non-integer values to be displayed with decimal points. But Excel can also display values with fractions. To enter a fractional value into a cell, leave a space between the whole number and the fraction. For example, to enter $6\frac{7}{8}$, enter **6 7/8** and then press Enter. When you select the cell, 6.875 appears in the Formula bar, and the cell entry appears as a fraction. If you have a fraction only (for example, $\frac{1}{8}$), you must enter a zero first, like this—**0 1/8**—or Excel will likely assume that you're entering a date. When you select the cell and look at the Formula bar, you see 0.125. In the cell, you see $\frac{1}{8}$.

Using a form for data entry

Many people use Excel to manage lists in which the information is arranged in rows. Excel offers a simple way to work with this type of data using a data entry form that Excel can create automatically. This data form works with either a normal range of data or a range that has been designated as a table. (Choose Insert \Rightarrow Tables \Rightarrow Table.) Figure 2.7 shows an example.

FIGURE 2.7

Excel's built-in data form can simplify many data-entry tasks.

	A	B	C	D	E	F	G	H	I	J	K
1	Agent	Date Listed	Area	List Price	Bedrooms	Baths	SqFt	Type	Pool	Sold	
2	Adams	1/22/2022	Central	295,000							TRUE
3	Adams	4/10/2022	Central	384,000							FALSE
4	Adams	10/4/2022	Downtown	258,000							TRUE
5	Adams	1/19/2022	Downtown	271,000							TRUE
6	Adams	4/14/2022	Downtown	310,000							TRUE
7	Adams	3/17/2022	Downtown	312,000							TRUE
8	Adams	9/26/2022	Downtown	352,000							FALSE
9	Adams	1/10/2022	Downtown	372,000							TRUE
10	Adams	1/12/2022	Downtown	397,000							TRUE
11	Adams	9/22/2022	North	288,000							TRUE
12	Adams	4/15/2022	North	289,000							TRUE
13	Adams	9/12/2022	North	291,000							TRUE
14	Adams	5/21/2022	South	282,000							TRUE
15	Adams	4/16/2022	South	388,000							TRUE
16	Adams	11/22/2022	South	389,000							TRUE
17	Barnes	11/18/2022	Downtown	275,000							TRUE
18	Barnes	9/3/2022	Downtown	289,000	3	3	2,100	Split Level	TRUE		FALSE
19	Barnes	1/4/2022	Downtown	313,000	2	3	2,200	Ranch	FALSE		TRUE
20	Barnes	3/6/2022	Downtown	342,000	4	2	2,400	2 Story	FALSE		TRUE
21	Barnes	6/10/2022	Downtown	354,000	3	2	2,500	Bank	TRUE		TRUE

Sheet1

Agent: Adams

Date Listed: 1/22/2022

Area: Central

List Price: 295000

Bedrooms: 4

Baths: 3.5

SqFt: 2100

Type: Condo

Pool: TRUE

Sold: TRUE

?

X

1 of 79

New

Delete

Restore

Find Prev

Find Next

Criteria

Close

Unfortunately, the command to access the data form is not on the Ribbon. To use the data form, you must add it to your Quick Access Toolbar, add it to the Ribbon, or search for Form in the Search box. Here's how to add this command to your Quick Access Toolbar:

1. Right-click the Quick Access Toolbar and choose Customize Quick Access Toolbar. The Quick Access Toolbar panel of the Excel Options dialog box appears.
2. In the Choose Commands From drop-down list, choose Commands Not in the Ribbon.

3. In the list box on the left, select Form.
4. Click the Add button to add the selected command to your Quick Access Toolbar.
5. Click OK to close the Excel Options dialog box.

After you perform these steps, a new icon appears on your Quick Access Toolbar.

To use a data entry form, follow these steps:

1. Arrange your data so that Excel can recognize it as a table by entering headings for the columns into the first row of your data entry range.
2. Select any cell in the table and click the Form button on your Quick Access Toolbar. Excel displays a dialog box customized to your data (refer to the example in Figure 2.7).
3. Fill in the information. Press Tab to move between the text boxes. If a cell contains a formula, the formula result appears as text (not as an edit box). In other words, you can't modify formulas using the data entry form.
4. When you complete the data form, click the New button. Excel enters the data into a row in the worksheet and clears the dialog box for the next row of data.

You can also use the form to edit existing data.



This workbook is available on this book's website at www.wiley.com/go/excel365bible.
The file is named `data_form.xlsx`.

Entering the current date or time into a cell

If you need to date-stamp or time-stamp your worksheet, Excel provides two shortcut keys that do this task for you:

- Current date: Ctrl+; (semicolon)
- Current time: Ctrl+Shift+; (semicolon)

To enter both the date and time, press Ctrl+;, type a space, and then press Ctrl+Shift+;.

The date and time are from the system time in your computer. If the date or time isn't correct in Excel, use the Windows Settings to make the adjustment.

NOTE

When you use the Ctrl+; or the Ctrl+Shift+; shortcut to enter a date or time into your worksheet, Excel enters a static value into the worksheet. In other words, the date or time entered doesn't change when the worksheet is recalculated. In most cases, this setup is probably what you want, but you should be aware of this limitation. If you want the date or time display to update, use one of these functions:

```
=TODAY()  
=NOW()
```

Applying Number Formatting

Applying number formatting changes the appearance of values contained in cells. Excel provides a variety of number formatting options. In the following sections, you will see how to use many of Excel's formatting options to improve the appearance and readability of your worksheets quickly.

TIP

The formatting that you apply works with the selected cell or cells. Therefore, you need to select the cell (or range of cells) before applying the formatting. Also remember that changing the number format does not affect the underlying value. Number formatting affects only the appearance.

Values that you enter into cells normally are unformatted. In other words, they simply consist of a string of numerals. Typically, you want to format the numbers so that they're easier to read or are more consistent in terms of the number of decimal places shown.

Figure 2.8 shows a worksheet that has two columns of values. The first column consists of unformatted values. The cells in the second column are formatted to make the values easier to read. The third column describes the type of formatting applied.

FIGURE 2.8

Use numeric formatting to make it easier to understand what the values in the worksheet represent.

	A	B	C	D
1				
2	Unformatted	Formatted	Type	
3	1200	\$1,200.00	Currency	
4	0.231	23.1%	Percentage	
5	44703	5/22/2022	Short Date	
6	44703	Sunday, May 22, 2022	Long Date	
7	123439832	\$ 123,439,832.00	Accounting	
8	5559832	555-9832	Phone Number	
9	434988723	434-98-8723	Social Security Number	
10	0.552	1:14:53 PM	Time	
11	0.25	1/4	Fraction	
12	12332354090	1.23E+10	Scientific	
13				



This workbook is available on this book's website at www.wiley.com/go/excel365bible.
The file is named number_formatting.xlsx.

TIP

If you select a cell that has a formatted value, the Formula bar displays the value in its unformatted state because the formatting affects only the way the value appears in the cell—not the actual value contained in the cell. There are a few exceptions, however. When you enter a date or a time, Excel always displays the value as a date or a time, even though it's stored internally as a value. Also, values that use the Percentage format display with a percent sign in the Formula bar.

Using automatic number formatting

Excel is able to perform some formatting for you automatically. For example, if you enter **12.2%** into a cell, Excel knows that you want to use a percentage format and applies it for you automatically. If you use commas to separate thousands (such as 123,456), Excel applies comma formatting for you. And if you precede your value with a dollar sign, the cell is formatted for currency (assuming that the dollar sign is your system currency symbol).

Anything you enter that can possibly be construed as a date will be treated as such. And depending on how you enter it, Excel will choose a date format to match. If you enter **1/31/2022**, Excel will interpret that as a date and format the cell as 1/31/2022 (just as it was entered). If you enter **Jan 31, 2022**, Excel will format it as 31-Jan-22 (if you omit the comma, Excel won't recognize it as a date). The less obvious example of entering **1-31** causes Excel to display 31-Jan. If you need to enter **1-31** in a cell and it's not supposed to be a date, type an apostrophe (') first.

TIP

A handy default feature in Excel makes entering percentage values into cells easier. If a cell is formatted to display as a percent, you can simply enter a normal value (for example, **12.5** for 12.5%). To enter values less than 1%, precede the value with a zero (for example, **0.52** for 0.52%). If this automatic percent entry feature isn't working (or if you prefer to enter the actual value for percents), access the Excel Options dialog box and click the Advanced tab. In the Editing Options section, locate the Enable Automatic Percent Entry check box and add or remove the check mark.

Formatting numbers by using the Ribbon

The Home \Rightarrow Number group on the Ribbon contains controls that let you quickly apply common number formats.

The Number Format drop-down list contains 11 common number formats (see Figure 2.9). Additional options in the Home \Rightarrow Number group include an Accounting Number Format drop-down list (to select a currency format), a Percent Style button, and a Comma Style button. The group also contains a button to increase the number of decimal places and another to decrease the number of decimal places.

When you select one of these controls, the active cell takes on the specified number format. You also can select a range of cells (or even entire rows or columns) before clicking these buttons. If you select more than one cell, Excel applies the number format to all the selected cells.

Using shortcut keys to format numbers

Another way to apply number formatting is to use shortcut keys. Table 2.1 summarizes the shortcut-key combinations that you can use to apply common number formatting to the selected cells or range. Notice that these Ctrl+Shift keys are located together, in the lower left of your keyboard.

Part I: Getting Started with Excel

FIGURE 2.9

You can find number formatting commands in the Number group of the Home tab.

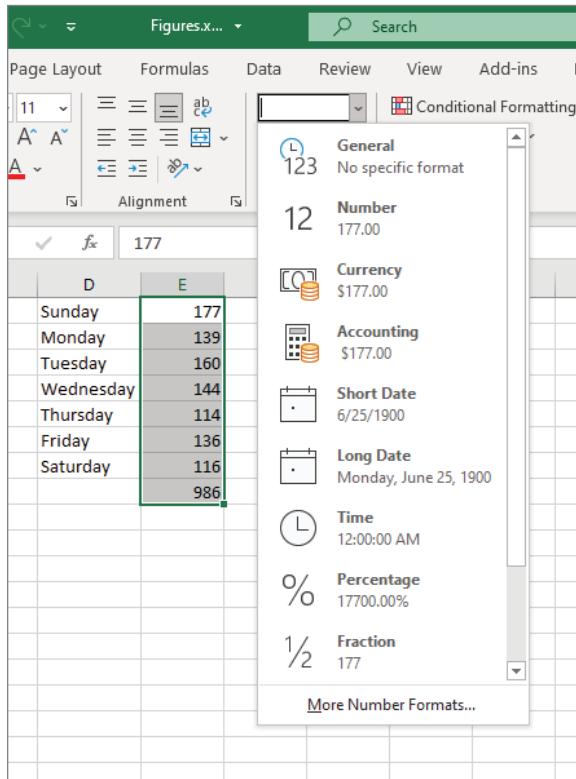


TABLE 2.1 Number Formatting Keyboard Shortcuts

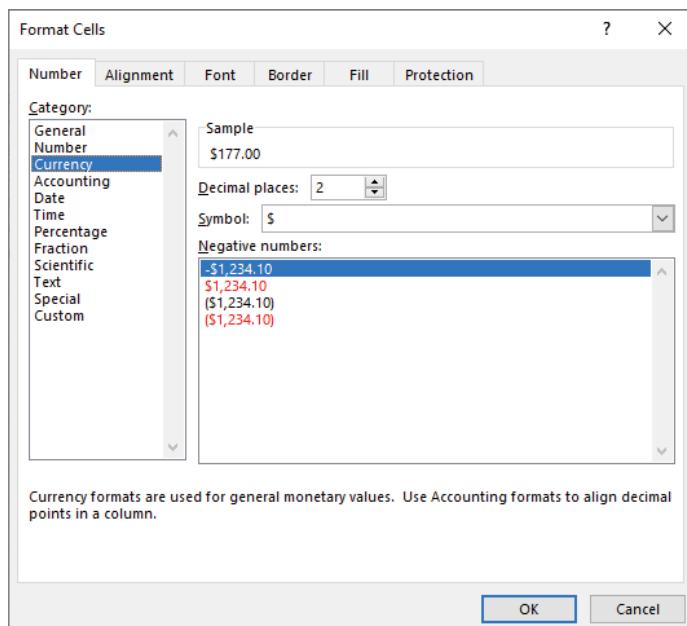
Key Combination	Formatting Applied
Ctrl+Shift+~	General number format (that is, unformatted values)
Ctrl+Shift+\$	Currency format with two decimal places (negative numbers appear in red and inside parentheses)
Ctrl+Shift+%	Percentage format, with no decimal places
Ctrl+Shift+^	Scientific notation number format, with two decimal places
Ctrl+Shift+#	Date format with the day, month, and year
Ctrl+Shift+@	Time format with the hour, minute, and AM or PM
Ctrl+Shift+!	Two decimal places, thousands separator, and a hyphen for negative values

Formatting numbers by using the Format Cells dialog box

In most cases, the number formats that are accessible from the Number group on the Home tab are just fine. Sometimes, however, you want more control over how your values appear. Excel offers a great deal of control over number formats using the Format Cells dialog box, as shown in Figure 2.10. For formatting numbers, you need to use the Number tab.

FIGURE 2.10

When you need more control over number formats, use the Number tab of the Format Cells dialog box.



You can bring up the Format Cells dialog box in several ways. Start by selecting the cell or cells that you want to format and then do one of the following:

- Choose Home \Rightarrow Number and click the small dialog box launcher icon (in the lower-right corner of the Number group).
- Choose Home \Rightarrow Number, click the Number Format drop-down list, and choose More Number Formats from the drop-down list.
- Right-click the cell, and choose Format Cells from the shortcut menu.
- Press **Ctrl+1**.

The Number tab of the Format Cells dialog box displays 12 categories of number formats. When you select a category from the list box, the right side of the tab changes to display options appropriate to that category.

The Number category has three options that you can control: the number of decimal places displayed, whether to use a thousands separator, and how you want negative numbers displayed. The Negative Numbers list box has four choices (two of which display negative values in red), and the choices change depending on the number of decimal places and whether you choose to separate thousands.

The top of the tab displays a sample of how the active cell will appear with the selected number format (visible only if a cell with a value is selected). After you make your choices, click OK to apply the number format to all the selected cells.

When Numbers Appear to Add Incorrectly

Applying a number format to a cell doesn't change the value—it only changes how the value appears in the worksheet. For example, if a cell contains 0.874543, you may format it to appear as 87%. If that cell is used in a formula, the formula uses the full value (0.874543), not the displayed value (87%).

In some situations, formatting may cause Excel to display calculation results that appear incorrect, such as when totaling numbers with decimal places. For example, if values are formatted to display two decimal places, you may not see the actual numbers used in the calculations. But because Excel uses the full precision of the values in its formula, the sum of the two values may appear to be incorrect.

Several solutions to this problem are available. You can format the cells to display more decimal places. You can use the ROUND function on individual numbers and specify the number of decimal places to which Excel should round. Or, you can instruct Excel to change the worksheet values to match their displayed format. To do so, access the Excel Options dialog box and click the Advanced tab. Check the Set Precision as Displayed check box (located in the When Calculating This Workbook section).

CAUTION

Selecting the Set Precision as Displayed option changes the numbers in your worksheets to match their appearance on-screen permanently. This setting applies to all sheets in the active workbook. Most of the time, this option is *not* what you want. Make sure you understand the consequences of using the Set Precision as Displayed option.

The following are the number format categories, along with some general comments:

General The default format; it displays numbers as integers, as decimals, or in scientific notation if the value is too wide to fit in the cell.

Number Enables you to specify the number of decimal places, whether to use a comma to separate thousands, and how to display negative numbers (with a minus sign, in red, in parentheses, or in red and in parentheses).

Currency Enables you to specify the number of decimal places, choose a currency symbol, and specify how to display negative numbers (with a minus sign, in red, in parentheses, or in red and in parentheses). This format always uses a comma to separate thousands.

Accounting Differs from the Currency format in that the currency symbols always align vertically at the left edge of the cell.

Date Enables you to choose from several different date formats.

Time Enables you to choose from several different time formats.

Percentage Enables you to choose the number of decimal places and always displays a percent sign.

Fraction Enables you to choose from among nine fraction formats.

Scientific Displays numbers in exponential notation (with an E): 2.00E+05 = 200,000; 2.05E+05 = 205,000. You can choose the number of decimal places to display to the left of E. The second example can be read as “2.05 times 10 to the fifth.”

Text When applied to a value, causes Excel to treat the value as text (even if it looks like a number). This feature is useful for such items as part numbers and credit card numbers.

Special Contains additional number formats. In the U.S. version of Excel, the additional number formats are ZIP Code, ZIP Code +4, Phone Number, and Social Security Number.

Custom Enables you to define custom number formats that aren’t included in any other category.

TIP

If a cell displays a series of hash marks (such as #####), it usually means that the column isn’t wide enough to display the value in the number format that you selected. Either make the column wider or change the number format. Hashmarks also indicate a negative time value or an invalid date (that is, a date prior to January 1, 1900).

Adding your own custom number formats

Sometimes you may want to display numerical values in a format that isn’t included in any of the other categories. If so, the answer is to create your own custom format. Basic custom number formats contain four sections separated by semicolons. Those four sections determine how a number will be formatted if it is a positive value, negative value, a zero, or text.

Using Excel on a Tablet

In recent years, Microsoft has released a mobile app called Office for use on iOS and Android devices. This app combines the Excel, Word, and PowerPoint apps previously released by Microsoft. These apps are free to download and use to view files, but if you want to save changes, you’ll need an Office 365 account.

Part I: Getting Started with Excel

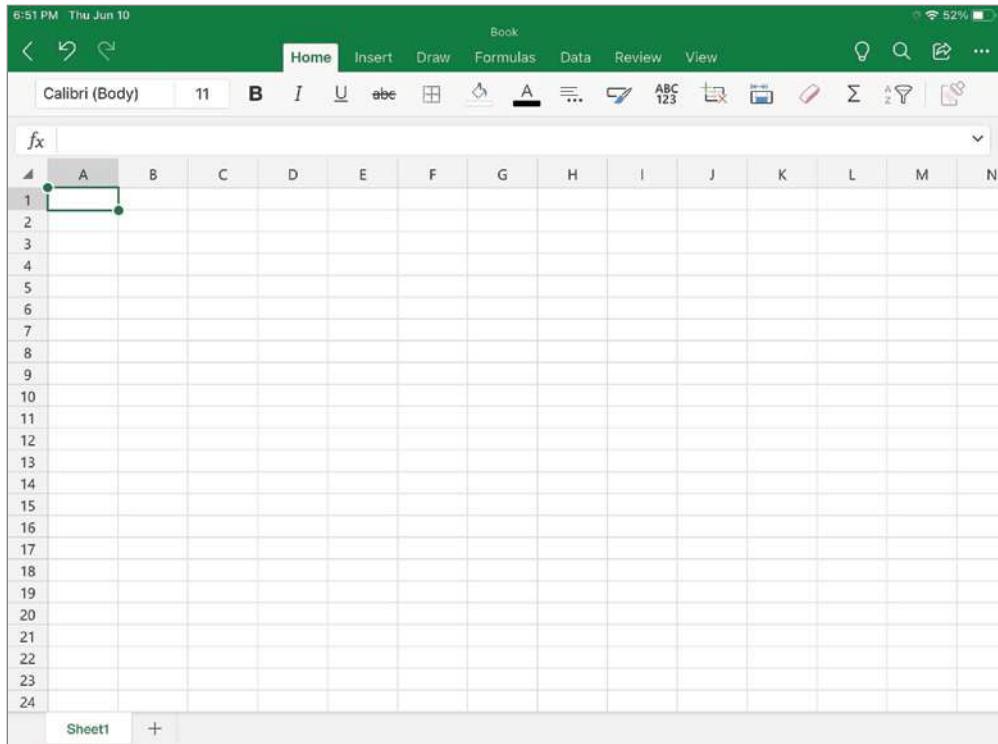
Excel on a tablet will feel quite familiar. You'll see the grid of cells, sheet tabs across the bottom, and a Ribbon at the top. There are, however, some differences that make Excel work better on a smaller screen. For example, the Ribbon is smaller to conserve screen space and the touch targets on the Ribbon and menus are larger when your finger is the pointing device.

Exploring Excel's tablet interface

Figure 2.11 shows a blank Excel workbook on a tablet. The Ribbon tabs are similar to the desktop version of Excel, but the controls are a narrower strip below the tabs. When there are too many controls to display, you can scroll by sliding left or right to see more.

FIGURE 2.11

A blank workbook on a tablet



The selected cell, cell A1 in Figure 2.11, has handles in the upper-left and lower-right corners. To select a single cell, simply tap it. To expand the range, drag one of the handles. When you expand a range, Excel displays a context menu above or below the selected range containing often used commands. To see the context menu for a single cell, long press the cell.

To edit a cell, double-tap the cell to enter Edit mode. You can edit directly in the cell or tap in the Formula bar and make your changes there. When you enter Edit mode, the onscreen keyboard will appear. Another way to edit a cell is to long press to get the context menu and tap the Edit control.

To commit your edits to the cell, you can tap return on the onscreen keyboard, tap another cell, or tap the check mark control to the right of the Formula bar. There is also a red x control to the right of the Formula bar to exit Edit mode without committing your changes.

Entering formulas on a tablet

The Formulas tab on the Ribbon contains a control for each of the formula categories you're used to from desktop Excel. Tapping one of these controls presents a scrollable menu with all the functions in that category. The Formulas tab also has an AutoSum control that displays a menu of common functions and a Recent control that displays a menu of functions you've used recently.

Another way to enter formulas is by tapping the fx control to the left of the Formula bar. It displays recently used functions at the top and a menu of function categories and works similarly to the buttons on Formulas tab on the Ribbon.

Yet a third way to enter a formula is to simply type an equal sign. With this method, you can type the first letter of the worksheet function you want and a menu with all the functions that begin with that letter is displayed. Or you can type the entire worksheet function, but you may find that tedious. When you type an equal sign or select a function, Excel is in Point mode and tapping cells adds their addresses to the formula, just as in desktop Excel.

Figure 2.12 shows a cell with an INDEX worksheet function in it. When you select a function from a menu, the arguments are shown as buttons. You can tap the buttons to select cells or enter values for that argument.

In this example, the first argument was completed by tapping B3 and dragging the handle down to B7. Then the second argument button was selected and is shown slightly darker than the unselected argument buttons. While you wouldn't want to build a large spreadsheet using these methods, Excel has provided a very nice interface for entering a formula or two.

Introducing the Draw Ribbon

On mobile devices, the Draw tab on the Ribbon is shown by default. It contains drawing controls like pens, pencils, and highlighters for drawing in Excel. Like shapes and illustrations, drawn objects are on the drawing layer above the cells and aren't inside any particular cell.

You can draw a shape and convert it into an actual shape using Draw \leftrightarrow Convert \leftrightarrow Ink to Shape. Use the selection tool to select the drawing and click the Ink to Shape control or the lightning bolt that appears next to the shape. Figure 2.13 shows a crudely drawn circle converted to an oval shape.

Part I: Getting Started with Excel

FIGURE 2.12

Function arguments are displayed as buttons.

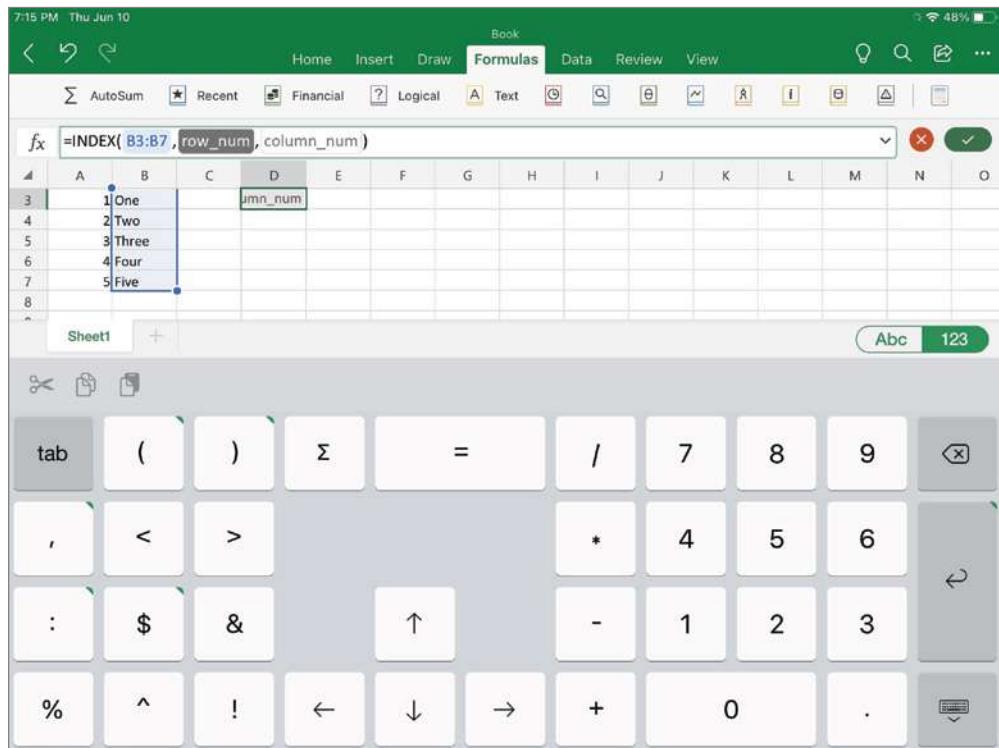
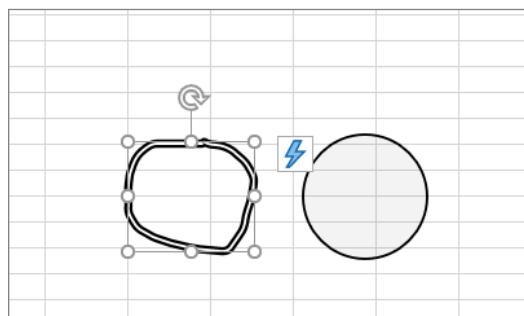


FIGURE 2.13

Convert a drawn shape into a native shape.



NOTE

To use the Draw tab in desktop Excel, you may have to display it first. If the Draw tab isn't displayed, right-click any Ribbon control and select Customize the Ribbon from the shortcut menu. The Customize Ribbon tab of the Excel Options dialog box appears. In the list box on the right, place a check mark next to Draw. Click OK to return to Excel. The Action Pen control doesn't appear on the Draw tab in desktop Excel.

The Draw tab on the Ribbon contains a special pen called an Action Pen. The Action Pen lets you handwrite text and formulas in Excel and converts them after a short pause. To enter a formula with the Action Pen, select the cell first, then click the Action Pen and start writing. Figure 2.14 shows the formula $=1+2$ handwritten in a cell using the Action Pen and Figure 2.15 shows the result after a brief pause.

FIGURE 2.14

Handwrite a formula with the Action Pen.

**FIGURE 2.15**

Excel converts a handwritten formula into an actual formula in the cell.

=1+2		
K	L	M
	3	

Performing Basic Worksheet Operations

IN THIS CHAPTER

Understanding Excel worksheet basics

Controlling your views

Manipulating the rows and columns

This chapter covers some basic information regarding workbooks, worksheets, and windows. You'll discover tips and techniques to help you take control of your worksheets and help you to work more efficiently.

Learning the Fundamentals of Excel Worksheets

In Excel, each file is called a *workbook*, and each workbook can contain one or more worksheets. You may find it helpful to think of an Excel workbook as a binder and worksheets as pages in the binder. As with a binder, you can view a particular sheet, add new sheets, remove sheets, rearrange sheets, and copy sheets.

A workbook can hold any number of sheets, and these sheets can be either worksheets (sheets consisting of rows and columns) or chart sheets (sheets that hold a single chart). A worksheet is what people usually think of when they think of a spreadsheet.

The following sections describe the operations that you can perform with windows and worksheets.

Working with Excel windows

Each Excel workbook file that you open is displayed in a window. A window is the operating system's container for that workbook. You can open as many Excel workbooks as necessary at the same time.

Each Excel window has five icons at the right side of its title bar. From left to right, they are Account, Ribbon Display Options, Minimize, Maximize (or Restore Down), and Close.

An Excel window can be in one of the following states:

Maximized Fills the entire screen. To maximize a window, click its Maximize button.

Minimized Hidden but still open. To minimize a window, click its Minimize button.

Restored Visible but smaller than the whole screen. To restore a maximized window, click its Restore Down button. To restore a minimized window, click its icon in the Windows taskbar. A window in this state can be resized and moved.

TIP

To increase the amount of information you can see in the worksheet, click the Ribbon Display Options button and choose Auto-Hide Ribbon. This maximizes the window and hides the Ribbon and status bar. In this mode, you can get temporary access to the Ribbon commands by clicking the title bar. To return to the default Ribbon view, click the Ribbon Display Options button and choose Show Tabs and Commands.

If you work with more than one workbook simultaneously (which is quite common), you need to know how to move, resize, close, and switch among the workbook windows.

Moving and resizing windows

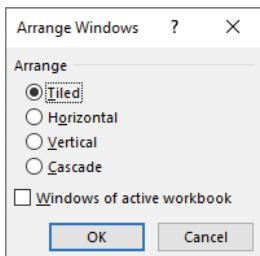
To move a window, click and drag its title bar with your mouse. If it's maximized, it will change to a restored state. If it's already in a restored state, it will maintain its current size.

To resize a window, click and drag any of its borders until it's the size that you want it to be. When you position the mouse pointer on a window's border, the mouse pointer changes to a double arrow, which lets you know that you can now click and drag to resize the window. To resize a window horizontally and vertically at the same time, click and drag any of its corners.

If you want all your workbook windows to be visible (that is, not obscured by another window), you can move and resize the windows manually, or you can let Excel do it for you. Choosing View \Rightarrow Window \Rightarrow Arrange All displays the Arrange Windows dialog box, as shown in Figure 3.1. This dialog box has four window arrangement options. Just select the one that you want and click OK. Windows that are minimized aren't affected by this command.

FIGURE 3.1

Use the Arrange Windows dialog box to arrange all open non-minimized workbook windows quickly.



Switching among windows

At any given time, one (and only one) workbook window is the active window. The active window accepts your input, and it is the window on which your commands work. The text in the active window's title bar is brighter than that of other windows. To work in a workbook in a different window, you need to make that window active. You can make a different window the active window in several ways.

- **Click another window if it's visible.** The window you click moves to the top and becomes the active window. This method isn't possible if the current window is maximized unless the other window is on a different monitor.
- **Press Ctrl+Tab to cycle through all open windows until the window that you want to work with appears on top as the active window.** Pressing Ctrl+Shift+Tab cycles through the windows in the opposite direction.
- **Press Alt+Tab or Alt+Shift+Tab to cycle through all open windows of all running programs.** Since more recent versions of Excel display each Excel window in its own Windows window, you can use this shortcut key combination to switch between them as you would switch between two programs, like between Excel and Word.
- **Choose View ▾ Window ▾ Switch Windows and select the window that you want from the drop-down list (the active window has a check mark next to it).** This menu can display as many as nine windows. If you have more than nine workbook windows open, choose More Windows (which appears below the nine window names).
- **Click the corresponding Excel icon in the Windows taskbar.**

You might be one of the many people who prefer to do most work with maximized workbook windows, which enables you to see more cells and eliminates the distraction of other workbook windows getting in the way. At times, however, viewing multiple windows is preferred. For example, displaying two windows is more efficient if you need to compare information in two workbooks or if you need to copy data from one workbook to another.

TIP

You also can display a single workbook in more than one window. For example, if you have a workbook with two worksheets, you may want to display each worksheet in a separate window to compare the two sheets. All of the window manipulation procedures described previously still apply. See "Viewing a worksheet in multiple windows" later in this chapter.

Closing windows

If you have multiple windows open, you may want to close those windows that you no longer need. Excel offers several ways to close the active window.

- Choose File ▾ Close.
- Click the Close button (the X icon) on the right side of the workbook window's title bar.
- Press Alt+F4.
- Press Ctrl+W.

When you close a workbook window, Excel checks whether you have made any changes since the last time you saved the file. If you have made changes, Excel prompts you to save the file before it closes the window. If you haven't, the window closes without a prompt from Excel.

Sometimes you will be prompted to save a workbook even if you've made no changes to it. This occurs if your workbook contains any *volatile* functions. Volatile functions recalculate every time the workbook recalculates. For example, if a cell contains =NOW(), you will be prompted to save the workbook because the NOW function updated the cell with the current date and time.

Activating a worksheet

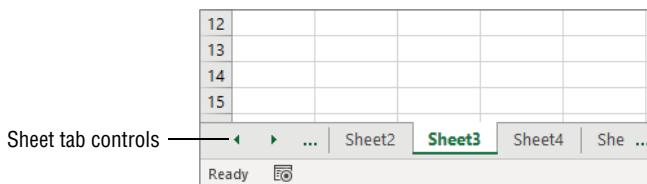
At any given time, one workbook is the active workbook and one sheet is the active sheet in the active workbook. To activate a different sheet, just click its sheet tab, which is located at the bottom of the workbook window. You also can use the following shortcut keys to activate a different sheet:

- **Ctrl+PgUp** activates the previous sheet if one exists.
- **Ctrl+PgDn** activates the next sheet if one exists.

If your workbook has many sheets, all its tabs may not be visible. Use the sheet tab controls (see Figure 3.2) to scroll the sheet tabs. Clicking the sheet tab controls scrolls one tab at a time, and Ctrl+clicking scrolls to the first or last sheet. The sheet tabs share space with the worksheet's horizontal scrollbar. You also can drag the tab split control (to the left of the horizontal scrollbar) to display more or fewer tabs. Dragging the tab split control simultaneously changes the number of visible tabs and the size of the horizontal scrollbar.

FIGURE 3.2

Use the sheet tab controls to activate a different worksheet or to see additional worksheet tabs.



TIP

When you right-click any of the tab scrolling controls, Excel displays a list of all sheets in the workbook. You can quickly activate a sheet by double-clicking it in the list.

Adding a new worksheet to your workbook

Worksheets can be an excellent organizational tool. Instead of placing everything on a single worksheet, you can use additional worksheets in a workbook to separate various workbook elements logically. For example, if you have several products whose sales you track individually, you may want to assign each product to its own worksheet and then use another worksheet to consolidate your results.

Here are four ways to add a new worksheet to a workbook:

- Click the New sheet button, which is the plus sign icon located to the right of the last visible sheet tab. A new sheet is added after the active sheet.
- Press Shift+F11. A new sheet is added before the active sheet.
- From the Ribbon, choose Home ➔ Cells ➔ Insert ➔ Insert Sheet. A new sheet is added before the active sheet.
- Right-click a sheet tab, choose Insert from the shortcut menu, and select the General tab of the Insert dialog box that appears. Then select the Worksheet icon and click OK. A new sheet is added before the active sheet.

Deleting a worksheet you no longer need

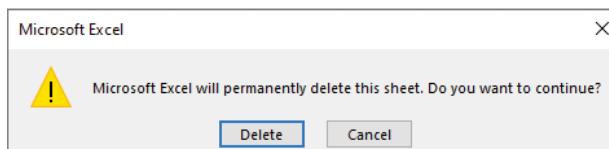
If you no longer need a worksheet or if you want to get rid of an empty worksheet in a workbook, you can delete it in either of two ways.

- Right-click its sheet tab and choose Delete from the shortcut menu.
- Activate the unwanted worksheet and choose Home ➔ Cells ➔ Delete ➔ Delete Sheet.

If the worksheet is not empty, Excel asks you to confirm that you want to delete the sheet (see Figure 3.3).

FIGURE 3.3

Excel's warning that you might be losing some data



TIP

You can delete multiple sheets with a single command by selecting the sheets that you want to delete. To select multiple sheets, press Ctrl while you click the sheet tabs that you want to delete. To select a group of contiguous sheets, click the first sheet tab, press Shift, and then click the last sheet tab (Excel displays the selected sheet names bold and underlined). Then use either method to delete the selected sheets.

CAUTION

When you delete a worksheet, it's gone for good. Deleting a worksheet is one of the few operations in Excel that can't be undone.

Changing the name of a worksheet

The default names that Excel uses for worksheets—Sheet1, Sheet2, and so on—are generic and nondescriptive. To make it easier to locate data in a multisheet workbook, you'll want to make the sheet names more descriptive.

These are three ways to change a sheet's name:

- From the Ribbon, choose Home \Rightarrow Cells \Rightarrow Format \Rightarrow Rename Sheet.
- Double-click the sheet tab.
- Right-click the sheet tab and choose Rename.

Excel highlights the name on the sheet tab so that you can edit the name or replace it with a new name. While you're editing a sheet name, all the normal text selection techniques work, such as Home, End, arrow keys, and Shift+arrow keys. Press Enter when you're finished editing and the focus will be back on the active cell.

Sheet names can contain as many as 31 characters, and spaces are allowed. However, you can't use the following characters in sheet names:

:	Colon
/	Slash
\	Backslash
[]	Square brackets
?	Question mark
*	Asterisk

Keep in mind that a longer worksheet name results in a wider tab, which takes up more space onscreen. Therefore, if you use lengthy sheet names, you won't be able to see as many sheet tabs without scrolling the tab list.

Changing a sheet tab color

Excel allows you to change the background color of your worksheet tabs. For example, you may prefer to color-code the sheet tabs to make identifying the worksheet's contents easier.

To change the color of a sheet tab, choose Home \Rightarrow Cells \Rightarrow Format \Rightarrow Tab Color, or right-click the tab and choose Tab Color from the shortcut menu. Then select the color from the color palette. You can't change the text color, but Excel will choose a contrasting color to make the text visible. For example, if you make a sheet tab black, Excel will display white text.

If you change a sheet tab's color, the tab shows a gradient from that color to white when the sheet is active. When a different sheet is active, the whole tab appears in the selected color.

Rearranging your worksheets

You may want to rearrange the order of worksheets in a workbook. If you have a separate worksheet for each sales region, for example, arranging the worksheets in alphabetical order might be helpful. You can also move a worksheet from one workbook to another and create copies of worksheets, either in the same workbook or in a different workbook.

You can move a worksheet in the following ways:

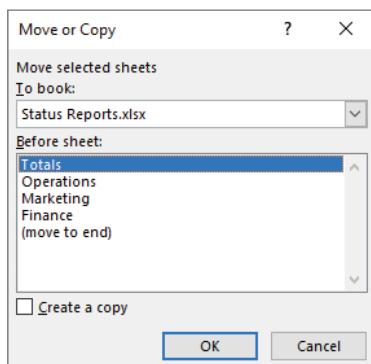
- Right-click the sheet tab and choose Move or Copy to display the Move or Copy dialog box (see Figure 3.4). Use this dialog box to specify the location for the sheet.

- From the Ribbon, choose Home \Rightarrow Cells \Rightarrow Format \Rightarrow Move or Copy Sheet. This shows the same dialog box as the previous method.
- Click the worksheet tab and drag it to its desired location. When you drag, the mouse pointer changes to a small sheet icon, and a small arrow indicates where the sheet will be placed when you release the mouse button. To move a worksheet to a different workbook by dragging, both workbooks must be visible.

Copying the worksheet is similar to moving it. If you use one of the options that shows the Move or Copy dialog box, select the Create a Copy check box. To drag and create a copy, hold down the Ctrl key while you drag the worksheet tab. The mouse pointer will change to a small sheet icon with a plus sign on it.

FIGURE 3.4

Use the Move or Copy dialog box to move or copy worksheets in the same or another workbook.



TIP

You can move or copy multiple sheets simultaneously. First, select the sheets by clicking their sheet tabs while holding down the Ctrl key. Then you can move or copy the set of sheets by using the preceding methods.

If you move or copy a worksheet to a workbook that already has a sheet with the same name, Excel changes the name to make it unique. For example, Sheet1 becomes Sheet1 (2). You probably want to rename the copied sheet to give it a more meaningful name. (See “Changing the name of a worksheet” earlier in this chapter.)

NOTE

When you move or copy a worksheet to a different workbook, any defined names and custom formats are also copied to the new workbook.

Hiding and unhiding a worksheet

In some situations, you may want to hide one or more worksheets. Hiding a sheet may be useful if you don't want others to see it or if you just want to get it out of the way. When a sheet is hidden, its sheet tab is also hidden. You can't hide all the sheets in a workbook; at least one sheet must remain visible.

To hide a worksheet, choose Home \Rightarrow Cells \Rightarrow Format \Rightarrow Hide & Unhide \Rightarrow Hide Sheet, or right-click its sheet tab and choose Hide. The active worksheet (or selected worksheets) will be hidden from view.

To unhide a hidden worksheet, choose Home \Rightarrow Cells \Rightarrow Format \Rightarrow Hide & Unhide \Rightarrow Unhide Sheet, or right-click any sheet tab and choose Unhide. Excel opens the Unhide dialog box, which lists all hidden sheets. Choose the sheet that you want to redisplay and click OK. You can select multiple sheets from this dialog box by holding down the Ctrl key and clicking the sheets to unhide. Hold down the Shift key and click a sheet to select a range of contiguous sheets. When you unhide a sheet, it appears in its previous position among the sheet tabs.

Preventing Sheet Actions

To prevent others from unhideing hidden sheets, inserting new sheets, renaming sheets, copying sheets, or deleting sheets, protect the workbook's structure.

1. Choose Review \Rightarrow Protect \Rightarrow Protect Workbook.
2. In the Protect Structure and Windows dialog box, select the Structure option.
3. Provide a password (optional) and click OK.

After you perform these steps, several commands will no longer be available from the Ribbon or when you right-click a sheet tab: Insert, Delete, Rename, Move or Copy, Tab Color, Hide, and Unhide. Be aware, however, that this is a weak security measure. Cracking this protection feature is relatively easy, especially if you don't add a password.

Controlling the Worksheet View

As you add more information to a worksheet, you may find that navigating and locating what you want becomes more difficult. Excel includes a few options that enable you to view your sheet, and sometimes multiple sheets, more efficiently. The following sections discuss a few additional worksheet options at your disposal.

Zooming in or out for a better view

Normally, everything you see onscreen is displayed at 100%. You can change the zoom percentage from 10% (very tiny) to 400% (huge). Zooming out by using a small zoom percentage can help you get a bird's-eye view of your worksheet to see how it's laid out. Zooming in by using a large zoom percentage is useful if you have trouble deciphering tiny type. Zooming doesn't change the font size specified for the cells, so it has no effect on printed output.



Excel contains separate options for changing the size of your printed output. (Use the controls in the Page Layout \Rightarrow Scale to Fit Ribbon group.) See Chapter 7, “Printing Your Work,” for details.

You can change the zoom factor of the active worksheet window by using any of these three methods:

- Use the Zoom slider located on the right side of the status bar. Click and drag the slider, and your screen transforms as you slide. You can also click anywhere on the Zoom slider to move the bar directly to that spot or use the Zoom In (+) or Zoom Out (-) buttons on either side of the slider to change the zoom by 10%.
- Press Ctrl and use the wheel button on your mouse to zoom in or out.
- Choose View \Rightarrow Zoom \Rightarrow Zoom, which displays a dialog box with some zoom options.

Also in the Zoom Ribbon group is a 100% button to return to 100% zoom quickly and a Zoom to Selection button to change the zoom so that whatever cells you have selected take up the whole window (but still limited to the 10%–400% zoom range).

TIP

Zooming affects only the active worksheet window, so you can use different zoom factors for different worksheets. Also, if you have a worksheet displayed in two different windows, you can set a different zoom factor for each of the windows.



If your worksheet uses named ranges (see Chapter 4, “Working with Excel Ranges and Tables”), zooming your worksheet to 39% or less displays the name of the range overlaid on the cells. Viewing named ranges in this manner is useful for getting an overview of how a worksheet is laid out.

3

Viewing a worksheet in multiple windows

Sometimes, you may want to view two different parts of a worksheet simultaneously—perhaps to make referencing a distant cell in a formula easier. Or you may want to examine more than one sheet in the same workbook simultaneously. You can accomplish either of these actions by opening a new view to the workbook, using one or more additional windows.

To create and display a new view of the active workbook, choose View \Rightarrow Window \Rightarrow New Window.

Excel displays a new window for the active workbook, similar to the one shown in Figure 3.5. In this case, each window shows a different worksheet in the workbook. Notice the text in the active window’s title bar is *FinancialInformation.xlsx - 2*. To help you keep track of the windows, Excel appends a hyphen and a number to each window.

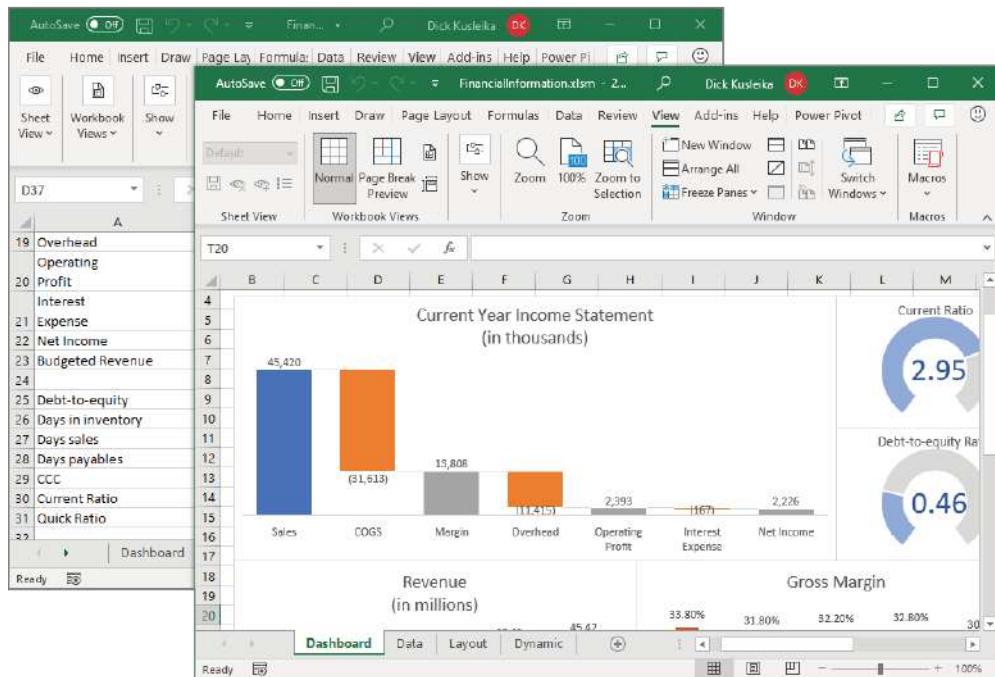
TIP

If the workbook is maximized when you create a new window, you may not even notice that Excel created the new window. If you look at the Excel title bar, though, you’ll see that the workbook title now has – 2 appended to the name. Choose View \Rightarrow Window \Rightarrow Arrange All and then choose one of the Arrange options in the Arrange Windows dialog box to display the open windows. If you select the Windows of Active Workbook check box, only the windows of the active workbook are arranged.

Part I: Getting Started with Excel

FIGURE 3.5

Use multiple windows to view different sections of a workbook at the same time.



A single workbook can have as many views (that is, separate windows) as you want. Each window is independent. In other words, scrolling to a new location in one window doesn't cause scrolling in the other window(s). However, if you make changes to the worksheet shown in a particular window, those changes are also made in all views of that worksheet.

You can close these additional windows when you no longer need them. For example, clicking the Close button on the active window's title bar closes the active window but doesn't close the other windows for the workbook. If you have unsaved changes, Excel will prompt you to save only when you close the last window.

TIP

Multiple windows make copying or moving information from one worksheet to another easier. You can use Excel's drag-and-drop procedures to copy or move ranges.

Comparing sheets side by side

In some situations, you may want to compare two worksheets that are in different windows. The View Side by Side feature makes this task a bit easier.

First, make sure that the two sheets are displayed in separate windows. (The sheets can be in the same workbook or in different workbooks.) If you want to compare two sheets in the same workbook, choose **View** \Rightarrow **Window** \Rightarrow **New Window** to create a new window for the active workbook. Activate the first window; then choose **View** \Rightarrow **Window** \Rightarrow **View Side by Side**. If more than two windows are open, you see a dialog box that lets you select the window for the comparison. The two windows are tiled to fill the entire screen.

When using the Compare Side by Side feature, scrolling in one of the windows also scrolls the other window. If you don't want this simultaneous scrolling, choose **View** \Rightarrow **Window** \Rightarrow **Synchronous Scrolling** (which is a toggle). If you have rearranged or moved the windows, choose **View** \Rightarrow **Window** \Rightarrow **Reset Window Position** to restore the windows to the initial side-by-side arrangement. To turn off the side-by-side viewing, choose **View** \Rightarrow **Window** \Rightarrow **View Side by Side** again.

Keep in mind that this feature is for manual comparison only. Unfortunately, Excel doesn't provide a way to identify the differences between two sheets automatically.

Splitting the worksheet window into panes

If you prefer not to clutter your screen with additional windows, Excel provides another option for viewing multiple parts of the same worksheet. Choosing **View** \Rightarrow **Window** \Rightarrow **Split** splits the active worksheet into two or four separate panes. The split occurs at the location of the active cell. If the active cell is in row 1 or column A, this command results in a two-pane split; otherwise, it gives you four panes. You can use the mouse to drag the split separating individual panes to resize them.

Figure 3.6 shows a worksheet split into four panes. Notice that row numbers aren't continuous. The top panes show rows 7 through 9, and the bottom panes show rows 40 through 57. In other words, splitting panes enables you to display widely separated areas of a worksheet in a single window. To remove the split panes, choose **View** \Rightarrow **Window** \Rightarrow **Split** again or double-click the split bar(s).

Keeping the titles in view by freezing panes

If you set up a worksheet with column headings or descriptive text in the first column, this identifying information won't be visible when you scroll down or to the right. Excel provides a handy solution to this problem: freezing panes. Freezing panes keeps the column or row headings visible while you're scrolling through the worksheet.

To freeze panes, start by moving the active cell to the cell below the row(s) that you want to remain visible while you scroll vertically and to the right of the column(s) that you want to remain visible while you scroll horizontally. Then choose **View** \Rightarrow **Window** \Rightarrow **Freeze Panes** and select the **Freeze Panes** option from the drop-down list. Excel inserts dark lines to indicate the frozen rows and columns. The frozen rows and columns remain visible while you scroll throughout the worksheet. To remove the frozen panes, choose **View** \Rightarrow **Window** \Rightarrow **Freeze Panes**, and select the **Unfreeze Panes** option from the drop-down list.

Figure 3.7 shows a worksheet with frozen panes. In this case, row 1 and column A are frozen in place. (Cell B2 was the active cell when the **View** \Rightarrow **Window** \Rightarrow **Freeze Panes** \Rightarrow **Freeze Panes** command was used.) This technique allows you to scroll down and to the right to locate some information while keeping the column titles and the column A entries visible.

Part I: Getting Started with Excel

FIGURE 3.6

You can split the worksheet window into two or four panes to view different areas of the worksheet at the same time.

	A	B	C	D	E	F	G
7	Other Long-term Assets	50,243	40,668	59,527	48,066	60,330	53,006
8	Accounts Payable	(2,836,512)	(2,929,327)	(3,018,834)	(2,914,237)	(2,891,749)	(2,985,140)
9	Accrued Expenses	(200,148)	(230,845)	(220,480)	(182,498)	(182,284)	(183,294)
40							
41							
42	Debt-to-equity	0.982109893	0.75924756	0.637722947	0.527408438	0.423740945	
43	Days in inventory	20.74269355	22.17958774	21.14420167	22.17274726	21.21260649	
44	Days sales	28.90	29.83	29.58	30.61	31.17	
45	Days payables	38.68	41.25	38.75	40.32	39.69	
46	CCC	10.97	10.76	11.97	12.46	12.69	
47	Current Ratio	1.90	1.86	2.21	2.44	2.61	
48	Quick Ratio	1.21	1.16	1.50	1.73	1.92	
49	Return on Equity	0.402644179	0.328682049	0.405086336	0.273459176	0.283311968	
50	Return on Assets	0.14	0.13	0.18	0.13	0.15	
51							
52	Gross Margin	30.70%	32.30%	30.20%	33.70%	32.50%	
53	Operating Margin	3.89%	3.80%	5.43%	4.25%	4.93%	
54	Net Margin	3.50%	3.40%	5.00%	3.90%	4.60%	
55	Budgeted Revenue		(41,224,000)	(40,812,000)	(42,036,000)	(41,616,000)	
56							
57							

FIGURE 3.7

Freeze certain columns and rows to make them remain visible while you scroll the worksheet.

	A	D	E	F	G	H
1		2017	2018	2019	2020	
8	Accounts Payable	(2,707,529)	(2,828,555)	(2,841,051)	(2,965,245)	
9	Accrued Expenses	(202,490)	(237,866)	(252,016)	(247,640)	
10	Debt	(3,916,585)	(3,600,261)	(3,558,208)	(3,394,742)	
11	Common Stock	(100,000)	(100,000)	(100,000)	(100,000)	
12	Additional Paid-in Capital	(1,547,889)	(1,547,889)	(1,547,889)	(1,547,889)	
13	Retained Earnings	(2,610,338)	(3,430,074)	(4,230,292)	(4,974,105)	
14		-	-	-	-	
15						
16	Sales	(39,653,640)	(42,429,395)	(42,429,395)	(43,702,277)	
17	COGS	27,123,090	29,530,859	29,530,859	29,236,823	
18	Margin	(12,530,550)	(12,898,536)	(12,898,536)	(14,465,454)	
19	Overhead	10,745,050	10,933,248	10,984,935	12,681,788	
20	Operating Profit	(1,785,500)	(1,965,288)	(1,913,601)	(1,783,666)	
21	Interest Expense	199,354	183,253	173,996	166,682	
22	Net Income	(1,586,146)	(1,782,035)	(1,739,605)	(1,616,984)	

Most of the time you'll want to freeze either the first row or the first column. The View \Rightarrow Window \Rightarrow Freeze Panes drop-down list has two additional options: Freeze Top Row and Freeze First Column. Using these commands eliminates the need to position the active cell before freezing panes.

TIP

If you designated a range to be a table (by choosing Insert \Rightarrow Tables \Rightarrow Table), you may not even need to freeze panes. When you scroll down, Excel displays the table column headings in place of the column letters. Figure 3.8 shows an example. The table headings replace the column letters only when a cell within the table is selected.

FIGURE 3.8

When using a table, scrolling down displays the table headings where the column letters normally appear.

Agent	Date Listed	Area	List Price	Bedrooms	Baths	SqFt	Type	Pool	Sold
37 Hamilton	12/12/2022	South	334,000	3	2	2,400	Split Level	TRUE	TRUE
38 Jenkins	2/19/2022	South	433,000	4	3	3,100	Split Level	TRUE	FALSE
39 Lang	1/3/2022	Downtown	215,000	3	2.5	1,500	2 Story	FALSE	TRUE
40 Peterson	7/27/2022	North	344,000	3	2.5	2,500	Ranch	TRUE	FALSE
41 Randolph	3/15/2022	North	286,000	4	3	2,000	2 Story	TRUE	TRUE
42 Robinson	2/20/2022	Downtown	372,000	2	2.5	2,700	Split Level	FALSE	FALSE
43 Romero	4/28/2022	South	252,000	4	2	1,800	Split Level	TRUE	FALSE
44 Shasta	9/10/2022	North	232,000	3	3	1,700	Condo	TRUE	FALSE
45 Chung	8/15/2022	North	341,000	3	3	2,400	2 Story	TRUE	FALSE
46 Daily	1/20/2022	Downtown	399,000	4	3	2,900	Ranch	FALSE	TRUE
47 Kelly	12/4/2022	Downtown	251,000	3	2.5	1,800	Split Level	TRUE	TRUE
48 Hamilton	3/9/2022	South	264,000	3	3	1,900	Ranch	TRUE	FALSE
49 Jenkins	10/11/2022	South	244,000	2	3	1,700	Ranch	FALSE	FALSE
50 Lang	2/2/2022	Downtown	413,000	4	3	3,000	2 Story	FALSE	FALSE
51 Peterson	3/2/2022	North	337,000	4	3	2,400	Split Level	FALSE	FALSE
52 Randolph	11/23/2022	Central	323,000	2	3	2,300	Ranch	FALSE	TRUE
53 Robinson	11/19/2022	North	366,000	4	3	2,600	Split Level	TRUE	TRUE

Monitoring cells with a Watch Window

In some situations, you may want to monitor the value in a particular cell as you work. As you scroll throughout the worksheet, that cell may disappear from view. A feature known as a Watch Window can help. A Watch Window displays the value of any number of cells in a handy window that's always visible.

To display the Watch Window, choose Formulas \Rightarrow Formula Auditing \Rightarrow Watch Window. The Watch Window is actually a task pane, and you can dock it to the side of the window or drag it and make it float over the worksheet.

To add a cell to watch, click Add Watch and specify the cell that you want to watch. The Watch Window displays the value in that cell. You can add any number of cells to the Watch Window. Figure 3.9 shows the Watch Window monitoring cells in different worksheets.

FIGURE 3.9

Use the Watch Window to monitor the value in one or more cells.

Book	Sheet	Name	Cell	Value	Formula
FinancialInformation.xlsxm	Data	C33	Cell C33	33.80%	=C18/C16
FinancialInformation.xlsxm	Dynamic	D13	Cell D13	(2,190,715)	=C13-SUM(D37:...
FinancialInformation.xlsxm	Data	C39	Cell C39	1.00	=C40/0.8-C40

TIP

Double-click a cell in the Watch Window to select that cell immediately. This works only if the watched cell is in the active workbook.

Working with Rows and Columns

The following sections discuss worksheet operations that involve complete rows and columns (rather than individual cells). Every worksheet has exactly 1,048,576 rows and 16,384 columns, and these values can't be changed.

NOTE

If you open a workbook that was created in a version of Excel prior to Excel 2007, the workbook is opened in Compatibility Mode. These workbooks have 65,536 rows and 256 columns. If you would like to increase the number of rows and columns, save the workbook as an Excel .xlsx file and then reopen it.

Selecting rows and columns

Select an entire row by clicking on the row header or using the Shift+spacebar keyboard shortcut. Select an entire column by clicking on the column header or using the Ctrl+spacebar keyboard shortcut.

If you need to select multiple, contiguous rows or columns, click on the header and drag to the adjacent row or column to extend the selection. You can also click on a row or column header, hold down the Shift key, and click on another row or column to extend the selection. To select rows or columns that aren't contiguous, hold down the Ctrl key while you click each header.

Inserting rows and columns

Although the number of rows and columns in a worksheet is fixed, you can still insert and delete rows and columns if you need to make room for additional information.

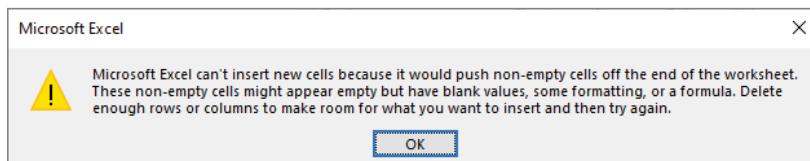
These operations don't change the number of rows or columns. Instead, inserting a new row moves down the other rows to accommodate the new row. The last row is simply removed from the worksheet if it's empty. Inserting a new column shifts the columns to the right, and the last column is removed if it's empty.

NOTE

If the last row isn't empty, you can't insert a new row. Similarly, if the last column contains information, Excel doesn't let you insert a new column. In either case, attempting to insert a row or column displays the dialog box shown in Figure 3.10.

FIGURE 3.10

You can't insert a new row or column if it causes nonblank cells to move off the worksheet.



To insert a new row or rows, use either of these methods:

- Select an entire row or multiple rows by clicking the row numbers in the worksheet border. Right-click and choose Insert from the shortcut menu.
- Select a cell in the row that you want to insert and then choose Home \Rightarrow Cells \Rightarrow Insert \Rightarrow Insert Sheet Rows. If you select multiple cells in the column, Excel inserts additional rows that correspond to the number of cells selected in the column and moves the rows below the insertion down.

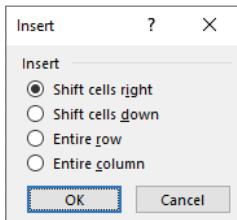
To insert a new column or columns, use either of these methods:

- Select an entire column or columns by clicking the column letters in the worksheet border. Right-click and choose Insert from the shortcut menu.
- Select a cell in the column that you want to insert and then choose Home \Rightarrow Cells \Rightarrow Insert \Rightarrow Insert Sheet Columns. If you select multiple cells in the row, Excel inserts additional columns that correspond to the number of cells selected in the row.

You can also insert cells rather than just entire rows or columns. Select the range into which you want to add new cells and then choose Home \Rightarrow Cells \Rightarrow Insert \Rightarrow Insert Cells (or right-click the selection and choose Insert). To insert cells, you must shift the existing cells to the right or down. Therefore, Excel displays the Insert dialog box shown in Figure 3.11 so that you can specify the direction in which you want to shift the cells. Notice that this dialog box also enables you to insert entire rows or columns.

FIGURE 3.11

You can insert partial rows or columns by using the Insert dialog box.



Deleting rows and columns

You may also want to delete rows or columns in a worksheet. For example, your sheet may contain old data that is no longer needed, or you may want to remove empty rows or columns.

To delete a row or rows, use either of these methods:

- Select an entire row or multiple rows by clicking the row numbers in the worksheet border. Right-click and choose Delete from the shortcut menu.
- Move the active cell to the row that you want to delete and then choose Home \Rightarrow Cells \Rightarrow Delete \Rightarrow Delete Sheet Rows. If you select multiple cells in the column, Excel deletes all rows in the selection.

To delete a column or columns, use either of these methods:

- Select an entire column or multiple columns by clicking the column letters in the worksheet border. Right-click and choose Delete from the shortcut menu.
- Move the active cell to the column that you want to delete and then choose Home \Rightarrow Cells \Rightarrow Delete \Rightarrow Delete Sheet Columns. If you select multiple cells in the row, Excel deletes all columns in the selection.

If you discover that you accidentally deleted a row or column, select Undo from the Quick Access Toolbar (or press Ctrl+Z) to undo the action.

TIP

You can use the shortcut keys Ctrl + (plus sign) and Ctrl - (minus sign) to insert and delete rows, columns, or cells. If you have an entire row or entire column selected, those shortcuts will insert or delete the entire row or column. If the selection isn't an entire row or column, the Insert dialog is displayed.

Changing column widths and row heights

Often, you'll want to change the width of a column or the height of a row. For example, you can make columns narrower to show more information on a printed page. Or you may want to increase row height to create a "double-spaced" effect.

Excel provides several ways to change the widths of columns and the height of rows.

Changing column widths

Column width is measured in terms of the number of characters of a monospaced font that will fit into the cell's width. By default, each column's width is 8.43 units, which equates to 64 pixels (px).

TIP

If hash symbols (#) fill a cell that contains a numerical value, the column isn't wide enough to accommodate the information in the cell. Widen the column to solve the problem.

Before you change the column width, you can select multiple columns so that the width will be the same for all selected columns. To select multiple columns, either click and drag in the column border or press Ctrl while you select individual columns. To select all columns, click the button where the row and column headers intersect. You can change column widths by using any of the following techniques:

- Drag the right-column border with the mouse until the column is the desired width.
- Choose Home \Rightarrow Cells \Rightarrow Format \Rightarrow Column Width and enter a value in the Column Width dialog box.
- Choose Home \Rightarrow Cells \Rightarrow Format \Rightarrow AutoFit Column Width to adjust the width of the selected column so that the widest entry in the column fits. Instead of selecting an entire column, you can just select cells in the column, and the column is adjusted based on the widest entry in your selection.
- Double-click the right border of a column header to set the column width automatically to the widest entry in the column.

TIP

To change the default width of all columns, choose Home \Rightarrow Cells \Rightarrow Format \Rightarrow Default Width. This command displays a dialog box into which you enter the new default column width. All columns that haven't been previously adjusted take on the new column width.

CAUTION

After you manually adjust a column's width, Excel will no longer automatically adjust the column to accommodate longer numerical entries. If you enter a long number that displays as hash symbols (#), you need to change the column width manually.

Changing row heights

Row height is measured in points (a standard unit of measurement in the printing trade—72 pts is equal to 1 inch). The default row height using the default font is 15 pts, or 20 pixels (px).

The default row height can vary, depending on the font defined in the Normal style. In addition, Excel automatically adjusts row heights to accommodate the tallest font in the row. So, if you change the font size of a cell to 20 pts, for example, Excel makes the row taller so that the entire text is visible.

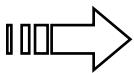
You can set the row height manually, however, by using any of the following techniques. As with columns, you can select multiple rows:

- Drag the lower row border with the mouse until the row is the desired height.
- Choose Home \Rightarrow Cells \Rightarrow Format \Rightarrow Row Height and enter a value (in points) in the Row Height dialog box.
- Double-click the bottom border of a row to set the row height automatically to the tallest entry in the row. You can also choose Home \Rightarrow Cells \Rightarrow Format \Rightarrow AutoFit Row Height for this task.

Changing the row height is useful for spacing out rows and is almost always preferable to inserting empty rows between lines of data.

Hiding rows and columns

In some cases, you may want to hide rows or columns. Hiding rows and columns may be useful if you don't want users to see particular information or if you need to print a report that summarizes the information in the worksheet without showing all the details.



Chapter 24, "Creating and Using Worksheet Outlines," discusses another way to summarize worksheet data without showing all the details—worksheet outlining.

To hide rows in your worksheet, select the row or rows that you want to hide by clicking in the row header on the left. Then right-click and choose Hide from the shortcut menu. Or you can use the commands on the Home \Rightarrow Cells \Rightarrow Format \Rightarrow Hide & Unhide menu.

To hide columns, select the column or columns that you want to hide. Then right-click and choose Hide from the shortcut menu. Or you can use the commands on the Home \Rightarrow Cells \Rightarrow Format \Rightarrow Hide & Unhide menu.

TIP

You can also drag the row or column's border to hide the row or column. You must drag the border in the row or column heading. Drag the bottom border of a row upward or the right border of a column to the left.

A hidden row is actually a row with its height set to zero. Similarly, a hidden column has a column width of zero. When you use the navigation keys to move the active cell, cells in hidden rows or columns are skipped. In other words, you can't use the navigation keys to move to a cell in a hidden row or column.

Notice, however, that Excel displays a narrow column heading for hidden columns and a narrow row heading for hidden rows. You can click and drag the column heading to make the column wider—and make it visible again. For a hidden row, click and drag the small row heading to make the row visible.

Another way to unhide a row or column is to choose Home \Rightarrow Editing \Rightarrow Find & Select \Rightarrow Go To (or use one of its two shortcut keys: F5 or Ctrl+G) to select a cell in a hidden row or column. For example, if column A is hidden, you can press F5 and go to cell A1 (or any other cell in column A) to move the active cell to the hidden column. Then you can choose Home \Rightarrow Cells \Rightarrow Format \Rightarrow Hide & Unhide \Rightarrow Unhide Columns.

Working with Excel Ranges and Tables

IN THIS CHAPTER

Understanding Excel cells and ranges

Selecting cells and ranges

Copying or moving ranges

Using names to work with ranges

Adding comments to cells

Adding notes to cells

Working with tables

Most of the work you do in Excel involves cells and ranges. Understanding how best to manipulate cells and ranges will save you time and effort. This chapter discusses a variety of techniques that are essential for Excel users.

Understanding Cells and Ranges

A *cell* is a single element in a worksheet that can hold a value, some text, or a formula. A cell is identified by its address, which consists of its column letter and row number. For example, cell D9 is the cell in the fourth column and the ninth row.

A group of one or more cells is called a *range*. You designate a range address by specifying its upper-left cell address and its lower-right cell address, separated by a colon.

Here are some examples of range addresses:

C24	A range that consists of a single cell.
A1:B1	Two cells that occupy one row and two columns.
A1:A100	100 cells in column A.
A1:D4	16 cells (four rows by four columns).
C1:C1048576	An entire column of cells; this range also can be expressed as C:C.
A6:XFD6	An entire row of cells; this range also can be expressed as 6:6.
A1:XFD1048576	All cells in a worksheet. This range also can be expressed as either A:XFD or 1:1048576.

Selecting ranges

To perform an operation on a range of cells in a worksheet, you must first select the range. For example, if you want to make the text bold for a range of cells, you must select the range and then choose Home \Rightarrow Font \Rightarrow Bold (or press Ctrl+B).

When you select a range, the cells appear highlighted. The exception is the active cell, which remains its normal color. Figure 4.1 shows an example of a selected range (A4 : D8) in a worksheet. Cell A4, the active cell, is in the selected range but not highlighted.

FIGURE 4.1

When you select a range, it appears highlighted, but the active cell within the range is not highlighted.

	A	B	C	D	E	F	G
1	Budget Summary						
2							
3		Q1	Q2	Q3	Q4	Year Total	
4	Salaries	286,500	286,500	286,500	290,500	1,150,000	
5	Travel	40,500	42,525	44,651	46,884	174,560	
6	Supplies	59,500	62,475	65,599	68,879	256,452	
7	Facility	144,000	144,000	144,000	144,000	576,000	
8	Total	530,500	535,500	540,750	550,263	2,157,013	
9							
10							
11							

You can select a range in several ways:

- Left-click and drag over the range. If you drag to the end of the window, the worksheet will scroll.
- Press the Shift key while you use the navigation keys or mouse to select a range.
- Press F8 to enter Extend Selection mode (Extend Selection appears in the status bar). In this mode, click the lower-right cell of the range or use the navigation keys to extend the range. Press F8 again to exit Extend Selection mode.
- Type the cell or range address into the Name box (located to the left of the Formula bar) and press Enter. Excel selects the cell or range that you specified.
- Choose Home \Rightarrow Editing \Rightarrow Find & Select \Rightarrow Go To (or press F5 or Ctrl+G) and enter a range's address manually in the Go To dialog box. When you click OK, Excel selects the cells in the range that you specified.

TIP

While you're selecting a range that contains more than one cell, Excel displays the number of rows and columns in your selection in the Name box (which is to the left of the Formula bar). When you finish making the selection, the Name box reverts to showing the address of the active cell.

Selecting complete rows and columns

Often, you'll need to select an entire row or column. For example, you may want to apply the same numeric format or the same alignment options to an entire row or column. You can select entire rows and columns in several ways, including the following:

- Click the row or column header to select a single row or column or click and drag for multiple rows or columns.
- To select multiple (nonadjacent) rows or columns, click the first row or column header and then hold down the Ctrl key while you click the additional row or column header that you want.
- Press Ctrl+spacebar to select the column(s) of the currently selected cells. Press Shift+spacebar to select the row(s) of the currently selected cells.

TIP

Press Ctrl+A to select all cells in the worksheet, which is the same as selecting all rows and all columns. If the active cell is within a contiguous range, Ctrl+A will just select that range. In that case, press Ctrl+A again to select all the cells in the worksheet. You can also click the area at the intersection of the row and column headers to select all cells.

Selecting noncontiguous ranges

Most of the time, the ranges that you select are contiguous—a single rectangle of cells. Excel also enables you to work with noncontiguous ranges, which consist of two or more ranges (or single cells) that aren't necessarily adjacent to each other. Selecting noncontiguous ranges is also known as a *multiple selection*. If you want to apply the same formatting to cells in different areas of your worksheet, one approach is to make a multiple selection. When the appropriate cells or ranges are selected, the formatting that you select is applied to all of them. Figure 4.2 shows a noncontiguous range selected in a worksheet. Two ranges are selected: B4 : C8 and F15 : F17.

You can select a noncontiguous range in the same ways that you select a contiguous range with a few minor differences. Instead of simply clicking and dragging for contiguous ranges, after you select the first range, hold down the Ctrl key while you click and drag to select additional ranges. If you're selecting a range using the arrow keys, select the first range, then press Shift+F8 to enter Add or Remove Selection mode (that term will appear in the status bar). Press Shift+F8 again to exit Add or Remove Selection mode. Anywhere you type the range manually, such as in the Name box or the Go To dialog box, simply separate the noncontiguous ranges with a comma. For example, typing A1:A10 , C5:C6 will select those two noncontiguous ranges.

NOTE

Noncontiguous ranges differ from contiguous ranges in several important ways. One major difference is that you can't use drag-and-drop methods (described later) to move or copy noncontiguous ranges.

FIGURE 4.2

Excel enables you to select noncontiguous ranges.

	A	B	C	D	E	F	G
1	Budget Summary						
2							
3		Q1	Q2	Q3	Q4	Year Total	
4	Salaries	286,500	286,500	286,500	290,500	1,150,000	
5	Travel	40,500	42,525	44,651	46,884	174,560	
6	Supplies	59,500	62,475	65,599	68,879	256,452	
7	Facility	144,000	144,000	144,000	144,000	576,000	
8	Total	530,500	535,500	540,750	550,263	2,157,013	
9							
10							
11							
12							
13							
14							
15				Salaries	-		
16				Travel	-		
17				Supplies	-		
18				Facility	-		
19				Total	-		
20							

Selecting multi-sheet ranges

In addition to two-dimensional ranges on a single worksheet, ranges can extend across multiple worksheets to be three-dimensional ranges.

Suppose you have a workbook set up to track budgets. One approach is to use a separate worksheet for each department, making it easy to organize the data. You can click a sheet tab to view the information for a particular department.

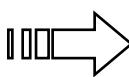
Figure 4.3 shows a simplified example. The workbook has four sheets: Totals, Operations, Marketing, and Manufacturing. The sheets are laid out identically. The only difference is the values. The Totals sheet contains formulas that compute the sum of the corresponding items in the three departmental worksheets.

FIGURE 4.3

The worksheets in this workbook are laid out identically.

The screenshot shows a Microsoft Excel window with three tabs at the bottom: **Totals**, **Operations**, and **Marketing**. The **Totals** tab is selected. The **Budget Summary** worksheet is displayed, containing the following data:

	A	B	C	D	E	F	G	H
1	Budget Summary							
2								
3		Q1	Q2	Q3	Q4	Year Total		
4	Salaries	286,500	286,500	286,500	290,500	1,150,000		
5	Travel	40,500	42,525	44,651	46,884	174,560		
6	Supplies	59,500	62,475	65,599	68,879	256,452		
7	Facility	144,000	144,000	144,000	144,000	576,000		
8	Total	530,500	535,500	540,750	550,263	2,157,013		
9								
10								
11								
12								
13								
14								
15				Salaries	-			
16				Travel	-			
17				Supplies	-			
18				Facility	-			
19				Total	-			
20								



This workbook, named `budget.xlsx`, is available on this book's website at www.wiley.com/go/excel365bible.

4

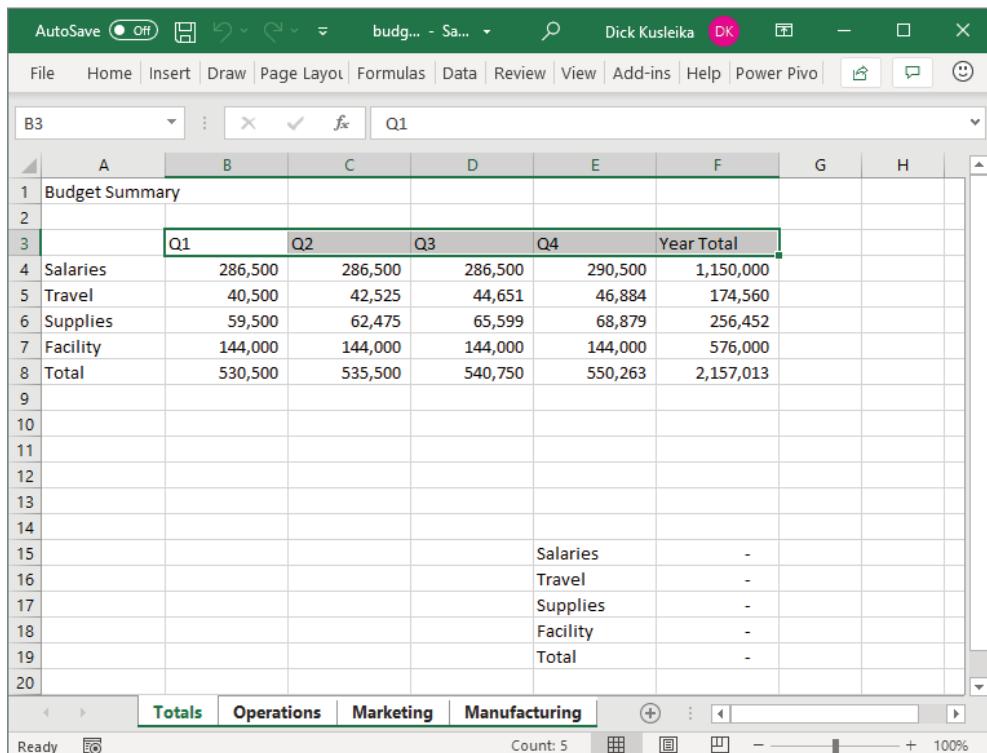
Assume that you want to apply formatting to the sheets—for example, to make the column headings bold with background shading. One (albeit not-so-efficient) approach is to format the cells in each worksheet separately. A better technique is to select a multi-sheet range and format the cells in all the sheets simultaneously. The following is a step-by-step example of multi-sheet formatting using the workbook shown in Figure 4.3:

1. **Activate the Totals worksheet by clicking its tab.**
2. **Select the range B3 : F3.**
3. **Press Shift and click the Manufacturing sheet tab.** This step selects all worksheets between the active worksheet (Totals) and the sheet tab that you click—in essence, a three-dimensional range of cells (see Figure 4.4). When multiple sheets are selected, the workbook window's title bar displays **Group** to remind you that you've selected a group of sheets and that you're in Group mode.

Part I: Getting Started with Excel

FIGURE 4.4

In Group mode, you can work with a three-dimensional range of cells that extend across multiple worksheets.



4. Choose Home \Rightarrow Font \Rightarrow Bold and then choose Home \Rightarrow Font \Rightarrow Fill Color to apply a colored background. Excel applies the formatting to the selected range across the selected sheets.
5. Click one of the other sheet tabs. This step selects the sheet and cancels Group mode; Group is no longer displayed in the title bar.

When a workbook is in Group mode, any changes that you make to cells in one worksheet also apply to the corresponding cells in all the other grouped worksheets. You can use this to your advantage when you want to set up a group of identical worksheets because any labels, data, formatting, or formulas you enter are automatically added to the same cells in all the grouped worksheets.

NOTE

When Excel is in Group mode, some commands are disabled and can't be used. For example, in the preceding example, you can't convert all these ranges to tables by choosing Insert \Rightarrow Tables \Rightarrow Table.

In general, selecting a multi-sheet range is a simple two-step process: select the range in one sheet and then select the worksheets to include in the range. To select a group of contiguous worksheets, select the first worksheet in the group and then press Shift and click the sheet tab of the last worksheet that you want to include in the selection. To select individual worksheets, select one of the worksheets in the group and then press Ctrl and click the sheet tab of each additional worksheet that you want to select. If all the worksheets in a workbook aren't laid out the same, you can skip the sheets that you don't want to format. When you make the selection, the sheet tabs of the selected sheets display in bold with underlined text, and Excel displays Group in the title bar.

TIP

To select all sheets in a workbook, right-click any sheet tab and choose Select All Sheets from the shortcut menu.

CAUTION

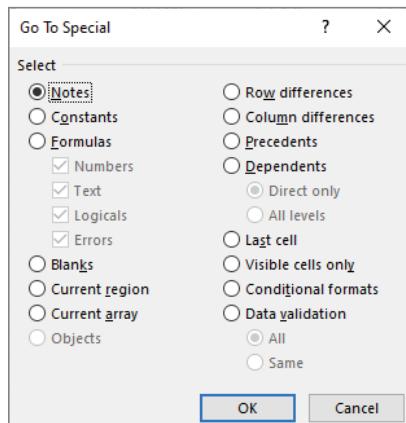
When sheets are grouped, you are making changes to sheets you can't see. Before you group sheets, be sure you understand what changes you intend to make and how that will affect all the sheets in the group. When you're done, don't forget to ungroup the sheets by selecting a sheet outside the group or selecting Ungroup Sheets from the tab's shortcut menu. You can overwrite data on the other sheets if you start typing on the active sheet while in Group mode.

Selecting special types of cells

As you use Excel, you may need to locate specific types of cells in your worksheets. For example, wouldn't it be handy to be able to locate every cell that contains a formula—or perhaps all the formula cells that depend on the active cell? Excel provides an easy way to locate these and many other special types of cells: select a range and choose Home \Rightarrow Editing \Rightarrow Find & Select \Rightarrow Go To Special to display the Go To Special dialog box, as shown in Figure 4.5.

FIGURE 4.5

Use the Go To Special dialog box to select specific types of cells.



Part I: Getting Started with Excel

After you make your choice in the dialog box, Excel selects the qualifying subset of cells in the current selection. Often, this subset of cells is a multiple selection. If no cells qualify, Excel lets you know with the message No cells were found.

TIP

If you bring up the Go To Special dialog box with only one cell selected, Excel bases its selection on the entire used area of the worksheet. Otherwise, the selection is based on the selected range.

Table 4.1 offers a description of the options available in the Go To Special dialog box.

TABLE 4.1 Go To Special Options

Option	What it does
Notes	Selects the cells that contain a note.
Constants	Selects all nonempty cells that don't contain formulas. Use the check boxes under the Formulas option to choose which types of nonformula cells to include.
Formulas	Selects cells that contain formulas. Qualify this by selecting the type of result: numbers, text, logical values (TRUE or FALSE), or errors.
Blanks	Selects all empty cells. If a single cell is selected when the dialog box is displayed, this option selects the empty cells in the used area of the worksheet.
Current Region	Selects a rectangular range of cells around the active cell. This range is determined by surrounding blank rows and columns. You can also press Ctrl+Shift+*.
Current Array	Selects the entire array. (See Chapter 10, "Understanding and Using Array Formulas," for more on arrays.)
Objects	Selects all embedded objects on the worksheet, including charts and graphics.
Row Differences	Analyzes the selection and selects cells that are different from other cells in each row.
Column Differences	Analyzes the selection and selects the cells that are different from other cells in each column.
Precedents	Selects cells that are referred to in the formulas in the active cell or selection (limited to the active sheet). You can select either direct precedents or precedents at all levels. (See Chapter 17, "Making Your Formulas Error-Free," for more information.)
Dependents	Selects cells with formulas that refer to the active cell or selection (limited to the active sheet). You can select either direct dependents or dependents at all levels. (See Chapter 17 for more information.)
Last Cell	Selects the bottom-right cell in the worksheet that contains data or formatting. For this option, the entire worksheet is examined, even if a range is selected when Excel displays the dialog box.

Option	What it does
Visible Cells Only	Selects only visible cells in the selection. This option is useful when dealing with a filtered list or a table.
Conditional Formats	Selects cells that have a conditional format applied (by choosing Home ▾ Styles ▾ Conditional Formatting). The All option selects all such cells. The Same option selects only the cells that have the same conditional formatting as the active cell.
Data Validation	Selects cells that are set up for data entry validation (by choosing Data ▾ Data Tools ▾ Data Validation). The All option selects all such cells. The Same option selects only the cells that have the same validation rules as the active cell.

TIP

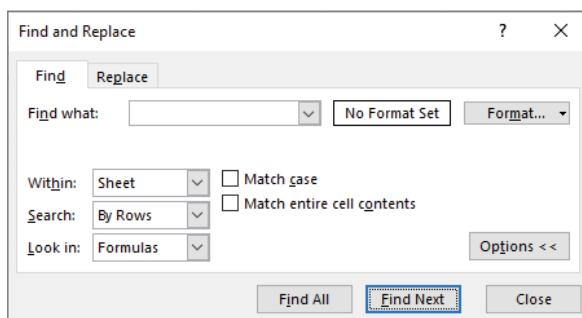
When you select an option in the Go To Special dialog box, be sure to note which suboptions become available. The placement of these suboptions can be misleading. For example, when you select Constants, the suboptions under Formulas become available to help you further refine the results. Likewise, the suboptions under Dependents also apply to Precedents, and those under Data Validation also apply to Conditional Formats.

Selecting cells by searching

Another way to select cells is to choose Home ▾ Editing ▾ Find & Select ▾ Find (or press Ctrl+F), which allows you to select cells by their contents. The Find and Replace dialog box is shown in Figure 4.6. This figure illustrates additional options that are available when you click the Options button.

FIGURE 4.6

The Find and Replace dialog box, with its options displayed

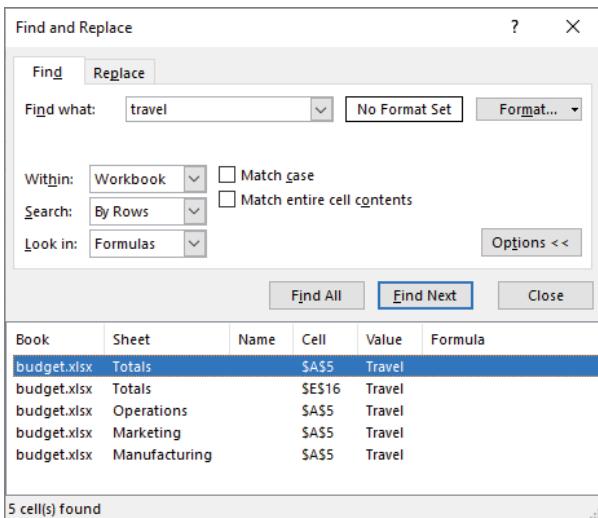


Part I: Getting Started with Excel

Enter the text you're looking for; then click Find All. The dialog box expands to display all the cells that match your search criteria. For example, Figure 4.7 shows the dialog box after Excel has located all cells that contain the text *travel*. You can click an item in the list, and the screen will scroll so that you can view the cell in context. To select all the cells in the list, first select any single item in the list. Then press Ctrl+A to select them all. You are limited to only selecting cells on the active sheet using this method.

FIGURE 4.7

The Find and Replace dialog box, with its results listed



NOTE

The Find and Replace dialog box allows you to return to the worksheet without dismissing the dialog box.

The Find and Replace dialog box supports two wildcard characters:

- | | |
|---|----------------------------------|
| ? | Matches any single character |
| * | Matches any number of characters |

Wildcard characters also work with values when the Match Entire Cell Contents option is selected. For example, searching for 3* locates all cells that contain a value that begins with 3. Searching for 1?9 locates all three-digit entries that begin with 1 and end with 9. Searching for *00 locates values that end with two zeros.

TIP

To search for a question mark or an asterisk, precede the character with a tilde (~). For example, the following search string finds the text *NONE*:

~*NONE~*

If you need to search for the tilde character, use two tildes.

If your searches don't seem to be working correctly, double-check these three options:

Match Case If this check box is selected, the case of the text must match exactly. For example, searching for smith does not locate Smith.

Match Entire Cell Contents If this check box is selected, a match occurs if the cell contains only the search string (and nothing else). For example, searching for Excel doesn't locate a cell that contains Microsoft Excel. When using wildcard characters, an exact match is not required.

Look In This drop-down list has four options: Values, Formulas, Notes, and Comments. The Formulas option looks only at the text that makes up the formula or the contents of the cell if there is no formula. The Values option looks at the cell value and the results, not the text, of the formula. If, for example, Formulas is selected, searching for 900 doesn't find a cell that contains the formula =899+1 but will find a cell with a value of 900. The Values option will find both of those cells.

Copying or Moving Ranges

As you create a worksheet, you may find it necessary to copy or move information from one location to another. Excel makes copying or moving ranges of cells easy. Here are some common things that you might do:

- Copy a cell to another location.
- Copy a cell to a range of cells. The source cell is copied to every cell in the destination range.
- Copy a range to another range.
- Move a cell or range of cells to another location.

The primary difference between copying and moving a range is the effect of the operation on the source range. When you copy a range, the source range is unaffected. When you move a range, the contents are removed from the source range.

NOTE

Copying a cell normally copies the cell's contents, any formatting that is applied to the original cell (including conditional formatting and data validation), the cell note (if it has one), and the cell comment (if it has one). When you copy a cell that contains a formula, the cell references in the copied formula are changed automatically to be relative to their new destination.

Copying or moving consists of two steps (although shortcut methods are available):

1. Select the cell or range to copy (the source range) and copy it to the Clipboard. To move the range instead of copying it, cut the range instead of copying it.
2. Select the cell or range that will hold the copy (the destination range) and paste the Clipboard contents.

CAUTION

When you paste information, Excel overwrites any cells that get in the way without warning you. If you find that pasting overwrote some essential cells, choose Undo from the Quick Access Toolbar (or press **Ctrl+Z**).

CAUTION

When you copy a cell or range, Excel surrounds the copied area with an animated, dashed border. As long as that border remains visible, the copied information is available for pasting. If you press **Esc** to cancel the border, Excel removes the information from the Clipboard.

Because copying (or moving) is used so often, Excel provides many different methods. We discuss each method in the following sections. Copying and moving are similar operations, so we point out only important differences between the two.

Copying by using Ribbon commands

Choosing **Home** \Rightarrow **Clipboard** \Rightarrow **Copy** transfers a copy of the selected cell or range to the Windows Clipboard and the Office Clipboard. After performing the copy part of this operation, select the destination cell and choose **Home** \Rightarrow **Clipboard** \Rightarrow **Paste**.

Instead of choosing **Home** \Rightarrow **Clipboard** \Rightarrow **Paste**, you can just activate the destination cell and press **Enter**. If you use this technique, Excel removes the copied information from the Clipboard so that it can't be pasted again.

If you're copying a range, you don't need to select an entire same-sized range before you click the Paste button. You only need to activate the upper-left cell in the destination range.

TIP

The **Home** \Rightarrow **Clipboard** \Rightarrow **Paste** control contains a drop-down arrow that, when clicked, gives you additional paste option icons. The paste preview icons are explained later in this chapter (see "Pasting in special ways").

About the Office Clipboard

Whenever you cut or copy information from a Windows program, Windows stores the information on the Windows Clipboard, which is an area of your computer's memory. Each time that you cut or copy information, Windows replaces the information previously stored on the Clipboard with the new information that you cut or copied. The Windows Clipboard can store data in a variety of formats. Because Windows manages information on the Clipboard, it can be pasted to other Windows applications, regardless of where it originated.

Microsoft Office has its own Clipboard (the Office Clipboard), which is available only in Office programs. To view or hide the Office Clipboard, click the dialog box launcher in the bottom-right corner of the Home ➔ Clipboard group.

Whenever you cut or copy information in an Office program, such as Excel or Word, the program places the information on both the Windows Clipboard and the Office Clipboard. However, the program treats information on the Office Clipboard differently from the way it treats information on the Windows Clipboard. Instead of replacing information on the Office Clipboard, the program appends the information to the Office Clipboard. With multiple items stored on the Clipboard, you can then paste the items either individually or as a group.

You can find out more about this feature (including an important limitation) in "Using the Office Clipboard to paste" later in this chapter.

Copying by using shortcut menu commands

If you prefer, you can use the following shortcut menu commands for copying and pasting:

- Right-click the range and choose Copy (or Cut) from the shortcut menu to copy the selected cells to the Clipboard.
- Right-click and choose Paste from the shortcut menu that appears to paste the Clipboard contents to the selected cell or range.

For more control over how the pasted information appears, right-click the destination cell and use one of the Paste icons in the shortcut menu (see Figure 4.8).

Instead of using Paste, you can just activate the destination cell and press Enter. If you use this technique, Excel removes the copied information from the Clipboard so that it can't be pasted again.

4

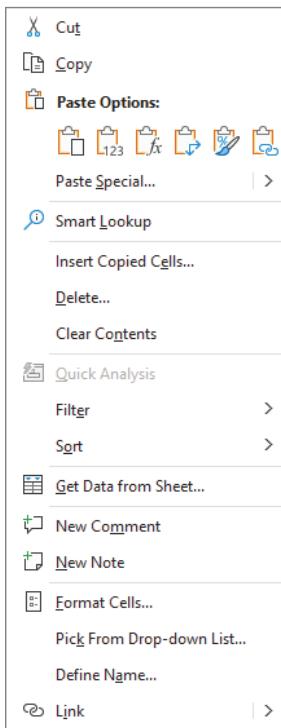
Copying by using shortcut keys

The copy and paste operations also have shortcut keys associated with them:

- Ctrl+C copies the selected cells to both the Windows Clipboard and the Office Clipboard.
- Ctrl+X cuts the selected cells to both the Windows Clipboard and the Office Clipboard.
- Ctrl+V pastes the Windows Clipboard contents to the selected cell or range.

FIGURE 4.8

The Paste icons on the shortcut menu provide more control over how the pasted information appears.

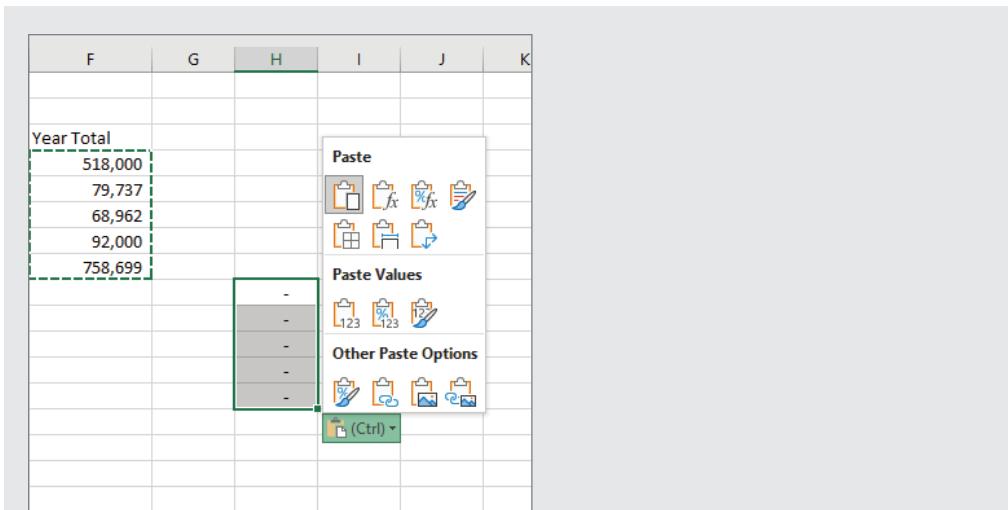


TIP

These are standard key combinations used by many other Windows applications.

Using Paste Options Buttons When Inserting and Pasting

Some cell and range operations—specifically inserting, pasting, and filling cells by dragging—result in the display of Paste Options buttons. For example, if you copy a range and then paste it to a different location using Home \Rightarrow Clipboard \Rightarrow Paste, a drop-down options list appears at the lower right of the pasted range. Click the list (or press Ctrl), and you see the options shown in the following figure. These options enable you to specify how the data should be pasted, such as values only or formatting only. In this case, using the Paste Options buttons is an alternative to using options in the Paste Special dialog box. (Read more about Paste Special in the upcoming section “Using the Paste Special dialog box.”)



To disable this feature, choose File ➔ Options and click the Advanced tab. Remove the check mark from the two options labeled Show Paste Options Button When Content Is Pasted and Show Insert Options Buttons.

Copying or moving by using drag-and-drop

Excel also enables you to copy or move a cell or range by dragging. Unlike other methods of copying and moving, dragging and dropping does not place any information on either the Windows Clipboard or the Office Clipboard.

CAUTION

The drag-and-drop method of moving does offer one advantage over the cut-and-paste method: Excel warns you if a drag-and-drop move operation will overwrite existing cell contents. Oddly, you do not get a warning if a drag-and-drop copy operation will overwrite existing cell contents.

To copy using drag-and-drop, select the cell or range that you want to copy, press Ctrl, and move the mouse to one of the selection's borders. (The mouse pointer is augmented with a small plus sign.) Then drag the selection to its new location while you continue to hold down the Ctrl key. The original selection remains behind, and Excel makes a new copy when you release the mouse button.

To move a range using drag-and-drop, don't press Ctrl while dragging the border.

TIP

If the mouse pointer doesn't turn into an arrow when you point to the border of a cell or range, you need to make a change to your settings. Choose File \Rightarrow Options to display the Excel Options dialog box, select the Advanced tab, and place a check mark on the option labeled Enable Fill Handle and Cell Drag-and-Drop.

Copying to adjacent cells

Often, you need to copy a cell to an adjacent cell or range. This type of copying is quite common when you're working with formulas. For example, if you're working on a budget, you might create a formula to add the values in column B. You can use the same formula to add the values in the other columns. Rather than re-enter the formula, you can copy it to the adjacent cells.

Excel provides additional options for copying to adjacent cells. To use these commands, activate the cell that you're copying and extend the cell selection to include the cells to which you're copying. Then issue the appropriate command from the following list for one-step copying:

- Home \Rightarrow Editing \Rightarrow Fill \Rightarrow Down (or Ctrl+D) copies the cell to the selected range below.
- Home \Rightarrow Editing \Rightarrow Fill \Rightarrow Right (or Ctrl+R) copies the cell to the selected range to the right.
- Home \Rightarrow Editing \Rightarrow Fill \Rightarrow Up copies the cell to the selected range above.
- Home \Rightarrow Editing \Rightarrow Fill \Rightarrow Left copies the cell to the selected range to the left.

None of these commands places information on either the Windows Clipboard or the Office Clipboard.

TIP

You also can use AutoFill to copy to adjacent cells by dragging the selection's fill handle (the small square in the bottom-right corner of the selected cell or range). Excel copies the original selection to the cells that you highlight while dragging. For more control over the AutoFill operation, click the AutoFill Options button that appears after you've released the mouse button, or drag the fill handle with the right mouse button. Each method shows a shortcut menu with additional options, although not the same menu.

Copying a range to other sheets

You can use the copy procedures described previously to copy a cell or range to another worksheet, even if the worksheet is in a different workbook. You must, of course, activate the other worksheet before you select the location to which you want to copy.

Excel offers a quicker way to copy a cell or range and paste it to other worksheets in the same workbook:

1. **Select the range to copy.**
2. **Press Ctrl and click the sheet tabs for the worksheets to which you want to copy the information.** Excel displays Group in the workbook's title bar.
3. **Choose Home ➔ Editing ➔ Fill ➔ Across Worksheets.** A dialog box appears to ask you what you want to copy (All, Contents, or Formats).
4. **Make your choice and then click OK.** Excel copies the selected range to the selected worksheets; the new copy occupies the same cells in the selected worksheets as the original occupies in the initial worksheet.

CAUTION

Be careful with the Home ➔ Editing ➔ Fill ➔ Across Worksheets command because Excel doesn't warn you when the destination cells contain information. You can quickly overwrite lots of cells with this command and not even realize it. So make sure that you check your work, and use Undo if the result isn't what you expected.

Using the Office Clipboard to paste

Whenever you cut or copy information in an Office program such as Excel, you can place the data on both the Windows Clipboard and the Office Clipboard. When you copy information to the Office Clipboard, you append the information to the Office Clipboard instead of replacing what is already there. With multiple items stored on the Office Clipboard, you can then paste the items either individually or as a group.

To use the Office Clipboard, you first need to open it. Use the dialog box launcher on the bottom right of the Home ➔ Clipboard group to toggle the Clipboard task pane on and off.

TIP

To make the Clipboard task pane open automatically, click the Options button near the bottom of the task pane and choose the Show Office Clipboard Automatically option.

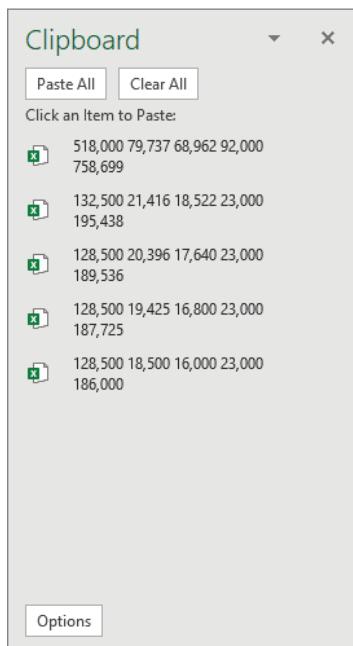
After you open the Clipboard task pane, select the first cell or range that you want to copy to the Office Clipboard and copy it by using any of the preceding techniques. Repeat this process, selecting the next cell or range that you want to copy. As soon as you copy the information, the Office Clipboard task pane shows you the number of items that you've copied and a

Part I: Getting Started with Excel

brief description (it will hold up to 24 items). Figure 4.9 shows the Office Clipboard with five copied items.

FIGURE 4.9

Use the Clipboard task pane to copy and paste multiple items.



When you're ready to paste information, select the cell into which you want to paste it. To paste an individual item, click it in the Clipboard task pane. To paste all the items that you've copied, click the Paste All button (which is at the top of the Clipboard task pane). The items are pasted, one after the other. The Paste All button is probably more useful in Word for situations in which you copy text from various sources and then paste it all at once.

You can clear the contents of the Office Clipboard by clicking the Clear All button.

The following items about the Office Clipboard and how it functions are worth noting:

- Excel pastes the contents of the Windows Clipboard (the last item you copied to the Office Clipboard) when you paste by choosing Home \Rightarrow Clipboard \Rightarrow Paste, by pressing Ctrl+V, or by right-clicking and choosing Paste from the shortcut menu.
- The last item that you cut or copied appears on both the Office Clipboard and the Windows Clipboard.
- Clearing the Office Clipboard also clears the Windows Clipboard.

CAUTION

The Office Clipboard has a serious problem that limits its usefulness for Excel users: if you copy a range that contains formulas, the formulas are not transferred when you paste from the Clipboard task pane to a different range. Only the values are pasted. Furthermore, Excel doesn't even warn you about this fact.

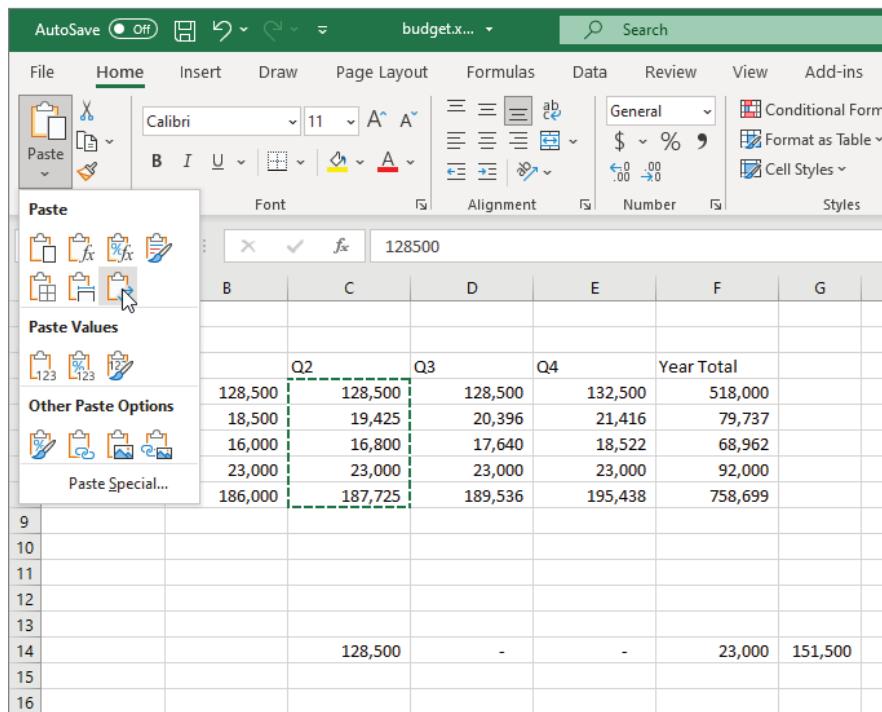
Pasting in special ways

You may not always want to copy everything from the source range to the destination range. For example, you may want to copy only the formula results rather than the formulas themselves. Or you may want to copy the number formats from one range to another without overwriting any existing data or formulas.

To control what is copied into the destination range, choose Home \Rightarrow Clipboard \Rightarrow Paste and use the drop-down menu shown in Figure 4.10. When you hover your mouse pointer over an icon, you'll see a preview of the pasted information in the destination range. Click the icon to use the selected paste option.

FIGURE 4.10

Excel offers several pasting options, with preview. Here, the information is copied from C4 : C8 and is being pasted beginning at cell C14 using the Transpose option.



The screenshot shows the Microsoft Excel interface with the 'Home' tab selected. In the top-left corner of the ribbon, there is a 'Paste' button with a dropdown arrow. A dropdown menu is open, showing several options under the 'Paste' heading, with the 'Transpose' option highlighted. The main area of the screen displays a table with financial data. The table has columns labeled Q2, Q3, Q4, and Year Total. The rows show values for various categories like 128,500, 18,500, etc. Row 14 is currently selected, and the value 128,500 is displayed in the cell C14. The entire range C4:C8 is highlighted with a green dashed border, indicating it is the source range for the paste operation.

	Q2	Q3	Q4	Year Total
128,500	128,500	128,500	132,500	518,000
18,500	19,425	20,396	21,416	79,737
16,000	16,800	17,640	18,522	68,962
23,000	23,000	23,000	23,000	92,000
186,000	187,725	189,536	195,438	758,699
9				
10				
11				
12				
13				
14	128,500	-	-	23,000 151,500
15				
16				

The paste options are as follows:

Paste (P) Pastes the cell's contents, formula, formats, and data validation from the Windows Clipboard.

Formulas (F) Pastes formulas but not formatting.

Formulas & Number Formatting (O) Pastes formulas and number formatting only.

Keep Source Formatting (K) Pastes formulas and all formatting.

No Borders (B) Pastes everything except borders that appear in the source range.

Keep Source Column Widths (W) Pastes formulas and duplicates the column width(s) of the copied cells.

Transpose (T) Changes the orientation of the copied range. Rows become columns, and columns become rows. Any formulas in the copied range are adjusted so that they work properly when transposed.

Merge Conditional Formatting (G) This icon is displayed only when the copied cells contain conditional formatting. When clicked, it merges the copied conditional formatting with any conditional formatting in the destination range.

Values (V) Pastes the results of formulas. The destination for the copy can be a new range or the original range. In the latter case, Excel replaces the original formulas with their current values.

Values & Number Formatting (A) Pastes the results of formulas plus the number formatting.

Values & Source Formatting (E) Pastes the results of formulas plus all formatting.

Formatting (R) Pastes only the formatting of the source range.

Paste Link (N) Creates formulas in the destination range that refer to the cells in the copied range.

Picture (U) Pastes the copied information as a picture.

Linked Picture (I) Pastes the copied information as a "live" picture that is updated if the source range is changed.

Paste Special Displays the Paste Special dialog box (described in the next section).

NOTE

After you paste, you're offered another chance to change your mind. A Paste Options button appears at the lower right of the pasted range. Click it (or press Ctrl), and you will see the paste option icons again.

Using the Paste Special dialog box

For yet another pasting method, choose Home \Rightarrow Clipboard \Rightarrow Paste \Rightarrow Paste Special to display the Paste Special dialog box (see Figure 4.11). You can also right-click and choose Paste Special from the shortcut menu to display this dialog box. This dialog box has several options,

some of which are identical to the buttons in the Paste drop-down menu. The options that are different are explained in the following list.

Comments and Notes Copies only the cell comments and cell notes from a cell or range. This option doesn't copy cell contents or formatting.

Validation Copies the validation criteria so that the same data validation will apply. Data validation is applied by choosing Data \Rightarrow Data Tools \Rightarrow Data Validation.

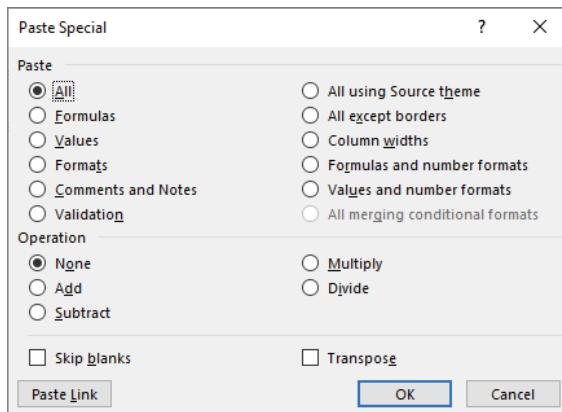
All using Source theme Pastes everything but uses the formatting from the document theme of the source. This option is relevant only if you're pasting information from a different workbook and the workbook uses a different document theme than the active workbook.

Column widths Pastes only column width information.

All merging conditional formats Merges the copied conditional formatting with any conditional formatting in the destination range. This option is enabled only when you're copying a range that contains conditional formatting.

FIGURE 4.11

The Paste Special dialog box



Note

Excel has several different Paste Special dialog boxes, each with different options. The one displayed depends on what's copied. This section describes the Paste Special dialog box that appears when a range or cell has been copied.

TIP

For the Paste Special command to be available, you need to copy a cell or range. (Choosing Home \Rightarrow Clipboard \Rightarrow Cut doesn't work.)

In addition, the Paste Special dialog box enables you to perform other operations, described in the following sections.

Performing mathematical operations without formulas

The option buttons in the Operation section of the Paste Special dialog box let you perform an arithmetic operation on values and formulas in the destination range. For example, you can copy a range to another range and select the Multiply operation. Excel multiplies the corresponding values in the source range and the destination range and replaces the destination range with the new values.

This feature also works with a single copied cell, pasted to a multicell range. Assume that you have a range of values, and you want to increase each value by 5 percent. Enter 105% into any blank cell and copy that cell to the Clipboard. Then select the range of values and bring up the Paste Special dialog box. Select the Multiply option, and each value in the range is multiplied by 105%.

CAUTION

If the destination range contains formulas, the formulas are also modified. In many cases, this is not what you want.

Skipping blanks when pasting

The Skip Blanks option in the Paste Special dialog box prevents Excel from overwriting cell contents in your paste area with blank cells from the copied range. This option is useful if you're copying a range to another area but don't want the blank cells in the copied range to overwrite existing data.

Transposing a range

The Transpose option in the Paste Special dialog box changes the orientation of the copied range. Rows become columns, and columns become rows. Any formulas in the copied range are adjusted so that they work properly when transposed. Note that you can use this check box with the other options in the Paste Special dialog box. Figure 4.12 shows an example of a horizontal range (A3 : F8) that was transposed to a different range (A11 : F16).

TIP

If you click the Paste Link button in the Paste Special dialog box, you create formulas that link to the source range. As a result, the destination range automatically reflects changes in the source range.

FIGURE 4.12

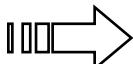
Transposing a range changes the orientation as the information is pasted into the worksheet.

	A	B	C	D	E	F
1	Operations					
2						
3		Q1	Q2	Q3	Q4	Year Total
4	Salaries	128,500	128,500	128,500	132,500	518,000
5	Travel	18,500	19,425	20,396	21,416	79,737
6	Supplies	16,000	16,800	17,640	18,522	68,962
7	Facility	23,000	23,000	23,000	23,000	92,000
8	Total	186,000	187,725	189,536	195,438	758,699
9						
10						
11		Salaries	Travel	Supplies	Facility	Total
12	Q1	128500	18500	16000	23000	186000
13	Q2	128500	19425	16800	23000	187725
14	Q3	128500	20396.25	17640	23000	189536.25
15	Q4	132500	21416.0625	18522	23000	195438.0625
16	Year Total	518000	79737.3125	68962	92000	758699.3125
17						

Using Names to Work with Ranges

Dealing with cryptic cell and range addresses can sometimes be confusing, especially when you work with formulas, which we cover in Chapter 9, “Introducing Formulas and Functions.” Fortunately, Excel allows you to assign descriptive names to cells and ranges. For example, you can give a cell a name such as `Interest_Rate`, or you can name a range `JulySales`. Working with these names (rather than cell or range addresses) has several advantages:

- A meaningful range name (such as `Total_Income`) is much easier to remember than a cell address (such as `AC21`).
- Entering a name is less error prone than entering a cell or range address, and if you type a name incorrectly in a formula, Excel will display a `#NAME?` error.
- You can quickly move to areas of your worksheet either by using the Name box, located at the left side of the Formula bar (click the arrow to drop down a list of defined names) or by choosing Home \Rightarrow Editing \Rightarrow Find & Select \Rightarrow Go To (or pressing F5 or Ctrl+G) and specifying the range name.
- Creating formulas is easier. You can paste a cell or range name into a formula by using Formula AutoComplete.



See Chapter 9 for information on Formula AutoComplete.

- Names make your formulas more understandable and easier to use. A formula such as `=Income-Taxes` is certainly more intuitive than `=D20-D40`.

Creating range names in your workbooks

Excel provides several methods that you can use to create range names. Before you begin, however, you should be aware of a few rules:

- Names can't contain spaces. You may want to use an underscore character to simulate a space (such as Annual_Total).
- You can use any combination of letters and numbers, but the name must begin with a letter, underscore, or backslash. A name can't begin with a number (such as 3rdQuarter) or look like a cell address (such as QTR3). If these are desirable names, though, you can precede the name with an underscore or a backslash, for example, _3rdQuarter and \QTR3.
- Symbols—except for underscores, backslashes, and periods—are not allowed.
- Names are limited to 255 characters, but it's a good practice to keep names as short as possible yet still meaningful.

CAUTION

Excel also uses a few names internally for its own use. Although you can create names that override Excel's internal names, you should avoid doing so. To be on the safe side, avoid using the following for names: Print_Area, Print_Titles, Consolidate_Area, and Sheet_Title. To learn how to delete a range name or rename a range, see "Managing names" later in this chapter.

Using the Name box

The fastest way to create a name is to use the Name box (to the left of the Formula bar). Select the cell or range to name, click the Name box, and type the name. Press Enter to create the name. (You must press Enter to actually record the name; if you type a name and then click in the worksheet, Excel doesn't create the name.)

If you type an invalid name (such as May21, which happens to be a cell address, MAY21), Excel activates that address and doesn't warn you that the name is not valid. If the name you type includes an invalid character, Excel displays an error message. If a name already exists, you can't use the Name box to change the range to which that name refers. Attempting to do so simply selects the range.

The Name box is a drop-down list and shows all names in the workbook. To choose a named cell or range, click the arrow on the right side of the Name box and choose the name. The name appears in the Name box, and Excel selects the named cell or range in the worksheet.

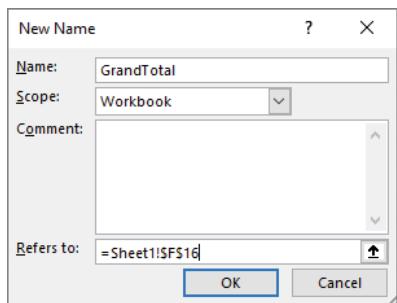
Using the New Name dialog box

For more control over naming cells and ranges, use the New Name dialog box. Start by selecting the cell or range that you want to name. Then choose Formulas \Rightarrow Defined Names \Rightarrow Define

Name. Excel displays the New Name dialog box, shown in Figure 4.13. Note that this is a resizable dialog box. Click and drag a border to change the dimensions.

FIGURE 4.13

Create names for cells or ranges by using the New Name dialog box.



Type a name in the Name text field (or use the name that Excel proposes, if any). The selected cell or range address appears in the Refers To text field. Use the Scope drop-down list to indicate the scope for the name. The scope indicates where the name will be valid, and it's either the entire workbook or the worksheet in which the name is defined. If you like, you can add a comment that describes the named range or cell. Click OK to add the name to your workbook and close the dialog box.

Using the Create Names from Selection dialog box

You may have a worksheet that contains text that you want to use for names for adjacent cells or ranges. For example, you may want to use the text in column A to create names for the corresponding values in column B. Excel makes this task easy.

To create names by using adjacent text, start by selecting the name text and the cells that you want to name. (These items can be individual cells or ranges of cells.) The names must be adjacent to the cells that you're naming. (A multiple selection is allowed.) Then choose Formulas ➔ Defined Names ➔ Create from Selection. Excel displays the Create Names from Selection dialog box, shown in Figure 4.14.

The check marks in the Create Names from Selection dialog box are based on Excel's analysis of the selected range. For example, if Excel finds text in the first row of the selection, it proposes that you create names based on the top row. If Excel didn't guess correctly, you can change the check boxes. Click OK, and Excel creates the names. Using the data in Figure 4.14, Excel creates the seven named ranges shown in Figure 4.15.

Part I: Getting Started with Excel

FIGURE 4.14

Use the Create Names from Selection dialog box to name cells using labels that appear in the worksheet.

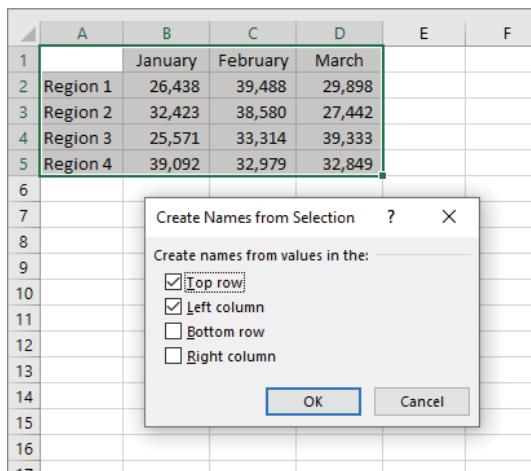
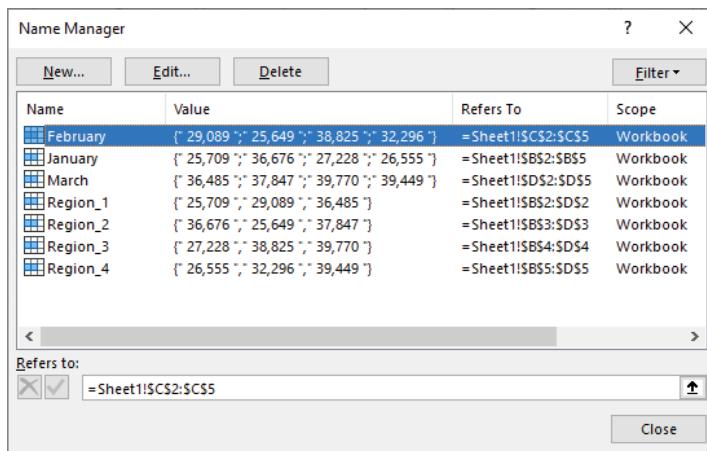


FIGURE 4.15

Use the Name Manager to work with range names.



Note

If the text contained in a cell would result in an invalid name, Excel modifies the name to make it valid. For example, if a cell contains the text Net Income (which is invalid for a name because it contains a space), Excel converts the space to an underscore character. If Excel encounters a value or a numeric formula where text should be, however, it doesn't convert it to a valid name. It simply doesn't create a name, and it does not inform you of that fact.

Caution

If the upper-left cell of the selection contains text and you choose the Top Row and Left Column options, Excel uses that text for the name of the entire range, excluding the top row and left column. So, after Excel creates the names, take a minute to make sure that they refer to the correct ranges. If Excel creates a name that is incorrect, you can delete or modify it by using the Name Manager (described next).

Managing names

A workbook can have any number of named cells and ranges. If your workbook has many names, you should know about the Name Manager, which is shown in Figure 4.15.

The Name Manager appears when you choose Formulas \Rightarrow Defined Names \Rightarrow Name Manager (or press Ctrl+F3). The Name Manager has the following features:

Displays information about each name in the workbook You can resize the Name Manager dialog box, widen the columns to show more information, and even rearrange the order of the columns. You can also click a column heading to sort the information by the column.

Allows you to filter the displayed names Clicking the Filter button lets you show only those names that meet certain criteria. For example, you can view only the worksheet-level names.

Provides quick access to the New Name dialog box Click the New button to create a new name without closing the Name Manager.

Lets you edit names To edit a name, select it in the list, and then click the Edit button or double-click the name. You can change the name itself, modify the Refers To range, or edit the comment.

Lets you quickly delete unneeded names To delete a name, select it in the list and click Delete.

Caution

Be extra careful when deleting names. If the name is used in a formula, deleting the name causes the formula to become invalid. (It displays #NAME?.) It seems logical that Excel would replace the name with its actual address, but that doesn't happen. However, deleting a name can be undone, so if you find that formulas return #NAME? after you delete a name, choose Undo from the Quick Access Toolbar (or press Ctrl+Z) to get the name back.

If you delete the rows or columns that contain named cells or ranges, the names contain an invalid reference. For example, if cell A1 on Sheet1 is named Interest and you delete row 1 or column A, the name Interest then refers to =Sheet1!#REF! (an erroneous reference). If you use the name Interest in a formula, the formula displays #REF!.

TIP

To create a list of names in a worksheet, first select a cell in an empty area of your worksheet. The list is created at the active cell position and overwrites any information at that location. Press F3 to display the Paste Name dialog box, which lists all the defined names, and then click the Paste List button. Excel creates a list of all names in the workbook and their corresponding addresses.

Adding Comments to Cells

In recent versions of Excel, Microsoft has added a new way to annotate cells. Prior to this change, you could add a comment to a cell; the comment consisted of a simple text box containing whatever text you wished to add. What used to be called a comment is now called a *note* and is discussed in the next section. The new method of annotation is called a *comment*. It is a threaded list of comments and is designed to facilitate a conversation among users.

To add a comment to a cell, select the cell and perform any of these commands:

- Choose Review \Rightarrow Comments \Rightarrow New Comment.
- Right-click the cell and choose New Comment from the shortcut menu.
- Press Ctrl+Shift+F2.

Excel attaches a new, blank comment to the cell, as shown in Figure 4.16. To complete the comment, type in the Start a conversation text box and click the Post button. If you don't type anything and click Post, the comment will not be created. When you post, Excel creates the comment including which user posted it and the date and time it was posted. Figure 4.17 shows a completed comment.

TIP

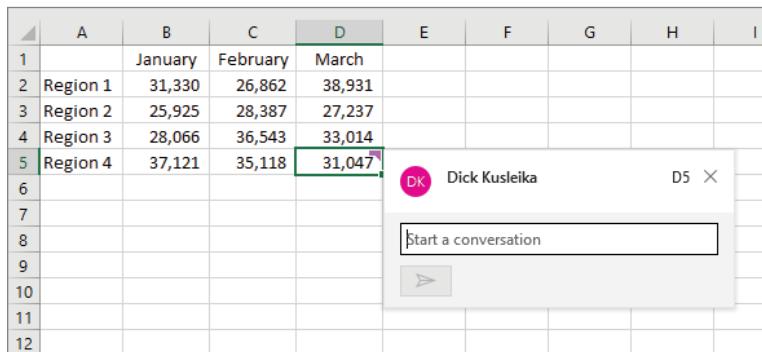
To post a comment or reply, you can press Ctrl+Enter instead of clicking the Post button.

TIP

To convert a cell note into a cell comment, select a cell with an existing note and choose Review \Rightarrow Notes \Rightarrow Convert to Comments. Excel doesn't time-stamp a converted comment the way it does with new ones.

FIGURE 4.16

Creating a comment attaches a new, blank comment to a cell.

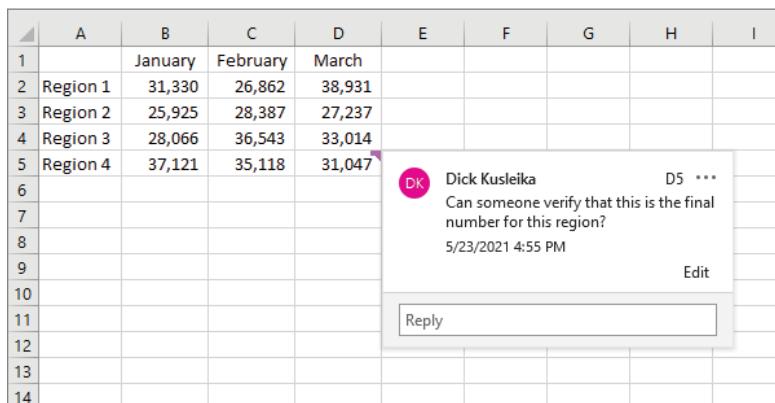


A screenshot of a Microsoft Excel spreadsheet. The data is organized into columns A through I and rows 1 through 12. The first four rows represent different regions with their respective January, February, and March values. Row 5 is highlighted in green and contains the value '31,047' in cell D5. A comment box is overlaid on the cell, showing the name 'Dick Kusleika' and the cell reference 'D5'. The comment text input field contains the placeholder 'Start a conversation'. There is also a small 'OK' button and a close 'X' button in the top right corner of the comment box.

	A	B	C	D	E	F	G	H	I
1		January	February	March					
2	Region 1	31,330	26,862	38,931					
3	Region 2	25,925	28,387	27,237					
4	Region 3	28,066	36,543	33,014					
5	Region 4	37,121	35,118	31,047	OK	Dick Kusleika	D5	X	
6									
7									
8									
9									
10									
11									
12									

FIGURE 4.17

Typing a comment and clicking the Post button finishes the comment creation process.



A screenshot of a Microsoft Excel spreadsheet, identical to Figure 4.16, showing a comment being created in cell D5. The comment box is open over cell D5, displaying the name 'Dick Kusleika', the cell reference 'D5', and the message 'Can someone verify that this is the final number for this region?'. Below the message is the timestamp '5/23/2021 4:55 PM'. At the bottom right of the comment box is an 'Edit' button. The 'Reply' input field is visible at the bottom of the comment box.

	A	B	C	D	E	F	G	H	I
1		January	February	March					
2	Region 1	31,330	26,862	38,931					
3	Region 2	25,925	28,387	27,237					
4	Region 3	28,066	36,543	33,014					
5	Region 4	37,121	35,118	31,047	OK	Dick Kusleika	D5	***	
6									
7									
8									
9									
10									
11									
12									
13									
14									

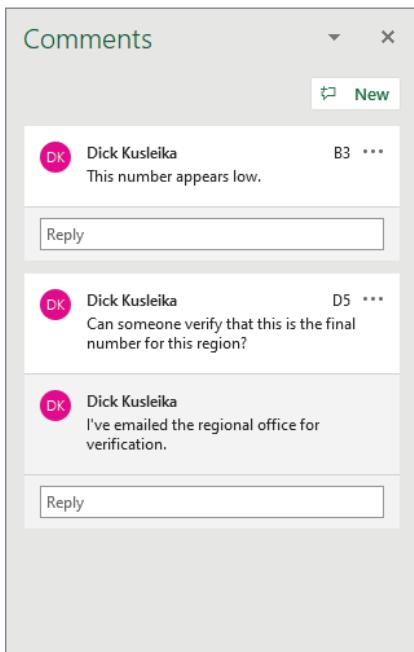
Showing comments

Cells with comments have a small mark in the upper-right corner. To view a comment thread for a cell, select the cell that contains it. Hovering your mouse over the cell will display the comment, but it will disappear when you move your mouse away. Clicking on a comment while it's displayed will select the cell. You can also right-click on the cell and choose Reply to Comment to view it.

To see all the comments on the worksheet, chose Review \Rightarrow Comments \Rightarrow Show Comments. This action opens the Comments task pane that lists all the comments on the active worksheet. Figure 4.18 shows two comment threads, one in cell B3 and one in cell D5 that also includes a reply.

FIGURE 4.18

Showing all comments in the Comments task pane.



The Comments task pane allows you to create new comments and reply to existing ones. The New button will create a new comment if the active cell doesn't already contain a comment or will put the cursor in the reply box for the comment already attached to the active cell.

The Ribbon contains two commands to cycle through all the comments on the worksheet: Choose Review \Rightarrow Comments \Rightarrow Previous Comment or Review \Rightarrow Comments \Rightarrow Next Comment to move through the comments. If the Comments task pane is visible, these actions will select the comments in the task pane. If it's not visible, the comments will appear next to the cells to which they're attached.

Replies to comments

Each comment has a Reply box for entering additional text, threaded underneath the comment you reply to. When you start a reply, the comment box displays Post and Cancel buttons to complete the reply. You can only reply to the last comment in the thread.

Editing comments and replies

The comment box contains an Edit button that allows you to change the text of the comment if you are the user that created it. Editing a comment does not update the date and time.

stamp on it. Replies also have an Edit button if you are the user who replied. The Edit button is only visible when that comment or reply is active or your mouse is hovering over it.

When you're editing a comment, Save and Cancel buttons are displayed to confirm your edits or revert to the old text.

Deleting comments and replies

The first comment in a thread has an ellipsis in the upper-right corner. This menu contains two entries: Delete Thread and Resolve Thread (discussed in the next section), as shown in Figure 4.19. Clicking Delete Thread deletes the comment and all comments below it. You can also right-click on the cell and choose Delete Comment or choose Review \Rightarrow Comments \Rightarrow Delete to delete the entire thread. If you delete a thread by mistake, choose Undo from the Quick Access Toolbar or press Ctrl+Z.

FIGURE 4.19

The first comment in a thread has an ellipsis menu to delete the thread.

26,862	38,931			
28,387	27,237			
36,543	33,014			
35,118	31,047			
		Dick Kusleika Can someone number for t 5/23/2021 4:5	D5 ...	
			Delete thread	
			Resolve thread	
			Edit	
		Dick Kusleika I've emailed the regional office for verification. 5/23/2021 5:04 PM		
		Reply		

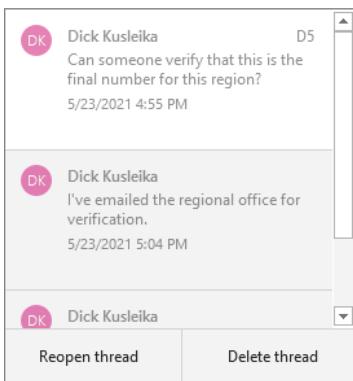
Replies to comments have a Delete button instead of the ellipsis menu. When you click Delete on a reply, only that reply is deleted. Any replies below the deleted reply remain in the thread.

Resolving comment threads

The other menu item in the ellipsis menu is Resolve Thread. When you resolve a thread, you prevent any more replies from being posted to it. Resolved threads are grayed out and two new buttons appear at the bottom of the thread: Reopen Thread and Delete Thread. Figure 4.20 shows a comment that has been resolved.

FIGURE 4.20

A resolved comment is grayed out and buttons to reopen or delete are displayed at the bottom.



Reopening a thread changes it back to its state prior to being resolved. That is, it's no longer grayed out and users can reply to it. Anyone with write access to the workbook can resolve or reopen threads, not just the user who created it. Deleting a resolved thread works just like deleting an open thread. The advantage of resolving a thread is to signal to other users that all the issues have been handled but to keep the thread there as an artifact. Resolved threads still show on the Comments task pane and are still displayed if you cycle through all the threads.

Adding Notes to Cells

What used to be called a comment is now called a note. Notes are simple text boxes that are attached to cells. While they are simple, they support resizing and some formatting. Unlike the new comment feature, you can't reply to a note, but you can add additional text to the bottom of it. Notes are useful when you need to describe a particular value or explain how a formula works.

To add a note to a cell, select the cell and use any of these actions:

- Choose Review \Rightarrow Notes \Rightarrow Notes \Rightarrow New Note.
- Right-click the cell and choose New Note from the shortcut menu.
- Press Shift+F2.

Excel inserts a note that points to the active cell. Initially, the note consists of your name, as specified in the General tab of the Excel Options dialog box (choose File \Rightarrow Options to display this dialog box). If you like, you can delete your name from the note. Enter the text for the cell note and then click anywhere in the worksheet to hide the note. You can change the size of the note by clicking and dragging any one of its eight sizing handles. Figure 4.21 shows a cell with a note.

FIGURE 4.21

You can add a note to a cell to help document your worksheets.

The screenshot shows a portion of an Excel spreadsheet. The columns are labeled A through E and the rows are numbered 1 through 14. Row 1 contains column headers: January, February, March. Rows 2 through 5 contain data for four regions: Region 1 (31,330, 26,862, 38,931), Region 2 (25,925, 28,387, 27,237), Region 3 (28,066, 36,543, 33,014), and Region 4 (37,121, 35,118, 31,047). A red triangle in the upper-right corner of cell D5 indicates a note is present. A callout box with a black border and yellow background is positioned over cell D5, containing the text "Dick Kusleika: This is an estimate".

	A	B	C	D	E
1		January	February	March	
2	Region 1	31,330	26,862	38,931	
3	Region 2	25,925	28,387	27,237	
4	Region 3	28,066	36,543	33,014	
5	Region 4	37,121	35,118	31,047	
6					
7					
8					
9					
10		Dick Kusleika: This is an estimate			
11					
12					
13					
14					

Cells that have a note display a small red triangle in the upper-right corner. When you move the mouse pointer over a cell that contains a note, the note becomes visible.

You can force a note to be displayed even when the mouse is not hovering over the cell. Right-click the cell and choose Show/Hide Note. The menu item toggles between always showing the note and returning to just the red triangle and only displaying the note when you hover over it.

TIP

You can control the way comments and notes are displayed. Choose File \Rightarrow Options and then select the Advanced tab of the Excel Options dialog box. In the Display section, select from one of the following three options:

- No Comments, Notes, or Indicators: This option hides everything. To re-display the notes after hiding them, choose Review \Rightarrow Notes \Rightarrow Notes \Rightarrow Show All Notes from Ribbon. This Ribbon command also sets the Excel Option to Indicators and notes, comments on hover (the third option). Or you can choose Review \Rightarrow Notes \Rightarrow Notes, and then choose any of Edit Note, Next Note, or Previous Note to display only one note. Choosing one of these commands does not change the setting under File \Rightarrow Options.
- Indicators Only, and Comments and Notes on Hover: This option is the default.
- Indicators and Notes, and Comments on Hover: This option always displays notes. There is no option to always display comments, but you can keep the Comments task pane visible, as discussed in the previous section.

4

Showing notes

If you want all cell notes to be visible (regardless of the location of the mouse), choose Review \Rightarrow Notes \Rightarrow Notes \Rightarrow Show All Notes. This command is a toggle; select it again to hide all cell notes.

To toggle the display of an individual note, select its cell and then choose Review \Rightarrow Notes \Rightarrow Notes \Rightarrow Show/Hide Note. The Notes button in the Notes group on the Ribbon also contains a Next Note and a Previous Note command. Choosing one of these commands displays the next or previous note and hides the current one, if any.



See Chapter 7, “Printing Your Work,” for information about including comments and notes on your printouts.

Formatting notes

If you don’t like the default look of cell notes, you can make some changes. Right-click the cell and choose Edit Note. Select the text in the note and use the commands of the Font and the Alignment groups on the Home tab to make changes to the note’s appearance.

For even more formatting options, choose Home \Rightarrow Cells \Rightarrow Format \Rightarrow Format Comment or right-click on the note’s border and choose Format Comment from the shortcut menu. Excel responds by displaying the Format Comment dialog box, which allows you to change many aspects of a comment’s appearance, including color, border, and margins.

TIP

Don’t be confused that the Ribbon and shortcut menu use the term *comment* instead of *note*. This is likely an oversight on Microsoft’s part, and we expect it will be corrected in a future release.

CAUTION

The Format Comment dialog box shows different tabs depending on whether you’re editing the note. When you’re editing the note, that is, when your cursor is an insertion point in the note’s text, the Format Comment dialog box only shows the Font tab. When the note is selected but not being edited, all the available tabs are shown.

TIP

You can also display an image inside a note. Right-click the cell and choose Edit Note. Then right-click the note’s border (not the note itself) and choose Format Comment. Select the Colors and Lines tab in the Format Comment dialog box. Click the Color drop-down list and select Fill Effects. In the Fill Effects dialog box, click the Picture tab and then click the Select Picture button to specify a graphics file. Figure 4.22 shows a note that contains a picture.

FIGURE 4.22

This note contains an image.

	A	B	C	D	E
1		January	February	March	
2	Region 1	31,330	26,862	38,931	
3	Region 2	25,925	28,387	27,237	
4	Region 3	28,066	36,543	33,014	
5	Region 4	37,121	35,118	31,047	
6					
7					
8					
9					
10					
11					
12					
13					
14					
15					
16					



An Alternative to Cell Notes

You can use Excel's Data Validation feature (see Chapter 23, "Using Data Validation") to add a different type of note to a cell. This type of note appears automatically when the cell is selected. Follow these steps:

1. Select the cell that will contain the note.
2. Choose **Data** ➔ **Data Tools** ➔ **Data Validation**. The Data Validation dialog box appears.
3. Click the **Input Message** tab.
4. Make sure that the **Show Input Message When Cell Is Selected** check box is selected.
5. Type your note in the **Input Message** box.
6. (Optional) Type a title in the **Title** box. This text will appear in bold at the top of the message.
7. Click **OK** to close the Data Validation dialog box.

After you perform these steps, the message appears when the cell is activated, and it disappears when any other cell is activated.

Note that this message isn't a "real" note. For example, a cell that contains this type of message doesn't display a note indicator, and it's not affected by any of the commands used to work with cell notes. In addition, you can't format these messages in any way, and you can't print them.

Editing notes

To edit the text in a note, activate the cell, right-click, and then choose Edit Note from the shortcut menu. Or select the cell and press Shift+F2. After you make your changes, click any cell.

Deleting notes

To delete a cell note, activate the cell that contains the note and then choose Review ➔ Comments ➔ Delete. Or right-click and then choose Delete Note from the shortcut menu.

Working with Tables

A *table* is a specially designated area of a worksheet. When you designate a range as a table, Excel gives it special properties that make certain operations easier and that help prevent errors.

The purpose of a table is to enforce some structure around your data. If you're familiar with a table in a database (like Microsoft Access), then you already understand the concept of structured data. If not, don't worry. It's not difficult.

In a table, each row contains information about a single entity. In a table that holds employee information, each row will contain information about one employee (such as name, department, and hire date). Each column contains the same piece of information for each employee. The same column that holds the hire date for the first employee holds the hire date for all the other employees.

Understanding a table's structure

Figure 4.23 shows a simple table. The various components of a table are described in the following sections.

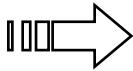
FIGURE 4.23

The areas that make up a table

The diagram illustrates the structure of a table in Excel. It shows a grid of data with various components labeled:

- Header row:** The top row of the table, containing column headers like EmployeeID, Name, Department, etc.
- Sort and filter control:** A set of small icons located at the top right of the header row, used for sorting and filtering data.
- Banded rows:** Rows 1 through 3 are highlighted with a light blue background, representing alternating row colors.
- Total row:** Row 4, which contains a summary value (\$ 190,000) and is also highlighted with a light blue background.

	A	B	C	D	E	F	G
1	EmployeeID	Name	Department	Hire Date	Salary	Days Employed	
2	00001	John Smith	Marketing	4/8/2001	\$ 80,000	7,358	
3	00002	Jane Burns	Accounting	12/12/1996	\$ 110,000	8,936	
4	Total				\$ 190,000		
5							
6							



This workbook, named EmployeeTable.xlsx, is available on this book's website at www.wiley.com/go/excel365bible.

The header row

The *header row* is generally colored differently than the other rows. The names in the header identify the columns. If you have a formula that refers to a table, the header row will determine how the column is referred to. For example, the Days Employed column (column F) contains a formula that refers to the Hire Date column (column D). The formula is =NOW() - [@ [Hire Date]]. If your table is longer than one screen, the header row will replace the normal column headers in Excel when you scroll down.



See Chapter 16, “Using Formulas with Tables and Conditional Formatting,” to learn about special table referencing in formulas.

The header also contains *Filter Buttons*. These drop-down buttons work exactly like Excel’s normal AutoFilter feature. You can use them to sort and filter the table’s data.

The data body

The *data body* is one or more rows of data. By default, the rows are *banded*, that is, formatted with alternating colors. When you add new data to the table, the formatting of the existing data is applied to the new data. For example, if a column is formatted as Text, that column in the new row will also be formatted as Text. The same is true for conditional formatting.

It’s not just formatting that applies to the new data. If a column contains a formula, that formula is automatically inserted into the new row. Data validation will also be transferred. You can make a robust data entry area knowing that the table structure will apply to new data.

One of the best features of tables is that as the data body expands, anything that refers to the table will expand automatically. If you were to base a PivotTable or a chart on your table, the PivotTable or chart would adjust as you added or deleted rows from the table.

The total row

The *total row* is not visible by default when you create a table. To show the total row, check the Total Row check box in the Table Style Options group on the Table Design Ribbon. When you show the total row, the text *Total* is placed in the first column. You can change this to another value or to a formula.

Each cell in the total row has a drop-down arrow with a list of common functions. It’s no accident that the list of functions resembles the arguments for the SUBTOTAL function. When you select a function from the list, Excel inserts a SUBTOTAL formula in the cell. The SUBTOTAL function ignores filtered cells, so the total will change if you filter the table.

In addition to the list of functions in the total row, there is a More Functions option at the bottom of the drop-down list. Selecting this option shows the Insert Function dialog box and makes all of Excel's functions available to you. Beyond that, you can simply type whatever formula you want in the total row.

The resizing handle

At the bottom right of the last cell in the table is the *resizing handle*. You can drag this handle to change the size of the table. Increasing the length of the table adds blank rows, copying down formatting, formulas, and data validation. Increasing the width of the table adds new columns with generic names like Column1, Column2, and so forth. You can change those names to something more meaningful.

Decreasing the size of the table simply changes what data is considered part of the table. It does not delete any data, formatting, formulas, or data validation. If you want to change what's in your table, you're better off deleting the columns and rows as you would any range rather than trying to do it with the resizing handle.

Creating a table

Most of the time, you'll create a table from an existing range of data. However, Excel also allows you to create a table from an empty range so that you can fill in the data later. The following instructions assume that you already have a range of data that's suitable for a table:

1. **Make sure the range doesn't contain any completely blank rows or columns; otherwise, Excel will not guess the table range correctly.**
2. **Select any cell within the range.**
3. **Choose Insert ▾ Tables ▾ Table (or press Ctrl+T).** Excel responds with its Create Table dialog box, shown in Figure 4.24. Excel tries to guess the range, as well as whether the table has a header row. Most of the time, it guesses correctly. If not, make your corrections before you click OK.

The range is converted to a table (using the default table style), and the Table Design tab of the Ribbon appears.

NOTE

Excel may not guess the table's dimensions correctly if the table isn't separated from other information by at least one empty row or column. If Excel guesses incorrectly, just specify the exact range for the table in the Create Table dialog box. Better yet, click Cancel and rearrange your worksheet such that the table is separated from your other data by at least one blank row and column.

To create a table from an empty range, select the range and choose Insert ▾ Tables ▾ Table. Excel creates the table, adds generic column headers (such as Column1 and Column2), and applies table formatting to the range. Almost always, you'll want to replace the generic column headers with more meaningful text.

FIGURE 4.24

Use the Create Table dialog box to verify that Excel guessed the table dimensions correctly.

Agent	Date Listed	Area	List Price	Bedrooms	Baths	SqFt	Type	Pool	Sold
Adams	1/22/2022	Central	295,000	4	3.5	2,100	Condo	TRUE	TRUE
Adams	4/10/2022	Central	384,000	4	2	2,700	Condo	FALSE	FALSE
Adams	10/4/2022	Dow			2.5	1,800	Split Level	FALSE	TRUE
Adams	1/19/2022	Dow			3	1,900	Ranch	FALSE	TRUE
Adams	4/14/2022	Dow			3.5	2,200	Split Level	FALSE	TRUE
Adams	3/17/2022	Dow			3	2,200	Ranch	FALSE	TRUE
Adams	9/26/2022	Dow			3	2,500	Condo	TRUE	FALSE
Adams	1/10/2022	Dow			2	2,700	Condo	FALSE	TRUE
Adams	1/12/2022	DOWNTOWN	327,000	3	3	2,800	2 Story	TRUE	FALSE
Adams	9/22/2022	North	288,000	3	2	2,100	2 Story	FALSE	FALSE
Adams	4/15/2022	North	289,000	2	2	2,100	Condo	FALSE	FALSE
								FALSE	FALSE

Adding data to a table

If your table doesn't have a total row, the easiest way to enter data is simply to start typing in the row just below the table. When you enter something in a cell, Excel automatically expands the table and applies the formatting, formulas, and data validation to the new row. You can also paste a value in the next row. In fact, you could paste several rows' worth of data and the table will expand to accommodate.

If your table has a total row, you can't use that technique. In that case, you can insert rows into a table just as you would insert a row into any range. To insert a row, select a cell or the entire row and choose Home \Rightarrow Cells \Rightarrow Insert. When the selected range is inside a table, you'll see new entries on the Insert menu that deal with tables specifically. When you use these, the table is changed, but the data outside the table is unaffected.

When the selected cell is inside a table, the shortcut keys Ctrl- (minus sign) and Ctrl+ (plus sign) work on the table only and not on data outside the table. Moreover, as opposed to when you're not in a table, those shortcuts work on the whole table row or column regardless of whether you've selected the whole row or column.

4

Sorting and filtering table data

Each item in the header row of a table contains a drop-down arrow known as a *Filter Button*. When clicked, the Filter Button displays sorting and filtering options (see Figure 4.25).

TIP

If you don't plan to sort or filter the data in a table, you can turn off the display of Filter Buttons in a table's header row. Choose Table Design \Rightarrow Table Style Options \Rightarrow Filter Button to display or hide the drop-down arrows.

FIGURE 4.25

Each column in a table has sorting and filtering options.

A	B	C	D	E	F	G	H	I	J	
1	Agent	Date Listed	Area	List Price	Bedrooms	Baths	SqFt	Type	Pool	Sold
2	Adams	1/22/2022	Central					Condo	TRUE	TRUE
3	Adams	4/10/2022	Central					Condo	FALSE	FALSE
4	Adams	10/4/2022	Downtown					Split Level	FALSE	TRUE
5	Adams	1/19/2022	Downtown					Ranch	FALSE	TRUE
6	Adams	4/14/2022	Downtown					Split Level	FALSE	TRUE
7	Adams	3/17/2022	Downtown					Ranch	FALSE	TRUE
8	Adams	9/26/2022	Downtown					Condo	TRUE	FALSE
9	Adams	1/10/2022	Downtown					Condo	FALSE	TRUE
10	Adams	1/12/2022	Downtown					2 Story	TRUE	FALSE
11	Adams	9/22/2022	North					2 Story	FALSE	FALSE
12	Adams	4/15/2022	North					Condo	FALSE	FALSE
13	Adams	9/12/2022	North					Ranch	FALSE	FALSE
14	Adams	5/21/2022	South					Condo	FALSE	FALSE
15	Adams	4/16/2022	South					2 Story	FALSE	TRUE
16	Adams	11/22/2022	South					Ranch	TRUE	TRUE
17	Barnes	11/18/2022	Downtown					Ranch	TRUE	TRUE
18	Barnes	9/3/2022	Downtown					Split Level	TRUE	FALSE
19	Barnes	1/4/2022	Downtown					Ranch	FALSE	TRUE
20	Barnes	3/6/2022	Downtown					2 Story	FALSE	TRUE
21	Barnes	9/18/2022	Downtown					Ranch	TRUE	TRUE
22	Barnes	8/21/2022	North					Ranch	FALSE	FALSE
23	Barnes	10/12/2022	North					Condo	FALSE	TRUE
24	Barnes	4/11/2022	North	389,000	3	2	2,800	2 Story	FALSE	TRUE
25	Barnes	7/23/2022	South	305,000	4	2	2,200	Condo	TRUE	TRUE

Sorting a table

Sorting a table rearranges the rows based on the contents of a particular column. You may want to sort a table to put names in alphabetical order. Or maybe you want to sort your sales staff by the total sales made.

To sort a table by a particular column, click the Filter Button in the column header and choose one of the sort commands. The exact command varies, depending on the type of data in the column.

You can also select Sort by Color to sort the rows based on the background or text color of the data. This option is relevant only if you've overridden the table style colors with custom formatting.

You can sort on any number of columns. The trick is to sort the least significant column first and then proceed until the most significant column is sorted last. For example, in the real estate table, you may want to sort the list by agent. And within each agent's group, sort the rows by area. Then within each area, sort the rows by list price in descending order. For this type of sort, first sort by the Agent column, then sort by the Area column, and then sort by the List Price column. Figure 4.26 shows the table sorted in this manner.

FIGURE 4.26

A table after performing a three-column sort

	A	B	C	D	E	F	G	H	I	J
1	Agent	Date Listed	Area	List Price	Bedrooms	Baths	SqFt	Type	Pool	Sold
2	Adams	4/10/2022	Central	384,000	4	2	2,700	Condo	FALSE	FALSE
3	Adams	1/22/2022	Central	295,000	4	3.5	2,100	Condo	TRUE	TRUE
4	Adams	1/12/2022	Downtown	397,000	3	3	2,800	2 Story	TRUE	FALSE
5	Adams	1/10/2022	Downtown	372,000	2	2	2,700	Condo	FALSE	TRUE
6	Adams	9/26/2022	Downtown	352,000	2	3	2,500	Condo	TRUE	FALSE
7	Adams	3/17/2022	Downtown	312,000	4	3	2,200	Ranch	FALSE	TRUE
8	Adams	4/14/2022	Downtown	310,000	3	3.5	2,200	Split Level	FALSE	TRUE
9	Adams	1/19/2022	Downtown	271,000	2	3	1,900	Ranch	FALSE	TRUE
10	Adams	10/4/2022	Downtown	258,000	2	2.5	1,800	Split Level	FALSE	TRUE
11	Adams	9/12/2022	North	291,000	3	2	2,100	Ranch	FALSE	FALSE
12	Adams	4/15/2022	North	289,000	2	2	2,100	Condo	FALSE	FALSE
13	Adams	9/22/2022	North	288,000	3	2	2,100	2 Story	FALSE	FALSE
14	Adams	11/22/2022	South	389,000	3	2	2,800	Ranch	TRUE	TRUE
15	Adams	4/16/2022	South	388,000	4	3	2,800	2 Story	FALSE	TRUE
16	Adams	5/21/2022	South	282,000	2	3	2,000	Condo	FALSE	FALSE
17	Barnes	9/18/2022	Downtown	354,000	2	2	2,500	Ranch	TRUE	TRUE
18	Barnes	3/6/2022	Downtown	342,000	4	2	2,400	2 Story	FALSE	TRUE
19	Barnes	1/4/2022	Downtown	313,000	2	3	2,200	Ranch	FALSE	TRUE
20	Barnes	9/3/2022	Downtown	289,000	3	3	2,100	Split Level	TRUE	FALSE
21	Barnes	11/18/2022	Downtown	275,000	3	2	2,000	Ranch	TRUE	TRUE
22	Barnes	4/11/2022	North	389,000	3	2	2,800	2 Story	FALSE	TRUE
23	Barnes	10/12/2022	North	300,000	2	2.5	2,100	Condo	FALSE	TRUE
24	Barnes	8/21/2022	North	280,000	4	2	2,000	Ranch	FALSE	FALSE
25	Barnes	12/25/2022	South	390,000	2	3	2,800	Split Level	TRUE	FALSE
26	Barnes	4/24/2022	South	343,000	3	2.5	2,500	Split Level	TRUE	FALSE

Note

When a column is sorted, the Filter Button in the header row displays a different graphic to remind you that the table is sorted by that column.

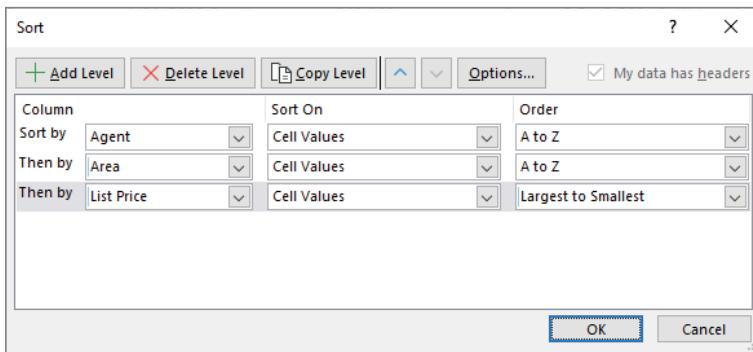
4

Another way of performing a multiple-column sort is to use the Sort dialog box (choose Home \Rightarrow Editing \Rightarrow Sort & Filter \Rightarrow Custom Sort). Or right-click any cell in the table and choose Sort \Rightarrow Custom Sort from the shortcut menu.

In the Sort dialog box, use the drop-down lists to specify the sort specifications. In this example, you start with Agent. Then click the Add Level button to insert another set of search controls. In this new set of controls, specify the sort specifications for the Area column. Then add another level and enter the specifications for the List Price column. Figure 4.27 shows the dialog box after entering the specifications for the three-column sort. This technique produces the same sort as described previously.

FIGURE 4.27

Using the Sort dialog box to specify a three-column sort



Filtering a table

Filtering a table refers to displaying only the rows that meet certain conditions. (The other rows are hidden.)

Note that entire worksheet rows are hidden. Therefore, if you have other data to the left or right of your table, that information may also be hidden when you filter the table. If you plan to filter your list, don't include any other data to the left or right of your table.

Using the real estate table, assume that you're interested only in the data for the Downtown area. Click the Filter Button in the Area row header and remove the check mark from Select All, which deselects everything. Then place a check mark next to Downtown and click OK. The table, shown in Figure 4.28, is now filtered to display only the listings in the Downtown area. Notice that some of the row numbers are missing. These rows are hidden and contain data that does not meet the specified criteria.

Also notice that the Filter Button in the Area column now shows a different graphic—an icon that indicates the column is filtered.

You can filter by multiple values in a column using multiple check marks. For example, to filter the table to show only Downtown and Central, place a check mark next to both values in the drop-down list in the Area row header.

You can filter a table using any number of columns. For example, you may want to see only the Downtown listings in which the Type is Condo. Just repeat the operation using the Type column. The table then displays only the rows in which the Area is Downtown and the Type is Condo.

For additional filtering options, select Text Filters (or Number Filters, if the column contains values). The options are fairly self-explanatory, and you have a great deal of flexibility in displaying only the rows in which you're interested. For example, you can display rows in which the List Price is greater than or equal to \$200,000 and less than \$300,000 (see Figure 4.29).

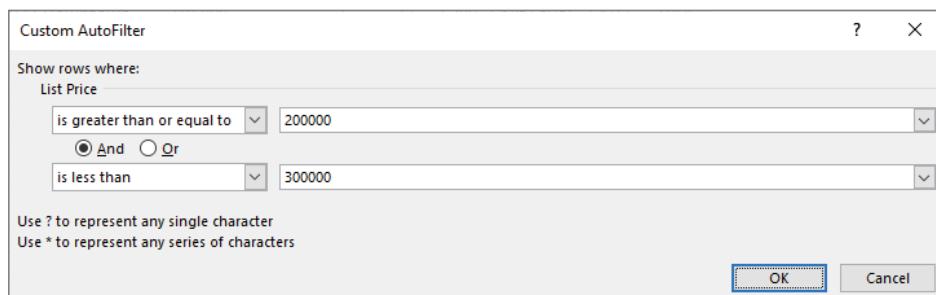
FIGURE 4.28

This table is filtered to show the information for only one area.

	A	B	C	D	E	F	G	H	I	J
1	Agent	Date Listed	Area	List Price	Bedrooms	Baths	SqFt	Type	Pool	Sold
4	Adams	1/12/2022	Downtown	397,000	3	3	2,800	2 Story	TRUE	FALSE
5	Adams	1/10/2022	Downtown	372,000	2	2	2,700	Condo	FALSE	TRUE
6	Adams	9/26/2022	Downtown	352,000	2	3	2,500	Condo	TRUE	FALSE
7	Adams	3/17/2022	Downtown	312,000	4	3	2,200	Ranch	FALSE	TRUE
8	Adams	4/14/2022	Downtown	310,000	3	3.5	2,200	Split Level	FALSE	TRUE
9	Adams	1/19/2022	Downtown	271,000	2	3	1,900	Ranch	FALSE	TRUE
10	Adams	10/4/2022	Downtown	258,000	2	2.5	1,800	Split Level	FALSE	TRUE
17	Barnes	9/18/2022	Downtown	354,000	2	2	2,500	Ranch	TRUE	TRUE
18	Barnes	3/6/2022	Downtown	342,000	4	2	2,400	2 Story	FALSE	TRUE
19	Barnes	1/4/2022	Downtown	313,000	2	3	2,200	Ranch	FALSE	TRUE
20	Barnes	9/3/2022	Downtown	289,000	3	3	2,100	Split Level	TRUE	FALSE
21	Barnes	11/18/2022	Downtown	275,000	3	2	2,000	Ranch	TRUE	TRUE
30	Bennet	3/19/2022	Downtown	385,000	4	3.5	2,800	2 Story	TRUE	FALSE
31	Bennet	5/15/2022	Downtown	379,000	2	3	2,700	Condo	TRUE	FALSE
32	Bennet	7/8/2022	Downtown	343,000	3	2	2,500	2 Story	FALSE	FALSE
37	Chung	8/6/2022	Downtown	336,000	3	2	2,400	Split Level	TRUE	TRUE
41	Daily	1/20/2022	Downtown	399,000	4	3	2,900	Ranch	FALSE	TRUE
42	Daily	12/5/2022	Downtown	281,000	2	3	2,000	2 Story	TRUE	FALSE
45	Hamilton	2/2/2022	Downtown	447,000	4	2	3,200	Condo	TRUE	FALSE
49	Jenkins	12/21/2022	Downtown	225,000	3	2.5	1,600	Ranch	TRUE	FALSE
53	Kelly	2/18/2022	Downtown	403,000	4	3	2,900	Split Level	FALSE	FALSE
54	Kelly	12/4/2022	Downtown	251,000	3	2.5	1,800	Split Level	TRUE	TRUE
57	Lang	2/2/2022	Downtown	413,000	4	3	3,000	2 Story	FALSE	FALSE

FIGURE 4.29

Specifying a more complex numeric filter



Also, you can right-click a cell and use the Filter command on the shortcut menu. This menu item leads to several additional filtering options that enable you to filter data based on the contents of the selected cell or by formatting.

Note

As you may expect, when you use filtering, the total row is updated to show the total only for the visible rows.

When you copy data from a filtered table, only the visible data is copied. In other words, rows that are hidden by filtering aren't copied. This filtering makes it easy to copy a subset of a larger table and paste it to another area of your worksheet. Keep in mind, though, that the pasted data is not a table—it's just a normal range. You can, however, convert the copied range to a table.

To remove filtering for a column, click the drop-down in the row header and select Clear Filter. If you've filtered using multiple columns, it may be faster to remove all filters by choosing Data \Rightarrow Sort & Filter \Rightarrow Clear.

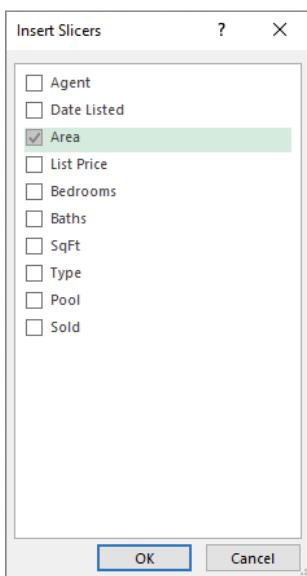
Filtering a table with slicers

Another way to filter a table is to use one or more *slicers*. This method is less flexible but more visually appealing. Slicers are particularly useful when the table will be viewed by novices or those who find the normal filtering techniques too complicated. Slicers are very visual, and it's easy to see exactly what type of filtering is in effect. A disadvantage of slicers is that they take up a lot of room on the screen.

To add one or more slicers, activate any cell in the table and choose Table Design \Rightarrow Tools \Rightarrow Insert Slicer. Excel responds with a dialog box that displays each header in the table (see Figure 4.30).

FIGURE 4.30

Use the Insert Slicers dialog box to specify which slicers to create.



Place a check mark next to the field(s) that you want to filter. You can create a slicer for each column, but that's rarely needed. In most cases, you'll want to be able to filter the table by only a few fields. Click OK, and Excel creates a slicer for each field you specified.

A slicer contains a button for every unique item in the field. In the real estate listing example, the slicer for the Agent field contains 14 buttons because the table has records for 14 different agents.

NOTE

Slicers may not be appropriate for columns that contain numeric data. For example, the real estate listing table has 78 different values in the List Price column. Therefore, a slicer for this column would have 78 buttons (and there's no way to group the values into numeric ranges). This is an example of how a slicer is not as flexible as normal filtering using Filter Buttons.

To use a slicer, just click one of the buttons. The table displays only the rows that have a value that corresponds to the button. You can also press Ctrl to select multiple buttons and press Shift to select a continuous group of buttons, which would be useful for selecting a range of List Price values. The slicer also has a multi-select button at the top. Click that button to toggle multi-select mode, and you won't have to hold down the Ctrl key.

If your table has more than one slicer, it's filtered by the selected buttons in each slicer. To remove filtering for a particular slicer, click the Clear Filter icon in the upper-right corner of the slicer.

Use the tools in the Slicer Ribbon to change the appearance or layout of a slicer. You have quite a bit of flexibility.

Figure 4.31 shows a table with two slicers. The table is filtered to show only the records for Adams, Barnes, and Chung in the Downtown area.

Changing the table's appearance

When you create a table, Excel applies the default table style. The actual appearance depends on which document theme is used in the workbook (see Chapter 5, "Formatting Worksheets"). If you prefer a different look, you can easily apply a different table style.

Select any cell in the table and choose Table Design \Rightarrow Table Styles. The Ribbon shows one row of styles, but if you click the More button at the bottom of the scrollbar to the right, Excel displays the Table Styles gallery, as shown in Figure 4.32. The styles are grouped into three categories: Light, Medium, and Dark. Notice that you get a "live" preview as you move your mouse among the styles. When you see one you like, just click to make it permanent. And yes, some are really ugly and practically illegible.

Part I: Getting Started with Excel

FIGURE 4.31

The table is filtered by two slicers.

1	A	B	C	D	E	F	G	H	I	J
4	Agent	Date Listed	Area	List Price	Bedrooms	Baths	SqFt	Type	Pool	Sold
5	Adams	1/12/2022	Downtown	397,000	3	3	2,800	2 Story	TRUE	FALSE
6	Adams	1/10/2022	Downtown	372,000	2	2	2,700	Condo	FALSE	TRUE
7	Adams	9/26/2022	Downtown	352,000	2	3	2,500	Condo	TRUE	FALSE
8	Adams	3/17/2022	Downtown	312,000	4	3	2,200	Ranch	FALSE	TRUE
9	Adams	4/14/2022	Downtown	310,000	3	3.5	2,200	Split Level	FALSE	TRUE
10	Adams	1/19/2022	Downtown	271,000	2	3	1,900	Ranch	FALSE	TRUE
17	Barnes	10/4/2022	Downtown							
18	Barnes	9/18/2022	Downtown							
19	Barnes	3/6/2022	Downtown							
20	Barnes	1/4/2022	Downtown							
21	Barnes	9/3/2022	Downtown							
37	Chung	11/18/2022	Downtown							
81										
82										
83										
84										
85										
86										
87										

Two filter dropdowns are open on the right side of the table:

- Agent**: Shows a list of agents: Adams, Barnes, Bennet, Chung, Daily, Hamilton, Jenkins, Kelly. Adams is selected.
- Area**: Shows a list of areas: Central, Downtown, North, South. Downtown is selected.

TIP

To change the default table style for the workbook, right-click the style in the Table Styles group and choose Set As Default from the shortcut menu. Subsequent tables that you create in that workbook will use that style.

For a different set of color choices, choose Page Layout \Rightarrow Themes \Rightarrow Themes to select a different document theme.



For more information about themes, see Chapter 5.

You can change some elements of the style by using the check box controls in the Table Design \Rightarrow Table Style Options group. These controls determine whether various elements of the table are displayed and whether some formatting options are in effect:

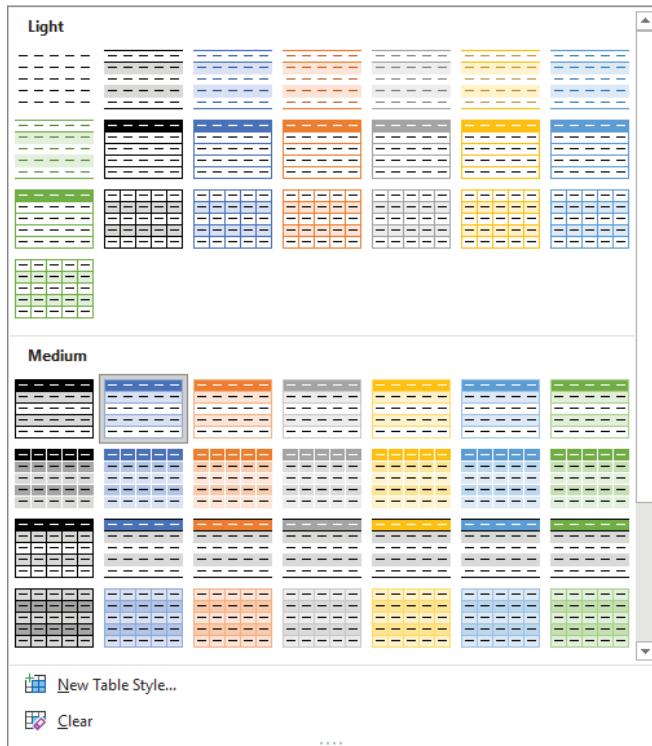
Header Row Toggles the display of the header row.

Total Row Toggles the display of the total row.

First Column Toggles special formatting for the first column. Depending on the table style used, this command might have no effect.

FIGURE 4.32

Excel offers many different table styles.



Last Column Toggles special formatting for the last column. Depending on the table style used, this command might have no effect.

Banded Rows Toggles the display of banded (alternating color) rows.

Banded Columns Toggles the display of banded columns.

Filter Button Toggles the display of the drop-down buttons in the table's header row.

TIP

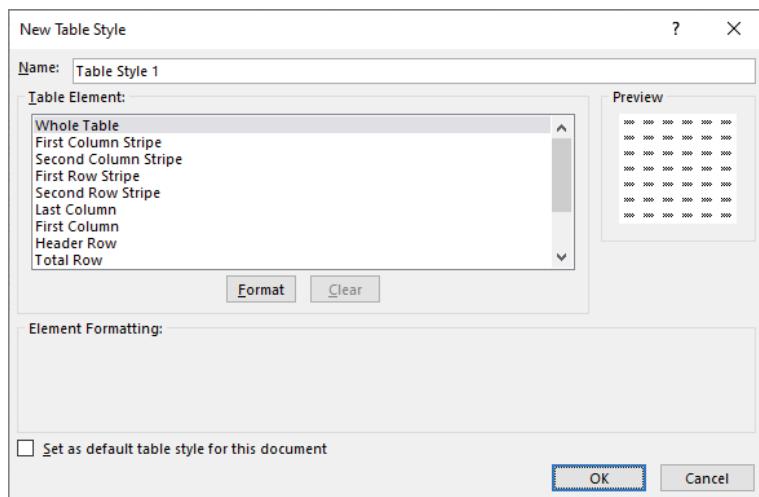
If applying table styles isn't working, it's probably because the range was already formatted before you converted it to a table. Table formatting doesn't override normal formatting. To clear existing background fill colors, select the entire table and choose Home ➔ Font ➔ Fill Color ➔ No Fill. To clear existing font colors, choose Home ➔ Font ➔ Font Color ➔ Automatic. To clear existing borders, choose Home ➔ Font ➔ Borders ➔ No Border. After you issue these commands, the table styles should work as expected.

If you'd like to create a custom table style, choose Table Design \Rightarrow Table Styles \Rightarrow New Table Style to display the New Table Style dialog box, shown in Figure 4.33. You can customize any or all of the 12 table elements. Select an element from the list, click Format, and specify the formatting for that element. When you're finished, give the new style a name and click OK. Your custom table style will appear in the Table Styles gallery in the Custom category.

Custom table styles are available only in the workbook in which they were created. However, if you copy a table that uses a custom style to a different workbook, the custom style will be available in the other workbook.

FIGURE 4.33

Use this dialog box to create a new table style.



TIP

If you want to make changes to an existing table style, locate it in the Ribbon and right-click and then choose Duplicate from the shortcut menu. Excel displays the Modify Table Style dialog box with all of the settings from the specified table style. Make your changes, give the style a new name, and click OK to save it as a custom table style.



Formatting Worksheets

IN THIS CHAPTER

Understanding how formatting can improve your worksheets

Getting to know the formatting tools

Using formatting in your worksheets

Using conditional formatting

Using named styles for easier formatting

Understanding document themes

Formatting your worksheet is more than just making your worksheet pretty. Proper formatting can help users understand the purpose of the worksheet and help prevent data entry errors.

Stylistic formatting isn't essential for every workbook that you develop, especially if it's only for your own use. On the other hand, it takes only a few moments to apply some simple formatting, and after you apply it, the formatting will remain in place without further effort on your part.

Getting to Know the Formatting Tools

Figure 5.1 shows how even simple formatting can significantly improve a worksheet's readability. The unformatted worksheet (on the left) is perfectly functional but not very readable compared to the formatted worksheet (on the right).

 This workbook is available on this book's website at www.wiley.com/go/excel365bible. The file is named loan_payments.xlsx.

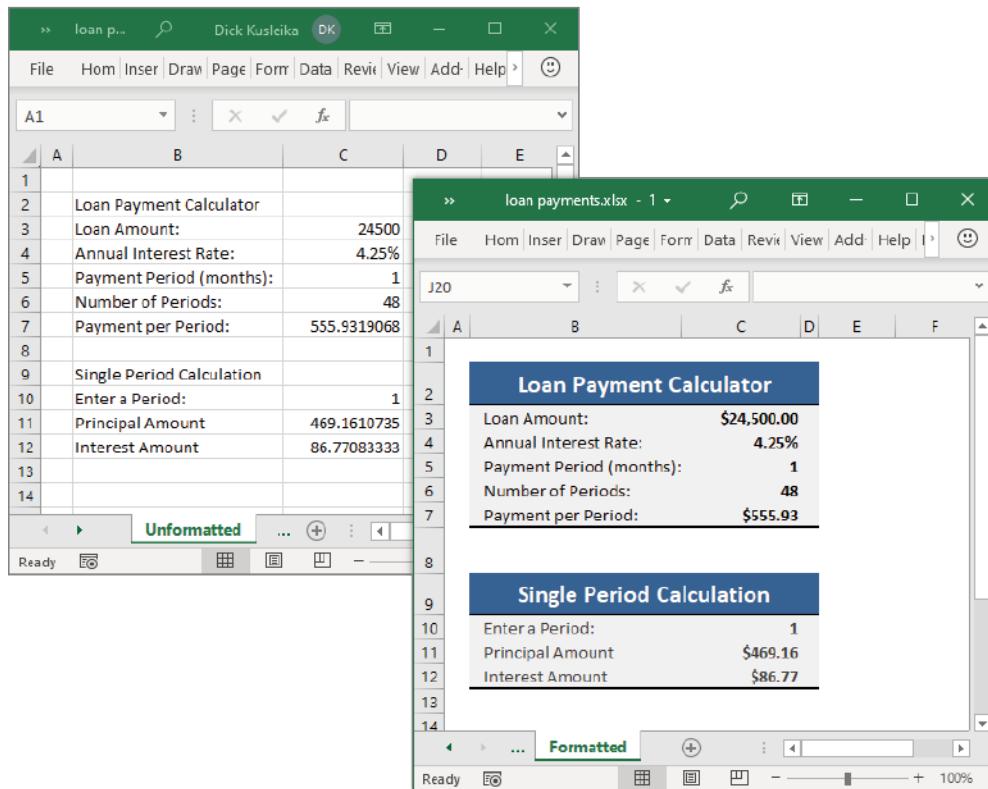
The Excel cell formatting tools are available in three locations.

- On the Home tab of the Ribbon
- On the Mini toolbar that appears when you right-click a selected range or a cell
- From the Format Cells dialog box

Part I: Getting Started with Excel

FIGURE 5.1

Simple formatting can greatly improve the appearance of your worksheet.



In addition, many common formatting commands have keyboard shortcuts.

Using the formatting tools on the Home tab

The Home tab of the Ribbon provides quick access to the most commonly used formatting options. Start by selecting the cell or range you want to format. Then use the appropriate tool in the Font, Alignment, or Number group.

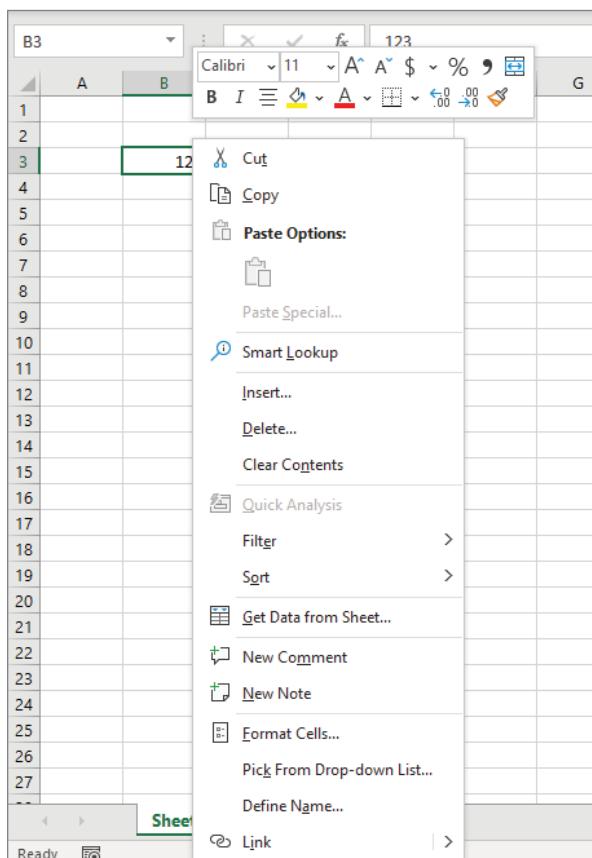
Using these tools is intuitive, and the best way to familiarize yourself with them is to experiment. Enter some data, select some cells, and then click the controls to change the appearance. Note that some of these controls are actually drop-down lists. Click the small arrow on the button, and the button expands to display your choices.

Using the Mini toolbar

When you right-click a cell or a range selection, you get a shortcut menu. In addition, the Mini toolbar appears above or below the shortcut menu. Figure 5.2 shows how this toolbar looks. The Mini toolbar for cell formatting contains the most commonly used controls from the Home tab of the Ribbon.

FIGURE 5.2

The Mini toolbar appears above or below the right-click shortcut menu.



If you use a tool on the Mini toolbar, the shortcut menu disappears, but the toolbar remains visible so that you can apply other formatting to the selected cells. To hide the Mini toolbar, just click in any cell or press Esc.

Using the Format Cells dialog box

The formatting controls available on the Home tab of the Ribbon are sufficient most of the time, but some types of formatting require that you use the Format Cells dialog box. This tabbed dialog box lets you apply nearly any type of stylistic formatting and number formatting. The formats that you choose in the Format Cells dialog box apply to the selected cells. Later sections in this chapter cover some of the tabs of the Format Cells dialog box.

After selecting the cell or range to format, you can display the Format Cells dialog box by using any of the following methods:

- Press **Ctrl+1**.
- Click the dialog box launcher in Home \Rightarrow Font, Home \Rightarrow Alignment, or Home \Rightarrow Number. (The dialog box launcher is the small downward-pointing arrow icon displayed to the right of the group name in the Ribbon.) When you display the Format Cells dialog box using a dialog box launcher, the dialog box is displayed with the appropriate tab visible.
- Right-click the selected cell or range and choose **Format Cells** from the shortcut menu.
- Click the **More** command in some of the drop-down controls in the Ribbon. For example, the Home \Rightarrow Font \Rightarrow Border drop-down includes an item named **More Borders**.

The Format Cells dialog box contains six tabs: Number, Alignment, Font, Border, Fill, and Protection. The following sections contain more information about the formatting options available in this dialog box.

Formatting Your Worksheet

Excel offers most of the same formatting options as other Office applications like Word or PowerPoint. As you might expect, cell-related formatting like fill color and borders feature more prominently in Excel than some of the other applications.

Using fonts to format your worksheet

You can use different fonts, font sizes, or text attributes in your worksheets to make various parts stand out, such as the headers for a table. You also can adjust the font size. For example, using a smaller font allows for more information on a single screen or printed page.

By default, Excel uses the 11-point (pt) Calibri font. A font is described by its typeface (Calibri, Cambria, Arial, Times New Roman, Courier New, and so on) as well as by its size, measured in points. (Seventy-two points equal one inch.) Excel's row height, by default, is 15 pt. Therefore, 11-pt type entered into 15-pt rows leaves a small amount of blank space between the characters in vertically adjacent rows.

TIP

If you haven't manually changed a row's height, Excel automatically adjusts the row height based on the tallest text that you enter into the row.

TIP

If you plan to distribute a workbook to other users, remember that Excel does not embed fonts. Therefore, you should stick with the standard fonts that are included with Windows or Microsoft Office. If you open a workbook and your system doesn't have the font used in the workbook, Windows attempts to use a similar font. Sometimes this attempt works; other times it doesn't.

Use the Font and Font Size tools in the Font group on the Home tab of the Ribbon (or on the Mini toolbar) to change the font or size for selected cells.

You also can use the Font tab in the Format Cells dialog box to choose fonts, as shown in Figure 5.3. This tab enables you to control several other font attributes that aren't available elsewhere. Besides choosing the font and font size, you can change the font style (bold, italic), underlining, color, and effects (strikethrough, superscript, or subscript). If you select the Normal Font check box, Excel displays the selections for the font defined for the Normal style. We discuss styles later in this chapter (see "Using Named Styles for Easier Formatting").

FIGURE 5.3

The Font tab of the Format Cells dialog box gives you many additional font attribute options.

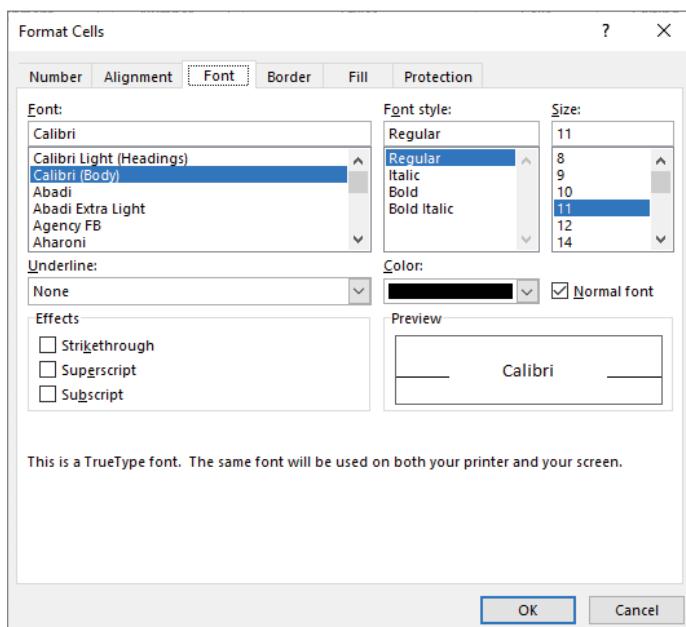


Figure 5.4 shows several examples of font formatting. In this figure, gridlines were turned off to make the underlining more visible. Notice, in the figure, that Excel provides four different underlining styles. In the two nonaccounting underline styles, only the cell contents are underlined. In the two accounting underline styles, the entire width of the cells is always underlined.

FIGURE 5.4

You can choose many different font formatting options for your worksheets.

	A	B
1	Large and Bold	Fonts come in many varieties
2	<i>Large and Italic</i>	Many varieties
3	<u>Single Underline</u>	Many varieties
4	<u>Double Underline</u>	FONTS COME IN MANY VARIETIES
5	<u>Single Accounting Underline</u>	Many varieties
6	<u>Double Accounting Underline</u>	<i>Fonts come in many varieties</i>
7	<i>Superscript x² and Subscript x₂</i>	Fonts come in many varieties
8	Strikethrough	Fonts come in many varieties
9		
10		

If you prefer to keep both hands on the keyboard, you can use the following shortcut keys to format a selected range quickly:

- Ctrl+B: Bold
- Ctrl+I: Italic
- Ctrl+U: Underline
- Ctrl+5: Strikethrough

These shortcut keys act as a toggle. For example, you can turn bold on and off by repeatedly pressing Ctrl+B.

Using Multiple Formatting Styles in a Single Cell

If a cell contains text (as opposed to a value or a formula), you can apply formatting to individual characters in the cell. To do so, switch to Edit mode (press F2, or double-click the cell) and then select the characters that you want to format. You can select characters either by dragging the mouse over them or by pressing the Shift key as you press the left or right arrow key.

This technique is useful if you need to apply superscript or subscript formatting to a few characters in the cell (refer to Figure 5.4 for examples).

After you select the characters to format, use any of the standard formatting techniques, including options in the Format Cells dialog box. To display the Format Cells dialog box when editing a cell, press Ctrl+1. The changes apply only to the selected characters in the cell. This technique doesn't work with cells that contain values or formulas.

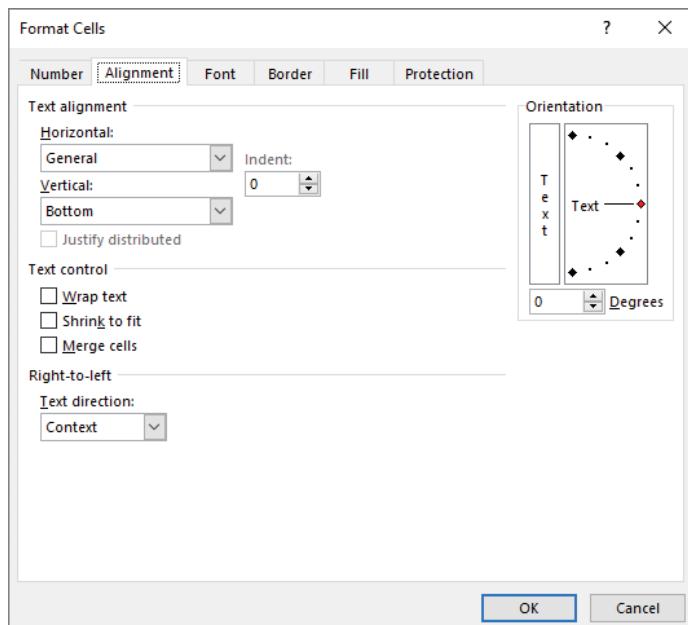
Changing text alignment

The contents of a cell can be aligned horizontally and vertically. By default, Excel aligns numbers to the right and text to the left. All cells use bottom alignment by default.

Overriding these defaults is a simple matter. The most commonly used alignment commands are in the Alignment group on the Home tab of the Ribbon. Use the Alignment tab of the Format Cells dialog box for even more options (see Figure 5.5).

FIGURE 5.5

The full range of alignment options is available on the Alignment tab of the Format Cells dialog box.



Choosing horizontal alignment options

Horizontal alignment options, which control how cell contents are distributed across the width of the cell (or cells), are available from the Format Cells dialog box.

General Aligns numbers to the right, aligns text to the left, and centers logical and error values. This option is the default horizontal alignment.

Left Aligns the cell contents to the left side of the cell. If the text is wider than the cell, the text spills over to the cell on the right. If the cell on the right isn't empty, the text is truncated and not completely visible. This is also available on the Ribbon.

Center Centers the cell contents in the cell. If the text is wider than the cell, the text spills over to cells on either side if they're empty. If the adjacent cells aren't empty, the text is truncated and not completely visible. This is also available on the Ribbon.

Right Aligns the cell contents to the right side of the cell. If the text is wider than the cell, the text spills over to the cell on the left. If the cell on the left isn't empty, the text is truncated and not completely visible. This is also available on the Ribbon.

Fill Repeats the contents of the cell until the cell's width is filled. If cells to the right also are formatted with Fill alignment, they also are filled.

Justify Justifies the text to the left and right of the cell. This option is applicable only if the cell is formatted as wrapped text and uses more than one line.

Center Across Selection Centers the text over the selected columns. This option is useful for centering a heading over several columns.

Distributed Distributes the text evenly across the cell, adding additional white space between words where necessary.

NOTE

If you choose Left, Right, or Distributed, you can also adjust the Indent setting, which adds horizontal space between the cell border and the text.

Figure 5.6 shows examples of text that uses three types of horizontal alignment: Left, Justify, and Distributed (with an indent).

FIGURE 5.6

The same text, displayed with three types of horizontal alignment

A	B	C	D	E	F
1					
2	Left	Justify	Distributed (with indent)		
3	Four score and seven years ago our fathers brought forth on this continent, a new nation, conceived in Liberty, and dedicated to the proposition that all men are created equal. Now we are engaged in a great civil war, testing whether that nation, or any nation so conceived and so dedicated, can long endure. We are met on a great battle-field of that war. We have come to dedicate a portion of that field, as a final resting place for those who here gave their lives that that nation might live. It is altogether fitting and proper that we should do this.	Four score and seven years ago our fathers brought forth on this continent, a new nation, conceived in Liberty, and dedicated to the proposition that all men are created equal.	Four score and seven years ago our fathers brought forth on this continent, a new nation, conceived in Liberty, and dedicated to the proposition that all men are created equal.	Four score and seven years ago our fathers brought forth on this continent, a new nation, conceived in Liberty, and dedicated to the proposition that all men are created equal.	Now we are engaged in a great civil war, testing whether that nation, or any nation so conceived and so dedicated, can long endure. We are met on a great battle-field of that war. We have come to dedicate a portion of that field, as a final resting place for those who here gave their lives that that nation might live. It is altogether fitting and proper that we should do this.
4					
5					



If you want to experiment with text alignment settings, this workbook is available at this book's website at www.wiley.com/go/excel365bible. The file is named `text_alignment.xlsx`.

Choosing vertical alignment options

Vertical alignment options typically aren't used as often as horizontal alignment options. In fact, these settings are useful only if you've adjusted row heights so that they're considerably taller than normal.

Here are the vertical alignment options available in the Format Cells dialog box:

Top Aligns the cell contents to the top of the cell. This is also available on the Ribbon.

Center Centers the cell contents vertically in the cell. This is also available on the Ribbon.

Bottom Aligns the cell contents to the bottom of the cell. This is also available on the Ribbon. This option is the default vertical alignment.

Justify Justifies the text vertically in the cell; this option is applicable only if the cell is formatted as wrapped text and uses more than one line. This setting can be used to increase the line spacing.

Distributed Distributes the text evenly between the top and bottom of the cell, adding additional white space between lines where necessary. If there is only one line of text in the cell, it's identical to the Center option.

Wrapping or shrinking text to fit the cell

If you have text too wide to fit the column width but you don't want that text to spill over into adjacent cells, you can use either the Wrap Text option or the Shrink to Fit option to accommodate that text. The Wrap Text option is also available on the Ribbon.

The Wrap Text option displays the text on multiple lines in the cell, if necessary. Use this option to display lengthy headings without having to make the columns too wide and without reducing the size of the text.

The Shrink to Fit option reduces the size of the text so that it fits into the cell without spilling over to the next cell. The times that this command is useful seem to be rare. Unless the text is just slightly too long, the result is almost always illegible.

Note

If you apply Wrap Text formatting to a cell, you can't use the Shrink to Fit formatting.

Merging worksheet cells to create additional text space

A handy formatting option is the ability to merge two or more cells. When you merge cells, you don't combine the contents of cells. Rather, you combine a group of cells into a single cell that occupies the same space. The worksheet shown in Figure 5.7 contains four sets of merged

cells. Range C2 : I2 has been merged into a single cell and so have ranges J2 : P2, B4 : B8, and B9 : B13. In the latter two cases, the text orientation has also been changed (see “Displaying text at an angle” later in this chapter).

FIGURE 5.7

Merge worksheet cells to make them act as if they were a single cell.

	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P	Q
1																	
2																	
3																	
4																	
5																	
6																	
7																	
8																	
9																	
10																	
11																	
12																	
13																	
14																	
15																	

You can merge any number of cells occupying any number of rows and columns. In fact, you can merge all 17 billion cells in a worksheet into a single cell—although there probably isn’t a good reason to do so, except maybe to play a trick on a co-worker.

The range that you intend to merge should be empty, except for the upper-left cell. If any of the other cells that you intend to merge are not empty, Excel displays a warning. If you continue, all the data (except in the upper-left cell) will be deleted.

You can use the Alignment tab of the Format Cells dialog box to merge cells, but using the Merge & Center control in the Alignment group on the Ribbon (or on the Mini toolbar) is simpler. To merge cells, select the cells that you want to merge and then click the Merge & Center button. The cells will be merged, and the content in the upper-left cells will be centered horizontally. The Merge & Center button acts as a toggle. To unmerge cells, select the merged cells and click the Merge & Center button again.

After you merge cells, you can change the alignment to something other than Center by using the controls in the Home \Rightarrow Alignment group.

The Home \Rightarrow Alignment \Rightarrow Merge & Center control contains a drop-down list with these additional options:

Merge Across When a multirow range is selected, this command creates multiple merged cells—one for each row.

Merge Cells Merges the selected cells without applying the Center attribute.

Unmerge Cells Unmerges the selected cells.

Displaying text at an angle

In some cases, you may want to create more visual impact by displaying text at an angle within a cell. You can display text horizontally, vertically, or at any angle between 90 degrees up and 90 degrees down.

From the Home \Rightarrow Alignment \Rightarrow Orientation drop-down list, you can apply the most common text angles. For more control, use the Alignment tab of the Format Cells dialog box. In the Format Cells dialog box (refer to Figure 5.5), use the Degrees spinner control—or just drag the red pointer in the gauge. You can specify a text angle between -90 and +90 degrees.

Figure 5.8 shows an example of text displayed at a 45-degree angle.

FIGURE 5.8

Rotate text for additional visual impact.

NOTE

Rotated text may look a bit distorted onscreen, but the printed output is usually of much better quality.

Using colors and shading

Excel provides the tools to create some colorful worksheets. You can change the color of the text or add colors to the backgrounds of the worksheet cells. Prior to Excel 2007, workbooks were limited to a palette of 56 colors. Since then, Microsoft has increased the number of colors to more than 16 million.

You control the color of the cell's text by choosing Home \Rightarrow Font \Rightarrow Font Color. Control the cell's background color by choosing Home \Rightarrow Font \Rightarrow Fill Color. Both of these color controls are also available on the Mini toolbar, which appears when you right-click a cell or range.

Tip

To hide the contents of a cell, make the background color the same as the font text color. The cell contents are still visible in the Formula bar when you select the cell. Keep in mind, however, that some printers may override this setting, and the text may be visible when printed.

Even though you have access to a lot of colors, you might want to stick with the 10 theme colors (and their light/dark variations) displayed in the various color selection controls. In other words, avoid using the More Colors option, which lets you select a color. Why? First, those 10 colors were chosen because they “go together.” (Well, at least somebody thought they did.) Another reason involves document themes. If you switch to a different document theme for your workbook, nontheme colors aren’t changed. In some cases, the result may be less than pleasing aesthetically. (See “Understanding Document Themes” later in this chapter for more information about themes.)

Adding borders and lines

Borders (and lines within the borders) are another visual enhancement that you can add around groups of cells. Borders are often used to group a range of similar cells or to delineate rows or columns. Excel offers 13 preset styles of borders, as you can see in the Home \Rightarrow Font \Rightarrow Borders drop-down list shown in Figure 5.9. This control works with the selected cell or range and enables you to specify which, if any, border style to use for each border of the selection.

You may prefer to draw borders rather than select a preset border style. To do so, use the Draw Border or Draw Border Grid command from the Home \Rightarrow Font \Rightarrow Borders drop-down list. Selecting either command lets you create borders by dragging your mouse. Use the Line Color or Line Style command to change the color or style. When you’re finished drawing borders, press Esc to cancel the border-drawing mode.

Another way to apply borders is to use the Border tab of the Format Cells dialog box, which is shown in Figure 5.10. One way to display this dialog box is to select More Borders from the Borders drop-down list.

Before you display the Format Cells dialog box, select the cell or range to which you want to add borders. Then, in the Format Cells dialog box, choose a line style and color and then choose the border position for the line style by clicking one or more of the Border icons. (These icons are toggles.)

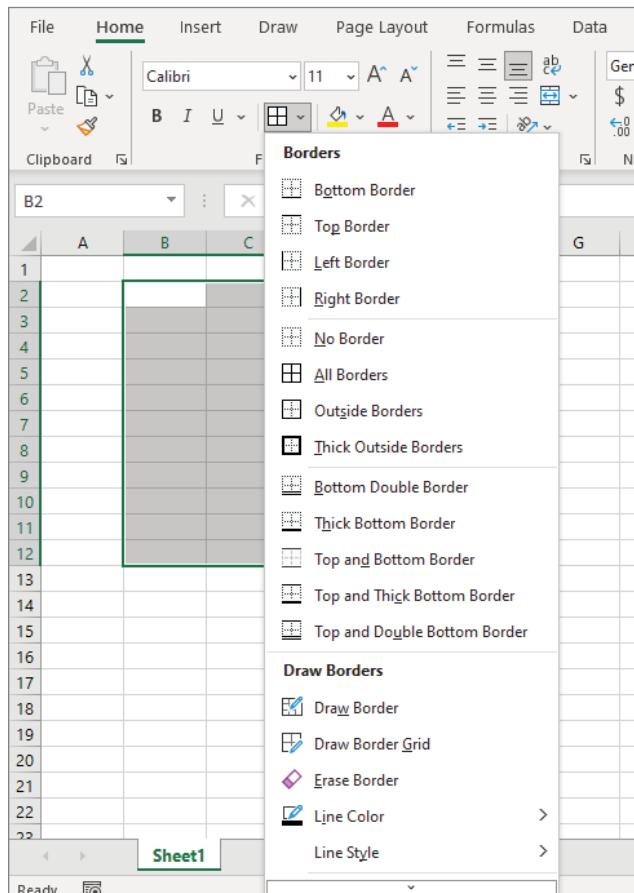
Notice that the Border tab has three preset icons, which can save you some clicking. If you want to remove all borders from the selection, click None. To put an outline around the selection, click Outline. To put borders inside the selection, click Inside.

Excel displays the selected border style in the dialog box; there is no live preview in the worksheet. You can choose different styles for different border positions; you can also choose a color for the border. Using this dialog box may require some experimentation, but you’ll get the hang of it.

When you apply two diagonal lines, the cells look like they’ve been crossed out.

FIGURE 5.9

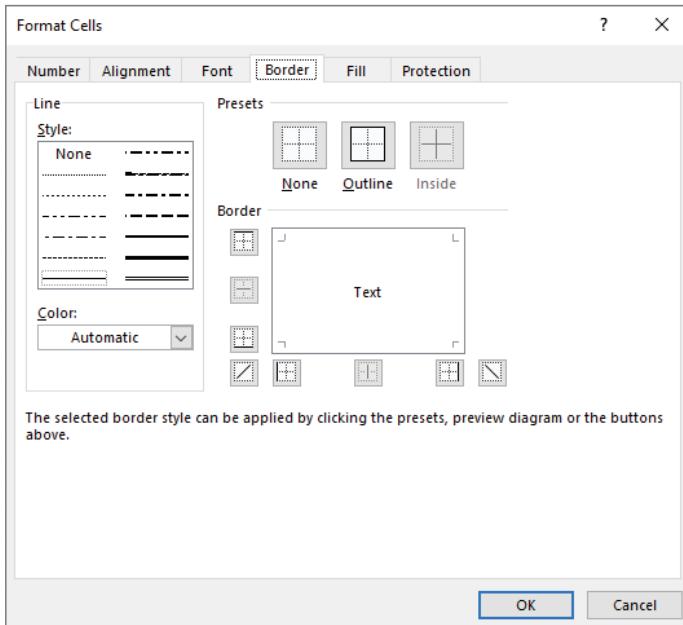
Use the Borders drop-down list to add lines around worksheet cells.

**TIP**

If you use border formatting in your worksheet, you may want to turn off the grid display to make the borders more pronounced. Choose View \Rightarrow Show \Rightarrow Gridlines to toggle the gridline display.

FIGURE 5.10

Use the Border tab of the Format Cells dialog box for more control over cell borders.



Copying Formats by Painting

Perhaps the quickest way to copy the formats from one cell to another cell or range is to use the Format Painter button (the button with the paintbrush image) of the Home \Rightarrow Clipboard group.

1. Select the cell or range that has the formatting attributes that you want to copy.
2. Click the **Format Painter** button. The mouse pointer changes to include a paintbrush.
3. Select the cells to which you want to apply the formats.
4. Release the mouse button, and Excel applies the same set of formatting options that were in the original range.

If you double-click the Format Painter button, you can paint multiple areas of the worksheet with the same formats. Excel applies the formats that you copy to each cell or range that you select. To get out of Paint mode, click the Format Painter button again (or press Esc).

Using Conditional Formatting

You can apply *conditional formatting* to a cell so that the cell looks different depending on its contents. Conditional formatting is a useful tool for visualizing numeric data. In some cases, conditional formatting may be a viable alternative to creating a chart.

Conditional formatting lets you apply cell formatting selectively and automatically, based on the contents of the cells. For example, you can apply conditional formatting in such a way that all negative values in a range have a light-yellow background color. When you enter or change a value in the range, Excel examines the value and checks the conditional formatting rules for the cell. If the value is negative, the background is shaded; otherwise, no formatting is applied.

Specifying conditional formatting

To apply a conditional formatting rule to a cell or range, select the cells and then use one of the commands from the Home \Rightarrow Styles \Rightarrow Conditional Formatting drop-down list to specify a rule. The choices are as follows:

Highlight Cells Rules Examples include highlighting cells that are greater than a particular value, are between two values, contain a specific text string, contain a date, or are duplicated.

Top/Bottom Rules Examples include highlighting the top 10 items, the items in the bottom 20 percent, and the items that are above average.

Data Bars Applies graphic bars directly in the cells, proportional to the cell's value.

Color Scales Applies background color, proportional to the cell's value.

Icon Sets Displays icons directly in the cells. The icons depend on the cell's value.

New Rule Enables you to specify other conditional formatting rules, including rules based on a logical formula.

Clear Rules Deletes all the conditional formatting rules from the selected cells.

Manage Rules Displays the Conditional Formatting Rules Manager dialog box in which you create new conditional formatting rules, edit rules, or delete rules.

Using graphical conditional formats

The following sections describe the three conditional formatting options that display graphics: data bars, color scales, and icon sets. These types of conditional formatting can be useful for visualizing the values in a range.

Using data bars

The data bars conditional format displays horizontal bars directly in the cell. The length of the bar is based on the value of the cell relative to the other values in the range.

Part I: Getting Started with Excel

Figure 5.11 shows an example of data bars. It's a list of tracks on 39 Bob Dylan albums, with the length of each track in column D. With data bar conditional formatting applied to column D, you can tell at a glance which tracks are longer.

FIGURE 5.11

The length of the data bars is proportional to the track length in the cell in column D.

A	B	C	D	E	F
1	Artist	Album	Title	Length	
2	Bob Dylan	Infidels	Jokerman	0:06:19	
3	Bob Dylan	Modern Times	When The Deal Goes Down	0:05:04	
4	Bob Dylan	Infidels	Don't Fall Apart On Me Tonight	0:05:56	
5	Bob Dylan	Blonde On Blonde	Leopard-Skin Pill-Box Hat	0:04:00	
6	Bob Dylan	Under the Red Sky	2 X 2	0:03:36	
7	Bob Dylan	Highway 61 Revisited	Just Like Tom Thumb's Blues	0:05:32	
8	Bob Dylan	Self Portrait	She Belongs to Me	0:02:43	
9	Bob Dylan	New Morning	New Morning	0:03:59	
10	Bob Dylan	Planet Waves	Forever young	0:04:57	
11	Bob Dylan	30th Anniversary Concert	Emotionally Yours	0:05:43	
12	Bob Dylan	Good As I Been to You	Canadee-I-O	0:04:23	
13	Bob Dylan	Down in the Groove	When Did You Leave Heaven	0:02:13	
14	Bob Dylan	Street Legal	We Better Talk This Over	0:04:04	
15	Bob Dylan	The Times They Are A-Changin'	North Country Blues	0:04:33	
16	Bob Dylan	30th Anniversary Concert	Foot Of Pride	0:08:47	
17	Bob Dylan	Time Out of Mind	Til I Fell in Love with You	0:05:17	
18	Bob Dylan	Knocked Out Loaded	Under Your Spell	0:03:56	
19	Bob Dylan	Another Side of Bob Dylan	To Ramona	0:03:52	
20	Bob Dylan	Time Out of Mind	Million Miles	0:05:52	
21	Bob Dylan	Another Side of Bob Dylan	I Don't Believe You	0:04:22	



The examples in this section are available on this book's website at www.wiley.com/go/excel365bible. The workbook is named `data bars examples.xlsx`.

TIP

When you adjust the column width, the bar lengths adjust accordingly. The differences among the bar lengths are more prominent when the column is wider.

Excel provides quick access to 12 data bar styles via Home \Rightarrow Styles \Rightarrow Conditional Formatting \Rightarrow Data Bars. For additional choices, click the More Rules option, which displays the New Formatting Rule dialog box. Use this dialog box to do the following:

- Show the bar only. (Hide the numbers.)
- Specify Minimum and Maximum values for the scaling.
- Change the appearance of the bars.
- Specify how negative values and the axis are handled.
- Specify the direction of the bars.

Note

Oddly, if you add data bars using one of the 12 data bar styles, the colors used for data bars are not theme colors. If you apply a new document theme, the data bar colors do not change. However, if you add the data bars by using the New Formatting Rule dialog box, the colors you choose are theme colors.

Using color scales

The color scale conditional formatting option varies the background color of a cell based on the cell's value relative to other cells in the range.

Figure 5.12 shows examples of color scale conditional formatting. The example on the left depicts monthly sales for three regions. Conditional formatting was applied to the range B4 : D15. The conditional formatting uses a three-color scale, with red for the lowest value, yellow for the midpoint, and green for the highest value. Values in between are displayed using a color within the gradient. It's clear that the Central region consistently has lower sales volumes, but the conditional formatting doesn't help identify monthly differences for a particular region.

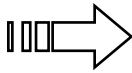
FIGURE 5.12

Two examples of color scale conditional formatting

	A	B	C	D	E	F	G	H	I	
1	A single conditional formatting rule					A separate rule for each region				
2										
3	Month	Western	Central	Eastern		Month	Western	Central	Eastern	
4	January	214,030	103,832	225,732		January	214,030	103,832	225,732	
5	February	204,476	105,777	239,316		February	204,476	105,777	239,316	
6	March	195,555	102,221	230,848		March	195,555	102,221	230,848	
7	April	191,417	102,222	233,342		April	191,417	102,222	233,342	
8	May	184,896	100,034	221,743		May	184,896	100,034	221,743	
9	June	181,829	94,401	212,820		June	181,829	94,401	212,820	
10	July	182,288	90,825	208,116		July	182,288	90,825	208,116	
11	August	187,710	93,860	221,596		August	187,710	93,860	221,596	
12	September	193,884	92,300	228,499		September	193,884	92,300	228,499	
13	October	184,089	88,193	214,495		October	184,089	88,193	214,495	
14	November	186,954	92,654	227,823		November	186,954	92,654	227,823	
15	December	185,778	95,806	237,911		December	185,778	95,806	237,911	
16										

The example on the right shows the same data, but conditional formatting was applied to each region separately. This approach facilitates comparisons within a region and can identify high or low sales months.

Neither one of these approaches is necessarily better. The way you set up conditional formatting depends entirely on what you're trying to visualize.



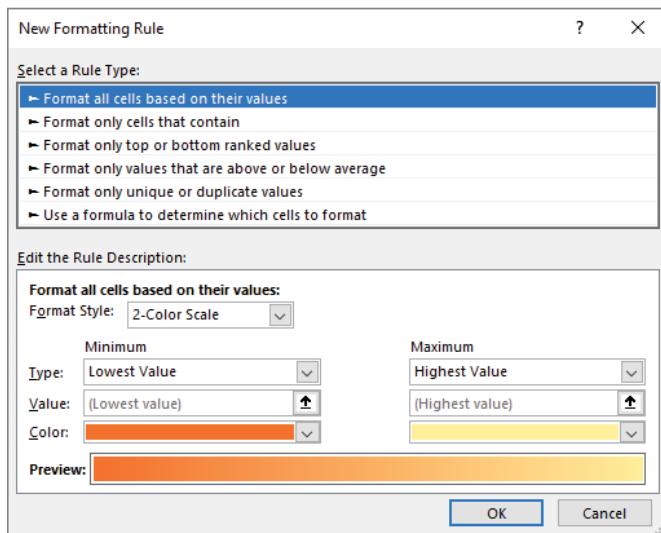
This workbook, named `color scale example.xlsx`, is available on this book's website at www.wiley.com/go/excel365bible.

Excel provides six two-color scale presets and six three-color scale presets, which you can apply to the selected range by choosing Home \Rightarrow Styles \Rightarrow Conditional Formatting \Rightarrow Color Scales.

To customize the colors and other options, choose Home \Rightarrow Styles \Rightarrow Conditional Formatting \Rightarrow Color Scales \Rightarrow More Rules, and the New Formatting Rule dialog box, shown in Figure 5.13, appears. Adjust the settings and watch the Preview box to see the effects of your changes.

FIGURE 5.13

Use the New Formatting Rule dialog box to customize a color scale.



Using icon sets

Yet another conditional formatting option is to display an icon in the cell. The icon displayed depends on the value of the cell.

To assign an icon set to a range, select the cells and choose Home \Rightarrow Styles \Rightarrow Conditional Formatting \Rightarrow Icon Sets. Excel provides 20 icon sets from which you may choose. The number of icons in the sets ranges from three to five. You can't create a custom icon set.

Figure 5.14 shows an example that uses an icon set. The symbols graphically depict the status of each project, based on the value in column C.

FIGURE 5.14

Using an icon set to indicate the status of projects

	A	B	C	D
1				
2				
Project Status Report				
3	Project	Pct Completed		
4	Project 1	!	95%	
5	Project 2	✓	100%	
6	Project 3	!	50%	
7	Project 4	✗	0%	
8	Project 5	!	20%	
9	Project 6	!	80%	
10	Project 7	✓	100%	
11	Project 8	✗	0%	
12	Project 9	✗	0%	
13	Project 10	!	50%	
14				



The icon set example in this section is available on this book's website at www.wiley.com/go/excel365bible. The workbook is named `icon set examples.xlsx`.

By default, the symbols are assigned using percentiles. For a three-symbol set, the items are grouped into three percentiles. For a four-symbol set, they're grouped into four percentiles. For a five-symbol set, the items are grouped into five percentiles.

If you would like more control over how the icons are assigned, choose Home \Rightarrow Styles \Rightarrow Conditional Formatting \Rightarrow Icon Sets \Rightarrow More Rules to display the New Formatting Rule dialog box. To modify an existing rule, choose Home \Rightarrow Styles \Rightarrow Conditional Formatting \Rightarrow Manage Rules. Then select the rule to modify and click the Edit Rule button.

Figure 5.15 shows how to modify the icon set rules such that only projects that are 100 percent complete get the check mark icons. Projects that are 0 percent complete get the X icon. All other projects get no icon.

Figure 5.16 shows the project status list after making this change.

Creating formula-based rules

The graphical conditional formats are generally used to show a cell in relation to other, nearby cells. Formula-based rules generally apply to one cell independently. The same rule may apply to many cells, but each cell is considered on its own.

The Highlight Cells Rules and Top/Bottom Rules options under the Conditional Formatting Ribbon control are commonly used shortcuts for formula-based rules. If you choose Home \Rightarrow Styles \Rightarrow Conditional Formatting \Rightarrow New Rule, Excel displays the New Formatting Rule dialog box. You saw this dialog box in the previous section when the built-in graphical conditional formats needed tweaking. The entry Format Only Cells that Contain is another shortcut for a formula-based rule.

Part I: Getting Started with Excel

FIGURE 5.15

Changing the icon assignment rule

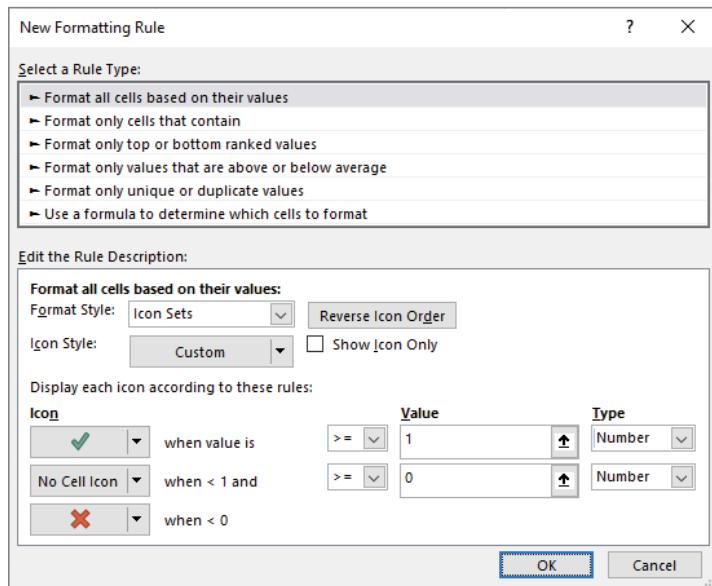


FIGURE 5.16

Using a modified rule and eliminating an icon makes the table more readable.

E	F	G
Project Status Report		
Project	Pct Completed	
Project 1		95%
Project 2	✓	100%
Project 3		50%
Project 4	✗	0%
Project 5		20%
Project 6		80%
Project 7	✓	100%
Project 8	✗	0%
Project 9	✗	0%
Project 10		50%

The last entry in the New Formatting Rule dialog box is Use a Formula to Determine Which Cells to Format. This is the entry you choose if none of the other shortcuts do what you want. It provides maximum flexibility for creating a rule.

Note

The formula must be a logical one that returns either TRUE or FALSE. If the formula evaluates to TRUE, the condition is satisfied and the conditional formatting is applied. If the formula evaluates to FALSE, the conditional formatting is not applied.



A workbook with examples from this section is available on this book's website at www.wiley.com/go/excel365bible. The file is named conditional formatting formulas.xlsx.

Understanding relative and absolute references

If the formula that you enter into the New Formatting Rule or Edit Formatting Rule dialog box contains a cell reference, that reference is considered a relative reference based on the upper-left cell in the selected range.

For example, suppose that you want to set up a conditional formatting condition that applies shading to cells in range A1:B10 only if the cell contains text. None of Excel's conditional formatting options can do this task, so you need to create a formula that will return TRUE if the cell contains text and FALSE otherwise. Follow these steps:

1. **Select the range A1:B10, and make sure that cell A1 is the active cell.**
2. **Choose Home ➔ Styles ➔ Conditional Formatting ➔ New Rule.** The New Formatting Rule dialog box appears.
3. **Click the Use a Formula to Determine Which Cells to Format rule type.**
4. **Enter the following formula into the Formula box:**
`=ISTEXT(A1)`
5. **Click the Format button.** The Format Cells dialog box appears.
6. **From the Fill tab, specify the cell shading that will be applied if the formula returns TRUE.**
7. **Click OK to return to the New Formatting Rule dialog box (see Figure 5.17).**
8. **Click OK to close the New Formatting Rule dialog box.**

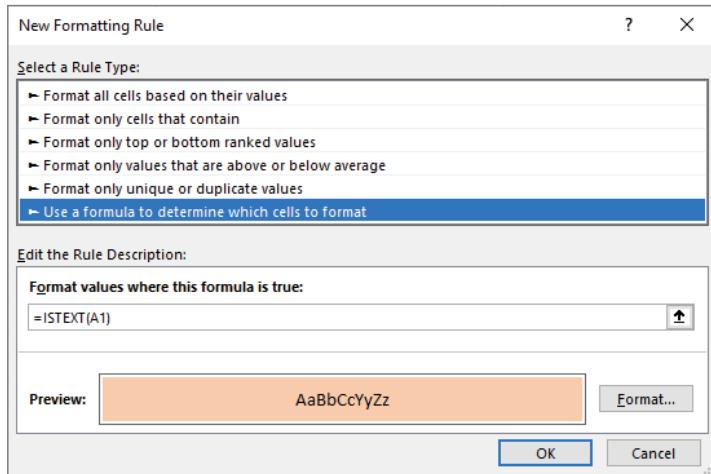
Notice that the formula entered in step 4 contains a relative reference to the upper-left cell in the selected range.

Generally, when entering a conditional formatting formula for a range of cells, you'll use a reference to the active cell, which is typically the upper-left cell in the selected range. One exception is when you need to refer to a specific cell. For example, suppose that you select range A1:B10 and you want to apply formatting to all cells in the range that exceed the value in cell C1. Enter this conditional formatting formula:

`=A1>C1`

FIGURE 5.17

Creating a conditional formatting rule based on a formula



In this case, the reference to cell C1 is an absolute reference; it will not be adjusted for the cells in the selected range. In other words, the conditional formatting formula for cell A2 looks like this:

=A2>\$C\$1

The relative cell reference is adjusted, but the absolute cell reference is not.

Conditional formatting formula examples

Each of these examples uses a formula entered directly into the New Formatting Rule dialog box after the Use a Formula to Determine Which Cells to Format rule type is selected. You decide the type of formatting that you apply conditionally.

Identifying weekend days

Excel provides a number of conditional formatting rules that deal with dates, but it doesn't let you identify dates that fall on a weekend. Use this formula to identify weekend dates:

=OR(WEEKDAY(A1)=7, WEEKDAY(A1)=1)

This formula assumes that a range is selected and that cell A1 is the active cell.

Highlighting a row based on a value

Figure 5.18 shows part of a worksheet that contains a conditional formula in the range A3 : G28. If a name entered in cell B1 is found in the first column, the entire row for that name is highlighted.

FIGURE 5.18

Highlighting a row, based on a matching name

	A	B	C	D	E	F	G	H
1	Name:	Noel						
2								
3	Alice	7	118	61	55	85	26	
4	Bob	198	134	180	3	132	63	
5	Carl	2	46	59	63	59	26	
6	Denise	190	121	12	26	68	97	
7	Elvin	174	42	176	68	124	14	
8	Francis	129	114	83	103	129	129	
9	George	9	128	24	44	139	108	
10	Harald	168	183	200	167	134	83	
11	Ivan	165	141	95	91	100	144	
12	June	116	171	109	84	148	15	
13	Kathy	131	43	197	82	103	163	
14	Larry	139	30	171	122	34	196	
15	Mary	31	171	185	162	171	17	
16	Noel	78	126	190	78	123	2	
17	Oliver	157	98	100	75	137	10	
18	Patrick	120	144	106	39	39	119	
19	Quincey	156	200	58	74	37	76	
20	Raul	58	147	160	182	11	79	
21	Shiela	79	183	5	161	104	23	

The conditional formatting formula is:

`=A3=B1`

Notice that a mixed reference is used for cell A3. Because the column part of the reference is absolute, the comparison is always done using the contents of column A.

Displaying alternate-row shading

The conditional formatting formula that follows was applied to the range A1:D18, as shown in Figure 5.19, to apply shading to alternate rows:

`=MOD(ROW(), 2)=0`

Alternate row shading can make your spreadsheets easier to read. If you add or delete rows within the conditional formatting area, the shading is updated automatically.

This formula uses the ROW function (which returns the row number) and the MOD function (which returns the remainder of its first argument divided by its second argument). For cells in even-numbered rows, the MOD function returns 0, and cells in that row are formatted.

For alternate shading of columns, use the COLUMN function instead of the ROW function.

FIGURE 5.19

Using conditional formatting to apply formatting to alternate rows

	A	B	C	D	E
1	433	516	157	26	
2	291	31	362	168	
3	625	265	619	551	
4	548	900	569	375	
5	178	279	270	369	
6	782	250	346	139	
7	741	15	602	472	
8	136	840	898	655	
9	421	953	790	94	
10	708	151	437	725	
11	45	598	179	133	
12	128	481	238	325	
13	835	623	367	32	
14	400	881	779	113	
15	138	153	321	886	
16	303	84	990	503	
17	801	745	840	999	
18	809	318	679	246	
19					

Creating checkerboard shading

The following formula is a variation on the example in the preceding section. It applies formatting to alternate rows and columns, creating a checkerboard effect.

=MOD(ROW(), 2) =MOD(COLUMN(), 2)

Shading groups of rows

Here's another row shading variation. The following formula shades alternate groups of rows. It produces four shaded rows, followed by four unshaded rows, followed by four more shaded rows, and so on.

=MOD(INT((ROW() - 1) / 4) + 1, 2) = 1

Figure 5.20 shows an example.

For different sized groups, change the 4 to some other value. For example, use this formula to shade alternate groups of two rows:

=MOD(INT((ROW() - 1) / 2) + 1, 2) = 1

Working with conditional formats

The following sections describe some additional information about conditional formatting that you may find useful.

FIGURE 5.20

Conditional formatting produces these groups of alternating shaded rows.

	A	B	C	D	E
1	301	458	76	437	
2	70	197	261	454	
3	292	289	282	48	
4	412	88	205	426	
5	269	402	457	34	
6	215	213	86	228	
7	306	424	67	328	
8	236	7	390	356	
9	372	354	188	450	
10	132	400	109	338	
11	364	19	250	324	
12	79	195	415	373	
13	75	100	42	296	
14	377	441	106	474	
15	83	164	154	220	
16	3	1	214	68	
17	251	81	432	415	
18	180	114	387	481	
19	35	46	316	99	
20	45	155	443	454	
21	424	200	161	378	

Managing rules

The Conditional Formatting Rules Manager dialog box is useful for checking, editing, deleting, and adding conditional formats. First select any cell in the range that contains conditional formatting. Then choose Home \Rightarrow Styles \Rightarrow Conditional Formatting \Rightarrow Manage Rules.

You can specify as many rules as you like by clicking the New Rule button. Cells can even use data bars, color scales, and icon sets at the same time.

Copying cells that contain conditional formatting

Conditional formatting information is stored with a cell much like standard formatting information is stored with a cell. As a result, when you copy a cell that contains conditional formatting, you also copy the conditional formatting.

TIP

To copy only the formatting (including conditional formatting), copy the cells and then use the Paste Special dialog box and select the Formats option, or choose Home \Rightarrow Clipboard \Rightarrow Paste \Rightarrow Formatting (R).

If you insert rows or columns within a range that contains conditional formatting, the new cells have the same conditional formatting.

Deleting conditional formatting

When you press Delete to delete the contents of a cell, you don't delete the conditional formatting for the cell (if any). To remove all conditional formats (as well as all other cell formatting), select the cell and then choose Home \Rightarrow Editing \Rightarrow Clear \Rightarrow Clear Formats. Or, choose Home \Rightarrow Editing \Rightarrow Clear \Rightarrow Clear All to delete the cell contents and the conditional formatting.

To remove only conditional formatting (and leave the other formatting intact), choose Home \Rightarrow Styles \Rightarrow Conditional Formatting \Rightarrow Clear Rules and choose one of the available options.

Locating cells that contain conditional formatting

You can't always tell, just by looking at a cell, whether it contains conditional formatting. You can, however, use the Go To Special dialog box to select such cells.

1. Choose Home \Rightarrow Editing \Rightarrow Find & Select \Rightarrow Go To Special. The Go To Special dialog box appears.
2. In the Go To Special dialog box, select the Conditional Formats option.
3. To select all cells on the worksheet containing conditional formatting, select the All option. To select only the cells that contain the same conditional formatting as the active cell, select the Same option.
4. Click OK. Excel selects the cells for you.

Note

The Excel Find and Replace dialog box includes a feature that allows you to search your worksheet to locate cells that contain specific formatting. This feature does not locate cells that contain formatting resulting from conditional formatting.

Using Named Styles for Easier Formatting

One of the most underutilized features in Excel is named styles. *Named styles* make it easy to apply a set of predefined formatting options to a cell or range. In addition to saving time, using named styles helps to ensure a consistent look.

A style can consist of settings for up to six attributes:

- Number format
- Alignment (vertical and horizontal)
- Font (type, size, and color)
- Borders

- Fill
- Cell protection (locked and hidden)

The real power of styles is apparent when you change a component of a style. All cells that use that named style automatically incorporate the change. Suppose that you apply a particular style to a dozen cells scattered throughout your worksheet. Later, you realize that these cells should have a font size of 14 pt rather than 12 pt. Rather than change each cell, simply edit the style. All cells with that particular style change automatically.

Applying styles

Excel includes a good selection of predefined named styles that work in conjunction with document themes. Figure 5.21 shows the effect of choosing Home \Rightarrow Styles \Rightarrow Cell Styles. Note that this display is a live preview—as you move your mouse over the style choices, the selected cell or range temporarily displays the style. When you see a style you like, click it to apply the style to the selection.

NOTE

If Excel's window is wide enough, you won't see the Cell Styles command in the Ribbon. Instead, you'll see four or more formatted style boxes. Click the drop-down arrow to the right of these boxes to display all of the defined styles.

NOTE

By default, all cells use the Normal style. If you modify the Normal style, all cells that haven't been assigned a different style will reflect the new formatting.

After you apply a style to a cell, you can apply additional formatting to it by using any formatting method discussed in this chapter. Formatting modifications that you make to the cell don't affect other cells that use the same style.

You have quite a bit of control over styles. In fact, you can do any of the following:

- Modify an existing style.
- Create a new style.
- Merge styles from another workbook into the active workbook.

The following sections describe these procedures.

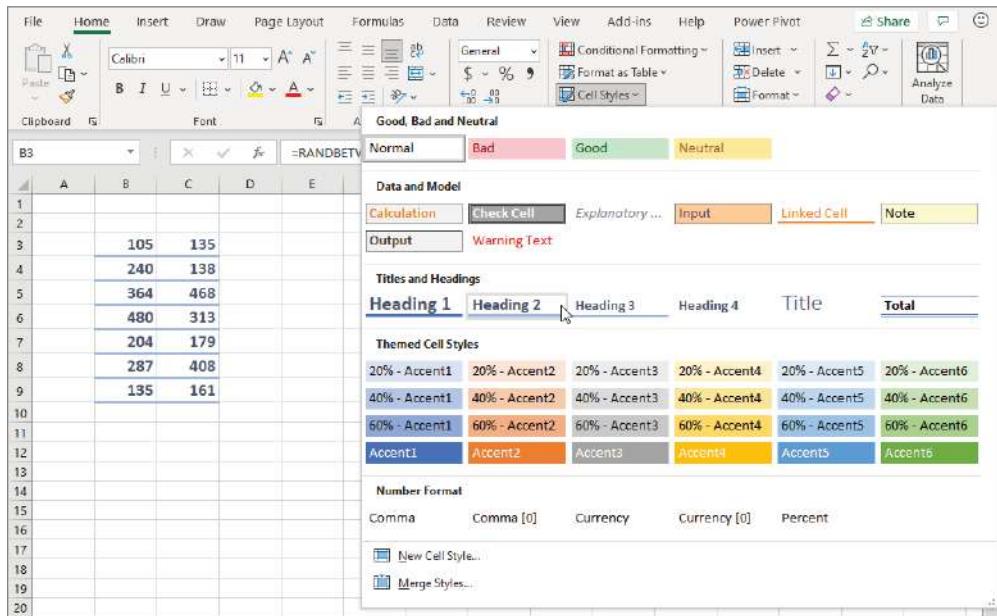
Modifying an existing style

To change an existing style, choose Home \Rightarrow Styles \Rightarrow Cell Styles. Right-click the style that you want to modify and choose Modify from the shortcut menu. Excel displays the Style dialog box, as shown in Figure 5.22. In this example, the Style dialog box shows the settings for the Office theme Normal style—which is the default style for all cells. The style definitions vary, depending on which document theme is active.

Part I: Getting Started with Excel

FIGURE 5.21

Excel displays samples of predefined cell styles.



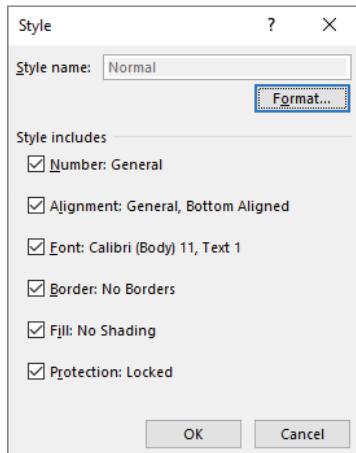
Here's a quick example of how you can use styles to change the default font used throughout your workbook:

1. **Choose Home ▾ Styles ▾ Cell Styles.** Excel displays the list of styles for the active workbook.
2. **Right-click Normal and choose Modify.** Excel displays the Style dialog box (refer to Figure 5.22), with the current settings for the Normal style.
3. **Click the Format button.** Excel displays the Format Cells dialog box.
4. **Click the Font tab and choose the font and size that you want as the default.**
5. **Click OK to return to the Style dialog box.** Notice that the Font item displays the font choice you made.
6. **Click OK again to close the Style dialog box.**

The font for all cells that use the Normal style changes to the font that you specified. You can change any formatting attributes for any style.

FIGURE 5.22

Use the Style dialog box to modify named styles.



Creating new styles

In addition to using Excel's built-in styles, you can create your own styles. This feature can be quite handy because it enables you to apply your favorite formatting options quickly and consistently.

To create a new style, follow these steps:

- 1. Select a cell and apply all the formatting that you want to include in the new style.** You can use any of the formatting that is available in the Format Cells dialog box.
- 2. After you format the cell to your liking, choose Home ➔ Styles ➔ Cell Styles, and choose New Cell Style.** Excel displays its Style dialog box (refer to Figure 5.22), along with a proposed generic name for the style. Note that Excel displays the words By Example to indicate that it's basing the style on the current cell.
- 3. Enter a new style name in the Style Name field.** The check boxes display the current formats for the cell. By default, all check boxes are selected.
- 4. (Optional) If you don't want the style to include one or more format categories, remove the check(s) from the appropriate check box(es).**
- 5. Click OK to create the style and to close the dialog box.**

After you perform these steps, the new custom style is available when you choose Home ➔ Styles ➔ Cell Styles. To delete a custom style, right-click it in the Styles gallery and choose Delete. Custom styles are available only in the workbook in which they were created. To copy your custom styles to another workbook, see the section that follows.

Note

The Protection option in the Style dialog box controls whether users will be able to modify cells for the selected style. This option is effective only if you've also turned on worksheet protection by choosing Review \Rightarrow Protect \Rightarrow Protect Sheet.

Merging styles from other workbooks

Custom styles are stored with the workbook in which they were created. If you've created some custom styles, you probably don't want to go through all of the work required to create copies of those styles in each new Excel workbook. A better approach is to merge the styles from a workbook in which you previously created them.

To merge styles from another workbook, open both the workbook that contains the styles that you want to merge and the workbook that will contain the merged styles. Activate the second workbook, choose Home \Rightarrow Styles \Rightarrow Cell Styles, and then choose Merge Styles. Excel displays the Merge Styles dialog box that shows a list of all open workbooks. Select the workbook that contains the styles you want to merge and click OK. Excel copies custom styles from the workbook that you selected into the active workbook.

Controlling styles with templates

When you start Excel, it loads with several default settings, including the settings for stylistic formatting. If you spend a lot of time changing the default elements for every new workbook, you should know about templates.

Here's an example. You may prefer that gridlines aren't displayed in worksheets. And maybe you prefer Wrap Text to be the default setting for alignment. Templates provide an easy way to change defaults.

The trick is to create a workbook with the Normal style modified in the way you want it. Then save the workbook as a template (with an .xltx extension). After doing so, you can choose this template as the basis for a new workbook.



Refer to Chapter 6, "Understanding Excel Files and Templates," for more information about templates.

Understanding Document Themes

In an attempt to help users create more professional-looking documents, the Office designers incorporated a feature known as *document themes*. Using themes is an easy (and almost fool-proof) way to specify colors, fonts, and a variety of graphic effects in a document. Best of all, changing the entire look of your document is a breeze. A few mouse clicks is all that it takes to apply a different theme and change the look of your workbook.

Importantly, the concept of themes is incorporated into other Office applications. Therefore, a company can easily create a standard and consistent look for all of its documents.

Note

Themes don't override specific formatting that you apply. For example, assume that you apply the Accent1 named style to a range. Then you change the font color for a few cells in that range. If you change to a different theme, the manually applied font colors won't be modified to use the new theme font colors. Bottom line: If you plan to take advantage of themes, stick with the default formatting choices.

Figure 5.23 shows a worksheet that contains a SmartArt diagram, a table, and a chart, and a range formatted with the Title named style, and a range formatted with the Explanatory Text named style. These items all use the default theme, which is the Office theme.

FIGURE 5.23

The elements in this worksheet use the default theme.

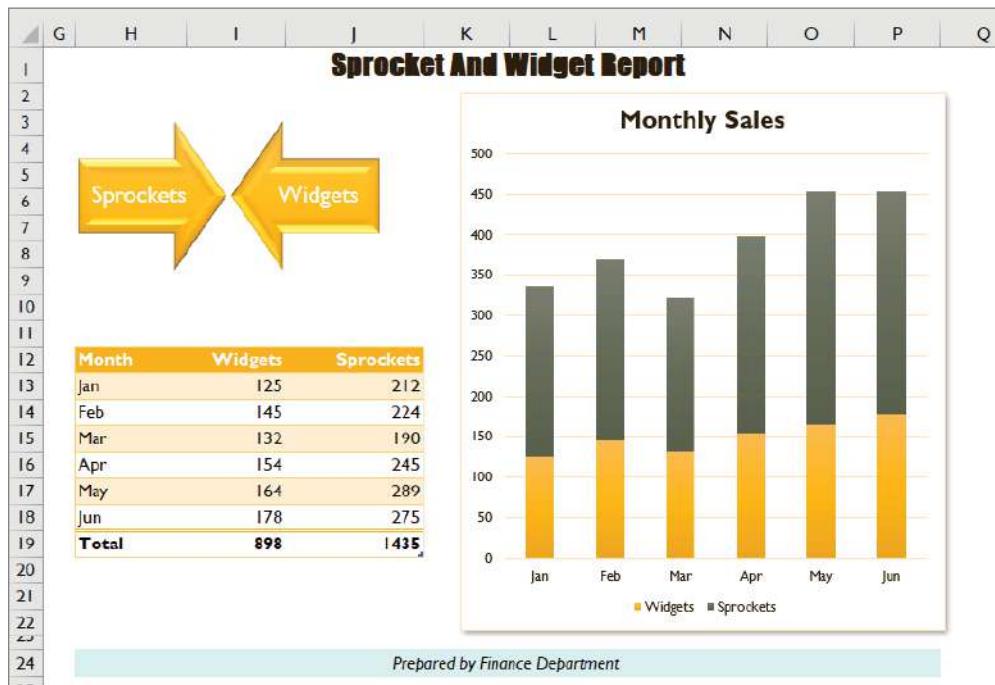


Figure 5.24 shows the same worksheet after applying a different document theme. The different theme changed the fonts, the colors (which may not be apparent in the figure), and the graphics effects for the SmartArt diagram.

FIGURE 5.24

The worksheet after applying a different theme



If you'd like to experiment with using various themes, the workbook shown in Figure 5.23 and Figure 5.24 is available on this book's website at www.wiley.com/go/excel365bible. The file is named theme examples.xlsx.

Applying a theme

Figure 5.25 shows the theme choices that appear when you choose Page Layout → Themes → Themes. This display is a live preview. As you move your mouse over the theme choices, the active worksheet displays the theme. When you see a theme you like, click it to apply the theme to all worksheets in the workbook.

NOTE

A theme applies to the entire workbook. You can't use different themes on different worksheets within a workbook.

FIGURE 5.25

Built-in Excel theme choices



When you specify a particular theme, the gallery choices for various elements reflect the new theme. For example, the chart styles that you can choose from vary, depending on which theme is active.

NOTE

Because themes use different fonts (which can vary in size), changing to a different theme may affect the layout of your worksheet. For example, after you apply a new theme, a worksheet that printed on a single page may spill over to a second page. Therefore, you may need to make some adjustments after you apply a new theme.

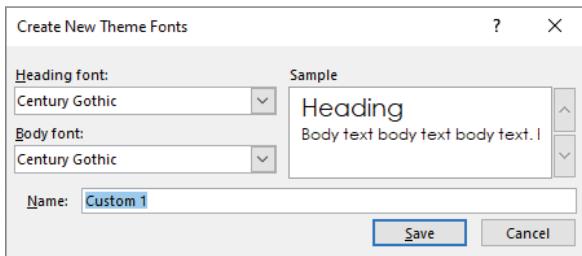
Customizing a theme

Notice that the Themes group on the Page Layout tab contains three other controls: Colors, Fonts, and Effects. You can use these controls to change just one of the three components of a theme. For example, you might like the colors and effects in the Office theme but would prefer different fonts. To change the font set, apply the Office theme and then specify your preferred font set by choosing Page Layout \Rightarrow Themes \Rightarrow Fonts.

Each theme uses two fonts (one for headers and one for the body), and in some cases these two fonts are the same. If none of the theme choices is suitable, choose Page Layout \Rightarrow Themes \Rightarrow Fonts \Rightarrow Customize Fonts to specify the two fonts that you prefer (see Figure 5.26).

FIGURE 5.26

Use this dialog box to specify two fonts for a theme.



TIP

When you choose Home \Rightarrow Font \Rightarrow Font, the two fonts for the current theme are listed first in the drop-down list.

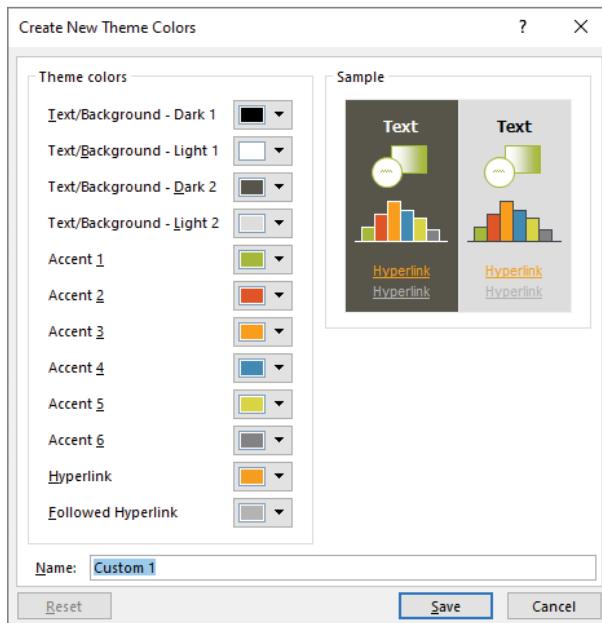
Choose Page Layout \Rightarrow Themes \Rightarrow Colors to select a different set of colors. And, if you're so inclined, you can even create a custom set of colors by choosing Page Layout \Rightarrow Themes \Rightarrow Colors \Rightarrow Customize Colors. This command displays the Create New Theme Colors dialog box, as shown in Figure 5.27. Note that each theme consists of twelve colors. Four of the colors are for text and backgrounds, six are for accents, and two are for hyperlinks. As you specify different colors, the preview panel in the dialog box updates.

NOTE

Theme effects operate on graphics elements, such as SmartArt, shapes, and charts. You can choose a different set of theme effects, but you can't customize theme effects.

FIGURE 5.27

If you're feeling creative, you can specify a set of custom colors for a theme.



If you've customized a theme using different fonts or colors, you can save the new theme by choosing Page Layout \Rightarrow Themes \Rightarrow Save Current Theme. Your customized themes appear in the Themes list in the Custom category. Other Office applications, such as Word and PowerPoint, can use these theme files. If you need to delete your custom theme, right-click it in the Themes gallery and choose Delete.

Understanding Excel Files and Templates

IN THIS CHAPTER

Creating a new workbook

Opening an existing workbook

Saving and closing workbooks

Working with templates

This chapter describes the operations that you perform with workbook files: opening, saving, closing, and so on. It discusses how Excel uses files and provides an overview of the various types of files. Most of the file operations discussed here occur in the *Backstage view*, the screen that you see when you click the File button above the Excel Ribbon. It also discusses templates, a special kind of workbook file.

Creating a New Workbook

When you start Excel, it displays a Start screen that lists recently used files and shows templates that you can use as the basis for a new workbook. One of the template options is Blank Workbook, which gives you an empty workbook. Figure 6.1 shows a portion of the Start screen.

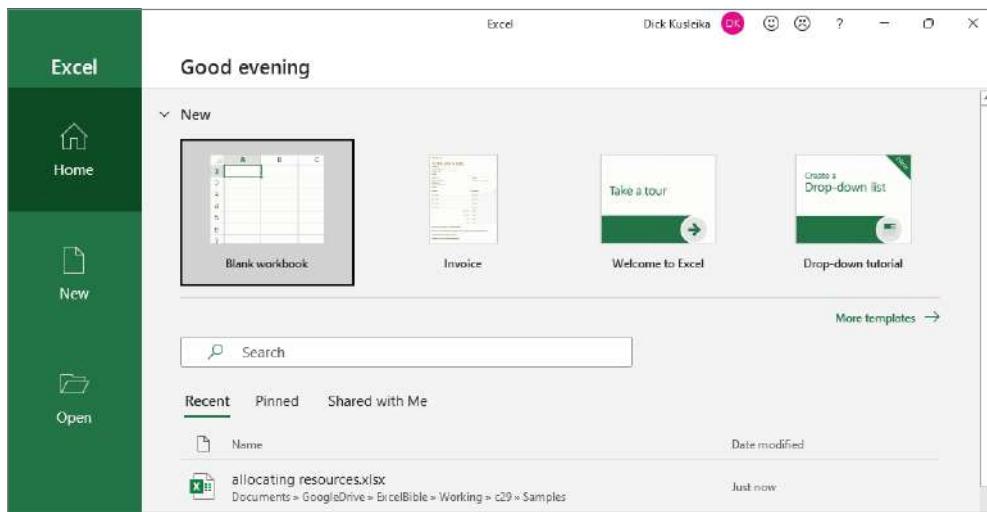
TIP

If you prefer to skip the Start screen and always start with an empty workbook, choose File \Rightarrow Options. In the Excel Options dialog box, click the General tab and remove the check mark from the option labeled Show the Start Screen When This Application Starts.

After you start Excel and create a blank workbook, the empty workbook is called *Book1*. This workbook exists only in memory, and it hasn't been saved to disk. By default, this workbook contains one worksheet named *Sheet1*. If you're starting a project from scratch, you can use this blank workbook. By the way, you can change the default number of sheets in a new workbook by using the General tab of the Excel Options dialog box.

FIGURE 6.1

Choosing Blank workbook from Excel's Start screen



While you're working in Excel, you can create a new (empty) workbook at any time. Excel provides two ways to create a new workbook:

- Choose File \Rightarrow New, which displays a screen that lets you create a blank workbook or a workbook based on a template. To create a new blank workbook, click Blank workbook.
- Press Ctrl+N. This shortcut is the fastest way to start a new workbook if you're not using a template.

Opening an Existing Workbook

Here are some of the ways to open a workbook that's already been saved:

- Choose File \Rightarrow Open \Rightarrow Recent and then select the file that you want from the list on the right. Only the most recently used files are listed. You can specify the number of files to display (maximum of 50) in the Advanced section of the Excel Options dialog box.
- Choose File \Rightarrow Open and choose a location from the list on the left. The locations will vary depending on which "places" you've set up. You may see cloud-based options. One of the options is always This PC. You can navigate to your files directly using the list, or you can click Browse to display the Open dialog box, which gives you many more options.

- Locate the Excel workbook file via a File Explorer file list. Just double-click the file-name (or icon), and the workbook opens in Excel. If Excel isn't running, Windows automatically starts Excel and loads the workbook file.

TIP

When you choose File ➔ Open ➔ Recent, each file in the recent workbooks list displays a pushpin icon on the right when you hover the mouse pointer over the filename. Click the pushpin icon, and that file becomes “pinned” to the list and will always appear at the top of the list. This handy feature ensures that important files always appear on the recent workbooks list—even if you haven’t opened the file recently.

Also, notice that you can right-click a workbook in the list and choose Remove from List. Choose Clear Unpinned Items to clear the list and start fresh.

To open a workbook from the Open dialog box, use the folder tree display on the left to locate the folder that contains the file and then select the workbook file from the list on the right. You can resize the Open dialog box by using the control in the lower-right corner. After you locate and select the file, click Open, and the file opens. Or just double-click the filename to open it.

Notice that the Open button is actually a drop-down list. Click the arrow, and you will see these additional options:

Open Opens the file normally.

Open Read-Only Opens the selected file in read-only mode. When a file is opened in this mode, you can’t save changes with the original filename.

Open as Copy Opens a copy of the selected file. If the file is named budget.xlsx, the workbook that opens is named Copy (1) budget.xlsx.

Open in Browser Opens the file in your default web browser. If the file can’t be opened in a browser, this option is disabled.

Open in Protected View Opens the file in a special mode that doesn’t allow editing. In this view, most of the Excel Ribbon commands are disabled. Read more about this feature in the nearby sidebar “About Protected View.”

Open and Repair Attempts to open a file that may be damaged and recover information contained in it.

Show Previous Versions Applies to documents stored where version histories are maintained, as on OneDrive or SharePoint Online.

TIP

In the Open dialog box, you can hold down the Ctrl key and select multiple workbooks. When you click Open, all the selected workbook files open.

Right-clicking a filename in the Open dialog box displays a shortcut menu with many extra Windows commands. For example, you can copy, delete, or rename the file; modify its properties; and so on. Almost anything you can do in a File Explorer window, you can do in the Open dialog box.

About Protected View

Excel 2010 introduced a security feature known as Protected View. Although it might seem like Excel is trying to keep you from opening your own files, Protected View is all about protecting you from malware. Malware refers to something that can harm your system. Hackers have figured out several ways to manipulate Excel files in a way that harmful code can be executed. Protected View essentially prevents these types of attacks by opening a file in a protected environment (a *sandbox*).

If you open an Excel workbook that you downloaded from outside your local network (like from the Internet), you'll see a colorful message above the Formula bar. In addition, the Excel title bar displays [Protected View]. Choose File \Rightarrow Info to find out why Excel opened the file in Protected View.

If you're certain that the file is safe, click Enable Editing. If you don't enable editing, you'll be able to view the contents of the workbook, but you won't be able to make any changes to it.

If the workbook contains macros, you'll see another message after you enable editing: Security Warning. Macros have been disabled. If you're sure that the macros are harmless, click Enable Content.

By default, Protected View kicks in for the following:

- Files downloaded from the Internet
- Attachments opened from Outlook
- Files opened from potentially unsafe locations, such as your Temporary Internet Files folder
- Files that are blocked by the File Block Policy (a Windows feature that allows administrators to define potentially dangerous files)
- Files that have a digital signature that has expired

In some situations, you don't care about working with the document. You just want to print it. In that case, choose File \Rightarrow Print and then click the Enable Printing button.

Also, note that you can copy a range of cells from a workbook in Protected View and paste it into a different workbook.

You have some control over the types of files that trigger Protected View. To change the settings, choose File \Rightarrow Options and click Trust Center. Then click the Trust Center Settings button and click the Protected View tab in the Trust Center dialog box.

Filtering filenames

Near the bottom-right corner of the Open dialog box is a drop-down list. When the Open dialog box is displayed, this list defaults to All Excel Files (and a long list of filename extensions).

The Open dialog box displays only those files that match the filename extensions. In other words, you see only standard Excel files.

If you want to open a file of a different type, click the arrow in the drop-down list and select the file type that you want to open. This changes the filtering and displays only files of the type that you specify.

You can also type a filter directly in the File Name box. For example, typing ***.xlsx** will display only files that have an **.xlsx** filename extension. Press Enter after typing the filter.

Choosing your file display preferences

The Open dialog box can display your workbook filenames in several styles: as a list, with complete details, as icons, and so on. You control the style by clicking the More Options arrow (near the upper-right corner) and then selecting a display style from the drop-down list.

Opening Workbooks Automatically

Many people work on the same workbooks each day. If this describes you, you'll be happy to know that Excel can open specific workbook files automatically whenever you start Excel. Any workbooks placed in the XLStart folder open automatically.

The location of the XLStart folder varies, depending on your Windows version. To determine the location of the XLStart folder on your system, follow these steps:

1. Choose File \Rightarrow Options, and select the Trust Center tab.
2. Click the Trust Center Settings button. The Trust Center dialog box appears.
3. In the Trust Center dialog box, select the Trusted Locations tab. You'll see a list of trusted locations.
4. Look in the path for the location described as User Startup. The path might look something like this:

C:\Users\<username>\AppData\Roaming\Microsoft\Excel\XLSTART\

Another XLStart folder may be located here:

C:\Program Files\Microsoft Office\root\Office16\XLSTART\

Any workbook files (excluding template files) stored in either of these XLStart folders open automatically when Excel starts. If one or more files open automatically from an XLStart folder, Excel won't start with the Start screen or a blank workbook.

You can specify an alternate startup folder in addition to the XLStart folder. Choose File \Rightarrow Options and select the Advanced tab. Scroll down to the General section and enter the path to the new folder name in the At Startup, Open All Files In field. Then when you start Excel, it automatically opens all workbook files in both the XLStart folders and the alternate folder that you specified.

Saving a Workbook

When you're working in Excel, your workbook is vulnerable to day-ruining events such as power failures and system crashes. Therefore, you should save your work often. Saving a file takes only a few seconds, but re-creating lost work can take many hours.

Part I: Getting Started with Excel

Excel provides four ways to save your workbook:

- Click the Save icon on the Quick Access Toolbar. (It looks like an old-fashioned floppy disk.)
- Press Ctrl+S.
- Press Shift+F12.
- Choose File ➔ Save.

CAUTION

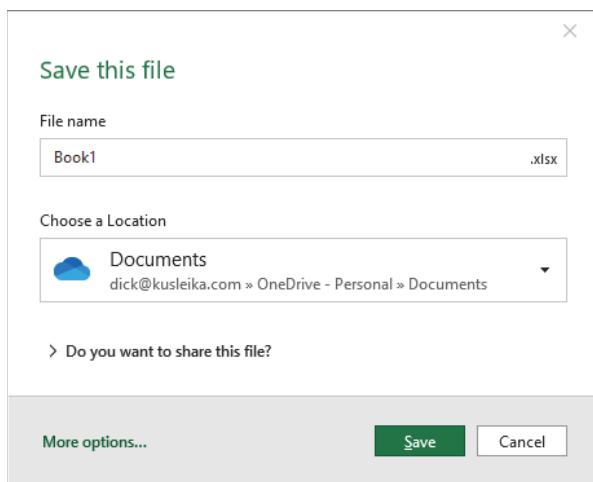
Saving a file overwrites the previous version of the file on your hard drive. If you open a workbook and then completely mess it up, don't save the file. Instead, close the workbook without saving it and then reopen the good copy.

If your workbook has already been saved, it's saved again using the same filename in the same location. If you want to save the workbook to a new file or to a different location, choose File ➔ Save As (or press F12).

If your workbook has never been saved, Excel displays a Save This File dialog box, as shown in Figure 6.2. You can enter a filename and choose a location, such as on your PC or in your OneDrive account—if you have one. A new (unsaved) workbook has a default name, such as Book1 or Book2. Although Excel allows you to use these generic workbook names for filenames, you'll almost always want to specify a more descriptive filename. This dialog box has a More Options link that takes you to the Save As portion of Backstage view.

FIGURE 6.2

Saving a previously unsaved workbook



The Save As portion of Backstage view has most of the same information as the Save This File dialog box, just in a different layout. There are locations on the left side and folders on the right side. The Browse option on the left opens the classic Save As dialog box. The Save As dialog box is similar to the Open dialog box. Select the desired folder in the folder list on the left. After you select the folder, enter the filename in the File Name field. You don't need to specify a filename extension—Excel adds it automatically, based on the file type specified in the Save as Type field. By default, files are saved in the standard Excel file format, which uses an .xlsx filename extension. If the file contains macros, Excel defaults to the .xlsm filename extension.

TIP

To change the default file format for saving files, choose File ➔ Options to access the Excel Options dialog box. Click the Save tab and change the setting for the Save Files in This Format option. For example, if your workbooks must be compatible with older versions of Excel (versions before Excel 2007), you can change the default format to Excel 97–2003 Workbook (*.xls). Doing so eliminates the need to select the older file type every time you save a new workbook.

If a file with the same name already exists in the location that you specify, Excel asks whether you want to overwrite that file with the new file. Be careful! You can't recover the previous file after you overwrite it.

File-Naming Rules

Excel workbook files are subject to the same rules that apply to other Windows files. A filename can be up to 255 characters, including spaces. This length enables you to give meaningful names to your files. You can't, however, use any of the following characters in your filenames:

\ (slash)	? (question mark)
: (colon)	* (asterisk)
" (quote)	< (less than)
> (greater than)	(vertical bar)

You can use uppercase and lowercase letters in your names to improve readability, but the filenames aren't case sensitive. For example, My 2022 Budget.xlsx and MY 2022 BUDGET.xlsx are equivalent names.

Using AutoRecover

If you've used computers for any length of time, you've probably lost some work. You forgot to save a file, or maybe the power went out and your unsaved work was lost. Or maybe you were working on something and didn't think it was important, so you closed the file without

saving. Of course, later you realized that it was *indeed* important. Excel's AutoRecover feature might make these types of "d'oh!" moments less frequent.

As you work in Excel, your work is periodically saved automatically. It happens in the background so you don't even know that it's happening. If necessary, you can access these autosaved versions of your work. This even applies to workbooks that you never explicitly saved.

The AutoRecover feature consists of two components:

- Versions of a workbook are saved automatically, and you can view them.
- Workbooks that you closed without saving are saved as draft versions.

Recovering versions of the current workbook

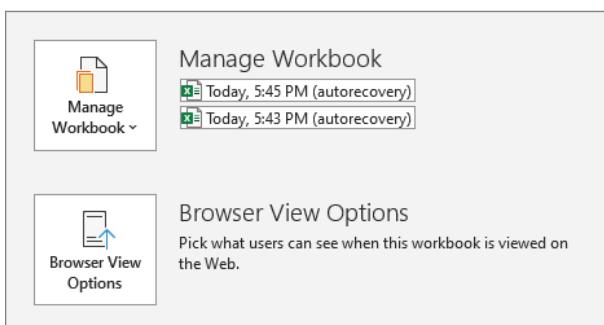
To see whether any previous versions of the active workbook are available, choose File \Rightarrow Info. The Manage Workbook section lists the available old versions (if any) of the current workbook. In some cases, more than one autosaved version will be listed. In other cases, no autosaved versions will be available. Figure 6.3 shows a workbook with two recovery points.

You can open an autosaved version by clicking its name. Remember that opening an autosaved version won't automatically replace the current version of your workbook. Therefore, you can decide whether the autosaved version is preferable to the current version. Or you can just copy some information that may have been accidentally deleted and paste it to your current workbook.

When you close the workbook, the autosaved versions are deleted.

FIGURE 6.3

You can recover older versions of your workbook.



Recovering unsaved work

When you close a workbook without saving your changes, Excel asks if you're sure. If that unsaved workbook has an autosaved version, the Are You Sure dialog box informs you of that fact.

To recover a workbook that you closed without saving, choose File \Rightarrow Info \Rightarrow Manage Workbook \Rightarrow Recover Unsaved Workbooks. You'll see a list of all draft versions of your workbooks. You can open them and (if you're lucky) recover something that you needed. Note that the unsaved workbooks are stored in the XLSB file format and are read-only files. If you want to save one of these files, you need to provide a new name.

Draft versions are deleted after four days or when you edit the file.

Configuring AutoRecover

Normally, AutoRecover files are saved every 10 minutes. You can adjust the AutoRecover save time on the Save tab of the Excel Options dialog box. You can specify a save interval of from 1 to 120 minutes.

If you work with sensitive documents, you might prefer that previous versions aren't saved automatically on your computer. The Save tab of the Excel Options dialog box lets you disable this feature completely or disable it just for a specific workbook.

Password-Protecting a Workbook

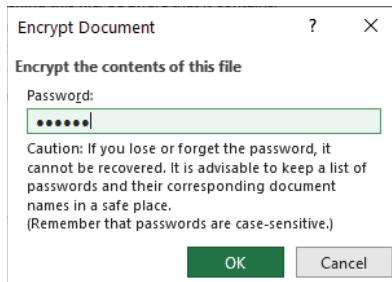
In some cases, you may want to specify a password for your workbook. When a user attempts to open a password-protected workbook, a password must be entered before the file is opened.

To set a password for a workbook, follow these steps:

1. Choose File \Rightarrow Info and click the Protect Workbook button. This button displays some additional options in a drop-down list.
2. Choose Encrypt with Password from the list. Excel displays the Encrypt Document dialog box, as shown in Figure 6.4.

FIGURE 6.4

The Encrypt Document dialog box is where you specify a password for your workbook.



3. Enter the password, click OK, and then enter it again.
4. Click OK and save the workbook.

When you reopen the workbook, you'll be prompted for a password.

CAUTION

Passwords are case sensitive. Be careful when using password protection because it's impossible to open the workbook (using normal methods) if you forget the password. Also, remember that Excel passwords can be cracked, so it's not a perfect security measure.

Organizing Your Files

If you have hundreds of Excel files, you might have a problem locating the workbook that you need. Using descriptive filenames can help. Using folders and subfolders (with descriptive names) also makes it easier to find the particular file you need. In some cases, though, that's not enough.

Fortunately, Excel lets you assign a variety of descriptive information (sometimes known as *metadata*) to a workbook. These are known as document properties. This information includes such items as the author, tags, and categories.

When you choose File \Rightarrow Info, you can view (or modify) the document properties for the active workbook. This information is shown on the right side of the screen. Click the link at the bottom to toggle between showing more or fewer properties.

TIP

To access even more properties for your workbook, click the down arrow on Properties and choose Advanced Properties.

Other Workbook Info Options

The Info pane of Backstage view displays more file-related options. To display this pane, choose File \Rightarrow Info. These options, described in the following sections, may be useful if you plan to distribute your workbook to others. Note that not all workbooks display all the options described in the following sections. Only the relevant options are shown.

Protect Workbook options

The File \Rightarrow Info \Rightarrow Protect Workbook drop-down list contains the following options:

Always Open Read-Only Use this option to save the file as read-only to prevent changes.

Encrypt with Password Use this command to specify a password that is required to open the workbook. (See "Password-Protecting a Workbook" earlier in this chapter.)

Protect Current Sheet This command lets you protect various elements of a worksheet. It displays the same dialog box as the Review \Rightarrow Protect \Rightarrow Protect Sheet command.

Protect Workbook Structure This command lets you protect the structure of a workbook. It displays the same dialog box as Review \Rightarrow Protect \Rightarrow Protect Workbook.

Restrict Access If your organization uses the Azure Rights Management System, you can connect your workbook to it and get more granular options for protecting it.

Add a Digital Signature This command allows you to “sign” a workbook digitally so that users can trust that you signed it.

Mark as Final Use this option to designate the workbook as “final.” The document is saved as a read-only file to prevent changes. This isn’t a security feature. Rather, the Mark as Final command is useful to let others know that you’re sharing a completed version of a workbook.



See Chapter 31, “Protecting Your Work,” for more information about protecting worksheets, protecting workbooks, and using digital signatures.

Check for Issues options

The File \Rightarrow Info \Rightarrow Check for Issues drop-down list contains the following options:

Inspect Document This command displays the Document Inspector dialog box. The Document Inspector can alert you to some potentially private information that may be contained in your workbook—perhaps information that’s contained in hidden rows or columns or hidden worksheets. If you plan to make a workbook available to a large audience, it’s an excellent idea to use the Document Inspector for a final check.

Check Accessibility This command checks the workbook for potential problems that might occur for people with disabilities. The results of the check are displayed in a task pane in the workbook.

Check Compatibility This command is useful if you need to save your workbook in an older file format. It displays a helpful Compatibility Checker dialog box that lists potential compatibility problems. This dialog box also appears when you save a workbook using an older file format.

Version History

If you save your files to OneDrive, Microsoft will maintain a version history every time you save the file. You can view or restore prior versions.

Manage Workbook options

If Excel automatically saved previous versions of your workbook, you can recover one of the earlier versions.

Browser View options

If your workbook will be viewed in a web browser, you can specify which sheets and other objects will be viewable.

Compatibility Mode section

If the active workbook is an old workbook opened in compatibility mode, you'll see the Compatibility Mode section in the Info pane. To convert the workbook to the current Excel file format, click the Convert button.

CAUTION

Be aware that this command deletes the original version of the file, which seems like a rather drastic measure. It's probably wise to make a copy of your workbook before you use this command.

Closing Workbooks

After you're finished with a workbook, you can close it to free the memory that it uses. Other workbooks will remain open. When you close the last open workbook, you also close Excel.

You can close a workbook by using any of the following methods:

- Choose File ➔ Close.
- Click the Close button (the X) in the right corner of the window's title bar.
- Press Ctrl+F4.
- Press Ctrl+W.

If you've made any changes to your workbook since it was last saved, Excel asks if you want to save the changes to the workbook before closing it.

Safeguarding Your Work

Nothing is more frustrating than spending hours creating a complicated Excel workbook only to have it destroyed by a power failure, a hard drive crash, or even human error. Fortunately, protecting yourself from these disasters is not a difficult task.

Earlier in the chapter, we discussed the AutoRecover feature that makes Excel save a backup copy of your workbook at regular intervals (see "Using AutoRecover"). AutoRecover is a good idea, but it certainly isn't the only backup protection you should use. If a workbook is important, you need to take extra steps to ensure its safety. The following backup options help ensure the safety of individual files:

Keep a backup copy of the file on the same drive. Although this option offers some protection if you make a mess of the workbook, it won't do you any good if the entire hard drive crashes.

Keep a backup copy on a different hard drive. This method assumes, of course, that your system has more than one hard drive. This option offers more protection than the preceding method because the likelihood that both hard drives will fail is remote. If the entire system is destroyed or stolen, however, you're out of luck.

Keep a backup copy on a network server. This method assumes that your system is connected to a server on which you can write files. This method is fairly safe. If the network server is located in the same building, however, you're at risk if the entire building burns down or is otherwise destroyed.

Keep a backup copy on an Internet backup site. Several websites specialize in storing backup files.

Keep a backup copy on a removable medium. This is probably the safest method. Using a removable medium, such as a USB drive, enables you to physically take the backup to another location. So if your system (or the entire building) is damaged, your backup copy remains intact.

Working with Templates

A *template* is essentially a model that serves as the basis for something else. An Excel template is a special type of workbook that's used as the basis to create other workbooks. The following sections discuss some of the templates available from Microsoft, and it describes how to create your own template files. Creating a template takes some time, but in the long run, doing so may save you a lot of work.

Exploring Excel templates

The best way to become familiar with Excel template files is to jump in and try a few. Excel gives you quick access to hundreds of template files.

TIP

Examining templates is also a good way to learn about Excel. You may discover some techniques that you can incorporate into your own work.

Viewing templates

To explore the Excel templates, choose File \Rightarrow New. The template thumbnails displayed on the screen that appears are just a small sampling of those that are available. Click one of the suggested search terms, or enter a descriptive word and search for more.

NOTE

The searching is done at Microsoft Office Online, so you must be connected to the Internet to search for templates.

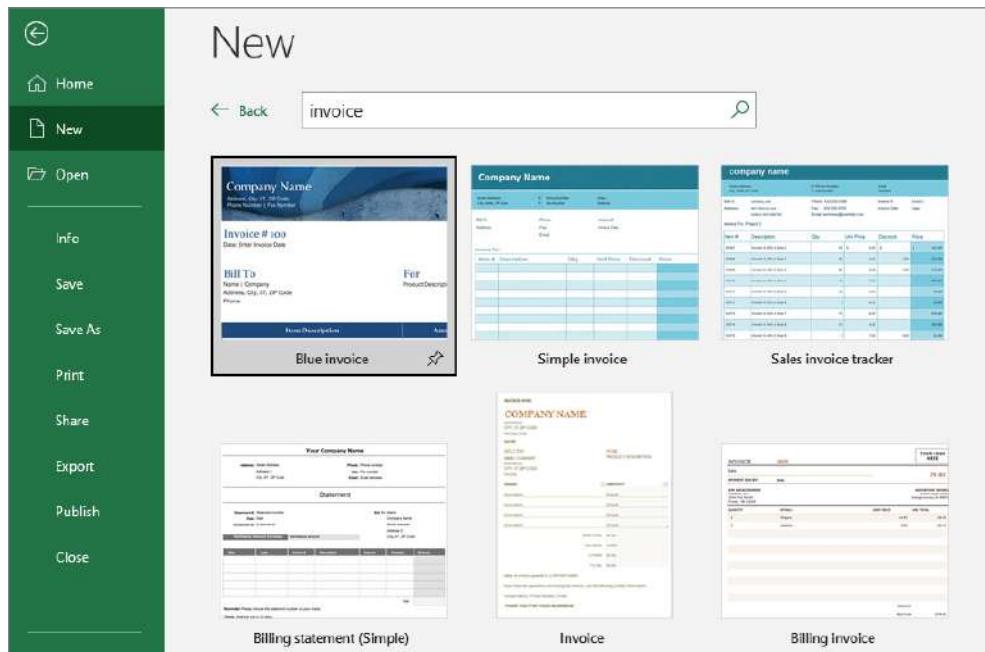
Part I: Getting Started with Excel

For example, enter **invoice** and click the Search button. Excel displays many thumbnails.

Figure 6.5 shows the results of a template search for **invoice**.

FIGURE 6.5

The New page in Backstage view allows you to search for templates.



Note

Microsoft Office Online has a variety of templates. Some are better than others, so if you download a few duds, don't give up. Even though a template may not be perfect, you may be able to modify it to meet your needs. Modifying an existing template is often easier than creating a workbook from scratch.

Creating a workbook from a template

To create a workbook based on a template, just locate a template that looks like it might do the job and click the thumbnail. Excel displays a box with a larger image, the source for the template, and some additional information. If it still looks good, click the Create button. Otherwise, click one of the arrows to view details for the next (or previous) template in the list.

When you click the Create button, Excel downloads the template and then creates a new workbook based on that template.

What you do next depends on the template. Every template is different, but most are self-explanatory. Some workbooks require customization. Just replace the generic information with your own information.

Note

It's important to understand that you're not working with the template file. Instead, you're working with a workbook that was created from the template file. If you make any changes, you're not changing the template—you're changing the workbook that's based on the template. After you download a template from Microsoft Office Online, that template is saved for future use. (You won't have to download it again.) Downloaded templates appear as thumbnails when you choose File \Rightarrow New.

Figure 6.6 shows a workbook created from a template. This workbook needs to be customized in several areas. But if this template will be used again, it's more efficient to customize the template rather than every workbook created from the template.

FIGURE 6.6

A workbook created from a template

	A	B	C	D
1		INVOICE #100		
2		COMPANY NAME		
3	ADDRESS			
4	CITY, ST ZIP CODE			
5	PHONE FAX			
6	DATE			
7	BILL TO		FOR	
8	NAME COMPANY		PRODUCT DESCRIPTION	
9	ADDRESS			
10	CITY, ST ZIP CODE			
11	PHONE			
12	Details	<input type="checkbox"/> AMOUNT		
	Description		Amount	
	Description		Amount	

If you want to save the newly created workbook, click the Save button. Excel proposes a filename based on the template's name, but you can use any filename you like.

Modifying a template

A template file that you download is just like a workbook file. You can open a template file, make changes to it, and then resave it. For example, with the invoice template shown in Figure 6.6, you may want to modify it so that it shows your company information and logo and uses your actual sales tax rate. Then, when you use that template in the future, the workbook created from it will already be customized.

To open a template for editing, choose File \Rightarrow Open (not File \Rightarrow New) and locate the template file (it will have an .xltx, .xltm, or .xlt filename extension). When you open a template file by choosing File \Rightarrow Open, you're opening the actual template file—you are *not* creating a workbook from the template file.

One way to find the location of your downloaded template files is to look at your trusted locations list:

1. **Choose File \Rightarrow Options.** The Excel Options dialog box appears.
2. **Choose Trust Center and click the Trust Center Settings button.** The Trust Center dialog box appears.
3. **In the Trust Center dialog box, choose Trusted Locations.** You'll see a list of trusted locations. Downloaded templates are stored in the location described as User Templates. If you want to modify (or delete) a downloaded template, this is where you'll find it.

On our systems, downloaded templates are stored here:

C:\Users\<username>\AppData\Roaming\Microsoft\Templates\

After you've made changes to the template, use File \Rightarrow Save to save the template file. Future workbooks that you create from this template will use the modified version of the template.

Using default templates

Excel supports three types of templates:

The default workbook template This type is used as the basis for new workbooks.

The default worksheet template This type is used as the basis for new worksheets inserted into a workbook.

Custom workbook templates Usually, these ready-to-run workbooks include formulas, but they can be as simple or as complex as you like. Typically, these templates are set up so that a user can simply plug in values and get immediate results. The Microsoft Office Online templates (discussed earlier in this chapter) are examples of this type of template.

The default workbook and default worksheet templates are discussed in the following sections. Custom workbook templates are discussed later.

Using the workbook template to change workbook defaults

Every new workbook that you create starts out with some default settings. For example, the workbook has one worksheet, the worksheet has gridlines, the page header and footer are blank, text appears in the font defined in the default Normal style, columns are 8.43 units wide, and so on. If you're not happy with any of the default workbook settings, you can change them by creating a workbook template.

Making changes to Excel's default workbook is fairly easy to do, and it can save you lots of time in the long run. Here's how to change Excel's workbook defaults:

1. Open a new workbook.
2. Add or delete sheets to give the workbook the number of worksheets that you want.
3. Make any other changes that you want to make, which can include column widths, named styles, page setup options, and many of the settings that are available in the Excel Options dialog box. To change the default formatting for cells, choose Home \Rightarrow Styles \Rightarrow Cell Styles and then modify the settings for the Normal style. For example, you can change the default font, size, or number format.
4. When your workbook is set up to your liking, choose File \Rightarrow Save As \Rightarrow Browse. The Save As dialog box appears.
5. Select Excel Template (*.xltx) from the Save as Type list. If your template contains VBA macros, select Excel Macro-Enabled Template (*.xltm).
6. Enter book for the filename.

CAUTION

Excel will offer a name, such as Book1.xltx. You must change this name to book.xltx (or book.xltm) if you want Excel to use your template to set the workbook defaults.

7. Save the file in your XLStart folder (not in the Templates folder that Excel proposes).
8. Close the template file.

After you perform the preceding steps, the new default workbook is based on the book.xltx (or book.xltm) workbook template. You can create a workbook based on your template by using any of these methods:

- Press Ctrl+N.
- Open Excel without first selecting a workbook to open. This option works only if you disable the option to show the Start screen when Excel starts. This option is specified in the General tab of the Excel Options dialog box. (Choose File \Rightarrow Options to display the Excel Options dialog box.)

NOTE

The book.xltx template is not used if you choose File \Rightarrow New and choose Blank Workbook from the list of templates. That command results in a default workbook and provides a way to override the custom book.xltx template if you need to do so.

Creating a worksheet template

You can also create a single sheet template named sheet.xltx. Use the same procedure described for book.xltx. The sheet.xltx template is used when you insert a new worksheet.

Editing your template

After you create your book.xltx template, you may discover that you need to change it. You can open the template file and edit it just as you would any other workbook. After you make your changes, save the file to its original location and close it.

Resetting the default workbook

If you create a book.xltx (or book.xltm) file and then decide that you'd rather use the standard default settings, simply delete (or rename) the book.xltx (or book.xltm) template file. Excel then uses its built-in default settings for new workbooks.

Using custom workbook templates

The book.xltx template discussed in the preceding sections is a special type of template that determines default settings for new workbooks. We will now discuss other types of templates, referred to as *custom workbook templates*, which are simply workbooks that you set up as the basis for new, specific types of workbooks.

Creating custom templates

Creating a custom workbook template can eliminate repeating work. Assume that you create a monthly sales report that consists of your company's sales by region, plus several summary calculations and charts. You can create a template file that consists of everything except the input values. Then when it's time to create your report, you can open a workbook based on the template, fill in the blanks, and be finished.

NOTE

You could, of course, just use the previous month's workbook and save it with a different name. This is prone to errors, however, because you easily can forget to use the Save As command and accidentally overwrite the previous month's file. Another option is to choose File \Rightarrow Open and choose Open as Copy in the Open dialog box. (This command appears when you click the arrow on the Open button.) Opening a file as a copy creates a new workbook from an existing one, but it uses a different name to ensure that the old file is not overwritten.

When you create a workbook that's based on a template, the default workbook name is the template name with a number appended. For example, if you create a new workbook based on a template named Sales Report.xlsx, the workbook's default name is Sales Report1.xlsx. The first time you save a workbook created from a template, Excel displays the Save As dialog box so that you can give the workbook a different name if you want to.

A custom template is essentially a normal workbook. It can use any Excel feature, such as charts, formulas, and macros. Usually, a template is set up so that the user can enter values and get immediate results. In other words, most templates include everything but the data, which the user enters.

NOTE

If your template contains macros, it must be saved as an Excel Macro-Enabled Template, with an .xlsm filename extension.

Locking Formula Cells in a Template File

If novices are going to use the template, you might consider locking all the formula cells to make sure that the formulas aren't deleted or modified. By default, all cells are locked and can't be changed when the worksheet is protected. The following steps describe how to unlock the nonformula cells:

1. Choose Home \Rightarrow Editing \Rightarrow Find & Select \Rightarrow Go To Special. The Go To Special dialog box appears.
2. Select Constants and click OK. This step selects all nonformula cells.
3. Press Ctrl+1. The Format Cells dialog box appears.
4. Select the Protection tab.
5. Remove the check mark from the Locked check box.
6. Click OK to close the Format Cells dialog box.
7. Choose Review \Rightarrow Protect \Rightarrow Protect Sheet. The Protect Sheet dialog box appears.
8. Specify a password (optional) and click OK.

After you perform these steps, you can't modify the formula cells unless the sheet is unprotected.

Saving your custom templates

To save a workbook as a template, choose File \Rightarrow Save As \Rightarrow Browse and select Excel Template (*.xlsx) from the Save as Type drop-down list. If the workbook contains any VBA macros, select Excel Macro-Enabled Template (*.xlsm). Save the template in your Templates folder—which Excel automatically suggests—or a folder within that Templates folder.

If you later discover that you want to modify the template, choose File \Rightarrow Open to open and edit the template.

Using custom templates

To create a workbook based on a custom template, choose File \Rightarrow New and then click Personal (above the Search box). You'll see thumbnails of all your custom worksheet templates (plus others). Click a template, and Excel creates a workbook based on the template.

Printing Your Work

IN THIS CHAPTER

Changing your worksheet view

Adjusting your print settings for better results

Preventing some cells from being printed

Using the Custom Views feature

Creating PDF files

Despite predictions of the “paperless office,” paper remains an excellent way to carry information around and share it with others, particularly if there’s no electricity or Wi-Fi where you’re going. Some of the worksheets that you develop with Excel will end up as hard-copy reports, and you’ll want them to look as good as possible. You’ll find that printing from Excel is quite easy and that you can generate attractive, well-formatted reports with minimal effort. In addition, Excel has many options that give you a great deal of control over the printed page. We explain these options in this chapter.

Doing Basic Printing

If you want to print a copy of a worksheet with no fuss and bother, use the Quick Print option. One way to access this command is to choose File \Rightarrow Print (which displays the Print pane of Backstage view) and then click the Print button. The keyboard shortcut Ctrl+P has the same effect as File \Rightarrow Print. When you use Ctrl+P to show the Backstage view, the Print button has the focus, so you can simply press Enter to print.

If you like the idea of one-click printing, take a few seconds to add a new button to your Quick Access Toolbar. Click the downward-pointing arrow on the right of the Quick Access Toolbar and then choose Quick Print from the drop-down list. Excel adds the Quick Print icon to your Quick Access Toolbar.

Clicking the Quick Print button prints the current worksheet on the currently selected printer, using the default print settings. If you've changed any of the default print settings (by using the Page Layout tab), Excel uses the new settings; otherwise, it uses the following default settings:

- Prints the active worksheet (or all selected worksheets), including any embedded charts or objects
- Prints one copy
- Prints the entire active worksheet
- Prints in portrait mode
- Doesn't scale the printed output
- Uses letter-size paper with 0.75-inch margins for the top and bottom and 0.70-inch margins for the left and right margins (for the U.S. version)
- Prints with no headers or footers
- Doesn't print cell notes or comments
- Prints with no cell gridlines
- For wide worksheets that span multiple pages, prints down and then over

When you print a worksheet, Excel prints only the active area of the worksheet. In other words, it won't print all 17 billion cells—just those that have data in them. If the worksheet contains any embedded charts or other graphic objects (such as SmartArt or shapes), they're also printed.

Using Print Preview

When you choose File \Rightarrow Print (or press Ctrl+P), Backstage view displays a preview of your printed output, exactly as it will be printed. Initially, Excel displays the first page of your printed output. To view subsequent pages, use the page controls along the bottom of the preview pane (or use the vertical scrollbar along the right side of the screen).

The Print Preview window has a few other commands (at the bottom) that you can use while previewing your output. For a multipage printout, use the page number controls to jump quickly to a particular page. The Show Margins button toggles the display of margins, and Zoom to Page ensures that a complete page is displayed.

When the Show Margins option is in effect, Excel adds markers to the preview that indicate column borders and margins. You can drag the column or margin markers to make changes that appear onscreen. Changes that you make to column widths in preview mode are also made in the actual worksheet.

Print Preview is certainly useful, but you may prefer to use Page Layout view to preview your output (see "Changing Your Page View").

Changing Your Page View

Page Layout view shows your worksheet divided into pages. In other words, you can visualize your printed output while you work.

Page Layout view is one of three worksheet views, which are controlled by the three icons on the right side of the status bar. You could also use the commands in the View \Rightarrow Workbook Views group on the Ribbon to switch views. The three view options are as follows:

- Normal** The default view of the worksheet. This view may or may not show page breaks.
- Page Layout** Shows individual pages.
- Page Break Preview** Allows you to adjust page breaks manually.

Just click one of the icons to change the view. You can also use the Zoom slider to change the magnification from 10% (a very tiny, bird's-eye view) to 400% (very large, for showing fine detail).

The following sections describe how these views can help with printing.

Normal view

Most of the time when you work in Excel, you use Normal view. Normal view can display page breaks in the worksheet. The page breaks are indicated by horizontal and vertical dotted lines. These page break lines adjust automatically if you change the page orientation, insert or delete rows or columns, change row heights, change column widths, and so on. For example, if you find that your printed output is too wide to fit on a single page, you can adjust the column widths (keeping an eye on the page break display) until the columns are narrow enough to print on one page.

NOTE

Page breaks aren't displayed until you print (or preview) the worksheet at least one time. Page breaks are also displayed if you set a print area by choosing Page Layout \Rightarrow Page Setup \Rightarrow Print Area.

TIP

If you'd prefer not to see the page break in Normal view, choose File \Rightarrow Options and select the Advanced tab. Scroll down to the Display Options for This Worksheet section and remove the check mark from Show Page Breaks. This setting applies only to the active worksheet and is grayed out if you're in anything but Normal view. Unfortunately, the option to turn off page break display is not on the Ribbon, and it's not even available for inclusion on the Quick Access Toolbar.

Figure 7.1 shows a part of a worksheet in Normal view, zoomed to about 50%, and with gridlines turned off. You can see the dotted line between columns D and E that indicates a page break.

FIGURE 7.1

In Normal view, dotted lines indicate page breaks.

	A	B	C	D	E	F
1	Gabriel	Palmer	Gizmonic Institute	Elizabeth	Virginia	
2	Mason	Carpenter	Monks Diner	Beaumont	Virginia	
3	Matthew	Gray	Extensive Enterprise	Denver	South Dakota	
4	Jasmine	Morales	The Frying Dutchman	Inglewood	Texas	
5	Sophia	Rose	Plow King	Athens	Florida	
6	Natalie	Wilson	Big T Burgers and Fries	Anaheim	Utah	
7	Landon	Stone	Carry's Candles	Portland	Colorado	
8	Brady	Alexander	The Legitimate Businessmen's Club	Hollywood	Pennsylvania	
9	Alexa	Mitchell	Widget Corp	Boston	Ohio	
10	Kylie	Ortiz	Gadgetron	Hollywood	Virginia	
11	Ian	Cunningham	Demo Company	Newport News	Michigan	
12	Ariana	Armstrong	Globo Gym American Corp	High Point	Texas	
13	Alexa	Hunter	Zovo Toys	Tampa	District of Columbia	
14	Ryan	Williams	Galaxy Corp	Orlando	Mississippi	
15	Victoria	Crawford	Three Waters	Palmdale	Kentucky	
16	Jasmine	Lawrence	Roxxon	Newark	Nevada	
17	Joshua	Johnson	Sample, inc	Laredo	Alaska	
18	Brian	Carpenter	Spade and Archer	Nashville	Kansas	
19	Tristan	Harris	Taggart Transcontinental	New York	Ohio	
20	Adrian	Murphy	Omni Consumer Products	Daly City	Arkansas	
21	Zoey	Gibson	Spade and Archer	Independence	Hawaii	
22	Alexandra	Dunn	Ankh-Sto Associates	Moreno Valley	Arkansas	
23	Julia	Wilson	Sto Plains Holdings	Salt Lake City	South Carolina	
24	Cal	Anderson	Waterford Dams	Salt Lake City	Oregon	

Page Layout view

Unlike the preview in Backstage view (choose File \Rightarrow Print), Page Layout view is not a view-only mode. You have complete access to all Excel commands. In fact, you can use Page Layout view all the time if you like.

Figure 7.2 shows a worksheet in Page Layout view, zoomed to about 40% to show multiple pages. Contrary to Normal view, you can see the margins, the page header and footer (if any), and even some space separating each page. If you've specified any repeated rows and columns, they are also displayed—giving you a true preview of the printed output.

TIP

If you move the mouse to the corner of a page while in Page Layout view, you can click to hide the white space in the margins. Doing so gives you all the advantages of Page Layout view, but you can see more information onscreen because the unused margin space is hidden.

FIGURE 7.2

In Page Layout view, the worksheet resembles printed pages.

The screenshot shows a Microsoft Excel spreadsheet in Page Layout view. The main area contains a large grid of data with columns A through J and rows 1 through 47. The data includes names like Gabriel, Darren, United Fried Chicken, and various cities and states. Several floating callout boxes are visible, each containing the text 'Add header' or 'Add footer'. One 'Add header' box is positioned above the first data row, and another is below it. There are also 'Add header' boxes on the right side of the screen, one above the second column and another above the third column. The 'Add footer' boxes are located at the bottom of the grid. The top of the screen shows a ribbon with tabs like Home, Insert, Page Layout, and View. The status bar at the bottom indicates 'Page 1 of 1'.

Page Break Preview

Page Break Preview displays the worksheet and the page breaks. Figure 7.3 shows an example. This view mode is different from Normal view mode with page breaks turned on. The key difference is that you can drag the page breaks. You can also drag the edges of the print area to change its size (if you've set a print area). Unlike Page Layout view, Page Break Preview does not display margins, headers, or footers.

Part I: Getting Started with Excel

FIGURE 7.3

Page Break Preview allows you to drag page breaks and print area borders.

	A	B	C	D	E	F	G
19	Tristan	Harris	Taggart Transcontinental	New York	Ohio		
20	Adrian	Murphy	Omni Consumer Products	Daly City	Arkansas		
21	Zoey	Gibson	Spade and Archer	Independence	Hawaii		
22	Alexandra	Dunn	Ankh-Sto Associates	Moreno Valley	Arkansas		
23	Julia	Wilson	Sto Plains Holdings	Salt Lake City	South Carolina		
24	Cole	Anderson	Water and Power	Salt Lake City	Georgia		
25	Brooklyn	Reynolds	Water and Power	Downey	Oregon		
26	Adrian	Richardson	Zev Toys	Augusta	Washington		
27	Adam	Coleman	Axis Chemical Co.	Hartford	Pennsylvania		
28	Brian	Myers	Three Waters	Montgomery	Minnesota		
29	Autumn	Rice	Sample Company	Yonkers	New York		
30	Benjamin	Washington	Virtucon	Lexington	Illinois		
31	Lauren	Cox	Gadgetron	Stockton	Arkansas		
32	Brody	Harris	Smith and Co.	Baltimore	Kansas		
33	Nathan	Henderson	General Services Corporation	Vallejo	Hawaii		
34	Evan	Ellis	Taco Grande	Roseville	Massachusetts		
35	Aiden	Mills	ABC Telecom	Winston-Salem	Iowa		
36	Brian	Harrison	Galaxy Corp	Huntington Beach	New Jersey		
37	Sara	Hunter	ZiffCorp	Antioch	Massachusetts		
38	Elizabeth	Dunn	Stay Puft Corporation	Carrollton	South Carolina		
39	Kaden	Jenkins	Cogswell Cogs	Aurora	Michigan		
40	Genesis	Stone	Transworld Consortium	Victorville	Michigan		
41	Landon	Gordon	The New Firm	Reno	New Jersey		
42	Hayden	Warren	Megadodo Publications	Virginia Beach	New York		
43	Ian	Holmes	Powell Motors	Billings	Wisconsin		
44	Connor	Nelson	Charles Townsend Agency	Oceanside	Alaska		
45	Emma	Grant	Mammoth Pictures	Vallejo	Montana		
46	Samuel	White	The New Firm	Pasadena	Arizona		
47	Madison	Collins	Sto Plains Holdings	Hialeah	New Hampshire		
48	Caroline	Carpenter	Acme Corp	Berkeley	District of Columbia		
49	Addison	Spencer	Incom Corporation	Fremont	Mississippi		
50	Bailey	Turner	ZiffCorp	Birmingham	Michigan		
51	Allison	Baker	Sonky Rubber Goods	Pittsburgh	Wisconsin		
52	Carson	Parker	Widget Corp	Killeen	Nevada		
53	Riley	Ruiz	Spacey Sprockets	Pueblo	Oklahoma		
54	Joseph	Stewart	Intrude	Tucson	Illinois		

When you enter Page Break Preview, Excel performs the following:

- Changes the zoom factor so that you can see more of the worksheet
- Displays the page numbers overlaid on the pages
- Displays the current print range with a white background; nonprinting areas appear with a gray background
- Displays all page breaks as draggable dashed lines

When you change the page breaks by dragging, Excel automatically adjusts the scaling so that the information fits on the pages, per your specifications.

Tip

In Page Break Preview, you still have access to all Excel commands. You can change the zoom factor if you find the text to be too small.

To exit Page Break Preview, just click one of the other View icons on the right side of the status bar.

Adjusting Common Page Setup Settings

Clicking the Quick Print button (or choosing File ➔ Print ➔ Print) may produce acceptable results in many cases, but a little tweaking of the print settings can often improve your printed reports. You can adjust print settings in three places:

- The Print settings screen in Backstage view, displayed when you choose File ➔ Print.
- The Page Layout tab of the Ribbon.
- The Page Setup dialog box is displayed when you click the dialog box launcher in the lower-right corner of the Page Layout ➔ Page Setup group on the Ribbon. You can also access the Page Setup dialog box from the Print settings screen in Backstage view.

Table 7.1 summarizes the locations where you can make various types of print-related adjustments in Excel.

TABLE 7.1 Where to Change Printer Settings

Setting	Print Settings Screen	Page Layout Tab of Ribbon	Page Setup Dialog Box
Number of copies	X		
Printer to use	X		
What to print	X		
Pages to print	X		
Specify worksheet print area		X	X
1-sided or 2-sided	X		
Collated	X		
Orientation	X	X	X
Paper size	X	X	X
Adjust margins	X	X	X
Specify manual page breaks		X	
Specify repeating rows or columns			X
Set print scaling	X	X	X
Print or hide gridlines		X	X
Print or hide row and column headings		X	X
Specify the first page number			X

Continues

TABLE 7.1 (continued)

Setting	Print Settings Screen	Page Layout Tab of Ribbon	Page Setup Dialog Box
Center output on page			X
Specify header/footers and options			X
Specify how to print cell notes or comments			X
Specify page order			X
Specify black-and-white output			X
Specify how to print error cells			X
Launch Printer Properties dialog box	X		X

Table 7.1 might make printing seem more complicated than it really is. The key point to remember is this: if you can't find a way to make a particular adjustment, it's probably available from the Page Setup dialog box.

Choosing your printer

To switch to a different printer or output device, choose File \Rightarrow Print, and use the drop-down control in the Printer section to select a different installed printer.

Note

To adjust printer settings, click the Printer Properties link on the Print settings screen in Backstage view to display a Properties dialog box for the selected printer. The exact dialog box that you see depends on the printer. The Properties dialog box lets you adjust printer-specific settings, such as the print quality and the paper source. In most cases, you won't have to change any of these settings, but if you're having print-related problems, you may want to check the settings.

Specifying what you want to print

Sometimes you may want to print only part of the worksheet rather than the entire used area. Or you may want to reprint selected pages of a report without printing all the pages. Choose File \Rightarrow Print and use the controls in the Settings section to specify what to print.

You have several options:

Print Active Sheets Prints the active sheet or sheets that you selected. (This option is the default.) You can select multiple sheets to print by pressing Ctrl and clicking the sheet tabs. If you select multiple sheets, Excel begins printing each sheet on a new page.

Print Entire Workbook Prints every sheet in the workbook, including chart sheets.

Print Selection Prints only the range that you selected before choosing File \Rightarrow Print.

Print Selected Chart Appears only if a chart is selected. If this option is chosen, only the chart will be printed.

Print Selected Table Appears only if the active cell is within a table (created by choosing Insert \Rightarrow Tables \Rightarrow Table) when the Print settings screen is displayed. If this option is chosen, only the table will be printed.

TIP

You can also choose Page Layout \Rightarrow Page Setup \Rightarrow Print Area \Rightarrow Set Print Area to specify the range(s) to print. Before you choose this command, select the range(s) that you want to print. To clear the print area, choose Page Layout \Rightarrow Page Setup \Rightarrow Print Area \Rightarrow Clear Print Area. To override the print area, select the Ignore Print Area check box in the list of Print What options.

NOTE

The print area does not have to be a single range. You can select multiple areas before you set the print area. Each area will print on a separate page.

If your printed output uses multiple pages, you can select which pages to print by indicating the number of the first and last pages to print by using Pages controls in the Settings section. You can either use the spinner controls or type the page numbers in the text boxes.

Changing page orientation

Page orientation refers to the way output is printed on the page. Choose Page Layout \Rightarrow Page Setup \Rightarrow Orientation \Rightarrow Portrait to print tall pages (the default) or Page Layout \Rightarrow Page Setup \Rightarrow Orientation \Rightarrow Landscape to print wide pages. Landscape orientation is useful when you have a wide range that doesn't fit on a vertically oriented page.

If you change the orientation, the onscreen page breaks adjust automatically to accommodate the new paper orientation.

Page orientation settings are also available when you choose File \Rightarrow Print.

Specifying paper size

Choose Page Layout \Rightarrow Page Setup \Rightarrow Size to specify the paper size you're using. The paper size settings are also available when you choose File \Rightarrow Print.

NOTE

Even though Excel displays a variety of paper sizes, your printer may not be capable of using all of them.

Printing multiple copies of a report

Use the Copies control at the top of the Print tab in Backstage view to specify the number of copies to print. Just enter the number of copies you want and then click Print.

TIP

If you're printing multiple copies of a report, make sure that the Collated option is selected so that Excel prints the pages in order for each set of output. If you're printing only one page, Excel ignores the Collated setting.

Adjusting the page margins

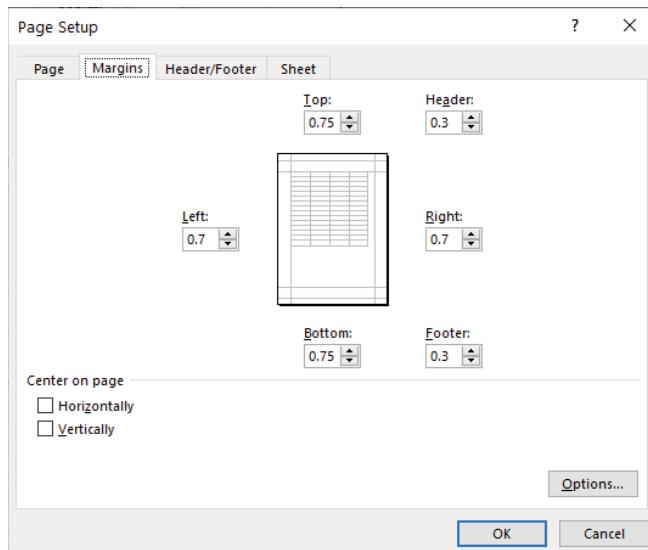
Margins are the unprinted areas along the sides, top, and bottom of a printed page. Excel provides four "quick margin" settings; you can also specify the exact margin size you require. All printed pages have the same margins. You can't specify different margins for different pages.

In Page Layout view, a ruler is displayed above the column headers and to the left of the row headers. Use your mouse to drag the margins in the ruler. Excel adjusts the page display immediately. Use the horizontal ruler to adjust the left and right margins and use the vertical ruler to adjust the top and bottom margins.

From the Page Layout \Rightarrow Page Setup \Rightarrow Margins drop-down list, you can select Normal, Wide, Narrow, or the Last Custom Setting if you previously customized the margins. These options are also available when you choose File \Rightarrow Print. If none of these settings does the job, choose Custom Margins to display the Margins tab of the Page Setup dialog box, as shown in Figure 7.4.

FIGURE 7.4

The Margins tab of the Page Setup dialog box



To change a margin, click the appropriate spinner (or you can enter a value directly). The margin settings that you specify in the Page Setup dialog box will then be available in the Page Layout \Rightarrow Page Setup \Rightarrow Margins drop-down list, referred to as Last Custom Setting.

NOTE

The Preview box in the center of the Page Setup dialog box is a bit deceiving because it doesn't really show you how your changes look in relation to the page; instead, it displays a darker line to let you know which margin you're adjusting.

You can also adjust margins in the preview window in Backstage view (choose File \Rightarrow Print). Click the Show Margins button in the bottom-right corner to display the margins in the preview pane. Then drag the margin indicators to adjust the margins.

In addition to the page margins, you can adjust the distance of the header from the top of the page and the distance of the footer from the bottom of the page. These settings should be less than the corresponding margin; otherwise, the header or footer may overlap with the printed output.

By default, Excel aligns the printed page at the top and left margins. If you want the output to be centered vertically or horizontally, select the appropriate check box in the Center on Page section of the Margins tab.

Understanding page breaks

When printing lengthy reports, controlling where pages break is often important. For example, you probably don't want a row to print on a page by itself, nor do you want a table header row to be the last line on a page. Fortunately, Excel gives you precise control over page breaks.

Excel handles page breaks automatically, but sometimes you may want to force a page break—either a vertical or a horizontal one—so that the report prints the way you want. For example, if your worksheet consists of several distinct sections, you may want to print each section on a separate sheet of paper.

Inserting a page break

To insert a horizontal page break line, select the cell that will begin the new page. Make sure you select a cell in column A, though; otherwise, you'll insert a vertical page break and a horizontal page break. For example, if you want row 14 to be the first row of a new page, select cell A14. Then choose Page Layout \Rightarrow Page Setup \Rightarrow Breaks \Rightarrow Insert Page Break.

NOTE

Page breaks are visualized differently, depending on which view mode you're using. (See "Changing Your Page View" earlier in this chapter.)

To insert a vertical page break line, select the cell that will begin the new page. In this case, though, make sure to select a cell in row 1. Choose **Page Layout** \Rightarrow **Page Setup** \Rightarrow **Breaks** \Rightarrow **Insert Page Break** to create the page break.

Removing manual page breaks

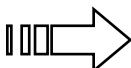
To remove a page break that you've added, select a cell in the first row beneath (or the first column to the right of) the manual page break and then choose **Page Layout** \Rightarrow **Page Setup** \Rightarrow **Breaks** \Rightarrow **Remove Page Break**.

To remove all manual page breaks in the worksheet, choose **Page Layout** \Rightarrow **Page Setup** \Rightarrow **Breaks** \Rightarrow **Reset All Page Breaks**.

Printing row and column titles

If your worksheet is set up with titles in the first row and descriptive names in the first column, it can be difficult to identify data that appears on printed pages where those titles don't appear. To resolve this problem, you can choose to print selected rows or columns as titles on each page of the printout.

Row and column titles serve pretty much the same purpose on a printout as frozen panes do in navigating within a worksheet. Keep in mind, however, that these features are independent of each other. In other words, freezing panes doesn't affect the printed output.



See Chapter 3, “Performing Basic Worksheet Operations,” for more information on freezing panes.

CAUTION

Don't confuse print titles with headers; these are two different concepts. Headers appear at the top of each page and contain information such as the worksheet name, date, or page number. Row and column titles describe the data being printed, such as field names in a database table or list.

You can specify rows to repeat at the top of every printed page or columns to repeat at the left of every printed page. To do so, choose **Page Layout** \Rightarrow **Page Setup** \Rightarrow **Print Titles**. Excel displays the **Sheet** tab of the **Page Setup** dialog box, as shown in Figure 7.5.

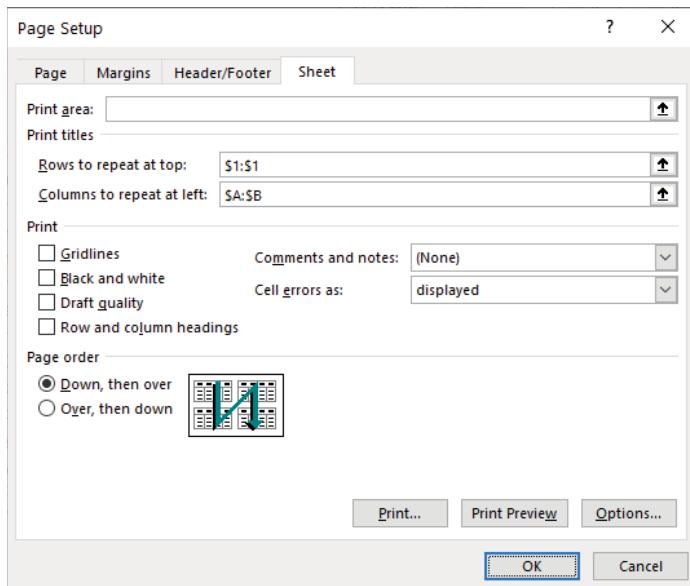
Figure 7.5 shows that row 1 will repeat at the top of the page and columns A and B will repeat on the left of the page. Even if you only want one row or column, you have to include the colon. To set these properties, activate the appropriate box (either **Rows to Repeat at Top** or **Columns to Repeat at Left**) and then select the rows or columns in the worksheet. Or you can enter these references manually. For example, to specify rows 1 and 2 as repeating rows, enter **1:2**.

NOTE

When you specify row and column titles and use **Page Layout** view, these titles will repeat on every page (just as when the document is printed). However, the cells used in the title can be selected only on the page in which they first appear.

FIGURE 7.5

Use the Sheet tab of the Page Setup dialog box to specify rows or columns that will appear on each printed page.



Scaling printed output

In some cases, you may need to force your printed output to fit on a specific number of pages. You can do so by enlarging or reducing the size. To enter a scaling factor, choose Page Layout \Rightarrow Scale to Fit \Rightarrow Scale. You can scale the output from 10% up to 400%. To return to normal scaling, enter **100%**.

To force Excel to print using a specific number of pages, choose Page Layout \Rightarrow Scale to Fit \Rightarrow Width and Page Layout \Rightarrow Scale to Fit \Rightarrow Height. When you change either one of these settings, the corresponding scale factor is displayed in the Scale control.

CAUTION

Excel doesn't ensure legibility. It will gladly scale your output to be so small that no one can read it.

Printing cell gridlines

Typically, cell gridlines aren't printed. If you want your printout to include the gridlines, choose Page Layout \Rightarrow Sheet Options \Rightarrow Gridlines \Rightarrow Print.

Alternatively, you can insert borders around some cells to simulate gridlines. Change the border color to White, Background 1, 25% Darker to get a pretty good simulation of a gridline.

To change the color, choose Home \Rightarrow Font \Rightarrow Borders \Rightarrow More Borders. Make sure you change the color before you apply the border.



See Chapter 5, “Formatting Worksheets,” for information about borders.

Printing row and column headers

By default, row and column headers for a worksheet are not printed. If you want your printout to include these items, choose Page Layout \Rightarrow Sheet Options \Rightarrow Headings \Rightarrow Print.

Using a background image

Would you like to have a background image on your printouts? Unfortunately, you can't. You may have noticed the Page Layout \Rightarrow Page Setup \Rightarrow Background command. This button displays a dialog box that lets you select an image to display as a background. Placing this control among the other print-related commands is misleading. Background images placed on a worksheet are never printed.

TIP

In lieu of a true background image, you can insert WordArt, a shape, or a picture on your worksheet and then adjust its transparency. Then copy the image to all printed pages. Alternatively, you can insert an object in a page header or footer. (See the sidebar “Inserting a Watermark.”)

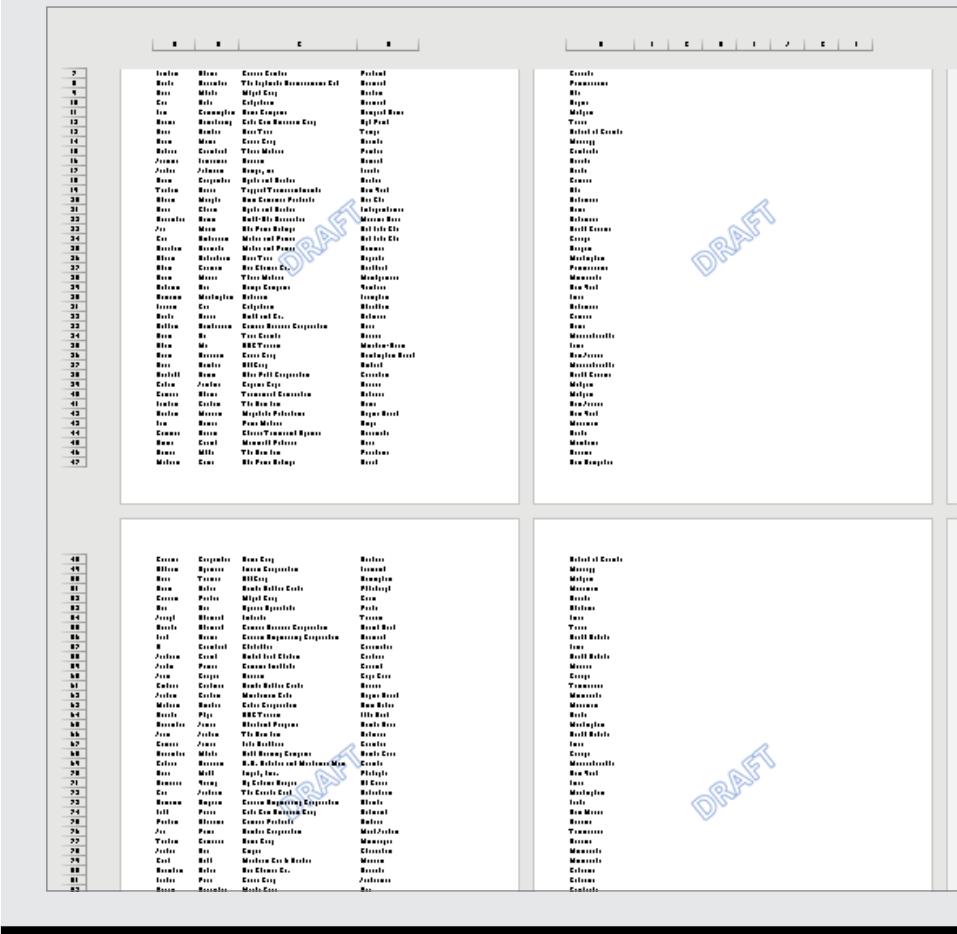
Inserting a Watermark

A watermark is an image (or text) that appears on each printed page. A watermark can be a faint company logo or a word such as DRAFT. Excel doesn't have an official command to print a watermark, but you can add a watermark by inserting a picture in the page header or footer. Here's how:

1. Locate an image on your hard drive that you want to use for the watermark.
2. Choose View \Rightarrow Workbook Views \Rightarrow Page Layout View.
3. Click the center section of the header.
4. Choose Header & Footer \Rightarrow Header & Footer Elements \Rightarrow Picture. The Insert Pictures dialog box appears.
5. Click Browse and locate the image from step 1 (or locate a suitable image from other sources listed).
6. Click outside the header to see your image.
7. To center the image in the middle of the page, click the center section of the header and add some carriage returns before the `&[Picture]` code. You'll need to experiment to determine the number of carriage returns required to push the image into the body of the document.

- 8. If you need to adjust the image (for example, make it lighter), click the center section of the header and then choose Header & Footer ➔ Header & Footer Elements ➔ Format Picture. Use the Image controls in the Picture tab of the Format Picture dialog box to adjust the image.** You may need to experiment with the settings to make sure that the worksheet text is legible.

The accompanying figure shows an example of a header image (the word DRAFT made with SmartArt) used as a watermark. You can do a similar thing with plain text, but you don't get the same formatting controls, such as controlling the brightness and contrast.



Adding a Header or a Footer to Your Reports

A header is information that appears at the top of each printed page. A footer is information that appears at the bottom of each printed page. By default, new workbooks have space for headers or footers but nothing in them.

You can specify headers and footers by using the Header/Footer tab of the Page Setup dialog box. Or you can simplify the task by switching to Page Layout view, where you can click the section labeled Add Header or Add Footer.

Note

If you're working in Normal view, you can choose Insert \Rightarrow Text \Rightarrow Header & Footer. Excel switches to Page Layout view and activates the center section of the page header.

You can then type the information and apply any type of formatting you like. Note that headers and footers consist of three sections: left, center, and right. For example, you can create a header that prints your name at the left margin, the worksheet name centered in the header, and the page number at the right margin.

TIP

If you want a consistent header or footer for all your documents, create a book .xltx template with your headers or footers specified. A book .xltx template is used as the basis for new workbooks.



See Chapter 6, “Understanding Excel Files and Templates,” for details on creating a template.

When you activate the header or footer section in Page Layout view, the Ribbon displays a new contextual tab: Header & Footer. Use the controls on this tab to work with headers and footers.

Selecting a predefined header or footer

You can choose from a number of predefined headers or footers by using either of the two drop-down lists in the Header & Footer \Rightarrow Header & Footer group. Notice that some items in these lists consist of multiple parts, separated by a comma. Each part goes into one of the three header or footer sections (left, center, or right). Figure 7.6 shows an example of a header that uses all three sections.

Understanding header and footer element codes

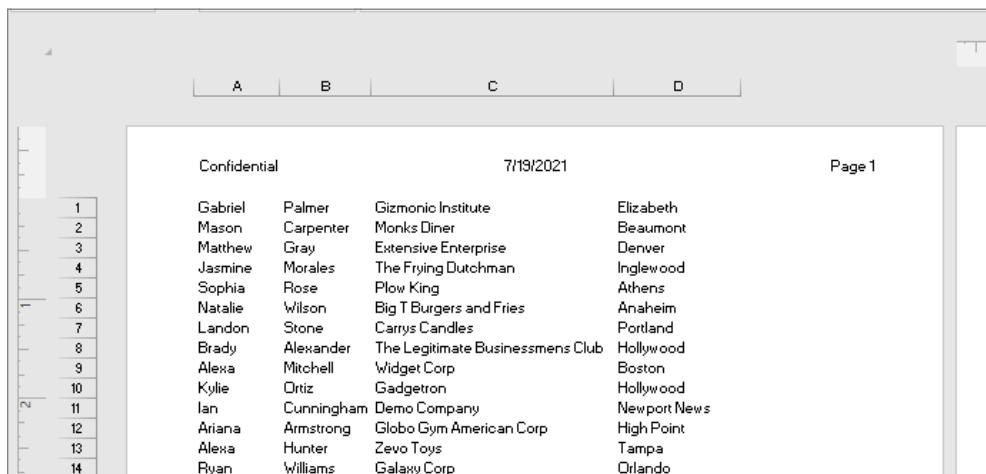
When a header or footer section is activated, you can type whatever text you like into the section. Or to insert variable information, you can insert any of several element codes by clicking a button in the Header & Footer \Rightarrow Header & Footer Elements group. Each button inserts a code into the selected section. For example, to insert the current date, click the Current Date button. Table 7.2 lists the buttons and their functions.

TABLE 7.2 Header and Footer Buttons and Their Functions

Button	Code	Function
Page Number	& [Page]	Displays the page number
Number of Pages	& [Pages]	Displays the total number of pages to be printed
Current Date	& [Date]	Displays the current date
Current Time	& [Time]	Displays the current time
File Path	& [Path] & [File]	Displays the workbook's complete path and filename
File Name	& [File]	Displays the workbook name
Sheet Name	& [Tab]	Displays the sheet's name
Picture	& [Picture]	Enables you to add a picture
Format Picture	Not applicable	Enables you to change an added picture's settings

FIGURE 7.6

This three-part header is one of Excel's predefined headers.



You can combine text and codes and insert as many codes as you like into each section.

Note

If the text that you enter uses an ampersand (&), you must enter the ampersand twice (because Excel uses an ampersand to signal a code). For example, to enter the text Research & Development into a section of a header or footer, type Research && Development.

You can also use different fonts and sizes in your headers and footers. Just select the text you want to change and then use the formatting tools in the Home \Rightarrow Font group. Or use the controls on the Mini toolbar, which appears automatically when you select the text. If you don't change the font, Excel uses the font defined for the Normal style.

TIP

You can use as many lines as you like. Press Enter to force a line break for multiline headers or footers. If you use multiline headers or footers, you may need to adjust the top or bottom margin so that the text won't overlap with the worksheet data. (See "Adjusting the page margins" earlier in this chapter.)

Unfortunately, you can't print the contents of a specific cell in a header or footer. For example, you may want Excel to use the contents of cell A1 as part of a header. To do so, you need to enter the cell's contents manually—or write a VBA macro to perform this operation before the sheet is printed.

Exploring other header and footer options

When a header or footer is selected in Page Layout view, the Header & Footer \Rightarrow Options group contains controls that let you specify other options:

Different First Page If this option is checked, you can specify a different header/footer for the first printed page.

Different Odd & Even Pages If this option is checked, you can specify a different header/footer for odd and even pages.

Scale with Document If this option is checked, the font size in the header and footer will be sized accordingly if the document is scaled when printed. This option is enabled by default.

Align with Page Margins If this option is checked, the left header and footer will be aligned with the left margin, and the right header and footer will be aligned with the right margin. This option is enabled by default.

NOTE

If you check either of the Different First Page or Different Odd & Even Pages check boxes, you can no longer use the predefined headers and footers. You must use the Custom Header and Custom Footer buttons in the Page Setup dialog box.

Exploring Other Print-Related Topics

The following sections cover some additional topics related to printing from Excel.

Copying Page Setup settings across sheets

Each Excel worksheet has its own print setup options (orientation, margins, headers and footers, and so on). These options are specified in the Page Setup group of the Page Layout tab.

When you add a new sheet to a workbook, it contains the default page setup settings. Here's an easy way to transfer the settings from one worksheet to additional worksheets:

1. **Activate the sheet that contains the desired setup information.** This is the source sheet.
2. **Select the target sheets, and Ctrl+click the sheet tabs of the sheets that you want to update with the settings from the source sheet.**
3. **Click the dialog box launcher in the lower-right corner of the Page Layout ↴ Page Setup group.**
4. **When the Page Setup dialog box appears, click OK to close it.**
5. **Ungroup the sheets by right-clicking any selected sheet and choosing Ungroup Sheets from the shortcut menu.** Because multiple sheets are selected when you close the Page Setup dialog box, the settings of the source sheet will be transferred to all target sheets.

Note

Two settings located on the Sheet tab of the Page Setup dialog box are not transferred: Print Area and Print Titles. In addition, pictures in the header or footer are not transferred.

Preventing certain cells from being printed

If your worksheet contains confidential information, you may want to print the worksheet but not the confidential parts. You can use several techniques to prevent certain parts of a worksheet from printing:

Hide rows or columns When you hide rows or columns, the hidden rows or columns aren't printed. Choose Home ↴ Cells ↴ Format drop-down list to hide the selected rows or columns.

Hide cells or ranges by making the text color the same color as the background color You can hide cells or ranges by making the text color the same color as the background color. Be aware, however, that this method may not work for all printers.

Hide cells or ranges by using a custom number format You can hide cells by using a custom number format that consists of three semicolons (; ; ;). See Chapter 2, "Entering and Editing Worksheet Data," for more information about using custom number formats.

Mask an area You can mask a confidential area of a worksheet by covering it with a rectangle shape. Choose Insert ↴ Illustrations ↴ Shapes and click the Rectangle shape. You'll probably want to adjust the fill color to match the cell background and remove the border.

If you find that you must regularly hide data before you print certain reports, consider using the Custom Views feature, discussed later in this chapter. (See “Creating custom views of your worksheet.”) This feature allows you to create a named view that doesn’t show the confidential information.

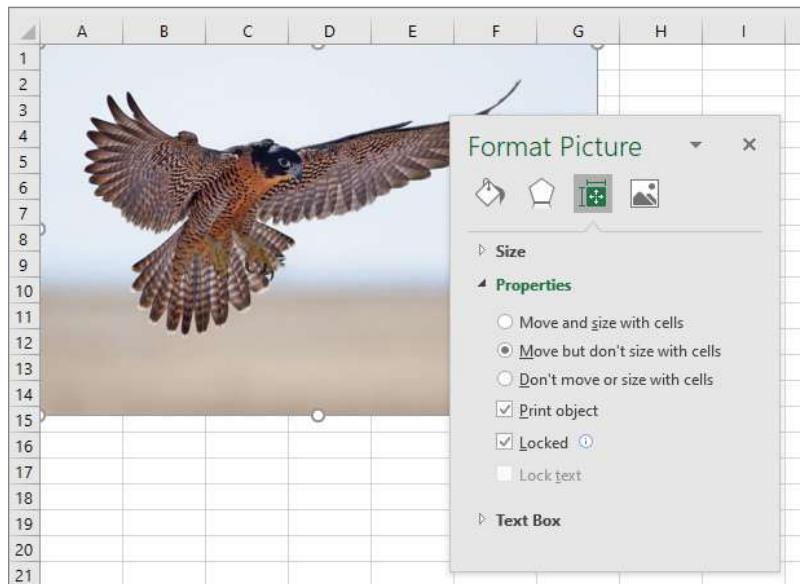
Preventing objects from being printed

To prevent objects on the worksheet (such as charts, shapes, and SmartArt) from being printed, you need to access the Properties tab of the object’s Format task pane (see Figure 7.7):

1. Right-click the object and choose Format xxxx from the shortcut menu. (xxxx varies, depending on the object.)
2. In the Format task pane that opens for the object, click the Size & Properties icon.
3. Expand the Properties section of the task pane.
4. Remove the check mark for Print Object.

FIGURE 7.7

Use the Properties section in the object’s Format task pane to prevent objects from printing.



Note

For a chart, you must right-click the chart's chart area (the background of the chart). Or double-click the chart's border to display the Format Chart Area task pane. Then expand the Properties section and remove the check mark from Print Object.

Creating custom views of your worksheet

If you need to create several different printed reports from the same Excel workbook, setting up the specific settings for each report can be a tedious job. For example, you may need to print a full report in landscape mode for your boss. Another department may require a simplified report using the same data but with some hidden columns in portrait mode. You can simplify the process by creating custom named views of your worksheets that include the proper settings for each report.

The Custom Views feature enables you to give names to various views of your worksheet. You can quickly switch among these named views. A view includes settings for the following:

- Print settings, as specified in the Page Layout \Rightarrow Page Setup, Page Layout \Rightarrow Scale to Fit, and Page Layout \Rightarrow Sheet Options groups
- Hidden rows and columns
- The worksheet view (Normal, Page Layout, Page Break Preview)
- Selected cells and ranges
- The active cell
- The zoom factor
- Window sizes and positions
- Frozen panes

If you find that you're constantly fiddling with these settings before printing and then changing them back, using custom named views can save you some work.

CAUTION

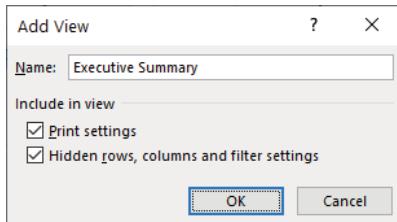
Unfortunately, the Custom Views feature doesn't work if the workbook (not just the worksheet) contains at least one table, created using Insert \Rightarrow Tables \Rightarrow Table. When a workbook that contains a table is active, the Custom Views command is disabled. This severely limits the usefulness of the Custom Views feature.

To create a named view, follow these steps:

1. **Set up the view settings the way you want them.** For example, hide some columns.
2. **Choose View \Rightarrow Workbook Views \Rightarrow Custom Views.** The Custom Views dialog box appears.
3. **Click the Add button.** The Add View dialog box (shown in Figure 7.8) appears.

FIGURE 7.8

Use the Add View dialog box to create a named view.



4. **Provide a descriptive name.** You can also specify what to include in the view by using the two check boxes. For example, if you don't want the view to include print settings, remove the check mark from Print Settings.
5. **Click OK to save the named view.**

Then when you're ready to print, open the Custom Views dialog box to see all named views. To select a particular view, just select it from the list and click the Show button. To delete a named view from the list, click the Delete button.

Creating PDF files

The PDF file format is widely used as a way to present information in a read-only manner with precise control over the layout. If you need to share your work with someone who doesn't have Excel, creating a PDF is often a good solution. Free software to display PDFs is available from a number of sources.

Note

Excel can create PDFs, but it can't open them. Word can create and open PDFs.

XPS is another “electronic paper” format, developed by Microsoft as an alternative to the PDF format. Currently, there is little third-party support for the XPS format.

To save a worksheet in PDF or XPS format, choose File \Rightarrow Export \Rightarrow Create PDF/XPS Document \Rightarrow Create PDF/XPS. Excel displays the Publish as PDF or XPS dialog box, in which you can specify a filename and location and set some other options.



Customizing the Excel User Interface

IN THIS CHAPTER

Customizing the Quick Access Toolbar

Customizing the Ribbon

A software program's user interface consists of all the ways that the user interacts with the software. In Excel, the user interface consists of the following:

- The Ribbon
- The Quick Access Toolbar
- Right-click shortcut menus
- Dialog boxes
- Task panes
- Keyboard shortcuts

This chapter describes how to make changes to two Excel user interface components: the Quick Access Toolbar and the Ribbon. You might want to customize these elements to make Excel more suited to the way you use it.

Customizing the Quick Access Toolbar

The Quick Access Toolbar is always visible, regardless of which Ribbon tab is selected. After you customize the Quick Access Toolbar, your frequently used commands will always be one click away.

NOTE

The only situation in which the Quick Access Toolbar is not visible is in full-screen mode, which is enabled by clicking the Ribbon Display Options button in the Excel title bar and choosing Auto-Hide Ribbon. To display the Quick Access Toolbar (and Ribbon) temporarily in full-screen mode, click the title bar or press the Alt key. To cancel full-screen mode, click the Ribbon Display Options button in the Excel title bar and choose Show Tabs or Show Tabs and Commands.

About the Quick Access Toolbar

By default, the Quick Access Toolbar is located on the left side of the Excel title bar, above the Ribbon (see Figure 8.1). Unless you customize it, this toolbar includes four tools; five if Excel detects you're on a touch-enabled device.

AutoSave Periodically saves workbooks that have already been saved to OneDrive or SharePoint.

Save Saves the active workbook.

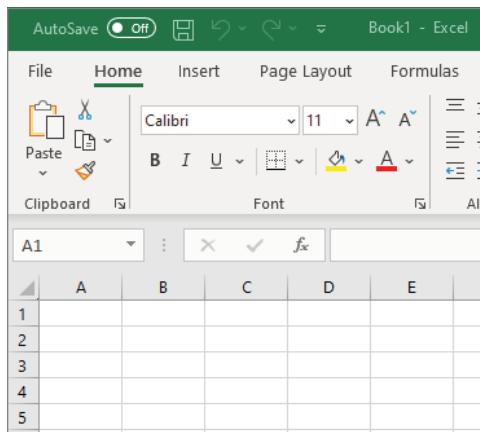
Undo Reverses the effect of the last action.

Redo Reverses the effect of the last undo.

Touch/Mouse Mode For touch-enabled devices, toggles between touch mode and mouse mode. Touch mode provides larger touch targets.

FIGURE 8.1

The default location for the Quick Access Toolbar is on the left side of the Excel title bar.



If you prefer, you can move the Quick Access Toolbar below the Ribbon. To do so, right-click the Quick Access Toolbar and choose Show Quick Access Toolbar Below the Ribbon. Moving the Quick Access Toolbar below the Ribbon uses additional vertical space on your screen. In other words, you'll be able to see one or two fewer rows of your worksheet if you move the Quick Access Toolbar from its default location. Unlike traditional toolbars, the Quick Access Toolbar cannot be made free floating, so you can't move it to a convenient location. It always appears either above or below the Ribbon.

Commands on the Quick Access Toolbar always appear as small icons with no text. An exception to this rule is drop-down controls that display text. For example, if you add the Font control from the Home \leftrightarrow Font group, it appears as a drop-down control in the Quick Access

Toolbar. When you hover your mouse pointer over an icon, you see a ScreenTip showing the name of the command and a brief description.

You can customize the Quick Access Toolbar by adding or removing commands. If you find that you use some Excel commands frequently, you can make these commands easily accessible by adding them to your Quick Access Toolbar. You can also rearrange the order of the icons.

As far as we can tell, there is no limit to the number of commands that you can add. The Quick Access Toolbar always displays only a single line of icons. If the number of icons exceeds the Excel window width, it displays an additional icon at the end: More Controls. Click the More Controls icon, and the hidden Quick Access Toolbar icons appear in a pop-up window.

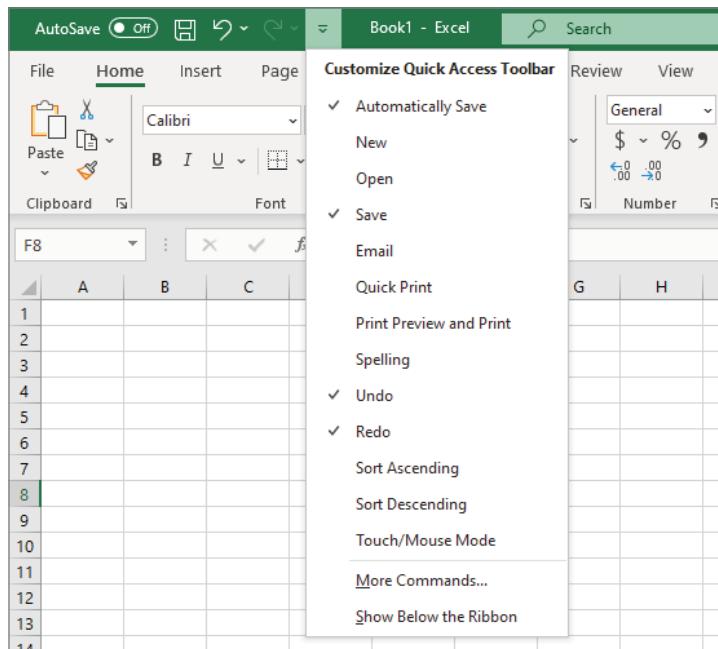
Adding new commands to the Quick Access Toolbar

You can add a new command to the Quick Access Toolbar in three ways:

- **Click the Customize Quick Access Toolbar drop-down control, which is located on the right side of the Quick Access Toolbar (see Figure 8.2).** The list contains a few commonly used commands. Select a command from the list, and Excel adds it to your Quick Access Toolbar.

FIGURE 8.2

This drop-down list is one way to add a new command to the Quick Access Toolbar.

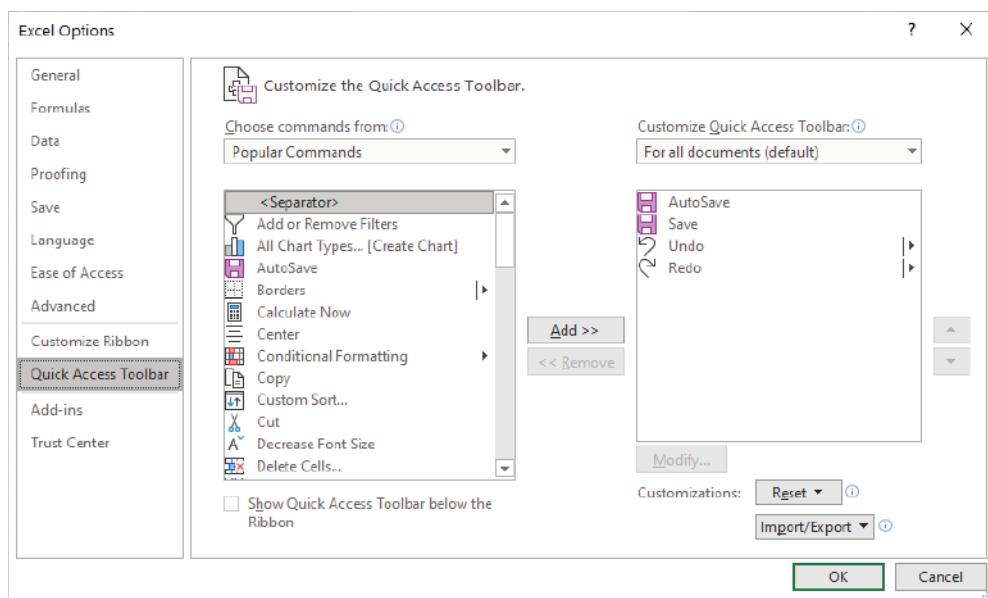


- **Right-click any control on the Ribbon and choose Add to Quick Access Toolbar.** The control is added to your Quick Access Toolbar, positioned to the right of the last control.
- **Use the Quick Access Toolbar tab in the Excel Options dialog box.** A quick way to access this dialog box is to right-click the Quick Access Toolbar and choose Customize Quick Access Toolbar.

The remainder of this section discusses the Quick Access Toolbar tab of the Excel Options dialog box, shown in Figure 8.3.

FIGURE 8.3

Use the Quick Access Toolbar tab in the Excel Options dialog box to customize the Quick Access Toolbar.



The left side of the dialog box displays a list of Excel commands, and the right side shows the commands currently on your Quick Access Toolbar. Above the command list on the left is the Choose Commands From drop-down list, from which you can filter the list. Select an item from the drop-down list, and the list displays only the commands for that item. In Figure 8.3, the list shows commands in the Popular Commands category.

Some of the items in the drop-down list are as follows:

Popular Commands Displays commands that Excel users commonly use.

Commands Not in the Ribbon Displays a list of commands that you can't access from the Ribbon. Many, but not all, are obsolete or not very useful.

All Commands Displays a complete list of Excel commands.

Macros Displays a list of all available macros.

File Tab Displays the commands available in the Backstage view.

Home Tab Displays all commands available when the Home tab is active.

In addition, the Choose Commands From drop-down list contains items for every other tab, including the contextual tabs (for example, the additional tabs that are displayed when a chart is selected). To add an item to your Quick Access Toolbar, select it from the list on the left and then click Add. The command appears in the list on the right. At the top of each list is an item called <Separator>. Adding this item to your Quick Access Toolbar results in a vertical bar to help you group commands.

The commands are listed in alphabetical order. Sometimes you need to do some guessing to find a particular command.

TIP

By default, Quick Access Toolbar customizations are visible for all documents. You can create a Quick Access Toolbar configuration that's specific to a particular workbook. In other words, the commands on the Quick Access Toolbar appear only when a particular workbook is active. Start by activating the workbook and then display the Quick Access Toolbar tab of the Excel Options dialog box. When you add a command to the Quick Access Toolbar, use the drop-down list in the upper right to specify whether the change is for all workbooks or just the active workbook.

8

When you select Macros from the Choose Commands From drop-down, Excel lists all available macros. You can attach a macro to a Quick Access Toolbar icon so that when you click the icon, the macro is executed. If you add a macro to your Quick Access Toolbar, you can click the Modify button to change the text and choose a different icon for the macro.

When you finish making your Quick Access Toolbar customizations, click OK to close the Excel Options dialog box. The new icon(s) will appear on the Quick Access Toolbar.

TIP

The only time you ever need to use the Quick Access Toolbar tab of the Excel Options dialog box is when you want to add a command that's not on the Ribbon, add a command that executes a macro, or rearrange the order of the icons. In all other situations, it's much easier to locate the command in the Ribbon, right-click the command, and choose Add to Quick Access Toolbar.

Other Quick Access Toolbar actions

Other Quick Access Toolbar actions include the following:

Rearranging the Quick Access Toolbar icons If you want to change the order of your Quick Access Toolbar icons, you can do so from the Quick Access Toolbar tab of the Excel Options dialog box. Select the command and then use the Move Up and Move Down arrow buttons on the right to move the icon.

Removing Quick Access Toolbar icons The easiest way to remove an icon from your Quick Access Toolbar is to right-click the icon and choose Remove from Quick Access Toolbar. You can also use the Quick Access Toolbar tab of the Excel Options dialog box. Just select the command in the list on the right and click Remove.

Resetting the Quick Access Toolbar If you want to return the Quick Access Toolbar to its default state, display the Quick Access Toolbar tab of the Excel Options dialog box and click the Reset button. Then choose Reset Only Quick Access Toolbar. The Quick Access Toolbar then displays its default commands.

CAUTION

You can't undo resetting the Quick Access Toolbar.

Sharing User Interface Customizations

In the Excel Options dialog box, the Quick Access Toolbar tab and the Customize Ribbon tab both have an Import/Export button. You can use this button to save and open files that contain user interface customizations. For example, you might create a new Ribbon tab and want to share it with your officemates.

Click the Import/Export button, and you get two options:

Import Customization File You're prompted to locate the file. Before you load a file, you're asked whether you want to replace all existing Ribbon and Quick Access Toolbar customizations.

Export All Customizations You're prompted to provide a filename and location for the file. The information is stored in a file that has an `exportedUI` filename extension.

Unfortunately, importing and exporting are not implemented very well. Excel doesn't allow you to save or load only the Quick Access Toolbar customizations or only the Ribbon customizations. Both types of customizations are exported and imported. Therefore, you can't share your Quick Access Toolbar customizations without also sharing your Ribbon customizations.

TIP

Commands on the Quick Access Toolbar are assigned numbers for shortcut keys. For example, the shortcut Alt+1 executes the first command on the Quick Access Toolbar. After the ninth command, the shortcuts change to 09, 08, 07.... After the 18th command, the shortcuts change to 0M, 0N, 0O.... And after Alt+OZ, Excel stops assigning shortcuts.

Customizing the Ribbon

The Ribbon is Excel's primary user interface component. It consists of tabs along the top. When you click a tab, it displays a set of relevant commands, and the commands are arranged in groups.

Why you may want to customize the Ribbon

Most users have no need to customize the Ribbon. If you find that you tend to use the same command over and over, though—and you’re constantly clicking tabs to access these commands—you might benefit from customizing the Ribbon in such a way that the commands you need are on the same tab.

What can be customized

You can customize tabs on the Ribbon in the following ways:

- Adding new custom tabs
- Deleting custom tabs
- Changing the order of the tabs
- Changing the name of tabs
- Hiding built-in tabs

You can customize groups on the Ribbon in the following ways:

- Adding new custom groups
- Adding commands to custom groups
- Removing commands from custom groups
- Removing groups from tabs
- Moving a group to a different tab
- Changing the order of the groups within a tab
- Changing the name of a group

Those are fairly comprehensive lists of customization options, but there are some actions that you can’t do:

- Remove built-in tabs (but you can hide them)
- Remove specific commands from built-in groups (but you *can* remove entire groups)
- Change the order of commands in a built-in group

Note

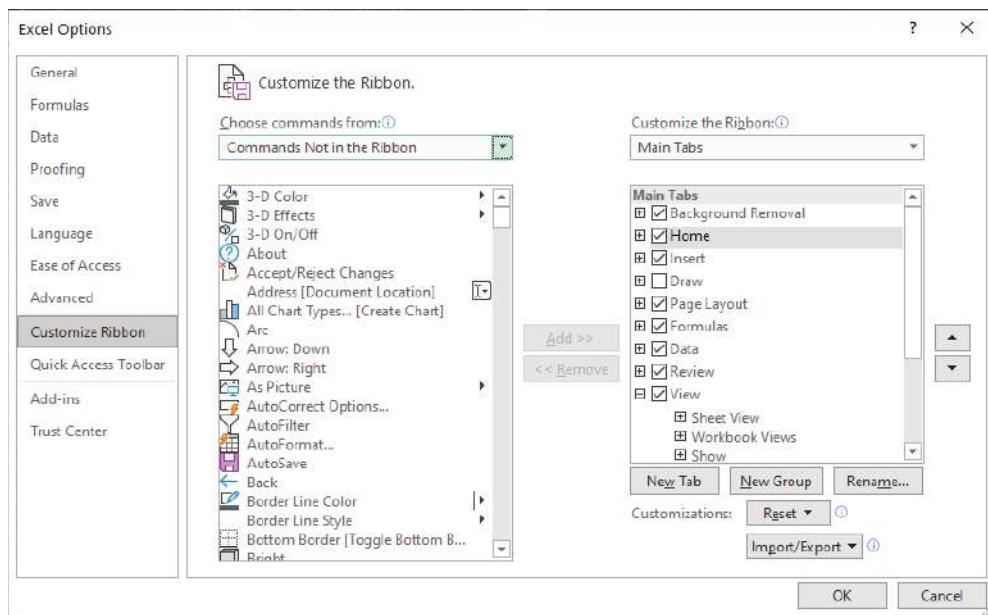
Unfortunately, you can’t use VBA macros to customize the Ribbon (or Quick Access Toolbar). However, developers can write RibbonX code and store it in workbook files. When the file is opened, the Ribbon is modified to display new commands. Writing RibbonX code is relatively complicated and beyond the scope of this book.

How to customize the Ribbon

Customizing the Ribbon is done via the Customize Ribbon tab of the Excel Options dialog box (see Figure 8.4). The quickest way to display this dialog box is to right-click anywhere on the Ribbon and choose Customize the Ribbon.

FIGURE 8.4

The Customize Ribbon tab of the Excel Options dialog box



Creating a new tab

If you'd like to create a new tab, click the New Tab button. Excel creates a tab named New Tab (Custom) and a new group in the tab named New Group (Custom).

You'll almost always want to give the tab (and the group) better names. Select the item and click Rename. Use the Move Up and Move Down arrow buttons on the right to reposition the new tab, if necessary.

NOTE

You don't need to add a new tab just because you want to add new commands to the Ribbon. You can create a new group for an existing tab.

Creating a new group

To create a new group, select the tab that will hold the new group and click the New Group button. Excel creates a new group named New Group (Custom). Use the Rename button to provide a more descriptive name, and use the Move Up and Move Down arrow buttons on the right to reposition the group within the tab.

Adding commands to a new group

Adding commands to the Ribbon is similar to adding commands to the Quick Access Toolbar, which we described earlier in this chapter. Commands that you add must be placed in a new group. Here's the general procedure:

1. Use the **Choose Commands From** drop-down list on the left to display various groups of commands.
2. Select the command in the list box on the left.
3. Use the **Customize the Ribbon** drop-down list on the right to choose a group of tabs. Main Tabs refers to the tabs that are always visible; Tool Tabs refers to the contextual tabs that appear when a particular object is selected.
4. In the list box on the right, select the tab and the group where you want to put the command. You'll need to click the plus-sign control to expand the tab name so that it displays its group names.

NOTE

You can add commands only to groups that you've created.

5. Click the **Add** button to add the selected command from the left to the selected group on the right.

To rearrange the order of tabs, groups, or commands, select the item and use the Move Up and Move Down buttons on the right. Note that you can move a group into a different tab.

8

NOTE

Although you can't remove a built-in tab, you can hide the tab by clearing the check box next to its name.

Figure 8.5 shows a part of a customized Ribbon. In this case, two groups were added to the View tab (to the right of the Macros group): Extra Commands (with three new commands) and Speech (with five new commands).

FIGURE 8.5

The View tab, with two new groups added



Resetting the Ribbon

To restore all or part of the Ribbon to its default state, right-click any part of the Ribbon and choose Customize the Ribbon from the shortcut menu. Excel displays the Customize Ribbon tab of the Excel Options dialog box. Click the Reset button to display two options: Reset Only Selected Ribbon Tab and Reset All Customizations. If you choose the latter, the Ribbon will be returned to its default state, and you'll lose any Quick Access Toolbar customizations that you made.

Part II

Working with Formulas and Functions

Formulas and worksheet functions are essential to manipulating data and obtaining useful information from your Excel workbooks. The chapters in this part start with a solid foundation in using Excel formulas and gradually introduce more sophisticated concepts such as the new dynamic array functionality in Excel.

IN THIS PART

Chapter 9

Introducing Formulas and Functions

Chapter 10

Understanding and Using Array Formulas

Chapter 11

Using Formulas for Common Mathematical Operations

Chapter 12

Using Formulas to Manipulate Text

Chapter 13

Using Formulas with Dates and Times

Chapter 14

Using Formulas for Conditional Analysis

Chapter 15

Using Formulas for Matching and Lookups

Chapter 16

Using Formulas with Tables and Conditional Formatting

Chapter 17

Making Your Formulas Error-Free

Introducing Formulas and Functions

IN THIS CHAPTER

- Understanding formula basics
- Entering formulas and functions into your worksheets
- Understanding how to use cell references in formulas
- Using formula variables
- Building formulas with Excel table objects
- Correcting common formula errors
- Getting tips for working with formulas

Excel is, at its core, a calculation engine. Like a calculator, Excel accepts questions in the form of a formula (such as $=2+2$) and returns an answer. Formulas allow you to perform not only mathematical operations, but also a myriad of other complex actions. With formulas, you can parse textual values, look up data based on specific criteria, and perform conditional analysis. If it weren't for formulas, a spreadsheet would simply be a fancy word processing document that has great support for tabular information.

If you want to leverage the full power of Excel, it's important to understand how formulas work. This chapter introduces formulas and functions and helps you to get up to speed with writing your own formulas and functions.

Understanding Formula Basics

Each cell in a worksheet is given a name based on its position. The top-leftmost cell is called A1. It's located in column A, row 1. When you enter a value in cell A1 (for example, the number 5), that cell's value becomes 5. Entering the number 10 in cell B1 will make that cell's value 10. You can then use these values in a formula.

For instance, you can click in cell C1, and begin typing `=A1+B1`. After you press the Enter key, Excel recognizes what you're asking and performs the calculation that gives you the result of 15 ($5+10=15$).

Part II: Working with Formulas and Functions

A formula always begins with an equal sign and can contain any of these elements:

- **Constants:** Hard-coded numbers can be used directly in a formula. For example, you can enter `=5+10` directly into a cell to get the answer 15.
- **Operators:** These include symbols that designate mathematical operations such as `+` (for addition) and `*` (for multiplication). Some operators compare values (`>`, `<`, `=`), while others join values together (`&`). As you may have guessed, you can mix constants and operators to create a formula. For instance, entering `=15>10` into a cell would return TRUE as the result because 15 is indeed greater than 10.
- **Cell references:** These include any value that points back to a single cell or a range of cells. When you enter `=A1+B1` into a cell, you are essentially using two cell references. Cell references tell Excel to use the values in the cells being referenced in the formula.
- **Text strings:** Any text string can be used as an argument in a formula as long as it's wrapped in quotes. For instance, entering `="Microsoft"&"Excel"` in any cell will result in the joined text Microsoft Excel.
- **Worksheet functions:** You can use a virtually unlimited combination of Excel functions (such as `SUM` and `AVERAGE`) to build your formulas. For example, clicking on cell A1 and entering `=SUM(A1:A10)` will add all the values in cells A1 through A10.

After you enter a formula, the cell containing your formula displays the calculated result. While your formula cell displays the result, the Formula bar will always show the formula.

Here are a few examples of formulas:

<code>=150*.05</code>	Multiples 150 times 0.05. This formula uses only values, and it always returns the same result. You could just enter the value 7.5 into the cell, but using a formula provides information on how the value was calculated.
<code>=A3</code>	Displays the value from cell A3. No calculation is performed on A3.
<code>=A1+A2</code>	Adds the values in cells A1 and A2.
<code>=Income-Expenses</code>	Subtracts the value in the cell named Expenses from the value in the cell named Income.
<code>=SUM(A1:A12)</code>	Adds the values in the range A1:A12, using the <code>SUM</code> function.
<code>=A1=C12</code>	Compares cell A1 with cell C12. If the cells are the same, the formula returns TRUE; otherwise, it returns FALSE.

Again, note that every formula begins with an equal sign (`=`). The initial equal sign allows Excel to distinguish a formula from plain text.

Using operators in formulas

Excel formulas support a variety of operators. Operators are symbols that indicate what mathematical (or logical) operation you want the formula to perform. Table 9.1 lists the operators that Excel recognizes. In addition to these, Excel has many built-in functions that enable you to perform additional calculations.

TABLE 9.1 Operators Used in Formulas

Operator	Name
+	Addition
-	Subtraction
*	Multiplication
/	Division
^	Exponentiation
&	Concatenation
=	Logical comparison (equal to)
>	Logical comparison (greater than)
<	Logical comparison (less than)
>=	Logical comparison (greater than or equal to)
<=	Logical comparison (less than or equal to)
<>	Logical comparison (not equal to)

You can, of course, use as many operators as you need to perform the desired calculation.

Here are some examples of formulas that use various operators:

Formula	What It Does
= "Part-" & "23A"	Joins (concatenates) the two text strings to produce Part-23A.
=A1&A2	Concatenates the contents of cell A1 with cell A2. Concatenation works with values as well as text. If cell A1 contains 123 and cell A2 contains 456, this formula would return the text 123456. Note that the result of concatenation is always formatted as text.
=6^3	Raises 6 to the third power (216).
=216^(1/3)	Raises 216 to the power of 1/3. This is mathematically equivalent to calculating the cube root of 216, which is 6.

Part II: Working with Formulas and Functions

Formula	What It Does
=A1<A2	Returns TRUE if the value in cell A1 is less than the value in cell A2. Otherwise, it returns FALSE. Logical comparison operators also work with text. If A1 contains Bill and A2 contains Julia, the formula would return TRUE because Bill comes before Julia in alphabetical order.
=A1<=A2	Returns TRUE if the value in cell A1 is less than or equal to the value in cell A2. Otherwise, it returns FALSE.

Understanding operator precedence in formulas

When Excel calculates the value of a formula, it uses certain rules to determine the order in which the various parts of the formula are calculated. You need to understand these rules so that your formulas produce accurate results.

Table 9.2 lists the Excel operator precedence. This table shows that exponentiation has the highest precedence (performed first) and logical comparisons have the lowest precedence (performed last).

TABLE 9.2 Operator Precedence in Excel Formulas

Symbol	Operator	Precedence
$^$	Exponentiation	1
*	Multiplication	2
/	Division	2
+	Addition	3
-	Subtraction	3
&	Concatenation	4
=	Equal to	5
<	Less than	5
>	Greater than	5

You can use parentheses to override Excel's built-in order of precedence. Expressions within parentheses are always evaluated first. For example, the following formula uses parentheses to control the order in which the calculations occur. In this case, cell B3 is subtracted from cell B2, and the result is multiplied by cell B4:

$$= (B2 - B3) * B4$$

If you enter the formula without the parentheses, Excel computes a different answer. Because multiplication has a higher precedence, cell B3 is multiplied by cell B4. This result is then subtracted from cell B2, which may not be what was intended.

The formula without parentheses looks like this:

```
=B2 - B3 * B4
```

TIP

It's good practice to use parentheses even when they aren't strictly necessary. Doing so helps to clarify what the formula is intended to do. For example, the following formula makes it perfectly clear that B3 should be multiplied by B4 and the result subtracted from cell B2. Without the parentheses, you would need to remember Excel's order of precedence.

```
=B2 - (B3 * B4)
```

You can also nest parentheses within formulas—that is, put them inside other parentheses. If you do so, Excel evaluates the most deeply nested expressions first—and then works its way out. Here's an example of a formula that uses nested parentheses:

```
= ( (B2 * C2) + (B3 * C3) + (B4 * C4) ) * B6
```

This formula has four sets of parentheses—three sets are nested inside the fourth set. Excel evaluates each nested set of parentheses and then sums the three results. This result is then multiplied by the value in cell B6.

Although the preceding formula uses four sets of parentheses, only the outer set is really necessary in this example. If you understand operator precedence, it should be clear that you can rewrite this formula as follows:

```
= (B2 * C2 + B3 * C3 + B4 * C4) * B6
```

Note that operators at the same level of precedence, such as multiplication and division, are evaluated from left to right (unless parentheses would indicate a different order).

Every left parenthesis, of course, must have a matching right parenthesis. If you have many levels of nested parentheses, keeping them straight can sometimes be difficult. If the parentheses don't match, Excel displays a message explaining the problem, and it won't let you enter the formula.

CAUTION

In some cases, if your formula contains mismatched parentheses, Excel may propose a correction to your formula. Figure 9.1 shows an example of a proposed correction. You may be tempted simply to accept Excel's suggestion, but be careful—in many cases, the proposed formula, although syntactically correct, isn't the formula you intended, and it will produce an incorrect result.

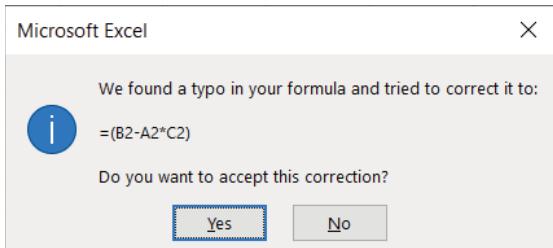
TIP

When you're editing a formula, Excel lends a hand in helping you match parentheses by displaying matching parentheses in the same color.

Part II: Working with Formulas and Functions

FIGURE 9.1

Excel sometimes suggests a syntactically correct formula, but not the formula you had in mind.



Using functions in your formulas

Many formulas that you create use worksheet functions. These functions enable you to greatly enhance the power of your formulas and perform calculations that are difficult (or even impossible) if you use only the operators discussed previously. For example, you can use the TAN function to calculate the tangent of an angle. You can't do this complicated calculation by using the mathematical operators alone.

Examples of formulas that use functions

A worksheet function can simplify a formula significantly.

Here's an example. To calculate the average of the values in 10 cells (A1 : A10) without using a function, you'd have to construct a formula like this:

```
= (A1+A2+A3+A4+A5+A6+A7+A8+A9+A10) / 10
```

Not very pretty, is it? Even worse, you would need to edit this formula if you added another cell to the range. Fortunately, you can replace this formula with a much simpler one that uses one of Excel's built-in worksheet functions, AVERAGE:

```
=AVERAGE (A1:A10)
```

The following formula demonstrates how using a function can enable you to perform calculations that are not otherwise possible. Say you need to determine the largest value in a range. A formula can't tell you the answer without using a function. Here's a formula that uses the MAX function to return the largest value in the range A1:D100:

```
=MAX (A1:D100)
```

Functions also can sometimes eliminate manual editing. Assume that you have a worksheet that contains 1,000 names in cells A1 : A1000 and the names appear in all-uppercase letters. Your boss sees the listing and informs you that the use of all-uppercase letters is not acceptable; for example, JOHN F. SMITH must now appear as John F. Smith.

With Excel's PROPER function, there's no need to manually correct each of the 1,000 names. You can simply enter the following formula in cell B1 then copy it down to the next 999 rows.

```
=PROPER(A1)
```

Several hours of work completed with one Excel function.

TIP

You can also use Excel's Flash Fill feature to make this type of transformation, without formulas. See Chapter 22, "Importing and Cleaning Data," for more about Flash Fill.

One last example should convince you of the power of functions. Suppose you have a worksheet that calculates sales commissions. If the salesperson sold \$100,000 or more of product (cell A1), the commission rate is 7.5 percent; otherwise, the commission rate is 5.0 percent. You can write a formula using the IF function to automatically calculate the correct commission.

The IF function allows you to build formulas that return a different result based on some logic check. For instance, in this scenario, you could use:

```
=IF(A1<100000,A1*5%,A1*7.5%)
```

The formula checks the value of cell A1, which contains the sales amount. If this value is less than 100,000, the formula returns cell A1 multiplied by 5 percent. Otherwise, it returns what's in cell A1 multiplied by 7.5 percent. This example uses three arguments, separated by commas.

In the next section, we'll go into more detail on the concept of arguments.



You'll dive into the IF function in more detail with Chapter 14, "Using Formulas for Conditional Analysis."

Function arguments

In the preceding examples, you may have noticed that all of the functions used parentheses. The information inside the parentheses is the list of arguments.

9

Functions vary in the way they use arguments. Depending on what it's designed to do, a function may use one of the following:

- No arguments
- One argument
- A fixed number of arguments
- An indeterminate number of arguments
- Optional arguments

An example of a function that doesn't use an argument is the NOW function, which returns the current date and time. Even if a function doesn't use an argument, you must still provide a set of empty parentheses, like this:

```
=NOW()
```

If a function uses more than one argument, separate each argument with a comma. The examples at the beginning of the chapter used cell references for arguments. Excel is quite flexible when it comes to function arguments, however. An argument can consist of a cell reference, literal values, literal text strings, expressions, and even other functions. Here are some examples of functions that use various types of arguments:

Cell reference: =SUM(A1:A24)

Literal value: =SQRT(121)

Literal text string: =PROPER("john f. smith")

Expression: =SQRT(183+12)

Other functions: =SQRT(SUM(A1:A24))

NOTE

A comma is the list separator character for the U.S. version of Excel. Some other language versions may use a semicolon. The list separator is a Windows setting that can be adjusted in the Windows Control Panel (the Region dialog box).

More about functions

All told, Excel includes more than 450 built-in functions. And if that's not enough, you can download or purchase additional specialized functions from third-party suppliers—and even create your own custom functions (by using VBA) if you're so inclined.

Some users feel a bit overwhelmed by the sheer number of functions, but you'll probably find that you use only a dozen or so on a regular basis. And as you'll see, the Excel Insert Function dialog box (described later in this chapter) makes it easy to locate and insert a function, even if it's not one that you use frequently.



You'll find many examples of Excel's built-in functions throughout Part II of this book, "Working with Formulas and Functions." Chapter 40, "Creating Custom Worksheet Functions," covers the basics of creating custom functions with Visual Basic for Applications (VBA).

Entering Formulas into Your Worksheets

Every formula must begin with an equal sign to inform Excel that the cell contains a formula rather than text. Excel provides two ways to enter a formula into a cell: manually or by pointing to cell references.

Entering a formula manually starts by clicking the cell where you want your formula to reside. In the selected cell, you type an equal sign (=) followed by the formula. For example, you can enter the following formula in cell C1 to add the contents of cells A1 and B1:

=A1+B1

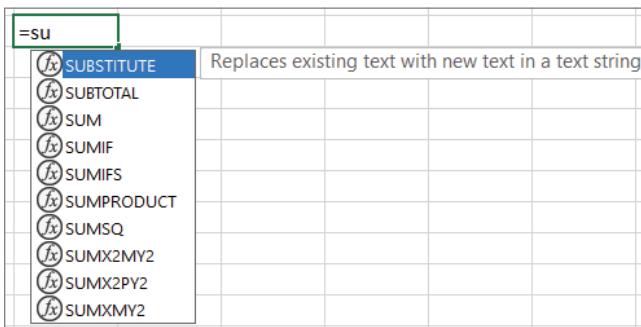
As you type, the characters appear in the cell and in the Formula bar. You can, of course, use all of the normal editing keys when entering a formula.

Excel provides additional assistance when you create formulas by displaying a drop-down list that contains function names and range names. The items displayed in the list are determined by what you've already typed.

For example, if you click a cell and type an equal sign (=) followed by the letters **su**, you'll see the drop-down list shown in Figure 9.2. If you type an additional letter, the list is shortened to show only the matching functions. To have Excel AutoComplete an entry in that list, use the navigation keys to highlight the entry and then press Tab. Notice that highlighting a function in the list also displays a brief description of the function. See the next sidebar, "Using Formula AutoComplete," for an example of how this feature works.

FIGURE 9.2

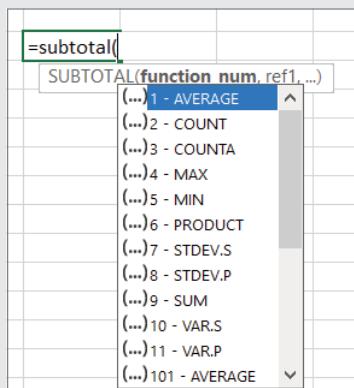
Excel displays a drop-down list when you enter a formula.



Using Formula AutoComplete

The Formula AutoComplete feature makes entering formulas easier than ever. Simply start typing your formula and Excel will help by presenting you with a list of the possible options and arguments available. In this example, Excel is presenting the options for the **SUBTOTAL** function.

9



Continues

continued

Formula AutoComplete includes the following items (and each type is identified by an icon):

- Excel built-in functions.
- User-defined functions (functions defined by the user through VBA or other methods).
- Defined names (cells or ranges named using the Formulas \Rightarrow Defined Names \Rightarrow Define Name command).
- Enumerated arguments that use a value to represent an option. (Only a few functions use such arguments, and SUBTOTAL is one of them.)
- Table structure references (used to identify portions of a table).

Entering formulas by pointing

Even though you can enter formulas by typing in the entire formula, Excel provides another method of entering formulas that is generally easier, faster, and less prone to errors. This method still involves some manual typing, but you can simply point to the cell references instead of typing their values manually. For example, to enter the formula =A1+A2 into cell A3, follow these steps:

1. **Select cell A3.**
2. **Type an equal sign (=) to begin the formula.** Notice that Excel displays Enter in the status bar (lower left of your screen).
3. **Press the up arrow twice.** As you press this key, Excel displays a dashed border around cell A1, and the cell reference appears in cell A3 and in the Formula bar. In addition, Excel displays Point in the status bar.
4. **Type a plus sign (+).** A solid color border replaces the dashed border of A1, and Enter reappears in the status bar.
5. **Press the up arrow again.** The dashed border encompasses cell A2 and adds that cell address to the formula.
6. **Press Enter to complete the formula.**

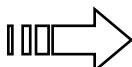
TIP

When creating a formula by pointing, you can also point to the data cells by using your mouse.

Pasting range names into formulas

If your formula uses named cells or ranges, you can either type the name in place of the address or choose the name from a list and have Excel insert the name for you automatically. Three ways to insert a name into a formula are available:

- Select the name from the drop-down list. To use this method, you must know at least the first character of the name. When you're entering the formula, type the first character and then select the name from the drop-down list and press Tab.
- Press F3. The Paste Name dialog box appears. Select the name from the list and then click OK (or just double-click the name). Excel enters the name into your formula. If no names are defined, pressing F3 has no effect.
- Click the Use in Formula drop-down on the Formulas tab (Defined Names group). This command is available while you are in edit mode, and it allows you to select from the available range names.



See Chapter 4, "Working with Excel Ranges and Tables," for information about creating names for cells and ranges.

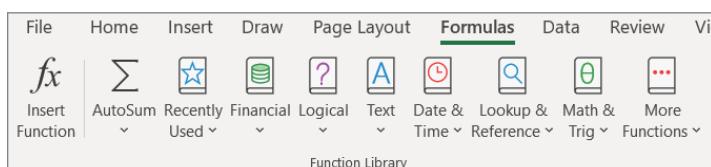
Inserting functions into formulas

The easiest way to enter a function into a formula is to select a cell, type the equal sign (=) and then type the first few letters of the function, letting the AutoComplete feature guide you the rest of the way.

Alternatively, you can insert a function using the Function Library group on the Formulas tab on the Ribbon (see Figure 9.3). This method is especially useful if you can't remember which function you need. When entering a formula, click the function category (Financial, Logical, Text, and so on) to get a list of the functions in that category. Click the function that you want, and Excel displays its Function Arguments dialog box. This is where you enter the function's arguments. In addition, you can click the Help on This Function link to learn more about the selected function.

FIGURE 9.3

You can insert a function by selecting it from one of the function categories.



9

Yet another way to insert a function into a formula is to use the Insert Function dialog box (see Figure 9.4). You can access this dialog box in several ways:

- Choose Formulas → Function Library → Insert Function.
- Use the Insert Function command, which appears at the bottom of each drop-down list in the Formulas → Function Library group.
- Click the Insert Function icon, which is directly to the left of the Formula bar. This button displays fx.
- Press Shift+F3.

Part II: Working with Formulas and Functions

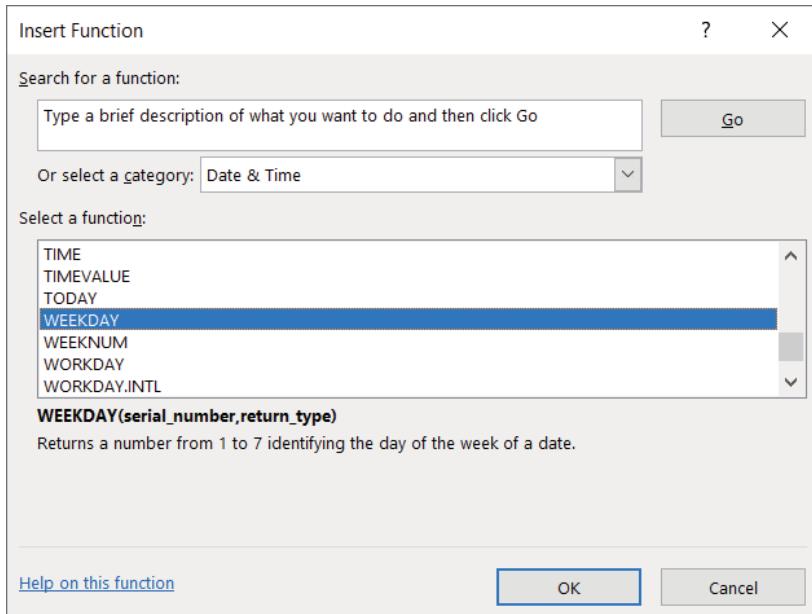
The Insert Function dialog box shows a drop-down list of function categories. Select a category, and the functions in that category are displayed in the list box.

If you're not sure which function you need, you can search for the appropriate function by using the Search for a Function field at the top of the dialog box.

1. **Enter your search terms and click Go.** You get a list of relevant functions. When you select a function from the Select a Function list, Excel displays the function (and its argument names) in the dialog box along with a brief description of what the function does.

FIGURE 9.4

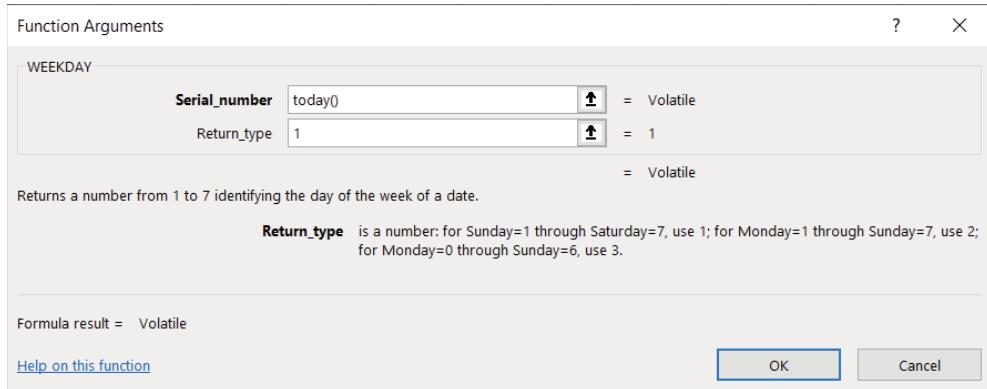
The Insert Function dialog box



2. **When you locate the function you want to use, highlight it and click OK.** Excel then displays its Function Arguments dialog box, as shown in Figure 9.5.

FIGURE 9.5

The Function Arguments dialog box



3. **Specify the arguments for the function.** The Function Arguments dialog box will vary depending on the function that you are inserting, and it will show one text box for each of the function's arguments. To use a cell or range reference as an argument, you can enter the address manually or click inside the argument box and then select (that is, point to) the cell or range in the sheet.
4. **After you specify all the function arguments, click OK.**

TIP

When you're entering or editing a formula, the space typically occupied by the Name box displays a list of the functions that you've used most recently (if any). You can quickly add any of the functions listed simply by selecting one. After you select a function from this list, Excel will display the Function Arguments dialog box.

Function entry tips

Here are some additional tips to keep in mind when you use the Insert Function dialog box to enter functions:

- You can use the Insert Function dialog box to insert a function into an existing formula. Just edit the formula and move the insertion point to the location at which you want to insert the function. Then open the Insert Function dialog box (using any of the methods described earlier) and select the function.
- You can also use the Function Arguments dialog box to modify the arguments for a function in an existing formula. Click the function in the Formula bar and then click the Insert Function button (the fx button, to the left of the Formula bar).

- If you change your mind about entering a function, click the Cancel button.
- The number of boxes that you see in the Function Arguments dialog box depends on the number of arguments used in the function you selected. If a function uses no arguments, you won't see any boxes. If the function uses a variable number of arguments (such as the AVERAGE function), Excel adds a new box every time you enter an optional argument. Note that required arguments appear in bold, while optional arguments are not bold.
- As you provide arguments in the Function Arguments dialog box, the value of each argument is displayed to the right of each box.
- A few functions, such as INDEX, have more than one form. If you choose such a function, Excel displays another dialog box that lets you choose which form you want to use.
- As you become familiar with the functions, you can bypass the Insert Function dialog box and type the function name directly. Excel prompts you with argument names as you enter the function.

Editing Formulas

After you enter a formula, you can (of course) edit it. You may need to edit a formula if you make some changes to your worksheet and then have to adjust the formula to accommodate the changes. Or the formula may return an error value, in which case you need to edit the formula to correct the error.

Here are some of the ways to get into cell edit mode:

- Double-click the cell, which enables you to edit the cell contents directly in the cell.
- Press F2, which enables you to edit the cell contents directly in the cell.
- Select the cell that you want to edit and then click in the Formula bar. This enables you to edit the cell contents in the Formula bar.
- If the cell contains a formula that returns an error, Excel will display a small triangle in the upper-left corner of the cell. Activate the cell, and you'll see an error indicator. Click the error indicator, and you can choose one of the options for correcting the error. (The options will vary according to the type of error in the cell.)

TIP

You can control whether Excel displays these formula error indicators in the Formulas section of the Excel Options dialog box. To display this dialog box, choose File ➔ Options. If you remove the check mark from Enable Background Error Checking, Excel no longer displays these error indicators.

While you're editing a formula, you can select multiple characters either by dragging the mouse cursor over them or by pressing Shift while you use the navigation keys. You may notice that while editing a formula, Excel will color-code the range addresses in the formula. This helps you quickly spot the cells that are used in a formula.

TIP

If you have a formula that you can't seem to edit correctly, you can convert the formula to text and tackle it again later. To convert a formula to text, just remove the initial equal sign (=). When you're ready to try again, type the initial equal sign to convert the cell contents back to a formula.

Using Cell References in Formulas

Imagine you go to cell C1 and enter the formula =A1+B1. Your human eyes will see that formula as add the value in A1 to the value in B1. However, Excel doesn't see it that way. Excel translates the formula like this: take the value in the cell two spaces to the left and add it to the value one space to the left. If you would copy the formula in cell C1 and paste it into cell C2, you'll see the formula in cell C2 would read =A2+B2. In short, Excel doesn't use the actual column and row coordinates you see. Instead, it evaluates cell references in terms of where they are relative to the cell in which the formula resides.

By default, Excel uses what is known as relative references. If you want to ensure that Excel does not adjust cell references when a formula is copied, you can lock the references by turning them into absolute references. This is done by adding a dollar symbol (\$) before the column and row reference. For instance, click cell C1 and enter =\$A\$1+\$B\$1. The dollar symbols ensure that the references in your formula remain fixed (or absolute). If you would copy the formula in cell C1 and paste it into cell C2, you'll see the formula in cell C2 would remain =\$A\$1+\$B\$1.

Using relative, absolute, and mixed references

Excel gives you the flexibility to make any part of your formula a relative reference, absolute reference, or even a mixed reference.

Relative Reference The row and column references are not locked by dollar symbols (\$) and will change when you copy the formula to another cell relative to the cell in which the formula resides.

Absolute Reference The row and column references are locked by the dollar symbol (\$) and will not change when you copy the formula. An absolute reference uses two dollar signs in its address: one for the column letter and one for the row number (for example, \$A\$5).

Mixed Reference Either the row or the column reference is relative, and the other is absolute. Only one of the address parts is absolute (for example, \$A4 or A\$4).

The type of cell reference is important only if you plan to copy the formula to other cells. The following examples illustrate this point.

Part II: Working with Formulas and Functions

Figure 9.6 shows a simple worksheet. The formula in cell D2, which multiplies the quantity by the price, is

=B2*C2

FIGURE 9.6

Copying a formula that contains relative references

	A	B	C	D
1	Item	Quantity	Price	Total
2	Chair	4	\$125.00	\$500.00
3	Desk	4	\$695.00	\$2,780.00
4	Lamp	3	\$39.95	\$119.85

This formula uses relative cell references. Therefore, when the formula is copied to the cells below it, the references adjust in a relative manner. For example, the formula in cell D3 is

=B3*C3

But what if the cell references in D2 contained absolute references, like this?

=\$B\$2*\$C\$2

In this case, copying the formula to the cells below would produce incorrect results. The formula in cell D3 would be the same as the formula in cell D2.

Now we'll extend the example to calculate sales tax, which is stored in cell B7 (see Figure 9.7). In this situation, the formula in cell D2 is

= (B2*C2) *\$B\$7

FIGURE 9.7

Formula references to the sales tax cell should be absolute.

	A	B	C	D	E
1	Item	Quantity	Price	Sales Tax	Total
2	Chair	4	\$125.00	\$37.50	
3	Desk	4	\$695.00		
4	Lamp	3	\$39.95		
5					
6					
7	Sales Tax:	7.50%			

The quantity is multiplied by the price, and the result is multiplied by the sales tax rate stored in cell B7. Notice that the reference to B7 is an absolute reference. When the formula in D2 is copied to the cell below it, cell D3 will contain this formula:

= (B3*C3) *\$B\$7

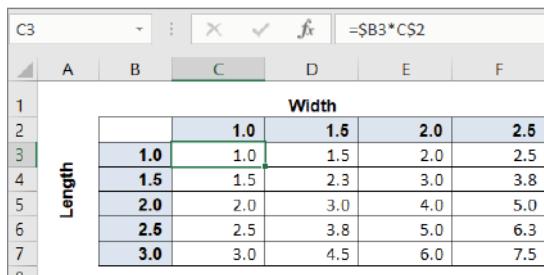
Here, the references to cells B2 and C2 were adjusted, but the reference to cell B7 was not, which is exactly what you want because the address of the cell that contains the sales tax never changes.

Figure 9.8 demonstrates the use of mixed references. The formulas in the C3 : F7 range calculate the area for various lengths and widths. Here's the formula in cell C3:

`=\$B3 * C$2`

FIGURE 9.8

Using mixed cell references

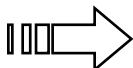


	B	C	D	E	F
1		Width			
2		1.0	1.5	2.0	2.5
3	1.0	1.0	1.5	2.0	2.5
4	1.5	1.5	2.3	3.0	3.8
5	2.0	2.0	3.0	4.0	5.0
6	2.5	2.5	3.8	5.0	6.3
7	3.0	3.0	4.5	6.0	7.5

Notice that both cell references are mixed. The reference to cell B3 uses an absolute reference for the column (\$B), and the reference to cell C2 uses an absolute reference for the row (\$2). As a result, this formula can be copied down and across, and the calculations will be correct. For example, the formula in cell F7 is

`=\$B7 * F$2`

If C3 used either absolute or relative references, copying the formula would produce incorrect results.



A workbook that demonstrates the various types of references is available on this book's website at www.wiley.com/go/excel365bible. The file is named `cell_references.xlsx`.

9

Note

When you copy and paste a formula, Excel will automatically adjust the cell references. However, if you cut and paste a formula, Excel will assume you want to keep the same cell references and will not adjust them.

Changing the types of your references

You can enter nonrelative references (that is, absolute or mixed) manually by inserting dollar signs in the appropriate positions of the cell address. Or you can use a handy shortcut: the F4 key. When you've entered a cell reference (by typing it or by pointing), you can press F4 repeatedly to have Excel cycle through all four reference types.

Part II: Working with Formulas and Functions

For example, if you enter =A1 to start a formula, pressing F4 converts the cell reference to =\$A\$1. Pressing F4 again converts it to =A\$1. Pressing it again displays =\$A1. Pressing it one more time returns to the original =A1. Keep pressing F4 until Excel displays the type of reference that you want.

NOTE

When you name a cell or range, Excel (by default) uses an absolute reference for the name. For example, if you give the name SalesForecast to B1:B12, the Refers To box in the New Name dialog box lists the reference as \$B\$1:\$B\$12. If you copy a cell that has a named reference in its formula, the copied formula contains a reference to the original name.

Referencing cells outside the worksheet

Formulas can also refer to cells in other worksheets—and the worksheets don't even have to be in the same workbook. Excel uses a special type of notation to handle these types of references.

Referencing cells in other worksheets

To use a reference to a cell in another worksheet in the same workbook, use this format:

=SheetName!CellAddress

In other words, precede the cell address with the worksheet name followed by an exclamation point. Here's an example of a formula that uses a cell on the Sheet2 worksheet:

=A1*Sheet2!A1

This formula multiplies the value in cell A1 on the current worksheet by the value in cell A1 on Sheet2.

TIP

If the worksheet name in the reference includes one or more spaces, you must enclose it in single quotation marks. (Excel does that automatically if you use the point-and-click method when creating the formula.) For example, here's a formula that refers to a cell on a sheet named All Depts:

=A1*'All Depts'!A1

Referencing cells in other workbooks

To refer to a cell in a different workbook, use this format:

=[WorkbookName]SheetName!CellAddress

In this case, the workbook name (in square brackets), the worksheet name, and an exclamation point precede the cell address. The following is an example of a formula that uses a cell reference in the Sheet1 worksheet in a workbook named Budget:

=[Budget.xlsx]Sheet1!A1

If the workbook name in the reference includes one or more spaces, you must enclose it (and the sheet name and square brackets) in single quotation marks. For example, here's a formula that refers to a cell on Sheet1 in a workbook named Budget For 2022:

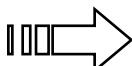
```
=A1*' [Budget For 2022.xlsx]Sheet1'!A1
```

When a formula refers to cells in a different workbook, the other workbook doesn't have to be open. If the workbook is closed, however, you must add the complete path to the reference so that Excel can find it. Here's an example:

```
=A1*'C:\My Documents\ [Budget For 2022.xlsx]Sheet1'!A1
```

A linked file can also reside on another system that's accessible on your corporate network. The following formula refers to a cell in a workbook in the files directory of a computer named DataServer:

```
='\DataServer\files\ [budget.xlsx]Sheet1'!$D$7
```



See Chapter 25, "Linking and Consolidating Worksheets," for more information about linking workbooks.

TIP

To create formulas that refer to cells in a different worksheet, open the workbook you're referencing and simply point to the cells needed instead of entering their references manually. Excel takes care of the details regarding the workbook and worksheet references.

NOTE

If you point to a different worksheet or workbook when creating a formula, you'll notice that Excel always inserts absolute cell references. Therefore, if you plan to copy the formula to other cells, make sure that you change the cell references to relative before you copy.

Introducing Formula Variables

As you build formulas, you'll notice it sometimes becomes necessary to repeat the same calculation more than once in a formula. This is often the case when writing nested formulas; formulas that use the results of other functions as arguments.

For instance, the following formula uses a nested IF function to check the result of A1+B1. If the result is <50 then "Below Fifty" will be displayed. If the value is >100, then "Above One Hundred" will be displayed. If neither the first nor the second condition is true, then the formulas will display "In Between."

```
=IF(A1+B1<50, "Below Fifty", IF(A1+B1>100, "Above One Hundred", "In Between"))
```

You can see Excel has to calculate A1+B1 twice; once to check the first IF condition and once to check the second condition. In this simple example, calculating A1+B1 twice won't bring Excel to a standstill, but the fact is Excel is forced to perform a redundant calculation. In more complex formulas with many nested functions, the need to perform the same calculation over and over can affect performance. In addition to performance considerations, nested calculations can be difficult to read.

In the following sections, you'll discover how formula variables can be leveraged to simplify your formulas and potentially improve performance.

Understanding the LET function

The LET function allows you to create a container of sorts that holds the results of a function or calculation. This container is called a formula variable.

The LET function requires a minimum of three arguments:

- **A name for your variable:** You'll want this name to be a friendly name you can easily recognize and understand. Note that your variable name can't start with a number.
- **A value for your variable:** Once you have a name for your variable, you need to specify the value you want the variable to hold. The name and value combination together make up what is known as a value pair.
- **A formula that uses the variable:** The final argument is a formula expression that uses the variable.

To build a basic LET function, click cell A1 and enter the following:

```
=LET(MyVariable, 5*2, MyVariable*10)
```

The result of this formula is 100. It first fills MyVariable with the result of $5*2$ (in this case 10), then multiplies MyVariable by 10. This particular formula has one value pair. Again, a value pair is a name and value combination.

The LET function shown here uses two value pairs. MyVariable1, which contains 10 (the result of $5*2$), and MyVariable2, which contains 50. The final argument multiplies the variables by 10, resulting in 5000.

```
=LET(MyVariable1, 5*2, MyVariable2, 50, MyVariable1*MyVariable2*10)
```

TIP

You can use up to 126 value pairs in a single LET function.

CAUTION

When naming variables for your LET function, it's best practice to avoid using named ranges defined in the Name Manager. If you use a name already defined in the Name Manager, LET will effectively ignore and override the named range. Refer to Chapter 4, "Working with Excel Ranges and Tables," for information on creating and managing named ranges.

Formula variables in action

With the basics under your belt, let's take a moment to witness formula variables in action.

Figure 9.9 illustrates two formulas. The first formula in cell C6 is a traditional nested IF statement that evaluates the value of C3/D3. You can see in this first formula that C3/D3 is performed twice.

```
=IF((C3/D3)>1,"Goal Met",IF((C3/D3)<1,"Below Goal","Flat"))
```

Leveraging the LET function allows us to gain efficiency by holding the value of C3/D3 in the variable called X. Excel performs that division only one time and then passes the answer to the rest of the formula.

```
=LET(X,C3/D3,IF(X>1,"Goal Met",IF(X<1,"Below Goal","Flat")))
```

FIGURE 9.9

Using the LET function to simplify a nested IF statement

	A	B	C	D	E
1					
2					
3		Month	Projected	Actual	
4		Jan	4,000	3,255	
5					
6	Nested IF	Goal Met	=IF((C3/D3)>1,"Goal Met",IF((C3/D3)<1,"Below Goal","Flat"))		
7					
8	LET Function	Goal Met	=LET(X,C3/D3,IF(X>1,"Goal Met",IF(X<1,"Below Goal","Flat")))		

Figure 9.10 demonstrates a more complex SWITCH function that calculates the number of days between the Due Date and Invoice Date then returns text to represent the category of day's aging of an invoice. The following formula is listed in cell J5:

```
=SWITCH(TRUE,(I3-H3)>90,"90+",(I3-H3)>30,"31-60", (I3-H3)>0,"")
```

This function calculates the difference in days between the Due Date (cell I3) and the Invoice Date (cell H3) and then returns the appropriate aging category based on the difference in days. Again, not only is this formula difficult to read, but Excel will need to calculate the result of I3-H3 up to three times to fully calculate the correct answer.

FIGURE 9.10

Using the LET function to simplify a SWITCH statement

	H	I	J
1			
2	Invoice Date	Due Date	
3	5/5/2021	7/12/2021	
4			
5	Aging Days	31 - 60 Days	=SWITCH(TRUE,(I3-H3)>90,"90+",(I3-H3)>30,"31-60", (I3-H3)>0,"")
6			
7	LET Function	31 - 60 Days	=LET(dys,I3-H3,SWITCH(TRUE,dys>90,"90+",dys>30,"31-60",dys>0,""))

Part II: Working with Formulas and Functions

The improved LET function shown in cell J7 (see Figure 9.10) uses the variable dys to store the result I3-H3.

```
=LET(dys,I3-H3,SWITCH(TRUE,dys>90,"90+",dys>30,"31-60",dys>0,""))
```

Instead of repeating the initial calculation to get the difference in days, the dys variable filled one time, then passed to the SWITCH function.



You'll discover the SWITCH function along with other logic functions in Chapter 14, "Using Formulas for Conditional Analysis."

Using Formulas in Tables

A table is a specially designated range of cells set up with column headers. In the following sections, we describe how formulas work with tables.



See Chapter 4, "Working with Excel Ranges and Tables," for an introduction to the Excel table features.

Summarizing data in a table

Figure 9.11 shows a simple table with three columns. We entered the data and then converted the range to a table by choosing Insert ▾ Tables ▾ Table. Note that we didn't define any names, but the table is named Table1 by default.

FIGURE 9.11

A simple table with three columns of information

	A	B	C	D
1		Month	Projected	Actual
2		Jan	4,000	3,255
3		Feb	4,000	4,102
4		Mar	4,000	3,982
5		Apr	5,000	4,598
6		May	5,000	5,873
7		Jun	5,000	4,783
8		Jul	5,000	5,109
9		Aug	6,000	5,982
10		Sep	6,000	6,201
11		Oct	7,000	6,833
12		Nov	8,000	7,983
13		Dec	9,000	9,821
14				



This workbook is available on this book's website at www.wiley.com/go/excel365bible. It is named table formulas.xlsx.

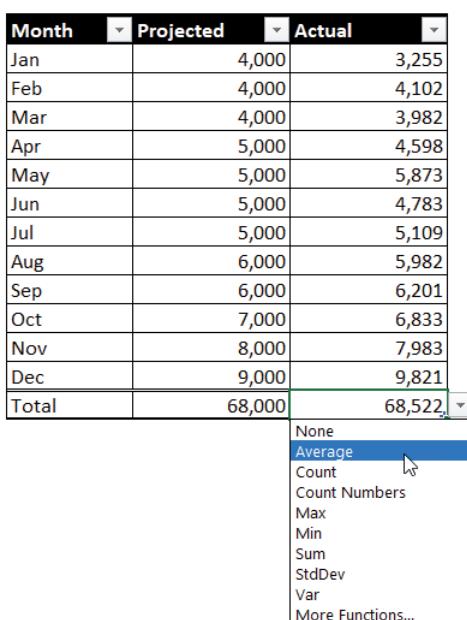
If you'd like to calculate the total projected and total actual sales, you don't even need to write a formula. Simply click a button to add a row of summary formulas to the table:

1. **Activate any cell in the table.**
2. **Place a check mark next to Table Design \Rightarrow Table Style Options \Rightarrow Total Row.** Excel adds a total row to the table and displays the sum of each numeric column.
3. **To change the type of summary formula, activate a cell in the total row and use the drop-down list to change the type of summary formula to use (see Figure 9.12).** For example, to calculate the average of the Actual column, select AVERAGE from the drop-down list in cell D15. Excel creates this formula:

=SUBTOTAL(101, [Actual])

FIGURE 9.12

A drop-down list enables you to select a summary formula for a table column.



The figure shows a Microsoft Excel table with three columns: Month, Projected, and Actual. The Actual column contains numerical values ranging from 3,255 to 9,821, with a total of 68,522. The last row of the table is a total row, also labeled 'Total' in the first column, with values 68,000 and 68,522 in the second and third columns respectively. A dropdown menu is open over the cell containing the value 68,522 in the Actual column. The menu lists various summary functions: None, Average (which is highlighted), Count, Count Numbers, Max, Min, Sum, StdDev, Var, and More Functions... The menu is displayed in a white box with a black border.

Month	Projected	Actual
Jan	4,000	3,255
Feb	4,000	4,102
Mar	4,000	3,982
Apr	5,000	4,598
May	5,000	5,873
Jun	5,000	4,783
Jul	5,000	5,109
Aug	6,000	5,982
Sep	6,000	6,201
Oct	7,000	6,833
Nov	8,000	7,983
Dec	9,000	9,821
Total	68,000	68,522

For the SUBTOTAL function, 101 is an enumerated argument that represents AVERAGE. The second argument for the SUBTOTAL function is the column name in square brackets. Using the column name within brackets creates “structured” references within a table. (We discuss this further in the upcoming section “Referencing data in a table.”)

NOTE

You can toggle the total row display via Table Design \Rightarrow Table Style Options \Rightarrow Total Row. If you turn it off, the summary options you selected will be displayed again when you turn it back on.

Using formulas within a table

In many cases, you'll want to use formulas within a table to perform calculations that use other columns in the table. For example, in the table shown in Figure 9.12, you may want a column that shows the difference between the Actual and Projected amounts. To add this formula, follow these steps:

1. **Activate cell E2 and type Difference for the column header.** After you press Enter, Excel automatically expands the table for you to include the new column.
2. **Move to cell E3 and type an equal sign to signify the beginning of a formula.**
3. **Press the left arrow key.** Excel displays `[@Actual]`, which is the column heading, in the Formula bar.
4. **Type a minus sign and then press the left arrow key twice.** Excel displays `[@Projected]` in your formula.
5. **Press Enter to end the formula.** Excel copies the formula to all rows in the table.

Figure 9.13 shows the table with the new column.

FIGURE 9.13

The Difference column contains a formula.

Month	Projected	Actual	Difference
Jan	4,000	3,255	-745
Feb	4,000	4,102	102
Mar	4,000	3,982	-18
Apr	5,000	4,598	-402
May	5,000	5,873	873
Jun	5,000	4,783	-217
Jul	5,000	5,109	109
Aug	6,000	5,982	-18
Sep	6,000	6,201	201
Oct	7,000	6,833	-167
Nov	8,000	7,983	-17
Dec	9,000	9,821	821
Total	68,000	68,522	

Examine the table, and you will find this formula for all cells in the Difference column:

=[@Actual] - [@Projected]

Although the formula was entered into the first row of the table, that's not necessary. Any time a formula is entered into an empty table column, it will automatically fill all of the cells in that column. If you need to edit the formula, Excel will automatically copy the edited formula to the other cells in the column.

NOTE

The at symbol (@) that precedes the column header represents "this row." So [@Actual] means "the value in the Actual column in this row."

These steps use the pointing technique to create the formula. Alternatively, you could have entered the formula manually using standard cell references rather than column headers. For example, you could have entered the following formula in cell E3:

=D3 - C3

If you type the cell references, Excel will still copy the formula to the other cells automatically.

One thing should be clear, however, about formulas that use the column headers instead of cell references—they're much easier to understand.

TIP

To override the automatic column formulas, access the Proofing tab of the Excel Options dialog box. Click AutoCorrect Options and then select the AutoFormat As You Type tab in the AutoCorrect dialog box. Remove the check mark from Fill Formulas in Tables to Create Calculated Columns.

Referencing data in a table

Excel offers some other ways to refer to data that's contained in a table by using the table name and column headers.

You can, of course, use standard cell references to refer to data in a table, but using the table name and column headers has a distinct advantage: the names adjust automatically if the table size changes by inserting or deleting rows. In addition, formulas that use table names and column headers will adjust automatically if you change the name of the table or give a new name to a column.

To refer to the table shown in Figure 9.13, you can use the table name. For instance, to calculate the sum of all the data in the table, enter this formula into a cell outside the table:

=SUM(Table1)

Part II: Working with Formulas and Functions

This formula will return the sum of all the data (excluding calculated total row values, if any), even if rows or columns are added or deleted. And if you change the name of Table1, Excel will adjust formulas that refer to that table automatically. For example, if you renamed Table1 to AnnualData, the preceding formula would change to

```
=SUM(AnnualData)
```

Most of the time, a formula will refer to a specific column in the table. The following formula returns the sum of the data in the Actual column:

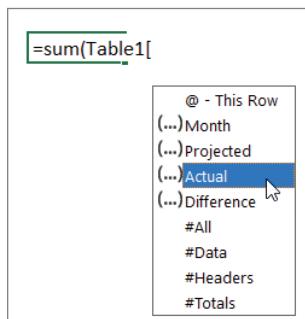
```
=SUM(Table1[Actual])
```

Notice that the column name is enclosed in square brackets. Again, the formula adjusts automatically if you change the text in the column heading.

Even better, Excel provides some helpful assistance when you create a formula that refers to data within a table. Figure 9.14 shows the formula AutoComplete feature helping to create a formula by showing a list of the elements in the table. Notice that, in addition to the column headers in the table, Excel lists other table elements that you can reference: #All, #Data, #Headers, #Totals, and @ - This Row.

FIGURE 9.14

The formula AutoComplete feature is useful when creating a formula that refers to data in a table.



NOTE

Remember that you don't need to create names for tables and columns. The data in the table itself has a range name, which is created automatically when you create the table (for example, Table1). You can find the name of a table object by first selecting any cell in the table and then choosing the Table Design tab. The Table Name text box in the Properties group displays the name of the table. You can update the name by replacing the value in the Table Name text box.

Correcting Common Formula Errors

Sometimes when you enter a formula, Excel displays a value that begins with a hash mark (#). This is a signal that the formula is returning an error value. You have to correct the formula (or correct a cell that the formula references) to get rid of the error display.

TIP

If the entire cell is filled with hash-mark characters, the column isn't wide enough to display the value. You can either widen the column or change the number format of the cell.

In some cases, Excel won't even let you enter an erroneous formula. For example, the following formula is missing the closing parenthesis:

=A1* (B1+C2

If you attempt to enter this formula, Excel informs you that you have unmatched parentheses, and it proposes a correction. Often, the proposed correction is accurate, but you can't count on it.

Table 9.3 lists the types of error values that may appear in a cell that has a formula. Formulas may return an error value if a cell to which they refer has an error value. This is known as the *ripple effect*—a single error value can make its way into lots of other cells that contain formulas that depend on that one cell.

TABLE 9.3 Excel Error Values

Error Value	Explanation
#DIV/0!	The formula is trying to divide by zero. Because Excel applies a value of 0 to empty cells, this error also occurs when the formula attempts to divide by a cell that is blank or has a value of 0.
#NAME?	The formula uses a name that Excel doesn't recognize. This can happen if you delete a name that's used in the formula, if you misspell a name and then press Enter, or if you have unmatched quotes when using text.
#N/A	The formula is referring (directly or indirectly) to a cell that uses the NA function to signal that data is not available. For instance, the following formula returns an #N/A error if cell A1 is empty: =IF(A1="",NA(),A1). Some lookup functions (for example, VLOOKUP and MATCH) can also return #N/A when they do not find a match.
#NULL!	The formula uses an intersection of two ranges that don't intersect. (This concept is described later in the chapter.)
#NUM!	A problem with a value exists; for example, you specified a negative number as an argument where a positive number is required.
#REF!	The formula refers to a cell that isn't valid. This can happen if the cell has been deleted from the worksheet.
#VALUE!	The formula includes an argument or operand of the wrong type. (An operand is a value or cell reference that a formula uses to calculate a result.)

Handling circular references

When you're entering formulas, you may occasionally see a warning message indicating that the formula you just entered will result in a circular reference. A *circular reference* occurs when a formula refers to its own cell—either directly or indirectly. For example, you create a circular reference if you enter =A1+A2+A3 into cell A3 because the formula in cell A3 refers to cell A3. Every time the formula in A3 is calculated, it must be calculated again because A3 has changed. The calculation could go on forever.

When you get the circular reference message after entering a formula, Excel gives you two options:

- Click OK to enter the formula as is.
- Click Help to see a help screen about circular references.

Regardless of which option you choose, Excel displays a message on the left side of the status bar to remind you that a circular reference exists.

CAUTION

Excel won't tell you about a circular reference if the Enable Iterative Calculation setting is in effect. You can check this setting in the Formulas section of the Excel Options dialog box. If Enable Iterative Calculation is turned on, Excel performs the circular calculation exactly the number of times specified in the Maximum Iterations field (or until the value changes by less than 0.001 or whatever value is in the Maximum Change field). In a few situations, you may use a circular reference intentionally. In these cases, the Enable Iterative Calculation setting must be on. However, it's best to keep this setting turned off so that you're warned of circular references. Usually, a circular reference indicates an error that you must correct.

Often, a circular reference is quite obvious and easy to identify and correct. But when a circular reference is indirect (as when a formula refers to another formula that refers to yet another formula that refers to the original formula), it may require a bit of detective work to get to the problem.

Specifying when formulas are calculated

You've probably noticed that Excel calculates the formulas in your worksheet immediately. If you change any cells that the formula uses, Excel displays the formula's new result with no effort on your part. All of this happens when Excel's Calculation mode is set to Automatic. In Automatic Calculation mode (which is the default mode), Excel follows these rules when it calculates your worksheet:

- When you make a change—enter or edit data or formulas, for example—Excel calculates immediately those formulas that depend on new or edited data.
- If Excel is in the middle of a lengthy calculation, it temporarily suspends the calculation when you need to perform other worksheet tasks; it resumes calculating when you're finished with your other worksheet tasks.
- Formulas are evaluated in a natural sequence. In other words, if a formula in cell D12 depends on the result of a formula in cell D24, Excel calculates cell D24 before calculating cell D12.

Sometimes, however, you may want to control when Excel calculates formulas. For example, if you create a worksheet with thousands of complex formulas, you may find that processing can slow to a snail's pace while Excel does its thing. In such a case, set Excel's calculation mode to Manual, which you can do by choosing **Formulas** ➔ **Calculation** ➔ **Calculation Options** ➔ **Manual**.

TIP

Large ranges of data, which Excel calls Data Tables, can calculate notoriously slowly when containing formulas. If your worksheet uses any large data tables, you may want to select the **Automatic Except for Data Tables** option.

By the way, don't let the term *data table* confuse you. It refers to a normal range of cells that has *not* been converted to a Table object by choosing **Insert** ➔ **Tables** ➔ **Table**.



See Chapter 28, “Performing Spreadsheet What-If Analysis,” for more on data tables.

When you’re working in Manual Calculation mode, Excel displays **Calculate** in the status bar when you have any uncalculated formulas. You can use the following shortcut keys to recalculate the formulas:

F9: Calculates the formulas in all open workbooks.

Shift+F9: Calculates only the formulas in the active worksheet. Other worksheets in the same workbook aren’t calculated.

Ctrl+Alt+F9: Forces a complete recalculation of all formulas.

Ctrl+Alt+Shift+F9: Rebuilds the calculation dependency tree and performs a complete recalculation.

Note

Excel’s Calculation mode isn’t specific to a particular worksheet. When you change the Calculation mode, it affects all open workbooks, not just the active workbook.

Using Advanced Naming Techniques

Using range names can make your formulas easier to understand and modify and even help prevent errors. Dealing with a meaningful name such as `AnnualSales` is much easier than dealing with a range reference, such as `AB12 : AB68`.



See Chapter 4, “Working with Excel Ranges and Tables,” for basic information regarding working with names.

Excel offers a number of advanced techniques that make using names even more useful. We discuss these techniques in the sections that follow. This information is for those who are interested in exploring some of the aspects of Excel that most users don't even know about.

Using names for constants

Many Excel users don't realize that you can give a name to an item that doesn't appear in a cell. For example, if formulas in your worksheet use a sales tax rate, you would probably insert the tax rate value into a cell and use this cell reference in your formulas. To make things easier, you would probably also name this cell something similar to SalesTax.

Here's how to provide a name for a value that doesn't appear in a cell:

1. Choose **Formulas** \Rightarrow **Defined Names** \Rightarrow **Define Name**. The New Name dialog box appears.
2. Enter the name (in this case, `SalesTax`) into the Name field.
3. Select a scope in which the name will be valid (either the entire workbook or a specific worksheet).
4. Click the Refers To text box, delete its contents, and replace the old contents with a value (such as `.075`).
5. (Optional) Use the Comment box to provide a comment about the name.
6. Click **OK** to close the New Name dialog box and create the name.

You just created a name that refers to a constant rather than a cell or range. Now if you type `=SalesTax` into a cell that's within the scope of the name, this simple formula returns `0.075`—the constant that you defined. You can also use this constant in a formula, such as `=A1*SalesTax`.

TIP

A constant also can be text. For example, you can define a constant for your company's name.

NOTE

Named constants don't appear in the Name box or in the Go To dialog box. This makes sense because these constants don't reside anywhere tangible. They do appear in the drop-down list that's displayed when you enter a formula. This is handy because you use these names in formulas.

Using names for formulas

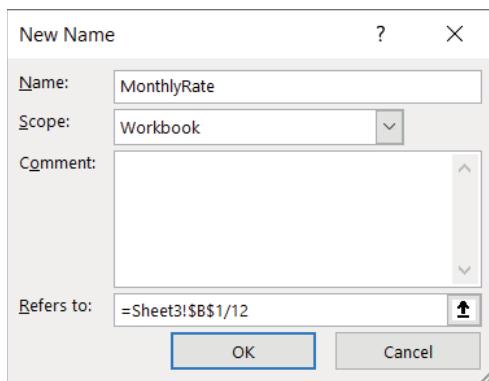
In addition to creating named constants, you can create named formulas. Like a named constant, a named formula doesn't reside in a cell.

You create named formulas the same way that you create named constants—by using the New Name dialog box. For example, you might create a named formula that calculates the monthly interest rate from an annual rate; Figure 9.15 shows an example. In this case, the name `MonthlyRate` refers to the following formula:

`=Sheet3!B1/12`

FIGURE 9.15

Excel allows you to name a formula that doesn't exist in a worksheet cell.



When you use the name `MonthlyRate` in a formula, it uses the value in B1 divided by 12 on Sheet3. Notice that the cell reference is an absolute reference.

Naming formulas gets more interesting when you use relative references rather than absolute references. When you use the pointing technique to create a formula in the Refers To field of the New Name dialog box, Excel always uses absolute cell references—which is unlike its behavior when you create a formula in a cell.

For example, activate cell B1 on Sheet1 and create the name `Cubed` for the following formula:

`=Sheet1!A1^3`

In this example, the relative reference points to the cell to the left of the cell in which the name is used. Therefore, make certain that cell B1 is the active cell before you open the New Name dialog box; this is important. The formula contains a relative reference. When you use this named formula in a worksheet, the cell reference is always relative to the cell that contains the formula. For example, if you enter `=Cubed` into cell D12, cell D12 displays the contents of cell C12 raised to the third power. (C12 is the cell directly to the left of cell D12.)

Using range intersections

This section describes a concept known as *range intersections* (individual cells that two ranges have in common). Excel uses an intersection operator—a space character—to determine the overlapping references in two ranges. Figure 9.16 shows a simple example.

Part II: Working with Formulas and Functions

The formula in cell B9 is

=C1:C6 A3:D3

This formula returns 13, the value in cell C3—that is, the value at the intersection of the two ranges.

The intersection operator is one of three reference operators used with ranges. Table 9.4 lists these operators.

FIGURE 9.16

You can use a range intersection formula to determine values.

A	B	C	D	E
1	99	94	51	54
2	90	75	89	46
3	58	89	13	60
4	37	95	84	19
5	11	37	56	26
6	65	11	62	36
7				
8				
9	A3:D3			

TABLE 9.4 Reference Operators for Ranges

Operator	What It Does
:	(colon) Specifies a range.
,	(comma) Specifies the union of two ranges. This operator combines multiple range references into a single reference.
Space	Specifies the intersection of two ranges. This operator produces cells that are common to two ranges.

The real value of knowing about range intersections is apparent when you use names. Examine Figure 9.17, which shows a table of values. We selected the entire table and then chose **Formulas** \Rightarrow **Defined Names** \Rightarrow **Create from Selection** to create names automatically using the top row and the left column.

Excel created the following eight names:

North	=Sheet1!\$B\$2:\$E\$2	Quarter1	=Sheet1!\$B\$2:\$B\$5
South	=Sheet1!\$B\$3:\$E\$3	Quarter2	=Sheet1!\$C\$2:\$C\$5
East	=Sheet1!\$B\$4:\$E\$4	Quarter3	=Sheet1!\$D\$2:\$D\$5
West	=Sheet1!\$B\$5:\$E\$5	Quarter4	=Sheet1!\$E\$2:\$E\$5

FIGURE 9.17

Creating names for all values in a table

With these names defined, you can create formulas that are easy to read and use. For example, to calculate the total for Quarter 4, just use this formula:

```
=SUM(Quarter4)
```

To refer to a single cell, use the intersection operator. Move to any blank cell and enter the following formula:

```
=Quarter1 West
```

This formula returns the value for the first quarter for the West region. In other words, it returns the value that exists where the Quarter1 range intersects with the West range. Naming ranges in this manner can help you create very readable formulas.

Applying names to existing references

When you create a name for a cell or a range, Excel doesn't automatically use the name in place of existing references in your formulas. For example, suppose you have the following formula in cell F10:

```
=A1-A2
```

If you later define a name Income for A1 and Expenses for A2, Excel won't automatically change your formula to =Income-Expenses. Replacing cell or range references with their corresponding names is fairly easy, however.

To apply names to cell references in formulas after the fact, start by selecting the range that you want to modify. Then choose Formulas \Rightarrow Defined Names \Rightarrow Define Name \Rightarrow Apply Names. The Apply Names dialog box appears. Select the names you want to apply by clicking them and then click OK. Excel replaces the range references with the names in the selected cells.

Working with Formulas

In the following sections, we offer a few additional tips and pointers relevant to formulas.

Not hard-coding values

When you create a formula, think twice before you use any specific value in the formula. For example, if your formula calculates sales tax (which is 6.5%), you may be tempted to enter a formula, such as the following:

=A1*.065

A better approach is to insert the sales tax rate in a cell—and use the cell reference. Or you can define the tax rate as a named constant, using the technique presented earlier in this chapter. Doing so makes modifying and maintaining your worksheet easier. For example, if the sales tax rate changed to 6.75%, you would have to modify every formula that used the old value. If you store the tax rate in a cell, however, you simply change that one cell, and Excel updates all the formulas.

Using the Formula bar as a calculator

If you need to perform a quick calculation, you can use the Formula bar as a calculator. For example, enter the following formula but don't press Enter:

= (145*1.05) /12

If you press Enter, Excel enters the formula into the cell. But because this formula always returns the same result, you may prefer to store the formula's result rather than the formula itself. To do so, press F9 and watch the result appear in the Formula bar. Press Enter to store the result in the active cell. (This technique also works if the formula uses cell references or worksheet functions.)

Making an exact copy of a formula

When you copy a formula, Excel adjusts its cell references when you paste the formula to a different location. Sometimes you may want to make an exact copy of the formula. One way to do this is to convert the cell references to absolute values, but this isn't always desirable. A better approach is to select the formula in Edit mode and then copy it to the Clipboard as text. You can do this in several ways. Here's a step-by-step example of how to make an exact copy of the formula in A1 and copy it to A2:

1. **Double-click A1 (or press F2) to get into Edit mode.**
2. **Drag the mouse to select the entire formula.** You can drag from left to right or from right to left. To select the entire formula with the keyboard, press End, followed by Shift+Home.
3. **Choose Home ➔ Clipboard ➔ Copy (or press Ctrl+C).** This copies the selected text (which will become the copied formula) to the Clipboard.

4. Press Esc to leave Edit mode.
5. Select cell A2.
6. Choose Home \Rightarrow Clipboard \Rightarrow Paste (or press Ctrl+V) to paste the text into cell A2.

You can also use this technique to copy just part of a formula if you want to use that part in another formula. Just select the part of the formula that you want to copy by dragging the mouse and then use any of the available techniques to copy the selection to the Clipboard. You can then paste the text to another cell.

Formulas (or parts of formulas) copied in this manner won't have their cell references adjusted when they're pasted into a new cell. That's because the formulas are being copied as text, not as actual formulas.

TIP

You can also convert a formula to text by adding an apostrophe (') in front of the equal sign. Then copy the cell as usual and paste it to its new location. Remove the apostrophe from the pasted formula, and it will be identical to the original formula. Don't forget to remove the apostrophe from the original formula.

Converting formulas to values

In certain scenarios, you may want to use formulas to get to an answer and then convert those formulas to actual values. For example, if you use the RANDBETWEEN function to create a set of random numbers and you don't want Excel to recalculate those random numbers each time you press Enter, you can convert the formulas to values. Just follow these steps:

1. Select A1:A20.
2. Choose Home \Rightarrow Clipboard \Rightarrow Copy (or press Ctrl+C).
3. Choose Home \Rightarrow Clipboard \Rightarrow Paste Values (V).
4. Press Esc to cancel Copy mode.

Understanding and Using Array Formulas

IN THIS CHAPTER

Understanding legacy array formulas

Introduction to dynamic arrays

Understanding spill ranges

Exploring dynamic array functions

The ability to work with arrays is one of Excel's most powerful features. Unlike standard formulas where a single formula results in a single answer, array formulas can return a series of answers to a series of cells on your worksheet. While array formulas have always been part of Excel, Microsoft has introduced dynamic arrays, paving the way for a paradigm shift in how we leverage the power of array formulas.

In this chapter, you'll first be introduced to the concept of arrays and get a foundational understanding of how the old legacy array formulas work. Then you'll explore Excel's exciting new dynamic array functions and discover some of the ways these new functions can help you go far beyond the capabilities of traditional formulas.



Most of the examples in this chapter are available on this book's website at www.wiley.com/go/excel365bible. The filename is `Array Formulas.xlsx`.

NOTE

For the purposes of this chapter, the term *legacy* refers to traditional array formulas created by pressing Ctrl+Shift+Enter. Microsoft refers to these kinds of formulas as CSE formulas, but to avoid being overly verbose, we'll stick with the term *legacy array formulas*.

Before the release of the new dynamic arrays, the only way to tell Excel that a formula was an array formula was to wrap it in curly brackets. However, you couldn't simply enter the curly brackets manually. Pressing Ctrl+Shift+Enter was the only way Excel would recognize your formula as an array formula.

Understanding Legacy Array Formulas

The easiest way to grasp the concept of arrays is to imagine a collection of rows and columns. Picture a table containing five rows and one column. On the worksheet, we call that a range of cells, but if we somehow took that range and stored it in memory, it ceases to be a range and becomes an array. While in memory, any formula operation performed on the array is performed as a collective. That is to say, all items in that array would be affected by the formula in various ways (depending on the formula you're using).

In Excel, an array can be one-dimensional or two-dimensional. These dimensions correspond to rows and columns. For example, a one-dimensional array can be stored in a range that consists of one row (a horizontal array) or one column (a vertical array). A two-dimensional array can be stored in a rectangular range of cells (multiple rows and multiple columns).

Example of a legacy array formula

Figure 10.1 shows a simple worksheet set up to calculate product sales. Normally, you'd calculate the value in column D (total sales per product) with a formula such as the one that follows, and then you'd copy this formula down the column. This would obviously work, but your worksheet would require six formulas in column D.

=B2*C2

FIGURE 10.1

Column D contains formulas to calculate the total revenue for each product.

	A	B	C	D
1	Product	Units Sold	Unit Price	
2	AR-998	3	\$50	=B2*C2
3	BZ-011	10	\$100	1000
4	MR-919	5	\$20	100
5	TR-811	9	\$10	90
6	TS-333	3	\$60	180
7	ZL-001	1	\$200	200

We can reduce these six formulas down to just one by using a legacy multicell array formula. Follow the steps below to create a legacy array formula that calculates all six values shown in D2:D7 (refer to Figure 10.1).

1. Select an empty range of cells to hold the results of your formula. In this case, select the range D2:D7. Because you can't display more than one value in a single cell, you select six cells to make this array work.
2. Type the following formula:
`=B2:B7*C2:C7`
3. Press **Ctrl+Shift+Enter** to enter the formula. Yes, you normally press **Enter** to enter a standard formula, but because this is an array formula, you need to press **Ctrl+Shift+Enter**.

You'll notice the formula entered automatically populates all six selected cells. If you examine the Formula bar, you see the following:

{=B2:B7*C2:C7}

Curly brackets around a formula in a worksheet means that formula is a legacy array formula.

Editing legacy array formulas

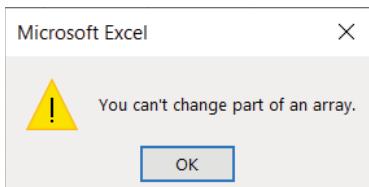
You can manually select the cells that contain a legacy array formula by using the normal cell selection procedures, or you can use either of the following methods:

- **Activate any cell in the array formula range.** Choose Home \Rightarrow Editing \Rightarrow Find & Select \Rightarrow Go To, or just press F5. The Go To dialog box appears. In the Go To dialog box, click the Special button and then choose the Current Array option. Click OK to close the dialog box.
- **Activate any cell in the array formula range and press Ctrl+/ (forward slash) to select the cells that make up the array.**

If an array formula occupies multiple cells, you must edit the entire range as though it were a single cell. The key point to remember is that you can't change just one element of a multicell array formula. If you attempt to do so, Excel displays the message shown in Figure 10.2.

FIGURE 10.2

Excel's warning message reminds you that you can't edit just one cell of a multicell array formula.



To edit a legacy array formula, select all the cells in the array range and activate the Formula bar as usual. (Click it or press F2.) Excel removes the curly brackets from the formula while you edit it. Edit the formula and then press Ctrl+Shift+Enter to enter the changes. All of the cells in the array will reflect your editing changes (and the curly brackets reappear).

CAUTION

If you accidentally press Ctrl+Enter (instead of Ctrl+Shift+Enter) after editing an array formula, the formula will be entered into each selected cell, but it will no longer be an array formula, and it will probably return an incorrect result or a #SPILL! error. Just select the cells again, press F2, and then press Ctrl+Shift+Enter.

Legacy Array Formulas: The Downside

Legacy array formulas are one of the least understood features of Excel. Consequently, if you plan to share a workbook with someone who may need to make modifications, you should probably avoid using these array formulas altogether. Someone who encounters one of these array formulas can easily destroy your work by attempting to edit them without remembering to use the magical Ctrl+Shift+Enter key combination. Outside of logical errors, this is probably the most common problem that users have with the legacy array formulas.

Another potential problem with array formulas is that they can sometimes slow your worksheet's recalculations, especially if you use very large arrays. On a faster system, this delay in speed may not be a problem. But, conversely, using an array formula is almost always faster than using a custom VBA function. See Chapter 40, "Creating Custom Worksheet Functions," for more information about creating custom VBA functions.

Introducing Dynamic Arrays

In many ways, legacy array formulas have always been impenetrable for the everyday Excel user. Sure, if you were a guru, you could make them perform magic. But the power of array formulas escaped the average Excel user; until now.

With dynamic arrays, Excel moves into a new era where you no longer need to be a formula guru to leverage the power of arrays. The underlying changes to Excel's calculation engine make legacy array formulas obsolete. You no longer need to differentiate array formulas with the magic Ctrl+Shift+Enter key combination.

To illustrate this new paradigm shift, Figure 10.3 shows a simple formula that references cells A2 : C7.

FIGURE 10.3

A simple formula that references a range

	A	B	C	D	E
1	Product	Units Sold	Unit Price		
2	AR-998	3	\$50	=A2:C7	
3	BZ-011	10	\$100		
4	MR-919	5	\$20		
5	TR-811	9	\$10		
6	TS-333	3	\$60		
7	ZL-001	1	\$200		
8					
9					

Simply pressing Enter will automatically propagate the formula to the surrounding cells (see Figure 10.4). The initial formula provides the array dimensions of six rows and three columns. Excel takes that information and outputs the results into a grid equivalent to the dimensions provided.

FIGURE 10.4

Excel automatically spills the results into the surrounding cells.

Product	Units Sold	Unit Price	E	F	G
AR-998	3	\$50	AR-998	3	50
BZ-011	10	\$100	BZ-011	10	100
MR-919	5	\$20	MR-919	5	20
TR-811	9	\$10	TR-811	9	10
TS-333	3	\$60	TS-333	3	60
ZL-001	1	\$200	ZL-001	1	200

Take a look at the Formula bar in Figure 10.4. Notice there are no curly brackets. There was no need to press the Ctrl+Shift+Enter key combination to tell Excel this is an array formula; it already knows!

This behavior of automatically propagating results to surrounding cells is called spilling. The area around an array formula is called the spill range. The spill range is determined by dimensions specified in your array formula. In the example illustrated in Figure 10.4, the array formula references an array that is six rows by three columns. Therefore, no matter where we put the formula, the spill range will be fixed to a grid that is six rows by three columns.

It's important to note that dynamic array behavior is fundamentally part of Excel's calculation engine. When any function uses an array that returns multiple values, the results will be output to a spill range. This even includes older functions that weren't designed to output arrays. For instance, the Formula bar in Figure 10.5 shows the following:

```
= SUM(B14:B19*C14:C19)
```

This formula uses a simple SUM function to multiply the sum of B14 : B19 by the sum of C14 : C19. Because dynamic arrays are now inherent to Excel's calculation engine, the arrays are automatically processed without the need to press Ctrl+Shift+Enter.

FIGURE 10.5

Dynamic arrays work with any traditional Excel function that accepts arrays as arguments.

Product	Units Sold	Unit Price	>>	E	F
AR-998	3	\$50			
BZ-011	10	\$100			
MR-919	5	\$20			
TR-811	9	\$10			
TS-333	3	\$60			
ZL-001	1	\$200			

Dynamic Arrays and Compatibility

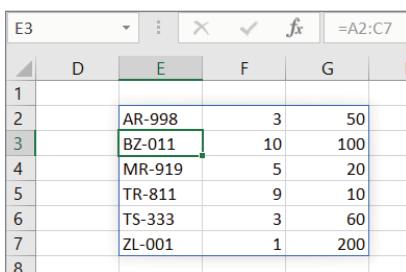
As of this writing, dynamic arrays will only be available to those with Office 365 subscriptions and those using the stand-alone (perpetual license) version of Office/Excel 2022 or later. Dynamic arrays will not be available to those using Excel 2019 or earlier versions. If you save a worksheet with dynamic arrays and open that workbook in an older version of Excel, the dynamic arrays will be converted to legacy array formulas. Excel does its best to retain fidelity, but be aware that the conversion may cause a loss of functionality or unexpected behavior. If your audience is using older versions of Excel, you'll want to fully test your dynamic formulas in their version of Excel to determine how they behave when converted.

Understanding spill ranges

When you enter a dynamic array, the results spill into adjacent cells. Excel visually demarcates a spill range with a solid blue line. All cells inside the spill range are effectively disabled except for the point of entry (the cell containing the original formula). Figure 10.6 shows an active spill range containing the results of a dynamic array. The original formula was entered into cell E2. None of the values in the spill range, except for cell E2, can be deleted, moved, or edited in any way.

FIGURE 10.6

Spill ranges will visually show a blue line around them.



A screenshot of a Microsoft Excel spreadsheet. The formula bar at the top shows '=A2:C7'. The main area contains a table with columns D, E, F, G, and H. Rows 1 through 7 have data, while row 8 is empty. Cell E2 contains the value 'AR-998'. Cell E3 contains the value 'BZ-011' and is highlighted with a green border, indicating it is the active cell. Cells E2 and E3 are part of a spill range, which is visually represented by a thick blue border around both cells. Other cells in the range, such as F2, F3, G2, and G3, contain numerical values (3, 10, 50, 100, 5, 20, 9, 10, 3, 60, 1, 200) and are also part of the spill range but do not have a green border.

E3	D	E	F	G	H
1					
2		AR-998	3	50	
3		BZ-011	10	100	
4		MR-919	5	20	
5		TR-811	9	10	
6		TS-333	3	60	
7		ZL-001	1	200	
8					

Attempting to enter data inside a spill range will cause what is known as a #SPILL! error. Figure 10.7 demonstrates a #SPILL! error caused by entering a value in the spill range. Removing the obstructing data from the spill range will immediately bring back the array values.

FIGURE 10.7

A spill error caused by an obstruction in the spill range

E2	D	E	F	G	H
1					
2		#SPILL!			
3					
4					
5		New Data Here			
6					
7					
8					

There are several reasons you may see a #SPILL! error. Take a moment to review the following list of scenarios in which a #SPILL! error can occur:

- **The spill range contains an obstruction:** This is the most common cause of a #SPILL! error. The spill range must be blank for your dynamic arrays to work. Any cell containing data will cause an error if it is in the path of the spill range.
- **The spill range extends beyond the worksheet end:** Try selecting cell A2 and entering the formula =B:B. You will get a #SPILL! error because the spill range needs more room than the worksheet allows.
- **Attempting to use dynamic arrays in a table object:** Excel table objects won't allow other objects inside them. Since spill ranges are essentially a kind of auto-expanding object themselves, you can't use dynamic arrays inside tables.
- **Out of memory:** If your dynamic array causes Excel to run out of memory, you'll get a #SPILL! error. In these cases, you'll need to rethink your formula and reference a smaller array.
- **The spill range hits a merged cell:** Your spill range cannot include merged cells.
- **Indeterminate size:** As mentioned earlier in this chapter, dynamic arrays allow you to use any existing Excel functions within them. However, dynamic array formulas have trouble keeping up with volatile functions such as RAND and RANDBETWEEN. Dynamic arrays inherently trigger calculation refreshes until the formula is fully resolved. Meanwhile, volatile functions will resize on each calculation pass, in turn causing the dynamic array to continue calculating. This causes a kind of continuous loop until Excel finally throws a #SPILL! error.
- **Unrecognized error:** In rare cases, Excel will encounter an error it can't reconcile, so it throws a #SPILL! error. In these cases, it's best to review your formula to make sure it has all the required arguments.

Referencing spill ranges

It's often useful to reference a spill range in other formulas. However, simply pointing subsequent formulas to a single cell in the spill range won't capture the entirety of the range. You need to use the spill range operator, designated by a hashtag (#). To explore this concept, see Figure 10.8. In the example shown, we want to capture the length of each string value in the spill range found in E2 : E7. You can see that entering the formula =LEN(E2) only gives us the length for the first value in the spill range.

FIGURE 10.8

Referencing a single cell in a spill range won't allow us to capture all values.

	D	E	F	G	H
1				String Length	
2		AR-90018		8	=Len(E2)
3		BZ-011			
4		MR-9198			
5		TR-81			
6		TS-3333			
7		ZL-001			
8					

In Figure 10.9, you can see in the Formula bar that we've added the spill range operator (#) to the cell reference, telling Excel to use the entire spill range. This has the effect of carrying the spill range over and effectively creating a new spill range for the LEN function.

FIGURE 10.9

Using the spill range operator to apply a function to the entire spill range

	D	E	F	G	H
1				String Length	
2		AR-90018		8	
3		BZ-011		6	
4		MR-9198		7	
5		TR-81		5	
6		TS-3333		7	
7		ZL-001		6	
8					

Figure 10.10 illustrates another example. Here, we're using the spill range operator with the COUNTA function to count all the values in the referenced spill range. Without the spill range operator, we would get a count of 1 for the first cell in the spill range.

FIGURE 10.10

Using the spill range operator to count all values in the referenced spill range

G10	D	E	F	G	H
9				Count all Values	
10		AR-90018		6	
11		BZ-011			
12		MR-9198			
13		TR-81			
14		TS-3333			
15		ZL-001			

Why do I see an @ symbol in my formulas?

Traditional Excel formulas had an inherent calculation behavior called implicit intersection. Implicit intersection means that a single value would be returned from an array of values. This ensured formulas would always return just one value because, before dynamic arrays came along, a cell could only contain one value.

Now that we have dynamic arrays, Excel is no longer limited to returning single values, so those of us with Office 365 subscriptions and those using the stand-alone (perpetual license) version of Office/Excel 2022 won't have implicit intersection applied.

However, if you open a workbook created in an older version of Excel, you may see the @ symbol in the formulas. The @ symbol is known as the implicit intersection operator. The @ operator is effectively a visual representation of the previously invisible implicit intersection behavior. Generally speaking, functions that return arrays (INDEX and OFFSET for instance), will be prefixed with @ if they were authored in an older version of Excel.

If the formula containing the @ operator returns a single value, you can safely remove the @ without effect. If, however, the formula returns an array, removing the @ will cause it to spill to the neighboring cells.

If you remove the @ from a formula that returns an array and later open the workbook in an older version of Excel, the formula will be converted to a legacy array formula.

Exploring Dynamic Array Functions

With the introduction of dynamic arrays, Microsoft released several new functions that leverage dynamic arrays to improve the ability to carry out complex formula operations with ease. These new functions can remove duplicates, extract unique values, filter data, dynamically sort data, and perform lookups. The following sections provide an overview of each of the new dynamic array functions.

The SORT function

The `SORT` function sorts the values in a given range in ascending or descending order with a formula. Results will be output into a spill range that automatically updates when values in the source range change. The `SORT` function takes four arguments: `[array]`, `[sort_index]`, `[sort_order]`, and `[by_col]`.

The `[array]` argument defines the source range to sort and is the only required argument. The following formula will sort the values in A2 : A15 in ascending order:

```
=SORT(A2:A15)
```

The `[sort_index]` argument allows you to specify the column to sort by. By default, the `SORT` function uses the first column in the given range. You can set the `[sort_index]` argument to a different column based on the column number. The following formula will sort the values in A2 : B15 by column B in ascending order:

```
=SORT(A2:B15, 2)
```

The `SORT` function will sort in ascending order by default. You can sort in descending order by adding the `[sort_order]` argument. Use `-1` as the `[sort_order]` argument to sort in descending order. The following formula sort the values in A2 : B15 by column B in descending order:

```
=SORT(A2:B15, 2, -1)
```

Figure 10.11 shows the `SORT` function in action. Here, we are sorting students in descending order by the change in test scores.

FIGURE 10.11

Using the `SORT` function to sort students in descending order by the change in test scores

The screenshot shows a Microsoft Excel spreadsheet with data in columns A through I. The data includes student names, pre-test scores, post-test scores, and the change in scores. The formula bar at the top shows =SORT(A2:D15,4,-1). The data is sorted by the 'Change' column in descending order. The sorted rows are as follows:

	A	B	C	D	E	F	G	H	I
1	Student	Pre-Test	Post-Test	Change					
2	Andy	56	67	11	Linda	45	68	23	
3	Beth	59	74	15	Michelle	71	92	21	
4	Cindy	98	92	-6	Eddy	81	100	19	
5	Duane	78	79	1	Beth	59	74	15	
6	Eddy	81	100	19	Isabel	54	69	15	
7	Francis	92	94	2	Andy	56	67	11	
8	Georgia	100	100	0	Kent	80	88	8	
9	Hilda	92	99	7	Hilda	92	99	7	
10	Isabel	54	69	15	Francis	92	94	2	
11	Jack	91	92	1	Duane	78	79	1	
12	Kent	80	88	8	Jack	91	92	1	
13	Linda	45	68	23	Georgia	100	100	0	
14	Michelle	71	92	21	Cindy	98	92	-6	
15	Nancy	94	83	-11	Nancy	94	83	-11	

The `SORT` function will sort rows by default. The last argument of the `SORT` function is the `[by_col]` argument. You can set the `[by_col]` argument to 1 (see Figure 10.12) to sort by columns instead of rows.

FIGURE 10.12

Sorting by columns

M	N	O	P	Q	R	S	T	U	V	W	X	Y	Z	AA
Student	Andy	Beth	Cindy	Duane	Eddy	Francis	Georgia	Hilda	Isabel	Jack	Kent	Linda	Michelle	Nancy
Score	56	59	98	78	81	92	100	92	54	91	80	45	71	94
	Georgia	Cindy	Nancy	Francis	Hilda	Jack	Eddy	Kent	Duane	Michelle	Beth	Andy	Isabel	Linda
	100	98	94	92	92	91	81	80	78	71	59	56	54	45

The `SORTBY` function

The `SORTBY` function sorts the contents of a range based on the values from other ranges. This function comes in handy when you need to apply sorting based on multiple columns. The `SORTBY` function takes three arguments: `[array]`, `[by_array]`, and `[sort_order]`.

The `[array]` argument defines the range to sort, while the `[by_array]` argument specifies a range to sort by. You can sort ascending or descending using the `[sort_order]` argument. Use 1 as the `[sort_order]` argument to sort in ascending order. Use -1 as the `[sort_order]` argument to sort in descending order. For instance, the following formula sorts the values in A2:C18 by the values in C2:C18 in descending order:

```
=SORTBY(A2:C18,C2:C18,-1)
```

That's neat, but it's not different than using this `SORT` function to sort the same range by the third column (column C) in descending order:

```
=SORT(A2:C18,3,-1)
```

Well, the real power of the `SORTBY` function is the ability to sort by multiple columns. The formula in Figure 10.13 shows the following:

```
=SORTBY(A2:C18,A2:A18,1,C2:C18,-1,B2:B18,1)
```

In this formula, we are sorting the range first by Market ascending (A2:C18), then by Sales Amount descending (C2:C18), then by Quarter ascending (B2:B18). The result gives us an output containing markets in alphabetical order, sorted by the quarters with the most sales amount.

Part II: Working with Formulas and Functions

FIGURE 10.13

Using the `SORTBY` function to apply a multi-column sort

The screenshot shows an Excel spreadsheet with two main sections. On the left, there is a data table with columns labeled 'Market' (A), 'Quarter' (B), 'Sales' (C), and 'Amount' (D). The data consists of 18 rows, each containing a market name and its corresponding quarter and sales amounts. On the right, there is a second section with columns E, F, and G. This section contains the same data as the first section, but it is sorted by the 'Sales' column in descending order. The formula bar at the top of the screen displays the formula `=SORTBY(A2:C18, A2:A18,1, C2:C18,-1, B2:B18,1)`. The 'Market' column header is highlighted in red, indicating it is the primary sort key.

Market	Quarter	Sales	Amount			
BUFFALO	Q2	66,845		C2:C18,-1, B2:	Q3	500,297
BUFFALO	Q3	500,297		BUFFALO	Q1	283,337
BUFFALO	Q1	283,337		BUFFALO	Q2	66,845
CALIFORNIA	Q2	137,401		CALIFORNIA	Q1	1,138,579
CALIFORNIA	Q3	78,755		CALIFORNIA	Q2	137,401
CALIFORNIA	Q1	1,138,579		CALIFORNIA	Q3	78,755
CANADA	Q3	192,717		CANADA	Q1	683,529
CANADA	Q1	683,529		CANADA	Q3	192,717
CHARLOTTE	Q2	1,170,341		CHARLOTTE	Q2	1,170,341
CHARLOTTE	Q3	126,723		CHARLOTTE	Q1	593,458
CHARLOTTE	Q1	593,458		CHARLOTTE	Q3	126,723
DALLAS	Q2	318,807		DALLAS	Q1	352,632
DALLAS	Q3	295,650		DALLAS	Q2	318,807
DALLAS	Q1	352,632		DALLAS	Q3	295,650
DENVER	Q2	4,312,564		DENVER	Q2	4,312,564
DENVER	Q3	116,030		DENVER	Q1	516,989
DENVER	Q1	516,989		DENVER	Q3	116,030

NOTE

The `SORTBY` function does not require that the columns you sort by (specified by the `[by_array]` arguments) be part of the source data. You can sort by another range in a separate table if you wish. However, the range you choose to sort by must have compatible dimensions. For instance, if your source data has 15 rows, the range by which you sort must also have 15 rows.

The UNIQUE function

The `UNIQUE` function extracts a list of distinct values from a range or array using just three arguments: `[array]`, `[by_col]`, and `[exactly_once]`. The `[array]` argument is the only required argument. The following formula illustrates the most basic use of the `UNIQUE` function. This formula will extract the unique values from cells A1 : A10 to a spill range that automatically updates when values in the source range change:

```
=UNIQUE(A1:A10)
```

The `UNIQUE` function works on rows by default. You can, however, tell Excel to extract unique values from columns by setting the `[by_col]` argument to 1. In the following example, the `UNIQUE` function is used to extract the unique values from A1:J1 into a new spill range:

```
=UNIQUE(A1:J1,1)
```

The last argument, [exactly_once], tells Excel to extract values that appear only once in the given array. Figure 10.14 illustrates the difference between using a basic UNIQUE and using UNIQUE with the [exactly_once] argument. As you can see, the values in column G show only those markets that appear one time in the source range (CANADA and TULSA).

FIGURE 10.14

Adding the [exactly_once] argument extracts values that appear only one time in the given range.

A	B	C	D	E	F	G
Market	Quarter	Sales Amount		=UNIQUE(A2:A13)		=UNIQUE(A2:A13,,1)
BUFFALO	Q2	66,845	>>	BUFFALO	>>	CANADA
BUFFALO	Q1	283,337		CALIFORNIA		TULSA
CALIFORNIA	Q2	137,401		CANADA		
CALIFORNIA	Q3	78,755		CHARLOTTE		
CANADA	Q3	192,717		DALLAS		
CHARLOTTE	Q3	126,723		DENVER		
CHARLOTTE	Q1	593,458		TULSA		
DALLAS	Q2	318,807				
DALLAS	Q1	352,632				
DENVER	Q2	4,312,564				
DENVER	Q3	116,030				
TULSA	Q1	516,989				

The RANDARRAY function

The RANDARRAY function generates an array of random numbers. This function comes in handy if you need to generate random numbers for modeling and simulation exercises. The RANDARRAY function takes five arguments: [rows], [columns], [min], [max], and [integer]. None of the arguments are required. Entering the following formula in cell A1 will return a single random decimal between 0 and 1:

```
=RANDARRAY()
```

Passing the number 10, the [rows] argument will generate 10 rows of random decimal numbers between 0 and 1:

```
=RANDARRAY(10)
```

The following formula adds the [columns] argument to generate 10 rows and 5 columns of random decimals between 0 and 1:

```
=RANDARRAY(10, 5)
```

Use the [min] and [max] arguments to specify the minimum and maximum random number to generate. The following formula will generate 10 rows of random decimals ranging from 1 to 5. Note that the second argument, [columns], is not used:

```
=RANDARRAY(10, ,1,5)
```

You can tell Excel to return whole numbers instead of decimals by setting the last argument, [integer], to 1. The following formula generates 10 rows and 5 columns of random *whole* numbers ranging from 20 to 50:

```
=RANDARRAY(10,5,20,50,1)
```

The SEQUENCE function

The SEQUENCE function can be used to generate a list of sequential numbers. This function can be used on its own to create sequential numbers directly in a spill range or used within another formula to perform more complex operations. The SEQUENCE function accepts four arguments: [rows], [columns], [start], and [step].

The [rows] argument is the only required argument, telling Excel the number of rows of sequential numbers to generate. Without any other arguments, Excel will generate a list of numbers that starts at 1, by default, and ends when the number of rows specified is reached. For example, entering the following in cell A1 will generate a list of numbers from 1 to 12 (the number of rows specified):

```
=SEQUENCE(12)
```

You can add the [columns] argument if you want to create multiple columns of numbers. The following formula will generate a sequential list of numbers in two columns. Again, the default behavior dictates that the list starts at 1 and continues until the number of rows and columns specified is reached:

```
=SEQUENCE(12,2)
```

For a more complex set of sequential numbers, you can use the [start] argument to dictate at which number the sequential list starts and the [step] argument to specify how the list progresses as each row is generated. Entering the following formula will generate 10 rows of numbers, starting at 5 and progressing by fives until all 10 rows are generated. To see this formula in action, check out column F in Figure 10.15.

```
=SEQUENCE(10,,5,5)
```

You can enter a negative number in the [step] argument to generate a sequential list that runs backward. Column H in Figure 10.15 demonstrates this using a SEQUENCE function that counts backward from 50.

As mentioned earlier in this section, the SEQUENCE function can be used within another formula to perform more complex operations. Take Figure 10.16, for example. You can see in column D that the SEQUENCE function is being used in conjunction with the LARGE function to extract the three largest numbers in A19:A28.

Figure 10.17 demonstrates another useful application of the SEQUENCE function. Here, we're using the SEQUENCE function within the DATE function to generate a list of months in the year 2022.

FIGURE 10.15

Example usage of the SEQUENCE function

A	B	C	D	E	F	G	H
1	=SEQUENCE(12)	=SEQUENCE(12,2)		=SEQUENCE(10,,5,5)		=SEQUENCE(10,,50,-5)	
2	1	1	2		5	50	
3	2	3	4		10	45	
4	3	5	6		15	40	
5	4	7	8		20	35	
6	5	9	10		25	30	
7	6	11	12		30	25	
8	7	13	14		35	20	
9	8	15	16		40	15	
10	9	17	18		45	10	
11	10	19	20		50	5	
12	11	21	22				
13	12	23	24				

FIGURE 10.16

Getting a list of the three largest numbers in A19:A28

D19		X	✓	fx	=LARGE(A19:A28,SEQUENCE(3))
A	B	C	D	E	F
19	12	>>	13		
20	-5		12		
21	3		8		
22	2				
23	0				
24	6				
25	13				
26	7				
27	4				
28	8				

FIGURE 10.17

Getting a list of dates representing each month in 2022

=DATE(2022,SEQUENCE(12),1)
1/1/2022
2/1/2022
3/1/2022
4/1/2022
5/1/2022
6/1/2022
7/1/2022
8/1/2022
9/1/2022
10/1/2022
11/1/2022
12/1/2022

The FILTER function

The FILTER function extracts matching records from a dataset based on criteria specified in arguments. Results are output to a spill range and are automatically updated when the source data changes. This function becomes quite useful when you want to report on a subset of a larger data table. The FILTER function takes three arguments: [array], [include], and [if_empty].

The [array] argument points to the source data you are extracting from, the [include] argument specifies the criteria that source records must meet in order to be extracted, and the [if_empty] argument defines what should be returned if no data matches the criteria.

Take a moment to examine Figure 10.18 to see how the FILTER function works. Looking at the Formula bar, you'll see we are extracting all records from A2:D15 (the [array] argument) if the value D2:D15 is greater than 10 (the [include] argument). If no records match, then the formula will return "No Matches" (the [if_empty] argument).

FIGURE 10.18

Filtering records where the change value is greater than 10

The screenshot shows a Microsoft Excel spreadsheet. The formula bar at the top displays the formula =FILTER(A2:D15,D2:D15>10,"No Matches"). The main area contains a data table with columns labeled 'Student', 'Pre-Test', 'Post-Test', and 'Change'. Rows 1 through 15 contain data for various students. Row 15 is highlighted in blue. The 'Change' column shows values such as 11, -6, 23, etc. The formula bar also shows the current cell is F3. The data table starts at row 2:

	A	B	C	D	E	F	G	H	I
1	Student	Pre-Test	Post-Test	Change					
2	Andy	56	67	11					
3	Beth	59	74	15	Andy	56	67	11	
4	Cindy	98	92	-6	Beth	59	74	15	
5	Duane	78	79	1	Linda P.	45	68	23	
6	Linda P.	45	68	23	Michelle	71	92	21	
7	Michelle	71	92	21	Linda J.	81	100	19	
8	Nancy	94	83	-11	Isabel	54	69	15	
9	Linda J.	81	100	19					
10	Francis	92	94	2					
11	Georgia	100	100	0					
12	Roland	91	92	1					
13	Kent	80	88	8					
14	Hilda	92	99	7					
15	Isabel	54	69	15					

TIP

Although the [if_empty] argument is technically optional, it's a best practice to always specify a value to return if the FILTER function returns no matches. Omitting the [if_empty] argument can potentially lead to a #CALC! error if none of the records in your source range meets the specified criteria.

Instead of hard-coding the [include] argument, you may find it useful to reference a cell that will hold the criteria for your FILTER function. For example, the formula used in Figure 10.19 is essentially the same as the one in Figure 10.18 except we are getting the

criteria from F1. With this setup, we can enter different values in F1 to see the results in the FILTER dynamically change.

FIGURE 10.19

Getting the FILTER criteria from cell F1

=FILTER(A2:D15,D2:D15>F1,"No Matches")				
E	F	G	H	I
	15			
Linda P.	45	68	23	
Michelle	71	92	21	
Linda J.	81	100	19	

It's often useful to wrap the FILTER function inside of a SORT function to sort the filtered results. Figure 10.20 demonstrates sorting filtered results by the first column in the filtered array.

FIGURE 10.20

Combining SORT with FILTER to sort results

=SORT(FILTER(A2:D15,D2:D15>F1,"No Matches"),1)				
E	F	G	H	I
	10			
Andy	56	67	11	
Beth	59	74	15	
Isabel	54	69	15	
Linda J.	81	100	19	
Linda P.	45	68	23	
Michelle	71	92	21	

Using multiple conditions with the FILTER function

There may be situations in which you'll need to filter a dataset based on multiple criteria. To do so, you simply need to wrap each criterion in parentheses and separate them with an asterisk (*). For instance, Figure 10.21 expands on the formula from Figure 10.19 by adding a second criteria for name. Note that each criterion is wrapped in its own set of parentheses.

FIGURE 10.21

Using multiple filter conditions

=FILTER(A2:D15,(D2:D15>F1)*(A2:A15=H1),"No Matches")						
E	F	G	H	I	J	
	10		Beth			
Beth	59	74	15			

Part II: Working with Formulas and Functions

So, what's exactly going on here? Well, behind the scenes, the two conditions are evaluated for each row of the dataset, resulting in a TRUE or FALSE for each respective condition. In Figure 10.22, we've broken the conditions to separate formulas to clearly illustrate the results from each.

FIGURE 10.22

Criteria evaluation behind the scenes

	A	B	C	D	E	F	G	H
1	Student	Pre-Test	Post-Test	Change		10		Beth
2	Andy	56	67	11	=	=D2:D15>F1)		=A2:A15=H1)
3	Beth	59	74	15		TRUE		FALSE
4	Cindy	98	92	-6		TRUE		TRUE
5	Duane	78	79	1		FALSE		FALSE
6	Linda P.	45	68	23		FALSE		FALSE
7	Michelle	71	92	21		TRUE		FALSE
8	Nancy	94	83	-11		TRUE		FALSE
9	Linda J.	81	100	19		FALSE		FALSE
10	Francis	92	94	2		TRUE		FALSE
11	Georgia	100	100	0		FALSE		FALSE
12	Roland	91	92	1		FALSE		FALSE
13	Kent	80	88	8		FALSE		FALSE
14	Hilda	92	99	7		FALSE		FALSE
15	Isabel	54	69	15		FALSE		FALSE
16						TRUE		FALSE

In Excel, TRUE is equivalent to 1 and FALSE is equivalent to 0. The asterisk in the FILTER function (see Figure 10.21) will effectively multiply the TRUE and FALSE (1 and 0) results for each row, resulting in a 1 only for those records that show TRUE for both conditions. The asterisk can be thought of as an AND statement. Return results where the first condition is TRUE *and* the second condition is TRUE.

To return results if the first condition is TRUE *or* the second condition is TRUE, use the plus (+) operator. In the formula shown in Figure 10.23, Excel returns all records where the change value is greater than 20 *or* the student is Beth.

Filtering records that contain a search term

You can adjust the condition in your FILTER function to find records that contain a search term. To do so, you can use the SEARCH function to find all values where your search term is contained in the column you're evaluating. In Figure 10.24, you can see an example of this technique. Here, we're filtering for all records where the student name contains Linda. Wrapping the SEARCH function inside ISNUMBER converts the search results into a number the dynamic array can read.

FIGURE 10.23

Using the + operator to return results if the first condition is TRUE or the second condition is TRUE

	A	B	C	D	E	F	G	H	I	J
1	Student	Pre-Test	Post-Test	Change		20		Beth		
2	Andy	56	67	11						
3	Beth	59	74	15	Beth	59	74	15		
4	Cindy	98	92	-6	Linda P.	45	68	23		
5	Duane	78	79	1	Michelle	71	92	21		
6	Linda P.	45	68	23						
7	Michelle	71	92	21						
8	Nancy	94	83	-11						
9	Linda J.	81	100	19						
10	Francis	92	94	2						
11	Georgia	100	100	0						
12	Roland	91	92	1						
13	Kent	80	88	8						
14	Hilda	92	99	7						
15	Isabel	54	69	15						

FIGURE 10.24

Filtering records where values are contained in a search term

	A	B	C	D	E	F	G	H	I	
1	Student	Pre-Test	Post-Test	Change		Linda				
2	Andy	56	67	11						
3	Beth	59	74	15	Beth	59	74	15		
4	Cindy	98	92	-6	Linda P.	45	68	23		
5	Duane	78	79	1	Linda J.	81	100	19		
6	Linda P.	45	68	23						
7	Michelle	71	92	21						
8	Nancy	94	83	-11						
9	Linda J.	81	100	19						
10	Francis	92	94	2						
11	Georgia	100	100	0						
12	Roland	91	92	1						
13	Kent	80	88	8						
14	Hilda	92	99	7						
15	Isabel	54	69	15						

The XLOOKUP function

The XLOOKUP function is designed to be the successor to the traditional VLOOKUP and HLOOKUP functions you may be used to. With more flexible options like approximate matching and wildcard matching, this function is truly an improved version of Excel's lookup

Part II: Working with Formulas and Functions

capability. XLOOKUP accepts a whopping six arguments: [lookup_value], [lookup_array], [return_array], [not_found], [match_mode], and [search_mode].

The minimum arguments you need for XLOOKUP to work are [lookup_value], [lookup_array], and [return_array]. To understand these arguments, take a gander at Figure 10.25. In this example, we are attempting to populate the Customer Type column (column C) based on the matching values in the table found in columns E and F. A quick look at the Formula bar shows the following formula:

```
=XLOOKUP(B2:B19,E2:E5,F2:F5,"No Match")
```

This formula is telling Excel to look up the values in column B (the [lookup_value]), match them to the values in column E (the [lookup_array]), and return the matching value in column F (the [return_array]). The final argument ([not_found]) defines what should be returned if no data matches the lookup criteria. Although the [not_found] argument is optional, it's typically best practice to include it to avoid the ugly #N/A errors you get when a value is not found.

FIGURE 10.25

A basic XLOOKUP function to find customer type based on revenue

	A	B	C	D	E	F
1	CustomerName	Revenue	Customer Type		Revenue	Customer Type
2	Communication Connections	\$5,000	C		30000	A
3	Aaron Fitz Electrical	\$30,000	A		10000	B
4	Astor Suite	\$10,000	B		5000	C
5	Blue Yonder Airlines	\$9,306	No Match		1000	D
6	Central Communications LTD	\$11,975	No Match			
7	Computerized Phone Systems	\$59	No Match			
8	Country View Estates	\$19	No Match			
9	Lawrence Telemarketing	\$50,024	No Match			
10	Leisure & Travel Consultants	\$92	No Match			
11	Magnificent Office Images	\$9,489	No Match			
12	Mahler State University	\$11,161	No Match			

NOTE

When building your XLOOKUP function, keep in mind that the [lookup_array] argument must have the same dimensions as the [return_array] argument, that is, the same number of rows and columns.

As you look at Figure 10.25, you'll notice the XLOOKUP function found a customer type only for those records that exactly match the revenue in the [lookup_array] (column E). In many cases, you won't find an exact match, especially when attempting to match numbers like revenue. In our example, we want to perform an approximate match on our records to catch

even those revenues that do not exactly match the values in our lookup table. To do so, we can add the [match_mode] argument.

Entering a comma after the [not_found] argument will bring up a list of choices (see Figure 10.26). The one you choose will depend on what exactly you're attempting to achieve. It's best to experiment with each option to see which option is right for you. In this case, we want to apply the option for Exact Match or Next Smaller Item. Selecting the appropriate option will immediately retrieve the valid customer types. Figure 10.27 shows the result of the applied approximate matching.

FIGURE 10.26

Enter a comma to see choices for the next argument.

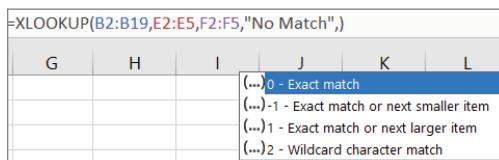


FIGURE 10.27

XLOOKUP results using approximate matching

	A	B	C	D	E	F
1	CustomerName	Revenue	Customer Type		Revenue	Customer Type
2	Communication Connections	\$5,000	C		30000	A
3	Aaron Fitz Electrical	\$30,000	A		10000	B
4	Astor Suite	\$10,000	B		5000	C
5	Blue Yonder Airlines	\$9,306	C		1000	D
6	Central Communications LTD	\$11,975	B			
7	Computerized Phone Systems	\$59	No Match			
8	Country View Estates	\$19	No Match			
9	Lawrence Telemarketing	\$50,024	A			
10	Leisure & Travel Consultants	\$92	No Match			
11	Magnificent Office Images	\$9,489	C			
12	Mahler State University	\$11,161	B			
13	Metropolitan Fiber Systems	\$10,686	B			

In rare cases, you can use the [search_mode] argument to dictate how Excel performs the search for the XLOOKUP. For this argument, you have the following options:

- 1:** This is the default search mode and will be the most appropriate in most of your lookup scenarios.
- 1:** Search from last value in the array up (essentially reverse order of the default behavior).

- 2:** Binary search on values that are already sorted in ascending order. This mode is used to improve performance on very large arrays. Binary searches are fast, but if your data is not sorted as prescribed, your XLOOKUP can return invalid results.
- 2:** Binary search on values that are already sorted in descending order. Again, this mode is used to improve performance, but you'll want to ensure your data is not sorted as prescribed or your XLOOKUP will return invalid results.

NOTE

When using XLOOKUP to match across workbooks, both workbooks must be open or Excel will return a #REF! error.

XLOOKUP with wildcards

One of the options for the [match_mode] argument is wildcard character match (see Figure 10.26). Wildcards are special characters that enable complex searches using approximate matching. Excel allows the use of three different wildcard characters, each with its own purpose.

Asterisk (*): The asterisk wildcard is used to tell your formula to look for a portion of text no matter what comes before or after. For example, searching for *hotel will look for any value that ends with the word *hotel*, no matter what comes before. Searching for *hotel** will look for any value that starts with *hotel*, no matter what comes after. And as you may have guessed, searching for **hotel** will look for any value containing the word *hotel*.

Question mark (?): The question mark wildcard is used to tell a formula to look for any one character within the text. For instance, searching for p?ace would return both *peace* and *place*. You can use multiple question mark wildcards in a search (?onder and sm?itt?n are both valid searches). In addition, you can combine the asterisk and question mark wildcards to create a more complex search. For example, searching for *vis??* would return any text string containing the word *visor* or *vision*.

Tilde (~): The tilde wildcard is useful when you need to include a character that is itself a wildcard. For example, if you need to search for any text string that ends in a question mark, you'll need to preface the question mark with a tilde because the question mark itself is a wildcard (*~?). The same goes for the asterisk. To search for any text string that starts with an asterisk, you would use ~**.

Figure 10.28 demonstrates the use of wildcard characters. In column J, XLOOKUP searches for the Revenue (B2:B19) where the Customer Name contains the letters LTD. In column L, XLOOKUP searches for the Revenue where the Customer Name contains f?tz.

FIGURE 10.28

Using wildcard characters to perform complex searches

	A	B	D	J	K	L	M
1	CustomerName	Revenue		*LTD*	*f?tz*		
2	Communication Connections	\$5,000		\$11,975		\$30,000	
3	Aaron Fitz Electrical	\$30,000					
4	Astor Suite	\$10,000					
5	Blue Yonder Airlines	\$9,306					
6	Central Communications LTD	\$11,975					
7	Computerized Phone Systems	\$59					

Using Formulas for Common Mathematical Operations

IN THIS CHAPTER

Calculating percentages

Rounding numbers

Counting values in a range

Most Excel analysts working in the corporate world will be asked to perform mathematical operations that provide insight into key operational metrics. In this chapter, you'll get to explore a few mathematical operations commonly used in the world of business analytics.

Calculating Percentages

Calculations such as percent of totals, variance to budget, and running totals are the cornerstone of any basic business analysis. In the following sections, you'll explore some examples of formulas that will help with these types of analyses.



This book's website, www.wiley.com/go/excel365bible, includes a copy of the sample workbook for this chapter. The file is named Mathematical Formulas.xlsx.

Calculating percent of goal

When someone asks you to calculate a percent of goal, they are simply asking you to compare the actual performance against a stated goal. The math involved in this calculation is simple: divide the actual by the goal. This will give you a percentage value that represents how much of the goal has been achieved. For instance, if your goal is to sell 100 widgets and you sell 80 widgets, your percent of goal is 80% ($80/100$).

Figure 11.1 shows a list of regions with a column for goals and a column for actuals. Note that the formula in cell E5 simply divides the value in the Actual column by the value in the Goal column.

=D5/C5

Part II: Working with Formulas and Functions

FIGURE 11.1

Calculating the percent of goal

	A	B	C	D	E
1					
2					
3					
4	Region	Goal	Actual	Percent of Goal	
5	North	\$509,283	\$553,887	=D5/C5	
6	South	\$483,519	\$511,115	106%	
7	East	\$640,603	\$606,603	95%	
8	West	\$320,312	\$382,753	119%	

There isn't much to this formula. You're simply using cell references to divide one value by another. You can enter the formula one time in the first row (cell E5 in this case) and then copy that formula down to every other row in your table.

Alternatively, if you need to compare actuals to a common goal, you can set up a model like the one shown in Figure 11.2. In this model, each region does not have its own goal. Instead, you're comparing the values in the Actual column to a single goal found in cell B3.

=C6 / \$B\$3

FIGURE 11.2

Calculating the percent of goal using a common goal

	A	B	C	D
1				
2		Common Goal		
3		\$700,000		
4				
5	Region	Actual	Percent of Goal	
6	North	\$553,887	=C6/\$B\$3	
7	South	\$511,115	73%	
8	East	\$606,603	87%	
9	West	\$382,753	55%	

Note that the cell reference to the common goal is entered as an absolute reference (\$B\$3). Using the dollar symbols locks the reference to the goal in place, ensuring that the cell reference pointing to the common goal does not adjust as you copy the formula down.



See Chapter 9, “Introducing Formulas and Functions,” for more information on absolute and relative cell references.

Calculating percent variance

A *variance* is an indicator of the difference between one number and another. To understand this, imagine that you sold 120 widgets one day, and on the next day, you sold 150 widgets.

Chapter 11: Using Formulas for Common Mathematical Operations

The difference in sales in actual terms is easy to see; you sold 30 more widgets on the second day, and 150 widgets minus 120 widgets gives you a unit variance of +30.

So, what is a percent variance? This is essentially the percentage difference between the benchmark number (120) and the new number (150). You calculate the percent variance by subtracting the benchmark number from the new number and then dividing that result by the benchmark number; In this example, $(150-120)/120 = 25\%$. The percent variance tells us that you sold 25% more widgets than on the previous day.

Figure 11.3 demonstrates how to translate this into a formula. The formula in E4 calculates the percent variance between current year sales and previous year sales.

$$= (D4 - C4) / C4$$

FIGURE 11.3

Calculating the percent variance between current year sales and previous year sales

	A	B	C	D	E
1					
2					
3	Region	Prior Year	Current Year	Percent Variance	
4	North	\$509,283	\$553,887	= (D4-C4)/C4	
5	South	\$483,519	\$511,115	6%	
6	East	\$640,603	\$606,603	-5%	
7	West	\$320,312	\$382,753	19%	

The one thing to note about this formula is the use of parentheses. By default, Excel's order of operations states that division must be done before subtraction. But if we let that happen, we would get an erroneous result. Wrapping the first part of the formula in parentheses ensures that Excel performs subtraction before the division.

You can simply enter the formula one time in the first row (cell E4 in this case) and then copy that formula down to every other row in your table.



See Chapter 9, "Introducing Formulas and Functions," for a detailed explanation of the order of operator precedence.

An alternative formula for calculating percent variance is simply to divide the current year's sales by the previous year's sales and then subtract 1. Because Excel performs division operations before subtraction, you don't have to use parentheses with this alternative formula.

$$= D4 / C4 - 1$$

Calculating percent variance with negative values

In the previous section, "Calculating percent variance," you discovered how to calculate a percent variance. This works beautifully in most cases. However, when the benchmark value is a zero or less, the formula breaks down.

Part II: Working with Formulas and Functions

For example, imagine you're starting a business and expect to take a loss during the first year. So, you give yourself a budget of negative \$10,000. Now imagine that after your first year, you made money, earning \$12,000. Calculating the percent variance between your actual revenue and budgeted revenue would give you -220%. You can try it on a calculator. 12,000 minus -10,000 divided by -10,000 equals -220%.

How can you say that your percent variance is -220% when you clearly made money? Well, the problem is that when your benchmark value is a negative number, the math inverts the results, causing numbers to look wacky. This is a real problem in the corporate world, where budgets can often be negative values.

The fix is to leverage the ABS function to negate the negative benchmark value:

$$= (C4 - B4) / \text{ABS}(B4)$$

Figure 11.4 uses this formula in cell E4, illustrating the different results you get when using the standard percent variance formula and the improved percent variance formula.

FIGURE 11.4

Using the ABS function will give you an accurate percent variance when dealing with negative values.

Excel Formula Bar: =(C4 - B4) / ABS(B4)				
A	B	C	D	E
1				
2				
3	Budget	Actual	Standard Percent Variance	Improved Percent Variance
4	-10,000	12,000	-220%	220%

Excel's ABS function returns the absolute value for any number you pass to it. Entering =ABS(-100) into cell A1 would return 100. The ABS function essentially makes any number a non-negative number. Using ABS in this formula negates the effect of the negative benchmark (the negative 10,000 budget in our example) and returns the correct percent variance.

TIP

You can safely use this formula for all your percent variance needs because it works with any combination of positive and negative numbers.

Calculating a percent distribution

Percent distribution is a measure of how a metric (such as total revenue) is distributed among the components that make up the total. As you can see in Figure 11.5, the calculation is relatively simple. You divide each component by the total. In this example, we have a cell that contains the total revenue (cell C9). We then divide each region's revenue by the total to get a percent distribution for each region.

Chapter 11: Using Formulas for Common Mathematical Operations

FIGURE 11.5

Calculating a percent distribution of revenue across regions

A	B	C	D
1			
2	Region	Revenue	Percent of Total
3	North	\$7,626	=C3/\$C\$9
4	South	\$3,387	18%
5	East	\$1,695	9%
6	West	\$6,457	34%
7			
8			
9	Total	\$19,165	

There isn't much to this formula. You're simply using cell references to divide each component value by the total. The one thing to note is that the cell reference to the total is entered as an absolute reference (`C9`). Using the dollar symbols locks the reference in place, ensuring that the cell reference pointing to the total value does not adjust as you copy the formula down.

You don't have to dedicate a separate cell to an actual total value. You can simply calculate the total on the fly within the percent distribution formula. Figure 11.6 demonstrates how you can use the `SUM` function in place of a cell dedicated to holding a total value. The `SUM` function adds together any numbers you pass to it.

FIGURE 11.6

Calculating percent distribution with the SUM function

A	B	C	D
1			
2	Region	Revenue	Percent of Total
3	North	\$7,626	=C3/SUM(\$C\$3:\$C\$6)
4	South	\$3,387	18%
5	East	\$1,695	9%
6	West	\$6,457	34%

Again, note the use of absolute references in the `SUM` function. This ensures that the `SUM` range stays locked as you copy the formula down.

`=C3 / SUM (C3 : C6)`

Calculating a running total

Some organizations like to see a running total as a mechanism to analyze the changes in a metric as a period of time progresses. Figure 11.7 illustrates a running total of units sold for January through December. The formula used in cell D3 is copied down to each month.

`=SUM (C3 : C3)`

Part II: Working with Formulas and Functions

FIGURE 11.7

Calculating a running total

	A	B	C	D
1				
2			Units Sold	Running Total
3		January	78	=SUM(\$C\$3:C3)
4		February	63	141
5		March	38	179
6		April	17	196
7		May	84	280
8		June	63	343
9		July	32	375
10		August	20	395
11		September	98	493
12		October	63	556
13		November	75	631
14		December	75	706

In this formula, the SUM function is used to add up all the units from cell C3 to the current row. The trick to this formula is the absolute reference (\$C\$3). Placing an absolute reference in the reference for the first value of the year locks that value down. This ensures that as the formula is copied down, the SUM function will always capture and add the units from the first value to the value on the current row.

Applying a percent increase or decrease to values

A common task for an Excel analyst is to apply a percentage increase or decrease to a given number. For instance, when applying a price increase to a product, you would typically raise the original price by a certain percent. When giving a customer a discount, you would decrease that customer's rate by a certain percent.

Figure 11.8 illustrates how to apply a percent increase and decrease using a simple formula. In cell E5, we're applying a 10% price increase to Product A. In cell E9, we're giving a 20% discount to Customer A.

FIGURE 11.8

Applying a percent increase and decrease using a simple formula

	A	B	C	D	E	F
1						
2						
3						
4			Unit Cost	Price Increase	Final Price	
5		Product A	100	10%	=C5*(1+D5)	
6						
7						
8			Cost per Service	Percent Discount	Discounted Cost	
9		Customer A	1000	20%	=C9*(1-D9)	
10						

To increase a number by a percentage amount, multiply the original amount by 1 plus the percent of increase. In the example in Figure 11.8, Product A is getting a 10% increase. So, we first add 1 to the 10%. This gives us 110%. We then multiply the original price of 100 by 110%. This calculates to the new price of 110.

To decrease a number by a percentage amount, multiply the original amount by 1, which is the percent discount. In the example in Figure 11.8, Customer A is getting a 20% discount. So, we first subtract 20% from 1. This gives us 80%. We then multiply the original 1000 cost per service by 80%. This calculates to the new rate of 800.

Note the use of parentheses in the formulas. By default, Excel's order of operations states that multiplication must be done before addition or subtraction. But if we let that happen, we would get an erroneous result. Wrapping the second part of the formula in parentheses ensures that Excel performs the multiplication last.

Dealing with divide-by-zero errors

In mathematics, division by zero is impossible. One way to understand why it's impossible is to consider what happens when you divide a number by another.

Division is really nothing more than fancy subtraction. For example, 10 divided by 2 is the same as starting with 10 and continuously subtracting 2 as many times as needed to get to zero. In this case, you would need to continuously subtract 2 five times.

$$10 - 2 = 8$$

$$8 - 2 = 6$$

$$6 - 2 = 4$$

$$4 - 2 = 2$$

$$2 - 2 = 0$$

So, $10/2 = 5$.

Now if you tried to do this with 10 divided by 0, you would never get anywhere because $10 - 0$ is 10 all day long. You'd be sitting there subtracting 0 until your calculator dies.

$$10 - 0 = 10$$

$$10 - 0 = 10$$

$$10 - 0 = 10$$

$$10 - 0 = 10$$

. . .infinity

Mathematicians call the result you get when dividing any number by zero *undefined*. Software like Excel simply gives you an error when you try to divide by zero. In Excel, when you divide a number by zero, you get the #DIV/0! error.

Part II: Working with Formulas and Functions

You can avoid this by telling Excel to skip the calculation if your denominator is a zero. Figure 11.9 illustrates how to do this by wrapping the division operation in Excel's IF function.

=IF(C4=0, 0, D4/C4)

FIGURE 11.9

Using the IF function to avoid a division-by-zero error

	A	B	C	D	E
1					
2			Budget	Actual	Percent to Budget
3		Jim	200	200	100%
4		Tim	0	100	=IF(C4=0, 0, D4/C4)
5		Kim	300	350	117%

The IF function has three arguments: the condition, what to do if the condition is true, and what to do if the condition is false.

The condition argument in this example is that the budget in C4 is equal to zero (C4=0). Condition arguments must be structured to return TRUE or FALSE, which usually means there is a comparison operation (such as an equal sign or greater-than sign) or another worksheet function that returns TRUE or FALSE (like ISERR or ISBLANK).

If our condition argument returns TRUE, the second argument of the IF function is returned to the cell. Our second argument is 0, meaning that we simply want a zero displayed if the budget number in cell C4 is a zero.

If the condition argument is not zero, the third argument takes effect. In our third argument, we tell Excel to perform the division calculation (D4/C4).

So, this formula states that if C4 equals 0, then return a 0; otherwise, return the result of D4 / C4.

TIP

Note that the second argument for the IF function can be any value or text. For instance, you could use two quotes ("") as the second argument to return a blank cell instead of a 0.

Rounding Numbers

Often, your customers will want to look at clean, round numbers. Inundating a user with decimal values and unnecessary digits for the sake of precision can make your reports harder to read. For this reason, you may consider using Excel's rounding functions.

In the following sections, you'll explore some of the techniques that you can leverage to apply rounding to your calculations.

Rounding numbers using formulas

Excel's ROUND function is used to round a given number to a specified number of digits. The ROUND function takes two arguments: the original value and the number of digits to round to.

Passing 0 as the second argument tells Excel to remove all decimal places and round the integer portion of the number based on the first decimal place. For instance, this formula rounds to 94:

```
=ROUND(94.45, 0)
```

Passing a 1 as the second argument tells Excel to round to one decimal based on the value of the second decimal place. For example, this formula rounds to 94.5:

```
=ROUND(94.45, 1)
```

You can also pass a negative number to the second argument, telling Excel to round based on values to the left of the decimal point. The following formula, for example, returns 90:

```
=ROUND(94.45, -1)
```

You can force rounding in a particular direction using the ROUNDUP or ROUNDDOWN function.

This ROUNDDOWN formula rounds 94.45 down to 94:

```
=ROUNDDOWN(94.45, 0)
```

This ROUNDUP formula rounds 94.45 up to 95:

```
=ROUNDUP(94.45, 0)
```

Rounding to the nearest penny

In some industries, it is common practice to round a dollar amount to the nearest penny.

Figure 11.10 demonstrates how rounding a dollar amount up or down to the nearest penny can affect the resulting number.

FIGURE 11.10

Rounding to the nearest penny

	A	B	C	D
1				
2		Dollar Amount	Round up to Nearest Penny	Round Down to the Nearest Penny
3	\$ 34.243	\$34.25	\$34.24	
4				
5		=CEILING(B3,0.01)	=FLOOR(B3,0.01)	

You can round to the nearest penny by using the CEILING or FLOOR function.

The CEILING function will round a number up to the nearest multiple of significance that you pass to it. This comes in handy when you need to override the standard rounding protocol with your own business rules. For instance, you can force Excel to round 123.222 to 124 by using the CEILING function with a significance of 1:

```
=CEILING(123.222, 1)
```

So, passing a .01 as the significance tells the CEILING function to round up to the nearest penny.

If you wanted to round up to the nearest nickel, you can use .05 as the significance. For instance, the following formula returns 123.15:

```
=CEILING(123.11, .05)
```

The FLOOR function works the same way except it forces a rounding down to the nearest significance. The following example function rounds 123.19 down to the nearest nickel, giving you 123.15 as the result:

```
=FLOOR(123.19, .05)
```

Rounding to significant digits

In some financial reports, figures are presented in significant digits. The idea is that when you're dealing with numbers in the millions, there is no need to inundate a report with superfluous numbers for the sake of showing precision down to the tens, hundreds, and thousands place.

For instance, instead of showing the number 883,788, you could choose to round the number to one significant digit. This would mean displaying the same number as 900,000. Rounding 883,788 to two significant digits would show the number as 880,000.

In essence, you're deeming that a particular number's place is significant enough to show. The rest of the number can be replaced with zeros. This may feel like it could introduce problems, but when dealing with large enough numbers, any number below a certain significance would be inconsequential.

Figure 11.11 demonstrates how you can implement a formula that rounds numbers to a given number of significant digits.

Let's take a moment to see how this works.

Excel's ROUND function is used to round a given number to a specified number of digits. The ROUND function takes two arguments: the original value and number of digits to round to.

Passing a negative number to the second argument tells Excel to round based on significant digits to the left of the decimal point. The following formula, for example, returns 9500:

```
=ROUND(9489, -2)
```

FIGURE 11.11

Rounding numbers to one significant digit

A	B	C	D	E
1				
2				
3				
4	Raw Number	Significant Digits		Significant Digits
5	605,390	=ROUND(B5,LEN(INT(ABS(B5)))*-1+\$E\$3)		2
6	900,942	900,000		
7	591,007	590,000		
8	491,235	490,000		
9	883,788	880,000		
10	952,687	950,000		
11	(332,602)	-330,000		

Changing the significant digits argument to -3 will return a value of 9000.

`=ROUND(9489, -3)`

This works great, except what if we have numbers on differing scales? In other words, what if some of our numbers are millions while others are hundreds of thousands? If we wanted to display all our numbers in one significant digit, we would need to build a different ROUND function for each number to account for the differing significant digits argument that we would need for each type of number.

To help solve this, we can replace our hard-coded significant digits argument with a formula that calculates what that number should be.

Imagine that our number is -2330.45 . We can use this formula as the significant digits argument in our ROUND function:

`LEN(INT(ABS(-2330.45))) * -1 + 2`

This formula first wraps our number within the ABS function, effectively removing any negative symbol that may exist. It then wraps that result in the INT function, stripping out any decimals that may exist. Finally, it wraps that result in the LEN function to get a measure of how many digits are in the number without any decimals or negation symbols.

In the example, this part of the formula results in 4. If you take the number -2330.45 and strip away the decimals and negative symbol, you have four digits left.

This number is then multiplied by -1 to make it a negative number, and it is then added to the number of significant digits we are seeking. In this example, $4 * -1 + 2 = -2$.

Again, this formula will be used as the second argument for our ROUND function. Enter this formula into Excel, and you'll round this number to -2300 (two significant digits).

`=ROUND(-2330.45, LEN(INT(ABS(-2330.45))) * -1 + 2)`

Part II: Working with Formulas and Functions

You can then replace this formula with cell references that point to the source number and cell that holds the number of desired significant digits. This is what you saw in Figure 11.11.

=ROUND(B5, LEN(INT(ABS(B5))) * -1 + \$E\$3)

Counting Values in a Range

Excel provides several functions to count the values in a range: COUNT, COUNTA, and COUNTBLANK. Each of these functions provides a different method of counting based on whether the values are numbers, numbers and text, or blank.

Figure 11.12 illustrates the different kinds of counting that you can perform. In row 12, we are using the COUNT function to count only the exams where students have passed. In column H, we are using the COUNTA function to count all the exams taken by a student. In column I, we are using the COUNTBLANK function to count only those exams that have not yet been taken.

FIGURE 11.12

A demonstration of counting cells

	A	B	C	D	E	F	G	H	I
1									
2									
3									
4									
5									
6									
7									
8									
9									
10									
11									
12									

The table shows student exam results and counts. Row 12 contains formulas:

- Cell C12: =COUNT(C4:C8) (counts 4 passed exams)
- Cell H4: =COUNTA(C4:F4) (counts 10 total exams)
- Cell H8: =COUNTBLANK(C4:F4) (counts 6 blank exams)

Below the table, a callout box labeled "How many students passed each exam." contains a table:

	Math	English	Art	History
	2	3	4	1

The COUNT function will count only numeric values in a given range. It requires only a single argument where you pass a range of cells. For example, this formula will count only those cells in range C4 : C8 that contain a numeric value:

=COUNT(C4 : C8)

The COUNTA function will count any cell that is not blank. This function can be used when counting cells that contain any combination of numbers and text. It requires only a single argument where you pass a range of cells. For instance, this formula will count all the non-blank cells in range C4 : F4:

=COUNTA(C4 : F4)

The COUNTBLANK function will count only the blank cells in a given range. It requires only a single argument where you pass a range of cells. For instance, this formula will count all the blank cells in range C4 : F4:

```
=COUNTBLANK(C4:F4)
```

Using Excel's Conversion Functions

You may work at a company where it's important to know how many cubic yards can be covered by a gallon of material or how many cups are needed to fill an Imperial gallon.

You can use Excel's CONVERT function to produce a conversion table containing every possible type of conversion you need for a set of measures. Figure 11.13 illustrates a conversion table created using nothing but Excel's CONVERT function.

FIGURE 11.13

Creating a unit-of-measure conversion table

	D	E	F	G	H	I
1						
2			Teaspoon			
3			tsp			
4	Teaspoon	tsp	=CONVERT(1,\$E4,F\$3)	0.33	0.17	0.02
5	Tablespoon	tbs	3.00	1.00	0.50	0.06
6	Fluid ounce	oz	6.00	2.00	1.00	0.13
7	Cup	cup	48.00	16.00	8.00	1.00
8	U.S. pint	us_pt	96.00	32.00	16.00	2.00
9	U.K. pint	uk_pt	115.29	38.43	19.22	2.40
10	Quart	qt	192.00	64.00	32.00	4.00
11	Imperial quart	uk qt	230.58	76.86	38.43	4.80
12	Gallon	gal	768.00	256.00	128.00	16.00

With this table, you can get a quick view of the conversions from one unit of measure to another. You can see that it takes 48 teaspoons to make a cup, 2.4 cups to make an English pint, and so forth.

The CONVERT function requires three arguments: a number value, the unit you're converting from, and the unit you're converting to. For instance, to convert 100 miles into kilometers, you can enter this formula to get the answer 160.93:

```
=CONVERT(100, "mi", "km")
```

You can use the following formula to convert 100 gallons into liters. This will give you the result of 378.54:

```
=CONVERT(100, "gal", "l")
```

Part II: Working with Formulas and Functions

You'll notice the conversion codes for each unit of measure. These codes are specific and must be entered exactly how Excel expects to see them. Entering a CONVERT formula using gallon or GAL instead of the expected gal will return an error.

Luckily, Excel provides a ScreenTip as you start entering your CONVERT function, letting you pick the correct unit codes from a list.

You can refer to Excel's help files on the CONVERT function to get a list of valid units of measure conversion codes.

Once you have the codes in which you are interested, you can enter them in a matrix-style table like the one you saw in Figure 11.13. In the top-left cell in your matrix, enter a formula that points to the appropriate conversion code for the matrix row and matrix column.

Be sure to include the absolute references necessary to lock the references to the conversion codes. For the codes located in the matrix row, lock to the column reference. For the codes located in the matrix column, lock the row reference.

```
=CONVERT(1,$E4,F$3)
```

At this point, you can simply copy your formula across the entire matrix.

Using Formulas to Manipulate Text

IN THIS CHAPTER

- Seeing how Excel handles text entered into cells
- Exploring Excel worksheet functions that handle text
- Getting examples of advanced text formulas

Often, the work you do with Excel not only involves calculating numbers but also includes transforming and shaping data to fit your data models. Many of these activities include manipulating text strings. This chapter will highlight some of the common text transformation exercises an Excel analyst performs and, in the process, give you a sense of some of the text-based functions that Excel has to offer.

Working with Text

When you enter data into a cell, Excel immediately goes to work and determines whether you're entering a formula, a number (including a date or time), or anything else. That "anything else" is considered text.

NOTE

You may hear the term *string* used instead of *text*. You can use these terms interchangeably. Sometimes they even appear together, as in *text string*.

A single cell can hold up to 32,000 characters, but Excel is not a word processor, and there's no reason why anyone would need to come even close to that number of characters in a cell.

If you need to display lots of text in a worksheet, consider using a text box. Choose Insert ➔ Text ➔ Text Box, click the worksheet to create the text box, and then start typing. Working with large amounts of text in a text box is easier than editing cells. In addition, you can easily move, resize, or change the dimensions of a text box. However, if you need to work with the text using formulas and functions, the text must reside in cells.

When a Number Isn't Treated as a Number

If you import data into Excel, you may be aware of a common problem: sometimes the imported values are treated as text.

Depending on your error-checking settings, Excel may display error indicators to identify numbers stored as text. An error indicator appears as a green triangle in the upper-left corner of cells. In addition, an icon appears next to the cell. Activate the cell and click the icon, which expands to show a list of options. To force the number to be treated as an actual number, select Convert to Number from the list of options.

To control which error-checking rules are in effect, choose File ➔ Options and then select the Formulas tab. You can enable or disable any or all of the rules listed under Error Checking Rules.

Here's another way to convert these non-numbers to actual values. Activate any empty cell and choose Home ➔ Clipboard ➔ Copy (or press Ctrl+C). Then select the range that contains the values you need to fix. Choose Home ➔ Clipboard ➔ Paste Special. In the Paste Special dialog box, select the Add operation and then click OK. This procedure essentially adds zero to each cell—and, in the process, forces Excel to treat the non-numbers as actual values.

Using Text Functions

Excel has an excellent assortment of worksheet functions that can handle text. You can access these functions just where you'd expect: from the Text control in the Function Library group of the Formulas tab.

Many of the text functions are not limited to text: They can also operate with cells that contain numeric values. You'll find that Excel is very accommodating when it comes to treating numbers as text.

The examples discussed in the following sections demonstrate some common (and useful) things you can do with text. You may need to adapt some of these examples for your own use.



This book's website, www.wiley.com/go/excel365bible, includes a copy of the sample workbook for this chapter. The file is named `Text Formulas.xlsx`.

Joining text strings

One of the more basic text manipulation actions that you can do is joining text strings together. In the example shown in Figure 12.1, we are creating a full name column by joining together first and last names.

This example illustrates the use of the ampersand (&) operator. The ampersand operator tells Excel to concatenate values with one another. As you can see in Figure 12.1, you can join cell values with text of your own. In this example, we are joining the values in cells B3 and C3 separated by a space (created by entering a space in quotes).

FIGURE 12.1

Joining first and last names

A	B	C	D	E
1				
2	FirstName	LastName	Full Name	Middle Initial
3	Guy	Gilbert	=B3&" "&C3	H.
4	Kevin	Brown	Kevin Brown	P.
5	Roberto	Tamburello	Roberto Tamburello	B.
6	Rob	Walters	Rob Walters	A.
7	Thierry	Alexander	Thierry Alexander	D.
8	David	Bradley	David Bradley	

The TEXTJOIN function provides an easier way to handle more complex scenarios. This function requires just a few arguments:

```
TEXTJOIN(delimiter, ignore_empty_values, text)
```

The first argument is the character that you want placed between the cells that you are joining. If you type a comma as the “delimiter,” the function will place a comma between the joined values.

The second argument determines what to do when Excel encounters an empty cell. You can either set this argument to TRUE, telling Excel to ignore empty cells, or set it to FALSE. The best way to think about this argument is how you want Excel to place your chosen delimiter. Setting this argument to TRUE will ensure that Excel does not add extra commas between your joined text because of blank cells in the selected range.

The third argument is the text to be joined. This can be a simple text string, or it can be an array of strings like a range of cells. The TEXTJOIN function requires at least one value or cell reference in this argument.

Figure 12.2 demonstrates how you can use the TEXTJOIN function to easily pull together the first name, middle initial, and last name for each person in the table.

FIGURE 12.2

Using the TEXTJOIN function

A	B	C	D	E
13				
14				
15	FirstName	Middle Initial	LastName	TEXTJOIN
16	Guy	H.	Gilbert	=TEXTJOIN(" ",TRUE,B16:D16)
17	Kevin	P.	Brown	Kevin P. Brown
18	Roberto	B.	Tamburello	Roberto B. Tamburello
19	Rob	A.	Walters	Rob A. Walters
20	Thierry	D.	Alexander	Thierry D. Alexander
21	David		Bradley	David Bradley
22	JoLynn		Dobney	JoLynn Dobney

Setting text to sentence case

Excel provides three useful functions to change the text to upper-, lower-, or proper case. As you can see in rows 6, 7, and 8 illustrated in Figure 12.3, these functions require nothing more than a pointer to the text that you want to be converted. As you might guess, the **UPPER** function converts text to all uppercase, the **LOWER** function converts text to all lowercase, and the **PROPER** function converts text to title case (the first letter of every word is capitalized).

FIGURE 12.3

Converting text into upper, lower, proper, and sentence case

A	B	C
4		The QUICK brown FOX JUMPS over the lazy DOG.
5		
6	=UPPER(C4)	THE QUICK BROWN FOX JUMPS OVER THE LAZY DOG.
7	=LOWER(C4)	the quick brown fox jumps over the lazy dog.
8	=PROPER(C4)	The Quick Brown Fox Jumps Over The Lazy Dog.
9		
10	=UPPER(LEFT(C4,1))&LOWER(RIGHT(C4,LEN(C4)-1))	The quick brown fox jumps over the lazy dog.

What Excel lacks is a function to convert text to sentence case (where only the first letter of the first word is capitalized). But as you can see in Figure 12.3, you can use the following formula to force text into sentence case:

=UPPER(LEFT(C4,1))&LOWER(RIGHT(C4,LEN(C4)-1))

If you look at this formula closely, you can see that it's made up of two parts that are joined by the ampersand.

The first part uses Excel's **LEFT** function.

UPPER(LEFT(C4,1))

The **LEFT** function allows you to extract a given number of characters from the left of a given text string. The **LEFT** function requires two arguments: the text string you are evaluating and the number of characters you need to be extracted from the left of the text string. In this case, we are extracting one character from the left of the text in cell C4. We are then making it uppercase by wrapping it in the **UPPER** function.

The second part is a bit trickier. Here we are using Excel's **RIGHT** function:

LOWER(RIGHT(C4,LEN(C4)-1))

Like the **LEFT** function, the **RIGHT** function requires two arguments: the text you are evaluating and the number of characters you need to be extracted from the right of the text string. In this case, however, we can't just give the **RIGHT** function a hard-coded number for the second argument. We have to calculate that number by subtracting 1 from the entire length of the text string. We subtract 1 to account for the first character, which is already uppercase thanks to the first part of the formula.

The `LEN` function is used to get the entire length of the text string. We subtract 1 from that, and we have the number of characters needed for our `RIGHT` function.

We can finally pass all of that to the `LOWER` function to make everything but the first character lowercase.

Joining the two parts together gives us our sentence case:

```
=UPPER(LEFT(C4, 1)) & LOWER(RIGHT(C4, LEN(C4) - 1))
```

Removing spaces from a text string

If you pull data in from external databases and legacy systems, you will no doubt encounter text that contains extra spaces. Sometimes these extra spaces are found at the beginning of the text, at the end of the text, or even between text strings (as in cell B6 shown in Figure 12.4).

FIGURE 12.4

Removing excess spaces from text

	A	B	C	D
3		Original Text	Trimmed Text	
4		ABCD	ABCD	
5		A B C D	A B C D	
6		Alan Jones	Alan Jones	
7		ABCD	=TRIM(B7)	
8				

Extra spaces are generally evil because they can cause problems in lookup formulas, charting, column sizing, and printing.

Figure 12.4 illustrates how you can remove superfluous spaces by using the `TRIM` function.

The `TRIM` function is relatively straightforward. Simply give it some text and it will remove all spaces from the text except for single spaces between words.

As with other functions, you can nest the `TRIM` function in other functions to clean up your text while applying some other manipulation. For instance, this function trims the text in cell A1 and converts it to uppercase all in one step:

```
=UPPER(TRIM(A1))
```

It's important to note that the `TRIM` function was designed to trim only the ASCII space character from text. The ASCII space character has a code value of 32. In the Unicode character set, however, there is an additional space character called the *nonbreaking space character*. This character is commonly used in web pages and has the Unicode value of 160.

The `TRIM` function is designed to handle only `CHAR(32)` space characters. It cannot, by itself, handle `CHAR(160)` space characters. To handle this kind of space, you'll need to utilize the `SUBSTITUTE` function to find `CHAR(160)` space characters and replace them with

Part II: Working with Formulas and Functions

CHAR (32) space characters so that the TRIM function can fix them. You can accomplish this all at one time in the following formula:

```
=TRIM(SUBSTITUTE(A4, CHAR(160), CHAR(32)))
```

For a detailed look at the SUBSTITUTE function, see the section “Substituting text strings” later in this chapter.

Extracting parts of a text string

One of the most important techniques for manipulating text in Excel is the ability to extract specific portions of text. Using Excel’s LEFT, RIGHT, and MID functions, you can do things like this:

- Convert nine-digit postal codes into five-digit postal codes
- Extract phone numbers without the area code
- Extract parts of employee or job codes for use somewhere else

Figure 12.5 demonstrates how using the LEFT, RIGHT, and MID functions can help easily accomplish these tasks.

FIGURE 12.5

Using the LEFT, RIGHT, and MID functions

	A	B
Convert these 9 digit postal codes into 5 digit postal codes.		
2	Zip	Zip
3	70056-2343	70056
4	75023-5774	75023
Extract the phone number without the area code.		
7	Phone	Phone
8	(214)887-7765	887-7765
9	(703)654-2180	654-2180
Extract the 4th character of each Job Code.		
12	Job Code	Job Level
13	2214001	4
14	5542075	2
15	1113543	3

The LEFT function allows you to extract a given number of characters from the left of a given text string. The LEFT function requires two arguments: the text you are evaluating and the number of characters you need to be extracted from the left of the text string. In this example, we are extracting the left five characters from the value in cell A4:

```
=LEFT(A4, 5)
```

The **RIGHT** function allows you to extract a given number of characters from the right of a given text string. The **RIGHT** function requires two arguments: the text string you are evaluating and the number of characters you need to be extracted from the right of the text string. In this example, we are extracting the right eight characters from the value in cell A9:

```
=RIGHT(A9, 8)
```

The **MID** function allows you to extract a given number of characters from the middle of a given text string. The **MID** function requires three arguments: the text string you are evaluating, the character position in the text string from which to start the extraction, and the number of characters you need to be extracted. In this example, we are starting at the fourth character in our text string and extracting one character:

```
=MID(A14, 4, 1)
```

Finding a particular character in a text string

Excel's **LEFT**, **RIGHT**, and **MID** functions work great for extracting text, but only if you know the exact position of the characters you are targeting. What can you do when you don't know exactly where to start the extraction? For example, if you had the following list of product codes, how would you go about extracting all the text after the hyphens?

PRT-432

COPR-6758

SVCCALL-58574

The **LEFT** function wouldn't work because you need the right few characters. The **RIGHT** function alone wouldn't work because you need to tell it exactly how many characters to extract from the right of the text string. Any number that you give will pull either too many or too few characters from the text. The **MID** function alone wouldn't work because you need to tell it exactly where in the text to start extracting. Again, any number you give will pull either too many or too few characters from the text.

The reality is that you will often need to find specific characters to get the appropriate starting position for the extraction. This is where Excel's **FIND** function comes in handy. With the **FIND** function, you can get the position number of a particular character and use that character position in other operations.

In the example shown in Figure 12.6, we use the **FIND** function in conjunction with the **MID** function to extract the middle numbers from a list of product codes. As you can see from the formula, we find the position of the hyphen and use that to feed the **MID** function.

```
=MID(B3, FIND(" - ", B3) + 1, 2)
```

The **FIND** function has two required arguments. The first argument is the text you want to find. The second argument is the text you want to search. By default, the **FIND** function will return the position number of the character you are trying to find. If the text you are

Part II: Working with Formulas and Functions

searching contains more than one of your search characters, the FIND function will return the position number of the first encounters.

FIGURE 12.6

Using the FIND function to extract data based on the position of the hyphen

	A	B	C
1			
2		Product Code	Extract the Numbers
3	PWR-16-Small	=MID(B3,FIND("-",B3)+1,2)	
4	PW-18-Medium	18	
5	PW-19-Large	19	
6	CWS-22-Medium	22	
7	CWTP-44-Large	44	

For instance, the following formula will search for a hyphen in the text string PWR-16-Small. The result will be a number 4 because the first hyphen it encounters is the fourth character in the text string.

```
=FIND ("-", "PWR-16-Small")
```

You can use the FIND function as an argument in a MID function to extract a set number of characters after the position number returned by the FIND function.

Entering this formula in a cell will give you the two numbers after the first hyphen found in the text. Note the +1 in the formula. This ensures that you move over one character to get to the text after the hyphen.

```
=MID ("PWR-16-Small", FIND ("-", "PWR-16-Small") +1, 2)
```

Finding the second instance of a character

By default, the FIND function returns the position number of the first instance of the character for which you are searching. If you want the position number of the second instance, you can use the optional Start_Num argument. This argument lets you specify the character position in the text string to start the search.

For example, this formula will return the position number of the second hyphen because we're telling the FIND function to start searching at position 5 (after the first hyphen):

```
=FIND ("-", "PWR-16-Small", 5)
```

To do this dynamically (without knowing where to start the search), you can nest a FIND function as the Start_Num argument in another FIND function. You can enter this formula into Excel to get the position number of the second hyphen:

```
=FIND ("-", "PWR-16-Small", FIND ("-", "PWR-16-Small") +1)
```

Figure 12.7 demonstrates a real-world example of this concept. Here we are extracting the size attribute from the product code by finding the second instance of the hyphen and using that

position number as the starting point in the MID function. The formula shown in cell C3 is as follows:

```
=MID(B3, FIND(" - ", B3, FIND(" - ", B3) + 1) + 1, 10000)
```

FIGURE 12.7

Nesting the FIND function to extract everything after the second hyphen

A	B	C
1		
2	Product Code	Extract the Size Designation
3	PWR-16-Small	=MID(B3,FIND(" - ",B3,FIND(" - ",B3)+1)+1,10000)
4	PW-18-Medium	Medium
5	PW-19-Large	Large
6	CWS-22-Medium	Medium
7	CWTP-44-Large	Large

This formula tells Excel to find the position number of the second hyphen, move over one character, and then extract the next 10,000 characters. Of course, there aren't 10,000 characters, but this ensures that everything after the second hyphen is pulled.

Substituting text strings

There are cases when it's helpful to substitute some text with other text. One such case is where you encounter the annoying apostrophe S ('S) quirk that you get with the PROPER function. To see what we mean, enter this formula into Excel:

```
=PROPER("STARBUCK'S COFFEE")
```

This formula is meant to convert the given text into title case (where the first letter of every word is capitalized). The actual result of the formula is this:

Starbuck'S Coffee

Note how the PROPER function capitalizes the S after the apostrophe—annoying to say the least.

However, with a little help from Excel's SUBSTITUTE function, you can avoid this annoyance. Figure 12.8 shows the fix using the following formula:

```
=SUBSTITUTE(PROPER(SUBSTITUTE(B4, "'''", "qzx")), "qzx", "'''")
```

FIGURE 12.8

Fixing the apostrophe S issue with the SUBSTITUTE function

A	B	C	D	E
2				
3	Company	Bad Proper Case		Better Proper Case
4	STARBUCK'S COFFEE	Starbuck'S Coffee		Starbuck's Coffee
5	MCDONALD'S	Mcdonald'S		Mcdonald's
6	MICHAEL'S DELI	Michael'S Deli		Michael's Deli

Part II: Working with Formulas and Functions

Our formula uses the `SUBSTITUTE` function, which requires three arguments: the target text, the old text you want to be replaced, and the new text to use as the replacement.

As you look at the full formula, you'll note there are two `SUBSTITUTE` functions in use. This formula actually consists of two formulas (one nested in the other). The first formula is the part that reads as follows:

```
PROPER(SUBSTITUTE(B4, " ", "qzx"))
```

In this part, we are using the `SUBSTITUTE` function to replace the apostrophe (') with `qzx`. This may seem like a crazy thing to do, but there is some method here. The `PROPER` function will essentially capitalize any letter coming directly after a symbol. Here, we are tricking the `PROPER` function by substituting the apostrophe with a benign set of letters that are unlikely to be strung together in the original text.

The second formula wraps the first. This formula substitutes the benign `qzx` with an apostrophe:

```
=SUBSTITUTE(PROPER(SUBSTITUTE(B4, " ", "qzx")), "qzx", " ")
```

So, the entire formula replaces the apostrophe with `qzx`, performs the `PROPER` function, and then reverts the `qzx` to an apostrophe.

Counting specific characters in a cell

A useful trick is to be able to count the number of times a specific character exists in a text string. The technique for doing this in Excel is relatively clever. If you wanted to count the number of times the letter *s* appears in the word *Mississippi*, for example, you could count them by hand, of course, but systematically, you could follow these general steps:

1. Measure the character length of the word *Mississippi* (11 characters).
2. Measure the character length after removing every letter *s* (7 characters).
3. Subtract the adjusted length from the original length.

After performing these steps, you'll accurately conclude that the letter *s* occurs four times in the word *Mississippi*.

A real-world use for this technique of counting specific characters is to calculate a word count in Excel. In Figure 12.9, you'll see the following formula used to count the number of words entered in cell B4 (nine words in this case):

```
=LEN(B4) - LEN(SUBSTITUTE(B4, " ", "")) + 1
```

This formula essentially follows the steps mentioned earlier in the setup for this section. The formula first uses the `LEN` function to measure the length of the text in cell B4:

```
LEN(B4)
```

It then uses the `SUBSTITUTE` function to remove the spaces from the text:

```
SUBSTITUTE(B4, " ", "")
```

FIGURE 12.9

Calculating the number of words in a cell

A	B	C
1		
2		
3		Get Word Count
4	The Quick Brown Fox Jumps Over The Lazy Dog.	=LEN(B4)-LEN(SUBSTITUTE(B4," ",""))+1
5		9

Wrapping the SUBSTITUTE function in a LEN function gives us the length of the text without the spaces. Note that we must add one (+1) to that answer to account for the fact that the last word will not have an associated space:

`LEN(SUBSTITUTE(B4, " ", "")) +1`

Subtracting the adjusted length from the original length gives us our word count:

`=LEN(B4) -LEN(SUBSTITUTE(B4, " ", "")) +1`

Adding a line break within a formula

When creating charts in Excel, it's sometimes useful to force line breaks to create better visualizations. Take the chart shown in Figure 12.10, for example. Here you see a chart where the x-axis labels include the data value for each sales rep. This comes in handy when you don't want to inundate your chart with data labels.

FIGURE 12.10

The x-axis labels in this chart include a line break and a reference to the data values.



The secret to this trick is the use of the CHAR() function within a formula that makes up your chart labels (see Figure 12.11).

Part II: Working with Formulas and Functions

Every character in Excel has an associated American National Standards Institute (ANSI) character code. The ANSI character code is a Windows system code set that defines the characters you see on your screen. The ANSI character set consists of 255 characters, numbering from 1 to 255. The uppercase letter A is character 65. The number 9 is character 57.

Even nonprinting characters have codes. The code for a space is 32. The code for a line break is 10.

FIGURE 12.11

Using the CHAR () function to force a line break between the sales rep name and the data value

	A	B	C
1			
2			Units Sold
3	Mike	=A3&CHAR(10)&C3	843
4	Mary	Mary573	573
5	Ethan	Ethan858	858
6	Emma	Emma890	890
7	Ava	Ava602	602
8	Mia	Mia984	984

You can call up any character in a formula by using the CHAR () function. In the example shown in Figure 12.11, we are calling up the line break character and joining it with the values in cells A3 and C3:

=A3&CHAR(10)&C3

The cell itself won't show the line break unless you have applied wrap text. But even if you haven't, any chart using this kind of formula will display the data returned by the formula with the line breaks.

Cleaning strange characters from text fields

When importing data from external data sources like text files or web feeds, strange characters may come in with your data. Instead of trying to clean these up manually, you can use Excel's CLEAN function (see Figure 12.12).

FIGURE 12.12

Cleaning data with the CLEAN function

	A	B	C
1			
2	Store	Cleaned Text	
3	Detroit (Store #1)	=TRIM(CLEAN(B3))	
4	Detroit (Store #2)	Detroit (Store #2)	
5	Detroit (Store #3)	Detroit (Store #3)	
6	Charlotte (Store #1)	Charlotte (Store #1)	
7	Charlotte (Store #2)	Charlotte (Store #2)	
8	Charlotte (Store #3)	Charlotte (Store #3)	
9			

The CLEAN function removes nonprintable characters from any text you pass to it. You can wrap the CLEAN function within the TRIM function to remove unprintable characters and excess spaces at the same time.

```
=TRIM(CLEAN(B3))
```

Padding numbers with zeros

In many cases, the work you do in Excel ends up in other database systems within the organization. Those database systems often have field length requirements that mandate a certain number of characters. A common technique for ensuring that a field consists of a set number of characters is to pad data with zeros.

Padding data with zeros is a relatively easy concept. If you have a Customer ID field that must be 10 characters long, you essentially will need to add enough zeros to fulfill that requirement. So, Customer ID 2345 would need to be padded with six zeros, making that ID 2345000000.

Cell C4 shown in Figure 12.13 uses this formula to pad the Customer ID fields with zeros:

```
=LEFT(B4&"0000000000",10)
```

FIGURE 12.13

Padding Customer ID fields to 10 characters

A	B	C
2		
3	Customer ID	Pad to 10 characters
4	5381656	=LEFT(B4&"0000000000", 10)
5	832	8320000000
6	23	2300000000
7	290	2900000000
8	2036	2036000000
9	5965	5965000000
10	6	6000000000
11	7457	7457000000
12	2903	2903000000

The formula shown in Figure 12.13 first joins the value in cell B4 and a text string composed of 10 zeros. This effectively creates a new text string that guarantees a Customer ID value with 10 zeros.

We then use the LEFT function to extract the left 10 characters of that new text string.

Formatting the numbers in a text string

It's not uncommon to have reporting that joins text with numbers. For example, you may have a requirement to display a line in your report that summarizes a salesperson's results like this:

John Hutchison: \$5,000

Part II: Working with Formulas and Functions

The problem is that when you join numbers in a text string, the number formatting does not follow. Look at Figure 12.14, for example. Note how the numbers in the joined string do not adopt the formatting of the source cells.

FIGURE 12.14

Numbers joined with text do not inherently adopt number formatting.

	A	B	C	D	E
1					
2		Rep	Revenue		Rep and Revenue
3		Gilbert	\$6,820		=B3&": "&C3
4	Brown		\$5,205		Brown: 5205
5	Tamburello		\$246		Tamburello: 246
6	Walters		\$7,136		Walters: 7136
7	Alexander		\$2,921		Alexander: 2921
8	Bradley		\$8,225		Bradley: 8225
9	Dobney		\$5,630		Dobney: 5630
10	Ellerbrock		\$7,994		Ellerbrock: 7994
11	Hartwig		\$6,676		Hartwig: 6676
12	Campbell		\$5,716		Campbell: 5716

To solve this problem, you'll have to wrap the cell reference for your number value in the TEXT function. Using the TEXT function, you can apply the needed formatting on the fly. The formula shown in Figure 12.15 resolves the issue.

=B3&": "&TEXT(C3, "\$0,000")

FIGURE 12.15

Using the TEXT function lets you format numbers joined with text.

	A	B	C	D	E
1					
2		Rep	Revenue		Rep and Revenue
3		Gilbert	\$6,820		=B3&": "&TEXT(C3, "\$0,000")
4	Brown		\$5,205		Brown: \$5,205
5	Tamburello		\$246		Tamburello: \$0,246
6	Walters		\$7,136		Walters: \$7,136
7	Alexander		\$2,921		Alexander: \$2,921
8	Bradley		\$8,225		Bradley: \$8,225
9	Dobney		\$5,630		Dobney: \$5,630
10	Ellerbrock		\$7,994		Ellerbrock: \$7,994
11	Hartwig		\$6,676		Hartwig: \$6,676
12	Campbell		\$5,716		Campbell: \$5,716

The TEXT function requires two arguments: a value and a valid Excel format. You can apply any formatting that you want to a number as long as it's a format that Excel recognizes. For example, you can enter this formula into Excel to display \$99:

=TEXT(99.21, "\$#,###")

You can enter this formula into Excel to display 9921%:

```
=TEXT(99.21, "0%")
```

You can enter this formula into Excel to display 99.2:

```
=TEXT(99.21, "0.0")
```

An easy way to get the syntax for a particular number format in which you're interested is to look at the Number tab in the Format Cells dialog box. Follow these steps:

1. Right-click any cell and select Format Cells.
2. On the Number tab, select the formatting you need.
3. Select Custom from the Category list on the left of the Number tab in the Format Cells dialog box.
4. Copy the syntax found in the Type input box.

Using the DOLLAR function

If the number value you're joining with text is a dollar figure, you can use the simpler DOLLAR function. This function applies the regional currency format to the given text.

The DOLLAR function has two basic arguments: the number value and the number of decimals you want to display.

```
=B3&" : "&DOLLAR(C3, 0)
```


Using Formulas with Dates and Times

IN THIS CHAPTER

Getting an overview of dates and times in Excel

Using Excel date-related functions

Working with Excel time-related functions

Many worksheets contain dates and times in cells. For example, you might track information by date or create a schedule based on time. Beginners often find that working with dates and times in Excel can be frustrating. To work with dates and times, you need a good understanding of how Excel handles time-based information. This chapter provides the information you need to create powerful formulas that manipulate dates and times.

NOTE

The dates in this chapter correspond to the U.S. English language date format: month/day/year. For example, the date 3/1/1952 refers to March 1, 1952, not January 3, 1952. We realize that this setup may seem illogical, but that's the way Americans have been trained. We trust that the non-American readers of this book can make the adjustment.

Understanding How Excel Handles Dates and Times

The following sections present a quick overview of how Excel deals with dates and times. It covers Excel's date and time serial number system. It also provides some tips for entering and formatting dates and times.

Understanding date serial numbers

To Excel, a date is simply a number. More precisely, a date is a serial number that represents the number of days since the fictitious date of January 0, 1900. A serial number of 1 corresponds to January 1, 1900; a serial number of 2 corresponds to January 2, 1900; and so on. This system makes it possible to create formulas that perform calculations with dates. For example, you can create a formula to calculate the number of days between two dates (just subtract one from the other).

Excel supports dates from January 1, 1900, through December 31, 9999 (serial number = 2958465).

You may wonder about January 0, 1900. This nondate (which corresponds to date serial number 0) is actually used to represent times that aren't associated with a particular day. This concept becomes clear later in this chapter. (See "Entering times.")

To view a date serial number as a date, you must format the cell as a date. Choose Home ↴ Number ↴ Number Format. This drop-down control provides you with two date formats. To select from additional date formats, see "Formatting dates and times" later in this chapter.

Choose Your Date System: 1900 or 1904

Excel supports two date systems: the 1900 date system and the 1904 date system. Which system you use in a workbook determines what date serves as the basis for dates. The 1900 date system uses January 1, 1900, as the day assigned to date serial number 1. The 1904 date system uses January 1, 1904, as the base date. By default, Excel for Windows uses the 1900 date system, and pre-2011 versions of Excel for Mac use the 1904 date system.

Excel for Windows supports the 1904 date system for compatibility with older Mac files. You can choose the date system for the active workbook in the Advanced section of the Excel Options dialog box. (It's in the When Calculating This Workbook section.) Generally, you should use the default 1900 date system. And you should exercise caution if you use two different date systems in workbooks that are linked. For example, assume that Book1 uses the 1904 date system and contains the date 1/15/1999 in cell A1. Assume that Book2 uses the 1900 date system and contains a link to cell A1 in Book1. Book2 displays the date as 1/14/1995. Both workbooks use the same date serial number (34713), but they're interpreted differently.

One advantage to using the 1904 date system is that it enables you to display negative time values. With the 1900 date system, a calculation that results in a negative time (for example, 4:00 PM–5:30 PM) cannot be displayed. When using the 1904 date system, the negative time displays as –1:30 (that is, a difference of 1 hour and 30 minutes).

Entering dates

You can enter a date directly as a serial number (if you know the serial number) and then format it as a date. More often, you enter a date by using any of several recognized date formats. Excel automatically converts your entry into the corresponding date serial number (which it uses for calculations), and it applies the default date format to the cell so that it displays as an actual date rather than as a cryptic serial number.

For example, if you need to enter June 18, 2022, into a cell, you can enter the date by typing **June 18, 2022** (or any of several different date formats). Excel interprets your entry and stores the value 44730, the date serial number for that date. It also applies the default date format, so the cell contents may not appear exactly as you typed them.

Note

Depending on your regional settings, entering a date in a format such as June 18, 2022, may be interpreted as a text string. In such a case, you need to enter the date in a format that corresponds to your regional settings, such as 18 June 2022.

When you activate a cell that contains a date, the Formula bar shows the cell contents formatted by using the default date format—which corresponds to your system's short date format. The Formula bar doesn't display the date's serial number. If you need to find out the serial number for a particular date, format the cell with the General format.

TIP

To change the default date format, you need to change a systemwide setting. From the Windows Control Panel, select Clock and Region and then click Region to open the Region dialog box. The exact procedure varies, depending on the version of Windows you use. Look for the drop-down list that lets you change the Short Date format. The setting you choose determines the default date format that Excel uses to display dates in the Formula bar.

Excel is rather flexible when it comes to recognizing dates entered into a cell. It's not perfect, however. If you attempt to enter a date that lies outside the supported date range, Excel interprets it as text. If you attempt to format a serial number that lies outside the supported range as a date, the value displays as a series of hash marks #####.

13

Searching for Dates

If your worksheet uses many dates, you may need to search for a particular date by using the Find and Replace dialog box (Home \Rightarrow Editing \Rightarrow Find & Select \Rightarrow Find, or Ctrl+F). Excel is rather picky when it comes to finding dates. You must enter the date as it appears in the Formula bar. For example, if a cell contains a date formatted to display as June 19, 2022, the date appears in the Formula bar using your system's short date format (for example, 6/19/2022). Therefore, if you search for the date as it appears in the cell, Excel won't find it. But it will find the cell if you search for the date in the format that appears in the Formula bar.

Understanding time serial numbers

When you need to work with time values, you extend the Excel date serial number system to include decimals. In other words, Excel works with time by using fractional days. For example, the date serial number for June 1, 2022, is 44713. Noon (halfway through the day) is represented internally as 44713.5.

Part II: Working with Formulas and Functions

The serial number equivalent of one minute is approximately 0.00069444. The following formula calculates this number by multiplying 24 hours by 60 minutes and dividing the result into 1. The denominator consists of the number of minutes in a day (1,440).

$$=1 / (24 * 60)$$

Similarly, the serial number equivalent of one second is approximately 0.00001157, obtained by the following formula:

$$=1 / (24 * 60 * 60)$$

In this case, the denominator represents the number of seconds in a day (86,400).

In Excel, the smallest unit of time is one 1,000th of a second. The time serial number shown here represents 23:59:59.999 (one 1,000th of a second before midnight):

0.99999999

Table 13.1 shows various times of day along with each associated time serial number.

TABLE 13.1 Times of Day and Their Corresponding Serial Numbers

Time of Day	Time Serial Number
12:00:00 AM (midnight)	0.00000000
1:30:00 AM	0.06250000
7:30:00 AM	0.31250000
10:30:00 AM	0.43750000
12:00:00 PM (noon)	0.50000000
1:30:00 PM	0.56250000
4:30:00 PM	0.68750000
6:00:00 PM	0.75000000
9:00:00 PM	0.87500000
10:30:00 PM	0.93750000

Entering times

As with entering dates, you normally don't have to worry about the actual time serial numbers. Just enter the time into a cell using a recognized format. Table 13.2 shows some examples of time formats that Excel recognizes.

TABLE 13.2 Time Entry Formats Recognized by Excel

Entry	Excel Interpretation
11:30 AM	11:30 AM
11:30:00 AM	11:30 AM
11:30 PM	11:30 PM
11:30	11:30 AM
13:30	1:30 PM

Because the preceding samples don't have a specific day associated with them, Excel uses a date serial number of 0, which corresponds to the nonday January 0, 1900. Often, you'll want to combine a date and time. Do so by using a recognized date entry format, followed by a space, and then a recognized time entry format. For example, if you enter `6/1/2022 11:30` into a cell, Excel interprets it as 11:30 AM on June 1, 2022. Its date/time serial number is 44713.47917.

When you enter a time that exceeds 24 hours, the associated date for the time increments accordingly. For example, if you enter `25:00:00` into a cell, it's interpreted as 1:00 AM on January 1, 1900. The day part of the entry increments because the time exceeds 24 hours. Keep in mind that a time value without a date uses January 0, 1900, as the date.

Similarly, if you enter a date and a time (and the time exceeds 24 hours), the date you entered is adjusted. If you enter `9/18/2022 25:00:00`, for example, it's interpreted as 9/19/2022 1:00:00 AM.

If you enter a time only (without an associated date) into an unformatted cell, the maximum time you can enter into a cell is 9999:59:59 (just less than 10,000 hours). Excel adds the appropriate number of days. In this case, 9999:59:59 is interpreted as 3:59:59 PM on 02/19/1901. If you enter a time that exceeds 10,000 hours, the entry is interpreted as a text string rather than a time.

Formatting dates and times

You have a great deal of flexibility in formatting cells that contain dates and times. For example, you can format the cell to display the date part only, the time part only, or both the date and time parts.

You format dates and times by selecting the cells and then using the Number tab of the Format Cells dialog box. To display this dialog box, click the dialog box launcher icon in the Number group of the Home tab, or click the Number Format control and choose More Number Formats from the list that appears.

The Date category shows built-in date formats, and the Time category shows built-in time formats. Some formats include both date and time displays. Just select the desired format from the Type list and then click OK.

TIP

When you create a formula that refers to a cell containing a date or a time, Excel sometimes automatically formats the formula cell as a date or a time. Often, this automation is helpful; at other times, it's completely inappropriate and downright annoying. To return the number formatting to the default General format, choose Home \Rightarrow Number \Rightarrow Number Format and choose General from the drop-down list. Or just press **Ctrl+Shift+~** (tilde).

Problems with dates

Excel has some problems when it comes to dates. Many of these problems stem from the fact that Excel was designed many years ago. Excel designers basically emulated the Lotus 1-2-3 program's limited date and time features, which contain a nasty bug that was duplicated intentionally in Excel (described next). If Excel were being designed from scratch today, I'm sure it would be much more versatile in dealing with dates. Unfortunately, users are currently stuck with a product that leaves much to be desired in the area of dates.

Excel's leap year bug

A leap year, which occurs every four years, contains an additional day (February 29). Specifically, years that are evenly divisible by 100 are not leap years, unless they are also evenly divisible by 400. Although the year 1900 was not a leap year, Excel treats it as such. In other words, when you type **2/29/1900** into a cell, Excel interprets it as a valid date and assigns a serial number of 60.

If you type **2/29/1901**, however, Excel correctly interprets it as a mistake and doesn't convert it to a date. Instead, it simply makes the cell entry a text string.

How can a product used daily by millions of people contain such an obvious bug? The answer is historical. The original version of Lotus 1-2-3 contained a bug that caused it to treat 1900 as a leap year. When Excel was released some time later, the designers knew about this bug and chose to reproduce it in Excel to maintain compatibility with Lotus 1-2-3 worksheet files.

Why does this bug still exist in later versions of Excel? Microsoft asserts that the disadvantages of correcting this bug outweigh the advantages. If the bug were eliminated, it would mess up millions of existing workbooks. In addition, correcting this problem would possibly affect compatibility between Excel and other programs that use dates. As it stands, this bug really causes very few problems because most users don't use dates prior to March 1, 1900.

Pre-1900 dates

The world, of course, didn't begin on January 1, 1900. People who use Excel to work with historical information often need to work with dates before January 1, 1900. Unfortunately, the

only way to work with pre-1900 dates is to enter the date into a cell as text. For example, you can enter July 4, 1776, into a cell, and Excel won't complain.

TIP

If you plan to sort information by old dates, you should enter your text dates with a four-digit year, followed by a two-digit month, and then a two-digit day—for example, 1776-07-04. You won't be able to work with these text strings as dates, but this format will enable accurate sorting.

Using text as dates works in some situations, but the main problem is that you can't perform manipulation on a date that's entered as text. For example, you can't change its numeric formatting, you can't determine which day of the week this date occurred, and you can't calculate the date that occurs seven days later.

Inconsistent date entries

You need to be careful when entering dates by using two digits for the year. When you do so, Excel has some rules that determine which century to use.

Two-digit years between 00 and 29 are interpreted as 21st-century dates, and two-digit years between 30 and 99 are interpreted as 20th-century dates. For example, if you enter 12/15/28, Excel interprets your entry as December 15, 2028. But if you enter 12/15/30, Excel sees it as December 15, 1930, because Windows uses a default boundary year of 2029. You can keep the default as is or change it via the Windows Control Panel. From the Region dialog box, click the Additional Settings button to display the Customize Format dialog box. Select the Date tab and then specify a different year.

TIP

The best way to avoid any surprises is to simply enter all years using all four digits for the year.

Using Excel's Date and Time Functions

Excel has quite a few functions that work with dates and times. These functions are accessible by choosing **Formulas** \Rightarrow **Function Library** \Rightarrow **Date & Time**.

These functions leverage the fact that, beneath the covers, dates and times are nothing more than a numbering system. This opens the door for all kinds of cool formula-driven analyses. In the following sections, you'll walk through some of these cool analyses. Along the way, you'll pick up a few techniques that will help you create your own formulas.



This book's website, www.wiley.com/go/excel365bible, includes a copy of the sample workbook for this chapter. The file is named **Dates and Times.xlsx**.

Getting the current date and time

Instead of typing the current date and time, you can use one of two Excel functions. The TODAY function returns the current date:

```
=TODAY()
```

The NOW() function returns the current date along with the current time:

```
=NOW()
```

Both the TODAY and NOW functions return date serial numbers that represent the current system date and time. The TODAY function assumes 12:00 AM as the time, while the NOW function returns the actual time.

It's important to note that both of these functions will automatically recalculate each time you change or open your workbook, so don't use these functions as a timestamp of record.

TIP

If you want to enter a static date that won't change, press Ctrl+; (semicolon) on your keyboard. This will insert a static date in the active cell.

You can use the TODAY function as part of a text string by wrapping it in the TEXT function with some date formatting. This formula will display text that will return today's date in Month Day, Year format.

```
= "Today is "&TEXT(TODAY(), "mmmm d, YYYY")
```



For more details on using the TEXT function, see the section called “Formatting the numbers in a text string” in Chapter 12.

Calculating age

One of the easiest ways to calculate the age of anything is to use Excel's DATEDIF function. This function makes calculating any kind of date comparisons a breeze.

To calculate a person's age using the DATEDIF function, you can enter a formula like this:

```
=DATEDIF("5/16/1972", TODAY(), "Y")
```

You can, of course, reference a cell that contains a date:

```
=DATEDIF(B4, TODAY(), "Y")
```

The DATEDIF function calculates the number of days, months, or years between two dates. It requires three arguments—a start date, an end date, and a time unit.

The time units are defined by a series of codes listed in Table 13.3.

TABLE 13.3 DATEDIF Time Unit Codes

Code	What It Returns
"y"	The number of complete years in the period.
"m"	The number of complete months in the period.
"d"	The number of days in the period.
"md"	The difference between the days in start_date and end_date. The months and years of the dates are ignored.
"ym"	The difference between the months in start_date and end_date. The days and years of the dates are ignored.
"yd"	The difference between the days of start_date and end_date. The years of the dates are ignored.

Using these time codes, you can easily calculate the number of years, months, and days between two dates. If someone was born on May 16, 1972, you could find that person's age in years, months, and days using these respective formulas:

```
=DATEDIF("5/16/1972", TODAY(), "y")
=DATEDIF("5/16/1972", TODAY(), "m")
=DATEDIF("5/16/1972", TODAY(), "d")
```

Calculating the number of days between two dates

One of the most common date calculations performed in the corporate world is figuring the number of days between two dates. Project management teams use it to measure performance against a milestone, HR departments use it to measure time to fill a requisition, and finance departments use it to track aging receivables. Luckily, it's one of the easiest calculations to perform thanks to the handy DATEDIF function.

Figure 13.1 demonstrates a sample report that uses the DATEDIF function to calculate the number of days outstanding for a set of invoices.

FIGURE 13.1

Calculating the number of days between today and the invoice date

A	B	C	D
2			
3	Invoice Date		Days Outstanding
4	25-Jul-21	=DATEDIF(C4,TODAY(),"d")	
5	04-Jul-21	57	
6	04-May-21	118	
7	22-May-21	100	
8	03-Apr-21	149	
9	28-Apr-21	124	

Part II: Working with Formulas and Functions

Looking at Figure 13.1, you'll see the formula in cell D4 is as follows:

```
=DATEDIF(C4, TODAY(), "d")
```

This formula uses the DATEDIF function with the time code d. This tells Excel to return the number of days based on the start date (C4) and the end date (TODAY).

Calculating the number of workdays between two dates

Oftentimes when reporting on the elapsed number of days between a start date and end date, it's not appropriate to count the weekends in the final number of days. Operations are typically shut down on the weekends, so you would want to avoid counting those days.

You can use Excel's NETWORKDAYS function to calculate the number of days between a start date and end date excluding weekends.

As you can see in Figure 13.2, the NETWORKDAYS function is used in cell E4 to calculate the number of workdays between 1/1/2022 and 12/31/2022.

FIGURE 13.2

Calculating the number of workdays between two dates

A	B	C	D	E
1				
2				
3	Start Date	End Date	Net Work Days	
4	1/1/2022	12/31/2022	=NETWORKDAYS(B4,C4)	
5				

This formula is fairly straightforward. The NETWORKDAYS function has two required arguments: a start date and an end date. If your start date is in cell B4 and your end date is in cell C4, this formula would return the number of workdays (excluding Saturdays and Sundays):

```
=NETWORKDAYS(B4, C4)
```

Using NETWORKDAYS.INTL

The one drawback to using the NETWORKDAYS function is that it defaults to excluding Saturday and Sunday. But what if you work in a region where the weekends are actually Fridays and Saturdays? Or worse yet, what if your weekends only include Sundays?

Excel has you covered with the NETWORKDAYS.INTL function. In addition to the required start and end dates, this function has an optional third argument—a weekend code. The weekend code allows you to specify which days to exclude as a weekend day.

As you enter the NETWORKDAYS.INTL function, Excel displays a menu as soon as you go into the third argument (see Figure 13.3). Simply select the appropriate weekend code and press Enter.

FIGURE 13.3

`NETWORKDAY.INTL` allows you to specify which days to exclude as weekend days.

A	B	C	D	E	F
1					
2					
3	Start Date	End Date	Net Work Days		
4	1/1/2022	12/31/2022	=NETWORKDAYS.INTL(C4,D4,11)	(...)	
5				1 - Saturday, Sunday	
6				(...)2 - Sunday, Monday	
7				(...)3 - Monday, Tuesday	
8				(...)4 - Tuesday, Wednesday	
9				(...)5 - Wednesday, Thursday	
10				(...)6 - Thursday, Friday	
11				(...)7 - Friday, Saturday	
12				(...)11 - Sunday only	
13				(...)12 - Monday only	
14				(...)13 - Tuesday only	

Generating a list of business days excluding holidays

When creating dashboards and reports in Excel, it's often useful to have a helper table that contains a list of dates that represent business days (that is, dates that are not weekends or holidays). This kind of a helper table can help assist in calculations like revenue per business day, units per business day, and so on.

One of the easiest ways to generate a list of business days is to use the `WORKDAY.INTL` function. Start with a spreadsheet that contains the last date of the previous year and a list of your organization's holidays. As you can see in Figure 13.4, your list of holidays should be formatted as dates.

FIGURE 13.4

Start with a sheet containing the last date of the previous year and a list of holidays.

A	B	C	D
1			
2			
3	12/31/2021	Holidays	
4		1/1/2022	
5		1/17/2022	
6		3/17/2022	
7		5/30/2022	
8		7/4/2022	
9		9/1/2022	
10		10/10/2022	
11		11/24/2022	
12		12/24/2022	
13		12/25/2022	
14		12/31/2022	

Part II: Working with Formulas and Functions

In the cell beneath the last date of the previous year, enter this formula:

=WORKDAY.INTL(B3, 1, 1, \$D\$4 : \$D\$14)

At this point, you can copy the formula down to create as many business days as you need (see Figure 13.5).

FIGURE 13.5

Creating a list of business days

A	B	C	D
1			
2			
3	12/31/2021		
4	=WORKDAY.INTL(B3,1,1,\$D\$4:\$D\$14)	Holidays	
5	1/4/2022	1/1/2022	
6	1/5/2022	1/17/2022	
7	1/6/2022	3/17/2022	
8	1/7/2022	5/30/2022	
9	1/10/2022	7/4/2022	
10	1/11/2022	9/1/2022	
11	1/12/2022	10/10/2022	
12	1/13/2022	11/24/2022	
13	1/14/2022	12/24/2022	
14	1/18/2022	12/25/2022	
15	1/19/2022	12/31/2022	
16	1/20/2022		

The WORKDAY.INTL function returns a workday date based on the number of days you tell it to increment. This function has two required arguments and two optional arguments:

Start Date (required) This argument is the date to start from.

Days (required) This argument is the number of days from the start date that you want returned.

Weekends (optional) By default, the WORKDAY.INTL function excludes Saturdays and Sundays, but this third argument allows you to specify which weekdays to exclude as a weekend day. As you enter the WORKDAY.INTL function, Excel displays a menu where you can select the appropriate weekend code.

Holidays (optional) This argument allows you to give Excel a list of dates to exclude in addition to the weekend days.

In this example, we are telling Excel to start from 12/31/2021 and then increment up 1 to give us the next business day after our start date. For our optional arguments, we specify that we need to exclude Saturdays and Sundays, along with the holidays listed in cells \$D\$4 : \$D\$14.

=WORKDAY.INTL(B3, 1, 1, \$D\$4 : \$D\$14)

Be sure to lock down the range for your list of holidays with absolute references so that it remains locked as you copy your formula down.

Extracting parts of a date

Although it may seem trivial, it's often helpful to pick out a specific part of a date. For example, you may need to filter all records that have order dates within a certain month or all employees who have time allocated to Saturdays. In these situations, you would need to pull out the month and workday number from the formatted dates.

Excel provides a simple set of functions to parse dates out into their component parts. These functions are as follows:

YEAR extracts the year from a given date.

MONTH extracts the month from a given date.

DAY extracts the month day number from a given date.

WEEKDAY returns the weekday number for a given date.

WEEKNUM returns the week number for a given date.

Figure 13.6 demonstrates the use of these functions to parse the date in cell C3 into its component parts.

FIGURE 13.6

Extract the parts of a date.

	A	B	C
1			
2			
3			5/16/2021
4			
5	=YEAR(C3)	2021	
6	=MONTH(C3)	5	
7	=DAY(C3)	16	
8	=WEEKDAY(C3)	1	
9	=WEEKNUM(C3)	21	
10			

These functions are fairly straightforward.

The **YEAR** function returns a four-digit number that corresponds to the year of a specified date. This formula returns 2021.

```
=YEAR ("5/16/2021")
```

The **MONTH** function returns a number between 1 and 12 that corresponds to the month of a specified date. This formula returns 5.

```
=MONTH ("5/16/2021")
```

Part II: Working with Formulas and Functions

The DAY function returns a number between 1 and 31 that corresponds to the day of the month represented in a specified date. This formula returns 16.

```
=DAY("5/16/2021")
```

The WEEKDAY function returns a number from 1 to 7 that corresponds to the day of the week (Sunday through Saturday) on which the given date falls. If the date falls on a Sunday, the number 1 is returned. If the date falls on a Monday, the number 2 is returned, and so on. This formula returns 1 because 5/16/2021 falls on a Sunday.

```
=WEEKDAY("5/16/2021")
```

This function actually has an optional `return_type` argument that lets you define which day of the week holds the first position. As you enter the WEEKDAY function, Excel displays a menu where you can select the appropriate `return_type` code.

You can adjust the formula so that the return values 1 through 7 represent Monday through Sunday.

```
=WEEKDAY("5/16/2021", 2)
```

The WEEKNUM function returns the week number in the year for the week in which the specified date occurs. This formula returns 21 because 5/16/2021 falls within week 21 in 2021.

```
=WEEKNUM("5/16/2021")
```

This function actually has an optional `return_type` argument that lets you specify which day of the week defines the start of the week. By default, the WEEKNUM function defines the start of the week as Sunday. As you enter the WEEKNUM function, Excel displays a menu where you can select a different `return_type` code.

Calculating number of years and months between dates

In some cases, you'll be asked to express the difference between two dates in years and months. In these cases, you can create a text string using two DATEDIF functions.

Cell C4 shown in Figure 13.7 contains the following formula:

```
=DATEDIF(A4, B4, "Y") & " Years, " & DATEDIF(A4, B4, "YM") & " Months"
```

We accomplish this task by using two DATEDIF functions joined in a text string with the ampersand (&) operator.

The first DATEDIF function calculates the number of years between the start and end dates by passing the year time unit (Y):

```
DATEDIF(A4, B4, "Y")
```

The second DATEDIF function uses the YM time unit to calculate the number of months ignoring the year portion of the date:

```
DATEDIF(A4, B4, "YM")
```

FIGURE 13.7

Showing the years and months between dates

<code>=DATEDIF(A4,B4,"Y") & " Years, " & DATEDIF(A4,B4,"YM") & " Months"</code>			
Internal			
A	B	C	D
1			
2			
3	Start Date	End Date	Number of Years and Months
4	11/23/1960	8/30/2021	60 Years, 9 Months
5	10/25/1944	8/30/2021	76 Years, 10 Months
6	4/14/1920	8/30/2021	101 Years, 4 Months
7	8/28/1940	8/30/2021	81 Years, 0 Months
8	8/5/1987	8/30/2021	34 Years, 0 Months
9	8/24/1982	8/30/2021	39 Years, 0 Months
10	3/17/1959	8/30/2021	62 Years, 5 Months
11	4/6/1961	8/30/2021	60 Years, 4 Months
12	6/5/1944	8/30/2021	77 Years, 2 Months
13	3/15/1930	8/30/2021	91 Years, 5 Months
14	9/29/1921	8/30/2021	99 Years, 11 Months
15	5/10/1953	8/30/2021	68 Years, 3 Months

We join these two functions with some text of our own to let the users know which number represents years and which represents months:

13

$$=DATEDIF(A4, B4, "Y") & " Years, " & DATEDIF(A4, B4, "YM") & " Months"$$

Converting dates to Julian date formats

Julian dates are often used in manufacturing environments as a timestamp and quick reference for batch number. This type of date coding allows retailers, consumers, and service agents to identify when a product was made and thus the age of the product. Julian dates are also used in programming, the military, and astronomy.

Different industries have their own variations on Julian dates, but the most commonly used variation is made up of two parts: a two-digit number representing the year and the number of elapsed days in the year. For example, the Julian date for 1/1/1960 would be 601. The Julian date for 5/10/21 would be 21130.

Excel has no built-in function to convert a standard date to a Julian date, but Figure 13.8 illustrates how a formula can be used to accomplish the task.

$$=RIGHT(YEAR(A4), 2) & A4 - DATE(YEAR(A4), 1, 0)$$

This formula is really two formulas joined as a text string using the ampersand (&).

The first formula uses the RIGHT function to extract the right two digits of the year number. Note we are using the YEAR function to pull out the year portion from the actual date.

$$=RIGHT(YEAR(A4), 2)$$


For more details on using the RIGHT function, see the section called “Extracting parts of a text string” in Chapter 12.

Part II: Working with Formulas and Functions

FIGURE 13.8

Converting standard dates into Julian dates

A	B
Standard Date	Julian Date
1/1/1960	=RIGHT(YEAR(A4),2)&A4-DATE(YEAR(A4),1,0)
10/25/1944	44299
4/14/1920	20105
8/28/1940	40241
8/5/1987	87217
8/24/1982	82236
3/17/1959	5976
4/6/1961	6196
6/5/1944	44157
3/15/1930	3074
9/29/2000	00273
5/10/2021	21130

The second formula is a bit trickier. Here we have to find out how many days have elapsed since the beginning of the year. For this, we first need to subtract the target date from the last day of the previous year.

`A4 - DATE (YEAR (A4) , 1 , 0)`

You'll note the use of the DATE function. The DATE function allows us to build a date on the fly using three arguments: the year, the month, and the day. The year can be any whole number from 1900 to 9999. The month and date can be any positive or negative number. For example, this formula would return the date serial number for December 1, 2013:

`=DATE (2013 , 12 , 1)`

Note that in our Julian date formula, we are using a zero as our day argument. When you use 0 as the day argument, you are telling Excel that you want the day before the 1st of the given month. In this example, the day before January 1 is December 31. For instance, entering this formula into a blank cell will return December 31, 1959:

`=DATE (1960 , 1 , 0)`

Joining our two formulas together with an ampersand brings the Julian date together:

`=RIGHT (YEAR (A4) , 2) &A4 - DATE (YEAR (A4) , 1 , 0)`

Calculating the percent of year completed and remaining

When building Excel reports and dashboards, oftentimes it's beneficial to calculate the percent of the year that has elapsed and what percent remains. These percentages can be used in other calculations or simply as a notification for your audience.

Figure 13.9 shows a sample of this concept. Note in the Formula bar that we are using the YEARFRAC function.

FIGURE 13.9

Calculating the percent of the year completed

B7	B	C
A	Start Date	End Date
1		
2	1/1/2021	8/30/2021
3		
4		
5		
6	Pcnt of this Year Complete	Percent of this Year Left
7	66%	34%

The YEARFRAC function simply requires a start date and an end date. Once it has those two variables, it calculates the fraction of the year representing the number of days between the start date and end date.

=YEARFRAC(B3, C3)

To get the percent remaining, as shown in cell C7 of Figure 13.9, simply subtract 1 from the YEARFRAC formula:

=1 - YEARFRAC(B3, C3)

Returning the last date of a given month

A common need when working with dates is to calculate dynamically the last date in a given month. Although the last day for most months is fixed, the last day for February varies depending on whether the given year is a leap year.

Figure 13.10 illustrates how to get the last date in February for each date given in order to see which years are leap years.

FIGURE 13.10

Calculating the last day of each date

A	B	C
2	First Day of February	Last Day of February
3	2/1/1999	=DATE(YEAR(B3),MONTH(B3)+1,0)
4	2/1/2000	2/29/2000
5	2/1/2001	2/28/2001
6	2/1/2002	2/28/2002
7	2/1/2003	2/28/2003
8	2/1/2004	2/29/2004
9	2/1/2005	2/28/2005
10	2/1/2006	2/28/2006
11	2/1/2007	2/28/2007
12	2/1/2008	2/29/2008
13	2/1/2009	2/28/2009

Part II: Working with Formulas and Functions

The DATE function lets you build a date on the fly using three arguments: the year, the month, and the day. The year can be any whole number from 1900 to 9999. The month and date can be any positive or negative number.

For example, this formula would return the date serial number for December 1, 2013:

```
=DATE(2013,12,1)
```

When you use 0 as the day argument, you are telling Excel that you want the day before the 1st of the month. For instance, entering this formula into a blank cell will return February 29, 2000:

```
=DATE(2000,3,0)
```

In our example, instead of hard-coding the year and month, we use the YEAR function to get the desired year and the MONTH function to get the desired month. We add 1 to the month so that we go into the next month. This way, when we use 0 as the day, we get the last day of the month in which we're actually interested.

```
=DATE(YEAR(B3),MONTH(B3)+1,0)
```

As you look at Figure 13.10, keep in mind that you can use the formula to get the last day of any month, not just February.

Using the EOMONTH function

The EOMONTH function is an easy alternative to using the DATE function. With the EOMONTH function, you can get the last date of any future or past month. All you need is two arguments: a start date and the number of months in the future or past.

For example, this formula will return the last day of April 2022:

```
=EOMONTH("1/1/2022",3)
```

Specifying a negative number of months will return a date in the past. This formula will return the last day of October 2021:

```
=EOMONTH("1/1/2022",-3)
```

You can combine the EOMONTH function with the TODAY function to get the last day of the current month.

```
=EOMONTH(TODAY(),0)
```

Calculating the calendar quarter for a date

Believe it or not, there is no built-in function to calculate quarter numbers in Excel. If you need to calculate into which calendar quarter a specific date falls, you'll need to create your own formula.

Figure 13.11 demonstrates the following formula used for calculating calendar quarters:

=ROUNDUP(MONTH(B3)/3, 0)

FIGURE 13.11

Calculating calendar quarters

A	B	C
1		
2	Date	Calendar Quarter
3	1/1/2022	=ROUNDUP(MONTH(B3)/3,0)
4	1/21/2022	1
5	3/29/2022	1
6	3/31/2022	1
7	5/31/2022	2
8	7/4/2022	3
9	9/1/2022	3
10	10/14/2022	4
11	11/28/2022	4
12	12/24/2022	4
13	12/25/2022	4
14	12/31/2022	4

The secret to this formula is simple math. Here you're dividing the month number for the given month by 3 and then rounding that number up to the nearest integer. For instance, let's say you want to calculate the quarter into which August falls. Since August is the 8th month of the year, you could divide 8 by 3. That would give you the answer 2.66. Round that number up, and you get 3. Thus, August is in the third quarter of the calendar year.

The following formula does the same thing. We're using the MONTH function to extract the month number from the given date and the ROUNDUP function to force rounding up.

=ROUNDUP(MONTH(B3)/3, 0)

Calculating the fiscal quarter for a date

Many of us work in organizations where the fiscal year does not start in January. Instead, it starts in October or April or any other month. In these organizations, the fiscal quarters can't be calculated in the same way as calendar quarters.

Figure 13.12 demonstrates a clever formula for converting a date into a fiscal quarter using the CHOOSE function. In this example, we're calculating the fiscal quarters when our fiscal year starts in April. The formula seen in the Formula bar shows the following:

=CHOOSE(MONTH(B3), 4, 4, 4, 1, 1, 1, 2, 2, 2, 3, 3, 3)

The CHOOSE function returns an answer from a list of choices based on a position number. If you were to enter the formula =CHOOSE(2, "Gold", "Silver", "Bronze", "Coupon"), you would get Silver because Silver is the second choice in your list of choices. Replace the 2 with a 4, and you would get Coupon, the fourth choice.

Part II: Working with Formulas and Functions

FIGURE 13.12

Calculating fiscal quarters

A	B	C
1		
2	Date	Fiscal Quarter (Fiscal Year Starts in April)
3	1/1/2022	=CHOOSE(MONTH(B3),4,4,4,1,1,1,2,2,2,3,3,3)
4	1/21/2022	4
5	3/29/2022	4
6	3/31/2022	4
7	5/31/2022	1
8	7/4/2022	2
9	9/1/2022	2
10	10/14/2022	3
11	11/28/2022	3
12	12/24/2022	3
13	12/25/2022	3
14	12/31/2022	3

The CHOOSE function's first argument is a required index number. This argument is a number from 1 to as many choices as you list in the next set of arguments. The index number determines which of the next arguments is returned.

The next 254 arguments (only the first one is required) defines your choices and determines what is returned when an index number is provided. If the index number is 1, the first choice is returned. If the index number is 2, the second choice is returned.

The idea here is to use the CHOOSE function to pass a date to a list of quarter numbers:

```
=CHOOSE(MONTH(B3),4,4,4,1,1,1,2,2,2,3,3,3)
```

The formula shown in cell C3 (see Figure 13.12) tells Excel to use the month number for the given date and select a quarter that corresponds to that number. In this case, since the month is January, Excel returns the first choice (January is the first month). The first choice happens to be a 4. January is in the fourth fiscal quarter.

Let's say that your company's fiscal year starts in October instead of April. You can easily compensate for this by simply adjusting your list of choices to correlate with your fiscal year's start month.

```
=CHOOSE(MONTH(B3),2,2,2,3,3,3,4,4,4,1,1,1)
```

Returning a fiscal month from a date

In some organizations, the operationally recognized months don't start on the 1st and end on the 30th or 31st. Instead, they have specific days marking the beginning and end of a month. For instance, you may work in an organization where each fiscal month starts on the 21st and ends on the 20th of the next month. In these organizations, it's important to be able to translate a standard date into their own fiscal months.

Figure 13.13 demonstrates a formula for converting a date into a fiscal month using the EOMONTH function in conjunction with the TEXT function. In this example, we're calculating the fiscal month where our fiscal month starts on the 21st and ends on the 20th of the next month. The formula in cell C3 shows the following:

```
=TEXT(EOMONTH(B3-20,1), "mmm")
```

FIGURE 13.13

Calculating fiscal months

A	B	C
1		Fiscal Month (Starts on the 21st and ends on the 20th of the Next Month)
2	Date	
3	1/1/2022	=TEXT(EOMONTH(B3-20,1),"mmm")
4	1/1/2022	Jan
5	1/21/2022	Feb
6	3/20/2022	Mar
7	3/31/2022	Apr
8	4/21/2022	May
9	6/20/2022	Jun
10	6/21/2022	Jul
11	7/21/2022	Aug

In this formula, we're first taking our date (in B3) and going back 20 days by subtracting 20. Then we are using that new date in the EOMONTH function to get the last day next month:

```
EOMONTH(B3-20,1)
```

We then wrap that in a TEXT function to format the resulting date into a three-letter month name:

```
=TEXT(EOMONTH(B3-20,1), "mmm")
```

Calculating the date of the Nth weekday of the month

Many analytical processes rely on knowing the dates of specific events. For example, if payroll processing occurs the second Friday of every month, it's beneficial to know which dates in the year represent the second Friday of each month.

Using the date functions covered thus far in this chapter, you can build dynamic date tables that automatically provide you with the key dates you need.

Figure 13.14 illustrates such a table. In this table, formulas calculate the Nth weekday for each month listed. The idea is to fill in the years and months you need and then tell it what number occurrence of each weekday you are seeking. In this example, cell B2 shows that we are looking for the second occurrence of each weekday.

Part II: Working with Formulas and Functions

FIGURE 13.14

A dynamic date table calculating the Nth occurrence of each weekday

	A	B	C	D	E	F	G	H	I
1		Nth Occurrence							
2		2							
3									
4			1	2	3	4	5	6	7
5	YEAR	MONTH	Nth Sun of the Month	Nth Mon of the Month	Nth Tues of the Month	Nth Wed of the Month	Nth Thur of the Month	Nth Fri of the Month	Nth Sat of the Month
6	2021	1	1/10/2021	1/11/2021	1/12/2021	1/13/2021	1/14/2021	1/8/2021	1/9/2021
7	2021	2	2/14/2021	2/8/2021	2/9/2021	2/10/2021	2/11/2021	2/12/2021	2/13/2021
8	2021	3	3/14/2021	3/8/2021	3/9/2021	3/10/2021	3/11/2021	3/12/2021	3/13/2021
9	2021	4	4/11/2021	4/12/2021	4/13/2021	4/14/2021	4/8/2021	4/9/2021	4/10/2021
10	2021	5	5/9/2021	5/10/2021	5/11/2021	5/12/2021	5/13/2021	5/14/2021	5/8/2021
11	2021	6	6/13/2021	6/14/2021	6/8/2021	6/9/2021	6/10/2021	6/11/2021	6/12/2021
12	2021	7	7/11/2021	7/12/2021	7/13/2021	7/14/2021	7/8/2021	7/9/2021	7/10/2021
13	2021	8	8/8/2021	8/9/2021	8/10/2021	8/11/2021	8/12/2021	8/13/2021	8/14/2021
14	2021	9	9/12/2021	9/13/2021	9/14/2021	9/8/2021	9/9/2021	9/10/2021	9/11/2021
15	2021	10	10/10/2021	10/11/2021	10/12/2021	10/13/2021	10/14/2021	10/8/2021	10/9/2021
16	2021	11	11/14/2021	11/8/2021	11/9/2021	11/10/2021	11/11/2021	11/12/2021	11/13/2021
17	2021	12	12/12/2021	12/13/2021	12/14/2021	12/8/2021	12/9/2021	12/10/2021	12/11/2021

Cell C6 (see Figure 13.14) contains the following formula:

```
=DATE ($A6 , $B6 , 1) +C$4 -WEEKDAY (DATE ($A6 , $B6 , 1) ) + ($B$2 - (C$4 >=WEEKDAY  
 (DATE ($A6 , $B6 , 1) ))) *7
```

This formula applies some basic math to calculate which date within the month should be returned given a specific week number and occurrence.

To use the table in Figure 13.14, simply enter the years and months you are targeting, starting in cells A6 and B6. Then adjust the occurrence number you need in cell B2.

If you are looking for the first Monday of each month, enter a 1 in cell B2 and look in the Monday column. If you are looking for the third Thursday of each month, enter a 3 in cell B2 and look in the Thursday column.

Calculating the date of the last weekday of the month

You can leverage the functions covered in this chapter thus far to build a dynamic date table that automatically provides you with the last instance of a given weekday. For instance, Figure 13.15 illustrates a table that calculates the last Sunday, Monday, Tuesday, and so forth for each month listed.

Cell C4 (see Figure 13.15) contains the following formula:

```
=DATE ($A4 , $B4+1 , 1) -WEEKDAY (DATE ($A4 , $B4+1 , C$2) )
```

This formula applies some basic math to calculate which date within the month should be returned given a specific year, month, and week number.

FIGURE 13.15

A dynamic date table calculating the last weekday in each month

	A	B	C	D	E	F	G	H	I
2			7	6	5	4	3	2	1
3	YEAR	MONTH	Last Sun of the Month	Last Mon of the Month	Last Tues of the Month	Last Wed of the Month	Last Thurs of the Month	Last Fri of the Month	Last Sat of the Month
4	2021	1	1/31/2021	1/25/2021	1/26/2021	1/27/2021	1/28/2021	1/29/2021	1/30/2021
5	2021	2	2/28/2021	2/22/2021	2/23/2021	2/24/2021	2/25/2021	2/26/2021	2/27/2021
6	2021	3	3/28/2021	3/29/2021	3/30/2021	3/31/2021	3/25/2021	3/26/2021	3/27/2021
7	2021	4	4/25/2021	4/26/2021	4/27/2021	4/28/2021	4/29/2021	4/30/2021	4/24/2021
8	2021	5	5/30/2021	5/31/2021	5/25/2021	5/26/2021	5/27/2021	5/28/2021	5/29/2021
9	2021	6	6/27/2021	6/28/2021	6/29/2021	6/30/2021	6/24/2021	6/25/2021	6/26/2021
10	2021	7	7/25/2021	7/26/2021	7/27/2021	7/28/2021	7/29/2021	7/30/2021	7/31/2021
11	2021	8	8/29/2021	8/30/2021	8/31/2021	8/25/2021	8/26/2021	8/27/2021	8/28/2021
12	2021	9	9/26/2021	9/27/2021	9/28/2021	9/29/2021	9/30/2021	9/24/2021	9/25/2021
13	2021	10	10/31/2021	10/25/2021	10/26/2021	10/27/2021	10/28/2021	10/29/2021	10/30/2021
14	2021	11	11/28/2021	11/29/2021	11/30/2021	11/24/2021	11/25/2021	11/26/2021	11/27/2021
15	2021	12	12/26/2021	12/27/2021	12/28/2021	12/29/2021	12/30/2021	12/31/2021	12/25/2021

To use the table in Figure 13.15, simply enter the years and months that you are targeting, starting in cells A4 and B4. The idea is to use this table in your Excel models as a place to which you can link or simply copy from to get the dates you need.

Extracting parts of a time

It's often helpful to pick out a specific part of a time. Excel provides a simple set of functions to parse time out into its component parts. These functions are as follows:

HOUR extracts the hour portion of a given time value.

MINUTE extracts the minute portion of a given time value.

SECOND extracts the second portion of a given time value.

Figure 13.16 demonstrates the use of these functions to parse the time in cell C3 into its component parts.

FIGURE 13.16

Extract the parts of a time.

	A	B	C
1			
2			
3			6:15:27 AM
4			
5	=HOUR(C3)	6	
6	=MINUTE(C3)	15	
7	=SECOND(C3)	27	

Part II: Working with Formulas and Functions

These functions are fairly straightforward.

The HOUR function returns a number between 0 and 23 corresponding to the hour of a given time. This formula returns 6.

```
=HOUR ("6:15:27 AM")
```

The MINUTE function returns a number between 0 and 59 corresponding to the minutes of a given time. This formula returns 15.

```
=MINUTE ("6:15:27 AM")
```

The SECOND function returns a number between 0 and 59 corresponding to the seconds of a given time. This formula returns 27.

```
=SECOND ("6:15:27 AM")
```

Calculating elapsed time

One of the more common calculations done with time values is calculating elapsed time, that is, how many hours and minutes between a start time and an end time.

The table in Figure 13.17 shows a list of start and end times, along with calculated elapsed times. Look at Figure 13.17, and you can see that the formula in cell D4 is as follows:

```
=IF(C4 < B4, 1 + C4 - B4, C4 - B4)
```

To get the elapsed time between a start and end time, all you need to do is to subtract the start time from the end time. There is a catch, however. If the end time is less than the start time, you have to assume that the clock has been running for a full 24-hour period, effectively looping back the clock.

In these cases, you have to add a 1 to the time to represent a full day. This ensures that you don't have negative elapsed times.

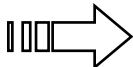
In our elapsed time formula, we use an IF function to check whether the end time is less than the start time. If so, we add a 1 to our simple subtraction. If not, we can just do the subtraction.

```
=IF(C4 < B4, 1 + C4 - B4, C4 - B4)
```

FIGURE 13.17

Calculating elapsed time

A	B	C	D
1			
2			
3	Start Time	End Time	Elapsed Minutes:Seconds
4	8:57:50 AM	10:04:39 AM	=IF(C4 < B4, 1 + C4 - B4, C4 - B4)
5	4:35:20 PM	4:23:23 PM	23:48
6	8:24:35 AM	4:14:36 PM	7:50
7	3:10:39 PM	9:50:59 AM	18:40
8	2:33:22 PM	2:01:49 PM	23:28
9	8:42:35 AM	11:16:31 AM	2:33
10	11:20:24 AM	9:36:17 AM	22:15
11	3:56:53 PM	2:05:17 PM	22:08
12	3:33:16 PM	10:46:08 AM	19:12



For more details on the `IF` function, see “Checking if a simple condition is met” in Chapter 14.

Rounding time values

It's often necessary to round time to a particular increment. For instance, if you're a consultant, you may always want to round time up to the next 15-minute increment or down to 30-minute increments.

Figure 13.18 demonstrates how you can round to 15- and 30-minute increments.

FIGURE 13.18

Rounding time values to 15- and 30-minute increments

A	B	C	D	E	F
1					
2					
3	Start Time	End Time	Elapsed Minutes:Seconds	Round Up to Nearest 15 Minutes	Round Down to Nearest 30 minutes
4	8:49:12 AM	10:03:56 AM	1:14	10:15	10:00
5	10:58:31 AM	10:18:55 AM	23:20	10:30	10:00
6	9:23:37 AM	1:48:26 PM	4:24	14:00	13:30
7	8:39:16 AM	4:40:22 PM	8:01	16:45	16:30

The formula in cell E4 is as follows:

```
=ROUNDUP (C4*24/0.25, 0) * (0.25/24)
```

The formula in cell F4 is as follows:

```
=ROUNDDOWN (C4*24/0.5, 0) * (0.5/24)
```

You can round a time value to the nearest hour by multiplying the time by 24, passing that value to the `ROUNDUP` function, and then dividing the result by 24. For instance, this formula would return 7:00:00 AM:

```
=ROUNDUP ("6:15:27" * 24, 0) / 24
```

To round up to 15-minute increments, you simply divide 24 by 0.25 (a quarter). This formula would return 6:30:00 AM:

```
=ROUNDUP ("6:15:27" * 24 / 0.25, 0) * (0.25/24)
```

To round down to 30-minute increments, divide 24 by 0.5 (a half). This formula would return 6:00:00 AM:

```
=ROUNDDOWN ("6:15:27" * 24 / 0.5, 0) * (0.5/24)
```



For more details on the `ROUNDDOWN` and `ROUNDUP` functions, see Chapter 11, “Using Formulas for Common Mathematical Operations.”

Converting decimal hours, minutes, or seconds to a time

It's not uncommon to get a feed from an external source where the times are recorded in decimal hours. For example, for 1 hour and 30 minutes, you see 1.5 instead of the standard 1:30. You can easily correct this by dividing the decimal hour by 24 and then formatting the result as a time.

Figure 13.19 shows some example decimal hours and the converted times.

FIGURE 13.19

Converting decimal hours to hours and minutes

	A	B	C
1			
2	Decimal Hours	Hours:Minutes	
3	11.50	=B3/24	
4	13.75	13:45	
5	18.25	18:15	
6	11.35	11:21	
7	12.45	12:27	
8	15.60	15:36	
9	18.36	18:21	

Dividing the decimal hour by 24 will result in a decimal that Excel recognizes as a time value.

To convert decimal minutes into time, divide the number by 1440. This formula will return 1:04 (one hour and four minutes):

=64.51/1440

To convert decimal seconds into time, divide the number by 86400. This formula will return 0:06 (six minutes):

=390.45/86400

Adding hours, minutes, or seconds to a time

Because time values are nothing more than a decimal extension of the date serial number system, you can add two time values together to get a cumulative time value. In some cases, you may want to add a set number of hours and minutes to an existing time value. In these situations, you can use the TIME function.

Cell D4 in Figure 13.20 contains this formula:

=C4+TIME(5,30,0)

In this example, we're adding 5 hours and 30 minutes to all of the times in the list.

FIGURE 13.20

Adding a set number of hours and minutes to an existing time value

A	B	C	D
1			
2			
3	Start Time	End time if working 5 hours and 30 minutes	
4	9:23:46 AM	=C4+TIME(5,30,0)	
5	9:18:00 AM	2:48:00 PM	
6	11:41:51 AM	5:11:51 PM	
7	9:25:12 AM	2:55:12 PM	
8	11:59:51 AM	5:29:51 PM	
9	4:57:30 PM	10:27:30 PM	
10	3:56:53 PM	9:26:53 PM	
11	2:02:50 PM	7:32:50 PM	
12	4:37:55 PM	10:07:55 PM	
13	12:45:30 PM	6:15:30 PM	

The TIME function lets you build a time value on the fly using three arguments: hour, minute, and second. For example, this formula would return the time value 2:30:30 PM:

```
=TIME(14,30,30)
```

To add a certain number of hours to an existing time value, simply use the TIME function to build a new time value and then add them together. This formula adds 30 minutes to the existing time, resulting in a time value of 3:00:00 PM:

```
="2:30:00 PM"+TIME(0,30,0)
```


Using Formulas for Conditional Analysis

IN THIS CHAPTER

Understanding conditional analysis

Performing conditional calculations

Excel provides several worksheet functions for performing conditional analysis. You'll use some of those functions in this chapter. Conditional analysis means performing different actions depending on whether a condition is met.

Understanding Conditional Analysis

A condition is a value or expression that returns TRUE or FALSE. Based on the value of the condition, a formula can branch into two separate calculations. That is, when the condition returns TRUE, one value or expression is evaluated while the other is ignored. A FALSE condition reverses the flow of the formula, and the first value or expression is ignored while the other is evaluated.

In the following sections, you'll explore some of the logical functions available in Excel.

 This book's website, www.wiley.com/go/excel365bible, includes a copy of the sample workbook for this chapter. The file is named `Conditional Analysis.xlsx`.

Checking if a simple condition is met

Figure 14.1 shows a list of states with monthly gas prices. Say you want to determine whether a state's price in a given month is above or below the average of all the states for the same month. For higher-than-average prices, you will report "High" and for lower than average, "Low." A grid below the data will be used to report the results.

```
=IF(C3>AVERAGE(C$3:C$11), "High", "Low")
```

Part II: Working with Formulas and Functions

FIGURE 14.1

Monthly gas prices by state

C14	A	B	C	D	E	F	G
1							
2	State	Aug-13	Sep-13	Oct-13	Nov-13	Dec-13	
3	California	3.919	3.989	3.829	3.641	3.642	
4	Colorado	3.569	3.582	3.410	3.231	3.122	
5	Florida	3.614	3.558	3.388	3.377	3.516	
6	Massachusetts	3.761	3.703	3.518	3.419	3.520	
7	Minnesota	3.577	3.540	3.318	3.143	3.113	
8	New York	3.933	3.879	3.700	3.633	3.736	
9	Ohio	3.542	3.512	3.317	3.231	3.281	
10	Texas	3.509	3.383	3.180	3.104	3.171	
11	Washington	3.855	3.767	3.567	3.373	3.348	
12							
13							
14	California	High	High	High	High	High	
15	Colorado	Low	Low	Low	Low	Low	
16	Florida	Low	Low	Low	High	High	

The `IF` function is the most basic conditional analysis function in Excel. It has three arguments: the condition, what to do if the condition is true, and what to do if the condition is false.

The condition argument in this example is `C3>AVERAGE (C$3 : C$11)`. Condition arguments must be structured to return TRUE or FALSE, and this usually means that there is a comparison operation (like an equal sign or greater-than sign) or another worksheet function that returns TRUE or FALSE (like `ISERR` or `ISBLANK`). The example condition has a greater-than sign and compares the value in C3 to the average of all the values in C3 : C11.

If our condition argument returns TRUE, the second argument of the `IF` function is returned to the cell. The second argument is `High`, and since the value in C3 is indeed larger than the average, cell C14 shows the word `High`.

Cell C15 compares the value in C4 to the average. Because it is lower, the condition argument returns FALSE, and the third argument is returned. Cell C15 shows `Low`, the third argument of the `IF` function.

Checking for multiple conditions

Simple conditions like the one shown in Figure 14.1 can be strung together. These are known as *nesting* functions. The `value_if_true` and `value_if_false` arguments can contain simple conditions of their own. This allows you to test more than one condition where subsequent conditions are dependent on the first one.

Figure 14.2 shows a spreadsheet with two user input fields for the type of automobile and a property of that automobile type. The properties are listed in two ranges below the user input

fields. When the user selects the type and property, we want a formula to report whether the user has identified a coupe, a sedan, a pickup, or an SUV.

```
=IF(E2="Car", IF(E3="2-door", "Coupe", "Sedan"), IF(E3="Has Bed", "Pickup", "SUV"))
```

FIGURE 14.2

A model for selecting an automobile

	D	E	F	G	H	I	J	K
1								
2	Auto Type:	Truck			Which Auto:	Pickup		
3	Auto Property:	Has Bed						
4								
5		Car						
6		2-door						
7		4-door						
8								
9		Truck						
10		Has Bed						
11		No Bed						

With some conditional analysis, the result of the first condition causes the second condition to change. In this case, if the first condition is Car, the second condition is 2-door or 4-door. But if the first condition is Truck, the second condition changes to either Has Bed or No Bed.

You've seen that Excel provides the `IF` function to perform conditional analysis. You can also nest `IF` functions; that is, you can use another `IF` function as an argument to the first `IF` function when you need to check more than one condition. In this example, the first `IF` checks the value of E2. Rather than returning a value if TRUE, the second argument is another `IF` formula that checks the value of cell E3. Similarly, the third argument doesn't simply return a value of FALSE, but it contains a third `IF` function that also evaluates cell E3.

In Figure 14.2, the user has selected Truck. The first `IF` returns FALSE because E2 doesn't equal Car and the FALSE argument is evaluated. In that argument, E3 is seen to be equal to Has Bed and the TRUE condition (Pickup) is returned. If the user had selected No Bed, the FALSE condition (SUV) would have been the result.

Validating conditional data

The user input fields in Figure 14.2 are data validation lists. The user can make selections from a drop-down box rather than typing in the values. The data validation in cell E3 uses an interesting technique with an `INDIRECT` function to change its list depending on the value in E2.

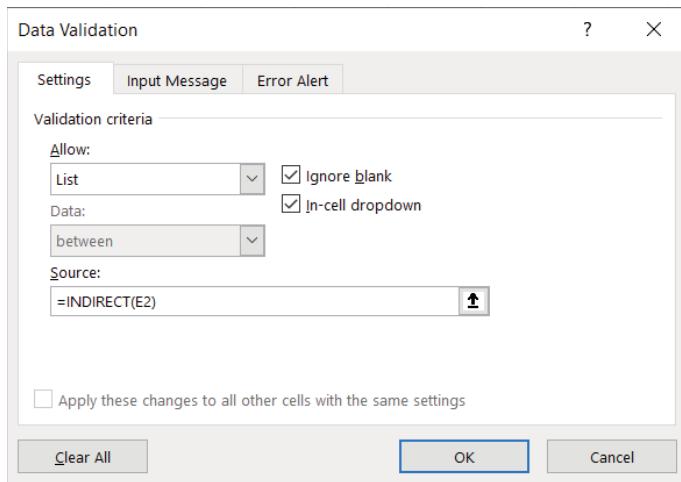
There are two named ranges in the worksheet. The range named Car points to E6 : E7, and the range named Truck points to E10 : E11. The names are identical to choices in the

Part II: Working with Formulas and Functions

E2 data validation list. Figure 14.3 shows the Data Validation dialog box for cell E3. The source is an `INDIRECT` function with E2 as the argument.

FIGURE 14.3

Data validation using `INDIRECT`



The `INDIRECT` function takes a text argument that it resolves into a cell reference. In this case, since E2 is `Truck`, the formula becomes `=INDIRECT ("Truck")`. Because `Truck` is a named range, `INDIRECT` returns a reference to E10:E11, and the values in those cells become the choices. If E2 contained `Car`, `INDIRECT` would return E6:E7, and those values become the choices.

One problem with this type of conditional data validation is that when the value in E2 is changed, the value in E3 does not change. The choices in E3 change, but the user still has to select from the available choices or your formulas may return inaccurate results.

Looking up values

When you have too many nested `IF` functions, your formulas can become long and hard to manage. Figure 14.4 shows a slightly different setup to the auto-selector model. Instead of hard-coding the results in nested `IF` functions, we enter the results into the cells next to their properties (for example, `Sedan` is entered in the cell next to `4-door`).

The new formula is as follows:

```
=IF(E2="Car", VLOOKUP(E3, E6:F7, 2, FALSE), VLOOKUP(E3, E10:F11, 2, FALSE))
```

This formula can now be used to return the automobile. The `IF` condition is the same, but now a `TRUE` result looks up the proper value in E6:F7, and a `FALSE` result looks it up in E10:F11. You can learn more about `VLOOKUP` in Chapter 15.

FIGURE 14.4

A different auto-selector model

The screenshot shows a Microsoft Excel spreadsheet with the following data and formula:

	D	E	F	G	H	I	J
1							
2	Auto Type:	Car		Which Auto:	Sedan		
3	Auto Property:	4-door					
4							
5		Car					
6		2-door	Coupe				
7		4-door	Sedan				
8							
9		Truck					
10		Has Bed	Pickup				
11		No Bed	SUV				

The formula in cell H2 is: =IF(E2="Car",VLOOKUP(E3,E6:F7,2,FALSE),VLOOKUP(E3,E10:F11,2,FALSE))

Checking if Condition1 AND Condition2 are met

In addition to nesting conditional functions, we can evaluate together inside an AND function. This is useful when two or more conditions need to be evaluated at the same time to determine where the formula should branch.

Figure 14.5 shows a listing of inventory items, their quantities, and the discount that applies when they are sold. The inventory items are structured with three sections divided by hyphens. The first section is the department; the second section determines whether the item is a part, a subassembly, or a final assembly; and the third condition is a unique four-digit number. We want to assign a discount of 10 percent only to those items that are in Department 202 and are final assemblies. All other items have no discount.

```
=IF(AND(LEFT(B3, 3) = "202", MID(B3, 5, 3) = "FIN"), 10%, 0%)
```

The IF function returns 10 percent if TRUE and 0 percent if FALSE. For the condition argument (the first argument), you need an expression that returns TRUE if both the first section of the item number is 202 and the second section is FIN. Excel provides the AND function to accomplish this. The AND function takes up to 255 logical arguments separated by commas. Logical arguments are expressions that return either TRUE or FALSE. For this example, we're using only two logical arguments.

The first logical argument, LEFT(B3, 3) = "202", returns TRUE if the first three characters of B3 are equal to 202. The second logical argument, MID(B3, 5, 3) = "FIN", returns TRUE if the three digits starting at the fifth position are equal to FIN. Text manipulation functions are discussed in Chapter 12.

With the AND function, all logical arguments must return TRUE for the entire function to return TRUE. If even one of the logical arguments returns FALSE, then the AND function returns FALSE. Table 14.1 shows the results of the AND function with two logical arguments.

Part II: Working with Formulas and Functions

FIGURE 14.5

An inventory listing

	A	B	C	D	E	F	G
1							
2		Inventory Item	Quantity	Discount			
3		202-PRT-3013	76	0%			
4		201-FIN-1452	69	0%			
5		202-FIN-8206	12	10%			
6		201-FIN-8238	79	0%			
7		203-FIN-8882	16	0%			
8		202-PRT-9587	87	0%			
9		203-FIN-4614	97	0%			
10		201-PRT-2478	25	0%			
11		202-SLR-1055	14	0%			

TABLE 14.1 A Truth Table for the AND Function

First Logical Argument	Second Logical Argument	Result of AND Function
TRUE	TRUE	TRUE
TRUE	FALSE	FALSE
FALSE	TRUE	FALSE
FALSE	FALSE	FALSE

In cell D3, the first logical condition returns TRUE because the first three characters of the item number are 202. The second logical condition returns FALSE because the middle section of the item number is PRT, not FIN. According to Table 14.1, a TRUE condition and a FALSE condition return FALSE, and 0 percent is the result. Cell D5, on the other hand, returns TRUE because both logical conditions return TRUE.

Referring to logical conditions in cells

The AND function in Figure 14.5 includes two logical conditions that evaluate to TRUE or FALSE. The arguments to AND can also reference cells as long as those cells evaluate to TRUE or FALSE. When building a formula with the AND function, it can be useful to break out the logical conditions into their own cells. In Figure 14.6, the inventory listing is modified to show two extra columns. These columns can be inspected to understand why a particular item does or does not get the discount.

With these modifications, the result doesn't change, but the formula becomes

=IF(AND(D3, E3), 10%, 0%)

Checking if Condition1 OR Condition2 are met

In Figure 14.6, we applied a discount to certain products based on their item number. In this example, we want to expand the number of products eligible for the discount. As before, only final assembly products will get the discount, but the departments will be expanded to include both department 202 and department 203. Figure 14.7 shows the inventory list and the new discount schedule.

```
=IF(AND(OR(LEFT(B3,3)="202",LEFT(B3,3)="203"),MID(B3,5,3)="FIN"),10%,0%)
```

FIGURE 14.6

A modified inventory listing

F3	A	B	C	D	E	F
1						
2	Inventory Item	Quantity	IsDept202	IsFinalAssembly	Discount	
3	202-PRT-3013	76	TRUE	FALSE	0%	
4	201-FIN-1452	69	FALSE	TRUE	0%	
5	202-FIN-8206	12	TRUE	TRUE	10%	
6	201-FIN-8238	79	FALSE	TRUE	0%	
7	203-FIN-8882	16	FALSE	TRUE	0%	
8	202-PRT-9587	87	TRUE	FALSE	0%	
9	203-FIN-4614	97	FALSE	TRUE	0%	
10	201-PRT-2478	25	FALSE	FALSE	0%	
11	202-SUB-1955	14	TRUE	FALSE	0%	

FIGURE 14.7

A revised discount scheme

F3	A	B	C	D	E	F	G
2	Inventory Item	Quantity	Discount				
3	202-PRT-3013	76	0%				
4	201-FIN-1452	69	0%				
5	202-FIN-8206	12	10%				
6	201-FIN-8238	79	0%				
7	203-FIN-8882	16	10%				
8	202-PRT-9587	87	0%				
9	203-FIN-4614	97	10%				
10	201-PRT-2478	25	0%				
11	202-SUB-1955	14	0%				

We've expanded the conditional argument to the IF function to account for the changes in the discount scheme. The AND function is restrictive because all of the arguments must be TRUE for AND to return TRUE. Conversely, the OR function is inclusive. With OR, if any one

of the arguments is TRUE, the entire function returns TRUE. In this example, we've nested an OR function inside the AND function—we've made it one of the arguments. Table 14.2 shows a truth table for how our nested functions work.

Cell D9 in Figure 14.7 shows a previously undiscounted product that receives a discount under the new scheme. The OR section, OR (LEFT(B9, 3) = "202", LEFT(B9, 3) = "203"), returns TRUE because one of its arguments returns TRUE.

TABLE 14.2 A Truth Table for an OR Function Nested in an AND Function

OR Logical 1	OR Logical 2	OR Result	AND Logical 2	Final Result
TRUE	TRUE	TRUE	TRUE	TRUE
TRUE	FALSE	TRUE	TRUE	TRUE
FALSE	TRUE	TRUE	TRUE	TRUE
FALSE	FALSE	FALSE	TRUE	FALSE
TRUE	TRUE	TRUE	FALSE	FALSE
TRUE	FALSE	TRUE	FALSE	FALSE
FALSE	TRUE	TRUE	FALSE	FALSE
FALSE	FALSE	FALSE	FALSE	FALSE

Performing Conditional Calculations

Simple conditional functions like IF generally work on only one value or cell at a time. Excel provides some different conditional functions for aggregating data, such as summing or averaging.

In the following sections, you'll dive into some of the techniques for applying calculations based on a given set of conditions.

Summing all values that meet a certain condition

Figure 14.8 shows a listing of accounts with positive and negative values. We want to sum all of the negative balances, which we will later compare to the sum of all the positive balances to ensure that they are equal. Excel provides the SUMIF function to sum values based on a condition.

=SUMIF(C3:C12, "<0")

SUMIF takes each value in C3 : C12 and compares it to the condition (the second argument in our function). If the value is less than zero, it meets the condition and is included in the sum. If it is zero or greater, the value is ignored. Text values and blank cells are also ignored. For the example in Figure 14.8, cell C3 is evaluated first. Because it is greater than zero, it is ignored. Next, cell C4 is evaluated. It meets our condition of being less than zero, so it is added to the total. This continues for each cell. When it's complete, cells C4, C7, C8, C9, and C11 are included in the sum, and the others are not.

The second argument of SUMIF, the condition to be met, has quotes around it. Because we're using a less-than sign for this example, we have to create a string that represents the expression.

FIGURE 14.8

Summing values less than zero

C14	B	C
	Account	Balance
2		
3	1510 Equipment	9,863.00
4	1540 Accumulated Depreciation	(9,502.00)
5	1690 Land	5,613.00
6	1915 Other Assets	8,653.00
7	2320 Wages Payable	(6,937.00)
8	2420 Current Portion of Long-term Debt	(6,826.00)
9	2440 Deposits from Customers	(3,717.00)
10	5800 Cost of Goods Sold, Other	73.00
11	5900 Purchase Returns and Allowances	(4,443.00)
12	6300 Charitable Contributions Expense	7,223.00
13		
14	Negative Balances	(31,425.00)
15	Positive Balances	31,425.00

The SUMIF function has an optional third argument called the `sum_range`. So far, we've been applying the condition to the very numbers we are summing. By using the third argument, we can sum a range of numbers but apply our conditions to a different range. Figure 14.9 shows a listing of regions and their associated sales. To sum the sales for the East region, use the formula `=SUMIF(B2:B11, "East", C2:C11)`.

FIGURE 14.9

List of regions and sales values

A	B	C
2	South	2,714,745.31
3	South	1,434,322.83
4	North	1,983,811.70
5	East	929,430.78
6	East	3,154,066.47
7	South	1,264,430.64
8	North	4,674,274.42
9	East	940,684.25
10	South	2,497,381.24
11	South	1,728,260.75
12		
13	Sum on the East	5,024,181.50

Summing greater than zero

Figure 14.8 also shows the total of all the positive balances. The formula for that calculation is `=SUMIF(C3:C12, ">0")`. Note that the only difference between this formula and our example formula is the expression string. Instead of "`<0`" as the second argument, this formula has "`>0`".

We don't have to include zero in our calculation because we're summing and zero never changes a sum. If, however, we were interested in summing numbers greater or less than 1,000, we couldn't simply use "`<1000`" and "`>1000`" as our second arguments because we would never include anything that was exactly 1,000.

When you use greater than or less than a nonzero number in a SUMIF, make either the greater than a greater than or equal to, such as "`>=1000`", or the less than a less than or equal to, such as "`<=1000`". Don't use the equal sign for both; just use one. This will ensure that you include any numbers that are exactly 1,000 in one or the other calculation but not both.

The syntax needed to use your comparison operators can be tricky. Table 14.3 lists a set of simple rules that can help you get it right every time.

TABLE 14.3 Simple Rules for Using Comparison Operators

To Set a Condition Follow These Rules	For Example
Equal to a number or cell reference	Don't use an equal sign or any double quotes.	<code>=SUMIF(A1:A10, 3)</code>
Equal to a string	Don't use an equal sign, but put the string in quotes.	<code>=SUMIF(A1:A10, "book")</code>
Nonequal comparison to a number	Put both the operator and the number in double quotes.	<code>=SUMIF(A1:A10, ">=50")</code>
Nonequal comparison to a string	Put both the operator and the string in double quotes.	<code>=SUMIF(A1:A10, "<>Payroll")</code>
Nonequal comparison to a cell reference or formula	Put the operator in double quotes and concatenate the cell reference or formula with the ampersand (&).	<code>=SUMIF(A1:A10, "<"&C1)</code>

You can use the TODAY function (to get the current date), or most other functions, in the second argument. Figure 14.10 shows a listing of dates and values. To sum a range of numbers that correspond to today, use the formula `=SUMIF(B3:B11, TODAY(), C3:C11)`. To sum only those values that are today or before, concatenate the less-than-or-equal-to sign to the function, such as `=SUMIF(B3:B11, "<=" & TODAY(), C3:C11)`.

FIGURE 14.10

SUMIF using the TODAY function

A	B	C	D
	Date	Value	
2			
3	2/1/2021	10	
4	2/2/2021	20	
5	2/3/2021	30	
6	3/9/2021	40	
7	7/5/2021	50	
8	8/30/2021	60	
9	9/9/2021	70	
10	10/5/2021	80	
11	2/10/2021	90	
12			
13	Sum today	60	
14	Sum today and prior	300	

There are two wildcard characters that you can use in the condition argument to SUMIF. The question mark (?) represents any single character, and the asterisk (*) represents zero, one, or any number of characters. The formula =SUMIF(B2:B11, "?o*", C2:C11) will sum all the values in C2 : C11 that correspond to the values in B2 : B11 where the second character is a lowercase o. If we apply that formula to the data in Figure 14.9, we will get the sum for sales in both the North and South regions because both have a lowercase o as the second letter and East does not.

Summing all values that meet two or more conditions

The limitation of SUMIF shown in Figure 14.9 is that it works with only one condition. The SUMIFS function can be used when more than one condition is needed.

Figure 14.11 shows a partial listing of countries and their gross domestic product (GDP) from 2000 to 2009. We want to total Brazil's GDP from 2003 to 2006. Excel's SUMIFS worksheet function is used to sum values where two or more conditions must be met, such as Country and Year in this example.

```
=SUMIFS(D3:D212, B3:B212, G3, C3:C212, ">=" & G4, C3:C212, "<=" & G5)
```

SUMIFS arguments start with the range that contains the value that you want to sum. The remaining arguments are in pairs that follow the pattern *criteria_range*, *criteria*. Because of the way that the arguments are laid out, SUMIFS will always have an odd number of arguments. The first criteria pair is required—without at least one condition, SUMIFS would be no different than SUM. The remaining pairs of conditions, up to 126 of them, are optional.

```
=SUMIFS(D3:D212, B3:B212, G3, C3:C212, ">=" & G4, C3:C212, "<=" & G5)
```

Part II: Working with Formulas and Functions

FIGURE 14.11

A partial listing of countries and their gross domestic product

	A	B	C	D	E	F	G
2		Country	Year	GDP			
3		Australia	2000	399,594		Country	Brazil
4		Australia	2001	377,207		Start Year	2003
5		Australia	2002	423,676		End Year	2006
6		Australia	2003	539,162			
7		Australia	2004	654,968	Total GDP		3,187,415
8		Australia	2005	730,729			
9		Australia	2006	777,933			
10		Australia	2007	945,364			
11		Australia	2008	1,051,261			
12		Australia	2009	993,349			
13		Belgium	2000	233,354			
14		Belgium	2001	232,686			
15		Belgium	2002	253,689			
16		Belgium	2003	312,285			
17		Belgium	2004	362,160			
18		Belgium	2005	378,006			
19		Belgium	2006	400,337			
20		Belgium	2007	460,280			
21		Belgium	2008	509,765			
22		Belgium	2009	474,580			
23		Brazil	2000	644,734			
24		Brazil	2001	554,185			
25		Brazil	2002	506,043			
26		Brazil	2003	552,383			
27		Brazil	2004	663,734			
28		Brazil	2005	882,043			
29		Brazil	2006	1,089,255			
30		Brazil	2007	1,366,854			
31		Brazil	2008	1,653,538			
32		Brazil	2009	1,622,311			

In this example, each cell in D3 : D212 is added to the total only if the corresponding values in B3 : B212 and C3 : C212 meet their respective conditions. The condition for B3 : B212 is that it matches whatever is in cell G3. There are two year conditions because we need to define the lower bound and upper bound of our year range. The lower bound is in cell G4, and the upper bound is in cell G5. Those two cells are concatenated with greater than or equal to and less than or equal to, respectively, to create the year conditions. Only if all three conditions are true is the value included in the total.

Summing if values fall between a given date range

One way that you can use SUMIF with two or more conditions is to add or subtract multiple SUMIF calculations. If the two conditions operate on the same range, this is an effective way to use multiple conditions. When you want to test different ranges, the formulas get tricky because you have to make sure you don't double count values.

Figure 14.12 shows a list of dates and amounts. We want to find the sum of the values that are between June 23 and June 29, inclusive. The starting and ending dates will be put in cells F4 and F5, respectively. Cell F7 contains the following formula:

```
=SUMIF(B3:B20, "<="&F5, C3:C20) - SUMIF(B3:B20, "<"&F4, C3:C20)
```

FIGURE 14.12

Summing values that are between two dates

F7	=SUMIF(B3:B20,"<="&F5,C3:C20)-SUMIF(B3:B20,"<"&F4,C3:C20)					
A	B	C	D	E	F	G
1						
2	Date	Amount				
3	6/20/2020	843.77				
4	6/21/2020	400.60	Start Date	6/23/2020		
5	6/22/2020	396.54	End Date	6/29/2020		
6	6/23/2020	656.56				
7	6/24/2020	249.77	Sum	4,321.42		
8	6/25/2020	318.04				

This technique subtracts one SUMIF from another to get the desired result. The first SUMIF, `=SUMIF(B3:B20, "<="&F5, C3:C20)`, returns the sum of the values less than or equal to the date in F5, June 29 in this example. The conditional argument is the less-than-or-equal-to operator concatenated to the cell reference F5. If that was the whole formula, the result would be 5,962.33. But we want only values that are also greater than or equal to June 23. That means we want to exclude values that are less than June 23. The second SUMIF achieves that. Sum everything less than or equal to the later date, and subtract everything less than the earlier date to get the sum of values between the two dates.

Using SUMIFS

You may even find SUMIFS to be more intuitive than the subtraction technique. The formula `=SUMIFS(C3:C20, B3:B20, "<="&F5, B3:B20, ">="&F4)` sums the values in C3:C20 that correspond to the values in B3:B20 that meet the criteria pairs. The first criteria pair is identical to the first SUMIF criteria, `"<="&F5`. The second criteria pair limits the dates to greater than or equal to the start date.

Getting a count of values that meet a certain condition

Summing values isn't the only aggregation you can do in Excel. Like SUMIF and SUMIFS, Excel provides functions for conditionally counting values in a range.

In Figure 14.13, there is a partial listing of countries and their gross domestic product from 2000 to 2009. We want to know how many times the GDP was greater than or equal to 1 million. The criterion to be applied will be in cell G3.

```
=COUNTIF(D3:D212, G3)
```

Part II: Working with Formulas and Functions

FIGURE 14.13

A partial listing of countries and their gross domestic product

	A	B	C	D	E	F	G
1							
2	Country	Year	GDP				
3	Australia	2000	399,594	Criterion	>=1000000		
4	Australia	2001	377,207				
5	Australia	2002	423,676	Count	96		
6	Australia	2003	539,162				
7	Australia	2004	654,968	Between	500000		
8	Australia	2005	730,729	And	1000000		
9	Australia	2006	777,933				
10	Australia	2007	945,364	Count	62		

The COUNTIF function works in a similar manner as the SUMIF function from Figure 14.9. The obvious difference, as the name suggests, is that it counts entries that meet the criteria rather than sum them. Another difference is that there is no optional third argument as there is in SUMIF. With SUMIF, you can sum a range that's different from the range to which the criterion is applied. But with COUNTIF, that wouldn't make sense because counting a different range would get the same result.

The formula in this example uses a slightly different technique to construct the criteria argument. The string concatenation occurs all in cell G3 rather than in the function's second argument. We could have also entered the second argument as ">=1000000" or ">="&G3 rather than just pointing to G3.

TIP

You may notice in the sample file that the formula in G3, =">="&10^6, uses the exponent operator, or caret (^), to calculate 1 million. Representing large numbers using the caret can help reduce errors caused by miscounting the number of zeros that you typed.

Getting a count of values that meet two or more conditions

The SUMIF function has its COUNTIF cousin. Of course, Microsoft couldn't introduce SUMIFS for summing multiple conditions without also introducing COUNTIFS to count them.

In Figure 14.14, there is a list of Alpine Skiing medalists from the 1972 Winter Olympics. We would like to know how many silver medalists have an ö in their name. The letter we're looking for is typed in cell I3 and the type of medal in cell I4.

```
=COUNTIFS(C3:C20,"*"&I3&"*",F3:F20,I4)
```

FIGURE 14.14

1972 Alpine Skiing Olympic medalists

I5	B	C	D	E	F	G	H	I
1	Event	Athlete	Country	Result	Medal			
3	Downhill Men	Bernhard Russi	SUI	01:51.4	GOLD	Name contains	ö	
4	Downhill Men	Roland Collombin	SUI	01:52.1	SILVER	Medal Won	SILVER	
5	Downhill Men	Heini Messner	AUT	01:52.4	BRONZE	Count	3	
6	Slalom Men	Francisco Fernández	ESP	01:49.3	GOLD			
7	Slalom Men	Gustav Thöni	ITA	01:50.3	SILVER			
8	Slalom Men	Roland Thöni	ITA	01:50.3	BRONZE	T	84	
9	Giant Slalom Men	Gustav Thöni	ITA	03:09.6	GOLD	h	104	
10	Giant Slalom Men	Edmund Brugmann	SUI	03:10.7	SILVER	ö	246	
11	Giant Slalom Men	Werner Mattle	SUI	03:11.0	BRONZE	n	110	
12	Downhill Women	Marie-Thérèse Nadig	SUI	01:36.7	GOLD	i	105	
13	Downhill Women	Annemarie Moser-Pröll	AUT	01:37.0	SILVER			

The criteria_range and criteria arguments come in pairs just like in SUMIFS. Whereas SUMIFS will always have an odd number of arguments, COUNTIFS will always have an even number.

The first criteria_range argument is the list of athlete names in C3 : C20. The matching criteria argument, " * " & I3 & " * ", surrounds whatever is in I3 with asterisks. Asterisks are wildcard characters in COUNTIFS that stand for zero, one, or more characters of any kind. By including an asterisk both before and after the character, we're asking Excel to count all of the names that include that character anywhere within the name. That is, we don't care if there are zero, one, or more characters before ö, and we don't care if there are zero, one, or more characters after ö as long as that character is in there somewhere.

The second criteria_range, criteria argument pair counts those entries in F3 : F20 that are SILVER (the value typed into I4). Only those rows where both the first argument pair and the second argument pair match (only rows where the athlete's name contains ö and the medal won was silver) are counted. In this example, Gustav Thöni won the silver in the Men's Slalom, and Annemarie Moser-Pröll placed in both the Women's Downhill and the Women's Giant Slalom for a count of three.

Finding nonstandard characters

The ö was typed into cell I3 by holding down the Alt key and typing 0246 on the numeric keyboard. Don't try to type those numbers on the number keys across the top of your keyboard because it won't work. The number 0246 is the ASCII code that represents ö. You can get the same character by entering the following formula:

```
=CHAR(246)
```

The CHAR Function will give you the character associated with an ASCII code. Every character has an associated ASCII code. So how do you know which ASCII code to use? You can find a character's ASCII code by using the CODE function.

Another look at Figure 14.14 shows some characters in cells H8 : I12. You can see a small table of characters and their codes. In cell I8 is the formula =CODE(H8). The CODE worksheet function returns the ASCII code for the letter that's passed in. In this example, we can see that a capital T is ASCII code 84, a lowercase i is ASCII code 105, and ö is ASCII code 246. Armed with that knowledge, we can use the CHAR function to get any character you need.

Getting the average of all numbers that meet a certain condition

After summing and counting, taking an average of a range of numbers is the next most common aggregator. The average, also known as the *arithmetic mean*, is the sum of the numbers divided by the count of the numbers.

Figure 14.15 once again shows the medalists' results from the 1972 Winter Olympics. We want to determine the average result but only for those skiers from Switzerland. The country code is entered in cell I3 so that it can be easily changed to a different country.

```
=AVERAGEIF(D3:D20, I3, E3:E20)
```

FIGURE 14.15

Averaging results based on a country

			=AVERAGEIF(D3:D20,I3,E3:E20)		
1	D	E	F	G	H
2	Country	Result	Medal		
3	SUI	01:51.4	GOLD	Country	SUI
4	SUI	01:52.1	SILVER		
5	AUT	01:52.4	BRONZE	Average Result	02:12.0
6	ESP	01:49.3	GOLD	SUMIF and COUNTIF	02:12.0
7	ITA	01:50.3	SILVER		
8	ITA	01:50.3	BRONZE		
9	ITA	03:09.6	GOLD		
10	SUI	03:10.7	SILVER		
11	SUI	03:11.0	BRONZE		

Excel provides the AVERAGEIF function to accomplish just what we want. Like its cousin, the SUMIF function, AVERAGEIF has a criteria_range and a criteria argument. The final argument is the range to average. In this example, each cell in E3 : E20 is either included in or excluded from the average depending on whether the corresponding cell in D3 : D20 meets the criteria.

If no rows meet the criteria in AVERAGEIF, the function returns the #DIV/0! error.

Getting the average of all numbers that meet two or more conditions

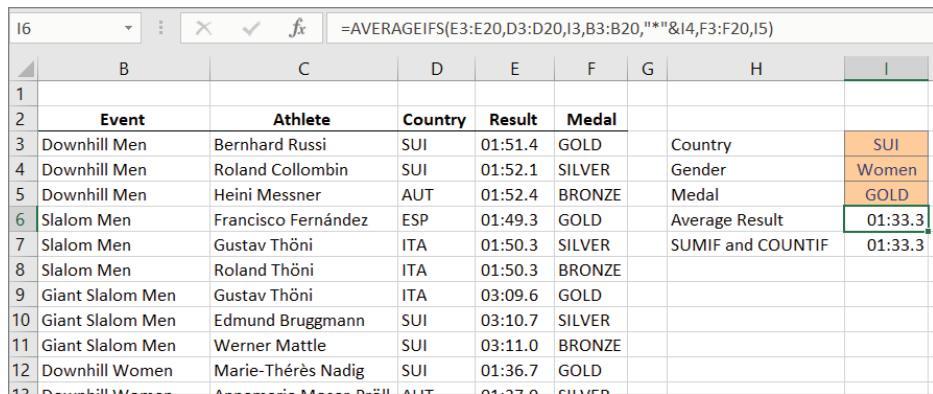
Microsoft introduced AVERAGEIFS along with SUMIFS and COUNTIFS to allow you to average a range of numbers based on more than one condition.

Continuing our analysis of skiing times, Figure 14.16 shows some results of the 1972 Winter Olympics. In this case, we want to determine the average time based on more than one condition. The country, gender, and medal are entered into cells I3:I5. We want to average only those results that meet all three criteria.

=AVERAGEIFS (E3:E20, D3:D20, I3, B3:B20, "*"&I4, F3:F20, I5)

FIGURE 14.16

Averaging on three conditions



	B	C	D	E	F	G	H	I
1								
2	Event	Athlete	Country	Result	Medal	Country	Gender	Medal
3	Downhill Men	Bernhard Russi	SUI	01:51.4	GOLD			SUI
4	Downhill Men	Roland Collombin	SUI	01:52.1	SILVER		Women	
5	Downhill Men	Heini Messner	AUT	01:52.4	BRONZE			GOLD
6	Slalom Men	Francisco Fernández	ESP	01:49.3	GOLD	Average Result		01:33.3
7	Slalom Men	Gustav Thöni	ITA	01:50.3	SILVER	SUMIF and COUNTIF		01:33.3
8	Slalom Men	Roland Thöni	ITA	01:50.3	BRONZE			
9	Giant Slalom Men	Gustav Thöni	ITA	03:09.6	GOLD			
10	Giant Slalom Men	Edmund Bruggmann	SUI	03:10.7	SILVER			
11	Giant Slalom Men	Werner Mattle	SUI	03:11.0	BRONZE			
12	Downhill Women	Marie-Thérèse Nadig	SUI	01:36.7	GOLD			
13	Downhill Women		AUT	01:37.0	SILVER			

The AVERAGEIFS function is structured similarly to the SUMIFS function. The first argument is the range to average, and it's followed by up to 127 pairs of criteria_range/criteria arguments. The three criteria pairs are as follows:

- D3:D20, I3 includes only those rows where the country code is SUI.
- B3:B20, "*"&I4 includes only those rows where the event name ends with the word Women.
- F3:F20, I5 includes only those rows where the medal is GOLD.

When all three conditions are met, the time in the Result column is averaged.

Using Formulas for Matching and Lookups

IN THIS CHAPTER

- Introducing formulas that look up values in a table
- Identifying the worksheet functions used to perform lookups
- Delving into more sophisticated lookup formulas

This chapter discusses various techniques that you can use to look up a value in a range of data. Excel has three worksheet functions (`LOOKUP`, `VLOOKUP`, and `HLOOKUP`) designed for this task, but you may find that these functions don't quite cut it.

In this chapter you'll explore many lookup examples, including alternative techniques that go well beyond the Excel program's normal lookup capabilities.

NOTE

This chapter focuses on the traditional lookup functions available in Excel. With Office 365, Microsoft has introduced new dynamic array functionality (along with the `XLOOKUP` function) designed to replace many of the lookup functions and techniques found in this chapter. You will, however, likely run into these traditional functions and techniques in many of the workbooks you'll touch. It's worth taking time to discover the workings of these functions. To explore the new dynamic array capabilities of Excel and the new `XLOOKUP` function, feel free to visit Chapter 10, "Understanding and Using Array Formulas."

Introducing Lookup Formulas

A lookup formula returns a value from a table by looking up another related value. A common telephone directory (remember those?) provides a good analogy. If you want to find a person's telephone number, you first locate the name (look it up) and then retrieve the corresponding number.

NOTE

We use the term *table* to describe any rectangular range of data. The range does not necessarily need to be an "official" table, as created by choosing `Insert` \Rightarrow `Tables`.

Part II: Working with Formulas and Functions

Several Excel functions are useful when writing formulas to look up information in a table. Table 15.1 describes these functions.

TABLE 15.1 Functions Used in Lookup Formulas

Function	Description
CHOOSE	Returns a specific value from a list of values supplied as arguments.
HLOOKUP	Horizontal lookup. Searches for a value in the top row of a table and returns a value in the same column from a row that you specify in the table.
IF	Returns one value if a condition you specify is TRUE and returns another value if the condition is FALSE.
IFERROR	If the first argument returns an error, the second argument is evaluated and returned. If the first argument does not return an error, then it is evaluated and returned.
INDEX	Returns a value (or the reference to a value) from within a table or range.
LOOKUP	Returns a value from either a one-row or one-column range. Another form of the LOOKUP function works like VLOOKUP but is restricted to returning a value from the last column of a range.
MATCH	Returns the relative position of an item in a range that matches a specified value.
OFFSET	Returns a reference to a range that is a specified number of rows and columns from a cell or range of cells.
VLOOKUP	Vertical lookup. Searches for a value in the first column of a table and returns a value in the same row from a column that you specify in the table.

Leveraging Excel's Lookup Functions

Finding data in a list or table is central to many Excel formulas. Excel provides several functions to assist in looking up data vertically, horizontally, left to right, and right to left. By nesting some of these functions, you can write a formula that looks up the correct data even after the layout of your table changes.

Let's take a look at some of the more common ways to utilize Excel's lookup functions.



This book's website, www.wiley.com/go/excel365bible, includes a copy of the sample workbook for this chapter. The file is named Performing_Lookups.xlsx.

Looking up an exact value based on a left lookup column

Many tables are arranged so that the key piece of data, the data that makes a certain row unique, is in the far-left column. While Excel has many lookup functions, VLOOKUP was designed for just that situation. Figure 15.1 shows a table of employees. We want to fill out a simplified paystub form by pulling the information from this table when an employee's ID is selected.

FIGURE 15.1

A table of employee information

	A	B	C	D	E	F	G	H	I
1									
2	ID	Employee Name	Address	Frequency	Salary	Tax Rate	Insurance	401k	
3	154	Paige Jones	427 John A. Creighton Boulevard A	26	42,900.00	15%	100.00	8%	
4	240	Elijah Ward	888 192nd Street Anytown, USA 12	26	64,600.00	16%	200.00	7%	
5	319	Elizabeth Marshall	530 Dodge Street Anytown, USA 12	52	72,300.00	24%	300.00	3%	
6	331	Cooper Smith	271 Dodge Street Anytown, USA 12	12	99,700.00	20%	300.00	5%	
7	428	Isabella Harris	715 136th Street Anytown, USA 12	26	57,600.00	25%	100.00	5%	
8	451	Kaylee Perez	772 North 30th Street Anytown, US	12	82,800.00	23%	200.00	7%	
9	527	Kimberly Hall	652 Regency Parkway Drive Anytov	12	41,700.00	17%	300.00	4%	
10	540	Jesus Clark	803 Fontenelle Boulevard Anytowr	52	83,100.00	18%	200.00	5%	
11	665	Kylie Woods	245 Fontenelle Boulevard Anytowr	26	70,400.00	21%	100.00	1%	
12	981	Jackson Stephens	827 Harrison Street Anytown, USA	52	50,200.00	18%	100.00	8%	

The user will select an employee ID from a data validation list in cell L3 (see Figure 15.2). From that piece of data, the employee's name, address, and other information will be pulled into the form. The formulas for the paystub form in Figure 15.2 are shown here:

Employee Name

```
=VLOOKUP ($L$3,$B$3:$I$12,2,FALSE)
```

Pay

```
=VLOOKUP ($L$3,$B$3:$I$12,5,FALSE)/VLOOKUP ($L$3,$B$3:$I$12,4,FALSE)
```

Taxes

```
= (M7-O8-O9) *VLOOKUP ($L$3,$B$3:$I$12,6,FALSE)
```

Insurance

```
=VLOOKUP ($L$3,$B$3:$I$12,7,FALSE)
```

Retirement

```
=M7*VLOOKUP ($L$3,$B$3:$I$12,8,FALSE)
```

Total

```
=SUM (O7:O10)
```

Net Pay

```
=M7-O11
```

The formula to retrieve the employee's name uses the VLOOKUP function. VLOOKUP takes four arguments: lookup value, lookup range, column, and match. VLOOKUP will search down

Part II: Working with Formulas and Functions

the first column of the lookup range until it finds the lookup value. Once the lookup value is found, VLOOKUP returns the value in the column identified by the column argument. In this case, the column argument is 2, and VLOOKUP returns the employee's name from the second column of the lookup range.

FIGURE 15.2

A simplified paystub form

M3	K	L	M	N	O
1					
2	Employee ID	Employee Name	Check #	Net Pay	
3	319	Elizabeth Marshall	164	\$ 796.99	
4	Employee Address				
5	530 Dodge Street Anytown, USA 12345				
6		Pay	Deductions		
7		1,390.38	Taxes	251.68	
8			Insurance	300.00	
9			Retirement	41.71	
10			Total	\$ 593.39	
11					

NOTE

All the VLOOKUP functions in this example have FALSE as the final argument. A FALSE in the match argument tells VLOOKUP to return a value only if it finds an exact match. If it doesn't find an exact match, VLOOKUP returns #N/A. Figure 15.7, later in this chapter, shows an example of using TRUE to get an approximate match.

The other formulas also use VLOOKUP with a few twists. The address and insurance formulas work just like the employee name formula, but they pull from a different column. The pay formula uses two VLOOKUPS; one is divided by the other. The employee's annual pay is pulled from the fifth column, and that is divided by the frequency from the fourth column, resulting in the pay for one paystub.

The retirement formula pulls the percentage from the eighth column and multiplies that by the gross pay to calculate the deduction. Finally, the taxes formula deducts both insurance and retirement from gross pay and multiplies that by the tax rate, found with VLOOKUP pulling from the sixth column.

Of course, payroll calculations are a little more complex than this, but once you understand how VLOOKUP works, you can build ever more complex models.

Looking up an exact value based on any lookup column

Unlike the table used in Figure 15.1, not all tables have the value that you want to look up in the leftmost column. Fortunately, Excel provides some functions for returning values that are to the right of the value you're looking up.

Figure 15.3 shows the locations of our stores by city and state. We want to return the city and store number when the user selects the state from a drop-down list.

City: =INDEX(B3:D25, MATCH(G4, C3:C25, FALSE), 1)

Store: =INDEX(B3:D25, MATCH(G4, C3:C25, FALSE), 3)

FIGURE 15.3

A list of stores with their city and state locations

G5	A	B	C	D	E	F	G
1							
		City	State	Store #			
2		Chandler	AZ	6493			
3		Glendale	CA	4369	State:	NH	
4		Fort Collins	CO	4505	City:	Manchester	
5		Gainesville	FL	8745	Store:	2608	
6		Peoria	IL	6273			
7		Indianapolis	IN	9384	LOOKUP City:	Manchester	
8		Lafayette	LA	5654	LOOKUP Store:	2608	
9		Grand Rapids	MI	3972			
10		St. Louis	MO	8816			
11		Billings	MT	3331			
12		Raleigh	NC	3335			
13		Manchester	NH	2608			
14		Albuquerque	NM	4422			

The INDEX function returns the value from a particular row and column of a range. In this case, we pass it our table of stores, a row argument in the form of a MATCH function, and a column number. For the City formula, we want the first column, so the column argument is 1. For the Store formula, we want the third column, so the column argument is 3.

Unless the range you use starts in A1, the row and column will not match the row and column in the spreadsheet. They relate to the top-left cell in the range, not the spreadsheet as a whole. A formula like =INDEX(G2:P10, 2, 2) would return the value in cell H3. The cell H3 is in the second row and the second column of the range G2:P10.

TIP

The second argument of the MATCH function can only be a range that is one row tall or one column wide. If you send it a range that's a rectangle, MATCH returns the #N/A error.

Part II: Working with Formulas and Functions

To get the correct row, we use a MATCH function. The MATCH function returns the position in the list in which the lookup value is found. It has three arguments.

Lookup value: The value we want to find.

Lookup array: The single column or single row to look in.

Match type: For exact matches only, set this argument to FALSE or 0.

The value we want to match is the state in cell G4, and we're looking for it in the range C3 : C25, our list of states. MATCH looks down the range until it finds NH. It finds it in the 12th position, so 12 is used by INDEX as the row argument.

With MATCH computed, INDEX now has all it needs to return the right value. It goes to the 12th row of the range and gets the value from either the first column (for City) or the third column (for Store #).

NOTE

If you pass INDEX a row number that is more rows than are included in the range or a column number that is more columns than are included in the range, INDEX returns the #REF! error.

Looking up values horizontally

If the data is structured in such a way that your lookup value is in the top row rather than the first column and you want to look down the rows for data rather than across the columns, Excel has a function just for you.

Figure 15.4 shows a table of cities and their temperatures. The user will select a city from a drop-down box, and you will return the temperature to the cell just below it.

=HLOOKUP (C5 , C2 : L3 , 2 , FALSE)

FIGURE 15.4

A table of cities and temperatures

C6	A	B	C	D	E	F	
1							
2	City	Chongqing	São Paulo	Kathmandu	Casablanca	Lis	
3	Temp	50	73	48	55		
4							
5	Your City:	Casablanca					
6	Your Temp:	55					
7							
8	INDEX:	55					

The HLOOKUP function has the same arguments as VLOOKUP. The *H* in HLOOKUP stands for “horizontal,” while the *V* in VLOOKUP stands for “vertical.” Instead of looking down the first column for the *lookup_value* argument, HLOOKUP looks across the first row. When it finds a match, it returns the value from the second row of the matching column.

Hiding errors returned by lookup functions

So far, we’ve used FALSE for the last argument of our lookup functions so that we return only exact matches. When we force a lookup function to return an exact match and it can’t find one, it returns the #N/A error.

The #N/A error is useful in Excel models because it alerts you when a match cannot be found. But you may be using all or a portion of your model for reporting, and #N/As are ugly. Excel has functions to see those errors and return something different.

Figure 15.5 shows a list of companies and CEOs. The other list shows CEOs and salaries. A VLOOKUP function is used to combine the two tables. But we obviously don’t have salary information for all of the CEOs, so we have a lot of #N/A errors.

=VLOOKUP (C3 , \$F\$3 : \$G\$11 , 2 , FALSE)

FIGURE 15.5

A report of CEO salaries

B	C	D	E	F	G
Company	CEO	Salary		Name	2012 Salary
Activision Blizzard Inc	Robert A. Kotick	#N/A		David M. Cote	33,247,178
CBS Corp	Leslie Moonves	62,157,026		Gregory B. Maffei	45,302,040
Cheniere Energy Inc	Charif Souki	#N/A		John H. Hammergren	51,744,999
Credit Acceptance Corp	Brett A. Roberts	#N/A		Leslie Moonves	62,157,026
Discovery Communications Inc	David M. Zaslav	#N/A		Mario J. Gabelli	68,970,486
Disney (Walt) Co	Robert A. Iger	#N/A		Marissa A. Mayer	36,615,404
Exxon Mobil Corp	R. W. Tillerson	40,266,501		Mark G. Parker	35,212,678
Gamco Investors Inc	Mario J. Gabelli	68,970,486		R. W. Tillerson	40,266,501
HCA Holdings Inc	Richard M. Bracken	#N/A		Ralph Lauren	36,325,782
Honeywell International Inc	David M. Cote	33,247,178			

In Figure 15.6, the formula has been changed to use the IFERROR function to return a blank if there’s no information available.

=IFERROR (VLOOKUP (C3 , \$F\$3 : \$G\$11 , 2 , FALSE) , " ")

The IFERROR function accepts a value or formula for its first argument and an alternate return value for its second argument. When the first argument returns an error, the second argument is returned. When the first argument is not an error, the results of the first argument are returned.

Part II: Working with Formulas and Functions

FIGURE 15.6

A cleaner report

	C	D	E	F	G
1					
2	CEO	Salary		Name	Salary
3	Robert A. Kotick		David M. Cote	33,247,178	
4	Leslie Moonves	62,157,026	Gregory B. Maffei	45,302,040	
5	Charif Souki		John H. Hammergren	51,744,999	
6	Brett A. Roberts		Leslie Moonves	62,157,026	
7	David M. Zaslav		Mario J. Gabelli	68,970,486	
8	Robert A. Iger		Marissa A. Mayer	36,615,404	
9	R. W. Tillerson	40,266,501	Mark G. Parker	35,212,678	
10	Mario J. Gabelli	68,970,486	R. W. Tillerson	40,266,501	
11	Richard M. Bracken		Ralph Lauren	36,325,782	
12	David M. Cote	33,247,178			
13	Richard B. Handler				

In this example, we've made our alternate return value an empty string (two double quotes with nothing between them). That keeps the report nice and clean. But you could return anything you want, such as "No info" or zero.

TIP

The `IFERROR` function checks for every error that Excel can return, including `#N/A`, `#DIV/0!`, and `#VALUE`. Note that you can't restrict which errors `IFERROR` catches.

Excel provides three other error trapping functions:

- `ISERROR` returns TRUE if its argument returns any error.
- `ISERR` returns TRUE if its argument returns any error except `#N/A`.
- `ISNA` returns TRUE if its argument returns `#N/A` and returns FALSE for anything else, including other errors.

All of these error-trapping functions return TRUE or FALSE and are most commonly used with an `IF` function.

Finding the closest match from a list of banded values

The `VLOOKUP`, `HLOOKUP`, and `MATCH` functions allow the data to be sorted in any order. Each of them has a final argument that will force the function to find an exact match or return an error if it cannot.

These functions also work on sorted data for the times when you want only an approximate match. Figure 15.7 shows a method for calculating income tax withholding. The withholding table doesn't have every possible value; rather, it has bands of values. You first determine into

which band the employee's pay falls, and then you use the information on that row to compute the withholding in cell D16.

```
=VLOOKUP(D15, B3 : E10, 3, TRUE) + (D15 - VLOOKUP(D15, B3 : E10, 1, TRUE))  
*VLOOKUP(D15, B3 : E10, 4, TRUE)
```

FIGURE 15.7

Computing income tax withholding

	B	C	D	E
2	Wages over	But not over	Base amount	Percentage
3	-	325	-	0.0%
4	325	1,023	-	10.0%
5	1,023	3,163	69.80	15.0%
6	3,163	6,050	390.80	25.0%
7	6,050	9,050	1,112.56	28.0%
8	9,050	15,906	195.56	33.0%
9	15,906	17,925	4,215.03	35.0%
10	17,925		4,921.68	39.6%
11				
12	Bi-weekly wage:		2,307.69	
13	Withholding allowances:		2	
14	Allowance:		303.80	
15	Wage less allowance:		2,003.89	
16	Withholding amount:		216.93	

The formula uses three VLOOKUP functions to get three pieces of data from the table. The final argument for each VLOOKUP formula is set to TRUE, indicating that we want only an approximate match.

To get a correct result when using a final argument of TRUE, the data in the lookup column (column B in Figure 15.7) must be sorted lowest to highest. VLOOKUP looks down the first column and stops when the *next* value is higher than the lookup value. In that way, it finds the largest value that is not larger than the lookup value.

CAUTION

Finding an approximate match with a lookup function does not find the closest match. Rather, it finds the largest match that's not larger than the lookup value, even if the next highest value is closer to the lookup value.

If the data in the lookup column isn't sorted lowest to highest, you may not get an error, but you will likely get an incorrect result. The lookup functions use a *binary search* to find an approximate match. A binary search basically starts in the middle of the lookup column and determines whether the match will be in the first half or the second half of the values. Then it splits that half in the middle and looks either forward or backward depending on the middle value. That process is repeated until the result is found.

You can see with a binary search that unsorted values could cause the lookup function to choose the wrong half to look in and return bad data.

Part II: Working with Formulas and Functions

In the example in Figure 15.7, VLOOKUP stops at row 5 because 1,023 is the largest value in the list that's not larger than our lookup value of 2,003.89. The three sections of the formula work as follows:

- The first VLOOKUP returns the base amount in the third column, or 69.80.
- The second VLOOKUP subtracts the “Wages over” amount (from the first column) from the total wages.
- The last VLOOKUP returns the percentage in the fourth column. This percentage is multiplied by the “excess wages,” and the result is added to the base amount.

When all three VLOOKUP functions are evaluated, the formula computes as follows:

$$=69.80 + (2,003.89 - 1,023.00) * 15.0\%$$

TIP

The method the lookup functions use to find an approximate match is much faster than an exact match. For an exact match, the function has to look at every single value in the lookup column. If you know your data will always be sorted lowest to highest and will always contain an exact match, you can decrease calculation time by setting the last argument to TRUE. An approximate match lookup will always find an exact match if it exists and the data is sorted.

Finding the closest match with the INDEX and MATCH functions

As with all of our lookup formulas, the INDEX and MATCH combination can be substituted. Like VLOOKUP and HLOOKUP, MATCH also has a final argument to find approximate matches. MATCH has the added advantage of being able to work with data that is sorted highest to lowest (see Figure 15.8).

The VLOOKUP-based formula from Figure 15.7 returns #N/A as shown in cell D16 in Figure 15.8. This is because VLOOKUP looks at the middle of the lookup column, determines that it is higher than the lookup value, and then looks only at values before the middle value. Because our data is sorted in descending order, there are no values before the middle value that are lower than the lookup value.

The INDEX and MATCH formula in cell D18 of Figure 15.8 returns the correct result and is shown here:

```
=INDEX(B3:E10,MATCH(D15,B3:B10,-1)+1,3)+  
(D15-INDEX(B3:E10,MATCH(D15,B3:B10,-1)+1,1))*  
INDEX(B3:E10,MATCH(D15,B3:B10,-1)+1,4)
```

The final argument of MATCH can be -1, 0, or 1.

- -1 is used for data that is sorted highest to lowest. It finds the smallest value in the lookup column that is larger than the lookup value. There is no equivalent method using VLOOKUP or HLOOKUP.
- 0 is used for unsorted data to find the exact match. It is equivalent to setting the final argument of VLOOKUP or HLOOKUP to FALSE.

- 1 is used for data that is sorted lowest to highest. It finds the largest value in the lookup column that is smaller than the lookup value. It is equivalent to setting the final argument of VLOOKUP or HLOOKUP to TRUE.

With a final argument of -1, MATCH finds a value that is larger than the lookup value, and the formula adds 1 to the result to get the proper row.

FIGURE 15.8

The same withholding table as Figure 15.7 except the data is sorted in descending order

	B	C	D	E
2	Wages over	But not over	Base amount	Percentage
3	17,925		4,921.68	39.6%
4	15,906	17,925	4,215.03	35.0%
5	9,050	15,906	195.56	33.0%
6	6,050	9,050	1,112.56	28.0%
7	3,163	6,050	390.80	25.0%
8	1,023	3,163	69.80	15.0%
9	325	1,023	-	10.0%
10	-	325	-	0.0%
11				
12	Bi-weekly wage:		2,307.69	
13	Withholding allowances:		2	
14	Allowance:		303.80	
15	Wage less allowance:		2,003.89	
16	Withholding amount:		#N/A	
17				
18	INDEX and MATCH:		216.93	

Looking up values from multiple tables

Sometimes the data that you want to look up can come from more than one table, depending on a choice that the user makes. In Figure 15.9, a withholding calculation similar to Figure 15.7 is shown. The difference is that the user can select whether the employee is single or married. If the user chooses Single, the data is looked up in the single-person table, and if the user chooses Married, the data is looked up in the married table.

In Excel, we can use *named ranges* and the INDIRECT function to direct our lookup to the appropriate table. Before we can write our formula, we need to name two ranges: Married for the married-person table and Single for the single-person table. Follow these steps to create the named ranges:

1. Select the range B4:E11.
2. Choose Define Name from the Formulas tab on the Ribbon. The New Name dialog box shown in Figure 15.10 is displayed.
3. Change the Name box to Married.
4. Click OK.

Part II: Working with Formulas and Functions

5. Select the range B15 : E22.
6. Choose Define Name from the Formulas tab on the Ribbon.
7. Change the Name box to Single.
8. Click OK.

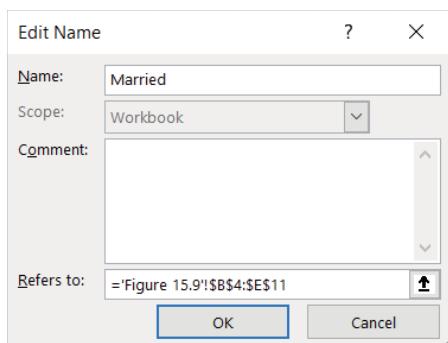
FIGURE 15.9

Computing income tax withholding from two tables

	B	C	D	E
2	<i>Married person</i>			
3	Wages over	But not over	Base amount	Percentage
4	-	325	-	0.0%
5	325	1,023	-	10.0%
6	1,023	3,163	69.80	15.0%
7	3,163	6,050	390.80	25.0%
8	6,050	9,050	1,112.56	28.0%
9	9,050	15,906	195.56	33.0%
10	15,906	17,925	4,215.03	35.0%
11	17,925		4,921.68	39.6%
12				
13	<i>Single person</i>			
14	Wages over	But not over	Base amount	Percentage
15	-	87	-	0.0%
16	87	436	-	10.0%
17	436	1,506	34.90	15.0%
18	1,506	3,523	195.40	25.0%
19	3,523	7,254	699.65	28.0%
20	7,254	15,667	1,744.33	33.0%
21	15,667	15,731	4,520.62	35.0%
22	15,731		4,543.02	39.6%
23				
24				
25	Married or Single:	Single		
26	Bi-weekly wage:	4,038.46		
27	Withholding allowances:	3		
28	Allowance:	455.70		
29	Wage less allowance:	3,582.76		
30	Withholding amount:	716.38		

FIGURE 15.10

The New Name dialog box



There is a data validation drop-down in cell D25 in Figure 15.9. The drop-down box contains the terms *Married* and *Single*, which are identical to the names that we just created. We'll be using the value in D25 to determine in which table we'll look, so the values must be identical.

The revised formula for computing the withholding follows:

```
=VLOOKUP(D29, INDIRECT(D25), 3, TRUE) +
(D29-VLOOKUP(D29, INDIRECT(D25), 1, TRUE))
*VLOOKUP(D29, INDIRECT(D25), 4, TRUE)
```

The formula in this example is strikingly similar to the one shown in Figure 15.7. The only difference is that an `INDIRECT` function is used in place of the table's location.

`INDIRECT` takes an argument named `ref_text`. The `ref_text` argument is a text representation of a cell reference or a named range. In Figure 15.9, cell D25 contains the text *Single*. `INDIRECT` attempts to convert that into a cell or range reference. If `ref_text` is not a valid range reference, as in our case, `INDIRECT` checks the named ranges to see whether there is a match. Had we not already created a range named *Single*, `INDIRECT` would return the `#REF!` error.

`INDIRECT` has a second optional argument named `a1`. The `a1` argument is `TRUE` if `ref_text` is in the A1 style of cell references and `FALSE` if `ref_text` is in the R1C1 style of cell references. For named ranges, `a1` can be either `TRUE` or `FALSE`, and `INDIRECT` will return the correct range.

CAUTION

`INDIRECT` can also return ranges from other worksheets or even other workbooks. However, if it references another workbook, that workbook must be open. `INDIRECT` does not work on closed workbooks.

Looking up a value based on a two-way matrix

A two-way matrix is a rectangular range of cells. That is, it's a range with more than one row and more than one column. In other formulas, we've used the `INDEX` and `MATCH` combination as an alternative to some of the lookup functions. But `INDEX` and `MATCH` were made for two-way matrices.

Figure 15.11 shows a table of sales figures by region and year. Each row represents a region, and each column represents a year. We want the user to select a region and a year and return the sales figure at the intersection of that row and column.

```
=INDEX(C4:F9, MATCH(C13, B4:B9, FALSE), MATCH(C14, C3:F3, FALSE))
```

By now you're no doubt familiar with `INDEX` and `MATCH`. Unlike with other formulas, we're using two `MATCH` functions within the `INDEX` function. The second `MATCH` function returns the column argument of `INDEX` as opposed to hard-coding a column number.

Part II: Working with Formulas and Functions

FIGURE 15.11

Sales data by region and year

The screenshot shows a Microsoft Excel spreadsheet titled "Regional Sales Report". The table has columns for Region, 2017, 2018, 2019, 2020, and Total. Row 10 is highlighted in green and contains the formula =INDEX(C4:F9,MATCH(C13,B4:B9,0),MATCH(C14,C3:F3,0)). Rows 11 and 12 are empty. Row 13 contains "Region:" followed by a dropdown menu showing "North". Row 14 contains "Year:" followed by a dropdown menu showing "2018". Row 15 contains "Sales:" followed by the value 4,441,886, which is also highlighted in green.

B	C	D	E	F	G	
2	Regional Sales Report					
3	Region	2017	2018	2019	2020	Total
4	South	1,525,017	1,504,678	1,227,847	1,019,616	5,277,158
5	Northeast	2,704,237	2,135,564	1,411,782	716,535	6,968,118
6	North	3,563,687	4,441,886	4,805,431	3,716,674	16,527,678
7	West	4,489,700	2,651,064	796,330	2,898,601	10,835,695
8	Mid-Atlantic	2,167,319	1,357,850	776,850	3,024,542	7,326,561
9	Canada	1,861,239	3,578,280	4,069,389	1,475,301	10,984,209
10	Total	16,311,199	15,669,322	13,087,629	12,851,269	57,919,419
11						
12						
13	Region:	North				
14	Year:	2018				
15	Sales:	4,441,886				

Recall that MATCH returns the position in a list of the matched value. In Figure 15.11, the North region is matched, so MATCH returns 3 because it's the third item in the list. That becomes the row argument for INDEX. The year is matched across the header row, and since 2018 is the second item, MATCH returns 2. INDEX then takes the 3 and 2 returned by the MATCH functions to return the proper value.

Using default values for match

Let's add a twist to our sales lookup formula. We'll change the formula to allow the user to select only a region, to select only a year, or to select neither. If one of the selections is omitted, we'll assume that the user wants the total. If neither is selected, we'll return the total for the whole table.

```
=INDEX (C4 : G10 ,  
IFERROR (MATCH (C13 , B4 : B10 , FALSE) , COUNTA (B4 : B10)) ,  
IFERROR (MATCH (C14 , C3 : G3 , FALSE) , COUNTA (C3 : G3)))
```

The overall structure of the formula is the same, but we've changed a few details. The range we're using for INDEX now includes row 10 and column G. Each MATCH function's range is also extended. Finally, both MATCH functions are surrounded by an IFERROR function that will return the Total row or column.

The alternate value for IFERROR is a COUNTA function. COUNTA counts both numbers and text and, in effect, returns the position of the last row or column in our range. We could have hard-coded those values, but if we happen to insert a row or column, COUNTA will adjust to always return the last one.

Figure 15.12 shows the same sales table, but the user has left the Year input blank. Since there are no blanks in the column headers, MATCH returns #N/A. When it encounters that error, IFERROR passes control to the value_if_error argument, and the last column is passed to INDEX.

FIGURE 15.12

Returning totals from the sales data

The screenshot shows a Microsoft Excel spreadsheet. The formula bar at the top contains the formula: =INDEX(C4:G10,IFERROR(MATCH(C13,B4:B10,FALSE),COUNTA(B4:B10)),IFERROR(MATCH(C14,C3:G3,FALSE),COUNTA(C3:G3)))

The table below is titled "Regional Sales Report". It has columns for Region (Region), years (2017, 2018, 2019, 2020), and a Total column. Row 10 is labeled "Total". Row 13 is a filter row with "Region:" and "Mid-Atlantic" selected. Row 14 is a filter row with "Year:" and "2018" selected. Row 15 is a result row with "Sales:" and "1,357,850" displayed.

B	C	D	E	F	G
Regional Sales Report					
3	Region	2017	2018	2019	2020
4	South	1,525,017	1,504,678	1,227,847	1,019,616
5	Northeast	2,704,237	2,135,564	1,411,782	716,535
6	North	3,563,687	4,441,886	4,805,431	3,716,674
7	West	4,489,700	2,651,064	796,330	2,898,601
8	Mid-Atlantic	2,167,319	1,357,850	776,850	3,024,542
9	Canada	1,861,239	3,578,280	4,069,389	1,475,301
10	Total	16,311,199	15,669,322	13,087,629	12,851,269
11					
12					
13	Region:	Mid-Atlantic			
14	Year:				
15	Sales:	1,357,850			

Finding a value based on multiple criteria

Figure 15.13 shows a table of departmental budgets. When the user selects a region and department, we want a formula to return the budget. We can't use VLOOKUP for this formula because it accepts only one lookup value. You need two values because the regions and departments appear multiple times.

You can use the SUMPRODUCT function to get the row that contains both lookup values as follows:

```
=SUMPRODUCT( ($B$3:$B$45=H5) * ($C$3:$C$45=H6) * ($E$3:$E$45) )
```

SUMPRODUCT compares every cell in a range with a value and returns an array of TRUEs and FALSEs depending on the result. When multiplied with another array, TRUE becomes 1 and FALSE becomes 0. The third parenthetical section in the SUMPRODUCT function does not contain a comparison because that range contains the value we want to return.

If either the Region comparison or the Department comparison is FALSE, the total for that line will be 0. A FALSE result is converted to zero, and anything times zero is zero. If both Region and Department match, both comparisons return 1. The two 1s are multiplied with the corresponding row in column E, and that's the value returned.

Part II: Working with Formulas and Functions

FIGURE 15.13

A table of departmental budgets

	A	B	C	D	E	F	G	H
1								
2	Region	Department	Manager	Budget				
3	Mid-Atlantic	Legal	Jose Palmer	4,406,018				
4	Canada	Legal	Audrey Washington	2,564,165				
5	North	Logistics	Samantha Allen	1,443,535		Region:	South	
6	Mid-Atlantic	Customer Service	Katherine Nichols	2,834,014		Department:	Accounting	
7	Pacific Northwest	Quality Assurance	David Gordon	1,119,596		Budget:	697,697	
8	Mid-Atlantic	Operational	Lillian Hart	2,949,401		Manager:	Brooke Bailey	
9	Mid-Atlantic	Sales	Claire Peterson	2,371,246				
10	East	Inventory	Hannah Porter	3,043,499				
11	Mid-Atlantic	Services	Victoria Gomez	1,621,716				
12	South	Accounting	Brooke Bailey	697,697				
13								

In the example shown in Figure 15.13, when SUMPRODUCT gets to row 12, it multiplies $1 * 1 * 697,697$. That number is summed with the other rows, all of which are zero because they contain at least one FALSE. The resulting SUM is the value 697,697.

Returning text with SUMPRODUCT

SUMPRODUCT works this way only when we want to return a number. If we want to return text, all of the text values would be treated as zero, and SUMPRODUCT would always return zero.

However, we can pair SUMPRODUCT with the INDEX and ROW functions to return text. If we want to return the manager's name, for example, we could use the following formula:

```
=INDEX(D:D, SUMPRODUCT((($B$3:$B$45=H5)*($C$3:$C$45=H6)*  
*(ROW($E$3:$E$45))), 1))
```

Instead of including the values from column E, the ROW function is used to include the row numbers in the array. SUMPRODUCT now computes $1 * 1 * 12$ when it gets to row 12. The 12 is then used for the row argument in INDEX against the entire column D:D. Because the ROW function returns the row in the worksheet and not the row in our table, INDEX uses the whole column as its range.

Finding the last value in a column

Figure 15.14 shows an unsorted list of invoices. We want to find the last invoice in the list. A simple way to find the last item in the column is to use the INDEX function and count the items in the list to determine the last row.

```
=INDEX(B:B, COUNTA(B:B)+1)
```

FIGURE 15.14

A list of invoices

	B	C	D	E	F	G
2	Invoice #	Item Count	Amount			
3	IN6787	53	555.73			
4	IN4374	160	940.56			
5	IN5061	40	3,026.10		Last Invoice	Last Amount
6	IN4305	146	4,885.94		IN6513	3,326.98
7	IN1477	84	969.46			
8	IN5552	97	2,979.33			
9	IN8685	200	2,950.74			
10	IN1491	40	3,970.50			
11	IN2408	155	3,332.94			
12	IN6513	75	3,326.98			

The INDEX function when used on a single column needs only a row argument. The third argument indicates the column isn't necessary. COUNTA is used to count the nonblank cells in column B. That count is increased by 1 because we have a blank cell in the first row. The INDEX function returns the 12th row of column B.

CAUTION

COUNTA counts numbers, text, dates, and anything except blanks. If there are blank rows in your data, COUNTA will not return the desired result.

Finding the last number using LOOKUP

INDEX and COUNTA are great for finding values when there are no blank cells in the range. If you have blanks and the values you're searching for are numbers, you can use LOOKUP and a really large number. The formula in cell G5 of Figure 15.14 uses this technique.

=LOOKUP(9.99E+307,D:D)

The lookup value is the largest number that Excel can handle (just under 1 with 308 zeros behind it). Since LOOKUP won't find a value that large, it stops at the last value it does find, and that's the value returned.

TIP

A number like 9.99E+307 is written in *exponential notation*. The number before the *E* has one number to the left of the decimal and two to the right. The number after the *E* is how many places are required to move the decimal point to show the number in regular notation (307 in this case). A positive number means to move the decimal to the right, and a negative number means to move it to the left. A number like 4.32E-02 is equivalent to 0.0432.

This LOOKUP method has the additional advantage of returning the last number, even if there are text, blanks, or errors in the range.

Using Formulas with Tables and Conditional Formatting

IN THIS CHAPTER

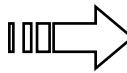
Highlighting cells that meet certain criteria

Highlighting differences between data sets

Conditional formatting based on dates

Conditional formatting is the term given to the functionality of Excel dynamically changing the formatting of a value, cell, or range of cells based on a set of conditions that you define. Conditional formatting allows you to look at your Excel reports and make split-second determinations on which values are “good” and which are “bad,” all based on formatting.

In this chapter, we’ll give you a few examples of how the conditional formatting feature in Excel can be used in conjunction with formulas to add an extra layer of visualizations to your analyses.



This book’s website, www.wiley.com/go/excel365bible, includes a copy of the sample workbook for this chapter. The file is named Using Functions with Conditional Formatting.xlsx.

NOTE

Feel free to revisit Chapter 5, “Formatting Worksheets,” for a refresher on conditional formatting.

Highlighting Cells That Meet Certain Criteria

One of the more basic conditional formatting rules that you can create is the highlighting of cells that meet some business criteria. This first example demonstrates the formatting of cells with values that are lower than a hard-coded value of 4000 (see Figure 16.1).

Part II: Working with Formulas and Functions

FIGURE 16.1

The cells in this table are conditionally formatted to show a red background for values less than 4000.

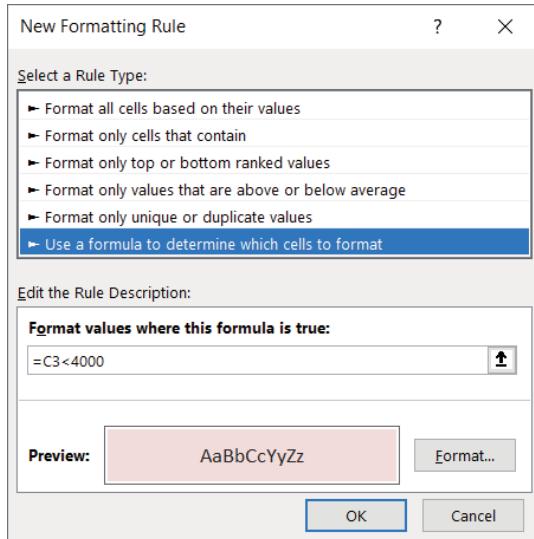
	A	B	C
1			
2			
3			
4			
5			
6			
7			
8			
9			
10			
11			
12			
13			
14			

To build this basic formatting rule, follow these steps:

1. Select the data cells in your target range (cells C3 : C14 in this example).
2. Click the Home tab of the Excel Ribbon and then select Conditional Formatting ↴ New Rule. This will open the New Formatting Rule dialog box, shown in Figure 16.2.

FIGURE 16.2

Configure the New Formatting Rule dialog box to apply the needed formula rule.



3. In the list box at the top of the dialog box, click the option called **Use a formula to determine which cells to format**. This selection evaluates values based on a formula you specify. If a particular value evaluates to TRUE, then the conditional formatting is applied to that cell.
4. In the **formula input box**, enter the formula shown here. Note that we are simply referencing the first cell in our target range. There is no need to reference the entire range.

=C3<4000

CAUTION

Note in the formula that we exclude the absolute reference dollar symbols (\$) for the target cell (C3). If you click cell C3 with your mouse instead of typing it in, Excel will automatically make your cell reference absolute. It's important that you don't include the absolute reference dollar symbols in your target cell, as you need Excel to apply this formatting rule based on each cell's own value.

5. Click the **Format button** and choose your desired formatting. This will open the Format Cells dialog box where you'll have a full set of options for formatting the font, border, and fill for your target cell.
6. Click the **OK** button once you've completed choosing your formatting options.
7. Click the **OK** button again to confirm your formatting rule in the New Formatting Rule dialog box.

TIP

If you need to edit your conditional formatting rule, simply select any of the data cells within your formatted range and then go to the Home tab and select Conditional Formatting \Rightarrow Manage Rules. This will open the Conditional Formatting Rules Manager dialog box. Click the rule that you want to edit and then click the Edit Rule button.

Highlighting cells based on the value of another cell

In many cases, the formatting rule for your cells will be based on how they compare to the value of another cell. Take the example illustrated in Figure 16.3. Here the cells are conditionally highlighted if their respective values fall below the Prior Year Average value shown in cell B3.

To build this basic formatting rule, follow these steps:

1. Select the data cells in your target range (cells E3 : E14 in this example).
2. Click the Home tab of the Excel Ribbon and then select Conditional Formatting \Rightarrow New Rule. This will open the New Formatting Rule dialog box shown in Figure 16.4.

Part II: Working with Formulas and Functions

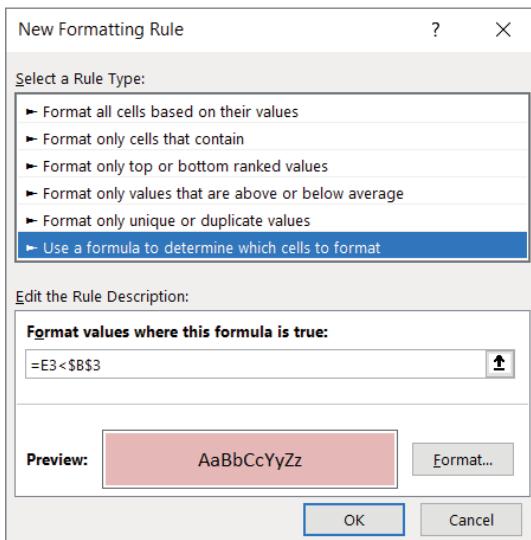
FIGURE 16.3

The cells in this table are conditionally formatted to show a red background for values falling below the Prior Year Average value.

	A	B	C	D	E
1					
2		Prior Year Average		Month	Units Sold
3		3500	January		2661
4			February		3804
5			March		5021
6			April		1001
7			May		4375
8			June		2859
9			July		7659
10			August		3061
11			September		2003
12			October		5147
13			November		4045
14			December		1701

FIGURE 16.4

Compare our target cell (E3) with the value in the comparison cell (\$B\$3).



3. In the list box at the top of the dialog box, click the option called **Use a formula to determine which cells to format**. This selection evaluates values based on a formula you specify. If a particular value evaluates to TRUE, then the conditional formatting is applied to that cell.

4. In the formula input box, enter the formula shown here. Note that we are simply comparing our target cell (E3) with the value in the comparison cell (\$B\$3). As with standard formulas, you'll need to ensure that you use absolute references so that each value in your range is compared to the appropriate comparison cell.
=E3<\$B\$3
5. Click the Format button and choose your desired formatting. This will open the Format Cells dialog box, where you'll have a full set of options for formatting the font, border, and fill for your target cell.
6. Click the OK button once you've completed choosing your formatting options.
7. Click the OK button again to confirm your formatting rule in the New Formatting Rule dialog box.

Highlighting Values That Exist in List1 but Not List2

You'll often be asked to compare two lists and pick out the values that are in one list but not the other. Conditional formatting is an ideal way to present your findings. Figure 16.5 illustrates a conditional formatting exercise that compares customers from 2020 and 2021, highlighting those customers in 2021 that are new customers, that is, those customers who did not exist in 2020.

FIGURE 16.5

You can conditionally format the values that exist in one list but not the other.

A	B	C	D	E	F
1					
2	2020			2021	
3	Customer_Name	Revenue		Customer_Name	Revenue
4	GKNEAS Corp.	\$2,333.60		JAMSEA Corp.	\$2,324.36
5	JAMSEA Corp.	\$2,324.36		JAMWUS Corp.	\$2,328.53
6	JAMWUS Corp.	\$2,328.53		JAYKA Corp.	\$2,328.53
7	JAYKA Corp.	\$2,328.53		JUSDAN Corp.	\$3,801.86
8	MAKUTE Corp.	\$2,334.01		MAKUTE Corp.	\$2,334.01
9	MOSUNC Corp.	\$2,311.70		MALEBO Corp.	\$3,099.45
10	NCUANT Corp.	\$2,311.79		MOSUNC Corp.	\$2,311.70
11	OSADUL Corp.	\$2,311.50		NCUANT Corp.	\$2,311.79
12	RRCAR Corp.	\$2,315.14		OSADUL Corp.	\$2,311.50
13	RULLAN Corp.	\$2,332.94		PUNSKE Corp.	\$7,220.80
14	SMATHE Corp.	\$2,336.59		REBUST Corp.	\$14,224.84
15	SOFANU Corp.	\$2,333.60		RRCAR Corp.	\$2,315.14
16	SUMTUK Corp.	\$2,321.61		RULLAN Corp.	\$2,332.94
17	TULUSS Corp.	\$2,311.96		RUTANS Corp.	\$4,175.75
18	UDGUWU Corp.	\$2,328.58		SCHOUL Corp.	\$5,931.46

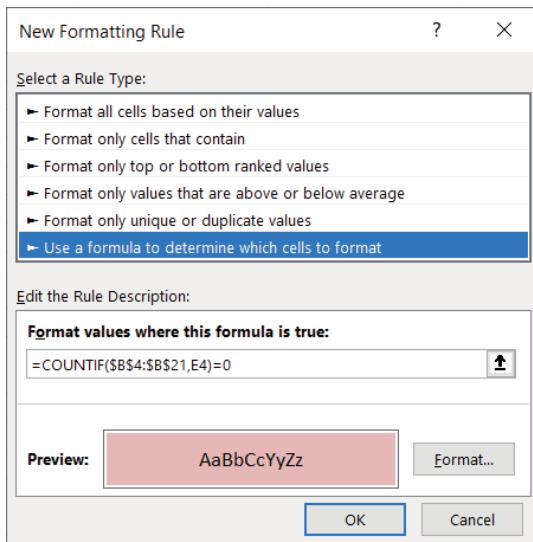
Part II: Working with Formulas and Functions

To build this basic formatting rule, follow these steps:

1. **Select the data cells in your target range (cells E4 : E28 in this example).**
2. **Click the Home tab of the Excel Ribbon and then select Conditional Formatting ↗ New Rule.** This will open the New Formatting Rule dialog box shown in Figure 16.6.

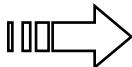
FIGURE 16.6

Apply the conditional format if there are zero instances of the value in the target cell (E4) found in our comparison range (\$B\$4 : \$B\$21).



3. **In the list box at the top of the dialog box, click the option called Use a formula to determine which cells to format.** This selection evaluates values based on a formula you specify. If a particular value evaluates to TRUE, then the conditional formatting is applied to that cell.
4. **In the formula input box, enter the formula shown here.** Note we're using the COUNTIF function to evaluate whether the value in the target cell (E4) is found in our comparison range (\$B\$4 : \$B\$21). If the value is not found, the COUNTIF function will return a 0, thus triggering the conditional formatting. As with standard formulas, you'll need to ensure that you use absolute references so that each value in your range is compared to the appropriate comparison cell.
 $=COUNTIF($B$4:$B$21,E4)=0$
5. **Click the Format button and choose your desired formatting.** This will open the Format Cells dialog box, where you'll have a full set of options for formatting the font, border, and fill for your target cell.

6. Click the OK button once you've completed choosing your formatting options.
7. Click the OK button again to confirm your formatting rule in the New Formatting Rule dialog box.



For more detail on the COUNTIF function, see Chapter 14, “Using Formulas for Conditional Analysis.”

Highlighting Values That Exist in List1 and List2

Sometimes, you'll need to compare two lists and pick out only the values that exist in *both* lists. Again, conditional formatting is an ideal way to present your findings. Figure 16.7 illustrates a conditional formatting exercise that compares customers from 2020 and 2021, highlighting those customers in 2021 who are in both lists.

FIGURE 16.7

You can conditionally format the values that exist in both lists.

	A	B	C	D	E	F
1						
2		2020			2021	
3		Customer_Name	Revenue		Customer_Name	Revenue
4	GKNEAS Corp.	\$2,333.60		JAMSEA Corp.	\$2,324.36	
5	JAMSEA Corp.	\$2,324.36		JAMWUS Corp.	\$2,328.53	
6	JAMWUS Corp.	\$2,328.53		JAYKA Corp.	\$2,328.53	
7	JAYKA Corp.	\$2,328.53		JUSDAN Corp.	\$3,801.86	
8	MAKUTE Corp.	\$2,334.01		MAKUTE Corp.	\$2,334.01	
9	MOSUNC Corp.	\$2,311.70		MALEBO Corp.	\$3,099.45	
10	NCUANT Corp.	\$2,311.79		MOSUNC Corp.	\$2,311.70	
11	OSADUL Corp.	\$2,311.50		NCUANT Corp.	\$2,311.79	
12	RRCAR Corp.	\$2,315.14		OSADUL Corp.	\$2,311.50	
13	RULLAN Corp.	\$2,332.94		PUNSKE Corp.	\$7,220.80	
14	SMATHE Corp.	\$2,336.59		REBUST Corp.	\$14,224.84	
15	SOFANU Corp.	\$2,333.60		RRCAR Corp.	\$2,315.14	
16	SUMTUK Corp.	\$2,321.61		RULLAN Corp.	\$2,332.94	

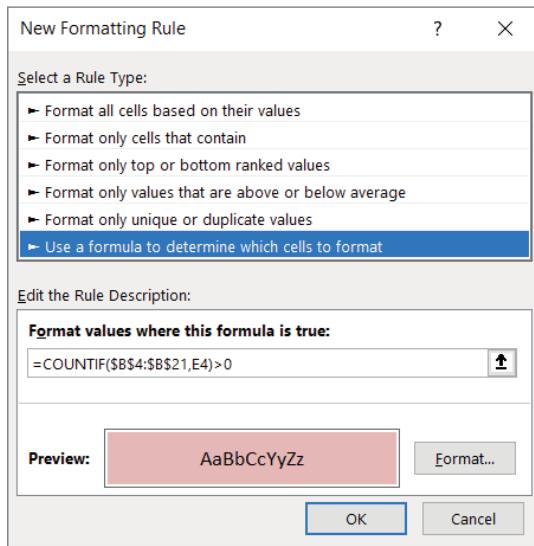
To build this basic formatting rule, follow these steps:

1. Select the data cells in your target range (cells E4 : E28 in this example).
2. Click the Home tab of the Excel Ribbon and then select Conditional Formatting ↗ New Rule. This will open the New Formatting Rule dialog box shown in Figure 16.8.
3. In the list box at the top of the dialog box, click the option called Use a formula to determine which cells to format. This selection evaluates values based on a formula you specify. If a particular value evaluates to TRUE, then the conditional formatting is applied to that cell.

Part II: Working with Formulas and Functions

FIGURE 16.8

Apply the conditional format if there is at least one instance (>0) of the value in the target cell (E4) found in our comparison range (\$B\$4 : \$B\$21).



4. In the formula input box, enter the formula shown here. Note we're using the COUNTIF function to evaluate whether the value in the target cell (E4) is found in our comparison range (\$B\$4 : \$B\$21). If the value is found, the COUNTIF function will return a number greater than 0, thus triggering the conditional formatting. As with standard formulas, you'll need to ensure that you use absolute references so that each value in your range is compared to the appropriate comparison cell.
`=COUNTIF (B4 : B21 , E4) >0`
5. Click the Format button and choose your desired formatting. This will open the Format Cells dialog box, where you'll have a full set of options for formatting the font, border, and fill for your target cell.
6. Click the OK button once you've completed choosing your formatting options.
7. Click the OK button again to confirm your formatting rule in the New Formatting Rule dialog box.

Highlighting Based on Dates

You may find it useful to indicate visually when certain dates trigger a certain scenario. For instance, when working with timecards and scheduling, it is often beneficial to be able to easily pinpoint any dates that fall on weekends. The conditional formatting rule illustrated in Figure 16.9 highlights all of the weekend dates in the list of values.

Chapter 16: Using Formulas with Tables and Conditional Formatting

FIGURE 16.9

You can conditionally format any weekend dates in a list of dates.

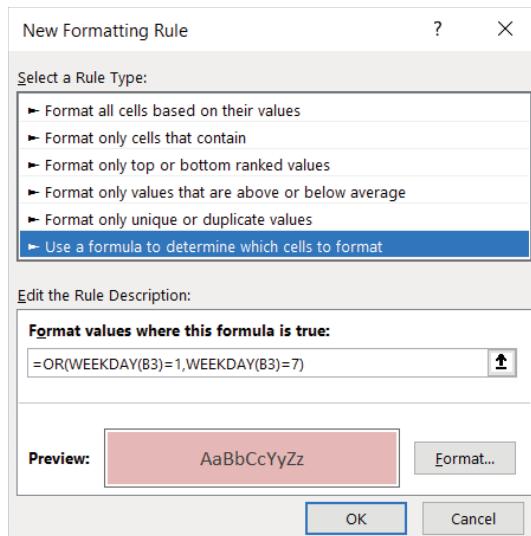
A	B
1	
2	Highlight Weekends
3	1/23/2020
4	12/28/2020
5	9/26/2021
6	12/8/2020
7	4/25/2021
8	11/7/2021
9	7/31/2020
10	11/24/2020
11	12/28/2019
12	7/28/2020
13	12/17/2020
14	8/3/2021
15	5/1/2021
16	4/2/2020
17	7/18/2021
18	8/12/2021

To build this basic formatting rule, follow these steps:

1. Select the data cells in your target range (cells B3:B18 in this example).
2. Click the Home tab of the Excel Ribbon and then select Conditional Formatting \Rightarrow New Rule. This will open the New Formatting Rule dialog box, shown in Figure 16.10.

FIGURE 16.10

Using the WEEKDAY function to evaluate the weekday number of the target cell (B3)



3. In the list box at the top of the dialog box, click the option called **Use a formula to determine which cells to format**. This selection evaluates values based on a formula you specify. If a particular value evaluates to TRUE, then the conditional formatting is applied to that cell.
4. In the formula input box, enter the formula shown here. Note we're using the **WEEKDAY** function to evaluate the weekday number of the target cell (B3). If the target cell returns as weekday 1 or 7, it means that the date in B3 is a weekend date. In this case, the conditional formatting will be applied.
`=OR(WEEKDAY(B3)=1, WEEKDAY(B3)=7)`
5. Click the **Format button** and choose your desired **formatting**. This will open the Format Cells dialog box, where you'll have a full set of options for formatting the font, border, and fill for your target cell.
6. Click the **OK button** once you've completed choosing your **formatting options**.
7. Click the **OK button** again to confirm your **formatting rule** in the **New Formatting Rule dialog box**.

Highlighting days between two dates

Some analysis requires the identification of dates that fall within a certain time period. Figure 16.11 demonstrates how you can apply conditional formatting that highlights dates based on a start date and end date. As the start and end dates are adjusted, the conditional formatting adjusts with them.

FIGURE 16.11

You can conditionally format dates that fall between a start and end date.

B	C	D	E
Start	End	Highlight Days within 2020 and 2021	
1/1/2020	12/31/2021	1/23/2020	
		12/28/2021	
		9/26/2019	
		12/8/2018	
		4/25/2020	
		11/7/2019	
		7/31/2021	
		11/24/2021	
		12/28/2018	
		7/28/2019	
		12/17/2020	
		8/3/2019	
		5/1/2021	
		4/2/2020	
		7/17/2019	
		8/12/2018	

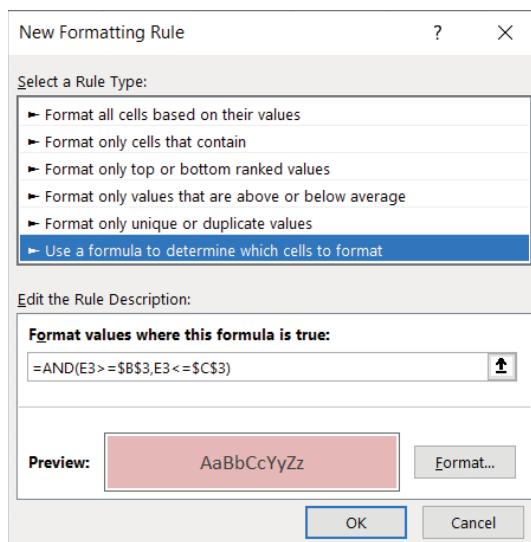
Chapter 16: Using Formulas with Tables and Conditional Formatting

To build this basic formatting rule, follow these steps:

1. Select the data cells in your target range (cells E3 : E18 in this example), click the Home tab of the Excel Ribbon, and then select Conditional Formatting \Rightarrow New Rule. This will open the New Formatting Rule dialog box shown in Figure 16.12.

FIGURE 16.12

Using the AND function to compare the date in our target cell (E3) to both the start and end dates found in cells \$B\$3 and \$C\$3



2. In the list box at the top of the dialog box, click the option called Use a formula to determine which cells to format. This selection evaluates values based on a formula you specify. If a particular value evaluates to TRUE, then the conditional formatting is applied to that cell.
3. In the formula input box, enter the formula shown here. Note we're using the AND function to compare the date in our target cell (E3) to both the start and end dates found in cells \$B\$3 and \$C\$3, respectively. If the target cell falls within the start and end dates, the formula will evaluate to TRUE, thus triggering the conditional formatting.
 $=AND(E3>=$B$3, E3<=$C$3)$
4. Click the Format button and choose your desired formatting. This will open the Format Cells dialog box, where you'll have a full set of options for formatting the font, border, and fill for your target cell.
5. Click the OK button once you've completed choosing your formatting options.
6. Click the OK button again to confirm your formatting rule in the New Formatting Rule dialog box.

Highlighting dates based on a due date

The example shown in Figure 16.13 demonstrates that you can conditionally format dates that are past due by a given number of days. In this scenario, the dates that are over 90 days overdue are highlighted with a red background.

FIGURE 16.13

You can conditionally format dates based on due date.

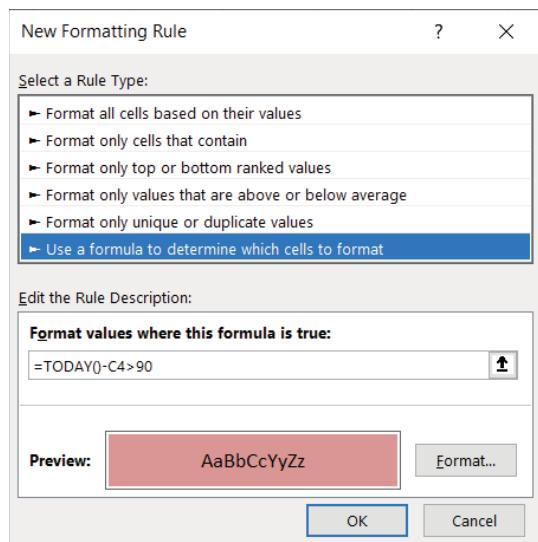
B	C
	Due Date
	04/25/20
	05/04/21
	05/04/21
	03/28/20
	04/22/21
	03/31/21

To build this basic formatting rule, follow these steps:

1. Select the data cells in your target range (cells C4 : C9 in this example), click the Home tab of the Excel Ribbon, and then select Conditional Formatting \Rightarrow New Rule. This will open the New Formatting Rule dialog box, shown in Figure 16.14.

FIGURE 16.14

Evaluating whether today's date is greater than 90 days past the date in our target cell (C4)



2. In the list box at the top of the dialog box, click the option called **Use a formula to determine which cells to format**. This selection evaluates values based on a formula you specify. If a particular value evaluates to TRUE, then the conditional formatting is applied to that cell.
3. In the formula input box, enter the formula shown here. In this formula, we're evaluating whether today's date is greater than 90 days past the date in our target cell (C4). If so, the conditional formatting will be applied.
`=TODAY() - C4 > 90`
4. Click the **Format button** and choose your desired formatting. This will open the Format Cells dialog box, where you'll have a full set of options for formatting the font, border, and fill for your target cell.
5. Click the **OK** button once you've completed choosing your formatting options.
6. Click the **OK** button again to confirm your formatting rule in the New Formatting Rule dialog box.

Making Your Formulas Error-Free

IN THIS CHAPTER

Identifying and correcting common formula errors

Using Excel auditing tools

Using formula AutoCorrect

Tracing cell relationships

Checking spelling and related features

It goes without saying that you want your Excel worksheets to produce accurate results. Unfortunately, it's not always easy to be certain that the results are correct, especially if you deal with large, complex worksheets. This chapter introduces the tools and techniques available to help identify, correct, and prevent errors.

Finding and Correcting Formula Errors

Making a change in a worksheet—even a relatively minor change—may produce a ripple effect that introduces errors in other cells. For example, accidentally entering a value into a cell that previously held a formula is all too easy to do. This simple error can have a major impact on other formulas, and you may not discover the problem until long after you make the change—if you discover the problem at all.

Formula errors tend to fall into one of the following general categories:

Syntax errors You have a problem with the syntax of a formula. For example, a formula may have mismatched parentheses, or a function may not have the correct number of arguments.

Logical errors A formula doesn't return an error, but it contains a logical flaw that causes it to return an incorrect result.

Incorrect reference errors The logic of the formula is correct, but the formula uses an incorrect cell reference. As a simple example, the range reference in a SUM formula may not include all of the data that you want to sum.

Semantic errors An example is a function name that is spelled incorrectly. Excel will attempt to interpret it as a name and will display the #NAME? error.

Circular references A circular reference occurs when a formula refers to its own cell, either directly or indirectly. Circular references are useful in a few cases, but most of the time, a circular reference indicates a problem.

Dynamic array #SPILL! error Excel displays #SPILL! errors when a dynamic array formula returns results that cannot be written to a specified return range. #SPILL! can typically be quickly resolved by ensuring the return range is blank, is not overlapping a Table object, and does not include merged cells.

Incomplete calculation errors The formulas simply aren't calculated fully. To ensure that your formulas are fully calculated, press Ctrl+Alt+Shift+F9.



Refer to Chapter 10, “Understanding and Using Array Formulas,” for an introduction to dynamic array formulas.

Syntax errors are usually the easiest to identify and correct. In most cases, you'll know when your formula contains a syntax error. For example, Excel won't permit you to enter a formula with mismatched parentheses. Other syntax errors also usually result in an error display in the cell.

The following sections describe common formula problems and offer advice on identifying and correcting them.

Mismatched parentheses

In a formula, every left parenthesis must have a corresponding right parenthesis. If your formula has mismatched parentheses, Excel usually won't permit you to enter it. An exception to this rule involves a simple formula that uses a function. For example, if you enter the following formula (which is missing a closing parenthesis), Excel accepts the formula and provides the missing parenthesis:

```
=SUM(A1:A500
```

A formula may have an equal number of left and right parentheses, but the parentheses may not match properly. For example, consider the following formula, which converts a text string such that the first character is uppercase and the remaining characters are lowercase. This formula has five pairs of parentheses, and they match properly:

```
=UPPER(LEFT(A1)) & RIGHT(LOWER(A1), LEN(A1) - 1)
```

The following formula also has five pairs of parentheses, but they're mismatched. The result displays a syntactically correct formula that simply returns the wrong result:

```
=UPPER(LEFT(A1) & RIGHT(LOWER(A1), LEN(A1) - 1))
```

Often, parentheses that are in the wrong location will result in a *syntax error*, which is usually a message that tells you that you entered too many or too few arguments for a function.

TIP

Excel can help you with mismatched parentheses. When you're editing a formula and you move the cursor over a parenthesis, Excel displays it (and its matching parenthesis) in bold for about one-half second. In addition, Excel color-codes pairs of nested parentheses while you're editing a formula.

Using Formula AutoCorrect

When you enter a formula that has a syntax error, Excel attempts to determine the problem and offers a suggested correction.

Be careful when accepting corrections for your formulas from Excel because it doesn't always guess correctly. For example, imagine that you entered the following formula (which has mismatched parentheses):

```
=AVERAGE (SUM (A1 :A12 , SUM (B1 :B12) )
```

Excel then proposes the following correction to the formula:

```
=AVERAGE (SUM (A1 :A12 , SUM (B1 :B12) ))
```

You may be tempted to accept the suggestion without even thinking. In this case, the proposed formula is syntactically correct, but it's not what you intended. The correct formula is as follows:

```
=AVERAGE (SUM (A1 :A12) , SUM (B1 :B12) )
```

Cells are filled with hash marks

A cell is filled with a series of hash marks (#) for one of two reasons:

- **The column is not wide enough to accommodate the formatted numeric value.** To correct it, you can make the column wider or use a different number format. (See Chapter 21, "Visualizing with Custom Number Formats and Shapes.")
- **The cell contains a formula that returns an invalid date or time.** For example, Excel doesn't support dates prior to 1900 or the use of negative time values. A formula that returns either of these values results in a cell filled with hash marks. Widening the column won't fix it.

Blank cells are not blank

Some Excel users have discovered that by pressing the spacebar, the contents of a cell seem to erase. Actually, pressing the spacebar inserts an invisible space character, which isn't the same as erasing the cell.

For example, the following formula returns the number of nonempty cells in range A1 :A10. If you "erase" any of these cells by using the spacebar, these cells are included in the count, and the formula returns an incorrect result:

```
=COUNTA (A1 :A10)
```

Part II: Working with Formulas and Functions

If your formula doesn't ignore blank cells the way that it should, check to make sure that the blank cells are really blank. Here's how to search for cells that contain only blank characters:

1. **Press Ctrl+F.** The Find and Replace dialog box appears.
2. **Click the Options button to expand the dialog box so that it displays additional options.**
3. **In the Find What box, enter * *. That's an asterisk, followed by a space, followed by another asterisk.**
4. **Make sure that the Match Entire Cell Contents check box is selected.**
5. **Click Find All.** If any cells that contain only space characters are found, Excel lists the cell address at the bottom of the Find and Replace dialog box.

Extra space characters

If you have formulas or use procedures that rely on comparing text, be careful that your text doesn't contain additional space characters. Adding an extra space character is particularly common when data has been imported from another source.

Excel automatically removes trailing spaces from values that you enter, but trailing spaces in text entries are not deleted. It's impossible to tell just by looking at a cell whether it contains one or more trailing space characters.

You can leverage the TRIM function to identify values that contain leading spaces, trailing spaces, and multiple spaces within a text string. For instance, this formula will return FALSE if the text in cell A1 contains leading spaces, trailing spaces, or multiple spaces.

=TRIM(A1)=A1

Formulas returning an error

A formula may return any of the following error values:

- #DIV/0!
- #N/A
- #NAME?
- #NULL!
- #NUM!
- #REF!
- #SPILL!
- #VALUE!

The following sections summarize possible problems that may cause these errors.

TIP

Excel allows you to choose the way error values are printed. To access this feature, display the Page Setup dialog box and select the Sheet tab. You can choose to print error values as displayed (the default) or as blank cells, dashes, or #N/A. To display the Page Setup dialog box, click the dialog box launcher of the Page Layout \leftrightarrow Page Setup group.

Tracing Error Values

Often, an error in one cell is the result of an error in a precedent cell. For help in identifying the cell causing an error value to appear, activate the cell that contains the error and then choose Formulas ➔ Formula Auditing ➔ Error Checking ➔ Trace Error. Excel draws arrows to indicate which cell is the source of the error.

After you identify the error, choose Formulas ➔ Formula Auditing ➔ Remove Arrows to get rid of the arrow display.

#DIV/0! errors

Division by zero is not a valid operation. If you create a formula that attempts to divide by zero, Excel displays its familiar #DIV/0! error value.

Because Excel considers a blank cell to be zero, you also get this error if your formula divides by a missing value.

To avoid the error display, you can use an IF function to check for a blank cell. For instance, this formula will display an empty string if cell B4 is blank or contains 0; otherwise, it displays the calculated value:

```
=IF(B4=0, "", C4/B4)
```

Another approach is to use an IFERROR function to check for *any* error condition. The following formula, for example, displays an empty string if the formula results in any type of error:

```
=IFERROR(C4/B4, "")
```

#N/A errors

The #N/A error occurs if any cell referenced by a formula displays #N/A.

NOTE

Some users like to use =NA() or #N/A explicitly for missing data. This method makes it perfectly clear that the data is not available and hasn't been deleted accidentally.

The #N/A error also occurs when a LOOKUP function (HLOOKUP, LOOKUP, MATCH, or VLOOKUP) can't find a match.

If you would like to display an empty string instead of #N/A, use the IFNA function in a formula like this:

```
=IFNA(VLOOKUP(A1, C1:F50, 4, FALSE), "")
```

#NAME? errors

The #NAME? error occurs under these conditions:

- The formula contains an undefined range or cell name.
- The formula contains text that Excel *interprets* as an undefined name. A misspelled function name, for example, generates a #NAME? error.
- The formula contains text that isn't enclosed in quotation marks.
- The formula contains a range reference that omits the colon between the cell addresses.
- The formula uses a worksheet function that's defined in an add-in, and the add-in is not installed.

CAUTION

Excel has a bit of a problem with range names. If you delete a name for a cell or a range and the name is used in a formula, the formula continues to use the name, even though it's no longer defined. As a result, the formula displays #NAME?. You might expect Excel automatically to convert the names to their corresponding cell references, but this doesn't happen.

#NULL! errors

A #NULL! error occurs when a formula attempts to use an intersection of two ranges that don't actually intersect. Excel's intersection operator is a space. The following formula, for example, returns #NULL! because the two ranges don't intersect:

```
=SUM(B5:B14 A16:F16)
```

The following formula doesn't return #NULL! but displays the contents of cell B9, which represents the intersection of the two ranges:

```
=SUM(B5:B14 A9:F9)
```

You also see a #NULL! error if you accidentally omit an operator in a formula. For example, this formula is missing the second operator:

```
= A1+A2 A3
```

#NUM! errors

A formula returns a #NUM! error if any of the following occurs:

- You pass a nonnumeric argument to a function when a numeric argument is expected (for example, \$1,000 instead of 1000).
- You pass an invalid argument to a function, such as attempting to calculate the square root of a negative number. This formula returns #NUM!:

```
=SQRT(-12)
```

- A function that uses iteration can't calculate a result. Examples of functions that use iteration are IRR and RATE.
- A formula returns a value that is too large or too small. Excel supports values between -1E-307 and 1E+307.

#REF! errors

A #REF! error occurs when a formula uses an invalid cell reference. This error can occur in the following situations:

- You delete the row or column of a cell that is referenced by the formula. For example, the following formula displays a #REF! error if row 1, column A, or column B is deleted:
`=A1/B1`
- You delete the worksheet of a cell that is referenced by the formula. For example, the following formula displays a #REF! error if Sheet2 is deleted:
`=Sheet2!A1`
- You copy a formula to a location that invalidates the relative cell references. For example, if you copy the following formula from cell A2 to cell A1, the formula returns #REF! because it attempts to refer to a nonexistent cell:
`=A1-1`
- You cut a cell (choose Home \Rightarrow Clipboard \Rightarrow Cut) and then paste it to a cell that's referenced by a formula. The formula will display #REF!.

#SPILL! errors

A #SPILL! error occurs when entering a dynamic array formula where the specified return range (the cells that display the result of the dynamic array) is not blank. Common reasons for a #SPILL! are as follows:

- The return range already contains data.
- The return range overlaps a Table or PivotTable object.
- The return range includes merged cells.
- The return range extends beyond the edge of the worksheet.
- The dynamic array formula uses a volatile function that changes the return range. By default, changes in an array while calculating trigger additional calculation passes. Excel performs these passes to ensure the spreadsheet is fully calculated. Volatile functions such as RANDARRAY, RAND and RANDBETWEEN can trigger changes in the array size, preventing Excel from resolving the formula.

#VALUE! errors

A #VALUE! error is common and can occur under the following conditions:

- An argument for a function is of an incorrect data type, or the formula attempts to perform an operation using incorrect data. For example, a formula that adds a value to a text string returns the #VALUE! error.
- A function's argument is a range when it should be a single value.
- A custom worksheet function is not calculated. You can press Ctrl+Alt+F9 to force a recalculation.
- A custom worksheet function attempts to perform an operation that is not valid. For example, custom functions can't modify the Excel environment or make changes to other cells.

Pay Attention to the Colors

When you edit a cell that contains a formula, Excel color-codes the cell and range references in the formula. Excel also outlines the cells and ranges used in the formula by using corresponding colors. Therefore, you can see at a glance the cells that are used in the formula.

You also can manipulate the colored outline to change the cell or range reference. To change the references used in a formula, drag the outline's border or fill handle (at the lower right of the outline). This technique is often easier than editing the formula.

Operator precedence problems

As noted in Chapter 9, Excel has some straightforward rules about the order in which mathematical operations are performed. When in doubt (or when you simply need to clarify your intentions), you should use parentheses to ensure that operations are performed in the correct order. For example, the following formula multiplies A1 by A2 and then adds 1 to the result. The multiplication is performed first because it has a higher order of precedence:

=1+A1*A2

The following is a clearer version of this formula. The parentheses aren't necessary, but in this case, the order of operations is perfectly obvious:

=1+(A1*A2)

Notice that the negation operator symbol is the same as the subtraction operator symbol. This, as you may expect, can cause some confusion. Consider these two formulas:

=-3^2
=0-3^2

The first formula, as expected, returns 9. The second formula, however, returns -9. Squaring a number always produces a positive result, so how is it that Excel can return the -9 result?

In the first formula, the minus sign is a *negation* operator and has the highest precedence. However, in the second formula, the minus sign is a *subtraction* operator, which has a lower precedence than the exponentiation operator. Therefore, the value 3 is squared, and then the result is subtracted from 0 (zero), which produces a negative result.

Using parentheses, as shown in the following formula, causes Excel to interpret the operator as a minus sign rather than a negation operator. This formula returns -9:

= - (3 ^ 2)

Formulas are not calculated

If you use custom worksheet functions written in VBA, you may find that the formulas that use these functions fail to be recalculated and may display incorrect results. For example, assume that you wrote a VBA function that returns the number format of a referenced cell. If you change the number format, the function will continue to display the previous number format. That's because changing a number format doesn't trigger a recalculation.

To force a single formula to be recalculated, select the cell, press F2, and then press Enter. To force a recalculation of all formulas, press Ctrl+Alt+F9.

Problems with decimal precision

By their nature, computers don't have infinite precision. Excel stores numbers in binary format by using 8 bytes, which can handle numbers with 15-digit accuracy. Some numbers can't be expressed precisely by using 8 bytes, so the number is stored as an approximation.

To demonstrate how this lack of precision may cause problems, enter the following formula into cell A1:

= (5 . 1 - 5 . 2) + 1

The result should be 0 . 9. However, if you format the cell to display 15 decimal places, you discover that Excel calculates the formula with a result of 0 . 899999999999999. This result occurs because the operation in parentheses is performed first, and this intermediate result stores in binary format by using an approximation. The formula then adds 1 to this value, and the approximation error is propagated to the final result.

In many cases, this type of error doesn't present a problem. However, if you need to test the result of that formula by using a logical operator, it *may* present a problem. For example, the following formula (which assumes that the previous formula is in cell A1) returns FALSE:

=A1=.9

One solution to this type of error is to use the ROUND function. The following formula, for example, returns TRUE because the comparison is made by using the value in A1 rounded to one decimal place:

=ROUND(A1,1)=0.9

Part II: Working with Formulas and Functions

Here's another example of a "precision" problem. Try entering the following formula:

```
= (1.333 - 1.233) - (1.334 - 1.234)
```

This formula should return 0, but it actually returns $-2.22045E-16$ (a number very close to zero).

If that formula is in cell A1, the following formula returns Not Zero:

```
=IF(A1=0, "Zero", "Not Zero")
```

One way to handle these "very close to zero" rounding errors is to use a formula like this:

```
=IF(ABS(A1)<1E-6, "Zero", "Not Zero")
```

This formula uses the less-than operator ($<$) to compare the absolute value of the number with a very small number. This formula returns Zero.

"Phantom link" errors

You may open a workbook and see a message asking whether you want to update links in a workbook. This message sometimes appears even when a workbook contains no linked formulas. Often, these phantom links are created when you copy a worksheet that contains names.

First, try choosing File \Rightarrow Info \Rightarrow Edit Links to Files to display the Edit Links dialog box. Then select each link and click Break Link. If that doesn't solve the problem, this phantom link may be caused by an erroneous name. Choose Formulas \Rightarrow Defined Names \Rightarrow Name Manager and scroll through the list of names in the Name Manager dialog box. If you see a name that refers to #REF!, delete the name. The Name Manager dialog box has a Filter button that lets you filter the names. For example, you can filter the lists to display only the names with errors.

Using Excel Auditing Tools

Excel includes several tools that can help you track down formula errors. The following sections describe the auditing tools built into Excel.

Identifying cells of a particular type

The Go To Special dialog box is a handy tool that enables you to locate cells of a particular type. To display this dialog box, choose Home \Rightarrow Editing \Rightarrow Find & Select \Rightarrow Go To Special.

NOTE

If you select a multicell range before displaying the Go To Special dialog box, the command operates only within the selected cells. If a single cell is selected, the command operates on the entire worksheet.

You can use the Go To Special dialog box to select cells of a certain type, which can often help you identify errors. For example, if you choose the Formulas option, Excel selects all of the

cells that contain a formula. If you zoom the worksheet out to a small size, you can get a good idea of the worksheet's organization.

To zoom a worksheet, use the zoom controls on the right side of the status bar or press Ctrl while you move the scroll wheel on your mouse.

TIP

Selecting the formula cells may also help you spot a common error, namely, a formula that has been replaced accidentally with a value. If you find a cell that's not selected amid a group of selected formula cells, chances are good that the cell previously contained a formula that has been replaced by a value.

Viewing formulas

You can become familiar with an unfamiliar workbook by displaying the formulas rather than the results of the formulas. To toggle the display of formulas, choose Formulas \Rightarrow Formula Auditing \Rightarrow Show Formulas.

You may want to create a second window for the workbook before issuing this command. This way, you can see the formulas in one window and the results of the formula in the other window. Choose View \Rightarrow Window \Rightarrow New Window to open a new window.

TIP

You can also use the keyboard to toggle between Formula view and Normal view by pressing Ctrl + the grave accent key (`).



See Chapter 4, "Working with Excel Ranges and Tables," for more information about this command.

Tracing cell relationships

To understand how to trace cell relationships, you need to familiarize yourself with the following two concepts:

Cell precedents Applicable only to cells that contain a formula, a formula cell's precedents are all of the cells that contribute to the formula's result. A *direct precedent* is a cell that you use directly in the formula. An *indirect precedent* is a cell that isn't used directly in the formula but is used by a cell that you refer to in the formula.

Cell dependents These formula cells depend upon a particular cell. A cell's dependents consist of all formula cells that use the cell. Again, the formula cell can be a direct dependent or an indirect dependent.

For example, consider this simple formula entered into cell A4:

=SUM(A1 : A3)

Cell A4 has three precedent cells (A1, A2, and A3), which are all direct precedents. Cells A1, A2, and A3 all have at least one dependent cell (cell A4).

Identifying cell precedents for a formula cell often sheds light on why the formula isn't working correctly. Conversely, knowing which formula cells depend on a particular cell is also helpful. For example, if you're about to delete a formula, you may want to check whether it has any dependents.

Identifying precedents

You can identify cells used by a formula in the active cell in a number of ways:

- **Press F2.** The cells that are used directly by the formula are outlined in color, and the color corresponds to the cell reference in the formula. This technique is limited to identifying cells on the same sheet as the formula.
- **Choose Home ▷ Editing ▷ Find & Select ▷ Go To Special to display the Go To Special dialog box.** Select the Precedents option and then select either Direct Only (for direct precedents only) or All Levels (for direct and indirect precedents). Click OK and Excel selects the precedent cells for the formula. This technique is limited to identifying cells on the same sheet as the formula.
- **Press Ctrl+].** This selects all direct precedent cells on the active sheet.
- **Press Ctrl+Shift+{.** This selects all precedent cells (direct and indirect) on the active sheet.
- **Choose Formulas ▷ Formula Auditing ▷ Trace Precedents.** Excel will draw arrows to indicate the cell's precedents. Click this button multiple times to see additional levels of precedents. Choose Formulas ▷ Formula Auditing ▷ Remove Arrows to hide the arrows.

Identifying dependents

You can identify formula cells that use a particular cell in a number of ways:

- **Choose Home ▷ Editing ▷ Find & Select ▷ Go To Special to display the Go To Special dialog box.** Select the Dependents option and then select either Direct Only (for direct dependents only) or All Levels (for direct and indirect dependents). Click OK. Excel selects the cells that depend on the active cell. This technique is limited to identifying cells on the active sheet only.
- **Press Ctrl+].** This selects all direct dependent cells on the active sheet.
- **Press Ctrl+Shift+}.** This selects all dependent cells (direct and indirect) on the active sheet.
- **Choose Formulas ▷ Formula Auditing ▷ Trace Dependents.** Excel will draw arrows to indicate the cell's dependents. Click this button multiple times to see additional levels of dependents. Choose Formulas ▷ Formula Auditing ▷ Remove Arrows to hide the arrows.

Tracing error values

If a formula displays an error value, Excel can help you identify the cell that is causing that error value. An error in one cell is often the result of an error in a precedent cell. Activate a cell that contains an error value and then choose **Formulas** ➔ **Formula Auditing** ➔ **Error Checking** ➔ **Trace Error**. Excel draws arrows to indicate the error source. Choose **Formulas** ➔ **Formula Auditing** ➔ **Remove Arrows** to hide the arrows.

Fixing circular reference errors

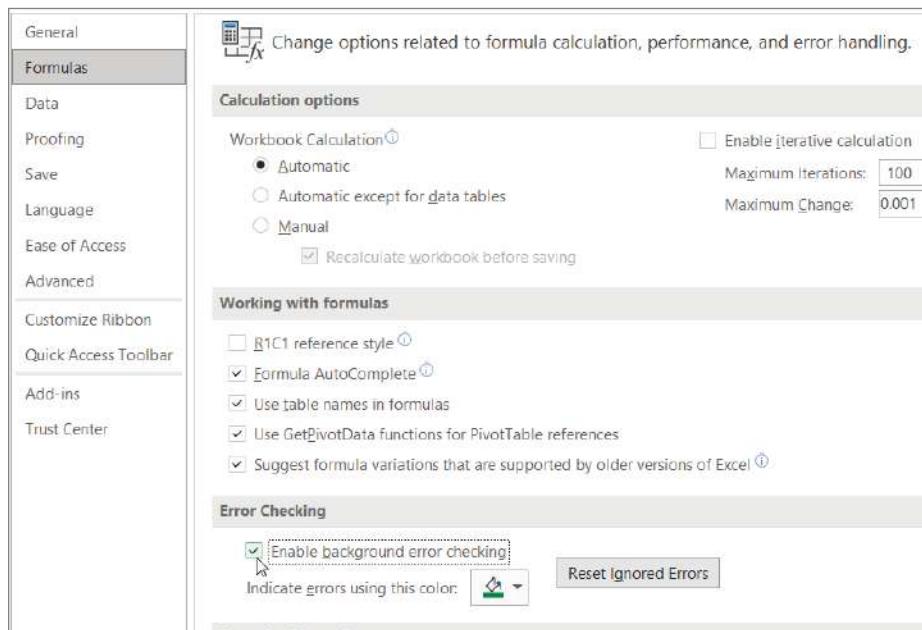
If you accidentally create a circular reference formula, Excel displays a warning message—**Circular Reference**—with the cell address in the status bar. It also draws arrows on the worksheet to help you identify the problem. If you can't figure out the source of the problem, choose **Formulas** ➔ **Formula Auditing** ➔ **Error Checking** ➔ **Circular References**. This command displays a list of all cells that are involved in the circular references. Start by selecting the first cell listed and then work your way down the list until you figure out the problem.

Using the background error-checking feature

Some people may find it helpful to take advantage of the Excel automatic error-checking feature. This feature is enabled or disabled via the **Enable Background Error Checking** check box, found on the **Formulas** tab of the Excel Options dialog box, shown in Figure 17.1. In addition, you can use the check boxes in the **Error Checking Rules** section to specify which types of errors to check.

FIGURE 17.1

Excel can check your formulas for potential errors.



Part II: Working with Formulas and Functions

When error checking is turned on, Excel continually evaluates the formulas in your worksheet. If a potential error is identified, Excel places a small green triangle in the upper-left corner of the cell. When the cell is activated, a drop-down control appears. Clicking this drop-down control provides you with options. The options vary, depending on the type of error.

In many cases, you'll choose to ignore an error by selecting the Ignore Error option. Selecting this option eliminates the cell from subsequent error checks. However, all previously ignored errors can be reset so that they appear again. (Use the Reset Ignored Errors button on the Formulas tab of the Excel Options dialog box.)

You can choose Formulas \Rightarrow Formula Auditing \Rightarrow Error Checking to display a dialog box that describes each potential error cell in sequence, much like using a spell-checking command.

CAUTION

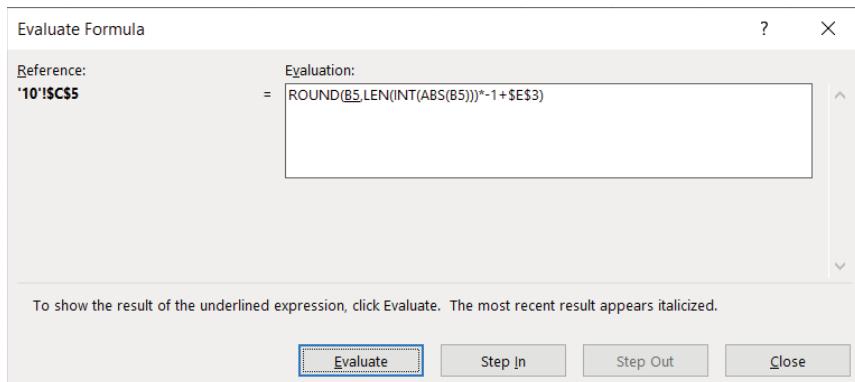
The error-checking feature isn't perfect. In fact, it's not even close to perfect. In other words, you can't assume that you have an error-free worksheet simply because Excel doesn't identify potential errors!

Using Formula Evaluator

Formula Evaluator lets you see the various parts of a nested formula evaluated in the order in which the formula is calculated. To use Formula Evaluator, select the cell that contains the formula and then choose Formulas \Rightarrow Formula Auditing \Rightarrow Evaluate Formula to display the Evaluate Formula dialog box (see Figure 17.2).

FIGURE 17.2

The Evaluate Formula dialog box shows a formula being calculated one step at a time.



Click the Evaluate button to show the result of calculating the expressions within the formula. Each click of the button performs another calculation. This feature may seem a bit complicated

at first, but if you spend some time working with it, you'll understand how it works and see the value.

Excel provides another way to evaluate a part of a formula:

1. **Select the cell that contains the formula.**
2. **Press F2 to get into Cell Edit mode.**
3. **Use your mouse to highlight the portion of the formula you want to evaluate, or press Shift and use the navigation keys.**
4. **Press F9.**

The highlighted portion of the formula displays the calculated result. You can evaluate other parts of the formula or press Esc to cancel and return your formula to its previous state.

CAUTION

Be careful when using this technique because if you press Enter (rather than Esc), the formula will be modified to use the calculated values.

Searching and Replacing

Excel has a powerful search-and-replace feature that makes it easy to locate information in a worksheet or across multiple worksheets in a workbook. As an option, you can also search for text and replace it with other text.

To access the Find and Replace dialog box, start by selecting the range that you want to search. If you select any single cell, Excel searches the entire sheet. Choose Home \Rightarrow Editing \Rightarrow Find & Select \Rightarrow Find (or press Ctrl+F).

If you're simply looking for information in the worksheet, select the Find tab. If you want to replace existing text with new text, use the Replace tab. Also note that you can use the Options button to display (or hide) additional options. The dialog box shown in Figure 17.3 displays these additional options.

Searching for information

Enter the information to search for in the Find What text box and then specify any of the following options:

- Within drop-down list** Specify where to search (the current sheet or the entire workbook).
- Search drop-down list** Specify the direction (by rows or by columns).
- Look In drop-down list** Specify what cell parts to search (formulas, values, notes, or comments).

Part II: Working with Formulas and Functions

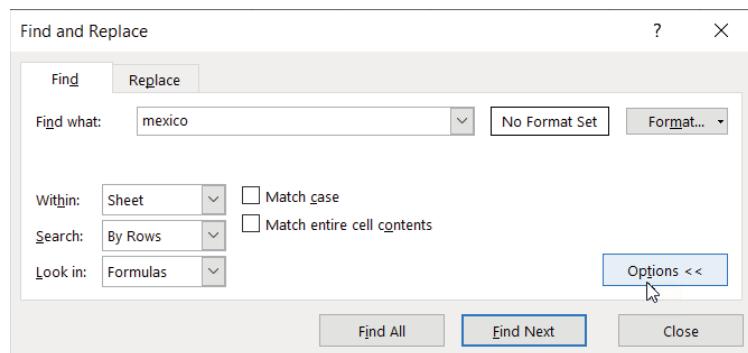
Match Case check box Specify whether the search should be case sensitive.

Match Entire Cell Contents check box Specify whether the entire cell contents must be matched.

Format button Click to search for cells that have a particular formatting. (See the upcoming section “Searching for formatting.”)

FIGURE 17.3

Use the Find and Replace dialog box to locate information in a worksheet or workbook.



Click Find Next to locate the matching cells one at a time or click Find All to locate all matches. If you use the Find All button, the Find and Replace dialog box expands to display the addresses of all matching cells in a list. When you select an entry in this list, Excel scrolls the worksheet so that you can view it in context.

TIP

After using Find All, press Ctrl+A to select all the found cells in the worksheet.

NOTE

The Find and Replace dialog box allows you to access the worksheet and make changes without the need to dismiss the dialog box.

Replacing information

To replace text with other text, use the Replace tab in the Find and Replace dialog box. Enter the text to be replaced in the Find What field and then enter the new text in the Replace With field. Specify other options as described in the previous section.

Click Find Next to locate the first matching item and then click Replace to do the replacement. When you click the Replace button, Excel locates the next matching item. To override the replacement, click Find Next. To replace all items without verification, click Replace All. If the replacement didn't occur as you planned, you can use the Undo button on the Quick Access Toolbar (or press Ctrl+Z).

TIP

To delete information, enter the text to be deleted in the Find What field, but leave the Replace With field empty.

Searching for formatting

From the Find and Replace dialog box, you can also locate cells that contain a particular type of formatting. As an option, you can replace that formatting with another type of formatting. For example, assume that you want to locate all cells that are formatted as bold and then change that formatting to bold and italic. Follow these steps:

1. **Choose Home ➔ Editing ➔ Find & Select ➔ Replace (or press Ctrl+H).** The Find and Replace dialog box appears.
2. **Make sure that the Replace tab is displayed.** If necessary, click the Options button to expand the dialog box.
3. **If the Find What and Replace With fields are not empty, delete their contents.**
4. **Click the top Format button.** The Find Format dialog box appears. This dialog box resembles the standard Format Cells dialog box.
5. **Select the Font tab.**
6. **Select Bold in the Font Style list and then click OK.**
7. **Click the bottom Format button.** The Replace Format dialog box appears.
8. **Select the Font tab.**
9. **Select Bold Italic from the Font Style list and then click OK.**
10. **In the Find and Replace dialog box, click Replace All.** Excel locates all cells that have bold formatting and changes the formatting to bold italic.

You can also find formatting based on a particular cell. In the Find Format dialog box, click the Choose Format from Cell button and then click the cell that contains the formatting you're looking for.

CAUTION

The Find and Replace dialog box can't find background color formatting in tables that was applied using table styles or formatting that is applied based on conditional formatting.

Spell-checking your worksheets

If you use a word-processing program, you probably take advantage of its spell-checker feature. Spelling mistakes can be just as embarrassing when they appear in a spreadsheet. Fortunately, Microsoft includes a spell-checker with Excel.

To access the spell-checker, choose Review \Rightarrow Proofing \Rightarrow Spelling, or press F7. To check the spelling in just a particular range, select the range before you activate the spell-checker.

NOTE

The spell-checker checks cell contents, text in graphics objects and charts, and page headers and footers. Even the contents of hidden rows and columns are checked.

The Spelling dialog box works similarly to other spell-checkers with which you may be familiar. If Excel encounters a word that isn't in the current dictionary or that is misspelled, it offers a list of suggestions. You can respond by clicking one of these buttons:

Ignore Once Ignore the word and continue the spell-check.

Ignore All Ignore the word and all subsequent occurrences of it.

Add to Dictionary Add the word to the dictionary.

Change Change the word to the selected word in the Suggestions list.

Change All Change the word to the selected word in the Suggestions list and change all subsequent occurrences of it without asking.

AutoCorrect Add the misspelled word and its correct spelling (which you select from the list) to the AutoCorrect list.

Using AutoCorrect

AutoCorrect is a handy feature that automatically corrects common typing mistakes. You can also add words to the list that Excel corrects automatically. The AutoCorrect dialog box appears in Figure 17.4. To access this feature, choose File \Rightarrow Options. In the Excel Options dialog box, select the Proofing tab and then click the AutoCorrect Options button.

This dialog box has several options:

Correct TWO INital CApitals Automatically corrects words with two initial uppercase letters. For example, BUdget is converted to Budget. This mistake is common among fast typists. You can click the Exceptions button to specify a list of exceptions to this rule.

Capitalize First Letter of Sentences Capitalizes the first letter in a sentence. All other letters are unchanged.

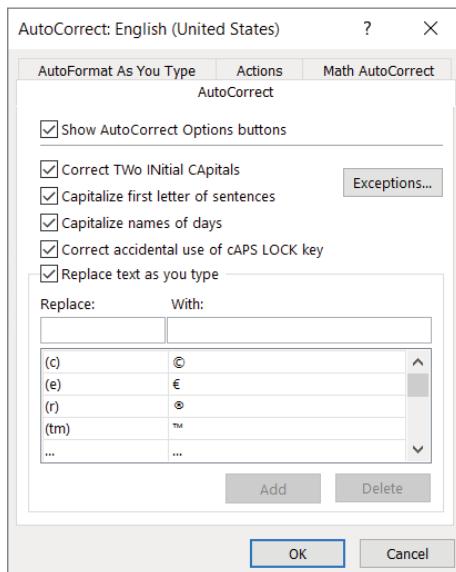
Capitalize Names of Days Capitalizes the days of the week. If you enter monday, Excel converts it to Monday.

Correct Accidental Use of cAPS LOCK Key Corrects errors caused if you accidentally pressed the Caps Lock key while typing.

Replace Text as You Type AutoCorrect automatically changes incorrect words as you type them.

FIGURE 17.4

Use the AutoCorrect dialog box to control the spelling corrections Excel makes automatically.



Excel includes a long list of AutoCorrect entries for commonly misspelled words. In addition, it has AutoCorrect entries for some symbols. For example, (c) is replaced with®, and (r) is replaced with®. You can also add your own AutoCorrect entries. For example, if you find that you frequently misspell the word *January* as *Janriuary*, you can create an AutoCorrect entry so that it's changed automatically. To create a new AutoCorrect entry, enter the misspelled word in the Replace box and the correctly spelled word in the With field. You can also delete entries that you no longer need.

Part II: Working with Formulas and Functions

TIP

You can use the AutoCorrect feature to create shortcuts for commonly used words or phrases. For example, if you work for a company named Consolidated Data Processing Corporation, you can create an AutoCorrect entry for an abbreviation, such as *cdp*. Then, whenever you type *cdp*, Excel automatically changes it to Consolidated Data Processing Corporation. Just make sure that you don't use a combination of characters that might normally appear in your text and be replaced erroneously.

NOTE

In some cases, you may want to override the AutoCorrect feature. For example, you may literally need to enter (c) rather than a copyright symbol. Just click the Undo button on the Quick Access Toolbar or press **Ctrl+Z**.

You can use the AutoFormat As You Type tab of the AutoCorrect dialog box to control a few other automatic settings in Excel.

The Actions tab enables what were formerly known as *Smart Tags* for certain types of data in your worksheets. The types of actions that Excel recognizes vary depending on the types of software that are installed on your system.

Part III

Creating Charts and Other Visualizations

The five chapters in this part deal with charts and visualizations. You'll discover how to use Excel's graphics capabilities to display your data in a chart or as Sparkline graphics. In addition, you'll learn how to use Excel's other drawing and graphics tools to enhance your worksheets with meaningful visualizations of your data.

IN THIS PART

- Chapter 18**
Getting Started with Excel Charts
- Chapter 19**
Using Advanced Charting Techniques
- Chapter 20**
Creating Sparkline Graphics
- Chapter 21**
Visualizing with Custom Number Formats and Shapes

Getting Started with Excel Charts

IN THIS CHAPTER

- How Excel handles charts
- The parts of a chart
- The basic steps for creating a chart
- Working with charts
- Looking at examples of chart types

Charts offer a visual representation of numeric values; they are at-a-glance views that allow you to specify relationships between data values, point out differences, and observe business trends. Few mechanisms allow you to absorb data faster than a chart, which can be a key component in your dashboard.

When most people think of a spreadsheet product such as Excel, they think of crunching rows and columns of numbers. But Excel also has a plethora of tools to present data visually in the form of a chart. In this chapter, we present an overview of Excel's charting capabilities and show you how to create and customize your own charts using Excel.



The sample workbook for this chapter (`Intro_to_Charts.xlsx`) can be found at www.wiley.com/go/excel365bible.

What Is a Chart?

A *chart* is a visual representation of numeric values. Charts (also known as *graphs*) have been an integral part of spreadsheets since the early days of Lotus 1-2-3. Charts generated by early spreadsheet products were extremely crude by today's standards. Over the years, however, the quality and flexibility have improved significantly. You'll find that Excel provides you with the tools to create a wide variety of highly customizable charts that can help you effectively communicate your message.

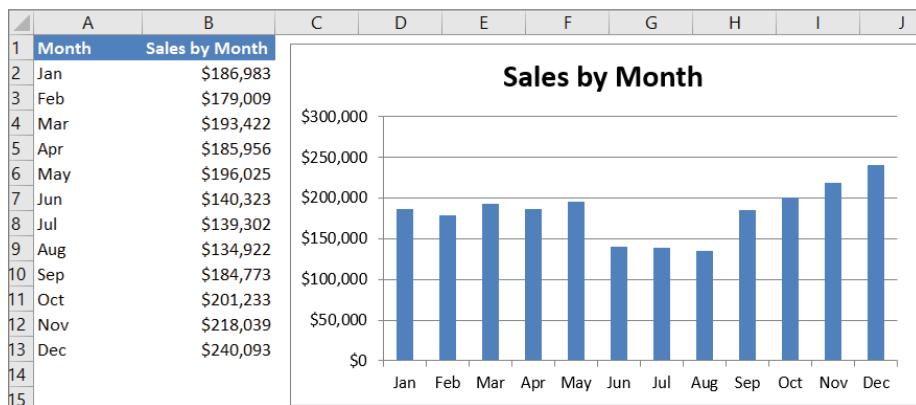
Displaying data in a well-conceived chart can make your numbers more understandable. Because a chart presents a picture, they are particularly useful for summarizing a series of numbers and their interrelationships. Making a chart can often help you spot trends and patterns that might otherwise go unnoticed.

Part III: Creating Charts and Other Visualizations

Figure 18.1 shows a worksheet that contains a simple column chart that depicts a company's sales volume by month. Viewing the chart makes it apparent that sales were off in the summer months (June through August), but they increased steadily during the final four months of the year. You could, of course, arrive at this same conclusion simply by studying the numbers. But viewing the chart makes the point much more quickly. A column chart is just one of many different types of charts that you can create with Excel.

FIGURE 18.1

A simple column chart depicts the sales volume for each month.



TIP

Charts can reside in either of two locations in a workbook: on a worksheet (these charts are known as embedded charts) or on a separate chart sheet.

To create an embedded chart quickly, you can select your data and then press Alt+F1. Alternatively, you can press F11 to create a chart on its own chart sheet.

How Excel handles charts

Before you can create a chart, you must have some numbers (data). The data, of course, is stored in the cells in a worksheet. Normally, the data that is used by a chart resides in a single worksheet, but that's not a strict requirement. A chart can use data that's stored in any number of worksheets, and the worksheets can even be in different workbooks. The decision to use data from one sheet or multiple sheets really depends on your data model, the nature of your data sources, and the interactivity that you want to give your dashboard.

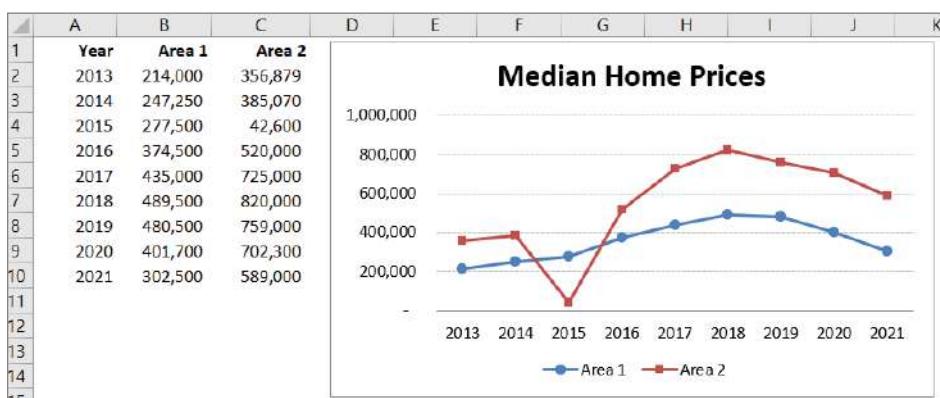
A chart is essentially an “object” that Excel creates upon request. This object consists of one or more *data series*, displayed graphically. The appearance of the data series depends on the selected *chart type*. For example, if you create a line chart that uses two data series, the chart contains two lines, and each line represents one data series. The data for each series is stored

in a separate row or column. Each point on the line is determined by the value in a single cell and is represented by a marker. You can distinguish the lines by their thickness, line style, color, or data markers.

Figure 18.2 shows a line chart that plots two data series across a nine-year period. The series are identified by using different data markers (squares versus circles), shown in the *legend* at the bottom of the chart. The lines also use different colors, which is not apparent in the gray-scale figure.

FIGURE 18.2

This line chart displays two data series.



A key point to keep in mind is that charts are dynamic. In other words, a chart series is linked to the data in your worksheet. If the data changes, the chart is updated automatically to reflect those changes so that your dashboard can show the most current information.

After you've created a chart, you can always change its type, change the formatting, add new data series to it, or change an existing data series so that it uses data in a different range.

Embedded charts

An *embedded chart* basically floats on top of a worksheet, on the worksheet's drawing layer. The charts shown previously in this chapter are both embedded charts.

As with other drawing objects (such as a text box or a shape), you can move an embedded chart, resize it, change its proportions, adjust its borders, and add effects such as a shadow. Using an embedded chart enables you to view the chart next to the data that it uses. Alternatively, you can place several embedded charts together so that they print on a single page.

When you create a chart, it always starts off as an embedded chart. The exception to this rule is when you select a range of data and press F11 to create a default chart. Such a chart is created on a chart sheet.

Part III: Creating Charts and Other Visualizations

To make any changes to the actual chart in an embedded chart object, you must click it to *activate* the chart. When a chart is activated, Excel displays the two contextual tabs shown in Figure 18.3: Chart Design and Format.

FIGURE 18.3

Activating a chart displays additional tabs on the Excel Ribbon.

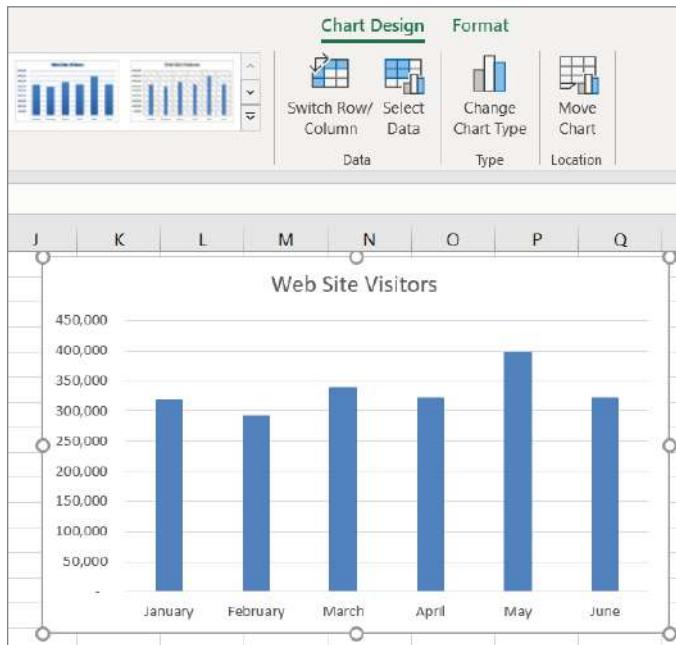


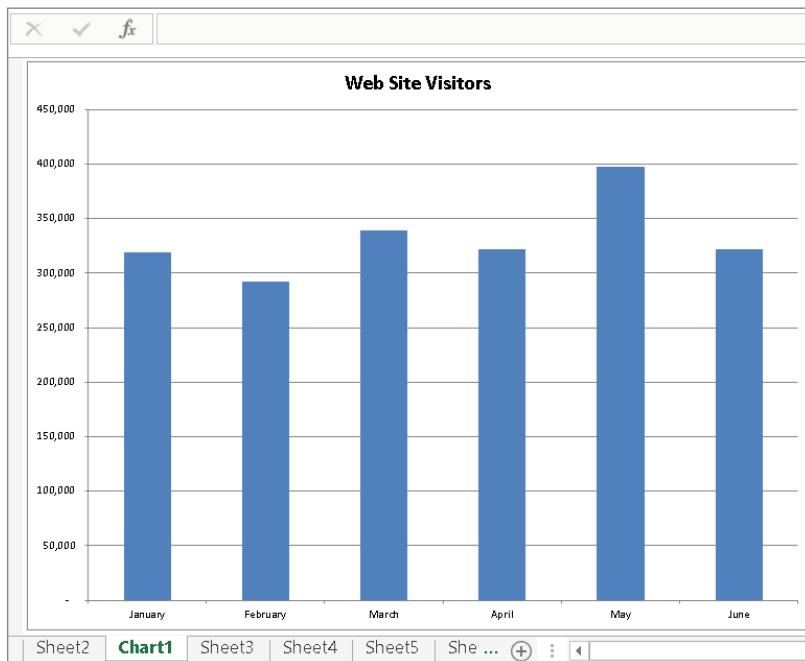
Chart sheets

You can move an embedded chart to its own chart sheet so that you can view it by clicking a sheet tab. When you move a chart to a chart sheet, the chart occupies the entire sheet. If you plan to print a chart on a page by itself, using a chart sheet is often your better choice. If you have many charts to create, you may want to put each one on a separate chart sheet to avoid cluttering your worksheet. This technique also makes locating a particular chart easier because you can change the name each chart sheets tab to an easy-to-find friendly name. Although chart sheets are not typically used in traditional dashboards, they can come in handy when producing reports that will be viewed in a multi-tab workbook.

Figure 18.4 illustrates a chart sheet. Note the only object on a chart sheet is the chart itself; there is no worksheet area.

FIGURE 18.4

A chart on a chart sheet



Parts of a chart

A chart is made up of many different elements, and all these elements are optional. Yes, you can create a chart that contains no chart elements—an empty chart. It's not very useful, but Excel allows it.

Refer to the chart in Figure 18.5 as you read the following description of the chart's elements.

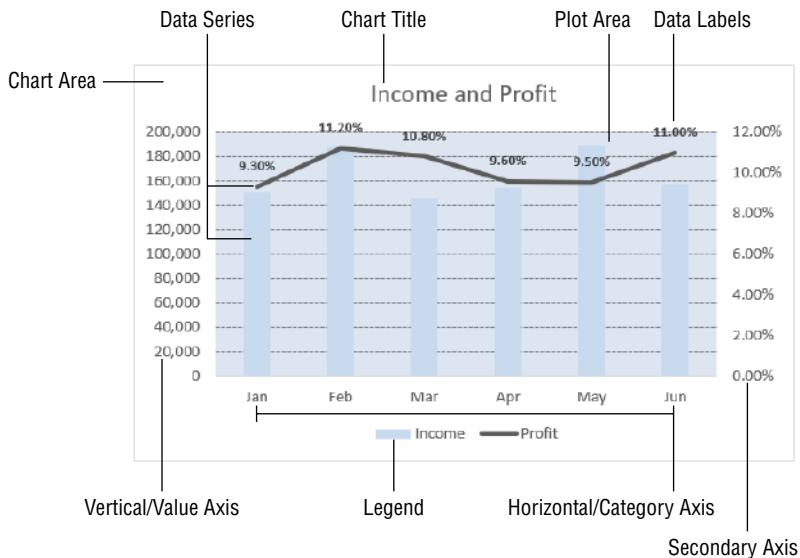
This particular chart is a combination chart that displays both columns and a line. The chart has two *data series*: Income and Profit Margin. Income is plotted as vertical columns, and Profit Margin is plotted as a line. Each bar represents a single *data point* (the value in a cell).

The chart has a horizontal axis, known as the *category axis*. This axis represents the category for each data point (January, February, and so on). This axis doesn't have a label because the category units are obvious.

Part III: Creating Charts and Other Visualizations

FIGURE 18.5

Parts of a chart



Notice that this chart has two vertical axes. These are known as *value axes*, and each one has a different scale. The axis on the left is for the column series (Income), and the secondary axis on the right is for the line series (Profit Margin).

Both vertical axes also display scale values. The axis on the left displays scale values from 0 to 200,000, in major unit increments of 20,000. The value axis on the right uses a different scale: 0 percent to 12 percent, in increments of 2 percent.

It's appropriate that the chart in Figure 18.5 have two value axes because the two data series vary dramatically in scale. If the Profit Margin data was plotted using the left axis, the line would not even be visible.

If a chart has more than one data series, you'll usually need a way to identify the data series or data points. A *legend*, for example, is often used to identify the various series in a chart. In this example, the legend appears at the bottom of the chart. Some charts also display *data labels* to identify specific data points. The example chart displays data labels for the Profit Margin series, but not for the Income series. In addition, most charts (including the example chart) contain a *chart title* and additional labels to identify the axes or categories.

The example chart also contains horizontal *gridlines* (which correspond to the values on the left axis). Gridlines are basically extensions of the value axis scale, which makes it easier for the viewer to determine the magnitude of the data points.

In addition, all charts have a *chart area* (the entire background area of the chart) and a *plot area* (the part that shows the actual chart, including the plotted data, the axes, and the axis labels).

Charts can have additional parts or fewer parts, depending on the chart type. For example, a pie chart (see Figure 18.6) has “slices” and no axes. A 3-D chart may have *walls* and a *floor* (see Figure 18.7).

Several other types of items can be added to a chart. For example, you can add a *trend line* or display *error bars*.

FIGURE 18.6

A pie chart

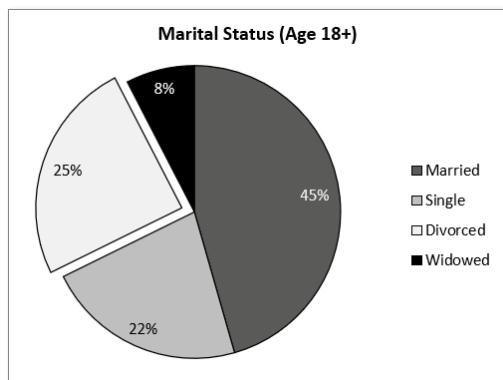


FIGURE 18.7

A 3-D column chart

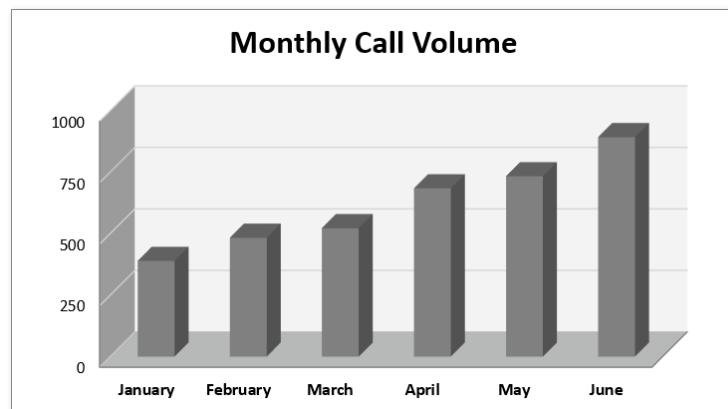


Chart limitations

As with most features in Excel, charts do have limits on the amount of data that they can handle and present. Table 18.1 lists the limitations of Excel charts.

TABLE 18.1 Limitations of Excel Charts

Item	Limitation
Charts in a worksheet	Limited by available memory
Worksheets referred to by a chart	255
Data series in a chart	255
Data points in a data series	32,000
Data points in a data series (3-D charts)	4,000
Total data points in a chart	256,000

Basic Steps for Creating a Chart

Creating a chart is relatively easy. The following sections describe how to create and then customize a basic chart to best communicate your business goals.

Creating the chart

Follow these general steps to create a chart using the data in Figure 18.8:

1. **Select the data you want to use in the chart.** Make sure you select the column and/or row headers if the data has them (in this example, you would select A1 : C4). Another option is to select a single cell within a range of data. Excel then uses the entire data range for the chart.

FIGURE 18.8

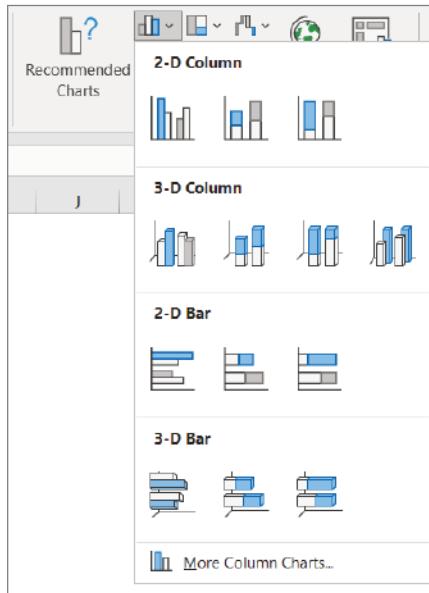
This data would make a good chart.

	A	B	C
1		Projected	Actual
2	Jan	2,000	1,895
3	Feb	2,500	2,643
4	Mar	3,500	3,648
c			

2. **Click the Insert tab and then click a Chart icon in the Charts group.** The icon expands into a gallery list that shows chart subtypes for the selected chart type (see Figure 18.9).

FIGURE 18.9

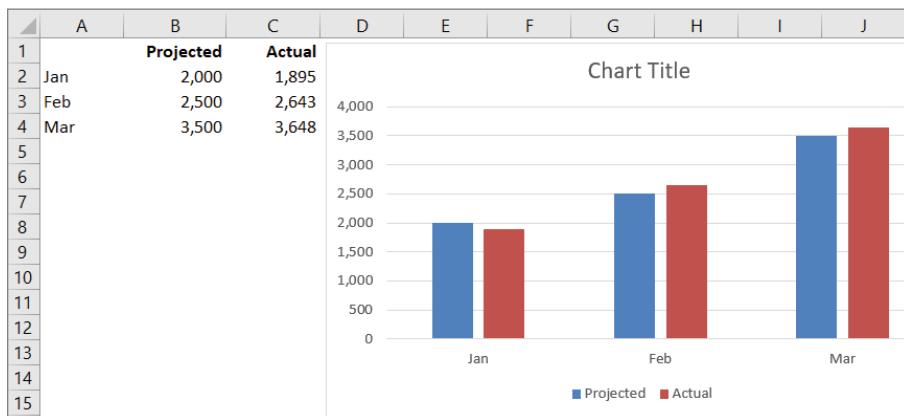
The icons in the Insert \Rightarrow Charts group expand to show a gallery of chart subtypes.



3. Click a Chart subtype, and Excel then creates the chart of the specified type.
Figure 18.10 shows a column chart created from the data.

FIGURE 18.10

A column chart with two data series



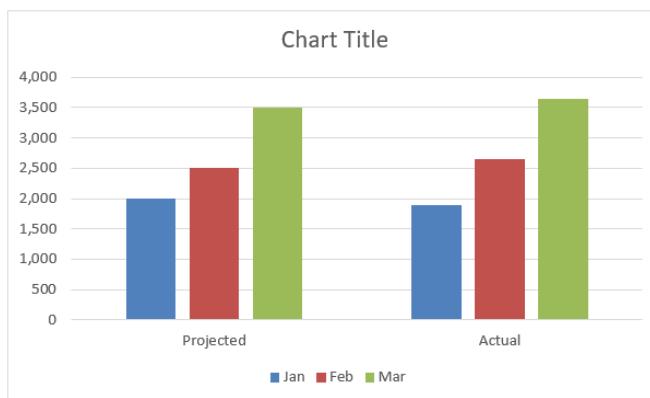
Switching the row and column orientation

When Excel creates a chart, it uses an algorithm to determine whether the data is arranged in columns or in rows. Most of the time Excel guesses correctly, but if it creates the chart using the wrong orientation, you can quickly change it by selecting the chart and choosing **Chart Design** \Rightarrow **Data** \Rightarrow **Switch Row/Column**. This command is a toggle, so if changing the data orientation doesn't improve the chart, just choose the command again (or click the Undo button found on the Quick Access Toolbar).

The orientation of the data has a drastic effect on the look (and, perhaps, the understandability) of your chart. Figure 18.11 shows the column chart in Figure 18.10 after changing the orientation. Notice that the chart now has three data series, one for each month. If the goal of your dashboard is to compare actual with projected values for each month, this version of the chart is much more difficult to interpret because the relevant columns are not adjacent.

FIGURE 18.11

The column chart, after swapping the row/column orientation



Changing the chart type

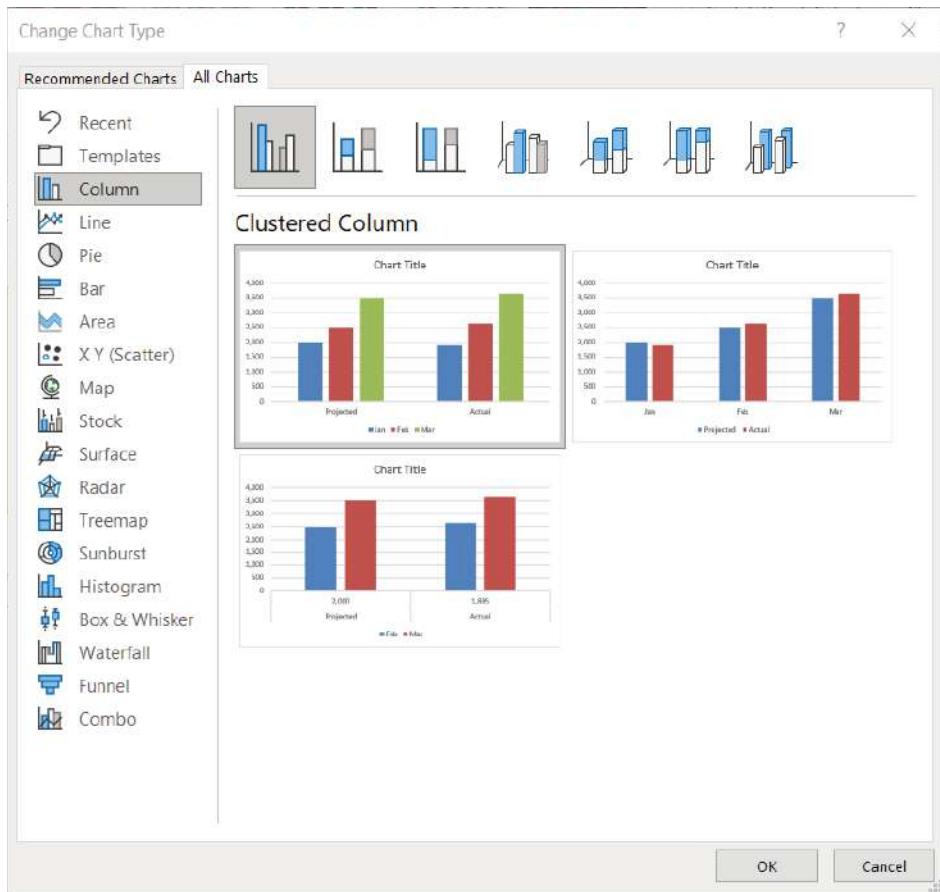
After you've created a chart, you can easily change the chart type. Although a column chart may work well for a particular data set, there's no harm in checking out some other chart types. You can choose **Chart Design** \Rightarrow **Type** \Rightarrow **Change Chart Type** to display the Change Chart Type dialog box and experiment with other chart types. Figure 18.12 shows the Change Chart Type dialog box.

Note

If your chart uses more than one data series, make sure that a chart element other than a data series is selected when you choose the **Chart Design** \Rightarrow **Type** \Rightarrow **Change Chart Type** command. If you select a series, the command changes the chart type of the selected series only.

FIGURE 18.12

The Change Chart Type dialog box



18

In the Change Chart Type dialog box, the main categories are listed on the left, and the sub-types are shown as icons. Select an icon and click OK, and Excel displays the chart using the new chart type. If you don't like the result, click the Undo button.

TIP

If the chart is an embedded chart, you can also change a chart's type by using the icons in the Insert \Rightarrow Charts group. In fact, this method is more efficient because it doesn't involve a dialog box.

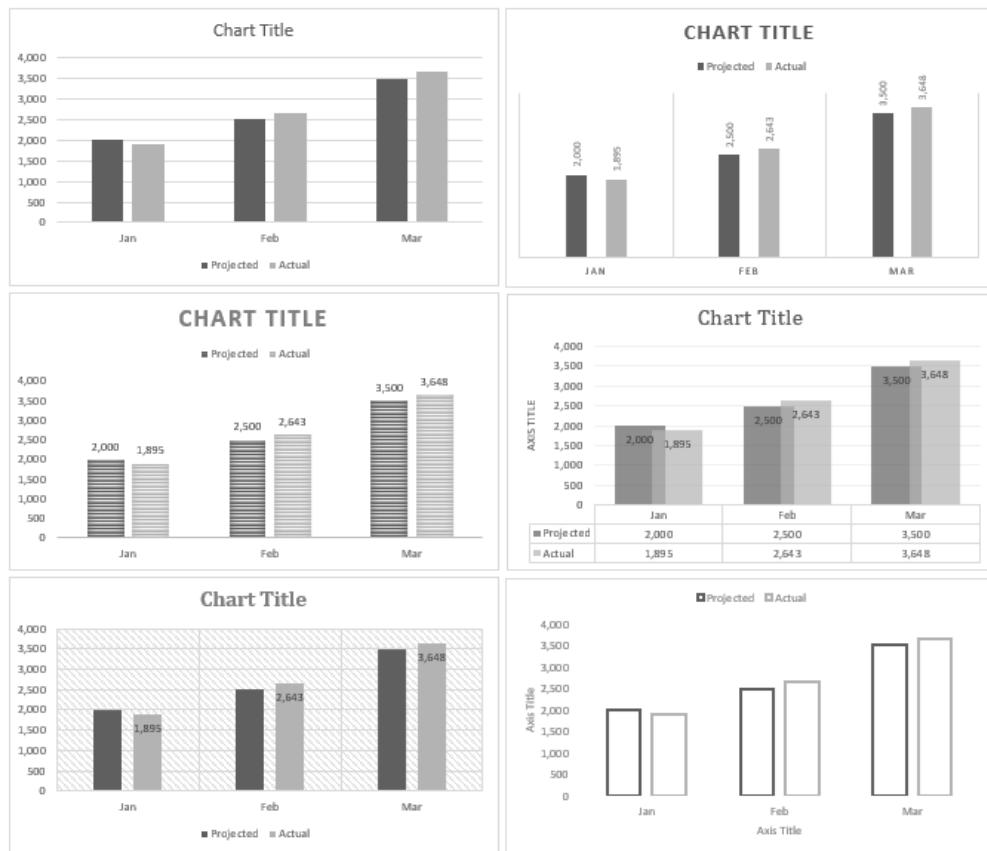
Applying a chart layout

Each chart type has a number of prebuilt layouts that you can apply with a single mouse click. A layout contains additional chart elements, such as a title, data labels, axes, and so on. This step is optional, but one of the prebuilt designs might be just what you need. Even if the layout isn't exactly what you want, it may be close enough that you need to make only a few adjustments.

To apply a layout, select the chart and use the Chart Design \Rightarrow Chart Layouts \Rightarrow Quick Layout gallery. Figure 18.13 shows how a column chart would look using various layouts.

FIGURE 18.13

One-click design variations of a column chart



Applying a chart style

The Chart Design \Rightarrow Chart Styles gallery contains quite a few styles that you can apply to your chart. The styles consist of various color choices and some special effects. Again, this step is optional.

TIP

The styles displayed in the gallery depend on the workbook's theme. When you choose Page Layout \Rightarrow Themes to apply a different theme, you'll see a new selection of chart styles designed for the selected theme.

Adding and deleting chart elements

In some cases, applying a chart layout or chart style (as described previously) gives you a chart with all of the elements that you need. Most of the time, however, you'll need to add or remove some chart elements and fine-tune the layout. You do this using the controls on the Chart Design and Format tabs.

For example, to give a chart a title, click the chart, then choose Chart Design \Rightarrow Chart Layouts \Rightarrow Add Chart Element \Rightarrow Chart Title. The control displays some options that determine where the title is placed. Excel inserts a title with the text "Chart Title." Click the text and replace it with your actual chart title.

Formatting chart elements

Every element in a chart can be formatted and customized in many ways. Many users are content with charts that are created using the steps described earlier in this chapter. Because you're reading this book, however, you probably want to find out how to customize charts for maximum impact.



Chapter 19 provides a detailed look into formatting and customizing charts.

Excel provides two ways to format and customize individual chart elements. Each of the following methods requires that you select the chart element first:

- Click the target chart element and then use the Ribbon controls on the Format tab.
- Click the chart element and then press **Ctrl+1** to display the Format task pane that's specific to the selected chart element.

You can also double-click a chart element to display the Format task pane for the element.

Assume that you'd like to change the color of the columns for one of the series in the chart. Click any column in the series (which selects the entire series). Then choose Format \Rightarrow Shape Styles \Rightarrow Shape Fill and choose a color from the displayed list.

To change the properties of the outline around the columns, use the Format \Rightarrow Shape Styles \Rightarrow Shape Outline control.

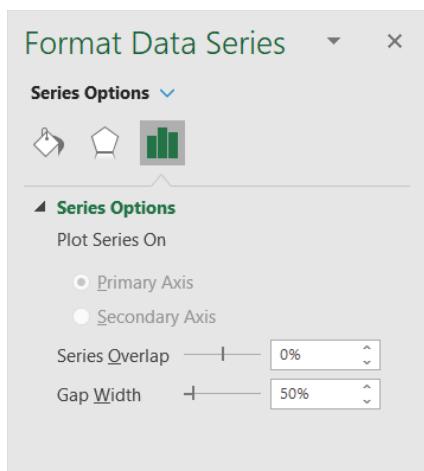
Part III: Creating Charts and Other Visualizations

To change the effects used in the columns (for example, to add a shadow), use the Format \Rightarrow Shape Styles \Rightarrow Shape Effects control.

Alternatively, you can select a series in the chart, press Ctrl+1, and use the Format Data Series task pane, as shown in Figure 18.14.

FIGURE 18.14

Using the Format Data Series task pane



Notice the icons in Figure 18.14. These icons are essentially tabs that expose various formatting options. Click an icon along the top and then expand a section on the left side to view additional controls. It's also a docked pane, so you can click another element in the chart. In other words, you don't have to close the task pane to see the changes you specify.

Modifying and Customizing Charts

The sections that follow cover common chart modifications.

Note

Before you can modify a chart, you must activate it. To activate an embedded chart, click it. Doing so activates the chart and also selects the element that you click. To activate a chart on a chart sheet, just click its sheet tab.

Moving and resizing a chart

If your chart is an embedded chart, you can freely move and resize it with your mouse. Click the chart's border and then drag the border to move the chart. Drag any of the eight "handles"

to resize the chart. The handles consist of white circles that appear on the chart's corners and edges when you click the chart's border. When the mouse pointer turns into a double arrow, click and drag to resize the chart.

Alternatively, after clicking the chart's border, you can use the Format \Rightarrow Size controls for more precise adjustments to the height and width of the chart. Simply select or type the dimensions directly into the Height and Width controls.

To move an embedded chart, just click its border at any location except one of the eight resizing handles. Then drag the chart to its new location. You also can use standard cut-and-paste techniques to move an embedded chart. Select the chart and choose Home \Rightarrow Clipboard \Rightarrow Cut (or press $Ctrl+X$). Then activate a cell near the desired location and choose Home \Rightarrow Clipboard \Rightarrow Paste (or press $Ctrl+V$).

The new location can be in a different worksheet or even in a different workbook. If you paste the chart to a different workbook, it will be linked to the data in the original workbook.

Another way to move a chart to a different location is to choose Chart Design \Rightarrow Location \Rightarrow Move Chart. This command displays the Move Chart dialog box, which lets you specify a new sheet for the chart (either a chart sheet or a worksheet).

Converting an embedded chart to a chart sheet

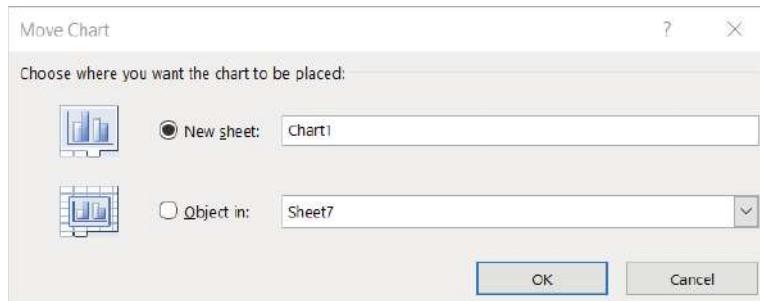
When you create a chart using the icons in the Insert \Rightarrow Charts group, the result is always an embedded chart. If you'd prefer that your chart be located on a chart sheet, you can easily move it.

To convert an embedded chart to a chart on a chart sheet, select the chart and choose Chart Design \Rightarrow Location \Rightarrow Move Chart to display the Move Chart dialog box shown in Figure 18.15. Select the New Sheet option and (optionally) provide a different name for the chart sheet.

To convert a chart on a chart sheet to an embedded chart, activate the chart sheet and then choose Chart Design \Rightarrow Location \Rightarrow Move Chart to display the Move Chart dialog box. Select the Object In option and specify the sheet by using the drop-down control.

FIGURE 18.15

Use the Move Chart dialog box to move an embedded chart to a chart sheet (or vice versa).



Copying a chart

To make an exact copy of an embedded chart, select the chart and choose Home \Rightarrow Clipboard \Rightarrow Copy (or press Ctrl+C). Then activate a cell near the desired location and choose Home \Rightarrow Clipboard \Rightarrow Paste (or press Ctrl+V). The new location can be in a different worksheet or even in a different workbook. If you paste the chart to a different workbook, it will be linked to the data in the original workbook.

To copy a chart on a chart sheet, press Ctrl while you drag the sheet tab to a new location in the tab list.

Deleting a chart

To delete an embedded chart, simply click the chart and then press Delete. When the Ctrl key is pressed, you can select multiple charts and then delete them all with a single press of the Delete key.

To delete a chart sheet, right-click its sheet tab and choose Delete from the shortcut menu. To delete multiple chart sheets, select them by pressing Ctrl while you click the sheet tabs.

Adding chart elements

To add new elements to a chart (such as a title, legend, data labels, or gridlines), activate the chart and use the controls in the Chart Elements + icon, which appears to the right of the chart. Note that each item expands to display additional options.

You can also use the Add Chart Element control on the Chart Design tab, in the Chart Layouts group.

Moving and deleting chart elements

Some elements within a chart can be moved: titles, legend, and data labels. To move a chart element, simply click it to select it and then drag its border.

The easiest way to delete a chart element is to select it and then press Delete. You can also use the controls on the Chart Elements icon, which appears to the right of the chart, to reposition chart elements.

Note

A few chart elements consist of multiple objects. For example, the data labels element consists of one label for each data point. To move or delete one data label, click once to select the entire element and then click a second time to select the specific data label. You can then move or delete the single data label.

Formatting chart elements

Many users are content to stick with the predefined chart styles and layouts. For more precise customizations, Excel allows you to work with individual chart elements and apply additional

formatting. You can use the Ribbon commands for some modifications, but the easiest way to format chart elements is to right-click the element and choose Format <Element> from the shortcut menu. The exact command depends on the element that you select. For example, if you right-click the chart's title, the shortcut menu command is Format Chart Title.

The Format command displays a task pane with options for the selected element. Changes that you make are displayed immediately in the chart. When you select a new chart element, the task pane changes to display the properties for the newly selected element. You can keep this task pane displayed while you work on the chart. It can be docked alongside the left or right part of the window or made free-floating and sizable.

TIP

If the Format task pane isn't displayed, you can double-click a chart element to display it.

Refer to the sidebar "Exploring the Format Task Pane" for an explanation of how the Format task panes work.

TIP

If you apply formatting to a chart element and decide that it wasn't such a good idea, you can revert to the original formatting for the particular chart style. Right-click the chart element and choose Reset to Match Style from the shortcut menu. To reset the entire chart, select the chart area when you issue the command.

18

Exploring the Format Task Pane

The Format task pane can be a bit deceiving. It contains many options that aren't visible, and you sometimes have to do quite a bit of clicking to find the formatting option you want.

The name of a given Format task pane depends on which chart element is selected. For instance, Figure 18.14 shows the task pane for a data series. The options available on the Format task pane will vary quite a bit, depending on which chart element you selected.

Regardless of the element that you select, however, the Format task pane will present you with icons at the top. Each icon has its own set of controls, which can be expanded to expose a set of formatting and customization options.

At first, the Format task pane will seem complicated and confusing. But as you get better acquainted with it, using the task pane gets much easier.

Copying a chart's formatting

If you create a nicely formatted chart and realize that you need to create several more charts that have the same formatting, you have these three choices:

- **Make a copy of the original chart and then change the data used in the copied chart.** One way to change the data used in a chart is to choose the Chart Design \Rightarrow Data \Rightarrow Select Data command and make the changes in the Select Data Source dialog box.
- **Create the other charts, but don't apply any formatting.** Then activate the original chart and press Ctrl+C. Select one of the other charts, and choose Home \Rightarrow Clipboard \Rightarrow Paste \Rightarrow Paste Special. In the Paste Special dialog box, click the Formats option and then click OK. Repeat for each additional chart.
- **Create a chart template and then use the template as the basis for the new charts, or you can apply the new template to existing charts.**



You will dive deep into customizing and formatting charts as well as chart templates in Chapter 19.

Renaming a chart

When you activate an embedded chart, its name appears in the Name box (located to the left of the Formula bar). To change the name of an embedded chart, simply select the chart and then enter the desired name in the Name box.

Why rename a chart? If a worksheet has many charts, you may prefer to activate a particular chart by name. Just type the chart's name in the Name box and press Enter. It's much easier to remember a chart named Monthly Sales as opposed to a chart named Chart 9.

NOTE

When you rename a chart, Excel allows you to use a name that already exists for another chart. Normally, it doesn't matter if multiple charts have the same name, but it can cause problems if you use VBA macros that select a chart by name.

Printing charts

Printing embedded charts is nothing special; you print them the same way that you print a worksheet. As long as you include the embedded chart in the range that you want to print, Excel prints the chart as it appears onscreen. When printing a sheet that contains embedded charts, it's a good idea to preview first (or use Page Layout view) to ensure that your charts don't span multiple pages. If you created the chart on a chart sheet, Excel always prints the chart on a page by itself.

TIP

If you select an embedded chart and choose File \Rightarrow Print, Excel prints the chart on a page by itself and does *not* print the worksheet data.

If you don't want a particular embedded chart to appear on your printout, access the Format Chart Area task pane and select the Size & Properties icon. Then expand the Properties section and clear the Print Object check box.

Understanding Chart Types

People who create charts usually do so to make a point or to communicate a specific message. Often, the message is explicitly stated in the chart's title or in a text box within the chart. The chart itself provides visual support.

Choosing the correct chart type is often a key factor in the effectiveness of the message. Therefore, it's often well worth your time to experiment with various chart types to determine which one conveys your message best.

In almost every case, the underlying message in a chart is some type of comparison. Examples of some general types of comparisons include the following:

- **Comparing an item to other items:** A chart may compare sales in each of a company's sales regions.
- **Comparing data over time:** A chart may display sales by month and indicate trends over time.
- **Making relative comparisons:** A common pie chart can depict relative proportions in terms of pie "slices."
- **Comparing data relationships:** An XY chart is ideal for this comparison. For example, you might show the relationship between monthly marketing expenditures and sales.
- **Comparing frequency:** You can use a common histogram, for example, to display the number (or percentage) of students who scored within a particular grade range.
- **Identifying outliers or unusual situations:** If you have thousands of data points, creating a chart may help identify data that isn't representative.

18

Choosing a chart type

A common question among Excel users is "How do I know which chart type to use for my data?" Unfortunately, this question has no cut-and-dried answer. Perhaps the best answer is a vague one: use the chart type that gets your message across in the simplest way. A good starting point is Excel's recommended charts. Select your data and choose Insert \Rightarrow Charts \Rightarrow Recommended Charts to see the chart types that Excel suggests. Remember that these suggestions are not always the best choices.

NOTE

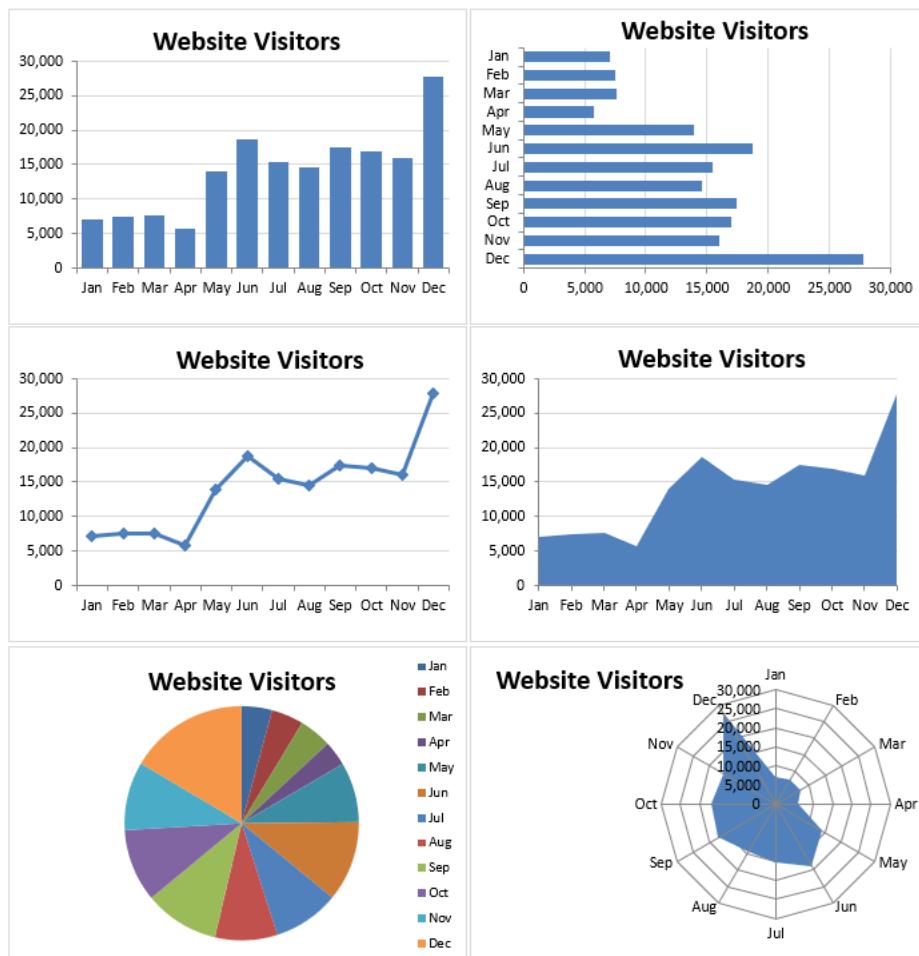
In the Ribbon, the Charts group of the Insert tab shows the Recommended Charts button, plus nine other drop-down buttons. All of these drop-down buttons display multiple chart types. For example, column and bar charts are all available from a single drop-down button. Similarly, scatter charts and bubble charts share a single button. Probably the easiest way to choose a particular chart type is to select Insert \Rightarrow Charts \Rightarrow Recommended Charts, which displays the Insert Chart dialog box with the Recommended Charts tab displayed. Select the All Charts tab, and you'll have a concise list of all chart and subchart types.

Part III: Creating Charts and Other Visualizations

Figure 18.16 shows the same set of data plotted by using six different chart types. Although all six charts represent the same information (monthly website visitors), they look quite different from one another.

FIGURE 18.16

The same data, plotted by using six chart types



The column chart (upper left) is probably the best choice for this particular set of data because it clearly shows the information for each month in discrete units. The bar chart (upper right) is similar to a column chart, but since the axes are swapped and most people are more

accustomed to seeing time-based information extend from left to right rather than from top to bottom, this isn't the optimal choice.

The line chart (middle left) may not be the best choice because it can imply that the data is continuous—that points exist in between the 12 actual data points. This same argument may be made against using an area chart (middle right).

The pie chart (lower left) is simply too confusing and does nothing to convey the time-based nature of the data. Pie charts are most appropriate for a data series in which you want to emphasize proportions among a relatively small number of data points. If you have too many data points, a pie chart can be impossible to interpret.

The radar chart (lower right) is clearly inappropriate for this data. People aren't accustomed to viewing time-based information in a circular direction!

Fortunately, changing a chart's type is easy, so you can experiment with various chart types until you find the one that represents your data accurately, clearly, and as simply as possible.

The remainder of this chapter contains more information about the various Excel chart types. The examples and discussion may give you a better handle on determining the most appropriate chart type for your data.

Column charts

Probably the most common chart type is the *column chart*, which displays each data point as a vertical column, the height of which corresponds to the value. The value scale is displayed on the vertical axis, which is usually on the left side of the chart. You can specify any number of data series, and the corresponding data points from each series can be stacked on top of each other. Typically, each data series is depicted in a different color or pattern.

Column charts are often used to compare discrete items, and they can depict the differences between items in a series or items across multiple series. Excel offers seven column-chart subtypes.

Figure 18.17 shows an example of a clustered column chart that depicts monthly sales for two products. From this chart, it's clear that Sprocket sales have always exceeded Widget sales. In addition, Widget sales have been declining over the five-month period, whereas Sprocket sales are increasing.

The same data, in the form of a stacked column chart, is shown in Figure 18.18. This chart has the added advantage of depicting the combined sales over time. It shows that total sales have remained fairly steady each month, but the relative proportions of the two products have changed.

Part III: Creating Charts and Other Visualizations

FIGURE 18.17

This clustered column chart compares monthly sales for two products.

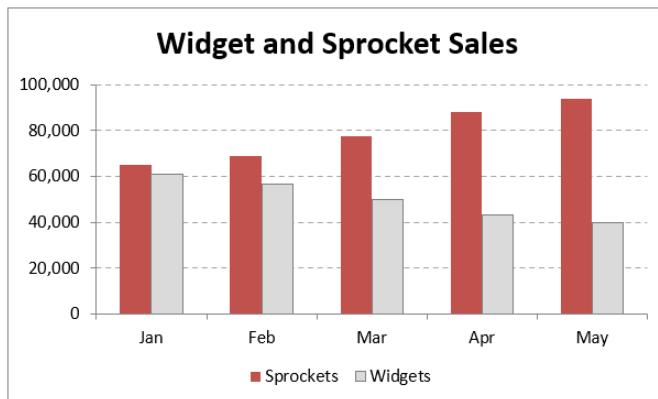


FIGURE 18.18

This stacked column chart displays sales by product and depicts the total sales.

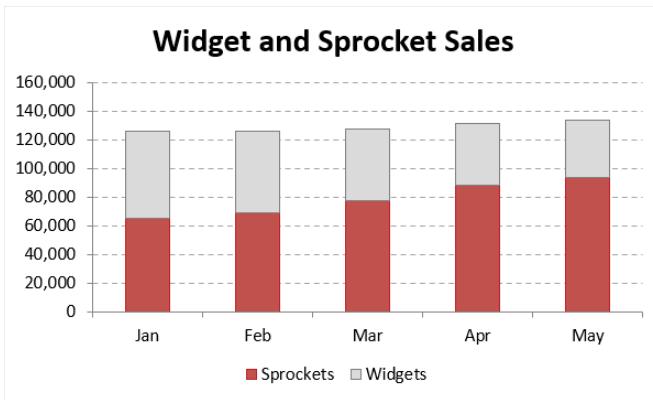


Figure 18.19 shows the same sales data plotted as a 100% stacked column chart. This chart type shows the relative contribution of each product by month. Notice that the vertical axis displays percentage values, not sales amounts. This chart provides no information about the actual sales volumes, but such information could be provided using data labels. This type of chart is often a good alternative to using several pie charts. Instead of using a pie to show the relative sales volume in each year, the chart uses a column for each year.

FIGURE 18.19

This 100% stacked column chart displays monthly sales as a percentage.

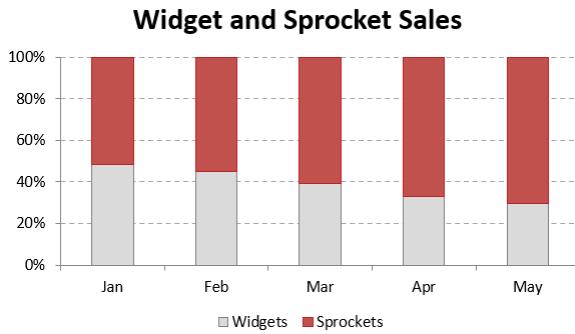
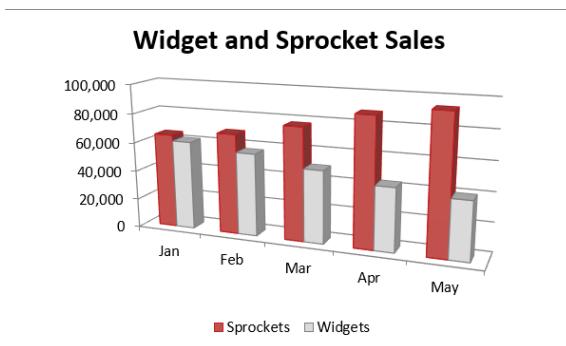


Figure 18.20 shows the same data plotted with a 3-D clustered column chart. Many people use this type of chart to add some visual pizzazz, but it's generally considered bad practice to use 3-D charts because precise comparisons are difficult due to the distorted perspective view.

FIGURE 18.20

A 3-D column chart



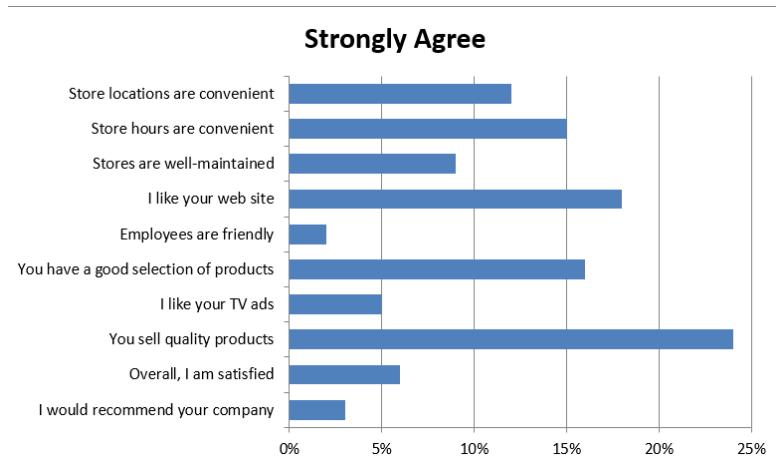
Bar charts

A *bar chart* is essentially a column chart that has been rotated 90 degrees clockwise. One distinct advantage to using a bar chart is that the category labels may be easier to read. Figure 18.21 shows a bar chart that displays a value for each of 10 survey items. The category labels are lengthy, and displaying them legibly with a column chart would be difficult. Excel offers six bar chart subtypes.

Part III: Creating Charts and Other Visualizations

FIGURE 18.21

If you have lengthy category labels, a bar chart may be a good choice.



You can include any number of data series in a bar chart. In addition, the bars can be “stacked” from left to right.

Line charts

Line charts are often used to plot continuous data and are useful for identifying trends. For example, plotting daily sales as a line chart may enable you to identify sales fluctuations over time. Normally, the category axis for a line chart displays equal intervals. Excel supports seven line chart subtypes.

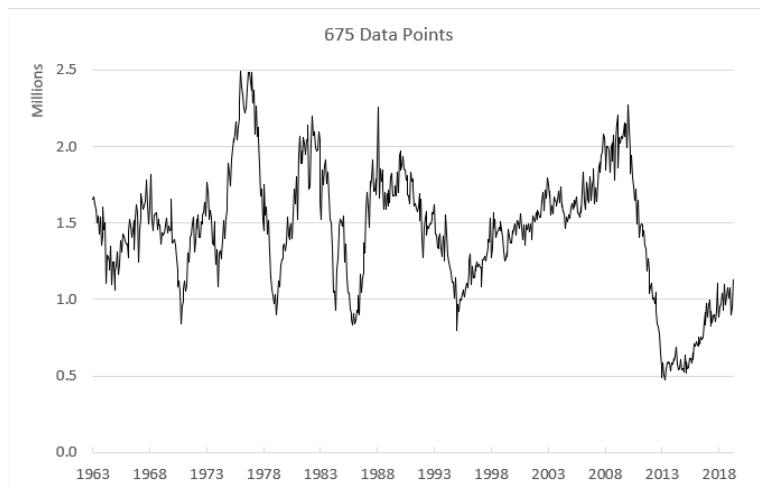
See Figure 18.22 for an example of a line chart that depicts monthly data (675 data points). Although the data varies quite a bit monthly, the chart clearly depicts the cycles.

A line chart can use any number of data series, and you distinguish the lines by using different colors, line styles, or markers. Figure 18.23 shows a line chart that has three series. The series are distinguished by markers (circles, squares, and triangles) and different line colors. When the chart is printed on a noncolor printer, the markers are the only way to identify the lines.

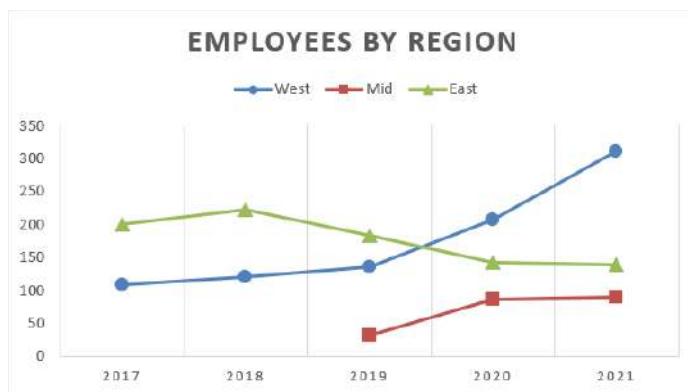
The final line chart example, shown in Figure 18.24, is a 3-D line chart. Although it has a nice visual appeal, it's certainly not the clearest way to present the data. In fact, it's fairly worthless.

FIGURE 18.22

A line chart often can help you spot trends in your data.

**FIGURE 18.23**

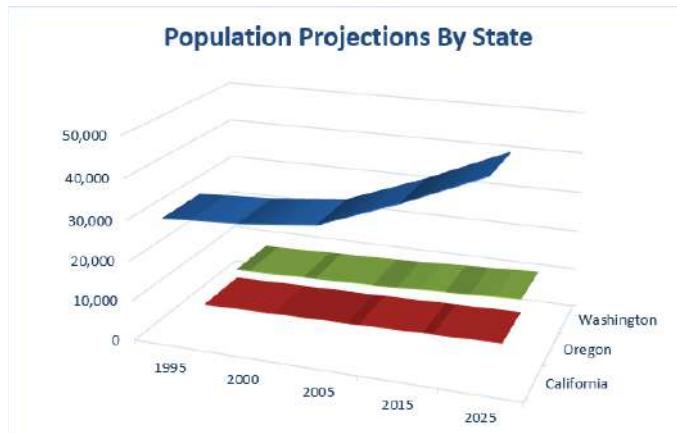
This line chart displays three series.



Part III: Creating Charts and Other Visualizations

FIGURE 18.24

This 3-D line chart does not present the data very well.



Pie charts

A *pie chart* is useful when you want to show relative proportions or contributions to a whole. A pie chart uses only one data series. Pie charts are most effective with a small number of data points. Generally, a pie chart should use no more than five or six data points (or slices). A pie chart with too many data points can be difficult to interpret.

CAUTION

All of the values in a pie chart must be positive numbers. If you create a pie chart that uses one or more negative values, the negative values will be converted to positive values, which is probably not what you intended.

You can “explode” one or more slices of a pie chart for emphasis (see Figure 18.25). Activate the chart and click any pie slice to select the entire pie. Then click the slice that you want to explode and drag it away from the center.

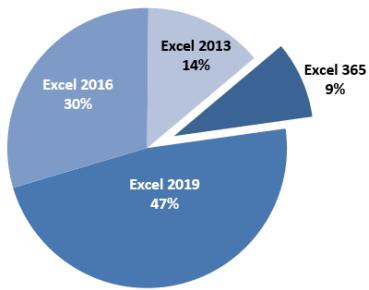
The “pie of pie” and “bar of pie” chart subtypes enable you to display a secondary chart that provides more detail for one of the pie slices. Figure 18.26 shows an example of a bar of pie chart. The pie chart shows the breakdown of four expense categories: Rent, Supplies, Utilities, and Salary. The secondary bar chart provides an additional regional breakdown of the Salary category.

The data used in the chart resides in A2 : B8. When the chart was created, Excel made a guess as to which categories belong to the secondary chart. In this case, the guess was to use the last three data points for the secondary chart—and the guess was incorrect.

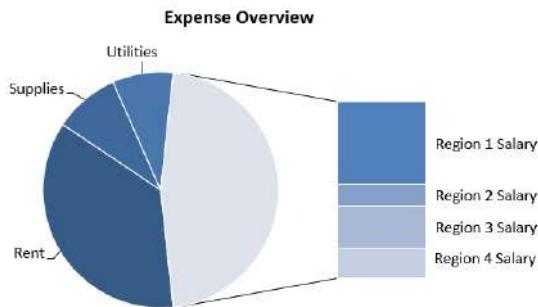
FIGURE 18.25

A pie chart with one slice exploded

Excel Usage At XYZ Corporation

**FIGURE 18.26**

A bar of pie chart that shows detail for one of the pie slices



To correct the chart, right-click any of the pie slices and choose Format Data Series. In the Format Data Series task pane, select the Series Options icon and make the changes. In this example, we chose Split Series by Position and specified that the second plot contains four values in the series.

Another pie chart subtype is called a *doughnut chart*. It's basically a pie chart with a hole in the middle. Unlike a pie chart, however, a doughnut chart can display multiple series.

XY (scatter) charts

Another common chart type is an *XY chart* (also known as scatter grams or scatter plots). An XY chart differs from most other chart types in that both axes display values. (An XY chart has no category axis.)

This type of chart often is used to show the relationship between two variables. Figure 18.27 shows an example of an XY chart that plots the relationship between weeks in a weight loss program (horizontal axis) and weight lost (vertical axis).

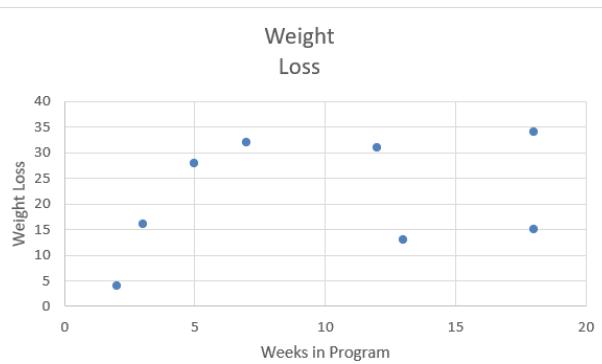
Part III: Creating Charts and Other Visualizations

Note

Although these data points correspond to time, the chart doesn't convey any time-related information. In other words, the data points are plotted based only on their two values.

FIGURE 18.27

An XY chart shows the relationship between two variables.



Area charts

Think of an *area chart* as a line chart in which the areas below the lines have been filled. Figure 18.28 shows an example of a stacked area chart. Stacking the data series enables you to see the total clearly, as well as the contribution by each series.

FIGURE 18.28

A stacked area chart

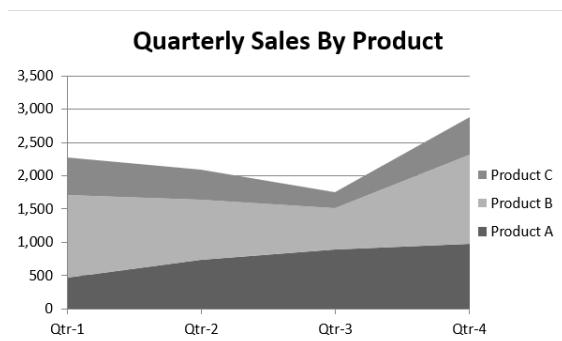
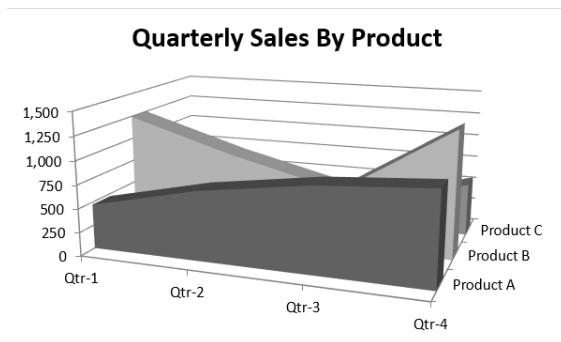


Figure 18.29 shows the same data, plotted as a 3-D area chart. As you can see, it's not an example of an effective chart. The data for products B and C is partially obscured. In some cases, the problem can be resolved by rotating the chart or using transparency. But usually the best way to salvage a chart like this is to select a new chart type.

FIGURE 18.29

This 3-D area chart is not a good choice.



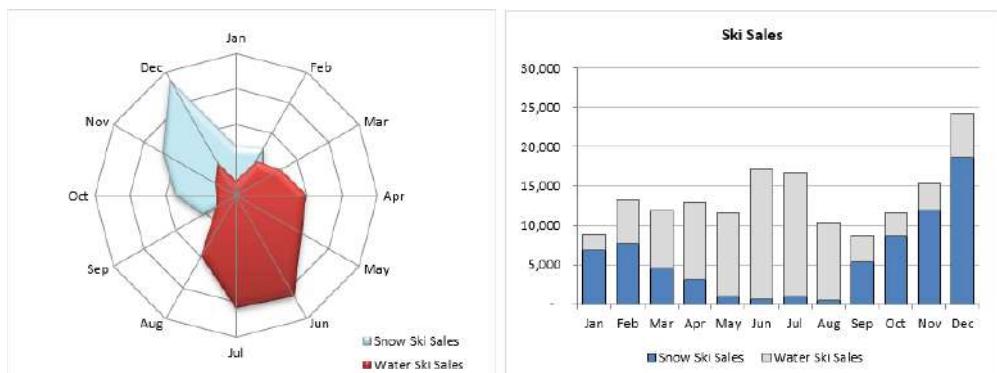
Radar charts

You may not be familiar with this type of chart. A *radar chart* is a specialized chart that has a separate axis for each category, and the axes extend outward from the center of the chart. The value of each data point is plotted on the corresponding axis.

Figure 18.30 shows an example of a radar chart on the left. This chart plots two data series across 12 categories (months) and shows the seasonal demand for snow skis versus water skis. Note that the water-ski series partially obscures the snow-ski series.

FIGURE 18.30

Plotting ski sales using a radar chart with 12 categories and two series



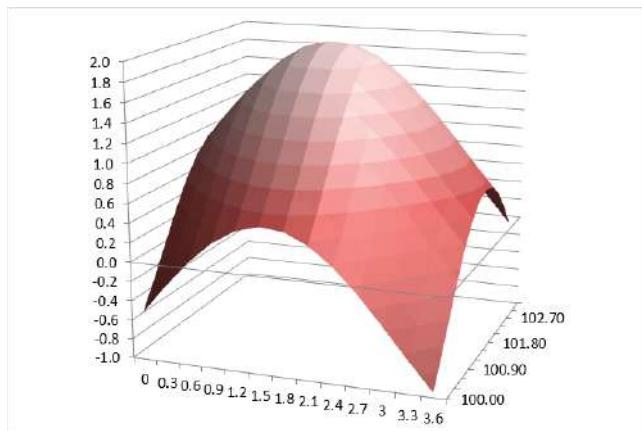
Using a radar chart to show seasonal sales may be an interesting approach, but it's certainly not the best chart type. As you can see, the stacked bar chart on the right shows the information much more clearly.

Surface charts

Surface charts display two or more data series on a surface. As Figure 18.31 shows, these charts can be quite interesting. Unlike other charts, Excel uses color to distinguish values, not to distinguish the data series. The number of colors used is determined by the major unit scale setting for the value axis. Each color corresponds to one major unit.

FIGURE 18.31

A surface chart



NOTE

A surface chart does not plot 3-D data points. The series axis for a surface chart, as with all other 3-D charts, is a category axis—not a value axis. In other words, if you have data that is represented by x, y, and z coordinates, it can't be plotted accurately on a surface chart unless the x and y values are equally spaced.

Bubble charts

Think of a *bubble chart* as an XY chart that can display an additional data series, which is represented by the size of the bubbles. As with an XY chart, both axes are value axes. (There is no category axis.)

Figure 18.32 shows an example of a bubble chart that depicts the results of a weight-loss program. The horizontal value axis represents the original weight, the vertical value axis shows the number of weeks in the program, and the size of the bubbles represents the amount of weight lost.

FIGURE 18.32

A bubble chart



Stock charts

Stock charts are most useful for displaying stock-market information. These charts require three to five data series, depending on the subtype.

Figure 18.33 shows an example of each of the four stock chart types. The two charts on the bottom display the trade volume and use two value axes. The daily volume, represented by columns, uses the axis on the left. The *up-bars*, sometimes referred to as *candlesticks*, are the vertical lines that depict the difference between the opening and closing price. A black up-bar indicates that the closing price was lower than the opening price.

Stock charts aren't just for stock price data. Figure 18.34 shows a stock chart that depicts the high, low, and average temperatures for each day in May. This is a high-low-close chart.

Newer Chart Types for Excel

If you haven't been able to upgrade since Excel 2013, you'll see a few new chart types introduced with Excel 2016 and Excel 2019. The following sections present examples of each new chart type along with an explanation of the type of data required.

Histogram charts

Histogram charts display the count of data items in each of several discrete bins. Those of us who have used the Analysis ToolPak in the past to create histogram charts (see Chapter 30, "Analyzing Data with the Analysis ToolPak") will tell you that it took some effort to create a nice-looking chart. However, the new histogram chart type makes short work of the task.

Part III: Creating Charts and Other Visualizations

FIGURE 18.33

The four stock chart subtypes

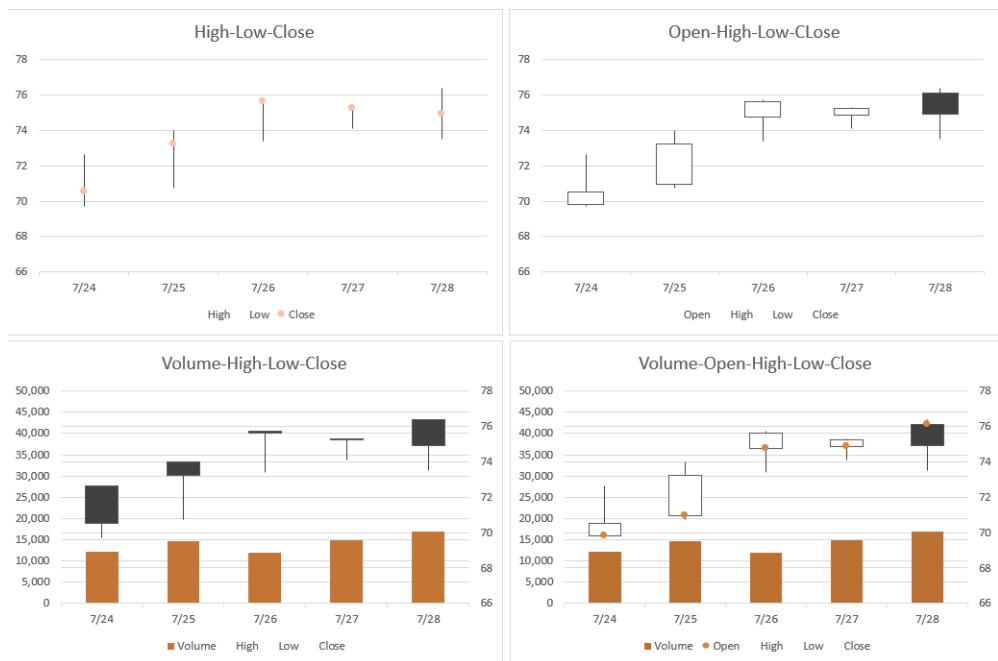


FIGURE 18.34

Plotting temperature data with a stock chart

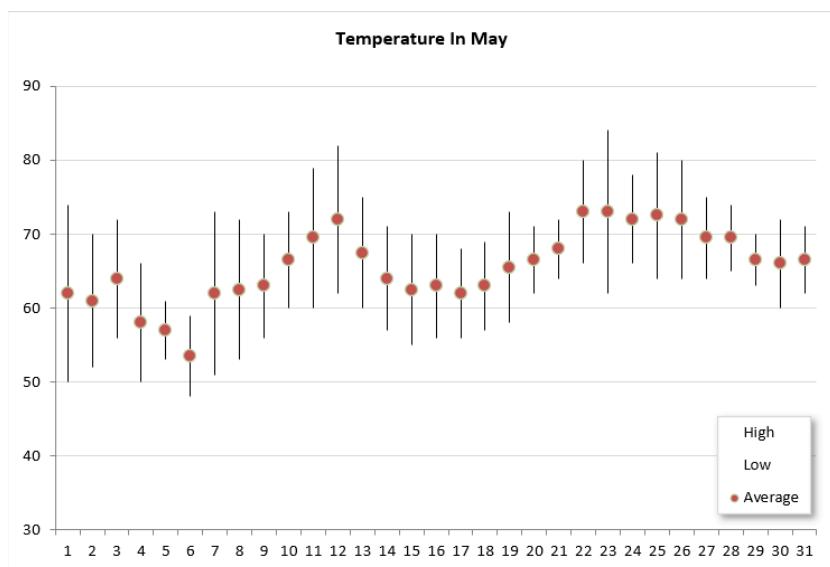
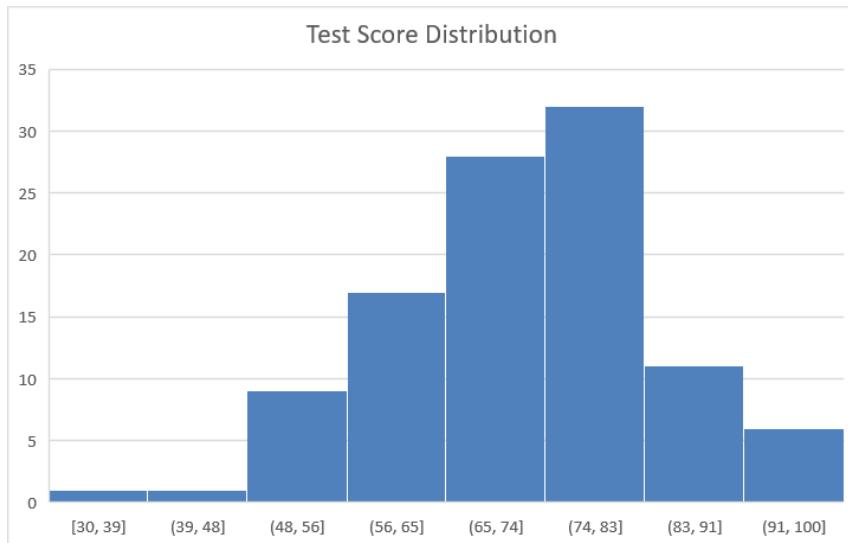


Figure 18.35 shows a histogram created from 105 student test scores. The bins are displayed as category labels. You control the number of bins by using the Axis Options section of the Format Axis task pane. In this example, we specified eight bins, and Excel took care of all of the details.

FIGURE 18.35

Displaying a student grade distribution using a histogram chart



Pareto charts

The *Pareto chart* type is a combination chart in which columns are displayed in descending order, and the columns use the left axis. The line shows the cumulative percentage and uses the right axis.

Figure 18.36 shows a Pareto chart created from the data in the range A2 : B14. Notice that Excel sorted the items in the chart. The line shows, for example, that approximately 50 percent of all complaints are in the top three categories.

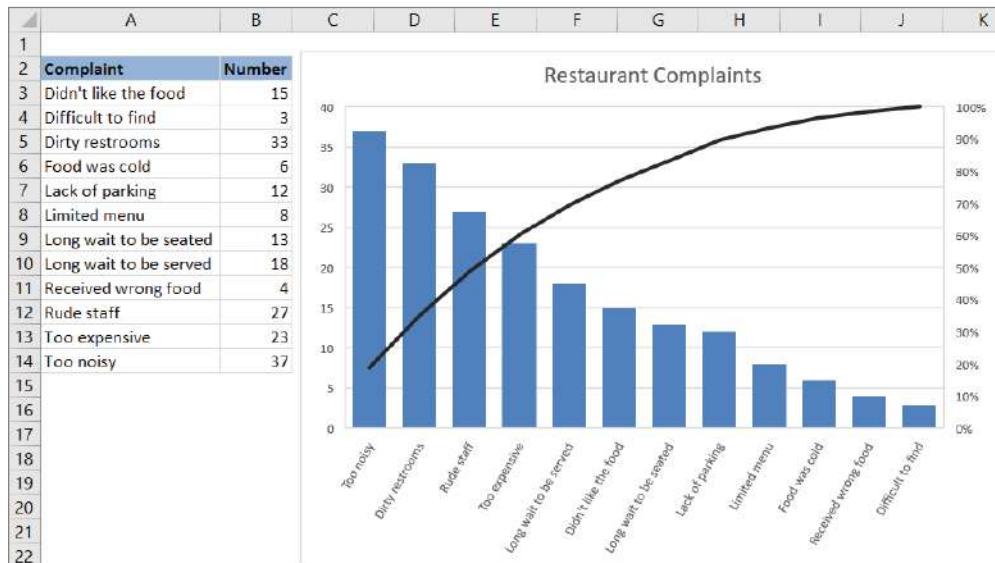
Note

The Pareto chart type is actually a Histogram subtype. You can find the Pareto chart type under Insert \Rightarrow Charts \Rightarrow Insert Statistic Chart.

Part III: Creating Charts and Other Visualizations

FIGURE 18.36

A Pareto chart displays the number of complaints graphically.



Waterfall charts

A *waterfall chart* is used to show the cumulative effect of a series of numbers, usually both positive and negative numbers. The result is a staircase-like display. Waterfall charts were introduced in Excel 2016.

Figure 18.37 shows a waterfall chart that uses the data in column D. Waterfall charts typically display the ending total as the last bar, with its origin at zero. To display the total column correctly, select the column, right-click, and choose Set as Total from the shortcut menu.

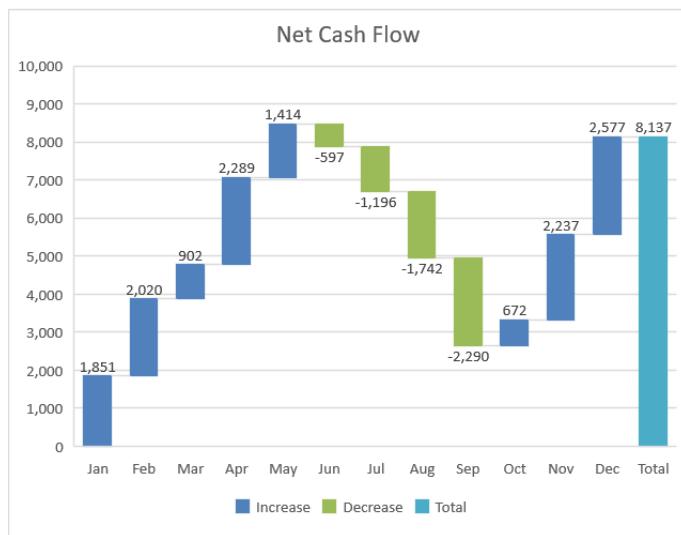
Box & whisker charts

A *box & whisker chart* is often used to summarize data visually. In the past, it was possible to create such charts using Excel, but it required quite a bit of setup work. Now, it's simple.

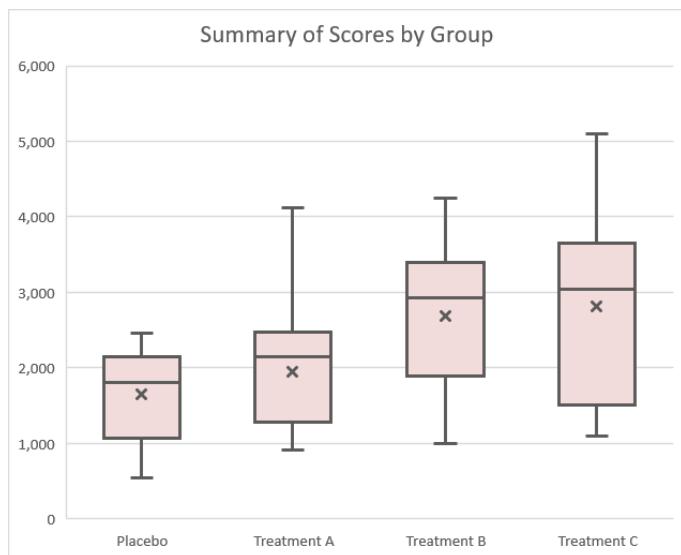
Figure 18.38 shows a box & whisker chart created for four groups of subjects. The data is in a two-column table. In the chart, the vertical lines extending from the box represent the numerical range of the data (minimum and maximum values). The "boxes" represent the 25th through the 75th percentile. The horizontal line inside the box is the median value (or 50th percentile), and the X is the average. This type of chart enables the viewer to make quick comparisons among groups of data.

FIGURE 18.37

A waterfall chart showing positive and negative net cash flows

**FIGURE 18.38**

A box & whisker chart that summarizes data for four groups



Part III: Creating Charts and Other Visualizations

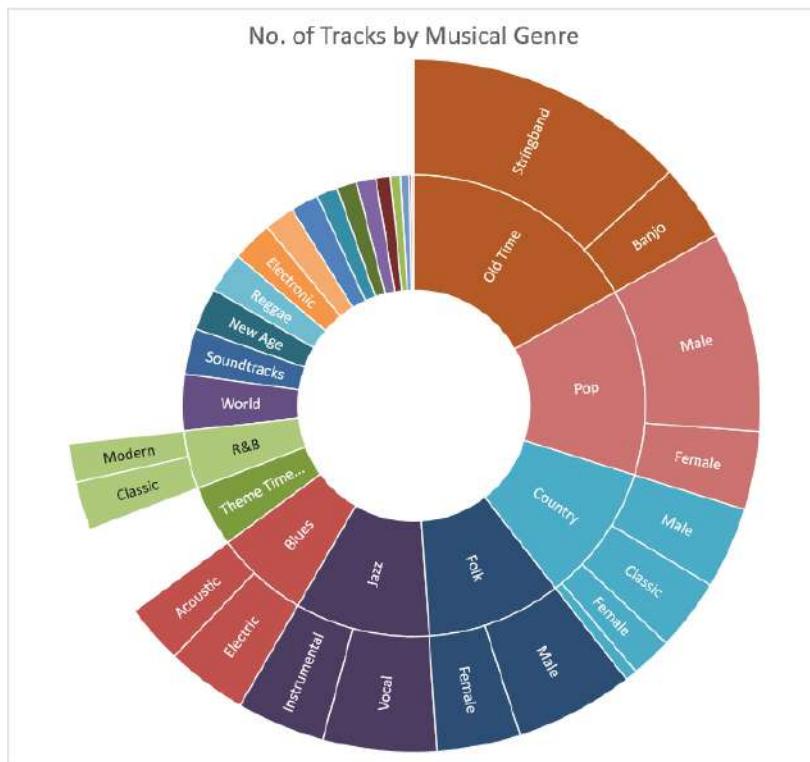
The Series Options section of the Format Data Series task pane contains some options for this chart type.

Sunburst charts

Sunburst charts are like pie charts with multiple concentric layers. This chart type is most useful for data that's organized hierarchically. Figure 18.39 shows an example of a sunburst chart that depicts a music collection. It shows the number of tracks by genre and subgenre. For example, you'll see in Figure 18.39 that jazz is a genre of music, and within that genre, there are vocal and instrumental subgenres. Also note that some genres have no subgenres.

FIGURE 18.39

A sunburst chart that depicts a music collection by genre and subgenre



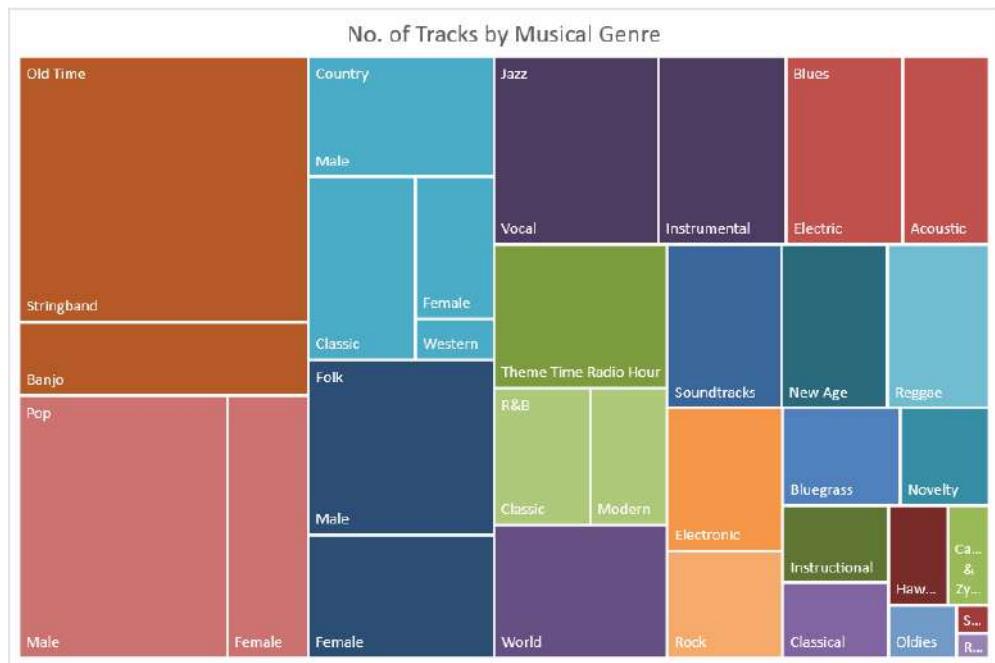
A potential problem with this chart type is that some slices are so small that the data labels can't be displayed.

Treemap charts

With the release of Excel 2016, *treemap charts* were included as a chart type. Like a sunburst chart, a treemap chart is suited for hierarchical data. The data, however, is represented as rectangles. Figure 18.40 shows the data from the previous example, plotted as a treemap chart.

FIGURE 18.40

A treemap chart that depicts a music collection by genre and subgenre



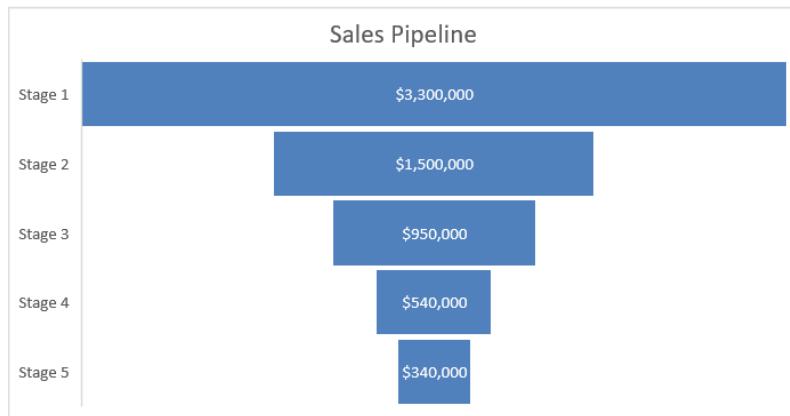
Funnel charts

Funnel charts (introduced with Excel 2019) are ideal for representing the relative values in each stage of a process. These charts are typically used to visualize sales pipelines (see Figure 18.41).

Part III: Creating Charts and Other Visualizations

FIGURE 18.41

A funnel chart visualizing the value in each stage of a sales pipeline

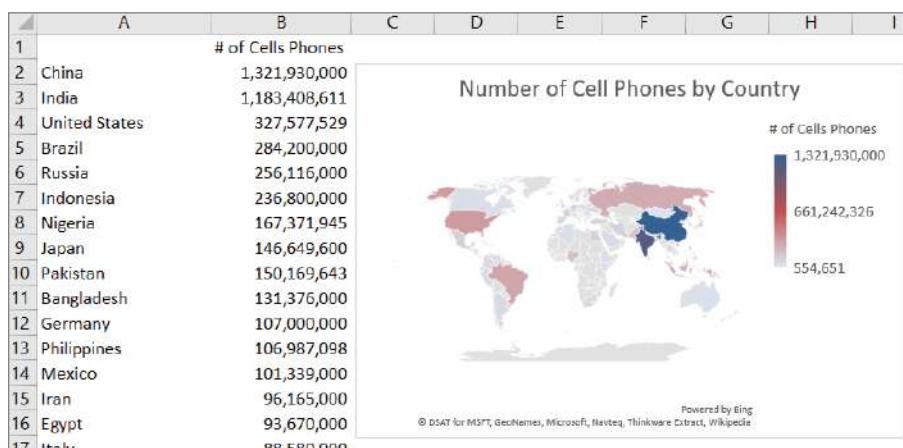


Map charts

Excel 2019 also introduced a new *map chart* type. Map charts leverage Bing maps to render location-based visualizations. As you can see in Figure 18.42, all you need to provide are location indicators (in this case, country names), and then Excel does the rest.

FIGURE 18.42

Map charts are ideal for visualizing location-based data.

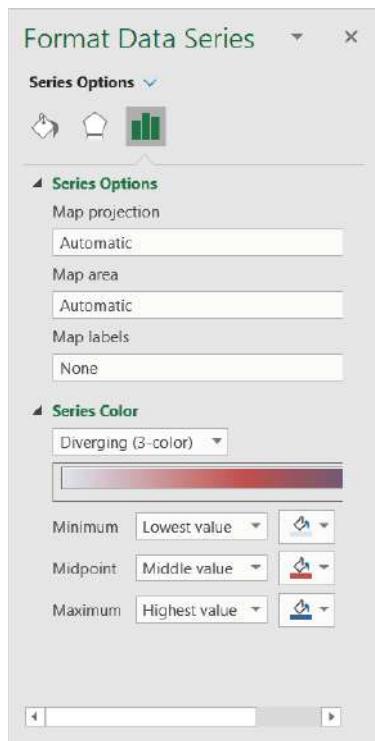


Map charts are remarkably flexible, allowing you create a chart based on province names, county names, cities, and even ZIP codes. As long as Bing can recognize the values that you are using to identify geography, your chart will render seamlessly.

Double-clicking the map will active the Format Data Series task pane (see Figure 18.43), exposing a handful of unique formatting options. These options allow you to change the projection (flat or curved map), the area (show all locations or just those that have data), and the series color (apply color banding based on values).

FIGURE 18.43

Map charts come with unique customization options.



Using Advanced Charting Techniques

IN THIS CHAPTER

Understanding chart customization

Changing basic chart elements

Working with data series

Discovering some chart-making tricks

Excel makes creating a basic chart easy. Select your data, choose a chart type, and you're finished. You may take a few extra seconds and select one of the prebuilt chart styles and maybe even select one of the chart layouts. But if your goal is to create the most effective chart possible, you probably want to take advantage of the additional customization techniques available in Excel.

Customizing a chart involves changing its appearance, as well as possibly adding new elements to it. These changes can be purely cosmetic (such as changing colors, modifying line widths, or adding a shadow) or quite substantial (say, changing the axis scales or adding a second value axis). Chart elements that you might add include such features as a data table, a trend line, or error bars.

The preceding chapter introduced charting in Excel and described how to create basic charts. This chapter takes the topic to the next level. You learn how to customize your charts to the maximum so that they look exactly like you want. You also pick up some slick charting tricks that will make your charts even more impressive.

Selecting Chart Elements

Modifying a chart is similar to everything else you do in Excel: first you make a selection (in this case, select a chart element), and then you issue a command to do something with your selection.

You can select only one chart element (or one group of chart elements) at a time. For example, if you want to change the font for two axis labels, you must work on each set separately.

Excel provides three ways, described in the following sections, to select a particular chart element:

- Mouse
- Keyboard
- Chart Elements control

Selecting with the mouse

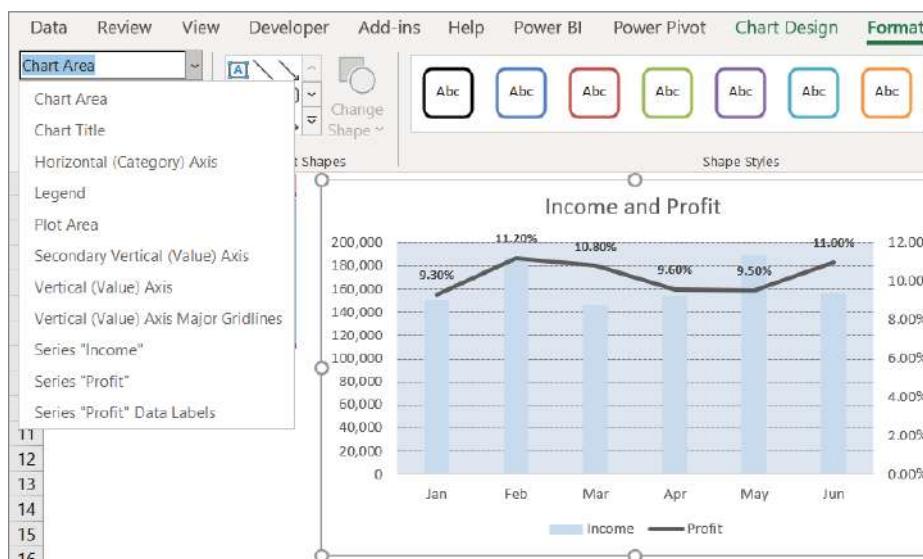
To select a chart element with your mouse, just click the element. The chart element appears with small circles at the corners.

TIP

Some chart elements can be a bit tricky to select. To ensure that you select the chart element that you intended to select, view the Chart Elements control, located in the Format \Rightarrow Current Selection group of the Ribbon (see Figure 19.1). Or, if the Format task pane is displayed, you can identify the selected chart element by the title of the task pane. Press **Ctrl+1** to display the Format task pane.

FIGURE 19.1

The Chart Elements control (in the upper-left corner) displays the name of the selected chart element. In this example, the “Chart Area” is selected.



When you move the mouse over a chart, a small *chart tip* displays the name of the chart element under the mouse pointer. When the mouse pointer is over a data point, the chart tip also displays the value of the data point.

TIP

If you find these chart tips annoying, you can turn them off. Choose File \Rightarrow Options, and select the Advanced tab in the Excel Options dialog box. Locate the Chart section, and clear either the Show Chart Element Names on Hover or the Show Data Point Values on Hover check box (you can also clear both).

Some chart elements (such as a series, a legend, and data labels) consist of multiple items. For example, a chart series element is made up of individual data points. To select a particular data point, click twice: first click the series to select it, and then click the specific element within the series (for example, a column or a line chart marker). Selecting the element enables you to apply formatting to only a particular data point in a series.

You may find that some chart elements are difficult to select with the mouse. If you rely on the mouse for selecting a chart element, you may have to click it several times before the desired element is actually selected. Fortunately, Excel provides other ways to select a chart element, and it's worth your while to be familiar with them. Keep reading to see how.

Selecting with the keyboard

When a chart is active, you can hold down Ctrl on your keyboard while pressing the up-arrow and down-arrow keys to cycle among the chart's elements. Again, keep your eye on the Chart Elements control to ensure that the selected chart element is what you think it is.

- **When a chart series is selected:** Use the left-arrow and right-arrow keys to select an individual item within the series.
- **When a set of data labels is selected:** You can select a specific data label by using the left-arrow or right-arrow key.
- **When a legend is selected:** Select individual elements within the legend by using the left-arrow or right-arrow key.

Selecting with the Chart Elements control

You can find the Chart Elements control located in the Format \Rightarrow Current Selection group. This control displays the name of the currently selected chart element. It's a drop-down control, and you can use it to select a particular element in the active chart.

The Chart Elements control also appears in the Mini toolbar, which is displayed when you right-click a chart element (see Figure 19.2).

FIGURE 19.2

Using the Chart Elements control in the Mini toolbar



The Chart Elements control enables you to select only the top-level elements in the chart. To select an individual data point within a series, for example, you need to select the series and then use the navigation keys (or your mouse) to select the desired data point.

Note

When a single data point is selected, the Chart Elements control *will* display the name of the selected element even though it's not actually available for selection from the drop-down list.

TIP

If you frequently work with charts, you may want to add the Chart Elements control to your Quick Access Toolbar. That way, it will always be visible regardless of which Ribbon tab is showing. To add the control to your Quick Access Toolbar, right-click the down arrow in the Chart Elements control in the Ribbon and choose Add to Quick Access Toolbar.

Exploring the User Interface Choices for Modifying Chart Elements

You have four main ways of working with chart elements: the Format task pane, the icons that display to the right of the chart, the Ribbon, and the Mini toolbar.

Using the Format task pane

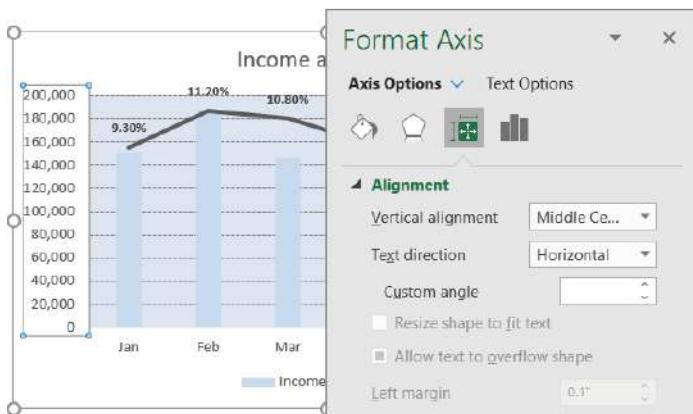
When a chart element is selected, use the element's Format task pane to format or set options for the element. Each chart element has a unique Format task pane that contains controls specific to the element (although many Format task panes have controls in common). To access the Format task pane, use any of these methods:

- Double-click the chart element.
- Right-click the chart element and then choose Format xxxx from the shortcut menu (where xxxx is the name of the element).
- Select a chart element and then choose Format ↴ Current Selection ↴ Format Selection.
- Select a chart element and press Ctrl+1.

Any of these actions displays the Format task pane from which you can make many changes to the selected chart element. For example, Figure 19.3 shows the task pane that appears when a chart's value axis is selected. The task pane is free-floating, not docked.

FIGURE 19.3

Use the Format task pane to set the properties of a selected chart element—in this case, the chart's value axis.

**TIP**

Normally, the Format task pane is docked on the right side of the window. But you can click the title and drag it anywhere that you like and resize it. To re-dock the task pane, maximize the Excel window, and then drag the task pane to the right side of the window. If you select a different chart element, the Format task pane changes to display the options appropriate for the new element.

Using the chart customization buttons

When a chart is selected, three buttons appear to the right of the chart (see Figure 19.4). These buttons, when clicked, expand to show various options. The icons are as follows:

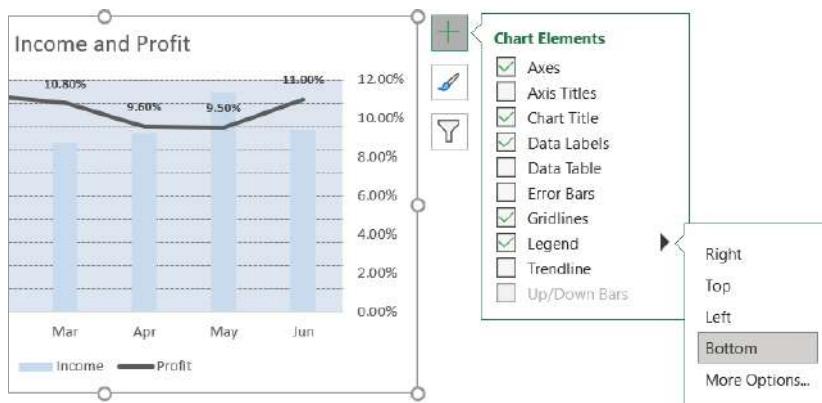
Chart Elements Use these tools to hide or display specific elements in the chart. Note that each item can be expanded to show additional options. To expand an item in the Chart Elements list, hover your mouse over the item and click the arrow that appears.

Chart Styles Use this icon to select from prebuilt chart styles or change the color scheme of the chart.

Chart Filters Use this icon to hide or display data series and specific points in a data series or hide and display categories. Some chart types do not display the Chart Filters button.

FIGURE 19.4

Chart customization buttons



Using the Ribbon

The controls in the Format tab are used to change the appearance of the selected chart element. Although the commands found on the Ribbon do *not* comprise a comprehensive set of tools for chart elements, you can modify the Shape Fill, Shape Outline, and Shape Effects settings for most elements in a chart.

Here are a few formatting actions that you can take right from the Ribbon (after you have selected your chart element):

Change the fill color: Choose Format \Rightarrow Shape Styles \Rightarrow Shape Fill and select a color.

Change the border or outline color: Choose Format \Rightarrow Shape Styles \Rightarrow Shape Outline and select a color.

Change the width and style of lines: Choose Format \Rightarrow Shape Styles \Rightarrow Shape Outline and select a weight or line style.

Change the visual effect: Choose Format \Rightarrow Shape Styles \Rightarrow Shape Effects and add one or more effects.

Change the font color: Choose Home \Rightarrow Font \Rightarrow Font Color.

Using the Mini toolbar

When you right-click an element in a chart, Excel displays a shortcut menu and the Mini toolbar. The Mini toolbar contains icons (Style, Fill, Outline) that, when clicked, display formatting options. For some chart elements, the Style icon isn't relevant, so the Mini toolbar displays the Chart Elements control (which you can use to select another chart element).

Modifying the Chart Area

The *Chart Area* is an object that contains all other elements in the chart. You can think of it as a chart's master background or container.

The only modifications that you can make to the Chart Area are cosmetic. You can change its fill color, outline, or effects such as shadows and soft edges.

NOTE

If you set the Chart Area of an embedded chart to use No Fill, the chart becomes transparent and the underlying cells are visible. In addition to specifying fill color, you can add effects such as shadow to chart elements. Figure 19.5 shows a chart where the Plot Area has been given a fill color and a shadow effect, making it appear to float.

The Chart Area element also controls all the fonts used in the chart. For example, if you want to change every font in the chart, you don't need to format each text element separately. Just select the Chart Area and then make the change from the options of the Home \Rightarrow Font group or the Format Chart Area task pane.

FIGURE 19.5

The Plot Area has been formatted to stand out from the rest of the chart.



19

Resetting Chart Element Formatting

If you go overboard formatting a chart element, you can always reset it to its original state. Just select the element and choose Format \Rightarrow Current Selection \Rightarrow Reset to Match Style. Or right-click the chart element and choose Reset to Match Style from the shortcut menu.

To reset all formatting changes in the entire chart, select the Chart Area before you issue the Reset to Match Style command.

Modifying the Plot Area

The *Plot Area* is the part of the chart that contains the actual chart. More specifically, the Plot Area is a container for the chart series.

TIP

If you set the **Shape Fill** property to **No Fill**, the Plot Area will be transparent. Therefore, the fill color applied to the Chart Area will show through.

You can move and resize the Plot Area. Select the Plot Area and then drag a border to move it. To change the size of the Plot Area, drag one of the corner handles.

Different chart types vary in the way that they respond to changes in the Plot Area dimensions. For example, you can't change the relative dimensions of the Plot Area of a pie chart or a radar chart. The Plot Area of these charts is always square. With other chart types, though, you can change the aspect ratio of the Plot Area by changing either the height or the width.

Figure 19.6 shows a chart in which the Plot Area was resized to make room for an inserted shape that contains text.

FIGURE 19.6

Reducing the size of the Plot Area makes room for the shape.



In some cases, the size of the Plot Area changes automatically when you adjust other elements of your chart. For example, if you add a legend to a chart, the size of the Plot Area may be reduced to accommodate the legend.

TIP

Changing the size and position of the Plot Area can have a dramatic effect on the overall look of your chart. When you're fine-tuning a chart, you'll probably want to experiment with various sizes and positions for the Plot Area.

Working with Titles in a Chart

A chart can have several different types of titles:

- Chart title
- Category axis title
- Value axis title
- Secondary category axis title
- Secondary value axis title
- Depth axis title (for true 3-D charts)

The number of titles that you can use depends on the chart type. For example, a pie chart supports only a chart title because it has no axes.

The easiest way to add a chart title is to use the Chart Elements button (the plus sign), which appears to the right of the chart. Activate the chart, click the Chart Elements button, and enable the Chart Title item. To specify a location, move the mouse over the Chart Title item and click the arrow. You can then specify the location for the Chart Title. Click More Options to display the Format Chart Title task pane.

The same basic procedure applies to axis titles. You have additional options to specify which axis titles you want.

After you add a title, you can replace the default text and drag the title to a different position. However, you can't change the size of a title by dragging its borders. The only way to change the size of a title is to change the font size.

TIP

The chart title or any of the axis titles can also use a cell reference. For example, you can create a link so the chart always displays the text contained in cell A1 as its title. To create a link, select the title, click the formula bar, then type an equal sign (=) and select the cell containing the data you want to see in the title and press Enter. After you create the link, the Formula bar displays the cell reference when you select the title.

19

Adding Free-Floating Text to a Chart

Text in a chart is not limited to titles. In fact, you can add free-floating text anywhere you want. To do so, activate the chart and choose Insert \Rightarrow Text \Rightarrow Text Box. Click in the chart to create the text box and enter the text. You can resize the text box, move it, change its formatting, and so on. You can also add a shape to the chart and then add text to the shape (if the shape is one that accepts text). Refer to Figure 19.6 for an example of an inserted shape with text.

Working with a Legend

A chart's *legend* consists of text and keys that identify the data series in the chart. A *key* is a small graphic that corresponds to the chart's series (one key for each series).

To add a legend to your chart, activate the chart and click the Chart Elements icon to the right of the chart. Place a check mark next to Legend. To specify a location for the legend, click the arrow next to the Legend item and choose a location (Right, Top, Left, or Bottom). After you add a legend, you can drag it to move it anywhere you like.

TIP

If you move a legend manually, you may need to adjust the size of the chart's Plot Area.

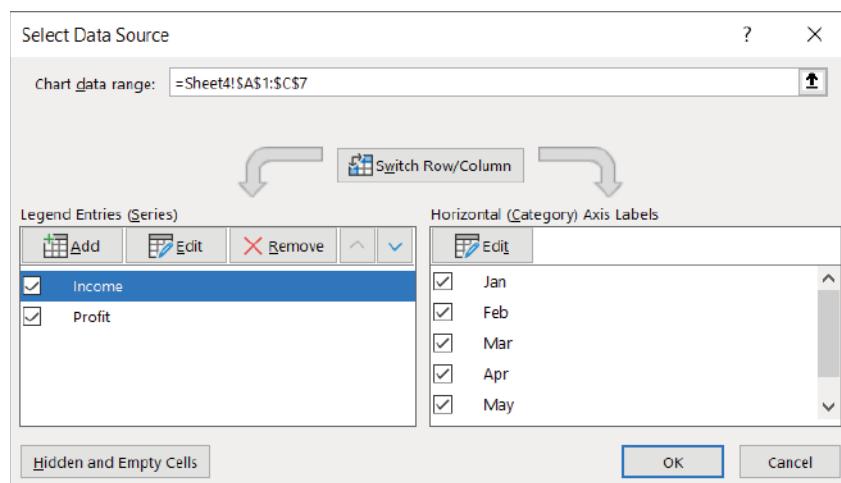
The quickest way to remove a legend is to select it and then press Delete.

You can select individual items within a legend and format them separately. For example, you may want to make the text bold to draw attention to a particular data series. To select an element in the legend, first select the legend and then click the desired element.

If you didn't include legend text when you originally selected the cells to create the chart, Excel displays Series 1, Series 2, and so on, in the legend. To add series names, choose Chart Design \Rightarrow Data \Rightarrow Select Data to display the Select Data Source dialog box (see Figure 19.7). Select the series name and click the Edit button. In the Edit Series dialog box, type the series name or enter a cell reference that contains the series name. Repeat for each series that needs to be named.

FIGURE 19.7

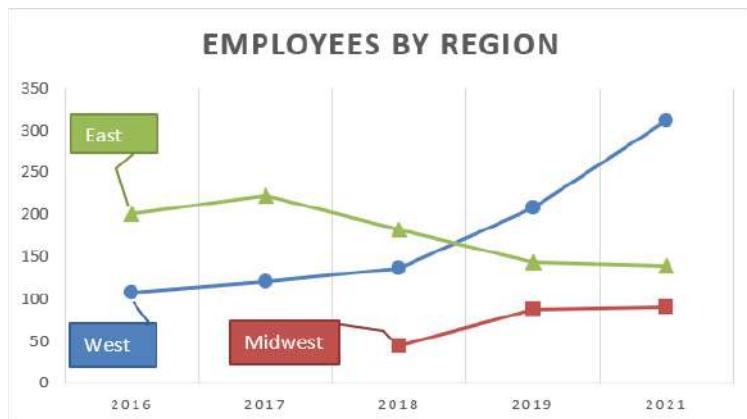
Use the Select Data Source dialog box to change the name of a data series.



In some cases, you may prefer to omit the legend and use callouts to identify the data series. Figure 19.8 shows a chart with no legend. Instead, it uses shapes to identify each series. These shapes are from the Callouts section of the Format \Rightarrow Insert Shapes gallery.

FIGURE 19.8

Using shapes as callouts in lieu of a legend



Copying Chart Formatting

You created a killer chart and spent hours customizing it. Now you need to create another one just like it but with a different set of data. What are your options? You have several choices:

Copy the formatting Create your new chart with the default formatting. Then select your original chart and choose Home \Rightarrow Clipboard \Rightarrow Copy (or press Ctrl+C). Click your new chart and choose Home \Rightarrow Clipboard \Rightarrow Paste \Rightarrow Paste Special. In the Paste Special dialog box, select the Formats option.

Copy the chart and change the data sources Press Ctrl while you click the original chart and drag. This creates an exact copy of your chart. Then choose Chart Design \Rightarrow Data \Rightarrow Select Data. In the Select Data Source dialog box, specify the data for the new chart in the Chart Data Range field.

Create a chart template Select your chart, right-click the Chart Area, and choose Save as Template from the shortcut menu. Excel prompts you for a name. When you create your next chart, use this template as the chart type. For more information about using chart templates, see “Creating Chart Templates” later in this chapter.

Working with Gridlines

Gridlines can help the viewer determine what the chart series represents numerically. Gridlines simply extend the tick marks on an axis. Some charts look better with gridlines, while others appear more cluttered. Sometimes horizontal gridlines alone are enough, although XY charts often benefit from both horizontal and vertical gridlines.

To add or remove gridlines, activate the chart, and click the Chart Elements button to the right of the chart. Place a check mark next to Gridlines. To specify the type of gridlines, click the arrow to the right of the Gridlines item.

NOTE

Each axis has two sets of gridlines: major and minor. Major units display a label. Minor units are located between the labels.

To modify the color or thickness of a set of gridlines, click one of the gridlines and use the commands from the Format \Rightarrow Shape Styles group. Alternatively, you can use the controls in the Format Major (or Format Minor) Gridlines task pane.

If gridlines seem too overpowering, consider changing them to a lighter color or use one of the dashed line options.

Modifying the Axes

Charts vary in the number of axes they use. Pie, doughnut, sunburst, and treemap charts have no axes. All 2-D charts have two axes but can have three (if you use a secondary value axis) or four (if you use a secondary category axis in an XY chart). True 3-D charts have three axes.

Excel gives you a great deal of control over these axes via the Format Axis task pane. The content of this task pane varies depending on the type of axis selected.

Modifying the value axis

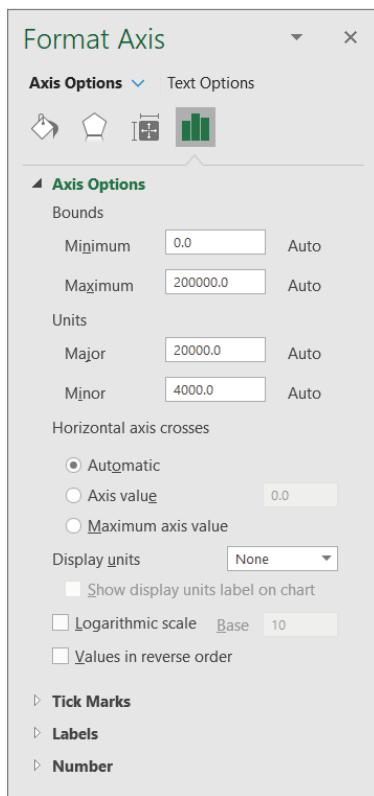
To change a value axis, right-click it and choose Format Axis. Figure 19.9 shows one panel (Axis Options) of the Format Axis task pane for a value axis. In this case, the Axis Options section is expanded, and the other three sections are contracted. The other icons along the top of this task pane deal with cosmetic and number formatting for the axis.

By default, Excel determines the minimum and maximum axis values automatically, based on the numerical range of the data. To override this automatic axis scaling, enter your own minimum and maximum values in the Bounds section. If you change these values, the word *Auto* changes to a Reset button. Click Reset to revert to automatic axis scaling.

Excel also adjusts the major and minor axis units automatically. Again, you can override Excel's choice and specify different units.

FIGURE 19.9

The Format Axis task pane for a value axis



Adjusting the bounds of a value axis can affect the chart's appearance. In some cases, manipulating the scale can present a false picture of the data. Figure 19.10 shows two line charts that depict the same data. The chart on the left uses Excel's default (Auto) axis bounds values. In the chart on the right, the Minimum bound value was set to 20,000, and the Maximum bound value was set to 400,000. The first chart makes the differences in the data seem more prominent. The second chart gives the impression that there isn't much change over time.

The actual scale you use depends on the situation. There are no hard-and-fast rules regarding setting scale values, except that you shouldn't misrepresent data by manipulating the chart to prove a point that doesn't exist.

TIP

If you're preparing several charts that use similarly scaled data, keeping the bounds the same is a good idea so that the charts can be compared more easily.

Part III: Creating Charts and Other Visualizations

FIGURE 19.10

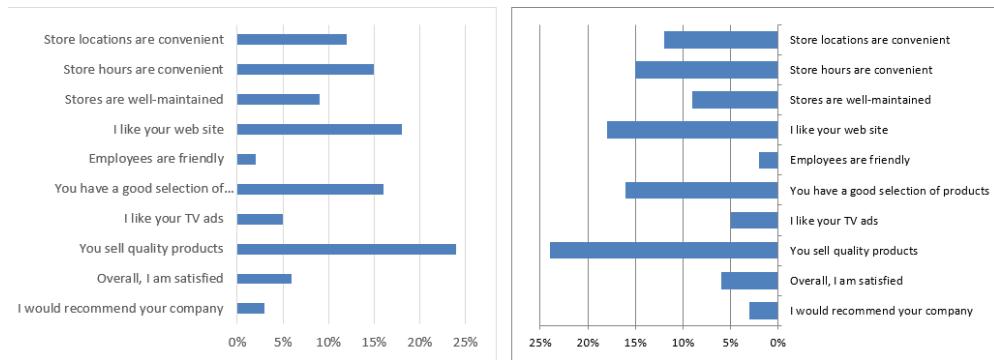
These two charts show the same data but use different value axis bounds.



Another option in the Format Axis task pane is Values in Reverse Order. The left chart in Figure 19.11 uses default axis settings. The right chart uses the Values in Reverse Order option, which reverses the scale's direction. Notice that the category axis is on the right. If you would prefer that the category axis labels remain in their original position when reversing an axis, select the Maximum Axis Value option for the Horizontal Axis Crosses setting.

FIGURE 19.11

The right chart uses the Values in Reverse Order option.



If the values to be plotted cover a large numerical range, you may want to use a logarithmic scale for the value axis. A log scale is most often used for scientific applications. Figure 19.12 shows two charts. The top chart uses a standard scale, and the bottom chart uses a logarithmic scale.

Note

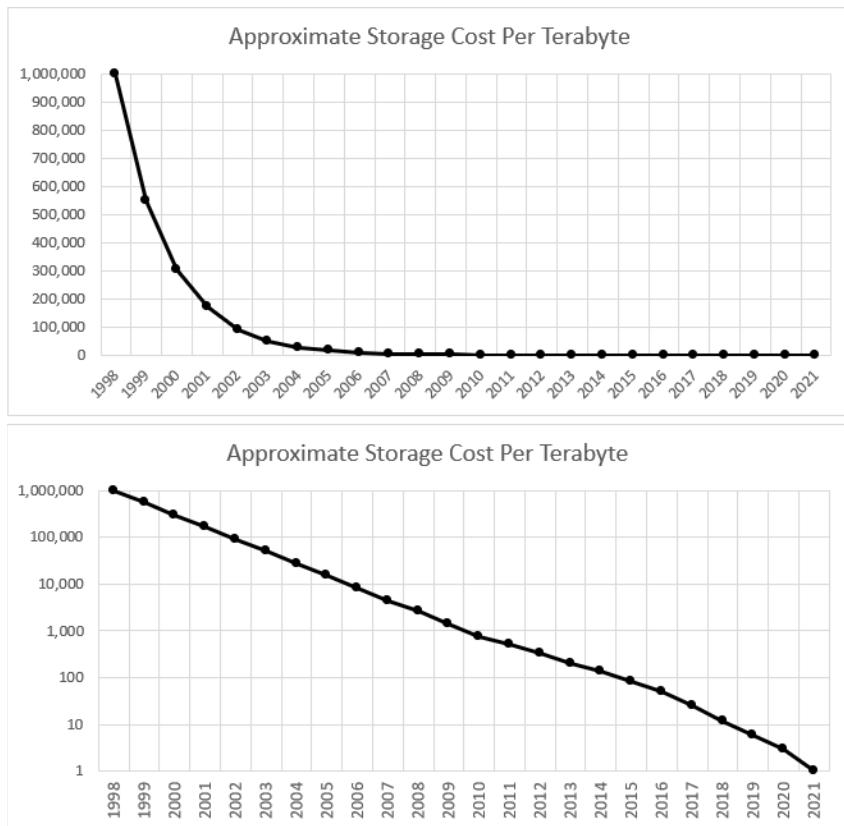
For a logarithmic scale, the Base setting is 10, so each scale value in the chart is 10 times greater than the one below it. Increasing the major unit to 100 results in a scale in which each tick mark value is 100 times greater than the one below it. You can specify a base value between 2 and 1,000.



This workbook, `log_scale.xlsx`, is available on this book's website at www.wiley.com/go/excel365bible.

FIGURE 19.12

These charts display the same data, but the bottom chart uses a logarithmic scale.

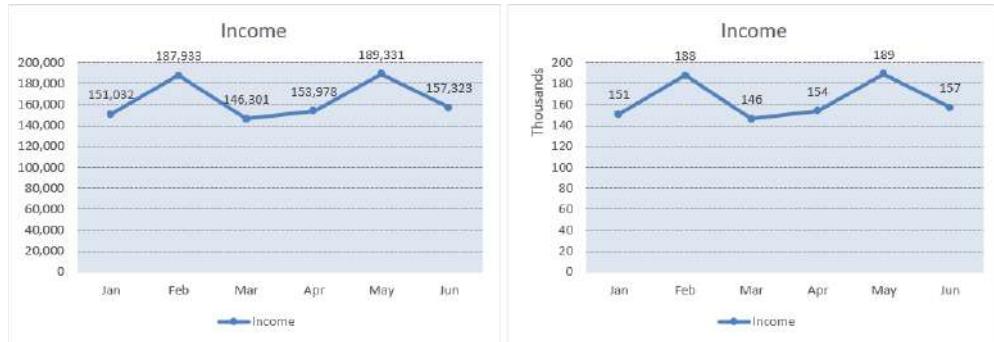


Part III: Creating Charts and Other Visualizations

If your chart uses large numbers, you may want to change the Display Units setting. Figure 19.13 shows a chart (left) that uses large numbers. In the chart on the right, the Display Units setting is set to Thousands, with the option Show Display Units Label on Chart selected.

FIGURE 19.13

The chart on the right uses display units of thousands.



To adjust the tick marks displayed on an axis, click the Tick Marks section of the Format Axis task pane to expand that section. The Major Type and Minor Type Tick Marks options control the way the tick marks are displayed. *Major Type tick marks* are the axis tick marks that normally have labels next to them. *Minor Type tick marks* fall between the major tick marks.

If you expand the Labels section, you can position the axis labels at three distinct locations: Next to Axis, High, and Low.

Refer to Figure 19.14. When you combine these settings with the Horizontal Axis Crosses option, you have a great deal of flexibility.

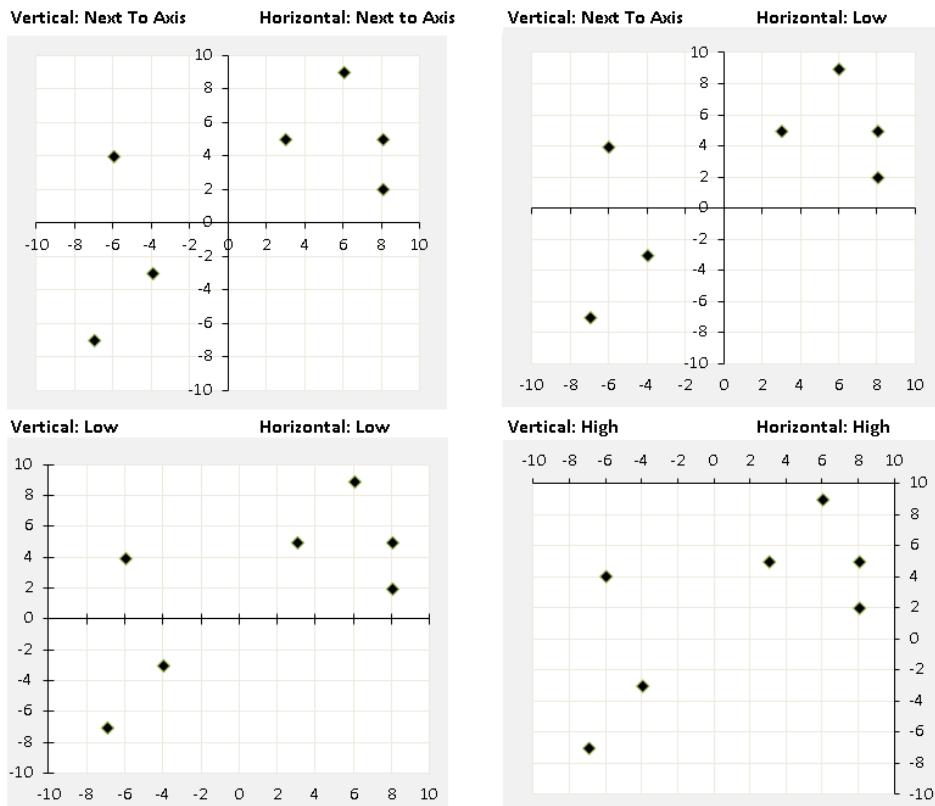
The last section of the task pane, Number, lets you specify the number formatting for the value axis. Normally, the number formatting is linked to the source data, but you can override that.

Modifying the category axis

Figure 19.15 shows part of the Axis Options section of the Format Axis task pane when a category axis is selected. Some options are the same as those for a value axis.

FIGURE 19.14

Excel offers a great deal of flexibility in how you can display axis labels and crossing points.



An important setting is **Axis Type: Text or Date**. When you create a chart, Excel recognizes whether your category axis contains date or time values. If it does, it uses a Date category axis. Figure 19.16 shows a simple example. Column A contains dates, and column B contains the values plotted in the column chart. The data consists of values for only 10 dates, yet Excel created the chart with 30 intervals on the category axis. It recognized that the category axis values were dates and created an equal-interval scale.

Part III: Creating Charts and Other Visualizations

FIGURE 19.15

Some of the options available for a category axis

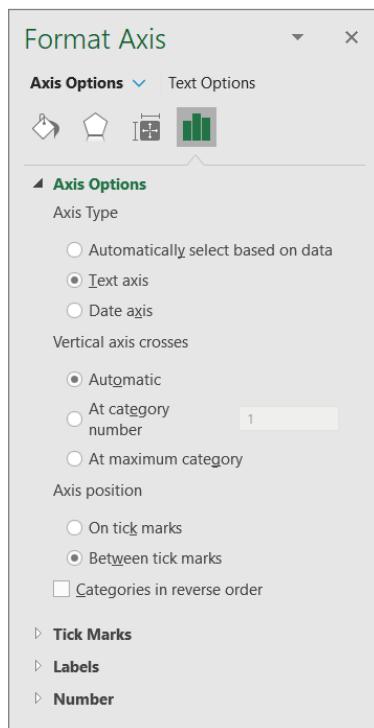
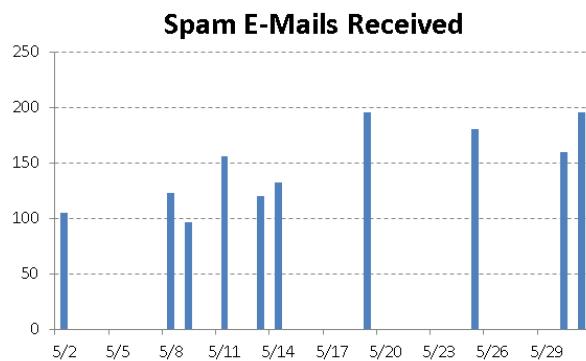


FIGURE 19.16

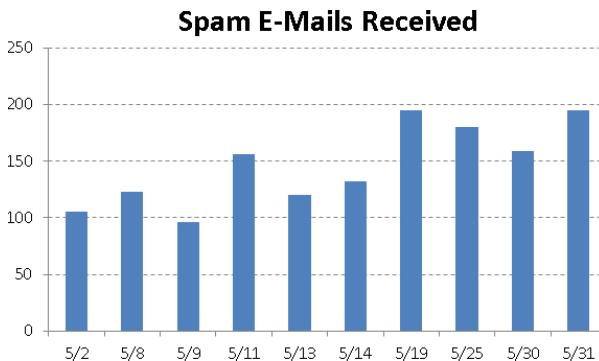
Excel recognizes dates and creates a time-based category axis.



You can override Excel's decision to use a Date category axis by choosing the Text Axis option for Axis Type. Figure 19.17 shows the chart after making this change. In this case, using a time-based category axis (as shown in Figure 19.16) presents a much truer picture of the data.

FIGURE 19.17

Overriding the Excel time-based category axis



Excel chooses the way to orient the category labels, but you can override its choice.

Figure 19.18 shows an XY chart with time labels. Because of the lengthy category labels, Excel displays the text at an angle. If you make the chart wider, the labels will then appear horizontally. You can also adjust the labels using the Alignment controls in the Size & Properties section of the Format Axis task pane.

Keep in mind that category axis labels can consist of more than one column. Figure 19.19 shows a chart that displays three columns of text for the category axis. We selected the range A1 : E10 and created a column chart, and Excel figured out the category axis.

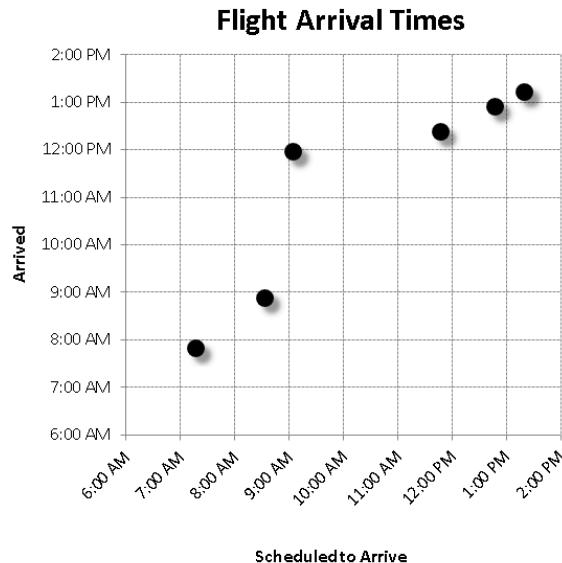
Don't Be Afraid to Experiment (But on a Copy)

We'll let you in on a secret: the key to mastering charts in Excel is experimentation, otherwise known as *trial and error*. Excel's charting options can be overwhelming, even to experienced users. This book doesn't even pretend to cover all of the charting features and options. Your job, as a potential charting master, is to dig deep and try the various options in your charts. With a bit of creativity, you can create original-looking charts.

After you create a basic chart, make a copy of the chart for your experimentation. That way, if you mess it up, you can always revert to the original and start again. To make a copy of an embedded chart, click the chart and press Ctrl+C. Then activate a cell and press Ctrl+V. To make a copy of a chart sheet, press Ctrl while you click the sheet tab and then drag it to a new location among the other tabs.

FIGURE 19.18

Excel determines the way to display category axis labels.



Working with Data Series

Every chart consists of one or more data series. This data translates into chart columns, bars, lines, pie slices, and so on. The following sections discuss some common operations that involve a chart's data series.

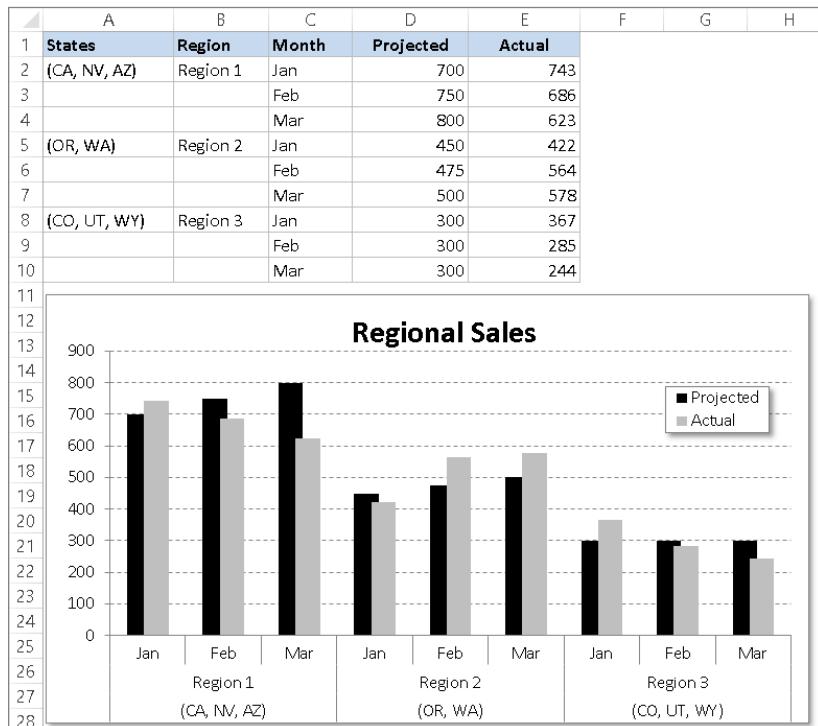
When you select a data series in a chart, Excel does the following:

- Displays the series name in the Chart Elements control (located in the Format \leftrightarrow Current Selection group)
- Displays the SERIES formula in the Formula bar
- Highlights the cells used for the selected series by outlining them in color

You can make changes to a data series by using the options on the Ribbon or from the Format Data Series task pane. This task pane varies depending on the type of data series on which you're working (column, line, pie, and so on).

FIGURE 19.19

This chart uses three columns of text for the category axis labels.



CAUTION

If the Format task pane isn't already displayed, the easiest way to display the Format Data Series task pane is to double-click the chart series. Be careful, however: if a data series is already selected, double-clicking brings up the Format Data Point task pane. Changes that you make affect only one point in the data series. To edit the entire series, make sure that a chart element other than the data series is selected before you double-click the data series. Or just press Ctrl+1 to display the task pane.

Deleting or hiding a data series

To delete a data series in a chart, select the data series and press Delete. The data series disappears from the chart. The data in the worksheet, of course, remains intact.

Note

You can delete all data series from a chart. If you do so, the chart appears empty. It retains its settings, however. Therefore, you can add a data series to an empty chart, and it again looks like a chart.

To hide a data series temporarily, activate a chart and click the Chart Filters button on the right. Remove the check mark from the data series that you want to hide, click Apply, and that data series is hidden—but it's still associated with the chart, so you can unhide it later. You can't hide all the series, though. At least one must be visible. The Chart Filters button also lets you hide individual points in a series. Note that the new chart types introduced in Excel 2016 and Excel 2019 do not display a Chart Filters button.

Adding a new data series to a chart

If you want to add another data series to an existing chart, one approach is to re-create the chart and include the new data series. However, adding the data to the existing chart is usually easier, and your chart retains any customization that you've made.

Excel provides three ways to add a new data series to a chart:

- **Use the Select Data Source dialog box:** Activate the chart and choose Chart Design ➔ Data ➔ Select Data. In the Select Data Source dialog box, click the Add button and Excel displays the Edit Series dialog box. Specify the series name (as a cell reference or text) and the range that contains the series values. The Select Data Source dialog box is also accessible from the shortcut menu displayed by right-clicking many elements in a chart.
- **Drag the range outline:** If the data series to be added is contiguous with other data in the chart, you can click the Chart Area in the chart. Excel highlights and outlines the data in the worksheet. Click one of the corners of the outline in the worksheet and drag to highlight the new data. This method works only for embedded charts.
- **Copy and paste:** Select the range to add and press Ctrl+C to copy it to the Clipboard. Then activate the chart and press Ctrl+V to paste the data into the chart.

Tip

If the chart was originally made from data in a table (created via Insert ➔ Tables ➔ Table), the chart is updated automatically when you add new rows or columns to the table (or remove rows or columns). If you have a chart that is updated frequently with new data, you can save time and effort by creating the chart from data in a table.

Changing data used by a series

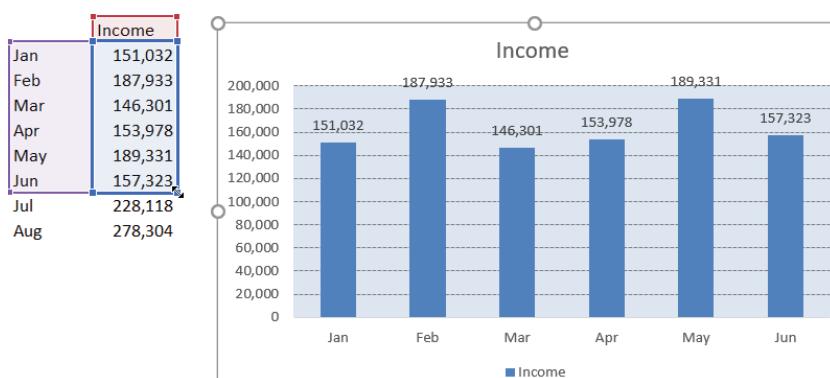
You may find that you need to modify the range that defines a data series. For example, say you need to add new data points or remove old ones from the data set. The following sections describe several ways to change the range used by a data series.

Changing the data range by dragging the range outline

If you have an embedded chart, the easiest way to change the data range for a data series is to drag the range outline. When you select a series in a chart, Excel outlines the data range used by that series. You can drag the small dot in the lower-right corner of the range outline to extend or contract the data series. In Figure 19.20, the range outline will be dragged to include two additional data points.

FIGURE 19.20

Changing a chart's data series by dragging the range outline



You can also click and drag one of the sides of the outline to move the outline to a different range of cells.

In some cases, you'll also need to adjust the range that contains the category labels. The labels are also outlined, and you can drag the outline to expand or contract the range of labels used in the chart.

If your chart is on a chart sheet, you need to use one of the two methods described next. These methods also work with embedded charts.

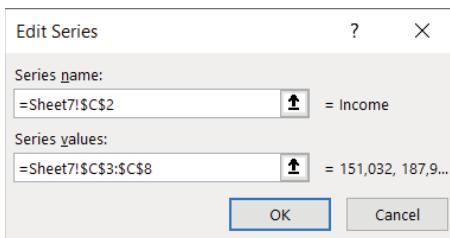
Using the Edit Series dialog box

Another way to update the chart to reflect a different data range is to use the Edit Series dialog box. A quick way to display this dialog box is to right-click the series in the chart and then choose Select Data from the shortcut menu. The Select Source Data dialog box appears. Select the data series in the list and click Edit to display the Edit Series dialog box, as shown in Figure 19.21.

Part III: Creating Charts and Other Visualizations

FIGURE 19.21

The Edit Series dialog box



You can change the entire data range used by the chart by adjusting the range references in the Chart Data Range field in the Select Data Source dialog box. Or you can select a series from the list and click Edit to modify the selected series.

Editing the Series formula

Every data series in a chart has an associated SERIES formula, which appears in the Formula bar when you select a data series in a chart. If you understand the way a SERIES formula is constructed, you can edit the range references in the SERIES formula directly to change the data used by the chart.

Note

The SERIES formula is not a real formula. In other words, you can't use it in a cell, and you can't use worksheet functions within the SERIES formula. You can, however, edit the arguments in the SERIES formula.

A SERIES formula has the following syntax:

```
=SERIES(series_name,category_labels,values,order,sizes)
```

The arguments that you can use in the SERIES formula include the following:

- **series_name (optional):** A reference to the cell that contains the series name used in the legend. If the chart has only one series, the name argument is used as the title. This argument can also consist of text in quotation marks. If this argument is omitted, Excel creates a default series name (for example, Series 1).
- **category_labels (optional):** A reference to the range that contains the labels for the category axis. If this argument is omitted, Excel uses consecutive integers beginning with 1. For XY charts, this argument specifies the X values. A noncontiguous range reference is also valid. The ranges' addresses are separated by commas and enclosed in parentheses. The argument could also consist of an array of comma-separated values (or text in quotation marks) enclosed in curly brackets.
- **values (required):** A reference to the range that contains the values for the series. For XY charts, this argument specifies the Y values. A noncontiguous range reference is also valid. The range addresses are separated by a comma and enclosed in

parentheses. The argument could also consist of an array of comma-separated values enclosed in curly brackets.

- **order (required):** An integer that specifies the plotting order of the series. This argument is relevant only if the chart has more than one series. Using a reference to a cell is not allowed.
- **sizes (only for bubble charts):** A reference to the range that contains the values for the size of the bubbles in a bubble chart. A noncontiguous range reference is also valid. The range addresses are separated by commas and enclosed in parentheses. The argument can also consist of an array of values enclosed in curly brackets.

Range references in a **SERIES** formula are always absolute (contain two dollar signs), and they always include the sheet name. For example, note the following:

```
=SERIES (Sheet1!$B$1,,Sheet1!$B$2:$B$7,1)
```

TIP

You can substitute range names for the range references. If you do so, Excel changes the reference in the **SERIES** formula to include the workbook name (if it's a workbook-level name) or to include the worksheet name (if it's a sheet-level name). For example, if you use a workbook-level range named **MyData** (in a workbook named **budget.xlsx**), the **SERIES** formula looks like this:

```
=SERIES (Sheet1!$B$1,,budget.xlsx!MyData,1)
```



For more information about named ranges, see Chapter 4, “Working with Excel Ranges and Tables.”

Displaying data labels in a chart

Sometimes, you may want your chart to display the actual numerical value for each data point. To add labels to data series in a chart, select the series and click the Chart Elements button on the right side of the chart. Place a check mark next to Data Labels. Click the arrow next to the Data Labels item to specify the position for the labels.

To add data labels for all series, use the same procedure, but start by selecting something other than a data series.

Figure 19.22 shows three minimalist charts with data labels.

To change the type of information that appears in data labels, select the data labels for a series and use the Format Data Labels task pane. (If the task pane isn't visible, press **Ctrl+1**.) Then use the Label Options section to customize the data labels. For example, you can include the series name and the category name along with the value.

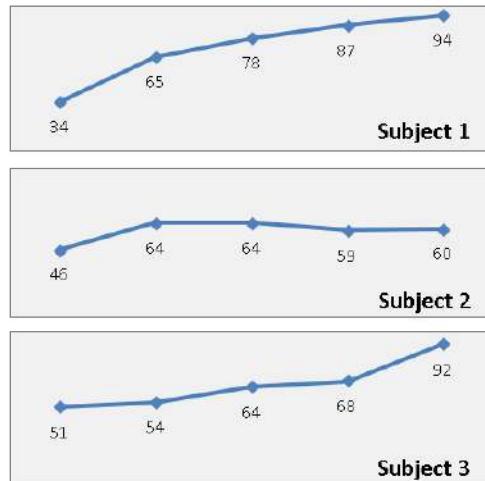
The data labels are linked to the worksheet, so if your data changes, the labels also change. If you want to override the data label with other text, select the label and enter the new text.

Part III: Creating Charts and Other Visualizations

FIGURE 19.22

These charts use data labels and don't display axes.

	Subject 1	Subject 2	Subject 3
1	34	46	51
2	65	64	54
3	78	64	64
4	87	59	68
5	94	60	92

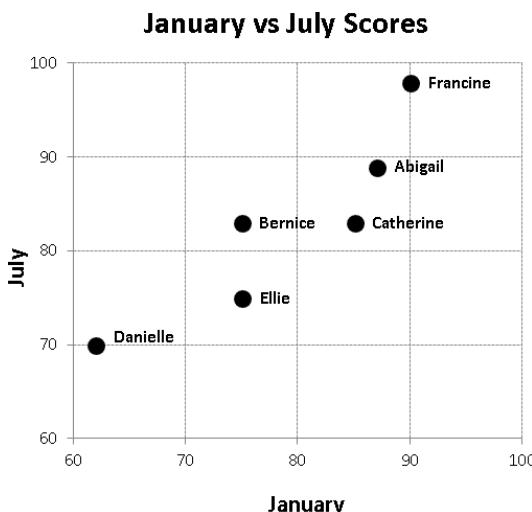


In some situations, you may need to use other data labels for a series. In the Format Data Labels task pane, select Value From Cells (in the Label Options section), and then in the Data Label Range dialog box, specify the range that contains the data point labels.

Figure 19.23 shows an XY chart that uses data labels stored in a range.

FIGURE 19.23

Data labels linked to text in an arbitrary range



TIP

Often, the data labels aren't positioned properly. For example, a label may be obscured by another data point or another label. If you select an individual data label, you can drag the label to a better location. To select an individual data label, click once to select them all and then click the single data label.

Handling missing data

Sometimes, data that you're charting may be missing one or more data points. As shown in Figure 19.24, Excel offers three ways to handle the missing data:

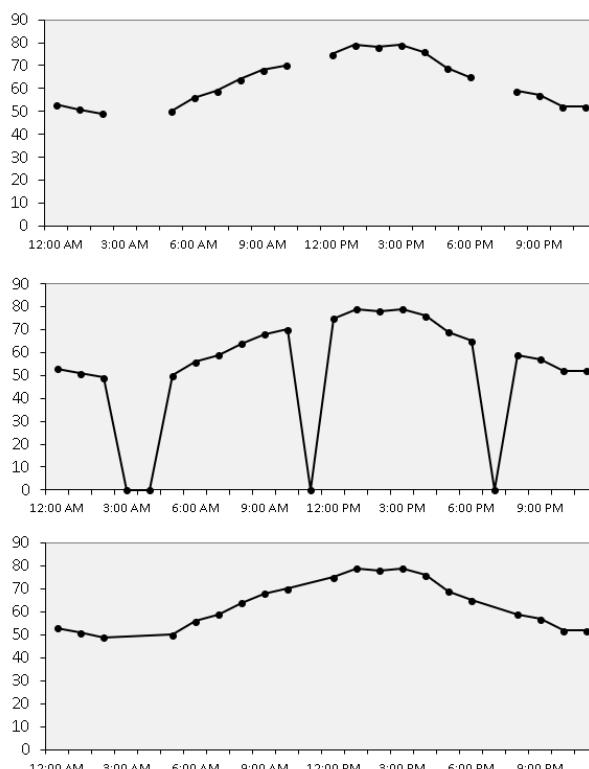
Gaps Missing data is simply ignored, and the data series will have a gap. This is the default.

Zero Missing data is treated as zero.

Connect Data Points with Line Missing data is interpolated, calculated by using data on either side of the missing point(s). This option is available for line charts, area charts, and XY charts only.

FIGURE 19.24

Three options for dealing with missing data



To specify how to deal with missing data for a chart, choose Chart Design \Rightarrow Data \Rightarrow Select Data. In the Select Data Source dialog box, click the Hidden and Empty Cells button. Excel displays its Hidden and Empty Cell Settings dialog box. Make your choice in the dialog box. The option that you choose applies to the entire chart, and you can't set a different option for different series in the same chart.

TIP

Normally, a chart doesn't display data that's in a hidden row or column. You can use the Hidden and Empty Cell Settings dialog box to force a chart to use hidden data, though.

Adding error bars

Some chart types support error bars. *Error bars* often are used to indicate "plus or minus" information that reflects uncertainty in the data. Error bars are appropriate for area, bar, column, line, and XY charts only.

To add error bars, select a data series and then click the Chart Elements icon to the right of the chart. Add a check mark next to Error Bars. Click the arrow next to the Error Bars item to specify the type of error bars. If necessary, you can fine-tune the error bar settings from the Format Error Bars task pane. The types of error bars are as follows:

Fixed value The error bars are fixed by an amount that you specify.

Percentage The error bars are a percentage of each value.

Standard deviation(s) The error bars are in the number of standard deviation units that you specify. (Excel calculates the standard deviation of the data series.)

Standard error The error bars are one standard error unit. (Excel calculates the standard error of the data series.)

Custom You set the error bar units for the upper or lower error bars. You can enter either a value or a range reference that holds the error values that you want to plot as error bars.

The chart shown in Figure 19.25 displays error bars based on percentage.

TIP

A data series in an XY chart can have error bars for both the X values and the Y values.



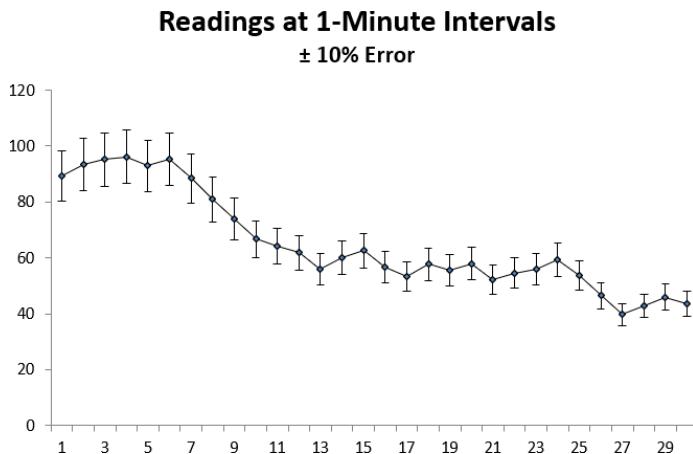
A workbook with several other error bar examples is available on this book's website at www.wiley.com/go/excel365bible. The filename is `error bars example.xlsx`.

Adding a trendline

When you're plotting data over time, you may want to display a trendline that describes the data. A *trendline* points out general trends in your data. In some cases, you can forecast future data with trendlines.

FIGURE 19.25

This line chart series displays error bars based on percentage.

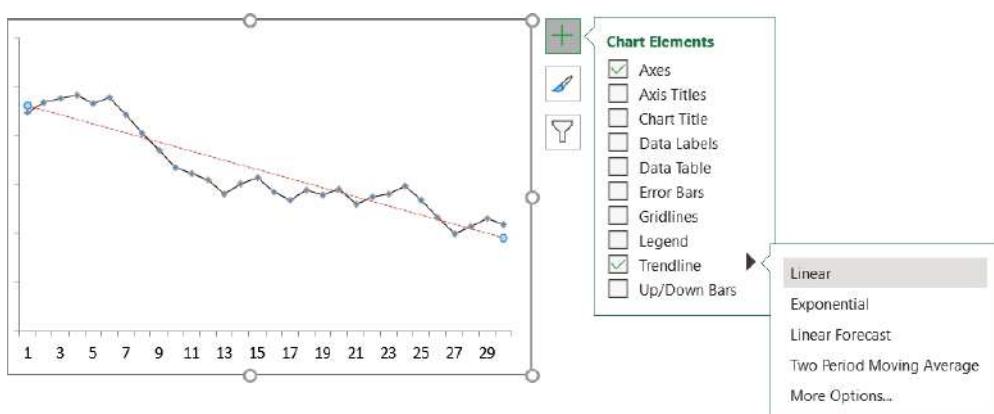


To add a trendline, select the data series and click the Chart Elements button to the right of the chart. Place a check mark next to Trendline. Alternatively, you can right-click the data series and choose Add Trendline.

To specify the type of trendline, click the arrow to the right of the Trendline item. The type of trendline that you choose depends on your data. Linear trends (see Figure 19.26) are most commonly used and the easiest to apply.

FIGURE 19.26

Applying a trend line

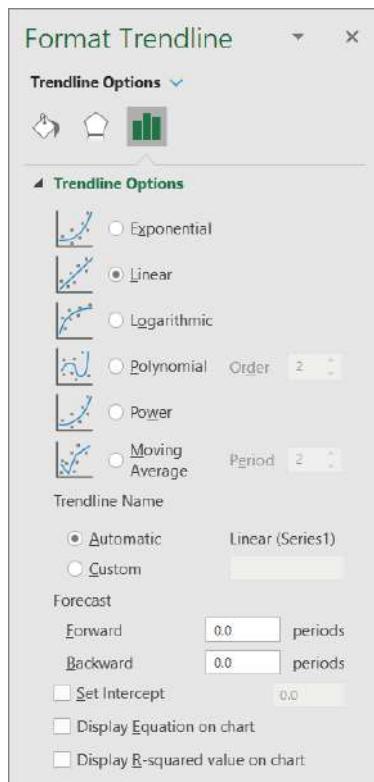


Part III: Creating Charts and Other Visualizations

Right-click the trendline, and then choose Format Trendline. This will activate the Format Trendline task pane illustrated in Figure 19.27. Here you can apply more advanced trendlines such as logarithmic, exponential, and moving average trendlines.

FIGURE 19.27

Use the Format Trendline task pane to customize your trendline.



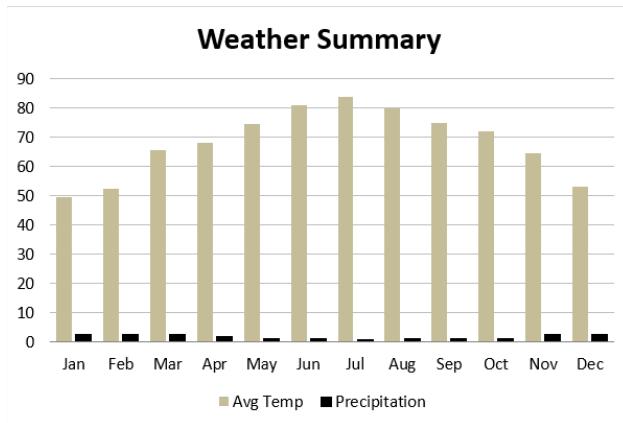
Creating combination charts

A *combination chart* is a single chart that consists of series that use different chart types. A combination chart may also include a second value axis. For example, you may have a chart that shows both columns and lines, with two value axes. The value axis for the columns is on the left, and the value axis for the line is on the right. A combination chart requires at least two data series.

Figure 19.28 shows a column chart with two data series. The values for the Precipitation series are so low that they're barely visible on the Value Axis scale. This is a good candidate for a combination chart.

FIGURE 19.28

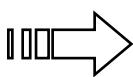
The Precipitation series is barely visible.



The following steps describe how to convert this chart into a combination chart (column and line) that uses a second value axis:

1. **Activate the chart, and choose Chart Design \Rightarrow Type \Rightarrow Change Chart Type.** The Change Chart Type dialog box appears.
2. **Select the All Charts tab.**
3. **In the list of chart types, click Combo.**
4. **For the Avg Temp series, specify Clustered Column as the chart type.**
5. **For the Precipitation series, specify Line as the chart type and click the Secondary Axis check box.**
6. **Click OK to apply these changes to the chart.**

Figure 19.29 shows the Change Chart Type dialog box after specifying the parameters for each series.



This workbook is available on this book's website at www.wiley.com/go/excel365bible. The file-name is weather combination chart.xlsx.

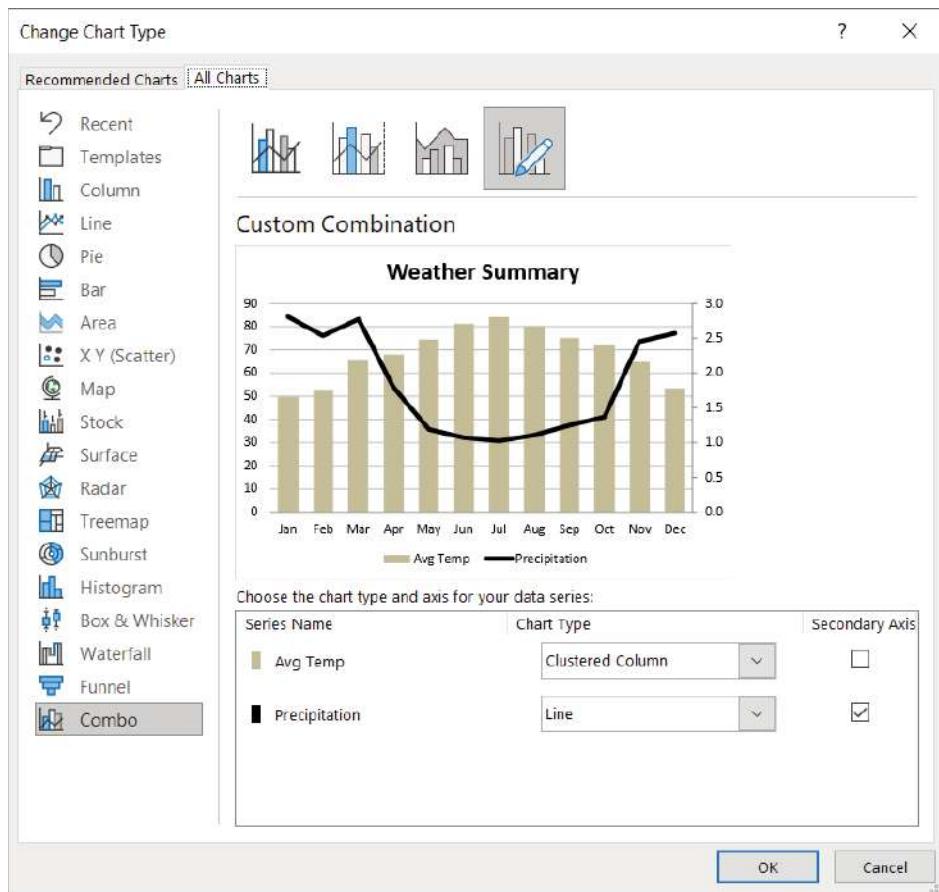
Note

In some cases, you can't combine chart types. For example, you can't create a combination chart that involves a bubble chart or a 3-D chart. In the Change Chart Type dialog box, Excel displays only the chart types that can be used.

Part III: Creating Charts and Other Visualizations

FIGURE 19.29

Using the Change Chart Type dialog box to convert a chart into a combination chart



Displaying a data table

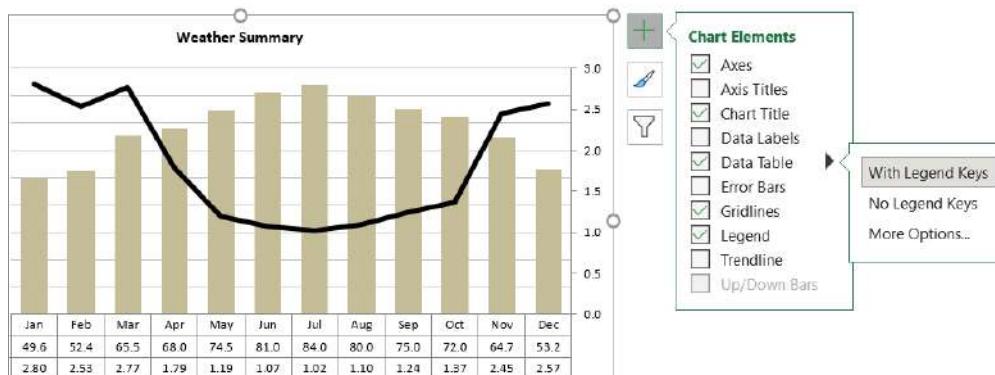
There may be situations where it's valuable to show all of the data values along with the plotted data points. However, adding data labels can inundate your audience with a bevy of numbers that muddle the chart.

Instead of using data labels, you can attach a data table to your Excel chart. A data table allows you to see the data values for each plotted data point beneath the chart, showing the data without overcrowding the chart itself.

To add a data table to a chart, activate the chart and click the Chart Elements button to the right of the chart. Place a check mark next to Data Table. Click the arrow to the right of the Data Table item for a few options. Figure 19.30 shows a combination chart that includes a data table.

FIGURE 19.30

This combination chart includes a data table that displays the values of the data points.



Not all chart types support data tables. If the Data Table option isn't available, this means the chart doesn't support this feature.

TIP

Using a data table is probably best suited for charts on chart sheets. If you need to show the data used in an embedded chart, you can do so using data in cells, which gives you more formatting flexibility.

Creating Chart Templates

This section describes how to create custom chart templates. A template includes customized chart formatting and settings. When you create a new chart, you can choose to use your template rather than a built-in chart type.

If you find that you're continually customizing your charts in the same way, you can probably save some time by creating a template. Or, if you create lots of combination charts, you can create a combination chart template and avoid making the manual adjustments required for a combination chart.

Part III: Creating Charts and Other Visualizations

To create a chart template, follow these steps:

1. **Create a chart to serve as the basis for your template.** The data you use for this chart isn't critical, but for best results it should be typical of the data that you'll eventually be plotting with your custom chart type.
2. **Apply any formatting and customizations that you like.** This step determines the appearance of the charts created from the template.
3. **Activate the chart, right-click the Chart Area or the Plot Area and choose Save as Template from the shortcut menu.** The Save Chart Template dialog box appears.
4. **Provide a name for the template and click Save.** Make sure you don't change the proposed directory for the file.

To create a chart based on a template, follow these steps:

1. **Select the data to be used in the chart.**
2. **Choose Insert \Rightarrow Charts \Rightarrow Recommended Charts.** The Insert Chart dialog box appears.
3. **Select the All Charts tab.**
4. **From the left side of the Insert Chart dialog box, select Templates.** Excel displays a thumbnail for each custom template that has been created.
5. **Click the thumbnail that represents the template you want to use and then click OK.** Excel creates the chart based on the template you selected.

NOTE

You can also apply a template to an existing chart. Select the chart and choose Chart Design \Rightarrow Type \Rightarrow Change Chart Type to display the Change Chart Type dialog box, which is identical to the Insert Chart dialog box.

Creating Sparkline Graphics

IN THIS CHAPTER

Introducing the Sparkline graphics feature

Adding Sparklines to a worksheet

Customizing Sparklines

Making a Sparkline display only the most recent data

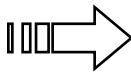
A *Sparkline* is a small chart that's displayed in a single cell. A Sparkline allows you to spot time-based trends or variations quickly in data. Because they're so compact, Sparklines are almost always used in a group.

Although Sparklines look like miniature charts (and can sometimes take the place of a chart), this feature is separate from the charting feature. For example, charts are placed on a worksheet's draw layer, and a single chart can display several series of data. A Sparkline is displayed inside a cell and displays only one series of data.



See Chapter 18, "Getting Started with Excel Charts," and Chapter 19, "Using Advanced Charting Techniques," for information about *real* charts.

This chapter introduces Sparklines and presents examples that demonstrate how to use them in your worksheets.



All examples in this chapter are available on this book's website at www.wiley.com/go/excel365bible. The filename is sparkline_examples.xlsx.

Sparkline Types

Excel supports three types of Sparklines. Figure 20.1 shows examples of each, displayed in the last column. Each Sparkline depicts the six data points to the left.

Part III: Creating Charts and Other Visualizations

FIGURE 20.1

Three groups of Sparklines

Line Sparklines

Fund Number	Jan	Feb	Mar	Apr	May	Jun	Sparklines
A-13	103.98	98.92	88.12	86.34	75.58	71.2	
C-09	212.74	218.7	202.18	198.56	190.12	181.74	
K-88	75.74	73.68	69.86	60.34	64.92	59.46	
W-91	91.78	95.44	98.1	99.46	98.68	105.86	
M-03	324.48	309.14	313.1	287.82	276.24	260.9	

Column Sparklines

Fund Number	Jan	Feb	Mar	Apr	May	Jun	Sparklines
A-13	103.98	98.92	88.12	86.34	75.58	71.2	
C-09	212.74	218.7	202.18	198.56	190.12	181.74	
K-88	75.74	73.68	69.86	60.34	64.92	59.46	
W-91	91.78	95.44	98.1	99.46	98.68	105.86	
M-03	324.48	309.14	313.1	287.82	276.24	260.9	

Win/Loss Sparklines

Fund Number	Jan	Feb	Mar	Apr	May	Jun	Sparklines
A-13	#N/A	-5.06	-10.8	-1.78	-10.76	-4.38	
C-09	#N/A	5.96	-16.52	-3.62	-8.44	-8.38	
K-88	#N/A	-2.06	-3.82	-9.52	4.58	-5.46	
W-91	#N/A	3.66	2.66	1.36	-0.78	7.18	
M-03	#N/A	-15.34	3.96	-25.28	-11.58	-15.34	

Line Similar to a line chart. As an option, the line can appear with a marker for each data point. The first group in Figure 20.1 shows line Sparklines with markers. A quick glance reveals that, with the exception of Fund Number W-91, the funds have been losing value over the six-month period.

Column Similar to a column chart. The second group in Figure 20.1 shows the same data displayed with column Sparklines.

Win/Loss A “binary” type chart that displays each data point as a high block or a low block. The third group shows Win/Loss Sparklines. Notice that the data is different. Each cell displays the change from the previous month. In the Sparkline, each data point is depicted as a high block (win) or a low block (loss). In this example, a positive change from the previous month is a win, and a negative change from the previous month is a loss.

Why Sparklines?

The term *Sparkline* was coined by Edward Tufte in his book *Beautiful Evidence* (Graphics Press, 2006). He describes it as follows:

Sparklines: Intense, simple, word-sized graphics.

In the case of Excel, Sparklines are cell-sized graphics. As you see in this chapter, Sparklines aren't limited to lines.

Creating Sparklines

Figure 20.2 shows some data to be summarized with Sparklines.

FIGURE 20.2

Data to be summarized with Sparklines

	A	B	C	D	E	F	G	H	I	J	K	L	M
1	Average Monthly Precipitation (Inches)												
2		Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct	Nov	Dec
4	ASHEVILLE, NC	4.06	3.83	4.59	3.5	4.41	4.38	3.87	4.3	3.72	3.17	3.82	3.39
5	BAKERSFIELD, CA	1.18	1.21	1.41	0.45	0.24	0.12	0	0.08	0.15	0.3	0.59	0.76
6	BATON ROUGE, LA	6.19	5.1	5.07	5.56	5.34	5.33	5.96	5.86	4.84	3.81	4.76	5.26
7	BILLINGS, MT	0.81	0.57	1.12	1.74	2.48	1.89	1.28	0.85	1.34	1.26	0.75	0.67
8	DAYTONA BEACH, FL	3.13	2.74	3.84	2.54	3.26	5.69	5.17	6.09	6.61	4.48	3.03	2.71
9	EUGENE, OR	7.65	6.35	5.8	3.66	2.66	1.53	0.64	0.99	1.54	3.35	8.44	8.29
10	HONOLULU, HI	2.73	2.35	1.89	1.11	0.78	0.43	0.5	0.46	0.74	2.18	2.26	2.85
11	ST. LOUIS, MO	2.14	2.28	3.6	3.69	4.11	3.76	3.9	2.98	2.96	2.76	3.71	2.86
12	TUCSON, AZ	0.99	0.88	0.81	0.28	0.24	0.24	2.07	2.3	1.45	1.21	0.67	1.03

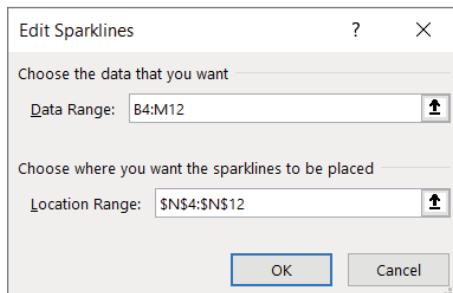
To create Sparkline graphics, follow these steps:

1. **Select the data that will be depicted (data only, not column or row headings); if you're creating multiple Sparklines, select all the data.** In this example, start by selecting B4 : M12.
2. **With the data selected, choose Insert ↴ Sparklines, and click one of the three Sparkline types: Line, Column, or Win/Loss.** The Create Sparklines dialog box, shown in Figure 20.3, appears.

Part III: Creating Charts and Other Visualizations

FIGURE 20.3

Use the Create Sparklines dialog box to specify the data range and the location for the Sparkline graphics.



3. **Specify the location for the Sparklines.** Typically, you'll put the Sparklines next to the data, but that's not a requirement. Most of the time, you'll use an empty range to hold the Sparklines. However, Excel doesn't prevent you from inserting Sparklines into cells that already contain data. The Sparkline location that you specify must match the source data in terms of number of rows or number of columns. For this example, specify N4 : N12 as the location range.
4. **Click OK.** Excel creates the Sparkline graphics of the type you specified.

The Sparklines are linked to the data, so if you change any of the values in the data range, the Sparkline graphic will update. Often, you'll want to increase the column width or row height to improve the readability of the Sparklines.

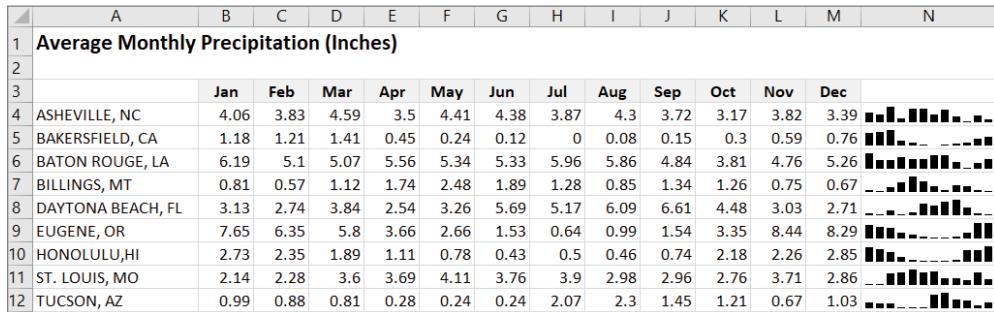
TIP

Most of the time, you'll create Sparklines on the same sheet that contains the data. If you want to create Sparklines on a different sheet, start by activating the sheet where the Sparklines will be displayed. Then, in the Create Sparklines dialog box, specify the source data either by pointing or by typing the complete sheet reference (for example, Sheet1!A1 : C12). The Create Sparklines dialog box lets you specify a different sheet for the Data Range but not for the Location Range. Or you can just create the Sparklines on the same sheet as the data and then cut and paste the cells to a different worksheet.

Figure 20.4 shows column Sparklines for the precipitation data.

FIGURE 20.4

Column Sparklines summarize the precipitation data for nine cities.



Understanding Sparkline Groups

In most situations, you'll probably create a group of Sparklines—one for each row or column of data. A worksheet can hold any number of Sparkline groups. Excel remembers each group, and you can work with the group as a single unit. For example, you can select one Sparkline in a group and then modify the formatting of all Sparklines in the group. When you select one Sparkline cell, Excel displays an outline of all the other Sparklines in the group.

You can, however, perform some operations on an individual Sparkline in a group:

- To change the Sparkline's data source: Select any of the Sparkline cells and choose Sparkline ➔ Edit Data ➔ Edit Single Sparkline's Data. Excel displays a dialog box that lets you change the data source for the selected Sparkline.
- To delete, select the Sparkline cell and choose Sparkline ➔ Group ➔ Clear ➔ Clear Selected Sparklines.

You can also ungroup a set of Sparklines by selecting any Sparkline in the group and choosing Sparkline ➔ Group ➔ Ungroup. After you ungroup a set of Sparklines, you can work with each Sparkline individually.

Note that you can also delete an entire Sparkline group (should you want to start from scratch). Simply select Sparkline ➔ Group ➔ Clear ➔ Clear Selected Sparkline Groups.

TIP

It's worth noting all the previously mentioned Sparkline operations are available via the shortcut menu. Select a Sparkline cell and right-click. Then, select the Sparklines option to display the commands to edit, delete, and manage your Sparklines.

Customizing Sparklines

When you activate a cell that contains a Sparkline, Excel displays an outline around all the Sparklines in its group. You can then use the commands on the Sparkline tab to customize the group of Sparklines.

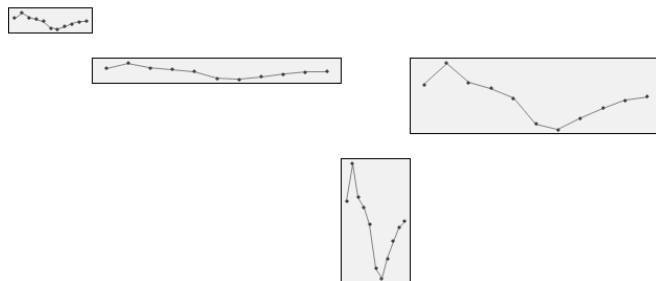
Sizing Sparkline cells

When you change the width or height of a cell that contains a Sparkline, the Sparkline adjusts accordingly. In addition, you can insert a Sparkline into merged cells.

Figure 20.5 shows the same Sparkline displayed at four sizes resulting from changing the column width or row height and merging cells. As you can see, the size and proportions of the cell (or merged cells) make a big difference in the appearance.

FIGURE 20.5

A Sparkline at various sizes



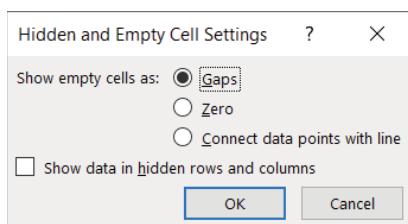
Handling hidden or missing data

By default, if you hide rows or columns that are used in a Sparkline graphic, the hidden data does not appear in the Sparkline. Also, missing data (an empty cell) is displayed as a gap in the graphic.

To change these settings, choose **Sparkline** \Rightarrow **Sparkline** \Rightarrow **Edit Data** \Rightarrow **Hidden & Empty Cells**. In the Hidden and Empty Cell Settings dialog box that appears (see Figure 20.6), specify how to handle hidden data and empty cells.

FIGURE 20.6

The Hidden and Empty Cell Settings dialog box



Changing the Sparkline type

As mentioned earlier, Excel supports three Sparkline types: Line, Column, and Win/Loss. After you create a Sparkline or group of Sparklines, you can easily change the type by selecting the Sparkline and clicking one of the three icons in the Sparkline \Rightarrow Type group. If the selected Sparkline is part of a group, all Sparklines in the group are changed to the new type.

TIP

If you've customized the appearance, Excel remembers your customization settings for each type if you switch among Sparkline types.

Changing Sparkline colors and line width

After you've created a Sparkline, changing the color is easy. Use the controls in the Sparkline \Rightarrow Style group.

Be aware that changing colors used in Sparkline graphics ties the colors to the document theme. Therefore, if you change the theme (by choosing Page Layout \Rightarrow Themes \Rightarrow Themes), the Sparkline colors will change to the new theme colors.

NOTE

Changing the document theme will not affect Sparklines that have their default colors. That is to say, switching the document theme will only alter the color of those Sparklines to which you've already applied a new color.



See Chapter 5, "Formatting Worksheets," for more information about document themes.

For Line Sparklines, you can also specify the line width. Choose Sparkline \Rightarrow Style \Rightarrow Sparkline Color \Rightarrow Weight.

Highlighting certain data points

Use the commands in the Sparkline \Rightarrow Show group to customize the Sparklines to highlight certain aspects of the data. The options are as follows:

High Point Apply a different color to the highest data point in the Sparkline.

Low Point Apply a different color to the lowest data point in the Sparkline.

Negative Points Apply a different color to negative values in the Sparkline.

First Point Apply a different color to the first data point in the Sparkline.

Last Point Apply a different color to the last data point in the Sparkline.

Markers Show data markers in the Sparkline. This option is available only for Line Sparklines.

Part III: Creating Charts and Other Visualizations

You control the color of the highlighting by using the Marker Color control in the Sparkline \Rightarrow Style group. Unfortunately, you can't change the size of the markers in Line Sparklines.

Figure 20.7 shows some Line Sparklines with various types of highlighting applied.

FIGURE 20.7

Highlighting options for Line Sparklines



Adjusting Sparkline axis scaling

As with a chart, Excel creates your Sparklines with a default axis scale. In other words, the minimum and maximum vertical axis values are determined automatically for each Sparkline in the group, based on the numeric range of the data used by the Sparkline.

The Sparkline \Rightarrow Group \Rightarrow Axis command lets you override this automatic behavior and control the minimum and maximum value for each Sparkline or for a group of Sparklines. For even more control, you can use the Custom Value option and specify the minimum and maximum for the Sparkline group.

NOTE

Although Sparklines have an axis scale, they don't display a vertical axis as you would see in a chart. You're essentially adjusting an invisible axis.

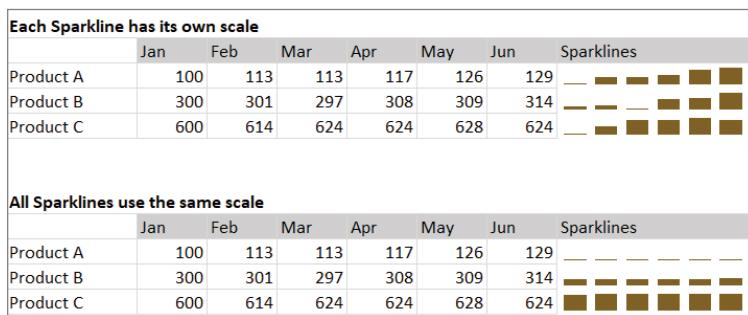
Figure 20.8 shows two groups of Sparklines. The group at the top uses the default axis settings (Automatic for Each Sparkline). Each Sparkline shows the six-month trend for the product, but there is no indication of the magnitude of the values.

For the Sparkline group at the bottom (which uses the same data), we changed the vertical axis minimum and maximum to use the Same for All Sparklines setting. With these settings in effect, the magnitude of the values across the products is apparent, but the trend across the months within a product is not.

The axis scaling option that you choose depends upon what aspect of the data you want to emphasize.

FIGURE 20.8

The bottom group of Sparklines shows the effect of using the same axis minimum and maximum values for all Sparklines in a group.



Faking a reference line

One useful feature that's missing in the Excel implementation of Sparklines is a reference line. For example, it might be useful to show performance relative to a goal. If the goal is displayed as a reference line in a Sparkline, the viewer can quickly see whether the performance for a period exceeded the goal.

You can, however, transform the data and then use a Sparkline axis as a fake reference line. Figure 20.9 shows an example. Students have a monthly reading goal of 500 pages. The range of data shows the actual pages read, with Sparklines in the last column. The Sparklines show the six-month page data, but it's impossible to tell who exceeded the goal and when they did it.

FIGURE 20.9

Sparklines display the number of pages read per month.

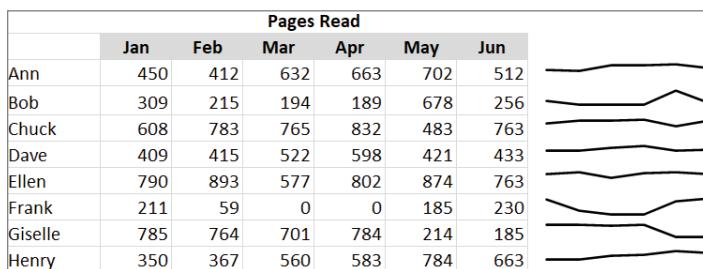


Figure 20.10 shows another approach: transforming the data such that meeting the goal is expressed as a 1 and failing to meet the goal is expressed as a -1. We used the following formula (in cell B18) to transform the original data:

$$=IF(B6>C2, 1, -1)$$

Part III: Creating Charts and Other Visualizations

FIGURE 20.10

Using Win/Loss Sparklines to display goal achievement

	A	B	C	D	E	F	G	H
1	Pages Read							
2	Monthly Goal:	500						
3								
4		Pages Read						
5		Jan	Feb	Mar	Apr	May	Jun	
6	Ann	450	412	632	663	702	512	
7	Bob	309	215	194	189	678	256	
8	Chuck	608	783	765	832	483	763	
9	Dave	409	415	522	598	421	433	
10	Ellen	790	893	577	802	874	763	
11	Frank	211	59	0	0	185	230	
12	Giselle	785	764	701	784	214	185	
13	Henry	350	367	560	583	784	663	
14								
15								
16		Pages Read (Did or Did Not Meet Goal)						
17		Jan	Feb	Mar	Apr	May	Jun	
18	Ann	-1	-1	1	1	1	1
19	Bob	-1	-1	-1	-1	1	-1
20	Chuck	1	1	1	1	-1	1
21	Dave	-1	-1	1	1	-1	-1
22	Ellen	1	1	1	1	1	1
23	Frank	-1	-1	-1	-1	-1	-1
24	Giselle	1	1	1	1	-1	-1
25	Henry	-1	-1	1	1	1	1

We copied this formula to the other cells in the B18 : G25 range.

Using the transformed data, we created Win/Loss Sparklines to visualize the results. This approach is better than the original, but it doesn't convey any magnitude differences. For example, you can't tell whether the student missed the goal by 1 page or by 500 pages.

Figure 20.11 shows a better approach. Here, we transformed the original data by subtracting the goal from the pages read. The formula in cell B31 is as follows:

=B6 - \$C\$2

We copied this formula to the other cells in the B31 : G38 range and created a group of Line Sparklines with the axis turned on. We also enabled the Negative Points option so that negative values (failure to meet the goal) clearly stand out.

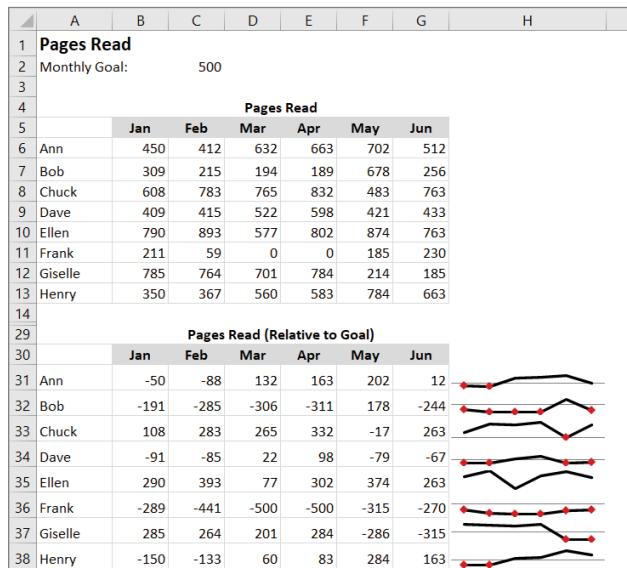
Specifying a Date Axis

Normally, data displayed in a Sparkline is assumed to be at equal intervals. For example, a Sparkline might display a daily account balance, sales by month, or profits by year. But what if the data isn't at equal intervals?

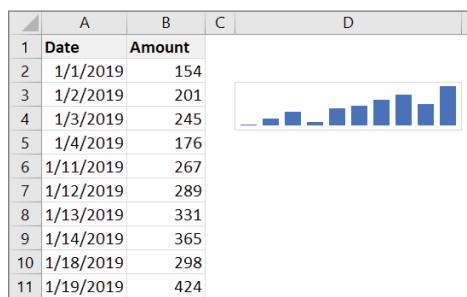
Figure 20.12 shows data, by date, along with a Sparklines graphic created from column B. Notice that some dates are missing, but the Sparkline shows the columns as if the values were spaced at equal intervals.

FIGURE 20.11

The axis in the Sparklines represents the goal.

**FIGURE 20.12**

The Sparkline displays the values as if they are at equal time intervals.

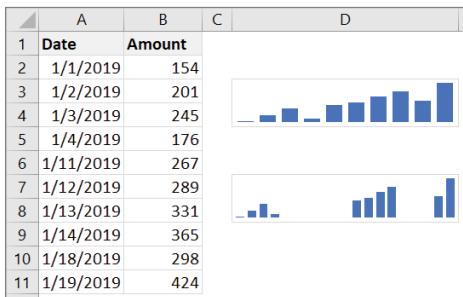


To better depict the data, the solution is to specify a date axis. Select the Sparkline and choose Sparkline \Rightarrow Group \Rightarrow Axis \Rightarrow Date Axis Type. Excel displays the Sparkline Date Range dialog box, asking for the range that contains the dates. In this example, specify range A2 : A11. Click OK, and the Sparkline displays gaps for the missing dates (see Figure 20.13).

Part III: Creating Charts and Other Visualizations

FIGURE 20.13

After a date axis is specified, the Sparkline shows the values accurately.



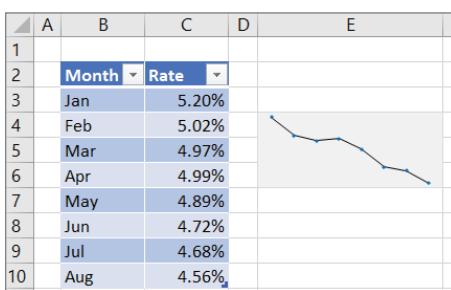
Auto-Updating Sparklines

If a Sparkline uses data in a normal range of cells, adding new data to the beginning or end of the range does not force the Sparkline to use the new data. You need to use the Edit Sparklines dialog box to update the data range. (Choose Sparkline \Rightarrow Sparkline \Rightarrow Edit Data.) But if the Sparkline data is in a column within a Table object (created by choosing Insert \Rightarrow Tables \Rightarrow Table), the Sparkline will use new data that's added to the end of the Table.

Figure 20.14 shows an example. The Sparkline was created using the data in the Rate column of the table. When you add the new rate for September, the Sparkline will automatically update its Data Range.

FIGURE 20.14

Creating a Sparkline from data in a table



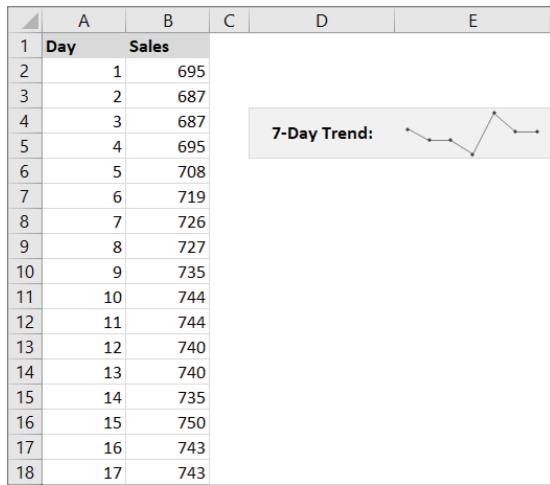
Displaying a Sparkline for a Dynamic Range

The example in this section describes how to create a Sparkline that displays only the most recent data points in a range. Figure 20.15 shows a worksheet that tracks daily sales. The Sparkline, in merged cells E4 : E5, displays only the seven most recent data points in

column B. When new data is added to column B, the Sparkline will adjust to show only the most recent seven days of sales.

FIGURE 20.15

Using a dynamic range name to display only the last seven data points in a Sparkline



We started by creating a dynamic range name. Here's how:

1. Choose **Formulas** \Rightarrow **Defined Names** \Rightarrow **Define Name**. Give your range a name (**Last7** in this case). Then, enter the following formula in the Refers To field:

```
=OFFSET ($B$2, COUNTA ($B:$B) -7-1, 0, 7, 1)
```

This formula calculates a range by using the **OFFSET** function. The first argument is the first cell in the range (B2). The second argument is the number of cells in the column (minus the number to be returned and minus 1 to accommodate the label in B1). The name always refers to the last seven nonempty cells in column B. To display a different number of data points, change both instances of 7 to a different value.

2. Choose **Insert** \Rightarrow **Sparklines** \Rightarrow **Line**. The Create Sparklines dialog box appears.
3. In the Data Range field, type **Last7** (name of the dynamic range you just created). Next specify cell **E4** as the Location Range. The Sparkline shows the data in range **B11:B17**.
4. Add new data to column B. The Sparkline adjusts to display only the last seven data points.

Visualizing with Custom Number Formats and Shapes

IN THIS CHAPTER

- Custom number formatting
- Creating visualizations with shapes and icons
- Creating your own infographic elements
- An overview of Excel's other graphics tools

Visualization is the presentation of abstract concepts or data in visual terms through some sort of graphical imagery. A traffic light, for example, is a visualization of the abstract concepts of stop and go.

In the business world, visualizations help us to communicate and process the meaning of data faster than simple tables of numbers. Excel offers a wide array of features that can be used to add visualizations to dashboards and reports.

In this chapter, you explore some of the formatting techniques that you can leverage to add layers of visualizations and turn your data into meaningful views.



All of the examples in this chapter are available on this book's website at www.wiley.com/go/excel365bible. The filename is `Visualizations.xlsx`.

Visualizing with Number Formatting

When you enter a number into a cell, you can display that number in a variety of different formats. Excel has quite a few built-in number formats, but sometimes there are none that are exactly what you need.

This chapter describes how to create custom number formats, and it provides many examples that you can use as is or adapt to your needs.

Doing basic number formatting

The Number group on the Home tab of the Ribbon contains several controls for applying common number formats quickly. The Number Format drop-down control gives you quick access to 11 common

Part III: Creating Charts and Other Visualizations

number formats. In addition, the Number group contains some buttons. When you click one of these buttons, the selected cells take on the specified number format. Table 21.1 summarizes the formats that these buttons perform in the U.S. English version of Excel.

TABLE 21.1 Number-Formatting Buttons on the Ribbon

Button Name	Formatting Applied
Accounting Number Format	Adds a dollar sign to the left, separates thousands with a comma, and displays the value with two digits to the right of the decimal point. This is a drop-down control, so you can select other common currency symbols.
Percent Style	Displays the value as a percentage, with no decimal places. This button applies a style to the cell.
Comma Style	Separates thousands with a comma and displays the value with two digits to the right of the decimal point. It's like the Accounting Number Format but without the currency symbol. This button applies a style to the cell.
Increase Decimal	Increases the number of digits to the right of the decimal point by one.
Decrease Decimal	Decreases the number of digits to the right of the decimal point by one.

NOTE

Some of these buttons apply predefined styles to the selected cells. You can access Excel's styles by using the Cell Styles gallery in the Styles group on the Home tab. You can modify the styles by right-clicking the style name and choosing Modify from the shortcut menu. See Chapter 5, "Formatting Worksheets," for details.

Using shortcut keys to format numbers

Another way to apply number formatting is to use shortcut keys. Table 21.2 summarizes the shortcut key combinations that you can use to apply common number formatting to the selected cells or range. Notice that these are the shifted versions of the number keys along the top of a typical keyboard.

TABLE 21.2 Number-Formatting Keyboard Shortcuts

Key Combination	Formatting Applied
Ctrl+Shift+~	General number format (that is, unformatted values).
Ctrl+Shift+!	Two decimal places, thousands separator, and a hyphen for negative values.
Ctrl+Shift+@	Time format with the hour, minute, and AM or PM.
Ctrl+Shift+#	Date format with the day, month, and year.

Key Combination	Formatting Applied
Ctrl+Shift+\$	Currency format with two decimal places. (Negative numbers appear in parentheses.)
Ctrl+Shift+%	Percentage format with no decimal places.
Ctrl+Shift+^	Scientific notation number format with two decimal places.

Using the Format Cells dialog box to format numbers

For maximum control of number formatting, you can use the Number tab in the Format Cells dialog box. You can access this dialog box in any of several ways:

- Click the dialog box launcher at the bottom right of the Home ▷ Number group.
- Choose Home ▷ Number ▷ Number Format ▷ More Number Formats.
- Press Ctrl+1.
- Right-click a cell or range and select the Format Cells command.

The Number tab in the Format Cells dialog box contains 12 categories of number formats from which to choose. When you select a category from the list box, the right side of the dialog box changes to display appropriate options.

Here are the number format categories, along with some general comments:

General The default format displays numbers as integers, as decimals, or in scientific notation if the value is too wide to fit into the cell.

Number Specify the number of decimal places, whether to use your system's thousands separator (for example, a comma) to separate thousands, and how to display negative numbers.

Currency Specify the number of decimal places, choose a currency symbol, and display negative numbers. This format always uses the system thousands separator symbol (for example, a comma) to separate thousands.

Accounting Differs from the Currency format in that the currency symbols always line up vertically, regardless of the number of digits displayed in the value.

Date Choose from a variety of date formats and select the locale for your date formats.

Time Choose from a number of time formats and select the locale for your time formats.

Percentage Choose the number of decimal places; always displays a percent sign.

Fraction Choose from among nine fraction formats.

Scientific Displays numbers in exponential notation (with an E): 2 .00E+05 = 200,000. You can choose the number of decimal places to display to the left of E.

Text When applied to a value, causes Excel to treat the value as text (even if it looks like a value). This feature is useful for such items as numerical part numbers and credit card numbers.

Special Contains additional number formats. The list varies, depending on the locale you choose. For the English (United States) locale, the formatting options are ZIP Code, ZIP Code + 4, Phone Number, and Social Security Number.

Custom Defines custom number formats not included in any of the other categories.

Getting fancy with custom number formatting

When you apply number formatting, you are giving Excel instructions using a format string. A number format string is a code that tells Excel how you want the values in a given range to appear cosmetically.

To see this code, follow these steps to apply basic number formatting:

1. **Right-click a cell or range and select Format Cells.** The Format Cells dialog box appears.
2. **Click the Number tab; then choose the Number category, select the Use 1000 Separator check box, 0 decimal places, and enclose negative numbers in parentheses.**
3. **Click the Custom category as shown in Figure 21.1.** Excel takes you to a screen that exposes the syntax that makes up the format you selected.

Number formatting strings consist of different individual number formats separated by semicolons. In this case, you see two different formats: the format to the left of the semicolon and the format to the right of the semicolon.

#,##0_);(#,##0)

By default, any formatting to the left of the first semicolon is applied to positive numbers, and any formatting to the right of the first semicolon is applied to negative numbers. So, with this choice, positive numbers will be formatted as a simple number, and negative numbers will be formatted with parentheses, like the following:

1,982
(1,890)

NOTE

The syntax for the positive formatting in the previous example ends with `_`. This tells Excel to leave a space the width of a parenthesis character at the end of positive numbers. This ensures that positive and negative numbers align nicely when negative numbers are wrapped in parentheses.

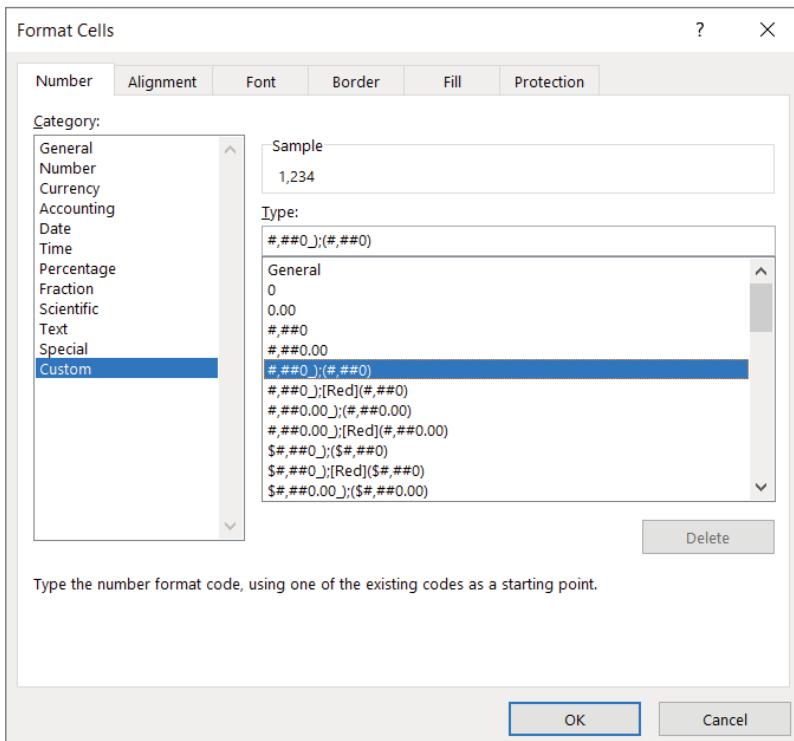
You can edit the syntax in the Type input box so that the numbers are formatted differently. For example, try changing the syntax to the following:

+#,##0;-#,##0

Chapter 21: Visualizing with Custom Number Formats and Shapes

FIGURE 21.1

The Type input box allows you to customize the syntax for the number format.



When this syntax is applied, positive numbers will start with the + symbol, and negative numbers will start with a – symbol, like so:

+1,200
-15,000

This comes in handy when formatting percentages. For instance, you can apply a custom percent format by entering the following syntax into the Type input box:

+0% ; -0%

This syntax gives you percentages that look like the following:

+43%
-54%

You can get fancy and wrap your negative percentages with parentheses using the following syntax:

0%_) ; (0%)

Part III: Creating Charts and Other Visualizations

This syntax gives you percentages that look like the following:

43%
(54%)

Note

If you include only one format syntax, meaning that you don't add a second formatting option with the use of a semicolon separator, that one format will be applied to all numbers, negative or positive.

Formatting numbers in thousands and millions

Formatting numbers to appear in thousands or millions helps present values without inundating your audience with overlarge numbers. To show your numbers in thousands, highlight them, right-click, and select Format Cells.

After the Format Cells dialog box opens, click the Custom category to get to the screen shown in Figure 21.1. In the Type input box, enter the following syntax:

#,##0,

After you confirm your changes, your numbers will automatically appear in the thousands place.

The beautiful thing here is that this technique doesn't change or truncate your numeric values in any way. Excel is simply applying a cosmetic effect to the number. Figure 21.2 illustrates this. Compare the values you see in the formula bar versus what is shown in the cells.

FIGURE 21.2

Formatting numbers applies only a cosmetic look. Look in the formula bar to see the real, unformatted number.

	A	B	C	D	E
4			X ✓ fx	117943.60578	
5					
6					
7					
8					
9					

Formatting Large Numbers

	January	February	March
North	118k	380k	463k
Northeast	24k	803k	328k
East	313k	780k	904k
Southeast	397k	466k	832k

The selected cell has been formatted to show in thousands. You see 118k. But if you look in the formula bar above it, you'll see the real unformatted number (117943.60578). The 118k you are seeing in the cell is a cosmetically formatted version of the real number shown in the formula bar.

Note

Custom number formatting has obvious advantages over using other techniques to format numbers to thousands. For instance, you could convert numbers to thousands by dividing them by 1,000 in a formula. But that changes the integrity of the number dramatically. When you perform a mathematical operation in a cell, you are literally changing the value represented in that cell. This forces you to carefully keep track of and maintain the formulas you introduced simply to achieve a cosmetic effect. Using custom number formatting avoids that by changing only how the number looks, keeping the actual number intact.

If needed, you can indicate that the number is in thousands by adding "k" to the number syntax.

```
#,##0,"k"
```

This would show your numbers like this:

```
118k  
318k
```

You can use this technique on both positive and negative numbers.

```
#,##0,"k"; (#,##0,"k")
```

After applying this syntax, your negative numbers also appear in thousands.

```
118k  
(318k)
```

Need to show numbers in millions? Easy. Simply add two commas to the number format syntax in the Type input box.

```
#,##0.00,, "m"
```

Note the use of the extra decimal places (.00). When converting numbers to millions, it's often useful to show additional precision points, as in the following example:

```
24.65 m
```

Hiding and suppressing zeros

In addition to formatting positive and negative numbers, Excel allows you to provide a format for zeros. You do this by adding another semicolon to your custom number syntax. By default, any format syntax placed after the second semicolon is applied to any number that evaluates to zero.

For example, the following syntax applies a format that shows n/a for any cells that contain zeros:

```
#,##0_); (#,##0); "n/a"
```

Part III: Creating Charts and Other Visualizations

You can also use this to suppress zeros entirely. If you add the second semicolon but don't follow it with any syntax, cells containing zeroes will appear blank.

```
#,\#\#0_) ; (#,\#\#0) ;
```

Again, custom number formatting affects only the cosmetic look of the cell. The actual data in the cell is not affected. Figure 21.3 demonstrates this. The selected cell is formatted so that zeros appear as n/a, but if you look at the Formula bar, you can see the actual unformatted cell contents.

FIGURE 21.3

Custom number formatting that shows zeros as n/a

	A	B	C	D	E
39					
40					
		Jim	Tim	Kim	
41	Printers	37,000	64,000	24,000	
42	Copiers	18,000	29,000	58,000	
43	Scanners	n/a	77,000	88,000	
44	Service Contracts	16,000	12,000	n/a	
45	Warranties	65,000	88,000	16,000	

Applying custom format colors

In addition to controlling the look of your numbers with custom number formatting, you can control their color. For instance, to format percentages so that positive percentages appear blue with a + symbol while negative percentages appear red with a - symbol, enter this syntax in the Type input box:

```
[Blue]+0%;[Red]-0%
```

Notice that all it takes to apply a color is to enter the color name wrapped in square brackets: [].

There are only certain colors (the eight Visual Basic colors) you can call out by name, as shown here. These colors make up the first eight colors of Excel's legacy palette (the standard 56 colors that were the default in versions pre-2007).

```
[Black]  
[Blue]  
[Cyan]  
[Green]  
[Magenta]  
[Red]  
[White]  
[Yellow]
```

Chapter 21: Visualizing with Custom Number Formats and Shapes

Although you would typically specify a custom color by name, not all the Visual Basic colors are pleasing. Visual Basic Green is notoriously difficult to look at (bright neon green). See for yourself by entering the following code into the Type input box:

[Green] +0% ; [Red] -0%

The good news is that there are 56 colors defined in the standard color palette by number. Every color in the standard 56-color palette is represented by a number. To call up a color by number, you would use [ColorN], where N represents a number from 1 to 56.

In this example, you can use [Color10] to present a much more acceptable green.

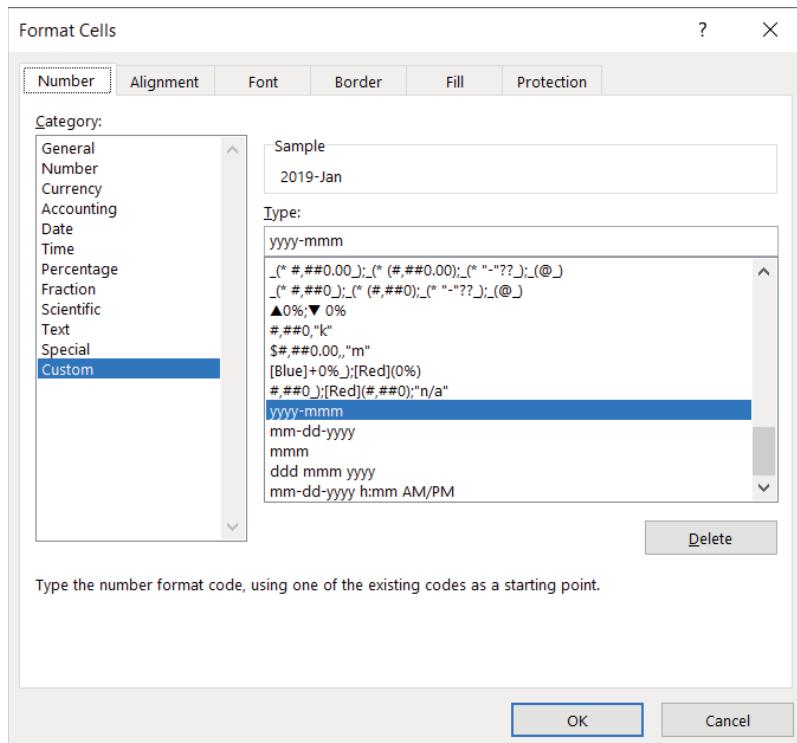
[Color10] +0% ; [Red] -0%

Formatting dates and times

Custom number formatting isn't just for numbers. You can also format dates and times. As Figure 21.4 illustrates, you use the same dialog box to apply date and time formats using the Type input box.

FIGURE 21.4

Dates and times can also be formatted using the Format Cells dialog box.



Part III: Creating Charts and Other Visualizations

The code used for date and time formatting is fairly intuitive. For example, `ddd` is the syntax for the three-letter day, `mmm` is the syntax for the three-letter month, and `yyyy` is the syntax for the four-digit year.

There are several variations on the format for days, months, years, hours, and minutes. It's worthwhile to take some time and experiment with different combinations of syntax strings.

Table 21.3 lists some common date and time format codes that you can use as starter syntax for your reports and dashboards.

TABLE 21.3 Common Date and Time Format Codes

Format Code	1/31/2022 7:42:53 PM is Displayed As
m	1
mm	01
mmm	Jan
mmmm	January
mmmmmm	J
dd	31
ddd	Mon
ddd	Monday
yy	22
yyyy	2022
mmm-yy	Jan-22
dd/mm/yyyy	31/01/2022
ddd mmm yyyy	Monday Jan 2022
mm-dd-yyyy h:mm AM/PM	01-31-2022 7:42 PM
h AM/PM	7 PM
h:mm AM/PM	7:42 PM
h:mm:ss AM/PM	7:42:53 PM

Using symbols to enhance reporting

Symbols are essentially tiny graphics, not unlike those you see when you use the Wingdings, Webdings, or other fancy fonts. However, symbols are not really fonts. They're Unicode characters. *Unicode characters* are a set of industry-standard text elements designed to provide a reliable character set that remains viable on any platform regardless of international font differences.

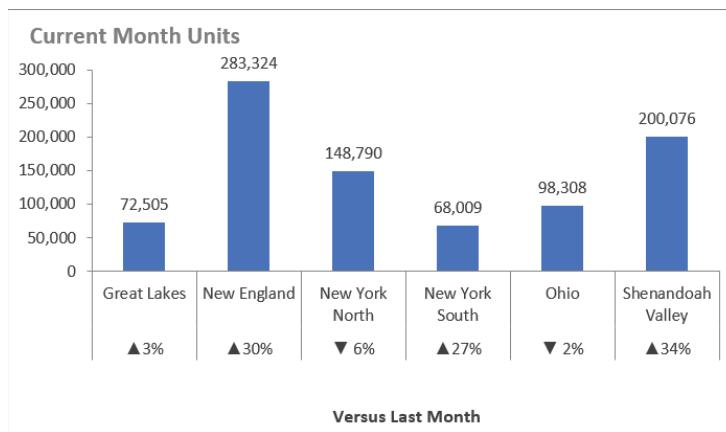
Chapter 21: Visualizing with Custom Number Formats and Shapes

One example of a commonly used symbol is the copyright symbol (©). This symbol is a Unicode character. You can use this symbol on a Chinese, Turkish, French, or English (United States) PC, and it will reliably be available with no international differences.

In terms of Excel presentations, Unicode characters (or symbols) can be used in places where conditional formatting cannot. For instance, in the chart labels that you see in Figure 21.5, the x-axis shows some trending arrows that allow for an extra layer of analysis. This couldn't be done with conditional formatting.

FIGURE 21.5

Use symbols to add an extra layer of analysis to charts.



Let's take some time to review the steps that led to the chart in Figure 21.5.

Start with the data shown in Figure 21.6. Note that you have a designated cell to hold any symbols that you're going to use (C1 in this case). This cell isn't really all that important. It's just a holding cell for the symbols that you'll insert.

FIGURE 21.6

Our starting data with a holding cell for our symbols

	A	B	C	D
1		Symbols>>		
2				
3	vs. Prior Month	Market	Current Month	
4	3%	Great Lakes	72,505	
5	30%	New England	283,324	
6	-6%	New York North	148,790	
7	27%	New York South	68,009	
8	-2%	Ohio	98,308	
9	34%	Shenandoah Valley	200,076	

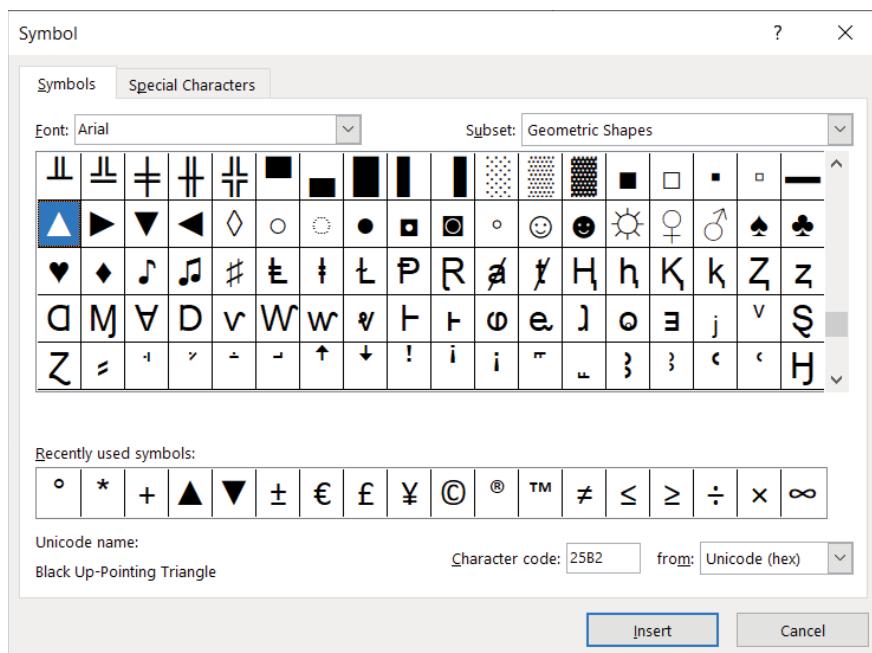
Part III: Creating Charts and Other Visualizations

Now, follow these steps:

1. **Click in C1 and then select Symbols ↴ Symbol on the Insert tab.** The Symbol dialog box shown in Figure 21.7 opens. The idea here is to scroll through the symbols to find the one you want. Note that you can use the Font and Subset drop-down selectors at the top of the dialog box to quickly jump to a particular set of symbols (in this case, Arial and Geometric Shapes).

FIGURE 21.7

Use the Symbol dialog box to insert the desired symbols into your holding cell.



2. **Find and select your desired symbols by clicking the Insert button after each symbol.** In this scenario, select the upward-pointing triangle and click Insert. Then click the downward-pointing triangle and click Insert. Close the dialog box when you're done. At this point, you have the up triangle and down triangle symbols in cell C1, as shown in Figure 21.8.
3. **Click the C1 cell, go to the Formula bar, and copy the two symbols by highlighting them and pressing Ctrl+C on your keyboard.**
4. **Go to your data table, right-click the percentages, and then select Format Cells.**

Chapter 21: Visualizing with Custom Number Formats and Shapes

FIGURE 21.8

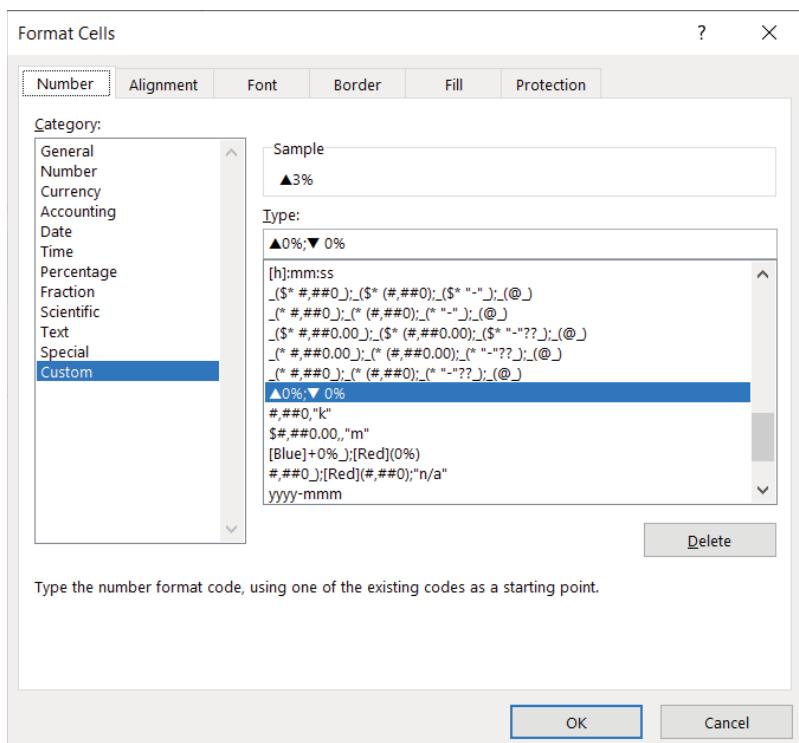
Copy the newly inserted symbols to the Clipboard.

A	B	C	D
1	Symbols>>	▲▼	
2	vs. Prior Month	Market	Current Month
3	3%	Great Lakes	72,505
4	30%	New England	283,324
5	-8%	New York North	148,790
6	27%	New York South	68,009
7	-2%	Ohio	98,308
8	34%	Shenandoah Valley	200,076
9			

5. In the Format Cells dialog box, create a new custom format by pasting the up-and down-triangle symbols into the appropriate syntax parts (see Figure 21.9). In this case, any positive percentage will be preceded with the up-triangle symbol, and any negative percentage will be preceded with the down-triangle symbol.

FIGURE 21.9

Create a custom number format using the symbols.



Part III: Creating Charts and Other Visualizations

6. Click OK. The symbols are now part of your number formatting.

Figure 21.10 illustrates what your percentages look like. Change any number from positive to negative (or vice versa), and Excel automatically applies the appropriate symbol.

FIGURE 21.10

Your symbols are now part of your number formatting.

A	B	C	D
1	Symbols>>	▲▼	
2			
3	vs. Prior Month	Market	Current Month
4	▲ 3%	Great Lakes	72,505
5	▲ 30%	New England	283,324
6	▼ 6%	New York North	148,790
7	▲ 27%	New York South	68,009
8	▼ 2%	Ohio	98,308
9	▲ 34%	Shenandoah Valley	200,076

Because charts automatically adopt number formatting, a chart created from this data will show the symbols as part of the labels. Simply use this data as the source for the chart.

This is just one way to use symbols in your reporting. With this basic technique, you can insert symbols to add visual appeal to tables, PivotTables, formulas, or any other object you can think of.

Using Shapes and Icons as Visual Elements

Microsoft Office, including Excel, provides access to a variety of customizable graphics images known as *shapes*. You might want to insert shapes to create simple diagrams, display text, or just add some visual appeal to a worksheet.

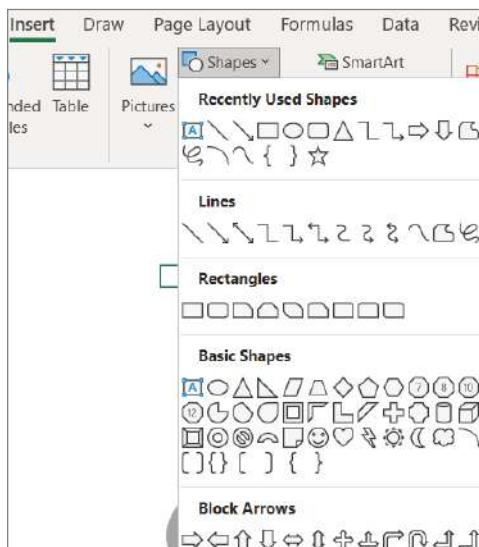
Keep in mind that shapes can add unnecessary clutter to a worksheet. Perhaps the best advice is to use shapes sparingly. Ideally, shapes can help draw attention to some aspect of your worksheet. They shouldn't be the main attraction.

Inserting a shape

You can add a shape to a worksheet by choosing Insert \Rightarrow Illustrations \Rightarrow Shapes. The Shapes gallery, shown in Figure 21.11, opens to show you the choices.

FIGURE 21.11

The Shapes gallery



Shapes are organized into categories, and the category at the top displays the shapes that you've used most recently. To insert a shape into a worksheet, you can do one of the following:

- **Click the shape in the Shapes gallery and then click in the worksheet.** A default-sized shape is added to your worksheet.
- **Click the shape and then drag in the worksheet.** This allows you to create a larger or a smaller shape or a shape with different proportions than the default.

Here are a few tips to keep in mind when creating shapes:

- Every shape has a descriptive name (for example, Line Arrow). To change the name of a shape, select it, type a new name in the Name box (to the left of the formula bar), and press Enter.
- To select a specific shape on a worksheet, just click it.
- When you create a shape by dragging, hold down the Shift key to maintain the object's default proportions.
- You can control the way objects appear onscreen in the Advanced tab of the Excel Options dialog box. (Choose File → Options.) This setting appears in the Display Options for This Workbook section. Normally, the All option is selected under For Objects Show. You can hide all objects by choosing Nothing (Hide Objects). Hiding objects may speed things up if your worksheet contains complex objects that take a long time to redraw.

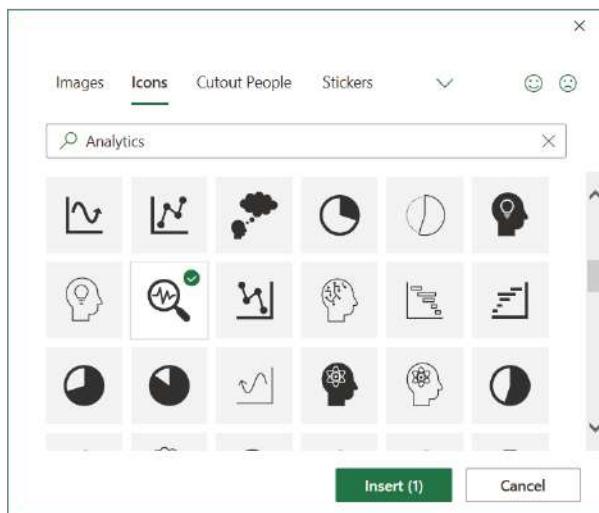
Inserting SVG icon graphics

Excel 365 includes an icon library that offers free Scalable Vector Graphics (SVG) icons. SVG graphics can be sized and formatted without losing image quality. These icon graphics are essentially a modern set of graphic files that can be used to add spicy visual elements to your Excel dashboards and infographics.

To add an icon to your worksheet, select Insert \Rightarrow Illustrations \Rightarrow Icons. This activates the dialog box shown in Figure 21.12. Here you can browse by category or search for a topic. When you find the graphic you want, select it and then click Insert. You can also simply double-click the graphic. Excel will insert it into your workbook.

FIGURE 21.12

The Microsoft Office Icons Library



Inserting 3D models

3D models are just what they sound like, graphics that can be dynamically rotated to see them from all sides. To add a 3D model to your worksheet, select Insert \Rightarrow Illustrations \Rightarrow 3D Models.

Just as with the previously mentioned icons, you can use the Online 3D Models dialog box to browse or search for certain objects. When you find the graphic you want, select it, and then click Insert or double-click the graphic. Excel will insert the graphic into your workbook.

Figure 21.13 illustrates how a 3D model can be rotated and viewed from all angles.

FIGURE 21.13

A 3D model in action



You'll quickly find that the stock 3D models are frankly underwhelming and irrelevant to most Excel users. However, the real benefit comes to those who regularly use their own 3D models. Excel allows the use of custom 3D model files outside of the stock options provided. You can use any of the following valid file types:

- .3mf** - 3D Manufacturing Format
- .fbx** - Filmbox Format
- .glb** - Binary GL Transmission Format
- .obj** - Object Format
- .ply** - Polygon Format
- .stl** - STereoLithography Format

To insert your own custom 3D model graphics, select **Insert** \Rightarrow **Illustrations** and click the drop-down arrow next to the **3D Models** command. Select the option labeled **This Device**. Excel opens a dialog box allowing you to browse your own device for the file you want to insert.

Selecting and Hiding Shape Objects

The Selection task pane allows you to easily find and manage shape objects. It's particularly useful when you've got lots of shape objects on a worksheet and you want to work with a select few. Activate the Selection task pane by choosing **Home** \Rightarrow **Editing** \Rightarrow **Find & Select** \Rightarrow **Selection Pane**.

Each object on the active worksheet is listed in the Selection task pane. Just click the object's name to select it. To select multiple objects, press **Ctrl** while you click the names.

To hide an object, click the "eye" icon to the right of its name. Use the buttons at the top of the task pane to hide (or show) all items quickly.

Formatting shapes and icons

Although icons feel and behave like standard shapes, they have different contextual tabs.

When you select a shape, the Shape Format contextual tab is available, with the following groups of commands:

Insert Shapes Insert new shapes; change a shape to a different shape.

Shape Styles Change the overall style of a shape; modify the shape's fill, outline, or effects.

WordArt Styles Modify the appearance of the text within a shape.

Accessibility Provide alternate text descriptions for the visually impaired.

Arrange Adjust the "stack order" of shapes, align shapes, group multiple shapes, and rotate shapes.

Size Change the size of a shape by typing dimensions.

When you select an icon, the Graphics Format contextual tab is available with the following groups of commands:

Change Graphic Replace an existing graphic with a new graphic from a file, the Microsoft icons gallery, or an online source.

Graphics Styles Change the overall style of a graphic; modify the graphic's fill, outline, or effects.

Accessibility Provide alternate text descriptions for the visually impaired.

Arrange Adjust the "stack order" of graphics, align graphics, group multiple graphics, and rotate graphics.

Size Change the size of a graphic by typing dimensions.

When you select a 3D model, the 3D Model contextual tab is available with the following groups of commands:

Play 3D Control the animation of the 3D model.

Adjust Replace the 3D model with another one or reset the 3D model to its original position.

3D Model Views Quickly position the 3D model to preset angles.

Accessibility Provide alternate text descriptions for the visually impaired.

Arrange Adjust the "stack order" of 3D models, align 3D models, and group multiple 3D models.

Size Change the size of a 3D model by typing dimensions; pan and zoom to focus on a specific area of the 3D model.

Chapter 21: Visualizing with Custom Number Formats and Shapes

As an alternative to the Ribbon, you can right-click the shape, icon, or 3D model and choose the Format option. This displays a task pane containing formatting options. Any changes made via the task pane appear immediately, and you can keep the Format task pane open while you work.

You could read 20 pages about formatting shapes, icons, and 3D models, but it wouldn't be a very efficient way of learning. The best way, by far, to learn about formatting shapes, icons, and 3D models is to experiment. The formatting commands are intuitive, and you can always use Undo if a command doesn't do what you expected it to do.

Enhancing Excel reports with shapes

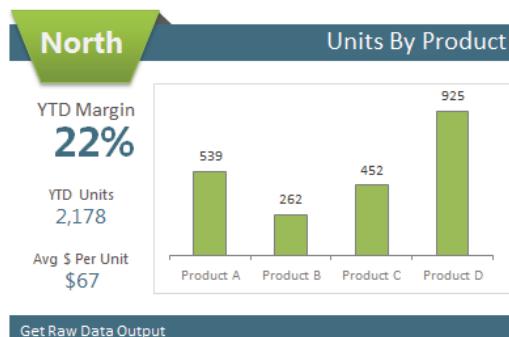
Most of us think of Excel shapes as mildly useful objects that can be added to a worksheet if we need to show a square, some arrows, a circle, and so forth. But if you use your imagination, you can leverage Excel shapes to create stylized interfaces that can really enhance your dashboards. Here are a few examples of how Excel shapes can spice up your dashboards and reports.

Creating visually appealing containers with shapes

A *peekaboo tab* lets you tag a section of your dashboard with a label that looks like it's wrapping around your dashboard components. In the example illustrated in Figure 21.14, a peekaboo tab is used to label this group of components as belonging to the North region.

FIGURE 21.14

Peekaboo tab



As you can see in Figure 21.15, there is no real magic here. It's just a set of shapes and text boxes that are cleverly arranged to give the impression that a label is wrapping around to show the region name.

Want to draw attention to handful of key metrics? Try wrapping your key metrics with a peekaboo banner. The banner shown in Figure 21.16 goes beyond boring text labels, allowing you to create the feeling that a banner is wrapping around your numbers. Again, this effect is achieved by layering a few Excel shapes so that they fall nicely on top of each other, creating a cohesive effect.

Part III: Creating Charts and Other Visualizations

FIGURE 21.15

Deconstructed view of the peekaboo tab

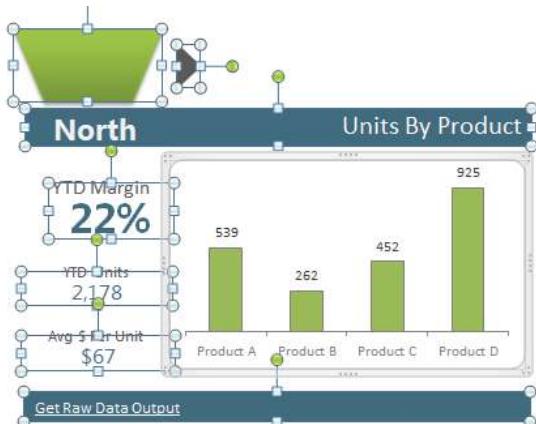
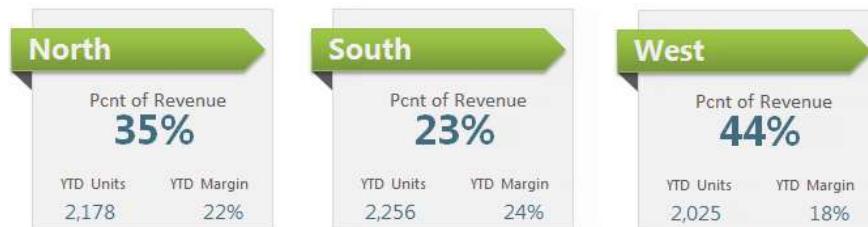


FIGURE 21.16

A visual banner made with shapes



North

YTD Units
2,178

Pcnt of Revenue
35%

West

Pcnt of Revenue
44%

YTD Units
2,025

South

Pcnt of Revenue
23%

YTD Units
2,256

Layering shapes to save space

Here's an idea to get the most out of your dashboard real estate. You can layer pie charts with column charts to create a unique set of views (see Figure 21.17).

FIGURE 21.17

Combine shapes with a chart to save dashboard real estate.



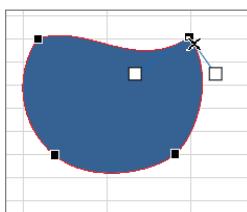
Each pie chart represents the percent of total revenue and a column chart showing some level of detail for the region. Simply layer your pie chart on top of a circle shape and a column chart.

Constructing your own infographic widgets with shapes

Excel offers a way to alter shapes by editing their anchor points. This opens the possibility of creating your own infographic widgets. Right-click a shape and select Edit Points. This places little points all around the shape (see Figure 21.18). You can then drag the points to reconfigure the shape.

FIGURE 21.18

Use the Edit Points feature to construct your own shape.

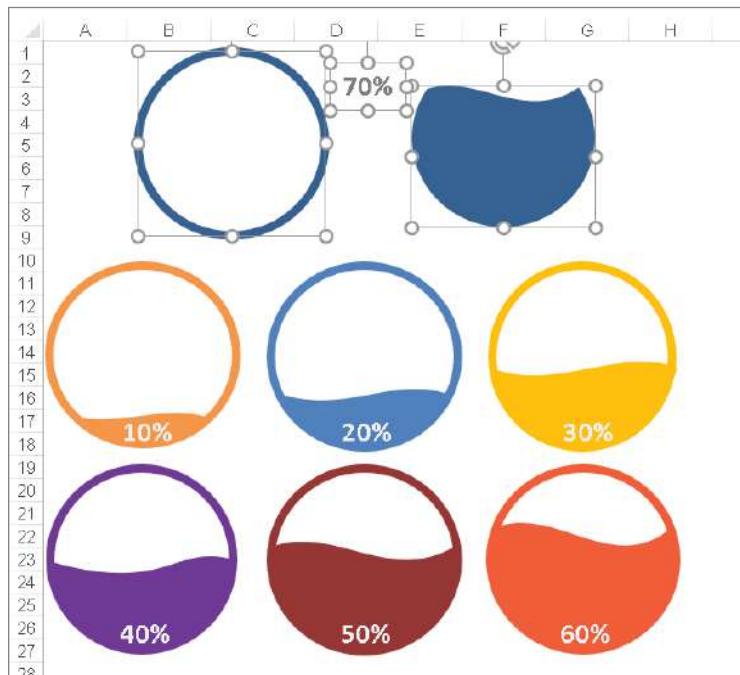


Constructed shapes can be combined with other shapes to create interesting infographic elements that can be used in your Excel dashboards. In Figure 21.19, a newly constructed shape is combined with a standard oval and text box to create nifty infographic widgets.

Part III: Creating Charts and Other Visualizations

FIGURE 21.19

Using a newly constructed shape to create custom infographic elements



Creating dynamic labels

Dynamic labeling is less a function in Excel than it is a concept. *Dynamic labels* are labels that change to correspond to the data you're viewing.

Figure 21.20 illustrates one example of this concept. The selected Text Box shape is linked to cell C3 (note the formula in the Formula bar). As the value in cell C3 changes, the text box displays the updated value.

NOTE

Be aware that Text Box shape objects cannot display more than 255 characters.

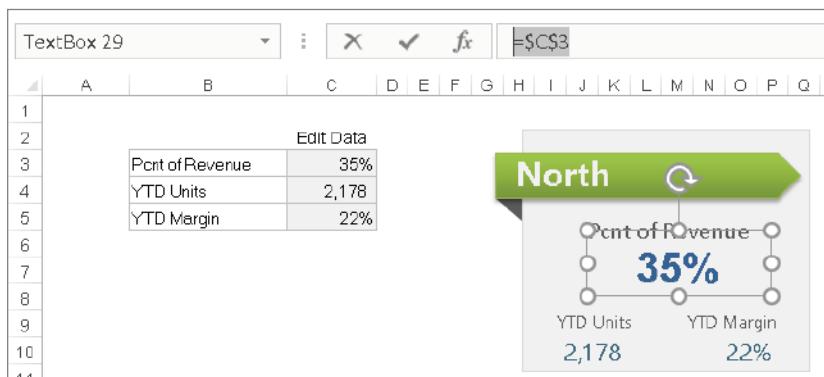
Creating linked pictures

A *linked picture* is a special kind of shape that displays a live picture of everything in a given range. Think of a linked picture as a camera that monitors a range of cells.

Chapter 21: Visualizing with Custom Number Formats and Shapes

FIGURE 21.20

Text Box shapes can be linked to cells.



To “take a picture” of a range, follow these steps:

1. Select the range.
2. Press **Ctrl+C** to copy the range.
3. Activate another cell.
4. Choose **Home** → **Clipboard** → **Paste** → **Linked Picture** (see Figure 21.21).

The result is a live picture of the range you selected in step 1.

Linked pictures give you the freedom to test different layouts and chart sizes without the need to work around column widths, hidden rows, or other such nonsense. In addition, linked pictures have access to the Picture Format options. When you click on a linked picture, you can go to the Picture Format contextual tab and play around with the picture styles there.

Figure 21.22 illustrates two linked pictures displaying the contents of the ranges on the left. As those ranges change, the linked pictures on the right will update. These can be moved, resized, and even placed on a completely different sheet.

TIP

The Excel Camera tool provides an easier way to create a linked picture. Unfortunately, it's not on the Ribbon interface. You can save some time by adding the Excel Camera tool to your Quick Access Toolbar as follows:

1. Right-click the Quick Access Toolbar and choose **Customize Quick Access Toolbar** from the shortcut menu that appears. The Excel Options dialog box appears, with the Quick Access Toolbar tab selected.
2. Select **Commands Not in the Ribbon** from the drop-down list on the left.
3. Select **Camera** from the list and click **Add**.
4. Click **OK** to close the Excel Options dialog box.

After you add the Camera tool to your Quick Access Toolbar, you can select a range of cells and click the Camera tool to take a “picture” of the range. Then click in the worksheet, and Excel places a live picture of the selected range on the worksheet's draw layer. If you make changes to the original ranges, the changes are shown in the picture of the range.

Part III: Creating Charts and Other Visualizations

FIGURE 21.21

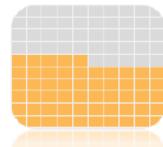
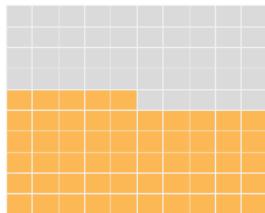
Pasting a linked picture



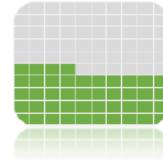
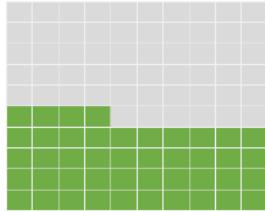
FIGURE 21.22

Using linked pictures to enhance visualizations

Metric 1
55%



Metric 2
44%



Using SmartArt and WordArt

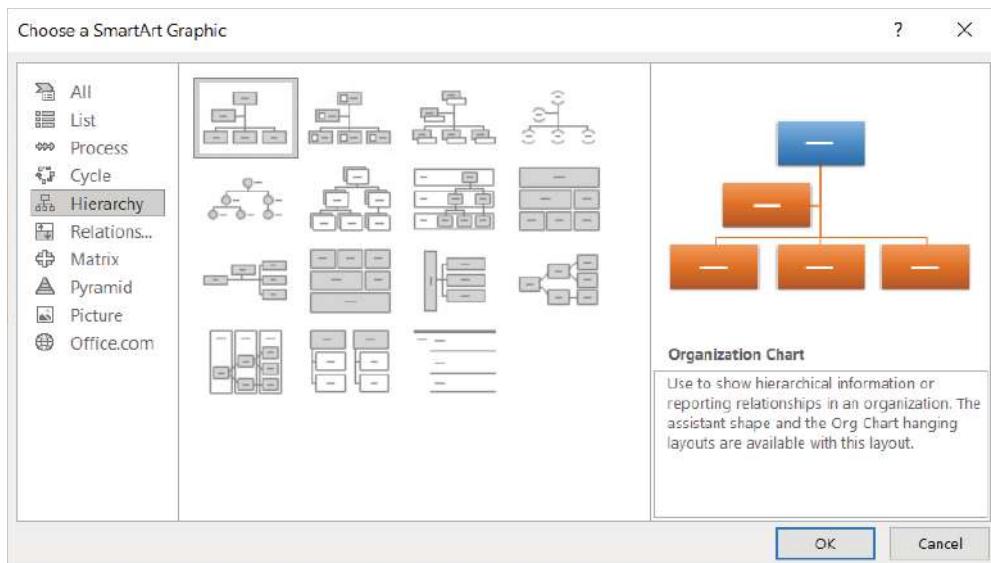
Using *SmartArt*, you can insert a variety of highly customizable diagrams into a worksheet. This feature is probably most useful in PowerPoint, as diagrams are more conducive to a presentation style. However, Excel offers SmartArt as yet another method of adding visualizations to your reporting solutions.

SmartArt basics

To insert SmartArt into a worksheet, choose the SmartArt command found under Insert \Rightarrow Illustrations. Excel displays the Choose a SmartArt Graphic dialog box shown in Figure 21.23. The diagrams are arranged in categories along the left. When you find one that looks appropriate, click it for a larger view in the panel on the right. When you find one that looks appropriate, click OK to insert the graphic.

FIGURE 21.23

Inserting a SmartArt graphic



NOTE

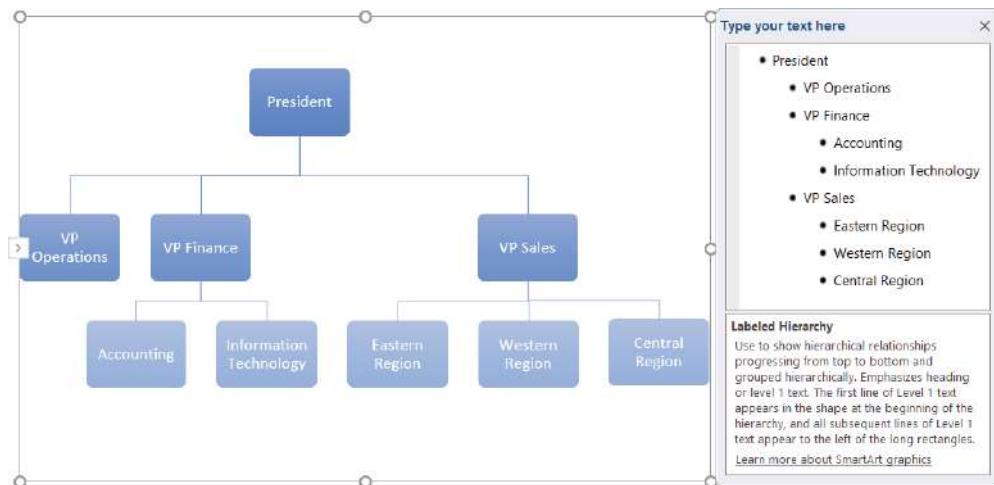
Don't be concerned about the number of elements contained in SmartArt graphics. You can customize the SmartArt to display the number of elements you need.

When you insert or select a SmartArt diagram, Excel displays the Text Pane, the Type Your Text Here window that guides you through entering text (see Figure 21.24).

Part III: Creating Charts and Other Visualizations

FIGURE 21.24

Entering text for an organizational chart



To add an element to the SmartArt graphic, choose SmartArt Design \Rightarrow Create Graphic
 \Rightarrow Add Shape.

When working with SmartArt, keep in mind that you can move, resize, or format individually any element within the graphic. Select the element and then use the tools on the Format tab.

You can easily change the layout of a SmartArt diagram. Select the object and then choose SmartArt Design \Rightarrow Layouts. Any text that you've entered remains intact.

After you decide on a layout, you may want to consider other styles or colors available in the SmartArt Design \Rightarrow SmartArt Styles group.

TIP

The SmartArt styles available vary depending on the document theme assigned to the workbook. To change a workbook's theme, choose Page Layout \Rightarrow Themes \Rightarrow Themes. Switching to a different theme can have a dramatic impact on the appearance of SmartArt diagrams.

WordArt basics

You can use *WordArt* to create graphical effects in text.

To insert a WordArt graphic on a worksheet, choose Insert \Rightarrow Text \Rightarrow WordArt and then select a style from the gallery. Excel inserts an object with the placeholder text Your text here. Replace that text with your own, resize it, and apply other formatting if you like.

When you select a WordArt image, Excel displays the Shape Format contextual tab. Use the controls to vary the look of your WordArt.

Working with Other Graphics Types

Excel can import a variety of graphics into a worksheet. You have several choices:

Inserting an image from your computer If the graphics image that you want to insert is available in a file, you can easily import the file into your worksheet. Choose Insert \Rightarrow Illustrations \Rightarrow Pictures \Rightarrow This Device. The Insert Picture dialog box appears, allowing you to browse for the file. Oddly, you can't drag and drop an image from this dialog box into a worksheet, but you can sometimes drag an image from your web browser and drop it into a worksheet.

Inserting an image from an online source Choose Insert \Rightarrow Illustrations \Rightarrow Pictures \Rightarrow Online Pictures. The Online Pictures window appears, allowing you to search for an image. The image is inserted on the current worksheet along with a text box denoting the origin of the picture.

Copying and pasting If an image is on the Windows Clipboard, you can paste it into a worksheet by choosing Home \Rightarrow Clipboard \Rightarrow Paste (or by pressing Ctrl+V).

About graphics files

Graphics files come in two main categories:

Bitmap Bitmap images are made up of discrete dots. They usually look pretty good at their original size but often lose clarity if you increase the size. Examples of common bitmap file formats include BMP, PNG, JPEG, TIFF, and GIF.

Vector Vector-based images, on the other hand, are composed of points and paths that are represented by mathematical equations, so they retain their crispness regardless of their size. Examples of common vector file formats include CGM, WMF, and EPS.

CAUTION

Using bitmap graphics in a worksheet can dramatically increase the size of your workbook (depending on the number and size of the graphics you use).

When you insert a picture on a worksheet, you can modify the picture in a number of ways from the Picture Format contextual tab, which becomes available when you select a picture object. For example, you can adjust the color, contrast, and brightness. In addition, you can add borders, shadows, reflections, and so on—similar to the operations available for shapes.

In addition, you can right-click and choose Format Picture to use the controls in the Format Picture task pane.

Part III: Creating Charts and Other Visualizations

An interesting feature is *Artistic Effects*. This command can apply a number of Photoshop-like effects to an image. To access this feature, select an image and choose Picture Format \Rightarrow Adjust \Rightarrow Artistic Effects. Each effect is somewhat customizable, so if you're not happy with the default effect, try adjusting some options.

Inserting screenshots

Excel can also capture and insert a screenshot of any program currently running on your computer (including another Excel window). To use the screenshot feature, follow these steps:

1. **Make sure the window you want to use displays the content you want.**
2. **Choose Insert \Rightarrow Illustrations \Rightarrow Screenshot.** You'll see a gallery that contains thumbnails of all windows open on your computer (except the current Excel window).
3. **Click the image you want.** Excel inserts it into your worksheet.

You can use any of the normal picture tools to work with screenshots.

Displaying a worksheet background image

If you want to use a graphics image for a worksheet's background (similar to wallpaper on the Windows Desktop), choose Page Layout \Rightarrow Page Setup \Rightarrow Background and select a graphics file. The selected graphics file is tiled on the worksheet.

Unfortunately, worksheet background images are for onscreen display only. These images do not appear when the worksheet is printed.

Using the Equation Editor

The final section in this chapter deals with the Equation Editor. Use this feature to insert a nicely formatted mathematical equation as a graphics object. Start by choosing Insert \Rightarrow Symbols \Rightarrow Equation. Excel creates a text box and displays the Shape Format and Equation contextual tabs.

As you can see in Figure 21.25, the Equation tab looks complicated, but one quickly gets the hang of it. The idea is to "write" your equation by simply choosing the symbols you need from the options shown in Figure 21.25.

FIGURE 21.25

Use the symbols on the Equation tab to write your equation.



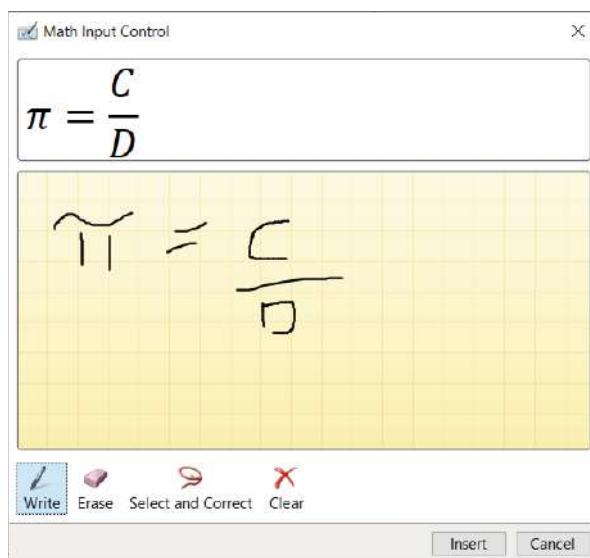
Chapter 21: Visualizing with Custom Number Formats and Shapes

Generally, you add a structure and then edit the various parts by adding text or symbols. You can put structures inside structures, and there is no limit to the complexity of the equations. Because the equation objects you work with are similar to the other shape objects you've covered, it won't take long to get the hang of building your own equations.

If all that seems much too daunting, you can use the nifty Ink Equation feature, which allows you to simply draw your equation. Select Tools \Rightarrow Ink Equation from the Equation tab to display the Math Input Control dialog box (see Figure 21.26). Here, you can draw your equation and have it translated to text. As you can see in Figure 21.26, Excel is quite accepting of even the shabbiest of penmanship.

FIGURE 21.26

You can save time and simply draw your equation.



Part IV

Managing and Analyzing Data

Excel is a superb data-analysis tool—if you know how to extract the information that you really need. In this part, you'll learn how to obtain, clean up, and analyze data in Excel. As you'll see, many of the data-analysis capabilities in Excel are both surprisingly powerful and easy to use.

IN THIS PART

- Chapter 22**
Importing and Cleaning Data
- Chapter 23**
Using Data Validation
- Chapter 24**
Creating and Using Worksheet Outlines
- Chapter 25**
Linking and Consolidating Worksheets
- Chapter 26**
Introducing PivotTables
- Chapter 27**
Analyzing Data with PivotTables
- Chapter 28**
Performing Spreadsheet What-If Analysis
- Chapter 29**
Analyzing Data Using Goal Seeking and Solver
- Chapter 30**
Analyzing Data with the Analysis ToolPak
- Chapter 31**
Protecting Your Work

Importing and Cleaning Data

IN THIS CHAPTER

Importing data into Excel

Manipulating and cleaning data

Using Flash Fill to extract and concatenate data

Reviewing a checklist for data cleaning

Exporting data to other formats

Data is everywhere. For example, if you run a website, you're collecting data continually, and you may not even know it. Every visit to your site generates information that is stored in a file on your server. This file contains lots of useful information if you take the time to examine it.

That's just one example of data collection. Virtually every automated system collects data and stores it. Most of the time, the system that collects the data is also equipped to verify and analyze the data—but not always. And, of course, data is also collected manually. A nonautomated telephone survey is a good example.

Excel is a good tool for analyzing data, and it's often used to summarize the information and display it in the form of tables and charts. But often, the data that's collected isn't perfect. For one reason or another, it needs to be cleaned up before it can be analyzed.

Excel is a great tool for cleaning up data. Cleaning data involves getting raw data into a worksheet and then manipulating it so that it conforms to your requirements. In the process, the data will be made consistent so that it can be properly analyzed.

This chapter describes the various ways to get data into a worksheet, and it provides some tips to help you clean it up.

Importing Data

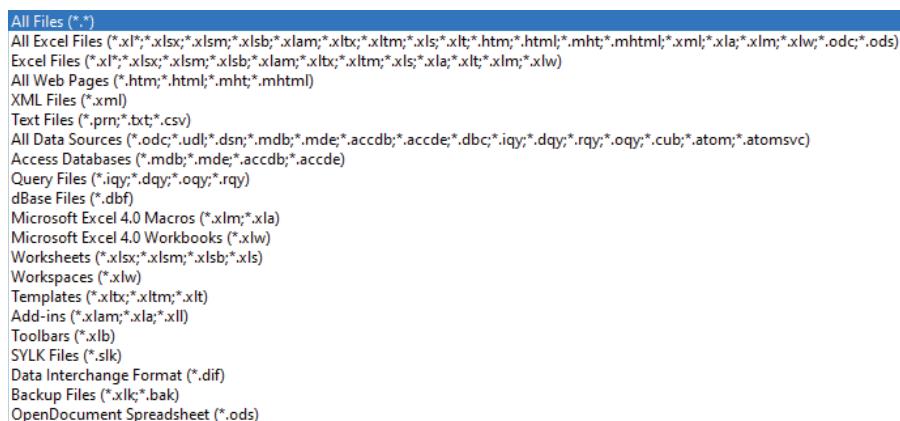
The first step is to get the data into a worksheet. Excel can import most common text file formats, including data from websites.

Importing from a file

The following sections describe the file types that Excel can open directly, using the File \Rightarrow Open command. Figure 22.1 shows the list of file filter options that you can specify in the Open dialog box.

FIGURE 22.1

Filtering by file extension in the Open dialog box



Spreadsheet file formats

In addition to the current file formats (XLSX, XLSM, XLSB, XLTX, XLTM, and XLAM), Excel can open workbook files from all previous versions of Excel:

XLS: Binary files created by Excel 4, Excel 95, Excel 97, Excel 2000, Excel 2002, and Excel 2003

XLM: Binary files that contain Excel 4 macros (no data)

XLT: Binary files for an Excel template

XLA: Binary files for an Excel add-in

Excel can also open ODS files (the OpenDocument Spreadsheet format) created by other spreadsheet products. ODS files are produced by a variety of “open” software programs, including Google Sheets, Apache OpenOffice, LibreOffice, and several others.

Database file formats

Excel can open the following database file formats:

Access files: The names of these files have various extensions, including .mdb and .accdb.

dBase files: These files are produced by dBase III and dBase IV and have names with a .dbf extension. Excel does not support dBase II files.

When you try to open database files using File ➔ Open, Excel doesn't actually open the file. Instead, it creates an external data connection to the table in the database that you select. Excel supports various types of database connections that enable you to access data selectively. For example, instead of "opening" a database and selecting a table, you can perform a query on a table to retrieve only the records that you need (rather than the entire table).

Text file formats

A text file contains raw characters with no formatting. Excel can open most types of text files:

CSV: This stands for "comma-separated values." Columns are delimited with a comma, and rows are delimited with a carriage return.

TXT: Columns are delimited with a tab, and rows are delimited with a carriage return.

PRN: Columns are delimited with multiple space characters, and rows are delimited with a carriage return. Excel imports this type of file into a single column.

DIF: This file format (DIF stands for Data Interchange Format) was originally used by the VisiCalc spreadsheet. This is rarely used.

SYLK: This file format (SYLK stands for Symbolic Link and the filenames have the .slk extension) was originally used by Multiplan. This is rarely used.

Most of these text file types have variants. For example, text files produced on a Mac have different end-of-row characters. Excel can usually handle the variants without a problem.

When you attempt to open a text file in Excel, the Text Import Wizard might kick in to help you specify how you want the data to be retrieved.

TIP

To bypass the Text Import Wizard, press Shift while you click the Open button in the Open dialog box.

When Excel Can't Open a File

If Excel doesn't support a particular file format, don't be too quick to give up. It's likely that others have had the same problem as you. Try searching the Web for the file extension, plus the word *Excel*. It's possible that a file converter is available, or maybe someone has figured out how to use an intermediary program to open the file and export it into a format that Excel recognizes.

HTML files

Excel can open most HTML files, which can be stored on your local drive or on a web server. The way that the HTML code renders in Excel varies considerably. Sometimes, the HTML file may look exactly as it does in a browser. At other times, it may bear little resemblance, especially if the HTML file uses Cascading Style Sheets (CSS) for its layout.

Part IV: Managing and Analyzing Data

In some cases, you can access data on the Web by using Power Query. We discuss this topic in Part V, “Understanding Power Pivot and Power Query.”

XML files

Extensible Markup Language (XML) is a text file format suitable for structured data. Data is enclosed in tags, which also serve to describe the data.

Excel can open XML files and simpler ones will be displayed with little or no effort. Complex XML files will require some work, however. A discussion of this topic is beyond the scope of this book. You’ll find information about getting data from XML files in Excel’s Help system and online.

Importing vs. opening

When you use File \Rightarrow Open to open a file that’s not in a traditional Excel format, you may be opening the file, or you may be importing it, depending on the file type. As mentioned, database files aren’t opened. Rather, a table from inside the database file is imported.

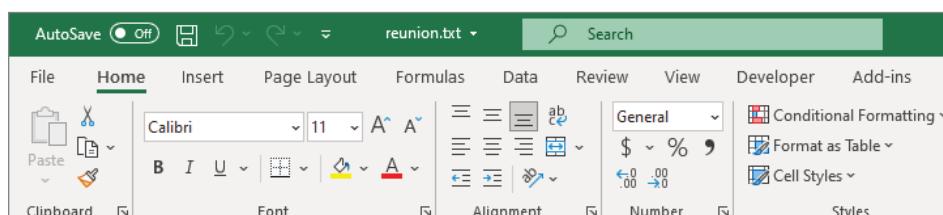
XML files are another example of a format that may not open directly. When you open an XML file, you have the choice of opening it as a read-only workbook or importing it into a table.

Text files are opened directly. Excel understands CSV files, so it doesn’t have to ask you any questions when you open them. For tab-delimited or fixed-width files, the Text Import Wizard will guide you in identifying where the data begins and ends.

When a file is opened directly, the Excel title bar will show the file’s name. Figure 22.2 shows the Excel title bar after a file named `reunion.txt` was opened. When the File \Rightarrow Open process actually imports the data, it’s imported into a new workbook. In that case, the title bar will have a generic new workbook name, like Book1.

FIGURE 22.2

Excel’s title bar displays the opened file’s name.



TIP

Prior to Excel 2019, when you worked with a CSV file in Excel, Excel would alert you that you were not working with a native Excel file every time you tried to save the file. If you added formats or formulas, for example, they would be lost if you saved the file as a CSV. Thankfully, Excel no longer prompts you for this change every time. Beginning with Excel 2019, an alert about possible data loss is presented when you open the file, and you can tell it not to show the alert again.

Importing a text file

One of the advantages of importing a text file instead of opening it is that you can put the data into a specific range in a worksheet rather than starting in cell A1. In this example, we'll show you how to import a text file to a specific range, and we'll walk through the steps of the Text Import Wizard.

Text files can be delimited or fixed-width. In delimited files, a character (like a comma or a tab) separates the columns. In fixed-width, each column occupies the same number of characters, and spaces are used to pad the data. For example, a fixed-width file whose first column is 20 characters long means that the second column always starts at the 21st character. The example in this section imports a file delimited by commas, but the process for importing fixed width is similar.

Beginning with Excel 2019, importing text files is done through Get & Transform rather than the legacy Text Import Wizard. Get & Transform is a powerful feature, and we discuss it at length in Part V, “Understanding Power Pivot and Power Query.” For this example, however, we’re going to use the legacy wizard. With all of Get & Transform’s power, it took away some flexibility like not having a header row. It’s valuable to know both ways of importing text files.

Before we begin, we have to enable the legacy wizard. To do that, choose File \Rightarrow Options \Rightarrow Data and check From Text (Legacy), as shown in Figure 22.3. This will add the necessary menu item for the next step.

Figure 22.4 shows a small CSV file. The following instructions describe how to import this file, named monthly.csv, beginning at cell C3.

1. Choose Data \Rightarrow Get & Transform Data \Rightarrow Get Data \Rightarrow Legacy Wizards \Rightarrow From Text (Legacy). The Import Text File dialog box appears.
2. Navigate to the folder that contains the text file.
3. Select the file from the list and then click the Import button. The Text Import Wizard appears, as shown in Figure 22.5.
4. Select Delimited and make sure that the My Data Has Headers check box is unchecked. Click Next to go to step 2.
5. Choose Comma as the delimiter and uncheck any others (see Figure 22.6).
6. Click the Finish button. The Import Data dialog box appears, as shown in Figure 22.7.
7. In the Import Data dialog box, specify the location for the imported data. It can be a cell in an existing worksheet or a new worksheet.
8. Click OK, and Excel imports the data (see Figure 22.8).

Part IV: Managing and Analyzing Data

FIGURE 22.3

Enabling the legacy import wizard

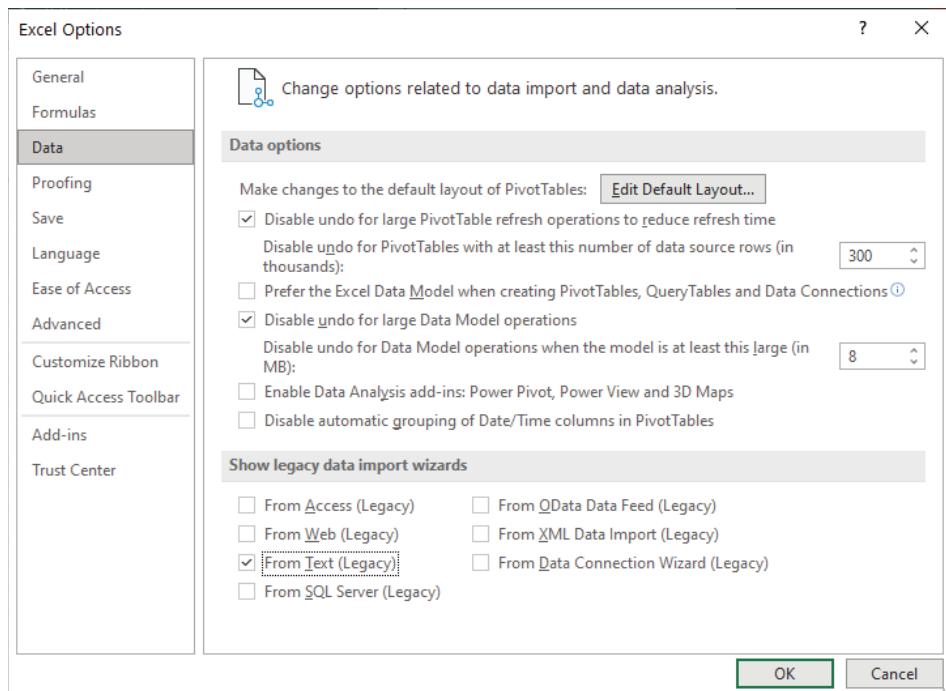


FIGURE 22.4

This CSV file will be imported.

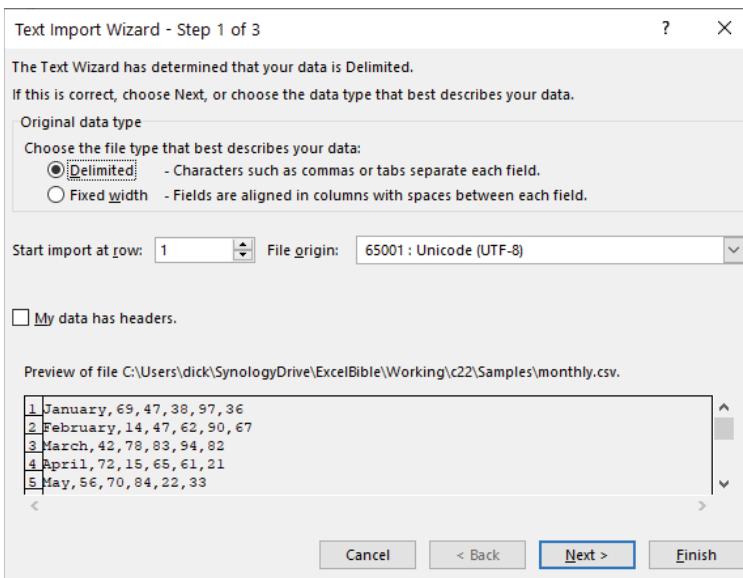
The screenshot shows a Notepad window titled 'monthly.csv - Notepad'. The window contains the following CSV data:

Month	Value 1	Value 2	Value 3	Value 4
January	69	47	38	97
February	14	47	62	90
March	42	78	83	94
April	72	15	65	61
May	56	70	84	22
June	37	38	97	39
July	18	96	76	91
August	13	59	79	43
September	70	36	33	77
October	29	58	25	99
November	56	83	88	39
December	93	99	52	33

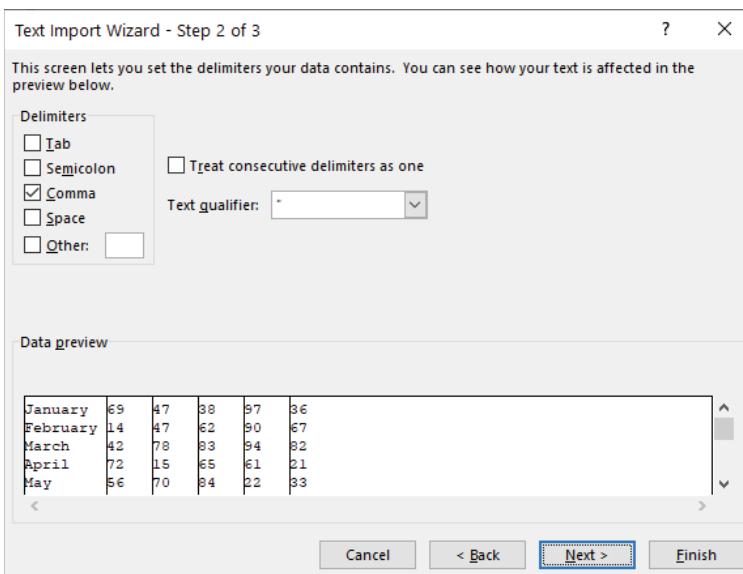
At the bottom of the Notepad window, status bar text includes 'Ln 1, Col 1', '100%', 'Windows (CRLF)', and 'UTF-8 with BOM'.

FIGURE 22.5

Step 1 of the Text Import Wizard

**FIGURE 22.6**

Select the delimiter in step 2 of the Text Import Wizard.



Part IV: Managing and Analyzing Data

FIGURE 22.7

Using the Import Data dialog box to import a CSV file

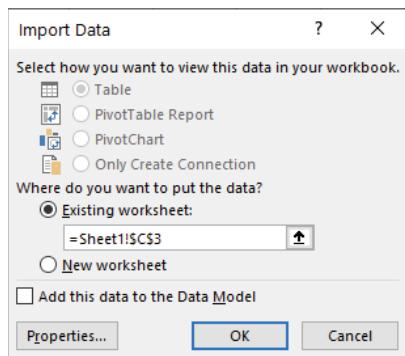


FIGURE 22.8

This range contains data imported directly from a CSV file.

	A	B	C	D	E	F	G	H	I
1									
2									
3		January	69 47 38 97 36						
4		February	14 47 62 90 67						
5		March	42 78 83 94 82						
6		April	72 15 65 61 21						
7		May	56 70 84 22 33						
8		June	37 38 97 39 34						
9		July	18 96 76 91 88						
10		August	13 59 79 43 80						
11		September	70 36 33 77 69						
12		October	29 58 25 99 87						
13		November	56 83 88 39 24						
14		December	93 99 52 33 97						
15									

Copying and pasting data

If all else fails, you can try standard copy-and-paste techniques. If you can copy data from an application (for example, a word-processing program or a document displayed in a PDF viewer), there's a good chance that you can paste it into an Excel workbook. For the best results, try pasting using the Home \Rightarrow Clipboard \Rightarrow Paste \Rightarrow Paste Special command, and try the various paste options listed. Usually, pasted data will require some cleanup.

Cleaning Up Data

The following sections discuss a variety of techniques that you can use to clean up data in a worksheet.

→ Chapter 12, “Using Formulas to Manipulate Text,” contains additional examples of text-related formulas that may be helpful when cleaning data.

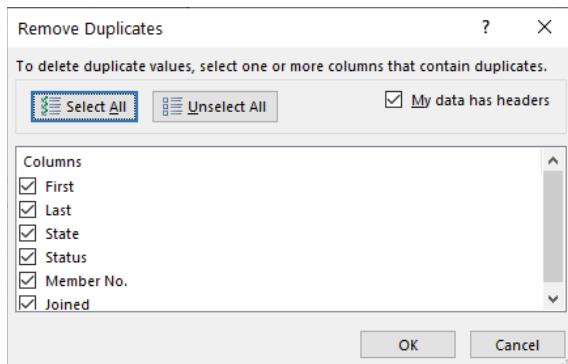
Removing duplicate rows

If data is compiled from multiple sources, it may contain duplicate rows. Most of the time, you will want to eliminate the duplicates. In the past, removing duplicate data was essentially a manual task—although it could be automated by using a confusing advanced filter technique. But now removing duplicate rows is easy, thanks to Excel’s Remove Duplicates command (introduced in Excel 2007).

Start by selecting any cell within your data range. Choose Data ➔ Data Tools ➔ Remove Duplicates, and the Remove Duplicates dialog box, shown in Figure 22.9, appears.

FIGURE 22.9

Use the Remove Duplicates dialog box to delete duplicate rows.



Note

If your data is in a table, you can also use Table Design ➔ Tools ➔ Remove Duplicates. These two commands work the same.

The Remove Duplicates dialog box lists all the columns in your data range or table. Place a check mark next to the columns that you want to be included in the duplicate search. Most of the time, you’ll want to select all the columns, which is the default. Click OK, and Excel weeds

out the duplicate rows and displays a message that tells you how many duplicates it removed. If Excel deleted too many rows, you can undo the procedure by clicking Undo (or by pressing Ctrl+Z).

When you select all columns in the Remove Duplicates dialog box, Excel will delete a row only if the content of every column is duplicated. In some situations, you may not care about matching some columns, so you would deselect those columns in the Remove Duplicates dialog box. For example, if each row has a unique ID code, Excel would never find any duplicate rows. So, you'd want to uncheck that column in the Remove Duplicates dialog box.

When duplicate rows are found, the first row is kept, and subsequent duplicate rows are deleted.

CAUTION

Duplicate values are determined by the value displayed in the cell—not necessarily the value stored in the cell. For example, assume that two cells contain the same date. One of the dates is formatted to be displayed as 5/15/2022, and the other is formatted to be displayed as May 15, 2022. When removing duplicates, Excel considers these dates to be different. Similarly, values that are formatted differently are considered to be different. For example, \$1,209.32 is not the same as 1209.32. Therefore, you might want to apply formatting to entire columns to ensure that duplicate rows are not overlooked just because of a formatting difference.

Identifying duplicate rows

If you would like to identify duplicate rows so that you can examine them without automatically deleting them, here's another method. Unlike the technique described in the previous section, this method looks at actual values, not formatted values.



A workbook that contains the examples in this section is available on this book's website at www.wiley.com/go/excel1365bible. The filename is remove duplicates.xlsx.

Create a formula to the right of your data that concatenates each of the cells to the left. The formulas that follow assume that the data is in columns A:F. The TEXTJOIN function in this example separates the columns with a pipe character (usually located above the Enter key on your keyboard).

Enter this formula into cell G2:

```
=TEXTJOIN(" | ", FALSE, A2:F2)
```

Add another formula in cell H2. This formula displays the number of times a duplicate value in column G occurs:

```
=COUNTIF(G:G, G2)
```

Copy these formulas down the column for each row of your data.

Column H displays the number of occurrences of that row. Unduplicated rows will display 1. Duplicated rows will display a number that corresponds to the number of times that row appears.

Figure 22.10 shows a simple example. If you don't care about a particular column, just omit it from the formula in column G. For example, if you want to find duplicates regardless of the Status column, change the formula in G2 to the following:

```
=TEXTJOIN(" | ", FALSE, A2 : C2, E2 : F2)
```

FIGURE 22.10

Using formulas to identify duplicate rows

	A	B	C	D	E	F	G	H	I
1	First	Last	State	Status	Member No.	Joined			
2	Joshua	Elliott	NJ	Active	99-3325	12/21/1909	Joshua Elliott NJ Active 99-3325 3643	1	
3	Brian	Dunn	MA	Inactive	25-7251	3/22/2007	Brian Dunn MA Inactive 25-7251 39163	1	
4	Sara	Diaz	MA	Inactive	62-6451	2/22/1976	Sara Diaz MA Inactive 62-6451 27812	2	
5	Adrian	Cunningham	MS	Active	72-1993	3/18/1939	Adrian Cunningham MS Active 72-1993 14322	1	
6	Zoey	Watkins	WA	Inactive	19-3615	12/30/1968	Zoey Watkins WA Inactive 19-3615 25202	1	
7	Hayden	Nichols	AR	Inactive	50-7290	5/31/1980	Hayden Nichols AR Inactive 50-7290 29372	2	
8	Julia	Walker	LA	Active	37-3660	2/25/1913	Julia Walker LA Active 37-3660 4805	1	
9	Cole	Carpenter	MN	Active	82-3786	8/18/1914	Cole Carpenter MN Active 82-3786 5344	1	
10	Brooklyn	Stone	MT	Active	43-4609	2/15/1967	Brooklyn Stone MT Active 43-4609 24518	1	
11	Adrian	Warren	NY	Inactive	85-4686	6/5/1941	Adrian Warren NY Inactive 85-4686 15132	1	
12	Adam	Ross	AK	Inactive	75-6928	6/22/1951	Adam Ross AK Inactive 75-6928 18801	1	
13	Brian	Jones	NM	Inactive	46-9552	7/14/1961	Brian Jones NM Inactive 46-9552 22476	1	

Splitting text

When importing data, you might find that multiple values are imported into a single column. Figure 22.11 shows an example of this type of import problem.

TIP

The data in Figure 22.11 is displayed in a fixed-width font (Courier New). With the default font, it was not apparent that the data lined up nicely in fixed-width columns.

If the text is all the same length (as in the example), you might be able to write a series of formulas that extract the information to separate columns. The LEFT, RIGHT, and MID functions are useful for this task.

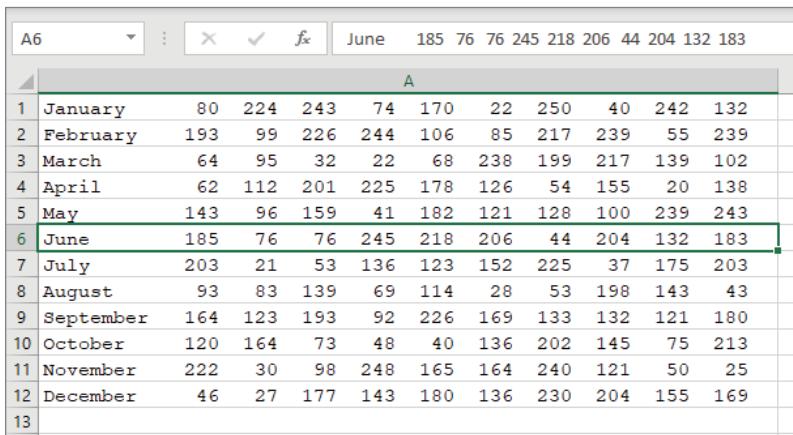


See Chapter 12 for examples of formulas that extract characters from text.

You should also be aware that Excel offers two nonformula methods to assist in splitting data so that it occupies multiple columns: Text to Columns and Flash Fill.

FIGURE 22.11

The imported data was put in one column rather than multiple columns.



A screenshot of a Microsoft Excel spreadsheet. The data is located in column A, starting at cell A6. The first few rows show the months from January to June with their corresponding values. Row 6, which contains the month "June" and its values, is highlighted with a green border. The data continues through December in row 12. The top of the screen shows the ribbon tabs and the formula bar with cell A6 selected. The data is as follows:

	June	185	76	76	245	218	206	44	204	132	183
1	January	80	224	243	74	170	22	250	40	242	132
2	February	193	99	226	244	106	85	217	239	55	239
3	March	64	95	32	22	68	238	199	217	139	102
4	April	62	112	201	225	178	126	54	155	20	138
5	May	143	96	159	41	182	121	128	100	239	243
6	June	185	76	76	245	218	206	44	204	132	183
7	July	203	21	53	136	123	152	225	37	175	203
8	August	93	83	139	69	114	28	53	198	143	43
9	September	164	123	193	92	226	169	133	132	121	180
10	October	120	164	73	48	40	136	202	145	75	213
11	November	222	30	98	248	165	164	240	121	50	25
12	December	46	27	177	143	180	136	230	204	155	169
13											

Using Text to Columns

The Text to Columns command can parse strings into their component parts.



This workbook is available on this book's website at www.wiley.com/go/excel365bible. The filename is `text to columns.xlsx`.

First, make sure the column that contains the data to be split up has enough empty columns to the right to accommodate the extracted data. Then select the data to be parsed and choose **Data** \Rightarrow **Data Tools** \Rightarrow **Text to Columns**. Excel displays the Convert Text to Columns Wizard, which consists of a series of dialog boxes that walk you through the steps to convert a single column of data into multiple columns. Figure 22.12 shows the initial step, in which you choose the type of data:

Delimited The data to be split is separated by delimiters such as commas, spaces, slashes, or other characters.

Fixed Width Each component occupies the same number of characters.

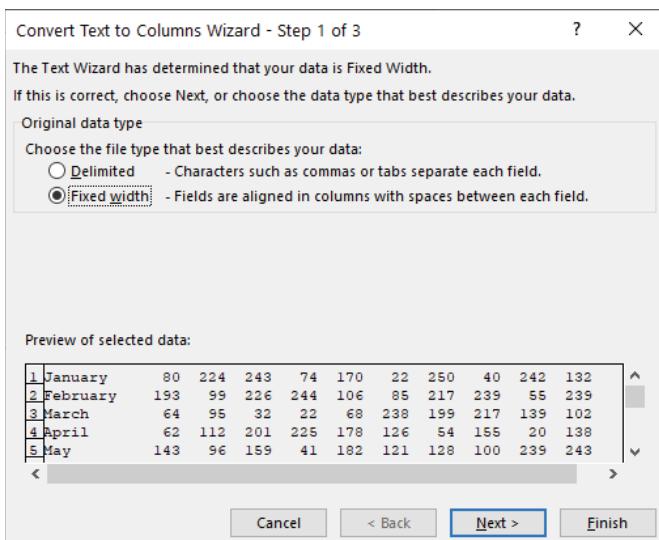
Make your choice and click **Next** to move on to step 2, which depends on the choice you made in step 1.

If you're working with delimited data, specify the delimiting character. You'll see a preview of the result. If you're working with fixed-width data, specify the column breaks directly in the preview window.

When you're satisfied with the column breaks, click **Next** to move to step 3. In this step, you can click a column in the preview window and specify formatting for the column. For example, if you have data that looks like a number but is really text, you can format the column as **Text** so that you preserve any leading zeros. Click **Finish**, and Excel splits the data as specified.

FIGURE 22.12

The first dialog box in the Convert Text to Columns Wizard



Using Flash Fill

The Text to Columns Wizard works well for many types of data. But sometimes you'll encounter data that can't be parsed by that wizard. For example, the Text to Columns Wizard is useless if you have variable-width data that doesn't have delimiters. In such a case, the Flash Fill feature might save the day. But keep in mind that Flash Fill works successfully only when the data is consistent.

Flash Fill uses pattern recognition to extract data (and also concatenate data). Just enter a few examples in a column that's adjacent to the data and choose Data \Rightarrow Data Tools \Rightarrow Flash Fill (or press Ctrl+E). Excel analyzes the examples and attempts to fill in the remaining cells. If Excel didn't recognize the pattern you had in mind, press Ctrl+Z, add another example or two, and try again.

By default, Excel has the Automatically Flash Fill option turned on. You can find that option under File \Rightarrow Options \Rightarrow Advanced \Rightarrow Automatically Flash Fill. If Excel can recognize the pattern, it will automatically fill the range. It usually recognizes the pattern when you type whole words containing only letters. If there are numbers or special characters, you have to invoke Flash Fill manually.

Figure 22.13 shows a worksheet with some text in a single column. The goal is to extract the numeric value from each text string and put the number into a separate cell. The Text to Columns Wizard can't do it because the space delimiters aren't consistent. It might be possible to write an array formula, but it would be complicated.

Part IV: Managing and Analyzing Data

FIGURE 22.13

The goal is to extract the numbers from column A.

	A	B
1	The box weighed 20 pounds.	
2	Bob is 6 feet tall.	
3	She drove for 9.5 hours straight.	
4	Pi is 3.14159	
5	He drank 5 cups of coffee.	
6	The sales tax was \$3.12 for that item.	
7	15 people showed up for jury duty	
8	He was in 7th heaven.	
9	The square root of 16 is four.	
10	Where is the 90210 zip code?	
11		



This workbook, which also includes other Flash Fill examples, is available on this book's website at www.wiley.com/go/excel365bible. The filename is `flash fill demo.xlsx`.

To try using Flash Fill, select cell B1 and type the first number (**20**). Move to B2 and type the second number (**6**). Can Flash Fill identify the remaining numbers and fill them in? Choose Data \Rightarrow Data Tools \Rightarrow Flash Fill (or press Ctrl+E), and Excel fills in the remaining cells in a flash. Figure 22.14 shows the result.

FIGURE 22.14

Using manually entered examples in B1 and B2, Excel's Flash Fill feature makes some incorrect guesses.

	A	B
1	The box weighed 20 pounds.	20
2	Bob is 6 feet tall.	6
3	She drove for 9.5 hours straight.	5
4	Pi is 3.14159	14159
5	He drank 5 cups of coffee.	5
6	The sales tax was \$3.12 for that item.	12
7	15 people showed up for jury duty	15
8	He was in 7th heaven.	7
9	The square root of 16 is four.	16
10	Where is the 90210 zip code?	90210
11		

As you can see, Excel identified most of the values. Accuracy increases if you provide more examples. For example, provide a decimal number. Delete the suggested values in column B, enter **3.12** in cell B6, and press Ctrl+E. This time, Flash Fill gets all of them correct (see Figure 22.15).

FIGURE 22.15

After you enter an example of a decimal number, Excel gets all of the values correct.

	A	B	
1	The box weighed 20 pounds.	20	
2	Bob is 6 feet tall.	6	
3	She drove for 9.5 hours straight.	9.5	
4	Pi is 3.14159	3.14159	
5	He drank 5 cups of coffee.	5	
6	The sales tax was \$3.12 for that item.	3.12	
7	15 people showed up for jury duty	15	
8	He was in 7th heaven.	7	7
9	The square root of 16 is four.	16	
10	Where is the 90210 zip code?	90210	
11			

This simple example demonstrates two important points:

- You must examine your data carefully after using Flash Fill. Just because the first few rows are correct, you can't assume that Flash Fill worked correctly for all rows.
- Flash Fill accuracy increases when you provide more examples.

Figure 22.16 shows another example: names in column A. The goal is to extract the first, last, and middle names (if it has one). In column B, Flash Fill successfully gets all the first names using only two examples (Mark and Tim). Plus, it successfully extracted all the last names (column C) using Russell and Colman. Extracting the middle names or initials (column D) didn't work until examples that included a space on either side of the middle name were included.

FIGURE 22.16

Using Flash Fill to split names

	A	B	C	D
1	Mark Russell	Mark	Russell	
2	Tim Colman	Tim	Colman	
3	Sam Daniel Bains	Sam	Bains	Daniel
4	Fred James Foster	Fred	Foster	James
5	James J. Wehr	James	Wehr	J.
6	Mitch Nicholls	Mitch	Nicholls	
7	Neal McCaslin	Neal	McCaslin	
8	Ned Poulakis	Ned	Poulakis	
9	Paul T. Wingfield	Paul	Wingfield	T.
10	Peter Gans	Peter	Gans	
11	Ron E. Hoffman	Ron	Hoffman	E.
12	Julia Hayes	Julia	Hayes	
13	Richard P Light	Richard	Light	P
14	Ray Walker	Ray	Walker	
15	Robert F. Mahaney	Robert	Mahaney	F.
16	Robert Fist	Robert	Fist	
17				

To summarize, Excel's Flash Fill is an interesting idea, but it works reliably only if the data is consistent. Even when you think it worked correctly, make sure that you examine the results carefully. And think twice before trusting it with important data because there's no way to document the way the data was extracted. But the main limitation is that (unlike formulas) Flash Fill is not a dynamic technique. If your data changes, the flash-filled columns do not update.

NOTE

You can also use the Flash Fill feature to create new data from multiple columns. Just provide a few examples of the way that you want the data to be combined, and Excel will figure out the pattern and fill in the column. Using Flash Fill to create data seems to work much better than using it to extract data. But then again, it's also easier to create formulas that create data from existing columns.

Changing the case of text

Often, you'll want to make text in a column consistent in terms of case. Excel provides no direct way to change the case of text, but it's easy to do with formulas. (See the sidebar "Transforming Data with Formulas.")

The three relevant functions are as follows:

UPPER: This converts the text to ALL UPPERCASE.

LOWER: This converts the text to all lowercase.

PROPER: This converts the text to Proper Case. (The first letter in each word is capitalized, as in a proper name.)

These functions are quite straightforward. They operate only on alphabetic characters and just ignore all other characters and return them unchanged.

If you use the PROPER function, you'll probably need to do some additional cleanup to handle exceptions. The following are examples of transformations that you probably would consider incorrect:

- The letter following an apostrophe is always capitalized (for example, Don'T). This is done, apparently, to handle names like O'Reilly.
- The PROPER function doesn't handle names with an embedded capital letter, such as McDonald.
- "Minor" words such as *and* as well as *the* are always capitalized. For example, some people would prefer that the third word in United States Of America not be capitalized.

You can correct some of these problems by using Find and Replace.

Transforming Data with Formulas

Many of the data cleanup examples in this chapter describe how to use formulas and functions to transform data in some way. For example, you can use the `UPPER` function to transform text into uppercase. When the data is transformed, you'll have two columns: the original data and the transformed data. Almost always, you'll want to replace the original data with the transformed data. Here's how to do it:

1. Insert a new temporary column for formulas to transform the original data.
2. Create your formulas in the temporary column, and make sure that the formulas do what they were intended to do.
3. Select the formula cells.
4. Choose Home \Rightarrow Clipboard \Rightarrow Copy (or press **Ctrl+C**).
5. Select the original data cells.
6. Choose Home \Rightarrow Clipboard \Rightarrow Paste \Rightarrow Values (V).

This procedure replaces the original data with the transformed data. Then you can delete the temporary column that holds the formulas.

Removing extra spaces

It's usually a good idea to ensure that data doesn't have extra spaces. It's difficult to spot a space character at the end of a text string. Extra spaces can cause lots of problems, especially when you need to compare text strings. The text *July* is not the same as the text *July* with a space appended to the end. The first is four characters long, and the second is five characters long.

The `TRIM` function removes all leading and trailing spaces and replaces interior multiple spaces with a single space. This formula uses the `TRIM` function. The formula returns **Fourth Quarter Earnings (with no excess spaces)**:

```
=TRIM("    Fourth    Quarter    Earnings    ")
```

Data that is imported from a web page often contains a different type of space: a nonbreaking space, indicated by ` ` in HTML code. In Excel, this character can be generated by this formula:

```
=CHAR(160)
```

You can use a formula like this to replace those spaces with normal spaces:

```
=SUBSTITUTE(A2, CHAR(160), " ")
```

Or you can use this formula to replace the nonbreaking space character with normal spaces and to remove excess spaces:

```
=TRIM(SUBSTITUTE(A2, CHAR(160), " "))
```

Removing strange characters

Often, data imported into an Excel worksheet contains strange (sometimes unprintable) characters. You can use the CLEAN function to remove all nonprinting characters from a string. If the data is in cell A2, this formula will do the job:

```
=CLEAN(A2)
```

NOTE

The CLEAN function can miss some nonprinting Unicode characters. It's programmed to remove the first 32 nonprinting characters in the 7-bit ASCII code. Consult the Excel Help system for information on how to remove the nonprinting Unicode characters. (Search Help for the CLEAN function.)

Converting values

In some cases, you may need to convert values from one system to another. For example, you may import a file that has values in fluid ounces and they need to be expressed in milliliters. Excel's handy CONVERT function can perform that and many other conversions.

If cell A2 contains a value in ounces, the following formula converts it to milliliters:

```
=CONVERT(A2, "oz", "ml")
```

This function is extremely versatile and can handle most common measurement units in the following categories: weight and mass, distance, time, pressure, force, energy, power, magnetism, temperature, volume, liquid, area, bits and bytes, and speed.

Excel can also convert between number bases. You may import a file that contains hexadecimal values, and you need to convert them to decimal. Use the HEX2DEC function to perform this conversion. For example, the following formula returns 1,279, the decimal equivalent of its hex argument:

```
=HEX2DEC("4FF")
```

Excel can also convert from binary to decimal (BIN2DEC) and from octal to decimal (OCT2DEC).

Functions that convert from decimal to another number base are DEC2HEX, DEC2BIN, and DEC2OCT and can be found in the Engineering category.

The BASE function converts a decimal number to any number base. Note that there is not a function that works in the opposite direction. Excel does not provide a function that converts any number base to decimal. You're limited to binary, octal, and hexadecimal.

Classifying values

Often, you may have values that need to be classified into a group. For example, if you have ages of people, you might want to classify them into groups such as 17 or younger, 18–24, 25–34, and so on.

The easiest way to perform this classification is with a lookup table. Figure 22.17 shows ages in column A and classifications in column B. Column B uses the lookup table in D2 : E9. The formula in cell B2 is as follows:

=VLOOKUP (A2 , \$D\$2 : \$E\$9 , 2)

FIGURE 22.17

Using a lookup table to classify ages into age ranges

	A	B	C	D	E	F
1	Age	Classification				
2	24	18-24		0	<18	
3	42	35-44		18	18-24	
4	44	35-44		25	25-34	
5	17	<18		35	35-44	
6	72	65-74		45	45-54	
7	51	45-54		55	55-64	
8	40	35-44		65	65-74	
9	51	45-54		75	75+	
10	34	25-34				
11	51	45-54				
12	81	75+				
13	18	18-24				
14	46	45-54				
15	60	55-64				
16	32	25-34				
17						

This formula was copied to the cells below.

You can also use a lookup table for nonnumeric data. Figure 22.18 shows a lookup table that is used to assign a region to a state.

The two-column lookup table is in the range D2 : E52. The formula in cell B2, which was copied to the cells below, is as follows:

=VLOOKUP (A2 , \$D\$2 : \$E\$52 , 2 , FALSE)



A workbook that contains the examples in this section is available on this book's website at www.wiley.com/go/excel365bible. The filename is classifying_data.xlsx.

TIP

A side benefit is that the VLOOKUP function will return #N/A if an exact match is not found—a good way to spot misspelled states, in this example. Using FALSE as the last argument in the function indicates that an exact match is required.

FIGURE 22.18

Using a lookup table to assign a region for a state

A	B	C	D	E	F
1 State	Region				
2 California	Region IX	Alabama	Region IV		
3 Iowa	Region VII	Alaska	Region X		
4 Missouri	Region VII	Arizona	Region IX		
5 Nevada	Region IX	Arkansas	Region VI		
6 New Hampshire	Region I	California	Region IX		
7 Illinois	Region V	Colorado	Region VIII		
8 New Mexico	Region VI	Connecticut	Region I		
9 California	Region IX	Delaware	Region III		
10 Oregon	Region X	District of Columbia	Region III		
11 Utah	Region VIII	Florida	Region IV		
12 Connecticut	Region I	Georgia	Region IV		
13		Hawaii	Region IX		
14		Idaho	Region X		
15		Illinois	Region V		
16		Indiana	Region V		
17		Iowa	Region VII		

Joining columns

To combine data from two or more columns, you can use the CONCAT function in a formula. For example, the following formula combines the contents of cells A1, B1, and C1:

```
=CONCAT(A1:C1)
```



A workbook that contains the examples in this section is available on this book's website at www.wiley.com/go/excel365bible. The filename is concatenating.xlsx.

Often, you'll need to insert spaces, or some other delimiter, between the cells—for example, if the columns contain a title, first name, and last name. Concatenating using the previous formulas would produce something like Mr.ThomasJones. To add spaces (to produce Mr. Thomas Jones), use the TEXTJOIN function:

```
=TEXTJOIN(" ", TRUE, A1:C1)
```

The first argument of TEXTJOIN is the delimiter that you want to insert between the cell values. The second argument is TRUE to ignore empty cells. If you set the second argument to FALSE and there are empty cells, you'll end up with two delimiters right next to each other.

Figure 22.19 shows three examples of TEXTJOIN. In the first example, there are no empty cells, so the second argument doesn't matter. In the second and third examples, the second argument is set to FALSE and TRUE, respectively, and the delimiter is changed from a space to a comma (so it's easier to see the duplication). Where empty cells are not ignored, two commas are shown together.

FIGURE 22.19

The TEXTJOIN function inserts delimiters between cell values.

	A	B	C	D	E	F	G
1	Title	First	Last	Ignore Empty?	Joined		
2	Mr	Thomas	Jones		Mr Thomas Jones	=TEXTJOIN(" ",TRUE,A2:C2)	
3	Mr		Jones	FALSE	Mr,,Jones	=TEXTJOIN(", ",D3,A3:C3)	
4	Mr		Jones	TRUE	Mr,Jones	=TEXTJOIN(", ",D4,A4:C4)	
5							
6							

You can also use the Flash Fill feature (discussed earlier in this chapter) to join columns without using formulas. Just provide an example or two in an adjacent column and press Ctrl+E. Excel will perform the concatenation for the other rows.

Rearranging columns

If you need to rearrange the columns in a worksheet, you could insert a blank column and then drag another column into the new blank column. But the moved column leaves a gap, which you need to delete.

Here's an easier way:

1. Click the column header of the column you want to move.
2. Choose Home \Rightarrow Clipboard \Rightarrow Cut.
3. Click the column header to the right of where you want the column to go.
4. Right-click and choose Insert Cut Cells from the shortcut menu.

Repeat these steps until the columns are in the order you want.

Randomizing the rows

If you need to arrange the rows in random order, here's a quick way to do it. In the column to the right of the data, insert this formula into the first cell and copy it down:

=RAND()

Then sort the data using this column as the sort key. The rows will be in random order, and you can delete the column.

Extracting a filename from a URL

In some cases, you may have a list of URLs and need to extract only the filename. The following formula returns the filename from a URL. Assume that cell A2 contains this URL:

http://example.com/assets/images/horse.jpg

The following formula returns horse.jpg:

```
=RIGHT(A2,LEN(A2)-FIND("/",SUBSTITUTE(A2,"/","*",LEN(A2)-LEN(SUBSTITUTE(A2,"/","")))))
```

Part IV: Managing and Analyzing Data

This formula returns all text that follows the last slash character. If cell A2 doesn't contain a slash character, the formula returns an error.

To extract the URL without the filename, use this formula:

```
=LEFT(A2, FIND ("*", SUBSTITUTE(A2, "/", "*",
LEN(A2) -LEN(SUBSTITUTE(A2, "/", "")))))
```

NOTE

This type of extraction is a good use for the Flash Fill feature. (See “Splitting text” earlier in this chapter.)

Matching text in a list

You may have some data that you need to check against another list. For example, you may want to identify the data rows in which data in a particular column appears in a different list. Figure 22.20 shows a simple example. The data is in columns A:C. The goal is to identify the rows in which the Member Num appears in the Resigned Members list in column F. These rows can then be deleted.

FIGURE 22.20

The goal is to identify member numbers that are in the Resigned Members list in column F.

	A	B	C	D	E	F	G
1	Name	Member Num	State			Resigned Members	
2	Alice Jones	39-5954	AZ			11-6587	
3	Jennifer Green	46-2010	UT			16-4523	
4	Rhoda Davis	93-1595	AZ			16-8075	
5	Rita Morris	35-5121	WA			21-5865	
6	Debra Hopkins	91-2687	UT			23-5078	
7	Marcela Garcia	93-4652	AZ	Resigned		36-9582	
8	Viola Jenkins	74-4701	CA			39-2953	
9	Charlotte Baker	21-5865	CA	Resigned		40-8172	
10	Angela Gonzalez	79-8010	AZ			42-6818	
11	Michelle Young	93-7380	WA			45-8343	
12	Linda Johnson	16-6377	AZ			58-2363	
13	Annette Williamson	94-2032	CA			58-8192	
14	Ruth McKinney	58-8192	WA	Resigned		65-3095	
15	Mary Gibson	27-3637	CO			67-5960	
16	Christine Warren	81-8640	AZ			78-4209	
17	Stacey Martin	82-8709	CO			78-8201	
18	Shirley Clarke	99-6607	AZ			81-1158	
19	Rosemary Ross	16-8075	CO	Resigned		86-7291	
20	Waltraud Adams	55-5367	AZ			87-2700	
21	Nancy Martinez	82-4869	CA			93-4652	
22	Dominique Jackson	28-9592	AZ			97-2586	
23	Deanne Elliott	14-3518	WA				
24	Vanessa Hill	31-8125	UT				



This workbook, named `match_names.xlsx`, is available on this book's website at www.wiley.com/go/excel365bible.

Here's a formula, entered into cell D2 and copied down, that will do the job:

```
=IF(COUNTIF($F$2:$F$22,B2)>0,"Resigned","")
```

This formula displays the word *Resigned* if the member number in column B is found in the Resigned Members list. If the member number is not found, it returns an empty string. If the list is sorted by column D, the rows for all resigned members will appear together and can be quickly deleted.

This technique can be adapted to other types of list-matching tasks.

Changing vertical data to horizontal data

Figure 22.21 shows a common type of data layout that you might see when importing a file. Each record consists of three consecutive cells in a single column: Name, Department, and Location. The goal is to convert this data so that each record appears in three columns.

FIGURE 22.21

Vertical data that needs to be converted to three columns

	A	B
1	Lori Howard	
2	Accounting	
3	Main Office	
4	Jacqueline Espinoza	
5	Sales	
6	West Branch	
7	Jose Collins	
8	Accounting	
9	Main Office	
10	Bernice Ryan	
11	Sales	
12	East Branch	
13	Diana Brown	
14	Accounting	
15	West Branch	
16	Jay J. Davis	
17	Research	
18	Main Office	
19	Shandi Johnson	
20	Accounting	
21	Main Office	
22		

Part IV: Managing and Analyzing Data

There are several ways to convert this type of data, but here's a method that's fairly easy. It requires a small amount of setup, but the work is done with a single formula, which is copied to a range.

Start by creating some numeric vertical and horizontal "headers," as shown in Figure 22.22. Column C contains numbers that correspond to the first row of each data item (in this case, the name). In this example, we put the following values in column C: 1, 4, 7, 10, 13, 16, and 19. You can use a simple formula to generate this series of numbers.

FIGURE 22.22

Headers that are used to convert the vertical data into rows

	A	B	C	D	E	F	G
1	Lori Howard			1	2	3	
2	Accounting		1				
3	Main Office		4				
4	Jacqueline Espinoza		7				
5	Sales		10				
6	West Branch		13				
7	Jose Collins		16				
8	Accounting		19				
9	Main Office						
10	Bernice Ryan						
11	Sales						
12	East Branch						

The horizontal range of headers consists of consecutive integers, starting with 1. In this example, each record contains three cells of data, so the horizontal header contains 1, 2, and 3.



This workbook, named `vertical data.xlsx`, is available on this book's website at www.wiley.com/go/excel365bible.

Here's the formula that goes into cell D2:

`=OFFSET(A1,$C2+D$1-2,0)`

Copy this formula across to the next two columns and down to the next six rows. The result is shown in Figure 22.23.

You can easily adapt this technique to work with vertical data that contains a different number of rows. For example, if each record contained 10 rows of data, the column C header values would be 1, 11, 21, 31, and so on. The horizontal headers would consist of values 1 through 10 rather than 1 through 3.

Notice that the formula uses an absolute reference to cell A1. That reference won't change when the formula is copied, so all the formulas use cell A1 as the base. If the data begins in a different cell, change \$A\$1 to the address of the first cell.

FIGURE 22.23

A single formula transforms the vertical data into rows.

	A	B	C	D	E	F	G
1	Lori Howard			1	2	3	
2	Accounting			1 Lori Howard	Accounting	Main Office	
3	Main Office			4 Jacqueline Espinoza	Sales	West Branch	
4	Jacqueline Espinoza			7 Jose Collins	Accounting	Main Office	
5	Sales			10 Bernice Ryan	Sales	East Branch	
6	West Branch			13 Diana Brown	Accounting	West Branch	
7	Jose Collins			16 Jay J. Davis	Research	Main Office	
8	Accounting			19 Shandi Johnson	Accounting	Main Office	
9	Main Office						
10	Bernice Ryan						
11	Sales						
12	East Branch						
13	Diana Brown						
14	Accounting						
15	West Branch						
16	Jay J. Davis						
17	Research						
18	Main Office						
19	Shandi Johnson						
20	Accounting						
21	Main Office						
22							

The formula also uses “mixed” referencing in the second argument of the `OFFSET` function. The C2 reference has a dollar sign in front of C, so column C is the absolute part of the reference. In the D1 reference, the dollar sign is before the 1, so row 1 is the absolute part of the reference.



See Chapter 9, “Introducing Formulas and Functions,” for more about using mixed references in formulas.

Filling gaps in an imported report

When you import data, you can sometimes end up with a worksheet that looks something like the one shown in Figure 22.24. This type of report formatting is common. As you can see, an entry in column A applies to several rows of data. If you sort this type of list, the missing data messes things up, and you can no longer tell who sold what when.

Part IV: Managing and Analyzing Data

FIGURE 22.24

This report contains gaps in the Sales Rep column.

	A	B	C	D	E
1					
2	Sales Rep	Month	Units Sold	Amount	
3	Jane	Jan	182	\$15,101	
4		Feb	3350	\$34,230	
5		Mar	114	\$9,033	
6	George	Jan	135	\$8,054	
7		Feb	401	\$9,322	
8		Mar	357	\$32,143	
9	Beth	Jan	509	\$29,239	
10		Feb	414	\$38,993	
11		Mar	53	\$309	
12	Dan	Jan	323	\$9,092	
13		Feb	283	\$12,332	
14		Mar	401	\$32,933	
15					

If the report is small, you can enter the missing cell values manually or by using a series of Home \Rightarrow Editing \Rightarrow Fill \Rightarrow Down commands (or the Ctrl+D shortcut). But if you have a large list that's in this format, here's a better way:

1. **Select the range that has the gaps (A3 : A14, in this example).**
2. **Choose Home \Rightarrow Editing \Rightarrow Find & Select \Rightarrow Go To Special.** The Go To Special dialog box appears.
3. **Select the Blanks option and click OK.** This action selects the blank cells in the original selection.
4. **In the formula bar, type an equal sign (=) followed by the address of the first cell with an entry in the column (=A3, in this example) and press Ctrl+Enter.**
5. **Reselect the original range and press Ctrl+C to copy the selection.**
6. **Choose Home \Rightarrow Clipboard \Rightarrow Paste \Rightarrow Paste Values to convert the formulas to values.**

After you complete these steps, the gaps are filled in with the correct information, and your worksheet looks similar to the one shown in Figure 22.25.



This workbook, named `fill in gaps.xlsx`, is available on this book's website at www.wiley.com/go/excel365bible.

FIGURE 22.25

The gaps are gone, and this list can now be sorted.

1	A	B	C	D	E
2	Sales Rep	Month	Units Sold	Amount	
3	Jane	Jan	182	\$15,101	
4	Jane	Feb	3350	\$34,230	
5	Jane	Mar	114	\$9,033	
6	George	Jan	135	\$8,054	
7	George	Feb	401	\$9,322	
8	George	Mar	357	\$32,143	
9	Beth	Jan	509	\$29,239	
10	Beth	Feb	414	\$38,993	
11	Beth	Mar	53	\$309	
12	Dan	Jan	323	\$9,092	
13	Dan	Feb	283	\$12,332	
14	Dan	Mar	401	\$32,933	
15					

Checking spelling

If you use a word-processing program, you probably take advantage of its spell-checker feature. Spelling mistakes can be embarrassing when they appear in a text document, but they can cause serious problems when they occur within your data. For example, if you tabulate data by month, a misspelled month name will make it appear that a year has 13 months.

To access the Excel spell-checker, choose Review \Rightarrow Proofing \Rightarrow Spelling, or press F7. To check the spelling in just a particular range, select the range before you activate the spell-checker.

If the spell-checker finds any words it doesn't recognize as correct, it displays the Spelling dialog box. The options are fairly self-explanatory.



See Chapter 17, “Making Your Formulas Error-Free,” for more about the Spelling dialog box.

Replacing or removing text in cells

You may need to replace (or remove) certain characters systematically in a column of data. For example, you may need to replace all backslash characters with forward slash characters. In many cases, you can use Excel's Find and Replace dialog box to accomplish this task. To remove text using the Find and Replace dialog box, just leave the Replace With field empty.

In other situations, you may need a formula-based solution. Consider the data shown in Figure 22.26. The goal is to replace the second hyphen character with a colon for the part numbers in column A. Using Find and Replace wouldn't work because there isn't a way to specify that only the second hyphen should be replaced.

Part IV: Managing and Analyzing Data

FIGURE 22.26

To replace only the second hyphen in these cells, Find and Replace is not an option.

	B5			
	A	B	C	D
1	Part Number	Modified		
2	ADC-983-2	ADC-983:2		
3	BG-8832-3	BG-8832:3		
4	QERP-9832-1	QERP-9832:1		
5	OPY-093-2	OPY-093:2		
6	RGNP-9932-4	RGNP-9932:4		
7	BB-221-2	BB-221:2		
8	PDR-9322-3	PDR-9322:3		
9				

In this case, the solution is a fairly simple formula that replaces the second occurrence of a hyphen with a colon:

=SUBSTITUTE(A2, "-", ":", 2)



A workbook that contains the examples in this section is available on this book's website at www.wiley.com/go/excel365bible. The filename is `substitute.xlsx`.

To remove the second occurrence of a hyphen, just omit the third argument for the SUBSTITUTE function:

=SUBSTITUTE(A2, "-", , 2)

This is another example where Flash Fill can also do the job.

NOTE

If you've worked with programming languages, you may be familiar with the concept of regular expressions. A regular expression is a way to match strings of text using concise (and often confusing) codes. Excel doesn't support regular expressions, but if you search the Web, you'll find ways to incorporate regular expressions in VBA, plus a few add-ins that provide this feature in the workbook environment.

Adding text to cells

If you need to add text to a cell, one solution is to use a new column of formulas. Here are some examples:

- The following formula adds ID: and a space to the beginning of a cell:

= "ID: " & A2

- The following formula adds .mp3 to the end of a cell:

= A2 & ".mp3"

- The following formula inserts a hyphen after the third character in a cell:

```
=LEFT(A2, 3) & "-" &RIGHT(A2, LEN(A2)-3)
```

You can also use the Flash Fill feature to add text to cells.

Fixing trailing minus signs

Imported data sometimes displays negative values with a trailing minus sign. For example, a negative value may appear as 3,498- rather than the more common -3,498. Excel does not convert these values. In fact, it considers them to be nonnumeric text.

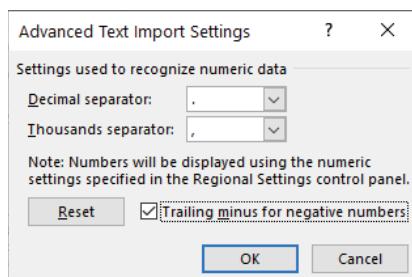
The solution is so simple it may even surprise you:

- Select the data that has the trailing minus signs.** The selection can also include positive values.
- Choose Data ➔ Data Tools ➔ Text to Columns.** The Text to Columns dialog box appears.
- Click Finish.**

This procedure works because of a default setting in the Advanced Text Import Settings dialog box (which you don't even see normally). To display this dialog box, which is shown in Figure 22.27, go to step 3 in the Text to Columns Wizard dialog box and click Advanced.

FIGURE 22.27

The Trailing Minus for Negative Numbers option makes it easy to fix trailing minus signs in a range of data.



Following a data cleaning checklist

This section contains a list of items that could cause problems with data. Not all of these are relevant to every set of data.

- Does each column have a unique and descriptive header?
- Is each column of data formatted consistently?
- Did you check for duplicate or missing rows?

- For text data, are the words consistent in terms of case?
- Did you check for spelling errors?
- Does the data contain any extra spaces?
- Are the columns arranged in the proper (or logical) order?
- Are there any blank cells that shouldn't be blank?
- Did you correct any trailing minus signs?
- Are the columns wide enough to display all data?

Exporting Data

This chapter began with sections on importing data, so it's only appropriate to end it with a discussion of exporting data to a file that's not a standard Excel file.

Exporting to a text file

When you choose File \leftrightarrow Save As, the Save As dialog box lets you choose from a variety of file formats. The three text file types are as follows:

CSV: Comma-separated value files

TXT: Tab-delimited files

PRN: Formatted text

We discuss these file types in the sections that follow.

CSV files

When you export a worksheet to a CSV file, the data is saved as displayed. In other words, if a cell contains 12.8312344 but is formatted to display with two decimal places, the value will be saved as 12.83.

Cells are delimited with a comma character, and rows are delimited with a carriage return and line feed.

NOTE

If you export a file using the Mac variant, rows are delimited with a carriage return only (no line feed character).

Note that if a cell contains a comma, the cell value is saved within quotation marks. If a cell contains a quotation mark character, that character appears twice.

TXT files

Exporting a workbook to a TXT file is almost identical to the CSV file format described earlier. The only difference is that cells are separated by a tab character rather than a comma.

If your worksheet contains any Unicode characters, you should export the file using the Unicode variant. Otherwise, Unicode characters will be saved as question mark characters.

PRN files

A PRN file is very much like a printed image of the worksheet. The cells are separated by multiple space characters. Also, a line is limited to 240 characters. If a line exceeds that limit, the remainder appears on the next line. PRN files are rarely used.

Exporting to other file formats

Excel also lets you save your work in several other formats:

Data Interchange Format These files have a .dif filename extension. These are not used very often.

Symbolic Link These files have an .slk filename extension. These are not used very often.

Portable Document Format These files have a .pdf filename extension. This is a common “read-only” file format.

XML Paper Specification Document These files have an .xps filename extension. This is Microsoft’s alternative to PDF files. It is not used very often.

Web Page These files have an .htm filename extension. Often, saving a file as a web-page will generate a directory of ancillary files required to render the page accurately.

OpenDocument Spreadsheet These files have an .ods filename extension. They’re compatible with various open-source spreadsheet programs.

Using Data Validation

IN THIS CHAPTER

Getting an overview of Excel's data validation feature

Looking at practical examples of using data validation formulas

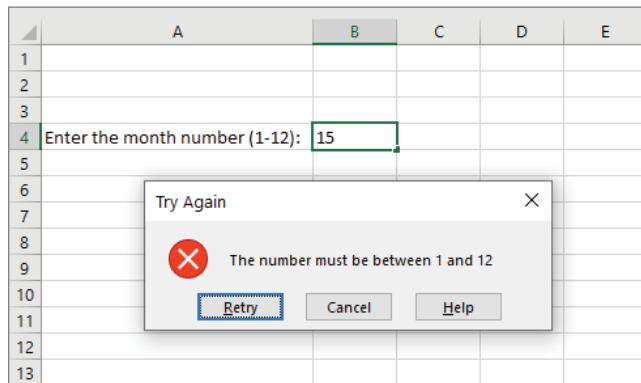
This chapter explores a useful Excel feature: data validation. Data validation enables you to add rules for what's acceptable in specific cells, and it allows you to add dynamic elements to your worksheet without using any macro programming.

About Data Validation

The Excel data validation feature allows you to set up rules that determine what can be entered into a cell. For example, you may want to limit data entry in a particular cell to whole numbers between 1 and 12. If the user makes an invalid entry, you can display a custom message, such as the one shown in Figure 23.1.

FIGURE 23.1

Displaying a message when the user makes an invalid entry



Part IV: Managing and Analyzing Data

Excel makes it easy to specify the validation criteria. You can also use a formula for more complex criteria.

CAUTION

The Excel data validation feature suffers from a potentially serious problem: if the user copies a cell that does not use data validation and pastes it into a cell that does use data validation, the data validation rules are deleted. In other words, the cell then accepts any type of data. This has always been a problem, and Microsoft still hasn't fixed it in Excel 2019.

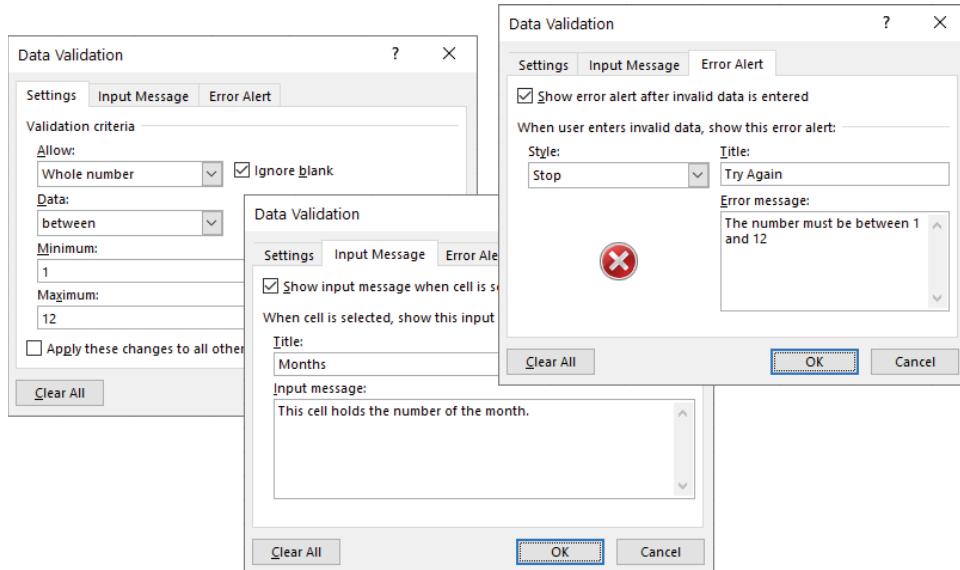
Specifying Validation Criteria

To specify the type of data allowable in a cell or range, follow these steps:

1. Select the cell or range of cells.
2. Choose Data ➦ Data Tools ➦ Data Validation. The Data Validation dialog box appears, as shown in Figure 23.2.

FIGURE 23.2

The three tabs of the Data Validation dialog box



3. **Select the Settings tab.**
4. **Choose an option from the Allow drop-down list.** The contents of the Data Validation dialog box change, displaying controls based on your choice. For example, to specify a formula, select Custom.
5. **Specify the conditions by using the displayed controls.** Your selection in step 4 determines the other controls that you can access.
6. **(Optional) Select the Input Message tab and specify which message to display when a user selects the cell.** You can use this optional step to tell the user what type of data is expected or to document the purpose of the cell. If this step is omitted, no message will appear when the user selects the cell.
7. **(Optional) Select the Error Alert tab and specify which error message to display when a user makes an invalid entry.** The selection for Style determines what choices users have when they make invalid entries. For example, choose Stop to prevent an invalid entry. If this step is omitted, a standard message will appear if the user makes an invalid entry.

CAUTION

Even with data validation in effect, a user can enter invalid data. If the Style setting on the Error Alert tab of the Data Validation dialog box is set to Warning or Information, invalid data can be entered. You can identify invalid entries by having Excel circle them (explained in the next section).

8. **Click OK.** The cell or range contains the validation criteria you specified.

Types of Validation Criteria You Can Apply

From the Settings tab of the Data Validation dialog box, you can specify a variety of data validation criteria. The following options are available from the Allow drop-down list. Keep in mind that the other controls on the Settings tab vary, depending on your choice from the Allow drop-down list.

Any Value Selecting this option removes any existing data validation. Note, however, that the input message, if any, still is displayed if the box is checked on the Input Message tab.

Whole Number The user must enter a whole number. You specify a valid range of whole numbers by using the Data drop-down list. For example, you can specify that the entry must be a whole number greater than or equal to 100.

Decimal The user must enter a number. You specify a valid range of numbers by refining the criteria from choices in the Data drop-down list. For example, you can specify that the entry must be between 0 and 1.

List The user must choose from a list of entries that you provide. Type a comma-separated list of values in the Source text box. This option is useful, and we discuss it in detail later in this chapter. (See “Creating a Drop-Down List.”)

Date The user must enter a date. You specify a valid date range from the choices in the Data drop-down list. For example, you can specify that the data entered must be greater than or equal to January 1, 2022.

Time The user must enter a time. You specify a valid time range from choices in the Data drop-down list. For example, you can specify that the data entered must be greater than 12:00 PM.

Text Length The length of the data (number of characters) is limited. You specify a valid length by using the Data drop-down list and the Length text box. For example, you can specify that the length of the data entered be 1 (a single alphanumeric character).

Custom To use this option, you must supply a logical formula that determines the validity of the user’s entry. (A logical formula returns either TRUE or FALSE.) You can enter the formula directly into the Formula control (which appears when you select the Custom option), or you can specify a cell reference that contains a formula. This chapter contains examples of useful formulas.

The Settings tab of the Data Validation dialog box contains these other check boxes:

Ignore Blank If this is selected, blank entries are not circled as invalid when using Data \Rightarrow Data Tools \Rightarrow Data Validation \Rightarrow Circle Invalid Data.

In-Cell Dropdown If you select List in the Allow drop-down list, you can choose to show or hide a drop-down arrow in the cell to aid the user in selecting a valid value.

Apply These Changes to All Other Cells with the Same Settings If this is selected, the changes that you make apply to all other cells that contain the original data validation criteria.

TIP

The Data \Rightarrow Data Tools \Rightarrow Data Validation drop-down list contains an item named Circle Invalid Data. When you select this item, circles appear around cells that contain incorrect entries. If you correct an invalid entry, the circle disappears. This is useful if you don’t use the Stop error style or if you apply data validation to cells that already have values. To get rid of the circles, choose Data \Rightarrow Data Tools \Rightarrow Data Validation \Rightarrow Clear Validation Circles. In Figure 23.3, valid entries are defined as values between 5 and 105. Values that are not within this numerical range are circled.

FIGURE 23.3

Excel can draw circles around invalid entries (in this case, cells that aren't between 5 and 105).

	A	B	C	D	E	F	G	H	I	J
1	0	3	105	75	67	9	0	83	14	
2	7	32	82	23	55	76	86	89	76	
3	104	101	5	24	74	89	70	51	-8	
4	44	16	70	16	-6	54	69	101	30	
5	106	-10	90	58	3	25	-2	51	80	
6	87	86	27	108	98	24	85	79	59	
7	57	40	59	43	40	29	102	-2	75	
8	58	77	46	35	26	74	71	57	22	
9	80	94	5	41	93	86	37	43	12	
10	4	79	65	45	78	12	105	108	9	
11	39	35	61	95	-2	1	98	12	34	
12	35	8	49	-3	15	8	27	53	22	
13	101	32	91	84	41	58	30	22	13	
14	-2	39	100	1	60	80	30	39	46	
15	67	88	31	7	71	0	6	40	94	
16	-2	40	39	55	81	94	20	-9	72	
17	2	-7	84	84	66	87	54	63	47	
18										

Creating a Drop-Down List

One of the most common uses of data validation is to create a drop-down list in a cell. Figure 23.4 shows an example that uses region names in cells A2 : A6 as the list source.

To create a drop-down list in a cell, follow these steps:

1. **Enter the list items into a single-row or single-column range.** These items will appear in the drop-down list.
2. **Select the cell that will contain the drop-down list and then choose Data ➔ Data Tools ➔ Data Validation.** The Data Validation dialog box appears.
3. **From the Settings tab, select the List option (from the Allow drop-down list) and specify the range that contains the list, using the Source control.** The range can be in a different worksheet, but it must be in the same workbook.

TIP

If you will be adding new items to the list, put the items in a single-column table that was created by using Insert ➔ Tables ➔ Table. Excel will then update the list of options when you add or remove items from the table column.

FIGURE 23.4

This drop-down list was created using data validation.

	A	B	C	D	E
1	Regions				
2	Northeast		Choose a Region:		
3	Southeast			Northeast	
4	Midwest			Southeast	
5	Northwest			Midwest	
6	Southwest			Northwest	
7				Southwest	

4. Make sure that the In-Cell Dropdown check box is selected.
5. Set any other Data Validation options as desired.
6. Click OK. The cell displays an input message (if specified) and a drop-down arrow when it's selected.
7. Click the arrow and choose an item from the list that appears.

TIP

If you have a short list, you can enter the items directly into the Source control on the Settings tab of the Data Validation dialog box. (This control appears when you choose the List option in the Allow drop-down list.) Just separate each item with list separators specified in your regional settings (a comma if you use the U.S. regional settings).

Unfortunately, you cannot control the font size used in drop-down lists. If the cell that displays the drop-down is formatted to show large text, the drop-down list does not use that formatting. If you zoom out on a worksheet, it may be difficult to read the items.

Using Formulas for Data Validation Rules

For simple data validation, the data validation feature is quite straightforward and easy to use. The real power of this feature, though, becomes apparent when you use the Custom option and supply a formula.

The formula that you specify must be a logical formula that returns either TRUE or FALSE. If the formula evaluates to TRUE, the data is considered valid and remains in the cell. If the formula evaluates to FALSE, a message box appears that displays the message that you specify on the Error Alert tab of the Data Validation dialog box.

Specify a formula in the Data Validation dialog box by selecting the Custom option from the Allow drop-down list of the Settings tab. Enter the formula directly into the Formula control or enter a reference to a cell that contains a formula. The Formula control appears on the Settings tab of the Data Validation dialog box when the Custom option is selected.

We present several examples of formulas used for data validation in the section “Data Validation Formula Examples” later in this chapter.

Understanding Cell References

If the formula that you enter in the Data Validation dialog box contains a cell reference, that reference is considered a relative reference, based on the upper-left cell in the selected range.

The following example clarifies this concept. Suppose you want to allow only an odd number to be entered into the range B2 : B10. None of the Excel data validation rules can limit entry to odd numbers, so a formula is required.

Follow these steps:

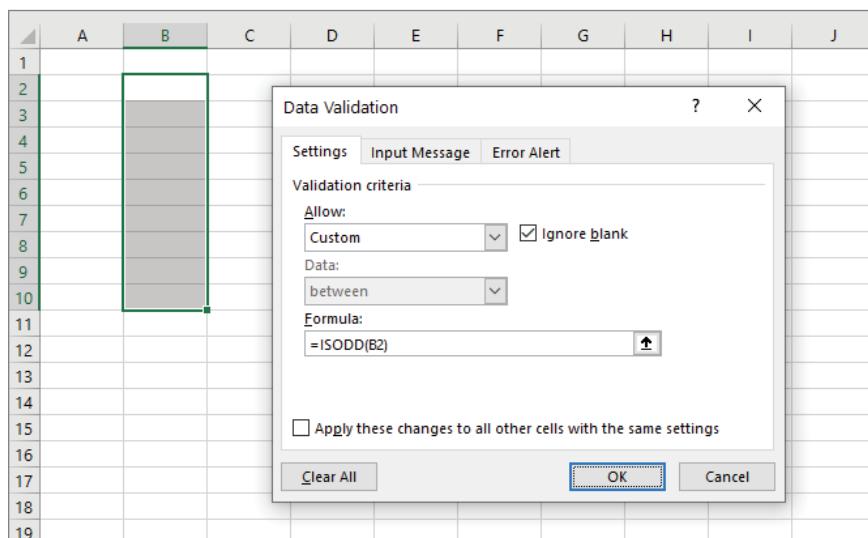
- Select the range (B2 : B10 in this example) and make sure that cell B2 is the active cell.**
- Choose Data ➔ Data Tools ➔ Data Validation.** The Data Validation dialog box appears.
- Select the Settings tab and select the Custom option (from the Allow drop-down list).**
- Enter the following formula in the Formula field, as shown in Figure 23.5:**

=ISODD(B2)

23

FIGURE 23.5

Entering a data validation formula



This formula uses the `ISODD` function, which returns TRUE if its numeric argument is an odd number. Notice that the formula refers to the active cell, which is cell B2.

5. On the Error Alert tab, choose Stop for the Style and then type An odd number is required here in the Error Message field.
6. Click OK to close the Data Validation dialog box.

Notice that the formula entered contains a reference to the upper-left cell in the selected range. This data validation formula was applied to a range of cells, so you might expect that each cell would contain the same data validation formula. Because you entered a relative cell reference as the argument for the `ISODD` function, Excel adjusts the formula for the other cells in the B2 : B10 range. To demonstrate that the reference is relative, select cell B5 and examine its formula displayed in the Data Validation dialog box. You'll see that the formula for this cell is:

=`ISODD(B5)`

NOTE

An alternative method is to enter the logical formula into a cell and then enter a cell reference in the Formula field in the Data Validation dialog box. For this example, cell C2 would contain =`ISODD(B2)`, and that formula would be copied down the column to cell C10. Then the Formula field in the Data Validation dialog box would have this formula: =`C2`. Most of the time, entering the formula into the Formula field is easier and more efficient.

Generally, when entering a data validation formula for a range of cells, you use a reference to the active cell, which is normally the upper-left cell in the selected range. An exception is when you need to refer to a specific cell. For example, suppose that you select range A1 : B10 and you want your data validation to allow only values that are greater than the value in cell C1. You would use this formula:

=`A1>C1`

In this case, the reference to cell C1 is an absolute reference; it will not be adjusted for the cells in the selected range, which is just what you want. The data validation formula for cell A2 looks like this:

=`A2>C1`

The relative cell reference is adjusted, but the absolute cell reference is not.

Data Validation Formula Examples

The following sections contain a few data validation examples that use a formula entered directly into the Formula control on the Settings tab of the Data Validation dialog box. These examples help you understand how to create your own data validation formulas.

CAUTION

All the examples in this section are available at this book's website at www.wiley.com/go/excel365bible. The file is named `data validation examples.xlsx`.

Accepting text only

Excel has a data validation option to limit the length of text entered into a cell, but it doesn't have an option to force text (rather than a number) into a cell. To force a cell or range to accept only text (no values), use the following data validation formula:

```
=ISTEXT(A1)
```

This formula assumes that the active cell in the selected range is cell A1.

Accepting a larger value than the previous cell

The following data validation formula enables the user to enter a value only if it's greater than the value in the cell directly above it:

```
=A2>A1
```

This formula assumes that A2 is the active cell in the selected range. Note that you can't use this formula for a cell in row 1.

Accepting nonduplicate entries only

The following data validation formula does not permit the user to make a duplicate entry in the range A1 : C20:

```
=COUNTIF($A$1:$C$20,A1)=1
```

This is a logical formula that returns TRUE if the value in the cell occurs only one time in the A1 : C20 range. Otherwise, it returns FALSE, and the Duplicate Entry dialog box is displayed.

This formula assumes that A1 is the active cell in the selected range. Note that the first argument for COUNTIF is an absolute reference. The second argument is a relative reference, and it adjusts for each cell in the validation range. Figure 23.6 shows this validation criterion in effect using a custom error alert message. The user is attempting to enter 12 into cell B5.

Accepting text that begins with a specific character

The following data validation formula demonstrates how to check for a specific character. In this case, the formula ensures that the user's entry is a text string that begins with the letter A (uppercase or lowercase):

```
=LEFT(A1)="a"
```

Part IV: Managing and Analyzing Data

FIGURE 23.6

Using data validation to prevent duplicate entries in a range

The screenshot shows a Microsoft Excel spreadsheet with data in columns A and B. Column A contains numbers from 1 to 20. Column B contains numbers from 32 to 51. Cell B5 contains the value 12. A data validation dialog box is open over the spreadsheet, titled "Data Validation" with the subtitle "Accepts only non-duplicated data". It displays a message "Duplicate Entry!" with a red X icon and the text "The entry you made already exists in range A1:C20. Please try again." There are three buttons at the bottom: "Retry" (highlighted with a blue border), "Cancel", and "Help".

	A	B	C	D	E	F	G	H	I	J
1	1	32								
2	2	33								
3	3	34								
4	4	35								
5	5	12								
6	6									
7	7									
8	8									
9	9									
10	10									
11	11									
12	12									
13	13									
14	14									
15	15									
16	16									
17	17									
18	18									
19	19									
20	20									
21										

This is a logical formula that returns TRUE if the first character in the cell is the letter A. Otherwise, it returns FALSE. This formula assumes that the active cell in the selected range is cell A1.

The following formula is a variation of this validation formula. It uses wildcard characters in the second argument of the COUNTIF function. In this case, the formula ensures that the entry begins with the letter A and contains exactly five characters:

```
=COUNTIF(A1, "A????")=1
```

Accepting dates by the day of the week

The following data validation formula assumes that the cell entry is a date, and it ensures that the date is a Monday:

```
=WEEKDAY(A1)=2
```

This formula assumes that the active cell in the selected range is cell A1. It uses the WEEKDAY function, which returns 1 for Sunday, 2 for Monday, and so on. Note that the WEEKDAY function accepts any non-negative value as an argument (not just dates).

Accepting only values that don't exceed a total

Figure 23.7 shows a simple budget worksheet, with the budget item amounts in the range B1 : B6. The planned budget is in cell E5, and the user is attempting to enter a value in cell B4 that would cause the total (cell E6) to exceed the budget. The following data validation formula ensures that the sum of the budget items does not exceed the budget:

=SUM(\$B\$1 : \$B\$6) <= \$E\$5

FIGURE 23.7

Using data validation to ensure that the sum of a range does not exceed a certain value

23

Creating a dependent list

As we described previously, you can use data validation to create a drop-down list in a cell (see “Creating a Drop-Down List” earlier in this chapter). This section explains how to use a drop-down list to control the entries that appear in a second drop-down list. In other words, the second drop-down list is dependent upon the value selected in the first drop-down list.

Figure 23.8 shows a simple example of a dependent list created by using data validation. Cell E2 contains data validation that displays a three-item list from the range A1 : C1 (Vegetables, Fruits, and Meats). When the user chooses an item from the list, the second list (in cell F2) displays the appropriate items.

This worksheet uses three named ranges:

Vegetables: A2 : A15

Fruits: B2 : B9

Meats: C2 : C5

Part IV: Managing and Analyzing Data

FIGURE 23.8

The items displayed in the list in cell F2 depend on the list item selected in cell E2.

	A	B	C	D	E	F	G
1	Vegetables	Fruits	Meats				
2	Arugula	Apple	Beef				
3	Asparagus	Cherry	Chicken				
4	Beet	Lemon	Lamb				
5	Broccoli	Orange	Pork				
6	Carrot	Peach					
7	Cauliflower	Pear					
8	Celery	Raspberry					
9	Corn	Strawberry					
10	Green Bean						
11	Onion						
12	Potato						
13	Radish						
14	Squash						
15	Turnip						
16							

Cell F2 contains data validation that uses this formula:

=INDIRECT(\$E\$2)

Therefore, the drop-down list displayed in F2 depends on the value displayed in cell E2.

Using Data Validation without Restricting Entry

The most common use of data validation is to prevent a user from entering invalid data. But data validation can also be used as a component of your spreadsheet's user interface without restricting the user's entry. The two examples of showing an input message and making suggestions would require quite a bit of VBA programming to accomplish, but it can be accomplished easily with data validation.

Showing an input message

Data validation provides a way to display a message when the user selects a cell. Normally, this message would tell the user what would be considered invalid data for that cell to prevent them from getting the error message if they enter invalid data. But you can use this message for anything.

Figure 23.9 shows an input message reminding the user to complete a previous step. The Allow drop-down list is set to Any Value, so the user is not prevented from entering anything into this cell. It's just a reminder set to a cell that the user will likely use early in the process of completing this workbook.

Making suggested entries

The default Style value on the Error Alert tab of the Data Validation dialog box is Stop. This not only displays a message, but it also prevents the user from completing the entry of invalid

data. There are two other options, Warning and Information, which will allow the user to enter any data. You can also deselect the check box and Excel will not show any message.

FIGURE 23.9

Data validation can be used to show messages to the user.

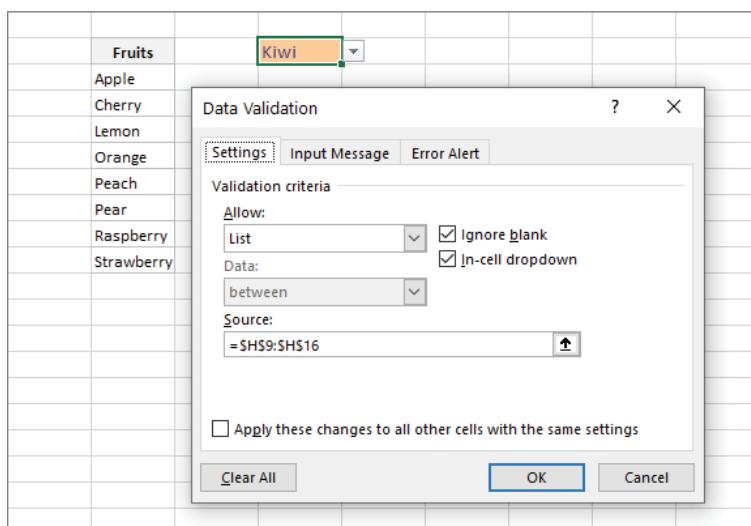
	A	B	C	D	E	F
1						
2						
3			Quarter:			
4						
5						
6						
7						
8						
9						
10						

Suppose that you have a field where the user is to enter the name of a fruit. To help the user, you provide a list of common fruits, but you want the user to be able to enter fruits that aren't on the list. The list is merely there to save the user some typing if they need to enter a common fruit.

To set up this scenario, select List and point to the list of fruits, as shown in Figure 23.10. Then uncheck the Show Error Alert After Invalid Data Is Entered check box on the Error Alert tab. Now the user can select from the list or type in a fruit that's not on the list. Of course, you can't prevent the user from typing in nonsense.

FIGURE 23.10

Data Validation dialog box



Creating and Using Worksheet Outlines

IN THIS CHAPTER

Introducing worksheet outlines

Creating an outline

Working with outlines

If you use a word processor, you may be familiar with the concept of an outline. Most word processors (including Microsoft Word) have an outline mode that lets you view only the headings and subheadings in your document. You can easily expand a heading to show the text below it. Using an outline makes visualizing the structure of your document easy.

Excel also has outlines. Understanding this feature can make working with certain types of worksheets much easier for you.

Introducing Worksheet Outlines

You'll find that some worksheets are more suitable for outlines than others. You can use outlines to create summary reports that don't show all the details. If your worksheet uses hierarchical data with subtotals, it's probably a good candidate for an outline.

The best way to understand how worksheet outlining works is to look at an example. Figure 24.1 shows a simple sales summary sheet without an outline. Formulas are used to calculate subtotals by region and by quarter.

Figure 24.2 shows the same worksheet after the outline was created by selecting only the rows and using Data ➔ Outline ➔ Group ➔ Auto Outline. Notice that Excel adds a new section to the left of the screen. This section contains outline symbols that enable you to determine which level to view. This particular outline has three levels: States, Regions (each region consists of states grouped into categories such as West, East, and Central), and Grand Total (the sum of each region's subtotal).

Figure 24.3 depicts the outline after clicking the level 2 symbol, which collapses everything below the second level of details. Now the outline shows only the totals for the regions. You can partially expand the outline to show the detail for a particular region by clicking one of the plus-sign buttons. Collapsing the outline to level 1 shows only the headers and the Grand Total row. To show all the rows after collapsing, click on the highest numbered symbol, level 3 in this example.

Part IV: Managing and Analyzing Data

FIGURE 24.1

A simple sales summary with subtotals

	A	B	C	D	E	F	G	H	I	J
1	State	Jan	Feb	Mar	Qtr-1	Apr	May	Jun	Qtr-2	Total
2	California	1,118	1,960	1,252	4,330	1,271	1,557	1,679	4,507	8,837
3	Washington	1,247	1,238	1,028	3,513	1,345	1,784	1,574	4,703	8,216
4	Oregon	1,460	1,954	1,726	5,140	1,461	1,764	1,144	4,369	9,509
5	Arizona	1,345	1,375	1,075	3,795	1,736	1,555	1,372	4,663	8,458
6	West Total	5,170	6,527	5,081	16,778	5,813	6,660	5,769	18,242	35,020
7	New York	1,429	1,316	1,993	4,738	1,832	1,740	1,191	4,763	9,501
8	New Jersey	1,735	1,406	1,224	4,365	1,706	1,320	1,290	4,316	8,681
9	Massachusetts	1,099	1,233	1,110	3,442	1,637	1,512	1,006	4,155	7,597
10	Florida	1,705	1,792	1,225	4,722	1,946	1,327	1,357	4,630	9,352
11	East Total	5,968	5,747	5,552	17,267	7,121	5,899	4,844	17,864	35,131
12	Kentucky	1,109	1,078	1,155	3,342	1,993	1,082	1,551	4,626	7,968
13	Oklahoma	1,309	1,045	1,641	3,995	1,924	1,499	1,941	5,364	9,359
14	Missouri	1,511	1,744	1,414	4,669	1,243	1,493	1,820	4,556	9,225
15	Illinois	1,539	1,493	1,211	4,243	1,165	1,013	1,445	3,623	7,866
16	Kansas	1,973	1,560	1,243	4,776	1,495	1,125	1,387	4,007	8,783
17	Central Total	7,441	6,920	6,664	21,025	7,820	6,212	8,144	22,176	43,201
18	Grand Total	18,579	19,194	17,297	55,070	20,754	18,771	18,757	58,282	113,352
19										

FIGURE 24.2

The worksheet after creating an outline

1 2 3		A	B	C	D	E	F	G	H	I	J
1	State	Jan	Feb	Mar	Qtr-1	Apr	May	Jun	Qtr-2	Total	
.	2	California	1,118	1,960	1,252	4,330	1,271	1,557	1,679	4,507	8,837
.	3	Washington	1,247	1,238	1,028	3,513	1,345	1,784	1,574	4,703	8,216
.	4	Oregon	1,460	1,954	1,726	5,140	1,461	1,764	1,144	4,369	9,509
.	5	Arizona	1,345	1,375	1,075	3,795	1,736	1,555	1,372	4,663	8,458
.	6	West Total	5,170	6,527	5,081	16,778	5,813	6,660	5,769	18,242	35,020
.	7	New York	1,429	1,316	1,993	4,738	1,832	1,740	1,191	4,763	9,501
.	8	New Jersey	1,735	1,406	1,224	4,365	1,706	1,320	1,290	4,316	8,681
.	9	Massachusetts	1,099	1,233	1,110	3,442	1,637	1,512	1,006	4,155	7,597
.	10	Florida	1,705	1,792	1,225	4,722	1,946	1,327	1,357	4,630	9,352
.	11	East Total	5,968	5,747	5,552	17,267	7,121	5,899	4,844	17,864	35,131
.	12	Kentucky	1,109	1,078	1,155	3,342	1,993	1,082	1,551	4,626	7,968
.	13	Oklahoma	1,309	1,045	1,641	3,995	1,924	1,499	1,941	5,364	9,359
.	14	Missouri	1,511	1,744	1,414	4,669	1,243	1,493	1,820	4,556	9,225
.	15	Illinois	1,539	1,493	1,211	4,243	1,165	1,013	1,445	3,623	7,866
.	16	Kansas	1,973	1,560	1,243	4,776	1,495	1,125	1,387	4,007	8,783
.	17	Central Total	7,441	6,920	6,664	21,025	7,820	6,212	8,144	22,176	43,201
.	18	Grand Total	18,579	19,194	17,297	55,070	20,754	18,771	18,757	58,282	113,352
.	19										

Excel can create outlines in both directions. In the preceding examples, the outline is a row (vertical) outline. Figure 24.4 shows the same model after a column (horizontal) outline was added by selecting the columns and using Data \Rightarrow Outline \Rightarrow Group \Rightarrow Auto Outline. When both outlines are in effect, Excel also displays outline symbols at the top.

FIGURE 24.3

The worksheet after collapsing the outline to the second level

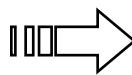
	A	B	C	D	E	F	G	H	I	J
1	State	Jan	Feb	Mar	Qtr-1	Apr	May	Jun	Qtr-2	Total
6	West Total	5,170	6,527	5,081	16,778	5,813	6,660	5,769	18,242	35,020
11	East Total	5,968	5,747	5,552	17,267	7,121	5,899	4,844	17,864	35,131
17	Central Total	7,441	6,920	6,664	21,025	7,820	6,212	8,144	22,176	43,201
18	Grand Total	18,579	19,194	17,297	55,070	20,754	18,771	18,757	58,282	113,352
19										
20										
21										

FIGURE 24.4

The worksheet after adding a column outline

	A	B	C	D	E	F	G	H	I	J
1	State	Jan	Feb	Mar	Qtr-1	Apr	May	Jun	Qtr-2	Total
2	California	1,118	1,960	1,252	4,330	1,271	1,557	1,679	4,507	8,837
3	Washington	1,247	1,238	1,028	3,513	1,345	1,784	1,574	4,703	8,216
4	Oregon	1,460	1,954	1,726	5,140	1,461	1,764	1,144	4,369	9,509
5	Arizona	1,345	1,375	1,075	3,795	1,736	1,555	1,372	4,663	8,458
6	West Total	5,170	6,527	5,081	16,778	5,813	6,660	5,769	18,242	35,020
7	New York	1,429	1,316	1,993	4,738	1,832	1,740	1,191	4,763	9,501
8	New Jersey	1,735	1,406	1,224	4,365	1,706	1,320	1,290	4,316	8,681
9	Massachusetts	1,099	1,233	1,110	3,442	1,637	1,512	1,006	4,155	7,597
10	Florida	1,705	1,792	1,225	4,722	1,946	1,327	1,357	4,630	9,352
11	East Total	5,968	5,747	5,552	17,267	7,121	5,899	4,844	17,864	35,131
12	Kentucky	1,109	1,078	1,155	3,342	1,993	1,082	1,551	4,626	7,968
13	Oklahoma	1,309	1,045	1,641	3,995	1,924	1,499	1,941	5,364	9,359
14	Missouri	1,511	1,744	1,414	4,669	1,243	1,493	1,820	4,556	9,225
15	Illinois	1,539	1,493	1,211	4,243	1,165	1,013	1,445	3,623	7,866
16	Kansas	1,973	1,560	1,243	4,776	1,495	1,125	1,387	4,007	8,783
17	Central Total	7,441	6,920	6,664	21,025	7,820	6,212	8,144	22,176	43,201
18	Grand Total	18,579	19,194	17,297	55,070	20,754	18,771	18,757	58,282	113,352
19										

If you create both a row outline and a column outline in a worksheet, you can work with each outline independently of the other. For example, you can show the row outline at the second level and the column outline at the first level. Figure 24.5 shows the model with both outlines collapsed at the second level. The result is a high-level summary table that displays regional totals by quarter.



You can find the workbook used in the preceding examples on this book's website at www.wiley.com/go/excel365bible. The file is named `outline example.xlsx`.

FIGURE 24.5

The worksheet with both outlines collapsed at the second level

	A	E	I	J
1	State	Qtr-1	Qtr-2	Total
2	West Total	16,778	18,242	35,020
3	East Total	17,267	17,864	35,131
4	Central Total	21,025	22,176	43,201
5	Grand Total	55,070	58,282	113,352
6				
7				
8				
9				
10				
11				
12				
13				
14				
15				
16				
17				
18				
19				

Keep in mind the following points about worksheet outlines:

- A worksheet can have only one outline. If you need to create more than one outline, use a different worksheet.
- You can either create an outline manually or have Excel do it for you automatically. If you choose the latter option, you may need to do some preparation to get the worksheet in the proper format. (The next sections cover both methods.)
- You can create an outline either for all data on a worksheet or just for a selected data range.
- You can remove an outline with a single command: Data \Rightarrow Outline \Rightarrow Ungroup \Rightarrow Clear Outline. See “Removing an outline” later in this chapter.
- You can hide the outline symbols (to free screen space) but retain the outline. We show you how in the section “Hiding the outline symbols” later in this chapter.
- An outline can have up to eight nested levels.

Worksheet outlines can be quite useful. If your main objective is to summarize a large amount of data, though, you may be better off using a PivotTable. A PivotTable is much more flexible and doesn’t require that you create the subtotal formulas; it does the summarizing for you automatically. The ultimate solution depends on your data source. If you’re entering data from scratch, the most flexible approach is to enter it in a normalized table format and create a PivotTable.



We discuss PivotTables (and normalized data) in Chapter 26, “Introducing PivotTables,” and Chapter 27, “Analyzing Data with PivotTables.”

Creating an Outline

The following sections describe the two ways to create an outline: automatically and manually. Before you create an outline, you need to ensure that data is appropriate for an outline and that the formulas are set up properly.

Preparing the data

What type of data is appropriate for an outline? Generally, the data should be arranged in a hierarchy, such as a budget that consists of an arrangement similar to the following:

```
Company
  Division
    Department
      Budget Category
        Budget Item
```

In this case, each budget item (for example, airfare and hotel expenses) is part of a budget category (for example, travel expenses). Each department has its own budget, and the departments are rolled up into divisions. The divisions make up the company. This type of arrangement is well suited for a row outline.

The data arrangement suitable for an outline is essentially a summary table of your data. In some situations, your data will be “normalized” data—one data point per row. You can easily create a PivotTable to summarize such data; a PivotTable is much more flexible than an outline.

After you create such an outline, you can view the information at any level of detail that you want by clicking the outline symbols. When you need to create reports for different levels of management, consider using an outline. For example, upper management may want to see only the division totals, division managers may want to see totals by department, and department managers will want to see the full details for their departments.

Keep in mind that using an outline isn’t a security feature. The data that’s hidden when an outline is collapsed can easily be revealed when the outline is expanded.

You can include time-based information that is rolled up into larger units (such as months and quarters) in a column outline. Column outlines work just like row outlines, however, and the levels don’t have to be time-based.

Before you create an outline, you need to make sure that all the summary formulas are entered correctly and consistently. In this context, *consistently* means that the formulas are in the same relative location. Generally, formulas that compute summary formulas (such as subtotals) are entered below the data to which they refer. In some cases, however, the summary formulas are entered above the referenced cells. Excel can handle either method, but you must be consistent throughout the range that you outline. If the summary formulas aren’t consistent, automatic outlining won’t produce the results you want.

Note

If your summary formulas aren’t consistent (that is, some are above and some are below the data), you still can create an outline, but you must do it manually.

Creating an outline automatically

Excel can create an outline for you automatically in a few seconds, whereas it may take you 10 minutes or more to do the same thing manually.

NOTE

If you've created a table for your data (by choosing Insert \Rightarrow Tables \Rightarrow Table), Excel can't create an outline automatically. You can create an outline from a table, but you must do so manually.

To have Excel create an outline, select any cell within the range of data that you're outlining. Then choose Data \Rightarrow Outline \Rightarrow Group \Rightarrow Auto Outline. Excel analyzes the formulas in the range and creates the outline. Depending on the formulas that you have, Excel creates a row outline, a column outline, or both.

If the worksheet already has an outline, Excel asks whether you want to modify the existing outline. Click OK to force Excel to remove the old outline and create a new one.

NOTE

Excel automatically creates an outline when you choose Data \Rightarrow Outline \Rightarrow Subtotal, which inserts subtotal formulas automatically.

Creating an outline manually

Usually, letting Excel create the outline is the best approach. It's much faster and less error-prone. If the outline that Excel creates isn't what you have in mind, however, you can create one manually.

When Excel creates a row outline, the summary rows must be all below the data or all above the data; they can't be mixed. Similarly, for a column outline, the summary columns must be all to the right of the data or all to the left of the data. If your worksheet doesn't meet these requirements, you have two choices:

- Rearrange the worksheet so that it does meet the requirements.
- Create the outline manually.

You also need to create an outline manually if the range doesn't contain formulas. You may have imported a file and want to use an outline to display it better. Because Excel uses the positioning of the formulas to determine how to create the outline, it can't make an outline without formulas.

Creating an outline manually consists of creating groups of rows (for row outlines) or groups of columns (for column outlines). To create a group of rows, follow these steps:

1. **Select the rows that you want to include in the group.** One way to do this is to click a row number and then drag to select other adjacent rows.

CAUTION

Don't select rows that have summary formulas like subtotals or grand totals. You don't want these rows to be included in the group.

2. **Choose Data ➔ Outline ➔ Group ➔ Group.** Excel displays outline symbols for the group.
3. **Repeat this process for each group that you want to create.** When you collapse the outline, Excel hides rows in the group, but the summary row, which isn't in the group, remains in view.

CAUTION

If you select a range of cells (rather than entire rows or columns) before you create a group, Excel displays a dialog box asking what you want to group. It then groups entire rows or columns based on the option you select.

You can also select groups of groups to create multilevel outlines. When you create multilevel outlines, always start with the innermost groupings and then work your way out. If you realize that you grouped the wrong rows, you can ungroup the group by selecting the rows and choosing Data ➔ Outline ➔ Ungroup ➔ Ungroup.

Here are keyboard shortcuts that you can use that speed up grouping and ungrouping:

Alt+Shift+right arrow: Groups selected rows or columns

Alt+Shift+left arrow: Ungroups selected rows or columns

Creating outlines manually can be confusing at first, but if you stick with it, you'll become a pro in no time.

Figure 24.6 shows a worksheet with a three-level outline of this book. We had to create it manually because it has no formulas, just text.



This workbook is available on this book's website at www.wiley.com/go/excel365bible. The file is named `book outline.xlsx`.

Part IV: Managing and Analyzing Data

FIGURE 24.6

An outline of this book, created manually

1	2	3	A	B	C	D	E
-	1		Excel Bible				
+	2		Part I: Getting Started With Excel				
+	11		Part II: Working with Formulas and Functions				
+	21		Part III: Creating Charts and Other Visualizations				
-	26		Part IV: Managing and Analyzing Data				
.	27		22. Importing and Cleaning Data				
.	28		23. Using Data Validation				
.	29		24. Creating and Using Worksheet Outlines				
.	30		25. Linking and Consolidating Worksheets				
.	31		26. Introducing PivotTables				
.	32		27. Analyzing Data with PivotTables				
.	33		28. Performing Spreadsheet What-if Analysis				
.	34		29. Analyzing Data Using Goal Seeking and Solver				
.	35		30. Analyzing Data with the Analysis ToolPak				
.	36		31. Protecting Your Work				
+	37		Part V: Understanding Power Pivot and Power Query				
+	45		Part VI: Automating Excel				
	53						

Working with Outlines

The following sections discuss the basic operations that you can perform with a worksheet outline.

Displaying levels

To display various outline levels, click the appropriate outline symbol. These symbols consist of headers with numbers on them (1, 2, and so on) or a plus sign (+) or minus sign (-). Refer to Figure 24.5, which shows these symbols for a row and column outline.

Clicking the 1 outline symbol collapses the outline so that it displays no detail (just the highest summary level of information), clicking the 2 outline symbol expands the outline to show one level, and so on. The number of headers depends on the number of outline levels. Choosing a level number displays the detail for that level, plus any levels with lower numbers. To display all levels (the most detail), click the highest-level number.

You can expand a particular section by clicking its plus-sign button, or you can collapse a particular section by clicking its minus-sign button. In short, you have complete control over the details that Excel exposes or hides in an outline.

If you prefer, you can use the Hide Detail and Show Detail commands on the Data \Rightarrow Outline group to hide and show details, respectively.

TIP

If you constantly adjust the outline to show different reports, consider using the Custom Views feature to save a particular view and give it a name. Then you can quickly switch among the named views. Choose View \Rightarrow Workbook Views \Rightarrow Custom Views.

Adding data to an outline

You may need to add additional rows or columns to an outline. In some cases, you may be able to insert new rows or columns without disturbing the outline, and the new rows or columns become part of the outline. In other cases, you'll find that the new row or column is not part of the outline. If you create the outline automatically, choose Data \Rightarrow Outline \Rightarrow Group \Rightarrow Auto Outline. Excel makes you verify that you want to modify the existing outline. If you create the outline manually, you need to make the adjustments manually as well.

Removing an outline

After you no longer need an outline, you can remove it by choosing Data \Rightarrow Outline \Rightarrow Ungroup \Rightarrow Clear Outline. Excel fully expands the outline by displaying all hidden rows and columns, and the outline symbols disappear. Be careful before you remove an outline, however. You can't make it reappear by clicking the Undo button; you must re-create the outline from scratch.

Adjusting the outline symbols

When you create a manual outline, Excel puts the outline symbols below the summary rows. This can be counterintuitive because you need to click the symbol in the row below the section that you want to expand.

If you prefer that the outline symbols appear in the same row as the summary row, click the dialog box launcher in the lower right of the Data \Rightarrow Outline group. Excel displays the dialog box shown in Figure 24.7. Remove the check mark from the Summary Rows Below Detail option and click OK. The outline will now display the outline symbols in a more logical position.

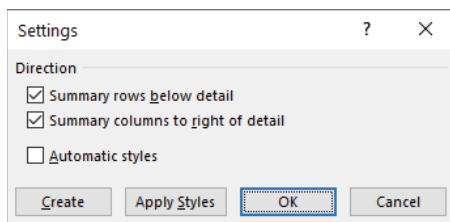
Hiding the outline symbols

The outline symbols that Excel displays when an outline is present take up quite a bit of space. (The exact amount depends on the number of levels.) If you want to see as much data as possible onscreen, you can temporarily hide these symbols without removing the outline. Press **Ctrl+8** to toggle the outline symbols on and off. When the outline symbols are hidden, you can't expand or collapse the outline.

Part IV: Managing and Analyzing Data

FIGURE 24.7

Use the Settings dialog box to adjust the position of the outline symbols.



NOTE

When you hide the outline symbols, the outline still is in effect, and the worksheet displays the data at the current outline level. (That is, some rows or columns may be hidden.)

The Custom Views feature, which saves named views of your outline, also saves the status of the outline symbols as part of the view, enabling you to name some views with the outline symbols and other views without them.

Linking and Consolidating Worksheets

IN THIS CHAPTER

Using various methods to link workbooks

Consolidating multiple worksheets

In this chapter, we discuss two procedures for using data from other workbooks or worksheets: linking and consolidating. *Linking* is the process of using references to cells in external workbooks to get data into your worksheet. *Consolidating* involves combining or summarizing information from two or more worksheets (which can be in multiple workbooks).

Linking Workbooks

Excel allows you to create formulas that contain references to other workbook files. The workbook that contains the external reference formulas is the dependent workbook because it contains formulas that depend on another workbook. The workbook that contains the information used in the external reference formula is the source workbook because that's where the information is.

When you consider linking workbooks, you may ask yourself the following question: If a workbook needs to access data that's in another workbook, why not just enter the data into the workbook that needs it? In some cases, you can. But the real value of linking becomes apparent when the source workbook is being continually updated by another person or group. Creating a link in the dependent workbook means you always have access to the most recent information in the source workbook. The dependent workbook is updated whenever the source changes.

Linking workbooks also can be helpful if you need to consolidate different files. For example, each regional sales manager may store data in a separate workbook. You can create a summary workbook that first uses link formulas to retrieve specific data from each manager's workbook and then calculates totals across all regions.

Linking also is useful to break up a large workbook into smaller files. You can create smaller workbooks that are linked with a few key external references.

Linking has its downside, however. External reference formulas are somewhat fragile, and accidentally severing the links that you create is relatively easy. You can prevent this mistake if you understand

how linking works. Later in the chapter, we discuss some problems that may arise and how to avoid them. (See “[Avoiding Potential Problems with External Reference Formulas](#).”)



The website for this book at www.wiley.com/go/excel365bible contains two linked files that you can use to get a feel for the way linking works. The files are named `source.xlsx` and `dependent.xlsx`. As long as these files remain in the same folder, the links will be maintained.

Creating External Reference Formulas

You can create an external reference formula by using several different techniques:

Type the cell references manually These references may be lengthy because they include workbook and sheet names and possibly even drive and path information. These references can also point to workbooks stored on the Internet. The advantage of manually typing the cell references is that the source workbook doesn't have to be open. The disadvantage is that it's error-prone. Mistyping a single character makes the formula return an error (or possibly return a wrong value from the workbook).

Point to the cell references If the source workbook is open, you can use the standard pointing techniques to create formulas that use external references.

Paste the links Copy your data to the Clipboard. Then, with the source workbook open, choose **Home** \Rightarrow **Clipboard** \Rightarrow **Paste** \Rightarrow **Paste Link (N)**. Excel pastes the copied data as external reference formulas.

Choose Data \Rightarrow Data Tools \Rightarrow Consolidate For more on this method, see the section “[Consolidating worksheets by using the Consolidate dialog box](#)” later in this chapter.

Understanding link formula syntax

Ideally, you won't have to enter too many external links manually. But it's good to know the structure of links in case you have to troubleshoot a problem. The general syntax for an external reference formula is as follows:

= [WorkbookName] SheetName!CellAddress

Start with the workbook name (in brackets) followed by the worksheet name and an exclamation point and, finally, the cell address. Here's an example of a formula that uses cell A1 in the worksheet named Sheet1 of a workbook named `Budget.xlsx`:

= [Budget.xlsx] Sheet1!A1

If the workbook name or the sheet name in the reference includes one or more spaces, you must enclose the text in single quotation marks. For example, here's a formula that refers to cell A1 on the worksheet Sheet1 in a workbook named `Annual Budget.xlsx`:

= ' [Annual Budget.xlsx] Sheet1' !A1

When a formula links to a different workbook, you don't need to open the other workbook. However, if the workbook is closed and not in the current folder, you must add the complete path to the reference. Here's an example:

```
= 'C:\Data\Excel\Budget\ [Annual Budget.xlsx]Sheet1'!A1
```

If the workbook is stored on the Internet, the formula will also include the URL. Here's an example:

```
= 'https://live.net/Docs/ [Annual Budget.xlsx]Sheet1'!A1
```

Note

Single quotes are always required when the link includes a path or a URL, even if the path or URL includes no spaces.

Creating a link formula by pointing

Entering external reference formulas manually usually isn't the best approach because you can easily make an error. Instead, have Excel build the formula for you, as follows:

1. Open the source workbook.
2. Select the cell in the dependent workbook that will hold the formula.
3. Type an equal sign (=).
4. Activate the source workbook, select the source worksheet, and then select the cell or range and press Enter. The dependent workbook is reactivated.

When you point to the cell or range, Excel automatically takes care of the details and creates a syntactically correct external reference. When you're using this method, the cell reference is always an absolute reference (such as \$A\$1). If you plan to copy the formula to create additional link formulas, you need to change the absolute reference to a relative reference by removing the dollar signs for the cell address.

As long as the source workbook remains open, the external reference doesn't include the path (or URL) to the workbook. If you close the source workbook, however, the external reference formulas change to include the full path (or URL).

Externally linked cells can also be an argument to a function, like SUM or VLOOKUP. Anywhere you would point to a cell on the same workbook as the function, you can point to a cell in another workbook. A SUM function with an external link looks like this:

```
=SUM( [source.xlsx]Sheet1!$B$3:$B$5)
```

Pasting links

Pasting links provides another way to create external reference formulas. This method is applicable when you want to create formulas that simply reference other cells rather than use links as part of a larger formula. Follow these steps:

1. **Open the source workbook.**
2. **Select the cell or range that you want to link, and then copy it to the Clipboard.** Ctrl+C is the quickest way.
3. **Activate the dependent workbook and select the cell in which you want the link formula to appear.** If you're pasting a copied range, just select the upper-left cell.
4. **Choose Home ➔ Clipboard ➔ Paste ➔ Paste Link (N).**

Working with External Reference Formulas

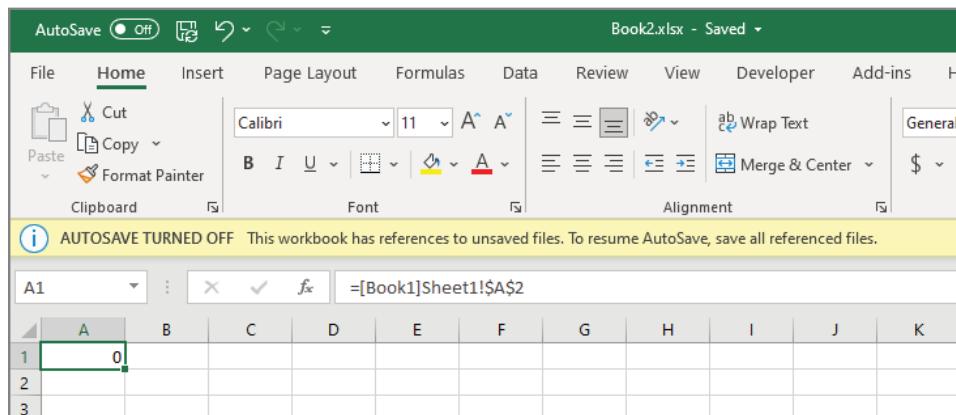
The following sections discuss some key points that you need to know when working with links. Understanding these details can help prevent some common errors.

Creating links to unsaved workbooks

Excel enables you to create link formulas to unsaved workbooks (and even to nonexistent workbooks). Assume that you have two workbooks open (Book1 and Book2) and you haven't saved either of them. If you create a link formula to Book1 in Book2 and then save Book2, Excel turns off AutoSave and displays a warning below the Ribbon like the one shown in Figure 25.1.

FIGURE 25.1

This message indicates that the workbook you saved contains references to a workbook that you haven't yet saved.



The warning appears even if you already had AutoSave turned off. Typically, you don't want to save a workbook that has links to an unsaved document. To avoid this prompt, save the source workbook first.

You can also create links to documents that don't exist. You may want to do so if you'll be using a source workbook from a colleague but the file hasn't yet arrived. When you enter an external reference formula that refers to a nonexistent workbook, Excel displays its Update Values dialog box, which resembles the Open dialog box. If you click Cancel, the formula retains the workbook name that you entered, but it returns a #REF! error.

When the source workbook becomes available, you can choose Data ➔ Queries & Connections ➔ Edit Links to update the link. (See "Updating links" later in this chapter.) After you do so, the error goes away, and the formula displays its proper value.

Opening a workbook with external reference formulas

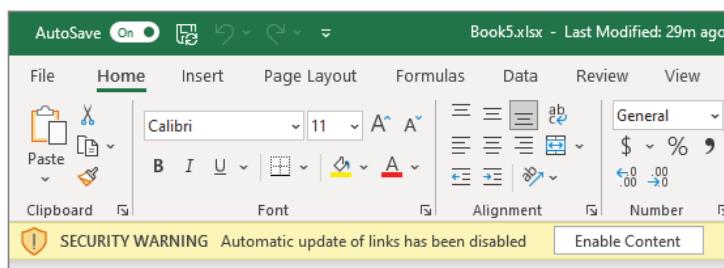
When you open a workbook that contains links, Excel displays a message below the Ribbon (shown in Figure 25.2) stating that the automatic update of links has been disabled. You can click the Enable Content button to update the links or inspect the formulas to make sure they're not linked to something malicious.

TIP

To prevent Excel from displaying the dialog box shown in Figure 25.2, open the Excel Options dialog box, select the Advanced tab, and in the General section, remove the check mark from Ask to Update Automatic Links. That disables the message for all workbooks.

FIGURE 25.2

Excel displays this message when you open a workbook that contains links to other files.

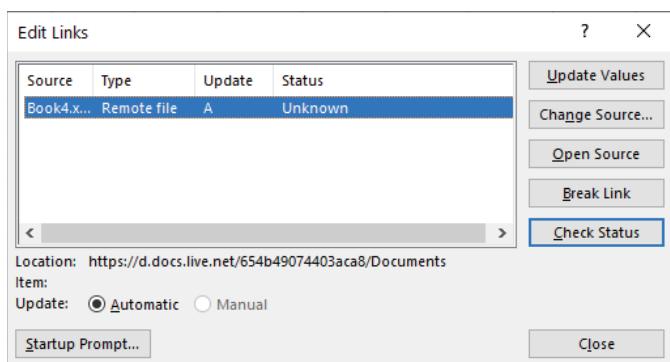


What if you choose to update the links but the source workbook is no longer available? If Excel can't locate a source workbook that's referred to in a link formula, it displays two dialog boxes informing you of that fact. You can edit each formula to point to the correct workbook, or to update all formulas at once, you can choose Data ➔ Queries & Connections ➔ Edit Links to display the Edit Links dialog box, shown in Figure 25.3 (see "Changing the link source" later in this chapter). The Edit Links dialog box lists all source workbooks plus other types of links to other documents.

Part IV: Managing and Analyzing Data

FIGURE 25.3

The Edit Links dialog box

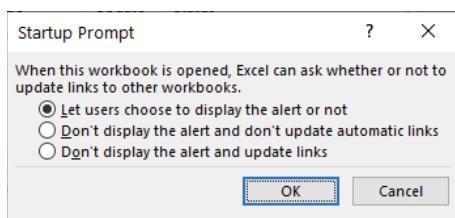


Changing the startup prompt

When you open a workbook that contains one or more external reference formulas, by default Excel disables updating the links automatically and displays a message (refer to Figure 25.2). You can eliminate this message by changing a setting in the Startup Prompt dialog box (see Figure 25.4).

FIGURE 25.4

Use the Startup Prompt dialog box to specify how Excel handles links when the workbook is opened.



To display the Startup Prompt dialog box, choose Data \Rightarrow Queries & Connections \Rightarrow Edit Links. The Edit Links dialog box (refer to Figure 25.3) appears. In the Edit Links dialog box, click the Startup Prompt button and then select the option that describes how you want to handle the links.

Updating links

If you want to ensure that your link formulas have the latest values from their source workbooks, you can force an update. For example, say that you just discovered that someone made changes to the source workbook and saved the latest version to your network server. In such a case, you may want to update the links to display the current data.

To update linked formulas with their current values, open the Edit Links dialog box (Data \Rightarrow Queries & Connections \Rightarrow Edit Links), choose the appropriate source workbook in the list, and then click the Update Values button (refer to Figure 25.3). Excel updates the link formulas with the latest version from the source workbook.

Note

Excel always sets worksheet links to the Automatic Update option in the Edit Links dialog box, and you can't change them to Manual, which means that Excel updates the links only when you open the workbook. Excel doesn't automatically update links when the source file changes (unless the source workbook is open).

Changing the link source

In some cases, you may need to change the source workbook for your external references. For example, say that you have a worksheet that has links to a file named Preliminary Budget, but you later receive a finalized version named Final Budget.

You can change the link source using the Edit Links dialog box (choose Data \Rightarrow Queries & Connections \Rightarrow Edit Links). Select the source workbook that you want to change and click the Change Source button (refer to Figure 25.3). Excel displays the Change Source dialog box from which you can select a new source file. After you select the file, all external reference formulas that referred to the old file are updated.

Severing links

If you have external references in a workbook and then decide that you no longer need the links, you can convert the external reference formulas to values, thereby severing the links. To do so, access the Edit Links dialog box (choose Data \Rightarrow Queries & Connections \Rightarrow Edit Links), select the linked file in the list, and then click Break Link (refer to Figure 25.3).

Caution

Excel prompts you to verify your intentions because you can't undo this operation.

Avoiding Potential Problems with External Reference Formulas

Using external reference formulas can be quite useful, but the links may be unintentionally severed. As long as the source file hasn't been deleted, you can almost always reestablish lost links. If you open the workbook and Excel can't locate the file, you see a dialog box that enables you to specify the workbook and re-create the links. You also can change the source file by clicking the Change Source button in the Edit Links dialog box (choose Data \Rightarrow Queries & Connections \Rightarrow Edit Links). The following sections discuss some pointers that you must remember when you use external reference formulas.

Renaming or moving a source workbook

If you or someone else renames the source document or moves it to a different folder, Excel won't be able to update the links. You need to use the Edit Links dialog box and specify the new source document. (See "Changing the link source" earlier in this chapter.)

NOTE

If the source and dependent files reside in the same folder, you can move both files to a different folder. In such a case, the links remain intact.

Using the Save As command

If both the source workbook and the dependent workbook are open, Excel doesn't display the full path to the source file in the external reference formulas. If you use the File \Rightarrow Save As command to give the source workbook a new name, Excel modifies the external references to use the new workbook name. In some cases, this change may be what you want. But in other cases, it may not.

Here's an example of how using File \Rightarrow Save As can cause a problem. You finished working on a source workbook and save the file. Then you decide to be safe and make a backup copy on a different drive, using File \Rightarrow Save As. The formulas in the dependent workbook now refer to the backup copy, not the original source file. This is not what you want.

Bottom line? Be careful when you choose File \Rightarrow Save As with a workbook that is the source of a link in another open workbook.

Modifying a source workbook

If you open a workbook that is the source for a link, be extremely careful if the workbook that contains the link isn't open. For example, if you insert a new row to the source workbook, all the cells move down one row. When you open the dependent workbook, it continues to use the old cell references, which is probably not what you want.

NOTE

It's easy to determine the source workbooks for a particular dependent workbook: just examine the files listed in the Edit Links dialog box (choose Data \Rightarrow Queries & Connections \Rightarrow Edit Links). However, it's not possible to determine whether a particular workbook is used as the source for another workbook.

You can avoid this problem by doing the following:

- **Always opening the dependent workbook(s) when you modify the source workbook:** If you do so, Excel adjusts the external references in the dependent workbook when you make changes to the source workbook.
- **Using names rather than cell references in your link formula:** This approach is the safest.
- **Using a formula to refer to a cell:** Your source workbook may not be structured to allow this method.

The following link formula refers to cell C21 on Sheet1 in the budget.xlsx workbook:

```
= [budget.xlsx] Sheet1!$C$21
```

If cell C21 is named Total, you can write the formula using that name:

```
=budget.xlsx!Total
```

Using a name ensures that the link retrieves the correct value, even if you insert or delete rows or columns in the source workbook.

By the way, notice that the filename isn't enclosed in brackets. That's because Total is assumed to be a workbook-level name and doesn't need to be qualified with a sheet name. If Total were a sheet-level name (defined on Sheet1), the formula would be as follows:

```
= [budget.xlsx] Sheet1!Total
```



See Chapter 4, "Working with Excel Ranges and Tables," for more information about creating names for cells and ranges.

If your source workbook contains a list of months and values and you want to return the July value, you can use a formula such as the following:

```
=VLOOKUP ("July", source.xlsx!MonthValues, 2, FALSE)
```

The source workbook has the workbook-level name MonthValues. If you used a normal cell reference, you would still have a problem if rows were inserted into the source workbook. However, by naming the entire range, you don't have to make a named range for each month separately.

Using Intermediary links

Excel doesn't place many limitations on the complexity of your network of external references. For example, Workbook A can contain external references that refer to Workbook B, which can contain an external reference that refers to Workbook C. In this case, a value in Workbook A can ultimately depend on a value in Workbook C. Workbook B is an intermediary link.

We don't recommend using intermediary links, but if you must use them, be aware that Excel doesn't update external reference formulas if the dependent workbook isn't open. In the preceding example, assume that Workbooks A and C are open. If you change a value in Workbook C, Workbook A won't reflect the change because you didn't open Workbook B (the intermediary link).

Consolidating Worksheets

The term *consolidation*, in the context of worksheets, refers to several operations that involve multiple worksheets or multiple workbooks. In some cases, consolidation involves creating link formulas. Here are two common examples of consolidation:

- The budget for each department in your company is stored in a single workbook, with a separate worksheet for each department. You need to consolidate the data and create a company-wide budget on a single sheet.

- Each department head submits a budget to you in a separate workbook file. Your job is to consolidate these files into a company-wide budget.

These types of tasks can be difficult or quite easy. The task is easy if the information is laid out in the same way in each worksheet. If the worksheets aren't laid out identically, they may be similar enough. In the second example, some budget files submitted to you may be missing categories that aren't used by a particular department. In this case, you can use a handy feature in Excel that matches data by using row and column titles. We discuss this feature in "Consolidating worksheets by using the Consolidate dialog box" later in this chapter.

If the worksheets bear little or no resemblance to each other, your best option may be to edit the sheets so that they correspond to one another. Or return the files to the department heads and ask that they submit them using a standardized format. Better yet, redesign your workflow to use normalized tables that can be used as the source for PivotTables.

You can use any of the following techniques to consolidate information from multiple workbooks:

- Use external reference formulas.
- Copy the data and choose Home \Rightarrow Clipboard \Rightarrow Paste \Rightarrow Paste Link (N).
- Use the Consolidate dialog box, which you get to by choosing Data \Rightarrow Data Tools \Rightarrow Consolidate.

Time to Rethink Your Consolidation Strategy?

If you're reading this chapter, there's a good chance that you're looking for a way to combine data from multiple sources. The consolidation methods we describe can work, but they may not be the most efficient way to approach the problem.

A typical budget is actually a summary. It's usually much easier to work with "normalized" data, which consists of one row per data item. Then you can use Excel's most sophisticated tool (a PivotTable) to consolidate and summarize the information.

For example, a budget for Region 1 might show a value for training expenses for the IT department for January. Instead of just entering this number into a grid, you gain a lot of flexibility by putting it into a table with multiple columns that describe the number. For example, this single item can be represented as a row in a normalized table with these six headings: Region, Department, Expense Description, Month, Year, and Budget Amount.

If each regional manager submitted his budget information in this format, it would be a simple matter to combine the data in a single worksheet and then create a PivotTable that displays a summary in just about any layout you want. See Chapter 26, "Introducing PivotTables," for more information.

Consolidating worksheets by using formulas

Consolidating with formulas simply involves creating formulas that use references to other worksheets or other workbooks. Here are the primary advantages to using this method of consolidation:

- If the values in the source worksheets change, the formulas are updated automatically.
- The source workbooks don't need to be open when you create the consolidation formulas.

If you're consolidating the worksheets in the same workbook and all the worksheets are laid out identically, the consolidation task is simple. You can just use standard formulas to create the consolidations. For example, to compute the total for cell A1 in worksheets named Sheet2 through Sheet10, enter the following formula:

```
=SUM(Sheet2:Sheet10!A1)
```

You can enter this formula manually or use the multi-sheet selection technique. You can then copy this formula to create summary formulas for other cells.



See Chapter 4, "Working with Excel Ranges and Tables," for more on multi-sheet selection.

If the consolidation involves other workbooks, you can use external reference formulas to perform your consolidation. For example, if you want to add the values in cell B2 from Sheet1 in two workbooks (named Region1 and Region2), you can use the following formula:

```
= [Region1.xlsx]Sheet1!B2 + [Region2.xlsx]Sheet1!B2
```

You can include any number of external references in this formula, up to the 8,000-character limit for a formula. However, if you use many external references, such a formula can be quite lengthy and confusing if you need to edit it.

If the worksheets that you're consolidating aren't laid out in the same way, you can still use formulas, but you need to ensure that each formula refers to the correct cell—a task that is both tedious and error-prone.

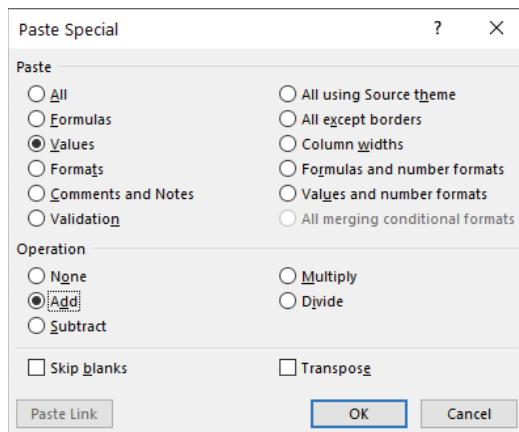
Consolidating worksheets by using Paste Special

Another method of consolidating information is to use the Paste Special dialog box. This technique takes advantage of the fact that the Paste Special dialog box can perform a mathematical operation when it pastes data from the Clipboard. For example, you can use the Add option to add the copied data to the selected range. Figure 25.5 shows the Paste Special dialog box.

This method is applicable only when all the worksheets that you're consolidating are open. The disadvantage is that the consolidation isn't dynamic. In other words, it doesn't generate formulas that refer to the original source data. So, if any data that was consolidated changes, the consolidation is no longer accurate.

FIGURE 25.5

Choosing the Add operation in the Paste Special dialog box



Here's how to use this method:

1. **Copy the data from the first source range.**
2. **Activate the dependent workbook and select a location for the consolidated data.**
A single cell is sufficient.
3. **Choose Home ➔ Clipboard ➔ Paste ➔ Paste Special.** The Paste Special dialog box appears.
4. **Choose the Values option and the Add operation and then click OK.**

Repeat these steps for each source range that you want to consolidate. Make sure that the consolidation location in step 2 is the same for each paste operation.

CAUTION

This method is probably the worst way of consolidating data. It can be rather error-prone, and the lack of formulas means that there is no “trail.” If an error is discovered, it may be difficult or impossible to determine the source of the error.

Consolidating worksheets by using the Consolidate dialog box

For the ultimate in data consolidation, use the Consolidate dialog box. This method is flexible, and in some cases it even works if the source worksheets aren't laid out identically. This technique can create consolidations that are static (no link formulas) or dynamic (with link formulas). The data consolidation feature supports the following methods of consolidation:

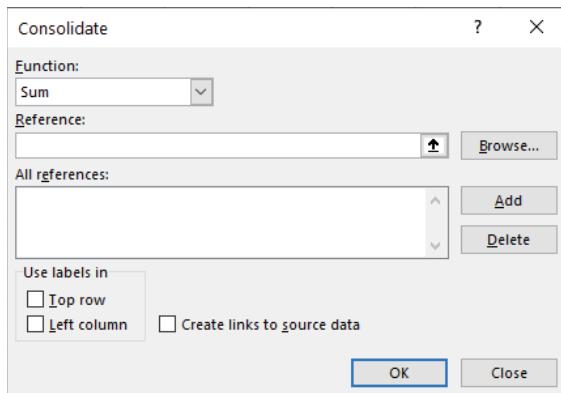
By position This method is accurate only if the worksheets are laid out identically.

By category Excel uses row and column labels to match data in the source worksheets. Use this option if the data is laid out differently in the source worksheets or if some source worksheets are missing rows or columns.

Figure 25.6 shows the Consolidate dialog box, which appears when you choose Data ➔ Data Tools ➔ Consolidate.

FIGURE 25.6

The Consolidate dialog box enables you to specify ranges to consolidate.



The following is a description of the controls in this dialog box:

Function drop-down list Specify the type of consolidation. SUM is the most commonly used consolidation function, but you can also select from 10 other options.

Reference box Specify a range from a source file that you want to consolidate. You can enter the range reference manually or use any standard pointing technique (if the workbook is open). Named ranges are also acceptable. After you enter the range in this box, click Add to add it to the All References list. If you consolidate by position, don't include labels in the range; if you consolidate by category, do include labels in the range.

All References list box This contains the list of references that you've added with the Add button.

Use Labels In check boxes Use this to instruct Excel to perform the consolidation by examining the labels in the top row, the left column, or both positions. Use these options when you consolidate by category.

Create Links to Source Data check box When you select this option, Excel adds summary formulas for each label and creates an outline. If you don't select this option, the consolidation doesn't use formulas, and an outline isn't created.

Browse button Click to display a dialog box that enables you to select a workbook to open. It inserts the filename in the Reference box, but you have to supply the range reference. You'll find that your job is much easier if all the workbooks to be consolidated are open.

Add button Click to add the reference in the Reference box to the All References list. Make sure that you click this button after you specify each range.

Delete button Click to delete the selected reference from the All References list.

Viewing a workbook consolidation example

The simple example in this section demonstrates the power of the data consolidation feature. Figure 25.7 shows three single-sheet workbooks that will be consolidated. These worksheets report three months of product sales. Notice, however, that all don't report on the same products. In addition, the products aren't listed in the same order. In other words, these worksheets aren't laid out identically. Creating consolidation formulas manually would be a tedious task.

FIGURE 25.7

Three workbooks to be consolidated

The figure displays three separate Excel windows side-by-side, each showing a different region's product sales data. The top window is titled 'region1.xlsx', the middle is 'region2.xlsx', and the bottom is 'region3.xlsx'. Each window contains a table with columns for Product ID, Jan, Feb, and Mar. The data in each window is as follows:

	Product ID	Jan	Feb	Mar
1	A-401	1,000	1,094	1,202
2	A-403	1,188	1,324	1,236
3	A-404	1,212	1,002	1,018
4	A-409	1,173		
5	A-412	1,298		
6	A-415	1,217		
7	A-503	1,285		
8	A-511	1,192		
9	A-536	1,202		
10				
11				
12				
13				

	Product ID	Jan	Feb	Mar
1	A-402	5,344	5,211	5,526
2	A-401	5,000	5,600	5,451
3	A-404	5,436		
4	A-408	5,336		
5	A-490	5,278		
6	A-415	5,497		
7	A-502	5,626		
8	A-505	5,497		
9	A-515	5,374		
10	A-523	5,597		
11	A-536	5,552		
12				
13				

	Product ID	Mar	Feb	Jan
1	A-407	3,301	3,478	3,453
2	A-401	3,224	3,246	3,000
3	A-405	3,299	3,221	3,039
4	A-406	3,263	3,255	3,282
5	A-512	3,023	3,217	3,218
6	A-514	3,011	3,024	3,177
7	A-523	3,209	3,482	3,348
8	A-533	3,447	3,252	3,327
9	A-535	3,074	3,026	3,426
10	A-536	3,489	3,087	3,205
11				
12				
13				



These workbooks are available on this book's website at www.wiley.com/go/excel365bible.
The files are named `region1.xlsx`, `region2.xlsx`, and `region3.xlsx`.

To consolidate this information, start with a new workbook. You don't need to open the source workbooks, but consolidation is easier if they're open. Follow these steps to consolidate the workbooks:

1. Choose **Data ➔ Data Tools ➔ Consolidate**. The Consolidate dialog box appears.
2. Select the **type of consolidation summary you want to use from the Function drop-down list**. Use **SUM** for this example.
3. **Enter the reference for the first worksheet to consolidate**. If the workbook is open, you can point to the reference; if it isn't open, click the **Browse** button to locate the file on disk. The reference must include a range. You can use a range that includes complete columns, such as A : K. This range is larger than the actual range to consolidate, but using this range ensures that the consolidation will still work if new rows and columns are added to the source file.
4. **When the reference in the Reference box is correct, click Add to add it to the All References list**.
5. **Enter the reference for the second worksheet**. You can point to the range in the Region2 workbook, or you can simply edit the existing reference by changing Region1 to Region2 and then clicking **Add**. This reference is added to the All References list.
6. **Enter the reference for the third worksheet**. Again, you can edit the existing reference by changing Region2 to Region3 and then clicking **Add**. This final reference is added to the All References list.
7. **Because the worksheets aren't laid out in the same way, select the Left Column and the Top Row check boxes to force Excel to match the data by using the labels**.
8. **Select the Create Links to Source Data check box to make Excel create an outline with external references**.
9. **Click OK to begin the consolidation**.

Excel creates the consolidation, beginning at the active cell. Notice that Excel created an outline, which is collapsed to show only the subtotals for each product. If you expand the outline (by clicking the number 2 or the plus-sign symbols in the outline), you can see the details. Examine it further, and you discover that each detail cell is an external reference formula that uses the appropriate cell in the source file. Therefore, the consolidated results are updated automatically when values are changed in any of the source workbooks.

Figure 25.8 shows the result of the consolidation, and Figure 25.9 shows the detail information (with the outline expanded).

Part IV: Managing and Analyzing Data

FIGURE 25.8

The result of consolidating the information in three workbooks

1	2	A	B	C	D	E	F
		1		Jan	Feb	Mar	
+	3	A-402		5,344	5,211	5,526	
+	5	A-407		3,453	3,478	3,301	
+	9	A-401		9,000	9,940	9,877	
+	11	A-403		1,188	1,324	1,236	
+	14	A-404		6,648	6,352	6,228	
+	16	A-409		1,173	1,116	1,110	
+	18	A-412		1,298	1,218	1,467	
+	20	A-408		5,336	5,358	5,653	
+	22	A-490		5,278	5,676	5,257	
+	25	A-415		6,714	6,612	6,617	
+	27	A-503		1,285	1,054	1,298	
+	29	A-511		1,192	1,408	1,010	
+	31	A-502		5,626	5,517	5,564	
+	33	A-505		5,497	5,239	5,348	
+	35	A-515		5,374	5,337	5,443	
+	37	A-405		3,039	3,221	3,299	
+	39	A-406		3,282	3,255	3,263	
+	41	A-512		3,218	3,217	3,023	
+	43	A-514		3,177	3,024	3,011	
+	46	A-523		8,945	8,851	8,537	
+	48	A-533		3,327	3,252	3,447	
+	50	A-535		3,426	3,026	3,074	
+	54	A-536		9,959	9,942	10,889	
	55						



For more information on Excel outlines, see Chapter 24, “Creating and Using Worksheet Outlines.”

Refreshing a consolidation

When you choose the option to create formulas, the external references in the consolidation workbook are created only for data that exists at the time of the consolidation. Therefore, if new rows are added to any of the original workbooks, the consolidation must be redone. Fortunately, the consolidation parameters are stored with the workbook, so it's a simple matter to rerun the consolidation if necessary. That's why specifying complete columns and including extra columns (step 3 in the preceding section) is a good idea.

Excel remembers the references that you entered in the Consolidate dialog box and saves them with the workbook. That way, if you want to refresh a consolidation, you won't have to re-enter the references. Just display the Consolidate dialog box, verify that the ranges are correct, and then click OK.

FIGURE 25.9

Expanding the outline to show more details

1	2	A	B	C	D	E	F
			Jan	Feb	Mar		
1			5,344	5,211	5,526		
2		Region2	5,344	5,211	5,526		
3		A-402	5,344	5,211	5,526		
4		Region3	3,453	3,478	3,301		
5		A-407	3,453	3,478	3,301		
6		Region1	1,000	1,094	1,202		
7		Region2	5,000	5,600	5,451		
8		Region3	3,000	3,246	3,224		
9		A-401	9,000	9,940	9,877		
10		Region1	1,188	1,324	1,236		
11		A-403	1,188	1,324	1,236		
12		Region1	1,212	1,002	1,018		
13		Region2	5,436	5,350	5,210		
14		A-404	6,648	6,352	6,228		
15		Region1	1,173	1,116	1,110		
16		A-409	1,173	1,116	1,110		
17		Region1	1,298	1,218	1,467		
18		A-412	1,298	1,218	1,467		
19		Region2	5,336	5,358	5,653		
20		A-408	5,336	5,358	5,653		
21		Region2	5,278	5,676	5,257		

CAUTION

When Excel has an outline applied, like after you consolidate, rerunning the consolidation can be unpredictable. The sure-fire way to rerun a consolidation is to clear the outline, delete the cell contents, display the Consolidation dialog box, and click OK.

Learning more about consolidation

Excel is flexible regarding the sources that you can consolidate. You can consolidate data from the following:

- Open workbooks.
- Closed workbooks. You need to enter the reference manually, but you can use the Browse button to get the filename part of the reference.
- The same workbook in which you're creating the consolidation.

And, of course, you can mix and match any of the preceding choices in a single consolidation.

If you perform the consolidation by matching labels, be aware that the matches must be exact. For example, *Jan* doesn't match *January*. The matching is not case sensitive, however, so *April* does match *APRIL*. In addition, the labels can be in any order, and they don't need to be in the same order in all the source ranges.

Part IV: Managing and Analyzing Data

If you don't select the Create Links to Source Data check box, Excel generates a static consolidation. (It doesn't create formulas.) Therefore, if the data on any of the source worksheets changes, the consolidation won't update automatically. To update the summary information, you need to choose Data ➔ Data Tools ➔ Consolidate again.



If you do select the Create Links to Source Data check box, Excel creates a standard worksheet outline that you can manipulate by using the techniques described in Chapter 24.

Introducing PivotTables

IN THIS CHAPTER

Understanding PivotTables

Identifying types of data appropriate for a PivotTable

Getting clear on PivotTable terminology

Creating PivotTables

Looking at PivotTable examples that answer specific questions about data

The PivotTable feature is perhaps the most technologically sophisticated component in Excel. With only a few mouse clicks, you can slice and dice data in dozens of different ways and produce just about any type of summary you can imagine.

If you haven't yet discovered the power of PivotTables, this chapter provides an introduction, and Chapter 27, "Analyzing Data with PivotTables," continues with many examples that demonstrate how easy it is to create powerful data summaries using PivotTables.

About PivotTables

A *PivotTable* is essentially a dynamic summary report generated from a database. The database can reside in a worksheet (in the form of a table) or in an external data file. A PivotTable can help transform endless rows and columns of numbers into a meaningful presentation of the data—and do it so quickly you'll be amazed.

Perhaps the most powerful aspect of a PivotTable is its interactivity. After you create a PivotTable, you can rearrange the information in almost any way imaginable and even insert special formulas that perform new calculations. You even can create post hoc groupings of summary items (for example, combine Northern Region totals with Western Region totals). With a few mouse clicks, you can apply formatting to a PivotTable to convert it into a report, albeit with a few limitations.

One minor drawback to using a PivotTable is that, unlike a formula-based summary report, a PivotTable does not update automatically when you change information in the source data. This drawback doesn't pose a serious problem, however, because a single click of the Refresh button forces a PivotTable to update itself with the latest data.

A PivotTable example

The best way to understand the concept of a PivotTable is to see one. Start with Figure 26.1, which shows a portion of the data used in creating the PivotTable in this chapter. This range happens to be in a table (created by using Insert \Rightarrow Tables \Rightarrow Table), but that's not a requirement for creating a PivotTable.

FIGURE 26.1

This table is used to create a PivotTable.

1	Date	Weekday	Amount	AcctType	OpenedBy	Branch	Customer	H
2	4/1/2022	Friday	5,000	IRA	New Accts	Central	Existing	
3	4/1/2022	Friday	10,000	CD	Teller	North County	Existing	
4	4/1/2022	Friday	500	Checking	New Accts	Central	Existing	
5	4/1/2022	Friday	11,779	CD	Teller	Central	New	
6	4/1/2022	Friday	4,623	Savings	New Accts	North County	Existing	
7	4/1/2022	Friday	8,721	Savings	New Accts	Westside	New	
8	4/1/2022	Friday	15,276	Savings	New Accts	North County	Existing	
9	4/1/2022	Friday	5,000	Savings	New Accts	Westside	Existing	
10	4/1/2022	Friday	12,000	CD	New Accts	Westside	Existing	
11	4/1/2022	Friday	13,636	CD	New Accts	North County	Existing	
12	4/1/2022	Friday	7,177	Savings	Teller	North County	Existing	
13	4/1/2022	Friday	6,837	Savings	New Accts	Westside	Existing	
14	4/1/2022	Friday	3,171	Checking	New Accts	Westside	Existing	
15	4/1/2022	Friday	50,000	Savings	New Accts	Central	Existing	
16	4/1/2022	Friday	4,690	Checking	New Accts	North County	New	
17	4/1/2022	Friday	12,438	Checking	New Accts	Central	Existing	
18	4/1/2022	Friday	5,000	Checking	New Accts	North County	Existing	
19	4/1/2022	Friday	7,000	Savings	New Accts	North County	New	
20	4/1/2022	Friday	11,957	Checking	New Accts	Central	Existing	
21	4/1/2022	Friday	14,571	CD	Teller	Central	New	
22	4/1/2022	Friday	15,000	CD	New Accts	Central	Existing	
23	4/1/2022	Friday	5,879	Checking	New Accts	Central	Existing	
24	4/1/2022	Friday	4,000	Savings	New Accts	Central	Existing	

This table consists of a month's worth of new account information for a three-branch bank. The table contains 712 rows, and each row represents a new account opened at the bank. The table has the following columns:

- The date the account was opened
- The day of the week the account was opened
- The opening deposit amount
- The account type (CD, checking, savings, or IRA)
- Who opened the account (a teller or a new-account representative)
- The branch at which it was opened (Central, Westside, or North County)
- The type of customer (an existing customer or a new customer)



This workbook, named bank_accounts.xlsx, is available on this book's website at www.wiley.com/go/excel365bible.

The bank accounts database contains quite a bit of information. In its current form, though, the data doesn't reveal much. To make the data more useful, you need to summarize it. Summarizing a database is essentially the process of arranging the data differently in order to answer questions about it. The following are a few of the questions that may be of interest to the bank's management:

- What is the daily total new deposit amount for each branch?
- Which day of the week accounts for the most deposits?
- How many accounts were opened at each branch, broken down by account type?
- How much money was used to open the accounts?
- What types of accounts do tellers open most often?
- In which branch do tellers open the most checking accounts for new customers?

You can, of course, spend time sorting the data and creating formulas to answer these questions. But almost always, a PivotTable is a better choice. Creating a PivotTable takes only a few seconds and doesn't require a single formula. In addition, creating PivotTables is much less prone to error than creating formulas.

Later in this chapter, you'll see several PivotTables that answer the preceding questions.

Figure 26.2 shows a PivotTable created from the bank data. This PivotTable shows the amount of new deposits, broken down by branch and account type. This particular summary is one of dozens of summaries that you can produce from this data.

FIGURE 26.2

A simple PivotTable

	A	B	C	D	E	F	G
1							
2							
3	Sum of Amount	Column I ▾					
4	Row Labels	CD	Checking	IRA	Savings	Grand Total	
5	Central	1,361,885	802,403	68,380	885,757	3,118,425	
6	North County	1,209,910	392,516	134,374	467,414	2,204,214	
7	Westside	650,237	292,995	10,000	336,088	1,289,320	
8	Grand Total	3,222,032	1,487,914	212,754	1,689,259	6,611,959	
9							

Figure 26.3 shows another PivotTable generated from the bank data. This PivotTable uses a drop-down Report Filter for the Customer item (in row 2). In the figure, the PivotTable displays the data only for existing customers. (The user can also select New or All from the drop-down control.)

Part IV: Managing and Analyzing Data

FIGURE 26.3

A PivotTable that uses a report filter

	A	B	C	D	E	F
1						
2	Customer	Existing				
3						
4	Sum of Amount	Column Labels				
5	Row Labels	Central	North	County	Westside	Grand Total
6	CD		974,112	845,522	356,079	2,175,713
7	Checking		505,822	208,375	144,391	858,588
8	IRA		68,380	125,374	10,000	203,754
9	Savings		548,198	286,891	291,728	1,126,817
10	Grand Total		2,096,512	1,466,162	802,198	4,364,872
11						

The orientation of the table changes between the two examples. For this PivotTable, branches appear as column labels, and account types appear as row labels. This change, which took about five seconds to make, is another example of the flexibility of a PivotTable.

Why Pivot?

Are you curious about the term *pivot*?

Pivot, as a verb, means to rotate or revolve. If you think of your data as a physical object, a PivotTable lets you rotate the data summary and look at it from different angles or perspectives. A PivotTable allows you to move fields around easily, nest fields within each other, and even create ad hoc groups of items.

If you were handed a strange object and asked to identify it, you'd probably look at it from several different angles in an attempt to figure it out. Working with a PivotTable is similar to investigating a strange object. In this case, the object happens to be your data. A PivotTable invites experimentation, so feel free to rotate and manipulate the PivotTable until you're satisfied. You may be surprised at what you discover.

Data appropriate for a PivotTable

A PivotTable requires that your data be in the form of a rectangular table. You can store the database in either a worksheet range (which can be a table or just a normal range) or an external database file. And although Excel can generate a PivotTable from any database, not all databases benefit.

Generally speaking, fields in a table consist of two types of information:

Data Contains a value or data to be summarized. For the bank account example, the Amount field is a data field.

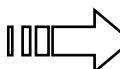
Category Describes the data. For the bank account data, the Date, Weekday, AcctType, OpenedBy, Branch, and Customer fields are category fields because they describe the data in the Amount field.

NOTE

A database table that's appropriate for a PivotTable is said to be *normalized*. In other words, each record (or row) contains information that describes the data.

A single table can have any number of data fields and category fields. When you create a PivotTable, you usually want to summarize one or more of the data fields. Conversely, the values in the category fields appear in the PivotTable as rows, columns, or filters.

Exceptions exist, however, and you may find the Excel PivotTable feature useful even for databases that don't contain actual numerical data fields.



Chapter 27 has an example of a PivotTable created from non-numeric data.

Figure 26.4 shows an example of an Excel range that is not appropriate for a PivotTable. You might recognize this data from the outline example in Chapter 24, “Creating and Using Worksheet Outlines.” Although the range contains descriptive information about each value, it does not consist of normalized data. In fact, this range actually resembles a PivotTable summary, but it’s much less flexible.

FIGURE 26.4

This range is not appropriate for a PivotTable.

	A	B	C	D	E	F	G	H	I	J
1	State	Jan	Feb	Mar	Qtr-1	Apr	May	Jun	Qtr-2	Total
2	California	1,118	1,960	1,252	4,330	1,271	1,557	1,679	4,507	8,837
3	Washington	1,247	1,238	1,028	3,513	1,345	1,784	1,574	4,703	8,216
4	Oregon	1,460	1,954	1,726	5,140	1,461	1,764	1,144	4,369	9,509
5	Arizona	1,345	1,375	1,075	3,795	1,736	1,555	1,372	4,663	8,458
6	West Total	5,170	6,527	5,081	16,778	5,813	6,660	5,769	18,242	35,020
7	New York	1,429	1,316	1,993	4,738	1,832	1,740	1,191	4,763	9,501
8	New Jersey	1,735	1,406	1,224	4,365	1,706	1,320	1,290	4,316	8,681
9	Massachusetts	1,099	1,233	1,110	3,442	1,637	1,512	1,006	4,155	7,597
10	Florida	1,705	1,792	1,225	4,722	1,946	1,327	1,357	4,630	9,352
11	East Total	5,968	5,747	5,552	17,267	7,121	5,899	4,844	17,864	35,131
12	Kentucky	1,109	1,078	1,155	3,342	1,993	1,082	1,551	4,626	7,968
13	Oklahoma	1,309	1,045	1,641	3,995	1,924	1,499	1,941	5,364	9,359
14	Missouri	1,511	1,744	1,414	4,669	1,243	1,493	1,820	4,556	9,225
15	Illinois	1,539	1,493	1,211	4,243	1,165	1,013	1,445	3,623	7,866
16	Kansas	1,973	1,560	1,243	4,776	1,495	1,125	1,387	4,007	8,783
17	Central Total	7,441	6,920	6,664	21,025	7,820	6,212	8,144	22,176	43,201
18	Grand Total	18,579	19,194	17,297	55,070	20,754	18,771	18,757	58,282	113,352
19										

Figure 26.5 shows the same data but normalized. This range contains 78 rows of data—one for each of the six monthly sales values for the 13 states. Notice that each row contains category information for the sales value. This table is an ideal candidate for a PivotTable, and it contains all the information necessary to summarize the information by region or quarter.

Part IV: Managing and Analyzing Data

FIGURE 26.5

This range contains normalized data and is appropriate for a PivotTable.

	A	B	C	D	E
1	State	Region	Month	Qtr	Sales
2	California	West	Jan	Qtr-1	1,118
3	California	West	Feb	Qtr-1	1,960
4	California	West	Mar	Qtr-1	1,252
5	California	West	Apr	Qtr-2	1,271
6	California	West	May	Qtr-2	1,557
7	California	West	Jun	Qtr-2	1,679
8	Washington	West	Jan	Qtr-1	1,247
9	Washington	West	Feb	Qtr-1	1,238
10	Washington	West	Mar	Qtr-1	1,028
11	Washington	West	Apr	Qtr-2	1,345
12	Washington	West	May	Qtr-2	1,784
13	Washington	West	Jun	Qtr-2	1,574
14	Oregon	West	Jan	Qtr-1	1,460
15	Oregon	West	Feb	Qtr-1	1,954
16	Oregon	West	Mar	Qtr-1	1,726
17	Oregon	West	Apr	Qtr-2	1,461
18	Oregon	West	May	Qtr-2	1,764
19	Oregon	West	Jun	Qtr-2	1,144
20	Arizona	West	Jan	Qtr-1	1,345
21	Arizona	West	Feb	Qtr-1	1,375
22	Arizona	West	Mar	Qtr-1	1,075
23	Arizona	West	Apr	Qtr-2	1,736
24	Arizona	West	May	Qtr-2	1,555

Figure 26.6 shows a PivotTable created from the normalized data. As you can see, it's similar to the nonnormalized data shown in Figure 26.4. Working with normalized data provides ultimate flexibility in designing reports.



This workbook, named **normalized data.xlsx**, is available on this book's website at www.wiley.com/go/excel365bible.

Creating a PivotTable Automatically

Creating a PivotTable is easy. This task requires practically no effort if your data is appropriately structured, and you choose a Recommended PivotTable.

If your data is in a worksheet, select any cell within the data range and choose **Insert** \Rightarrow **Tables** \Rightarrow **Recommended PivotTables**. Excel quickly scans your data, and the Recommended PivotTables dialog box presents thumbnails that depict some PivotTables from which you can choose.

Figure 26.7 shows the Recommended PivotTables dialog box for the bank account data.

FIGURE 26.6

A PivotTable created from normalized data

	A	B	C	D	E	F	G	H	I	J
1										
2	Sum of Sales	Col	▼							
3		Qtr-1			Qtr-1 Total	Qtr-2		Qtr-2 Total	Grand Total	
4	Row Labels	▼	Jan	Feb	Mar	Apr	May	Jun		
5	Central									
6	Illinois	1,539	1,493	1,211	4,243	1,165	1,013	1,445	3,623	7,866
7	Kansas	1,973	1,560	1,243	4,776	1,495	1,125	1,387	4,007	8,783
8	Kentucky	1,109	1,078	1,155	3,342	1,993	1,082	1,551	4,626	7,968
9	Missouri	1,511	1,744	1,414	4,669	1,243	1,493	1,820	4,556	9,225
10	Oklahoma	1,309	1,045	1,641	3,995	1,924	1,499	1,941	5,364	9,359
11	Central Total	7,441	6,920	6,664	21,025	7,820	6,212	8,144	22,176	43,201
12										
13	East									
14	Florida	1,705	1,792	1,225	4,722	1,946	1,327	1,357	4,630	9,352
15	Massachusetts	1,099	1,233	1,110	3,442	1,637	1,512	1,006	4,155	7,597
16	New Jersey	1,735	1,406	1,224	4,365	1,706	1,320	1,290	4,316	8,681
17	New York	1,429	1,316	1,993	4,738	1,832	1,740	1,191	4,763	9,501
18	East Total	5,968	5,747	5,552	17,267	7,121	5,899	4,844	17,864	35,131
19										
20	West									
21	Arizona	1,345	1,375	1,075	3,795	1,736	1,555	1,372	4,663	8,458
22	California	1,118	1,960	1,252	4,330	1,271	1,557	1,679	4,507	8,837
23	Oregon	1,460	1,954	1,726	5,140	1,461	1,764	1,144	4,369	9,509
24	Washington	1,247	1,238	1,028	3,513	1,345	1,784	1,574	4,703	8,216
25	West Total	5,170	6,527	5,081	16,778	5,813	6,660	5,769	18,242	35,020
26										
27	Grand Total	18,579	19,194	17,297	55,070	20,754	18,771	18,757	58,282	113,352

The PivotTable thumbnails use your actual data, and there's a good chance that one of them will be exactly what you're looking for—or at least close to what you're looking for. Select a thumbnail, click OK, and Excel creates the PivotTable on a new worksheet.

When any cell in a PivotTable is selected, Excel displays the PivotTable Fields task pane. You can use this task pane to make changes to the layout of the PivotTable.

Note

If your data is in an external database, start by selecting a blank cell. When you choose Insert \Rightarrow Tables \Rightarrow Recommended PivotTables, the Choose Data Source dialog box appears. Select Use an External Data Source and then click Choose Connection to specify the data source. Once you've selected the source, you'll see the thumbnails of the list of recommended PivotTables.

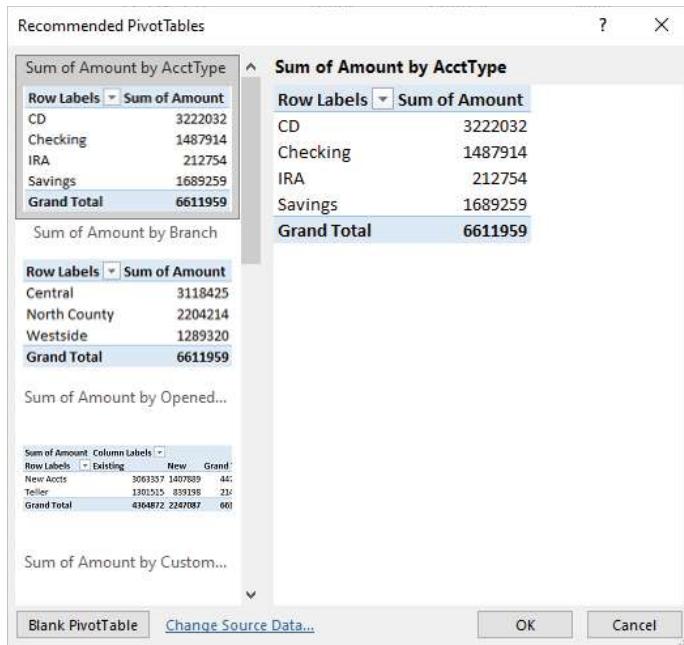
If none of the Recommended PivotTables is suitable, you have two choices:

- Create a PivotTable that's close to what you want and then use the PivotTable Fields task pane to modify it.
- Click the Blank PivotTable button (at the bottom of the Recommended PivotTables dialog box) and create a PivotTable manually.

Part IV: Managing and Analyzing Data

FIGURE 26.7

Selecting a Recommended PivotTable



Creating a PivotTable Manually

In the following sections, we describe the basic steps required to create a PivotTable using the bank account data described earlier in this chapter. Creating a PivotTable is an interactive process. It's not at all uncommon to experiment with various layouts until you find one with which you're satisfied. If you're unfamiliar with the elements of a PivotTable, see the sidebar "PivotTable Terminology."

Specifying the data

If your data is in a worksheet range, select any cell in that range and then choose **Insert** \Rightarrow **Tables** \Rightarrow **PivotTable**. The Create PivotTable dialog box, shown in Figure 26.8, appears.

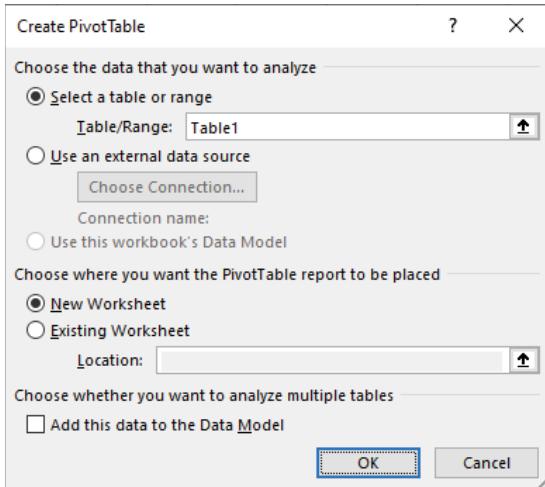
Excel attempts to guess the range, based on the location of the active cell. If you're creating a PivotTable from an external data source, you need to select that option and then click **Choose Connection** to specify the data source.

TIP

If you're creating a PivotTable from data in a worksheet, it's a good idea first to create a table for the range. (Choose **Insert** \Rightarrow **Tables** \Rightarrow **Table**.) Then, if you expand the table by adding new rows of data, the PivotTable will automatically adjust to cover the whole table without the need to indicate the new data range manually.

FIGURE 26.8

In the Create PivotTable dialog box, you tell Excel where the data is and where you want the PivotTable.



Specifying the location for the PivotTable

Use the bottom section of the Create PivotTable dialog box to indicate the location for your PivotTable. The default location is on a new worksheet, but you can specify any range on any worksheet, including the worksheet that contains the data.

Click OK, and Excel creates an empty PivotTable and displays a PivotTable Fields task pane, as shown in Figure 26.9.

TIP

The PivotTable Fields task pane is typically docked on the right side of the Excel window. Drag its title bar to move it anywhere you like. Also, if you click a cell outside the PivotTable, the task pane is temporarily hidden.

Laying out the PivotTable

Next, set up the actual layout of the PivotTable. You can do so by using any of the following techniques:

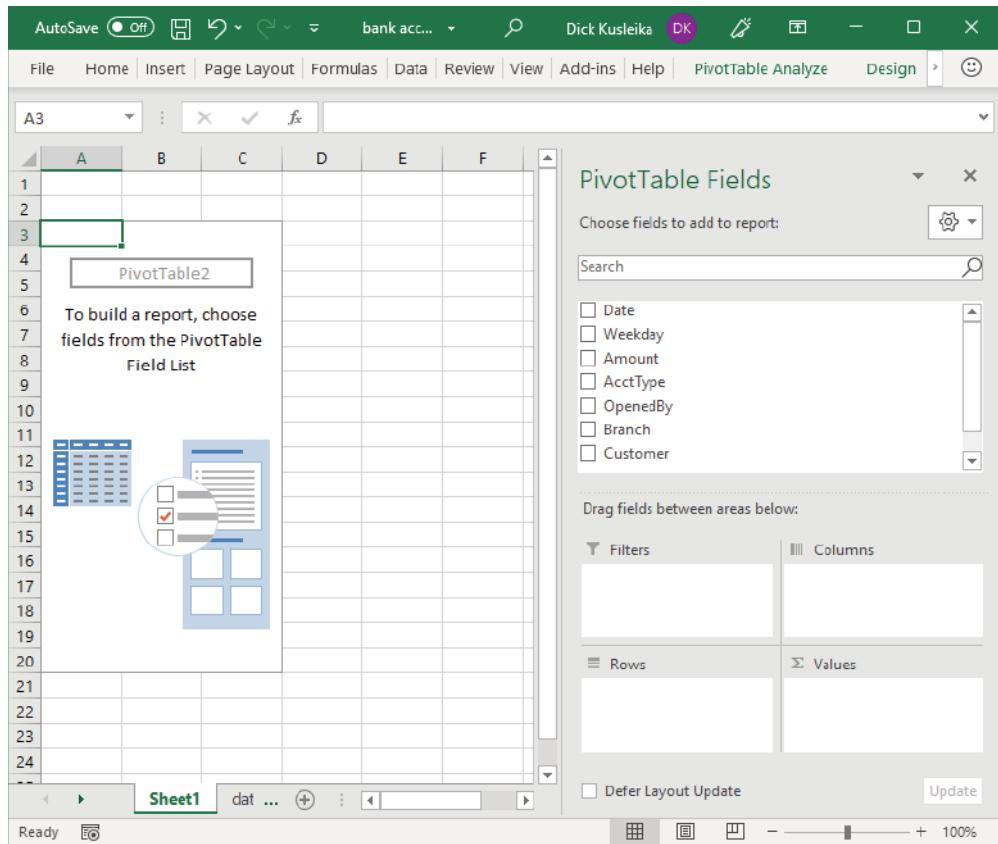
- Drag the field names (at the top of the PivotTable Fields task pane) to one of the four boxes at the bottom of the task pane.
- Place a check mark next to the item at the top of the PivotTable Fields task pane. Excel places the field into one of the four boxes at the bottom. You can drag it to a different box, if necessary.

Part IV: Managing and Analyzing Data

- Right-click a field name at the top of the PivotTable Fields task pane and choose its location from the shortcut menu (for example, Add to Row Labels).

FIGURE 26.9

Use the PivotTable Fields task pane to build the PivotTable.



The following steps create the PivotTable presented earlier in this chapter (see “A PivotTable example”). For this example, drag the items from the top of the PivotTable Fields task pane to the areas at the bottom of the PivotTable Fields task pane.

1. **Drag the Amount field into the Values area.** At this point, the PivotTable displays the total of all the values in the Amount column.
2. **Drag the AcctType field into the Rows area.** Now the PivotTable shows the total amount for each of the account types.

3. **Drag the Branch field into the Columns area.** The PivotTable shows the amount for each account type, cross-tabulated by branch. The PivotTable updates itself automatically with every change you make in the PivotTable Fields task pane.
4. **Right-click any cell in the PivotTable and choose Number Format.** Excel displays the Number tab of the Format Cells dialog box.
5. **Select a number format and click OK.** Excel applies the format to all numeric cells in the PivotTable.

Figure 26.10 shows the completed PivotTable.

FIGURE 26.10

After a few simple steps, the PivotTable shows a summary of the data.

The screenshot shows a Microsoft Excel spreadsheet titled "bank acc...". The PivotTable Fields task pane is open on the right side, showing fields for Amount, AcctType, Branch, and Customer. The PivotTable itself shows a summary of account types (Central, North County, Westside) and their Grand Totals, categorized by Branch (CD, Checking, IRA, Savings). The PivotTable Fields pane also shows filters for Branch and AcctType, and values for Sum of Amount.

AcctType	Branch	Central	North County	Westside	Grand Total
CD		1,361,885	1,209,910	650,237	3,222,032
Checking		802,403	392,516	292,995	1,487,914
IRA		68,380	134,374	10,000	212,754
Savings		885,757	467,414	336,088	1,689,259
Grand Total		3,118,425	2,204,214	1,289,320	6,611,959

PivotTable Terminology

Understanding the terminology associated with PivotTables is the first step in mastering this feature. Refer to the accompanying figure to get your bearings.

	A	B	C	D	E	F
1	OpenedBy	(All)				
2						
3	Sum of Amount		Column Labels			
4	Row Labels	AcctType	Existing	New	Grand Total	
5	Central					
6		CD	974,112	387,773	1,361,885	
7		Checking	505,822	296,581	802,403	
8		IRA	68,380		68,380	
9		Savings	548,198	337,559	885,757	
10	Central Total		2,096,512	1,021,913	3,118,425	
11						
12	North County					
13		CD	845,522	364,388	1,209,910	
14		Checking	208,375	184,141	392,516	
15		IRA	125,374	9,000	134,374	
16		Savings	286,891	180,523	467,414	
17	North County Total		1,466,162	738,052	2,204,214	
18						
19	Westside					
20		CD	356,079	294,158	650,237	
21		Checking	144,391	148,604	292,995	
22		IRA	10,000		10,000	
23		Savings	291,728	44,360	336,088	
24	Westside Total		802,198	487,122	1,289,320	
25						
26	Grand Total		4,364,872	2,247,087	6,611,959	
27						

Column labels A field that has a column orientation in the PivotTable. Each item in the field occupies a column. In the figure, Customer represents a column field that contains two items (Existing and New). You can have nested column fields.

Grand totals A row or column that displays totals for all cells in a row or column in a PivotTable. You can specify that grand totals be calculated for rows, columns, both, or neither. The PivotTable in the figure shows grand totals for both rows and columns.

Group A collection of items treated as a single item. You can group items manually or automatically (group dates into months, for example). The PivotTable in the figure does not have defined groups.

Item An element in a field that appears as a row or column header in a PivotTable. In the figure, Existing and New are items for the Customer field. The Branch field has three items: Central, North County, and Westside. AcctType has four items: CD, Checking, IRA, and Savings.

Refresh Recalculates the PivotTable after making changes to the source data.

Row labels A field that has a row orientation in the PivotTable. Each item in the field occupies a row. You can have nested row fields. In the figure, both Branch and AcctType represent row fields.

Source data The data used to create a PivotTable. It can reside in a worksheet or an external database.

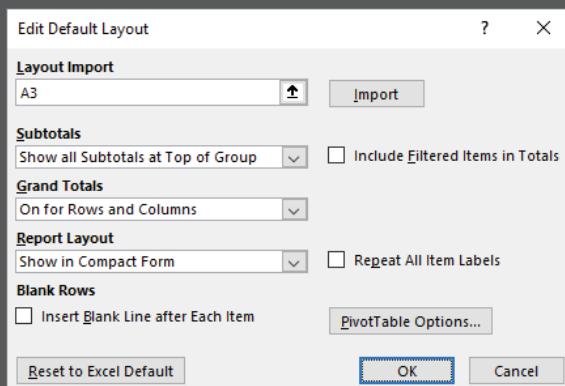
Subtotals A row or column that displays subtotals for detail cells in a row or column in a PivotTable. The PivotTable in the figure displays subtotals for each branch, below the data. You can also display subtotals above the data or hide subtotals. The label for a subtotal is the Item name that's being totaled and the word *Total*.

Table Filter A field that has a page orientation in the PivotTable, which is used to limit what data is summarized. You can display one item, multiple items, or all items in a page field at one time. In the figure, OpenedBy represents a page field that displays All (that is, not filtered).

Values area The cells in a PivotTable that contain the summary data. Excel offers several ways to summarize the data (sum, average, count, and so on).

Tip

If you find that you're making the same layout changes to PivotTables every time you create one, you can save certain layout options as your default. Choose File \Rightarrow Options \Rightarrow Data and click Edit Default Layout to show the Edit Default Layout dialog box.



Here, you can set certain options that all new PivotTables will inherit. Even easier, use the Import button to set all the options automatically based on the PivotTable you specify. If you want to go back to the layout that ships with Excel, click the Reset to Excel Default button.

Formatting the PivotTable

By default, PivotTables use General number formatting. To change the number format for all data, right-click any value and choose Number Format from the shortcut menu. Then use the Format Cells dialog box to change the number format for the displayed data.

Part IV: Managing and Analyzing Data

You can apply any of several built-in styles to a PivotTable. Select any cell in the PivotTable and then choose Design \Rightarrow PivotTable Styles to select a style. Fine-tune the display by using the controls in the Design \Rightarrow PivotTable Style Options group.

You can also use the controls from the Design \Rightarrow Layout group to control various elements in the PivotTable. You can adjust any of the following elements:

Subtotals Hide subtotals or choose where to display them (above or below the data).

Grand Totals Choose which types, if any, to display.

Report Layout Choose from three different layout styles (compact, outline, or tabular). You can also choose to hide repeating labels.

Blank Row Add a blank row between items to improve readability.

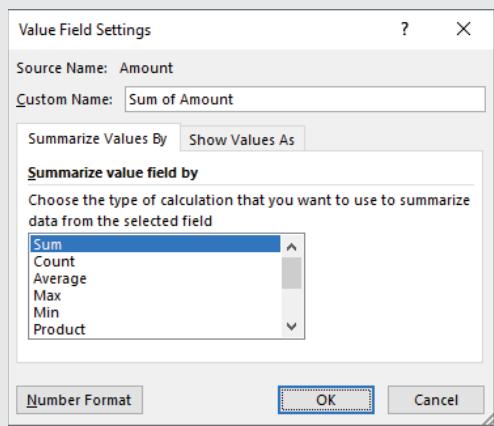
The PivotTable Analyze \Rightarrow Show group contains additional options that affect the appearance of your PivotTable. For example, you use the Field Headers button to toggle the display of the field headings.

Still more PivotTable options are available from the PivotTable Options dialog box. To display this dialog box, choose PivotTable Analyze \Rightarrow PivotTable \Rightarrow Options, or right-click any cell in the PivotTable and choose PivotTable Options from the shortcut menu.

The best way to become familiar with all these layout and formatting options is to experiment.

PivotTable Calculations

PivotTable data is most frequently summarized using a sum. However, you can display your data using a number of different aggregation functions, specified in the Value Field Settings dialog box. The quickest way to display this dialog box is to right-click any value in the PivotTable and choose Value Field Settings from the shortcut menu. This dialog box has two tabs: Summarize Values By and Show Values As.



Use the Summarize Values By tab to select a different summary function. Your choices are Sum, Count, Average, Max, Min, Product, Count Numbers, StdDev, StdDevP, Var, and Varp.

To display your values in a different form, use the drop-down control on the Show Values As tab. You have many options from which to choose, including as a percentage of the grand total or row total.

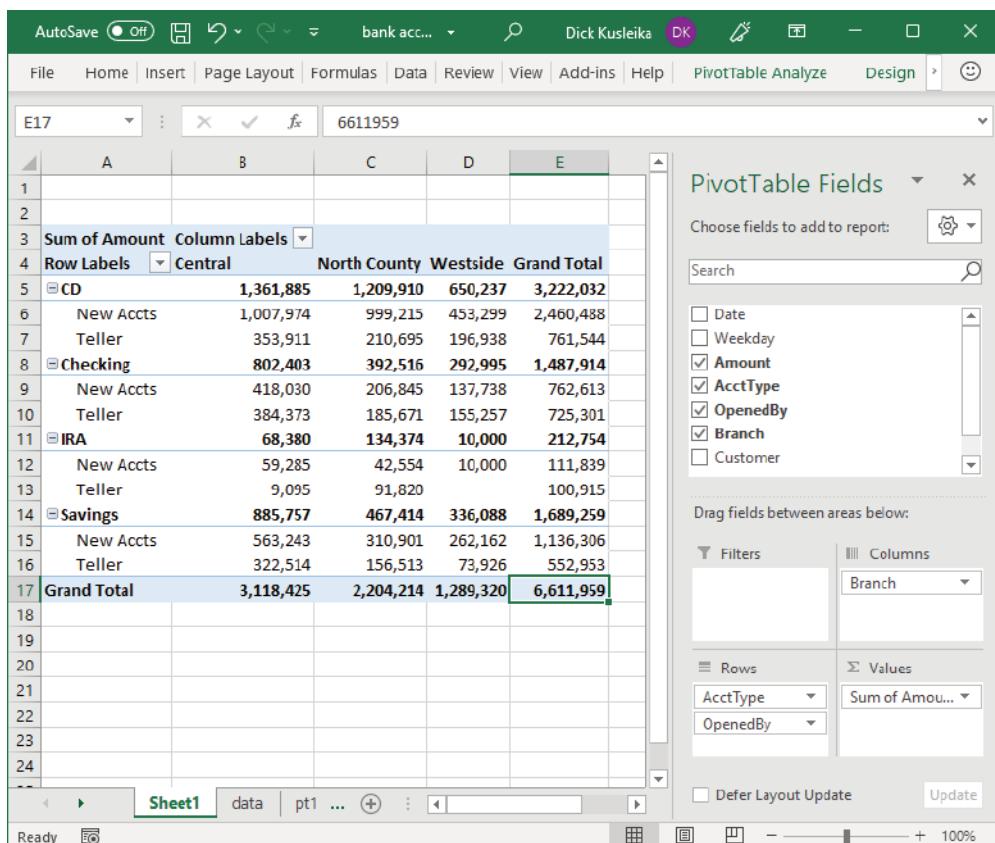
This dialog box also provides a way to apply a number format to the values. Just click the Number Format button and choose your number format.

Modifying the PivotTable

After you create a PivotTable, changing it is easy. For example, you can add further summary information by using the PivotTable Fields task pane. Figure 26.11 shows our PivotTable after we dragged a second field (OpenedBy) to the Rows section in the PivotTable Fields task pane.

FIGURE 26.11

Two fields are used for row labels.



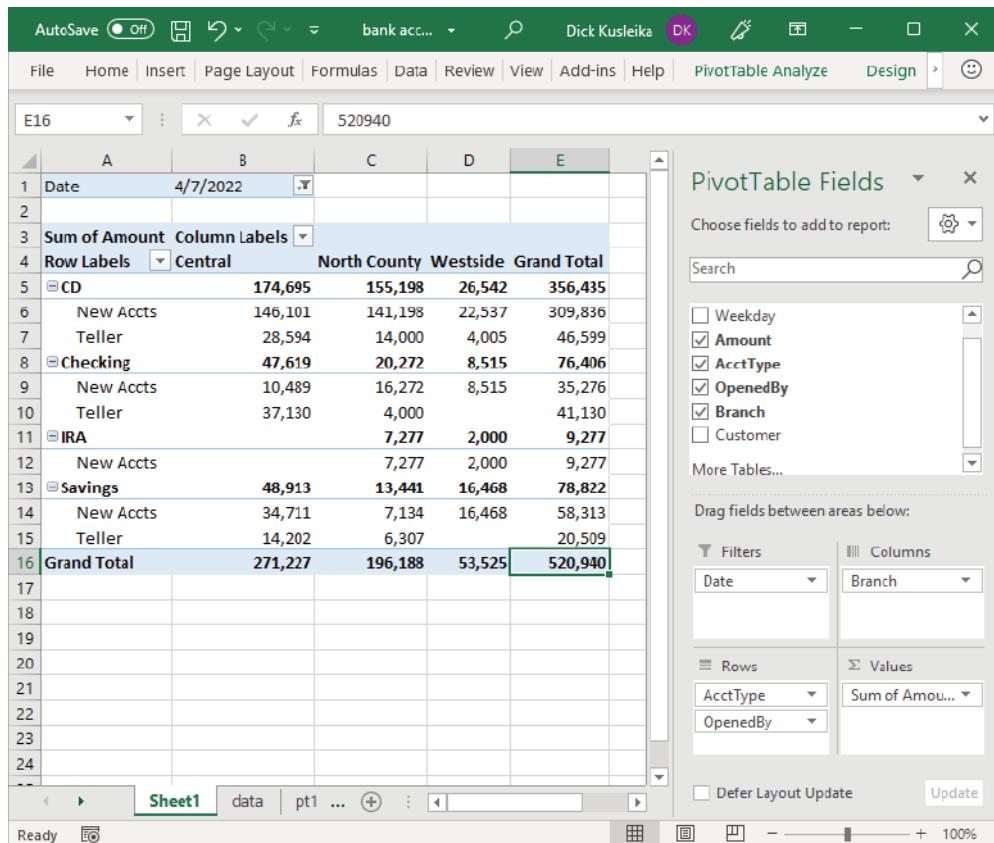
Part IV: Managing and Analyzing Data

Here are some tips on other PivotTable modifications you can make:

- To remove a field from the PivotTable, select it in the bottom part of the PivotTable Fields task pane and then drag it away.
- If an area has more than one field, you can change the order in which the fields are listed by dragging the field names. Doing so determines how nesting occurs, and it affects the appearance of the PivotTable.
- To remove a field temporarily from the PivotTable, remove the check mark from the field name in the top part of the PivotTable Fields task pane. The PivotTable is redisplayed without that field. Place the check mark back on the field name, and it appears in its previous section.
- If you add a field to the Filters section, the field items appear in a drop-down list, which allows you to filter the displayed data by one or more items. Figure 26.12 shows an example. The Date field was put into the Filters area. The PivotTable is now showing the data only for a single day (once it was selected from the drop-down list in cell B1).

FIGURE 26.12

The PivotTable is filtered by date.



Copying a PivotTable's Content

A PivotTable is flexible, but it does have some limitations. For example, you can't add new rows or columns, change any of the calculated values, or enter formulas within the PivotTable. If you want to manipulate a PivotTable in ways not normally permitted, make a copy of it so that it's no longer linked to its data source.

To copy a PivotTable, select the entire table and choose Home ➔ Clipboard ➔ Copy (or press Ctrl+C). Then activate a new worksheet and choose Home ➔ Clipboard ➔ Paste ➔ Paste Values. The PivotTable formatting is not copied—even if you repeat the operation and use the Formats option in the Paste Special dialog box.

To copy the PivotTable and its formatting, use the Office Clipboard to paste. If the Office Clipboard is not displayed, click the dialog box launcher in the bottom right of the Home ➔ Clipboard group.

The contents of the PivotTable are copied to the new location so that you can do whatever you like to them.

Note that the copied information is not a PivotTable, and it is no longer linked to the source data. If the source data changes, your copied PivotTable will not reflect these changes.

Seeing More PivotTable Examples

To demonstrate the flexibility of this feature, we created some additional PivotTables. The examples use the bank account data and answer the questions posed earlier in this chapter. (See “A PivotTable example.”)

What is the daily total new deposit amount for each branch?

Figure 26.13 shows the PivotTable that answers this question.

- The Branch field is in the Columns section.
- The Date field is in the Rows section.
- The Amount field is in the Values section and is summarized by Sum.

Note that the PivotTable can also be sorted by any column. For example, you can sort the Grand Total column in descending order to find out which day of the month had the largest amount of new funds. To sort, just right-click any cell in the column to sort and choose Sort from the shortcut menu.

Which day of the week accounts for the most deposits?

Figure 26.14 shows the PivotTable that answers this question.

- The Weekday field is in the Rows section.
- The Amount field is in the Values section and is summarized by Sum.

We added conditional formatting data bars to make it easier to see how the days compare. As you see, the largest deposit days are Fridays.

Part IV: Managing and Analyzing Data

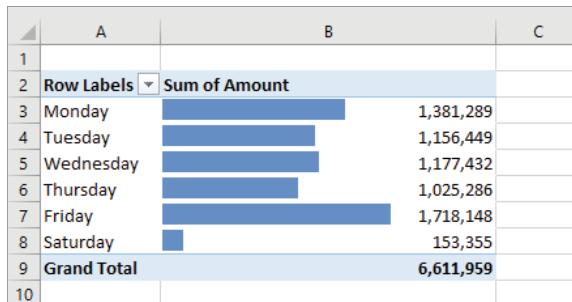
FIGURE 26.13

This PivotTable shows daily totals for each branch.

	A	B	C	D	E
1					
2	Sum of Amount	Branch			
3	Row Labels	Central North County	Westside	Grand Total	
4	4/1/2022	179,011	139,196	51,488	369,695
5	4/2/2022	72,256	27,805	7,188	107,249
6	4/4/2022	146,290	164,305	122,828	433,423
7	4/5/2022	101,480	50,294	97,415	249,189
8	4/6/2022	188,018	91,724	52,738	332,480
9	4/7/2022	271,227	196,188	53,525	520,940
10	4/8/2022	105,087	77,674	92,013	274,774
11	4/11/2022	172,920	43,953	89,258	306,131
12	4/12/2022	70,300	44,621	39,797	154,718
13	4/13/2022	143,921	176,698	29,075	349,694
14	4/14/2022	117,800	114,418	36,064	268,282
15	4/15/2022	191,611	62,787	85,015	339,413
16	4/18/2022	79,394	72,262	48,337	199,993
17	4/19/2022	209,916	213,728	53,721	477,365
18	4/20/2022	125,276	140,739	56,444	322,459
19	4/21/2022	79,355	35,753	3,419	118,527
20	4/22/2022	188,509	236,269	97,210	521,988
21	4/25/2022	218,889	137,025	85,828	441,742
22	4/26/2022	150,139	29,040	95,998	275,177
23	4/27/2022	56,379	72,948	43,472	172,799
24	4/28/2022	62,192	43,217	12,128	117,537
25	4/29/2022	170,279	22,570	19,429	212,278
26	4/30/2022	18,176	11,000	16,930	46,106
27	Grand Total	3,118,425	2,204,214	1,289,320	6,611,959

FIGURE 26.14

This PivotTable shows new account totals by day of the week.





See Chapter 16, “Using Formulas with Tables and Conditional Formatting,” for more information about conditional formatting.

How many accounts were opened at each branch, broken down by account type?

Figure 26.15 shows a PivotTable that answers this question.

- The AcctType field is in the Columns section.
- The Branch field is in the Rows section.
- The Amount field is in the Values section and is summarized by Count.

FIGURE 26.15

This PivotTable uses the Count function to summarize the data.

	A	B	C	D	E	F	G
1							
2	Count of Amount	AcctType					
3	Row Labels	CD	Checking	IRA	Savings	Grand Total	
4	Central	97	158	8	99	362	
5	North County	60	61	15	61	197	
6	Westside	54	59	5	35	153	
7	Grand Total	211	278	28	195	712	
8							

So far, all the PivotTable examples have used the Sum summary function. In this case, we changed the summary function to Count. To change the summary function to Count, right-click any cell in the Values area and choose Summarize Values By \Rightarrow Count from the shortcut menu.

How much money was used to open the accounts?

Figure 26.16 shows a PivotTable that answers this question. For example, 253 (or 35.53%) of the new accounts were for an amount of \$5,000 or less.

This PivotTable is unusual because it uses only one field: Amount.

- The Amount field is in the Rows section (grouped, to show dollar ranges).
- The Amount field is also in the Values section and is summarized by Count.
- A third instance of the Amount field is the Values section, summarized by Count and displayed as Percent of Column Total.

When you initially add the Amount field to the Rows section, the PivotTable will show one row for each unique dollar amount. To group the values, right-click one of the row items and choose Group from the shortcut menu. Then use the Grouping dialog box to set up bins of \$5,000 increments. Note that the Grouping dialog box does not appear if you select more than one row label.

Part IV: Managing and Analyzing Data

FIGURE 26.16

This PivotTable counts the number of accounts that fall into each value range.

	A	B	C	D
1				
2	Row Labels ▾	Count	Pct	
3	1-5000	253	35.53%	
4	5001-10000	192	26.97%	
5	10001-15000	222	31.18%	
6	15001-20000	19	2.67%	
7	20001-25000	3	0.42%	
8	25001-30000	1	0.14%	
9	30001-35000	3	0.42%	
10	40001-45000	3	0.42%	
11	45001-50000	5	0.70%	
12	60001-65000	2	0.28%	
13	70001-75000	5	0.70%	
14	75001-80000	1	0.14%	
15	85001-90000	3	0.42%	
16	Grand Total	712	100.00%	
17				

The second instance of the Amount field (in the Values section) is summarized by Count. To change from the default Sum, right-click any value and choose Summarize Data By \leftrightarrow Count from the shortcut menu.

Add another instance of Amount to the Values section and set it up to display as a percentage. Right-click a value in column C and choose Show Values As \leftrightarrow % of Column Total. This option is also available in the Show Values As tab of the Value Field Settings dialog box.

What types of accounts do tellers open most often?

The PivotTable in Figure 26.17 shows that the most common account opened by tellers is a checking account.

- The AcctType field is in the Rows section.
- The OpenedBy field is in the Filters section.
- The Amount field is in the Values section (summarized by Count).
- A second instance of the Amount field is in the Values section (shown as % of Column Total).

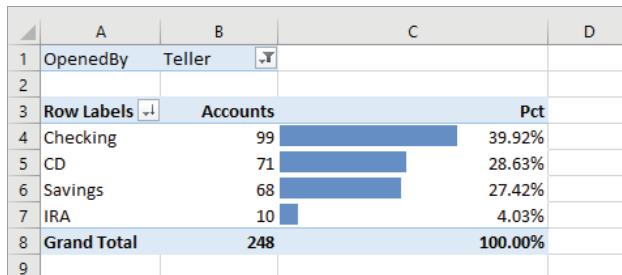
This PivotTable uses the OpenedBy field as a filter and is showing the data only for tellers. Sort the rows so that the largest value is at the top by right-clicking any of the values, choosing Sort, and choosing Sort Largest to Smallest. To display the data bars, use conditional formatting.



See Chapter 16 for more information about conditional formatting.

FIGURE 26.17

This PivotTable uses a filter to show only the teller data.



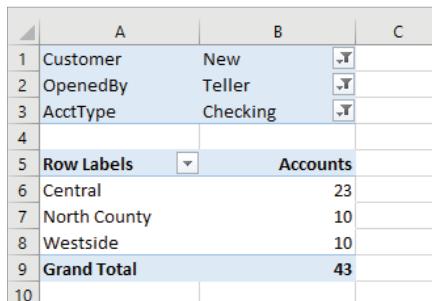
In which branch do tellers open the most checking accounts for new customers?

Figure 26.18 shows a PivotTable that answers this question. At the Central branch, tellers opened 23 checking accounts for new customers.

- The Customer field is in the Filters section.
- The OpenedBy field is in the Filters section.
- The AcctType field is in the Filters section.
- The Branch field is in the Rows section.
- The Amount field is in the Values section, summarized by Count.

FIGURE 26.18

This PivotTable uses three filters.



This PivotTable uses three report filters. The Customer field is filtered to show only New, the OpenedBy field is filtered to show only Teller, and the AcctType field is filtered to show only Checking.

Learning More

The examples in this chapter should give you an appreciation for the power and flexibility of Excel PivotTables. The next chapter drills down a bit deeper and covers some advanced features of PivotTables—with lots of examples.

Analyzing Data with PivotTables

IN THIS CHAPTER

- Creating a PivotTable from non-numeric data
- Grouping items in a PivotTable
- Creating a calculated field or a calculated item in a PivotTable
- Understanding the data model feature

The previous chapter introduced PivotTables. There we presented several examples to demonstrate the types of PivotTable summaries that you can generate from a set of data.

This chapter continues the discussion and explores the details of creating effective PivotTables. Creating a basic PivotTable is easy, and the examples in this chapter demonstrate additional PivotTable features that you may find helpful. We urge you to try these techniques with your own data. If you don't have suitable data, use the files available on this book's website.

Working with Non-Numeric Data

Most PivotTables are created to aggregate numeric data, but PivotTables are also useful with some types of non-numeric data. With non-numbers, you can count instead of summing or averaging.

Figure 27.1 shows a table and a PivotTable generated from the table. The table consists of a list of 400 employees, along with their location and gender. As you can see, the table has no numeric values, but you can create a useful PivotTable that counts the items rather than sums them. The PivotTable (in range E2 : H10) cross-tabulates the Location field by the Gender field for the 400 employees, and it shows the count for each combination of location and gender.



A workbook that demonstrates the PivotTable created from non-numeric data is available on this book's website at www.wiley.com/go/excel365bible. The file is named `employee_list.xlsx`.

Part IV: Managing and Analyzing Data

FIGURE 27.1

This table doesn't have any numeric fields, but you can use it to generate the PivotTable shown next to it.

The screenshot shows a Microsoft Excel spreadsheet. On the left is a table of 24 employees with columns for Employee, Location, and Gender. To the right is a PivotTable with three sections: a summary table, a detail table, and a grand total row.

	A	B	C	D	E	F	G	H	I
1	Employee	Location	Gender						
2	Al Grubbs	California	Male						
3	Sarah Parks	New York	Female						
4	Cheryl Cory	California	Female						
5	Gregory Steiger	California	Male						
6	Sheila Wigfall	California	Female						
7	Pedro H. Nicholson	Arizona	Male						
8	Howard Keach	California	Male						
9	Heather Lichtenstein	Washington	Female						
10	Janet Woodson	Arizona	Female						
11	Hosea Pierson	New York	Male						
12	Nadine Blankenship	New York	Female						
13	Roy Greene	New York	Male						
14	William N. Campbell	New York	Male						
15	Stephen Foster	New York	Male						
16	Charles S. Billings	Pennsylvania	Male						
17	Margaret Sirois	California	Female						
18	Phyllis Todd	Massachusetts	Female						
19	Mary Brinkmann	New York	Female						
20	Janie Little	Massachusetts	Female						
21	Bob Hunsberger	Pennsylvania	Male						
22	Cassandra Jackson	Arizona	Female						
23	Phyllis Todd	Massachusetts	Female						
24	Donald Jackson	California	Male						

Count

	Female	Male	Total
Arizona	5	15	20
California	44	64	108
Massachusetts	43	47	90
New York	51	40	91
Pennsylvania	17	29	46
Washington	16	29	45
Total	176	224	400

Here are the PivotTable Fields task pane settings we used for this PivotTable:

- The Gender field is used for the Columns.
- The Location field is used for the Rows.
- Location is also used for the Values, and it is summarized by Count.
- The PivotTable has the field headers turned off by using the Field Headers toggle control in the PivotTable Analyze \leftrightarrow Show group.

NOTE

The Employee field is not used. This example uses the Location field for the Values section, but you can use any of the three fields because the PivotTable is displaying a count.

Figure 27.2 shows the PivotTable after making some additional changes.

- To show percentages, add a second instance of the Location field to the Values section. Then right-click a value in that column and choose Show Values As \leftrightarrow % of Column Total.

- Change the field names in the PivotTable to **Ct** and **Pct** by simply selecting those cells and typing the new name.
- Select a PivotTable style that makes it easier to distinguish the columns.

FIGURE 27.2

Changing the PivotTable to show counts and percentages

Row Labels	Color	Female Ct	Male		Total Ct	Total Pct
	Ct		Ct	Pct		
Arizona	5	2.84%	15	6.70%	20	5.00%
California	44	25.00%	64	28.57%	108	27.00%
Massachusetts	43	24.43%	47	20.98%	90	22.50%
New York	51	28.98%	40	17.86%	91	22.75%
Pennsylvania	17	9.66%	29	12.95%	46	11.50%
Washington	16	9.09%	29	12.95%	45	11.25%
Grand Total	176	100.00%	224	100.00%	400	100.00%

Grouping PivotTable Items

One of the most useful features of a PivotTable is the ability to combine items into groups. You can group items that appear in the Rows or Columns section in the PivotTable Fields task pane. Excel offers two ways to group items:

Manually After creating the PivotTable, select the items to be grouped and then choose PivotTable Analyze \Rightarrow Group \Rightarrow Group Selection, or you can select the items, right-click, and choose Group from the shortcut menu.

Automatically If the items are numeric (or dates), use the Grouping dialog box to specify how you would like to group the items. Select any single item and then choose PivotTable Analyze \Rightarrow Group \Rightarrow Group Field, or right-click a single item and choose Group from the shortcut menu. In either case, the Grouping dialog box appears. Use this dialog box to specify how to group the items.

Note

If you plan on creating multiple PivotTables that use different groupings, make sure you read the sidebar “Multiple Groups from the Same Data Source” later in this chapter.

Grouping items manually

Figure 27.3 shows the PivotTable example from the previous sections, with two groups created from the row labels. To create the first group, hold down the Ctrl key and select Arizona, California, and Washington in the PivotTable. Then right-click the selected cells and choose Group from the shortcut menu. Select the three other states in the same manner and create a second group. Replace the default group names (Group 1 and Group 2) with more meaningful names (Western Region and Eastern Region).

Part IV: Managing and Analyzing Data

FIGURE 27.3

A PivotTable with two groups

Count			
	Female	Male	Total
Western Region			
Arizona	5	15	20
California	44	64	108
Washington	16	29	45
Eastern Region			
Massachusetts	43	47	90
New York	51	40	91
Pennsylvania	17	29	46
Total	176	224	400

You can create any number of groups and even create groups of groups.

Excel provides a number of layout options for displaying a PivotTable. You may want to experiment with these options when you use groups. These commands are on the Design tab of the Ribbon. There are no rules for choosing a particular option. The key is to try a few and see which makes your PivotTable look the best. In addition, try various style options in the Design tab. Often, the style that you choose can greatly enhance readability.

Figure 27.4 shows PivotTables using various options for displaying subtotals, grand totals, and styles.

FIGURE 27.4

PivotTables with options for subtotals and grand totals

	A	B	C	D	E	F	G	H	I	J	K
1				Count			Count				
2		■ Western Region	Arizona	20			■ Western Region				
3			California	108			Arizona	5	15	20	
4			Washington	45			California	44	64	108	
5							Washington	16	29	45	
6		■ Eastern Region	Massachusetts	90			■ Western Region Total	65	108	173	
7			New York	91							
8			Pennsylvania	46			■ Eastern Region				
9							Massachusetts	43	47	90	
10							New York	51	40	91	
11							Pennsylvania	17	29	46	
12							■ Eastern Region Total	111	116	227	
13				Count							
14		■ Western Region	Arizona	20			Total	176	224	400	
15			California	108							
16			Washington	45							
17											
18		■ Eastern Region	Massachusetts	90							
19			New York	91							
20			Pennsylvania	46							
21											
22											
23											
24											



A workbook that contains these grouping examples is available on this book's website at www.wiley.com/go/excel365bible. The file is named `grouping examples.xlsx`.

Grouping items automatically

When a field contains numbers, dates, or times, Excel can create groups automatically. The two examples in the following sections demonstrate automatic grouping.

Grouping by date

Figure 27.5 shows a portion of a simple table with two fields: Date and Sales. This table has 730 rows and covers the dates between January 1, 2022 and December 31, 2023. The goal is to summarize the sales information by month.

FIGURE 27.5

You can use a PivotTable to summarize the sales data by month.

	A	B	C
1	Date	Sales	
2	1/1/2022	3,830	
3	1/2/2022	3,763	
4	1/3/2022	4,362	
5	1/4/2022	3,669	
6	1/5/2022	3,942	
7	1/6/2022	4,488	
8	1/7/2022	4,416	
9	1/8/2022	3,371	
10	1/9/2022	3,628	
11	1/10/2022	4,548	
12	1/11/2022	5,493	
13	1/12/2022	5,706	
14	1/13/2022	6,579	
15	1/14/2022	6,333	
16	1/15/2022	6,101	
17	1/16/2022	5,289	
18	1/17/2022	5,349	
19	1/18/2022	5,814	
20	1/19/2022	6,501	



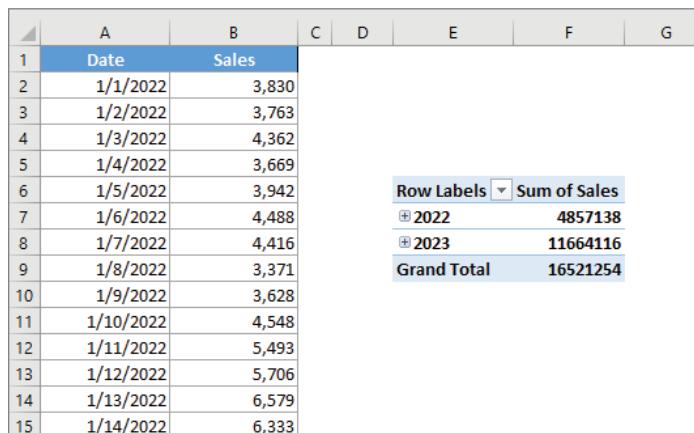
A workbook demonstrating how to group PivotTable items by date is available on this book's website at www.wiley.com/go/excel365bible. The file is `grouping sales by date.xlsx`.

Figure 27.6 shows a PivotTable (in columns E : F) created from the data. The Date field is in the Rows section, and the Sales field is in the Values section. When you add a date field to a PivotTable, Excel automatically groups the dates and adds the grouped fields to your PivotTable. If you prefer to not have your dates grouped, you can click Undo immediately after you add the date field to the PivotTable. You can also modify the grouping manually.

Part IV: Managing and Analyzing Data

FIGURE 27.6

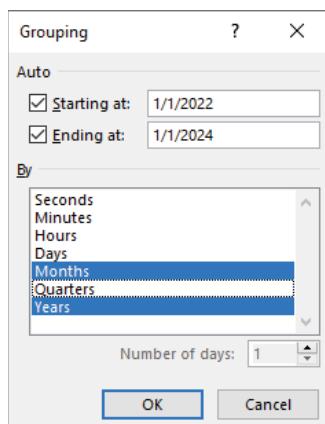
The PivotTable with Excel's automatic grouping



To group the items by month, select either of the years and choose PivotTable Analyze \Rightarrow Group \Rightarrow Group Field (or right-click and choose Group from the shortcut menu). The Grouping dialog box, shown in Figure 27.7, appears. Excel supplies values for the Starting At and Ending At fields. The values cover the entire range of data, and you can change them if you like.

FIGURE 27.7

Use the Grouping dialog box to group PivotTable items by dates.



In the By list box, select only Months and Years and deselect any other groupings Excel may have automatically created (such as Quarters). Verify that the starting and ending dates are

correct for your data. Click OK. The Date items in the PivotTable are grouped by years and by months, as shown in Figure 27.8.

FIGURE 27.8

The PivotTable, after grouping by month and year

D	E	F	G
Row Labels ▾ Sum of Sales			
	2022		
Jan	167624		
Feb	137825		
Mar	214896		
Apr	100872		
May	158005		
Jun	117649		
Jul	295248		
Aug	518966		
Sep	612673		
Oct	699854		
Nov	863085		
Dec	970441		
	2023		
Jan	974625		
Feb	969696		
Mar	1081596		
Apr	983306		
May	1044322		
Jun	930076		
Jul	961557		
Aug	938433		
Sep	975503		
Oct	948120		
Nov	950493		
Dec	906389		
	Grand Total		
	16521254		

NOTE

If you select only Months in the By list box in the Grouping dialog box, months in different years combine. For example, the January item would display the sum of sales for 2022 and 2023.

Figure 27.9 shows another view of the data, grouped by quarter and by year.

Part IV: Managing and Analyzing Data

FIGURE 27.9

This PivotTable shows sales by quarter and by year.

D	E	F	G
Row Labels ▾ Sum of Sales			
2022	4857138		
Qtr1	520345		
Qtr2	376526		
Qtr3	1426887		
Qtr4	2533380		
2023	11664116		
Qtr1	3025917		
Qtr2	2957704		
Qtr3	2875493		
Qtr4	2805002		
Grand Total	16521254		

Multiple Groups from the Same Data Source

If you create multiple PivotTables from the same data source, you may have noticed that grouping a field in one PivotTable affects the other PivotTables. Specifically, all the other PivotTables automatically use the same grouping. Sometimes, this is exactly what you want. At other times, it's not at all what you want. For example, you might like to see two PivotTable reports: one that summarizes data by month and year and another that summarizes the data by quarter and year.

Grouping affects other PivotTables because all the PivotTables are using the same PivotTable "cache." Unfortunately, there is no direct way to force a PivotTable to use a new cache. But there is a way to trick Excel into using a new cache. The trick involves giving multiple range names to the source data.

For example, name your source range Table1 and then give the same range a second name: Table2. The easiest way to name a range is to use the Name box to the left of the Formula bar. Select the range, type a name in the Name box, and press Enter. Then, with the range still selected, type a different name and press Enter. Excel will display only the first name, but you can verify that both names exist by choosing Formulas \Rightarrow Defined Names \Rightarrow Name Manager.

When you create the first PivotTable, specify Table1 as the Table/Range. When you create the second PivotTable, specify Table2 as the Table/Range. Each PivotTable will use a separate cache, and you can create groups in one PivotTable, independent of the other PivotTable.

You can use this trick with existing PivotTables. Make sure you give the data source a different name. Then select the PivotTable and choose PivotTable Analyze \Rightarrow Data \Rightarrow Change Data Source. In the Change PivotTable Data Source dialog box, type the new name that you gave to the range. This will cause Excel to create a new PivotCache for the PivotTable.

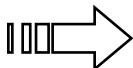
Grouping by time

Figure 27.10 shows a set of data in columns A:B. Each row is a reading from a measurement instrument, taken at one-minute intervals throughout an entire day. The table has 1,440 rows, each representing one minute. The PivotTable (in columns D:G) summarizes the data by hour.

FIGURE 27.10

This PivotTable is grouped by hours.

	A	B	C	D	E	F	G	H
1	Time	Reading		Average	Minimum	Maximum		
2	6/15/2022 0:00	105.32		12 AM	110.50	104.37	116.21	
3	6/15/2022 0:01	105.35		1 AM	118.57	112.72	127.14	
4	6/15/2022 0:02	104.37		2 AM	124.39	115.75	130.36	
5	6/15/2022 0:03	106.40		3 AM	122.74	112.85	132.90	
6	6/15/2022 0:04	106.42		4 AM	129.29	123.99	133.52	
7	6/15/2022 0:05	105.45		5 AM	132.91	125.88	141.04	
8	6/15/2022 0:06	107.46		6 AM	139.67	132.69	146.06	
9	6/15/2022 0:07	109.49		7 AM	128.18	117.53	139.65	
10	6/15/2022 0:08	110.54		8 AM	119.24	112.10	129.38	
11	6/15/2022 0:09	110.54		9 AM	134.36	129.11	142.79	
12	6/15/2022 0:10	110.55		10 AM	136.16	130.91	142.89	
13	6/15/2022 0:11	109.56		11 AM	122.79	108.63	138.10	
14	6/15/2022 0:12	107.60		12 PM	111.76	106.43	116.71	
15	6/15/2022 0:13	107.68		1 PM	104.91	98.48	111.86	
16	6/15/2022 0:14	109.69		2 PM	119.71	110.37	130.55	
17	6/15/2022 0:15	107.76		3 PM	131.83	121.92	139.65	
18	6/15/2022 0:16	107.81		4 PM	131.05	123.36	137.94	
19	6/15/2022 0:17	108.83		5 PM	138.90	133.05	145.06	
20	6/15/2022 0:18	109.85		6 PM	134.71	129.29	139.89	
21	6/15/2022 0:19	111.94		7 PM	123.09	113.97	135.23	
22	6/15/2022 0:20	114.04		8 PM	118.13	112.64	125.65	
23	6/15/2022 0:21	112.12		9 PM	112.64	108.09	117.72	
24	6/15/2022 0:22	112.21		10 PM	103.19	96.13	110.49	
25	6/15/2022 0:23	112.25		11 PM	106.01	100.03	111.76	
26	6/15/2022 0:24	113.34		Grand Total	123.11	96.13	146.06	
27	6/15/2022 0:25	112.41						



This workbook, named `time-based grouping.xlsx`, is available on this book's website at www.wiley.com/go/excel365bible.

Here are the settings to use for this PivotTable:

- The Values area has three instances of the Reading field, and each instance displays a different summary method (Average, Minimum, and Maximum). To change the summary method for a column, right-click any cell in the column and choose Summarize Values By and then the appropriate option.
- The Time field is in the Rows section. Use the Grouping dialog box to group by hours.

Using a PivotTable to Create a Frequency Distribution

Excel provides a number of ways to create a frequency distribution, but none of these methods is easier than using a PivotTable. Figure 27.11 shows part of a table of 221 students and the test score for each. The goal is to determine how many students are in each 10-point range (1–10, 11–20, and so on).

FIGURE 27.11

Creating a frequency distribution for these test scores is simple.

	A	B
1	Student	Score
2	Rhea Madsen	69
3	Jennifer Mendez	81
4	Brett Broyles	100
5	Shirley Smith	28
6	John Brown	93
7	Michael G. Welch	50
8	Donald Tse	100
9	Madeline Stevens	82
10	Howard Porter	86
11	Helen Craven	81
12	Lillie Schultz	75
13	Emily Li	80
14	Michael Long	71
15	Chris Herrman	88
16	Marshall Sherman	48
17	William Grindle	82
18	Pauline Haun	85
19	Lydia J. Evans	82
20	James Weaver	80



This workbook, named `frequency distribution.xlsx`, is available on this book's website at www.wiley.com/go/excel365bible.

The PivotTable is simple.

- The Score field is in the Rows section (grouped).
- Another instance of the Score field is in the Values section (summarized by Count).

The Grouping dialog box that generated the bins specified that the groups start at 1, end at 100, and are incremented by 10.

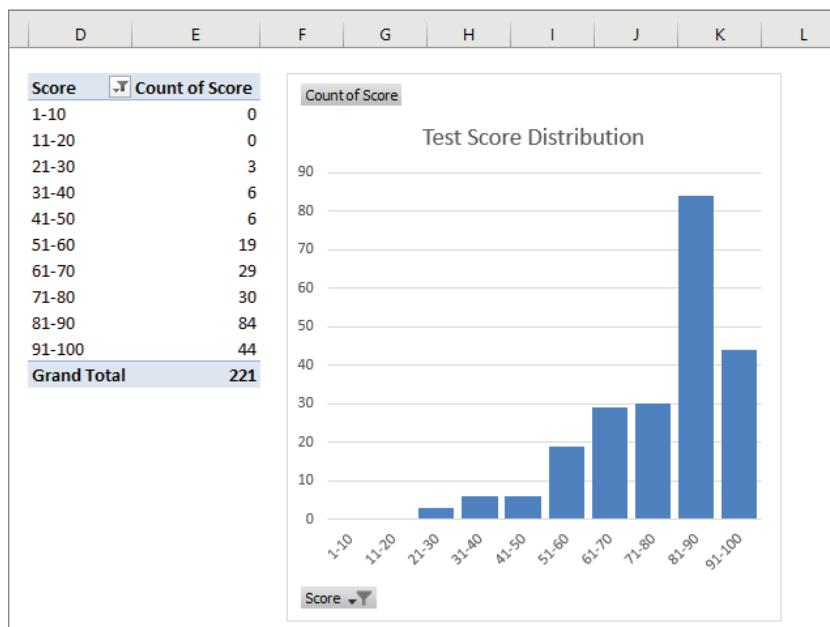
NOTE

By default, Excel does not display items with a count of zero. In this example, no test scores are less than 21, so the 1-10 and 11-20 range items were initially hidden when we created the PivotTable. To force the display of empty bins, right-click any cell and choose Field Settings from the shortcut menu. In the Field Settings dialog box, click the Layout & Print tab and select Show Items with No Data. Next, right-click and choose PivotTable Option and check the For empty cell show check box and type 0 in the input box. This will display a zero for the newly added rows instead of a blank.

Figure 27.12 shows the frequency distribution of the test scores, along with a PivotChart. (See “Creating PivotCharts” later in this chapter.) The Scores are filtered so that the PivotTable (and chart) do not show the <1 category and the >101 category.

FIGURE 27.12

The PivotTable and PivotChart show the frequency distribution for the test scores.

**NOTE**

This example uses the Excel Grouping dialog box to create the groups automatically. If you don't want to group in equal-sized bins, you can create your own groups. For example, you may want to assign letter grades based on the test score. Select the rows for the first group, right-click, and then choose Group from the shortcut menu. Repeat these steps for each additional group. Then replace the default group names with more meaningful names.

Creating a Calculated Field or Calculated Item

Perhaps the most confusing aspect of PivotTables is calculated fields versus calculated items. Many PivotTable users simply avoid dealing with calculated fields and items. However, these features can be useful, and they really aren't that complicated once you understand the way they work.

First, here are some basic definitions:

Calculated field A calculated field is a new field created from other fields in the PivotTable. If your PivotTable source is a worksheet table, an alternative to using a calculated field is to add a new column to the table and create a formula to perform the desired calculation. A calculated field must reside in the Values area of the PivotTable. You can't use a calculated field in the Columns area, in the Rows area, or in the Filters area.

Calculated item A calculated item uses the contents of other items within a field of the PivotTable. If your PivotTable source is a worksheet table, an alternative to using a calculated item is to insert one or more rows and write formulas that use values in other rows. A calculated item must reside in the Columns area, Rows area, or Filters area of a PivotTable. You can't use a calculated item in the Values area.

The formulas used to create calculated fields and calculated items aren't standard Excel formulas. In other words, you don't enter the formulas into cells. Rather, you enter these formulas into a dialog box, and they're stored along with the PivotTable data.

The examples in the following sections use the worksheet table shown in Figure 27.13. The table consists of 5 columns and 48 rows. Each row describes monthly sales information for a particular sales representative. For example, Amy is a sales rep for the North region, and she sold 239 units in January for total sales of \$23,040.

CAUTION

A workbook demonstrating calculated fields and items is available on this book's website at www.wiley.com/go/excel365bible. The file is named `calculated fields and items.xlsx`.

Figure 27.14 shows a PivotTable created from the data. This PivotTable shows Sales (Values area) cross-tabulated by Month (Rows area) and by SalesRep (Columns area).

The examples that follow create:

- A calculated field, to compute average sales per unit
- Four calculated items, to compute the quarterly sales commission

FIGURE 27.13

This data demonstrates calculated fields and calculated items.

	A	B	C	D	E
1	SalesRep	Region	Month	Sales	Units Sold
2	Amy	North	Jan	23,040	239
3	Amy	North	Feb	24,131	79
4	Amy	North	Mar	24,646	71
5	Amy	North	Apr	22,047	71
6	Amy	North	May	24,971	157
7	Amy	North	Jun	24,218	92
8	Amy	North	Jul	25,735	175
9	Amy	North	Aug	23,638	87
10	Amy	North	Sep	25,749	557
11	Amy	North	Oct	24,437	95
12	Amy	North	Nov	25,355	706
13	Amy	North	Dec	25,899	180
14	Bob	North	Jan	20,024	103
15	Bob	North	Feb	23,822	267
16	Bob	North	Mar	24,854	96
17	Bob	North	Apr	22,838	74
18	Bob	North	May	25,320	231
19	Bob	North	Jun	24,733	164
20	Bob	North	Jul	21,184	68

FIGURE 27.14

This PivotTable was created from the sales data.

Sum of Sales	Column Labels					
Row Labels	Amy	Bob	Chuck	Doug	Grand Total	
Jan		23,040	20,024	19,886	26,264	89,214
Feb		24,131	23,822	23,494	29,953	101,400
Mar		24,646	24,854	21,824	25,041	96,365
Apr		22,047	22,838	22,058	29,338	96,281
May		24,971	25,320	20,280	25,150	95,721
Jun		24,218	24,733	23,965	27,371	100,287
Jul		25,735	21,184	23,032	25,044	94,995
Aug		23,638	23,174	21,273	29,506	97,591
Sep		25,749	25,999	21,584	29,061	102,393
Oct		24,437	22,639	19,625	27,113	93,814
Nov		25,355	23,949	19,832	25,953	95,089
Dec		25,899	23,179	20,583	28,670	98,331
Grand Total		293,866	281,715	257,436	328,464	1,161,481

Creating a calculated field

Because a PivotTable is a special type of range, you can't insert new rows or columns within the PivotTable, which means you can't insert formulas to perform calculations with the data in a PivotTable. However, you can create calculated fields for a PivotTable. A calculated field consists of a calculation that can involve other fields.

A calculated field is basically a way to display new information (derived from other fields) in a PivotTable. It's an alternative to creating a new column in your source data. In many cases, you may find it easier to insert a new column in the source range with a formula that performs the desired calculation. A calculated field is most useful when the data comes from a source that you can't easily manipulate, such as an external database.

In the sales example, suppose you want to calculate the average sales amount per unit. You can compute this value by dividing the Sales field by the Units Sold field. The result shows a new field (a calculated field) for the PivotTable.

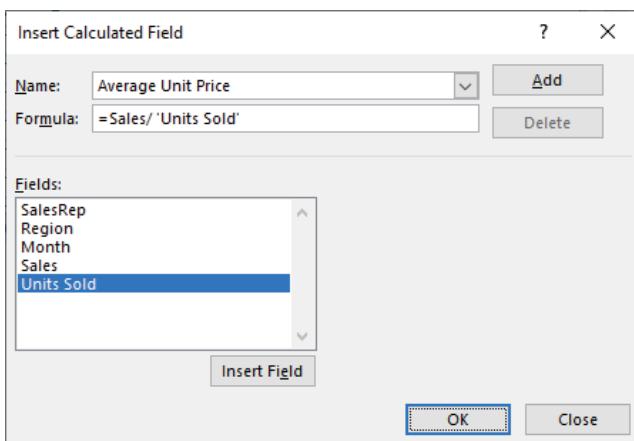
Use the following procedure to create a calculated field that consists of the Sales field divided by the Units Sold field:

1. Select any cell within the PivotTable.
2. Choose **PivotTable Analyze** \Rightarrow **Calculations** \Rightarrow **Fields, Items, & Sets** \Rightarrow **Calculated Field**. The Insert Calculated Field dialog box appears.
3. Enter a descriptive name in the **Name** box and specify the formula in the **Formula** box (see Figure 27.15). The formula can use worksheet functions and other fields from the data source. For this example, the calculated field name is Average Unit Price, and the formula is as follows:

=Sales/'Units Sold'

FIGURE 27.15

The Insert Calculated Field dialog box



4. Click Add to add this new field.
5. Click OK to close the Insert Calculated Field dialog box.

Note

You can create the formula manually by typing it or by double-clicking items in the Fields list box. Double-clicking an item transfers it to the Formula field. Because the Units Sold field contains a space, Excel adds single quotes around the field name.

After you create the calculated field, Excel adds it to the Values area of the PivotTable. (It also appears in the field list of the PivotTable Fields task pane.) You can treat it just like any other field, with one exception: you can't move it to the Rows, Columns, or Filters areas. It must remain in the Values area.

Figure 27.16 shows the PivotTable after adding the calculated field. The new field displayed Sum of Average Unit Price, but the label was shortened to Avg Price.

27

FIGURE 27.16

This PivotTable uses a calculated field.

	H	I	J	K	L	M	N	O	P	Q	R
Column Labels ▾											
Row Labels	Amy		Bob		Chuck		Doug		Total Sales		Total Avg Price
	Sales	Avg Price	Sales	Avg Price	Sales	Avg Price	Sales	Avg Price			
Jan	23,040	96	20,024	194	19,886	209	26,264	285	89,214	169	
Feb	24,131	305	23,822	89	23,494	159	29,953	35	101,400	75	
Mar	24,646	347	24,854	259	21,824	263	25,041	291	96,365	287	
Apr	22,047	311	22,838	309	22,058	230	29,338	132	96,281	208	
May	24,971	159	25,320	110	20,280	45	25,150	104	95,721	88	
Jun	24,218	263	24,733	151	23,965	32	27,371	288	100,287	90	
Jul	25,735	147	21,184	312	23,032	149	25,044	305	94,995	198	
Aug	23,638	272	23,174	203	21,273	28	29,506	286	97,591	91	
Sep	25,749	46	25,999	310	21,584	189	29,061	199	102,393	114	
Oct	24,437	257	22,639	87	19,625	236	27,113	226	93,814	168	
Nov	25,355	36	23,949	220	19,832	283	25,953	320	95,089	98	
Dec	25,899	144	23,179	50	20,583	116	28,670	145	98,331	96	
Grand Total	293,866	117	281,715	138	257,436	86	328,464	142	1,161,481	118	

TIP

The formulas that you develop can also use worksheet functions, but the functions can't refer to cells or named ranges.

Inserting a calculated item

The preceding section describes how to create a calculated field. Excel also enables you to create a calculated item for a PivotTable field. Keep in mind that a calculated field can be an alternative to adding a new field (column) to your data source. A calculated item, on the other hand, is an alternative to adding a new row to the data source—a row that contains a formula that refers to other rows.

In this example, you create four calculated items. Each item represents the commission earned on the quarter's sales, according to the following schedule:

- Quarter 1:** 10% of January, February, and March sales
- Quarter 2:** 11% of April, May, and June sales
- Quarter 3:** 12% of July, August, and September sales
- Quarter 4:** 12.5% of October, November, and December sales

NOTE

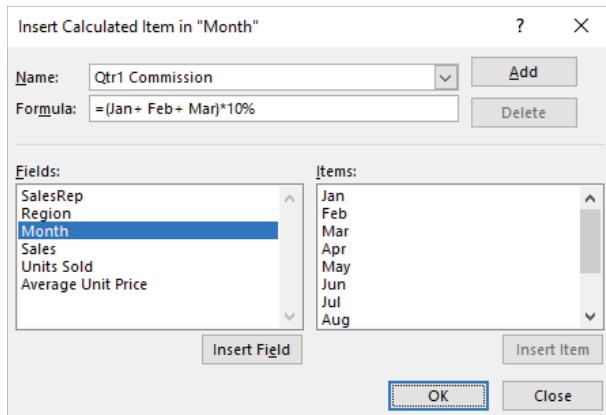
Modifying the source data to obtain this information would require inserting 16 new rows, each with formulas (four formulas for each sales rep). So, for this example, creating four calculated items may be an easier task.

To create a calculated item to compute the commission for January, February, and March, follow these steps:

1. **Select any cell in the Row Labels or Column Labels area of the PivotTable and choose PivotTable Analyze \Rightarrow Calculations \Rightarrow Fields, Items, & Sets \Rightarrow Calculated Item.** The Insert Calculated Item dialog box appears.
2. **Enter a name for the new item in the Name field and specify the formula in the Formula field (see Figure 27.17).** The formula can use items in other fields, but it can't use worksheet functions. For this example, the new item is named Qtr1 Commission, and the formula appears as follows:
 $= (Jan+Feb+Mar) *10\%$
3. **Click Add.**
4. **Repeat steps 2 and 3 to create three additional calculated items:**
Qtr2 Commission: $= (Apr+May+Jun) *11\%$
Qtr3 Commission: $= (Jul+Aug+Sep) *12\%$
Qtr4 Commission: $= (Oct+Nov+Dec) *12.5\%$
5. **Click OK to close the dialog box.**

FIGURE 27.17

The Insert Calculated Item dialog box



NOTE

A calculated item, unlike a calculated field, does not appear in the PivotTable Fields task pane. Only fields appear in the field list.

CAUTION

If you use a calculated item in your PivotTable, you may need to turn off the Grand Total display for columns to avoid double counting. In this example, the Grand Total includes the calculated items, so the commission amounts are included with the sales amounts. To turn off Grand Totals, choose Design \Rightarrow Layout \Rightarrow Grand Totals \Rightarrow Off for Rows and Columns.

After you create the calculated items, they appear in the PivotTable. Figure 27.18 shows the PivotTable after adding the four calculated items. Notice that the calculated items are added to the end of the Month items. You can rearrange the items by selecting the cell and dragging its border. Another option is to create two groups (manually): one for the sales numbers and one for the commission calculations. Figure 27.19 shows the PivotTable after creating the two groups and adding subtotals.

Part IV: Managing and Analyzing Data

FIGURE 27.18

This PivotTable uses calculated items for quarterly totals.

	H	I	J	K	L	M	N	O	P	Q	R	
	Column Labels ▾											
Row Labels	Amy		Bob		Chuck		Doug		Total Sales		Total Avg Price	
	Sales	Avg Price	Sales	Avg Price	Sales	Avg Price	Sales	Avg Price				
Jan	23,040	96	20,024	194	19,886	209	26,264	285	89,214	169		
Feb	24,131	305	23,822	89	23,494	159	29,953	35	101,400	75		
Mar	24,646	347	24,854	259	21,824	263	25,041	291	96,365	287		
Apr	22,047	311	22,838	309	22,058	230	29,338	132	96,281	208		
May	24,971	159	25,320	110	20,280	45	25,150	104	95,721	88		
Jun	24,218	263	24,733	151	23,965	32	27,371	288	100,287	90		
Jul	25,735	147	21,184	312	23,032	149	25,044	305	94,995	198		
Aug	23,638	272	23,174	203	21,273	28	29,506	286	97,591	91		
Sep	25,749	46	25,999	310	21,584	189	29,061	199	102,393	114		
Oct	24,437	257	22,639	87	19,625	236	27,113	226	93,814	168		
Nov	25,355	36	23,949	220	19,832	283	25,953	320	95,089	98		
Dec	25,899	144	23,179	50	20,583	116	28,670	145	98,331	96		
Qtr1 Commission	7,182	185	6,870	147	6,520	200	8,126	79	28,698	130		
Qtr2 Commission	7,836	223	8,018	155	7,293	51	9,004	146	32,152	110		
Qtr3 Commission	9,015	92	8,443	265	7,907	63	10,033	253	35,397	120		
Qtr4 Commission	9,461	77	8,721	84	7,505	181	10,217	205	35,904	113		
Grand Total	327,360	117	313,767	138	286,661	86	365,845	142	1,293,632	118		

FIGURE 27.19

The PivotTable, after creating two groups and adding subtotals

	H	I	J	K	L	M	N	O	P	
	Column Labels ▾									
Row Labels	Amy		Bob		Chuck		Doug			
	Sales	Avg Price	Sales	Avg Price	Sales	Avg Price	Sales	Avg Price		
Monthly Sales										
Jan	23,040	96	20,024	194	19,886	209	26,264	285		
Feb	24,131	305	23,822	89	23,494	159	29,953	35		
Mar	24,646	347	24,854	259	21,824	263	25,041	291		
Apr	22,047	311	22,838	309	22,058	230	29,338	132		
May	24,971	159	25,320	110	20,280	45	25,150	104		
Jun	24,218	263	24,733	151	23,965	32	27,371	288		
Jul	25,735	147	21,184	312	23,032	149	25,044	305		
Aug	23,638	272	23,174	203	21,273	28	29,506	286		
Sep	25,749	46	25,999	310	21,584	189	29,061	199		
Oct	24,437	257	22,639	87	19,625	236	27,113	226		
Nov	25,355	36	23,949	220	19,832	283	25,953	320		
Dec	25,899	144	23,179	50	20,583	116	28,670	145		
Monthly Sales Total	293,866	117	281,715	138	257,436	86	328,464	142		
Quarterly Commissions										
Qtr1 Commission	7,182	185	6,870	147	6,520	200	8,126	79		
Qtr2 Commission	7,836	223	8,018	155	7,293	51	9,004	146		
Qtr3 Commission	9,015	92	8,443	265	7,907	63	10,033	253		
Qtr4 Commission	9,461	77	8,721	84	7,505	181	10,217	205		
Quarterly Commissions Total	33,494	114	32,052	137	29,225	85	37,381	147		

A Reverse PivotTable

The Excel PivotTable feature creates a summary table from a list. But what if you want to perform the opposite operation? Often, you may have a two-way summary table, and it would be convenient if the data were in the form of a normalized list.

In the accompanying figure, range A1 : E13 contains a summary table with 48 data points. Notice that this summary table is similar to a PivotTable. Column G : I shows part of a 48-row table that was derived from the summary table. In other words, every value in the original summary table is converted to a row, which also contains the region name and month. This type of table is useful because it can be sorted and manipulated in other ways. And you can create a PivotTable from this transformed table.

	A	B	C	D	E	F	G	H	I	J
1		North	South	East	West		Col1	Col2	Col3	
2	Jan	132	233	314	441		Jan	North	132	
3	Feb	143	251	314	447		Jan	South	233	
4	Mar	172	252	345	450		Jan	East	314	
5	Apr	184	290	365	452		Jan	West	441	
6	May	212	299	401	453		Feb	North	143	
7	Jun	239	317	413	457		Feb	South	251	
8	Jul	249	350	427	460		Feb	East	314	
9	Aug	263	354	448	468		Feb	West	447	
10	Sep	291	373	367	472		Mar	North	172	
11	Oct	294	401	392	479		Mar	South	252	
12	Nov	302	437	495	484		Mar	East	345	
13	Dec	305	466	504	490		Mar	West	450	
14							Apr	North	184	
15	Select a cell in the summary table above, then click the Convert button to create a table with one row per data point.						Apr	South	290	
16	Replace the column headings to describe the fields.						Apr	East	365	
17	This macro can be used with any 2-way table						Apr	West	452	
18							May	North	212	
19							May	South	299	
20							May	East	401	
21							May	West	453	
22							Jun	North	239	
23							Jun	South	317	

The companion website contains a workbook, `reverse pivot.xlsxm`, which has a VBA macro that will convert any two-way summary table into a three-column normalized table.

Another way to perform this type of transformation is to use Get & Transform. See Chapter 36, “Transforming Data with Power Query,” for an example.

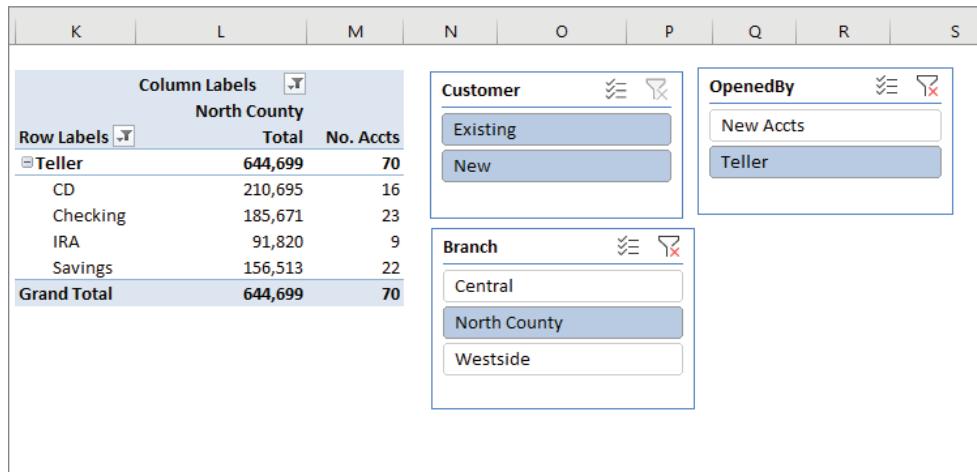
Filtering PivotTables with Slicers

A slicer is an interactive control that makes it easy to filter data in a PivotTable. Figure 27.20 shows a PivotTable with three slicers, each representing a particular field. In this case, the PivotTable is displaying data for new and existing customers whose accounts were opened by tellers at the North County branch.

Part IV: Managing and Analyzing Data

FIGURE 27.20

Using slicers to filter the data displayed in a PivotTable



The same type of filtering can be accomplished by using the field labels in the PivotTable, but slicers are intended for those who might not understand how to filter data in a PivotTable. Slicers can also be used to create an attractive and easy-to-use interactive “dashboard.”

To add one or more slicers to a worksheet, start by selecting any cell in a PivotTable. Then choose Insert \Rightarrow Filters \Rightarrow Slicer. The Insert Slicers dialog box appears, with a list of all fields in the PivotTable. Place a check mark next to the slicers you want and then click OK.

Slicers can be moved and resized, and you can change the look. To remove the effects of filtering by a particular slicer, click the Clear Filter icon in the slicer’s upper-right corner.

To use a slicer to filter data in a PivotTable, just click a button. To display multiple values, press Ctrl while you click the buttons in a slicer or click the Multi-Select icon at the top right of the slicer window. Press Shift and click to select a series of consecutive buttons.

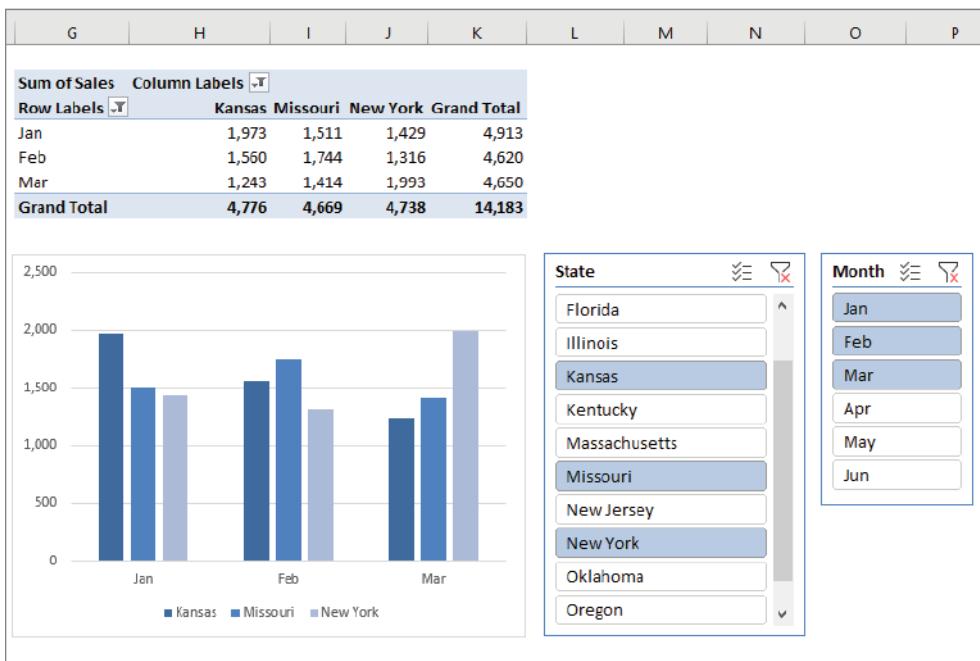
Figure 27.21 shows a PivotTable and a PivotChart. Two slicers are used to filter the data (by state and by month). In this case, the PivotTable and PivotChart show only the data for Kansas, Missouri, and New York for the months of January through March. Slicers provide a quick and easy way to create an interactive chart.



This workbook, named pivot table slicers.xlsx, is available on this book’s website at www.wiley.com/go/excel365bible.

FIGURE 27.21

Using slicers to filter a PivotTable and PivotChart by state and by month



Filtering PivotTables with a Timeline

A timeline is conceptually similar to a slicer, but this control is designed to simplify time-based filtering in a PivotTable.

A timeline is relevant only if your PivotTable has a field that's formatted as a date. This feature does not work with times. To add a timeline, select a cell in a PivotTable and choose Insert \Rightarrow Filters \Rightarrow Timeline. A dialog box appears listing all date-based fields. If your PivotTable doesn't have a field formatted as a date, Excel displays an error.

Figure 27.22 shows a PivotTable created from the data in columns A : E. This PivotTable uses a timeline set to allow date filtering by quarters. Click a button that corresponds to the quarter that you want to view, and the PivotTable is updated immediately. To select a range of quarters, drag the edges of an already-selected quarter forward or backward. Other filtering options (selectable from the drop-down in the upper-right corner) are Years, Months, and Days. In the figure, the PivotTable displays data from the second and third quarters of 2023.

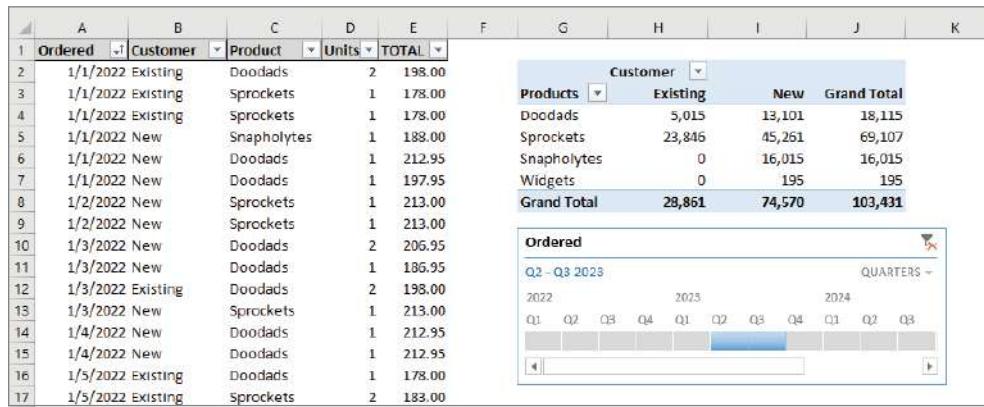


A workbook that uses a timeline is available on this book's website at www.wiley.com/go/excel365bible. The filename is pivot_table_timeline.xlsx.

Part IV: Managing and Analyzing Data

FIGURE 27.22

Using a timeline to filter a PivotTable by date



You can, of course, use both slicers and a timeline for a PivotTable. A timeline has the same type of formatting options as slicers, so you can create an attractive interactive dashboard that simplifies PivotTable filtering.

Referencing Cells within a PivotTable

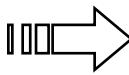
After you create a PivotTable, you may want to create formulas that reference one or more cells within it. Figure 27.23 shows a simple PivotTable that displays income and expense information for three years. In this PivotTable, the Month field is hidden, so the PivotTable shows the year totals.

FIGURE 27.23

The formulas in column F reference cells in the PivotTable.

The figure shows a PivotTable with three rows of data: 2020, 2021, and 2022. The columns are Row Labels, Sum of Income, Sum of Expenses, and Ratio. The Ratio column contains formulas: =Sum(B3)/Sum(D3) for 2020, =Sum(B4)/Sum(D4) for 2021, =Sum(B5)/Sum(D5) for 2022, and =Sum(B6)/Sum(D6) for Grand Total.

Row Labels	Sum of Income	Sum of Expenses	Ratio
2020	1,256,274	525,288	41.81%
2021	1,357,068	533,893	39.34%
2022	1,583,717	646,181	40.80%
Grand Total	4,197,059	1,705,362	40.63%



This workbook, named pivot table referencing.xlsx, is available on this book's website at www.wiley.com/go/excel365bible.

Column F contains formulas, and this column is not part of the PivotTable. These formulas calculate the expense-to-income ratio for each year. You can create these formulas by pointing to the cells in the PivotTable. You may expect to see this formula in cell F3:

=D3 /C3

In fact, the formula in cell F3 is:

```
=GETPIVOTDATA ("Sum of Expenses", $B$2, "Year", 2020)
/GETPIVOTDATA ("Sum of Income", $B$2, "Year", 2020)
```

When you use the pointing technique to create a formula that references a cell in a PivotTable, Excel replaces those simple cell references with a much more complicated GETPIVOTDATA function. If you type the cell references manually (instead of pointing to them), Excel doesn't use the GETPIVOTDATA function. Why use the GETPIVOTDATA function? Using the GETPIVOTDATA function helps to ensure that the formula will continue to reference the intended cells if the PivotTable layout is changed.

Figure 27.24 shows the PivotTable after expanding the years to show the month detail. As you can see, the formulas in column F still show the correct result even though the referenced cells are in a different location. Had we used simple cell references, the formulas would return incorrect results after expanding the Year field.

FIGURE 27.24

After expanding the Year field, formulas that use the GETPIVOTDATA function continue to display the correct result.

A	B	C	D	E	F	G
Row Labels	Sum of Income	Sum of Expenses	Ratio			
2020	1,256,274	525,288	41.81%			
Jan	98,085	42,874	39.34%			
Feb	98,698	44,167	40.80%			
Mar	102,403	43,349	40.63%			
Apr	106,044	43,102				
May	105,361	45,005				
Jun	105,729	44,216				
Jul	105,557	43,835				
Aug	109,669	41,952				
Sep	107,233	44,071				
Oct	105,048	43,185				
Nov	107,446	44,403				
Dec	105,001	45,129				
2021	1,357,068	533,893				
Jan	109,699	46,245				
Feb	109,146	45,672				

CAUTION

Using the GETPIVOTDATA function has one potential problem: The data that it retrieves must be visible. If you modify the PivotTable so that the value used by GETPIVOTDATA is no longer visible, the formula returns an error.

TIP

If you want to prevent Excel from using the GETPIVOTDATA function when you point to PivotTable cells when creating a formula, choose PivotTable Analyze \Rightarrow PivotTable \Rightarrow Options \Rightarrow Generate GetPivotData. This command is a toggle.

Creating PivotCharts

A PivotChart is a graphical representation of a data summary displayed in a PivotTable. If you're familiar with creating charts in Excel, you'll have no problem creating and customizing PivotCharts. All Excel charting features are available in a PivotChart.



See Chapter 18, “Getting Started with Excel Charts,” and Chapter 19, “Using Advanced Charting Techniques,” to learn more about charting in Excel.

Excel provides several ways to create a PivotChart:

- Select any cell in an existing PivotTable and then choose PivotTable Analyze \Rightarrow Tools \Rightarrow PivotChart.
- Select any cell in an existing PivotTable and then choose Insert \Rightarrow Charts \Rightarrow PivotChart.
- Choose Insert \Rightarrow Charts \Rightarrow PivotChart \Rightarrow PivotChart. If the active cell is not within a PivotTable, Excel prompts you for the data source and creates a PivotChart.
- Choose Insert \Rightarrow Charts \Rightarrow PivotChart \Rightarrow PivotChart & PivotTable. Excel prompts you for the data source and creates a PivotTable and a PivotChart. This command is available only when the active cell is not within a PivotTable.

A PivotChart example

Figure 27.25 shows part of a table that tracks daily sales by region. The Date field contains dates for the entire year (excluding weekends), the Region field contains the region name (Eastern, Southern, or Western), and the Sales field contains the sales amount.



This workbook, named sales by region pivot chart.xlsx, is available on this book's website at www.wiley.com/go/excel365bible.

Figure 27.26 shows a PivotTable created from this data. The Date field is in the Rows area, and the daily dates have been grouped into months. The Region field is in the Columns area. The Sales field is in the Values area.

FIGURE 27.25

This data will be used to create a PivotChart.

	A	B	C
1	Date	Region	Sales
2	1/1/2022	Eastern	10,909
3	1/1/2022	Southern	8,079
4	1/1/2022	Western	4,650
5	1/2/2022	Eastern	11,126
6	1/2/2022	Southern	8,131
7	1/2/2022	Western	4,521
8	1/5/2022	Eastern	11,224
9	1/5/2022	Southern	8,161
10	1/5/2022	Western	4,274
11	1/6/2022	Eastern	11,299
12	1/6/2022	Southern	8,071
13	1/6/2022	Western	4,365
14	1/7/2022	Eastern	11,265
15	1/7/2022	Southern	8,082

FIGURE 27.26

This PivotTable summarizes sales by region and by month.

	A	B	C	D
1				
2				
3	Sum of Sales	Column Label		
4	Row Labels	Eastern	Southern	Western
5	Jan	259,416	170,991	100,708
6	Feb	255,487	134,812	99,740
7	Mar	284,294	140,918	103,848
8	Apr	273,772	144,346	124,772
9	May	268,057	119,220	131,716
10	Jun	282,089	122,156	133,071
11	Jul	289,019	120,989	162,555
12	Aug	258,844	97,790	162,644
13	Sep	264,537	106,315	193,012
14	Oct	273,592	104,383	213,761
15	Nov	279,259	93,118	223,203
16	Dec	309,802	104,962	260,102
17	Grand Total	3,298,168	1,460,000	1,909,132
18				

The PivotTable is certainly easier to interpret than the raw data, but the trends would be easier to spot in a chart.

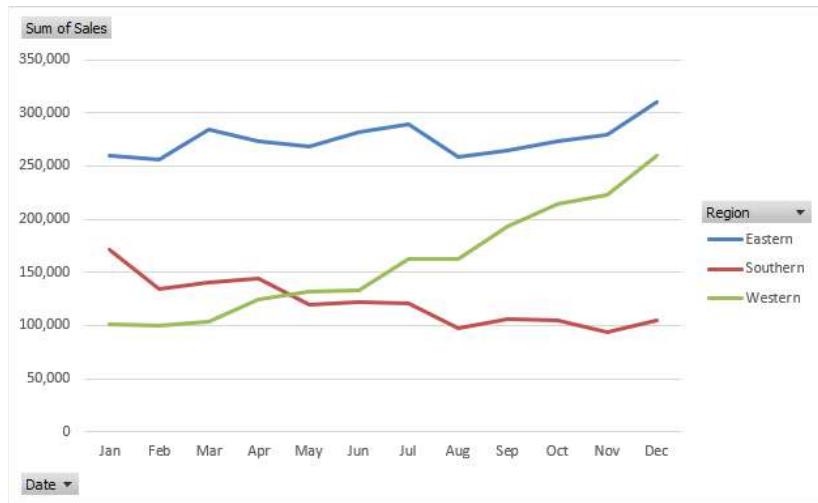
To create a PivotChart, select any cell in the PivotTable and choose PivotTable Analyze \Rightarrow Tools \Rightarrow PivotChart. The Insert Chart dialog box appears, from which you can choose a chart type. For this example, select a Line chart and then click OK. Excel creates the PivotChart shown in

Part IV: Managing and Analyzing Data

Figure 27.27. The chart makes it easy to see an upward sales trend for the Western division, a downward trend for the Southern division, and relatively flat sales for the Eastern division.

FIGURE 27.27

The PivotChart uses the data displayed in the PivotTable.



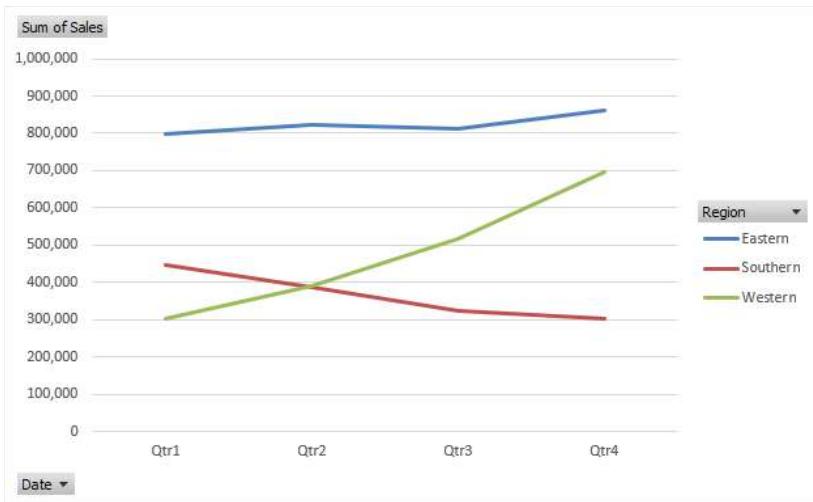
A PivotChart includes field buttons that let you filter the chart's data. To remove some or all the field buttons, select the PivotChart and use the Field Buttons control in the PivotChart Analyze \leftrightarrow Show/Hide group.

When you select a PivotChart, the Ribbon displays a new contextual tab: Format. The commands in the Format tab are virtually identical to those for a standard Excel chart, so you can manipulate the PivotChart any way you like. The Design tab contains different commands, more suited to charts than the Design tab you get when you click inside a PivotTable. The PivotChart Analyze tab contains selected commands from the PivotTable Analyze tab, such as Refresh and Insert Slicer.

If you modify the underlying PivotTable, the chart adjusts automatically to display the new summary data. Figure 27.28 shows the PivotChart after the Date grouping is changed to quarters.

FIGURE 27.28

If you modify the PivotTable, the PivotChart is also changed.



More about PivotCharts

Keep in mind these points when using PivotCharts:

- **A PivotTable and a PivotChart are joined in a two-way link.** If you make structural or filtering changes to one, the other is also changed.
- **When you activate a PivotChart, the PivotTable Fields task pane changes to the PivotChart Fields task pane.** In this task pane, Legend (Series) replaces the Columns area, and Axis (Categories) replaces the Rows area.
- **The field buttons in a PivotChart contain the same controls as the PivotTable's field headers.** These controls allow you to filter the data that's displayed in the PivotTable and PivotChart. If you make changes to the PivotChart using these buttons, those changes are also reflected in the PivotTable.
- **If you have a PivotChart linked to a PivotTable and you delete the underlying PivotTable, the PivotChart remains.** The PivotChart's Series formulas contain the original data, stored in arrays.
- **By default, PivotCharts are embedded in the sheet that contains the PivotTable.** To move the PivotChart to a different worksheet (or to a Chart sheet), choose PivotChart Analyze ➔ Actions ➔ Move Chart.

- You can create multiple PivotCharts from a PivotTable, and you can manipulate and format the charts separately. However, all the charts display the same data.
- A normal chart, when selected, displays the icons to the right: Chart Elements, Chart Styles, and Chart Filters. A PivotChart does not display the Chart Filters icon.
- Slicers and timelines also work with PivotCharts. See the examples earlier in this chapter.
- Don't forget about themes. You can choose Page Layout \Rightarrow Themes \Rightarrow Themes to change the workbook theme, and your PivotTable and PivotChart will both reflect the new theme.

Using the Data Model

So far, this chapter has focused exclusively on PivotTables that are created from a single table of data. A feature called the *data model* brings significantly more power to PivotTables. With the data model, you can use multiple tables of data in a single PivotTable. You'll need to create one or more “table relationships” so that the data can be tied together.

A workbook can have only one data model. Any data model set up in a particular workbook applies to all PivotTables that use the data model. You can't have one model for one PivotTable and another model for another PivotTable within the same workbook.

NOTE

The data model was introduced in Excel 2013, so workbooks that use this feature are not compatible with previous versions.



Part V, “Understanding Power Pivot and Power Query,” contains more information on the data model.

Figure 27.29 shows parts of three tables that are in a single workbook. (Each table is in its own worksheet and is shown in a separate window.) The tables are named Orders, Customers, and Regions. The Orders table contains information about product orders. The Customers table contains information about the company's customers. The Regions table contains a region identifier for each state.

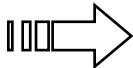
Notice that the Orders and Customers tables have a CustomerID column in common, and the Customers and Regions tables have a State column in common. The common columns will be used to form relationships among the tables.

These relationships are “one-to-many.” For every row in the Orders table, there is exactly one corresponding row in the Customers table, and that row is determined by the CustomerID column. For any row in the Customers table, there can be many rows in the Orders table. The Orders table is the many, and the Customers table is the one in the one-to-many relationship. Similarly, for every row in the Customers table, there is exactly one corresponding row in the Regions table, and that row is determined by the State column. And for any row in Regions there can be many matching rows in Customers.

FIGURE 27.29

These three tables will be used for a PivotTable, using the data model.

The screenshot shows three separate Excel workbooks stacked vertically. The top workbook is titled 'data model complete.xlsx' and contains the 'Orders' table. The middle workbook is titled 'data model.com...' and contains the 'Customers' table. The bottom workbook is also titled 'data model.com...' and contains the 'Regions' table. All three workbooks have their 'File' tabs selected. The 'Regions' table includes a column for state abbreviations with dropdown arrows, indicating they are part of a data model.



The example in this section is available on this book's website at www.wiley.com/go/excel365bible. You can use the workbook named `data model.xlsx` to follow along with the example. The workbook named `data model complete.xlsx` shows the final PivotTable.

NOTE

A PivotTable created using the data model has some restrictions, as opposed to a PivotTable created from a single table. Most notably, you can't group items or create calculated fields or calculated items.

For this example, the goal is to summarize sales by state, by region, and by year. Notice that the sales and date information is in the Orders table, the state information is in the Customers table, and the region names are in the Regions table. Therefore, all three tables will be used to generate this PivotTable.

Part IV: Managing and Analyzing Data

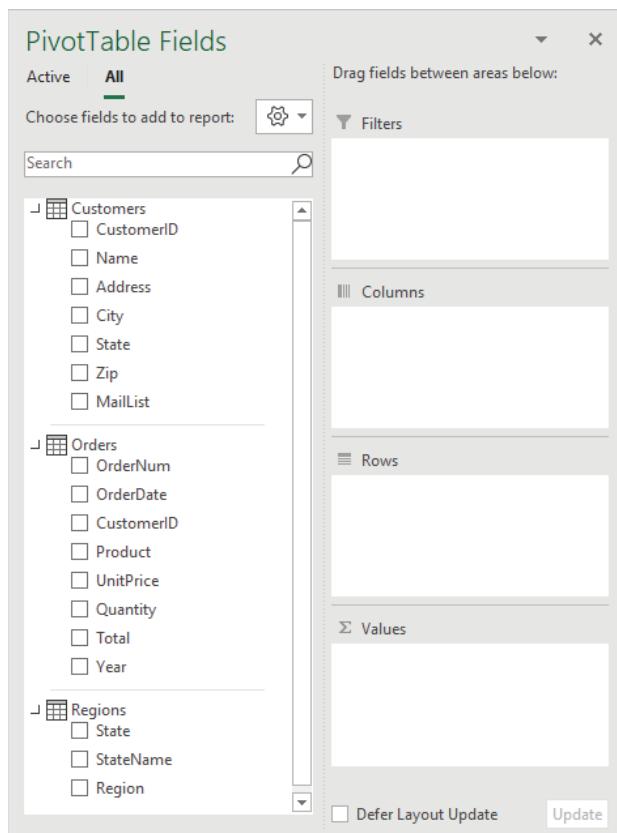
Start by creating a PivotTable (in a new worksheet) from the Orders table. Follow these steps:

1. **Select any cell within the Orders table and choose Insert ➔ Tables ➔ PivotTable.**
The Create PivotTable dialog box appears.
2. **Select the Add This Data to the Data Model check box and click OK.** The PivotTable Fields task pane looks a bit different when you're working with the data model. The task pane contains two tabs: Active and All. The Active tab lists only the Orders table. The All tab lists all of the tables in the workbook.

Figure 27.30 shows the All tab of the PivotTable Fields task pane, with all three tables expanded to show their column headers. To change the layout of the task pane, click the drop-down Tools control and choose Fields Section and Areas Section Side-by-Side.

FIGURE 27.30

The PivotTable Fields task pane for a data model

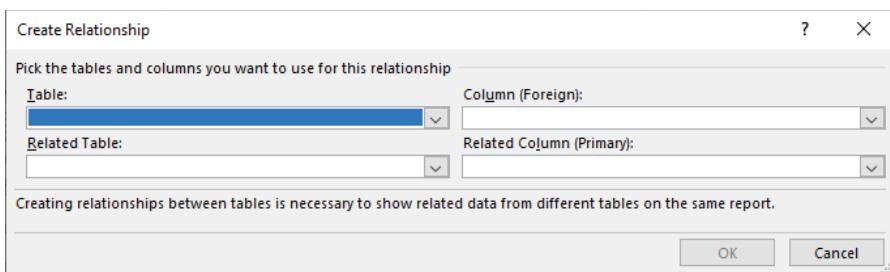


You made the Orders table part of the data model when you checked the box creating the PivotTable. The other two tables are not currently part of the data model.

3. Choose **Data** ➔ **Data Tools** ➔ **Relationships** and click the **New** button. Figure 27.31 shows the Create Relationship dialog box that appears when you click the New button. It's empty because we haven't set up our relationships yet.

FIGURE 27.31

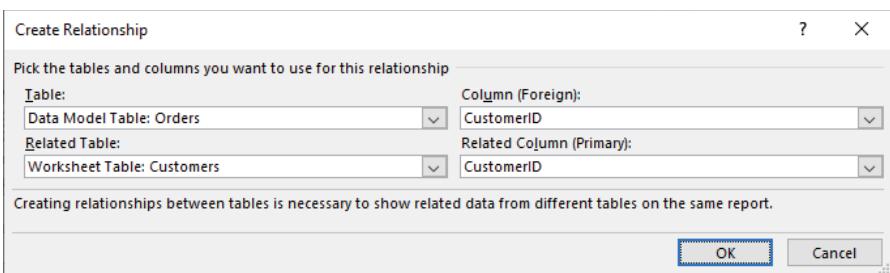
The Create Relationship dialog box



4. **Relate the Orders table to the Customers table.** Select Data Model Table: Orders and CustomerID in the top section and Worksheet Table: Customers and CustomerID in the bottom section. Compare your Create Relationship dialog box to Figure 27.32 and click OK.

FIGURE 27.32

Relating the Orders table to the Customers table

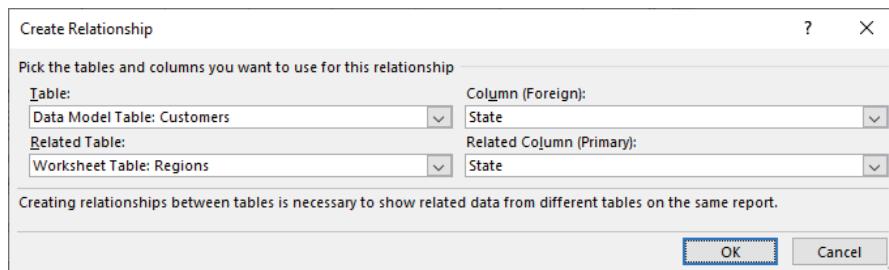


5. **Relate the Customers table to the Regions table.** Click the **New** button in the **Manage Relationships** dialog box and create a relationship similar to how you did it in the previous step (see Figure 27.33). Notice how the Customers table is now identified as being part of the data model. When you created the relationship in the previous step, you added the Customers table to the data model.

Part IV: Managing and Analyzing Data

FIGURE 27.33

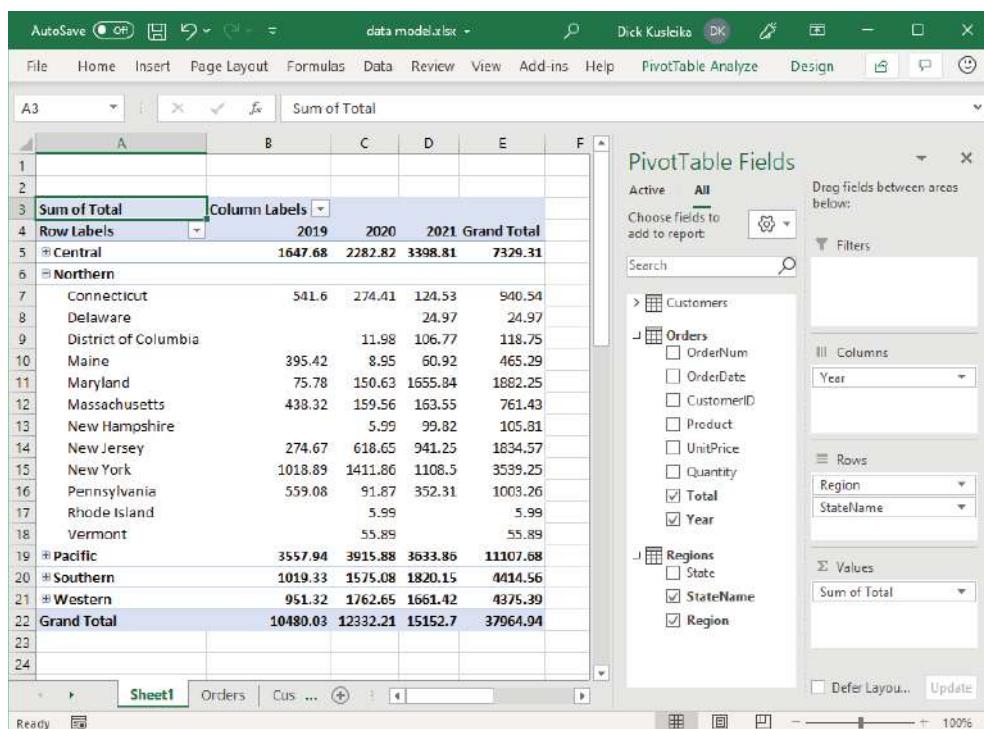
Relating the Customers and Regions tables by state



6. **Build the PivotTable.** Close the Manage Relationships dialog box. Move the Region and StateName fields to the Rows area, move the Year field to the Columns area, and move the Total field to the Values area. Figure 27.34 shows the PivotTable after this step. (Some of the data in Figure 27.34 was collapsed so that more of the PivotTable can be seen.)

FIGURE 27.34

A PivotTable based on the data model



The PivotTable based on three tables is complete. All that remains is some formatting. You can format your PivotTable to suit your taste. To make your PivotTable look like the one in `data model complete.xlsx`, follow these steps:

- 1. Add the slicers.** Choose PivotTable Analyze \Rightarrow Filter \Rightarrow Insert Slicer, select the All tab, and select Product and MailList.
- 2. Format the slicers.** Right-click the MailList slicer and choose Size and Properties. In the Position and Layout section of the Format Slicer pane, change the Number of Columns property to 2.
- 3. Format the values.** Right-click any value in the PivotTable and choose Number Format. Choose Number, set Decimal Places to 2, and select the Use 1000 Separator check box.
- 4. Add region subtotals.** Right-click any Region item (such as Central) and choose Field Settings. On the Subtotals and Filters tab, choose Automatic.
- 5. Add a blank line between regions.** On the Layout & Print tab of the Field Settings dialog box, check the Insert Blank Line after Each Item Label check box.

Figure 27.35 shows the final, formatted PivotTable.

FIGURE 27.35

Formatting applied to the data model PivotTable

A	B	C	D	E	F	G	H	I
2	Sum of Total	Column Labels	2019	2020	2021	Grand Total		
3	Row Labels							
4	Central							
6	Illinois	176.62	459.89	860.49	1,497.00			
7	Indiana	428.22	71.67	107.68	607.57			
8	Iowa		11.98	247.60	259.58			
9	Kentucky	55.89			55.89			
10	Michigan	291.23	436.20	853.32	1,580.75			
11	Minnesota	175.60	25.98	155.67	357.25			
12	Missouri	71.86	491.64	284.39	847.89			
13	Ohio	258.45	648.74	802.04	1,709.23			
14	Wisconsin	189.81	136.72	87.62	414.15			
15	Central Total	1,647.68	2,282.82	3,398.81	7,329.31			
16								
17	Northern							
18	Connecticut	541.60	274.41	124.53	940.54			
19	Delaware			24.97	24.97			
20	District of Columbia		11.98	106.77	118.75			
21	Maine	395.42	8.95	60.92	465.29			
22	Maryland	75.78	150.63	1,655.84	1,882.25			
23	Massachusetts	438.32	159.56	163.55	761.43			
24	New Hampshire		5.99	99.82	105.81			
25	New Jersey	274.67	618.65	941.25	1,834.57			
26	New York	1,018.89	1,411.86	1,108.50	3,539.25			
27	Pennsylvania	559.08	91.87	352.31	1,003.26			
28	Rhode Island		5.99		5.99			
29	Vermont		55.89		55.89			
30	Northern Total	3,303.76	2,795.78	4,638.46	10,738.00			

Part IV: Managing and Analyzing Data

TIP

When you create a PivotChart using the data model, you can convert the PivotTable to formulas. Converting a PivotTable to formulas is useful if you need the same summary information but you need it to update automatically rather than having to refresh a PivotTable. Select any cell in the PivotTable and choose PivotTable Analyze \Rightarrow Calculations \Rightarrow OLAP Tools \Rightarrow Convert to Formulas. The PivotTable is replaced by cells that use formulas. These formulas use CUBEMEMBER and CUBEVALUE functions. Although the range is no longer a PivotTable, the formulas update when the data changes.

Performing Spreadsheet What-If Analysis

IN THIS CHAPTER

- Considering a what-if example
- Identifying types of what-if analyses
- Looking at manual what-if analyses
- Creating one-input and two-input data tables
- Using Scenario Manager
- Analyzing data with artificial intelligence

One of the most appealing aspects of Excel is its ability to create dynamic models. A dynamic model uses formulas that instantly recalculate when you change values in cells that are used by the formulas. When you change values in cells in a systematic manner and observe the effects on specific formula cells, you're performing a type of what-if analysis.

What-if analysis is the process of asking such questions as "What happens if the interest rate on the loan changes to 7.5 percent?" and "What occurs if we raise our product prices by 5 percent?"

If you set up your worksheet properly, answering such questions is simply a matter of plugging in new values and observing the results of the recalculation. And if you're not sure what questions to ask, Excel can analyze the data and suggest ways to summarize it. Excel provides useful tools to assist you in your what-if endeavors.

Looking at a What-If Example

Figure 28.1 shows a simple worksheet model that calculates information pertaining to a mortgage loan. The worksheet is divided into two sections: the input cells and the result cells (which contain formulas).

Part IV: Managing and Analyzing Data

FIGURE 28.1

This simple worksheet model uses four input cells to produce the results.

	A	B	C	D	
Mortgage Loan Worksheet					
1					
2					
3		Input Cells			
4	Purchase Price:		\$385,500		
5	Down Payment:		10%		
6	Loan Term (Months):		360		
7	Interest Rate (APR):		3.10%		
8					
9		Result Cells			
10	Loan Amount:		\$346,950		
11	Monthly Payment:		\$1,482		
12	Total Payments:		\$533,352		
13	Total Interest:		\$186,402		
14					



This workbook is available on this book's website at www.wiley.com/go/excel365bible. The filename is mortgage loan.xlsx.

With this worksheet, you can easily answer the following what-if questions:

- What if I can negotiate a lower purchase price on the property?
- What if the lender requires a 20 percent down payment?
- What if I can get a 40-year mortgage?
- What if the interest rate increases to 3.50 percent?

You can answer these questions by simply changing the values in the cells in range C4 : C7 and observing the effects in the dependent cells (C10 : C13). You can, of course, vary any number of input cells simultaneously.

Avoid Hard-Coding Values in a Formula

The mortgage calculation example, simple as it is, demonstrates an important point about spreadsheet design: You should always set up your worksheet so that you have maximum flexibility to make changes. Perhaps the most fundamental rule of spreadsheet design is the following:

Do not hard-code values in a formula. Instead, store the values in separate cells and use cell references in the formula.

The term *hard-code* refers to using values, or constants, rather than cell references in a formula's arguments. In the mortgage loan example, the formulas only use cell references for their arguments, not actual input values.

You could use the value 360, for example, for the loan term argument of the PMT function in cell C11 of Figure 28.1 rather than referencing the value in C6. But using a cell reference has two advantages. First, you'll have no doubt about the values that the formula uses. That is, you don't have to read the formula to see the values. They're all out in the open and properly labeled. Second, you can easily change the value—typing a new value in the cell is easier than editing the formula.

Using values in formulas may not seem like much of an issue when only one formula is involved, but just imagine what would happen if this value were hard-coded into several hundred formulas that were scattered throughout a worksheet.

Exploring Types of What-If Analyses

Not surprisingly, Excel can handle much more sophisticated models than the preceding example. To perform a what-if analysis using Excel, you have three basic options:

Manual what-if analysis Plug in new values and observe the effects on formula cells.

Data tables Create a special type of table that displays the results of selected formula cells as you systematically change one or two input cells.

Scenario Manager Create named scenarios and generate reports that use outlines or PivotTables.

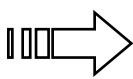
We discuss each of these types of what-if analysis in the rest of this chapter.

28

Performing manual what-if analysis

A manual what-if analysis doesn't require too much explanation. In fact, the example that opens this chapter demonstrates how it's done. Manual what-if analysis is based on the idea that you have one or more input cells that affect one or more formula cells. You change the value in the input cells and observe the formula calculations. You may want to print the results or save each scenario to a new workbook. The term *scenario* refers to a specific combination of values in the input cells.

Manual what-if analysis is common. People often use this technique without even realizing that they're doing a what-if analysis. This method of performing what-if analysis certainly has nothing wrong with it, but you should be aware of some other techniques.



If your input cells are not located near the formula cells, consider using a Watch Window to monitor the formula results in a movable window. We discuss this feature in Chapter 3, "Performing Basic Worksheet Operations."

Creating data tables

Another type of what-if analysis, a *data table*, is a dynamic range that summarizes formula cells for varying input cells. You can create a data table without much effort, but they have

Part IV: Managing and Analyzing Data

some limitations. In particular, a data table can deal with only one or two input cells at a time. This limitation becomes clear as you view the examples.

NOTE

Scenario Manager, discussed later in this chapter (see “Using Scenario Manager”), can produce a report that summarizes any number of input cells and result cells.

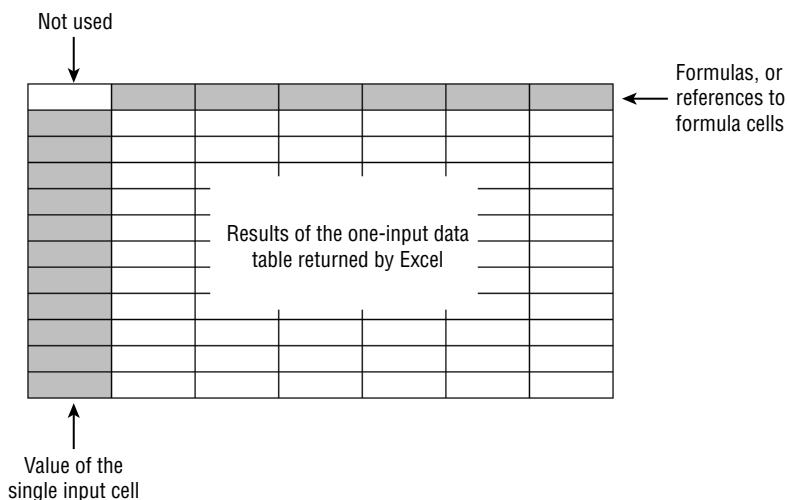
Don’t confuse a data table with a standard table (created by choosing Insert \Rightarrow Tables \Rightarrow Table). These two features are completely independent.

Creating a one-input data table

A one-input data table displays the results of one or more formulas for various values of a single input cell. Figure 28.2 shows the general layout for a one-input data table. You need to set up the table manually. This is not something that Excel will do for you.

FIGURE 28.2

How a one-input data table is set up



You can place the data table anywhere in a worksheet. The left column contains various values for the single input cell. The top row contains references to formulas located elsewhere in the worksheet. You can use a single formula reference or any number of formula references. The upper-left cell of the table remains empty. Excel calculates the values that result from each value of the input cell and places them under each formula reference.

This example uses the mortgage loan worksheet from earlier in the chapter (see “Looking at a What-If Example”). The goal of this exercise is to create a data table that shows the values of

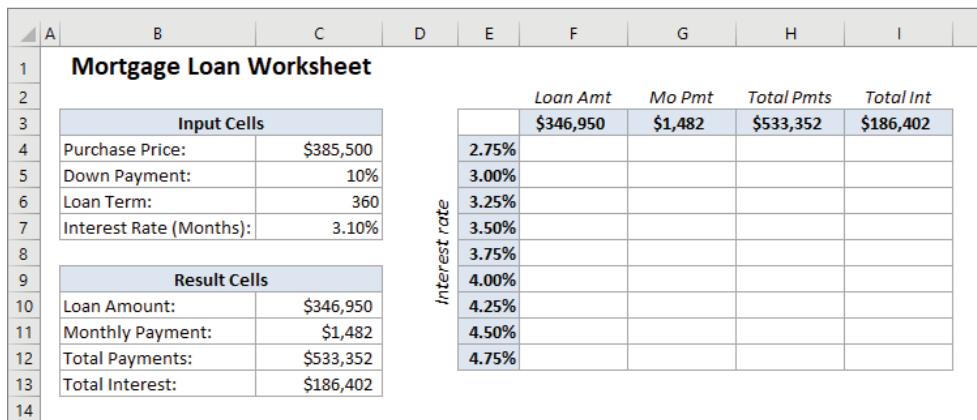
the four formula cells (loan amount, monthly payment, total payments, and total interest) for various interest rates ranging from 2.75% to 4.75%, in 0.25% increments.

 This workbook is available on this book's website at www.wiley.com/go/excel365bible. The file is named mortgage loan data table.xlsx.

Figure 28.3 shows the setup for the data table area. Row 3 consists of references to the formulas in the worksheet. For example, cell F3 contains the formula =C10, and cell G3 contains the formula =C11. Row 2 and column D contain optional descriptive labels that are not actually part of the data table. Column E contains the values of the single input cell (interest rate) that Excel will use in the table.

FIGURE 28.3

Preparing to create a one-input data table



	A	B	C	D	E	F	G	H	I
1		Mortgage Loan Worksheet							
2									
3		Input Cells							
4	Purchase Price:	\$385,500				Loan Amt	Mo Pmt	Total Pmts	Total Int
5	Down Payment:	10%				\$346,950	\$1,482	\$533,352	\$186,402
6	Loan Term:	360				2.75%			
7	Interest Rate (Months):	3.10%				3.00%			
8						3.25%			
9						3.50%			
10						3.75%			
11						4.00%			
12						4.25%			
13						4.50%			
14						4.75%			

	Interest rate	Loan Amt	Mo Pmt	Total Pmts	Total Int
	2.75%				
	3.00%				
	3.25%				
	3.50%				
	3.75%				
	4.00%				
	4.25%				
	4.50%				
	4.75%				

To create the table, select the entire data table range (in this case, E3 : I12) and then choose Data \Rightarrow Forecast \Rightarrow What-If Analysis \Rightarrow Data Table. The Data Table dialog box, shown in Figure 28.4, appears.

You must specify the worksheet cell that contains the input value. Because variables for the input cell appear in the left column in the data table, you place this cell reference in the Column Input Cell field. Enter C7 or click the cell in the worksheet. Leave the Row Input Cell field blank. Click OK, and Excel fills in the table with the calculated results (see Figure 28.5).

Part IV: Managing and Analyzing Data

FIGURE 28.4

The Data Table dialog box

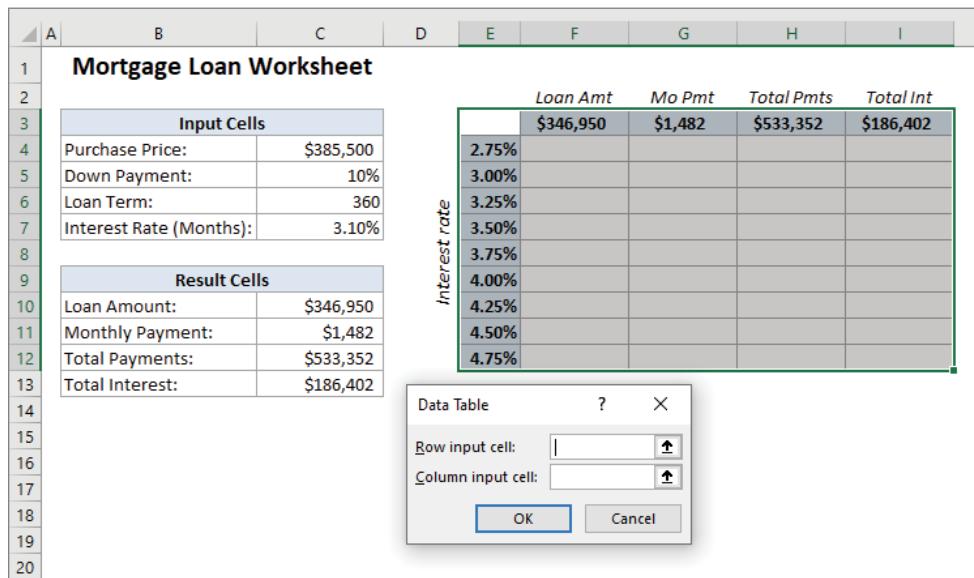
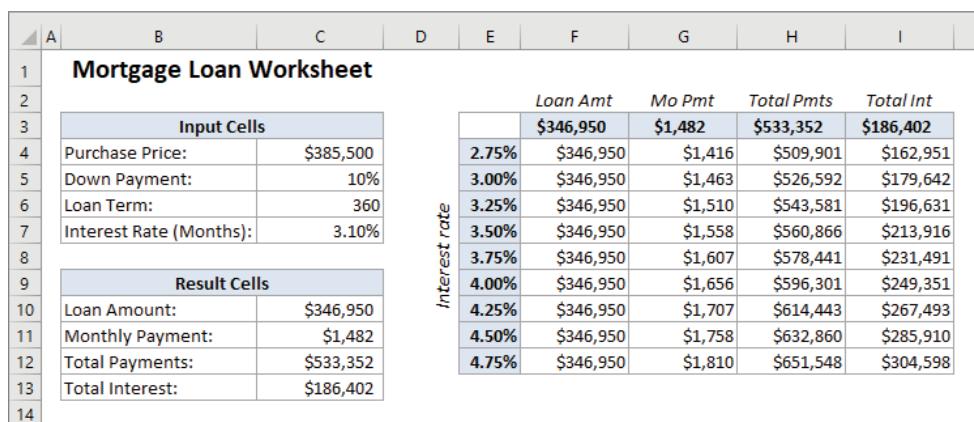


FIGURE 28.5

The result of the one-input data table



CAUTION

Using the Data Table command wipes out Excel's undo stack. Operations that you perform prior to using this command cannot be undone.

Using this table, you can now see the calculated loan values for varying interest rates. Notice that the Loan Amt column (column F) doesn't vary. That's because the formula in cell C10 doesn't depend on the interest rate.

If you examine the contents of the cells that Excel entered as a result of this command, you'll see that the data is generated with a multicell array formula:

```
{=TABLE(,C7)}
```

A multicell array formula is a single formula that can produce results in multiple cells (see Chapter 10, "Understanding and Using Array Formulas"). Because the table uses formulas, Excel updates the table that you produce if you change the cell references in the first row or plug in different interest rates in the first column.

NOTE

You can arrange a one-input table vertically (as in this example) or horizontally. If you place the values of the input cell in a row, you enter the input cell reference in the Row Input Cell field of the Data Table dialog box.

Creating a two-input data table

As the name implies, a two-input data table lets you vary two input cells. You can see the setup for this type of table in Figure 28.6. Although it looks similar to a one-input table, the two-input table has one critical difference: it can show the results of only one formula at a time. With a one-input table, you can place any number of formulas, or references to formulas, across the top row of the table. In a two-input table, this top row holds the values for the second input cell. The upper-left cell of the table contains a reference to the single result formula.

Using the mortgage loan worksheet, you could create a two-input data table that shows the results of a formula (say, monthly payment) for various combinations of two input cells (such as interest rate and down-payment percent). To see the effects on other formulas, you simply create multiple data tables—one for each formula cell that you want to summarize.

The example in this section uses the worksheet shown in Figure 28.7 to demonstrate a two-input data table. In this example, a company wants to conduct a direct-mail promotion to sell its product. The worksheet calculates the net profit from the promotion.

Part IV: Managing and Analyzing Data

FIGURE 28.6

The setup for a two-input data table

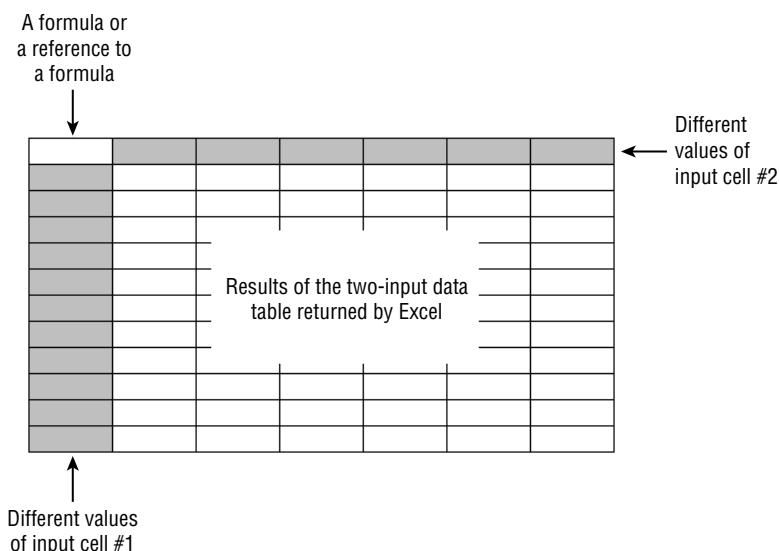


FIGURE 28.7

This worksheet calculates the net profit from a direct-mail promotion.

	A	B	C
Direct Mail Profit Model			
1			
2			
Input Cells			
3			
4	Number mailed:	275,000	
5	Response rate:	2.50%	
6			
Parameters			
7			
8	Printing costs per unit:	\$0.15	
9	Mailing costs per unit:	\$0.28	
10	Responses:	6,875	
11	Profit per response:	\$18.50	
12	Gross profit:	\$127,188	
13	Printing + mailing costs:	\$118,250	
14	Net Profit	\$8,937	
15			



This workbook, named direct mail data table.xlsx, is on this book's website at www.wiley.com/go/excel365bible.

This model uses two input cells: the number of promotional pieces mailed and the anticipated response rate. The following items appear in the Parameters area:

Printing costs per unit The cost to print a single mailer. The unit cost varies with the quantity: \$0.20 each for quantities less than 200,000; \$0.15 each for quantities of 200,001 through 299,000; and \$0.10 each for quantities of 300,000 or more. The following formula is used:

$$=IF(C4<200000, 0.2, IF(C4<300000, 0.15, 0.1))$$

Mailing costs per unit A fixed cost, \$0.28 per unit mailed.

Responses The number of responses, calculated from the response rate and the number mailed. The formula in this cell is as follows:

$$=C4 * C5$$

Profit per response A fixed value. The company knows that it will realize an average profit of \$18.50 per order.

Gross profit This is a simple formula that multiplies the profit-per-response by the number of responses:

$$=C10 * C11$$

Printing + mailing costs This formula calculates the total cost of the promotion:

$$=C4 * (C8+C9)$$

Net Profit This formula calculates the bottom line—the gross profit minus the printing and mailing costs:

$$=C12 - C13$$

If you enter values for the two input cells, you see that the net profit varies quite a bit, often going negative to produce a net loss.

Figure 28.8 shows the setup of a two-input data table that summarizes the net profit at various combinations of quantity (number mailed) and response rate; the table appears in the range B17:J27. Cell B17 contains a formula that references the Net Profit cell:

$$=C14$$

To create the data table, follow these steps:

1. Enter the response rate values in C17:J17.
2. Enter the number mailed values in B18:B27.
3. Enter the formula =C14 in cell B17.

Part IV: Managing and Analyzing Data

FIGURE 28.8

Preparing to create a two-input data table

	A	B	C	D	E	F	G	H	I	J	K
Direct Mail Profit Model											
Input Cells											
Number mailed: 275,000											
Response rate: 2.50%											
Parameters											
Printing costs per unit: \$0.15											
Mailing costs per unit: \$0.28											
Responses: 6,875											
Profit per response: \$18.50											
Gross profit: \$127,188											
Printing + mailing costs: \$118,250											
Net Profit \$8,937											
<i>Response Rate</i>											
Number Mailed	\$8,937	1.50%	1.75%	2.00%	2.25%	2.50%	2.75%	3.00%	3.25%		
	100,000										
	125,000										
	150,000										
	175,000										
	200,000										
	225,000										
	250,000										
	275,000										
	300,000										
	325,000										

4. Select the range B17 : J27 and choose **Data** \Rightarrow **Forecast** \Rightarrow **What-If Analysis** \Rightarrow **Data Table**. The Data Table dialog box appears.
5. Specify C5 as the Row input cell (the response rate) and C4 as the Column input cell (the number mailed).
6. Click **OK**. Excel fills in the data table.

Figure 28.9 shows the result. As you can see, quite a few of the combinations of response rate and quantity mailed result in a loss rather than a profit.

As with the one-input data table, this data table is dynamic. You can change the formula in cell B17 to refer to another cell (such as gross profit), or you can enter some different values for Response Rate and Number Mailed.

FIGURE 28.9

The result of the two-input data table

A	B	C	D	E	F	G	H	I	J	K
Direct Mail Profit Model										
Input Cells										
Number mailed: 275,000										
Response rate: 2.50%										
Parameters										
Printing costs per unit: \$0.15										
Mailing costs per unit: \$0.28										
Responses: 6,875										
Profit per response: \$18.50										
Gross profit: \$127,188										
Printing + mailing costs: \$118,250										
Net Profit \$8,937										
Response Rate										
\$8,937 1.50% 1.75% 2.00% 2.25% 2.50% 2.75% 3.00% 3.25%										
Number Mailed	100,000	(\$20,250)	(\$15,625)	(\$11,000)	(\$6,375)	(\$1,750)	\$2,875	\$7,500	\$12,125	
	125,000	(\$25,313)	(\$19,531)	(\$13,750)	(\$7,969)	(\$2,188)	\$3,594	\$9,375	\$15,156	
	150,000	(\$30,375)	(\$23,438)	(\$16,500)	(\$9,563)	(\$2,625)	\$4,312	\$11,250	\$18,188	
	175,000	(\$35,438)	(\$27,344)	(\$19,250)	(\$11,156)	(\$3,063)	\$5,031	\$13,125	\$21,219	
	200,000	(\$30,500)	(\$21,250)	(\$12,000)	(\$2,750)	\$6,500	\$15,750	\$25,000	\$34,250	
	225,000	(\$34,313)	(\$23,906)	(\$13,500)	(\$3,094)	\$7,312	\$17,719	\$28,125	\$38,531	
	250,000	(\$38,125)	(\$26,563)	(\$15,000)	(\$3,438)	\$8,125	\$19,688	\$31,250	\$42,813	
	275,000	(\$41,938)	(\$29,219)	(\$16,500)	(\$3,781)	\$8,937	\$21,656	\$34,375	\$47,094	
	300,000	(\$30,750)	(\$16,875)	(\$3,000)	\$10,875	\$24,750	\$38,625	\$52,500	\$66,375	
	325,000	(\$33,131)	(\$18,281)	(\$3,250)	\$11,781	\$26,813	\$41,844	\$56,875	\$71,906	

Using Scenario Manager

Data tables are useful, but they have a few limitations:

- You can vary only one or two input cells at a time.
- Setting up a data table is not intuitive.
- A two-input table shows the results of only one formula cell (although you can create additional tables for more formulas).
- In many situations, you're interested in a few select combinations, not an entire table that shows all possible combinations of two input cells.

The Scenario Manager is an easy way to automate some aspects of your what-if models. You can store different sets of input values (called *changing cells* in the terminology of Scenario Manager) for any number of variables and give a name to each set. You can then select a set of values by name, and Excel displays the worksheet by using those values. You can also generate a summary report that shows the effect of various combinations of values on any number of result cells. These summary reports can be an outline or a PivotTable.

Part IV: Managing and Analyzing Data

For example, your annual sales forecast may depend on several factors. Consequently, you can define three scenarios: best case, worst case, and most likely case. You then can switch to any of these scenarios by selecting the named scenario from a list. Excel substitutes the appropriate input values in your worksheet and recalculates the formulas.

Defining scenarios

To introduce you to Scenario Manager, this section starts with an example that uses a simplified production model, as shown in Figure 28.10.

FIGURE 28.10

A simple production model to demonstrate Scenario Manager

	A	B	C	D	E
1	Resource Cost Variables				
2	Hourly labor cost	30			
3	Material cost	57			
4					
5					
6		Product A	Product B	Product C	
7	Hours per unit	12	14	24	
8	Material per unit	6	9	14	
9	Cost to produce	\$702	\$933	\$1,518	
10	Sales price	\$795	\$1,295	\$2,195	
11	Unit profit	\$93	\$362	\$677	
12	Units produced	36	18	12	
13	Total profit per product	\$3,348	\$6,516	\$8,124	
14					
15	Total Profit	\$17,988			
16					



This workbook, named *production model scenarios.xlsx*, is available on this book's website at www.wiley.com/go/excel365bible.

This worksheet contains two input cells: the hourly labor cost (cell B2) and the unit cost for materials (cell B3). The company produces three products, and each product requires a different number of hours and a different amount of materials to produce.

Formulas calculate the total profit per product (row 13) and the total combined profit (cell B15). Management—trying to predict the total profit but uncertain what the hourly labor cost and material costs will be—has identified three scenarios, which are listed in Table 28.1.

TABLE 28.1 Three Scenarios for the Production Model

Scenario	Hourly Cost	Materials Cost
Best Case	30	57
Worst Case	38	62
Most Likely	34	59

The Best Case scenario has the lowest hourly cost and the lowest materials cost. The Worst Case scenario has high values for both the hourly cost and the materials cost. The third scenario, Most Likely, has intermediate values for both input cells. The managers need to be prepared for the worst case, however, and they're interested in what would happen under the Best Case scenario.

Choose Data \Rightarrow Forecast \Rightarrow What-If Analysis \Rightarrow Scenario Manager to display the Scenario Manager dialog box. When you first open this dialog box, it tells you that no scenarios are defined, which is not too surprising because you're just starting. As you add named scenarios, they appear in the Scenarios list in this dialog box.

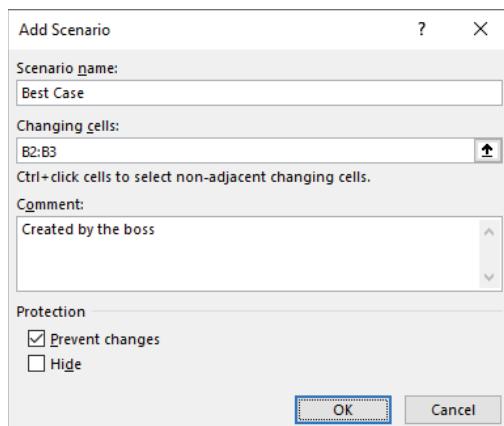
TIP

It's a good idea to create names for the changing cells and all the result cells that you want to examine. Excel uses these names in the dialog boxes and in the reports that it generates. If you use names, keeping track of what's going on is much easier; names also make your reports more readable.

To add a scenario, click the Add button in the Scenario Manager dialog box. Excel displays its Add Scenario dialog box, shown in Figure 28.11.

FIGURE 28.11

Use the Add Scenario dialog box to create a named scenario.



Part IV: Managing and Analyzing Data

This dialog box consists of four parts:

Scenario Name You can give the scenario any name you like.

Changing Cells These are the input cells for the scenario. You can enter the cell addresses directly or point to them. If you've created a name for the cells, type the name. Nonadjacent cells are allowed; if pointing to multiple cells, press Ctrl while you click the cells. Each named scenario can use the same set of changing cells or different changing cells. The number of changing cells for a scenario is limited to 32.

Comment By default, Excel displays the name of the person who created the scenario and the date it was created. You can change this text, add new text to it, or delete it. If you name the scenario well, you may not need much of a comment. However, some scenarios are so complex that more information will be useful both to you and to others who use your workbook.

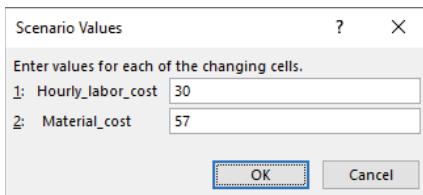
Protection The two Protection options (preventing changes and hiding a scenario) are in effect only when you protect the worksheet and choose the Scenario option in the Protect Sheet dialog box. Protecting a scenario prevents anyone from modifying it; a hidden scenario doesn't appear in the Scenario Manager dialog box.

In this example, define the three scenarios that are listed in Table 28.1. The changing cells are Hourly labor cost (B2) and Materials cost (B3).

After you enter the information in the Add Scenario dialog box, click OK. Excel then displays the Scenario Values dialog box, shown in Figure 28.12. This dialog box displays one field for each changing cell that you specified in the previous dialog box. Enter the values for each cell in the scenario. If you click OK, you return to the Scenario Manager dialog box, which then displays your named scenario in its list. If you have more scenarios to create, click the Add button to return to the Add Scenario dialog box.

FIGURE 28.12

You enter the values for the scenario in the Scenario Values dialog box.

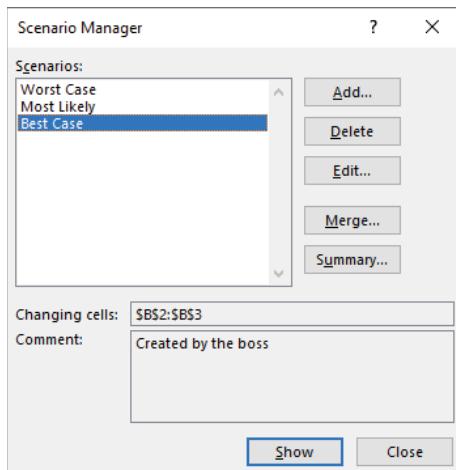


Displaying scenarios

After you define all the scenarios and return to the Scenario Manager dialog box, the dialog box displays the names of your defined scenarios. Select one of the scenarios and then click the Show button (or double-click the Scenario name). Excel inserts the corresponding values into the changing cells and calculates the worksheet to show the results for that scenario. Figure 28.13 shows an example of selecting a scenario.

FIGURE 28.13

Selecting a scenario to display



Using the Scenarios Drop-Down List

The Scenarios drop-down list shows all the defined scenarios and enables you to display a scenario quickly. Oddly, this useful tool doesn't appear on the Ribbon. But if you use Scenario Manager, you can add the Scenarios control to your Quick Access Toolbar. Here's how to do it:

1. Right-click the Quick Access Toolbar and choose Customize Quick Access Toolbar from the shortcut menu. The Excel Options dialog box appears with the Quick Access Toolbar tab selected.
2. From the Choose Commands From drop-down list, select Commands Not in the Ribbon.
3. Scroll down the list and select Scenario.
4. Click the Add button.
5. Click OK to close the Excel Options dialog box.

Alternatively, you can add the Scenarios control to the Ribbon. See Chapter 8, "Customizing the Excel User Interface," for additional details on customizing the Quick Access Toolbar and the Ribbon.

Modifying scenarios

After you've created scenarios, you may need to change them. To do so, follow these steps:

1. **Click the Edit button in the Scenario Manager dialog box to change one or more of the values for the changing cells of a scenario.**
2. **From the Scenarios list, select the scenario that you want to change and then click the Edit button.** The Edit Scenario dialog box appears.
3. **Click OK.** The Scenario Values dialog box appears.
4. **Make your changes and then click OK to return to the Scenario Manager dialog box.** Notice that Excel automatically updates the Comments box with new text that indicates when the scenario was modified.

Merging scenarios

In workgroup situations, you may have several people working on a spreadsheet model, and several people may have defined various scenarios. The marketing department, for example, may have its opinion of what the input cells should be, the finance department may have another opinion, and your CEO may have yet another opinion.

Excel makes it easy to merge these various scenarios into a single workbook. Before you merge scenarios, make sure that the workbook from which you're merging is open:

1. **Click the Merge button in the Scenario Manager dialog box.**
2. **From the Merge Scenarios dialog box that appears, choose the workbook that contains the scenarios you're merging in the Book drop-down list.**
3. **Choose the sheet that contains the scenarios you want to merge from the Sheet list box and click OK.** Notice that the bottom of the dialog box displays the number of scenarios in each sheet as you scroll through the Sheet list box.
4. **Click OK.** You return to the previous dialog box, which now displays the names of the scenarios that you merged from the other workbook.

Generating a scenario report

If you've created multiple scenarios, you may want to document your work by creating a scenario summary report. When you click the Summary button in the Scenario Manager dialog box, Excel displays the Scenario Summary dialog box.

You have a choice of report types:

Scenario Summary The summary report appears in the form of a worksheet outline.

Scenario PivotTable Report The summary report appears in the form of a PivotTable.



See Chapter 24, "Creating and Using Worksheet Outlines," for more information about outlines, and see Chapter 26, "Introducing PivotTables," for an introduction to PivotTables.

For simple cases of scenario management, a standard Scenario Summary report is usually sufficient. If you have many scenarios defined with multiple result cells, however, you may find that a Scenario PivotTable provides more flexibility.

The Scenario Summary dialog box also asks you to specify the result cells (the cells that contain the formulas in which you're interested). For this example, select B13:D13 and B15 on Sheet1 (a multiple selection) to make the report show the profit for each product plus the total profit.

NOTE

As you work with Scenario Manager, you may discover its main limitation, namely, that a scenario can use no more than 32 changing cells. If you attempt to use more cells, you get an error message.

Excel creates a new worksheet to store the summary table. Figure 28.14 shows the Scenario Summary form of the report. If you gave names to the changing cells and result cells, the table uses these names; otherwise, it lists the cell references.

FIGURE 28.14

A Scenario Summary report produced by Scenario Manager

	Current Values:	Worst Case	Most Likely	Best Case
Changing Cells:				
Hourly_labor_cost	38	38	34	30
Material_cost	62	62	59	57
Result Cells:				
ProductA_Profit	-\$1,188	-\$1,188	\$1,188	\$3,348
ProductB_Profit	\$3,690	\$3,690	\$5,184	\$6,516
ProductC_Profit	\$4,980	\$4,980	\$6,636	\$8,124
Total_Profit	\$7,482	\$7,482	\$13,008	\$17,988
Notes: Current Values column represents values of changing cells at time Scenario Summary Report was created. Changing cells for each scenario are highlighted in gray.				

Analyzing Data with Artificial Intelligence

Excel's Analyze Data feature uses artificial intelligence to analyze your data, make suggestions about ways to summarize your data, and allow you to ask questions about your data using natural language. This feature was formerly called Ideas and may still be labeled with that name in older versions.



This workbook, named bank accounts analyze.xlsx, is available on this book's website at www.wiley.com/go/excel365bible.

Using Excel's suggestions

Figure 28.15 shows part of a table containing data about bank accounts that were opened throughout a month. Data in a table, that is, structured data, gives Excel the best chance of producing useful results.

FIGURE 28.15

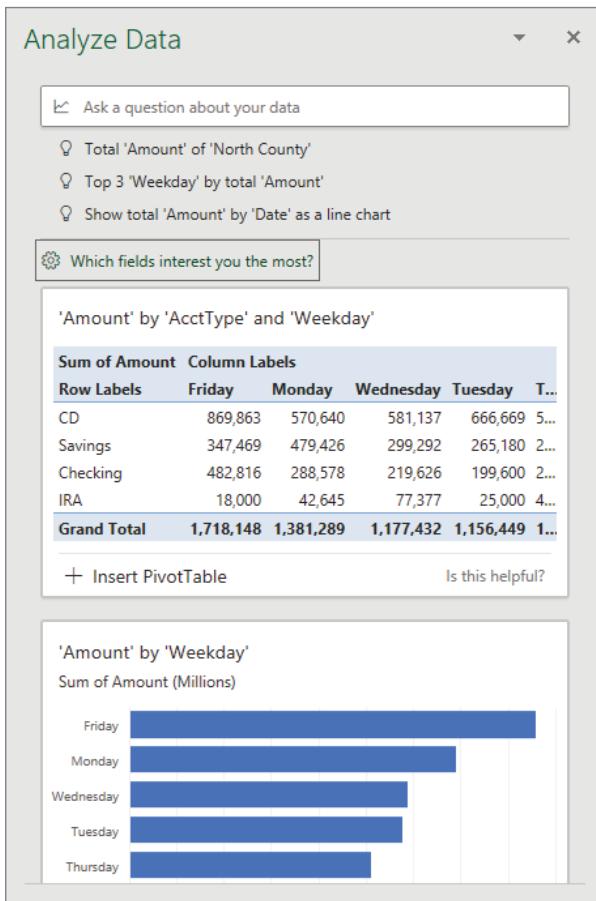
Bank accounts opened in a month

1	A Date	B Weekday	C Amount	D AcctType	E OpenedBy	F Branch	G Customer	H
2	4/1/2022	Friday	5,000	IRA	New Accts	Central	Existing	
3	4/1/2022	Friday	10,000	CD	Teller	North County	Existing	
4	4/1/2022	Friday	500	Checking	New Accts	Central	Existing	
5	4/1/2022	Friday	11,779	CD	Teller	Central	New	
6	4/1/2022	Friday	4,623	Savings	New Accts	North County	Existing	
7	4/1/2022	Friday	8,721	Savings	New Accts	Westside	New	
8	4/1/2022	Friday	15,276	Savings	New Accts	North County	Existing	
9	4/1/2022	Friday	5,000	Savings	New Accts	Westside	Existing	
10	4/1/2022	Friday	12,000	CD	New Accts	Westside	Existing	
11	4/1/2022	Friday	13,636	CD	New Accts	North County	Existing	
12	4/1/2022	Friday	7,177	Savings	Teller	North County	Existing	
13	4/1/2022	Friday	6,837	Savings	New Accts	Westside	Existing	
14	4/1/2022	Friday	3,171	Checking	New Accts	Westside	Existing	
15	4/1/2022	Friday	50,000	Savings	New Accts	Central	Existing	
16	4/1/2022	Friday	4,690	Checking	New Accts	North County	New	
17	4/1/2022	Friday	12,438	Checking	New Accts	Central	Existing	
18	4/1/2022	Friday	5,000	Checking	New Accts	North County	Existing	
19	4/1/2022	Friday	7,000	Savings	New Accts	North County	New	

To see Excel's suggestions, select any cell in the table and choose Home \Rightarrow Analysis \Rightarrow Analyze Data. Excel will briefly analyze the data and will then show the Analyze Data task pane, shown in Figure 28.16.

FIGURE 28.16

The Analyze Data task pane

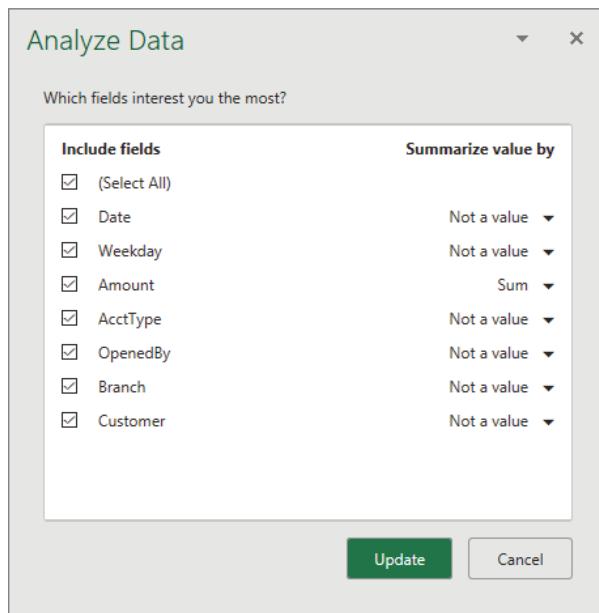


The top part of the Analyze Data task pane contains a text box where you can query your data. We discuss queries in the next section. Below the query section is the “Which fields interest you the most?” link. This link displays a list of fields from your data. You can uncheck fields that don’t interest you to narrow down Excel’s suggestions to focus on certain fields. Figure 28.17 shows the field list for the bank account data.

Part IV: Managing and Analyzing Data

FIGURE 28.17

Uncheck field names to narrow Excel's suggestions.



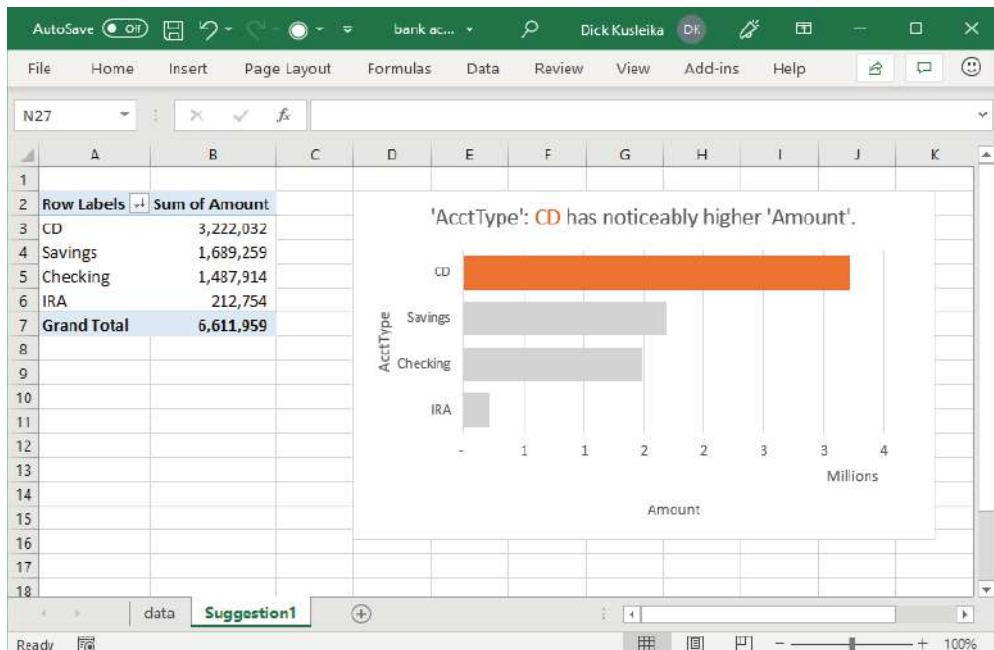
The bottom part of the Analyze Data task pane shows Excel's suggestions for summarizing your data. It's a list of PivotTables and PivotCharts that Excel thinks you might be interested in. Below each suggestion is a link that, when clicked, will insert that PivotTable or PivotChart into your worksheet. There is also a link called "Is this helpful?" where you can provide feedback on whether you think Excel's suggestion is useful. Figure 28.18 shows the result of inserting one of the suggestions. A new worksheet named Suggestion1 was created, and a PivotTable and PivotChart were added to the sheet.

You can see from this example that Excel isn't displaying every possible PivotTable you could create. It determined that CDs have an unusually high amount compared to other account types. It doesn't really know what an account type or an amount is, but its machine learning engine can uncover relationships without understanding the context.

The Analyze Data task pane shows a limited number of suggestions. However, there is a link at the bottom that expands to show you every suggestion available. You can add, however, as many suggestions to your workbook as you like.

FIGURE 28.18

A suggested PivotChart is inserted into a new worksheet.



28

Querying analyzed data

In addition to suggesting PivotTables and PivotCharts, the Analyze Data feature allows you to ask questions about your data using natural language. The term *natural language* means that you don't have to know any special keywords or syntax like, for example, the syntax you use when you write Excel formulas. You can simply type your questions as if you were speaking them.

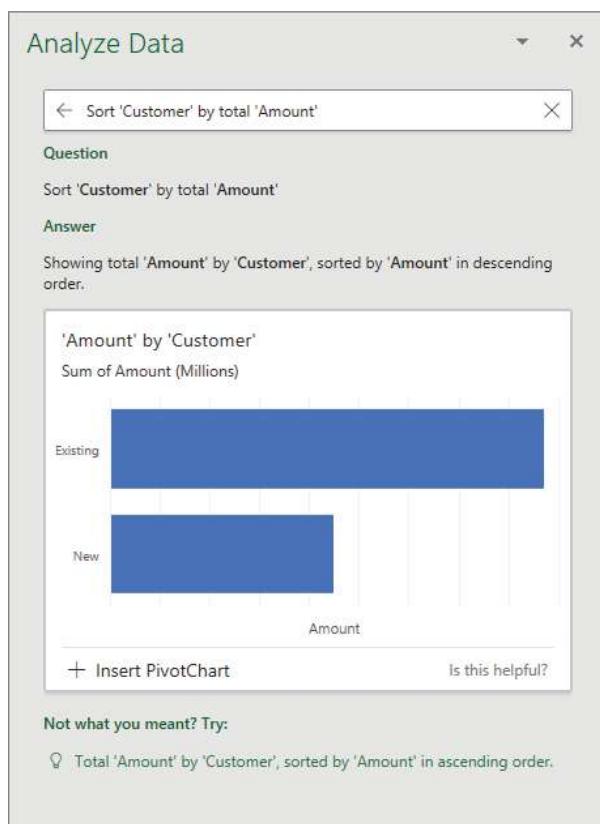
Below the query text box are several suggested queries. You can click on any of these suggestions to insert it in the text box. Figure 28.19 shows the Analyze Data task pane after the Sort Customer by Total Amount suggestion was clicked.

Based on that query, Excel suggests a PivotChart showing the difference in amounts between existing and new customers. You can insert that suggested PivotChart, type a new query in the text box, or return to the initial results using the left-pointing arrow in the text box.

Part IV: Managing and Analyzing Data

FIGURE 28.19

The Total Amount by Customer suggestion



In Chapter 26, “Introducing PivotTables”, we used PivotTables to answer questions about the bank account data. Figure 28.20 shows the results of entering the first question—“What is the daily total new deposit amount for each branch?”—in the query text box.

Analyze Data repeats the question, making words it thinks are significant bold (like *daily*) and striking through words it doesn’t understand (like *deposit*). You may understand the amount you pay when you open a new account is a deposit, but that word isn’t in any column headings or in the data, so Excel doesn’t understand it.

FIGURE 28.20

Using Analyze Data to query the data

The screenshot shows the 'Analyze Data' dialog box. At the top, a question is asked: 'What is the daily total new deposit amount for each branch?'. Below it, the word 'Answer' is displayed in green. A sub-question follows: 'Showing total 'Amount' by date of 'Date' and 'Branch' for 'Customer' 'New''. A table titled "'Amount' by date of 'Date' and 'Branch' for 'Customer' 'New'" is presented. The table has three columns: 'Date', 'Branch', and 'Sum of Amount'. The data is as follows:

Date	Branch	Sum of Amount
2022-04-01	Central	61,275
2022-04-01	Westside	8,721
2022-04-01	North County	16,190
2022-04-02	Central	44,932
2022-04-02	North County	22,049
2022-04-02	Westside	7,188
2022-04-04	North County	97,044
2022-04-04	Central	24,519
2022-04-04	Westside	27,429
2022-04-05	North County	8,807
...

At the bottom left is a '+ Insert' button, and at the bottom right is a link 'Is this helpful?'.

28

The results don't match the PivotTable from Chapter 26. The main difference is that Analyze Data suggested formulas rather than a PivotTable or PivotChart. Another difference is that Analyze Data is only summing amounts where the customer is new. You know from context that all these bank accounts were set up and are, therefore, new. But Excel doesn't know that and incorrectly assumes that we only want amounts for new customers.

Changing the query to "What is the daily total amount by branch?" doesn't result in a PivotTable either. Removing the word *new* from the query returned better results. It's unclear why Excel prefers hundreds of 750-character-long formulas over a PivotTable.

Analyzing Data Using Goal Seeking and Solver

IN THIS CHAPTER

Performing what-if analysis—in reverse

Single-cell goal seeking

Introducing Solver

Looking at Solver examples

The preceding chapter discussed what-if analysis—the process of changing input cells to observe the results on other dependent cells. This chapter looks at that process from the opposite perspective: finding the value of one or more input cells that produce a desired result in a formula cell.

Exploring What-If Analysis, in Reverse

Consider the following what-if question: “What is the total profit if sales increase by 20 percent?” If you set up your worksheet model properly, you can change the value in one or more cells to see what happens to the profit cell. The examples in this chapter take the opposite approach. If you know what a formula result should be, Excel can tell you the values that you need to enter in one or more input cells to produce that result. In other words, you can ask a question such as “How much do sales need to increase to produce a profit of \$1.2 million?” Excel provides two tools to accomplish this:

Goal Seek Determines the value that you need to enter in a single input cell to produce a result that you want in a dependent (formula) cell.

Solver Determines the values that you need to enter in multiple input cells to produce a result that you want. Moreover, because you can specify additional constraints to the problem, you gain significant problem-solving ability.

Using Single-Cell Goal Seeking

Single-cell goal seeking is a rather simple concept. Excel determines what value in an input cell produces a desired result in a formula cell. The following example shows you how single-cell goal seeking works.

Looking at a goal-seeking example

Figure 29.1 shows the mortgage loan worksheet used in the preceding chapter. This worksheet has four input cells (C4 : C7) and four formula cells (C10 : C13). Originally, this worksheet was used for a what-if analysis example. This example demonstrates the opposite approach. Rather than supply different input cell values to look at the calculated formulas, this example lets Excel determine one of the input values that will produce the desired result.

FIGURE 29.1

A mortgage calculator with input cells and formula cells

A	B	C	D
Mortgage Loan Worksheet			
Input Cells			
4	Purchase Price:	\$409,000	
5	Down Payment:	20%	
6	Loan Term (Months):	360	
7	Interest Rate (APR):	3.10%	
Result Cells			
10	Loan Amount:	\$327,200	
11	Monthly Payment:	\$1,397	
12	Total Payments:	\$502,991	
13	Total Interest:	\$175,791	
14			



This workbook is available on this book's website at www.wiley.com/go/excel365bible. The file is named `mortgage loan.xlsx`.

Assume that you're in the market for a new home and you know that you can afford a \$1,100 monthly mortgage payment. You also know that a lender can issue a 30-year fixed-rate mortgage loan for 3.10%, based on a 20% down payment. The question is, "What is the maximum purchase price I can afford?" In other words, what value in cell C4 (purchase price) causes the formula in cell C11 (monthly payment) to result in \$1,100? In this simple example, you could plug values into cell C4 until C11 displays \$1,100. With more complex models, Excel can usually determine the answer much more efficiently.

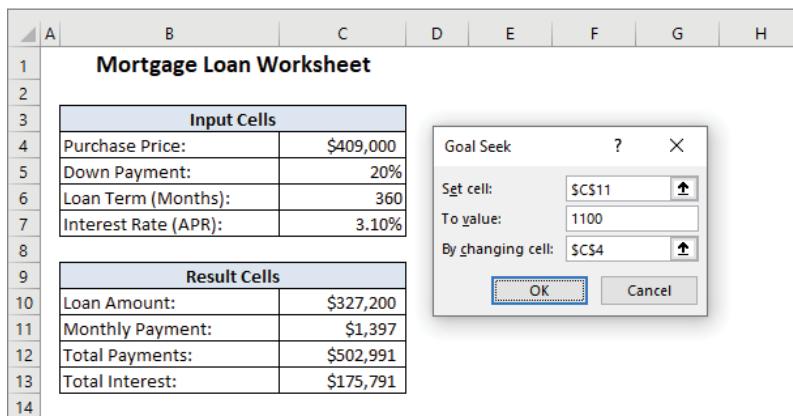
To answer the question posed in the preceding paragraph, first set up the input cells to match what you already know. Specifically:

1. Enter 20% into cell C5 (the down payment percent).
2. Enter 360 into cell C6 (the loan term, in months).
3. Enter 3.1% into cell C7 (the annual interest rate).

Next, choose Data \Rightarrow Forecast \Rightarrow What-If Analysis \Rightarrow Goal Seek. The Goal Seek dialog box appears. Completing this dialog box is similar to forming a sentence. You want to set cell C11 to 1100 by changing cell C4. Enter this information in the dialog box either by typing the cell references or by pointing with the mouse (see Figure 29.2). Click OK to begin the goal-seeking process.

FIGURE 29.2

The Goal Seek dialog box

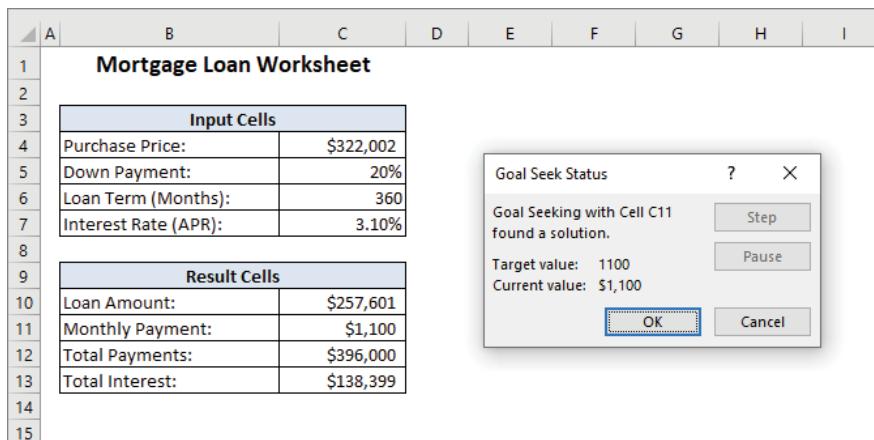


In less than a second, Excel displays the Goal Seek Status dialog box, shown in Figure 29.3. This shows the target value and the value that Excel calculated. In this case, Excel found an exact value. The worksheet now displays the found value in cell C4 (\$322,002). As a result of this value, the monthly payment amount is \$1,100. At this point, you have two options:

- Click OK to replace the original value with the found value. After doing so, you can use Undo to return to the original value.
- Click Cancel to restore your worksheet to the form that it had before you chose Goal Seek.

FIGURE 29.3

Goal Seek has found a solution.



Learning more about goal seeking

Excel can't always find a value that produces the result that you're seeking. Sometimes a solution simply doesn't exist. In such a case, the Goal Seek Status dialog box informs you of that fact.

At other times, however, Excel may report that it can't find a solution, but you're pretty sure that one exists. If that's the case, you can try the following options:

- Change the current value of the By Changing Cell field in the Goal Seek dialog box (refer to Figure 29.2) to a value that is closer to the solution and then reissue the command.
- Adjust the Maximum Iterations setting on the Formulas tab of the Excel Options dialog box (choose File \Rightarrow Options). Increasing the number of iterations (or calculations) makes Excel try more possible solutions.
- Double-check your logic and make sure that the formula cell does, indeed, depend on the specified changing cell.

NOTE

Like all computer programs, Excel has limited precision. To demonstrate this limitation, enter $=A1^2$ into cell A2. Then use the Goal Seek dialog box to find the value in cell A1 (which is empty) that makes the formula return 16. Excel comes up with a value of 4.00002269, which is close to the square root of 16 but certainly not exact. You can adjust the precision for goal seeking on the Formulas tab of the Excel Options dialog box. (Make the Maximum Change value smaller.)

NOTE

In some cases, multiple values of the input cell produce the same desired result. For example, the formula $=A1^2$ returns 16 if cell A1 contains either -4 or +4. If you use goal seeking when multiple solutions are possible, Excel gives you the solution that is closest to the current value.

Introducing Solver

The Excel Goal Seek feature is a useful tool, but it clearly has limitations. It can solve for only one adjustable cell, and it returns only a single solution. Excel's powerful Solver tool extends this concept by enabling you to do the following:

- Specify multiple adjustable cells.
- Specify constraints on the values that the adjustable cells can have.
- Generate a solution that maximizes or minimizes a particular worksheet cell.
- Generate multiple solutions to a problem.

Although goal seeking is a relatively simple operation, using Solver can be much more complicated. In fact, Solver is probably one of the most difficult (and potentially frustrating)

features in Excel, and it isn't for everyone. Most Excel users have no use for this feature. However, many users find that having this much power is worth spending the extra time to learn about it.

Looking at appropriate problems for Solver

Problems that are appropriate for Solver fall into a relatively narrow range. They typically involve situations that meet the following criteria:

- A target cell depends on other cells and formulas. Typically, you want to maximize or minimize this target cell or set it equal to some value.
- The target cell depends on a group of cells (called *changing cells*) that Solver can adjust to affect the target cell.
- The solution must adhere to certain limitations, or constraints.

After you set up your worksheet appropriately, you can use Solver to adjust the changing cells and produce the result that you want in your target cell while simultaneously meeting all the constraints that you defined.

No Solver Command?

You access Solver by choosing Data \Rightarrow Analyze \Rightarrow Solver. If this command isn't available, you need to install the Solver add-in. It's a simple process:

1. Choose File \Rightarrow Options. The Excel Options dialog box appears.
2. Select the Add-Ins tab.
3. At the bottom of the dialog box, select Excel Add-Ins from the Manage drop-down list and then click Go. The Add-Ins dialog box appears.
4. Place a check mark next to Solver Add-In and then click OK.

After you perform these steps, the Solver add-in loads whenever you start Excel.

29

Seeing a simple Solver example

We start with a simple example to introduce Solver and then present some increasingly complex examples to demonstrate what this feature can do.

Figure 29.4 shows a worksheet that is set up to calculate the profit for three products. Column B shows the number of units of each product, Column C shows the profit per unit for each product, and Column D contains formulas that calculate the total profit for each product by multiplying the units by the profit per unit.

Part IV: Managing and Analyzing Data

FIGURE 29.4

Use Solver to determine the number of units to maximize the total profit.

	A	B	C	D	E
1					
2	Units	Profit/Unit	Profit		
3	Product A	25	\$13	\$325	
4	Product B	25	\$18	\$450	
5	Product C	25	\$22	\$550	
6	Total	75		\$1,325	
7					



This workbook, named `three products.xlsx`, is available on this book's website at www.wiley.com/go/excel365bible.

You don't need an MBA to realize that the greatest profit comes from Product C. Therefore, to maximize total profit, the logical solution is to produce only Product C. If things were really this simple, you wouldn't need tools such as Solver. As in most situations, this company has some constraints that must be met:

- The combined production capacity is 300 total units.
- The company needs 50 units of Product A to fill an existing order.
- The company needs 40 units of Product B to fill an anticipated order.
- Because the market for Product C is relatively limited, the company doesn't want to produce more than 40 units of this product.

These four constraints make the problem more realistic and a bit more challenging. In fact, it's a perfect problem for Solver.

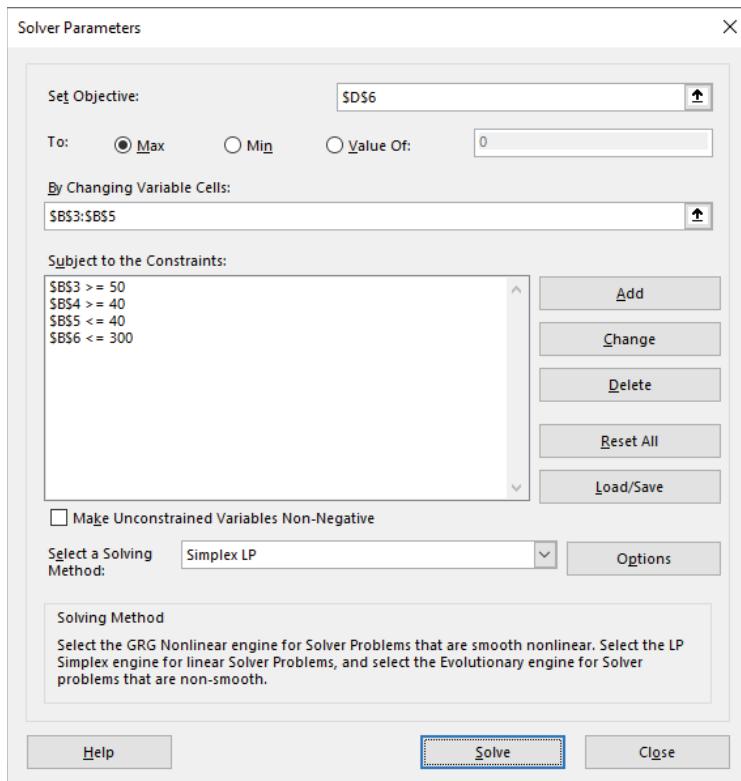
Here are the basic steps for using Solver:

1. **Set up the worksheet with values and formulas.** Make sure you format cells logically; for example, if you can't produce partial units of your products, format those cells to contain numbers with no decimal values.
2. **Choose Data ➔ Analyze ➔ Solver.** The Solver Parameters dialog box appears.
3. **Specify the target cell (also known as the objective).**
4. **Specify the range that contains the changing cells.**
5. **Specify the constraints.**
6. **Change the Solver options, if necessary.**
7. **Click Solve, and let Solver go to work.**

To start Solver to tackle this example, choose Data ➔ Analyze ➔ Solver. The Solver Parameters dialog box appears. Figure 29.5 shows this dialog box, set up to solve the problem.

FIGURE 29.5

The Solver Parameters dialog box



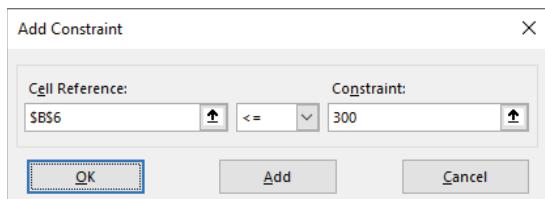
In this example, the target cell is D6—the cell that calculates the total profit for three products.

1. Enter D6 into the Set Objective field of the Solver Parameters dialog box.
2. Because the objective is to maximize this cell, select the Max option button.
3. Specify the changing cells (which are in the range B3 : B5) in the By Changing Variable Cells field. The next step is to specify the constraints on the problem. The constraints are added one at a time and appear in the Subject to the Constraints list.
4. To add a constraint, click the Add button. The Add Constraint dialog box, shown in Figure 29.6, appears. This dialog box has three parts: a Cell Reference value, an operator, and a Constraint value.

Part IV: Managing and Analyzing Data

FIGURE 29.6

The Add Constraint dialog box



5. To set the first constraint (that the total production capacity is 300 units), enter B6 as the Cell Reference value, choose less than or equal to (\leq) from the drop-down list of operators, and enter 300 as the Constraint value.
6. Click Add and enter the remaining constraints. Table 29.1 summarizes the constraints for this problem.

TABLE 29.1 Constraints Summary

Constraint	Expressed As
Capacity is 300 units	$B6 \leq 300$
At least 50 units of Product A	$B3 \geq 50$
At least 40 units of Product B	$B4 \geq 40$
No more than 40 units of Product C	$B5 \leq 40$

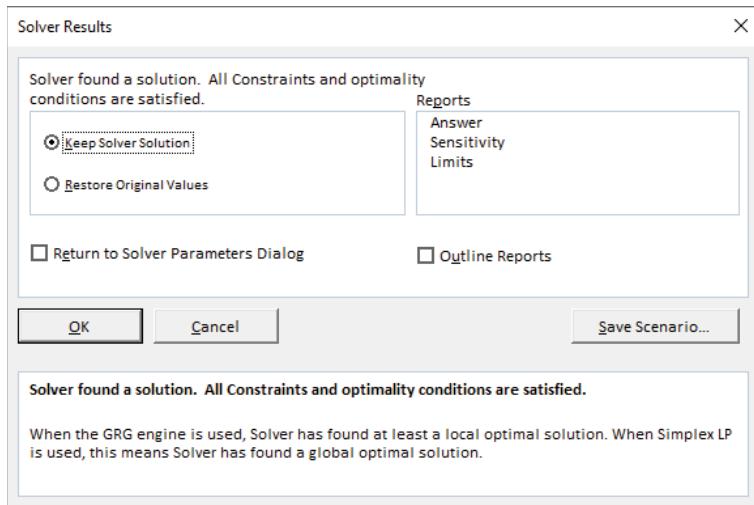
7. After you enter the last constraint, click OK to return to the Solver Parameters dialog box, which now lists the four constraints.
8. For the Solving Method, use Simplex LP.
9. Click the Solve button to start the solution process. You can watch the progress onscreen, and Excel soon announces that it has found a solution. The Solver Results dialog box is shown in Figure 29.7.

At this point, you have the following options:

- Keep the solution that Solver found.
- Restore the original changing cell values.
- Create any or all the three reports that describe what Solver did.
- Click the Save Scenario button to save the solution as a scenario so that Scenario Manager can use it.

FIGURE 29.7

Solver displays this dialog box when it finds a solution to the problem.



See Chapter 28, “Performing Spreadsheet What-If Analysis,” for more on Scenario Manager.

The Reports section of the Solver Results dialog box lets you select any or all of three optional reports. If you specify any report options, Excel creates each report on a new worksheet, with an appropriate name. Figure 29.8 shows an Answer Report. In the Constraints section of the report, three of the four constraints are binding, which means that these constraints were satisfied at their limit with no more room to change.

This simple example illustrates the way Solver works. The fact is, you could probably solve this particular problem manually by trial and error. That, of course, isn’t always the case.

CAUTION

When you close the Solver Results dialog box (by clicking either OK or Cancel), the Undo stack is cleared. In other words, you can’t undo any changes that Solver makes to your workbook.

Part IV: Managing and Analyzing Data

FIGURE 29.8

One of three reports that Solver can produce

	A	B	C	D	E	F	G	H	I	J	K
1											
2											
3											
4											
5											
6											
7											
8											
9											
10											
11											
12											
13											
14											
15											
16											
17											
18											
19											
20											
21											
22											
23											
24											
25											
26											
27											
28											
29											
30											
31											
32											

Selecting a Solving Method

Solver provides three solving methods, algorithms it uses to find the optimal solution. Mathematical optimization algorithms are beyond the scope of this book, but a brief explanation of each of the three methods follows:

- **Simplex LP:** This is the default method. It's best for linear programming problems. If you graph the inputs and outputs of your problem and it makes a straight line, you have a linear programming problem.
- **GRG Nonlinear:** This method is designed for smooth, nonlinear programming problems. A graph of these types of problems don't make a straight line, but the curved line it does make is continuous.
- **Evolutionary:** With this method, the best you can hope for is a good solution (instead of an optimal one). If you use functions like IF, CHOOSE, and VLOOKUP, it's an indicator you have a nonsmooth, nonlinear programming problem.

Exploring Solver options

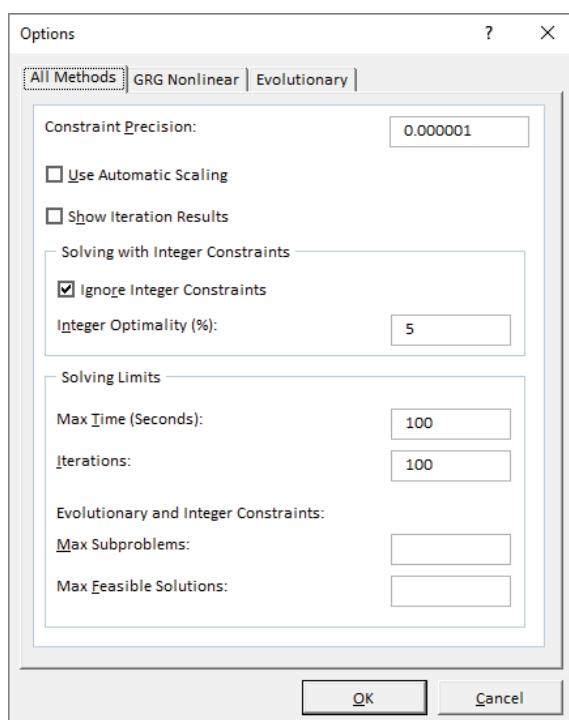
Before presenting more complex examples, this section discusses the Solver Options dialog box. From this dialog box, you control many aspects of the solution process, as well as load and save model specifications in a worksheet range.

Usually, you want to save a model only when you're using more than one set of Solver parameters with your worksheet. This is because Excel saves the first Solver model automatically with your worksheet (using hidden names). If you save additional models, Excel stores the information in the form of formulas that correspond to the specifications. (The last cell in the saved range is an array formula that holds the options settings.)

It's not unusual for Solver to report that it can't find a solution, even when you know that one should exist. Often, you can change one or more of the Solver options and try again. When you click the Options button in the Solver Parameters dialog box, the Solver Options dialog box, shown in Figure 29.9, appears.

FIGURE 29.9

You can control many aspects of the way Solver solves a problem.



This list describes Solver's options:

Constraint Precision Specify how close the Cell Reference and Constraint formulas must be to satisfy a constraint. Excel may solve the problem more quickly if you specify less precision.

Use Automatic Scaling Use when the problem involves large differences in magnitude—when you attempt to maximize a percentage, for example, by varying cells that are very large.

Show Iteration Results Instruct Solver to pause and display the results after each iteration by selecting this check box.

Ignore Integer Constraints When this check box is selected, Solver ignores constraints that specify that a particular cell must be an integer. Using this option may allow Solver to find a solution that can't be found otherwise.

Max Time Specify the maximum amount of time (in seconds) that you want Solver to spend on a problem. If Solver reports that it exceeded the time limit, you can increase the amount of time that it spends searching for a solution.

Iterations Enter the maximum number of trial solutions that you want Solver to perform.

Max Subproblems Use this for complex problems. Specify the maximum number of subproblems that may be explored by the Evolutionary algorithm.

Max Feasible Solutions Use this for complex problems. Specify the maximum number of feasible solutions that may be explored by the Evolutionary algorithm.

NOTE

The other two tabs in the Options dialog box contain additional options used by the GRG Nonlinear and Evolutionary algorithms.

Seeing Some Solver Examples

The remainder of this chapter consists of examples of using Solver for various types of problems.

Solving simultaneous linear equations

This example describes how to solve a set of three linear equations with three variables. Here's an example of a set of linear equations:

$$4x + y - 2z = 0$$

$$2x - 3y + 3z = 9$$

$$-6x - 2y + z = 0$$

The question that Solver will answer is this: What values of x , y , and z satisfy all three equations?

Figure 29.10 shows a workbook set up to solve this problem. This workbook has three named cells, which makes the formulas more readable:

x: C11

y: C12

z: C13

FIGURE 29.10

Solver will attempt to solve this series of linear equations.

B6	...	X	✓	f _x	= $(4*x)+(y)-(2*z)$
A	B	C	D		
1	$4x + y - 2z = 0$				
2	$2x - 3y + 3z = 9$				
3	$-6x -2y + z = 0$				
4					
5		Formula	Desired Value		
6	Equation 1:	3	0		
7	Equation 2:	2	9		
8	Equation 3:	-7	0		
9					
10	Variable	Value			
11	x:	1.00			
12	y:	1.00			
13	z:	1.00			
14					

All three named cells are initialized to 1.0 (which certainly doesn't solve the equations).



This workbook, named `linear equations.xlsx`, is available on this book's website at www.wiley.com/go/excel365bible.

29

The three equations are represented by formulas in the range B6:B8:

B6: = $(4*x)+(y)-(2*z)$

B7: = $(2*x)-(3*y)+(3*z)$

B8: = $-(6*x)-(2*y)+(z)$

These formulas use the values in the x , y , and z named cells. The range C6:C8 contains the desired result for these three formulas.

Part IV: Managing and Analyzing Data

Solver will adjust the values in x , y , and z —that is, the changing cells in C11:C13—subject to these constraints:

B6=C6
B7=C7
B8=C8

NOTE

This problem doesn't have a target cell because it's not trying to maximize or minimize anything. However, the Solver Parameters dialog box insists that you specify a formula for the Set Objective field. Therefore, just enter a reference to any cell that has a formula.

Figure 29.11 shows the solution. The x (0.75), y (-2.0), and z (0.50) values satisfy all three equations.

FIGURE 29.11

Solver finds a solution to the linear equations.

	A	B	C	D
1	4x + y - 2z =0			
2	2x - 3y + 3z =9			
3	-6x -2y + z = 0			
4				
5		Formula	Desired Value	
6	Equation 1:	0	0	
7	Equation 2:	9	9	
8	Equation 3:	0	0	
9				
10		Variable	Value	
11		x:	0.75	
12		y:	-2.00	
13		z:	0.50	
14				

NOTE

A set of linear equations may have one solution, no solution, or an infinite number of solutions.

Minimizing shipping costs

This example involves finding alternative options for shipping materials, while keeping total shipping costs at a minimum (see Figure 29.12). A company has warehouses in Los Angeles, St. Louis, and Boston. Retail outlets throughout the United States place orders, which the company then ships from one of the warehouses. The company wants to meet the product needs of all six retail outlets from available inventory and keep total shipping charges as low as possible.

FIGURE 29.12

This worksheet determines the least expensive way to ship products from warehouses to retail outlets.



This workbook, named **shipping costs.xlsx**, is available on this book's website at www.wiley.com/go/excel365bible.

The following is an explanation of each part of this somewhat complicated workbook:

Shipping Costs Table This table, in range B2 : E8, is a matrix that contains per-unit shipping costs from each warehouse to each retail outlet. The cost to ship a unit from Los Angeles to Denver, for example, is \$58.

Product needs of each retail store This information appears in C12:C17. For example, Denver needs 150 units, Houston needs 225, and so on. C18 contains a formula that calculates the total needed.

Number to ship from Range D12 : F17 holds the adjustable cells that Solver will change. All these cells are initialized with a value of 25 to give Solver a starting value. Column G contains formulas that sum the number of units the company needs to ship to each retail outlet.

Warehouse inventory Row 21 contains the amount of inventory at each warehouse, and row 22 contains formulas that subtract the amount shipped (row 18) from the inventory.

Calculated shipping costs Row 24 contains formulas that calculate the shipping costs.

Cell D24 contains the following formula, which is copied to the two cells to the right of cell D24:

=SUMPRODUCT(C3:C8,D12:D17)

Cell G24 is the total shipping costs for all orders.

Solver fills in values in the range D12 : F17 in such a way that minimizes shipping costs while still supplying each retail outlet with the desired number of units. In other words, the solution minimizes the value in cell G24 by adjusting the cells in D12 : F17, subject to the following constraints:

- **The number of units needed by each retail outlet must equal the number shipped.** (In other words, all the orders are filled.) These constraints are represented by the following specifications:

C12=G12 C14=G14 C16=G16

C13=G13 C15=G15 C17=G17

- **The number of units remaining in each warehouse's inventory must not be negative.** (In other words, they can't ship more than what's available.) This is represented by the following constraint specifications:

D22>=0 E22>=0 F22>=0

- **The adjustable cells can't be negative because shipping a negative number of units makes no sense.** The Solver Parameters dialog box has a handy option: Make Unconstrained Variables Non-Negative. Make sure that this check box is selected.

NOTE

Before you solve this problem with Solver, you may want to attempt to solve this problem manually by entering values in D12 : F17 that minimize the shipping costs. And, of course, you need to make sure that all the constraints are met. Doing so may help you better appreciate Solver.

Setting up the problem is the difficult part. For example, you must enter nine constraints. When you have specified all the necessary information, click the Solve button to put Solver to work. Solver displays the solution shown in Figure 29.13.

FIGURE 29.13

The solution that was created by Solver

	A	B	C	D	E	F	G	H
Shipping Costs Table								
1								
2								
3				L.A.	St. Louis	Boston		
4		Denver		\$58	\$47	\$108		
5		Houston		\$87	\$46	\$100		
6		Atlanta		\$121	\$30	\$57		
7		Miami		\$149	\$66	\$83		
8		Seattle		\$62	\$115	\$164		
9		Detroit		\$128	\$28	\$38		
10								
11								
12								
13								
14								
15								
16								
17								
18								
19								
20								
21								
22								
23								
24								
25								
Warehouse Inventory								
Starting Inventory:			400	350	500			
No. Remaining:			130	0	125			
Shipping Costs:			\$ 16,140	\$ 15,000	\$ 24,375	\$55,515		

The total shipping cost is \$55,515, and all the constraints are met. Notice that shipments to Miami come from both St. Louis and Boston and that St. Louis is now out of inventory. As you can imagine, you would never want to run a warehouse completely out of inventory. If that were the case, you could change the constraints for cells D22 : F22 to be greater than some minimum acceptable quantity.

Allocating resources

The example in this section is a common type of problem that's ideal for Solver. Essentially, problems of this sort involve optimizing the volumes of individual production units that use varying amounts of fixed resources. Figure 29.14 shows a simplified example for a toy company.

Part IV: Managing and Analyzing Data

FIGURE 29.14

Using Solver to maximize profit when resources are limited

	A	B	C	D	E	F	G	H	I	J
1	XYZ Toys Inc.									
2	Materials Needed									
3	Material	Toy A	Toy B	Toy C	Toy D	Toy E	Amt. Avail.	Amt. Used	Amt. Left	
4	Red Paint	0	1	0	1	3	625	250	375	
5	Blue Paint	3	1	0	1	0	640	250	390	
6	White Paint	2	1	2	0	2	1,100	350	750	
7	Plastic	1	5	2	2	1	875	550	325	
8	Wood	3	0	3	5	5	2,200	800	1,400	
9	Glue	1	2	3	2	3	1,500	550	950	
10	Unit Profit	\$15	\$30	\$20	\$25	\$25				
11	No. to Make	50	50	50	50	50				
12	Profit	\$750	\$1,500	\$1,000	\$1,250	\$1,250				
13	Total Profit	\$5,750								
14										



This workbook is available on this book's website at www.wiley.com/go/excel365bible. The file is named `allocating resources.xlsx`.

This company makes five different toys that use six different materials in varying amounts. For example, Toy A requires 3 units of blue paint, 2 units of white paint, 1 unit of plastic, 3 units of wood, and 1 unit of glue. Column G shows the current inventory of each type of material. Row 10 shows the unit profit for each toy.

The number of toys to make is shown in the range B11:F11. These are the values that Solver determines (the changing cells). The goal of this example is to determine how to allocate the resources to maximize the total profit (B13). In other words, Solver determines how many units of each toy to make. The constraints in this example are relatively simple:

- **Ensure that production doesn't use more resources than are available.** This can be accomplished by specifying that each cell in column I is greater than or equal to zero.
- **Ensure that the quantities produced aren't negative.** This can be accomplished by specifying the Make Unconstrained Variables Non-Negative option.

Figure 29.15 shows the results that are produced by Solver. It shows the product mix that generates \$12,365 in profit and uses all resources in their entirety, except for glue.

FIGURE 29.15

Solver determined how to use the resources to maximize the total profit.

	A	B	C	D	E	F	G	H	I	J
1	XYZ Toys Inc.									
2	Materials Needed									
3	Material	Toy A	Toy B	Toy C	Toy D	Toy E	Amt. Avail.	Amt. Used	Amt. Left	
4	Red Paint	0	1	0	1	3	625	625	0	
5	Blue Paint	3	1	0	1	0	640	640	0	
6	White Paint	2	1	2	0	2	1,100	1,100	0	
7	Plastic	1	5	2	2	1	875	875	0	
8	Wood	3	0	3	5	5	2,200	2,200	0	
9	Glue	1	2	3	2	3	1,500	1,353	147	
10	Unit Profit	\$15	\$30	\$20	\$25	\$25				
11	No. to Make	194	19	158	40	189				
12	Profit	\$2,903	\$573	\$3,168	\$1,008	\$4,713				
13	Total Profit	\$12,365								
14										

Optimizing an investment portfolio

This example demonstrates how to use Solver to help maximize the return on an investment portfolio. A portfolio consists of several investments, each of which has a different yield. In addition, you may have some constraints that involve reducing risk and diversification goals. Without such constraints, a portfolio problem becomes a no-brainer: put all your money in the investment with the highest yield.

This example involves a credit union (a financial institution that takes members' deposits and invests them in loans to other members, bank CDs, and other types of investments). The credit union distributes part of the return on these investments to the members in the form of dividends, or interest on their deposits.

This hypothetical credit union must adhere to some regulations regarding its investments, and the board of directors has imposed some other restrictions. These regulations and restrictions are the problem's constraints. Figure 29.16 shows a workbook set up for this problem.



This workbook is available on this book's website at www.wiley.com/go/excel365bible. The file is named `investment portfolio.xlsx`.

Part IV: Managing and Analyzing Data

FIGURE 29.16

This worksheet is set up to maximize a credit union's investments, given some constraints.

	A	B	C	D	E	F
1	Portfolio Amount:	\$5,000,000				
2						
3						
4	Investment	Pct Yield	Amount Invested	Yield	Pct. of Portfolio	
5	New Car Loans	6.90%	1,000,000	69,000	20.00%	
6	Used Car Loans	8.25%	1,000,000	82,500	20.00%	
7	Real Estate Loans	8.90%	1,000,000	89,000	20.00%	
8	Unsecured Loans	13.00%	1,000,000	130,000	20.00%	
9	Bank CDs	4.60%	1,000,000	46,000	20.00%	
10	TOTAL		\$5,000,000	\$416,500	100.00%	
11						
12		Total Yield:	8.33%			
13						
14		Auto Loans	40.00%			
15						

Allocating the \$5 million portfolio is subject to these constraints:

- **The amount that the credit union invests in new-car loans must be at least three times the amount that the credit union invests in used-car loans.** (Used-car loans are riskier investments.) This constraint is represented as:

$$C5 \geq C6 * 3$$

- **Car loans should make up at least 15% of the portfolio.** This constraint is represented as:

$$D14 \geq .15$$

- **Unsecured loans should make up no more than 25% of the portfolio.** This constraint is represented as:

$$E8 \leq .25$$

- **At least 10% of the portfolio should be in bank CDs.** This constraint is represented as:

$$E9 \geq .10$$

- **The total amount invested is \$5,000,000.**

- **All investments should be positive or zero.**

The changing cells are C5 : C9, and the goal is to maximize the total yield in cell D12. Starting values of 1,000,000 have been entered in the changing cells. When you run Solver with these parameters, it produces the solution shown in Figure 29.17, which has a total yield of 9.25%.

FIGURE 29.17

The results of the portfolio optimization

	A	B	C	D	E	F
1	Portfolio Amount:	\$5,000,000				
2						
3						
4	Investment	Pct Yield	Amount Invested	Yield	Pct. of Portfolio	
5	New Car Loans	6.90%	562,500	38,813	11.25%	
6	Used Car Loans	8.25%	187,500	15,469	3.75%	
7	Real Estate Loans	8.90%	2,500,000	222,500	50.00%	
8	Unsecured Loans	13.00%	1,250,000	162,500	25.00%	
9	Bank CDs	4.60%	500,000	23,000	10.00%	
10	TOTAL		\$5,000,000	\$462,281	100.00%	
11						
12			Total Yield:	9.25%		
13						
14			Auto Loans	15.00%		
15						

Learning More about Solver

Solver is a complex tool, and this chapter barely scratches the surface. If you'd like to learn more about Solver, visit the website for Frontline Systems (www.solver.com). Frontline Systems is the company that developed Solver for Excel. Its website has several tutorials and lots of helpful information, including a detailed manual that you can download. You can also find additional Solver products for Excel that can handle much more complex problems.

Analyzing Data with the Analysis ToolPak

IN THIS CHAPTER

Getting an overview of the Analysis ToolPak

Using the Analysis ToolPak

Meeting the Analysis ToolPak tools

Although Excel was designed primarily for business users, people in other disciplines, including education, research, statistics, and engineering, also use the software. One way Excel addresses these nonbusiness users is with its Analysis ToolPak add-in. However, many features in the Analysis ToolPak are valuable for business applications as well.

The Analysis ToolPak: An Overview

The Analysis ToolPak is an add-in that provides analytical capability that normally isn't available in Excel.

These analysis tools offer many features that may be useful to those in the scientific, engineering, and educational communities—not to mention business users whose needs extend beyond the normal spreadsheet fare.

The following list is of the types of analyses that you can perform with the Analysis ToolPak. We cover each of these tools in this chapter.

- Analysis of variance (three types)
- Correlation
- Covariance
- Descriptive statistics
- Exponential smoothing
- F-Test
- Fourier analysis
- Histogram
- Moving average
- Random number generation

- Rank and percentile
- Regression
- Sampling
- t-Test (three types)
- z-Test

As you can see, the Analysis ToolPak add-in brings a great deal of additional functionality to Excel. These procedures have limitations, however. In some cases, you may prefer to create your own formulas to do some calculations.

Installing the Analysis ToolPak Add-In

The Analysis ToolPak is implemented as an add-in. Before you can use it, though, you need to make sure that the add-in is installed. Select the Data tab. If you see an Analyze group, showing a Data Analysis button, the Analysis ToolPak is installed. If you can't access Data \Rightarrow Analysis \Rightarrow Data Analysis, install the add-in by following these steps:

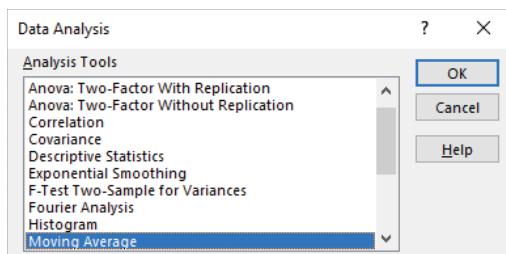
1. **Choose File \Rightarrow Options.** The Excel Options dialog box appears.
2. **Select the Add-Ins tab.**
3. **At the bottom of the dialog box, select Excel Add-Ins from the Manage drop-down list and then click Go.** The Add-Ins dialog box appears.
4. **Place a check mark next to Analysis ToolPak.** There is no need to check the add-in named Analysis ToolPak – VBA.
5. **Click OK to close the Add-Ins dialog box.**

Using the Analysis Tools

Using the procedures in the Analysis ToolPak add-in is relatively straightforward as long as you're familiar with the particular analysis type. To use any of these tools, choose Data \Rightarrow Analysis \Rightarrow Data Analysis, and the Data Analysis dialog box appears, as shown in Figure 30.1. Scroll through the list until you find the analysis tool that you want to use, select it, and then click OK. A dialog box specific to the procedure that you select appears.

FIGURE 30.1

Select your tool from the Data Analysis dialog box.



Usually, you need to specify one or more input ranges plus an output range. (Specifying the top-left cell of the output range is sufficient.) Alternatively, you can choose to place the results on a new worksheet or in a new workbook. The procedures vary in the amount of additional information required. In many dialog boxes, you may be able to indicate whether your data range includes labels. If so, you can specify the entire range, including the labels, and indicate to Excel that the first column (or row) contains labels. Excel then uses these labels in the tables that it produces. Most tools also provide different output options that you can select, based on your needs.

CAUTION

The Analysis ToolPak isn't consistent in the way that it generates its output. In some cases, the procedures use formulas, so you can change your data, and the results update automatically. For other procedures, Excel stores the results as values, so if you change your data, the results don't reflect your changes.

Introducing the Analysis ToolPak Tools

Each of the following sections describes a tool in the Analysis ToolPak and provides an example. We don't describe every available option in these procedures. If you need to use the advanced analysis tools, you probably already know how to use most of the options not covered here.

Before you use any of these tools, read the appropriate section in Excel's Help system.



This book's website at www.wiley.com/go/excel1365bible contains a workbook that shows output from all of the tools discussed in this chapter. The file is named `atp_examples.xlsx`. This workbook also contains some alternative formula-based solutions that are sometimes better than using the Analysis ToolPak.

Analysis of variance

Analysis of variance (sometimes abbreviated as *anova*) is a statistical test that determines whether two or more samples were drawn from the same population. Using tools in the Analysis ToolPak, you can perform three types of analysis of variance:

Single-factor A one-way analysis of variance, with only one sample for each group of data

Two-factor with replication A two-way analysis of variance, with multiple samples (or replications) for each group of data

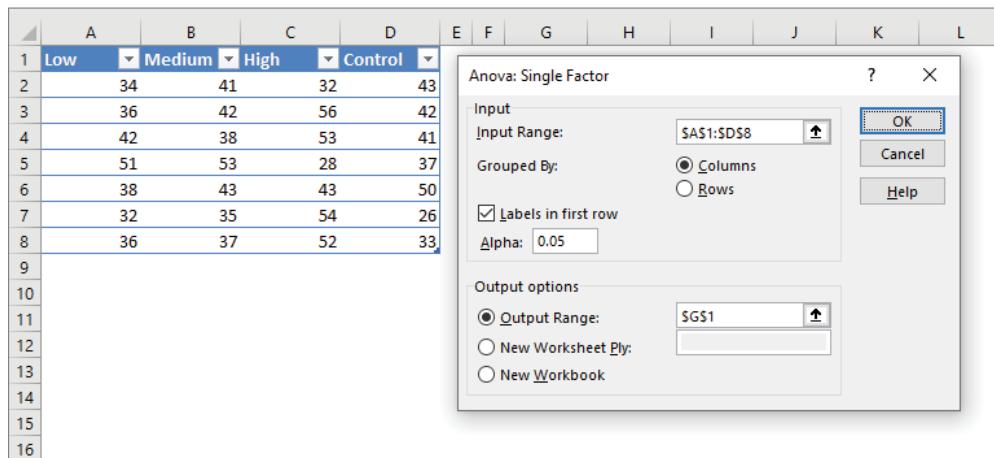
Two-factor without replication A two-way analysis of variance, with a single sample (or replication) for each group of data

Figure 30.2 shows the dialog box for a single-factor analysis of variance. Alpha represents the statistical confidence level for the test.

Part IV: Managing and Analyzing Data

FIGURE 30.2

Specifying parameters for a single-factor analysis of variance



The output for this test consists of the means and variances for each of the samples, the value of F, the critical value of F, and the significance of F (P-value).

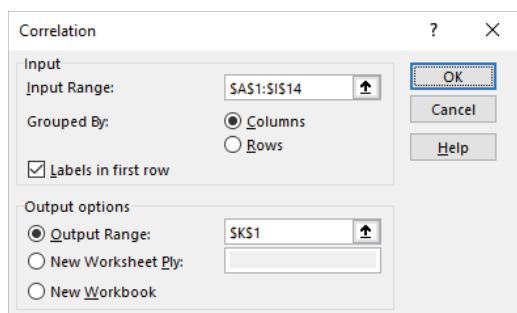
Correlation

Correlation is a widely used statistic that measures the degree to which two sets of values vary together. For example, if higher values in one data set are typically associated with higher values in the second data set, the two data sets have a positive correlation. The degree of correlation is expressed as a coefficient that ranges from -1.0 (a perfect negative correlation) to +1.0 (a perfect positive correlation). A correlation coefficient of 0 indicates that the two variables aren't correlated.

Figure 30.3 shows the Correlation dialog box. Specify the input range, which can include any number of variables, arranged in rows or columns.

FIGURE 30.3

The Correlation dialog box



The output consists of a correlation matrix that shows the correlation coefficient for each variable paired with every other variable.

Note

The resulting correlation matrix doesn't use formulas to calculate the results. Therefore, if any data changes, the correlation matrix isn't valid. You can use the `CORREL` function to create a correlation matrix that changes automatically when you change data.

Covariance

The covariance tool produces a matrix that is similar to the one generated by the correlation tool. *Covariance*, like correlation, measures the degree to which two variables vary together. Specifically, covariance is the average of the product of the deviations of each data point pair from their respective means.

Because the covariance tool does not generate formulas, you may prefer to calculate a covariance matrix using the `COVAR` function.

Descriptive statistics

The *descriptive statistics* tool produces a table that describes your data with some standard statistics. Figure 30.4 shows some sample output.

Because the output for this procedure consists of values (not formulas), you should use this procedure only when you're certain that your data isn't going to change; otherwise, you'll need to re-execute the procedure. You can generate all these statistics by using formulas.

Exponential smoothing

Exponential smoothing is a technique for predicting data that is based on the previous data point and the previously predicted data point. You can specify the damping factor (also known as a *smoothing constant*), which can range from 0 to 1. This factor determines the relative weighting of the previous data point and the previously predicted data point. You also can request standard errors and a chart.

The exponential smoothing procedure generates formulas that use the damping factor you specify. Therefore, if the data changes, Excel updates the formulas.

F-Test (two-sample test for variance)

An *F-Test* is a commonly used statistical test that enables you to compare two population variances. Figure 30.5 shows a small data set and F-Test output.

The output for this test consists of the means and variances for each of the two samples, the value of F, the critical value of F, and the significance of F.

Part IV: Managing and Analyzing Data

FIGURE 30.4

Descriptive statistics output

	A	B	C	D	E	F	G	H	I	J	K
	W. Coast Sample	Midwest Sample	E. Coast Sample		W. Coast Sample	Midwest Sample	E. Coast Sample				
1	35	41	52		Mean	39.25 Mean	46 Mean	41.35			
2	32	35	29		Standard Error	1.84801 Standard Error	2.10763 Standard Error	1.56487			
3	46	36	43		Median	37.5 Median	45.5 Median	41.5			
4	57	45	45		Mode	37 Mode	52 Mode	37			
5	45	44	28		Standard Deviation	8.26454 Standard Deviation	9.42561 Standard Deviation	6.99831			
6	28	62	35		Sample Variance	68.3026 Sample Variance	88.8421 Sample Variance	48.9763			
7	60	61	37		Kurtosis	1.47266 Kurtosis	-0.477 Kurtosis	-0.2803			
8	37	62	32		Skewness	1.18011 Skewness	0.14121 Skewness	-0.2486			
9	34	36	37		Range	32 Range	34 Range	26			
10	33	52	41		Minimum	28 Minimum	28 Minimum	28			
11	37	46	54		Maximum	60 Maximum	62 Maximum	54			
12	32	52	44		Sum	785 Sum	920 Sum	827			
13	38	38	42		Count	20 Count	20 Count	20			
14	41	28	48		Confidence	Confidence	Confidence	Confidence			
15	38	50	46		Level(95.0%)	3.86793 Level(95.0%)	4.41132 Level(95.0%)	3.27531			
16	42	52	47								
17	29	48	39								
18	40	38	40								
19	37	44	41								
20	44	50	47								
21											
22											

FIGURE 30.5

Output from the F-Test tool

	A	B	C	D	E	F	G
1	Group 1	Group 2		F-Test Two-Sample for Variances			
2	96	39			Group 1	Group 2	
3	78	53			Mean	75.44444	46.66667
4	72	51			Variance	109.5278	25
5	78	48			Observations	9	9
6	65	51			df	8	8
7	66	42			F	4.381111	
8	69	44			P(F<=f) one-tail	0.025855	
9	87	42			F Critical one-tai	3.438103	
10	68	50					
11							

Fourier analysis

Fourier analysis performs a “fast Fourier” transformation of a range of data. Using the Fourier analysis tool, you can transform a range limited to the following sizes: 1, 2, 4, 8, 16, 32, 64, 128, 256, 512, or 1,024 data points. This procedure accepts and generates complex numbers, which are represented as a text string (not numerical values).

Histogram

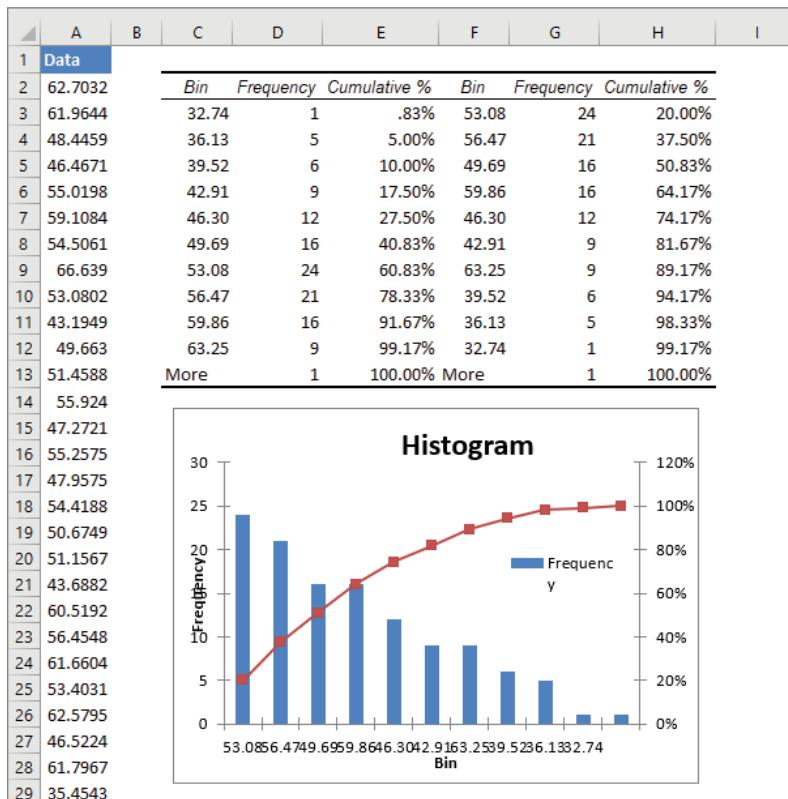
The *histogram* tool is useful for producing data distributions and histogram charts. It accepts an Input range and a Bin range. A Bin range is a range of values that specifies the limits for each column of the histogram. If you omit the Bin range, Excel creates 10 equal-interval bins for you. The size of each bin is determined by the following formula:

$$= (\text{MAX}(\text{input_range}) - \text{MIN}(\text{input_range})) / 10$$

Figure 30.6 shows output from the histogram tool. As an option, you can specify that the resulting histogram be sorted by frequency of occurrence in each bin.

FIGURE 30.6

Use the histogram tool to generate distributions and graphical output.



If you specify the Pareto (Sorted Histogram) option, the Bin range must contain values and can't contain formulas. If formulas appear in the Bin range, Excel doesn't sort properly, and your worksheet displays error values. Since the histogram tool doesn't use formulas, if you

change any of the input data, you need to repeat the histogram procedure to update the results.



For other ways of generating frequency distributions, see Chapter 27, “Analyzing Data with PivotTables.” Excel also supports two chart types: Histogram and Pareto. These work with a range of data and do not require a separate bin range. For an example, see Chapter 18, “Getting Started with Excel Charts.”

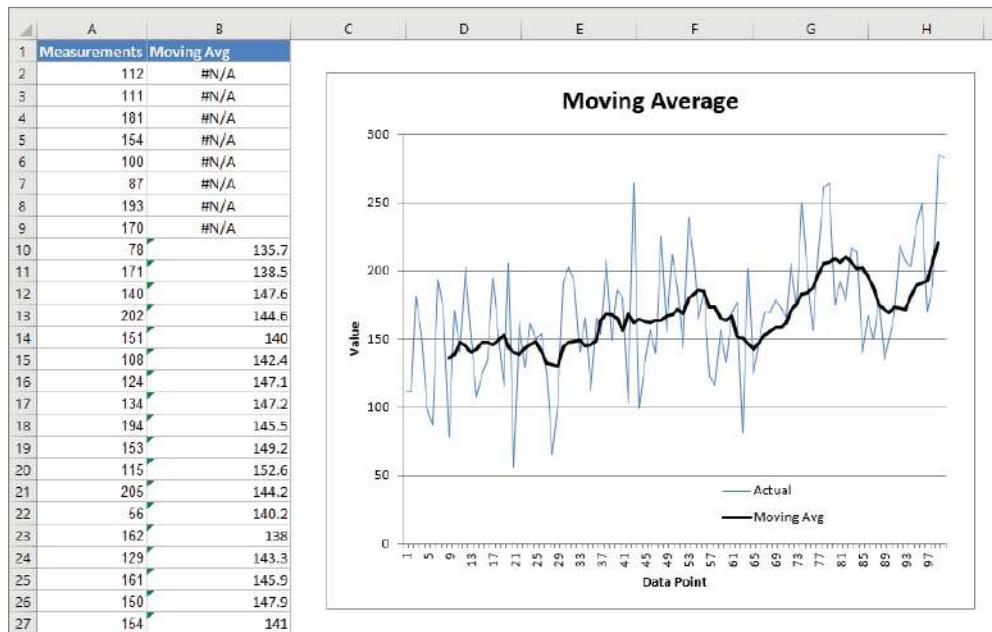
Moving average

The *moving average* tool helps you smooth out a data series that has a lot of variability. This procedure is often used in conjunction with a chart. Excel does the smoothing by computing a moving average of a specified number of values. In many cases, a moving average enables you to spot trends that otherwise would be obscured by noise in the data.

Figure 30.7 shows a chart generated by the moving average tool. You can, of course, specify the number of values that you want Excel to use for each average. If you select the Standard Errors check box in the Moving Average dialog box, Excel calculates standard errors and places formulas for these calculations next to the moving average formulas. The standard error values indicate the degree of variability between the actual values and the calculated moving averages.

FIGURE 30.7

A chart produced from data generated by the moving average tool



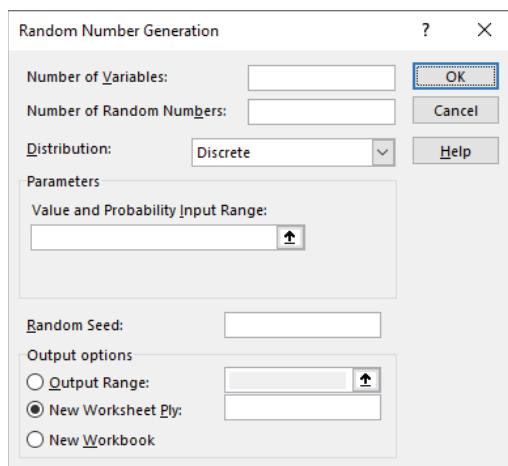
The first few cells in the output are #N/A because too few data points exist to calculate the average for these initial values.

Random number generation

Although Excel contains built-in functions to calculate random numbers, the *random number generation* tool is much more flexible because you can specify what type of distribution you want the random numbers to have. Figure 30.8 shows the Random Number Generation dialog box. The Parameters section varies, depending on the type of distribution that you select.

FIGURE 30.8

This dialog box enables you to generate a wide variety of random numbers.



Number of Variables refers to the number of columns that you want, and Number of Random Numbers refers to the number of rows that you want. For example, if you want 200 random numbers arranged in 10 columns of 20 rows, you specify 10 and 20, respectively, in these fields.

In the Random Seed field, you can specify the starting value that Excel uses in its random-number-generating algorithm. Usually, you leave this field blank. If you want to generate the same random number sequence, however, you can specify a seed between 1 and 32,767 (integer values only). You can create the following types of distributions via the Distribution drop-down list in the Random Number Generation dialog box:

Uniform Every random number has an equal chance of being selected. You specify the upper and lower limits.

Normal The random numbers correspond to a normal distribution. You specify the mean and standard deviation of the distribution.

Bernoulli The random numbers are either 0 or 1, determined by the probability of success that you specify.

Binomial This option returns random numbers based on a Bernoulli distribution over a specific number of trials, given a probability of success that you specify.

Poisson This option generates values in a Poisson distribution. A Poisson distribution is characterized by discrete events that occur in an interval, where the probability of a single occurrence is proportional to the size of the interval. The lambda parameter is the expected number of occurrences in an interval. In a Poisson distribution, lambda is equal to the mean, which also is equal to the variance.

Patterned This option doesn't generate random numbers. Rather, it repeats a series of numbers in steps that you specify.

Discrete This option enables you to specify the probability that specific values are chosen. It requires a two-column input range; the first column holds the values, and the second column holds the probability of each value being chosen. The sum of the probabilities in the second column must equal 100 percent.

Rank and percentile

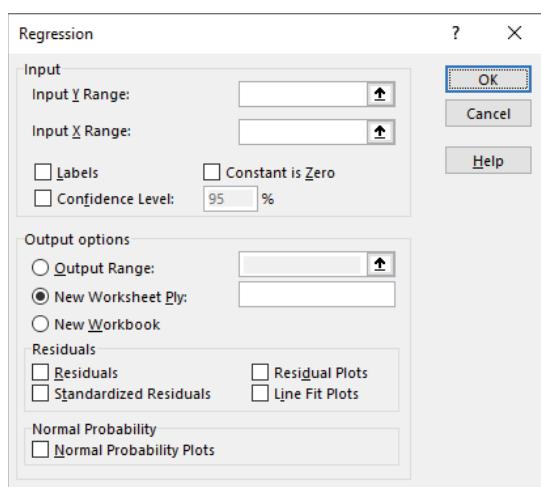
The *rank and percentile* tool creates a table that shows the ordinal and percentile ranking for each value in a range. You can also generate ranks and percentiles by using Excel functions (those that begin with RANK and PERCENTILE).

Regression

Use the *regression* tool (see Figure 30.9) to calculate a regression analysis from worksheet data. You can use regression to analyze trends, forecast the future, build predictive models, and, often, make sense out of a series of seemingly unrelated numbers.

FIGURE 30.9

The Regression dialog box



Regression analysis enables you to determine the extent to which one range of data (the dependent variable) varies as a function of the values of one or more other ranges of data (the independent variables). This relationship is expressed mathematically, using values that Excel calculates. You can use these calculations to create a mathematical model of the data and predict the dependent variable by using different values of one or more independent variables. This tool can perform simple and multiple linear regressions and calculate and standardize residuals automatically.

As you can see, the Regression dialog box offers many options:

Input Y Range The range that contains the dependent variable.

Input X Range One or more ranges that contain independent variables.

Constant Is Zero If selected, forces the regression to have a constant of 0 (which means that the regression line passes through the origin; when the X values are 0, the predicted Y value is 0).

Confidence Level The confidence level for the regression.

Residuals The four options in this section of the dialog box enable you to specify whether to include residuals in the output. Residuals are the differences between observed and predicted values.

Normal Probability Generates a chart for normal probability plots.

Sampling

The *sampling* tool generates a random sample from a range of input values. The sampling tool can help you work with a large database by creating a subset of it.

This procedure has two options: periodic and random. If you choose a periodic sample, Excel selects every nth value from the Input range, where n equals the period that you specify. With a random sample, you simply specify the size of the sample you want Excel to select; every value has an equal probability of being chosen.

t-Test

Use the *t-Test* tool to determine whether a statistically significant difference exists between two small samples. The Analysis ToolPak can perform three types of t-Tests:

Paired two-sample for means For paired samples in which you have two observations on each subject (such as a pretest and a posttest). The samples must be the same size.

Two-sample assuming equal variances For independent rather than paired samples. Excel assumes equal variances for the two samples.

Two-sample assuming unequal variances For independent rather than paired samples. Excel assumes unequal variances for the two samples.

Figure 30.10 shows output for the Paired Two Sample for Means t-Test. You specify the significance level (alpha) and the hypothesized difference between the two means (that is, the null hypothesis).

Part IV: Managing and Analyzing Data

FIGURE 30.10

Output from the paired t-Test dialog box

Student#	Pretest	Posttest	t-Test: Paired Two Sample for Means			
				Pretest	Posttest	
1	1	71	69			
2	2	63	61			
3	3	68	70			
4	4	67	68			
5	5	66	61			
6	6	63	60			
7	7	76	83			
8	8	70	72			
9	9	69	71			
10	10	73	77			
11	11	71	72			
12	12	66	66			
13	13	70	71			
14	14	78	86			
15	15	68	70			
16	16	75	78			
17	17	65	63			
18	18	67	68			
19	19	71	77			
20	20	74	78			
21	21	71	72			
22						
23						

z-Test (two-sample test for means)

The t-Test tool is used for small samples; the z-Test tool is used for larger samples or populations. You must know the variances for both input ranges. Generally, use a t-Test when your sample size is less than 30 and you don't know the standard deviation of the population. Otherwise, use the z-Test.

Protecting Your Work

IN THIS CHAPTER

- Protecting worksheets
- Protecting workbooks
- Protecting Visual Basic Projects
- Creating PDFs and checking documents

The concept of “protection” gets a lot of attention in the Excel forums. It seems that many users want to learn how to protect various workbook elements from being overwritten or copied. Excel has several protection-related features, and those features are covered in this chapter.

Types of Protection

Excel’s protection-related features fall into three categories:

- Worksheet protection** Protecting all or part of a worksheet from being modified or restricting the modifications to certain users
- Workbook protection** Protecting a workbook from having sheets inserted or deleted and requiring the use of a password to open the workbook
- Visual Basic (VB) protection** Using a password to prevent others from viewing or modifying your VBA code

CAUTION

Before we discuss these features, you should understand the limitations of security. Using a password to protect some aspect of your work doesn’t guarantee that it’s secure. Password-cracking utilities (and some simple tricks) have been around for a long time. Using passwords works in the majority of cases, but if unscrupulous individuals are truly intent on getting to your data, they can usually find a way. If security is critical, perhaps Excel isn’t the proper tool.

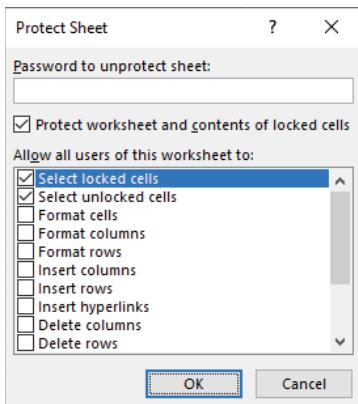
Protecting a Worksheet

You may want to protect a worksheet for a variety of reasons. One common reason is to prevent yourself or others from accidentally deleting formulas or other critical data. A typical scenario is to protect a worksheet so that the data can be changed but the formulas can't be.

To protect a worksheet, activate the worksheet and choose Review \Rightarrow Protect \Rightarrow Protect Sheet. Excel displays the Protect Sheet dialog box, shown in Figure 31.1. Providing a password is optional. If you enter a password, that password will be required to unprotect the worksheet. If you accept the default options in the Protect Sheet dialog box (and if you haven't unlocked any cells), none of the cells on the worksheet can be modified.

FIGURE 31.1

Use the Protect Sheet dialog box to protect a worksheet.



To unprotect a protected sheet, choose Review \Rightarrow Protect \Rightarrow Unprotect Sheet. If the sheet was protected with a password, you're prompted to enter that password.

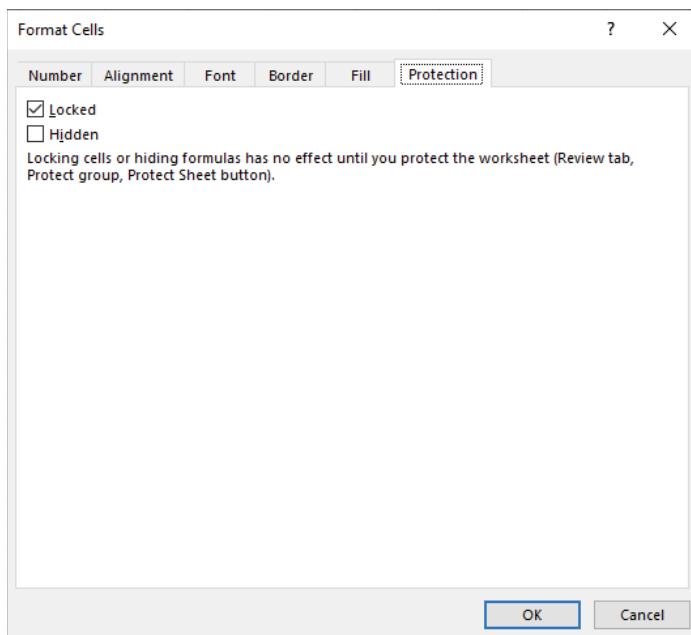
Unlocking cells

In many cases, you'll want to allow some cells to be changed when the worksheet is protected. For example, your worksheet may have some input cells that are used by formula cells. In such a case, you would want the user to be able to change the input cells but not the formula cells. Every cell has a Locked attribute, and that attribute determines whether the cell can be changed when the sheet is protected.

By default, all cells are locked, and worksheets are unprotected. To change the Locked attribute, select the cell or range, right-click, and choose Format Cells from the shortcut menu (or press **Ctrl+1**). Select the Protection tab of the Format Cells dialog box (see Figure 31.2), clear the Locked check box, and then click OK.

FIGURE 31.2

Use the Protection tab in the Format Cells dialog box to change the Locked attribute of a cell or range.

**NOTE**

The Protection tab of the Format Cells dialog box has another attribute: Hidden. If this check box is selected, the contents of the cell don't appear in the Formula bar when the sheet is protected. The cell isn't hidden in the worksheet. You may want to set the Hidden attribute for formula cells to prevent users from seeing the formula when the cell is selected.

After you unlock the desired cells, choose Review \Rightarrow Protect \Rightarrow Protect Sheet to protect the sheet. Then you can change the unlocked cells, but if you attempt to change a locked cell, Excel displays the message shown in Figure 31.3.

NOTE

The Locked attribute matters only when the worksheet is protected. On an unprotected worksheet, locked and unlocked cells behave the same way.

Part IV: Managing and Analyzing Data

FIGURE 31.3

Excel warns you if you attempt to change a locked cell.



Sheet protection options

The Protect Sheet dialog box (refer to Figure 31.1) has several options that determine what the user can do when the worksheet is protected:

Select Locked Cells Allows the user to select locked cells using the mouse or the keyboard (but not change them). This setting is enabled by default.

Select Unlocked Cells Allows the user to select unlocked cells using the mouse or the keyboard. This setting is enabled by default.

Format Cells Allows the user to apply formatting.

Format Columns Allows the user to hide or change the width of columns.

Format Rows Allows the user to hide or change the height of rows.

Insert Columns Allows the user to insert new columns.

Insert Rows Allows the user to insert new rows.

Insert Hyperlinks Allows the user to insert hyperlinks (even in locked cells).

Delete Columns Allows the user to delete columns.

Delete Rows Allows the user to delete rows.

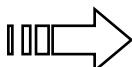
Sort Allows the user to sort data in a range as long as the range doesn't contain locked cells.

Use AutoFilter Allows the user to use existing AutoFilters.

Use PivotTable and PivotChart Allows the user to change the layout of PivotTables or create new PivotTables. This setting also applies to PivotCharts.

Edit Objects Allows the user to make changes to objects (such as shapes) and charts, as well as insert or delete notes.

Edit Scenarios Allows the user to use scenario management features.



See Chapter 28, “Performing Spreadsheet What-If Analysis,” for more on creating and using scenarios.

TIP

When the worksheet is protected and the Select Unlocked Cells option is selected, pressing Tab moves to the next unlocked cell (and skips locked cells), making data entry much easier.

Assigning user permissions

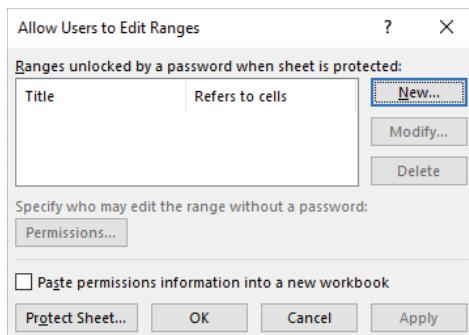
Excel also offers the ability to assign user-level permissions to different areas on a protected worksheet. You can specify which users can edit a particular range while the worksheet is protected. As an option, you can require a password to make changes.

This feature is rarely used, and the setup procedure is rather complicated. But if you need this level of protection, setting it up might be worth the effort.

1. **Unprotect the worksheet if it's protected.**
2. **Choose Review ➔ Protect ➔ Allow Edit Ranges.** The Allow Users to Edit Ranges dialog box appears, as shown in Figure 31.4.

FIGURE 31.4

The Allow Users to Edit Ranges dialog box



3. **Click the New button to show the New Range dialog box.**
4. **Fill out the Title, Refers to cells, and Range password fields, and then click the Permissions button.**
5. **Click the Add button to display the Select Users or Groups dialog box.**
6. **Type the Windows or domain username and click OK.**
7. **Accept the default Allow setting or change it to Deny and click OK.**
8. **Dismiss the remaining dialog boxes by clicking OK.**
9. **Protect the sheet.** You can click the Protect Sheet button on the Allow Users to Edit Ranges dialog box before you dismiss it or use the Protect Sheet command on the Review tab of the Ribbon.

Protecting a Workbook

Excel provides two ways to protect a workbook:

- Require a password to open the workbook.
- Prevent users from inserting sheets, deleting sheets, hiding sheets, and unhiding sheets.

These are not mutually exclusive, so both methods can be applied to a workbook. We discuss each of these methods in the sections that follow.

Requiring a password to open a workbook

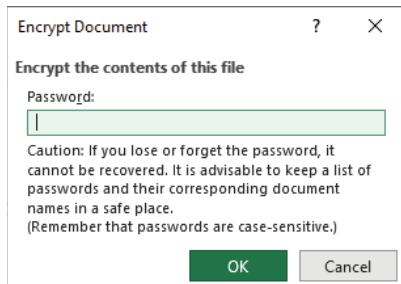
Excel lets you save a workbook with a password. After you do so, whoever tries to open the workbook must enter the password.

To add a password to a workbook, follow these steps:

1. Choose **File** \Rightarrow **Info** \Rightarrow **Protect Workbook** \Rightarrow **Encrypt with Password**. The Encrypt Document dialog box, shown in Figure 31.5, appears.

FIGURE 31.5

Specify a workbook password in the Encrypt Document dialog box.



2. Type a password and click **OK**.
3. Type the password again and click **OK**.
4. Save the workbook.

NOTE

You need to perform these steps only once. You don't need to specify the password every time you save the workbook.

To remove a password from a workbook, repeat the same procedure. In step 2, however, delete the existing password symbols from the Encrypt Document dialog box, click **OK**, and save your workbook.

Figure 31.6 shows the Password dialog box that appears when you try to open a file saved with a password.

FIGURE 31.6

Opening this workbook requires a password.



Excel provides another way to add a password to a document:

1. **Choose File ➔ Save As and click Browse.** The Save As dialog box appears.
2. **Click the Tools drop-down and choose General Options.** The General Options dialog box appears.
3. **Enter a password in the Password to Open field.**
4. **Click OK.** You're asked to re-enter the password before you return to the Save As dialog box.
5. **In the Save As dialog box, make sure that the filename, location, and type are correct; then click Save.**

Note

The General Options dialog box has another password field: Password to Modify. If you specify a password for this field, the file opens in read-only mode (it can't be saved under the same name) unless the user knows the password. If you use the Read-Only Recommended check box without a password, Excel suggests that the file be opened in read-only mode, but the user can override this suggestion.

Protecting a workbook's structure

To prevent others (or yourself) from performing certain actions in a workbook, you can protect the workbook's structure. When a workbook's structure is protected, the user may not do the following:

- Insert a sheet
- Delete a sheet
- Hide a sheet
- Unhide a sheet
- Rename a sheet
- Move a sheet

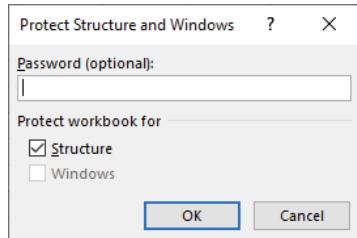
Part IV: Managing and Analyzing Data

To protect a workbook's structure, follow these steps:

1. Choose **Review** \Rightarrow **Protect** \Rightarrow **Protect Workbook**. The Protect Structure and Windows dialog box appears, as shown in Figure 31.7.

FIGURE 31.7

The Protect Structure and Windows dialog box



2. Select the **Structure** check box.
3. (Optional) Enter a password.
4. Click **OK**.

The Protect Workbook Ribbon tool changes color to indicate that you enabled workbook protection. To unprotect the workbook's structure, choose Review \Rightarrow Protect \Rightarrow Protect Workbook again. Workbook protection will be turned off, and the Ribbon tool will change back to its normal color. If the workbook's structure was protected with a password, you're prompted to enter the password.

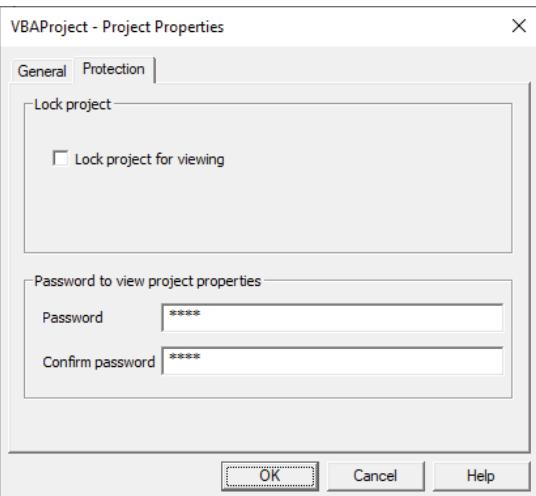
Protecting a VBA Project

If your workbook contains VBA macros, you may want to protect the VBA Project to prevent others from viewing or modifying your macros. To protect a VBA Project, follow these steps:

1. Press **Alt+F11** to activate the **VB Editor**.
2. Select your project in the **Projects** window.
3. Choose **Tools** \Rightarrow <Project Name> **Properties** (where <Project Name> corresponds to your Project name). The Project Properties dialog box appears.
4. Select the **Protection** tab (see Figure 31.8).
5. Select the **Lock Project for Viewing** check box.
6. Enter a password (twice).
7. Click **OK** and then save your file. When the file is closed and then reopened, a password will be required to view or modify the VBA code.

FIGURE 31.8

Protecting a VBA Project with a password



Part VI, "Automating Excel," discusses VBA macros.

Related Topics

In addition to Excel's built-in security features, there are other features and techniques related to protecting your Excel content.

Saving a worksheet as a PDF file

The Portable Document Format (PDF) file format is widely used as a way to present information in a read-only manner with precise control over the layout. Software to display PDF files is available from a number of sources. Excel can create PDF files, but it can't open them.

XPS is another "electronic paper" format, developed by Microsoft as an alternative to the PDF format. At this time, there is little third-party support for the XPS format.

Save a worksheet in PDF or XPS format by choosing File \Rightarrow Export \Rightarrow Create PDF/XPS Document \Rightarrow Create PDF/XPS. The Publish as PDF or XPS dialog box appears; here, you can specify a file-name and location and set some other options.

Marking a workbook as final

Excel lets you mark a workbook as final. This action makes two changes to the workbook:

- It makes the workbook read-only so that the file can't be saved using the same name.
- It makes the workbook view-only so that nothing may be changed.

Part IV: Managing and Analyzing Data

When you open a finalized document, you see a message below the Ribbon. You can override its final status by clicking the Edit Anyway button.

To finalize a workbook, choose File \Rightarrow Info \Rightarrow Protect Workbook \Rightarrow Mark as Final. A dialog box appears, where you can confirm your choice. If you click the OK button, another message box appears explaining what marking a workbook as final means. Once the workbook is marked, the status bar shows a sheet icon with a check mark to indicate the active workbook is marked as final.

CAUTION

Marking a document as final is not a robust security measure. Anyone who opens the workbook can cancel the mark-as-final designation. Therefore, this method doesn't guarantee that others will not change the workbook.

Inspecting a workbook

If you plan to distribute a workbook to others, you may want to have Excel check the file for hidden data and personal information. This tool can locate hidden information about you, your organization, or the workbook. In some situations, you may not want to share this information with others.

To inspect a workbook, choose File \Rightarrow Info \Rightarrow Check for Issues \Rightarrow Inspect Document. The Document Inspector dialog box appears, as shown in Figure 31.9. Click Inspect, and Excel displays the results of the inspection and gives you the opportunity to remove the items it finds.

CAUTION

If Excel identifies items in the Document Inspector, it doesn't necessarily mean that they should be removed. In other words, you shouldn't blindly use the Remove All buttons to remove the items that Excel locates. For example, you may have a hidden sheet that serves a critical purpose. Excel will identify that hidden sheet and make it easy for you to delete it. To be on the safe side, always make a backup copy of your workbook before running the Document Inspector.

Two other commands are available in the File \Rightarrow Info \Rightarrow Check for Issues menu:

Check Accessibility Checks the workbook for content that people with disabilities might find difficult to read

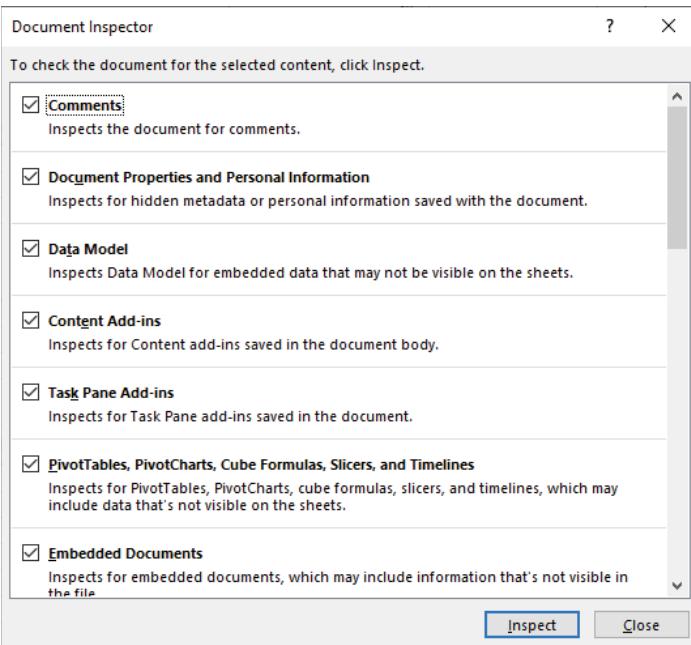
Check Compatibility Checks the workbook for the presence of features that may not work in previous versions of Excel



See Chapter 6, “Understanding Excel Files and Templates,” for more about checking file compatibility.

FIGURE 31.9

The Document Inspector dialog box identifies hidden and personal information in a workbook.



Using a digital signature

Excel lets you add a digital signature to a workbook. Using a digital signature is somewhat analogous to signing a paper document. A digital signature helps to ensure the authenticity of the workbook and ensure that the content hasn't been modified since it was signed.

After you sign a workbook, the signature is valid until you make changes and resave the file.

Getting a digital ID

To sign a workbook digitally, you must obtain a certificate from a certified authority that is able to verify the authenticity of your signature. Prices vary, depending on the certificate-granting company.

Signing a workbook

Excel supports two types of digital signatures: a visible signature and an invisible signature.

To add a visible digital signature, choose Insert \Rightarrow Text \Rightarrow Signature Line \Rightarrow Microsoft Office Signature Line. The Signature Setup dialog box appears, and you're prompted for the

Part IV: Managing and Analyzing Data

information for the signature. After you add the signature box, double-click it, and the Sign dialog box appears; here you actually sign the document either by typing your name or by uploading a scanned image of your signature. After signing the document, it will be marked as final. Any change to the file will invalidate the signature.

To add an invisible digital signature, choose File \Rightarrow Info \Rightarrow Protect Workbook \Rightarrow Add a Digital Signature. If the signed workbook is changed in any way, the digital signature is invalidated. Use a digital signature if you need to verify the integrity of the workbook but don't need a visible indication of that integrity for the user.

Part V

Understanding Power Pivot and Power Query

Part V introduces you to the world of Power Pivot and Power Query. Here, you'll discover how to develop powerful reporting with external data and the Power Pivot data model. You'll also explore how the rich set of tools Power Query offers can help you save time, automate data clean-up, and substantially enhance your data analysis and reporting capabilities.

IN THIS PART

Chapter 32
Introducing Power Pivot

Chapter 33
Working Directly with the Internal Data Model

Chapter 34
Adding Formulas to Power Pivot

Chapter 35
Introducing Power Query

Chapter 36
Transforming Data with Power Query

Chapter 37
Making Queries Work Together

Chapter 38
Enhancing Power Query Productivity

Introducing Power Pivot

IN THIS CHAPTER

- Getting started with Power Pivot
- Linking to Excel data
- Managing relationships
- Loading data from other data sources

Recognizing the importance of the business intelligence (BI) revolution and the place Excel holds within it, Microsoft has made substantial investments in improving Excel's BI capabilities, specifically focusing on Excel's self-service BI capabilities and its ability to manage and analyze information better from the increasing number of available data sources.

The key product of this endeavor is Power Pivot. With Power Pivot comes the ability to set up relationships between large, disparate data sources and to merge data sources with hundreds of thousands of rows into one analytical engine within Excel.

In this chapter, you'll get an overview of those capabilities, exploring the key features, benefits, and capabilities of Power Pivot.

Understanding the Power Pivot Internal Data Model

At its core, Power Pivot is essentially a SQL Server Analysis Services engine made available through an in-memory process that runs directly within Excel. The technical name for this engine is the *xVelocity analytics engine*. However, in Excel, it's referred to as the internal *data model*.

Every Excel workbook contains an internal data model, which is a single instance of the Power Pivot in-memory engine. The following sections outline how to leverage the Power Pivot internal data model to import and integrate disparate data sources.

The Power Pivot Ribbon

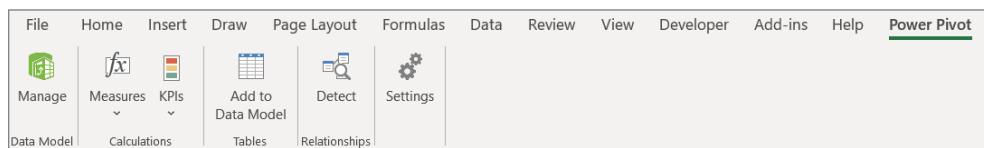
Anyone with an ongoing subscription to Office 365 Pro Plus will be able to see and activate the Power Pivot tab (see Figure 32.1). Those who are using any stand-alone version of Excel 2013 or later will also have access to the Power Pivot tab.

Note that organizations often install Excel in accordance to their own installation policies. In some organizations, Excel is installed without the Power Pivot add-in activated so the Power Pivot tab will not be visible. If you don't see the Power Pivot tab shown in Figure 32.1, you can follow these steps to activate it:

1. **Go up to the Excel Ribbon and click File>Options.**
2. **Choose the Add-ins option on the left and then select COM Add-ins from the Manage drop-down selector.**
3. **In the list of available COM add-ins, check the box next to Microsoft Office Power Pivot for Excel; then click OK.**
4. **Close and restart Excel.**

FIGURE 32.1

The Power Pivot Ribbon interface



A Word on Compatibility

Since Excel 2010 was released, Microsoft has made several versions of Power Pivot available. You'll need to be aware that there are different versions of Power Pivot being used, depending on the version of Excel. Be careful when sharing Power Pivot workbooks in environments where some of your audience is using earlier versions of Excel while others are using more recent versions of Excel. Opening and refreshing a workbook that contains a Power Pivot model created with an older version of the Power Pivot add-in will trigger an automatic upgrade of the underlying model. Once this happens, users with older versions of Excel will no longer be able to use the Power Pivot model in the workbook.

As a general rule, Power Pivot workbooks created in a version of Excel that is equal to or less than your version should give you no problems. However, you will not be able to use Power Pivot workbooks created in a later version of Excel than your version.

Linking Excel tables to Power Pivot

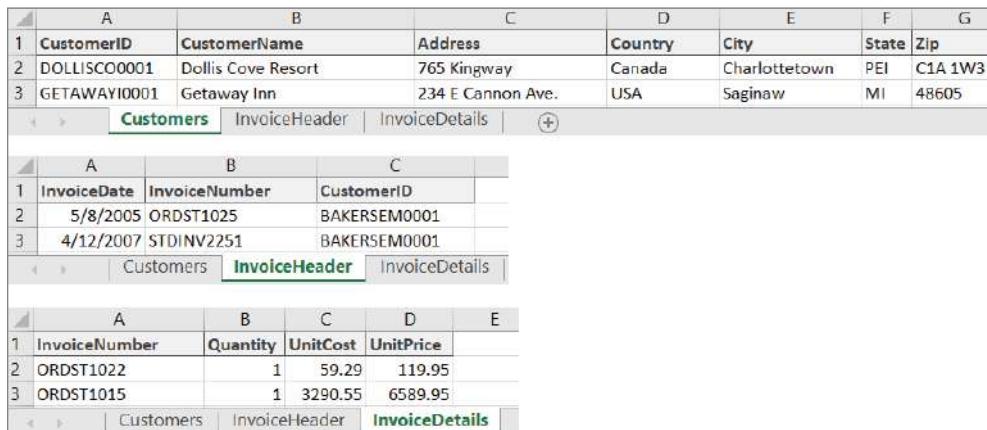
The first step in using Power Pivot is to fill it with data. You can either import data from external data sources or link to Excel tables in your current workbook. For now, let's start this walk-through by linking three Excel tables to Power Pivot.

 Most of the examples in this chapter are available on this book's website at www.wiley.com/go/excel365bible. The filename is Power Pivot Intro.xlsx.

In this scenario, we have three data sets in three different worksheets (see Figure 32.2): Customers, Invoice Header, and Invoice Details.

FIGURE 32.2

We want to use Power Pivot to analyze the data in the Customers, Invoice Header, and Invoice Details worksheets.



The screenshot shows the Power Pivot ribbon with three tabs: Customers, InvoiceHeader, and InvoiceDetails. The Customers tab is currently selected. Below the ribbon, there are three separate tables:

- Customers:** Contains columns CustomerID, CustomerName, Address, Country, City, State, and Zip. Data rows include DOLLISCO0001 (Dollis Cove Resort, Canada, PEI, C1A 1W3) and GETAWAY10001 (Getaway Inn, USA, MI, 48605).
- InvoiceHeader:** Contains columns InvoiceDate, InvoiceNumber, and CustomerID. Data rows include 5/8/2005 ORDST1025 (BAKERSEM0001) and 4/12/2007 STDINV2251 (BAKERSEM0001).
- InvoiceDetails:** Contains columns InvoiceNumber, Quantity, UnitCost, and UnitPrice. Data rows include ORDST1022 (1, 59.29, 119.95) and ORDST1015 (1, 3290.55, 6589.95).

The Customers data set contains basic information like Customer ID, Customer Name, Address, and so forth. The Invoice Header data set contains data that points specific invoices to specific customers. The Invoice Details data set contains the specifics of each invoice.

We want to analyze revenue by customer and month. It's clear that we'll somehow need to join these three tables together before we can do our analysis. In the past, we would have to go through a series of gyrations involving VLOOKUPS or other clever formulas. But with Power Pivot, we can build these relationships in just a few clicks.

Preparing your Excel tables

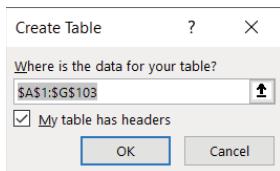
When linking Excel data to Power Pivot, it's a best practice first to convert your Excel data to explicitly named tables. Although not technically necessary, giving your tables friendly names helps track and manage your data in the Power Pivot data model. If you don't convert your data to tables first, Excel will do it for you and give your tables useless names like Table1, Table2, and so on.

Follow these steps to convert each data set into an Excel table:

1. **Go to the Customers tab and click anywhere inside the data range.**
2. **Press Ctrl+T on your keyboard.** This will activate the Create Table dialog box, shown in Figure 32.3.

FIGURE 32.3

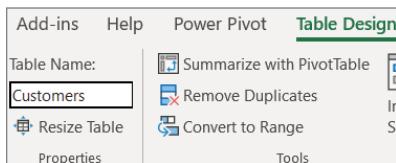
Convert your data range into an Excel table.



3. **In the Create Table dialog box, make sure the range for the table is correct and that the check box next to the "My table has headers" option is selected. Click the OK button.**
4. **You should now see a Table Design tab on the Ribbon.** Click that tab and use the Table Name input to give your table a friendly name (see Figure 32.4). This will ensure that you will be able to recognize the table when adding it to the internal data model.

FIGURE 32.4

Give your newly created Excel table a friendly name.



5. **Repeat steps 1–4 for the Invoice Header and Invoice Details data sets.**

Adding your Excel tables to the data model

Once you've converted your data to Excel tables, you're ready to add them to the Power Pivot data model. Follow these steps to add your newly created Excel tables to the data model using the Power Pivot tab:

1. Click any cell inside your Customers Excel table.
2. Click the Power Pivot tab on the Ribbon, and click the Add to Data Model command. Power Pivot creates a copy of your table and activates the Power Pivot window (shown in Figure 32.5).

FIGURE 32.5

The Power Pivot window shows all the data that currently exists in your data model.

CustomerID	CustomerName	Address	Country	City	State	Zip
1 DOLLISCO0001	Dollis Cove Resort	765 King...	Canada	Charl...	PEI	C1A ...
2 GETAWAYI0001	Getaway Inn	234 E Can...	USA	Sagin...	MI	48605
3 HOMEFURNO...	Home Furnishing...	234 Herit...	USA	Midla...	MI	4864...
4 JOHNSONKOO...	Johnson, Kimberly	5678 S. 4...	USA	Rockf...	IL	6112...
5 KELLYCONO001	Kelly Consulting	123 Yeo	Australia	Melb...	VIC	3000
6 KENSINGT0001	Kensington Garde...	12345 Re...	USA	Milw...	WI	5320...
7 HAMPTONV00...	Hampton Village ...	234 Hamp...	USA	Sprin...	IL	6270...
8 HEALTHYCO001	Healthy Concepts	1234 Wes...	USA	West ...	IA	5062...
9 LECLERC0001	LeClerc & Associa...	4321 Wes...	Canada	Mont...	PQ	H4Q ...
10 LEISURET0001	Leisure & Travel C...	City 123	Australia	Sydney	NSW	2000
11 LONDONBEO0...	Londonberry Nur...	987 Porta...	New Zeala...	Auckl...		

Although the Power Pivot window looks like Excel, it's actually a separate program altogether. You'll notice that the grid for the Customers table has row numbers but no column references. You'll also notice that you can't edit the data within the table. This data is simply a snapshot of the actual Excel table you imported.

Additionally, if you look at your Windows taskbar at the bottom of your screen, you'll see that Power Pivot has a separate window from Excel. You can switch between Excel and the Power Pivot window by clicking each respective program in the taskbar.

Repeat steps 1 and 2 for your other Excel tables: Invoice Header and Invoice Details. Once you have imported all your Excel tables into the data model, your Power Pivot window will show each data set on its own tab, as shown in Figure 32.6.

Part V: Understanding Power Pivot and Power Query

FIGURE 32.6

Each table you add to the data model will be placed on its own tab in Power Pivot.

CustomerID	CustomerName	Address	Country	City	State	Zip
1 DOLLISCO0001	Dollis Cove Resort	765 King...	Canada	Charl...	PEI	C1A ...
2 GETAWAY0001	Getaway Inn	234 E Can...	USA	Sagin...	MI	48605
3 HOMEFURNO...	Home Furnishing...	234 Herit...	USA	Midla...	MI	4864...
4 JOHNSONKOO...	Johnson, Kimberly	5678 S. 4...	USA	Rockf...	IL	6112...
5 KELLYCON0001	Kelly Consulting	123 Yeo	Australia	Melb...	VIC	3000
6 KENSINGT0001	Kensington Garde...	12345 Re...	USA	Milw...	WI	5320...
7 HAMPTONVO...	Hampton Village ...	234 Hamp...	USA	Sprin...	IL	6270...
8 HEALTHC0001	Healthy Concepts	1234 Wes...	USA	West ...	IA	5062...
9 LECLERC0001	LeClerc & Associa...	4321 Wes...	Canada	Mont...	PQ	H4Q ...
10 LEISURET0001	Leisure & Travel C...	City 123	Australia	Sydney	NSW	2000
11 LONDONBE00...	Londonberry Nur...	987 Porta...	New Zeala...	Auckl...		

NOTE

Because the data that you just imported into Power Pivot comes from an Excel table within the current workbook, Power Pivot will consider these linked tables. This means even though the data shown in Power Pivot is a snapshot at the time you added it, the data automatically updates when you edit the source table in Excel. Linked tables are the only kind of data source that automatically refreshes as the data changes. As you'll learn later in this chapter, data from outside sources will need to be explicitly refreshed.

Creating relationships between your PowerPivot tables

At this point in our walk-through, Power Pivot knows that we have three tables in the data model, but it has no idea how these three tables relate to one another. We'll need to connect these tables by defining relationships between the Customers, Invoice Header, and Invoice Details tables. We can do so directly within the Power Pivot window.

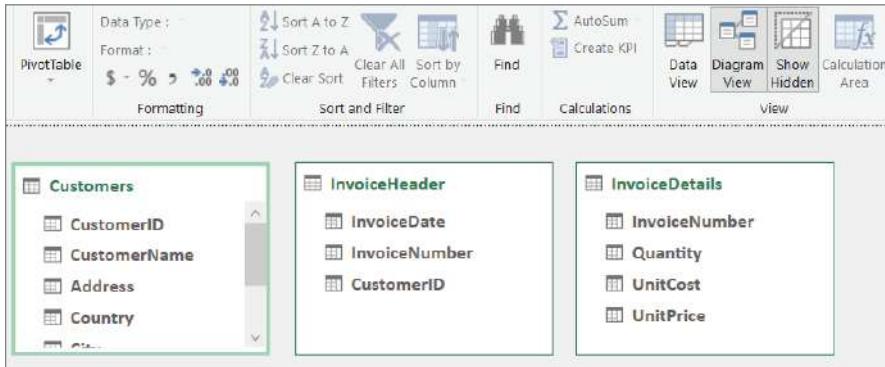
TIP

If you inadvertently closed the Power Pivot window, you can easily call it back up by clicking the Manage command button on the Power Pivot Ribbon tab. Alternatively, you can click the Manage Data Model command on the Data tab.

1. Activate the Power Pivot window and click the **Diagram View** command button on the **Home** tab. Power Pivot will display a screen that shows a visual representation of all the tables in the data model (see Figure 32.7).

FIGURE 32.7

The Diagram view allows you to see all the tables in your data model.

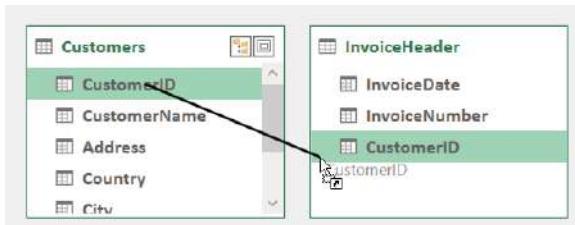


You can move the tables in the Diagram view around simply by clicking and dragging their title bars. The idea is to identify the primary index keys in each table and connect them. In this scenario, the Customers table and the Invoice Header table can be connected using the CustomerID field. The Invoice Header and Invoice Details tables can be connected using the InvoiceNumber field.

2. Click and drag a line from the **CustomerID** field in the **Customers** table to the **CustomerID** field in the **Invoice Header** table (as demonstrated in Figure 32.8).

FIGURE 32.8

To create a relationship, you simply click and drag a line between the fields in your tables.



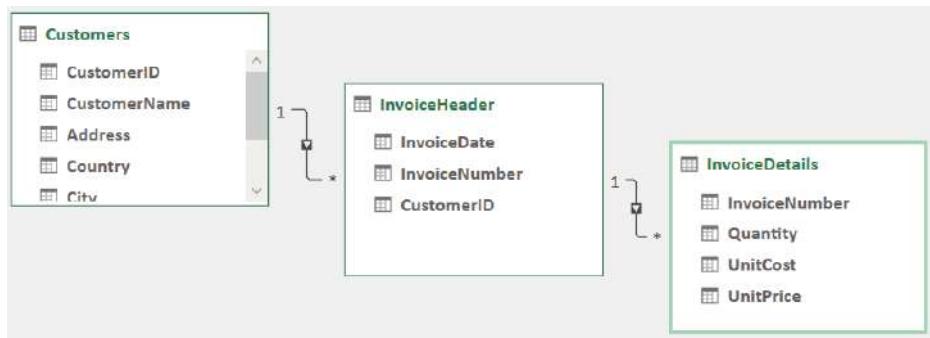
Part V: Understanding Power Pivot and Power Query

3. Click and drag a line from the **InvoiceNumber** field in the **Invoice Header** table to the **InvoiceNumber** field in the **Invoice Details** table.

At this point, your diagram will look similar to Figure 32.9. Notice that Power Pivot shows a line between the tables that you just connected. In database speak, these are referred to as *joins* or *join lines*.

FIGURE 32.9

When you create relationships, the Power Pivot diagram will show join lines between your tables.



The joins in Power Pivot are one-to-many joins. This means that when a table is joined to another, one of the tables has unique records with unique index numbers (CustomerID, for example), while the other can have many records where index numbers are duplicated.

Notice that the join lines have arrows pointing from a table to another table. The arrows in these join lines will always point to the table that has the duplicated indexes.

In this case, the Customers table contains a unique list of customers, each having its own unique identifier. No CustomerID in that table is duplicated. The Invoice Header table has many rows for each CustomerID; each customer can have many invoices.

TIP

To close the diagram and get back to seeing the data tables, click the Data View command button in the Power Pivot window.

Managing existing relationships

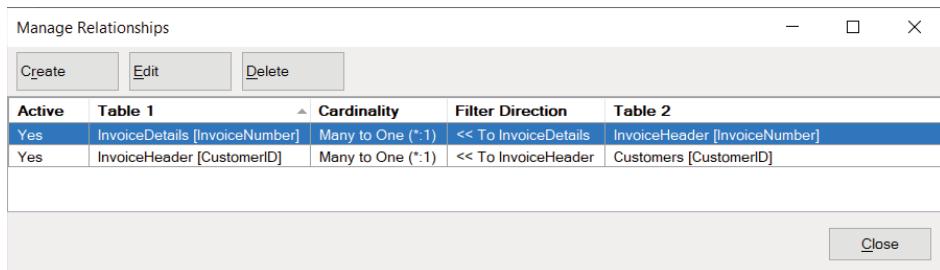
If you need to go back and edit or delete a relationship between two tables in your data model, you can do so by following these steps:

1. Activate the Power Pivot window, select the Design tab, and then click the Manage Relationships command.

2. In the Manage Relationships window, shown in Figure 32.10, click the relationship with which you want to work and select Edit or Delete.

FIGURE 32.10

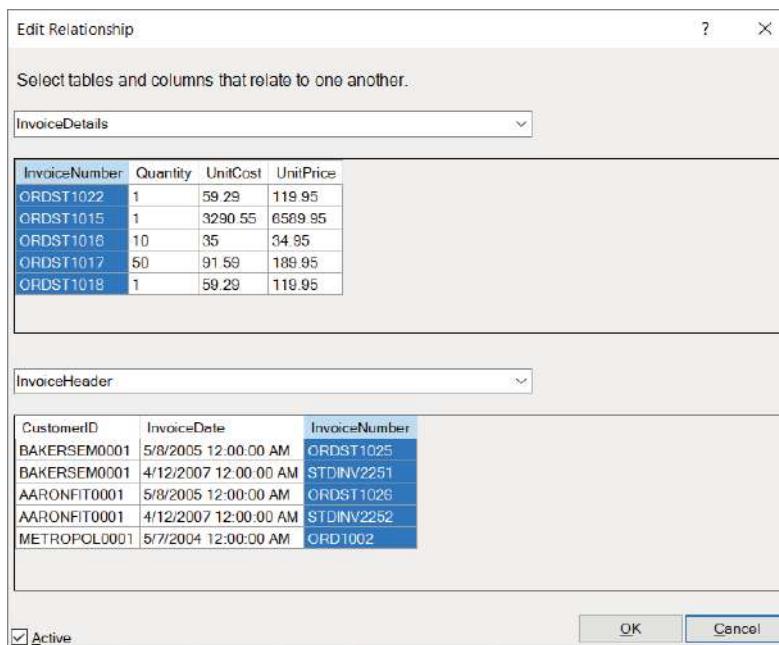
Use the Manage Relationships window to edit or delete existing relationships.



3. Clicking Edit will open the Edit Relationship dialog box, shown in Figure 32.11. You'll notice the columns used to form the relationship are highlighted. Here, you can redefine the relationship by simply selecting the appropriate columns. You can also use the Active check box control to disable or enable the relationship.

FIGURE 32.11

Use the Edit Relationship dialog box to adjust the tables and field names that define the selected relationship.



Part V: Understanding Power Pivot and Power Query

Using Power Pivot data in reporting

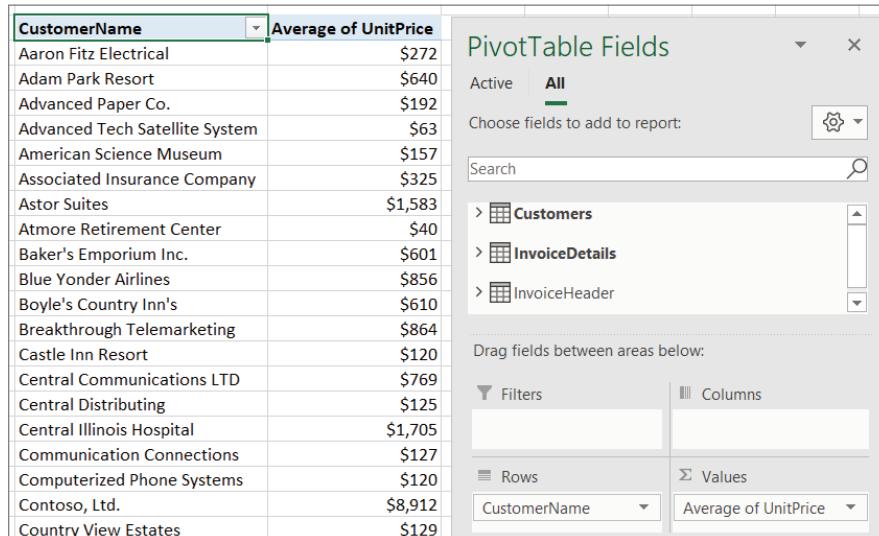
Once you define the relationships in your Power Pivot data model, it's essentially ready for action. In terms of Power Pivot, *action* basically means analysis with a PivotTable. In fact, all Power Pivot data is presented through the framework of PivotTables (PivotTables are covered in Chapter 26, "Introducing PivotTables," and Chapter 27, "Analyzing Data with PivotTables").

1. **Activate the Power Pivot window, select the Home tab, and then click the PivotTable command button.**
2. **Specify whether you want the PivotTable placed on a new worksheet or an existing sheet.**
3. **Build out your needed analysis just as you would any other standard PivotTable using the PivotTable Fields list.**

Notice in Figure 32.12 that our newly created PivotTable contains all the tables in our data model. Unlike a standard PivotTable where you can use fields from only one table, the relationships we've defined in our model allow us to use any of the fields from any of the tables. With this configuration, we essentially have a powerful cross-table analytical engine in the form of a familiar PivotTable. In this example, we are calculating the average unit price by customer.

FIGURE 32.12

You now have a Power Pivot–driven PivotTable that aggregates data across multiple tables.



In the days before Power Pivot, this analysis would have been tricky. You would have had to build VLOOKUP formulas to get from Customers to Invoice Header and then another set

of VLOOKUPS to get from Invoice Header to Invoice Details. After all that formula building, you still would have had to find a way to aggregate the data to average the unit price per customer.

Loading Data from Other Data Sources

As you'll discover in the following sections, you're not limited to using only the data that already exists in your Excel workbook. Power Pivot has the ability to reach outside the workbook and import data found in external data sources. Indeed, what makes Power Pivot so powerful is its ability to consolidate data from disparate data sources and build relationships between them. This means you can theoretically create a Power Pivot data model that contains some data from a SQL Server table, some data from a Microsoft Access database, and even data from a one-off text file.

32

Loading data from relational databases

One of the more common data sources used by Excel analysts are relational databases. It's not difficult to find an analyst who frequently uses data from Microsoft Access, SQL Server, or Oracle databases. In the following sections, we'll walk through the steps for loading data from external database systems.

Loading data from SQL Server

SQL Server databases are some of the most commonly used for storing enterprise-level data. Most SQL Server databases are managed and maintained by the IT department. To connect to a SQL Server database, you'll have to work with your IT department to obtain read access to the database from which you're trying to pull data.

Once you have access to the database, open Excel and activate the Power Pivot window by selecting Power Pivot \Rightarrow Manage.

Once the Power Pivot window is activated, select the From Other Sources command button on the Home tab. This will activate the Table Import Wizard dialog box, shown in Figure 32.13. Here, select the Microsoft SQL Server option and then click the Next button.

The Table Import Wizard will now ask for all the information it needs to connect to your chosen data source. This includes things such as server address, login credentials, and any other database name. The wizard will show you different fields based on the type of data source you select. The more common fields you'll see when connecting to an external data source are as follows:

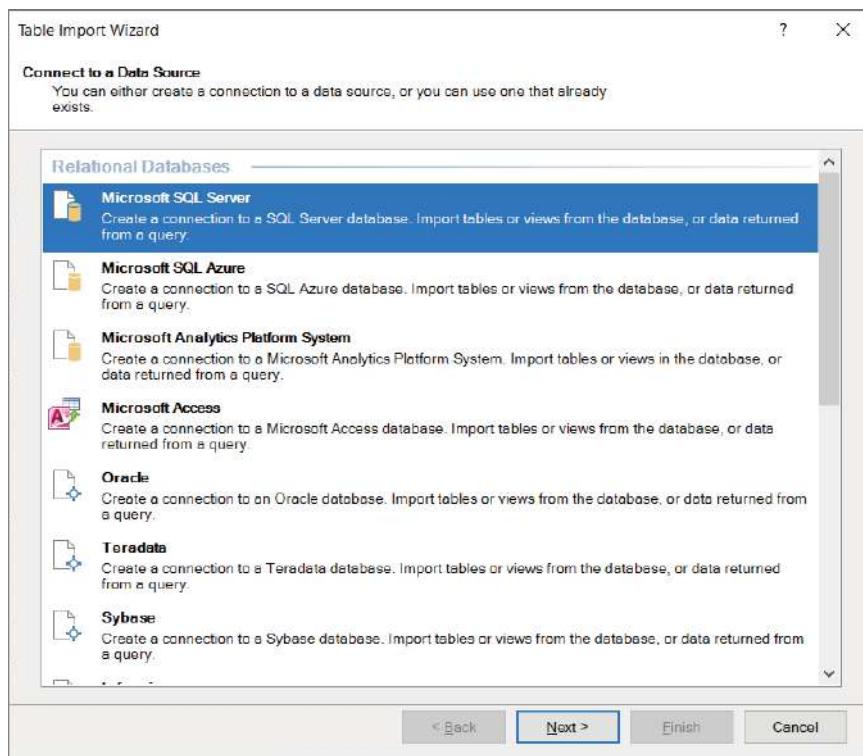
Friendly Connection Name The Friendly Connection Name field allows you to specify your own name for the external source. You typically enter a name that is descriptive and easy to read.

Server Name This is the name of the server that contains the database to which you are trying to connect. You will get this from your IT department when they give you access.

Part V: Understanding Power Pivot and Power Query

FIGURE 32.13

Activate the Table Import Wizard and select Microsoft SQL Server.



Log on to the Server These are your login credentials. Depending on how your IT department gives you access, you will select either Use Windows Authentication or Use SQL Server Authentication. The Use Windows Authentication option essentially means that the server will recognize you by your Windows login. The Use SQL Server Authentication option means that the IT department created a distinct username and password for you. If you select Use SQL Server Authentication, you will need to provide a username and password. Note that you will need at least READ privileges for the target database to pull the needed data.

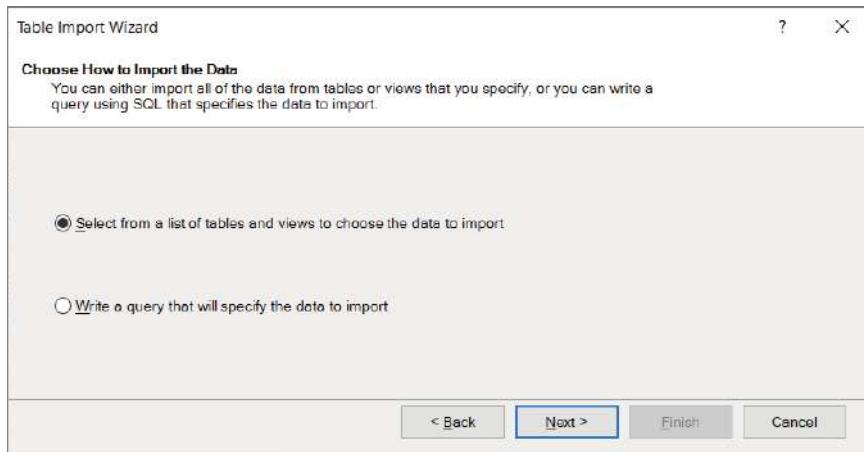
Save my Password You can select the check box next to Save my Password if you want your username and password to be stored in the workbook. This allows your connections to remain refreshable when being used by other people. There are obviously security issues with this option because anyone can view the connection properties and see your username and password. You should use this option only if your IT department set you up with an application account (that is, an account created specifically to be used by multiple people).

Database Name Every instance of SQL Server can contain multiple databases. Enter the name of the database to which you are connecting. You will get this from your IT department when they give you access.

Once you enter all the pertinent information, click the Next button to get to the subsequent screen, shown in Figure 32.14. Here you have the choice of selecting from a list of tables and views or writing your own custom query using SQL syntax. The latter involves writing your own SQL scripts. In most cases, you will choose the option to select from a list of tables.

FIGURE 32.14

Choose to select from a list of tables and views.



The Table Import Wizard reads the database and shows you a list of all available tables and views (see Figure 32.15). The tables will have an icon that looks like a grid, while views will have an icon that looks like a box on top of another box.

The idea is to place a check next to the tables and views that you want to import. The Friendly Name column allows you to enter a new name that will be used to reference the table in Power Pivot.

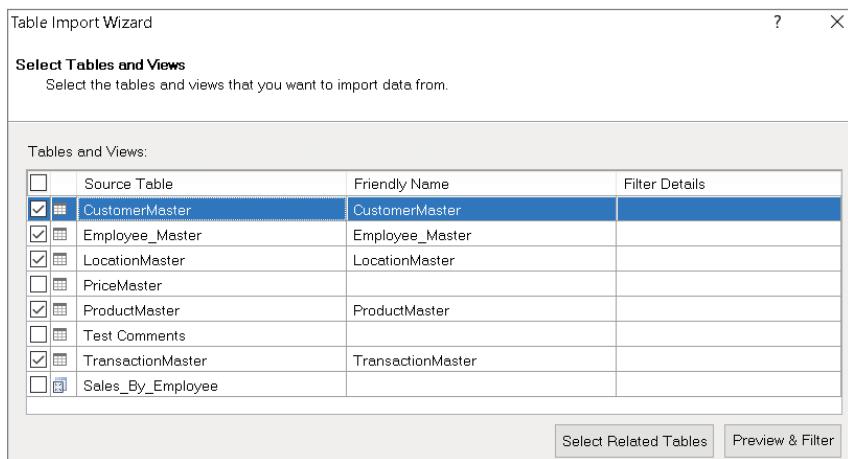
TIP

In Figure 32.15, you'll see a Select Related Tables button. After you select one or more tables, you can click this button to tell Power Pivot to scan for and automatically select any other tables that have a relationship with the table (or tables) you've already selected. This is a handy feature to have when sourcing large databases with dozens of tables.

Part V: Understanding Power Pivot and Power Query

FIGURE 32.15

The Table Import Wizard will display a list of tables and views.



It's important to remember that importing a table imports all of the columns and records for that table. This can have an impact on the size and performance of your Power Pivot data model. You will often find that you need only a handful of the columns from the tables you import. In these cases, you can use the Preview & Filter button.

Click the table name to highlight it in blue (as shown in Figure 32.15) and then click the Preview & Filter button. The Table Import Wizard will activate the preview screen illustrated in Figure 32.16. In this screen, you will see all the columns available in the table, with a sampling of rows.

Each column header has a check box next to it, indicating that the column will be imported with the table. Removing the check mark tells Power Pivot not to include that column in the data model.

You also have the option of filtering out certain records. Clicking one of the drop-down arrows illustrated in Figure 32.16 activates a filter menu that allows you to specify criteria to filter out unwanted records. This works just like the standard filtering in Excel.

Once you're done selecting your data and applying any needed filters, you can click the Finish button on the Table Import Wizard to start the import process. The import log shown in Figure 32.17 displays the progress of the import and summarizes the import actions taken after completion.

The final step in loading data from SQL Server is to review and create any needed relationships. Activate the Power Pivot window and click the Diagram View button on the Home tab. Power Pivot will display the diagram screen where you can view and edit relationships as needed (see "Creating relationships between your PowerPivot tables" earlier in this chapter).

FIGURE 32.16

The Preview & Filter screen allows you to exclude columns and filter for only data you need.

Table Import Wizard

Preview Selected Table

Use the checkbox to select specific columns. To filter the data in a column, use the drop-down arrow for the column to select values that should be included.

Table Name: CustomerMaster

	Customer_Number	FirstOfCustomer_N	FirstOfCity	FirstOfState
1	17214	HULCUM Corp.	SIOUX FALLS	SD
2	17216	SUASHD Corp.	PIERRE	SD
3	17218	BASLUS Corp.	SIOUX FALLS	SD
4	17236	GROHOM Corp.	WORTHINGTON	MN
5	17273	RAMPAR Corp.	CANISTOTA	SD
6	17286	REBANS Corp.	YANKTON	SD
7	17296	OSHUSM Corp.	SPENCER	IA
8	17304	SAUAXF Corp.	SIOUX FALLS	SD
9	17309	PORKVA Corp.	MONTEVIDEO	MN
10	17310	OMUSAC Corp.	MITCHELL	SD
11	17344	ULSUN Corp.	MARSHALL	MN
12	17347	CATCO Corp.	SIOUX FALLS	SD

FIGURE 32.17

The last screen of the Table Import Wizard shows you the progress of your import actions.

Table Import Wizard

Importing

The import operation might take several minutes to complete. To stop the import operation, click the Stop Import button.

Success

Details:

Work Item	Status
CustomerMaster	Success. 9,120 rows transferred.
Employee_Master	Success. 649 rows transferred.
LocationMaster	Success. 54 rows transferred.
ProductMaster	Success. 6 rows transferred.
TransactionMaster	Success. 88,614 rows transferred.
Data preparation	Completed

TIP

If you find that you need to adjust the filtering you've applied to your imported data, you can call up the Preview & Filter screen again. Simply select the target table in the Power Pivot window and activate the Edit Table Properties dialog box (Design \Rightarrow Table Properties). You'll note that this dialog box is basically the same Preview & Filter screen that you encountered with the Table Import Wizard (refer to Figure 32.16). From here, you can select columns that you originally filtered out, edit record filters, clear filters, or even use a different table/view.

Loading data from other relational database systems

Whether your data lives in Microsoft Access, Oracle, dBase, or MySQL, you can load data from virtually any relational database system. As long as you have the appropriate database drivers installed, you have a way to connect Power Pivot to your data.

NOTE

To connect to any database system, you must have that system's drivers installed on your PC. SQL Server and Access are Microsoft products whose drivers are virtually guaranteed to be installed on most machines that you'll encounter. The drivers for other database systems, however, need to be installed explicitly. This is typically done by the IT department either at the time the machine is loaded with corporate software or upon demand. If you don't see the needed drivers for your database system, contact your IT department.

Activate the Power Pivot window, and click the From Other Sources button on the Home tab. This will activate the same Table Import Wizard shown previously in Figure 32.13.

Select the appropriate relational database system from the plethora of options. If you need to import data from Oracle, select Oracle. If you need to import data from Sybase, select Sybase.

Connecting to any of these relational systems takes you through roughly the same steps you saw when importing SQL Server data earlier in this chapter. You may see some alternate dialog boxes, however, based on the needs of the database system you select.

Understandably, Microsoft cannot possibly create a named connection option for every database system out there. So, you may not find your database system listed. In this case, simply select the option called Others (OLEDB/ODBC). Selecting this option opens the Table Import Wizard starting with a screen asking you to enter the connection string for your database system.

Loading data from flat files

The term *flat file* refers to a file that contains some form of tabular data without any sort of structural hierarchy or relationship between records. The most common types of flat files are Excel files and text files. A ton of important data is maintained in flat files. In the following sections, you'll discover how to import these flat file data sources into the Power Pivot data model.

Loading data from external Excel files

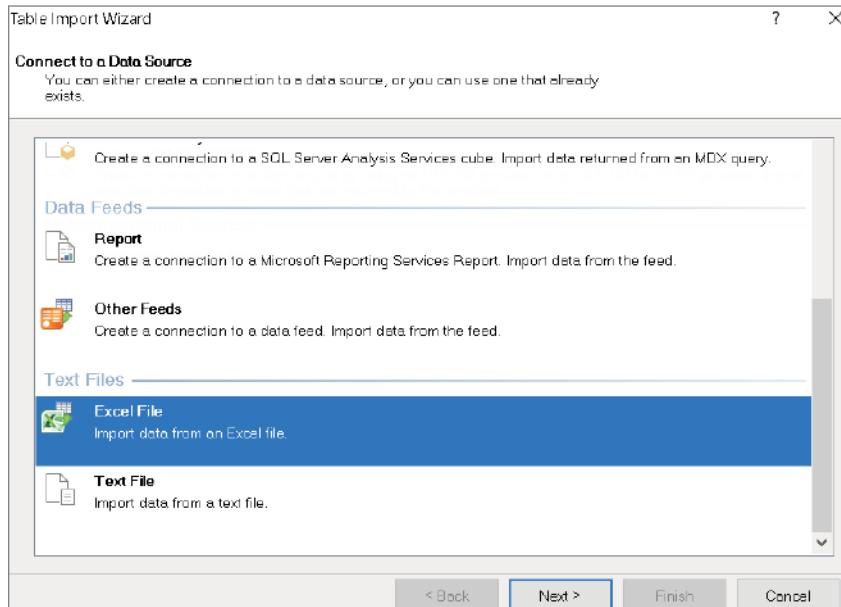
Earlier in this chapter, you created linked tables by loading Power Pivot with the data contained within the same workbook. Linked tables have a distinct advantage over other types of imported data in that they immediately respond to changes in the source data within the workbook. If you change the data in one of the tables in the workbook, the linked table within the Power Pivot data model automatically changes. The real-time interactivity you get with linked tables is really nice to have.

The drawback to linked tables is that the source data must be kept in the same workbook as the Power Pivot data model. This isn't always possible. You'll encounter plenty of scenarios where you'll need to incorporate Excel data into your analysis but that data lives in another workbook. In these cases, you can use Power Pivot's Table Import Wizard to connect to external Excel files.

Activate the Power Pivot window and click the From Other Sources button on the Home tab. This will activate the Table Import Wizard dialog box. Scroll down and select the Excel File option shown in Figure 32.18 and then click the Next button.

FIGURE 32.18

Activate the Table Import Wizard and select Excel File.



Part V: Understanding Power Pivot and Power Query

The Table Import Wizard will now ask for all the information it needs to connect to your target workbook. In this screen, you'll need to provide the following:

Friendly Connection Name The Friendly Connection Name field allows you to specify your own name for the external source. You typically enter a name that is descriptive and easy to read.

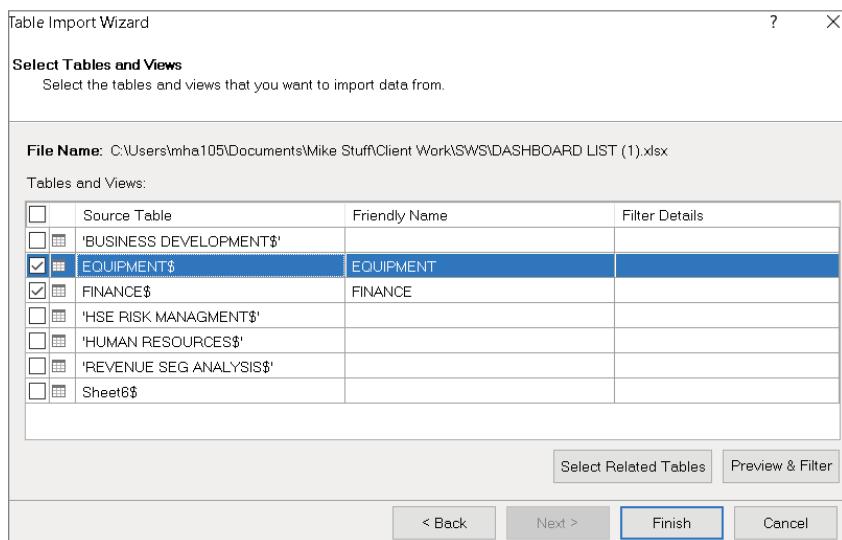
Excel File Path Enter the full path of your target Excel workbook. You can use the Browse button to search for and select the workbook from which you want to pull this information.

Use First Row as Column Headers In most cases, your Excel data will have column headers. Be sure to select the check box next to Use First Row as Column Headers to make sure that your column headers are recognized as headers when imported.

Once you enter all the pertinent information, click the Next button to get to the next screen, shown in Figure 32.19. Here you'll see a list of all the worksheets in the chosen Excel workbook. Place a check mark next to the worksheets you want to import. The Friendly Name column allows you to enter a new name that will be used to reference the table in Power Pivot.

FIGURE 32.19

Select the worksheets you want to import.



TIP

When reading from external Excel files, Power Pivot allows you to select named ranges as well as entire worksheets. This comes in handy when you want to import only a subset of data within a worksheet.

As discussed earlier in this chapter, you can use the Preview & Filter button to filter out unwanted columns and records if needed. Otherwise, continue with the Table Import Wizard to complete the import process.

As always, be sure to review and create relationships to any other tables you've loaded into the Power Pivot data model.

NOTE

Be aware that loading external Excel data does not give you the same interactivity you get with linked tables. Just as with importing database tables, the data you bring from an external Excel file is simply a snapshot. You'll need to refresh the data connection to see any new data that may have been added to the external Excel file (see "Manually refreshing your Power Pivot data" later in this chapter).

Loading data from text files

Text files are another type of flat file used to distribute data. These files are commonly outputs from legacy systems and websites. Excel has always been able to import text files. With Power Pivot, you can go further and integrate them with other data sources.

Activate the Power Pivot window, and click the From Other Sources button on the Home tab. This will activate the same Table Import Wizard dialog box illustrated in Figure 32.18. Select the Text File option and then click the Next button.

The Table Import Wizard will ask for all the information it needs to connect to the target text file. In this screen, you'll need to provide the following:

Friendly Connection Name The Friendly Connection Name field allows you to specify your own name for the external source. You typically enter a name that is descriptive and easy to read.

File Path Enter the full path of your target text file. You can use the Browse button to search for and select the file from which you want to pull this information.

Column Separator Select the character used to separate the columns in the text file.

Before you can do this, you'll need to know how the columns in your text file are delimited. For instance, a comma-delimited file will have commas separating the columns. A tab-delimited file will have tabs separating the columns. The drop-down list in the Table Import Wizard includes choices for the more common delimiters: Tab, Comma, Semicolon, Space, Colon, and Vertical Bar.

Use First Row as Column Headers If your text file contains header rows, be sure to select the check box next to Use First Row as Column Headers. This ensures that the column headers are recognized as headers when imported.

You'll get an immediate preview of the data in the text file. As with other data sources we've covered here, you can filter out any unwanted columns simply by removing the check mark next to the column names. You can also use the drop-down arrows next to each column to apply any record filters.

Clicking the Finish button will immediately start the import process. Upon completion, the data from your text file will be part of the Power Pivot data model. As always, be sure to review and create relationships to any other tables that you've loaded into Power Pivot.

Loading data from the Clipboard

Power Pivot includes an interesting option for loading data straight from the Clipboard, that is, pasting data you've copied from some other place. This option is meant to be used as a one-off technique to get useful information into the Power Pivot data model quickly.

As you consider this option, keep in mind that there is no real data source. It's just you manually copying and pasting. There is no way to refresh the data, and there is no way to trace back where you actually copied the data from.

Imagine that you received a Word document showing a list of branches in a table, as shown in Figure 32.20. You'd like to include this static list of branches in your Power Pivot data model.

FIGURE 32.20

You can copy data straight out of Microsoft Word.

	A	B	C
1	Branch_Number	Market	Region
2	101313	FLORIDA	SOUTH
3	101419	MICHIGAN	NORTH
4	102516	TULSA	MIDWEST
5	103516	TULSA	MIDWEST
6	173901	CHARLOTTE	SOUTH
7	201605	DENVER	MIDWEST
8	201709	SEATTLE	WEST
9	201714	PHOENIX	WEST
10	201717	CALIFORNIA	WEST
11	202600	DALLAS	SOUTH
12	202605	DENVER	MIDWEST
13	202714	PHOENIX	WEST

You can copy the table and then go to the Power Pivot window and click the Paste command on the Home tab. This activates the Paste Preview dialog box illustrated in Figure 32.21, where you can review what exactly will be pasted.

There aren't that many options here. You can specify the name that will be used to reference the table in Power Pivot, and you can specify whether the first row contains column headers.

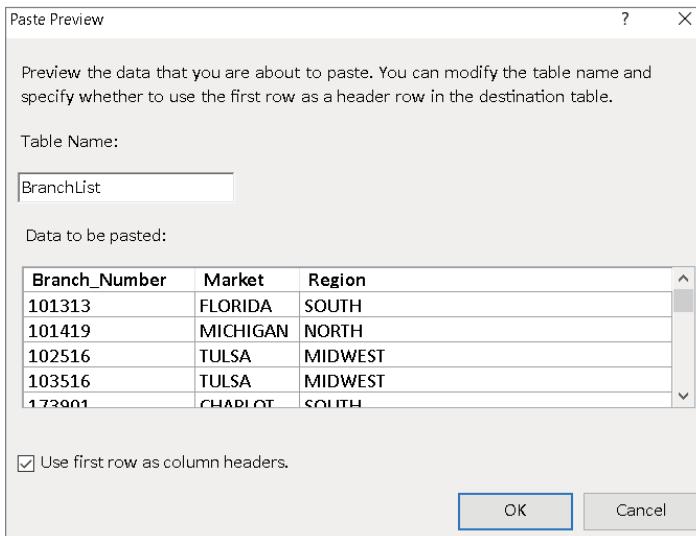
Clicking the OK button will import the pasted data into Power Pivot without a lot of fanfare. At this point, you can adjust the data formatting and create the needed relationships.

Refreshing and managing external data connections

When you load data from an external data source into Power Pivot, you essentially create a static snapshot of that data source at the time of creation. Power Pivot uses that static snapshot in its internal data model.

FIGURE 32.21

The Paste Preview dialog box gives you a chance to see what you're pasting.



As time goes by, the external data source may change and grow with newly added records. However, Power Pivot is still using its snapshot, so it can't incorporate any of the changes in your data source until you take another snapshot.

The action of updating the Power Pivot data model by taking another snapshot of your data source is called *refreshing your data*. You can refresh manually, or you can set up an automatic refresh.

Manually refreshing your Power Pivot data

On the Home tab of the Power Pivot window, you will see the Refresh command. Clicking the drop-down arrow below it will display two options: Refresh and Refresh All.

Use the Refresh option to refresh the PowerPivot table that's currently active. That is, if you are on the Customers tab in Power Pivot, clicking Refresh will reach out to the external data source and request an update for just the Customers table. This works nicely when you need to refresh strategically only certain data sources.

Use the Refresh All option to refresh all the tables in the Power Pivot data model.

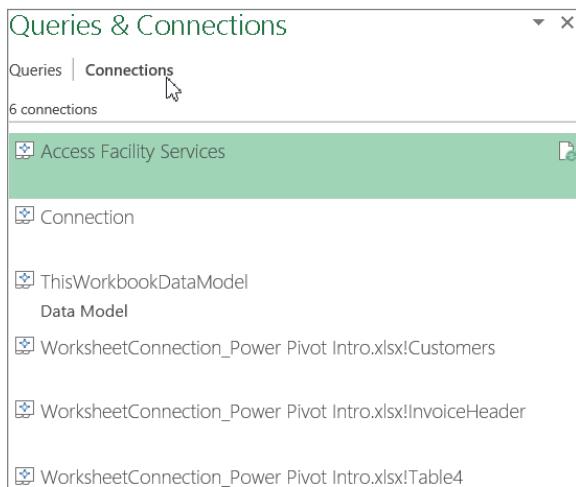
Setting up automatic refreshing

You can configure your data sources to pull the latest data and refresh Power Pivot automatically. Go to the Data tab in the Excel Ribbon and select the Queries & Connections command. This will activate the Queries & Connections task pane, illustrated in Figure 32.22.

Part V: Understanding Power Pivot and Power Query

FIGURE 32.22

The Queries & Connections task pane



Click Connections at the top of the task pane and then double-click the connection with which you want to work.

With the Connection Properties dialog box open (see Figure 32.23), select the Usage tab. Here you'll find the following options:

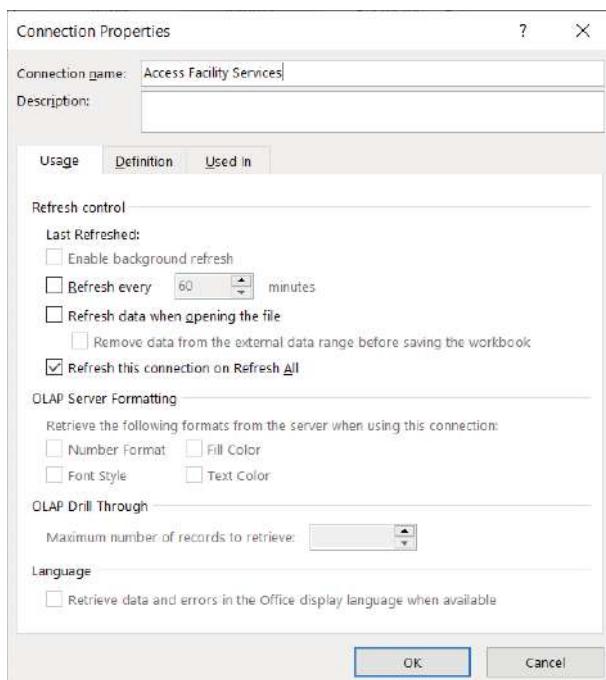
Refresh Every X Minutes Placing a check mark next to this option tells Excel to refresh the chosen data connection automatically after a specified number of minutes. This will refresh all tables associated with that connection.

Refresh data when opening the file Placing a check mark next to this option tells Excel to refresh the chosen data connection automatically upon opening the workbook. This will refresh all tables associated with that connection as soon as the workbook is opened.

Refresh this connection on Refresh All Earlier you discovered that you can refresh all the connections that feed Power Pivot by using the Refresh All command found on the Power Pivot Home tab. If the data connection in question imports millions of rows of data from an external data source, you may not want to slow down your machine each time you trigger a Refresh All. You can remove the check mark next to the Refresh this connection on Refresh All option. This essentially tells the connection to ignore the Refresh All command.

FIGURE 32.23

The Connection Properties dialog box lets you configure the chosen data connection to refresh automatically.



Editing your data connection

There may be instances where you need to edit your source data connection after you've already created it. Unlike refreshing, where you simply take another snapshot of the same data source, editing the source data connection allows you to go back and reconfigure the connection. Here are a few reasons that you'll need to edit a data connection:

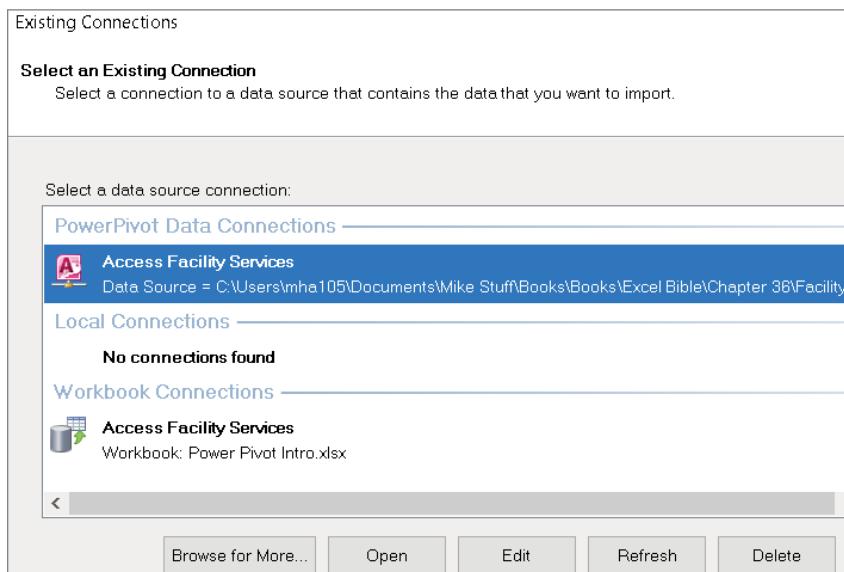
- The location or server or data source file has changed.
- The name of the server or data source file has changed.
- You need to edit your login credentials or authentication mode.
- You need to add tables that you left out during the initial import.

In the Power Pivot window, go to the Home tab and click the Existing Connections button. The Existing Connections dialog box, shown in Figure 32.24, will open. Your Power Pivot connections will be under the Power Pivot Data Connections subheading. Choose the data connection that needs editing.

Part V: Understanding Power Pivot and Power Query

FIGURE 32.24

Use the Existing Connections dialog box to reconfigure your Power Pivot source data connections.



Once your target data connection is selected, look to the Edit and Open buttons. The button you click depends on what you need to change.

The Edit button Lets you reconfigure the server address, file path, and authentication settings.

The Open button Lets you import a new table from the existing connection. This is handy when you inadvertently missed a table during the initial loading of data.

Working Directly with the Internal Data Model

IN THIS CHAPTER

- Directly feeding the internal data model
- Managing relationships in the internal data model
- Working with the Queries & Connections task pane

Every workbook comes with an internal data model that starts as an empty container. As you work with Power Pivot, the internal data model gets populated with the tables and connections you add to it.

Up until now, you used the Power Pivot Ribbon interface to work with the internal data model. As you will discover in this chapter, however, you can use a combination of PivotTables and Excel data connections to interact directly with the internal data model—without the Power Pivot Ribbon interface.



Most of the examples in this chapter are available on this book's website at www.wiley.com/go/excel365bible. The filename is Internal Data Model.xlsx.

Directly Feeding the Internal Data Model

Imagine that you have the Transactions table shown in Figure 33.1. On another worksheet, you have an Employees table (see Figure 33.2) that contains information about the employees.

You need to create an analysis that shows sales by job title. This would normally be difficult given the fact that sales and job title are in two separate tables. But with the internal data model, it's just a matter of a few clicks.

Part V: Understanding Power Pivot and Power Query

FIGURE 33.1

This table shows transactions by employee number.

	A	B	C	D
1	Sales_Rep	Invoice_Date	Sales_Amount	Contracted Hours
2	4416	1/5/2007	111.79	2
3	4416	1/5/2007	111.79	2
4	160006	1/5/2007	112.13	2
5	6444	1/5/2007	112.13	2
6	160006	1/5/2007	145.02	3
7	52661	1/5/2007	196.58	4
8	6444	1/5/2007	204.20	4
9	51552	1/5/2007	225.24	3
10	55662	1/6/2007	86.31	2
11	1336	1/6/2007	86.31	2
12	60224	1/6/2007	86.31	2
13	54564	1/6/2007	86.31	2
14	56146	1/6/2007	89.26	2
15	5412	1/6/2007	90.24	1

FIGURE 33.2

This table provides information on employees: last name, first name, and job title.

	A	B	C	D
1	Employee_Number	Last_Name	First_Name	Job_Title
2	21	SIOCAT	ROBERT	SERVICE REPRESENTATIVE 3
3	42	BREWN	DONNA	SERVICE REPRESENTATIVE 3
4	45	VAN HUILE	KENNETH	SERVICE REPRESENTATIVE 2
5	104	WIBB	MAURICE	SERVICE REPRESENTATIVE 2
6	106	CESTENGIAY	LUC	SERVICE REPRESENTATIVE 2
7	113	TRIDIL	ROCH	SERVICE REPRESENTATIVE 2
8	142	CETE	GUY	SERVICE REPRESENTATIVE 3
9	145	ERSINEILT	MIKE	SERVICE REPRESENTATIVE 2
10	162	GEBLE	MICHAEL	SERVICE REPRESENTATIVE 2
11	165	CERDANAL	ALAIN	SERVICE REPRESENTATIVE 3
12	201	GEIDRIOU	DOMINIC	TEAMLEAD 1

The first thing you will need to do is to convert your data tables into Excel Table objects that the internal data model can recognize:

1. **Click inside the TransactionMaster data table and select Insert \leftrightarrow Table.** This activates the Create Table dialog box.
2. **Select the table range and then click OK.** You will want to name your Table objects explicitly. This way, you can easily recognize your tables in the internal data model. If you don't name your tables, the internal data model shows them as Table1, Table2, and so on.
3. **Click any cell in the table and select the Table Design tab.** There you will find the Table Name input box in the Properties group. Enter a new name for your table (in this case, **Transactions**). At this point, you're ready to feed your table to the internal data model.

4. Select Data \Rightarrow Existing Connections.
5. In the Existing Connections dialog box, select the Tables tab. You will see a list of existing Table objects (see Figure 33.3).
6. Double-click the Transactions Table object. This will activate the Import Data dialog box shown in Figure 33.4.
7. Select the Only Create Connection option and then click OK.

Repeat the previous steps for the Employees table.

FIGURE 33.3

The Existing Connections dialog box lists all available Table objects.

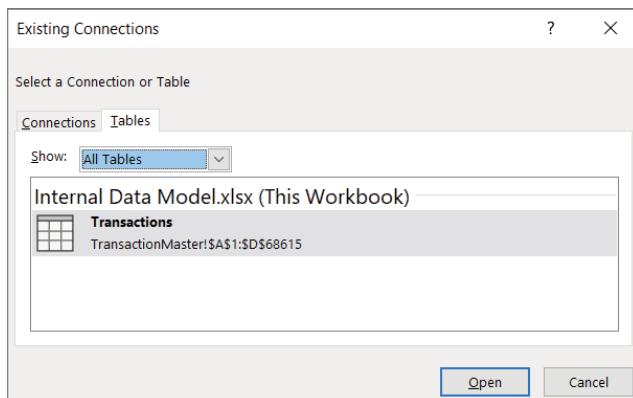
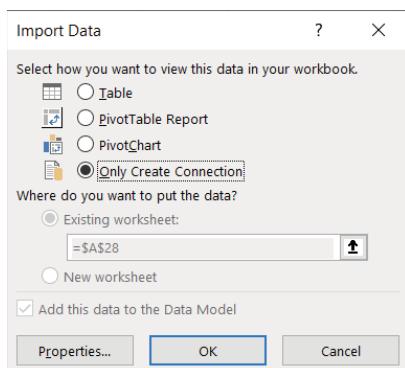


FIGURE 33.4

Use the Import Data dialog box to add the Table object to the internal data model.



Part V: Understanding Power Pivot and Power Query

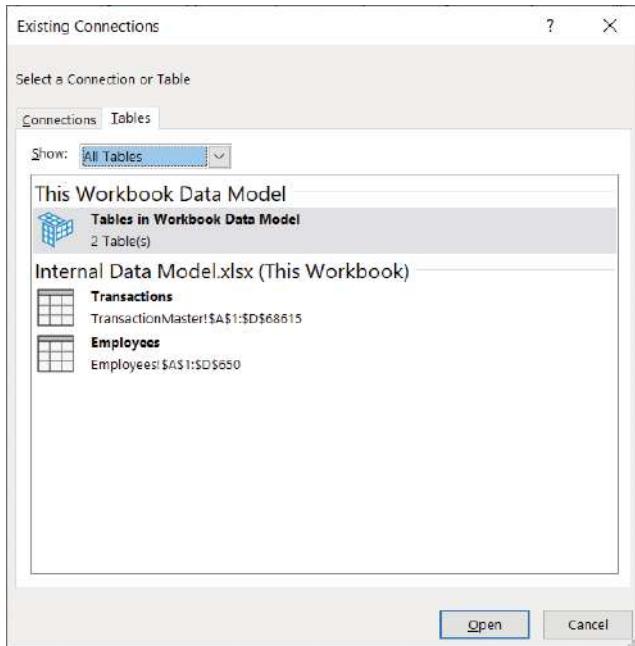
At this point, both tables are loaded into the internal data model. The internal data model can be seen in the Existing Connections dialog box. As Figure 33.5 illustrates, the internal model is called Tables in Workbook Data Model.

You can create the needed pivot table by following these steps:

1. **Select Data ↴ Get & Transform Data ↴ Existing Connections.**
2. **In the Existing Connections dialog box, select the Tables tab.**

FIGURE 33.5

You can explicitly choose the internal data model as the source for a PivotTable.



3. **Double-click the Tables in Workbook Data Model entry.**
4. **The Import Data dialog box will open.** As Figure 33.6 illustrates, you will choose to create a PivotTable report on a new worksheet. Click OK to confirm.
At this point, you can click the newly created PivotTable to activate the PivotTable Fields list, if necessary. Choose the All selector, as shown in Figure 33.7. This shows all available tables in the Fields list.
5. **Build out your PivotTable as normal.** In this case, Job_Title goes to the Rows area and Sales_Amount goes to the Values area.

FIGURE 33.6

Creating a new PivotTable on a new worksheet

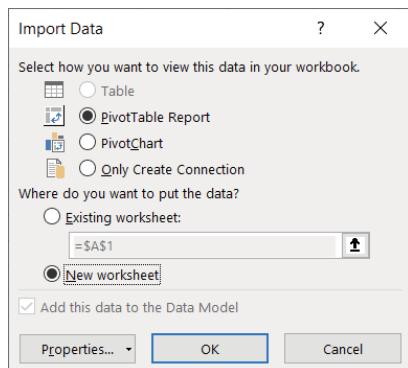
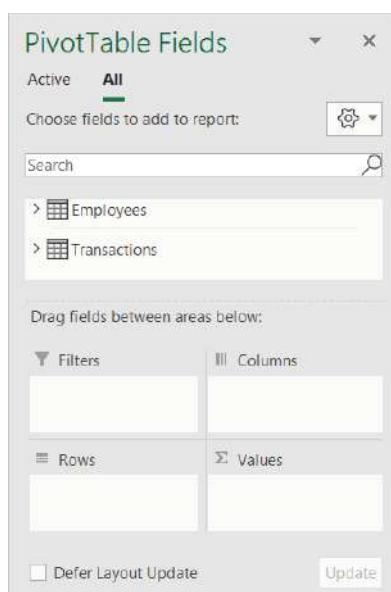


FIGURE 33.7

Select All in the PivotTable Fields list to see both tables in your internal data model.

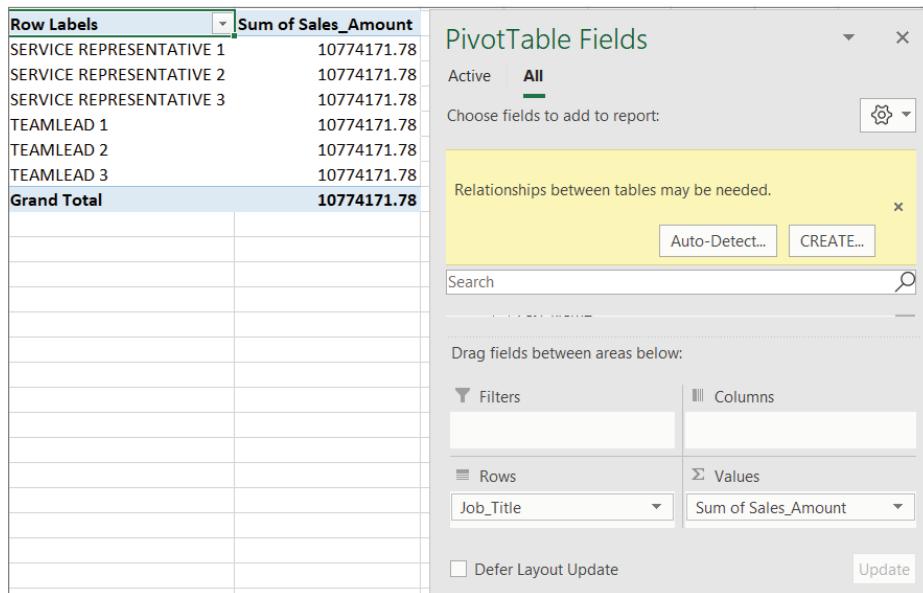


As you can see in Figure 33.8, your PivotTable results don't look quite right. The PivotTable is showing the same total sales amount for all job titles. This is because no relationship exists between the two tables we just imported.

Part V: Understanding Power Pivot and Power Query

FIGURE 33.8

If no relationships exist for the tables in your PivotTable, Excel prompts you to create the relationship between the two tables.



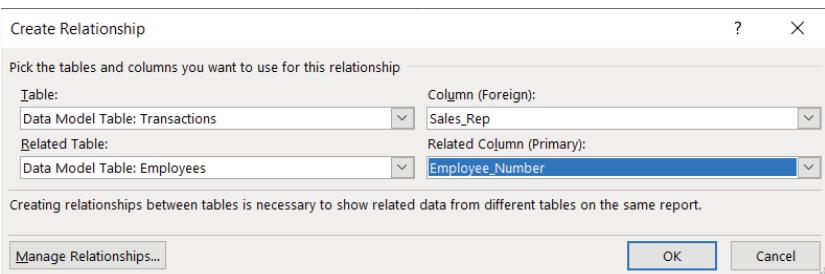
Excel immediately recognizes that you are using two tables from your internal data model and prompts you to create a relationship between them. You have the option of letting Excel auto-detect the relationships between your tables, or you can click the Create button. It's always best to create the relationships yourself to avoid any possibility of Excel getting it wrong.

6. **Click the Create button.** Excel activates the Create Relationship dialog box shown in Figure 33.9. Here you select the tables and fields that define the relationship. In Figure 33.9, you can see that the Transactions table has a Sales_Rep field. It is related to the Employees table via the Employee_Number field.

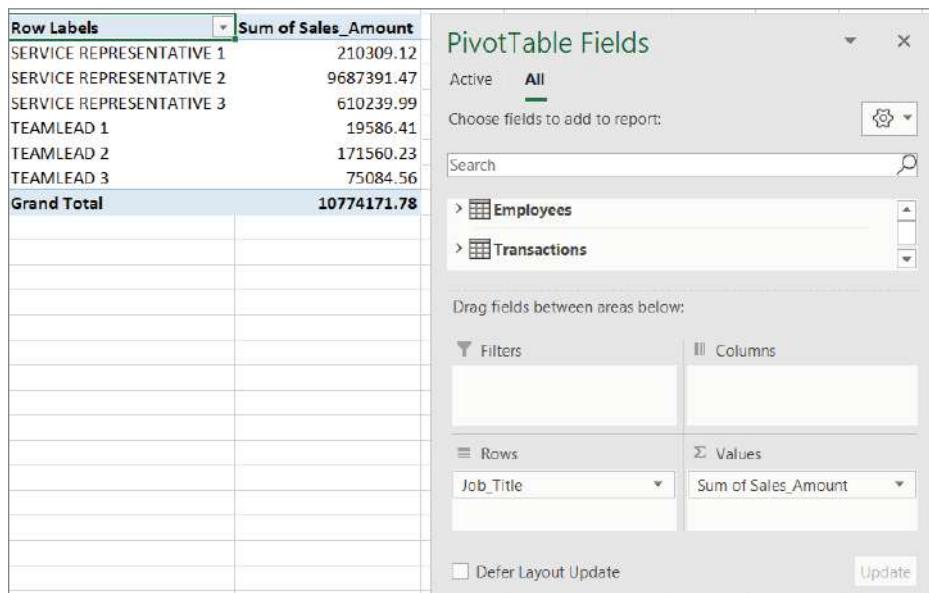
After you create your relationship, you have a single PivotTable that effectively uses data from both tables to create the analysis you need. Figure 33.10 illustrates the final PivotTable.

FIGURE 33.9

Build the appropriate relationship using the Table and Column drop-down lists.

**FIGURE 33.10**

You have achieved your goal of showing sales by job title.



NOTE

In Figure 33.9, you see that the lower-right drop-down is named Related Column (Primary). The term *primary* means that the internal data model uses this field from the associated table as the primary key.

A primary key is a field that contains only unique non-null values (no duplicates or blanks). Primary key fields are necessary in the data model to prevent aggregation errors and duplications. Every relationship that you create must have a field designated as the primary key.

So, the Employees table (in the scenario in Figure 33.9) must have all unique values in the Employee_Number field, with no blanks or null values. This is the only way that Excel can ensure data integrity when joining multiple tables.

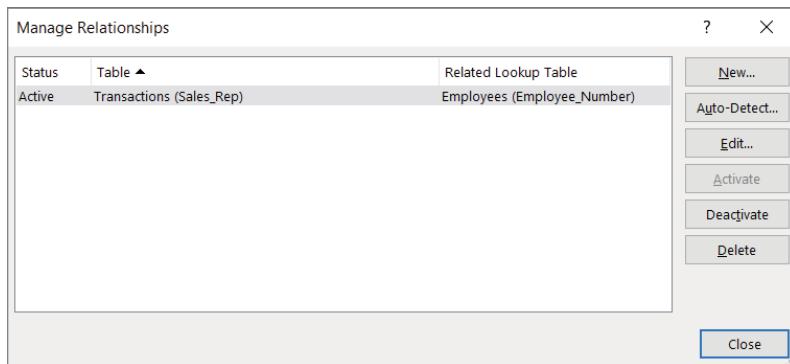
Managing Relationships in the Internal Data Model

After you assign tables to the internal data model, you might need to adjust the relationships between them. To make changes to the relationships used in a PivotTable, you'll need to activate the Manage Relationships dialog box.

Click anywhere inside your PivotTable, select the PivotTable Analyze tab on the Ribbon and then select the Relationships command. The dialog box shown in Figure 33.11 appears.

FIGURE 33.11

The Manage Relationships dialog box enables you to make changes to the relationships in the internal data model.



Here, you will find the following commands:

New Create a new relationship between two tables in the internal data model.

Auto-Detect Have Power Pivot automatically detect and create relationships based on the data in your tables.

Edit Alter the selected relationship.

Activate Enforce the selected relationship, telling Excel to consider the relationship when aggregating and analyzing the data in the internal data model.

Deactivate Turn off the selected relationship, telling Excel to ignore the relationship when aggregating and analyzing the data in the internal data model.

Delete Remove the selected relationship.

Managing Queries & Connections

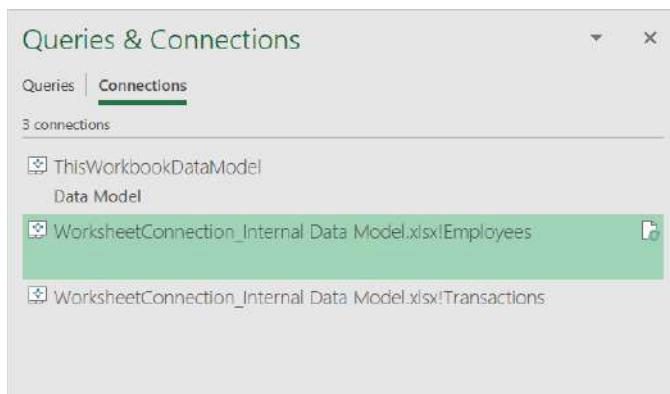
Click on the Data tab in the Ribbon and then select the Queries & Connections command. Excel will activate the Queries & Connections task pane (see Figure 33.12). This task pane shows all connections and queries within the current workbook.

If you receive a workbook that is unfamiliar to you, it's best practice to activate Queries & Connections just to see if you're dealing with any external connections or queries in the internal data model of the workbook.

Right-click any of the entries to expose a shortcut menu for that entry, allowing you to refresh the connection, delete the connection, or edit the properties.

FIGURE 33.12

Use the Queries & Connections task pane to manage the queries and connections in the internal data model.



Adding Formulas to Power Pivot

IN THIS CHAPTER

- Creating your own calculated columns
- Utilizing DAX to create calculated columns
- Creating calculated measures
- Using Cube functions to break out of PivotTables

When analyzing data with Power Pivot, you'll often find that you need to expand your analysis to include data based on calculations that are not in your original data set. Power Pivot has a robust set of functions called *Data Analysis Expressions*, or *DAX functions*, which allow you to perform mathematical operations, recursive calculations, data lookups, and much more.

This chapter introduces you to DAX functions, and it provides the ground rules for building your own calculations in Power Pivot data models.



Most of the examples in this chapter are available on this book's website at www.wiley.com/go/excel365bible. The filename is Power Pivot Formulas.xlsx.

Enhancing Power Pivot Data with Calculated Columns

Calculated columns are columns that you create to enhance a PowerPivot table with your own formulas. Calculated columns are entered directly in the Power Pivot window, becoming part of the source data you use to feed your PivotTable. Calculated columns work at the row level; that is, the formulas that you create in a calculated column perform their operations based on the data in each individual row. For example, imagine you have a Revenue column and a Cost column in your PowerPivot table. You could create a new column that calculated [Revenue] minus [Cost]. This calculation is simple and is valid for each row in the data set.

Calculated measures are used to perform more complex calculations that work on an aggregation of data. These calculations are applied directly to your PivotTable, creating a sort of virtual column that can't be seen in the Power Pivot window. Calculated measures are needed when you need to calculate based on an aggregated grouping of rows, such as the sum of [Year2] minus the sum of [Year1].

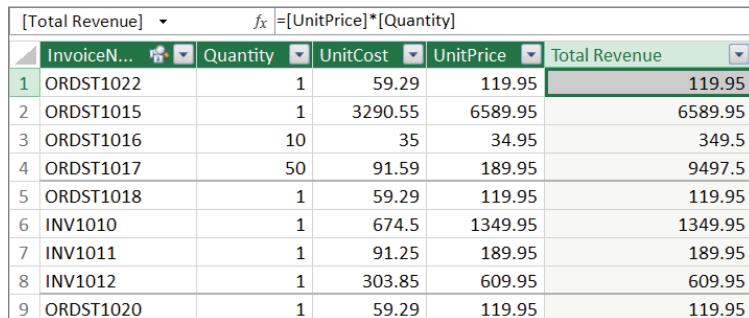
Creating your first calculated column

Creating a calculated column works very much like building formulas in an Excel table. Let's continue the walk-through with some sample data. Follow these steps to create a calculated column:

1. Open the Power Pivot Formulas.xlsx sample file, activate the Power Pivot window (by clicking the Manage button on the Power Pivot Ribbon tab), and then select the Invoice Details tab.
2. In the table, you will see an empty column on the far right labeled Add Column. Click the first blank cell in that column.
3. In the Formula bar (see Figure 34.1), enter the following formula:
`= [UnitPrice] * [Quantity]`
4. Press Enter to see your formula populate the entire column.
5. Power Pivot will automatically rename the column to Calculated Column
1. Double-click the column label and rename the column Total Revenue.

FIGURE 34.1

Start your calculated column by entering your desired operation in the Formula bar.



	InvoiceN...	Quantity	UnitCost	UnitPrice	Total Revenue
1	ORDST1022	1	59.29	119.95	119.95
2	ORDST1015	1	3290.55	6589.95	6589.95
3	ORDST1016	10	35	34.95	349.5
4	ORDST1017	50	91.59	189.95	9497.5
5	ORDST1018	1	59.29	119.95	119.95
6	INV1010	1	674.5	1349.95	1349.95
7	INV1011	1	91.25	189.95	189.95
8	INV1012	1	303.85	609.95	609.95
9	ORDST1020	1	59.29	119.95	119.95

NOTE

You can rename any column in the Power Pivot window by double-clicking the column name and entering a new name. Alternatively, you can right-click any column and choose the Rename Column option.

TIP

You can build your calculated columns by clicking instead of typing. For example, instead of manually entering `= [UnitPrice] * [Quantity]`, you can enter the equal sign (=), click the UnitPrice column, enter the asterisk (*), and then click the Quantity column. Note that you can also enter your own static data. For example, you can enter a formula to add a 10 percent tax by entering `= [UnitPrice] * 1.10`.

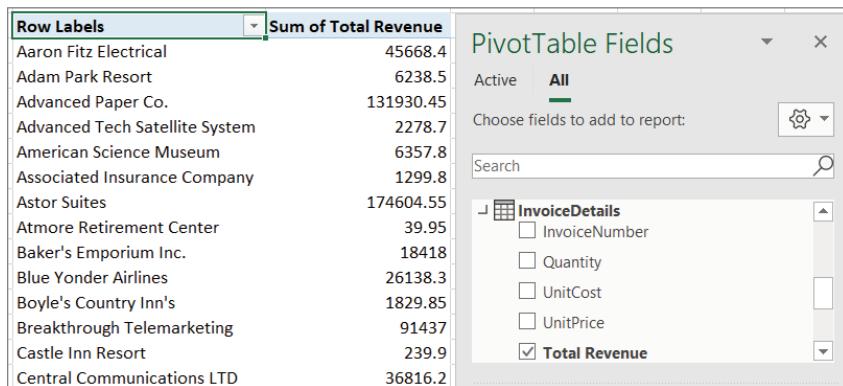
Each calculated column that you create will automatically be available in the field list of any PivotTable connected to the Power Pivot data model. Figure 34.2 illustrates the Total Revenue calculated column in the PivotTable Fields list. These calculated columns can be used just as you would any other field in your PivotTable.

TIP

If you don't see the PivotTable Fields list, simply right-click the PivotTable and select Show Field List.

FIGURE 34.2

Calculated columns automatically show up in your PivotTable Fields list.

**NOTE**

If you need to edit the formula in a calculated column, find the calculated column in the Power Pivot window, click the column, and then make your changes directly in the Formula bar.



See Chapter 32, "Introducing Power Pivot," for a refresher on how to create a PivotTable from Power Pivot.

34

Formatting your calculated columns

You'll often need to change the formatting of your Power Pivot columns to match the data within them appropriately. For example, you might want to show numbers as currency, remove decimal places, or display dates in a certain way.

You are by no means limited to formatting just calculated columns. These steps can be used to format any column that you see in your Power Pivot window:

1. In the Power Pivot window, click the column that you want to format.

Part V: Understanding Power Pivot and Power Query

2. Go to the Home tab of the Power Pivot window and find the Formatting group (see Figure 34.3).
3. Use the option to alter the formatting of the column as you see fit.

FIGURE 34.3

You can use the formatting tools found on the Power Pivot window's Home tab to format any column in the data model.

The screenshot shows the Microsoft Power Pivot window with the 'Home' tab selected in the ribbon. The ribbon also includes 'File', 'Design', and 'Advanced' tabs. Under the 'Home' tab, there are sections for 'Clipboard' (Paste Append, Paste Replace, Copy, Paste, From Database, From Service, From Other Sources, Existing Connections, Refresh, PivotTable), 'Formatting' (Data Type dropdown set to 'Currency', Format dropdown set to 'Currency', and buttons for '\$', '%', and '#,##0.00'), and a 'Get External Data' section. Below the ribbon, the formula bar shows '[Total Revenue] = [UnitPrice]*[Quantity]'. A PivotTable is displayed with the following data:

	InvoiceN...	Quantity	UnitCost	UnitPrice	Total Revenue	Add
1	ORDST1022	1	59.29	119.95	\$119.95	
2	ORDST1015	1	3290.55	6589.95	\$6,589.95	
3	ORDST1016	10	35	34.95	\$349.50	
4	ORDST1017	50	91.59	189.95	\$9,497.50	
5	ORDST1018	1	59.29	119.95	\$119.95	
6	INV1010	1	674.5	1349.95	\$1,349.95	
7	INV1011	1	91.25	189.95	\$189.95	
8	INV1012	1	303.85	609.95	\$609.95	

TIP

Veterans of Excel PivotTables know that changing PivotTable number formats one data field at a time is a pain. One fantastic feature of Power Pivot formatting is that any format you apply to your columns in the Power Pivot window will automatically be applied to all PivotTables connected to the data model!

Referencing calculated columns in other calculations

As with all calculations in Excel, Power Pivot allows you to reference a calculated column as a variable in another calculated column. Figure 34.4 illustrates this with a new calculated column called Gross Margin. Notice in the Formula bar that the calculation uses our previously created [Total Revenue] calculated column.

Hiding calculated columns from end users

Because calculated columns can reference each other, you can imagine creating columns simply as helper columns for other calculations. You may not want your end users to see these columns in your client tools. In this context, the term *client tools* refers to PivotTables, Power View dashboards, and Power Map.

FIGURE 34.4

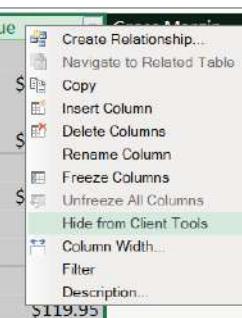
The new Gross Margin calculation uses the previously created [Total Revenue] calculated column.

	[Gross Margin]	fx	=[Total Revenue]-([UnitCost]*[Quantity])			
	InvoiceN...	Quantity	UnitCost	UnitPrice	Total Revenue	Gross Margin
1	ORDST1022	1	59.29	119.95	\$119.95	60.66
2	ORDST1015	1	3290.55	6589.95	\$6,589.95	3299.4
3	ORDST1016	10	35	34.95	\$349.50	-0.5
4	ORDST1017	50	91.59	189.95	\$9,497.50	4918
5	ORDST1018	1	59.29	119.95	\$119.95	60.66
6	INV1010	1	674.5	1349.95	\$1,349.95	675.45
7	INV1011	1	91.25	189.95	\$189.95	98.7
8	INV1012	1	303.85	609.95	\$609.95	306.1
9	ORDST1020	1	59.29	119.95	\$119.95	60.66
10	ORDST1021	1	59.29	119.95	\$119.95	60.66

Similar to hiding columns on an Excel worksheet, Power Pivot allows you to hide any column (it doesn't have to be a calculated column). To hide columns, simply select the columns you want hidden, right-click the selection, and then choose the Hide from Client Tools option (illustrated in Figure 34.5).

FIGURE 34.5

Right-click and select Hide from Client Tools.



	InvoiceN...	Quantity	UnitCost	UnitPrice	Total Revenue	
1	ORDST1022	1	59.29	119.95		
2	ORDST1015	1	3290.55	6589.95		
3	ORDST1016	10	35	34.95		
4	ORDST1017	50	91.59	189.95		
5	ORDST1018	1	59.29	119.95		
6	INV1010	1	674.5	1349.95		
7	INV1011	1	91.25	189.95		
8	INV1012	1	303.85	609.95		
9	ORDST1020	1	59.29	119.95		
10	ORDST1021	1	59.29	119.95		

When a column is hidden, it will not show as an available selection in your PivotTable Fields list. However, if the column you are hiding is already part of the PivotTable report, meaning that you already dragged it onto the PivotTable, hiding the column in the Power Pivot window will not automatically remove it from the report. Hiding merely affects the ability to see the column in the PivotTable Fields list.

You'll note in Figure 34.6 that Power Pivot recolors columns based on their attributes. Hidden columns are subdued and grayed out, while calculated columns that are not hidden have a darker (black) header.

FIGURE 34.6

Hidden columns are grayed out, and calculated columns have a darker header.

InvoiceN...	Quantity	UnitCost	UnitPrice	Total Revenue	Gross Margin
1 ORDST1022	1	59.29	119.95	\$119.95	60.66
2 ORDST1015	1	3290.55	6589.95	\$6,589.95	3299.4
3 ORDST1016	10	35	34.95	\$349.50	-0.5
4 ORDST1017	50	91.59	189.95	\$9,497.50	4918
5 ORDST1018	1	59.29	119.95	\$119.95	60.66
6 INV1010	1	674.5	1349.95	\$1,349.95	675.45
7 INV1011	1	91.25	189.95	\$189.95	98.7
8 INV1012	1	303.85	609.95	\$609.95	306.1

NOTE

To unhide columns, select the hidden columns in the Power Pivot window, right-click the selection, and then choose the Unhide from Client Tools option.

Utilizing DAX to Create Calculated Columns

Data Analysis Expressions (DAX) is essentially the formula language Power Pivot uses to perform calculations within its own construct of tables and columns. The DAX formula language comes with its own set of functions. Some of these functions can be used in calculated columns for row-level calculations, while others are designed to be used in calculated measures for aggregate operations.

In the following sections, we'll touch on some of the DAX functions that can be leveraged in calculated columns.

NOTE

There are more than 150 different DAX functions. The examples of DAX demonstrated in this chapter are meant to give you a sense of how calculated columns and calculated measures work. A full overview of DAX is beyond the scope of this book. If after reading this chapter you have a desire to learn more about DAX, consider picking up the book by Alberto Ferrari and Marco Russo called *The Definitive Guide to DAX* (Microsoft Press, 2019). Ferrari and Russo provide an excellent overview of DAX that is both easy to understand and comprehensive.

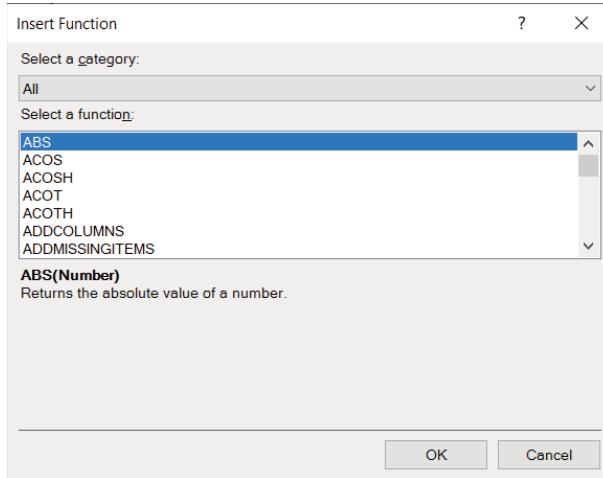
Identifying DAX functions safe for calculated columns

In the previous section, you used the Formula bar within the Power Pivot window to enter calculations. Next to that Formula bar, you may have noticed the Insert Function button labeled *fx*. This is similar to the Insert Function button found in Excel. Clicking this button will

activate the Insert Function dialog box shown in Figure 34.7. This dialog box allows you to browse, search for, and insert the available DAX functions.

FIGURE 34.7

The Insert Function dialog box shows you all available DAX functions.



As you look through the list of DAX functions, you'll notice that many of them look like the Excel functions with which you are already familiar. Make no mistake, however; these are *not* Excel functions. Where Excel functions work with cells and ranges, these DAX functions are designed to work at the table and column levels.

To understand what that means, start a new calculated column on the Invoice Details tab. Click the Formula bar, and enter a good old SUM function: =SUM([Gross Margin]). Figure 34.8 shows the result.

FIGURE 34.8

The DAX SUM function can only sum the column as a whole.

Entity	UnitCost	UnitPrice	Total Revenue	Gross Margin	Calculated Column 1
1	59.29	119.95	\$119.95	60.66	928378.069999998
1	3290.55	6589.95	\$6,589.95	3299.4	928378.069999998
10	35	34.95	\$349.50	-0.5	928378.069999998
50	91.59	189.95	\$9,497.50	4918	928378.069999998
1	59.29	119.95	\$119.95	60.66	928378.069999998
1	674.5	1349.95	\$1,349.95	675.45	928378.069999998
1	91.25	189.95	\$189.95	98.7	928378.069999998
1	303.85	609.95	\$609.95	306.1	928378.069999998

As you can see, the SUM function sums the entire column. This is because Power Pivot and DAX are designed to work with tables and columns. Power Pivot has no construct for cells and ranges. It doesn't even have column letters on its grid. Where you would normally reference a range, as in an Excel SUM function, DAX basically takes the entire column.

The bottom line is not all DAX functions can be used with calculated columns. Because a calculated column evaluates at the row level, only DAX functions that evaluate single data points can be used in a calculated column.

A good rule of thumb is if the function requires an array or a range of cells as an argument, then it's not viable in a calculated column. Therefore, functions such as SUM, MIN, MAX, AVERAGE, and COUNT don't work in calculated columns. Functions that require only single data point arguments work quite well in calculated columns, such as YEAR, MONTH, MID, LEFT, RIGHT, IF, and IFERROR.

Building DAX-driven calculated columns

To demonstrate the usefulness of employing a DAX function to enhance calculated columns, let's return to our walk-through example. Go to the Power Pivot window, and click the Invoice-Header tab. If you accidentally closed the Power Pivot window, you can activate it by clicking the Manage button on the Home tab of the Power Pivot Ribbon.

The InvoiceHeader table, shown in Figure 34.9, contains an InvoiceDate column. Although this column is valuable in the raw table, the individual dates aren't convenient when analyzing the data with a PivotTable. It would be beneficial to have a column for Month and a column for Year. This way, you could aggregate and analyze your data by month and year.

FIGURE 34.9

DAX functions can help enhance the InvoiceHeader data with Year and Month time dimensions.



InvoiceDate	InvoiceNu...	Custom...	Add Column
5/8/2005 12:00:00 AM	ORDST1025	BAKERSEM0001	
4/12/2007 12:00:00 AM	STDINV2251	BAKERSEM0001	
5/8/2005 12:00:00 AM	ORDT1026	AARONFIT0001	
4/12/2007 12:00:00 AM	STDINV2252	AARONFIT0001	
5/7/2004 12:00:00 AM	ORD1002	METROPOL0001	
2/10/2004 12:00:00 AM	INV1024	AARONFIT0001	
2/15/2004 12:00:00 AM	INV1025	AARONFIT0001	
5/10/2004 12:00:00 AM	ORDPH1005	LECLERC0001	

For this endeavor, you use the YEAR(), MONTH(), and FORMAT() DAX functions to add some time dimensions to your data model. Follow these steps:

1. In the InvoiceHeader table, click the first blank cell in the empty column labeled Add Column on the far right.
2. In the Formula bar, type = YEAR([InvoiceDate]), and then press Enter.
3. Power Pivot will automatically rename the column to Calculated Column
 1. Double-click the column label and rename the column Year.
4. Starting the next column, click the first blank cell in the empty column labeled Add Column on the far right.
5. In the Formula bar, type = MONTH([InvoiceDate]) and then press Enter.
6. Power Pivot will automatically rename the column to Calculated Column
 1. Double-click the column label and rename the column Month.
7. Starting the next column, click the first blank cell in the empty column labeled Add Column on the far right.
8. In the Formula bar, type = FORMAT([InvoiceDate],"mmm") and then press Enter.
9. Power Pivot will automatically rename the column to Calculated Column
 1. Double-click the column label and rename the column Month Name.

After going through these steps, you should have three new calculated columns similar to the ones shown in Figure 34.10.

FIGURE 34.10

Using DAX functions to supplement a table with Year, Month, and Month Name columns

	InvoiceDate	InvoiceNu...	Custom...	Year	Month	Month Name
1	5/8/2018 12:00:00 AM	ORDST1025	BAKERSEM0001	2018	5	May
2	4/12/2020 12:00:00 AM	STDINV2251	BAKERSEM0001	2020	4	Apr
3	5/8/2018 12:00:00 AM	ORDST1026	AARONFIT0001	2018	5	May
4	4/12/2020 12:00:00 AM	STDINV2252	AARONFIT0001	2020	4	Apr
5	5/7/2017 12:00:00 AM	ORDT002	METROPOL0001	2017	5	May
6	2/10/2017 12:00:00 AM	INV1024	AARONFIT0001	2017	2	Feb
7	2/15/2017 12:00:00 AM	INV1025	AARONFIT0001	2017	2	Feb
8	5/10/2017 12:00:00 AM	ORDPH1005	LECLERC0001	2017	5	May

As mentioned, creating calculated columns will automatically make them available through your PivotTable Fields list (see Figure 34.11).

Month sorting in Power Pivot-driven PivotTables

One of the more annoying things about Power Pivot is that it doesn't inherently know how to sort months. Unlike standard Excel, Power Pivot doesn't use the built-in custom lists that define the order of month names. When you create a calculated column like [Month Name]

Part V: Understanding Power Pivot and Power Query

and place it into your PivotTable, Power Pivot will put those months in alphabetical order. Figure 34.12 illustrates this in a PivotTable designed to show average revenue by month.

FIGURE 34.11

DAX calculations are immediately available in any connected PivotTable.

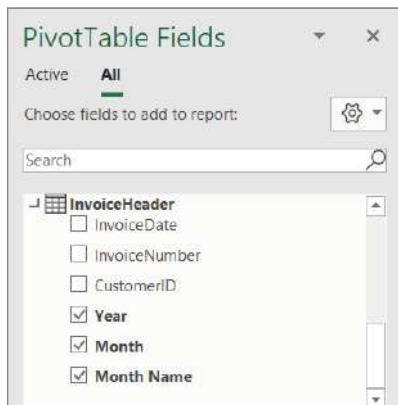


FIGURE 34.12

Month names in Power Pivot–driven PivotTables don't automatically sort in month order.

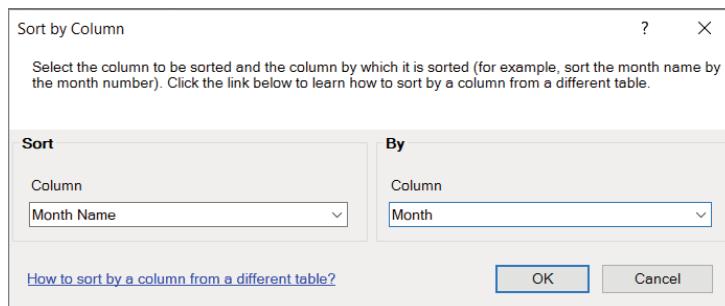
A screenshot of a Power Pivot-driven PivotTable. The table has 'Row Labels' and 'Average of Total Revenue' as columns. It shows data for four companies: Aaron Fitz Electrical, Adam Park Resort, Advanced Paper Co., and Advanced Tech Satellite System. The months listed are Apr, Feb, Jan, Mar, May, and Sep, which are not in chronological order.

Row Labels	Average of Total Revenue
Aaron Fitz Electrical	
Apr	\$1,121.88
Feb	\$826.75
Jan	\$396.49
Mar	\$248.14
May	\$1,829.83
Sep	\$59.95
Adam Park Resort	
Apr	\$839.92
Jan	\$599.50
May	\$59.90
Sep	\$2,399.95
Advanced Paper Co.	
Apr	\$13,049.13
Jan	\$359.80
Advanced Tech Satellite System	
Apr	\$49.95
Jul	\$138.24
May	\$949.75

The fix for this is easy. Activate the Power Pivot window and select the Home tab. There you will find the Sort by Column button. Upon clicking the button, the Sort by Column dialog box shown in Figure 34.13 opens.

FIGURE 34.13

The Sort by Column dialog box lets you define how your columns are sorted.



The idea is to select the column you want sorted and then select the column by which you want to sort. In this scenario, we want to sort Month Name by month.

Once you confirm the change, it will initially appear as though nothing happened. This is because the sort order that you defined is not for the Power Pivot window. The sort order is applied to your PivotTable. You can flip over to Excel and see the result in the PivotTable (see Figure 34.14).

FIGURE 34.14

Your month names now show in correct month order.

Row Labels	Average of Total Revenue
Aaron Fitz Electrical	
Jan	\$396.49
Feb	\$826.75
Mar	\$248.14
Apr	\$1,121.88
May	\$1,829.83
Sep	\$59.95
Adam Park Resort	
Jan	\$599.50
Apr	\$839.92
May	\$59.90
Sep	\$2,399.95
Advanced Paper Co.	
Jan	\$359.80
Apr	\$13,049.13
Advanced Tech Satellite System	
Apr	\$49.95
May	\$949.75
Jul	\$138.24

NOTE

PivotTables based on your data model will first inherit the formatting and sorting explicitly applied in the data model, then will apply any formatting you set in the PivotTable itself. In other words, any formatting you apply in the PivotTable itself will supersede the formatting and sorting applied via the Power Pivot window.

Referencing fields from other tables

Sometimes the operation you are trying to perform with a calculated column requires you to utilize fields from other tables within the Power Pivot data model. For instance, you may need to account for a customer-specific discount amount from the Customers table (see Figure 34.15) when creating a calculated column in the InvoiceDetails table.

FIGURE 34.15

The Discount Amount value in the Customers table can be used in a calculated column in another table.

	CustomerID	CustomerName	Discount Amount	Address
1	DOLLISCO0001	Dollis Cove Resort	11%	765 Kingway
2	GETAWAY0001	Getaway Inn	10%	234 E Cannon Ave.
3	HOMEFURN0001	Home Furnishings Limited	25%	234 Heritage Ave.
4	JOHNSONK0001	Johnson, Kimberly	12%	5678 S. 42nd Ave.
5	KELLYCON0001	Kelly Consulting	5%	123 Yeo
6	KENSINGT0001	Kensington Gardens Resort	13%	12345 Redmond Rd
7	HAMPTONV0001	Hampton Village Eatery	20%	234 Hampton Village
8	HEALTHYC0001	Healthy Concepts	11%	1234 Westown Road

To accomplish this, you can use a DAX function called RELATED. Similar to VLOOKUP in standard Excel, the RELATED function allows you to look up values from one table to use them in another.

Take a moment to follow these steps to create a new calculated column that displays a discounted revenue for each transaction in the InvoiceDetails table:

1. In the InvoiceDetails table, click the first blank cell in the empty column labeled Add Column on the far right.
2. In the Formula bar, type the following:

=RELATED (

As soon as you enter the open parenthesis, a menu of available fields (shown in Figure 34.16) will be displayed. Note that the items in the list represent the table name followed by the field name in brackets. In this case, we're interested in the Customers [Discount Amount] field.

FIGURE 34.16

Use the RELATED function to look up a field from another table.

The screenshot shows the Power Pivot ribbon with the formula bar containing the formula =RELATED(). A dropdown menu is open next to the formula, listing various fields from the 'Customers' table: Address, City, Country, CustomerID, CustomerName, Discount Amount, State, and Zip. The 'Discount Amount' option is highlighted.

InvoiceN...	Quantity	RELATED(Column Name)	Total Revenue
1 ORDST1022	1	Customers[Address]	\$119.95
2 ORDST1015	1	Customers[City]	\$589.95
3 ORDST1016	10	Customers[Country]	\$349.50
4 ORDST1017	50	Customers[CustomerID]	\$497.50
5 ORDST1018	1	Customers[CustomerName]	\$119.95
6 INV1010	1	Customers[Discount Amount]	\$349.95
7 INV1011	1	Customers[State]	\$189.95
8 INV1012	1	Customers[Zip]	\$609.95
9 ORDST1020	1	InvoiceHeader[CustomerID]	\$119.95
10 ORDST1021	1	InvoiceHeader[InvoiceDate]	\$119.95
11 INV1015	1	InvoiceHeader[InvoiceNumber]	\$359.95

3. Double-click the Customers[Discount Amount] field and then press Enter.
4. Power Pivot will automatically rename the column to Calculated Column
 1. Double-click the column label and rename the column Discount%.
5. Starting the next column, click the first blank cell in the empty column labeled Add Column on the far right.
6. In the Formula bar, type = [UnitPrice] * [Quantity] * (1 - [Discount%]) and then press Enter.
7. Power Pivot will automatically rename the column to Calculated Column
 1. Double-click the column label and rename the column Discounted Revenue.

The reward for your efforts will be a new column that uses the discount percent from the Customers table to calculate discounted revenue for each transaction. Figure 34.17 illustrates the new calculated column.

FIGURE 34.17

The final Discounted Revenue calculated column using the Discount% column from the Customers table

The screenshot shows the Power Pivot ribbon with the formula bar containing the formula =RELATED(Columns[Discount Amount]). A dropdown menu is open next to the formula, listing various fields from the 'Customers' table: Address, City, Country, CustomerID, CustomerName, Discount Amount, State, and Zip. The 'Discount Amount' option is highlighted.

	UnitCost	UnitPrice	Total Revenue	Discount%	Discounted Revenue
1	59.29	119.95	\$119.95	13.00%	\$104.36
1	3290.55	6589.95	\$6,589.95	11.00%	\$5,865.06
10	35	34.95	\$349.50	9.00%	\$318.05
50	91.59	189.95	\$9,497.50	6.00%	\$8,927.65
1	59.29	119.95	\$119.95	9.00%	\$109.15
1	674.5	1349.95	\$1,349.95	15.00%	\$1,147.46
1	91.25	189.95	\$189.95	17.00%	\$157.66
1	303.85	609.95	\$609.95	5.00%	\$579.45

Note

The RELATED function leverages the relationships that you defined when creating the data model to perform the lookup. So, this list of choices will contain only the fields that are available based on the relationships you defined.

Nesting functions

In the previous example, you first created a Discount% column using the RELATED function and then used that column in another calculated column to calculate the discounted revenue.

It's important to note that you don't necessarily have to create multiple calculated columns to accomplish a task like this. You could, instead, nest the RELATED function into the discounted revenue calculation. This would be the syntax for the nested calculation:

```
= [UnitPrice] * [Quantity] * (1-RELATED(Customers[Discount Amount]))
```

As you can see, nesting simply means embedding desired functions within a calculation. In this case, instead of using the RELATED function in a separate Discount% field, you can embed it directly in your discounted revenue calculation.

Nesting functions can definitely save time and even improve performance in larger data models. On the other hand, complicated nested functions can be harder to read and understand.

Understanding Calculated Measures

You can enhance the functionality of your Power Pivot reports with another kind of calculation called a *calculated measure*. Calculated measures are used to perform more complex calculations that work on an aggregation of data. These calculations are not applied to the Power Pivot window like calculated columns. Instead, they are applied directly to your PivotTable, creating a sort of virtual column that can't be seen in the Power Pivot window. You use calculated measures when you need to calculate based on an aggregated grouping of rows.

Imagine you wanted to show the difference in unit costs between the years 2020 and 2019 for each of your customers. Think about what technically has to be done to achieve this calculation. You'd have to figure out the sum of unit costs for 2020, then you'd have to get the sum of unit costs for 2019, and finally you'd have to subtract the sum of 2019 from the sum of 2020. This is a calculation that simply can't be done using calculated columns. Using calculated measures is the only way to get the cost variance between 2020 and 2019.

Follow these steps to create a calculated measure:

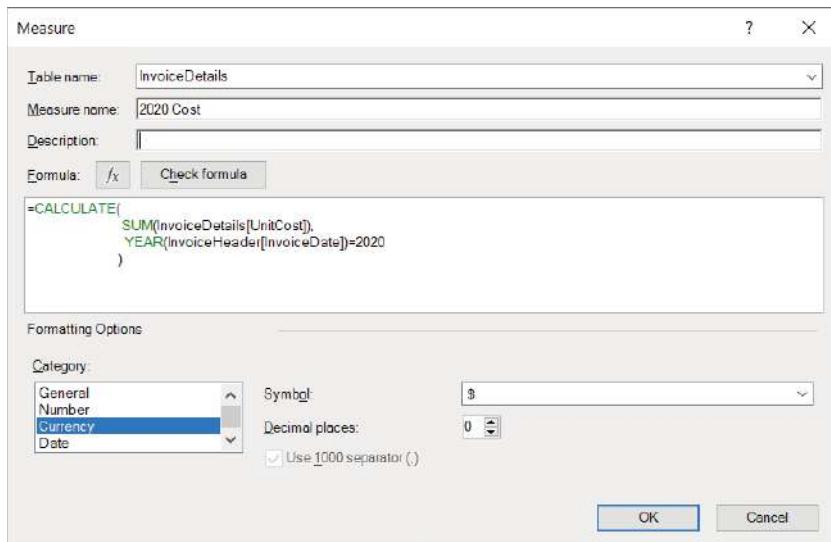
1. Start with a PivotTable created from a Power Pivot data model.
2. Click the Power Pivot tab in the Excel Ribbon and select Measures → New Measure. This will open the Measure dialog box shown in Figure 34.18.

NOTE

The sample file for this chapter contains a tab called Calculated Measures with a PivotTable already created.

FIGURE 34.18

Creating a new calculated measure



3. In the Measure dialog box, input the following information:

Table name Choose the table that you want to contain the calculated measure when looking at the PivotTable Fields list. Don't sweat this decision too much. The table you select has no bearing on how the calculation works. It's just a preference on where you want to see the new calculation within the PivotTable Fields list.

Measure name Give your calculated measure a descriptive name.

Formula Enter the DAX formula that will calculate the results of your new field. In this example, we are using the following DAX formula to calculate the sum of the UnitCost column from the InvoiceDetails table, where the InvoiceDate Year is equal to 2020:

```
=CALCULATE(
    SUM(InvoiceDetails[UnitCost]),
    YEAR(InvoiceHeader[InvoiceDate])=2020
)
```

Formatting Options Specify the formatting for the calculated measure results.

Part V: Understanding Power Pivot and Power Query

This formula uses the CALCULATE function to sum the UnitCost column from the InvoiceDetails table, where the Year column in the InvoiceHeader is equal to 2020.

4. **Click the Check Formula button to ensure that there are no syntax errors.** If your formula is well formed, you will see the message “No errors in formula.” If there are errors, you will see a full description of the errors.
5. **Click the OK button to confirm your changes and close the Measure dialog box.** You will immediately see your newly created calculated measure in the PivotTable.
6. **Repeat steps 2–5 for any other calculated measure that you need to create.**

In this example, we need a measure to show the 2019 cost:

```
=CALCULATE(  
    SUM(InvoiceDetails[UnitCost]),  
    YEAR(InvoiceHeader[InvoiceDate])=2019  
)
```

We will also need a measure to calculate 2020 vs. 2019:

```
=[2020 Cost] - [2019 Cost]
```

Note

You can use the results of a measure as an argument in a new measure. In this example, both [2020 Cost] and [2019 Cost] can be used in a new variance measure. Simply enter open square bracket ([) while entering your formula to see a list of previously created measures.

Figure 34.19 illustrates the newly created calculated measures. The calculated measures are applied to each customer, displaying the variance between their 2020 and 2019 costs. As you can see, each calculated measure is available for selection in the PivotTable Fields list.

FIGURE 34.19

Calculated measures can be seen in the PivotTable Fields list.

The screenshot shows the 'PivotTable Fields' pane in Excel. On the left, there is a table of data with columns labeled '2020 Cost', '2019 Cost', and '2020 vs 2019'. The '2020 vs 2019' column contains values like '\$4,683', '(\$64)', '\$0', '\$39', '\$324', '(\$2,731)', '\$1,859', '\$4,256', '\$416', '(\$6,163)', '\$8', '\$59', and '\$33,228'. To the right of the table is the 'PivotTable Fields' pane. At the top, it says 'Active' and 'All' with a green bar. Below that is a section titled 'Choose fields to add to report:' with a 'Search' input field. A list of measures is shown with checkboxes:

- fx 2020 Cost
- fx 2019 Cost
- fx 2020 vs 2019

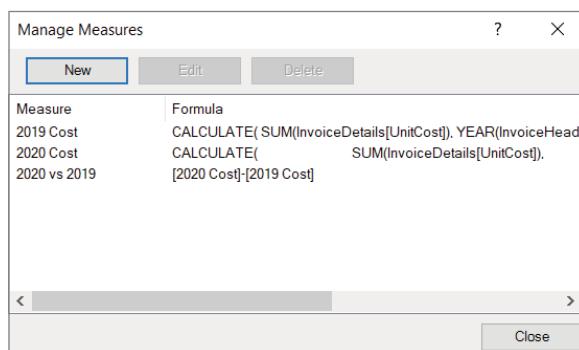
Editing and deleting calculated measures

You may find that you need either to edit or to delete a calculated measure. You can do so by following these steps:

1. Click anywhere inside your PivotTable and then click the Power Pivot tab in the Excel Ribbon and select Measures \Rightarrow Manage Measures. This will open the Manage Measures dialog box, shown in Figure 34.20.

FIGURE 34.20

The Manage Measures dialog box lets you edit or delete your calculated measures.



2. Select the target calculated measure and click either the Edit or Delete button.

Clicking the Edit button will activate the Measure dialog box where you can make changes to the calculation setting.

Clicking the Delete button will activate a message box asking you to confirm that you want to remove the measure. After you confirm with Yes, the calculated measure will be removed.

Using Cube Functions to Free Your Data

Cube functions are Excel functions that can be used to access the data in a Power Pivot data model outside the constraints of a PivotTable. Although Cube functions are technically not used to create calculations themselves, they can be used to free Power Pivot data so that it can be used with formulas that you may have in other parts of your Excel spreadsheet.

One of the easiest ways to start exploring Cube functions is to allow Excel to convert your PivotTable into Cube functions. The idea is to tell Excel to replace all cells in the PivotTable with a formula that connects back to the Power Pivot data model.

Follow these steps to create your first set of Cube functions:

1. Start with a PivotTable created from a Power Pivot model.

Part V: Understanding Power Pivot and Power Query

NOTE

The sample file for this chapter contains a tab called Cube Functions with a PivotTable already created.

2. Select any cell inside the PivotTable and then select PivotTable Analyze ↴ Calculations ↴ OLAP Tools ↴ Convert to Formulas.

After a second or two, the cells that used to house a PivotTable are now homes for Cube formulas. The Formula bar shown in Figure 34.21 illustrates a Cube function.

FIGURE 34.21

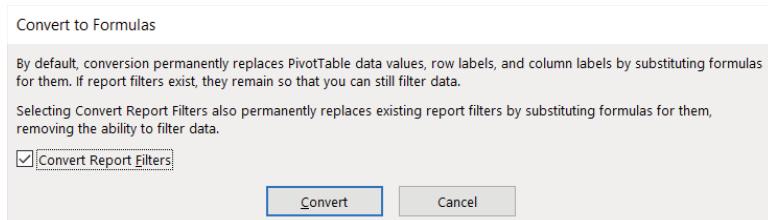
These cells are now a series of Cube functions!

A	B	C	D	E	F
3	Month Name	Sum of UnitCost	Sum of Quantity	Sum of UnitPrice	
4	Jan	\$44,072	420	\$88,034	
5	Feb	\$65,118	316	\$130,813	
6	Mar	\$116,467	816	\$233,011	
7	Apr	\$117,605	1198	\$230,592	
8	May	\$165,886	476	\$331,537	
9	Jul	\$477	20	\$999	
10	Sep	\$1,251	3	\$2,520	
11	Grand Total	\$510,875	3249	\$1,017,506	

If your PivotTable contains a report filter field, the dialog box shown in Figure 34.22 activates. This dialog box gives you the option of converting your filter drop-down selectors to Cube formulas. If you select this option, the drop-down selectors are removed, leaving a static formula.

FIGURE 34.22

Excel gives you the option of converting your report filter fields.



If you need to have your filter drop-down selectors intact so that you can continue to change the selections in the filter field interactively, be sure to leave the Convert Report Filters option unchecked when clicking the Convert button.

Now that the values you see are no longer part of a PivotTable object, you can insert rows and columns, you can add your own calculations, or you can combine the data with other formulas in your spreadsheet. For instance, in Figure 34.23, you can see we added two Quarter totals rows in the middle of the pivot data. We're only able to do this because the pivot data is actually no longer a PivotTable object but rather cube formulas.

FIGURE 34.23

Cube functions give you the flexibility of restructuring your pivot data without losing the link to your data model.

B	C	D	E
Month Name	Sum of UnitCost	Sum of Quantity	Sum of UnitPrice
Jan	\$44,072	420	\$88,034
Feb	\$65,118	316	\$130,813
Mar	\$116,467	816	\$233,011
Quarter totals	\$225,656	\$1,552	\$451,858
Apr	\$117,605	1198	\$230,592
May	\$165,886	476	\$331,537
Jul	\$477	20	\$999
Quarter totals	\$283,969	\$1,694	\$563,129

Introducing Power Query

IN THIS CHAPTER

- Understanding Power Query basics
- Understanding query steps
- Managing existing queries
- Getting an overview of query actions
- Using and managing external sources
- Data profiling with Power Query

In information management, ETL refers to the three separate functions typically required to integrate disparate data sources: extraction, transformation, and loading. The *extraction* function refers to the reading of data from a specified source and extracting a desired subset of data. The *transformation* function refers to the cleaning, shaping, and aggregating of data to convert it to the desired structure. The *loading* function refers to the actual importing or writing of the resulting data to a target location.

Excel analysts have been manually performing ETL processes for years—although they rarely call it ETL. Every day, millions of Excel users manually pull data from some source location, manipulate that data, and integrate it into their reporting. This amounts to lots of manual effort.

Power Query enhances the ETL experience by offering an intuitive mechanism to extract data from a wide variety of sources, perform complex transformations on that data, and then load the data into a workbook or the internal data model.

In this chapter, you'll explore the basics of Power Query and get a glimpse of how it helps you save time and automate the steps needed to ensure that clean data is imported into your reporting models.

Understanding Power Query Basics

To start this basic look at Power Query, let's walk through a simple example. Imagine that you need to import Microsoft Corporation stock prices for the past 30 days using Yahoo Finance. For this scenario, you need to perform a web query to pull the data needed from Yahoo Finance.

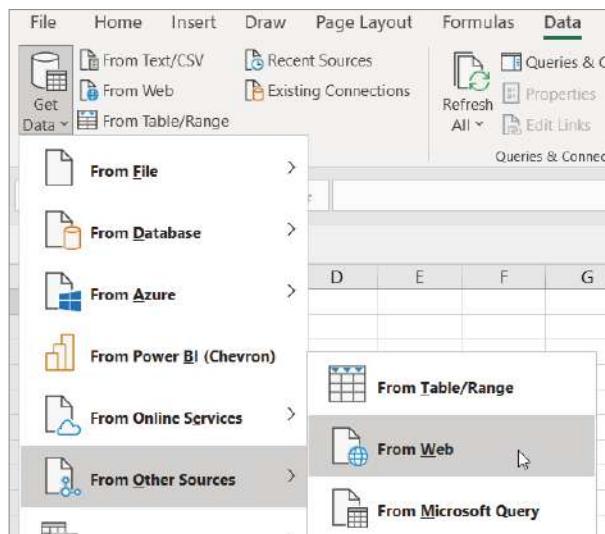
Part V: Understanding Power Pivot and Power Query

To start your query, follow these steps:

1. In Excel, select Data \Rightarrow Get & Transform Data \Rightarrow Get Data and then select From Other Sources \Rightarrow From Web (see Figure 35.1).

FIGURE 35.1

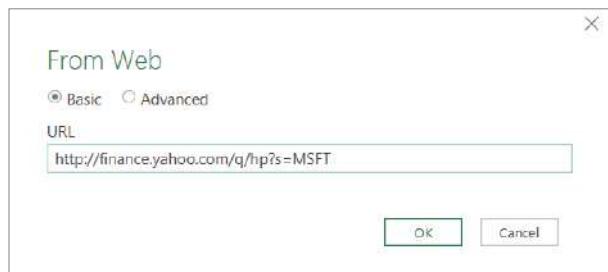
Starting a Power Query web query



2. In the dialog box that appears (see Figure 35.2), enter the URL for the data that you need, in this case <http://finance.yahoo.com/q/hp?s=MSFT>.

FIGURE 35.2

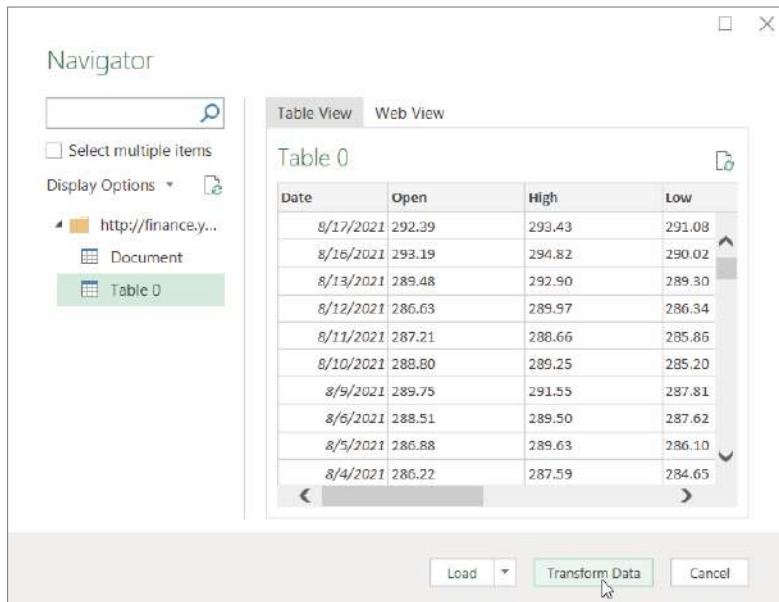
Enter the target URL for the source containing the data you need.



After you click OK, you'll notice a bit of gyrating, and then the Navigator pane, shown in Figure 35.3, appears. Here, you select the data source you want extracted. You can also click each table to see a preview of the data.

FIGURE 35.3

Select the correct data source and then click the Transform Data button.



3. In this case, the table labeled Table 0 holds the historical stock data you need, so click Table 0 and then click the Transform Data button.

NOTE

You may have noticed that the Navigator pane shown in Figure 35.3 offers a Load button (next to the Transform Data button). The Load button allows you to skip any editing and import your targeted data as is. If you are sure you will not need to transform or shape your data in any way, you can opt to click the Load button to import the data directly into the data model or a spreadsheet in your workbook.

CAUTION

Excel has another From Web command button on the Data tab next to the Get Data command. This unfortunate duplicate command is actually the legacy web scraping capability found in all Excel versions going back to Excel 2000.

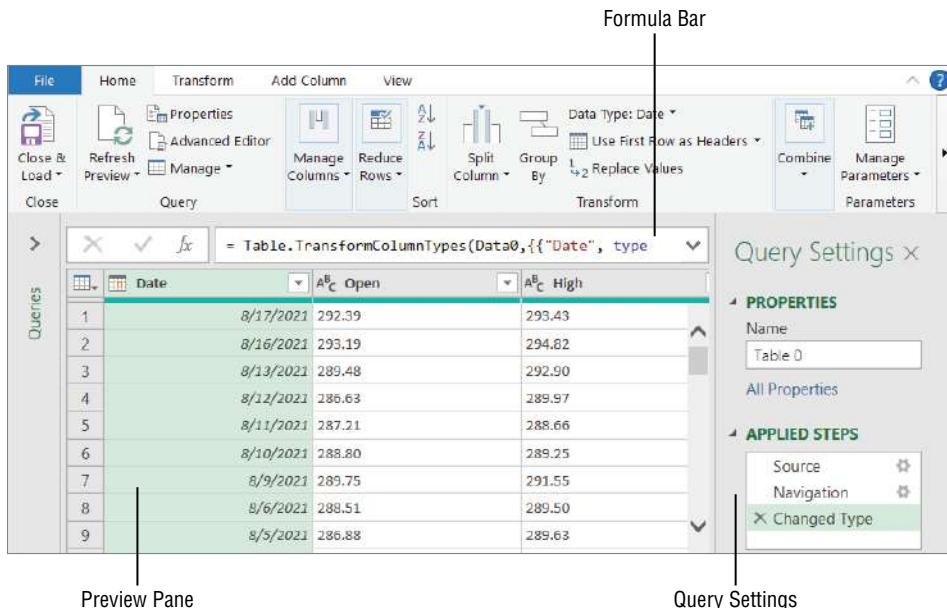
The Power Query version of the From Web command (found under the Get Data drop-down menu) goes beyond simple web scraping. Power Query is able to pull data from advanced web pages, and it is able to manipulate the data. Make sure you are using the correct feature when pulling data from the Web.

Part V: Understanding Power Pivot and Power Query

When you click the Transform Data button, Power Query activates a new Power Query Editor window, which contains its own Ribbon and a preview pane that shows a preview of the data (see Figure 35.4). Here, you can apply certain actions to shape, clean, and transform the data before importing.

FIGURE 35.4

The Power Query Editor window allows you to shape, clean, and transform data.

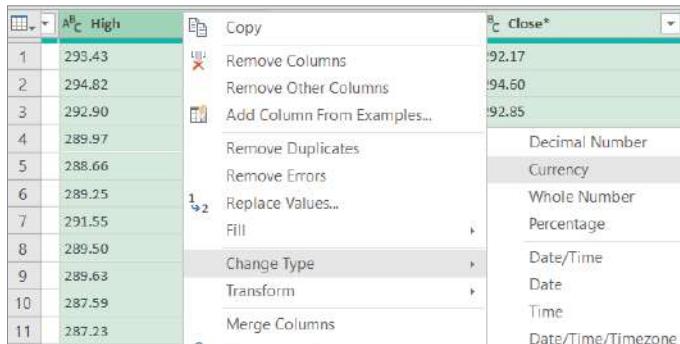


The idea is to work with each column shown in the Power Query Editor, applying the necessary actions that will give you the data and structure you need. You'll dive deeper into column actions later in this chapter. For now, you need to continue toward the goal of getting the last 30 days of stock prices for Microsoft Corporation.

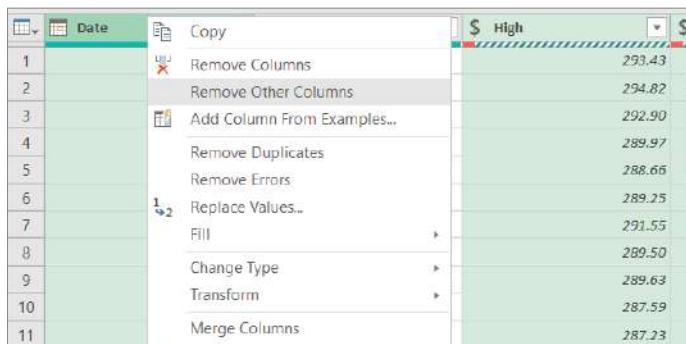
4. Click the High column, then hold down the Ctrl key on your keyboard while you click the Low and Close columns.
5. Right-click and select Change Type>Currency, as shown in Figure 35.5.
6. Power Query will ask if you want to replace the current step or add a new step. Choose the Add New Step button.
7. Remove all the columns that you do not need by right-clicking each one and clicking Remove. (Besides the Date column, the only other columns that you need are the High, Low, and Close columns.) Alternatively, you can hold down the Ctrl key on your keyboard, select the columns you want to keep, right-click any of the selected column headers, and then choose Remove Other Columns (see Figure 35.6).

FIGURE 35.5

Change the data type of the High, Low, and Close columns to currency format.

**FIGURE 35.6**

Select the columns you want to keep and then select Remove Other Columns to get rid of the other columns.



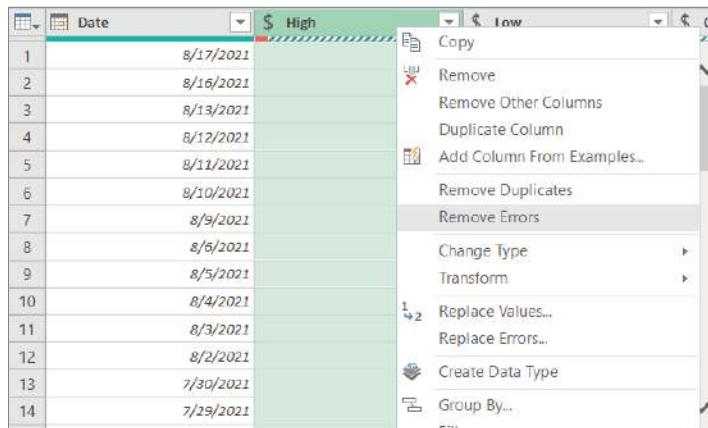
You may notice that some of the rows show the word *Error*. These are rows that contained text values that could not be converted.

8. Remove the Error rows by right-clicking the High column and selecting Remove Errors, as shown in Figure 35.7.
9. After all the errors are removed, add a Week Of column that displays the week to which each date in the table belongs. To do this, right-click the Date column and select the Duplicate Column option. A new column (named Date - Copy) is added to the preview.
10. Right-click the newly added column, select the Rename option, and then rename the column Week Of.

Part V: Understanding Power Pivot and Power Query

FIGURE 35.7

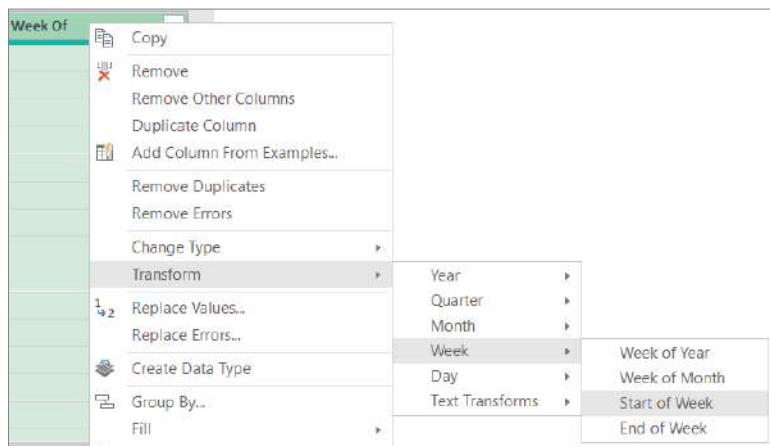
Removing errors caused by text values that could not be converted to currency



11. Right-click the Week Of column you just created and select Transform ↴ Week
↳ Start of Week, as shown in Figure 35.8. Excel transforms the date to display the start of the week for a given date.

FIGURE 35.8

The Power Query Editor can be used to apply transformation actions such as displaying the start of the week for a given date.

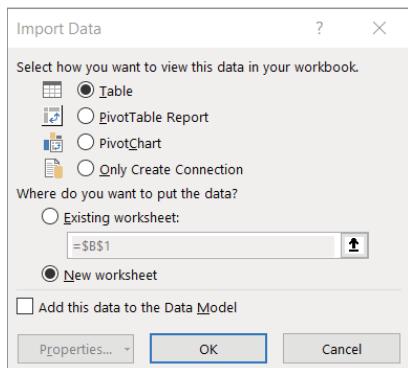


12. When you've finished configuring your Power Query feed, save and output the results. To do this, click the Close & Load drop-down found on the Home tab of the Power Query Editor Ribbon to reveal two options.

The Close & Load option saves your query and outputs the results as an Excel table to a new worksheet in your workbook. The Close & Load To option activates the Import Data dialog box (see Figure 35.9), where you can choose to output the results to a specific worksheet or to the internal data model.

FIGURE 35.9

The Import Data dialog box gives you more control over how the results of queries are used.



The Import Data dialog box also enables you to save the query as a query connection only, which means you will be able to use the query in various in-memory processes without actually needing to output the results anywhere. Select the New Worksheet option button to output your results as a table on a new worksheet in the active workbook.

At this point, you will have a table similar to the one shown in Figure 35.10, which can be used to produce the PivotTable you need.

Take a moment to appreciate what Power Query allowed you to do just now. With a few clicks, you searched the Internet, found some base data, shaped the data to keep only the columns you needed, and even manipulated that data to add an extra Week Of dimension to the base data. This is what Power Query is about: enabling you easily to extract, filter, and reshape data without the need for any programmatic coding skills.

TIP

You can get back to the Power Query Editor window for any query by activating the Queries & Connections task pane. On the Excel Ribbon, click Data ➔ Queries & Connections. From here, you can simply right-click your target query and select Edit.

Part V: Understanding Power Pivot and Power Query

FIGURE 35.10

Your final query pulled from the Internet: transformed, put into an Excel table, and ready to use in a PivotTable

	A	B	C	D	E
1	Date	High	Low	Close*	Week Of
2	8/17/2021	293.43	291.08	292.86	8/15/2021
3	8/16/2021	294.82	290.02	294.6	8/15/2021
4	8/13/2021	292.9	289.3	292.85	8/8/2021
5	8/12/2021	289.97	286.34	289.81	8/8/2021
6	8/11/2021	288.66	285.86	286.95	8/8/2021
7	8/10/2021	289.25	285.2	286.44	8/8/2021
8	8/9/2021	291.55	287.81	288.33	8/8/2021
9	8/6/2021	289.5	287.62	289.46	8/1/2021
10	8/5/2021	289.63	286.1	289.52	8/1/2021
11	8/4/2021	287.59	284.65	286.51	8/1/2021
12	8/3/2021	287.23	284	287.12	8/1/2021
13	8/2/2021	286.77	283.74	284.82	8/1/2021
14	7/30/2021	286.66	283.91	284.91	7/25/2021
15	7/29/2021	288.62	286.08	286.5	7/25/2021
16	7/28/2021	290.15	283.83	286.22	7/25/2021
17	7/27/2021	289.58	282.95	286.54	7/25/2021

Understanding query steps

Power Query uses its own formula language, known as the M language, to codify your queries. As with macro recording, each action you take when working with Power Query results in a line of code being written into a query step. Query steps are embedded M code, which allows your actions to be repeated each time you refresh your Power Query data.

To explore this concept, let's open the Power Query Editor for the table we just created. Right-click anywhere in the table shown in Figure 35.10 and select Table \leftrightarrow Edit Query. You can see the query steps for your queries in the Query Settings pane (see Figure 35.11).

Each query step represents an action that you took to get to a data table. You can click any step to see the underlying M code in the Power Query formula bar. For example, clicking the step called Removed Errors reveals the code for that step in the Formula bar.

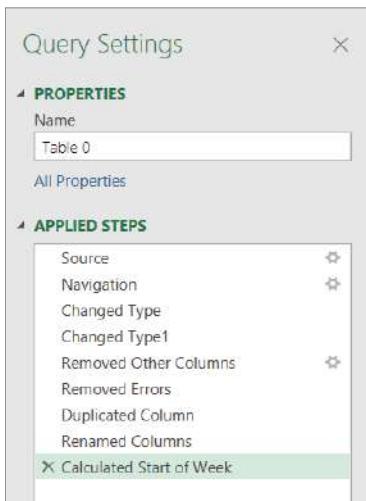
TIP

If you don't see the Query Settings pane, click the Query Settings command on the View tab of the Power Query Editor Ribbon. The View tab also contains the Formula Bar check box, allowing you to expose the Formula Bar that displays the M syntax for each given step.

When you click a query step, the data shown in the preview pane is a preview of what the data looked like up to and including the step you clicked. For example, clicking the step before the Removed Other Columns step in Figure 35.11 lets you see what the data looked like before you removed the nonessential columns.

FIGURE 35.11

Query steps can be viewed and managed in the Applied Steps section of the Query Settings pane.



You can right-click any step to see a menu of options for managing your query steps. Figure 35.12 illustrates the following options:

Edit Settings Edit the arguments or parameters that define the selected step.

Rename Give the selected step a meaningful name.

Delete Remove the selected step. Be aware that removing a step can cause errors if subsequent steps depend on the deleted step.

Delete Until End Remove the selected step and all following steps.

Insert Step After Insert a step after the selected step.

Move Before Move the selected step up in the order of steps.

Move After Move the selected step down in the order of steps.

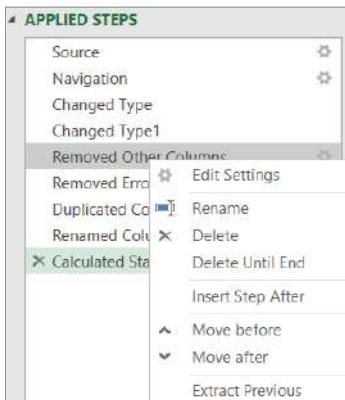
Extract Previous Extract all steps prior to this one into a new query.

Viewing the Advanced Query Editor

Power Query gives you the option of viewing and editing a query's embedded M code directly. While in the Power Query Editor window, click the View tab of the Ribbon and select Advanced Editor. The Advanced Editor dialog box is little more than a space for you to edit the existing M code or type your own M code. Advanced users can use the M language to extend the capabilities of Power Query by directly coding their own steps in the Advanced Editor.

FIGURE 35.12

Right-click any query step to edit, rename, delete, or move the step.



Refreshing Power Query data

It's important to note that Power Query data is not in any way connected to the source data used to extract it. A Power Query data table is merely a snapshot. In other words, as the source data changes, Power Query will not automatically keep up with the changes; you intentionally need to refresh your query.

If you chose to load your Power Query results to an Excel table in the existing workbook, you can manually refresh by right-clicking the table and selecting the Refresh option.

If you chose to load your Power Query data to the internal data model, you need to click Data \Rightarrow Queries & Connections and then right-click the target query and select the Refresh option.

To get a bit more automated with the refreshing of your queries, you can configure your data sources to refresh your Power Query data automatically. To do so, follow these steps:

1. **Go to the Data tab in the Excel Ribbon and click the Queries & Connections command.** The Queries & Connections task pane appears.
2. **Right-click the Power Query data connection that you want to refresh and then select the Properties option.**
3. **With the Query Properties dialog box open, select the Usage tab.**
4. **Set the options to refresh the chosen data connection.**

The refresh options are as follows:

Refresh Every X Minutes Placing a check next to this option tells Excel to refresh the chosen data automatically every specified number of minutes. Excel will refresh all tables associated with that connection.

Refresh Data When Opening the File Placing a check next to this option tells Excel to refresh the chosen data connection automatically upon opening the workbook. Excel will refresh all tables associated with that connection as soon as the workbook is opened.

These refresh options are useful when you want to ensure that your customers are working with the latest data. Of course, setting these options does not preclude the ability to refresh the data manually.

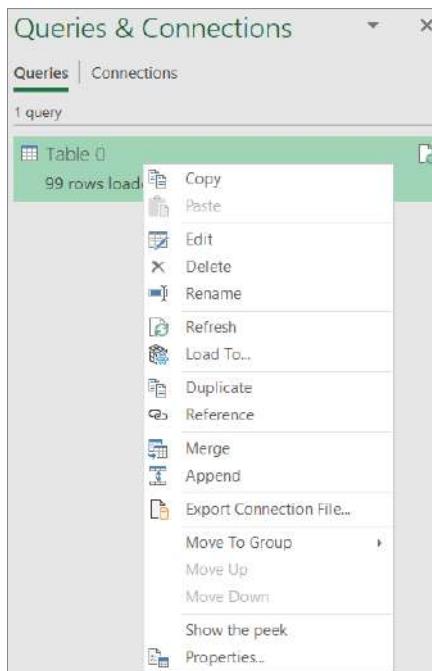
Managing existing queries

As you add various queries to a workbook, you will need a way to manage them. Excel accommodates this need by offering the Queries & Connections pane, which enables you to edit, duplicate, refresh, and generally manage all of the existing queries in the workbook. Activate the Queries & Connections pane by selecting the Queries & Connections command on the Data tab of the Excel Ribbon.

You need to find the query on which you want to work and then right-click it to take any one of the following actions (see Figure 35.13):

FIGURE 35.13

Right-click any query in the Queries & Connections pane to see the available management options.



- Edit** Open the Power Query Editor, where you can modify the query steps.
- Delete** Delete the selected query.
- Rename** Rename the selected query.
- Refresh** Refresh the data in the selected query.
- Load To** Activate the Import Data dialog box, where you can redefine where the selected query's results are used.
- Duplicate** Create a copy of the query.
- Reference** Create a new query that references the output of the original query.
- Merge** Merge the selected query with another query in the workbook by matching specified columns.
- Append** Append the results of another query in the workbook to the selected query.
- Send To Data Catalog** Publish and share the selected query via a Power BI server that your IT department sets up and manages.
- Export Connection File** Save an Office Data Connection (.odc) file with the connection credentials for the query's source data.
- Move To Group** Move the selected query into a logical group that you create for better organization.
- Move Up** Move the selected query up in the Queries & Connections pane.
- Move Down** Move the selected query down in the Queries & Connections pane.
- Show the Peek** Show a preview of the query results for the selected query.
- Properties** Rename the query and add a friendly description.

The Queries & Connections pane is especially useful when your workbook contains several queries. Think of it as a kind of table of contents that allows you to easily find and interact with the queries in your workbook.

Understanding column-level actions

Right-clicking a column in the Power Query Editor activates a shortcut menu that shows a full list of the actions you can take. You can also apply certain actions to multiple columns at one time by selecting two or more columns before right-clicking. Table 35.1 explains the commands you will see when right-clicking on a column within the Power Query Editor.

NOTE

Note that all the column-level actions available in Power Query are also available in the Power Query Editor Ribbon. Thus, you can either opt for the convenience of right-clicking to select an action quickly or choose to utilize the more visual Ribbon menu.

TABLE 35.1 Column-Level Actions

Action	Purpose	Available with Multiple Columns?
Remove	Remove the selected column from the Power Query data.	Yes
Remove Other Columns	Remove all non-selected columns from the Power Query data.	Yes
Duplicate Column	Create a duplicate of the selected column as a new column placed at the far right of the table. The name given to the new column is X - Copy, where X is the name of the original column.	No
Add Column From Examples	Create a custom column that combines data from other columns based on a few examples that you provide. Like the Flash Fill feature in Excel, Power Query's smart detection logic infers transformation logic based on your examples and then applies that logic to fill the new column.	Yes
Remove Duplicates	Remove all rows from the selected column where the values duplicate earlier values. The row with the first occurrence of a value is not removed.	Yes
Remove Errors	Remove rows containing errors in the selected column.	Yes
Change Type	Change the data type of the selected column.	Yes
Transform	Change the way values in the column are rendered. You can choose from the following options: Lowercase, Uppercase, Capitalize Each Word, Trim, Clean, Length, JSON, and XML. If the values in the column are date/time values, the options are as follows: Date, Time, Day, Month, Year, or Day of Week. If the values in the column are number values, the options are as follows: Round, Absolute Value, Factorial, Base-10 Logarithm, Natural Logarithm, Power, or Square Root.	Yes
Replace Values	Replace one value in the selected column with another specified value.	Yes
Replace Errors	Replace unsightly error values with your own friendlier text.	Yes
Create Data Type	Store multiple columns of data in one column as metadata, allowing you to expose all the data you need without taking up space in your worksheet. Excel formulas can interact with these rich data types to expose the stored data within.	Yes

Part V: Understanding Power Pivot and Power Query

TABLE 35.1 (continued)

Action	Purpose	Available with Multiple Columns?
Group By	Aggregate data by row values. For example, you can group by state and either count the number of cities in each state or sum the population of each state.	Yes
Fill	Fill empty cells in the column with the value of the first nonempty cell. You have the option of filling up or filling down.	Yes
Unpivot Columns	Transpose the selected columns from column-oriented to row-oriented or vice versa.	Yes
Unpivot Other Columns	Transpose the unselected columns from column-oriented to row-oriented or vice versa.	Yes
Unpivot Only Selected Columns	Transpose the selected columns from column-oriented to row-oriented or vice versa. This option also retains a columns list in the current step so that the same set of columns is unpivoted on future refresh operations.	Yes
Rename	Rename the selected column to a name you specify.	No
Move	Move the selected column to a different location in the table. You have these choices for moving the column: Left, Right, To Beginning, and To End.	Yes
Drill Down	Navigate to the contents of the column. This is used with tables that contain metadata representing embedded information.	No
Add as New Query	Create a new query with the contents of the column. This is done by referencing the original query in the new one. The name of the new query is the same as the column header of the selected column.	No
Split Column	Split the value of a single column into two or more columns, based on a number of characters or a given delimiter, such as a comma, semicolon, or tab. The Split Column action is available only when right-clicking a text column.	No
Merge Columns	Merge the values of two or more columns into a single column that contains a specified delimiter, such as a comma, semicolon, or tab. The Merge Columns action is available only when selecting two or more columns.	Yes

Understanding table actions

While you're in the Power Query Editor, Power Query allows you to apply certain actions to an entire data table. You can see the available table-level actions by clicking the Table Actions icon in the upper-left corner of the Power Query Editor preview pane, as shown in Figure 35.14.

FIGURE 35.14

Click the Table Actions icon to see the table-level actions you can use to transform the data.

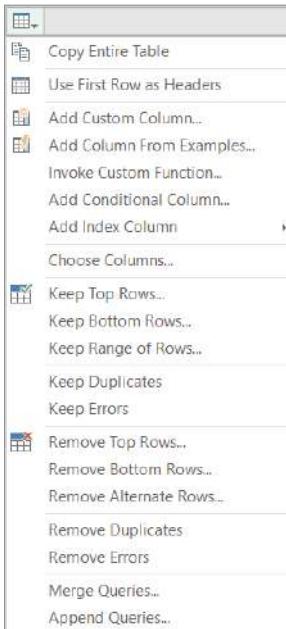


Table 35.2 lists the table-level actions and describes the primary purpose of each one.

TABLE 35.2 Table-Level Actions

Action	Purpose
Copy Entire Table	Copy the data within the current query to the Clipboard.
Use First Row as Headers	Replace each table header name with the value in the first row of each column.
Add Custom Column	Insert a new column after the last column of the table. The values in the new column are determined by the value or formula you define.
Add Column From Examples	Create a custom column that combines data from other columns based on a few examples that you provide. Like the Flash Fill feature in Excel, Power Query's smart detection logic infers transformation logic based on your examples and then applies that logic to fill the new column.
Invoke Custom Function	Insert a new column after the last column of the table and then run a user-defined function for each row in the column.
Add Conditional Column	Insert a new column after the last column of the table and then fill it with a conditional if-then-else statement that you define.

Part V: Understanding Power Pivot and Power Query

TABLE 35.2 (continued)

Action	Purpose
Add Index Column	Insert a new column containing a sequential list of numbers starting from 1, 0, or another specified value you define.
Choose Columns	Choose the columns you want to keep in the query results.
Keep Top Rows	Remove all but the top N number of rows. You specify the number threshold.
Keep Bottom Rows	Remove all but the bottom N number of rows. You specify the number threshold.
Keep Range of Rows	Remove all rows except the ones that fall within a range you specify.
Keep Duplicates	Remove all rows where the values in the selected columns are unique, enabling you to focus on the duplicate rows.
Keep Errors	Remove all rows that do not contain an error. This allows for the quick filtering of error values encountered during transformation.
Remove Top Rows	Remove the top N rows from the table.
Remove Bottom Rows	Remove the bottom N rows from the table.
Remove Alternate Rows	Remove alternate rows from the table, starting at the first row to remove and specifying the number of rows to remove and the number of rows to keep.
Remove Duplicates	Remove all rows where the values in the selected columns duplicate earlier values. The row with the first occurrence of a value set is not removed.
Remove Errors	Remove rows containing errors in the currently selected columns.
Merge Queries	Create a new query that merges the current table with another query in the workbook by matching specified columns.
Append Queries	Create a new query that appends the results of another query in the workbook to the current table.

NOTE

Note that all the table-level actions available in Power Query are also available in the Power Query Editor Ribbon. So, you can either opt for the convenience of right-clicking to select an action quickly or choose to utilize the more visual Ribbon menu.

Getting Data from External Sources

Microsoft has invested a great deal of time and resources in ensuring that Power Query has the ability to connect to a wide array of data sources. Whether you need to pull data from an

external website, a text file, a database system, Facebook, or a web service, Power Query can accommodate most, if not all, data sources.

You can see all the available connection types by clicking the Get Data drop-down on the Data tab in the Excel Ribbon. Power Query offers the ability to pull from a wide array of data sources.

From File Pulls data from a specified Excel file, text file, CSV file, XML file, JSON file, PDF file, or folder.

From Database Pulls data from databases like Microsoft Access, SQL Server, or SQL Server Analysis Services.

From Azure Pulls data from Microsoft's Azure Cloud Services.

From Power BI Pulls Power BI data sets to which your organization has granted you access.

From Online Services Pulls data from cloud application services such as Facebook, Salesforce, and Microsoft Dynamics online.

From Other Sources Pulls data from a wide array of Internet, cloud, and other ODBC data sources. Here, you will also find the Blank Query option. Selecting Blank Query will activate the Power Query Editor in the Advanced Editor view. This is handy when you want to copy and paste M code directly into the Power Query Editor.

In the rest of this chapter, you will explore the various connection types that can be leveraged to import external data.

Importing data from files

Organizational data is often kept in files such as text files, CSV files, and even other Excel workbooks. It's not uncommon to use these kinds of files as data sources for data analysis. Power Query offers several connection types that enable the importing of data from external files.

TIP

Keep in mind that the files you import don't necessarily have to be on your own PC. You can import files on network drives as well as cloud repositories such as Google Drive or Microsoft OneDrive.

NOTE

When starting with any of the data sources discussed here, Power Query will activate a set of dialog boxes customized for the selected connection. These dialog boxes ask for the basic parameters that Power Query needs to connect to the specified data source; they are parameters such as file path, URL, server name, credentials, and so forth.

Each connection type requires its own unique set of parameters, so their dialog boxes will be different. Luckily, Power Query rarely needs more than a handful of parameters to connect to any one data source, so the dialog boxes are relatively intuitive and hassle-free.

Getting data from Excel workbooks

You can import data from other Excel workbooks by going to the Excel Ribbon and selecting Data \Rightarrow Get & Transform Data \Rightarrow Get Data \Rightarrow From File \Rightarrow From Workbook.

Note that you can import any kind of Excel file, including macro-enabled workbooks and template workbooks. Power Query will not bring in charts, PivotTables, shapes, VBA code, or any other objects that may exist within a workbook. It simply imports the data found in the used cell ranges of the workbook.

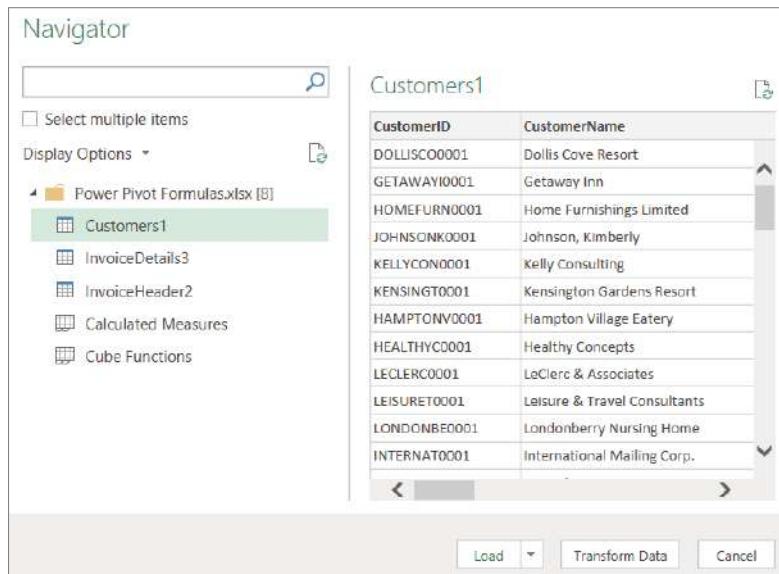
Once you've selected your file, the Navigator pane will activate, showing you all of the data sources available in the workbook. The idea here is to select the data source you want and then either load or transform the data using the buttons at the bottom of the Navigator pane. The Load button allows you to skip any editing and import your targeted data as is. Use the Transform Data button if you want to transform or shape the data before completing the import.

In terms of Excel workbooks, a data source is either a worksheet or a defined named range. The icons next to each data source let you distinguish those sources that are worksheets and those that are named ranges. In Figure 35.15, the source called Customers1 is a defined named range, while the source called Calculated Measures is a worksheet.

You can import multiple sources at once by clicking the Select Multiple Items check box and then placing a check next to each worksheet and named range that you want to import.

FIGURE 35.15

Select the data sources with which you want to work and then click the Load button.



Getting data from CSV and text files

Text files are commonly used to store and distribute data because of their inherent ability to hold many thousands of bytes of data without having an inflated file size. Text files can do this by forgoing all the fancy formatting, leaving only the text.

Comma-separated value (CSV) files are text files that contain commas to delimit (separate) values into columns of data.

To import a text or CSV file, go to the Excel Ribbon and select Data \Rightarrow Get & Transform Data \Rightarrow Get Data \Rightarrow From File \Rightarrow From Text/CSV. Excel will activate the Import Data dialog box where you can browse for and select a text or CSV file.

CAUTION

Excel has another From Text/CSV button on the Data tab next to the Get Data command. This duplicate command is actually the legacy import capability found in all Excel versions.

The Power Query version is much more powerful, allowing you to shape and transform text data before importing. Make sure you are using the correct Power Query version of the From Text/CSV feature.

Power Query will open the Power Query Editor to show you the contents of the text or CSV file that you just imported. The idea here is to apply any changes you want to make to the data and then click the Close & Load command on the Home tab to complete the import.

NOTE

Some text files are structured as tab-delimited files. Similar to CSV files, tab-delimited text files contain tab characters that separate text values into columns of data. Power Query will recognize tab-delimited text files and import these files into a table that contains a separate column for each tab delimiter.

Getting data from PDF files

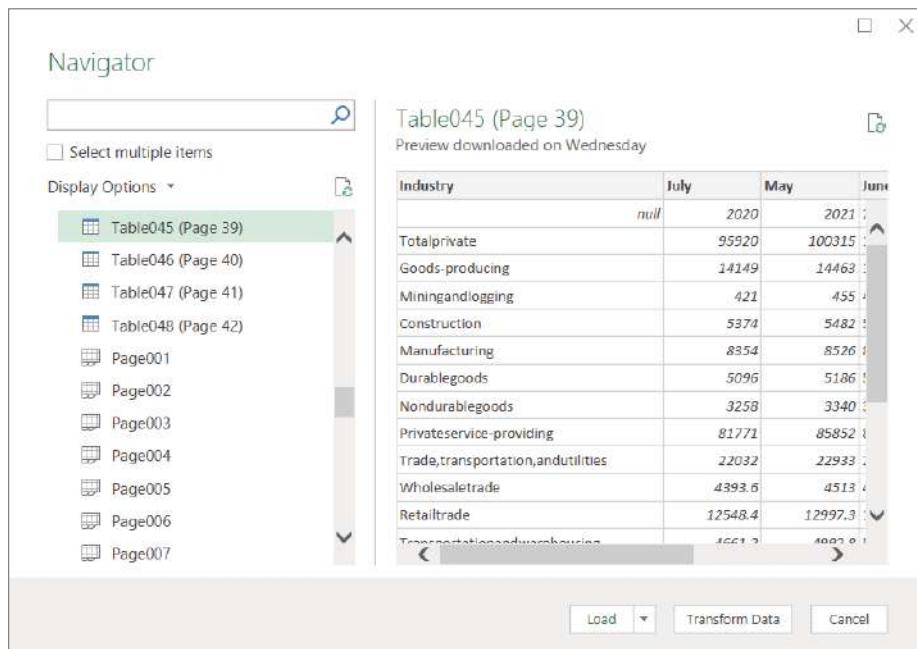
Power Query now offers the ability to import data from PDF files. You can access PDF data by going to the Excel Ribbon and selecting Data \Rightarrow Get & Transform Data \Rightarrow Get Data \Rightarrow From File \Rightarrow From PDF. The Import Data dialog box will activate, allowing you to browse for your target PDF file. After a few seconds, the Navigator pane shown in Figure 35.16 opens, showing you the available tables and pages found in your chosen file.

Notice that both structured tables and pages are shown, allowing you the option of importing a specific table or an entire page from the PDF. Simply click the item you want to import, and then click the Load button to import directly into your workbook or click the Transform Data button to clean the source data before importing.

You can even import multiple items from your PDF files by placing a check in the Select Multiple Items check box (see Figure 35.16).

FIGURE 35.16

The available tables and pages in the PDF are shown in the Navigator pane.



TIP

Rarely does the data from a PDF file come in clean. You'll find that you will almost always need to click the Transform Data button in order to clean up column names, remove empty spaces, and generally remove unwanted data elements.

Importing data from database systems

In larger organizations, the task of data management is not performed by Excel; rather, it is primarily performed by database systems such as Microsoft Access and SQL Server. Databases like these not only store millions of rows of data but also ensure data integrity, prevent redundancy, and allow for the rapid search and retrieval of data through queries and views.

Power Query offers options to connect to a wide array of database types. Microsoft has been keen to add connection types for as many commonly used databases as it can.

Importing data from relational and OLAP databases

Click Data \Rightarrow Get & Transform Data \Rightarrow Get Data \Rightarrow From Database, and you will see the list of databases to which you can connect. Power Query offers connection types for many of the popular database systems in use today: SQL Server, Microsoft Access, Oracle, MySQL, SQL Server Analysis Services, and so on.

Importing data from Azure databases

If your organization has a Microsoft Azure cloud database or a subscription to Microsoft Azure Marketplace, there is an entire set of connection types designed to import data from Azure. You can get to these connection types by clicking Data \Rightarrow Get & Transform Data \Rightarrow Get Data \Rightarrow From Azure.

Importing data using ODBC connections to nonstandard databases

Some of you may be using a unique nonstandard database system that isn't popular enough to be specifically included as an option under the Get Data command. Not to worry. As long as an ODBC connection string can be used to connect to your database system, Power Query can connect to it.

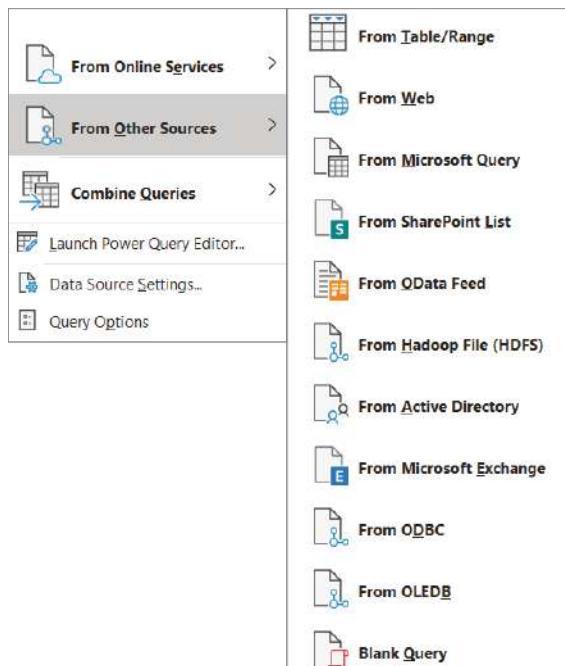
Click Data \Rightarrow Get & Transform Data \Rightarrow Get Data \Rightarrow From Other Sources to see a list of other connection types. Click the From ODBC option to start a connection to your unique database via an ODBC connection string.

Getting Data from Other Data Systems

In addition to ODBC, other kinds of data systems can be leveraged by Power Query, as shown in Figure 35.17.

FIGURE 35.17

Other systems Power Query can utilize as data sources



To get to the list shown in Figure 35.17, select Data \Rightarrow Get & Transform Data \Rightarrow Get Data \Rightarrow From Other Sources. Some of these data systems (SharePoint, Active Directory, and Microsoft Exchange) are popular ones that are used in many organizations to store data, track sales opportunities, and manage email. Other systems, like OData Feed and Hadoop, are less-common services used to work with large volumes of data. These are often mentioned in conversations about “Big Data.” Of course, the From Web option (demonstrated earlier in this chapter) is an integral connection type for any analyst who leverages data from the Internet.

Managing Data Source Settings

Each time you connect to any web-based data source or data source that requires some level of credentials, Power Query caches (stores) the settings for that data source.

For instance, let’s say you connected to a SQL Server database, entered all of your credentials, and imported the data you needed. At the moment of successful connection, Power Query caches information about that connection in a file located on your local PC. This includes connection string, username, password, privacy settings, and so on.

The purpose of all this caching is so that you don’t have to re-enter credentials each time you need to refresh your queries. That’s nifty, but what happens when your credentials are changed? Well, the short answer is those queries will fail until the data source settings are updated.

Editing data source settings

You can edit data source settings by activating the Data Source Settings dialog box. To do so, click Data \Rightarrow Get & Transform Data \Rightarrow Get Data \Rightarrow Data Source Settings.

The Data Source Settings dialog box, shown in Figure 35.18, contains a list of all credentials-based data sources previously used in queries. Select the data source you need to change and then click the Edit Permissions button.

Another dialog box will pop up, this one specific to the data source you selected (see Figure 35.19). This dialog box enables you to edit credentials as well as other data privacy settings.

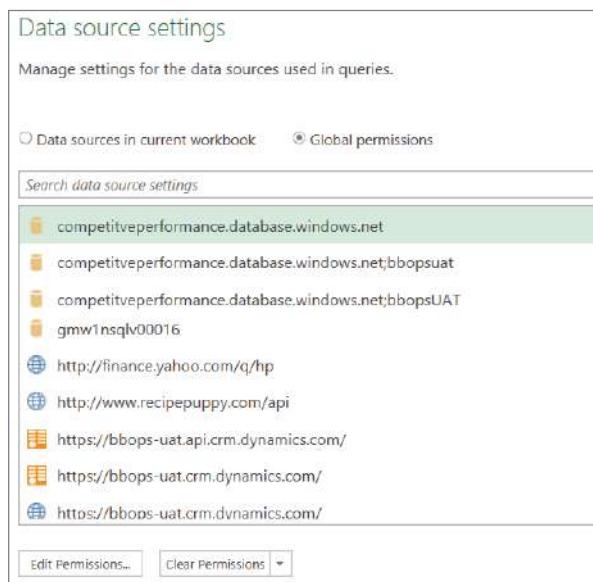
Click the Edit button in the Edit Permissions dialog box to make changes to the credentials for the data source. The credentials edit screen will differ based on the data source with which you are working, but again, the input dialog boxes are relatively intuitive and easy to update.

NOTE

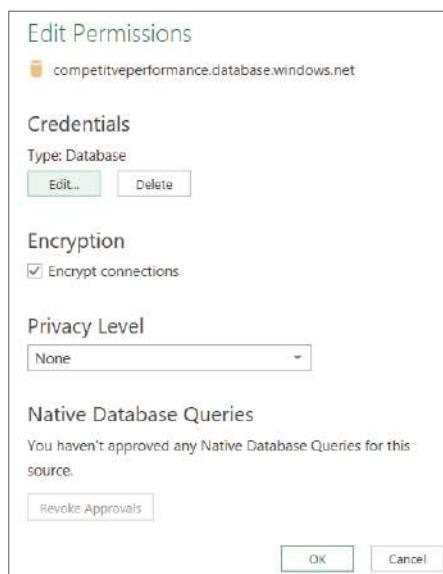
Power Query caches data source settings in a file located on your local PC. Even though you may have deleted a particular query, the data source setting is retained for possible future use. This can lead to a cluttered list of old and current data sources. You can clean out old items by selecting the data source in the Data Source Settings dialog box and clicking the Clear Permissions button.

FIGURE 35.18

Edit a data source by selecting it and clicking the Edit Permissions button.

**FIGURE 35.19**

The credentials edit screen for your selected data source



Data Profiling with Power Query

When importing a new data source, it's often useful to understand the intricacies and pitfalls of the data before you start working it. For instance, how many records are empty? How many unique values are there in a given column? What are the minimum and maximum values?

Power Query's data profiling capabilities allow you to know your data and identify potential issues before using it.

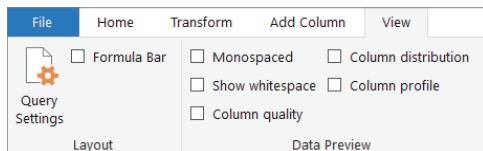
In the following sections, we'll explore some of the ways you can leverage data profiling in Power Query to get a better understanding of your data and address problem areas before they become a problem later in your reporting processes.

Data profiling options

While in the Power Query Editor window, click the View tab to see options for data profiling in the Data Preview group (see Figure 35.20).

FIGURE 35.20

Data Profiling options are found in the Data Preview group on the View tab.



Take a moment to review the purpose of each option:

Monospaced Converts the font in the preview pane to monospaced, making it easier to see differences in data.

Show whitespace Useful for calling out inline carriage returns and other invisible space characters.

Column Quality Displays the percentage of column values that are empty, the percentage that rendered as errors, and the percentage that are considered valid values. This option is the most powerful in terms of providing an at-a-glance view of your data.

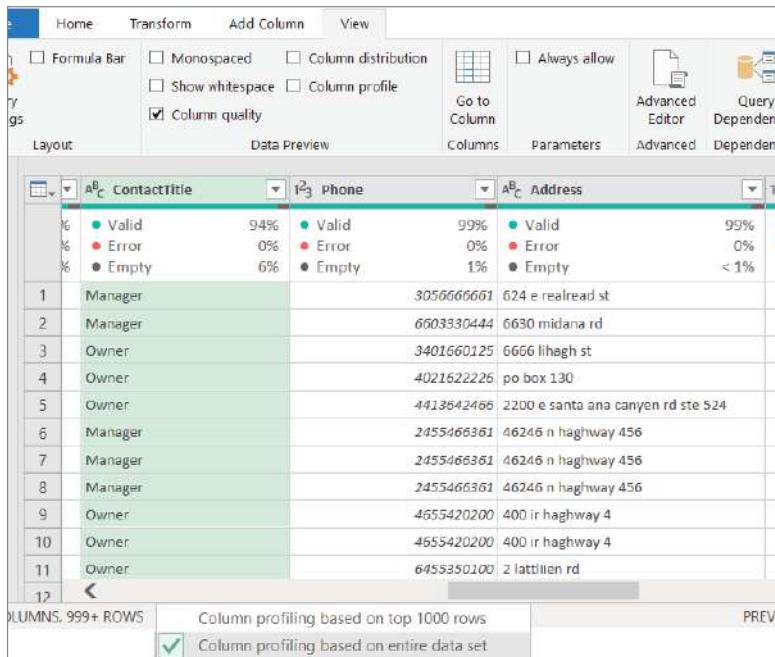
Column Distribution Provides a histogram visual displaying how many distinct and unique records are found in the values in each of the columns.

Column Profile Provides a useful way to see detailed descriptive statistics on a chosen column, such as the number of records with a 0 value, the minimum value in the column, the maximum value, the average value, and the standard deviation of all values in the column.

Be aware that the data profiler in Power Query, by default, only profiles the first 1,000 records. You can tell the profile to use the entire data set to get a more complete picture of your data. Figure 35.21 illustrates how to change the scope of data profiling to the entire data set.

FIGURE 35.21

Choose column profiling based on the entire data set to get a more complete picture of your data.



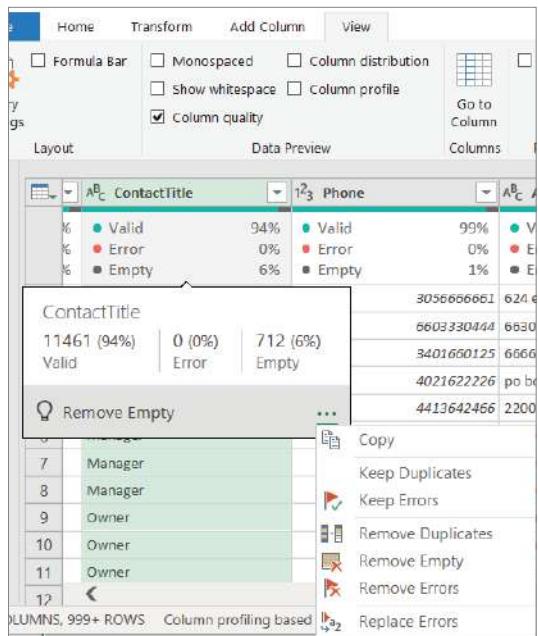
Data profiling quick actions

When selecting the Column Quality action, you'll see a set of figures that represent the percentage of values in a column that are valid, contain empty records, and rendered as an error. Hovering over these percentages exposes a pop-up containing an ellipsis (see Figure 35.22). Clicking the ellipsis activates a shortcut menu allowing you to apply quick actions such as Remove Errors, Remove Empty, and Remove Duplicates.

Part V: Understanding Power Pivot and Power Query

FIGURE 35.22

Exposing the quick actions for a column using the data column quality ellipsis



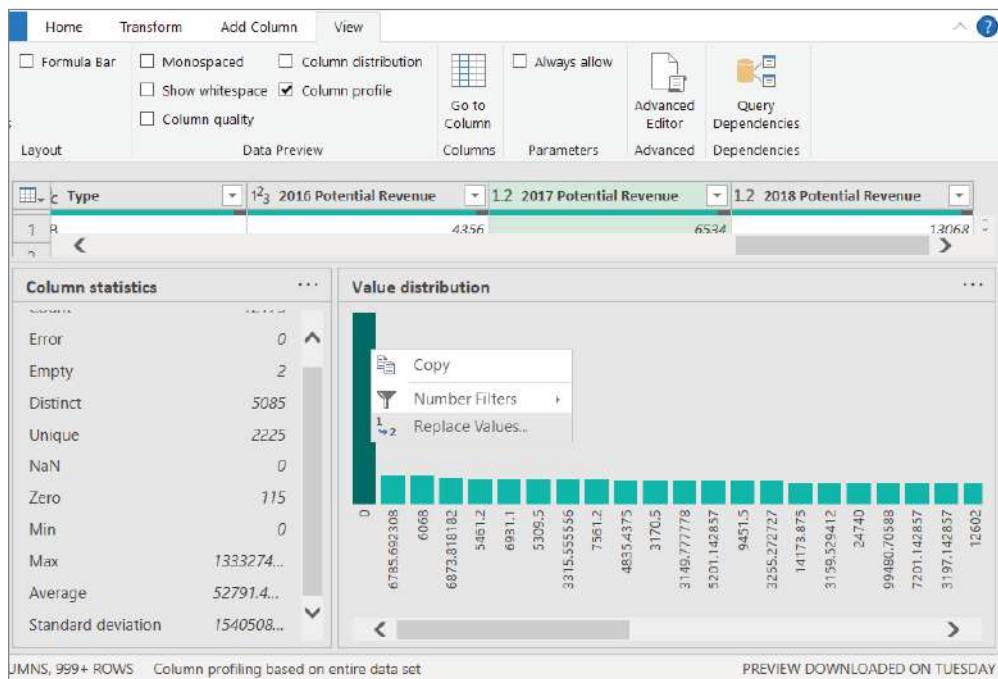
When selecting the Column Profile action, you'll see two new panes below the preview pane: Column statistics and Value distribution. As you can see in Figure 35.23, the Value distribution pane contains a histogram visual displaying the distribution of values. Right-clicking on any of the bars will reveal a quick action menu allowing you to apply transformations based on the type of data in that column. In this case, right-clicking on the bar for zero reveals the options for Number Filters and Replace Values.

TIP

The quick actions exposed via the data profiler are simply an easy way to find and apply needed transformations. They aren't any different than those found in the Power Query Editor Ribbon and those exposed by simply right-clicking a value in the preview pane.

FIGURE 35.23

Right-clicking a column profile histogram bar exposes the quick actions and the associated value.



Transforming Data with Power Query

IN THIS CHAPTER

Performing common transformations

Creating your own custom columns

Understanding data types

Understanding Power Query formulas

Applying conditional logic

Working with custom data types

Data transformation generally entails certain actions that are meant to “clean” your data—actions such as establishing a table structure, removing duplicates, cleaning text, removing blanks, and even adding your own calculations.

In this chapter, you will be introduced to some of the tools and techniques in Power Query that make it easy for you to clean and massage your data.

 You can follow along with the examples in this chapter by downloading the LeadList.txt sample file at www.wiley.com/go/excel365bible.

Once you've downloaded the file, you can import the sample file into Power Query. Open a new workbook, select Data ➔ Get & Transform Data ➔ Get Data ➔ From File ➔ From Text/CSV, browse to and import the LeadList.txt file, and then click the Transform Data button.

Performing Common Transformation Tasks

You will find that many of the unpolished data sets that come to you will require various types of transformation actions. The following sections cover some of the more common transformation tasks that you will have to perform.

Removing duplicate records

Duplicate records are absolute analysis killers. The effect that duplicate records have on your analysis can be far-reaching, corrupting almost every metric, summary, and analytical assessment that you

Part V: Understanding Power Pivot and Power Query

produce. It is for this reason that finding and removing duplicate records should be your first priority when you receive a new data set.

Before you jump into your data set to find and remove duplicate records, it's important to consider how you define a duplicate record. To demonstrate this point, look at the table shown in Figure 36.1, where you see 11 records. Out of the 11 records, how many are duplicates?

FIGURE 36.1

Are there duplicate records in this table? It depends on how you define one.

SicCode	PostalCode	CompanyNumber	DollarPotential	City	State	Address
1389	77032	11147805	\$9,517.00	houston	tx	6000 n sem heirten pkwy e
1389	77032	11147848	\$9,517.00	houston	tx	43410 e herdy rd
1389	77042	11160116	\$7,653.00	houston	tx	40642 rachmend ave ste 600
1389	77051	11165400	\$9,517.00	houston	tx	5646 helmis rd
1389	77057	11173241	\$9,517.00	houston	tx	2514 san filape st ste 6600
1389	77060	11178227	\$7,653.00	houston	tx	100 n sem heirten pkwy e ste 100
1389	77073	11190514	\$9,517.00	houston	tx	4660 rankan rd # 400
1389	77049	11218412	\$7,653.00	houston	tx	4541 mallir read 6
1389	77040	13398882	\$18,379.00	houston	tx	3643 wandfirn rd
1389	77040	13399102	\$18,379.00	houston	tx	3643 wandfirn rd
1389	77077	13535097	\$7,653.00	houston	tx	44160 wisthamir rd ste 100

If you were to define a duplicate record in Figure 36.1 as a duplication of just the SicCode, you would find 10 duplicate records. That is, out of the 11 records shown, 1 record has a unique SicCode, while the other 10 are duplications. Now, if you were to expand your definition of a duplicate record to a duplication of both SicCode and PostalCode, you would find only two duplicates: the duplication of PostalCodes 77032 and 77040. Finally, if you were to define a duplicate record as a duplication of the unique value of SicCode, PostalCode, and CompanyNumber, you would find no duplicates.

This example shows that having two records with the same value in a column does not necessarily mean you have a duplicate record. It's up to you to determine which field or combination of fields will best define a unique record in your data set.

Once you have a clear idea of what field, or fields, best make up a unique record in your table, you can remove duplicates easily by using the Remove Duplicates command.

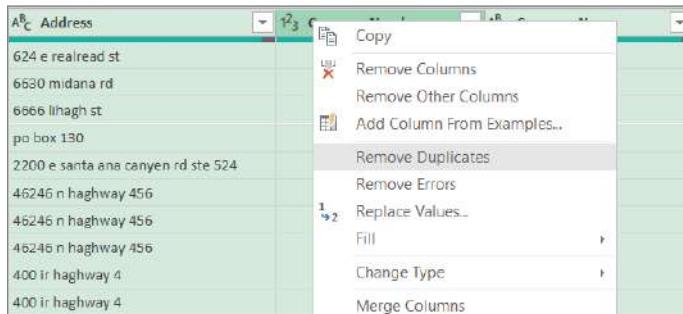
Figure 36.2 illustrates the removal of duplicate rows based on three columns. Note the importance of selecting the columns that define a duplicate. In this case, the combination of Address, CompanyNumber, and CompanyName defines a duplicate record. These columns are selected before right-clicking and selecting the Remove Duplicates command.

CAUTION

The Remove Duplicates command essentially looks for distinct values in the columns that you selected, and then it removes all of the records it needs to eliminate to end up with a unique list of values. If you select only one column before initiating the Remove Duplicates command, Power Query will use only the one column that you selected to determine the unique list of values. This will undoubtedly remove too many records—records that are not really duplicates. For this reason, it's important to make sure you select all of the columns that define a duplicate.

FIGURE 36.2

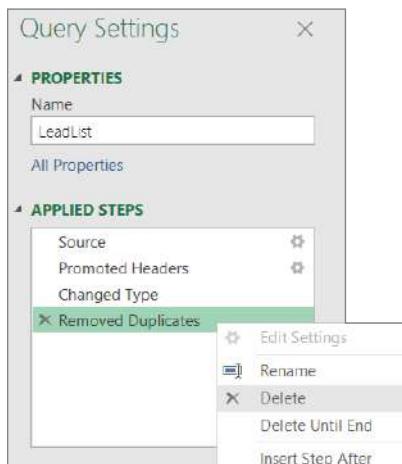
Removing duplicate records



If you made a mistake and removed duplicates based on the wrong set of columns, don't worry. You can always use the Query Settings pane to delete that step. Right-click the Removed Duplicates step and select Delete (see Figure 36.3).

FIGURE 36.3

Undo the removal of records by deleting the Removed Duplicates step.

**TIP**

Don't see a Query Settings pane in the Power Query Editor window? Select View \Rightarrow Layout \Rightarrow Query Settings to activate the Query Settings pane.

Filling in blank fields

It's important to note that there are actually two kinds of blank values: null and empty string. A **null** is essentially a numerical value of nothing, while an empty string is equivalent to entering two quotes (" ") in a cell.

Blank fields are not necessarily a bad thing, but an excessive number of blanks in your data can lead to unexpected problems when it comes time to analyze your data.

Your job is to decide whether to leave the blanks in your data set or fill them in with an actual value. When deciding this, you should consider the following best practices:

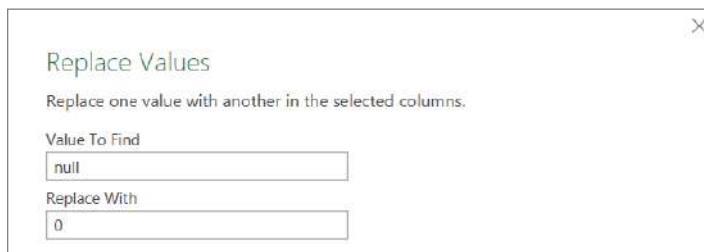
- **Use blanks sparingly:** Working with a data set is a much less daunting task when you don't have to test for blank values constantly.
- **Use alternatives when possible:** A good practice is to represent missing values with some logical missing value code whenever possible.
- **Never use null values in number fields:** Use zeros instead of null in a currency or number field that will be used in calculations.

Power Query will show the word **null** for any null value in your data. Replacing the null values is as simple as selecting the column or columns that you want to fix, right-clicking, and then selecting the Replace Values command.

The Replace Values dialog box, shown in Figure 36.4, will activate. The key here is to enter the word **null** in the Value To Find text box. You can then enter the value that you want to use instead. In this case, you can enter **0** as the Replace With value.

FIGURE 36.4

Replacing null values

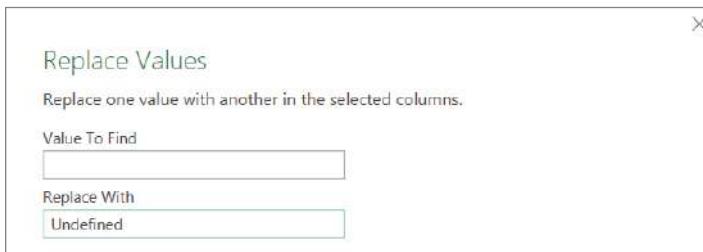


Filling in empty strings

It's always a best practice to represent missing values in a field with some logical value code whenever possible. For example, in Figure 36.5, we want to tag any record with a missing title with the word **Undefined**.

FIGURE 36.5

Replacing empty strings with the word **Undefined**



We can do so by right-clicking the ContactTitle field, selecting the Replace Values command, and then entering the word **Undefined** in the Replace With text box. As you can see in Figure 36.5, because we are replacing an empty string, there is no need to enter anything in the Value To Find text box.

TIP

If you need to adjust or correct the step where you replaced values, you can reactivate the Replace Values dialog box by clicking the gear icon next to the name for that step in the Query Settings pane. This is true for basically any action that requires a dialog box to complete. Clicking the gear icon next to any step name will activate the appropriate dialog box for that step.

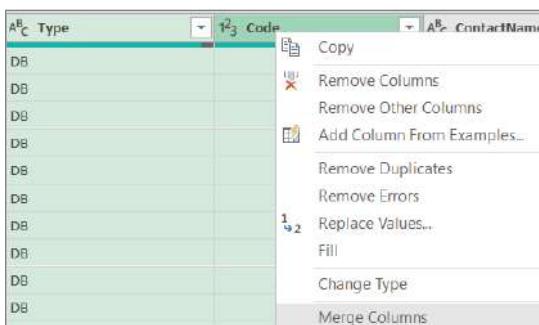
Concatenating columns

You can easily *concatenate* (join) the values in two or more columns. In Power Query, this is achieved by using the Merge Columns command. The Merge Columns command concatenates the values in two or more fields and outputs the newly merged values into a new column.

The idea is to select the columns that you want to concatenate, right-click, and select the Merge Columns command, as shown in Figure 36.6.

FIGURE 36.6

Merging the Type and Code fields

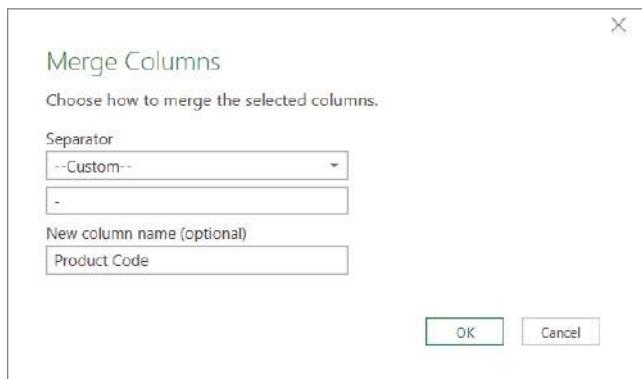


Part V: Understanding Power Pivot and Power Query

This will activate the Merge Columns dialog box. Here, you have the option of choosing a character that will act as a separator for the concatenated values. In addition to all the standard options such as comma, semicolon, space, and so on, you can specify a Custom character such as a hyphen, as shown in Figure 36.7. You also have the option of naming the new column that will be created.

FIGURE 36.7

The Merge Columns dialog box



The reward for your efforts will be a new field containing the concatenated values from the original column (see Figure 36.8).

FIGURE 36.8

The original columns are removed and replaced with a new merged column.

A ^B C Product Code	A ^B C ContactName
DB-100199	DAIMIRT, TAM, G.
DB-100199	THEMPSENJR, MAKE, G.
DB-100199	SCETT, ANDY, T.
DB-100199	MCKINZAE, DAVE, G.
DB-100199	NILSEN, REBIR, T.
DB-100199	KILLIRMAN, DAVAD, G.
DB-200	CELIMAN, TERRANCE, G.
DB-	SANSENE, TERRANCE, G.
DB-100199	GIRVES, STIPHIN, G.

This is nifty, but you will notice that Power Query removes the original Type and Code columns. There will definitely be instances when you want to concatenate values but still retain the source columns. In those instances, it's useful to first copy the column and perform the extraction on the duplicate column. You can create a copy of a column by right-clicking the

column and selecting Duplicate Column. When the duplicate column is created, it will be the last column (at the far right) of the table.

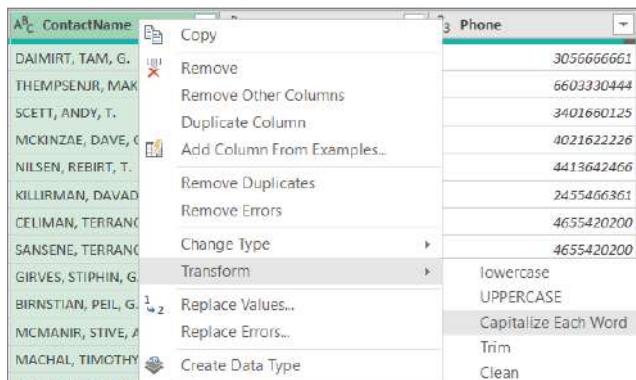
Changing case

Making sure that the text in your data has the correct capitalization may sound trivial, but it's important. Imagine that you receive a customer table that has an address field where all the addresses are lowercase. How is that going to look on labels, form letters, or invoices? Fortunately, Power Query has a few built-in functions that make changing the case of your text a snap.

For example, the ContactName field (see Figure 36.9) contains names that are formatted in all uppercase letters. To change these names to a more appropriate, proper case, you can right-click the column and then select Transform ➔ Capitalize Each Word.

FIGURE 36.9

Reformatting the ContactName field to proper case



Finding and replacing specific text

Imagine that you work in a company called BLVD, Inc. One day, the president of your company informs you that the abbreviation *blvd* on all addresses is now deemed an infringement on your company's trademarked name, and it must be changed to *Boulevard* as soon as possible. How would you go about meeting this new requirement?

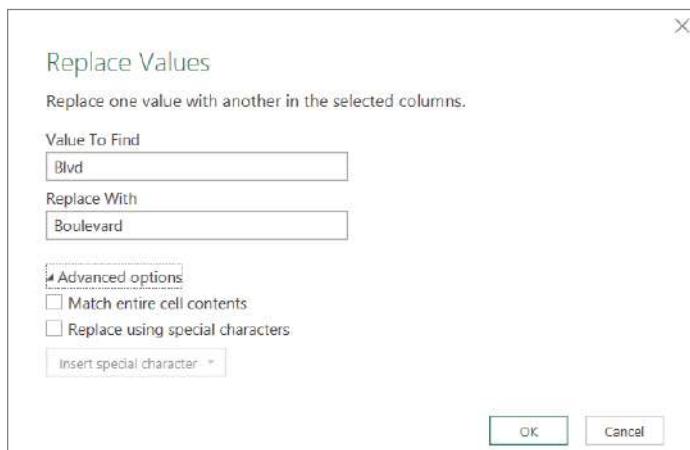
The Replace Values function is ideal in a situation like this. Right-click the Address field and then click the Replace Values command.

In the Replace Values dialog box (shown in Figure 36.10), simply fill in the Value To Find and the Replace With fields.

Part V: Understanding Power Pivot and Power Query

FIGURE 36.10

Replacing text values



Note that clicking Advanced Options exposes two optional settings:

Match Entire Cell Contents Checking this option tells Power Query to replace the specified value only if that value makes up the entire contents of the record. This is useful when you're attempting to replace a value such as 0 (zero) without replacing, for instance, all the zeros in the number 1,000.

Replace Using Special Characters Checking this option allows you to use special invisible characters such as line feed, carriage return, or tabs as replacement text. This is useful when you want to force an indent or reposition the text so that it shows up on two lines.

Trimming and cleaning text

When you receive a data set from a mainframe system, a data warehouse, or even a text file, it is not uncommon to have field values that contain leading and trailing spaces. These spaces can cause some abnormal results, especially when you're appending values with leading and trailing spaces to other values that are clean. To demonstrate this, look at the data set in Figure 36.11.

FIGURE 36.11

Leading spaces can cause issues in analysis.

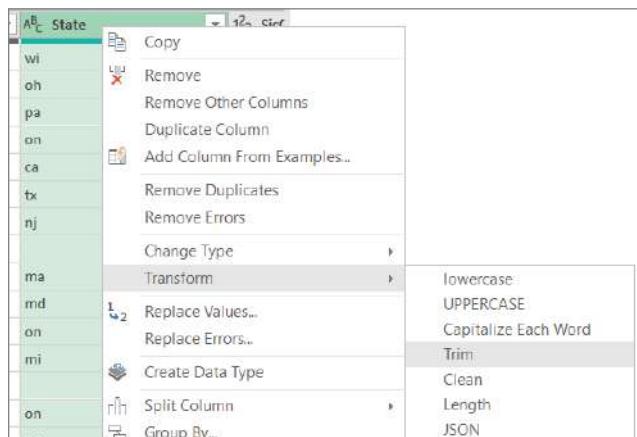
State	SumOfDollarPotential
ca	\$26,561,554.00
ny	\$7,483,960.00
tx	\$13,722,782.00
ca	\$12,475,489.00
ny	\$827,563.00
tx	\$7,669,208.00

This is intended to be an aggregate view that displays the sum of the dollar potential for California, New York, and Texas. However, the leading spaces are forcing each state into two sets, preventing you from discerning the accurate totals.

You can easily remove leading and trailing spaces by using the Trim function in Power Query. Figure 36.12 demonstrates how you would update a field to remove the leading and trailing spaces by right-clicking the column and using the Transform \Rightarrow Trim command.

FIGURE 36.12

The Trim command



Again, the Trim command will be applied to any column or columns you select. So, you can fix multiple columns at once simply by selecting them before selecting the Trim command.

Figure 36.12 also shows the Clean command (beneath Trim). While Trim removes leading and trailing spaces, the Clean command will remove any invisible characters such as carriage returns and other nonprintable characters that may slip in from external data source systems. These characters are typically rendered in Excel as a question mark or square box. But in Power Query, they show up as spaces.

If the source system that supplies your data has a nasty habit of including strange characters and leading spaces, you can apply the Trim and Clean functions to sanitize the data set.

Note

Some of you may know that the TRIM function in Excel removes the leading spaces, trailing spaces, and excess spaces within the given text. Power Query's TRIM function removes leading and trailing spaces but doesn't touch the excess spaces within the text. If excess spaces are an issue in your data, you can deal with them by using the Replace Values function to replace a given number of spaces with just one space.

Extracting the left, right, and middle values

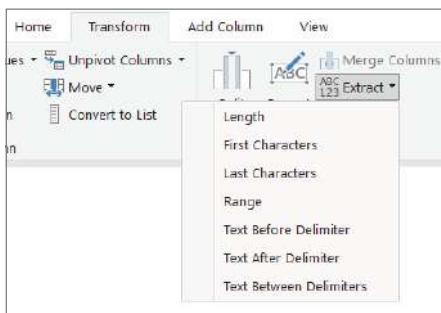
In Excel, we have the benefit of the RIGHT function, the LEFT function, and the MID function. These functions allow you to extract portions of a string starting from different positions:

- The LEFT function returns a specified number of characters starting from the leftmost character of the string. The required arguments for the LEFT function are the text you are evaluating and the number of characters that you want returned. For example, =LEFT ("70056-3504", 5) would return five characters starting from the leftmost character (70056).
- The RIGHT function returns a specified number of characters starting from the rightmost character of the string. The required arguments for the RIGHT function are the text you are evaluating and the number of characters you want returned. For example, =RIGHT ("Microsoft", 4) would return four characters starting from the rightmost character (soft).
- The MID function returns a specified number of characters starting from a specified character position. The required arguments for the MID function are the text that you are evaluating, the starting position, and the number of characters you want returned. For example, =MID ("Lonely", 2, 3) would return three characters starting from the second character, or character number 2 in the string (one).

Power Query has equivalent functions exposed through the Extract command found on the Transform tab (see Figure 36.13).

FIGURE 36.13

The Extract command allows you to pull out parts of the text found in a column.



The options under the Extract command are as follows:

Length Transforms a given column into numbers that represent the number of characters in each field (similar to Excel's LEN function).

First Characters Transforms a given column to show a specified number of characters from the beginning of text in each row (similar to Excel's LEFT function).

Last Characters Transforms a given column to show a specified number of characters from the end of text in each row (similar to Excel's RIGHT function).

Range Transforms a given column to show a specified number of characters starting from a specified character position (similar to Excel's MID function).

Text Before Delimiter Transforms a given column to show only the text that comes before a specified delimiter.

Text After Delimiter Transforms a given column to show only the text that comes after a specified delimiter.

Text Between Delimiters Transforms a given column to show only the text between two specified delimiters.

NOTE

Applying the Extract command to a column will effectively replace the original text with the results of the operation you choose to apply. That is, the original text will not be visible in the table after applying the Extract command. For this reason, it's often useful to first copy the column and perform the extraction on the duplicate column.

You can create a copy of a column by right-clicking the column and selecting Duplicate Column. When the duplicate column is created, it will be the last column (at the far right) of the table.

Extracting first and last characters

To extract the first N characters of text, highlight the column, select Extract \Rightarrow First Characters, and then use the dialog box shown in Figure 36.14 to specify the number of characters that you want to extract. In this case, the first three characters of the Phone field will be extracted.

FIGURE 36.14

Extracting the first three characters of the Phone field



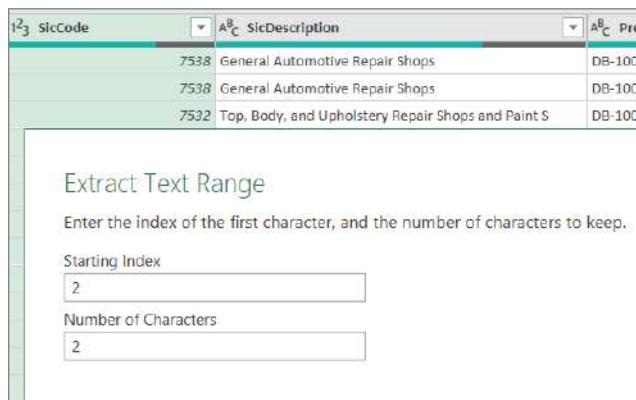
To extract the last N characters of text, highlight the column, select Extract \Rightarrow Last Characters, and then use the dialog box to specify the number of characters you want to extract.

Extracting middle characters

To extract the middle N characters of text, highlight the column, and select Extract \Rightarrow Range. The dialog box shown in Figure 36.15 will activate.

FIGURE 36.15

Extracting the two middle characters of the SicCode



The idea here is to tell Power Query to extract a specific number of characters starting from a certain position in the text. For example, the SicCode field is a four-digit field. If you wanted to extract the two middle numbers of the SicCode, you would tell Power Query to start at the second character and extract two characters from there.

As you can see in Figure 36.15, Starting Index is set to 2 (starting at the second character), and Number of Characters is set to 2 (extract two characters from the starting index).

Splitting columns using character markers

Have you ever gotten a data set where two or more distinct pieces of data were jammed into one field and separated by commas? For example, a field called Address may have a text value that represents "Address, City, State, ZIP." In a proper data set, this text would be split into four fields.

In Figure 36.16, you can see that the values in the ContactName field are strings that represent "Last name, First name, Middle initial." Imagine that you need to split this column string into three separate fields.

FIGURE 36.16

The Split Column command can easily split the ContactName field into three separate columns.

ContactName	ContactTitle	
DAIMIRT, TAM, G.	Manager	30
THEMPSENJR, MAKE, G.	Manager	65
SCKETT, ANDY, T.	Owner	34
MCKINZAE, DAVE, G.	Owner	40
NILSEN, REBIR, T.	Owner	44
KILLIRMAN, DAVAD, G.	Manager	24
CELMAN, TERRANCE, G.	Owner	46
SANSEME, TERRANCE, G.	Owner	46
GIRVES, STIPHIN, G.	Owner	64
BIRNSTIAN, PEIL, G.	Manager	14
MCMANIR, STIVE, A.	Plant Manager	40

Although this is not a straightforward undertaking in Excel, it can be done fairly easily with the Split Column command. Simply right-click the target column and then select the Split Column option. This will reveal seven options:

By Delimiter Split a column based on specific characters such as commas, semicolons, spaces, and so forth. This is useful for parsing names, addresses, or any field that contains multiple data points separated by delimiting characters.

By Number of Characters Split a column based on a specified number of characters. This is useful for parsing uniform text at a defined character position.

By Positions Split a column based on fixed numeric positions you specify.

By Lowercase to uppercase Split a column where the case changes from lowercase to uppercase.

By Uppercase to lowercase Split a column where the case changes from uppercase to lowercase.

By Digit to non-digit Split a column where the previous character is a digit, and the next consecutive character is a non-digit.

By Non-digit to digit Split a column where the previous character is a non-digit and the next consecutive character is a digit.

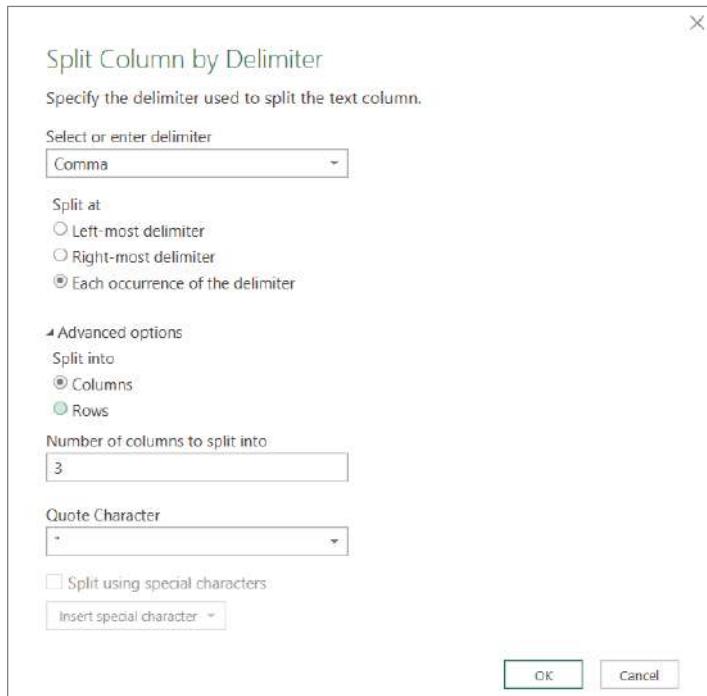
In the example shown in Figure 36.16, the contact names are made up of last names, first names, and middle initials, all separated (delimited) by commas. So, the By Delimiter option is the one we'll use.

Part V: Understanding Power Pivot and Power Query

You can highlight the ContactName field, right-click, and then select Split Column By Delimiter. This will activate the Split Column by Delimiter dialog box illustrated in Figure 36.17.

FIGURE 36.17

Splitting the ContactName column at every occurrence of a comma



The inputs here are as follows:

Select or enter delimiter Use the drop-down to choose the delimiter that will define where the values should be split. If your delimiter is not listed as a choice in the drop-down, you can select the Custom option and define your own.

Split Select how you want Power Query to use the specified delimiter. Power Query can split the column only on the first occurrence of the delimiter (the leftmost delimiter), effectively creating two columns. Alternatively, you can tell Power Query to split the column only on the last occurrence of the delimiter (the rightmost delimiter), again creating two columns. The third option is to tell Power Query to split the column at each occurrence of the delimiter.

Advanced options By default, selecting the option to split the column at each occurrence of the delimiter will create as many columns as there are delimiters. You can

use the Advanced options to override the default and limit the number of columns to create. You also have the advanced option to split your value into new rows instead of new columns.

Quote Character The quote character tells Excel not to consider any character inside quotes as delimiters. For example, when splitting a column using the comma delimiter, the quote character setting ensures that the commas in the text “Curly, Moe, and Larry” aren’t used to split the value into three columns.

Figure 36.18 demonstrates the new columns created after the ContactName column is split at each comma. As you can see, three new fields are created, and the original ContactName column is removed. You can rename the fields by right-clicking and selecting the Rename option.

FIGURE 36.18

The ContactName field has successfully been split into three columns.

ContactName.1	ContactName.2	ContactName.3	Contact
DAIMIRT	TAM	G.	Manager
THEMPSENJR	MAKE	G.	Manager
SCETT	ANDY	T.	Owner
MCKINZAE	DAVE	G.	Owner
NILSEN	REBIRT	T.	Owner
KILLIRMAN	DAVAD	G.	Manager
CELIMAN	TERRANCE	G.	Owner
SANSENE	TERRANCE	G.	Owner
GIRVES	STIPHIN	G.	Owner
BIRNSTIAN	PEIL	G.	Manager

Unpivoting columns

You often encounter data sets like the one shown in Figure 36.19, where important headings (such as the month) are spread across the top of the table, pulling double duty as column labels and actual data values. This matrix layout is easy to look at in a spreadsheet, but it causes issues when attempting to perform any kind of data analysis that requires aggregation, grouping, and so on.

Power Pivot offers an easy way to unpivot and pivot columns, allowing you to convert matrix-style tables to tabular data sets (and vice versa) quickly.

The Unpivot Columns command lets you select a set of columns and convert those columns into two columns: one column consisting of the column labels and another containing the column data.



You can follow along with this example by downloading the `UnpivotExample.xlsx` sample file at www.wiley.com/go/excel365bible. Open the file, then select **Data** ➔ **Get & Transform Data** ➔ **Get Data** ➔ **From Other Sources** ➔ **From Table/Range**.

Part V: Understanding Power Pivot and Power Query

FIGURE 36.19

Matrix layouts are problematic for data analysis.

Product_Description	Jan	Feb	Mar	Apr	May
Cleaning & Housekeeping Services	6219.66	4263.92	5386.12	6443.99	4360.14
Facility Maintenance and Repair	3255.82	9490	4409.23	4957.62	8851.12
Fleet Maintenance	5350.03	8924.71	6394.43	6522.46	9467.12
Green Plants and Foliage Care	2415.08	2579.61	2401.91	2981.01	2704.12
Landscaping/Grounds Care	5474.22	4500.52	5324.36	5705.68	5263.12
Predictive Maintenance/Preventative Maintenance	9810.95	10180.23	9626.31	11700.73	10947.12
Cleaning & Housekeeping Services	2840.76	2997.18	2096.78	4102.2	47.12
Facility Maintenance and Repair	16251.01	35878.99	18368.55	21843.53	2872.12
Fleet Maintenance	22574.77	36894.89	22016.38	27871.1	31985.12
Green Plants and Foliage Care	48250.9	90013.42	51130.17	75527.58	69418.12
Landscaping/Grounds Care	19401.16	21190.57	21292	20918.35	19465.12
Total	20712.74	10073.55	12040.05	16000.03	11000.12

The month columns in Figure 36.19 can be unpivoted by selecting the months and then right-clicking and selecting the Unpivot Columns command.

Figure 36.20 shows the resulting table. Note that the month labels are now entries in a new column called Attribute. The month values are now in a new column called Value. You can, of course, rename these columns to something like Month and Revenue.

FIGURE 36.20

All months are now in a tabular format.

Product_Description	Attribute	Value
Cleaning & Housekeeping Services	Jan	6219.66
Cleaning & Housekeeping Services	Feb	4263.92
Cleaning & Housekeeping Services	Mar	5386.12
Cleaning & Housekeeping Services	Apr	6443.99
Cleaning & Housekeeping Services	May	4360.14
Cleaning & Housekeeping Services	Jun	5097.46
Cleaning & Housekeeping Services	Jul	7566.19
Cleaning & Housekeeping Services	Aug	4263.92
Cleaning & Housekeeping Services	Sep	7245.64
Cleaning & Housekeeping Services	Oct	3847.15
Cleaning & Housekeeping Services	Nov	6540.21
Cleaning & Housekeeping Services	Dec	5610.45
Facility Maintenance and Repair	Jan	3255.82
Facility Maintenance and Repair	Feb	9490

Unpivoting other columns

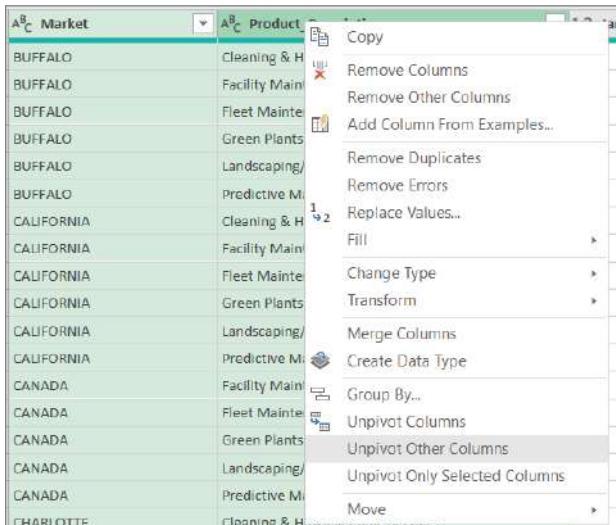
As nice as the Unpivot Columns command is, there is a flaw. The flaw is that you have to explicitly select the months you want unpivoted. But what if the number of columns is ever growing? What if you unpivot January through June, but next month, a new data set

will arrive with July and then August and then September? Because the Unpivot Columns command forces you essentially to hard-code the columns you want unpivoted, you'll have to redo the unpivot each and every month—that is, unless you use the Unpivot Other Columns command. This nifty command allows you to unpivot by selecting the columns you want to remain static and telling Power Query to unpivot all other columns.

For instance, Figure 36.21 demonstrates that instead of selecting the month columns, you can select the Market and Product Description columns, right-click, and then select Unpivot Other Columns.

FIGURE 36.21

Use Unpivot Other Columns when the number of matrix columns will be variable.



Now it doesn't matter how many new month columns are added or removed each month. Your query will always unpivot the correct columns.

TIP

It's a good idea always to use the Unpivot Other Columns option. Even if you don't anticipate new matrix columns, the Unpivot Other Columns option offers more flexibility for those unexpected changes in data.

Pivoting columns

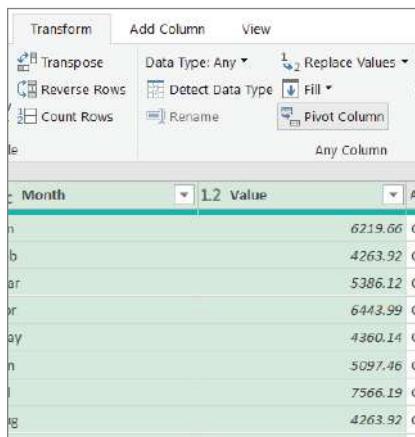
If you find that you need to transform your data from a tabular layout to a matrix-style layout, you can use the Pivot Column command.

Part V: Understanding Power Pivot and Power Query

Simply select the columns that will make up the header labels and values for the new matrix columns and then select the Pivot Column command from the Transform Ribbon tab. Figure 36.22 illustrates an example.

FIGURE 36.22

Pivoting the Month and Revenue columns



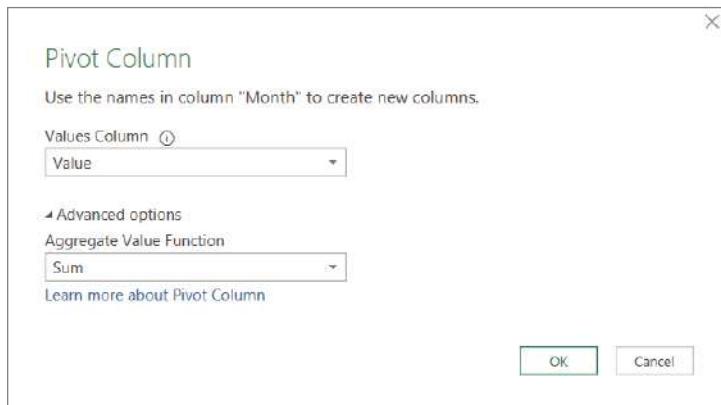
The screenshot shows the Power Query ribbon with the 'Transform' tab selected. Under the 'Transform' tab, the 'Pivot Column' button is highlighted with a yellow box. Below the ribbon, there is a table with two columns: 'Month' and 'Value'. The 'Month' column contains the values: Jan, Feb, Mar, Apr, May, Jun, Jul, Aug. The 'Value' column contains the corresponding revenue values: 6219.66, 4263.92, 5386.12, 6443.99, 4360.14, 5097.46, 7566.19, 4263.92.

Month	Value
Jan	6219.66
Feb	4263.92
Mar	5386.12
Apr	6443.99
May	4360.14
Jun	5097.46
Jul	7566.19
Aug	4263.92

Before finalizing the pivot operation, Power Query will activate a dialog box to confirm the value column and the aggregation method (see Figure 36.23). By default, Power Query will use the Sum operation to aggregate the data into the matrix format.

FIGURE 36.23

Confirm the aggregation operation to finalize the pivot transformation.



You can override this default by selecting a different operation (Count, Median, Minimum, and so on). You can even specify that you don't want any aggregation performed. Click OK to finalize the pivot operation.

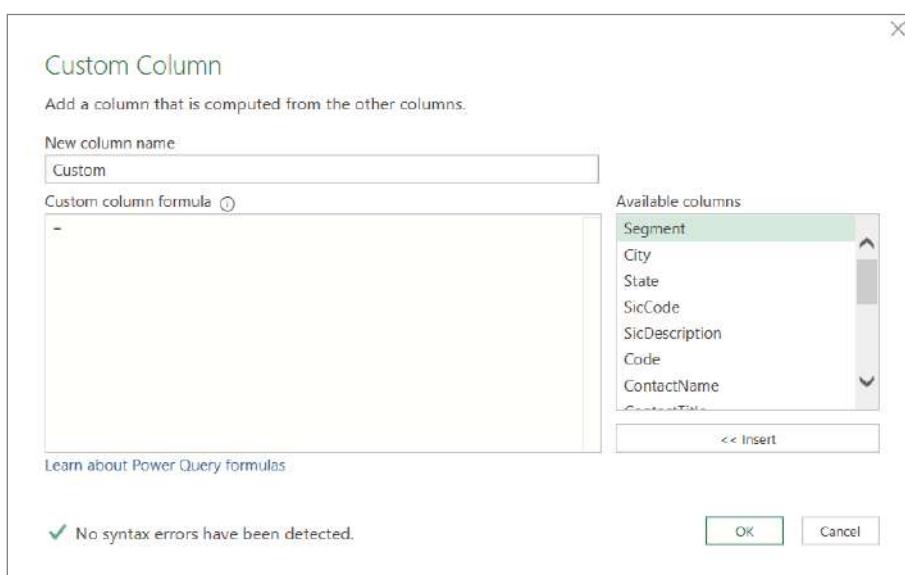
Creating Custom Columns

When transforming your data, you'll sometimes have to add your own columns to extract key data points, create new dimensions, or even create your own calculations.

You start a new custom column by going to the Add Column tab and clicking the Custom Column command. This activates the Custom Column dialog box (shown in Figure 36.24), where you will specify the contents of your new column through the use of Power Query formulas. When you add a new custom column, it won't do anything until you provide a formula that gives it some utility.

FIGURE 36.24

The Custom Column dialog box



As for the Custom Column dialog box, there's not much to it. The points of input are as follows:

- New column name** An input box where you enter a name for the column that you're creating.

Available columns A list box that contains the names of all the columns in your query. Double-click any column name in this list box to place it automatically in the formula area.

Custom column formula The area where you type your formula.

Just as in Excel, a formula can be as simple as =1 or as complicated as an IF statement that applies some conditional logic. Over the next few sections, you'll walk through a few examples of creating custom columns to go beyond the functionality provided through the user interface.

But before diving in to building Power Query formulas, it's important to understand how Power Query formulas differ from those in Excel. Here are some high-level differences of which you should be aware:

- **No cell references:** You can't point outside the Custom Column dialog box to select a range of cells. Power Query formulas work by referencing columns, not cells.
- **Excel functions don't work:** The Excel functions that you're used to won't work in Power Query. Power Query has many of the same kinds of functions as Excel, but it has its own formula language.
- **Everything is case sensitive:** In Excel, you can type in all-lowercase or all-uppercase letters and formulas will work. Not so in Power Query. To Power Query, *sum*, *Sum*, and *SUM* are three different things, and only one of them (*Sum*) is acceptable.
- **Data types matter:** Some fields are text fields, other fields are number fields, while others are dates. Excel does a good job at handling formulas that mix fields of differing data types. The Power Query formula language is extremely sensitive to data types. It doesn't have the built-in intelligence to handle data type mismatches gracefully. Data type issues are resolved with conversion functions, covered later in this chapter.
- **No ScreenTips or intelligence help:** Excel is quick to throw up a ScreenTip or a menu of options when you start entering a new formula. Power Query has none of that. The most Power Query currently offers is a "Learn about Power Query formulas" link, taking you to a Microsoft site dedicated to Power Query.

Don't worry. It's not as gloomy as it sounds. Let's start with a simple custom column.

Concatenating with a custom column

Earlier in this chapter, you discovered how to concatenate values from two or more columns by using the Merge Columns command. Although the Merge Columns command is easy to use, it results in the original source columns being removed. There will definitely be instances when you will want to concatenate values but still retain the source columns.

In these instances, you can create your own custom column. Follow these steps to create a new column that merges the Type and Code columns:

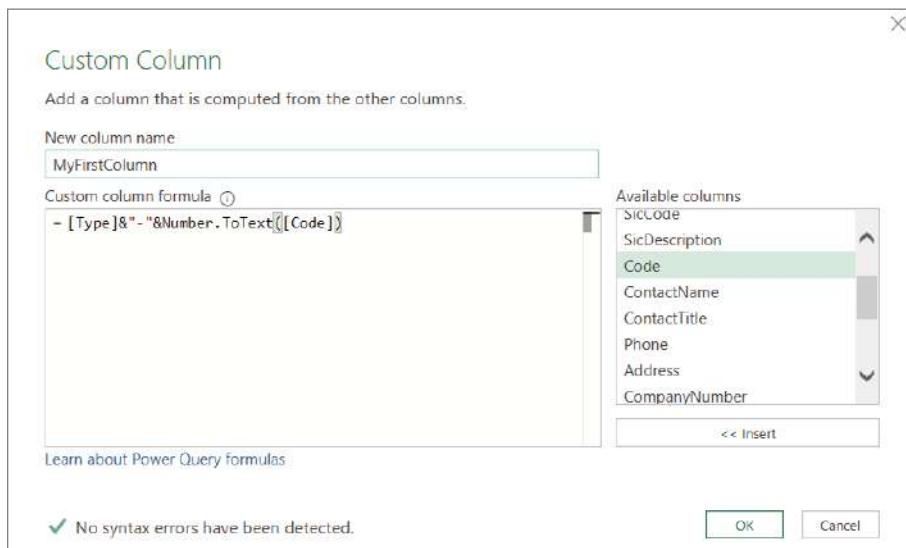
1. While in the Power Query Editor, click Add Column \Rightarrow General \Rightarrow Custom Column.
2. Place your cursor in the formula area, after the equal sign.

3. **Find the Type column in the Available Columns list and double-click it.** You will see [Type] pop into the formula area after the equal sign.
4. **After [Type], enter the following text: &" - "&.** This text will ensure that the values in the two columns are separated by a hyphen.
5. **Next, enter Number.ToText().** This Power Query function converts a number to text format on the fly so that it can be used with other text. In this case, since the Code field is formatted as a number, we'll need to convert it on the fly to join it to the Type field. You'll learn more about data type conversions later.
6. **Place your cursor between the parentheses for the Number.ToText function.**
7. **Next, find the Code column in the Available Columns list and double-click it.** You will see [Code] pop into the formula area between the parentheses.
8. **In the New Column Name input box, enter MyFirstColumn.**

At this point, your dialog box should look similar to the one shown in Figure 36.25. Note the message at the bottom of the dialog box stating that "No syntax errors have been detected." Each time you create or adjust a formula, you'll want to ensure that this message states that no errors have been detected.

FIGURE 36.25

A formula to merge the Type and Code columns



9. **Click OK to add your custom column.**

If all went well, you have a new custom column that concatenates two fields. With this example, you get the basic foundation for how Power Query formulas work.

Understanding data type conversions

When working with formulas in Power Query, you'll inevitably need to perform some action on fields that have differing data types. Take, for example, the previous exercise where we merged the Type column (a text field) with the Code column (a numeric field). In that example, we used a conversion function to change the data type of the Code field so that it can be temporarily treated as a text field.

Conversion functions are just what they sound like. They convert data from one data type to another.

Table 36.1 lists common conversion functions. As demonstrated in the previous section, you simply wrap these functions around the columns that need converting.

```
Number.Text([ColumnName])
```

TABLE 36.1 Common Conversion Functions

Convert From	To	Function
Date	Text	Date.Text()
Time	Text	Time.Text()
Number	Text	Number.Text()
Text	Number	Number.FromText()
Text Dates	Date	Date.FromText()
Numeric Dates	Date	Date.From()

All that being said, it's clear that you'll need to know the data types for the fields that you use in Power Query formulas. That way, you'll know which conversion functions to use.

To discover and change the data type for a field, place your cursor in the field and then select the Data Type drop-down on the Transform tab (see Figure 36.26). You can edit the data type for the field by selecting a new type from the drop-down list.

Spicing up custom columns with functions

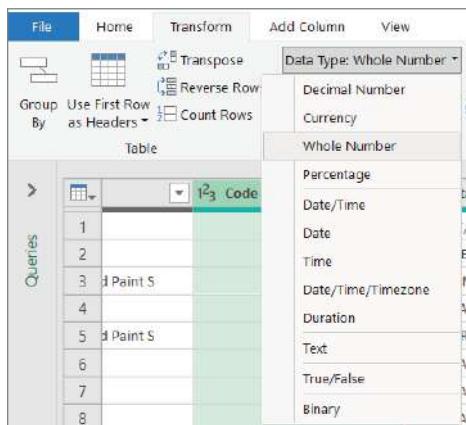
With a few basic fundamentals and a little knowledge of Power Query functions, you can create transformations that go beyond what you can do through the Power Query Editor. In this example, we'll use a custom column to pad numbers with zeros.

You may encounter a situation where key fields are required to be a certain number of characters for your data to be able to interface with peripheral platforms such as ADP or SAP. For example, imagine that the CompanyNumber field must be 10 characters long. Those

that are not 10 characters long must be padded with enough leading zeros to create a 10-character string.

FIGURE 36.26

Use the Data Type drop-down to discover and select the data type for a given field.



The secret to this trick is to add 10 leading zeros to every company number, regardless of the current length, and then pass them through a function similar to the RIGHT function that will extract only the right 10 characters.

For example, company number 29875764 would first be converted to 000000000029875764; then it would go into a RIGHT function that extracts only the right 10 characters. This would leave you with 0029875764.

Although this is essentially two steps, you can accomplish the same thing with just one custom column. Here's how to do it:

1. While in the Power Query Editor, click Add Column \Rightarrow General \Rightarrow Custom Column.
2. Place your cursor in the formula area, after the equal sign.
3. Enter 10 zeros in quotes (as in "0000000000") followed by an ampersand (&).
4. Next, enter Number.Text(). (Do not type the period.)
5. Place your cursor between the parentheses for the Number.Text function.
6. Next, find the CompanyNumber column in the Available Columns list and double-click it. You will see [CompanyNumber] pop into the formula area.

At this point, your formula area should contain this syntax:

```
"0000000000"&Number.Text([CompanyNumber])
```

Part V: Understanding Power Pivot and Power Query

This formula will result in nothing more than a concatenation of 10 zeros and the CompanyNumber. The goal is to go further and extract only the right 10 characters. The RIGHT function is an Excel function that won't work in Power Query. However, Power Query does have an equivalent function called Text . End (). Just like the RIGHT function, the Text . End function requires a couple of parameters: the text expression and the number of characters to extract.

Text . End ([MyText] , 10)

7. In this scenario, the text expression will be your formula, and the number of characters to extract will be 10. At the beginning of the formula, enter Text . End and a left parenthesis and then at the end of the formula, enter a comma and then the number 10 and a right parenthesis. The final syntax is shown here:

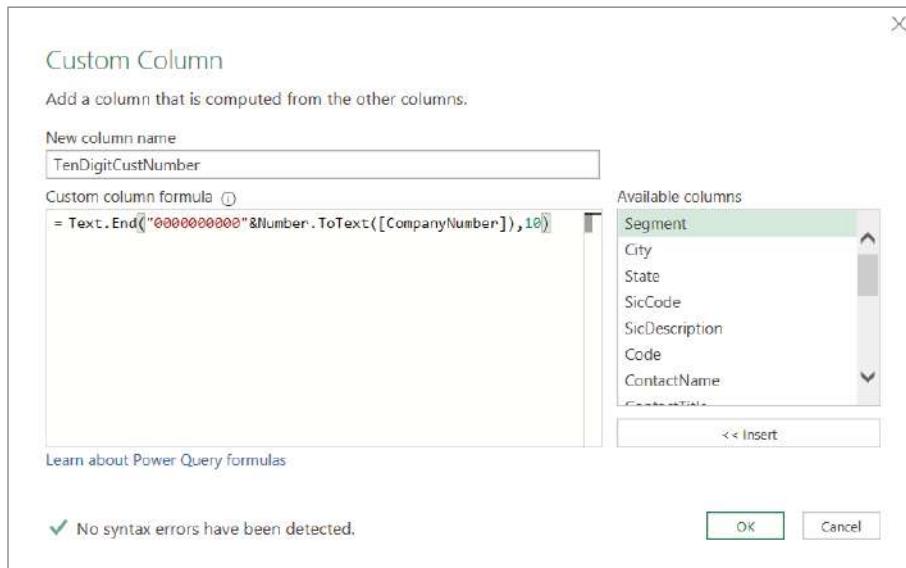
```
Text . End ("0000000000" &  
Number . ToText ( [CompanyNumber] ) , 10 )
```

8. In the New Column Name input box, enter TenDigitCustNumber.

At this point, your dialog box should look similar to the one shown in Figure 36.27. Again, make sure that the message at the bottom of the dialog box reads "No syntax errors have been detected."

FIGURE 36.27

A formula to create a consistent CompanyNumber padded with 10 digits



9. Click OK to apply your custom column.

Table 36.2 lists some other Power Query functions useful in extending the capabilities of custom columns. Take a moment to review the list of functions and note how they differ from their Excel equivalents. Remember that Power Query functions are case sensitive.

TABLE 36.2 Useful Transformation Functions

Excel Function	Power Query Function
LEFT([Text], [Number])	Text.Start([Text], [Number])
RIGHT([Text], [Number])	Text.End([Text], [Number])
MID([Text], [StartPosition], [Number])	Text.Range([Text], [StartPosition], [Number])
FIND([Find], [Within])	Text.PositionOf([Within], [Find])+1
IF([Expression], [Result1], [Result2])	if [Expression] then [Result1] else [Result2]
IFERROR([Procedure], [FailResult])	try [Procedure] otherwise [FailResult]

Adding conditional logic to custom columns

As you may have noticed from Table 36.2, Power Query has a built-in if function. The if function is designed to test for conditions and provide different outcomes based on the results of those tests. In this section, you'll see how you can control the output of your custom columns by utilizing Power Query's if function.

Just as in Excel, Power Query's if function evaluates a specific condition and returns a result based on a true or false determination:

```
if [Expression] then [Result1] else [Result2]
```

Note

In Excel, you think of the commas in an IF function as THEN and ELSE statements. In Power Query, you don't use commas. The Excel formula IF(Babies=2, "Twins", "Not Twins") would translate to the following:

If Babies equals 2 then "Twins" else "Not Twins".

Imagine that you need to tag customers as either large customers or small customers based on their dollar potential. You decide that you'll add a custom column that contains either "LARGE" or "SMALL" based on the revenue potential of the customer.

With the help of the if function, you can tag all customers with one custom column that uses the following formula:

```
if [2020 Potential Revenue]>=10000  
then "LARGE" else "SMALL"
```

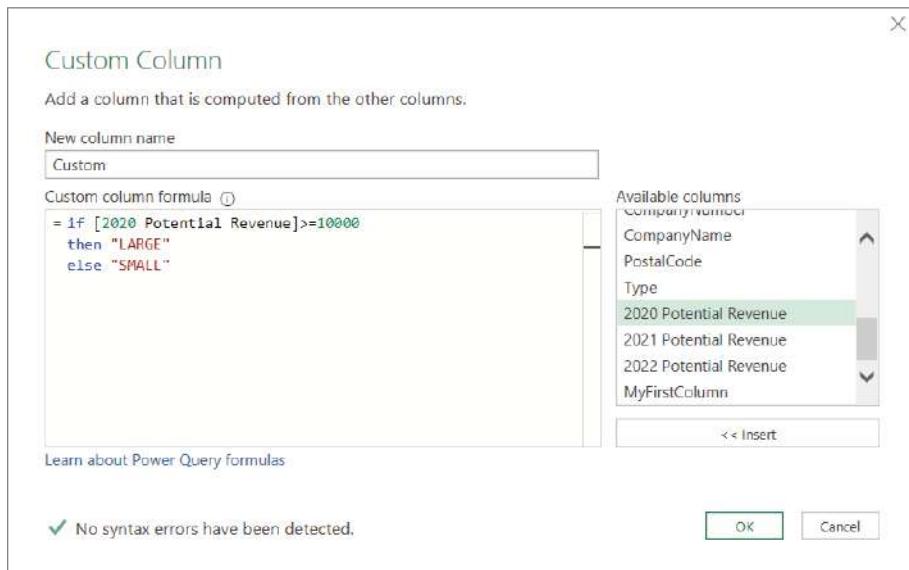
Part V: Understanding Power Pivot and Power Query

This function tells Power Query to evaluate the [2020 Potential Revenue] field for each record. If the potential revenue is greater than or equal to 10,000, use the word LARGE; if not, use the word SMALL.

Figure 36.28 demonstrates this if statement as it is applied in the Custom Column dialog box.

FIGURE 36.28

Applying an if statement in a custom column



TIP

Power Query pays no attention to white space. This means you can put as many spaces and carriage returns as you'd like. As long as the correct case and spelling is used, Power Query will not complain. Figure 36.28 illustrates how separating formulas into multiple lines can make them much easier to read.

Grouping and Aggregating Data

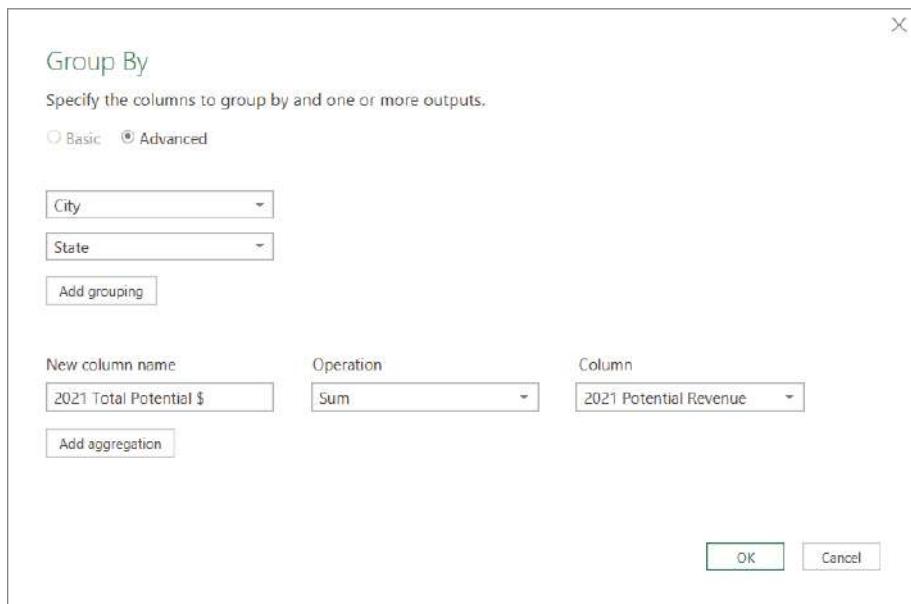
In some cases, you may need to transform your data set into compact groups to get it into a manageable size of unique values. You may even need to summarize numerical values into an aggregate view. An aggregate view is a grouped snapshot of your data that will show sums, averages, counts, and more.

Power Query offers a Group By feature that enables you to easily group data and create aggregate views.

While in the Power Query Editor, select the Group By command on the Transform tab. The Group By dialog box illustrated in Figure 36.29 will open.

FIGURE 36.29

The Group By dialog box, configured to sum potential revenue by State and City



You have the option of creating a Basic grouping with only one column to group by or creating an Advanced aggregation that uses multiple Group By columns.

Let's take a moment to walk through the steps for transforming data into an aggregate view by City and State:

1. **While in the Power Query Editor, activate the Group By dialog box (Transform ↴ Table ↴ Group By).**
2. **Click the option for Advanced grouping.** This exposes the fields that we need to group by multiple columns.
3. **Use the Group By drop-down list to select the column on which you want to aggregate. Then click the Add Grouping button to add columns to the grouping.** (In Figure 36.29, City and State were selected.)
4. **Use the New Column Name input box to give the new aggregate column a name.**

5. Use the Operation drop-down list to select the kind of aggregation you want to apply (Sum, Count Rows, Average, Min, Max, and so forth). In Figure 36.29, we have chosen to apply Sum.
6. Use the Column drop-down list to choose the column that will be aggregated. In this case, it's 2021 Potential Revenue.
7. Click the OK button to confirm and apply your changes.

Figure 36.30 illustrates the resulting output.

FIGURE 36.30

The resulting aggregate view by City and State

	City	State	Total Potential \$
1	adams	wi	6534
2	akron	oh	54754.34471
3	allentown	pa	77517.35
4	alliston	on	5757.260526
5	anaheim	ca	97655.01111
6	austin	tx	77736.85322
7	avenel	nj	34776.08889

Note

When you apply the Group By feature, Power Query removes all columns that were not used when configuring the Group By dialog box. This leaves you with a clean view of just your grouped data.

Working with Custom Data Types

The Custom Data feature of Power Query allows you to store multiple columns of data in one column as metadata. Excel formulas can then interact with these rich data types to expose the stored data within. In this section, you'll explore the basics of using Custom Data Types in your reporting.



You can follow along with this example by downloading the `CustomDataTypes.xlsx` sample file at www.wiley.com/go/excel365bible. Open the file, then select **Data** ➔ **Get & Transform Data** ➔ **Get Data** ➔ **From Other Sources** ➔ **From Table/Range**.

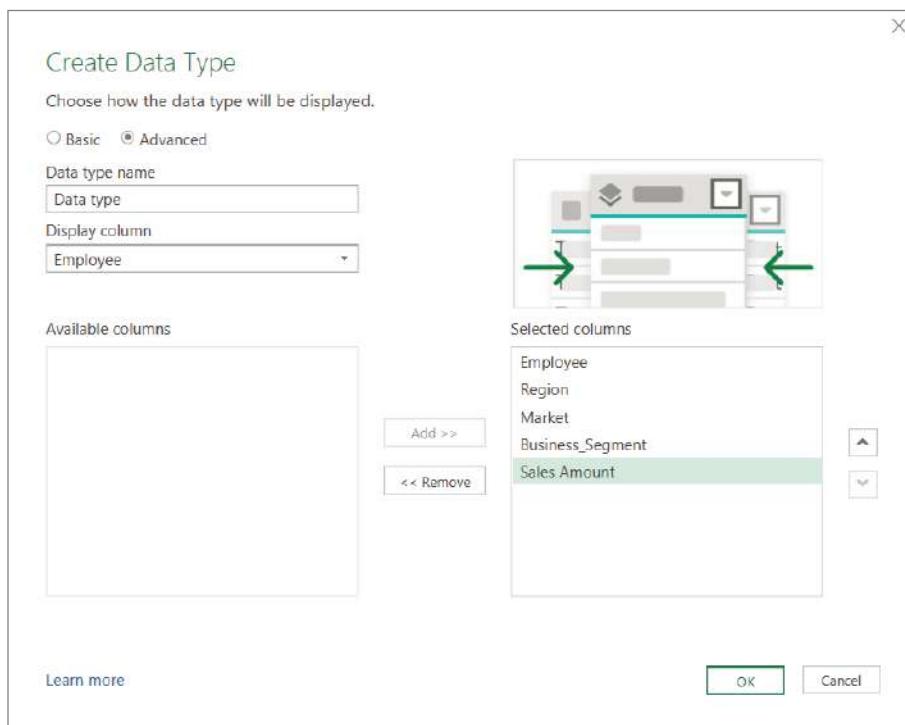
The term *Custom Data Type* is a bit unfortunate, as it doesn't really correlate with the traditional meaning of data type: number, date, currency, and so on. In Power Query, you can think of a Custom Data Type as a kind of container that allows you to store the data for many columns and then use it elsewhere in your workbook.

You can get a sense of the power of this feature by following these steps.

1. While in the Power Query Editor, right-click the Employee column and select Create Data Type. The Create Data Type dialog box will activate.
2. Click the Advanced option to reveal a list of available fields (see Figure 36.31). The idea here is the display column (Employee in this case) will be a kind of container that will hold data from other columns you specify.

FIGURE 36.31

Creating a Custom Data Type



3. Select the Region column in the list of available columns and then click the Add button. Repeat for each column in the list of available fields.
4. Click the OK button to confirm your changes.

At this point, the preview pane will show only the Employee column because that is the column you selected as your display column. As you can see in Figure 36.32, clicking on any value in the Employee column now results in a table below the preview pane, which shows all the data contained for that value.

Part V: Understanding Power Pivot and Power Query

FIGURE 36.32

Each value in a data type column contains the data for underlying columns, as seen below the preview pane.

Data type	
1	AAEMS, JOSEPH
2	ALCERO, ROBERT
3	ANDIRSEN, DORAN
4	ATKANS, TERRY
5	BECKMAN, ADRIAN
6	BEICHERD, SHANE
7	BERANSKI, BRANT

Employee	AAEMS, JOSEPH
Region	MIDWEST
Market	DENVER
Business_Segment	Housekeeping and Organization
Sales Amount	455.33

Click the Close & Load button on the Home tab of the Power Query Editor to send your results to a new worksheet. You can see in Figure 36.33, the output of this query is a list of values from the Employee column, each containing the Data Type icon designed to let you know that the value is part of a data type. Clicking the smart icon to the right of the table allows you to add any of the underlying values for that data type.

FIGURE 36.33

Data types will have a special icon next to each value and will allow you to see any values in underlying columns.

A	B	C	D	E	F
1 Data type					
2 ◊AAEMS, JOSEPH					
3 ◊ALCERO, ROBERT					
4 ◊ANDIRSEN, DORAN					
5 ◊ATKANS, TERRY					
6 ◊BECKMAN, ADRIAN					
7 ◊BEICHERD, SHANE					
8 ◊BERANSKI, BRANT					
9 ◊BERIN, WILLIAM					
10 ◊BEYD, MICHAEL					
11 ◊BILL, GERARD					
12 ◊BIOCH, RONALD					
13 ◊BIOGESS, ALAN					

Smart icon details:

- Business_Segment
- Employee
- Market
- Region
- Sales Amount
- Title

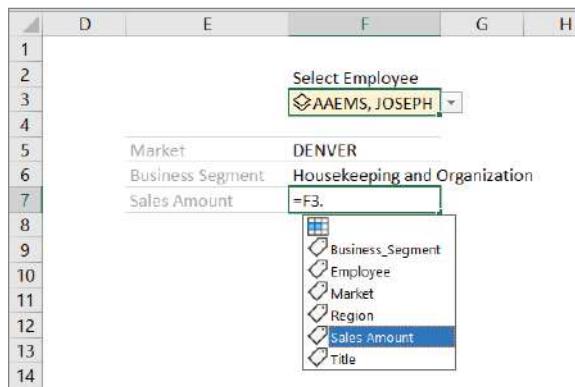
The real power of Custom Data Types is the ability to reference underlying values via simple formulas. To illustrate one way to use data types, take a look at Figure 36.34. Here, we loaded a data validation drop-down with the values from our Employee column (our data type header). You can see the icon next to the selected employee name, confirming the value is a

Chapter 36: Transforming Data with Power Query

data type. Now we can reference that value with a simple formula to see the available columns underneath.

FIGURE 36.34

Referencing a data type value and entering the dot (.) operator will allow you to select any of the underlying columns.



That's right. Reference a data type, enter the dot (.) operator, and you'll have access to any of the underlying column values. Selecting a new employee from the data validation drop-down will automatically pull that employee's data.



Feel free to review Chapter 23 “Using Data Validation” for a refresher on how to add data validation to your worksheets.

TIP

Remember that you can refresh your underlying Power Query. Refreshing your query will not only bring in any new data type headers, it will also automatically update any of the values resulting from a formula that references your data types.

Making Queries Work Together

IN THIS CHAPTER

- Reusing query steps
- Consolidating data with the Append feature
- Understanding join types
- Using the Merge feature to combine data
- Using fuzzy matching when merging queries

Data analysis is frequently done in layers, with each layer of analysis using or building on the previous layer. When you build a PivotTable using the results of a Power Query output, you are layering your analysis. When you build a query based on a table created by a SQL Server view, you are also creating a layered analysis.

You'll often find the need to build queries on top of other queries to get the results you're seeking. That's what this chapter is all about. Here you'll take a look at a few ways that you can advance your data analysis by making your queries work together.



You can follow along with the examples in this chapter by downloading these sample files from www.wiley.com/go/excel365bible:

- Sales_By_Employee.xlsx
- Appending_Data.xlsx
- Merging_Data.xlsx
- FuzzyMatch.xlsx

Reusing Query Steps

It's common to rely on the same main data table for all kinds of analysis. Even the simple table shown in Figure 37.1 can be used to create different views: sales by employee, sales by business segment, sales by region, and so forth.

Part V: Understanding Power Pivot and Power Query

Of course, you can build separate queries, each performing different grouping and aggregation steps, but that would mean repeating all of the data cleanup steps you needed before performing any kind of analysis.

FIGURE 37.1

This data can be used as the source for various levels of aggregated analysis.

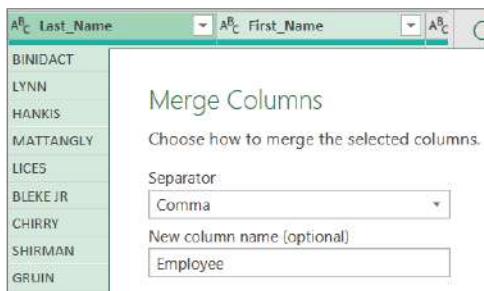
	A	B	C	D	E	F
1	Region	Market	Last_Name	First_Name	Business_Segment	Sales_Amount
2	MIDWEST	DENVER	AAEMS	JOSEPH	Housekeeping and Organization	\$465.33
3	MIDWEST	DENVER	AAEMS	JOSEPH	Landscaping and Area Beautificat	\$411.60
4	MIDWEST	DENVER	AAEMS	JOSEPH	Maintenance and Repair	\$760.31
5	MIDWEST	DENVER	BEALIY	CHRISTOPHER	Maintenance and Repair	\$2,125.38
6	MIDWEST	DENVER	BEALIY	CHRISTOPHER	Landscaping and Area Beautificat	\$5,909.14
7	MIDWEST	DENVER	BEALIY	CHRISTOPHER	Maintenance and Repair	\$39,829.79
8	MIDWEST	DENVER	BEWMAN	DIRK	Landscaping and Area Beautificat	\$319.18
9	MIDWEST	DENVER	BEWMAN	DIRK	Maintenance and Repair	\$119.38
10	MIDWEST	DENVER	BIHRINS	KURT	Landscaping and Area Beautificat	\$914.20
11	MIDWEST	DENVER	BIHRINS	KURT	Maintenance and Repair	\$17,645.38
12	MIDWEST	DENVER	BREWN	SCOTT	Maintenance and Repair	\$112.01
13	MIDWEST	DENVER	ROOKS	HENRY	Landscaping and Area Beautificat	\$605.66

To see what this means, take a moment to follow these steps:

1. Open the Sales By Employee.xlsx sample file.
2. Select any cell inside the table and then click Data \Rightarrow Get & Transform Data \Rightarrow Get Data \Rightarrow From Other Sources \Rightarrow From Table/Range. Power Query will open the Power Query Editor to show you a table that looks similar to the one in Figure 37.1.
3. Click the Filter drop-down for the Market field and filter out the Canada market (remove the check next to Canada).
4. Select the Last_Name and First_Name fields, right-click either column heading, then select Merge Columns.
5. Use the Merge Columns dialog box to create a new Employee field, joining Last_Name and First_Name separated by a comma (see Figure 37.2).

FIGURE 37.2

Merge the Last_Name and First_Name columns to create a new Employee field.

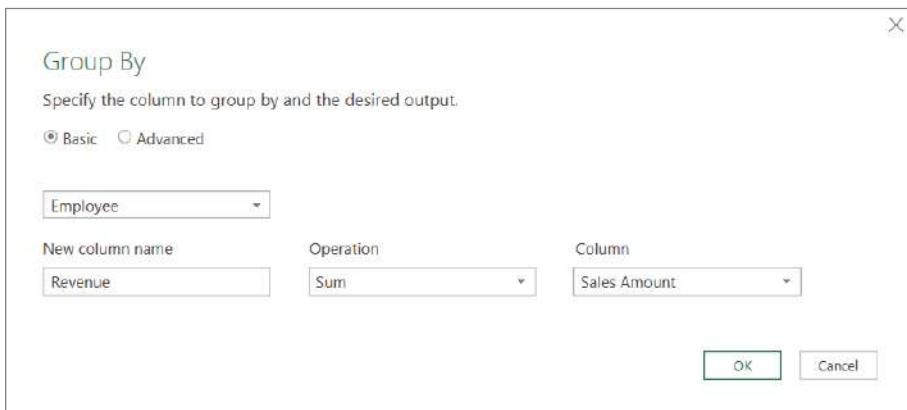


6. Select only the Employee column and click the Group By command on the Transform tab to display the Group By dialog box. The goal is to group by the Employee field to get the Sum of Sales Amount.
7. Name the new aggregated column Revenue. Select Sum in the Operation drop-down and Sales Amount in the Column drop-down.

Figure 37.3 illustrates the completed Group By dialog box.

FIGURE 37.3

Group the Employee field and sum the Sales Amount column to create a new Revenue column.



At this point, you've successfully created a view that shows total revenue by employee. As you can see in Figure 37.4, the query steps include all of the preparation work you did before grouping.

FIGURE 37.4

All of the query steps before Grouped Rows are needed to prepare the data for grouping.

	Employee	Revenue
1	ILLAS,CLYDE	8142.48
2	KALLAO,LEE	9703.44
3	HILTCIL,TOMMY	328.59
4	HERVIY,SHALAN	71427.16
5	HANKIS,OREN	109.53
6	HICKLIBIRRY,JERRY	19002.81
7	LENG,JAMES	693.52
8	LYNN,THEODORE	18889.85
9	MANDIVALLE,RICHARD	36386.95
10	LYNNIN,JOSHUA	11682.41
11	LICES,NORVEL	1713.56
12	LIWAS,JAMES	1888.37
13	HALL,TIMOTHY	677.75

So, what if you wanted to create another analysis using the same data? For instance, what if you wanted another view that shows employee sales by business segment?

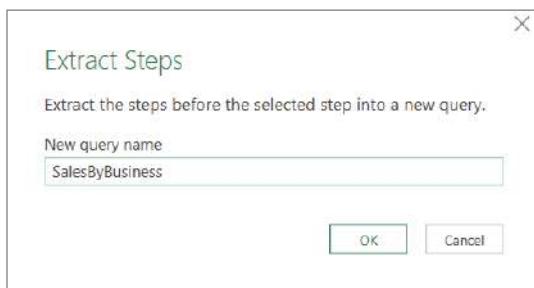
You could always start from step 1 and import another copy of the source data, but you'd have to repeat the preparation steps (the Filtered Rows and Merged Columns steps in this case).

A better way is to reuse the steps you've already created by extracting them into a new query. The idea is first to decide what steps you want to reuse and then right-click the step immediately below. In this scenario (see Figure 37.4), you want to keep all query steps up until Grouped Rows.

8. Right-click the **Grouped Rows** step and select **Extract Previous**.
9. Use the **Extract Steps** dialog box shown in Figure 37.5 to name the new query **SalesByBusiness**. Click **OK** to confirm.

FIGURE 37.5

Naming the new query SalesByBusiness



After you click OK, Power Query does two things. First, it moves all the extracted steps to the newly created query. Then it ties the original query to the new query. That is to say, both queries are sharing the first few query steps up until the Grouped Rows step.

You can see the new SalesByBusiness query (see Figure 37.6) in the pane on the left. When you select SalesByBusiness in the left pane, you'll notice that the query steps for the newly created SalesByBusiness query do not include the Grouped Row step. Power Query only moved the extracted steps (those before the Grouped Rows step).

This concept of extracting steps can be a bit confusing. The bottom line is that instead of starting from square one with a new query, you're telling Power Query that you want to create a new query that utilizes the steps you've already created.

NOTE

When two or more queries share extracted steps, the query that contains the extracted steps serves as the data source for the other queries. Because of this link, the query that contains the extracted steps cannot be deleted. You'll first have to delete all dependent queries before deleting the query that holds the extracted steps.

FIGURE 37.6

Your two queries are now sharing the extracted steps.

The screenshot shows the Power Query Editor interface. On the left, there's a list of 'Queries [2]': 'SalesByEmployee' and 'SalesByBusiness'. 'SalesByBusiness' is highlighted with a green background. To the right is a table preview with two columns: 'Region' and 'Market'. Both columns contain the value 'MIDWEST' for all 14 rows. The 'Market' column also contains the value 'TULSA' for all 14 rows. At the bottom of the preview, it says '5 COLUMNS, 999+ ROWS' and 'Column profiling based on top 1000 rows'. On the far right, there's a 'Query Settings' pane with sections for 'PROPERTIES' (Name: SalesByBusiness) and 'APPLIED STEPS'. Under 'APPLIED STEPS', there's a list: 'Source', 'Changed Type', 'Sorted Rows', and 'Merged Columns'. The 'Merged Columns' step is highlighted with a green background. At the bottom right of the editor, it says 'PREVIEW DOWNLOADED AT 9:58 AM'.

Understanding the Append Feature

Power Query's Append feature allows you to add the rows generated from one query to the results of another query. In other words, you're essentially copying records from one query and adding them to the end of another.

The Append feature comes in handy when you need to consolidate multiple identical tables into one table. For example, if you have tables from the North, Midwest, South, and West regions, you can consolidate the data from each region into one table using the Append feature.

To understand the Append feature better, let's walk through an exercise that consolidates data from four different regions into one table. In this walk-through, we'll use the regional data found on four different tabs in the Appending_Data.xlsx sample file (see Figure 37.7).

Creating the needed base queries

It's important to understand that the Append feature works only on existing queries. That is to say, no matter what kind of data sources you have, you'll need to import them into Power Query before you can append them together. In this case, it means getting all of the region tables into queries.

Follow these steps:

1. Go to the North Data worksheet, select any cell inside the table, and then click **Data** \Rightarrow **Get & Transform Data** \Rightarrow **Get Data** \Rightarrow **From Other Sources** \Rightarrow **From Table/Range**. The Power Query Editor will activate, showing you the contents of the table you just imported. To finalize the creation of the query, you'll need to use one of the Close & Load commands.

FIGURE 37.7

The data found on each region tab needs to be consolidated into one table.

	A	B	C	D	E	
1	Region	Market	Branch Number	Customer Name	State	Effect
2	NORTH	NEWYORK	801211	ALDUN Corp.	NJ	
3	NORTH	NEWYORK	801211	ALFRAN Corp.	NJ	
4	NORTH	NEWYORK	801211	ALLAUD Corp.	NJ	
5	NORTH	NEWYORK	801211	ALLAUD Corp.	NJ	
6	NORTH	NEWYORK	801211	ALLOMU Corp.	NJ	
7	NORTH	NEWYORK	801211	ALPHAP Corp.	NJ	
8	NORTH	NEWYORK	801211	AMPUSA Corp.	NY	
9	NORTH	NEWYORK	801211	ANAQAE Corp.	NY	
10	NORTH	NEWYORK	801211	ANDUVU Corp.	NJ	
11	NORTH	NEWYORK	801211	ANTUSS Corp.	NY	
12	NORTH	NEWYORK	801211	AQBANF Corp.	NJ	
13	NORTH	NEWYORK	801211	ARAWAC Corp.	NY	
14	NORTH	NEWYORK	801211	ATLANT Corp.	NJ	
15	NORTH	NEWYORK	801211	ATLANT Corp.	NJ	
16	NORTH	NEWYORK	801211	ATLANT Corp.	NJ	
17	NORTH	NEWYORK	801211	BALTMU Corp.	NY	
18	NORTH	NEWYORK	801211	BAQTAS Corp.	NY	
19	NORTH	NEWYORK	801211	BECUT Corp.	NY	
20	NORTH	NEWYORK	801211	BLECKP Corp.	NJ	

◀ ▶ North Data Midwest Data South Data West Data

Now, because you're creating this query simply for the purpose of appending it to other queries, there is no need to use the Close & Load command on the workbook. Instead, you can choose to close and load the data as a connection only.

2. On the Home tab of the Power Query Editor, click the drop-down arrow under the Close & Load button and select Close & Load To.
3. In the Import Data dialog box, choose the option Only Create Connection and then click OK.
4. Repeat steps 1 through 3 for the other worksheets in the workbook.

Once you've created queries for each region, activate the Queries & Connections pane (Data \leftrightarrow Queries & Connections) to see all of the queries. As you can see in Figure 37.8, each query is a connection-only query.

Now that your data tables have all been imported into queries, you can start appending the data.

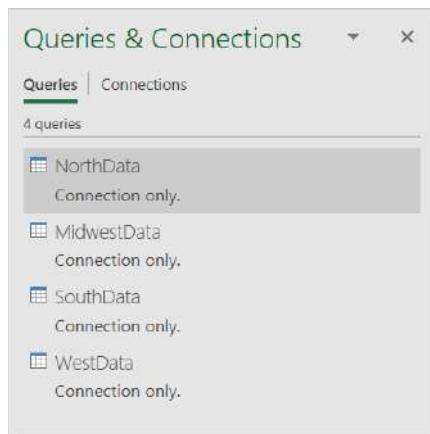
Appending the data

Follow these steps to append data from all other queries to the NorthData query:

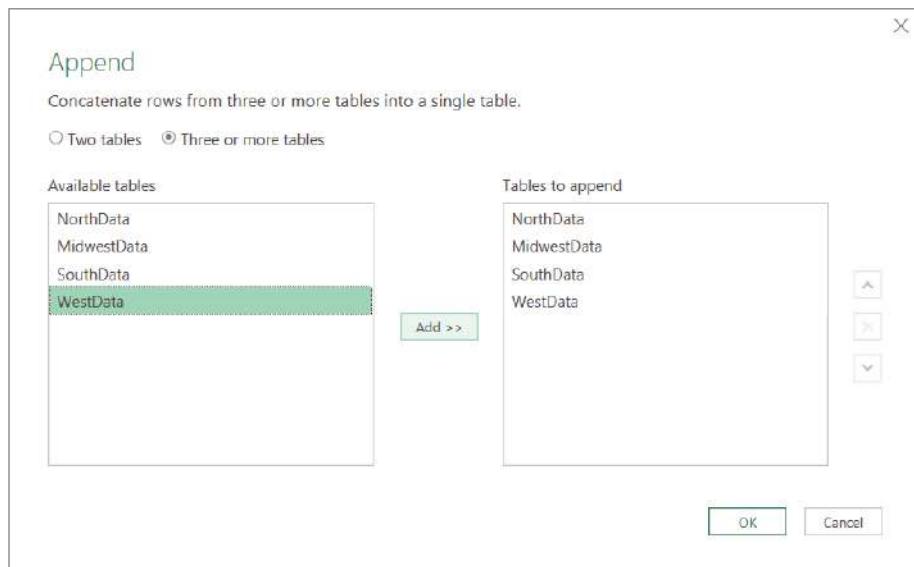
1. In the Queries & Connections pane, right-click the NorthData query and select Append to activate the Append dialog box, shown in Figure 37.9.

FIGURE 37.8

Create a connection-only query for each region.

**FIGURE 37.9**

Appending multiple queries to NorthData



2. Choose the Three or more tables option at the top of the dialog box. The Append dialog box reconfigures to show two list boxes. The list on the left includes all of the existing queries in your workbook. The list on the right contains the query to which you are currently appending data (the NorthData query in this scenario).

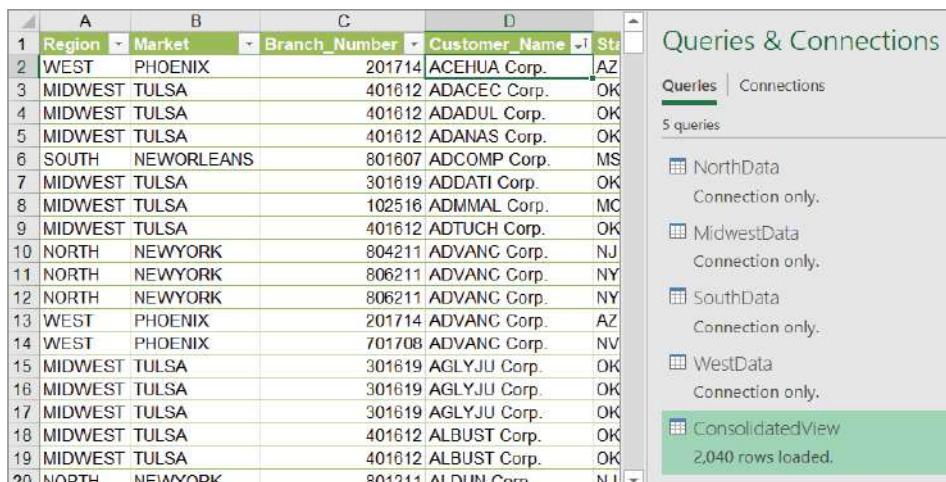
Part V: Understanding Power Pivot and Power Query

3. Select any query you want appended from the left and add it to the list box on the right.
4. Click OK to confirm your selections. The Power Query Editor will launch, giving you the opportunity to review and edit the results. You'll notice Power Query creates a new query called Append1.
5. Rename the query in the Query Settings by typing ConsolidatedView in the Name box.
6. Click the Close & Load button to save and exit the Power Query Editor.

Figure 37.10 illustrates the final output. You have successfully created a consolidated table of region data.

FIGURE 37.10

The final consolidated table of all regional data



A screenshot of the Microsoft Excel ribbon showing the 'Home' tab selected. To the right of the ribbon, there is a 'Queries & Connections' pane. This pane has a 'Queries' section with five entries: 'NorthData', 'MidwestData', 'SouthData', 'WestData', and 'ConsolidatedView'. The 'ConsolidatedView' entry is highlighted with a green background and shows '2,040 rows loaded'. Below the queries, there is a section for 'Connections'.

Region	Market	Branch_Number	Customer_Name	Status
WEST	PHOENIX	201714	ACEHUA Corp.	AZ
MIDWEST	TULSA	401612	ADACEC Corp.	OK
MIDWEST	TULSA	401612	ADADUL Corp.	OK
MIDWEST	TULSA	401612	ADANAS Corp.	OK
SOUTH	NEWORLEANS	801607	ADCOMP Corp.	MS
MIDWEST	TULSA	301619	ADDATI Corp.	OK
MIDWEST	TULSA	102516	ADMMAL Corp.	MO
MIDWEST	TULSA	401612	ADTUCH Corp.	OK
NORTH	NEWYORK	804211	ADVANC Corp.	NJ
NORTH	NEWYORK	806211	ADVANC Corp.	NY
NORTH	NEWYORK	806211	ADVANC Corp.	NY
WEST	PHOENIX	201714	ADVANC Corp.	AZ
WEST	PHOENIX	701708	ADVANC Corp.	NV
MIDWEST	TULSA	301619	AGLYJU Corp.	OK
MIDWEST	TULSA	301619	AGLYJU Corp.	OK
MIDWEST	TULSA	301619	AGLYJU Corp.	OK
MIDWEST	TULSA	401612	ALBUST Corp.	OK
MIDWEST	TULSA	401612	ALBUST Corp.	OK
NORTH	NEWYORK	801211	ALDUN Corp.	NJ

CAUTION

Note in Figure 37.9 that the NorthData query is on both the left and right list boxes. Be careful not to move the NorthData query to the right list box by mistake. If you do, you'll append the query to itself, effectively duplicating all of the records within the query. Unless you have some strange requirement where creating exact copies of records is beneficial, you will want to avoid appending the current query to itself.

Beware of Mismatched Column Labels

When you append one query to another, Power Query first scans the column labels for both queries to capture all column names. It then outputs all distinct column names and consolidates the data from both queries into the appropriate columns. It uses the column labels as a guide for knowing which data goes into which column.

If the column labels in your queries don't match, Power Query will consolidate data for any matching column and will leave null values in any columns that don't match.

For example, imagine you have one query with the column labels Region and Revenue and another query with the column labels Region and SalesAmount. Appending these two records will yield a final table with all three columns: Region, Revenue, and SalesAmount. The records from the first query will be input into the Region and Revenue fields. The records from the second query will be input into the Region and SalesAmount fields. This will essentially leave you with gaps in the Revenue and SalesAmount fields.

The bottom line is to make sure the column labels in your queries are identical before appending. As long as the column labels in each query are identical, Power Query will be able to append the data correctly. Even if the columns in each query are positioned in a different sequence, Power Query is able to use the column labels to get all the data into the correct columns.

Understanding the Merge Feature

You'll often find the need to build queries that join the data between two tables. For example, you may want to join an employee table to a transaction table to create a view that contains both transaction details and information on the employees who logged those transactions.

In the following sections, you'll discover how you can leverage the Merge feature in Power Query to join data from multiple queries.

Understanding Power Query joins

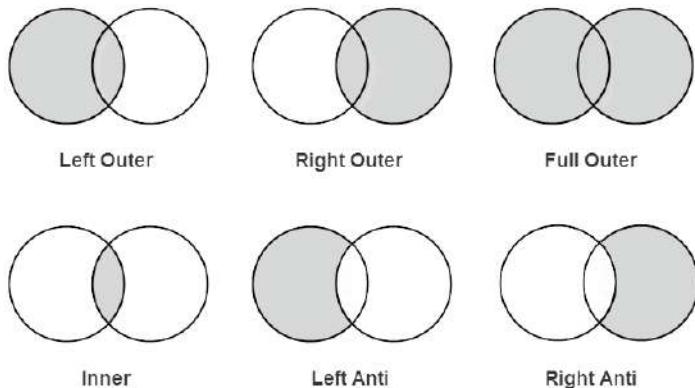
Similar to VLOOKUP in Excel, the Merge feature joins the records from one query to the records in another by matching on some unique identifier. An example of a unique identifier would be Customer ID or Invoice Number.

There are several ways to join two data sets. The kind of join you apply is important because it will determine which records are returned from each data set.

Power Query supports six kinds of joins. As you review each kind of join listed here, feel free to glance at Figure 37.11 to get a visual understanding of each one.

FIGURE 37.11

The kinds of joins supported by Power Query



Left Outer This join tells Power Query to return all of the records from the first query regardless of matching *and* only those records from the second query that have matching values in the joined field.

Right Outer This join tells Power Query to return all of the records from the second query, regardless of matching, *and* only those records from the first query that have matching values in the joined field.

Full Outer This join tells Power Query to return all of the records from both queries regardless of matching.

Inner This join tells Power Query to return only those records from both queries that have matching values.

Left Anti This join tells Power Query to return only those records from the first query that don't match any of the records from the second query.

Right Anti This join tells Power Query to return only those records from the second query that don't match any of the records from the first query.

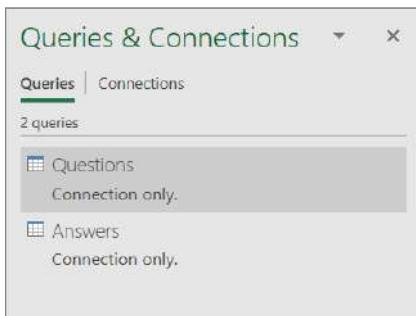
Merging queries

To understand the Merge feature better, let's walk through an exercise that merges interview questions and answers. In this walk-through, we'll use the interview predefined queries found in the `Merging_Data.xlsx` sample file.

As you can see in Figure 37.12, there are two existing queries in the Queries & Connections pane: Questions and Answers. These queries represent the questions and answers from the interview. The goal is to merge these two queries to create a new table showing the questions and answers side by side.

FIGURE 37.12

You need to merge the Questions and Answers queries into one table.

**NOTE**

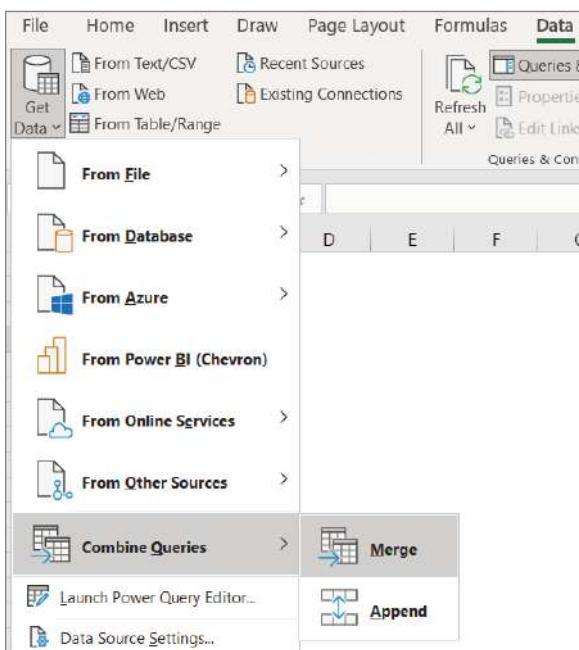
The Merge feature can be used only with existing queries. That is to say, no matter what kind of data sources you have, you'll need to import them into Power Query before you can use them in a merge.

Follow these steps to perform the merge:

1. Click Data \Rightarrow Get & Transform Data \Rightarrow Get Data \Rightarrow Combine Queries \Rightarrow Merge (see Figure 37.13). This will activate the Merge dialog box.

FIGURE 37.13

Activating the Merge dialog box

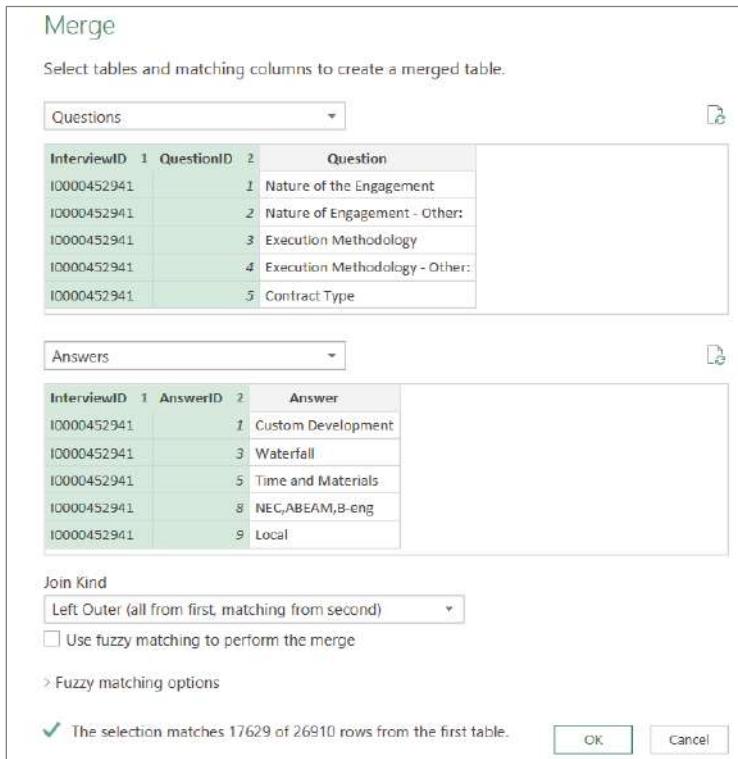


Part V: Understanding Power Pivot and Power Query

Figure 37.14 illustrates the Merge dialog box. The idea here is to use the drop-down boxes to select the queries you want to merge and then choose the columns that define the unique identifier for each record. In this case, the InterviewID and QuestionID/AnswerID fields make up the unique identifier for each record.

FIGURE 37.14

The completed Merge dialog box



2. Select the Questions query in the top drop-down box.
3. Hold down the Ctrl key on the keyboard and then click InterviewID and QuestionID—in that order.
4. Select the Answers query in the lower drop-down box.
5. Hold down the Ctrl key on the keyboard and then click InterviewID and AnswerID—in that order.
6. Use the Join Kind drop-down box to select the kind of join that you want Power Query to use. In this case, the default Left Outer works.
7. Click the OK button to finalize the merge and open the Power Query Editor.

CAUTION

Note the small numbers 1 and 2 in the InterviewID and QuestionID fields in Figure 37.14. These small numbers are assigned based on the order in which you selected them (see steps 3 and 5 earlier).

The order you selected the unique identifiers in each query matters. The two columns tagged with the small number 1 will be joined together regardless of column labels. The two columns tagged with the small number 2 will also be joined together.

NOTE

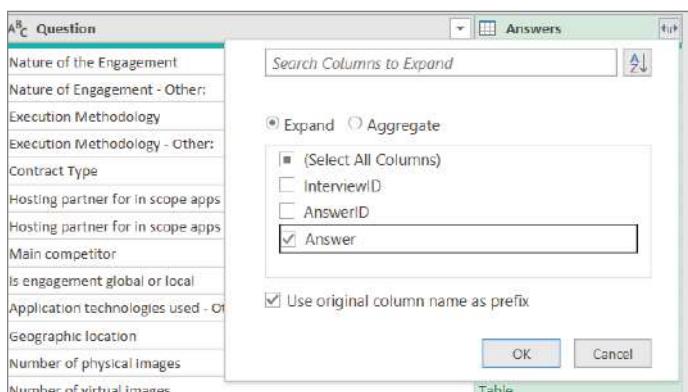
At the bottom of the Merge dialog box, Power Query shows you how many records from the lower query match the top query based on the unique identifiers that you selected. In this case, 17,629 answer records match the 26,910 question records.

Keep in mind there doesn't have to be a 100 percent match for the merge to be valid. There could well be a good reason why the records in the two queries don't all match up. In this case, not every question was answered in all interviews; therefore, the Answers query has fewer records.

- With the newly merged query open in the Power Query Editor, all that is left to do is to click the Expand icon in the newly added field and choose the fields you want to be included in the final output (demonstrated in Figure 37.15). In this case, all you need is the Answer file ID.

FIGURE 37.15

Expand the NewColumn field, and choose the merged fields you want output.



- At this point, you can apply more transformations if needed. When you're happy with the way things look, click the Close & Load command to output the results to the workbook.

Figure 37.16 illustrates the final merged query.

Part V: Understanding Power Pivot and Power Query

FIGURE 37.16

The final table with merged questions and answers

A	B	C	D
InterviewID	QuestionID	Question	Answers.Answer
2	I0000452941	1 Nature of the Engagement	Custom Development
3	I0000452941	2 Nature of Engagement - Other:	
4	I0000452941	3 Execution Methodology	Waterfall
5	I0000452941	5 Contract Type	Time and Materials
6	I0000452941	4 Execution Methodology - Other:	
7	I0000452941	8 Main competitor	NEC,ABEAM,B-eng
8	I0000452941	9 Is engagement global or local	Local
9	I0000452941	6 Hosting partner for in scope apps	
10	I0000452941	10 Application technologies used - Other:	Java
11	I0000452941	7 Hosting partner for in scope apps - Other:	
12	I0000453446	1 Geographic location	Client data center
13	I0000453446	2 Number of physical images	100
14	I0000453446	3 Number of virtual images	300

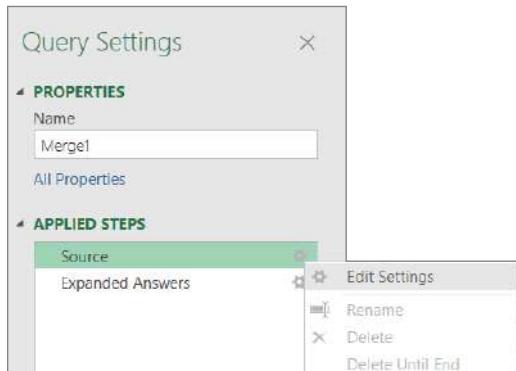
NOTE

When working in the Power Query Editor, you may see the word *null* wherever there are empty values. You don't need to take any special action to get rid of all of the instances of the word *null*. Excel automatically recognizes these as empty values, and it won't show them in the final worksheet.

If you find that you need to adjust or correct a merged query, right-click the query in the Queries & Connections pane and select Edit. In the Power Query Editor, right-click the Source query step and select Edit Settings (see Figure 37.17). Alternatively, you can simply click the gear icon next to the Source query step. This will activate the Merge dialog box, where you can make the necessary changes.

FIGURE 37.17

Right-click the Source query step and select Edit Settings to reactivate the Merge dialog box.



Understanding fuzzy matching

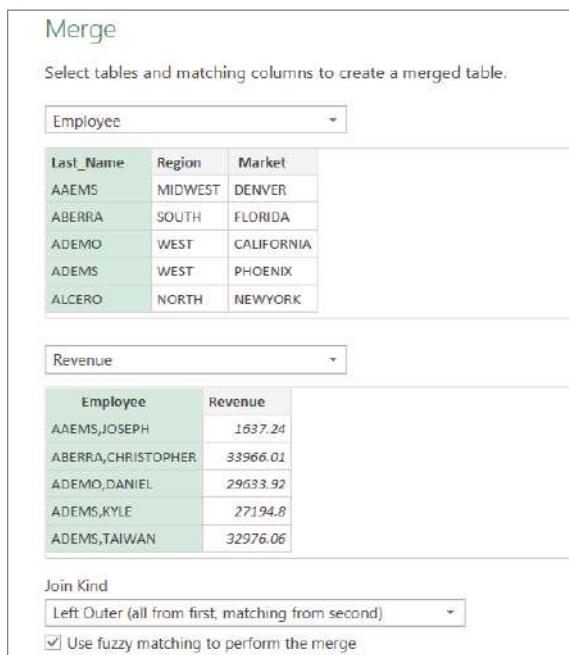
In some cases, the tables you need to merge won't have unique identifiers that match exactly. In these situations, you can leverage Power Query's Fuzzy Match feature. In this walk-through, we'll use the predefined Revenue and Employee queries found in the FuzzyMatch.xlsx sample file.

Follow these steps to perform the merge:

1. Click Data \Rightarrow Get & Transform Data \Rightarrow Get Data \Rightarrow Combine Queries \Rightarrow Merge.
2. Select the Employee query in the top drop-down box, then click the Last_Name column.
3. Select the Revenue query in the lower drop-down box and click the Employee column.
4. Select the Revenue query in the lower drop-down box and click the Employee column.
5. Use the Join Kind drop-down box to select the kind of join that you want Power Query to use. In this case, the default Left Outer works.
6. Click Use Fuzzy Matching to Perform the Merge. The Merge dialog box should look similar to the one shown in Figure 37.18.
7. Click the arrow next to Fuzzy Matching Options to reveal the available configuration options.

FIGURE 37.18

The Merge dialog box with fuzzy matching selected



Let's pause here to review the fuzzy matching options:

Similarity Threshold The Similarity Threshold tells Power Query how similar two values need to be in order to match. The default value (applied if left blank) is .80, which roughly translates to an 80% similarity. Entering a value of 1 in the Similarity Threshold option means you need the two values to match 100% in order to qualify as a match. The minimum value you can enter here is 0; however, this causes all values to match each other. The number you place here really depends on the data you're using. In most cases, the default of .80 is the safest bet, as it will match a decent number of records without resulting in too many false matches.

Ignore Case The Ignore Case option specifies what role character case (uppercase, lowercase, and so on) plays in the matching. The default behavior is case insensitive, which means case is ignored when matching values. If you need to match values taking character case into account, uncheck the Ignore Case option.

Match by Combining Text Parts You can tell Power Query to attempt to combine parts of text in each record in order to complete a match. For instance, if one of your tables contains the value "star light" while the other table contains the value "starlight," Power Query will attempt to match "starlight" with "star light." That is to say, in addition to its normal matching algorithm, it will combine text and try the match again.

Maximum Number of Matches This option defines the maximum number of matching rows that will be returned for each record. For example, if you only want to find one matching row for each record, you would specify a value of 1. The default behavior is to return all matches.

Transformation Table In some cases, you may already have a mapping table containing two values (in separate columns) you've determined should match automatically. You can import your mapping table in a separate query, then point to it in the Transformation Table option.

Continuing our walk-through example, take the following actions.

8. **Enter .60 as the Similarity Threshold.** Since we're dealing with names that are fairly unique, we can risk a lower similarity percent than the default.
9. **Enter 1 as the Maximum Number of Matches.** This will ensure that we only get one match per employee.
10. **Click the OK button to perform the merge and launch the Power Query Editor.**

11. Now you can click the Expand icon next to the Revenue column and choose the columns you want to be included in the final output. In this case, all you need is the Revenue column.
12. Click the Close & Load command to output the results to the workbook.

TIP

At the bottom of the Merge dialog box, Power Query shows you how many records from the lower query match the top query based on the unique identifiers that you selected. When using the fuzzy match feature, it's often useful to try different similarity thresholds in order to get more matches. You'll find that trial and error will often be necessary in order to find the right balance of matching as many records as possible without including too many false matches.

Enhancing Power Query Productivity

IN THIS CHAPTER

- Organizing your queries
- Saving time
- Configuring Power Query options
- Avoiding Power Query performance issues

This chapter offers practical tips for organizing queries and working more efficiently with Power Query. You will also discover a few pointers for optimizing query performance.

Implementing Some Power Query Productivity Tips

Microsoft has added countless features to Power Query. It has truly become a rich toolset with multiple ways to perform virtually any data transformation you can think of. This growth in functionality has paved the way for a good number of tips that can help you work more efficiently with your Power Query models.

Getting quick information about your queries

All of the Power Query queries that live in a particular workbook can be viewed in the Queries & Connections task pane. Click Data \Rightarrow Queries & Connections to activate it.

Here, you can get some quick information about a query simply by hovering over it. You can see the data source for the query and the last time the query was refreshed, and you can get a sneak peek of the data within the query. You can even click a column hyperlink to peek at a particular column (see Figure 38.1).

Part V: Understanding Power Pivot and Power Query

FIGURE 38.1

Hover over a query to get quick information including sneak peeks of column contents.

The screenshot shows the 'Queries & Connections' pane in Power BI. At the top, it displays the title 'Movies'. Below the title is a table with columns: 'Movie Title (click to view)', 'Studio', 'Total Gross /Theaters', and 'Total Gross /Theaters2'. The table lists several movies with their respective studio names and box office figures. A cursor is hovering over the 'Total Gross /Theaters2' column header. To the right of the table is a vertical scroll bar. Below the table, there are sections for 'Columns [9]', 'Last refreshed', 'Load status', and 'Data Sources [1]'. The 'Data Sources' section shows a connection to 'http://www.boxofficemojo.com/monthly'. At the bottom of the pane are buttons for 'VIEW IN WORKSHEET', 'EDIT', '...', and 'DELETE'.

Movie Title (click to view)	Studio	Total Gross /Theaters	Total Gross /Theaters2
Star Wars: The Force Awakens	BV	\$936,662,225	4,134
The Revenant	Fox	\$183,637,894	3,711
Daddy's Home	Par.	\$150,357,137	3,483
Sisters	Uni.	\$87,044,645	2,978
Alvin and the Chipmunks: The Road Chip	Fox	\$85,886,987	3,705
The Big Short	Par.	\$70,259,870	2,529
Joy	Fox	\$56,451,232	2,924
The Hateful Eight	Wein.	\$54,117,416	2,938
Krampus	Uni.	\$42,725,475	2,919

Organizing queries in groups

As you add queries to your workbook, your Queries & Connections pane may start to feel cluttered and disorganized. You can get organized by creating groups for your queries.

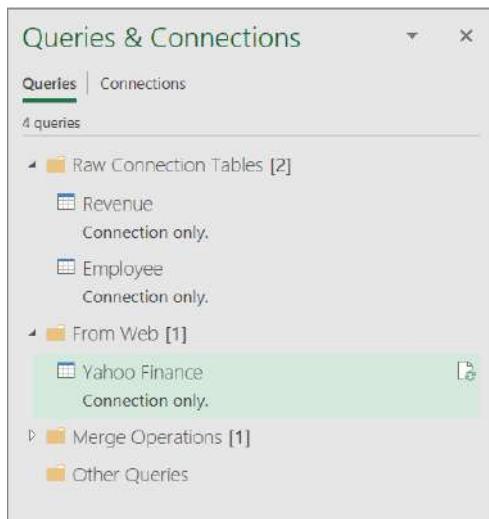
Figure 38.2 illustrates the kinds of groups that you can create. You can create a group for specific stages of data processing. Or how about a group for queries sourced from external databases? What about a group where you store small reference tables? Each group is collapsible, so you can neatly pack away queries that you aren't working with at present.

You can create a group by right-clicking a query in the Queries & Connections pane and selecting Move To Group ↳ New Group.

To move a query to an existing group, right-click the query in the Queries & Connections pane, hover over Move To Group, and then select the group in which you want to see the target query. Right-clicking the group name will expose a set of options for managing the group itself.

FIGURE 38.2

Queries can be organized into groups.



Selecting columns in your queries faster

When dealing with a large table with dozens of columns in the Power Query Editor, it can be a pain to find and select the right columns to work with. You can avoid all that scrolling back and forth by clicking the Choose Columns command on the Home tab of the Power Query Editor.

The Choose Columns dialog box in Figure 38.3 will activate and show you all of the available columns (including custom columns that you may have added). Here, you can easily find and select the columns you need.

Renaming query steps

Each time you apply an action in the Power Query Editor, a new entry is made in the Query Settings pane (shown in Figure 38.4). Query steps serve as a kind of audit trail for all the actions you've taken on the data.

Query steps are automatically given generic names like Uppercased text or Merged Columns. Why not take the time to add some clarity on what each step is doing? You can rename your steps by right-clicking each step and selecting Rename.

Part V: Understanding Power Pivot and Power Query

FIGURE 38.3

Use the Choose Columns command to find and select columns faster.

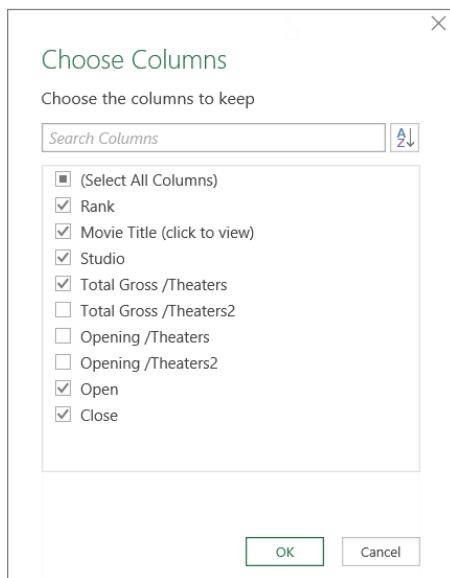
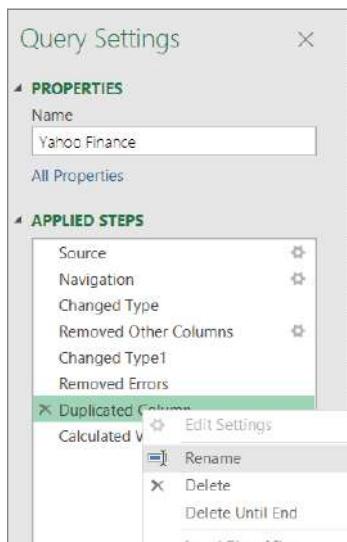


FIGURE 38.4

Right-click query steps to rename them.



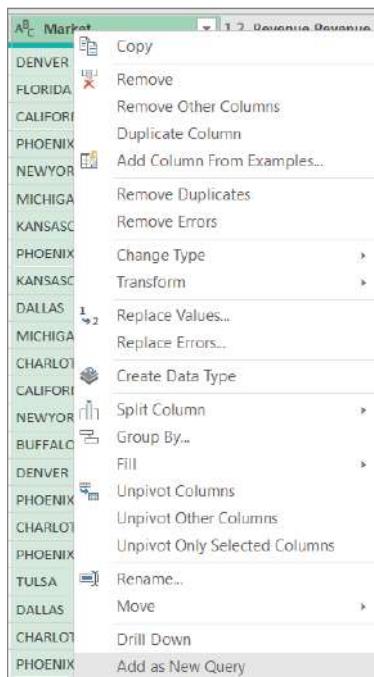
Quickly creating reference tables

There are always a handful of columns in a data set that make for fantastic reference tables. For instance, if your data set contains a column with a list of product categories, it would be useful to create a reference table of all the unique values in that column. Reference tables are often used to map data, feed menu selectors, serve as lookup values, and much more.

While in the Power Query Editor, you can right-click the column from which you want to create a reference table and then select Add as New Query (see Figure 38.5). A new query will be created, using the table that you just pulled from as the source. The Power Query Editor jumps into action, showing just the column you selected. From here, you can use the Power Query Editor to clean up duplicates, remove blanks, and so forth.

FIGURE 38.5

Create a new query from an existing column.



Copying queries to save time

It's always smart to reuse work whenever you can. Why reinvent the wheel when your Queries & Connections pane is full of wheels that you've already created?

Save time by duplicating the queries in your workbook. To do so, activate the Queries & Connections pane, right-click the query you want to copy, and then select Duplicate.

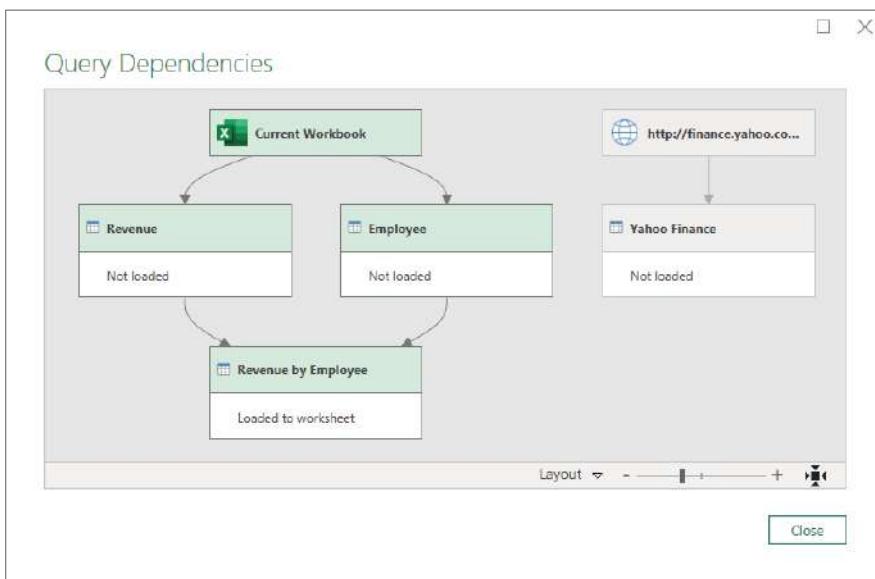
Viewing query dependencies

On the View tab of the Power Query Editor window, you'll see the Query Dependencies command. Clicking this command will activate the Query Dependencies dialog box (see Figure 38.6), where you will see a diagram displaying each of the queries in your workbook. Queries that rely on other queries will have a line connecting them.

The utility of this feature is a bit limited. You can change the zoom and layout of the diagram, but you won't be able to move the objects around. As far as printing goes, you'll have to resort to screen shots, as Power Query doesn't offer a print feature.

FIGURE 38.6

The Query Dependencies dialog box



Setting a default load behavior

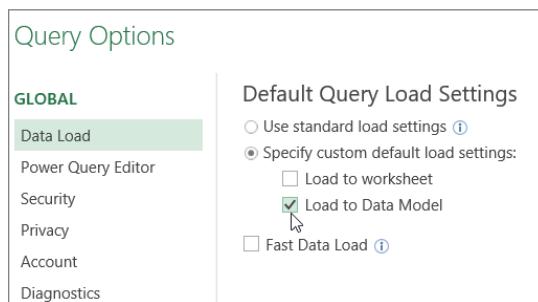
If you're working heavily with Power Pivot as well as Power Query, chances are that you're loading your Power Query queries to the internal data model a majority of the time. If you're one of those analysts who always loads to the data model, you can tweak the Power Query options to load to the data model automatically.

Click Data \Rightarrow Get & Transform Data \Rightarrow Get Data \Rightarrow Query Options to open the Query Options dialog box shown in Figure 38.7. Select Data Load under Global and then choose to specify a

custom default load setting. This enables the options to load to the worksheet or load to the data model by default.

FIGURE 38.7

Use the Global Data Load options to set a default load behavior.



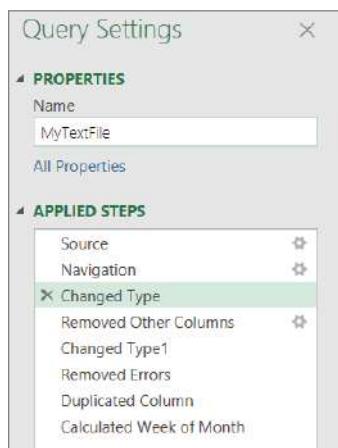
Preventing automatic data type changes

One of the more recent additions to Power Query is the ability to detect data types automatically and proactively change data types. This *type detection* is most often applied when new data is introduced to the query.

For instance, Figure 38.8 shows the query steps after importing a text file. Note the *Changed Type* step. This step was automatically performed by Power Query as part of its type detection feature.

FIGURE 38.8

Power Query automatically adds a step to change data types when data is imported.



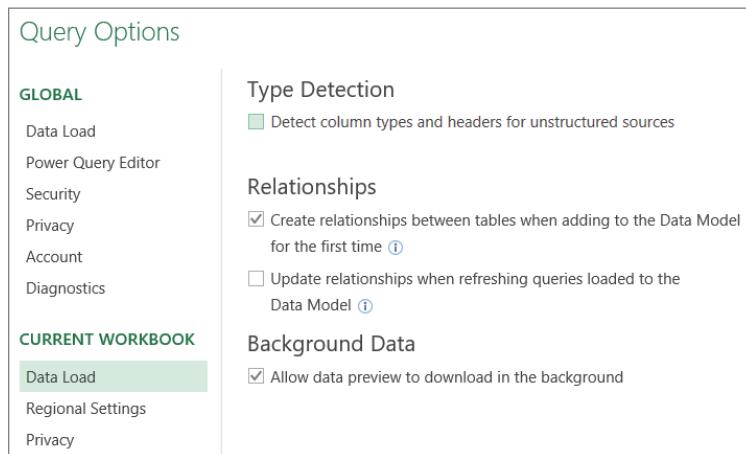
Part V: Understanding Power Pivot and Power Query

Although Power Query does a decent job at guessing what data types should be used, automatic data type changes can cause unexpected issues. Some veterans of Power Query frankly find the type detection feature annoying. If data types need to be changed, *they* want to be the ones to make that determination.

If you'd rather handle data type changes without the help from Power Query's type detection feature, you can turn it off. Click Data \Rightarrow Get & Transform Data \Rightarrow Get Data \Rightarrow Query Options to open the Query Options dialog box, shown in Figure 38.9. Select Data Load under Current Workbook and then uncheck the option to automatically detect column types and headers for unstructured sources.

FIGURE 38.9

Disabling the Type Detection feature



Avoiding Power Query Performance Issues

Because Power Query inherently paves the way for large amounts of data with fairly liberal restrictions, it's possible to end up with queries that are unbearably slow.

When you're wrangling a few thousand records, query performance is not an issue. However, when you are importing and crunching hundreds of thousands of records, performance becomes an issue. There is no getting around the fact that the larger the volume of data, the slower your queries will run. That being said, there are steps that you can take to optimize query performance.

Using views instead of tables

When connecting to an external database, Power Query allows you to import views as well as tables. A *view* is essentially a predefined query on the server itself.

While tables are more transparent, allowing you to see all of the raw unfiltered data, they come with all available columns and rows—whether you need them or not. This often forces you to take extra steps and processing power to remove columns and filter out data that you don't need.

Views not only provide you with cleaner, more user-friendly data, but they can help streamline your data model by limiting the amount of data you import.

Letting your back-end database servers do some crunching

Most Excel analysts who are new to Power Query have a tendency to pull raw data directly from the tables on their external database servers. Once the raw data is in Power Query, they then build transformation and aggregation steps as needed.

Why make Power Query perform transformations that the back-end server could have handled? The reality is that back-end database systems like SQL Server have the ability to shape, aggregate, clean, and transform data much more efficiently than Power Query. Why not utilize their powerful capabilities to massage and shape your data before importing into Power Query?

Instead of pulling raw table data, consider leveraging server-side functions and stored procedures to perform as much data transformation and aggregation work for you as possible. This reduces the amount of processing Power Query will have to do and will naturally improve performance.

Upgrading to 64-bit Excel

If you continue to run into performance issues, there is always the option of getting a better PC. Better, in this case, means moving to a 64-bit PC with 64-bit Excel installed.

The 64-bit version of Excel can access more of your PC's RAM, ensuring that it has the system resources needed to crunch through bigger data sets. In fact, Microsoft recommends 64-bit Excel for anyone working with data models made up of millions of rows.

Before you start installing 64-bit Excel, however, you'll need to consider a few things:

- **Do you already have 64-bit Excel installed?** To check, click File → Account → About Excel. A dialog box will activate showing either 32-bit or 64-bit at the top of the screen.
- **Are your data models large enough?** Unless you're working with large data models, the move to 64-bit may not produce a noticeable difference in your work. What's large? A general rule of thumb is that if you have workbooks that use the internal data model and have a file size upwards of 50 megabytes, you would definitely benefit from an upgrade.
- **Do you have a 64-bit operating system installed on your PC?** 64-bit Excel will not install on a 32-bit operating system. You can find out if you're running a 64-bit operating system by entering "My PC 64-bit or 32-bit" into your favorite search engine. You'll get loads of sites that will walk you through the steps to determine your version.

- **Will your other add-ins stop working?** If you're using other add-ins, be aware that some of them may not be compatible with 64-bit Excel. You don't want to install 64-bit Excel just to find out that your trusted add-ins don't work anymore. Contact your add-in providers to be sure that they are 64-bit compatible. This includes add-ins for all Office products, not just Excel. When you upgrade Excel to 64-bit, you'll have to upgrade the entire Office suite with it.

Disabling privacy settings to improve performance

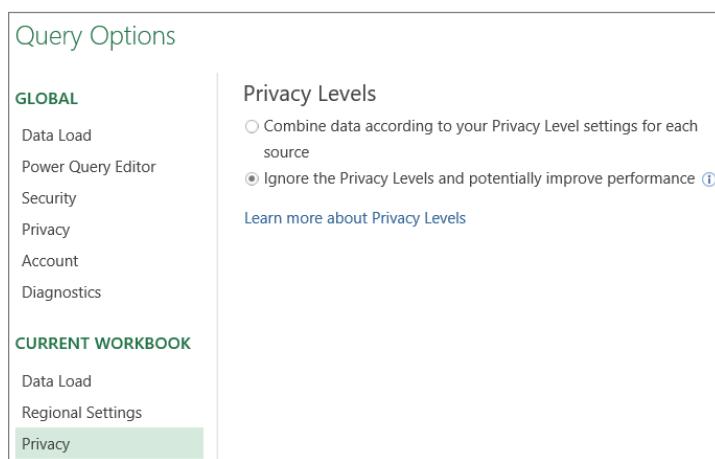
The privacy-level settings in Power Query are designed to protect organizational data as it gets combined with other sources. When you create a query that uses an external data source with an internal data source, Power Query stops the show to ask how you'd like to categorize the data privacy levels of each data source.

For a majority of analysts who deal solely with organizational data, the privacy level settings do little more than slow down queries and cause confusion. Fortunately, there is an option for ignoring privacy levels.

Click Data \Rightarrow Get & Transform Data \Rightarrow Get Data \Rightarrow Query Options to open the Query Options dialog box, shown in Figure 38.10. Select Privacy under Current Workbook and then choose the option to ignore the privacy levels.

FIGURE 38.10

Disabling the privacy level settings



Disabling relationship detection

When you're building a query and choosing Close & Load as the output, Power Query will, by default, attempt to detect relationships between queries and create those relationships within the internal data model. The relationships between queries are primarily driven by the defined

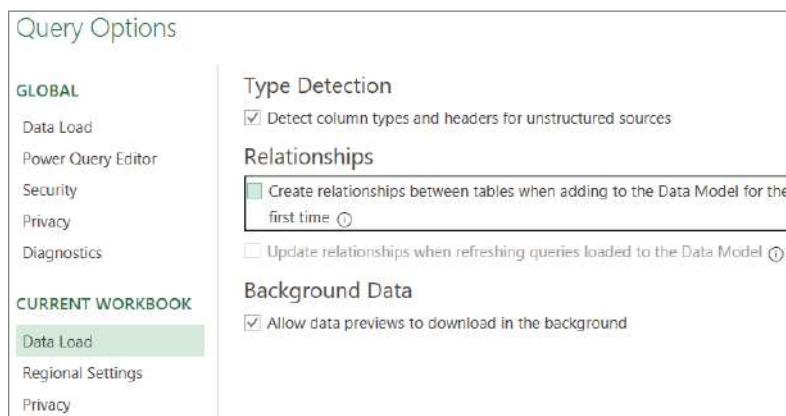
query steps. For instance, if you were to merge two queries and then load the result into the data model, a relationship is automatically created.

In larger data models with a dozen or so tables, Power Query relationship detection can affect performance and increase the time it takes to load the data model. You can avoid this hassle and even gain a performance boost by disabling relationship detection.

Click Data \Rightarrow Get & Transform Data \Rightarrow Get Data \Rightarrow Query Options to open the Query Options dialog box, shown in Figure 38.11. Select Data Load under Current Workbook and then uncheck this option: Create Relationships Between Tables When Adding to the Data Model for the First Time.

FIGURE 38.11

Disabling automatic table relationships



Part VI

Automating Excel

Visual Basic for Applications (VBA) is a powerful programming language built into Excel that you can use to automate routine or repetitive tasks, create custom worksheet formulas, and develop Excel-based applications for other users.

IN THIS PART

Chapter 39
Introducing Visual Basic for Applications

Chapter 40
Creating Custom Worksheet Functions

Chapter 41
Creating UserForms

Chapter 42
Using UserForm Controls in a Worksheet

Chapter 43
Working with Excel Events

Chapter 44
Seeing Some VBA Examples

Chapter 45
Creating Custom Excel Add-Ins

Introducing Visual Basic for Applications

IN THIS CHAPTER

Introducing VBA macros

Creating VBA macros

Recording VBA macros

Writing VBA code

Learning more about VBA

This chapter introduces the Visual Basic for Applications (VBA) macro language—a key component for users who want to customize and automate Excel. This chapter teaches you how to record macros and create simple macro procedures. Subsequent chapters expand upon the topics in this chapter.

Introducing VBA Macros

A macro is a sequence of instructions that automates some aspect of Excel so that you can work more efficiently and with fewer errors. You may create a macro, for example, to format and print your month-end sales report. After you write the macro, you can then execute it every month. Not only does this save you from having to repeat the formatting steps every month, it also ensures that the exact same formatting is applied.

You don't have to be a computer programmer to create and use simple VBA macros. Once they understand a few basics, even casual users can simply turn on Excel's macro recorder to record their actions and convert them into a VBA macro. When you execute a recorded macro, Excel performs the actions again. Once you get the hang of recorded macros, you can write macros from scratch, including code that performs tasks that can't be recorded. For example, you can write procedures that display custom dialog boxes or process data in a series of workbooks and even create special-purpose add-ins.

What You Can Do with VBA

VBA is an extremely rich programming language with thousands of uses. The following list contains just a few things you can do with VBA macros. (Not all of these tasks are covered in this book.)

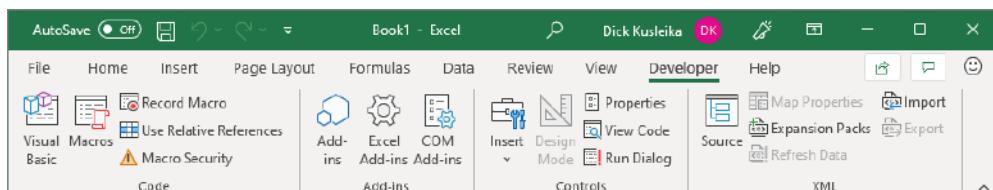
- **Insert boilerplate text.** If you need to enter standard text into a range of cells, you can create a macro to do the typing for you.
- **Automate procedures that you perform frequently.** For example, you may need to prepare a month-end summary. If the task is straightforward, you can develop a macro to do it for you.
- **Automate repetitive operations.** If you need to perform the same action in 12 different workbooks, you can record a macro while you perform the task once—and then let the macro repeat your action in the other workbooks.
- **Create custom commands.** For example, you can combine several Excel commands so that they're executed from a single keystroke or from a single mouse click.
- **Create a simplified “front end” for users who don’t know much about Excel.** For example, you can set up a foolproof data-entry template.
- **Develop new worksheet functions.** Although Excel includes a wide assortment of built-in functions, you can create custom functions that can greatly simplify your formulas.
- **Create complete macro-driven applications.** Excel macros can display custom dialog boxes and respond to new commands added to the Ribbon.
- **Create custom add-ins for Excel.** Add-ins are programs that extend Excel's capabilities.

Displaying the Developer Tab

If you plan to work with VBA macros, make sure you can see the Developer tab on the Excel Ribbon. The Developer tab, which does not appear by default, contains useful commands for VBA users (see Figure 39.1). To display this tab, follow these steps:

FIGURE 39.1

The Developer tab



1. Right-click any Ribbon control and select **Customize the Ribbon** from the **shortcut menu**. The Customize Ribbon tab of the Excel Options dialog box appears.
2. In the list box on the right, place a check mark next to **Developer**.
3. Click **OK** to return to Excel.

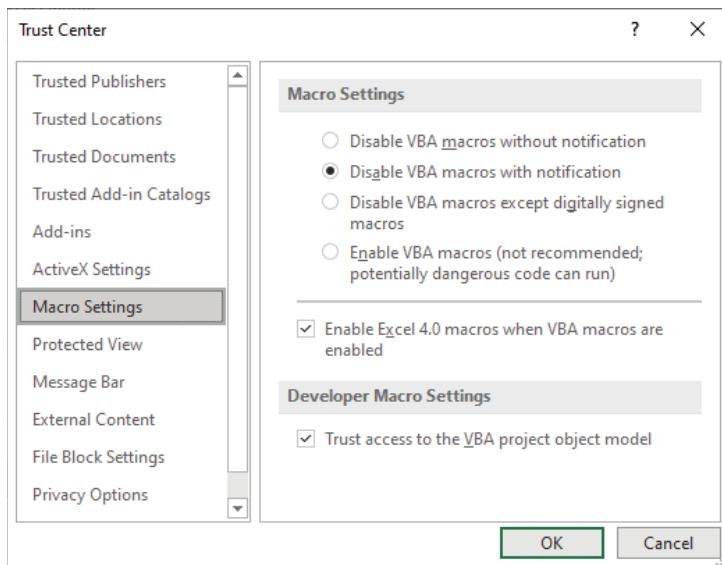
Learning about Macro Security

Macros have the potential to cause serious damage to your computer, such as erasing files or installing malware. Consequently, Microsoft has added macro-security features to help prevent macro-related problems.

Figure 39.2 shows the Macro Settings section of the Trust Center dialog box. To display this dialog box, choose Developer ▾ Code ▾ Macro Security.

FIGURE 39.2

The Macro Settings section of the Trust Center dialog box



By default, Excel uses the Disable VBA Macros with Notification option. With this setting in effect, if you open a workbook that contains macros (and the file is not already trusted), the macros will be disabled and Excel will display a security warning above the Formula bar (see Figure 39.3). If you're certain that the workbook comes from a trusted source, click the Enable Content button in the security warning area, and the macros will be enabled. Excel remembers your decision; if you enable the macros, you won't see the security warning the next time you open that file.

Note

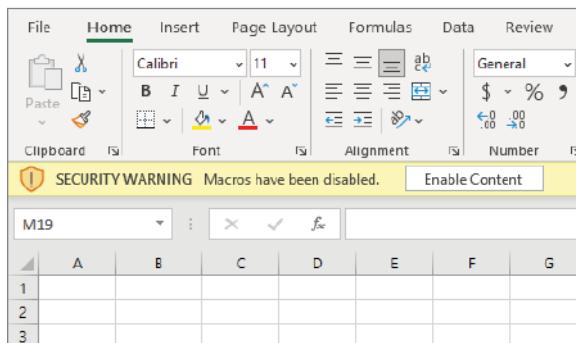
If the Visual Basic Editor (VBE) window is open when you open a workbook that contains macros, Excel does not display the security warning above the Formula bar. Instead, it displays a dialog box with two buttons: Enable Macros and Disable Macros.

Part VI: Automating Excel

Rather than deal with individual workbooks, you may prefer to designate one or more folders as “trusted locations.” Opening workbooks in a trusted location enables macros automatically. You designate trusted folders in the Trusted Locations section of the Trust Center dialog box.

FIGURE 39.3

Excel displays a security warning if a workbook contains macros.



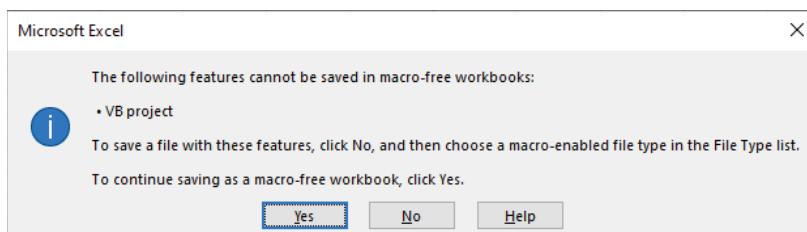
Saving Workbooks That Contain Macros

If you store one or more VBA macros in a workbook, you must save the file with an .xlsm file-name extension.

When you save a workbook that contains macros (or even an empty VBA module), the file format defaults to .xlsm. If you try to save a workbook with macros using the .xlsx file format, Excel will display the warning shown in Figure 39.4. You need to click No and then choose Excel Macro-Enabled Workbook (*.xlsm) from the drop-down list.

FIGURE 39.4

Excel warns you if your workbook contains macros and you attempt to save it as a regular Excel file.



NOTE

Alternatively, you can save the workbook in the old Excel 97–2003 Workbook format (which uses an .xls filename extension) or the Excel Binary Workbook format (which uses an .xlbs filename extension). Both file formats can contain macros.

Looking at Two Types of VBA Macros

A VBA macro (also known as a *procedure*) is usually one of two types: a Sub or a Function. The next two sections discuss the difference.

What's New in the Visual Basic Editor?

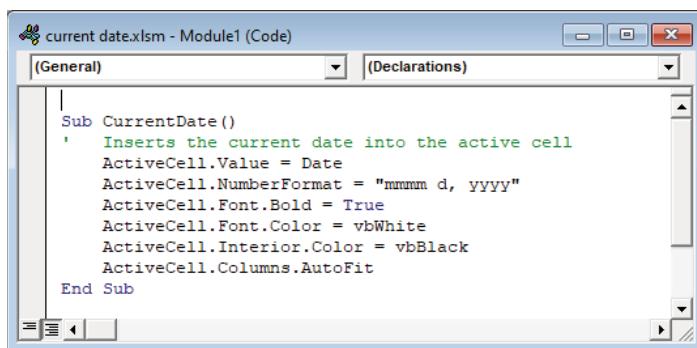
In a word, nothing. The Visual Basic Editor (VBE) has remained untouched for years. The VBA language has been updated to accommodate most of the new Excel features, but the VBE has no new features, and the toolbars and menus work as they always have.

VBA Sub procedures

You can think of a Sub procedure as a new command that either the user or another macro can execute. You can have any number of Sub procedures in an Excel workbook. Figure 39.5 shows a simple VBA Sub procedure. When this code is executed, VBA inserts the current date into the active cell, applies a number format, makes the cell bold, sets the text color to white, sets the background color to black, and adjusts the column width.

FIGURE 39.5

A simple VBA procedure



```

Sub CurrentDate()
    ' Inserts the current date into the active cell
    ActiveCell.Value = Date
    ActiveCell.NumberFormat = "mmmm d, yyyy"
    ActiveCell.Font.Bold = True
    ActiveCell.Font.Color = vbWhite
    ActiveCell.Interior.Color = vbBlack
    ActiveCell.Columns.AutoFit
End Sub

```



A workbook that contains this macro is available on this book's website at www.wiley.com/go/excel365bible. It also includes a button that makes it easy to execute the macro. The file is named *current date.xlsm*.

39

Sub procedures always start with the keyword Sub, the macro's name (every macro must have a unique name), and then a list of arguments inside of parentheses. The parentheses are required even if the procedure doesn't use arguments, like the example in Figure 39.5.

The End Sub statement signals the end of the procedure. The lines in between Sub and End Sub make up the procedure's code.

The CurrentDate macro also includes a comment. Comments are simply notes to yourself, and VBA ignores them. A comment line begins with an apostrophe. You can also put a comment in the same line as a statement. In other words, when VBA encounters an apostrophe, it ignores the rest of the text in the line.

You execute a VBA Sub procedure in any of the following ways:

- Choose Developer ➔ Code ➔ Macros (or press Alt+F8) to display the Macro dialog box. Select the procedure name from the list and then click Run.
- Assign the macro to a control in the Quick Access Toolbar or to a control in the Ribbon.
- Press the procedure's shortcut key combination (if it has one).
- Click a button or other shape that has a macro assigned to it.
- If the VBE is active, move the cursor anywhere within the code and press F5.
- Execute the procedure by calling it from another VBA procedure.
- Enter the procedure name in the Immediate window in the VBE.

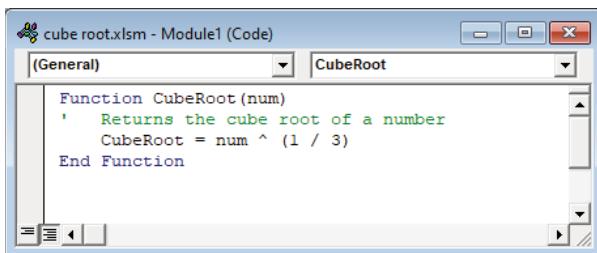
VBA functions

The second type of VBA procedure is a function. A function returns a value (just as a worksheet function always returns a value). A VBA function can be executed by other VBA procedures or used in worksheet formulas, just as you would use Excel's built-in worksheet functions.

Figure 39.6 shows a custom worksheet function. This function is named CubeRoot, and it requires a single argument. CubeRoot calculates the cube root of its argument and returns the result. A Function procedure looks much like a Sub procedure. Notice, however, that function procedures begin with the keyword Function and end with an End Function statement.

FIGURE 39.6

This VBA function returns the cube root of its argument.



A workbook that contains this function is available on this book's website at www.wiley.com/go/excel365bible. The file is named `cube_root.xlsm`.



Creating VBA functions that you use in worksheet formulas can simplify your formulas and enable you to perform calculations that otherwise may be impossible. Chapter 40, “Creating Custom Worksheet Functions,” discusses VBA functions in greater detail.

Some Definitions

If you’re new to VBA, you may be overwhelmed by the terminology. Here are some key definitions to help you keep the terms straight. These terms cover VBA and UserForms (custom dialog boxes)—two important elements that are used to customize Excel:

Code VBA instructions that are stored in a module when you record a macro or enter them manually.

Controls Objects on a UserForm (or in a worksheet) that are used to interact with the user. Examples include buttons, check boxes, and list boxes.

Function One of two types of VBA macros that you can create. (The other is a Sub procedure.) A function returns a value. You can use VBA functions in other VBA macros or in your worksheets.

Macro A set of VBA instructions.

Method An action taken on an object. For example, applying the Clear method to a Range object erases the contents and formatting of the cells.

Module A container for VBA code.

Object An element that you manipulate with VBA. Examples include ranges, charts, and drawing objects.

Procedure Another name for a macro. A VBA procedure can be a Sub procedure or a Function procedure.

Property A particular aspect of an object. For example, a Range object has properties, such as Height, Style, and Name.

Sub procedure One of two types of VBA macros that you can create. The other is a Function.

UserForm A container that holds controls for a custom dialog box and holds VBA code to manipulate the controls.

VBA Visual Basic for Applications. The macro language that is available in Excel as well as in the other Microsoft Office applications.

VBE Visual Basic Editor. The window (separate from Excel) that you use to create VBA macros and UserForms. Press Alt+F11 to toggle between Excel and the VBE.



Chapter 41, “Creating UserForms,” and Chapter 42, “Using UserForm Controls in a Worksheet,” explain UserForms in depth.

Creating VBA Macros

Excel provides two ways to create macros:

- Turn on the macro recorder and record your actions.
- Enter the code directly into a VBA module.

The following sections describe these methods.

Recording VBA macros

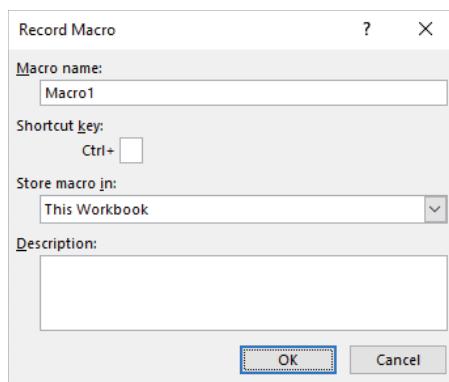
In the following sections, we describe the basic steps that you take to record a VBA macro. In most cases, you can record your actions as a macro and then simply replay the macro; you don't need to look at the code that's automatically generated. If simply recording and playing back macros is as far as you go with VBA, you don't need to be concerned with the language itself (although a basic understanding of how things work can be helpful).

Recording your actions to create VBA code: the basics

The Excel macro recorder translates your actions into VBA code. To start the macro recorder, choose Developer \Rightarrow Code \Rightarrow Record Macro (or click the Record Macro icon on the left side of the status bar). The Record Macro dialog box, shown in Figure 39.7, appears.

FIGURE 39.7

The Record Macro dialog box



The Record Macro dialog box presents several options:

Macro Name The name of the macro. Excel proposes generic names, such as Macro1, Macro2, and so on.

Shortcut Key You can specify a key combination that executes the macro. The key combination always uses the Ctrl key. You can also press Shift when you enter a letter.

For example, pressing Shift while you press H makes the shortcut key combination Ctrl+Shift+H.

CAUTION

Shortcut keys assigned to macros take precedence over built-in shortcut keys. For example, if you assign Ctrl+S to a macro, you can't use the key combination to save your workbook when that macro is available.

Store Macro In The location for the macro. Your choices are the current workbook, your Personal Macro Workbook (see “Storing macros in your Personal Macro Workbook” later in this chapter), or a new workbook.

Description A description of the macro (optional).

To begin recording your actions, click OK; your actions within Excel are converted to VBA code. When you finish recording the macro, choose Developer \Rightarrow Code \Rightarrow Stop Recording, or you can click the Stop Recording button on the status bar. This button replaces the Record Macro button while your macro is being recorded.

NOTE

Recording your actions always results in a new Sub procedure. You can't create a Function procedure by using the macro recorder. You must create function procedures manually.

Recording a macro: a simple example

This example demonstrates how to record a simple macro that inserts your name into the active cell.

To create the macro, start with a new workbook and follow these steps:

1. **Select an empty cell.**

NOTE

Select the cell before you start recording your macro. This step is important. If you select a cell while the macro recorder is turned on, the cell that you select will be recorded into the macro. In such a case, the macro would always use that particular cell, and it would not be a general-purpose macro.

2. **Choose Developer \Rightarrow Code \Rightarrow Record Macro.** The Record Macro dialog box appears. (Refer to Figure 39.7.)
3. **Enter a new single-word name for the macro to replace the default Macro1 name.** For example, type **MyName** as the name.
4. **Assign this macro to the shortcut key Ctrl+Shift+N by entering an uppercase N in the Shortcut Key field.**
5. **Make sure that This Workbook is selected in the Store Macro In field.**

6. Click OK to close the Record Macro dialog box and begin recording your actions.
7. Type your name into the selected cell and then press Enter.
8. Choose Developer \Rightarrow Code \Rightarrow Stop Recording (or click the Stop Recording button on the status bar).

Examining the macro

The macro was recorded in a new module named Module1. To view the code in this module, you must activate the VBE. You can activate the VBE in either of two ways:

- Press Alt+F11.
- Choose Developer \Rightarrow Code \Rightarrow Visual Basic.

In the VBE, the Project window displays a list of all open workbooks and add-ins. This list is displayed as a tree diagram, and you can expand or collapse elements of the tree. The code you recorded previously is stored in Module1 in the Modules folder of the current workbook. When you double-click Module1, the code in the module appears in the Code window.

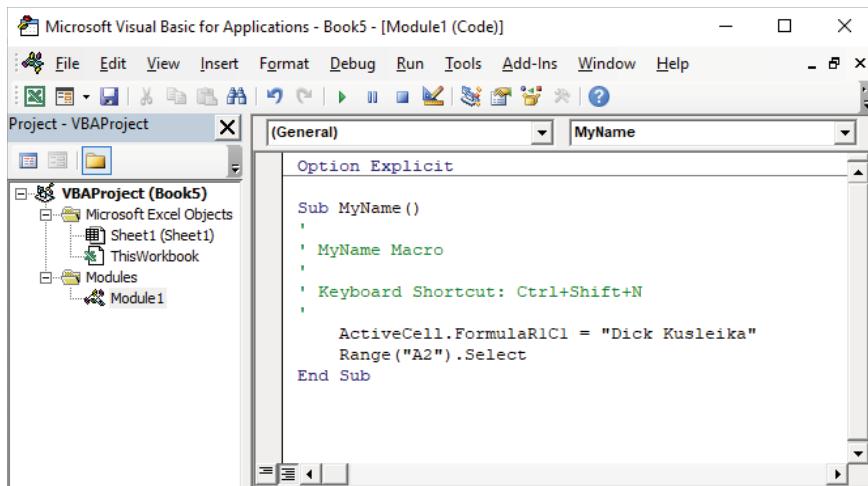
NOTE

If you don't see the list of open workbooks, choose Project Explorer from the View menu or press Ctrl+R to display the Project Explorer.

Figure 39.8 shows the recorded macro, as displayed in the Code window.

FIGURE 39.8

The MyName procedure was generated by the Excel macro recorder.



The macro should look something like this (with your name substituted in, of course):

```
Sub MyName ()  
'  
' MyName Macro  
'  
' Keyboard Shortcut: Ctrl+Shift+N  
'  
ActiveCell.FormulaR1C1 = "Dick Kusleika"  
Range("A2").Select  
End Sub
```

The macro recorded is a Sub procedure that is named MyName. The statements tell Excel what to do when the macro is executed.

Notice that Excel inserted some comments at the top of the procedure. These comments are based on information that appeared in the Record Macro dialog box. These comment lines (which begin with an apostrophe) aren't really necessary and deleting them has no effect on how the macro runs. If you ignore the comments, you'll see that this procedure has only two VBA statements:

```
ActiveCell.FormulaR1C1 = "Dick Kusleika"  
Range("A2").Select
```

The first statement causes the name that you typed while recording the macro to be inserted into the active cell. The FormulaR1C1 part is a property of the Range object, which we'll discuss later. When you press Enter in a cell, Excel moves down one cell (unless you've changed the default behavior). You can guess from this code that the active cell was A1 when the macro was recorded.

Testing the macro

Before you recorded this macro, you set an option that assigned the macro to the Ctrl+Shift+N shortcut key combination. To test the macro, return to Excel by using either of the following methods:

- Press Alt+F11.
- Click the View Microsoft Excel button on the VBE toolbar.

When Excel is active, activate a worksheet. It can be in the workbook that contains the VBA module or in any other workbook. Select a cell and press **Ctrl+Shift+N**. The macro immediately enters your name into the cell and selects cell A2.

Editing the macro

After you record a macro, you can make changes to it. For example, assume that you don't want to select A2, but rather you want to select the cell below the active cell. Press Alt+F11 to activate the VBE window. Then activate Module1 and change the second statement to the following:

```
ActiveCell.Offset(1, 0).Select
```

The edited macro appears as follows:

```
Sub MyName()
'
' MyName Macro
'
' Keyboard Shortcut: Ctrl+Shift+N
'
ActiveCell.FormulaR1C1 = "Dick Kusleika"
ActiveCell.Offset(1, 0).Select
End Sub
```

Test this new macro, and you will see that it performs as it should.

Relative versus absolute recording

If you're going to work with recorded macros, you need to understand the concept of relative versus absolute recording modes. In the previous example, we showed how even a simple macro could act unexpectedly because of an incorrect recording mode.

Normally, when you record a macro, Excel stores exact references to the cells that you select. That is, it uses absolute references. If you press Enter while in cell A1 and the active cell moves down one, the recorded macro will show that you selected cell A2. Similarly, if you were to select cells B1:B10 while you're recording a macro, for example, Excel records this selection as follows:

```
Range("B1:B10").Select
```

This VBA statement means exactly what it says: "Select the cells in the range B1 : B10." When you invoke the macro that contains this statement, the same cells are always selected, regardless of which cell is active.

In the Developer \Rightarrow Code group of the Ribbon there is a Use Relative References control. When you click this control, Excel changes its recording mode from absolute (the default) to relative. When recording in relative mode, selecting a range of cells is translated differently, depending on which cell is active. For example, if you're recording in relative mode and cell A1 is active, selecting the range B1:B10 generates the following statement:

```
ActiveCell.Offset(0, 1).Range("A1:A10").Select
```

This statement can be translated as "From the active cell, move 0 rows down and 1 column right. Then treat this new cell as if it were cell A1. Now select what would be A1 : A10." In other words, a macro that is recorded in relative mode uses the active cell as its starting point and then stores relative references to this cell. As a result, you get different results depending on the location of the active cell. When you replay this macro, the cells that are selected depend on the active cell. This macro selects a range that is 10 rows by 1 column, offset from the active cell by 0 rows and 1 column.

When Excel is recording in relative mode, the Use Relative References control appears with a background color. To return to absolute recording, click the Use Relative References control again and it displays its normal state, with no background color.

Another example

In the first example, the macro behaved oddly by selecting cell A2 after it entered the name. This odd behavior didn't cause any harm or cause the macro not to execute properly. This example demonstrates how choosing the wrong recording mode can cause the macro to work incorrectly. You will record a macro that inserts the current date and time into the active cell. To create the macro, follow these steps:

1. **Select an empty cell.**
2. **Choose Developer \Rightarrow Code \Rightarrow Record Macro.** The Record Macro dialog box appears.
3. **Enter a new, single-word name for the macro to replace the default Macro1 name.** A good name is TimeStamp.
4. **Assign this macro to the shortcut key Ctrl+Shift+T by entering an uppercase T in the Shortcut Key field.**
5. **Make sure that This Workbook is selected in the Store Macro In field.**
6. **Click OK to close the Record Macro dialog box.**
7. **Enter this formula into the selected cell:**
`=NOW()`
8. **With the date cell selected, click the Copy button (or press Ctrl+C) to copy the cell contents to the Clipboard.**
9. **Choose Home \Rightarrow Clipboard \Rightarrow Paste \Rightarrow Values (V).** This step replaces the formula with static text so that the date and time do not update when the worksheet is calculated.
10. **Press Esc to cancel Copy mode.**
11. **Choose Developer \Rightarrow Code \Rightarrow Stop Recording (or click the Stop Recording button on the status bar).**

Running the macro

Activate an empty cell and press Ctrl+Shift+T to execute the macro. There's a pretty good chance that the macro won't work!

The VBA code that is recorded in this macro depends on a setting on the Advanced tab of the Excel Options dialog box, namely, After Pressing Enter, Move Selection. If this setting is enabled (which is the default), the recorded macro won't work as intended because the active cell was changed when you pressed Enter. Even if you reactivated the date cell while recording (in step 7), the macro still fails.

Examining the macro

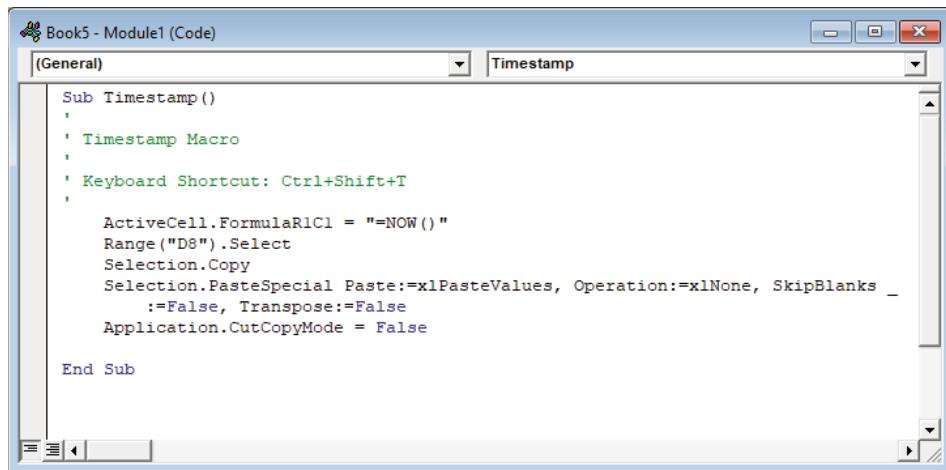
Activate the VBE and look at the recorded code. Figure 39.9 shows the recorded macro, as displayed in the Code window.

The procedure has five statements. The first inserts the NOW() formula into the active cell. The second statement selects cell D8—if the active cell moved to the next cell in step 7, then you had to reselect the active cell in step 8. The exact cell address depends on the location of the active cell when the macro was recorded.

The third statement copies the cell. The fourth statement, which is displayed on two lines (the underscore character means that the statement continues on the next line), pastes the Clipboard contents (as a value) to the current selection. The fifth statement cancels the dashed border around the copied cell.

FIGURE 39.9

The TimeStamp procedure was generated by the Excel macro recorder.



```
Sub Timestamp()
    ' Timestamp Macro
    ' Keyboard Shortcut: Ctrl+Shift+T
    ActiveCell.FormulaR1C1 = "=NOW()"
    Range("D8").Select
    Selection.Copy
    Selection.PasteSpecial Paste:=xlPasteValues, Operation:=xlNone, SkipBlanks _ :=False, Transpose:=False
    Application.CutCopyMode = False
End Sub
```

The problem is that the macro is hard-coded to select cell D8. If you execute the macro when a different cell is active, the code always selects cell D8 before it copies the cell. This is not what you intended, and it causes the macro to fail.

NOTE

You'll also notice that the macro recorded some actions that you didn't make. For example, it specified several options for the PasteSpecial operation. Recording actions that you don't specifically make is just a by-product of the method that Excel uses to translate actions into code.

Rerecording the macro

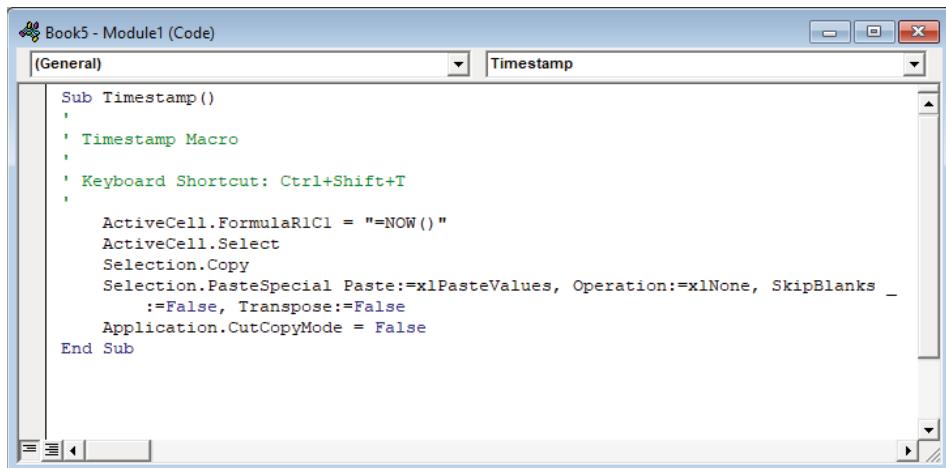
You can fix the macro in several ways. If you understand VBA, you can edit the code so that it works properly, or you can re-record the macro using relative references.

Activate the VBA Editor, delete the existing TimeStamp procedure, and re-record it. Before you start recording, click the Use Relative References command in the Code group of the Developer tab.

Figure 39.10 shows the new macro, recorded with relative references in effect.

FIGURE 39.10

This TimeStamp macro works correctly.



Note that the second line now says ActiveCell.Select instead of using a specific cell address. This solves the problem of copying and pasting the wrong cell, but why would it select a cell that's already selected? It's just one of the curiosities of recording macros. As you select cells, such as when you press Enter after keying a formula, the recorder dutifully registers every selection. In this example, when you pressed Enter in step 7, the recorder recorded this line:

```
ActiveCell.Offset(1, 0).Range("A1").Select
```

Since you didn't do anything while in that cell, however, it didn't save it. When you reselected the cell with the date in it, it replaced the line above with the one you see in the macro.

Testing the macro

When Excel is active, activate a worksheet. The worksheet can be in the workbook that contains the VBA module or in any other workbook. Select a cell and press Ctrl+Shift+T. The macro

immediately enters the current date and time into the cell. You may need to widen the column to see the date and time.

When the result of a macro requires additional manual intervention, it's a sign that the macro could be improved. To widen the column automatically, just add this statement to the end of the macro (before the `End Sub` statement):

```
ActiveCell.EntireColumn.AutoFit
```

More about recording VBA macros

If you followed along with the preceding examples, you should have a better feel for how to record macros and for problems that might occur with even simple macros. If you find the VBA code confusing, don't worry. You don't really have to be concerned with it as long as the macro that you record works correctly. If the macro doesn't work properly, rerecording the macro rather than editing the code is often easier.

A good way to learn about what is recorded is to set up your screen so that you can see the code that is being generated in the VBE window. To do so, make sure that the Excel window isn't maximized; then arrange the Excel window and the VBE window so that both are visible. While you're recording your actions, make sure that the VBE window is displaying the module in which the code is being recorded. (You may have to double-click the module name in the Project Explorer window.)

TIP

If you do a lot of work with VBA, consider adding a second monitor to your system. Then you can display Excel on one monitor and the VBE on the other.

Storing macros in your Personal Macro Workbook

Most user-created macros are designed for use in a specific workbook, but you may want to use some macros in all your workbooks. You can store these general-purpose macros in the Personal Macro Workbook so that they're always available to you. The Personal Macro Workbook is loaded whenever you start Excel. This file, named `personal.xlsx`, doesn't exist until you record a macro, using Personal Macro Workbook as the destination.

NOTE

The Personal Macro Workbook is normally located in a hidden window to keep it out of the way.

To record the macro in your Personal Macro Workbook, select the Personal Macro Workbook option in the Record Macro dialog box before you start recording. This option is in the Store Macro In drop-down box.

If you store macros in the Personal Macro Workbook, you don't have to remember to open the Personal Macro Workbook when you load a workbook that uses macros. When you want to exit, Excel asks whether you want to save changes to the Personal Macro Workbook.

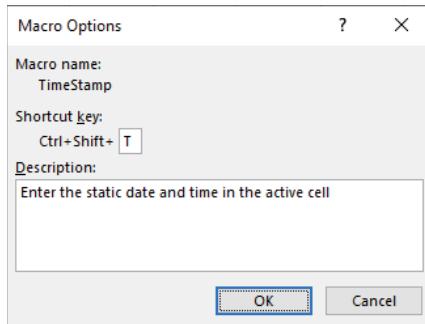
Assigning a macro to a shortcut key

When you begin recording a macro, the Record Macro dialog box gives you an opportunity to provide a shortcut key for the macro. Here's what to do if you'd like to change the shortcut key or provide a shortcut key for a macro that doesn't have one:

1. **Choose Developer ➔ Code ➔ Macros (or press Alt+F8).** The Macro dialog box appears.
2. **Select the macro name from the list.**
3. **Click the Options button.** The Macro Options dialog box, shown in Figure 39.11, appears.
4. **Specify the shortcut key.** Use a single letter (for a Ctrl+letter shortcut) or press Shift and enter an uppercase letter (for a Ctrl+Shift+letter shortcut).
5. **(Optional) Enter a description.**
6. **Click OK to return to the Macro dialog box.**
7. **Click Cancel to close the Macro dialog box.**

FIGURE 39.11

Use the Macro Options dialog box to add or change a shortcut key for a macro.



Assigning a macro to a button

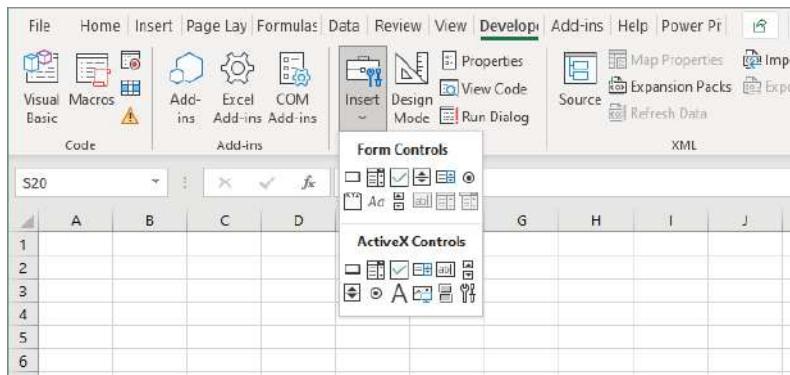
After you record a macro and test it, you may want to assign the macro to a button placed in a worksheet. You can follow these steps to do so:

1. **If the macro is a general-purpose one that you plan to use in more than a single workbook, make sure the macro is stored in your Personal Macro Workbook.**

2. Choose Developer \Rightarrow Controls \Rightarrow Insert and then click the icon identified as **Button (Form Control)**. Figure 39.12 shows the list of controls. Move your mouse pointer over the icons, and you will see a ScreenTip that describes the control.
3. Click the worksheet and drag to draw the button. When you release the mouse button, the Assign Macro dialog box appears.
4. Select the macro from the list.
5. Click OK to close the Assign Macro dialog box.
6. (Optional) Change the text that appears on the button to make it descriptive; right-click the button, choose Edit Text from the shortcut menu, and make your changes.

FIGURE 39.12

Adding a button to a worksheet so that it can be used to execute a macro



After you perform these steps, clicking the button executes the assigned macro.

Adding a macro to your Quick Access Toolbar

You can also assign a macro to a button on your Quick Access Toolbar:

1. Right-click the Quick Access Toolbar and choose **Customize Quick Access Toolbar** from the shortcut menu. The Quick Access Toolbar tab of the Excel Options dialog box appears.
2. Select Macros from the drop-down list on the left.
3. At the top of the list on the right, choose **For All Documents**. You can alternatively choose the workbook that contains the macro. If you do, the Quick Access Toolbar tool will be hidden when any other workbook is active.

4. Select your macro and click the Add button.
5. To change the icon or displayed text, click the Modify button, select an icon from the symbol list or type a new name, and click OK.
6. Click OK to close the Excel Options dialog box.

After performing these steps, your Quick Access Toolbar will display a button that executes your macro.

Writing VBA code

As we demonstrated in the preceding sections, the easiest way to create a simple macro is to record your actions. To develop more complex macros, however, you have to enter the VBA code manually—in other words, write a program. To save time (and assist in the learning process), you can often start with a recorded macro and edit it.

Before you begin writing VBA code, it helps to have a good understanding of such topics as objects, properties, and methods. As you progress, you will become familiar with common programming constructs, such as looping and If-Then statements.

The following sections are an introduction to VBA programming, which is essential if you want to write (rather than record) VBA macros. It isn't intended to be a complete instructional guide. Our book *Excel 2019 Power Programming with VBA* (Wiley, 2019) covers all aspects of VBA and advanced spreadsheet application development.

The basics: entering and editing code

Before you can enter code, you must insert a VBA module into the workbook. If the workbook already has a VBA module, you can use the existing module sheet for your new code.

Follow these steps to insert a new VBA module:

1. Press Alt+F11 to activate the VBE window. The Project window displays a list of all open workbooks and add-ins.
2. In the Project window, locate and select the workbook you're working in.
3. Choose Insert ↗ Module. VBA inserts a new (empty) module into the workbook and displays it in the code window.

A VBA module, which is displayed in a separate window, works like a text editor. You can move through the window, select text, insert text, copy, cut, paste, and so on.

VBA Coding Tips

When you enter code into a module, you're free to use indenting and blank lines to make the code more readable. In fact, this is an excellent habit.

After you enter a line of code (by pressing Enter), it's evaluated for syntax errors. If no errors are found, the line of code is reformatted and colors are added to keywords and identifiers. This automatic reformatting adds consistent spaces (before and after an equal sign, for example) and removes extra spaces that aren't needed. If a syntax error is found, you receive a pop-up message, and the line is displayed in a different color (red, by default). You need to correct your error before you can execute the macro.

A single statement can be as long as you need. However, you may want to break the statement into two or more lines. To do so, insert a space followed by an underscore (_). The following code, although written as two lines, is actually a single VBA statement:

```
Sheets ("Sheet1") .Range ("B1") .Value = _  
Sheets ("Sheet1") .Range ("A1") .Value
```

You can insert comments freely into your VBA code. The comment indicator is an apostrophe or single quote character ('). Any text that follows a single quote on that line is ignored. A comment can be a line by itself, or it can be inserted after a statement. The following examples show two comments:

```
' Assign the values to the variables  
Rate = .085 'Rate as of November 16
```

The Excel object model

VBA is a language designed to manipulate objects. Some objects are contained in the language itself, but most of the objects that you will use when programming VBA for Excel come from the Excel object model.

At the top of the object model is the Application object. This object represents Excel itself, and all other objects are below it in the hierarchy. One way to think about writing code is to ask which object you'd like to change and which property or method controls the aspect of the object you'd like to change. For example, if you want to force users to enter data using the Formula bar as opposed to entering data directly in cells, you can change the EditDirectlyInCell property of the Application object.

If you don't know which object or property to change—and when you're just starting out, you won't—you can use the macro recorder. Start recording a macro and make the changes. Then see what the recorder came up with. If you record a macro and choose File \Rightarrow Options \Rightarrow Advanced and uncheck Allow editing directly in cells, you will see this recorded macro.

```
Sub Macro1()  
'  
' Macro1 Macro
```

```
'  
' Application.EditDirectlyInCell = False  
End Sub
```

Now you know that the `EditDirectlyInCell` property of the `Application` object is where that setting is stored, and you can use that in your own code.

Objects and collections

In addition to the `Application` object, there are hundreds of other objects available to use in your code, such as `Range`, `Chart`, and `Shape` objects. These objects are arranged in a hierarchy with the `Application` object at the top.

Objects of the same type are contained in collections. (Collections are also an object.) Collection objects are named using the plural form of the objects that they contain. Each open workbook is a `Workbook` object, and all the open workbooks are in the `Workbooks` collection object. Similarly, the `Shapes` collection object contains all the objects of type `Shape`.

There are a few places where the plural collection naming convention breaks down. The `Range` object is an important exception that we'll discuss later in this chapter.

You reference a specific object by traversing the hierarchy. To reference cell A1, you might use code such as the following:

```
Application.Workbooks.Item("MyBook.xlsx").Worksheets.Item(1).Range("A1")
```

Fortunately, VBA provides some shortcuts. Since `Application` is at the top, you can omit it, and VBA will know what you want. VBA also provides default properties for some objects. All collection objects have a default property called `Item` that's used to access one of the collections. You can shorten your code as follows:

```
Workbooks("MyBook.xlsx").Worksheets(1).Range("A1")
```

When accessing an `Item` of a collection, you can ask for it by name or by number. For the `Workbooks` collection, we passed in the name of the workbook that we wanted, and it returned the `Workbook` object with that name. For the `Worksheets` collection, however, we asked for the first `Worksheet` object in the collection regardless of its name.

39

Properties

The objects you work with have properties, which you can think of as attributes of the objects. For example, a `Range` object has properties such as `Column`, `Row`, `Width`, and `Value`. A `Chart` object has properties such as `Legend`, `ChartTitle`, and so on. `ChartTitle` is also an object, with properties such as `Font`, `Orientation`, and `Text`. Excel has many objects, and each has its own set of properties. You can write VBA code to do the following:

- Examine an object's current property setting and take some action based on it.
- Change an object's property setting.

You refer to a property in your VBA code by placing a period (a dot) and the property name after the object's name. For example, the following VBA statement sets the Value property of a range named frequency to 15. (That is, the statement causes the number 15 to appear in the range's cells.)

```
Range("frequency").Value = 15
```

You might have noticed that we used the dot operator to traverse the hierarchy in the previous section, and we're using it again to access properties. That's not a coincidence. Properties can contain a lot of different values, and they can also contain other objects. When we used `Application.Workbooks("MyBook.xlsx")`, we were actually accessing the Workbooks property of the Application object. That property returns a Workbooks collection object.

Some properties are read-only, which means you can examine the property but you can't change the property. For a single-cell Range object, the Row and Column properties are read-only properties: you can determine where a cell is located (in which row and column), but you can't change the cell's location by changing these properties.

A Range object also has a Formula property, which is not read-only; that is, you can insert a formula into a cell by changing its Formula property. The following statement inserts a formula into cell A12 by changing the cell's Formula property:

```
Range("A12").Formula = "=SUM(A1:A10)"
```

NOTE

Contrary to what you may think, Excel doesn't have a Cell object. When you want to manipulate a single cell, you use the Range object and reference only one cell.

The Application object has several useful properties that refer to where the user is in the program:

`Application.ActiveWorkbook`: Returns the active workbook (a Workbook object) in Excel.

`Application.ActiveSheet`: Returns the active sheet (a Sheet object) of the active workbook.

`Application.ActiveCell`: Returns the active cell (a Range object) of the active window.

`Application.Selection`: Returns the object that is currently selected in the active window of the Application object. This can be a Range, a Chart, a Shape, or some other selectable object.

In many cases, you can refer to the same object in a number of different ways. Assume you have a workbook named `Sales.xlsx` and that it's the only workbook open. Furthermore, assume that this workbook has one worksheet, named `Summary`. Your VBA code can refer to the `Summary` sheet in any of the following ways:

```
Workbooks("Sales.xlsx").Worksheets("Summary")
Workbooks(1).Worksheets(1)
Workbooks(1).Sheets(1)
Application.ActiveWorkbook.ActiveSheet
ActiveWorkbook.ActiveSheet
ActiveSheet
```

The method that you use is determined by how much you know about the workspace. For example, if more than one workbook is open, the second or third method isn't reliable. If you want to work with the active sheet (whatever it may be), any of the last three methods would work. To be absolutely sure you're referring to a specific sheet on a specific workbook, the first method is your best choice.

Methods

Objects also have methods. You can think of a method as an action taken with an object. Generally, methods are used to interact with the computer outside the Excel application or to modify multiple properties at once. For example, Range objects have a Clear method. The following VBA statement clears a Range, an action that is equivalent to selecting the Range and then choosing Home \Rightarrow Editing \Rightarrow Clear \Rightarrow Clear All:

```
Range("A1:C12").Clear
```

The Clear method involves changing several properties of the Range object at once. It includes setting the Value property to Empty (clearing its contents), the Bold property of the Font object to False (clearing all formats), and the Comments property to Nothing (deleting the cell's comment). It's doing a few other things as well.

If your code interacts with files on a disk, printers, or other aspects of your computer outside of Excel, you'll probably use a method. Each Workbook object has a Name property that's read-only. You can't change the Name property by setting it directly as follows:

```
Workbooks(1).Name = "xyz.xlsx"
```

That will fail. However, you can change a Workbook's name by using the SaveAs method:

```
Workbooks(1).SaveAs "xyz.xlsx"
```

In addition to changing the name property, SaveAs changes a few other properties, and it also writes the file to the hard drive.

In VBA code, methods look like properties because they're connected to the object with a "dot." However, methods and properties are different concepts.

The Range object

The Range object is special. As you might imagine, it's central to the Excel object model. Workbooks and worksheets exist only to hold cells. But while the Worksheets collection holds a bunch of Worksheet objects and the Shapes collection holds a bunch of Shape objects, the Range object works differently.

A single cell is a Range object. A group of cells is also a Range object, but not a Ranges collection object. It's one of the few objects that breaks the plural naming convention of collections.

Most collection objects have a default property of Item. That allows you to write the following:

```
Workbooks(1)
```

instead of the following:

```
Workbooks.Item(1)
```

Generally, if an object has an Item property, that's the default. For objects that aren't collections, if they have a Value property, that's the default. For example, these two lines of VBA are identical because Value is the default property of Checkbox objects:

```
If Sheet1.CheckBox1.Value = True Then  
If Sheet1.CheckBox1 = True Then
```

The Range object has both an Item property and a Value property. In some contexts, the Item property is the default, while in others, it's the Value property. The good news is that VBA does a pretty good job of picking the right one.

Variables

Like all programming languages, VBA enables you to work with variables. In VBA (unlike in some languages), you don't need to declare variables explicitly before you use them in your code (although doing so is definitely a good practice).

NOTE

If your VBA module contains an Option Explicit statement at the top of the module, you must declare all variables in the module. Undeclared variables will result in a compile error, and your procedures will not run.

In the following example, the value in cell A1 on Sheet1 is assigned to a variable named Rate:

```
Rate = Worksheets("Sheet1").Range("A1").Value
```

After the statement is executed, you can work with the variable Rate in other parts of your VBA code.

Controlling execution

VBA uses many constructs that are found in other programming languages. These constructs are used to control the flow of execution. We'll now introduce a few of the more common programming constructs.

The If-Then construct

One of the most important control structures in VBA is the If-Then construct, which gives your applications decision-making capability. The basic syntax of the If-Then structure is as follows:

```
If condition Then statements [Else elsestatements]
```

In plain English, if a condition is true, then a group of statements will be executed. If you include the Else clause, then another group of statements will be executed if the condition is not true.

The following is an example (which doesn't use the optional Else clause). This procedure checks the active cell. If it contains a negative value, the cell's font color is changed to red. Otherwise, nothing happens:

```
Sub CheckCell()
    If ActiveCell.Value < 0 Then ActiveCell.Font.Color = vbRed
End Sub
```

Here's another multiline version of that procedure that uses an Else clause. Because it uses multiple lines, you must include an End If statement. This procedure colors the active cell text red if it's a negative value and green otherwise:

```
Sub CheckCell()
    If ActiveCell.Value < 0 Then
        ActiveCell.Font.Color = vbRed
    Else
        ActiveCell.Font.Color = vbGreen
    End If
End Sub
```

For-Next loops

You can use a For-Next loop to execute one or more statements a number of times. Here's an example of a For-Next loop:

```
Sub SumSquared()
    Total = 0
    For Num = 1 To 10
        Total = Total + (Num ^ 2)
    Next Num
    MsgBox Total
End Sub
```

This example has one statement between the For statement and the Next statement. This single statement is executed 10 times. The variable Num takes on successive values of 1, 2, 3, and so on, up to 10. The variable Total stores the sum of Num squared added to the previous value of Total. The result is a value that represents the sum of the first 10 integers squared. This result is displayed in a message box.

Do loops

For-Next loops execute a set of statements a particular number of times. Do loops execute a set of statements until a particular condition exists or stops existing.

```
Sub SumSquaredTo500()
    Total = 0
    num = 0
    Do
        num = num + 1
        Total = Total + (num ^ 2)
    Loop Until Total >= 500
    MsgBox num & Space(1) & Total
End Sub
```

This procedure keeps summing squares until the total reaches 500 or higher. With Do loops, you can check the condition on the Do line or the Loop line, but not both. The four options are as follows:

- Do Until
- Do While
- Loop Until
- Loop While

The With-End With construct

A construct that you sometimes encounter if you record macros is the With-End With construct. This is a shortcut way of dealing with several properties or methods of the same object. The following is an example:

```
Sub AlignCells()
    With Selection
        .HorizontalAlignment = xlCenter
        .VerticalAlignment = xlCenter
        .WrapText = False
        .Orientation = xlHorizontal
    End With
End Sub
```

The following macro performs the same operations but doesn't use the With-End With construct:

```
Sub AlignCells()
    Selection.HorizontalAlignment = xlCenter
```

```

    Selection.VerticalAlignment = xlCenter
    Selection.WrapText = False
    Selection.Orientation = xlHorizontal
End Sub

```

The Select Case construct

The Select Case construct is useful for choosing among two or more options. The following example demonstrates the use of a Select Case construct. In this example, the active cell is checked. If its value is less than 0, it's colored red. If it's equal to 0, it's colored blue. If the value is greater than 0, it's colored black:

```

Sub CheckCell()
    Select Case ActiveCell.Value
        Case Is < 0
            ActiveCell.Font.Color = vbRed
        Case 0
            ActiveCell.Font.Color = vbBlue
        Case Is > 0
            ActiveCell.Font.Color = vbBlack
    End Select
End Sub

```

Any number of statements can go below each Case statement, and they all are executed if the case is true.

A macro that can't be recorded

The following is a VBA macro that can't be recorded because it uses programming concepts that must be entered manually. This macro creates a list of all formulas on the active sheet. The list is stored on a new worksheet:

```

Sub ListFormulas()
    ' Create a range variable
    Set InputRange = ActiveSheet.UsedRange
    ' Add a new sheet and save in a variable
    Set OutputSheet = Worksheets.Add
    ' Variable for the output row
    OutputRow = 1
    ' Loop through the range
    For Each cell In InputRange
        If cell.HasFormula Then
            OutputSheet.Cells(OutputRow, 1) = "" & cell.Address
            OutputSheet.Cells(OutputRow, 2) = "" & cell.Formula
            OutputRow = OutputRow + 1
        End If
    Next Cell
End Sub

```



A workbook that contains this example is available on this book's website at www.wiley.com/go/excel365bible. The file is named `list formulas.xlsx`.

Although this macro may look complicated, it's fairly simple when you break it down. Here's how it works:

1. The macro creates an object variable named InputRange. This variable corresponds to the used range on the active sheet (avoiding the need to check every cell).
2. It then adds a new worksheet and assigns the worksheet to an object variable named OutputSheet. The OutputRow variable is set to 1. This variable is incremented later.
3. The For-Next loop examines each cell in the InputRange. If the cell has a formula, the cell's address and formula are written to the OutputSheet. The OutputRow variable is also incremented.

Figure 39.13 shows the part of the result of running this macro—a handy list of all formulas in the worksheet.

FIGURE 39.13

The ListFormulas macro creates a list of all formulas in a worksheet.

	A	B	C	D
1	\$G\$2	=B9/2		
2	\$A\$8	=SUM(A2:A7)		
3	\$B\$8	=SUM(B2:B7)		
4	\$C\$8	=SUM(C2:C7)		
5	\$D\$8	=SUM(D2:D7)		
6	\$A\$9	=AVERAGE(A2:A7)		
7	\$B\$9	=AVERAGE(B2:B7)		
8	\$C\$9	=AVERAGE(C2:C7)		
9	\$D\$9	=AVERAGE(D2:D7)		
10	\$E\$9	=AVERAGE(E2:E7)		
11	\$A\$15	=RANDBETWEEN(1,1000)		
12	\$B\$15	=RANDBETWEEN(1,1000)		
13	\$C\$15	=RANDBETWEEN(1,1000)		
14	\$D\$15	=RANDBETWEEN(1,1000)		
15	\$E\$15	=RANDBETWEEN(1,1000)		
16	\$A\$16	=RANDBETWEEN(1,1000)		
17	\$B\$16	=RANDBETWEEN(1,1000)		
18	\$C\$16	=RANDBETWEEN(1,1000)		
19	\$D\$16	=RANDBETWEEN(1,1000)		
20	\$E\$16	=RANDBETWEEN(1,1000)		

As macros go, this example is okay, but it's certainly not perfect. It's not very flexible, and it doesn't include error handling. For example, if the workbook structure is protected, trying to add a new sheet will cause an error.

Learning More

This chapter is a basic introduction to VBA. If this is your first exposure to VBA, you might be a bit overwhelmed by objects, properties, and methods. It's frustrating when you know what

you want to do but don't know which objects, properties, or methods you need to do it. Fortunately, several good ways are available to learn about objects, properties, and methods.

Read the rest of the book Subsequent chapters in this section contain additional information and many more examples.

Record your actions The best way to become familiar with VBA is to turn on the macro recorder and record the actions you make in Excel. You can then examine the code to gain some insights regarding the objects, properties, and methods.

Use the Help system The main source of detailed information about Excel's objects, methods, and procedures is the VBA Help system. Help is thorough and easy to access. When you're in a VBA module, just move the cursor to a property or method and press F1. You get help that describes the word that is under the cursor. All VBA Help is online, so you must be connected to the Internet to use the Help system.

Get another book Several books are devoted exclusively to using VBA with Excel. Our book *Excel 2019 Power Programming with VBA* (Wiley, 2019) is one of them.

Creating Custom Worksheet Functions

IN THIS CHAPTER

Getting an overview of VBA functions

Looking at function procedures

Focusing on function procedure arguments

Debugging custom functions

Pasting custom functions

As mentioned in the preceding chapter, you can create two types of VBA procedures: Sub procedures and Function procedures. This chapter focuses on Function procedures.

Introducing VBA Functions

Function procedures that you write in VBA are quite versatile. You can use these functions in two situations:

- You can call the function from a different VBA procedure.
- You can use the function in formulas that you create in a worksheet.

This chapter focuses on creating functions for use in your formulas.

Excel includes more than 450 predefined worksheet functions. With so many from which to choose, you may be curious as to why anyone would need to develop additional functions. The main reason is that creating a custom function can greatly simplify your formulas by making them shorter, and shorter formulas are more readable and easier to work with. For example, you can often replace a complex formula with a single function. Another reason is that you can write functions to perform operations that would otherwise be impossible.

NOTE

This chapter assumes that you're familiar with entering and editing VBA code in the Visual Basic Editor (VBE).



See Chapter 39, “Introducing Visual Basic for Applications,” for an overview of the VBE.

Seeing a Simple Example

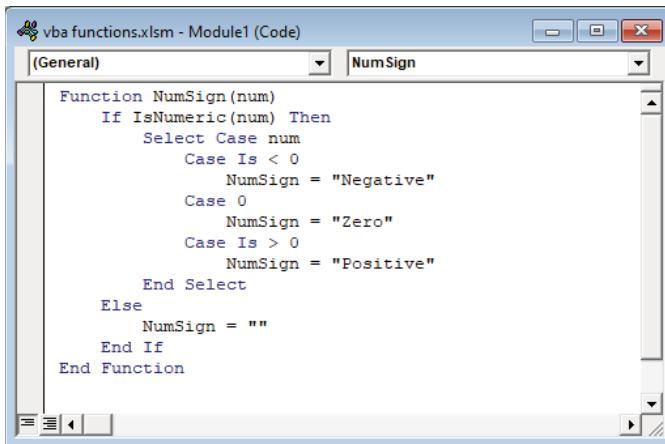
Creating custom functions is relatively easy after you understand VBA. Without further ado, here’s an example of a VBA function procedure. This function is stored in a VBA module, which is accessible from the VBE.

Creating a custom function

This example function, named NumSign, uses one argument. The function returns a text string of Positive if its argument is greater than zero, Negative if the argument is less than zero, and Zero if the argument is equal to zero. If the argument is nonnumeric, the function returns an empty string. Figure 40.1 shows the NumSign function.

FIGURE 40.1

A simple custom worksheet function



The screenshot shows the Microsoft VBE interface with the title bar "vba functions.xlsm - Module1 (Code)". The code editor window is open, showing the following VBA code:

```
Function NumSign(num)
    If IsNumeric(num) Then
        Select Case num
            Case Is < 0
                NumSign = "Negative"
            Case 0
                NumSign = "Zero"
            Case Is > 0
                NumSign = "Positive"
        End Select
    Else
        NumSign = ""
    End If
End Function
```

You can, of course, accomplish the same effect with the following worksheet formula, which uses nested IF functions:

```
=IF(ISNUMBER(A1), IF(A1=0, "Zero", IF(A1>0, "Positive", "Negative")) , "")
```

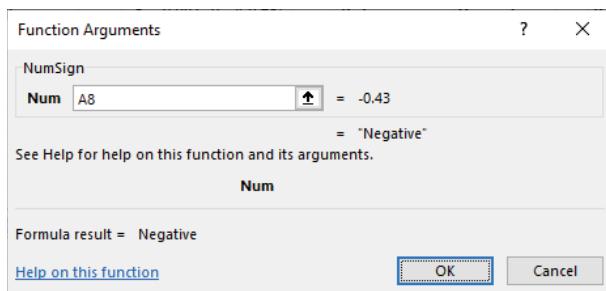
Many would agree that the custom function solution is easier to understand and to edit than the worksheet formula.

Using the function in a worksheet

When you enter a formula that uses the NumSign function, Excel executes the function to get the result. This custom function works just like any built-in worksheet function. You can insert it in a formula by choosing Formulas \Rightarrow Function Library \Rightarrow Insert Function, which displays the Insert Function dialog box. (Custom functions are listed in the User Defined category.) When you select the function from the list, you can then use the Function Arguments dialog box to specify the arguments for the function, as shown in Figure 40.2. You can also nest custom functions and combine them with other elements in your formulas.

FIGURE 40.2

Creating a worksheet formula that uses a custom function



Analyzing the custom function

This section describes the NumSign function. Here again is the code:

```
Function NumSign(num)
    If IsNumeric(num) Then
        Select Case num
            Case Is < 0
                NumSign = "Negative"
            Case 0
                NumSign = "Zero"
            Case Is > 0
                NumSign = "Positive"
        End Select
    Else
        NumSign = ""
    End If
End Function
```

Notice that the procedure starts with the keyword Function, followed by the name of the function (NumSign). This custom function uses one argument (num), and the argument's name

is enclosed in parentheses. The num argument represents the cell or value that is to be processed. When the function is used in a worksheet, the argument can be a cell reference (such as A1) or a literal value (such as -123). When the function is used in another procedure, the argument can be a numeric variable, a literal number, or a value that is obtained from a cell.

The first statement inside the function is an If statement. It starts what is known as an If block. An If block consists of an If statement, an End If statement, one or more optional Else If statements, and one optional Else statement. The previous code is indented in a way that makes it obvious that the Else and End If statements near the bottom of the function belong to the If statement near the top of the procedure. Indenting is optional, but you'll find your code is much easier to read if you do it.

The If statement contains the built-in function IsNumeric that returns a True if the argument is a number and False if it's not. Whenever a built-in function begins with Is or Has, it returns True or False (a Boolean value).

The NumSign function uses the Select Case construct (described in Chapter 39) to take a different action, depending on the value of num. If num is less than zero, NumSign is assigned the text Negative. If num is equal to zero, NumSign is Zero. If num is greater than zero, NumSign is Positive. The value returned by a function is always assigned to the function's name.

There is often more than one way to accomplish the same goal in VBA. Instead of using a Select Case construct, you could use an If block. The following code returns the same result as the original function but uses another If block with an Else If statement. Note how the indented code makes it easy to see which statements belong to which If blocks:

```
Function NumSignIfBlock(num)
    If IsNumeric(num) Then
        If num = 0 Then
            NumSign = "Zero"
        ElseIf num > 0 Then
            NumSign = "Positive"
        Else
            NumSign = "Negative"
        End If
    Else
        NumSign = ""
    End If
End Function
```

Learning about Function Procedures

A custom function has much in common with a Sub procedure. Function procedures have some important differences, however. Perhaps the key difference is that a function returns a value (such as a number, a date, or a text string). The value that's returned by the function is the value that has been assigned to the function's name when the function is finished executing.

To create a custom function, follow these steps:

1. **Activate the VB Editor.** (Press Alt+F11.)
2. **Select the workbook in the Project Explorer (choose View \Rightarrow Project Explorer if it's not already visible).**
3. **Choose Insert \Rightarrow Module to insert a VBA module, or you can use an existing code module.** However, it must be a standard VBA module.
4. **Enter the keyword Function followed by the function's name and a list of the arguments (if any) in parentheses.** If the function doesn't use an argument, the VBE adds a set of empty parentheses.
5. **Type the VBA code that performs the work—and make sure that the variable corresponding to the function's name has the appropriate value when the function ends.** This is the value that the function returns.
6. **End the function with an End Function statement.** The VBE adds this statement automatically when you type the function statement.

Note

Step 3 is important. If you put a function procedure in a code module for ThisWorkbook or a worksheet (for example, Sheet1), the function will not be recognized in a worksheet formula. Excel will display a #NAME? error. Putting a function procedure in the wrong type of code module is a common mistake.

Function names that are used in worksheet formulas must adhere to the same rules as variable names.

What a Function Can't Do

Almost everyone who starts creating custom worksheet functions using VBA makes a fatal mistake: they try to get the function to do more than is possible.

A worksheet function returns a value, and the function must be completely "passive." In other words, the function can't change anything on the worksheet. For example, you can't develop a worksheet function that changes the formatting of a cell. (Every VBA programmer has tried, and not one of them has been successful!) If your function attempts to perform an action that isn't allowed, the function simply returns an error.

Nevertheless, the preceding paragraph isn't absolutely true. There are a few cases in which a VBA function used in a formula can have an effect. For example, it's possible to create a custom worksheet function that adds or deletes cell comments. For the most part, however, functions used in formulas must be passive.

VBA functions that aren't used in worksheet formulas can do anything that a regular Sub procedure can do, including changing cell formatting.

Executing Function Procedures

You can execute a Sub procedure in many ways, but you can execute a function procedure in just two ways:

- Call it from another VBA procedure.
- Use it in a worksheet formula.

Calling custom functions from a procedure

You can call custom functions from a VBA procedure just as you call built-in VBA functions. For example, after you define a function called CalcTax, you can enter a statement such as the following:

```
Tax = CalcTax(Amount, Rate)
```

This statement executes the CalcTax custom function with Amount and Rate as its arguments. The function's result is assigned to the Tax variable.

Using custom functions in a worksheet formula

You use a custom function in a worksheet formula just as you use built-in functions. However, you must ensure that Excel can locate the function. If the function procedure is in the same workbook, you don't have to do anything special. If the function is defined in a different workbook, you may have to tell Excel where to find the function. The following are the three ways in which you can do this:

Precede the function's name with a file Reference For example, if you want to use a function called CountNames that's defined in a workbook named MyFunctions.xlsxm, you can use a reference such as the following:

```
=MyFunctions.xlsxm!CountNames(A1:A1000)
```

If the workbook name contains a space, you need to add single quotes around the workbook name. Here's an example:

```
= 'My Functions.xlsxm' !CountNames(A1:A1000)
```

If you insert the function with the Insert Function dialog box, the workbook reference is inserted automatically.

Set up a reference to the workbook If the custom function is defined in a referenced workbook, you don't need to precede the function name with the workbook name. You establish a reference to another workbook by choosing Tools ➔ References (in the VB Editor). You're presented with a list of references that includes all open workbooks. Place a check mark in the item that refers to the workbook containing the custom function. (Click the Browse button if the workbook isn't open.)

Create an add-in When you create an add-in from a workbook that has function procedures, you don't need to use the file reference when you use one of the functions in a formula; the add-in must be installed, however.



We discuss add-ins in Chapter 45, “Creating Custom Excel Add-Ins.”

Note

Function procedures don’t appear in the Macro dialog box because you can’t execute a function directly. As a result, you need to do extra, up-front work to test your functions while you’re developing them. One approach is to set up a simple Sub procedure that calls the function. If the function is designed to be used in worksheet formulas, you can enter a simple formula that uses the function to test it while you’re developing the function.

Using Function Procedure Arguments

Keep in mind the following about function procedure arguments:

- Arguments can be variables (including arrays), constants, literals, or expressions.
- Some functions do not have arguments.
- Some functions have a fixed number of required arguments (from 1 to 60).
- Some functions have a combination of required and optional arguments.

The following sections present a series of examples that demonstrate how to use arguments effectively with functions. Coverage of optional arguments is beyond the scope of this book.



The examples in this chapter are available on this book’s website at www.wiley.com/go/excel365bible. The file is named vba_functions.xlsm.

Creating a function with no arguments

Most functions use arguments, but that’s not a requirement. Excel, for example, has a few built-in worksheet functions that don’t use arguments, such as RAND, TODAY, and NOW.

The following is a simple example of a function that has no arguments. This function returns the UserName property of the Application object, which is the name that appears in the Personalize section on the General tab of the Excel Options dialog box. This function is simple, but it can be useful because there’s no built-in function that returns the user’s name:

```
Function User()
    ' Returns the name of the current user
    User = Application.UserName
End Function
```

When you enter the following formula into a worksheet cell, the cell displays the name of the current user:

```
=User()
```

As with Excel's built-in functions, when you use a function with no arguments, you must include a set of empty parentheses.

Creating a function with one argument

The function that follows takes a single argument and strips out any characters that aren't numbers:

```
Function NumbersOnly(txt)
    Dim i As Long

    For i = 1 To Len(txt)
        If IsNumeric(Mid(txt, i, 1)) Then
            NumbersOnly = NumbersOnly & Mid(txt, i, 1)
        End If
    Next i
End Function
```

One application of this function is to strip out alphabetic characters from invoice numbers. The formula that follows returns 759426 because those are the only numeric characters in the text that was passed to it.

```
=NumbersOnly("INV759426")
```

Another use is to isolate numbers from a sentence that you get externally. For example, if you copied the sentence "Net income for the quarter was \$15,267" from a web page and pasted it into cell A1, you could use the following formula to extract the numbers and use them in other calculations:

```
=NumbersOnly(A1)
```

Creating another function with one argument

This section contains a more complex function that is designed for a sales manager who needs to calculate the commissions earned by the sales force. The commission rate is based on the amount sold—those who sell more earn a higher commission rate. The function returns the commission amount, based on the sales made (which is the function's only argument—a required argument). The calculations in this example are based on the following table:

Monthly Sales	Commission Rate
0–\$9,999	8.0%
\$10,000–\$19,999	10.5%
\$20,000–\$39,999	12.0%
\$40,000+	14.0%

You can use any of several different methods to calculate commissions for various sales amounts that are entered into a worksheet. You could write a formula such as the following:

```
=IF(AND(A1>=0,A1<=9999.99),A1*0.08,IF(AND(A1>=10000,
A1<=19999.99),A1*0.105,IF(AND(A1>=20000,
A1<=39999.99),A1*0.12,IF(A1>=40000,A1*0.14,0))))
```

This approach isn't the best for a couple of reasons. First, the formula is overly complex and difficult to understand. Second, the values are hard-coded into the formula, making the formula difficult to modify if the commission structure changes.

A better solution is to use a lookup table function to compute the commissions; here's an example:

```
=VLOOKUP(A1,Table,2)*A1
```

Using the VLOOKUP function requires that you have a table of commission rates set up in your worksheet.

Another option is to create a custom function, such as the following:

```
Function Commission(Sales)
    ' Calculates sales commissions
    Tier1 = 0.08
    Tier2 = 0.105
    Tier3 = 0.12
    Tier4 = 0.14
    Select Case Sales
        Case 0 To 9999.99
            Commission = Sales * Tier1
        Case 10000 To 19999.99
            Commission = Sales * Tier2
        Case 20000 To 39999.99
            Commission = Sales * Tier3
        Case Is >= 40000
            Commission = Sales * Tier4
    End Select
End Function
```

After you define the Commission function in a VBA module, you can use it in a worksheet formula. Entering the following formula into a cell produces a result of 3,000. (The amount of 25,000 qualifies for a commission rate of 12%.)

```
=Commission(25000)
```

If the sales amount is in cell D23, the function's argument would be a cell reference, like this:

```
=Commission(D23)
```

Creating a function with two arguments

This example builds on the previous one. Imagine that the sales manager implements a new policy: the total commission paid is increased by 1 percent for every year that the salesperson has been with the company. For this example, the custom Commission function (defined in

the preceding section) has been modified so that it takes two arguments, both of which are required arguments. Call this new function Commission2:

```
Function Commission2(Sales, Years)
    ' Calculates sales commissions based on years in service
    Tier1 = 0.08
    Tier2 = 0.105
    Tier3 = 0.12
    Tier4 = 0.14
    Select Case Sales
        Case 0 To 9999.99
            Commission2 = Sales * Tier1
        Case 10000 To 19999.99
            Commission2 = Sales * Tier2
        Case 20000 To 39999.99
            Commission2 = Sales * Tier3
        Case Is >= 40000
            Commission2 = Sales * Tier4
    End Select
    Commission2 = Commission2 + (Commission2 * Years / 100)
End Function
```

The modification was quite simple. The second argument (Years) was added to the Function statement, and an additional computation was included that adjusts the commission before exiting the function.

The following is an example of how you write a formula using this function. It assumes that the sales amount is in cell A1 and that the number of years that the salesperson has worked is in cell B1:

```
=Commission2(A1,B1)
```

Creating a function with a range argument

The example in this section demonstrates how to use a worksheet range as an argument. It's not at all tricky; Excel takes care of the details behind the scenes.

Assume that you want to calculate the average of the five largest values in a range named Data. Excel doesn't have a function that can do this calculation, so you can write the following formula:

```
= (LARGE(Data,1)+LARGE(Data,2)+LARGE(Data,3) +
    LARGE(Data,4)+LARGE(Data,5)) / 5
```

This formula uses Excel's LARGE function, which returns the nth largest value in a range. The preceding formula adds the five largest values in the range named Data and then divides the result by 5. The formula works fine, but it's rather unwieldy. Plus, what if you need to compute the average of the top six values? You'd need to rewrite the formula and make sure that all copies of the formula also get updated.

Wouldn't it be easier if Excel had a function named TopAvg? For example, you could use the following (nonexistent) function to compute the average:

```
=TopAvg(Data, 5)
```

This situation is an example of when a custom function can make things much easier for you. The following is a custom VBA function, named TopAvg, which returns the average of the top n values in a range:

```
Function TopAvg(Data, Num)
    ' Returns the average of the highest Num values in Data
    Sum = 0
    For i = 1 To Num
        Sum = Sum + WorksheetFunction.Large(Data, i)
    Next i
    TopAvg = Sum / Num
End Function
```

This function takes two arguments: Data (which represents a range in a worksheet) and Num (the number of values to average). The code starts by initializing the Sum variable to 0. It then uses a For-Next loop to calculate the sum of the nth largest values in the range. (Note that Excel's LARGE function is used within the loop.) You can use an Excel worksheet function in VBA if you precede the function with WorksheetFunction and a period. Finally, TopAvg is assigned the value of Sum divided by Num.

You can use all Excel worksheet functions in your VBA procedures except those that have equivalents in VBA. For example, VBA has a Rnd function that returns a random number. Therefore, you can't use Excel's RAND function in a VBA procedure.

Creating a simple but useful function

Useful functions don't have to be complicated. The function in this section is essentially a wrapper for a built-in VBA function called Split. The Split function makes it easy to extract an element in a delimited string. The function is named ExtractElement:

```
Function ExtractElement(Txt, n, Separator)
    ' Returns the nth element of a text string, where the
    ' elements are separated by a specified separator character
    ExtractElement = Split(Application.Trim(Txt), Separator)(n - 1)
End Function
```

The function takes three arguments:

- Txt:** A delimited text string, or a reference to a cell that contains a delimited text string
- n:** The element number within the string
- Separator:** A single character that represents the separator

Here's a formula that uses the ExtractElement function:

```
=EXTRACTELEMENT("123-45-678", 2, "-")
```

The formula returns 45, the second element in the string that's delimited by hyphens.

The delimiter can also be a space character. Here's a formula that extracts the first name from the name in cell A1:

```
=EXTRACTELEMENT(A1,1," ")
```

Debugging Custom Functions

Debugging a function can be a bit more challenging than debugging a Sub procedure. If you develop a function to use in worksheet formulas, an error in the function simply results in an error display in the formula cell (usually #VALUE!). In other words, you don't receive the normal runtime error message that helps you locate the offending statement.

When you're debugging a worksheet formula, using only one instance of the function in your worksheet is the best technique. The following are three methods you may want to use in your debugging:

- Place MsgBox functions at strategic locations to monitor the value of specific variables. Fortunately, message boxes in function procedures pop up when the procedure is executed. But make sure you have only one formula in the worksheet that uses your function; otherwise, the message boxes appear for each formula that's evaluated.
- Test the procedure by calling it from a Sub procedure. Runtime errors display normally, and you can either fix the problem (if you know what it is) or jump right into the debugger.
- Set a breakpoint in the function, and then use the Excel debugger to step through the function. Press F9, and the statement at the cursor becomes a breakpoint. The code will stop executing, and you can step through the code line by line (by pressing F8). Consult the Help system for more information about using VBA debugging tools.

Inserting Custom Functions

The Excel Insert Function dialog box makes it easy to identify a function and insert it into a formula. This dialog box also displays custom functions written in VBA. After you select a function, the Function Arguments dialog box prompts you for the function's arguments.

Note

Function procedures that are defined with the Private keyword do not appear in the Insert Function dialog box. If you create a function that will be used only by other VBA procedures, you should declare the function by using the Private keyword.

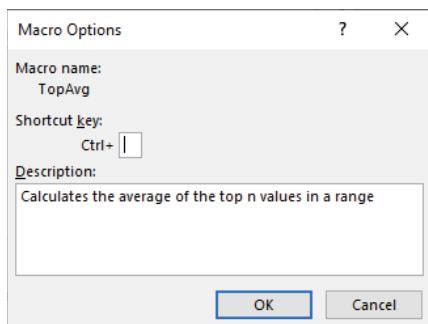
You also can display a description of your custom function in the Insert Function dialog box. To do so, follow these steps:

1. **Create the function in a module by using the VBE.**
2. **Activate Excel.**

3. Choose Developer ➔ Code ➔ Macros. The Macro dialog box appears.
4. Type the name of the function in the Macro Name field. Notice that functions don't appear in this dialog box, so you must enter the function name yourself.
5. Click the Options button. The Macro Options dialog box appears, as shown in Figure 40.3.
6. Enter a description of the function and then click OK. The Shortcut key field is irrelevant for functions.

FIGURE 40.3

Entering a description for a custom function. This description appears in the Insert Function dialog box.



The description that you enter appears in the Insert Function dialog box.

Another way to provide a description for a custom function is to execute a VBA statement that uses the MacroOptions method. The MacroOptions method also lets you assign your function to a specific category and even provide a description of the arguments. The argument descriptions display in the Function Arguments dialog box, which appears after you select the function in the Insert Function dialog box.

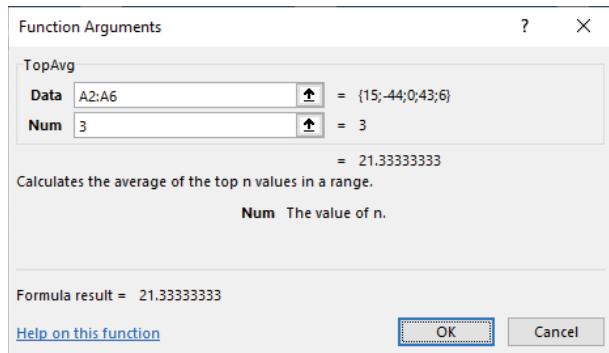
Figure 40.4 shows the Function Arguments dialog box, which prompts the user to enter arguments for a custom function (TopAvg). This function appears in function category 3 (Math and Trig). We've added the description, category, and argument descriptions by executing this Sub procedure:

```
Sub CreateArgDescriptions()
    Application.MacroOptions Macro:="TopAvg", _
        Description:= _
        "Calculates the average of the top n values in a range", _
        Category:=3, _
        ArgumentDescriptions:= _
        Array("The range that contains the data", "The value of n")
End Sub
```

The category numbers are listed in the VBA Help system. You execute this procedure only one time. After you execute it, the description, category, and argument descriptions are stored in the file.

FIGURE 40.4

Using the Function Arguments dialog box to insert a custom function



Learning More

The information in this chapter only scratches the surface when it comes to creating custom functions. It should be enough to get you started, however, if you're indeed interested in this topic.



See Chapter 44, “Seeing Some VBA Examples,” for more examples of useful VBA functions. You may be able to use the examples directly or adapt them for your needs.

Creating UserForms

IN THIS CHAPTER

Understanding why you may want to create UserForms

Identifying UserForm alternatives

Creating UserForms: An overview

Looking at some UserForm examples

Finding more information on creating UserForms

You can't use Excel very long without being exposed to dialog boxes. Excel, like most Windows programs, uses dialog boxes to obtain information, clarify commands, and display messages. If you develop VBA macros, you can create your own dialog boxes that work very much like those that are built into Excel. These dialog boxes are known as UserForms.

Understanding Why to Create UserForms

Some macros that you create behave the same every time that you execute them. For example, you may develop a macro that enters a list of your sales regions into a worksheet range. This macro always produces the same result and requires no additional user input. You may develop other macros, however, that perform differently under different circumstances or that offer options for the user. In such cases, the macro may benefit from a custom dialog box.

The following is an example of a simple macro that makes each cell in the selected range uppercase (but skips cells that have a formula). The procedure uses VBA's built-in StrConv function:

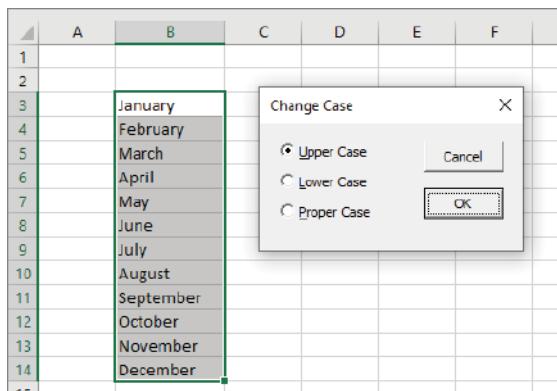
```
Sub ChangeCase()
    For Each cell In Selection
        If Not cell.HasFormula Then
            cell.Value = StrConv(cell.Value, vbUpperCase)
        End If
    Next Cell
End Sub
```

This macro is useful, but it can be improved. For example, the macro would be more helpful if it could also change the cells to lowercase or proper case (only the first letter of each word is uppercase).

This modification is not difficult to make, but if you make this change to the macro, you need some method of asking the user what type of change to make to the cells. The solution is to present a dialog box like the one shown in Figure 41.1. This dialog box is a UserForm that was created by using the Visual Basic Editor (VBE), and it's displayed by a VBA macro.

FIGURE 41.1

A UserForm that asks the user to select an option



Another solution is to develop three macros—one for each type of text case change. Combining these three operations into a single macro and using a UserForm is a more efficient approach, however. We discuss this example, including how to create the UserForm, in “Looking at Another UserForm Example” later in the chapter.

Exploring UserForm Alternatives

After you get the hang of it, developing UserForms isn't difficult. But sometimes using the tools that are built into VBA is easier. For example, VBA includes two functions (`InputBox` and `MsgBox`) that enable you to display simple dialog boxes without having to create a UserForm in the VBE. You can customize these dialog boxes in some ways, but they certainly don't offer the number of options that are available in a UserForm.

Using the InputBox function

The `InputBox` function is useful for obtaining a single input from the user. A simplified version of the function's syntax follows:

```
InputBox(prompt [,title] [,default] )
```

The elements are defined as follows:

prompt (Required): Text that is displayed in the input box

title (Optional): Text that appears in the input box's title bar

default (Optional): The value that appears in the input box's text area by default

The following is an example of how you can use the InputBox function:

```
CName = InputBox("Customer name?", "Customer Data")
```

When this VBA statement is executed, Excel displays the dialog box shown in Figure 41.2. Notice that this example uses only the first two arguments for the InputBox function and does not supply a default value. When the user enters a value and clicks OK, the value is assigned to the variable CName. Your VBA code can then use that variable.

FIGURE 41.2

This dialog box is displayed by the VBA InputBox function.



Using the MsgBox function

The VBA MsgBox function is a handy way to display information and to solicit simple input from users. We use the VBA MsgBox function in many of this book's examples to display a variable's value. A simplified version of the MsgBox syntax is as follows:

```
MsgBox(prompt [, buttons] [, title])
```

The elements are defined as follows:

prompt (Required): Text that is displayed in the message box

buttons (Optional): The code for the buttons that appear in the message box

title (Optional): Text that appears in the message box's title bar

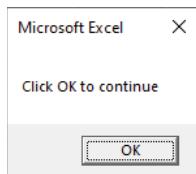
You can use the MsgBox function by itself or assign its result to a variable. If you use it by itself, don't include parentheses around the arguments. The following example displays a message and does not return a result:

```
Sub MsgBoxDemo()
    MsgBox "Click OK to continue"
End Sub
```

Figure 41.3 shows how this message box appears.

FIGURE 41.3

A simple message box, displayed with the VBA MsgBox function



To get a response from a message box, you can assign the result of the MsgBox function to a variable. The following code uses some built-in constants (described in Table 41.1) to make it easier to work with the values that are returned by MsgBox:

```
Sub GetAnswer()
    Ans = MsgBox("Continue?", vbYesNo)
    Select Case Ans
        Case vbYes
            '      ... [code if Ans is Yes] ...
        Case vbNo
            '      ... [code if Ans is No] ...
    End Select
End Sub
```

When this procedure is executed, the Ans variable contains a value that corresponds to vbYes or vbNo. The Select Case statement determines the action to take based on the value of Ans.

You can easily customize your message boxes because of the flexibility of the buttons argument. Table 41.1 lists the most common built-in constants that you can use for the buttons argument. You can specify which buttons to display, whether an icon appears, and which button is the default.

TABLE 41.1 Constants Used in the MsgBox Function

Constant	Value	Description
vbOKOnly	0	Displays OK button
vbOKCancel	1	Displays OK and Cancel buttons
vbAbortRetryIgnore	2	Displays Abort, Retry, and Ignore buttons
vbYesNoCancel	3	Displays Yes, No, and Cancel buttons
vbYesNo	4	Displays Yes and No buttons
vbRetryCancel	5	Displays Retry and Cancel buttons
vbCritical	16	Displays Critical Message icon

Constant	Value	Description
vbQuestion	32	Displays Query icon (a question mark)
VBExclamation	48	Displays Warning Message icon
vbInformation	64	Displays Information Message icon
vbDefaultButton1	0	First button is default
vbDefaultButton2	256	Second button is default
vbDefaultButton3	512	Third button is default

The following example uses a combination of constants to display a message box with a Yes button and a No button (vbYesNo) and a question mark icon (vbQuestion). The second button (the No button) is designated as the default button (vbDefaultButton2), which is the one that is executed if the user presses Enter. For simplicity, these constants are assigned to the Config variable, and Config is then used as the second argument in the MsgBox function.

```
Sub GetAnswer()
    Config = vbYesNo + vbQuestion + vbDefaultButton2
    Ans = MsgBox("Process the monthly report?", Config)
    If Ans = vbYes Then RunReport
    If Ans = vbNo Then Exit Sub
End Sub
```

Figure 41.4 shows how this message box appears when the GetAnswer procedure is executed. If the user clicks the Yes button, the routine executes the procedure named RunReport (which is not shown). If the user clicks the No button (or presses Enter), the procedure is ended with no action. Because the title argument was omitted in the MsgBox function, Excel uses the default title (Microsoft Excel).

FIGURE 41.4

The second argument of the MsgBox function determines what appears in the message box.



The procedure that follows is another example of using the MsgBox function:

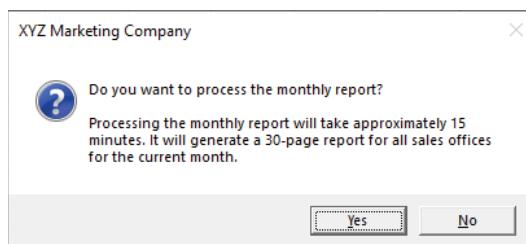
```
Sub GetAnswer2()
    Msg = "Do you want to process the monthly report?"
    Msg = Msg & vbCrLf & vbCrLf
    Msg = Msg & "Processing the monthly report will take approximately "
```

```
Msg = Msg & "15 minutes. It will generate a 30-page report for all "
Msg = Msg & "sales offices for the current month."
Title = "XYZ Marketing Company"
Config = vbYesNo + vbQuestion
Ans = MsgBox(Msg, Config, Title)
If Ans = vbYes Then RunReport
If Ans = vbNo Then Exit Sub
End Sub
```

This example demonstrates an efficient way to specify a longer message in a message box. A variable (Msg) and the concatenation operator (&) are used to build the message in a series of statements. vbNewLine is a constant that represents a line break character. (Using two line breaks will insert a blank line.) The title argument is also used to display a different title in the message box. The Config variable stores the constants that generate Yes and No buttons and a question mark icon. Figure 41.5 shows how this message box appears when the procedure is executed.

FIGURE 41.5

A message box with a longer message and a title



Creating UserForms: An Overview

The InputBox and MsgBox functions are adequate for many situations, but if you need to obtain more information, you need to create a UserForm.

The following is a list of the general steps that you typically take to create a UserForm:

1. **Determine how the dialog box is going to be used and where it will fit into your VBA macro.**
2. **Activate the VBE and insert a new UserForm.**
3. **Add the appropriate controls to the UserForm.**
4. **Create a VBA macro to display the UserForm.** This macro goes in a normal VBA module.
5. **Create event handler VBA procedures that are executed when the user manipulates the controls (for example, when the user clicks the OK button).** These procedures go in the code module for the UserForm.

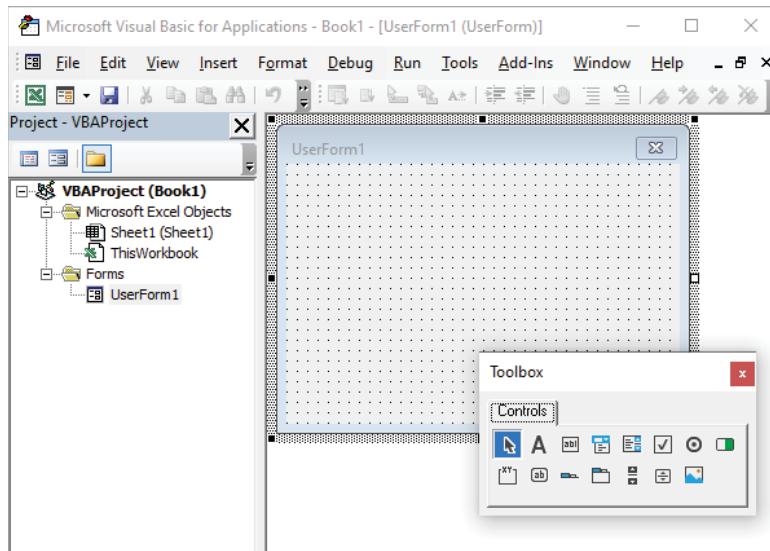
The following sections offer more details on creating a UserForm.

Working with UserForms

To create a dialog box, you must first insert a new UserForm in the VBE. To activate the VBE, choose Developer \Rightarrow Code \Rightarrow Visual Basic (or press Alt+F11). Make sure that the correct workbook is selected in the Project window and then choose Insert \Rightarrow UserForm. The VBE displays an empty UserForm, as shown in Figure 41.6. When you activate a UserForm, the VBE displays the Toolbox, which is used to add controls to the UserForm.

FIGURE 41.6

An empty UserForm



Adding controls

The Toolbox, also shown in Figure 41.6, contains ActiveX controls that you can add to your UserForm. If the Toolbox is not visible, choose View \Rightarrow Toolbox.

When you move the mouse pointer over a control in the Toolbox, the control's name appears. To add a control, select the control in Toolbox and either click in the form to get the default size or drag the desired size onto the form. After adding a control, you can move it or change its size.

Table 41.2 lists the Toolbox controls.



You can also place some of these controls directly on your worksheet. See Chapter 42, "Using UserForm Controls in a Worksheet," for details.

TABLE 41.2 Toolbox Controls

Control	Description
Select Objects	Lets you select other controls by dragging
Label	Adds a label (a container for text)
TextBox	Adds a text box, which allows the user to type text
ComboBox	Adds a combo box (a drop-down list)
ListBox	Adds a list box, which allows the user to select an item from a list
CheckBox	Adds a check box to control Boolean options
OptionButton	Adds an option button to allow the user to select from multiple options
ToggleButton	Adds a toggle button to control Boolean options
Frame	Adds a frame (a container for other objects)
CommandButton	Adds a command button (a clickable button)
TabStrip	Adds a tab strip (a container for other objects)
MultiPage	Adds a multipage control (a container for other objects)
ScrollBar	Adds a scrollbar, which allows the user to specify a value by dragging a bar
SpinButton	Adds a spin button, which allows the user to specify a value by clicking up or down arrows
Image	Adds a control that can contain an image

Changing the properties of a control

Every control that you add to a UserForm has several properties that determine the way the control looks and behaves. You can change some of these properties (such as Height and Width) by clicking and dragging the control's border. To change other properties, use the Properties window.

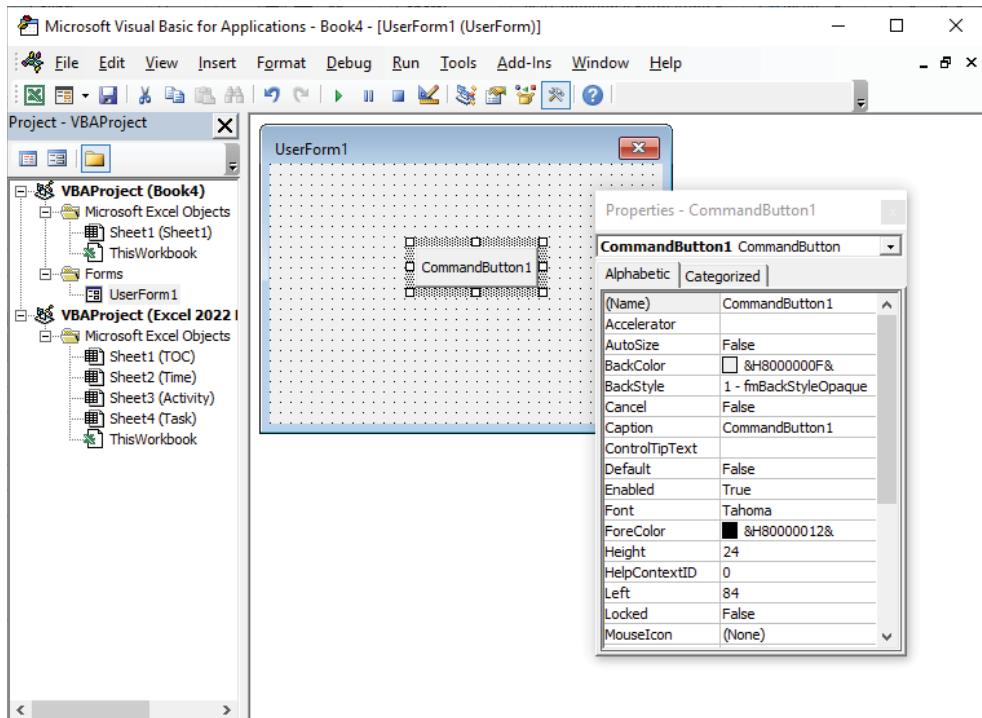
To display the Properties window, choose View \leftrightarrow Properties Window (or press F4). The Properties window displays a list of properties for the selected control. (Each type of control has a different set of properties.) If you click the UserForm itself, the Properties window displays properties for the form. Figure 41.7 shows the Properties window for a CommandButton control.

To change a property, select the property in the Properties window and then enter a new value. Some properties (such as BackColor) let you select a value from a list. The top of the Properties window contains a drop-down list of all the controls on the form. You can also click a control to select it and display its properties.

When you set properties via the Properties window, you're setting properties at design time. You can also use VBA to change the properties of controls while the UserForm is displayed (that is, at runtime).

FIGURE 41.7

The Properties window for a CommandButton control



A complete discussion of all the properties is well beyond the scope of this book—and it would indeed be dull reading. To find out about a particular property, select it in the Properties window and press F1 for help.

Handling events

When you insert a UserForm, that form can also hold VBA Sub procedures to handle the events that are generated by the UserForm. An event is something that occurs when the user manipulates a control. For example, clicking a button causes a Click event. Selecting an item in a list box control also triggers a Click event, as well as a Change event. To make a UserForm useful, you must write VBA code to do something when an event occurs.

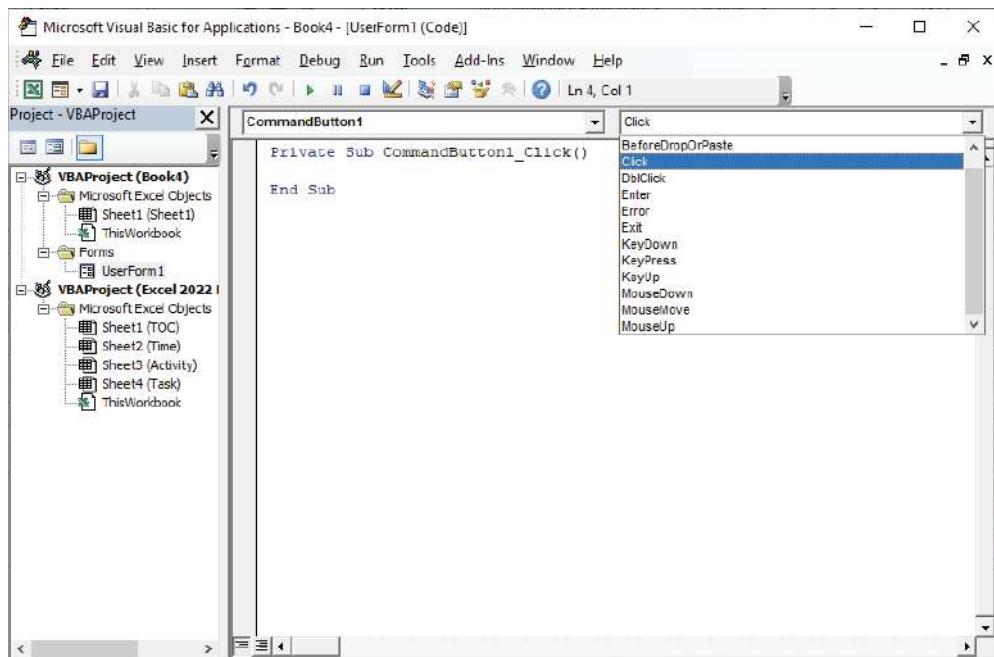
Event handler procedures have names that combine the control with the event. The general form is the control's name, followed by an underscore and then the event name. For example, the procedure that is executed when the user clicks a button named MyButton is MyButton_Click. You don't have to remember how they're named, however. Just right-click the control and choose View Code. The Private Sub and End Sub keywords will be inserted for you, and the

correct name of one of the control's events will already be constructed. Use the drop-downs at the top of the code pane to change the event from the default.

Figure 41.8 shows the result of right-clicking on a CommandButton control named CommandButton1 and choosing View Code. The CommandButton's default event, the Click event, is inserted into the code module. The figure also shows some of the other events available using the code pane's drop-down.

FIGURE 41.8

The code pane showing the Click event of a CommandButton and the drop-down listing other available events



Displaying a UserForm

You also need to write a procedure to display the UserForm. You use the Show method of the UserForm object. The following procedure displays the UserForm named UserForm1:

```
Sub ShowDialog()
    UserForm1.Show
End Sub
```

This procedure should be stored in a regular VBA module (not the code module for the UserForm). If your VB project doesn't have a regular VBA module, choose Insert → Module to add one.

When the ShowDialog procedure is executed, the UserForm is displayed. What happens next depends on the event handler procedures that you create.

Looking at a UserForm Example

The preceding sections are, admittedly, rudimentary. The following sections demonstrate, in detail, how to develop a UserForm. This example is rather simple. The UserForm displays a message to the user—something that can be accomplished more easily by using the MsgBox function. However, a UserForm gives you more flexibility in terms of formatting and layout of the message.



This workbook is available on this book's website at www.wiley.com/go/excel365bible. The file is named `show message.xlsm`.

Creating the UserForm

If you're following along on your computer, start with a new workbook. Then follow these steps:

1. Choose **Developer** \Rightarrow **Code** \Rightarrow **Visual Basic** (or press **Alt+F11**). The VBE appears.
2. Click your workbook's name in the **Project Explorer** to activate it.
3. Choose **Insert** \Rightarrow **UserForm**. The VBE adds an empty form named **UserForm1** and displays the **Toolbox**.
4. Press **F4** to display the **Properties** window, and then change the following properties of the **UserForm** object:

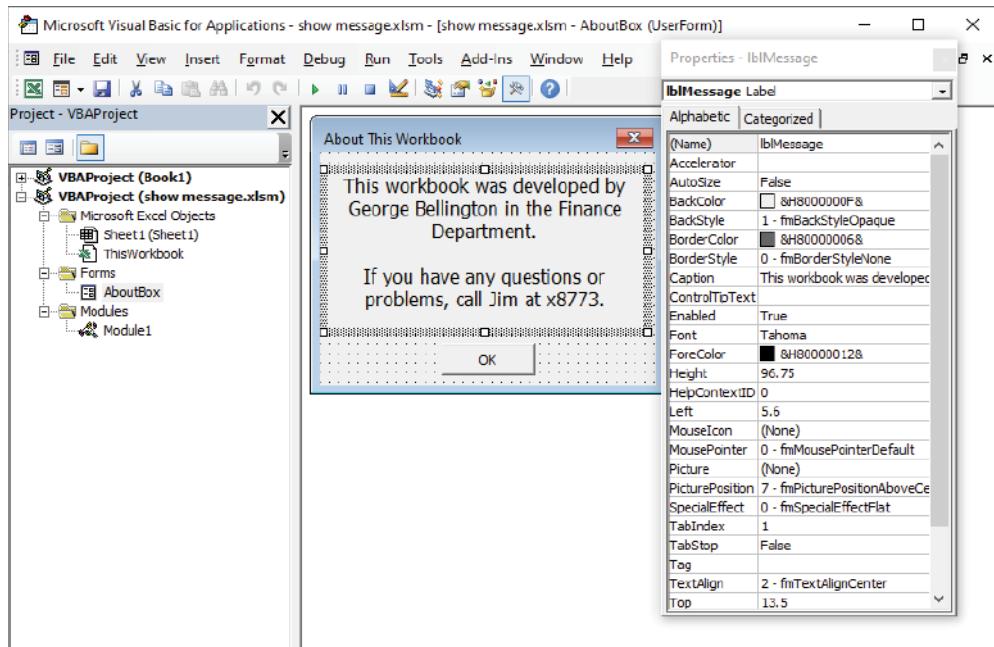
Property	Change to
Name	AboutBox
Caption	About This Workbook

5. Use the **Toolbox** to add a **Label** object to the **UserForm**. If the **Toolbox** is not visible, choose **View** \Rightarrow **Toolbox**.
6. Select the **Label** object, and in the **Properties** window change the **Name** property to **lblMessage** and enter any text you want in the **Caption** property.
7. In the **Properties** window, click the **Font** property and adjust the font. You can change the typeface, size, and so on. The changes then appear in the form. Figure 41.9 shows an example of a formatted Label control. In this example, the **TextAlign** property was set to the code that center aligns the text.

2 - `fmTextAlignCenter`

FIGURE 41.9

A Label control, after changing its Font properties



8. Use the Toolbox and add a CommandButton object to the UserForm. Then use the Properties window to change the following properties for the CommandButton:

Property	Change to
Name	OKButton
Caption	OK
Default	True

9. Make other adjustments so that the form looks good to you. You can change the size of the form or move or resize the controls.

Testing the UserForm

At this point, the UserForm has all the necessary controls. What's missing is a way to display the UserForm. While you're developing the UserForm, you can press F5 to display it and see how it looks. To close the UserForm, click the Close button (X) in the dialog box title bar.

This section explains how to write a VBA procedure to display the UserForm when Excel is active:

- 1. Insert a VBA module by choosing Insert \Rightarrow Module.**

- 2. In the empty module, enter the following code:**

```
Sub ShowAboutBox()
    AboutBox.Show
End Sub
```

- 3. Press Alt+F11 to activate Excel.**

- 4. Choose Developer \Rightarrow Code \Rightarrow Macros (or press Alt+F8).** The Macro dialog box appears.

- 5. Select ShowAboutBox from the list of macros, and then click Run.** The UserForm appears.

If you click the OK button, notice that it doesn't close the UserForm as you may expect. This button needs to have an event handler procedure for it to do anything when it's clicked. To dismiss the UserForm, click the Close button (X) in its title bar.



You may prefer to display the UserForm by clicking a CommandButton on your worksheet. See Chapter 42 for details on attaching a macro to a worksheet CommandButton.

Creating an event handler procedure

An event handler procedure is executed when an event occurs. In this case, you need a procedure to handle the Click event that's generated when the user clicks the OK button.

- 1. Press Alt+F11 to activate the VBE.**

- 2. Activate the AboutBox UserForm by double-clicking its name in the Project window.**

- 3. Double-click the CommandButton control.** The VBE activates the code module for the UserForm and inserts the Sub and End Sub statements for the button's Click event, as shown in Figure 41.10.

- 4. Insert the following statement before the End Sub statement:**

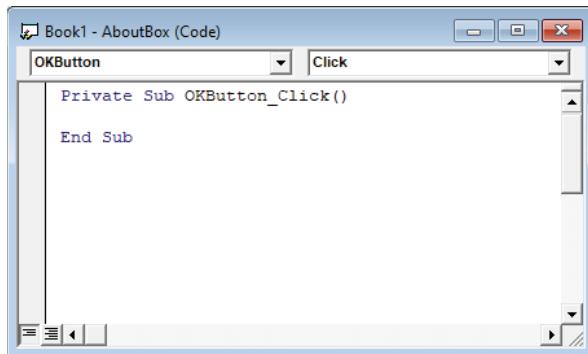
```
Unload Me
```

This statement simply dismisses the UserForm by using the Unload statement. The complete event handler procedure is as follows:

```
Private Sub OKButton_Click()
    Unload Me
End Sub
```

FIGURE 41.10

The code module for the UserForm



After you add the event procedure, clicking the OK button closes the form.

Note

The `Me` keyword when used in a UserForm's code module is a shortcut to refer to the UserForm itself. It's the same as coding `Unload AboutBox`, but if you decide to change the name of the form, using `Me` will continue to work.

Looking at Another UserForm Example

This example is an enhanced version of the `ChangeCase` procedure presented at the beginning of the chapter. Recall that the original version of this macro changes the text in the selected cells to uppercase characters. This modified version asks the user what type of case change to make: uppercase, lowercase, or proper case (initial capitals).



This workbook is available on this book's website at www.wiley.com/go/excel365bible. The file is `change case.xlsm`.

Creating the UserForm

This UserForm needs one piece of information from the user: the type of change to make to the text. Because only one option can be selected, `OptionButton` controls are appropriate. Start with an empty workbook and follow these steps to create the UserForm:

1. Press **Alt+F11** to activate the **VBE**.
2. In the **VBE**, choose **Insert** \Rightarrow **UserForm**. The VB Editor adds an empty form named `UserForm1` and displays the **Toolbox**.
3. Press **F4** to display the **Properties** window and then change the following properties of the **UserForm** object:

Property	Change to
Name	UChangeCase
Caption	Change Case

4. Add a CommandButton object to the UserForm and then change the following properties for the CommandButton:

Property	Change to
Name	OKButton
Caption	OK
Default	True

5. Add another CommandButton object and then change the following properties:

Property	Change to
Name	CancelButton
Caption	Cancel
Cancel	True

6. Add an OptionButton control and then change the following properties. (This option is the default, so its Value property should be set to True.)

Property	Change to
Name	OptionUpper
Caption	Upper Case
Value	True

7. Add a second OptionButton control and then change the following properties:

Property	Change to
Name	OptionLower
Caption	Lower Case

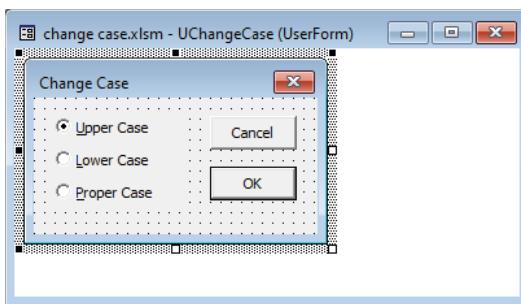
8. Add a third OptionButton control and then change the following properties:

Property	Change to
Name	OptionProper
Caption	Proper Case

9. Adjust the size and position of the controls and the form until your UserForm resembles the one shown in Figure 41.11. Make sure that the controls do not overlap.

FIGURE 41.11

The UserForm after adding controls and adjusting some properties



TIP

The VBE provides several useful commands to help you size and align the controls. For example, you can make a group of selected controls the same size or move them so that they're all aligned to the left. Select the controls with which you want to work and then choose a command from the Format menu. These commands tend to be self-explanatory, and the Help system has complete details.

Creating event handler procedures

The next step is to create two event handler procedures: one to handle the Click event for the CancelButton CommandButton and the other to handle the Click event for the OKButton CommandButton. Event handlers for the OptionButton controls are not necessary. The VBA code can determine which of the three OptionButton controls is selected, but it does not need to react when the choice is changed—only when OK or Cancel is clicked.

Event handler procedures are stored in the UserForm code module. To create the procedure to handle the Click event for the CancelButton, follow these steps:

1. Activate the UChangeCase UserForm form by double-clicking its name in the Project window.

2. **Double-click the CancelButton control.** The VBE activates the code module for the UserForm and inserts an empty procedure.

3. **Insert the following statement before the End Sub statement:**

```
Unload Me
```

That's all there is to it. The following is a listing of the entire procedure that's attached to the Click event for the CancelButton:

```
Private Sub CancelButton_Click()
    Unload Me
End Sub
```

This procedure is executed when the CancelButton is clicked. It consists of a single statement that unloads the form.

Next, add the code to handle the Click event for the OKButton control. Follow these steps:

1. **Select OKButton from the drop-down list at the top of the module or reactivate the UserForm and double-click the OKButton control.** The VBE creates a new procedure called OKButton_Click.
2. **Enter the following code.** The VBE has already entered the first and last statements for you:

```
Private Sub OKButton_Click()
    ' Exit if a range is not selected
    If TypeName(Selection) <> "Range" Then Exit Sub

    ' Upper case
    If Me.OptionUpper.Value Then
        For Each cell In Selection
            If Not cell.HasFormula Then
                cell.Value = StrConv(cell.Value, vbUpperCase)
            End If
        Next cell
    End If
    ' Lower case
    If Me.OptionLower.Value Then
        For Each cell In Selection
            If Not cell.HasFormula Then
                cell.Value = StrConv(cell.Value, vbLowerCase)
            End If
        Next cell
    End If
    ' Proper case
    If Me.OptionProper.Value Then
        For Each cell In Selection
            If Not cell.HasFormula Then
                cell.Value = StrConv(cell.Value, vbProperCase)
            End If
        Next cell
    End If
End Sub
```

```
    Next cell  
End If
```

```
Unload Me  
End Sub
```

The macro starts by checking the type of selection. If a range is not selected, the procedure ends. The remainder of the procedure consists of three separate blocks. Only one block is executed, determined by which OptionButton is selected. The selected OptionButton has a Value of True. Finally, the UserForm is unloaded (dismissed).

Showing the UserForm

At this point, the UserForm has all the necessary controls and event procedures. All that's left is a way to display the form. This section explains how to write a VBA procedure to display the UserForm:

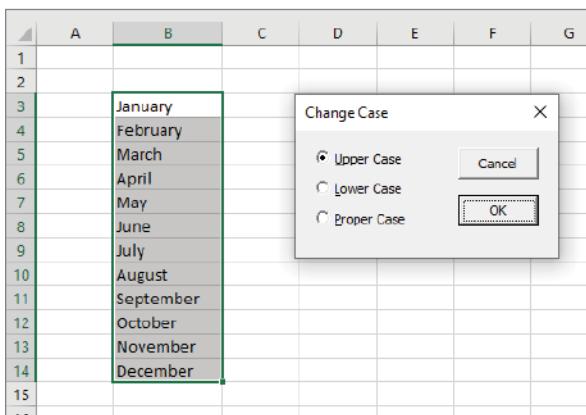
1. **Make sure the VBE window is activated.**
2. **Insert a module by choosing Insert \Rightarrow Module.**
3. **In the empty module, enter the following code:**

```
Sub ShowUserForm()  
    UChangeCase.Show  
End Sub
```

4. **Choose Run \Rightarrow Run Sub/UserForm (or press F5).** The Excel window is activated, and the new UserForm is displayed, as shown in Figure 41.12.

FIGURE 41.12

Displaying the UserForm



Testing the UserForm

To try the UserForm from Excel, follow these steps:

1. **Activate Excel.**
2. **Enter some text into a range of cells.**
3. **Select the range with the text.**
4. **Choose Developer \Rightarrow Code \Rightarrow Macros (or press Alt+F8).** The Macro dialog box appears.
5. **Select ShowUserForm from the list of macros and then click Run.** The UserForm appears, as shown in Figure 41.12.
6. **Make your choice and click OK.**

Try it with a few more selections, including noncontiguous cells. Notice that if you click Cancel, the UserForm is dismissed and no changes are made.

The code does have a problem, though: if you select one or more entire columns, the procedure processes every cell, which can take a long time. The version of the workbook on the website corrects this problem by working with a subset of the selection that intersects with the workbook's used range.

Making the macro available from a worksheet button

At this point, everything should be working properly. However, you have no quick and easy way to execute the macro. A good way to execute this macro would be from a button on the worksheet. You can use the following steps:

1. **Choose Developer \Rightarrow Controls \Rightarrow Insert and click the Button control in the Form Controls group.**
2. **Click and drag in the worksheet to create the button.** The Assign Macro dialog box appears.
3. **Select the ShowUserForm macro and then click OK.**
4. **(Optional) At this point, the button is still selected, so you can change the text to make it more descriptive.** You can also right-click the button and choose Edit Text at any time to change the text.

After you perform the preceding steps, clicking the button executes the macro and displays the UserForm.



The button in this example is from the Form Controls group. Excel also provides a button in the ActiveX Controls group. See Chapter 42 for more information about the ActiveX Controls group.

Making the macro available on your Quick Access Toolbar

If you'd like to use this macro while other workbooks are active, you may want to add a button to your Quick Access Toolbar. Follow these steps:

1. **Make sure the workbook containing the macro is open.**
2. **Right-click on the Quick Access Toolbar and choose Customize Quick Access Toolbar from the shortcut menu.** The Excel Options dialog box appears, with the Quick Access Toolbar tab selected.
3. **Choose Macros from the Choose Commands From drop-down menu on the left.** You'll see your macro listed.
4. **Select the macro's name and click Add to add the item to the list on the right.**
5. **(Optional) To change the icon, click Modify and choose a new image; then click OK.** You can also change the Display Name.
6. **Click OK to close the Excel Options dialog box.** The icon appears on your Quick Access Toolbar.

Enhancing UserForms

Creating UserForms can make your macros much more versatile. You can create custom commands that display dialog boxes that look exactly like those that Excel uses. The following sections contain some additional information to help you develop custom dialog boxes that work like those that are built into Excel.

Adding accelerator keys

All Excel dialog boxes work well with a mouse and a keyboard because each control has an associated accelerator key. The user can press Alt plus the accelerator key to work with a specific dialog box control.

Your custom dialog boxes should also have accelerator keys for all controls. You add accelerator keys in the Properties window by entering a character for the Accelerator property.

Your accelerator key can be any letter, number, or punctuation, regardless of whether that character appears in the control's caption. It's a good practice to use a letter that is in the control's caption, though, because that letter will be underlined—a visual cue for the user. (See Figure 41.12 for an example of option buttons with accelerator keys.) Another common convention is to use the first letter of the control's caption. But don't duplicate accelerator keys. If the first letter is already taken, use a different letter, preferably one that is easy to associate with the word (like a hard consonant). If you have duplicate accelerator keys, the accelerator key acts on the next control in the tab order of the UserForm. Then, pressing the accelerator key again takes you to the second control with that accelerator.

Some controls (such as text boxes) don't have a Caption property and other controls (such as labels) can't have the focus. You can assign an accelerator key to a label that describes the

control and put that label right before your target control in the tab order. Pressing the accelerator key for a control that can't take the focus activates the next control in the tab order.

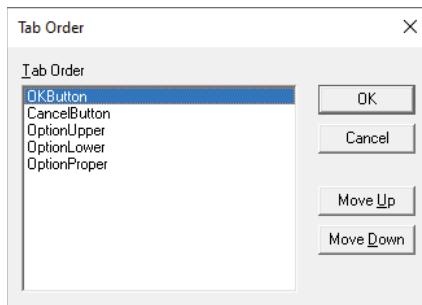
Controlling tab order

The previous section refers to a UserForm's tab order. When you're working with a UserForm, pressing Tab and Shift+Tab cycles through the dialog box's controls. When you create a UserForm, you should make sure that the tab order is correct. Usually, it means that tabbing should move through the controls in a logical sequence.

To view or change the tab order in a UserForm, choose View \leftrightarrow Tab Order to display the Tab Order dialog box (see Figure 41.13). You can then select a control from the list; use the Move Up and Move Down buttons to change the tab order for the selected control.

FIGURE 41.13

Adjusting the tab order in a UserForm



Learning More

Mastering UserForms takes practice. You should closely examine the dialog boxes that Excel uses to get a feel for the way dialog boxes are designed. You can duplicate many of the dialog boxes that Excel uses.

The best way to learn more about creating dialog boxes is by using the VBA Help system. Pressing F1 is the quickest way to display the Help window.

Using UserForm Controls in a Worksheet

IN THIS CHAPTER

Understanding why you may want to use controls on a worksheet

Using controls

Looking at the Controls Toolbox controls

Chapter 41, “Creating UserForms,” presented an introduction to UserForms. If you like the idea of using dialog box controls but you don’t like the idea of creating a custom dialog box, this chapter is for you. It explains how to enhance your worksheet with a variety of interactive controls, such as buttons, list boxes, and option buttons.

Understanding Why to Use Controls on a Worksheet

The main reason to use controls on a worksheet is to make it easier for the user to provide input. For example, if you create a model that uses one or more input cells, you can create controls to allow the user to select values for the input cells.

Adding controls to a worksheet requires much less effort than creating a dialog box. In addition, you may not have to create any macros because you can link a control to a worksheet cell. For example, if you insert a CheckBox control on a worksheet, you can link it to a particular cell. When the CheckBox is checked, the linked cell displays TRUE. When the CheckBox is not checked, the linked cell displays FALSE.

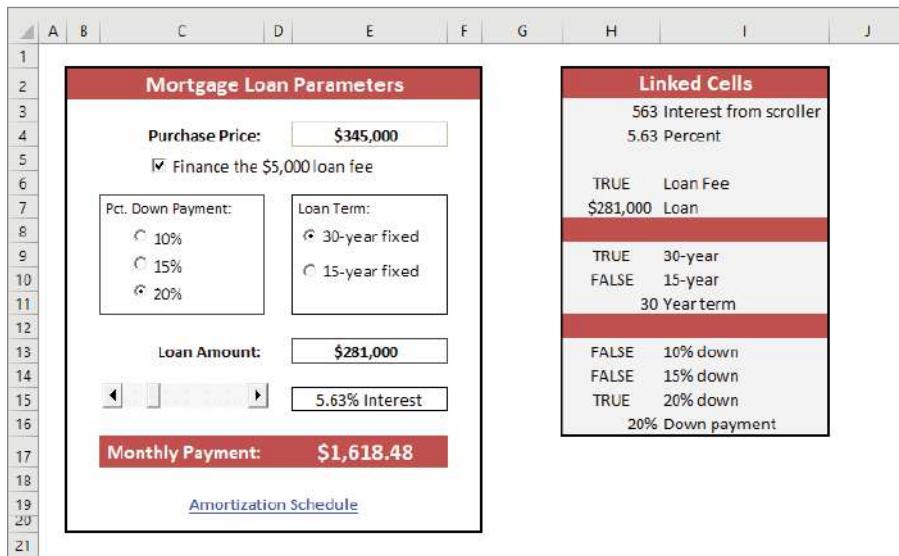
Figure 42.1 shows an example that uses three types of controls: a Checkbox, two sets of OptionButtons, and a ScrollBar. The user’s selections are used to display a loan amortization schedule on another worksheet. The workbook is interactive, but it uses no macros.



This workbook is available on this book’s website at www.wiley.com/go/excel365bible. The file is named Ch42 mortgage loan.xlsx.

FIGURE 42.1

This worksheet uses UserForm controls.



Adding controls to a worksheet can be a bit confusing because Excel offers two different sets of controls, both of which you access by choosing Developer \Rightarrow Controls \Rightarrow Insert:

Form controls These controls are unique to Excel.

ActiveX controls These controls are a subset of those that are available for use on UserForms.

Figure 42.2 shows the controls that appear when you choose Developer \Rightarrow Controls \Rightarrow Insert. When you move your mouse pointer over a control, Excel displays a ScreenTip that identifies the control.

To add to the confusion, many controls are available from both sources. For example, a control named ListBox is listed in both Form controls and ActiveX controls. However, they're two entirely different controls. In general, Form controls are easier to use, but ActiveX controls provide more flexibility.

NOTE

This chapter focuses exclusively on ActiveX controls. If the Developer tab isn't displayed in the Ribbon, right-click any Ribbon control and select *Customize the Ribbon* from the shortcut menu. The *Customize Ribbon* tab of the Excel Options dialog box appears. In the list box on the right, place a check mark next to *Developer*. Click *OK* to return to Excel.

FIGURE 42.2

Excel's two sets of worksheet controls

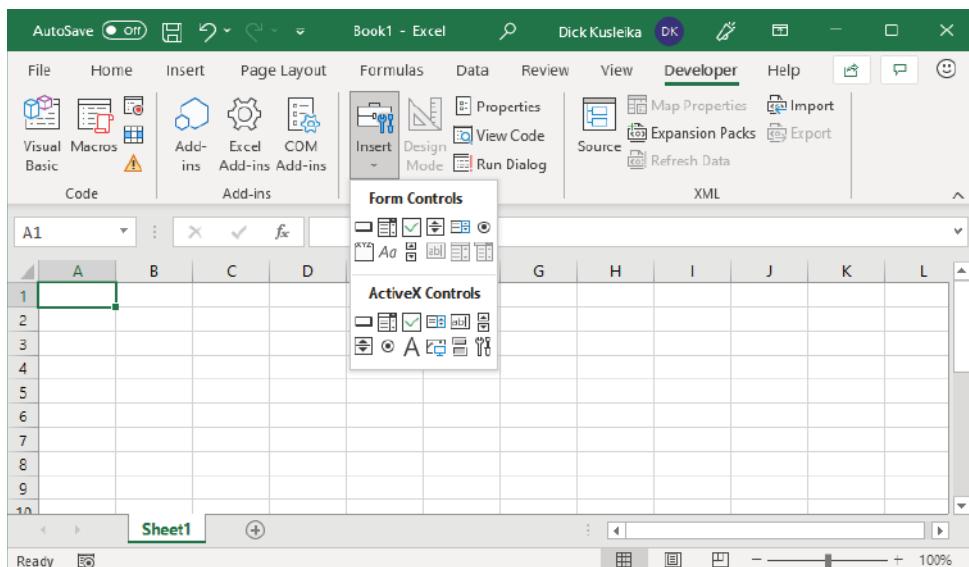


Table 42.1 describes the ActiveX controls.

TABLE 42.1 ActiveX Controls

Button	What It Does
Command Button	Inserts a CommandButton control (a clickable button).
Combo Box	Inserts a ComboBox control (a drop-down list).
Check Box	Inserts a CheckBox control (to control Boolean options).
List Box	Inserts a ListBox control (to allow the user to select an item from a list).
Text Box	Inserts a TextBox control (to allow the user to type text).
Scroll Bar	Inserts a ScrollBar control (to specify a value by dragging a bar).
Spin Button	Inserts a SpinButton control (to increment a value by clicking up or down).
Option Button	Inserts an OptionButton control (to allow a user to select from multiple options).
Label	Inserts a Label control (to display text).
Image	Inserts an Image control (to display an image).
Toggle Button	Inserts a ToggleButton control (to control Boolean options).
More Controls	Displays a list of other ActiveX controls that are installed on your system. Not all of these controls work with Excel.

Using Controls

Adding ActiveX controls in a worksheet is easy, but you need to learn a few basic facts about how to use them.

Adding a control

To add a control to a worksheet, choose Developer \Rightarrow Controls \Rightarrow Insert. From the Insert drop-down list, click the control you want to use and then drag in the worksheet to create the control. You don't need to be too concerned about the exact size or position of the control because you can modify those properties at any time.

CAUTION

Make sure that you select a control from the ActiveX controls, not from the Form controls. If you insert a Form control, the instructions in this chapter will not apply. When you choose Developer \Rightarrow Controls \Rightarrow Insert, the ActiveX controls appear in the lower half of the list.

Learning about Design mode

When you add a control to a worksheet, Excel goes into Design mode. In this mode, you can adjust the properties of any controls on your worksheet, add or edit the macro for the control, or change the control's size or position.

NOTE

When Excel is in Design mode, the Design Mode button in the Developer \Rightarrow Controls group appears highlighted. You can click this button to toggle Design mode on and off.

When Excel is in Design mode, the controls aren't enabled. To test the controls, you must exit Design mode by clicking the Design Mode button. When you're working with controls, you'll probably need to switch in and out of Design mode frequently.

Adjusting properties

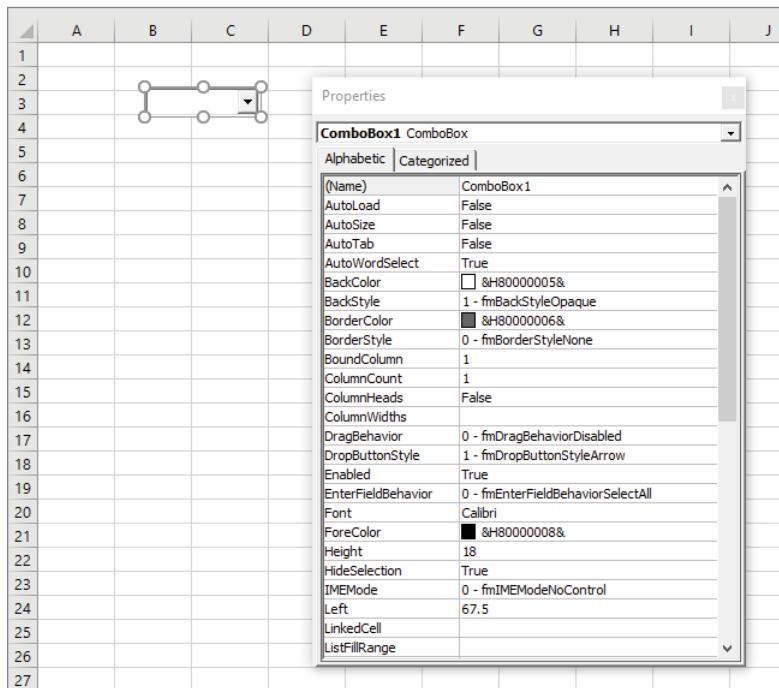
Every control that you add has various properties that determine how it looks and behaves. You can adjust these properties only when Excel is in Design mode. When you add a control to a worksheet, Excel enters Design mode automatically. If you need to change a control after you exit Design mode, click the Design Mode button in the Controls group of the Developer tab.

To change the properties for a control, follow these steps:

1. **Make sure Excel is in Design mode.**
2. **Click the control to select it.**
3. **If the Properties window isn't visible, click the Properties icon in the Controls group of the Developer tab.** The Properties window appears, as shown in Figure 42.3.
4. **Select the property and make the change.**

FIGURE 42.3

Use the Properties window to adjust the properties of a control—in this case, a ComboBox control.



The manner in which you change a property depends on the property. Some properties display a drop-down list from which you can select from a list of options. Others (such as Font) provide a button that displays a dialog box when clicked. Other properties require you to type the property value. When you change a property, the change takes effect immediately.

TIP

To find out about a particular property, select the property in the Properties window and press F1.

The Properties window has two tabs. The Alphabetic tab displays the properties in alphabetical order. The Categorized tab displays the properties by category. Both tabs show the same properties; only the order is different.

Using common properties

Each control has its own unique set of properties. However, many controls share properties. This section describes some properties that are common to all or many controls, as set forth in Table 42.2.

NOTE

Some ActiveX control properties are required (for example, the Name property). In other words, you can't leave the property empty. If a required property is missing, Excel will always tell you so by displaying an error message.

TABLE 42.2 Properties Shared by Multiple Controls

Property	Description
AutoSize	If True, the control resizes itself automatically based on the text in its caption.
BackColor	The background color of the control.
BackStyle	The style of the background (either transparent or opaque).
Caption	The text that appears on the control.
Left and Top	Values that determine the control's position.
LinkedCell	A worksheet cell that contains the current value of a control.
ListFillRange	A worksheet range that contains items displayed in a ListBox or ComboBox control.
Name	The name of the control. When you add a control, Excel assigns it a name based on the control type. You can change the name to any valid name. However, each control's name must be unique on the worksheet.
Picture	Enables you to specify a graphics image to display.
Value	The control's value.
Visible	If False, the control is hidden.
Width and Height	Values that determine the control's size.

Linking controls to cells

Often, you can use ActiveX controls in a worksheet without using macros. Many controls have a LinkedCell property, which specifies a worksheet cell that is linked to the control.

For example, you may add a SpinButton control and specify cell B1 as its LinkedCell property. After you do so, cell B1 contains the value of the SpinButton, and clicking the SpinButton changes the value in cell B1. You can, of course, use the value contained in the linked cell in your formulas.

NOTE

When specifying the LinkedCell property in the Properties window, you can't "point" to the linked cell in the worksheet. You must type the cell address or its name (if it has one).

Creating macros for controls

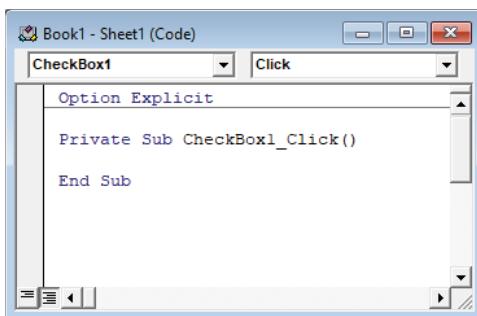
To create a macro for a control, you must use the Visual Basic Editor (VBE). The macros are stored in the code module for the sheet that contains the control. For example, if you place an ActiveX control on Sheet2, the VBA code for that control is stored in the Sheet2 code module. Each control can have a macro to handle any of its events. For example, a CommandButton control can have a macro for its Click event, its DblClick event, and various other events.

TIP

The easiest way to access the code module for a control is to double-click the control while in Design mode. Excel displays the VBE and creates an empty procedure for the control's default event. For example, the default event for a CheckBox control is the Click event. Figure 42.4 shows the autogenerated code for a control named CheckBox1, located on Sheet1.

FIGURE 42.4

Double-clicking a control in Design mode activates the VBE and enters an empty event-handler procedure.



The control's name appears in the upper-left portion of the code window, and the event appears in the upper-right area. If you want to create a macro that executes when a different event occurs, select the event from the list in the upper-right area.

The following steps demonstrate how to insert a CommandButton and to create a simple macro that displays a message when the button is clicked:

1. Choose **Developer** \Rightarrow **Controls** \Rightarrow **Insert**.
2. Click the **CommandButton** tool in the **ActiveX Controls** section.
3. Click and drag in the worksheet to create the button. Excel automatically enters Design mode.
4. Double-click the button. The VBE is activated, and an empty procedure is created for the button's Click event.

5. Enter the following VBA statement before the End Sub statement:

```
MsgBox "Hello, it's " & Time
```

6. Press Alt+F11 to return to Excel.

7. (Optional) Adjust any other properties for the CommandButton using the Properties window. Choose Developer \Rightarrow Controls \Rightarrow Properties if the Properties window isn't visible.

8. Click the Design Mode button in the Developer \Rightarrow Controls group to exit Design mode.

After you perform the preceding steps, when you click the CommandButton, the message box appears and displays the current time.

NOTE

You must enter the VBA code manually. You can't create macros for controls using the VBA macro recorder. However, you can record a macro and then execute it from an event procedure. For example, if you've recorded a macro named FormatCells, you can type a statement with the macro's name. When that statement is executed, your recorded macro will run. Or you can copy the recorded code and paste it to your event procedure.

Reviewing the Available ActiveX Controls

The following sections describe the ActiveX controls that are available for use in your worksheets.



This book's website at www.wiley.com/go/excel365bible contains a file that includes examples of all the ActiveX controls. This file is named `worksheet_controls.xlsb`.

CheckBox

A CheckBox control is useful for getting a binary choice: yes or no, true or false, on or off, and so on. The following is a description of the most useful properties of a CheckBox control:

Accelerator A letter that enables the user to change the value of the control by using the keyboard. For example, if the accelerator is *A*, pressing *Alt+A* changes the value of the CheckBox control. The accelerator letter is underlined in the caption of the control.

LinkedCell The worksheet cell that's linked to the CheckBox. The cell displays TRUE if the control is checked and FALSE if the control is not checked.

ComboBox

A ComboBox control is a combination of a TextBox and a ListBox. It acts as a TextBox because the user can type in it just like a TextBox, even if what they type isn't on the list. It acts as a ListBox because when you click its drop-down arrow, a list of available items appears.

FIGURE 42.5

A ComboBox control

	A	B	C
1	October		January
2			February
3			March
4	October		April
5	May		May
6	June		June
7	July		July
8	August		August
9	September		September
10	October		October
11	November		November
12	December		December
13			

Figure 42.5 shows a ComboBox control that uses the range D1:D12 for the ListFillRange and cell A1 for the LinkedCell.

The following is a description of the most useful properties of a ComboBox control:

- BoundColumn** If the ListFillRange contains multiple columns, this property determines which column contains the returned value.
- ColumnCount** This specifies the number of columns to display in the list.
- LinkedCell** This specifies the worksheet cell that displays the selected item.
- ListFillRange** This specifies the worksheet range that contains the list items.
- ListRows** This specifies the number of items to display when the list drops down.
- ListStyle** This determines the appearance of the list items.
- Style** This determines whether the control acts like a drop-down list or a ComboBox. A drop-down list doesn't allow the user to enter a new value.



You can also create a drop-down list directly in a cell by using data validation. See Chapter 23, “Using Data Validation,” for details.

CommandButton

A CommandButton control is used to execute a macro. When a CommandButton is clicked, it executes an event procedure with a name that consists of the CommandButton name, an underscore, and the word Click. For example, if a CommandButton is named MyButton, clicking it executes the macro named MyButton_Click. This macro is stored in the code module for the sheet that contains the CommandButton.

Image

An Image control is used to display an image. These are the most useful properties of an Image control:

AutoSize If this property is TRUE, the Image control is resized automatically to fit the image.

Picture This is the path to the image file. Click the button in the Properties window, and Excel displays a dialog box so that you can locate the image. Or, copy the image to the Clipboard, select the Picture property in the Properties window, and press Ctrl+V.

PictureSizeMode This determines how the picture is changed when the container size is different than the picture.

TIP

You can also insert an image on a worksheet by choosing Insert \Rightarrow Illustrations \Rightarrow Pictures.

Label

A Label control simply displays text. Like on a UserForm, it's used to describe other controls. You can also use its Click event to activate other controls with an accelerator key.

ListBox

A ListBox control presents a list of items, and the user can select an item (or multiple items). It's similar to a ComboBox. The main difference is that a ListBox doesn't require you to click a drop-down arrow to display more than one item at a time.

The following is a description of the most useful properties of a ListBox control:

BoundColumn If the list contains multiple columns, this property determines which column contains the returned value.

ColumnCount This specifies the number of columns to display in the list.

IntegralHeight This is TRUE if the height of the ListBox adjusts automatically to display full lines of text when the list is scrolled vertically. If this property is FALSE, the ListBox may display partial lines of text when it's scrolled vertically.

LinkedCell This specifies the worksheet cell that displays the selected item.

ListFillRange This specifies the worksheet range that contains the list items.

ListStyle This determines the appearance of the list items.

MultiSelect This determines whether the user can select multiple items from the list.

NOTE

If you use a MultiSelect ListBox, you can't specify a LinkedCell; you need to write a macro to determine which items are selected.

OptionButton

OptionButton controls are useful when the user needs to select from a small number of items. OptionButton controls are always used in groups of at least two.

The following are the most useful properties of an OptionButton control:

Accelerator A letter that lets the user select the option by using the keyboard. For example, if the accelerator for an OptionButton is *C*, pressing Alt+C selects the control.

GroupName A name that identifies an OptionButton as being associated with other OptionButtons with the same GroupName property.

LinkedCell The worksheet cell that's linked to the OptionButton. The cell displays TRUE if the control is selected or FALSE if the control isn't selected.

Note

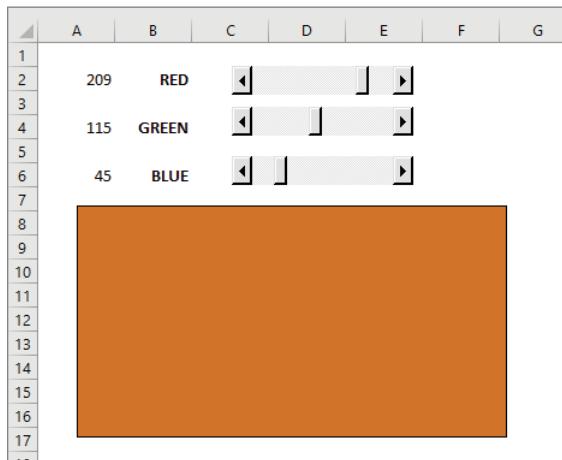
If your worksheet contains more than one set of OptionButton controls, you must ensure that each set of OptionButtons has a different GroupName property. Otherwise, all OptionButtons become part of the same set.

ScrollBar

A ScrollBar control is useful for specifying a cell value. Figure 42.6 shows a worksheet with three ScrollBar controls. These ScrollBar controls are used to change the color in the rectangle shape. The value of the ScrollBar control determines the red, green, or blue component of the rectangle's color. This example uses a few simple macros to change the colors.

FIGURE 42.6

This worksheet has three ScrollBar controls.



The following are the most useful properties of a ScrollBar control:

- Value** The current value of the control
- Min** The minimum value for the control
- Max** The maximum value for the control
- LinkedCell** The worksheet cell that displays the value of the control
- SmallChange** The amount that the control's value is changed by clicking the arrows
- LargeChange** The amount that the control's value is changed by clicking in the scroll area

The ScrollBar control is most useful for selecting a value that extends across a wide range of possible values.

SpinBox

A SpinButton control lets the user select a value by clicking the control, which has two arrows (one to increase the value and the other to decrease the value). A SpinButton can display either horizontally or vertically.

The following are the most useful properties of a SpinButton control:

- Value** The current value of the control.
- Min** The minimum value of the control.
- Max** The maximum value of the control.
- LinkedCell** The worksheet cell that displays the value of the control.
- SmallChange** The amount that the control's value is changed by a click. Usually, this property is set to 1, but you can make it any value.

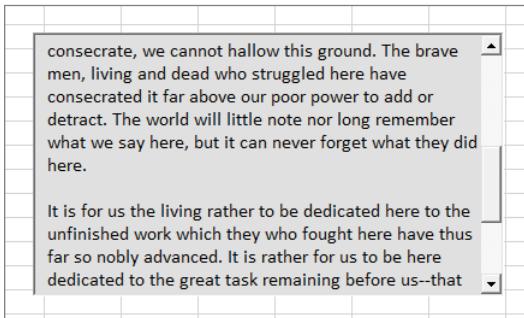
TextBox

On the surface, a TextBox control may not seem useful. After all, it simply contains text—you can usually use worksheet cells to get text input. In fact, TextBox controls are useful not so much for input control as for output control. Because a TextBox can have scrollbars, you can use a TextBox to display a great deal of information in a small area.

Figure 42.7 shows a TextBox control that contains Lincoln's Gettysburg Address. Notice the vertical scrollbar displayed using the ScrollBars property.

FIGURE 42.7

A TextBox control with a vertical scrollbar



The following are the most useful properties of a TextBox control:

AutoSize This determines whether the control adjusts its size automatically, based on the amount of text.

IntegralHeight If TRUE, the height of the TextBox adjusts automatically to display full lines of text when the list is scrolled vertically. If FALSE, the TextBox may display partial lines of text when it's scrolled vertically.

MaxLength This determines the maximum number of characters allowed in the TextBox. If 0, no limit exists on the number of characters.

MultiLine If TRUE, the TextBox can display more than one line of text. To create a new line in a MultiLine textbox, press Shift+Enter.

TextAlign This determines the way the text is aligned in the TextBox.

WordWrap If TRUE and the MultiLine property is also TRUE, lines that are longer than the width of the control are wrapped to the next line.

ScrollBars This determines the type of ScrollBars for the control: horizontal, vertical, both, or none.

ToggleButton

A ToggleButton control has two states: on and off. Clicking the button toggles between these two states, and the button changes its appearance to indicate its current state. Its value is either TRUE (pressed) or FALSE (not pressed). You can often use a ToggleButton in place of a CheckBox control.

Working with Excel Events

IN THIS CHAPTER

- Understanding events
- Using workbook-level events
- Working with worksheet events
- Using Special Application events

In the preceding chapters, we presented a few examples of VBA event-handler procedures for ActiveX controls. These procedures are the keys to making your Excel applications interactive. This chapter introduces the concept of events for Excel objects and includes many examples that you can adapt to meet your own needs.

Understanding Events

Excel monitors a variety of events and can execute your VBA code when any of these events occur. This chapter covers the following types of events:

Workbook events These occur for a particular workbook. Examples include Open (the workbook is opened or created), BeforeSave (the workbook is about to be saved), and NewSheet (a new sheet is added). You must store VBA code for workbook events in the ThisWorkbook code module.

Worksheet events These occur for a particular worksheet. Examples include Change (a cell on the sheet is changed), SelectionChange (the selection on the worksheet is changed), and Calculate (the worksheet is recalculated). VBA code for worksheet events must be stored in the code module for the worksheet (for example, the module named Sheet1).

Special Application events The final category consists of two useful application-level events: OnTime and OnKey. These are different from other events because the code isn't in a class module. Rather, you set the events by calling a method of the Application object.

Many events exist at both the worksheet and workbook levels. For example, Sheet1 has an event called Change that fires when any cell on Sheet1 is changed. The workbook has an event called SheetChange that fires every time any cell on any sheet is changed. The workbook version of this event has an additional argument that lets you know which sheet was affected.

Entering Event-Handler VBA Code

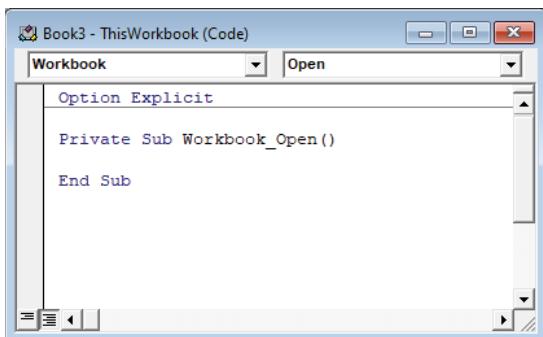
Every event-handler procedure must reside in a specific type of code module. Code for workbook-level events is stored in the ThisWorkbook code module. Code for worksheet-level events is stored in the code module for the particular sheet (for example, the code module named Sheet1). These modules are added to the project automatically by Excel. You can't add or remove the ThisWorkbook module. To add or remove a sheet's module, you must add or remove the sheet from the workbook.

In addition, every event-handler procedure has a predetermined name. You can declare the procedure by typing it, but a much better approach is to let the Visual Basic Editor (VBE) do it for you by using the two drop-down controls at the top of the window.

Figure 43.1 shows the code module for the ThisWorkbook object. Select this code module by double-clicking it in the Project window. To insert a procedure declaration, select Workbook from the objects list in the upper left of the code window and then select the event from the procedures list in the upper right. When you do, you get a procedure "shell" that contains the procedure declaration line and an End Sub statement.

FIGURE 43.1

The best way to create an event procedure is to let the VBE do it for you.



For example, if you select Workbook from the objects list and Open from the procedures list, the VBE inserts the following (empty) procedure:

```
Private Sub Workbook_Open()

End Sub
```

Your event-handler VBA code goes between these two lines.

Some event-handler procedures contain an argument list. For example, you may need to create an event-handler procedure to monitor the SheetActivate event for a workbook. (This event is

triggered when a user activates a different sheet.) If you use the technique described in the previous section, the VBE creates the following procedure:

```
Private Sub Workbook_SheetActivate(ByVal Sh As Object)
    End Sub
```

This procedure uses one argument (Sh), which represents the activated sheet. In this case, Sh is declared as an Object data type rather than a Worksheet data type because the activated sheet also can be a chart sheet.

Your code can, of course, use information passed as an argument. The following example displays the name of the activated sheet by accessing the argument's Name property. The argument becomes either a Worksheet object or a Chart object:

```
Private Sub Workbook_SheetActivate(ByVal Sh As Object)
    MsgBox Sh.Name & " was activated."
    End Sub
```

Several event-handler procedures use a Boolean argument named Cancel. For example, the declaration for a workbook's BeforePrint event is:

```
Private Sub Workbook_BeforePrint(Cancel As Boolean)
```

The value of Cancel passed to the procedure is False. However, your code can set Cancel to True, which cancels the printing. The following example demonstrates this:

```
Private Sub Workbook_BeforePrint(Cancel As Boolean)
    Msg = "Have you loaded the 5164 label stock?"
    Ans = MsgBox(Msg, vbYesNo, "About to print...")
    If Ans = vbNo Then Cancel = True
    End Sub
```

The Workbook_BeforePrint procedure executes before the workbook prints. This procedure displays a message box asking the user to verify that the correct paper is loaded. If the user clicks the No button, Cancel is set to True, and nothing prints.

Here's another procedure that uses the workbook's BeforePrint event. This example overcomes a deficiency in Excel's headers and footers: it's not possible to use the contents of a cell for a page header or footer. This simple procedure is triggered when the workbook is printed. It places the contents of cell A1 in the page header:

```
Private Sub Workbook_BeforePrint(Cancel As Boolean)
    ActiveSheet.PageSetup.CenterHeader =
        Worksheets(1).Range("A1")
    End Sub
```

Using Workbook-Level Events

Workbook-level events occur for a particular workbook. Table 43.1 lists the most used workbook events along with a brief description of each.

TABLE 43.1 Workbook Events

Event	Action That Triggers the Event
Activate	The workbook is activated.
AfterSave	The workbook was saved.
BeforeClose	The workbook is about to be closed.
BeforePrint	The workbook (or anything in it) is about to be printed.
BeforeSave	The workbook is about to be saved.
Deactivate	The workbook is deactivated.
NewSheet	A new sheet is created in the workbook.
Open	The workbook is opened.
SheetActivate	Any sheet in the workbook is activated.
SheetBefore-DoubleClick	Any worksheet in the workbook is double-clicked. This event occurs before the default double-click action.
Sheet-BeforeRightClick	Any worksheet in the workbook is right-clicked. This event occurs before the default right-click action.
SheetChange	Any worksheet in the workbook is changed by the user.
SheetDeactivate	Any sheet in the workbook is deactivated.
Sheet-SelectionChange	The selection on any worksheet in the workbook is changed.
WindowActivate	Any workbook window is activated.
WindowDeactivate	Any workbook window is deactivated.

The following sections present examples of using workbook-level events.

CAUTION

All the sample procedures that follow must be located in the code module for the `ThisWorkbook` object. If you put them into any other type of code module, they won't work—and you will not see an error message.

Using the Open event

One of the most common monitored events is a workbook's `Open` event. This event is triggered when the workbook opens and executes the `Workbook_Open` procedure. A `Workbook_Open` procedure is versatile, and it is often used for the following tasks:

- Displaying welcome messages
- Opening other workbooks
- Activating a specific sheet
- Ensuring that certain conditions are met; for example, a workbook may require that a particular add-in is installed

CAUTION

Keep in mind that there is no guarantee that your Workbook_Open procedure will be executed. For example, the user may choose to disable macros. And if the user holds down the Shift key while opening a workbook, the workbook's Workbook_Open procedure won't execute.

The following is a simple example of a Workbook_Open procedure. It uses the VBA Weekday function to determine the day of the week. If it's Friday, a message box appears to remind the user to perform a file backup. If it's not Friday, nothing happens.

```
Private Sub Workbook_Open()
    If Weekday(Now) = 6 Then
        Msg = "Make sure you do your weekly backup!"
        MsgBox Msg, vbInformation
    End If
End Sub
```

The following example performs a number of actions when the workbook is opened. It maximizes the workbook window, activates the sheet named DataEntry, selects the first empty cell in column A, and enters the current date into that cell. If a sheet named DataEntry does not exist, the code generates an error.

```
Private Sub Workbook_Open()
    ActiveWindow.WindowState = xlMaximized
    Worksheets("DataEntry").Activate
    Range("A1").End(xlDown).Offset(1, 0).Select
    ActiveCell.Value = Date
End Sub
```

Using the SheetActivate event

The following procedure executes whenever the user activates a sheet in the workbook. The code simply selects cell A1. Including the On Error Resume Next statement causes the procedure to ignore the error that occurs if the activated sheet is a chart sheet:

```
Private Sub Workbook_SheetActivate(ByVal Sh As Object)
    On Error Resume Next
    Range("A1").Select
End Sub
```

An alternative method to handle the case of a chart sheet is to check the sheet type. Use the Sh argument, which is passed to the procedure:

```
Private Sub Workbook_SheetActivate(ByVal Sh As Object)
    If TypeName(Sh) = "Worksheet" Then Range("A1").Select
End Sub
```

Using the NewSheet event

The following procedure executes whenever a new sheet is added to the workbook. The sheet is passed to the procedure as an argument. Because a new sheet can be either a worksheet or a chart sheet, this procedure determines the sheet type. If it's a worksheet, it inserts a date and time stamp in cell A1:

```
Private Sub Workbook_NewSheet(ByVal Sh As Object)
    If TypeName(Sh) = "Worksheet" Then
        Sh.Range("A1").Value = "Sheet added " & Now()
    End Sub
```

Using the BeforeSave event

The BeforeSave event occurs before the workbook is actually saved. As you know, choosing File \Rightarrow Save sometimes brings up the Save As dialog box—for example, when the file has never been saved or was opened in read-only mode.

When the Workbook_BeforeSave procedure executes, it receives an argument that enables you to identify whether the Save As dialog box will appear. The following example demonstrates this:

```
Private Sub Workbook_BeforeSave _
    (ByVal SaveAsUI As Boolean, Cancel As Boolean)

    If SaveAsUI Then
        MsgBox "Use the new file-naming convention."
    End If
End Sub
```

When the user attempts to save the workbook, the Workbook_BeforeSave procedure executes. If the save operation brings up the Save As dialog box, the SaveAsUI variable is True. The preceding procedure checks this variable and displays a message only if the Save As dialog box is displayed. In this case, the message is a reminder about how to name the file.

The BeforeSave event procedure also has a Cancel variable in its argument list. If the procedure sets the Cancel argument to True, the file is not saved.

Using the BeforeClose event

The BeforeClose event occurs before a workbook is closed. This event is often used in conjunction with a Workbook_Open event handler. For example, use the Workbook_Open procedure to initialize items in your workbook, and use the Workbook_BeforeClose procedure to clean up or restore settings to normal before the workbook closes.

When a user attempts to close a read-only workbook that has changes, Excel will prompt the user to save a copy of the workbook. In some cases, this is the desired behavior. But if the workbook is designed to be used and not saved, dismissing the Save As dialog is an annoyance and you might end up with copies of the workbook that you don't want because the user

became confused. You can use the BeforeClose event, as in the following example, to specify that the workbook has already been saved (even though it hasn't) and avoid that message.

```
Private Sub Workbook_BeforeClose(Cancel As Boolean)
    Me.Saved = True
End Sub
```

In Chapter 41, “Creating UserForms,” you saw that the Me keyword in the UserForm’s code module referred to the UserForm itself. It works the same in the ThisWorkbook and sheet code modules. In this case, Me refers to the workbook that contains the event procedure.

If you attempt to close a workbook that hasn’t been saved, Excel displays a prompt that asks whether you want to save the workbook before it closes.

CAUTION

A problem can arise from this event with unsaved workbooks. By the time the user sees the prompt to save before closing, the BeforeClose event has already executed. If the user cancels the close at the save prompt, the workbook remains open, but your event code has already run.

Working with Worksheet Events

The events for a Worksheet object are some of the most useful. As you’ll see, monitoring these events can make your applications perform feats that would otherwise be impossible.

Table 43.2 lists the more commonly used worksheet events with a brief description of each. Remember that these event procedures must be entered into the code module for the sheet. These code modules have default names like Sheet1, Sheet2, and so on.

TABLE 43.2 Worksheet Events

Event	Action That Triggers the Event
Activate	The worksheet is activated.
BeforeDoubleClick	The worksheet is double-clicked. This event occurs before the default double-click action.
BeforeRightClick	The worksheet is right-clicked. This event occurs before the default right-click action.
Change	Cells on the worksheet are changed by the user.
Deactivate	The worksheet is deactivated.
FollowHyperlink	A hyperlink on the worksheet was clicked.
SelectionChange	The selection on the worksheet is changed.

Using the Change event

A Change event is triggered when a user changes any cell in the worksheet. A Change event is not triggered when a calculation generates a different value for a formula or when an object (such as a chart or a shape) is added to the sheet.

When the Worksheet_Change procedure executes, it receives a Range object as its Target argument. This Range object corresponds to the changed cell or group of cells that triggered the event. The following example displays a message box that shows the address of the Target range:

```
Private Sub Worksheet_Change(ByVal Target As Range)
    MsgBox "Range " & Target.Address & " was changed."
End Sub
```

To get a feel for the types of actions that generate the Change event for a worksheet, enter the preceding procedure into the code module for a Worksheet object. After you enter this procedure, activate Excel and, using various techniques, make changes to the worksheet. Every time the Change event occurs, a message box displays the address of the range that changed.

Unfortunately, the Change event doesn't always work as expected, as in the following examples:

- Changing the formatting of a cell does not trigger the Change event (as expected), but choosing Home \Rightarrow Editing \Rightarrow Clear \Rightarrow Clear Formats *does*.
- Pressing Delete generates an event even if the cell is empty at the start.
- Cells changed via Excel commands may or may not trigger the Change event. For example, sorting and goal-seeking operations do not trigger the Change event. However, operations such as Find and Replace, using the AutoSum button, or adding a Total row to a table *do* trigger the event.
- If your VBA procedure changes a cell, it does trigger the Change event.

Monitoring a specific range for changes

Although the Change event occurs when any cell on the worksheet changes, most of the time you'll be concerned only with changes that are made to a specific cell or range. When the Worksheet_Change event-handler procedure is called, it receives a Range object as its argument. This Range object corresponds to the cell(s) that changed.

Assume that your worksheet has a range named InputRange and you want your VBA code to monitor changes to this range only. No Change event exists for a Range object, but you can perform a quick check within the Worksheet_Change procedure. The following procedure demonstrates this:

```
Private Sub Worksheet_Change(ByVal Target As Range)
    Dim VRange As Range
    Set VRange = Me.Range("InputRange")
```

```

If Union(Target, VRange).Address = _
    VRange.Address Then

    MsgBox "The changed cell is in the input range."
End If
End Sub

```

This example creates a Range object variable named VRange, which represents the worksheet range that you want to monitor for changes. The procedure uses the VBA Union function to determine whether VRange contains the Target range (passed to the procedure in its argument). The Union function returns an object that consists of all the cells in both its arguments. If the range address is the same as the VRange address, VRange contains Target, and a message box appears. Otherwise, the procedure ends, and nothing happens.

The preceding procedure has a potential flaw: Target may consist of a single cell or a range. For example, if the user changes more than one cell at a time, Target becomes a multicell range. As it is, all of the changed cells must be within InputRange. If you still want to act on cells within InputRange, even if some of the changes are not, you have to change the procedure to loop through all of the cells in Target. The following procedure checks each changed cell and displays a message box if the cell is within the desired range:

```

Private Sub Worksheet_Change(ByVal Target As Range)
    Set VRange = Me.Range("InputRange")
    For Each cell In Target.Cells
        If Union(cell, VRange).Address = _
            VRange.Address Then

            MsgBox "The changed cell is in the input range."
        End If
    Next cell
End Sub

```

One problem with the preceding procedure is that it displays a message for each cell in InputRange. If you change a lot of cells, you will be closing a lot of messages. You can use the Intersect method to get the cells that two ranges have in common, as in the following procedure.

```

Private Sub Worksheet_Change( _
    ByVal Target As Excel.Range)

    Dim VRange As Range
    Dim cell As Range
    Dim IRange As Range

    Set VRange = Me.Range("InputRange")
    Set IRange = Intersect(VRange, Target)

    If Not IRange Is Nothing Then
        MsgBox "The range " & IRange.Address & " was changed."
    End If
End Sub

```

If the two ranges don't have any cells in common, the IRange variable won't contain any cells. You test for that condition with the Is Nothing keywords. We want the If statement to return TRUE when IRange contain cells, so we use the Not keyword to test the opposite of Is Nothing. The construct If Not Object Is Nothing isn't the easiest code to read, but it's used often, and you'll likely encounter it if you read other people's code.



A workbook with this example is available on this book's website at www.wiley.com/go/excel365bible. The file is named `monitor a range.xlsm`.

Using the SelectionChange event

The following procedure demonstrates a SelectionChange event. It executes whenever the user makes a new selection on the worksheet:

```
Private Sub Worksheet_SelectionChange( _
    ByVal Target As Range)

    Me.Cells.Interior.ColorIndex = xlNone
    With Target
        .EntireRow.Interior.ColorIndex = 35
        .EntireColumn.Interior.ColorIndex = 35
    End With
End Sub
```

This procedure shades the row and column of the selected cells, making the selected cells easy to identify. The first statement removes the background color of all cells. Next, the entire rows and columns of the selected cells are shaded light green. Figure 43.2 shows the shading.



A workbook with this example is available on this book's website at www.wiley.com/go/excel365bible. The file is named `selection change.event.xlsm`.

CAUTION

You won't want to use this procedure if your worksheet contains background shading because the macro will erase it. However, if the shading is the result of a style applied to a table, the macro doesn't erase the table's background shading.

FIGURE 43.2

Selecting a cell causes the active cell's row and column to become shaded.

	A	B	C	D	E	F	G	H
1		Mary	Bill	Joe	Frank	Carol	Pete	Nancy
2	January	551	664	582	607	675	513	557
3	February	548	572	577	529	500	681	635
4	March	665	513	546	678	673	566	693
5	April	699	667	663	562	504	626	595
6	May	640	581	661	586	510	542	537
7	June	649	689	569	518	591	607	625
8	July	538	516	660	626	523	560	689
9	August	618	533	611	681	585	641	618
10	September	587	546	584	538	575	624	648
11	October	573	616	612	602	696	621	620
12	November	613	692	617	603	544	601	678
13	December	657	518	597	630	638	602	652
14								

Using the BeforeRightClick event

Normally, when the user right-clicks in a worksheet, a shortcut menu appears. If, for some reason, you want to prevent the shortcut menu from appearing, you can trap the RightClick event. The following procedure sets the Cancel argument to True, which cancels the RightClick event—and, thus, the shortcut menu. Instead, a message box appears:

```
Private Sub Worksheet_BeforeRightClick _
    (ByVal Target As Range, Cancel As Boolean)
    Cancel = True
    MsgBox "The shortcut menu is not available."
End Sub
```

43

Using Special Application Events

So far, the events discussed in this chapter are associated with an object like a worksheet. The following sections discuss two additional events: OnTime and OnKey. These events are not associated with an object. Instead, you access them by using methods of the Application object.

Note

Unlike the other events discussed in this chapter, you use a standard VBA module to program the On events in the following sections.

Using the OnTime event

The OnTime event occurs at a specified time. The following example demonstrates how to program Excel to beep and then display a message at 3 p.m.:

```
Sub SetAlarm()
    Application.OnTime TimeSerial(15, 0, 0), "DisplayAlarm"
End Sub

Sub DisplayAlarm()
    Beep
    MsgBox "Wake up. It's time for your afternoon break!"
End Sub
```

In this example, the SetAlarm procedure uses the OnTime method of the Application object to set up the OnTime event. This method takes two arguments: the time (the TimeSerial function is an easy way to get a time, and the hour argument of 15 is 3 p.m.) and the procedure to execute when the time occurs (DisplayAlarm in the example). In the example, after SetAlarm executes, the DisplayAlarm procedure is called at 3 p.m., bringing up the message.

You could also use VBA's TimeValue function to represent the time. TimeValue converts a string that looks like a time into a value that Excel can handle. The following statement shows another way to program an event for 3 p.m.:

```
Application.OnTime TimeValue("3:00:00 pm"), _
    "DisplayAlarm"
```

If you want to schedule an event that's relative to the current time—for example, 20 minutes from now—you can write an instruction like either of these:

```
Application.OnTime Now + TimeSerial(0, 20, 0), _
    "DisplayAlarm"

Application.OnTime Now + TimeValue("00:20:00"), _
    "DisplayAlarm"
```

You also can use the OnTime method to schedule a procedure on a particular day. Of course, you must keep your computer turned on, and Excel must be running.

To cancel an OnTime event, you must know the exact time the event was scheduled to run. Then you use the schedule argument of OnTime set to False. OnTime works to the nearest second. If you schedule something for 3 p.m., you can unschedule it with this code:

```
Application.OnTime TimeSerial(15, 0, 0), _
    "DisplayAlarm", , False
```

If you scheduled something relative to the current time and want to cancel it, you need to store that time. This code will schedule the event:

```
TimeToRun = Now + TimeSerial(0, 20, 0)
Application.OnTime TimeToRun, "DisplayAlarm"
```

The TimeToRun variable can be used to unschedule it, assuming that the variable is still in scope:

```
Application.OnTime TimeToRun, "DisplayAlarm", , False
```

Using the OnKey event

While you work, Excel constantly monitors what you type. As a result, you can set up a key-stroke or a key combination that, when pressed, executes a particular procedure.

The following example uses the OnKey method to set up an OnKey event. This event essentially reassigns the PgDn and PgUp keys. After the Setup_OnKey procedure executes, pressing PgDn executes the PgDn_Sub procedure, and pressing PgUp executes the PgUp_Sub procedure. The effect is that pressing PgDn moves down one row, and pressing PgUp moves up one row:

```
Sub Setup_OnKey()
    Application.OnKey "{PgDn}", "PgDn_Sub"
    Application.OnKey "{PgUp}", "PgUp_Sub"
End Sub

Sub PgDn_Sub()
    On Error Resume Next
    ActiveCell.Offset(1, 0).Activate
End Sub

Sub PgUp_Sub()
    On Error Resume Next
    ActiveCell.Offset(-1, 0).Activate
End Sub
```

Note

The key codes are enclosed in brackets, not parentheses. For a complete list of the keyboard codes, consult VBA Help. Search for OnKey.

TIP

The preceding examples used On Error Resume Next to ignore any errors generated. For example, if the active cell is in the first row, trying to move up one row causes an error. Furthermore, if the active sheet is a chart sheet, an error occurs because no such thing as an active cell exists in a chart sheet.

Part VI: Automating Excel

By executing the following procedure, you cancel the OnKey events, and the keys return to their normal functions:

```
Sub Cancel_OnKey()
    Application.OnKey "{PgDn}"
    Application.OnKey "{PgUp}"
End Sub
```

CAUTION

Contrary to what you may expect, using an empty string as the second argument for the OnKey method does not cancel the OnKey event. Instead, it causes Excel to ignore the keystroke and do nothing. For example, the following instruction tells Excel to ignore Alt+F4. (The percent sign represents the Alt key.)

```
Application.OnKey "%{F4}", ""
```

Seeing Some VBA Examples

IN THIS CHAPTER

Working with ranges

Working with charts

Modifying properties

Speeding up your VBA code

Our philosophy about learning to write Excel macros places heavy emphasis on examples. Often, a well-thought-out example communicates a concept much better than a lengthy description of the underlying theory. In this book, space limitations don't allow us to describe every nuance of VBA, so we prepared many examples. Don't overlook the VBA Help system for specific details. To get help while working in the Visual Basic Editor (VBE), press F1. For context-sensitive help, select a VBA keyword, object name, property, or method before you press F1.

This chapter consists of several examples that demonstrate common VBA techniques. You may be able to use some examples directly, but in most cases, you must adapt them to your own needs.

Working with Ranges

Most of what you do in VBA probably involves worksheet ranges. When you work with range objects, keep the following points in mind:

- Your VBA code doesn't need to select a range to work with the range.
- If your code does select a range, its worksheet must be active.
- The macro recorder doesn't always generate the most efficient code. Often, you can use the recorder to create your macro and then edit the code to make it more efficient.
- We recommend you use named ranges in your VBA code. For example, a reference such as Range ("Total") is better than Range ("D45"). In the latter case, you need to modify the macro if you add a row above row 45.
- When you record macros that select ranges, pay close attention to relative versus absolute recording mode. The recording mode that you choose can determine whether your macro works correctly.



See Chapter 39, “Introducing Visual Basic for Applications,” for more on recording modes.

- If you create a macro that loops through each cell in the currently selected range, be aware that the user can select entire columns or rows. In such a case, you need to create a subset of the selection that consists only of nonblank cells. Or, you can work with cells in the worksheet’s used range (by using the UsedRange property).
- Be aware that Excel allows you to select multiple ranges in a worksheet. For example, you can select a range, hold down the Ctrl key, and select another range. You can test for this in your macro using the Areas property and take appropriate actions.

The examples in the following sections demonstrate these points.

Copying a range

Copying a range is a frequent activity in macros. When you turn on the macro recorder (using absolute recording mode) and copy a range from A1 : A5 to B1 : B5, you get a VBA macro like this:

```
Sub CopyRange()  
    Range("A1:A5").Select  
    Selection.Copy  
    Range("B1").Select  
    ActiveSheet.Paste  
    Application.CutCopyMode = False  
End Sub
```

This macro works, but it’s not the most efficient way to copy a range. You can accomplish the same result with the following one-line macro:

```
Sub CopyRange2()  
    Range("A1:A5").Copy Range("B1")  
End Sub
```

This code takes advantage of the fact that the Copy method can use an argument that specifies the destination. Useful information about properties and methods is available in the Help system.

Note

Most of the examples in this chapter use unqualified object references. A *qualified object reference* is one where you explicitly tell VBA which object you want by identifying its Parent objects. For example, Range ("A1") is unqualified because we haven’t told VBA which worksheet it’s on. The fully qualified version is Application .Workbooks ("MyBook") .Worksheets ("MySheet") .Range ("A1") .

When you use an unqualified range reference in a standard module, VBA assumes you mean the ActiveSheet in the ActiveWorkbook. This saves you from having to type the whole string of parents if that is indeed what you mean. If you want or need to be explicit, consider using an object variable as described in “Simplifying object references” later in this chapter.

The example demonstrates that the macro recorder doesn't always generate the most efficient code. As you see, you don't have to select an object to work with it. Note that `CopyRange2` doesn't select a range; therefore, the active cell doesn't change when this macro is executed.

Copying a variable-size range

Often, you want to copy a range of cells in which the exact row and column dimensions are unknown.

Figure 44.1 shows a range on a worksheet. This range contains data that is updated weekly. Therefore, the number of rows changes. Because the exact range address is unknown at any given time, you can't hard-code the address in your VBA code.

FIGURE 44.1

The number of rows in this range can be different every time the macro is executed.

	A	B	C
1	Week Ending	Calls	Orders
2	4/5/2022	452	89
3	4/12/2022	546	102
4	4/19/2022	587	132
5	4/26/2022	443	65
6	5/3/2022	609	156
7	5/10/2022	592	92
8	5/17/2022	487	95
9	5/24/2022	601	105
10	5/31/2022	515	133
11	6/7/2022	540	122
12			

The macro that follows demonstrates how to copy this range to Sheet2 (beginning at cell A1). It uses the `CurrentRegion` property, which returns a Range object that corresponds to the block of used cells surrounding a particular cell. This is equivalent to choosing Home ➔ Editing ➔ Find & Select ➔ Go To, clicking the Special button, and then selecting the Current Region option:

```
Sub CopyCurrentRegion()
    Range("A1").CurrentRegion.Copy -
        Sheets("Sheet2").Range("A1")
End Sub
```

Another approach is to use a table to store the data. When you add new rows to a table, the table's range address adjusts automatically, so you can use a procedure like this:

```
Sub CopyTable()
    Range("Table1[#All]").Copy Sheets("Sheet2").Range("A1")
End Sub
```



A workbook that contains these macros is available on this book's website at www.wiley.com/go/excel365bible. The file is named `range_copy.xlsm`.

Selecting to the end of a row or column

You're probably in the habit of using key combinations, such as pressing **Ctrl+Shift+→** and **Ctrl+Shift+↓**, to select from the active cell to the end of a row or column. When you record these actions in Excel (using relative recording mode), you'll find that the resulting code works as you would expect.

The following VBA procedure selects the range that begins at the active cell and extends down to the last cell in the column (or to the cell above the first empty cell, whichever comes first). When the range is selected, you can do whatever you want with it—copy it, move it, format it, and so on:

```
Sub SelectDown()
    Range(ActiveCell, ActiveCell.End(xlDown)).Select
End Sub
```

Notice that the **Range** property has two arguments. These arguments represent the upper-left and lower-right cells in a range.

This example uses the **End** method of the **Range** object, which returns a **Range** object. The **End** method takes one argument, which can be any of the following constants: **xlUp**, **xlDown**, **xlToLeft**, or **xlToRight**.



A workbook that contains this macro is available on this book's website at www.wiley.com/go/excel365bible. The file is named `select_cells.xlsm`.

Selecting a row or column

The macro that follows demonstrates how to select the column of the active cell. It uses the **EntireColumn** property, which returns a range that consists of a column:

```
Sub SelectColumn()
    ActiveCell.EntireColumn.Select
End Sub
```

As you may suspect, an **EntireRow** property also is available, which returns a range that consists of a row.

If you want to perform an operation on all cells in the selected row or column, you don't need to select the row or column. For example, when the following procedure is executed, all cells in the row that contains the active cell are made bold:

```
Sub MakeRowBold()
    ActiveCell.EntireRow.Font.Bold = True
End Sub
```

Moving a range

Moving a range consists of cutting it to the Clipboard and then pasting it to another area. If you record your actions while performing a move operation, the macro recorder generates code as follows:

```
Sub MoveRange()
    Range("A1:C6").Select
    Selection.Cut
    Range("A10").Select
    ActiveSheet.Paste
End Sub
```

As demonstrated with copying earlier in this chapter (see “Copying a range”), this method is not the most efficient way to move a range of cells. In fact, you can do it with a single VBA statement as follows:

```
Sub MoveRange2()
    Range("A1:C6").Cut Range("A10")
End Sub
```

This statement takes advantage of the fact that the Cut method can use an argument that specifies the destination.



A workbook that contains this macro is available on this book's website at www.wiley.com/go/excel365bible. The file is named `range move.xlsm`.

Looping through a range efficiently

Many macros perform an operation on each cell in a range, or they may perform selective actions based on the content of each cell. These operations usually involve a For-Next loop that processes each cell in the range.

The following example demonstrates how to loop through all the cells in a range. In this case, the range is the current selection. In this example, Cell is a variable name that refers to the cell being processed. (Notice that this variable is declared as a Range object.) Within the For-Next loop, the single statement evaluates the cell. If the cell is negative, it's converted to a positive value:

```
Sub ProcessCells()
    Dim Cell As Range
    For Each Cell In Selection.Cells
        If Cell.Value < 0 Then Cell.Value = Cell.Value * -1
    Next Cell
End Sub
```

The preceding example works, but what if the selection consists of an entire column or an entire range? This is not uncommon because Excel lets you perform operations on entire

columns or rows. In this case, though, the macro seems to take forever because it loops through each cell—even those that are blank. What's needed is a way to process only the non-blank cells.

You can accomplish this task by using the SpecialCells method. In the following example, the SpecialCells method is used to create a new object: the subset of the selection that consists of cells with constants (as opposed to formulas). This subset is processed, with the net effect of skipping all blank cells and all formula cells:

```
Sub ProcessCells2()
    Dim ConstantCells As Range
    Dim Cell As Range
    ' Ignore errors
    On Error Resume Next
    ' Process the constants
    Set ConstantCells =
        Selection.SpecialCells(xlConstants, xlNumbers)
    For Each Cell In ConstantCells
        If Cell.Value < 0 Then Cell.Value = Cell.Value * -1
    Next Cell
End Sub
```

The ProcessCells2 procedure works fast, regardless of what is selected. For example, you can select the range, select all columns in the range, select all rows in the range, or even select the entire worksheet. In all of these cases, only the cells that contain constants are processed inside the loop. This procedure is a vast improvement over the ProcessCells procedure presented earlier in this section.

Notice that the following statement is used in the procedure:

```
On Error Resume Next
```

This statement causes Excel to ignore any errors that occur and simply process the next statement. This statement is necessary because the SpecialCells method produces an error if no cells qualify and because the numerical comparison will fail if a cell contains an error value. Normal error checking is resumed when the procedure ends. To return to normal error-checking mode inside a procedure, use the following statement:

```
On Error GoTo 0
```



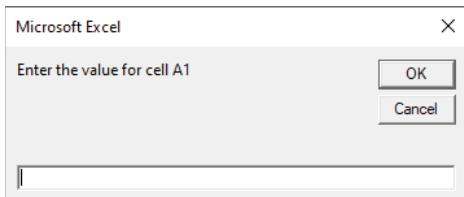
This macro is available on this book's website at www.wiley.com/go/excel365bible. The file is named skip blanks while looping.xlsm.

Prompting for a cell value

As discussed in Chapter 41, “Creating UserForms,” you can take advantage of the VBA InputBox function to ask the user to enter a value. Figure 44.2 shows an example.

FIGURE 44.2

Using the VBA InputBox function to get a value from the user



You can assign this value to a variable and use it in your procedure. Often, however, you want to place the value into a cell. The following procedure demonstrates how to ask the user for a value and place it into cell A1 of the active worksheet, using only one statement:

```
Sub GetValue()
    Range("A1").Value =
        InputBox("Enter the value for cell A1")
End Sub
```

This procedure has a problem, however: if the user clicks Cancel, the contents of cell A1 are replaced with an empty string. Here's a modified version in which the InputBox entry is assigned to a variable named UserVal. The code checks this variable and takes action only if the variable is not empty:

```
Sub GetValue()
    UserVal = InputBox("Enter the value for cell A1")
    If UserVal <> "" Then Range("A1").Value = UserVal
End Sub
```

Here's a variation that accepts only a numeric value. If the user enters a non-number, the InputBox keeps appearing until a number is entered. Only when a number is entered does the code exit the Do Loop, and the value is entered into A1. Another line inside the loop allows the user to click Cancel and get out of the procedure:

```
Sub GetValue()

    Do
        UserVal =
            InputBox("Enter a numeric value for cell A1")
        If UserVal = "" Then Exit Sub
    Loop Until IsNumeric(UserVal)

    Range("A1").Value = UserVal

End Sub
```

Determining the type of selection

If your macro is designed to work with a selected range, you need to determine that a range is actually selected. Otherwise, the macro most likely fails. The following procedure identifies the type of object selected:

```
Sub SelectionType()
    MsgBox TypeName(Selection)
End Sub
```



A workbook that contains this macro is available on this book's website at www.wiley.com/go/excel365bible. The file is named `selection type.xlsm`. For objects that let you enter text, the macro won't work in Edit mode. If you click the button and nothing happens, press the Esc key to exit Edit mode.

If a cell or a range is selected, the `MsgBox` displays `Range`. If your macro is designed to work only with ranges, you can use an `If` statement to ensure that a range is actually selected. The following is an example that displays a message if the current selection is not a `Range` object:

```
Sub CheckSelection()
    If TypeName(Selection) = "Range" Then
        ' ... [Other statements go here]
    Else
        MsgBox "Select a range."
    End If
End Sub
```

Identifying a multiple selection

Excel enables you to make a multiple selection by pressing `Ctrl` while you select objects or ranges. This method can cause problems with some macros; for example, you can't copy a multiple selection that consists of nonadjacent ranges. The following macro demonstrates how to determine whether the user has made a multiple selection:

```
Sub MultipleSelection()
    If Selection.Areas.Count = 1 Then
        ' ... [Other statements go here]
    Else
        MsgBox "Multiple selections not allowed."
    End If
End Sub
```

This example uses the `Areas` property, which returns a collection of all areas in the selection. An area is a `Range` object consisting of contiguous cells. The `Count` property returns the number of objects that are in the collection.

You may want to work with multiple selections rather than just ignore them—for example, if you want to copy the selection. You can loop through the Areas collection of the Range object as this example shows:

```
Sub LoopAreas()
    Dim Area As Range
    For Each Area In Selection.Areas
        'Copy each selection 10 columns to the right
        Area.Copy Area.Offset(0, 10)
    Next Area
End Sub
```

Counting selected cells

You can create a macro that works with the selected range of cells. Use the Count property of the Range object to determine how many cells are contained in a range selection (or any range, for that matter). For example, the following statement displays a message box that contains the number of cells in the current selection:

```
MsgBox Selection.Count
```

CAUTION

With the larger worksheet size introduced in Excel 2007, the Count property can generate an error. The Count property uses the Long data type, so the largest value that it can store is 2,147,483,647. For example, if the user selects 2,048 complete columns (2,147,483,648 cells), the Count property generates an error. Fortunately, Microsoft added a new property, CountLarge, which uses the Double data type that can handle values up to 1.79+E^308.

For more on VBA data types, see Table 44.1.

Bottom line? In most situations, the Count property will work fine. If there's a chance that you may need to count more cells (such as all cells in a worksheet), use CountLarge instead of Count.

44

If the active sheet contains a range named data, the following statement assigns the number of cells in the data range to a variable named CellCount:

```
CellCount = Range("data").Count
```

You can also determine how many rows or columns are contained in a range. The following expression calculates the number of columns in the currently selected range:

```
Selection.Columns.Count
```

And, of course, you can use the Rows property to determine the number of rows in a range. The following statement counts the number of rows in a range named data and assigns the number to a variable named RowCount:

```
RowCount = Range("data").Rows.Count
```

Working with Workbooks

The examples in the following sections demonstrate various ways to use VBA to work with workbooks.

Saving all workbooks

The following procedure loops through all workbooks in the Workbooks collection and saves each file that has been saved previously:

```
Public Sub SaveAllWorkbooks()
    Dim Book As Workbook
    For Each Book In Workbooks
        If Book.Path <> "" Then Book.Save
    Next Book
End Sub
```

Notice the use of the Path property. If a workbook's Path property is empty, the file has never been saved. (It's a new workbook.) This procedure ignores such workbooks and saves only the workbooks that have a nonempty Path property.

Saving and closing all workbooks

The following procedure loops through the Workbooks collection. The code saves and closes all workbooks:

```
Sub CloseAllWorkbooks()
    Dim Book As Workbook
    For Each Book In Workbooks
        If Book.Name <> ThisWorkbook.Name Then
            Book.Close SaveChanges:=True
        End If
    Next Book
    ThisWorkbook.Close SaveChanges:=True
End Sub
```

The procedure uses an If statement within the For-Next loop to determine whether the workbook is the one that contains the code. This is necessary because closing the workbook that contains the procedure would end the code, and subsequent workbooks would not be affected.

Creating a workbook

Use the Add method of the Workbooks collection object to create a new workbook, as in the following example:

```
Workbooks.Add
```

Some methods, like the Add method, return a value. In this case, the returned value is the workbook that was created. You can store the returned value in a variable and use it later in

your code. The following code declares a variable, creates a new workbook and stores it in the variable, and uses that variable to save the new workbook to disk:

```
Dim NewBook As Workbook
Set NewBook = Workbooks.Add
NewBook.SaveAs "test.xlsx"
```

Because the Add method returns a workbook, all the workbook's properties are available to you, such as the Sheets property. If you want to create a new workbook and interact with the first sheet in it, you can put the sheet in a variable rather than the workbook, as in the following example:

```
Dim NewSheet As Worksheet
Set NewSheet = Workbooks.Add.Sheets(1)
NewSheet.Range("A1").Value = Now
```

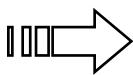
Working with Charts

Manipulating charts with VBA can be confusing, mainly because of the large number of objects involved. To get a feel for working with charts, turn on the macro recorder, create a chart, and perform some routine chart editing. You may be surprised by the amount of code that's generated.

When you understand the way that objects function in a chart, however, you can create some useful macros. The following sections present a few macros that deal with charts. When you write macros that manipulate charts, you need to understand some terminology. An *embedded chart* on a worksheet is a ChartObject object, and the ChartObject contains the actual Chart object. A *chart* on a chart sheet, on the other hand, does not have a ChartObject container.

It's often useful to create an object reference to a chart. (See "Simplifying object references" later in this chapter.) For example, the following statements declare an object variable (MyChart) and assign the embedded chart named Chart 1 on the active sheet to it:

```
Dim MyChart As Chart
Set MyChart = ActiveSheet.ChartObjects("Chart 1").Chart
```



These macros are available on this book's website at www.wiley.com/go/excel365bible. The file is named chart macros.xlsm.

Modifying the chart type

The following example changes the chart type of every embedded chart on the active sheet. It makes each chart a clustered column chart by adjusting the ChartType property of the Chart object. A built-in constant, xlColumnClustered, represents a standard column chart:

```
Sub ChartType()
    Dim ChtObj As ChartObject
    For Each ChtObj In ActiveSheet.ChartObjects
```

```
ChtObj.Chart.ChartType = xlColumnClustered  
Next ChtObj  
End Sub
```

The preceding example uses a For-Next loop to cycle through all the ChartObject objects on the active sheet. Within the loop, the chart type is assigned a new value, making it a clustered column chart.

The following macro performs the same function but works on all chart sheets in the active workbook:

```
Sub ChartType2()  
    Dim Cht As Chart  
    For Each Cht In ActiveWorkbook.Charts  
        Cht.ChartType = xlColumnClustered  
    Next Cht  
End Sub
```

Modifying chart properties

The following example changes the legend font for all charts that are on the active sheet. It uses a For-Next loop to process all ChartObject objects and sets the HasLegend property to True. The code then adjusts the properties of the Font object contained in the Legend object:

```
Sub LegendMod()  
    Dim ChtObj As ChartObject  
    For Each ChtObj In ActiveSheet.ChartObjects  
        ChtObj.Chart.HasLegend = True  
        With ChtObj.Chart.Legend.Font  
            .Name = "Arial"  
            .FontStyle = "Bold"  
            .Size = 8  
        End With  
    Next ChtObj  
End Sub
```

Applying chart formatting

This example applies several formatting types to the specified chart (in this case, Chart 1 on the active sheet):

```
Sub ChartMods()  
    With ActiveSheet.ChartObjects("Chart 1").Chart  
        .ChartType = xlColumnClustered  
        .ChartTitle.Text = "XYZ Corporation"  
        .ChartArea.Font.Name = "Arial"  
        .ChartArea.Font.FontStyle = "Regular"  
        .ChartArea.Font.Size = 9  
        .PlotArea.Interior.ColorIndex = 6
```

```

    .Axes(xlValue).TickLabels.Font.Bold = True
    .Axes(xlCategory).TickLabels.Font.Bold = True
End With
End Sub

```

The best way to learn about the objects, properties, and methods that you need to code for charts is to record a macro while you create or apply various changes to a chart.

VBA Speed Tips

VBA is fast, but it's often not fast enough. The following sections present programming examples that you can use to help speed your macros.

Turning off screen updating

You've probably noticed that when you execute a macro, you can watch everything that occurs in it. Sometimes this view is instructive, but after you get the macro working properly, it can be annoying and slow things considerably.

Fortunately, you can disable the normal screen updating that occurs when you execute a macro. Insert the following statement to turn off screen updating:

```
Application.ScreenUpdating = False
```

If at any point during the macro's execution you want the user to see the results of the macro, use the following statement to turn screen updating back on:

```
Application.ScreenUpdating = True
```

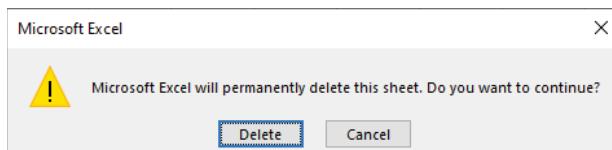
When the macro completes, Excel automatically turns screen updating back on.

Preventing alert messages

One benefit of using a macro is that you can perform a series of actions automatically. You can start a macro and then get a cup of coffee while Excel does its thing. Some operations cause Excel to display messages that must be addressed, however. For example, if your macro deletes a sheet, you see the message that is shown in the dialog box in Figure 44.3. These types of messages mean that you can't execute your macro unattended.

FIGURE 44.3

You can instruct Excel not to display these types of alerts while a macro is running.



To avoid these alert messages (and automatically choose the default response), insert the following VBA statement:

```
Application.DisplayAlerts = False
```

To turn alerts back on, use this statement:

```
Application.DisplayAlerts = True
```

As with screen updating, Excel turns alerts back on when the macro is done.

Simplifying object references

As you may have discovered, references to objects can get lengthy—especially if your code refers to an object that's not on the active sheet or in the active workbook. For example, a fully qualified reference to a Range object may look like this:

```
Workbooks("MyBook.xlsx").Worksheets("Sheet1")  
    .Range("IntRate")
```

If your macro uses this range frequently, you may want to use the Set command to create an object variable. For example, to assign this Range object to an object variable named Rate, use the following statement:

```
Set Rate = Workbooks("MyBook.xlsx")  
    .Worksheets("Sheet1").Range("IntRate")
```

After this variable is defined, you can use the object variable Rate instead of the lengthy reference. Here's an example:

```
Rate.Font.Bold = True  
Rate.Value = .0725
```

Besides simplifying your coding, using object variables speeds your macros quite a bit. We've seen complex macros execute twice as fast after creating object variables.

Declaring variable types

Usually, you don't have to worry about the type of data that's assigned to a variable. Excel handles all these details behind the scenes. For example, if you have a variable named MyVar, you can assign a number of any type to it. You can even assign a text string to it later in the procedure.

If you want your procedures to execute as fast as possible, though, you should tell Excel in advance what type of data is going to be assigned to each of your variables. Providing this information in your VBA procedure is known as declaring a variable's type.

Table 44.1 lists most of the data types that VBA supports. This table also lists the number of bytes that each type uses and the approximate range of possible values.

TABLE 44.1 VBA Data Types

Data Type	Bytes Used	Approximate Range of Values
Byte	1	0 to 255
Boolean	2	True or False
Integer	2	-32,768 to 32,767
Long (long integer)	4	-2,147,483,648 to 2,147,483,647
Single (single-precision floating-point)	4	-3.4E38 to -1.4E-45 for negative values; 1.4E-45 to 4E38 for positive values
Double (double-precision floating-point)	8	-1.7E308 to -4.9E-324 for negative values; 4.9E-324 to .7E308 for positive values
Currency (scaled integer)	8	-9.2E14 to 9.2E14
LongLong	8	-9.2E18 to 9.2E18
Decimal	14	+/-7.9E28 with no decimal point
Date	8	January 1, 100 to December 31, 9999
Object	4	Any object reference
String (variable length)	10 + string length	0 to approximately 2 billion
String (fixed length)	Length of string	1 to approximately 65,400
Variant (with numbers)	16	Any numeric value up to the range of a Double
Variant (with characters)	22 + string length	Same range as for variable-length String
User-defined (using Type)	Number required by elements	Range of each element is the same as the range of its data type

If you don't declare a variable, Excel uses the Variant data type. In general, the best approach is to use the data type that uses the smallest number of bytes yet one that can still handle all the data assigned to it. An exception is when you're performing floating-point calculations. In such a case, it's always best to use the Double data type (rather than the Single data type) to maintain maximum precision. Another exception involves the Integer data type. Although the Long data type uses more bytes, it usually results in faster performance.

When VBA works with data, execution speed is a function of the number of bytes that VBA has at its disposal. In other words, the fewer bytes that data uses, the faster that VBA can access and manipulate the data.

To declare a variable, use the Dim statement before you use the variable for the first time. For example, to declare the variable Units as a Long data type, use the following statement:

```
Dim Units As Long
```

To declare the variable UserName as a string, use the following statement:

```
Dim UserName As String
```

If you declare a variable within a procedure, the declaration is valid only within that procedure. If you declare a variable outside of any procedures (but before the first procedure), the variable is valid in all procedures in the module.

If you use an object variable (as described in “Simplifying object references” earlier in this chapter), you can declare the variable as the appropriate object data type. Here is an example:

```
Dim Rate As Range  
Set Rate = Workbooks("MyBook.xlsx").  
    Worksheets("Sheet1").Range("IntRate")
```

To force yourself to declare all the variables that you use, insert the following statement at the top of your module:

```
Option Explicit
```

If you use this statement and Excel encounters a variable that hasn't been declared, Excel displays an error message. After you get into the habit of correctly declaring all your variables, you'll find that it not only can speed up code execution, but it helps eliminate errors and simplifies spotting errors.

Creating Custom Excel Add-Ins

IN THIS CHAPTER

Understanding add-ins

Converting a workbook to an add-in

For developers, one of the most useful features in Excel is the capability to create add-ins. This chapter discusses this concept and provides a practical example of creating an add-in.

Understanding Add-Ins

An *add-in* provides software with additional functionality. Excel includes several add-ins, including the Analysis ToolPak and Solver. Ideally, the new features blend in well with the original interface so that they appear to be part of the program.

Excel's approach to add-ins is quite powerful: any knowledgeable Excel user can create add-ins from workbooks. The type of add-in covered in this chapter is basically a different form of a workbook file. Any Excel workbook can be converted into an add-in, but not every workbook is a good candidate for an add-in.

What distinguishes an add-in from a normal workbook? Add-ins, by default, have an .xlam extension. In addition, add-ins are always hidden, so you can't display worksheets or chart sheets that are contained in an add-in. However, you can access its VBA procedures and display dialog boxes that are contained on UserForms.

The following are some typical uses for Excel add-ins:

Store one or more custom worksheet functions When the add-in is loaded, you can use the functions as you would any built-in worksheet function.

Store Excel utilities VBA is ideal for creating general-purpose utilities that extend the power of Excel.

Store proprietary macros If you don't want end users to see (or modify) your macros, store the macros in an add-in and protect the VBA project with a password. Users can use the macros, but they can't view or change them unless they know the password. An additional benefit is that the add-in doesn't display a workbook window, which can be distracting.

As previously noted, Excel ships with several useful add-ins, and you can acquire other add-ins from third-party vendors or online. In addition, Excel includes the tools that enable you to create your own add-ins. We explain this process later in this chapter (see "Creating Add-Ins").

NOTE

There are four types of add-ins that work in Excel: VBA, COM, VSTO, and Office Add-ins. COM add-ins are ActiveX DLL files that are compiled instead of scripted like VBA. VSTO add-ins are created in Visual Studio with a set of development tools called Visual Studio Tools for Office. Both COM and VSTO add-ins are still supported, but the platforms aren't actively developed.

The latest add-in type, Office Add-ins, use web technologies like HTML, CSS, JavaScript, and TypeScript. The way an Office Add-in works with Excel is currently limited, but Microsoft is actively developing and improving the platform.

This chapter only deals with VBA add-ins. The other types are beyond the scope of this book.

Working with Add-Ins

The best way to work with add-ins is to use the Excel Add-In Manager. To display the Add-In Manager, follow these steps:

1. Choose **File** \Rightarrow **Options**. The Excel Options dialog box appears.
2. Select the **Add-Ins** category.
3. At the bottom of the dialog box, select **Excel Add-Ins** from the **Manage** list and then click **Go**.

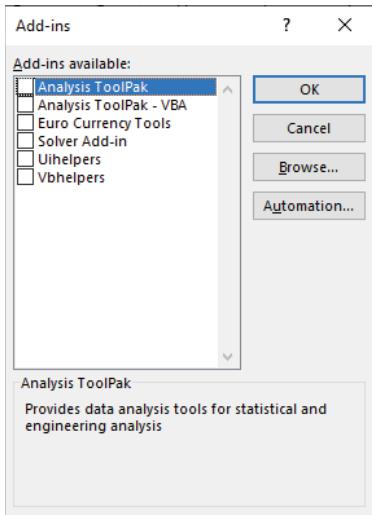
The Add-Ins dialog box appears, as shown in Figure 45.1. The dialog box contains all the add-ins that Excel knows about, which varies from computer to computer. The add-ins that are checked are open. You can open and close add-ins from this dialog box by selecting or deselecting the check boxes.

TIP

Pressing Alt+T followed by I is a much faster way to display the Add-Ins dialog box. Or, if the Developer tab is visible, choose **Developer** \Rightarrow **Add-Ins** \Rightarrow **Excel Add-Ins**.

FIGURE 45.1

The Add-Ins dialog box

**CAUTION**

You can also open most add-in files by opening the file using File \Rightarrow Open. After an add-in is opened, however, you can't choose File \Rightarrow Close to close it. The only way to remove the add-in is to exit and restart Excel or to write a macro to close the add-in. Therefore, you're usually better off opening the add-ins by using the Add-Ins dialog box.

The user interface for some add-ins (including those included with Excel) may be integrated into the Ribbon. For example, when you open the Analysis ToolPak add-in, you access these tools by choosing Data \Rightarrow Analysis \Rightarrow Data Analysis.

NOTE

If you open an add-in created in a version prior to Excel 2007 (an *.xla file), any user interface modifications made by the add-in will not appear as intended. Instead, you must access the user interface items (menus and toolbars) by choosing Add-Ins \Rightarrow Menu Commands or Add-Ins \Rightarrow Custom Toolbars. The Add-Ins tab will show on the Ribbon only if an add-in is loaded that uses the old menu and CommandBar user interface.

Understanding When to Create Add-Ins

Most Excel users have no need to create add-ins. However, if you develop spreadsheets for others—or if you simply want to get the most out of Excel—you may be interested in pursuing this topic further.

Here are some reasons why you may want to convert your Excel workbook application to an add-in:

To avoid confusion If an end user loads your application as an add-in, the file isn't visible in a window and, therefore, is less likely to confuse novice users or get in the way. Unlike a hidden workbook window, an add-in can't be unhidden.

To simplify access to worksheet functions Custom worksheet functions stored in an add-in don't require the workbook name qualifier. For example, if you have a custom function named MOVAVG stored in a workbook named Newfuncs.xlsx, you must use syntax such as the following to use this function in a different workbook:

```
=NEWFUNCS.XLSM!MOVAVG(A1:A50)
```

If this function is stored in an add-in file that's open, however, the syntax is much simpler because you don't need to include the file reference:

```
=MOVAVG(A1:A50)
```

To provide easier access After you identify the location of your add-in, it appears in the Add-Ins dialog box and can display a friendly name and a description of what it does.

To permit better loading control You can automatically open add-ins when Excel starts, regardless of the directory in which they're stored.

To omit prompts when unloading When an add-in is closed, the user never sees the Save Change In prompt because changes to add-ins aren't saved unless you specifically do so from the VB Editor window.

Creating Add-Ins

Technically, you can convert any workbook to an add-in. Not all workbooks benefit from this conversion, though. In fact, workbooks that consist only of worksheets (that is, no macros or custom dialog boxes) become unusable because add-ins are hidden.

Workbooks that benefit from conversion to an add-in are those with macros. For example, you may have a workbook that consists of general-purpose macros and functions. This type of workbook makes an ideal add-in.

The following steps describe how to create an add-in from a workbook:

1. **Develop your application, and make sure that everything works properly.**

2. **(Optional) Add a title and description for your add-in.** Choose File \Rightarrow Info and click Show All Properties at the bottom of the right panel. Enter a brief descriptive title in the Title field and then enter a longer description in the Comments field. Although this step isn't required, it makes installing and identifying the add-in easier.
3. **(Optional) Lock the VBA project.** This step protects the VBA code and UserForms from being viewed. You do this in the Visual Basic Editor (VBE) by choosing Tools \Rightarrow <Project Name> Properties (where <Project Name> corresponds to your VB project name). In the dialog box, click the Protection tab and select Lock Project for Viewing. If you like, you can specify a password to prevent others from viewing your code.
4. **Save the workbook as an add-in file by choosing File \Rightarrow Save As and selecting Excel Add-In (*.xlam) from the Save As Type drop-down list.** By default, Excel saves your add-in to your AddIns directory. You can override this location and choose any directory that you like.

NOTE

After you save the workbook as an add-in, the original (non-add-in) workbook remains active. If you're going to install the add-in and test it, you should close the original workbook file to avoid having two macros with the same name.

After you create the add-in, you need to install it:

1. **Choose File \Rightarrow Options \Rightarrow Add-Ins.**
2. **Select Excel Add-Ins from the Manage drop-down list and then click Go.** The Add-Ins dialog box appears.
3. **Click the Browse button to locate the XLAM file that you created, and click OK, which installs the add-in.** The Add-Ins dialog box uses the descriptive title that you provided for the Title field in the Show All Properties panel.

NOTE

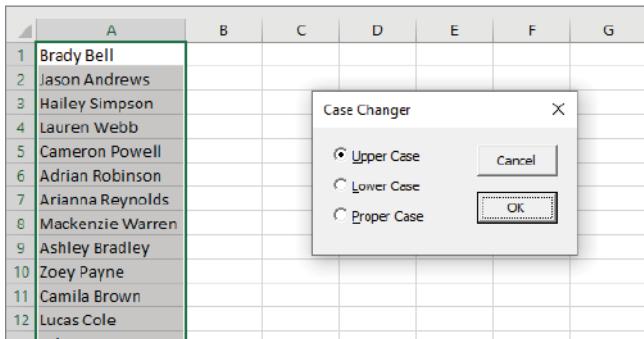
You can continue to modify the macros and UserForms in the XLAM version of your file. Because the add-in doesn't appear in the Excel window, you save your changes in the VB Editor by choosing File \Rightarrow Save.

Looking at an Add-In Example

The following sections discuss the steps required to create a useful add-in from the change case.xlsxm workbook that we covered in Chapter 41, "Creating UserForms." This workbook contains a UserForm that displays options that change the text case of selected cells (uppercase, lowercase, or proper case). Figure 45.2 shows the add-in in action.

FIGURE 45.2

This dialog box enables the user to change the case of text in the selected cells.



This book's website at www.wiley.com/go/excel365bible contains the original version of the workbook (change_case.xlsx) plus a version after it was converted to an add-in (called addin_change_case.xlam). Neither file is locked, so you have full access to the VBA code and UserForm.

This workbook contains one worksheet, which is empty. Although the worksheet is not used, it must be present because every workbook must have at least one sheet. It also contains one VBA module and one UserForm.

Learning about Module1

The Module1 code module contains one procedure that displays the UserForm. The ShowChangeCaseUserForm procedure checks the type of selection. If a range is selected, the dialog box in UserForm1 appears. If anything other than a range is selected, a message box is displayed:

```
Sub ShowChangeCaseUserForm ()  
    If TypeName(Selection) = "Range" Then  
        UserForm1.Show  
    Else  
        MsgBox "Select some cells."  
    End If  
End Sub
```



For information on how to navigate the Visual Basic Editor, including using the Project Explorer to find modules, see Chapter 39, “Introducing Visual Basic for Applications.”

Learning about the UserForm

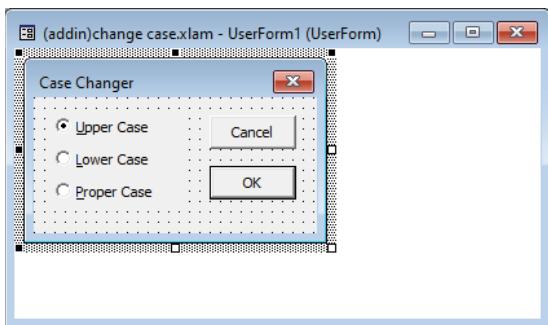
Figure 45.3 shows the UserForm1 form, which has five controls: three OptionButton controls and two CommandButton controls. The controls have descriptive names, and the Accelerator property is set so that the controls display an accelerator key (for keyboard users). The option

button with the Upper-Case caption has its Value property set to TRUE, which makes it the default option.

Refer to Chapter 41 for details about how the code works.

FIGURE 45.3

The custom dialog box



Testing the workbook

Before you convert a workbook to an add-in, test it when a different workbook is active to simulate what happens when the workbook is an add-in. Remember that an add-in is never the active workbook, and it never displays any of its worksheets.

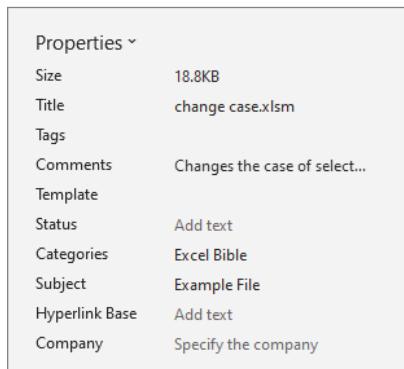
To test it, save the XLSM version of the workbook, close it, and then reopen it. With the workbook open, activate a different workbook, select some cells that contain text, and then press Alt+F8 to display the Macros dialog box. Execute the ShowChangeCaseUserForm macro and try all the options.

Adding descriptive information

Adding descriptive information is recommended but not necessary. Choose File ➔ Info and click Show All Properties at the bottom of the right panel (see Figure 45.4). Enter a title for the add-in in the Title field. This text appears in the Add-Ins dialog box. In the Comments field, enter a description. This information appears at the bottom of the Add-Ins dialog box when the add-in is selected.

FIGURE 45.4

Adding descriptive information about your add-in



Creating the user interface for your add-in macro

At this point, the future add-in is missing one key component: a way to execute the macro that displays the UserForm. The easiest solution is to provide a shortcut key that executes the macro. Ctrl+Shift+C is a good key combination. Here's how to do it:

1. **In Excel, choose Developer \Rightarrow Code \Rightarrow Macros (or press Alt+F8).** The Macro dialog box appears.
2. **In the Macro Name list, select the macro named ShowChangeCaseUserForm.**
3. **Click the Options button.** The Macro Options dialog box appears.
4. **Specify Ctrl+Shift+C as the shortcut key and click OK.**
5. **Click Cancel to close the Macro dialog box.**

Make sure you save the workbook after making this change.

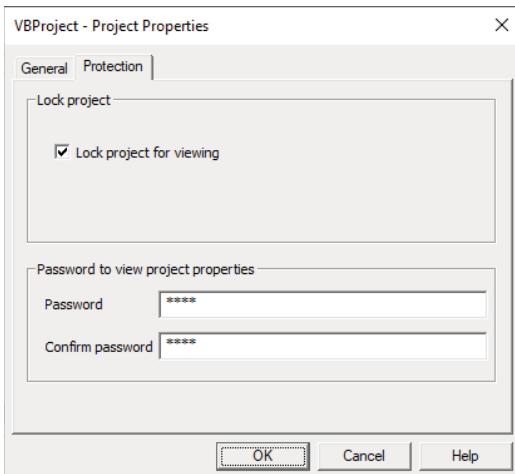
Protecting the project

In some situations (such as a commercial product), you may want to protect your project so that others can't see the source code. To protect the project, follow these steps:

1. **Activate the VBE.**
2. **In the Project window, click the project.**
3. **Choose Tools \Rightarrow <Project Name> Properties.** The VBE displays its Project Properties dialog box.
4. **Select the Protection tab (as shown in Figure 45.5).**

FIGURE 45.5

The Protection tab of the Project Properties dialog box



5. Select the **Lock Project for Viewing** check box.
6. Enter a password (twice) for the project.
7. Click **OK**.

Creating the add-in

To save the workbook as an add-in, follow these steps:

1. Switch to the Excel window and activate your workbook.
2. Choose **File** \Rightarrow **Save As**.
3. Select **Excel Add-In (*.xlam)** from the Save as Type drop-down list.
4. Enter a name for the add-in file and then click **Save**. By default, Excel saves the add-in in your AddIns directory, but you can choose a different directory if you like.

Installing the add-in

Now it's time to try the add-in. Make sure the XLSM version of the workbook is not open and then follow these steps:

1. Choose **File** \Rightarrow **Options** \Rightarrow **Add-Ins**.
2. Select **Excel Add-Ins** from the Manage drop-down list and click **Go**. The Add-Ins dialog box appears.

3. **Click the Browse button. Locate and select the `change_case.xlam` add-in that you just created. Click OK.** The Add-Ins dialog box displays the add-in in its list. Notice that the information that you provided in the Show All Properties panel appears here.
4. **Click OK to close the dialog box and open the add-in.**

When the add-in is installed, you can access it by pressing Ctrl+Shift+C. Another option is to add a new item to your Quick Access Toolbar or to the Ribbon.



See Chapter 8, “Customizing the Excel User Interface,” for information about customizing Excel’s user interface.

Index

Symbols

+ (addition) operator, 213, 214
& (ampersand), entering in text, 193
* (asterisk)
 in condition arguments, 339
 searching for, 82
 as a wildcard, 268
 in workbook names, 163
 in worksheet names, 58
\ (backslash), in worksheet names, 58
: (colon)
 in range addresses, 73, 242
 as range reference operator, 242
 in workbook names, 163
 in worksheet names, 58
, (comma)
 in automatic number formatting, 43
 as list separator, 218
 as range reference operator, 242
' (comment indicator), 888
& (concatenation) operator, 213, 214, 286
{ } (curly brackets)
 in array formulas, 247, 249
 in formula names, 247
/ (division) operator, 213, 214
= (equal sign)
 in condition arguments, 330
 as formula operator, 213
 precedence, 214
! (exclamation point)
 in link formula syntax, 578
 when referencing cells in other workbooks, 228
 when referencing cells in other worksheets, 228
^ (exponentiation) operator, 213, 214, 342
> (greater than) operator
 in condition arguments, 330
 as formula operator, 213, 214
 in workbook names, 163

>= (greater than or equal to) operator, 213
(hash) marks
 cells filled with, 47, 69, 381
 in formula errors, 236
<= (less than or equal to), 213
< (less than) symbol
 as formula operator, 213, 214
 in workbook names, 163
- (minus sign), 574
* (multiplication) operator, 213, 214, 918
- (negation) operator, 386, 387
<> (not equal to) operator, 213
() (parentheses), 273
+ (plus sign), in worksheet outlines, 574
? (question mark)
 in condition arguments, 339
 searching for, 82
 in workbook names, 163
 in worksheet names, 58
" (quote), in workbook filenames, 163
; (semicolon)
 in custom number formats, 47
 as list separator, 218
 in number formatting strings, 492
 when suppressing zeros, 496
/ (slash)
 in workbook names, 163
 in worksheet names, 58
[] (square brackets)
 in column names, 233, 236
 in link formula syntax, 578
 when applying custom format colors, 496
 in workbook names, 228
 in worksheet names, 58
- (subtraction) operator, 213, 214, 387
~ (tilde) wildcard, 268
| (vertical bar)
 as insertion point, 34
 in workbook names, 163

A

About Excel command, 863
ABS function, 274
absolute references
 about, 141–142, 225–227
 naming formulas, 241
.accdb file extension, 522
accelerator keys, 19
Accelerator property
 CheckBox control, 942
 OptionButton control, 945
Access files, 522
Accounting number format, 47
Accounting Number Format button, 490
Accounting Number Format dropdown list, 43
Action Pen, 51
Activate workbook event, 952
Activate worksheet event, 955
activating
 charts, 404
 Power Pivot Ribbon interface, 724
 worksheets, 56
active cells, 4–7
ActiveX controls
 adding to worksheets, 938
 CheckBox, 942
 code module for, accessing, 903
 ComboBox, 942–943
 CommandButton, 943
 Design mode, 938
 vs. Forms controls, 935–937
 Image, 944
 Label, 944
 linking to cells, 940
 ListBox, 944
 macros for, creating, 941–942
 OptionButton, 945
 properties
 adjusting, 938–939
 shared, 939–940
 ScrollBar, 945–946
 SpinButton, 946
 TextBox, 946–947
 ToggleButton, 947
Add a Digital Signature command, 720
Add Chart Element control, 416

Add Constraint dialog box, 681–682
Add Elements button
 Data Labels, 465
 Error Bars, 468
 Trendline, 469
Add Scenario dialog box, 663–664
Add to Quick Access Toolbar, 201–202
Add-In Manager, 980
add-ins
 adding descriptive language, 985–986
 closing, 980–981
 creating, 982–983
 creating UI for add-in macros, 986
 installing, 983
 opening, 980–981
 overview, 979
 protecting project, 986–987
 reasons for creating, 982
 saving workbooks as, 983, 987
 Solver (*See Solver*)
 testing workbooks, 985
 UserForm example, 983–988
 uses, 979–980
Add-Ins dialog box, 980–981, 985
Add-Ins tab, 10
addition (+) operator, 213, 214
Advanced tab (Excel Options dialog box), 34
AfterSave workbook event, 952
age, calculating, 308–309
aggregating data
 back-end databases and, 863
 calculated measures, 770–773
 conditional calculations, 336–345
 with Power Query, 830–832
 using DAX-driven calculated columns, 762–770
alert messages, preventing, 975–976
aligning cell contents, 127–131
Alignment tab, Format Cells dialog box, 127–131
All merging conditional formats option, Paste Special dialog box, 93
All using Source Theme option, Paste Special dialog box, 93
Allow Edit Ranges command, 713
Allow User to Edit Ranges dialog box, 713
Alt key, 19
ampersand (&), entering in text, 193
analysis of variance tool, Analysis ToolPak, 699–700

- Analysis ToolPak
 analysis of variance tool, 699–700
 correlation tool, 700–701
 covariance tool, 701
 descriptive statistics tool, 701
 exponential smoothing tool, 701
 Fourier analysis tool, 702
 F-test tool, 701–702
 histogram tool, 703–704
 installing, 698
 moving average tool, 704–705
 overview, 697–698
 random number generation tool, 705–706
 rank and percentile tool, 706
 regression tool, 706–707
 sampling tool, 707
 T-test tool, 707–708
 Z-test tool, 708
- Analyze Data feature, artificial intelligence (AI)
 and, 668–673
- AND function, 333–336
- appending Power Query data, 841–845
- application events
 OnKey, 959, 961–962
 OnTime, 959, 960–961
 overview, 949, 959–962
- Application object, 888–889, 890
- area charts, 428–429
- arguments
 AVERAGE function, 224
 AVERAGEIF function, 344
 AVERAGEIFS function, 345
 CalcTax function, 904
 CHOOSE function, 320
 CONVERT function, 283–284
 COUNT function, 282–283
 COUNTA function, 282–283
 COUNTBLANK function, 282–283
 COUNTIF function, 561
 creating functions
 with no arguments, 905–906
 with one argument, 906–907
 with a range argument, 908–909
 with two arguments, 907–908
 as wrapper for built-in VBA functions, 909–910
- DATE function, 315–316, 318
- DATEDIF function, 308–309
 DOLLAR function, 299
 EOMONTH function, 318
 FIND function, 291–292
 AND function, 335–336
 HLOOKUP function, 352–353
 IF function, 278, 330
 IFERROR function, 353
 INDEX function, 352
 INDIRECT function, 332, 357–359
 LEFT function, 288, 290, 814
 MATCH function, 351–352, 356–357
 MID function, 291, 814
 NETWORKDAYS function, 310
 NETWORKDAYS.INTL function, 310–311
 NumSign function, 900
 overview, 217–218
 RIGHT function, 288, 291, 814
 ROUND function, 279, 280–282
 SERIES function, 464–465
 SUBSTITUTE function, 294
 SUBTOTAL function, 233
 SUMIF function, 336–337
 SUMIFS function, 339
 syntax errors, 380
 TEXT function, 298
 TEXTJOIN function, 287, 540
 TIME function, 327
 VLOOKUP function, 350
 WORKDAY.INTL function, 312
- arithmetic mean, 344
- Arrange Windows dialog box, 54
- arrays, selecting, 80
- artificial intelligence (AI), analyzing data
 with, 668–673
- Artistic Effects command, 516
- asterisk (*)
 in condition arguments, 339
 searching for, 82
 as a wildcard, 268
 in workbook names, 163
 in worksheet names, 58
- @ symbol, 255
- auditing tools
 Excel's background error-checking feature, 391–392
 finding cells by type, 388–389

- auditing tools (*continued*)
 - Formula Evaluator, 392–393
 - tracing cell relationships, 389–390
 - tracing error values, 391
 - viewing formulas, 389
- Auto Outline, 567, 572
- AutoComplete, 37–38
- AutoCorrect
 - overview, 39, 396–398
 - and syntax errors, 381
- AutoFill
 - copying to adjacent cells, 88
 - overview, 37
- AutoFit Column Width, 69
- AutoFit Row Height, 70
- AutoRecover
 - adjusting, 26
 - configuring, 165
 - overview, 163–164
 - recovering unsaved work, 164–165
 - recovering versions of current workbook, 164
- AutoSize property, 940
 - Image control, 944
 - TextBox control, 947
- auto-updating Sparklines, 486
- AVERAGE function, 224
- AVERAGEIFS function, 344–345
- averages. *See* means
- axes
 - category axes, 405
 - for bubble charts, 430
 - modifying, 456–460
 - for radar charts, 429
 - for surface charts, 431
 - titles, 449
 - for XY charts, 427
- labels, positioning, 456
- value axes, 406
 - in bubble charts, 430
 - in combination charts, 470–471
 - modifying, 452–456
 - in stock charts, 431
- Azure databases, connecting to Power Query, 797
- B**
- BackColor property, 940
- background color, of sheet tabs, changing, 58
- background images
 - displaying, 516
 - printing, 190
- backslash (\), in worksheet names, 58
- Backstage view, 10, 26
- BackStyle property, 940
- backups
 - AutoRecover, adjusting, 26
 - overview, 168–169
- banded rows, 108, 109, 119
- bar charts, 423–424
- BASE function, 538
- BeforeClose workbook event, 952, 954–955
- BeforeDoubleClick worksheet event, 955
- BeforePrint workbook event, 952
- BeforeRightClick worksheet event, 955, 959
- BeforeSave workbook event, 952, 954
- Bin ranges, histograms and, 703–704
- BIN2DEC function, 538
- binary values, converting to decimal, 538
- bitmap images, 515
- blank cells, selecting, 80
- blank fields, filling in, 808
- Blank Workbook command, 20, 157, 158
- Blanks option, Go To Special dialog box, 80
- bold, applying, 126
- Boolean data type, 977
- Border tab, Format Cells dialog box, 132–134
- borders
 - adding, 132–134
 - simulating gridlines with, 189–190
- bottom text alignment, 129
- BoundColumn property
 - ComboBox control, 943
 - ListBox control, 944
- box & whisker charts, 434–436
- bubble charts, 430–431
- business intelligence. *See* Power Pivot
- buttons
 - chart customization buttons, 445–446
 - simple buttons, 12
 - spin buttons, 12
 - split buttons, 12
 - toggle buttons, 12
- Byte data type, 977

C

- calculated columns
 - creating, 758–759
 - formatting, 759–760
 - hiding from end users, 760–762
 - overview, 757
 - referencing in other calculations, 760
- calculated measures
 - creating, 770–772
 - deleting, 773
 - editing, 773
 - overview, 770–772
- calculator, Formula bar as, 244
- calendar quarters, calculating, 318–319
- Camera tool, 511
- Capitalize Each Word command, 811
- Caption property, 940
- category axis
 - for bubble charts, 430
 - overview, 405, 406
- category_labels argument, Series function, 464
- CEILING function, 280
- cell dependents, 389–390
- cell precedents, 389–390
- cell references
 - absolute references, 225–227, 241
 - to cells outside the worksheet, 228–229
 - changing types, 227–228
 - defined, 212
 - invalid, 385
 - mixed references, 225–227
 - names, applying, 243
 - overview, 225–229
 - in PivotTables, 638–640
- Cell Styles command, 146–150
- cells. *See also* ranges
 - active cells, 4–7
 - adding text, 548–549
 - aligning contents, 127–131
 - blank cells, 381–382
 - comments
 - deleting, 33
 - editing, 102–103
 - inserting, 100
 - replying to, 102
 - resolving threads, 103–104
 - showing, 101–102
- copying
 - to adjacent cells, 88
 - containing conditional formatting, 145–146
 - to other worksheets, 89
 - overview, 83–84
 - paste options, 91–95
 - using drag-and-drop, 87
 - using Office Clipboard, 89–91
 - using Ribbon commands, 84
 - using shortcut keys, 85–86
 - using shortcut menu commands, 85, 86
- data types, 27–29
- defined, 73
- deleting contents, 32–33
- editing
 - color-codes, 386
 - contents, 33–35
 - filled with hash (#) marks, 381
 - filling gaps, 545–547
 - finding, 388–389
 - gridlines, printing, 189–190
- highlighting
 - based on value of another cell, 367–369
 - with values lower than hard-coded value, 365–367
- input messages, displaying, 564
- inserting, 67
- merging, 129–130
- named styles, 146–150
- navigating, 183–191
- notes
 - adding to, 104–108
 - alternatives to, 107
 - deleting, 108
 - editing, 108
 - formatting, 106–107
 - showing, 105–106
- preventing from printing, 195–196
- removing text, 547–548
- replacing contents, 33
- replacing text, 549–550
- selected cell indicator, 5, 6
- tracing relationships, 389–390
- unique address, 4
- unlocking, 710–712

- center text alignment, 128, 129
- Change Chart Type dialog box, 410–411
- Change worksheet event, 955, 956
- changing cells, Scenario Manager
 - entering in Add Scenario dialog box, 664
 - maximum number per scenario, 667
 - modifying scenarios, 666
 - naming, 663
 - overview, 661
- CHAR () function, 295–296
- chart area, 406, 407, 447
- Chart Elements button
 - chart titles, adding, 449
 - gridlines, adding/removing, 452
 - overview, 445
- Chart Elements control, 442, 443–444
- Chart Filters button, 445
- chart sheets
 - converting chart on to embedded chart, 415
 - converting embedded chart to chart sheet, 415
 - copying charts on, 416
 - deleting, 416
 - overview, 404–405
- Chart Styles button, 445
- chart templates, creating, 473–474
- chart title, 406
- Chart Tools
 - Format, Insert Shapes, 451
 - Shape Styles, 446, 452
- Chart Tools Design
 - Chart Layouts
 - Add Chart Element, 416
 - Quick Layout, 412
 - Chart Styles, 413
 - Data, Switch Row/Column, 410
 - Location, Move Chart, 415
 - Type, Change Chart Type, 410
- Chart Tools Format
 - overview, 413–414
 - Shape Styles
 - Shape Effects, 414
 - Shape Fill, 413
 - Shape Outline, 413
 - Size, 415
- charts
 - activating, 404
 - adding elements, 413, 416
 - area charts, 428–429
 - bar charts, 423–424
 - box & whisker charts, 434–436
 - bubble charts, 430–431
 - column charts, 421–423
 - copying, 416
 - copying formatting, 417–418
 - creating
 - based on template, 473–474
 - basic steps, 408–409
 - combination charts, 470–472
 - customizing
 - adding labels, 465–467
 - adding trendlines, 468–470
 - category axis, 456–460
 - Chart Area, 447
 - Plot Area, 448
 - selecting elements, 441–444
 - using chart customization buttons, 445–446
 - using the Format task pane, 444–445
 - using the Mini toolbar, 446
 - using the Plot Area, 448
 - using the Ribbon, 446
 - deleting elements, 413, 416
 - elements of, 405–407
 - embedded (See embedded charts)
 - formatting
 - applying, 974–975
 - copying formatting, 451
 - overview, 413–414, 416–417
 - funnel charts, 437–438
 - histogram charts, 431–433
 - how Excel handles, 402–403
 - Insert Statistic Chart, 433
 - layout, applying, 412
 - legends, 406, 450–451
 - line breaks, adding, 295–296
 - line charts, 424–426
 - map charts, 438–439
 - moving, 414–415
 - overview, 401–402
 - Pareto charts, 433–434
 - pie charts, 426–427
 - PivotCharts, 640–64406
 - preventing from printing, 196–197

- printing, 418–419
properties, modifying, 974
radar charts, 429–430
renaming, 418
resizing, 414–415
Sparklines (See Sparklines)
stock charts, 431
styles, applying, 413
sunburst charts, 436
surface charts, 430
titles, 449
treemap charts, 437
types
 modifying, 410–411, 973–974
 overview, 419
 selecting, 419–421
waterfall charts, 434, 435
XY charts, 427–428
- Check Accessibility command**, 718
check box controls, 12
Check Compatibility command, 718
CheckBox ActiveX control, 942
checkerboard shading, creating, 144
CHOOSE function, 319–320, 348
circular reference errors, 238, 380, 391
CLEAN function, 296–297, 538
cleaning data
 adding text to cells, 548–549
 changing case of text, 536
 checklist for, 549–550
 classifying values, 538–540
 converting values, 538
 converting vertical data to horizontal data, 543–545
 extracting filenames from URLs, 541–542
 filling gaps in imported reports, 545–547
 fixing trailing minus signs, 549
 identifying duplicate rows, 530–531
 joining columns, 540–541
 matching text in lists, 542–543
 overview, 529
 randomizing rows, 541
 rearranging columns, 541
 removing duplicate rows, 529–530
 removing extra spaces, 537
 removing strange characters, 538
 removing text in cells, 547–548
 replacing text in cells, 547–548
 spell-checking, 547
 splitting text, 531–536
 from text fields, 296–297
cleaning text, 812–813
Clear All command, 146
Clear Formats command, 146
Clear Outline, 570
Clear Rules command, 146
Close button, 20
closing
 add-ins, 980–981
 task panes, 20
 UserForms, 924
 windows, 55–56
workbooks, 168
 BeforeClose event, 954–955
 using VBA, 972
- Collapse the Ribbon button**, 5
collections (of objects), 889
colon (:)
 in range addresses, 73, 242
 as range reference operator, 242
 in workbook names, 163
 in worksheet names, 58
- color**
 color scale conditional formatting, 137–138
 custom format colors, applying, 496–497
 document themes, customizing, 153–155
 formatting worksheets, 131–132
 sheet tab color, changing, 58
- column charts**, 421–423
- Column Differences option**, Go To Special dialog box, 80
- Column Distribution option**, 800
- Column Profile option**, 800
- Column Quality option**, 800
- Column Sparkline**, 476
- Column Width dialog box**, 69
- Column widths option**, Paste Special dialog box, 93
- ColumnCount property**
 ComboBox control, 943
 ListBox control, 944
- columns**
 column letters, 5
 concatenating, 809–811
 custom columns

- columns (*continued*)
 - adding, 823–824
 - concatenating, 824–826
 - conditional logic, 829–830
 - data type conversions, 826
 - transformation functions, 826–829
 - deleting, 68
 - duplicate
 - identifying, 530–531
 - removing, 529–530
 - headers, printing, 190
 - hiding, 70–71
 - inserting, 66–68
 - joining, 540–541
 - orientation, switching, 410
 - pivoting, 821–823
 - rearranging, 541
 - selecting, 66
 - splitting, 816–819
 - titles, printing, 188–189
 - unhiding, 70–71
 - unpivoting, 819–821
 - width, changing, 68–69
- combination charts, creating, 470–472
- ComboBox ActiveX control, 942–943
- comma (,)
 - in automatic number formatting, 43
 - as list separator, 218
 - as range reference operator, 242
- Comma Style button, 43, 490
- CommandButton ActiveX control, 943
- commands
 - adding to Quick Access toolbar, 16–17
 - not on the Ribbon, 16–17
 - Quick Access toolbar, 199–203
 - screen tips, 12
 - searching for, 14
 - touch-related, 7
- comment indicator ('), 888
- comments
 - in cells
 - deleting, 33
 - inserting, 100
 - replying to, 102
 - resolving threads, 103–104
 - showing, 101–102
 - in macros, 874
 - in VBA code, 888
- Comments and Notes option, Paste Special dialog box, 93
- Compare Side by Side feature, 63
- comparing worksheets, 62–63
- compatibility, dynamic arrays and, 252
- CONCAT function, 540–541
- concatenating columns, 809–811
- concatenation (&) operator, 213, 214, 286, 918
- conditional analysis
 - checking for multiple conditions, 330–331, 333–336
 - checking for simple conditions, 329–330
 - counting values that meet certain condition, 341–342
 - counting values that meet two or more conditions, 342–344
 - getting averages of numbers, 344–345
 - overview, 329
 - summing values between date range, 340–341
 - summing values that meet certain condition, 336–339
 - summing values that meet two or more conditions, 339–340
 - validating conditional data, 331–333
- Conditional Formats option, Go To Special dialog box, 81
- conditional formatting
 - absolute references, 141–142
 - applying, 135
 - color scales, 137–138
 - copying cells containing, 145–146
 - data bars, 135–137
 - deleting, 146
 - examples, 142–144
 - formula-based rules, 139–142
 - graphical formats, 135–139
 - highlighting cells
 - based on value of another cell, 367–369
 - with values lower than hard-coded value, 365–367
 - highlighting values
 - based on dates, 372–377
 - in one list and another, 371–372
 - in one list but not another, 369–371
 - icon sets, 138–139

- locating cells containing, 146
- managing rules, 145
- overview, 135
- relative references, 141–142
- selecting cells containing, 81, 146
- Conditional Formatting Rules Manager dialog box, 135, 145
- Confidence Level option, Regression dialog box, 707
- Consolidate dialog box, 588–589
- consolidating worksheets
 - example, 590–592
 - overview, 585–594
 - refreshing consolidation, 592–593
 - sources, 593–594
 - using external reference formulas, 587
 - using the Consolidate dialog box, 588–589
 - using the Paste Special dialog box, 587–588
- Constant Is Zero option, Regression dialog box, 707
- constants
 - defined, 212
 - named, 240
- Constants option, Go To Special dialog box, 80
- Constraint Precision option, Solver Options dialog box, 686
- contextual tabs, 11
- controls
 - Forms controls, 936
 - UserForm controls
 - accelerator keys, 932–933
 - adding, 919–920
 - changing properties of, 920–921
 - tab order, 933
- conversion functions, 283–284
- CONVERT function
 - unit conversions, 283–284
 - value conversions, 538
- Convert Text to Columns Wizard, 532–533
- copy and paste
 - cells/ranges
 - to adjacent cells, 88
 - from cell to cell, 33
 - containing conditional formatting, 145–146
 - to the Office Clipboard, 89–91
 - to other worksheets, 89
 - overview, 83–84, 964–965
 - paste options, 91–95
- using drag-and-drop, 87
- using Paste Options buttons, 86–87
- using Ribbon commands, 84
- using shortcut keys, 85–86
- using shortcut menu commands, 85
- using the Paste Special dialog box, 92–95
- chart formatting, 417–418, 451
- charts, 417–418
- formulas, 244–245
- overview, 529
- page setup settings, 195
- ranges, 964–965
- worksheets, 58–59
- Correlation dialog box, 700–701
- correlation tool, Analysis ToolPak, 700–701
- COUNT function, 282–283
- COUNTA function, 282–283
- COUNTBLANK function, 282–283
- COUNTIF function
 - accepting non-duplicate entries, 561
 - accepting text beginning with specific character, 561–562
- overview, 341–343, 561
- COUNTIFS function, 342
- counting
 - characters in text string, 294–295
 - values in a range, 282–283
- COVAR function, 701
- covariance tool, Analysis ToolPak, 701
- Create Names from Selection dialog box, 97–98
- Create PDF/XPS command, 717
- Create PivotTable dialog box, 602–603
- Create Sparklines dialog box, 477–478
- CSV files
 - exporting data to Power Query, 795
 - exporting to, 550
- overview, 523
- Cube functions, 773–775
- curly brackets ({})
 - in array formulas, 247, 249
 - in formula names, 247
- currency
 - applying regional currency format to text, 299
 - Currency data type, 977
 - Currency number format, 47
- Current Array option, Go To Special dialog box, 80

- Current Region option, Go To Special dialog box, 80
cursor, in dialog box controls, 19
custom error bars, 468
Custom number format, 47
Custom Views, 197–198
Customize Ribbon tab, Excel Options dialog box, 205–207
customizing
charts
 adding labels, 465–467
 adding trendlines, 468–470
 category axis, 456–460
 Chart Area, 447
 Plot Area, 448
 selecting elements, 441–444
 using chart customization buttons, 445–446
 using the Format task pane, 444–445
 using the Mini toolbar, 446
 using the Ribbon, 446
document themes, 153–155
Quick Access toolbar
 adding commands, 201–202
 adding Ribbon commands, 16–17
 overview, 16–17, 199–204
 rearranging icons, 203
 removing icons, 204
 resetting to default state, 204
 sharing customizations, 204
 for specific workbooks, 203
Ribbon
 adding commands to groups, 207
 creating new groups, 206
 creating new tabs, 206
 reasons for, 205
 restoring to default state, 208
 sharing customizations, 204
 what can be customized, 205
- D**
- d time unit code, DATEDIF function, 309
Data
 Analyze
 Data Analysis, 698
 Solver, 679, 680
 Data Tools
 Data Validation, 554–555
 Flash Fill, 533–536
Relationships, 647
Text to Columns, 532–533
Forecast, What-If Analysis, 655, 663
Get Data
 From Azure, 797
 Data Source Settings, 798
 From Database, 796
 From File, 794, 795
 From Other Sources, 797, 798
Queries & Connections, 786–7886
data analysis
 Analysis ToolPak
 analysis of variance tool, 699–700
 correlation tool, 700–701
 covariance tool, 701
 descriptive statistics tool, 701
 exponential smoothing tool, 701
 Fourier analysis tool, 702
 F-test tool, 701–702
 histogram tool, 703–704
 installing, 698
 moving average tool, 704–705
 overview, 697–698
 random number generation tool, 705–706
 rank and percentile tool, 706
 regression tool, 706–707
 sampling tool, 707
 T-test tool, 707–708
 Z-test tool, 708
 with artificial intelligence (AI), 668–673
DAX (Data Analysis Expression) functions
 building DAX-driven calculated columns, 762–770
 identifying for calculated columns, 762–770
 nesting, 770
 referencing fields from other tables, 768–770
Data Analysis dialog box, 698–699
data bars, displaying, 135–137
data body, 109
data entry
 AutoComplete, 37–38
 AutoCorrect, 39
 AutoFill, 37
 data forms, 40–41
 dates, 31–32
 Enter mode, 31
 numbers

- decimal points, 36
formatting, 42–47
fractional values, 40
overview, 30
text, 30–31
times, 32
- Data Interchange Format, exporting to, 551
- data labels, 406, 472
- data models
for PivotTables, 644–650
Power Pivot data model
adding tables, 727–728
feeding, 747–754
loading data from flat files, 738–742
managing queries and connections, 755
managing relationships, 754–755
preparing tables, 726
referencing fields from other tables, 768–770
refreshing, 742–745
using Cube functions in, 773–775
- data points, 406
- data profiling, with Power Query, 800–803
- data series
adding to charts, 462
in area charts, 428
in bar charts, 424
in bubble charts, 430
changing data used by, 462–465
in column charts, 421
combination charts, creating, 470–472
data tables, displaying, 472–473
deleting, 461–462
error bars, adding, 468
hiding, 461–462
identifying, 406
labels, adding, 465–467
limitations, 408
in line charts, 424
missing data, handling, 467–468
moving averages, 704–705
overview, 402–403, 460–461
in pie charts, 421, 426
in radar charts, 429
in stock charts, 431
in surface charts, 430
trendlines, adding, 468–470
- Data Source Settings command, 798
- Data tab, 9
- Data Table dialog box, 655–656
- data tables
adding to charts, 472–473
creating, 653–661
one-input data tables, 654–657
overview, 653–654
two-input data tables, 657–661
- data types
error values, 29
formulas, 29
numeric values, 27–28
overview, 27–29
text, 28
VBA data types, 976–978
- data validation
cell references, 559–560
criteria, specifying, 554–557
drop-down lists, creating, 557–558
formulas for
accepting dates by day of week, 562
accepting larger value than previous cell, 561
accepting non-duplicate entries only, 561
accepting text beginning with specific character, 561–562
accepting text only, 561
accepting value not exceeding a total, 563
creating dependent lists, 563–564
overview, 558–559
overview, 553–554
without restricting data entry, 564–565
- Data Validation dialog box
Error Alert tab, 555, 565
Input Message tab, 555
overview, 554–555
Settings tab, 554–555, 555–556, 558–559
- Data Validation option, Go To Special dialog box, 81
- databases
Azure databases, connecting to Power Query, 797
back-end databases, 863
file formats, 522–523
relational databases
connecting to Power Query, 796
loading data into Power Pivot data model, 733–746

Index

- Date category axis, 459
Date data type, 977
DATE function
 converting dates to Julian format, 315–316
 returning last day of month, 317–318
Date number format, 47
DATEDIF function
 calculating age, 308–309
 calculating days between two dates, 309–310
 calculating number of years and months
 between dates, 314–315
date(s)
 age, calculating, 308–309
 business days excluding holidays, listing, 311–313
 calendar quarters, calculating, 318–319
 current, returning, 308
 custom number formatting, 497–498
 entering, 302–303
 extracting parts of, 313–314
 fiscal month, returning, 320–321
 fiscal quarter, calculating, 319–320
 formatting, 305–306
 highlighting cells based on, 372–377
 Julian format, converting to, 315–316
 last day of month, returning, 317–318
 last weekday of month, calculating, 322–323
 of Nth weekday, calculating, 321–322
 number of days between, calculating, 309–310
 number of years and months between,
 calculating, 314–315
 percent of year completed/remaining,
 calculating, 316–317
 problems with, 306–307
 regional settings and, 32
 searching for, 303
 serial numbers, 301–302
 workdays between, calculating, 310–311
date-stamping worksheets, 41
DAX (Data Analysis Expression) functions
 building DAX-driven calculated columns, 762–770
 identifying for calculated columns, 762–764
 nesting, 770
 referencing fields from other tables, 768–770
DAY function, 313–314
dBase files, 522
Deactivate worksheet event, 955
Deactivate workbook event, 952
debugging custom functions, 910
DEC2BIN function, 538
DEC2OCT function, 538
Decimal data type, 977
decimals
 converting hours/minutes/seconds to a time, 326
 converting to other number bases, 538
 entering decimal points automatically, 36
 precision problems, 387–388
Decrease Decimal button, 23, 490
default
 column width, 69
 font, 124
 print settings, 17
 Quick Access toolbar tools, 200
 row height, 69
 templates, 172–174
 worksheet names, 57
Define Name command, 220, 240
Defined Names
 Create from Selection, 242
 Define Name, 220, 240
Delete Columns option, Protect Sheet dialog box, 712
Delete key, 32
Delete Rows option, Protect Sheet dialog box, 712
Delete Sheet Columns command, 68
Delete Sheet Rows command, 68
deleting
 calculated measures, 773
 cell contents, 32–33
 chart elements, 413
 chart sheets, 416
 columns, 68
 comments from cells, 33
 conditional formatting, 146
 data series, 461–462
 embedded charts, 416
 hyperlinks from cells, 33
 rows, 68
 worksheets, 57
dependent workbooks. *See also* linking workbooks
 intermediary links issue, 585
 Save As command issue, 584
dependents, selecting, 80
Dependents option, Go To Special dialog box, 80

- descriptive statistics tool, Analysis ToolPak, 701
 Design mode, 938
 Developer tab, 10
 Code
 Macro Security, 871
 Record Macro, 876
 Stop Recording, 877
 Use Relative References, 880
 Controls, Insert, 936, 938
 displaying on Ribbon, 870, 936
 DEX2HEX function, 538
 dialog box launchers, 12
 dialog boxes
 displaying
 using InputBox function, 914–915
 using MsgBox function, 915–918
 hot keys, 19
 modal, 18
 modeless, 18
 navigating, 18–1920
 tabbed, 19–20
 UserForms (*See* UserForms)
 DIF files, 523, 551
 digital signatures, adding to workbooks, 719–720
 direct precedents, 389
 distributed text alignment, 128, 129
 #DIV/0! errors, 237, 383
 divide-by-zero errors, 277–278
 division (/) operator, 213, 214
 Do loops, 894
 Document Inspector dialog box, 718–719
 document themes
 applying, 152–153
 customizing, 153–155
 overview, 150–152
 DOLLAR function, 299
 Double data type, 977
 drag-and-drop
 copying cells/ranges, 87
 moving cells/ranges, 87
 Draw Border command, 132
 Draw Border Grid command, 132
 drawing layer, 4
 Drawing Tools, Format
 Arrange, 506
 Insert Shapes, 506
 Shape Styles, 506
 WordArt Styles, 506
 drop-down lists, creating, 557–558
 duplicate records, removing, 530, 805–807
 dynamic arrays
 compatibility and, 252
 functions
 FILTER, 262–265
 overview, 255
 RANDARRAY, 259–260
 SEQUENCE, 260–261
 SORT, 256–257
 SORTBY, 257–258
 UNIQUE, 258–259
 XLOOKUP, 265–269
 overview, 250–251
 spill ranges, 252–255
 dynamic labels, creating, 510
 dynamic models, 651. *See also* what-if analysis

E

- Edit mode
 changing to Enter mode, 36
 entering, 33–35
 Edit Objects option, Protect Sheet dialog box, 712
 Edit Points, 509
 Edit Scenarios option, Protect Sheet dialog box, 712
 Edit Series dialog box, 462, 463–464
 Edit Sparklines dialog box, 486
 editing
 calculated measures, 773
 cells
 cell contents, 33–35
 color-codes, 386
 formulas, 224–225
 legacy array formulas, 249
 PivotTables, 609–610
 Power Pivot data connections, 745–746
 scenarios, 666
 templates, 172
 text, 102–103
 elapsed time, calculating, 324–325
 embedded charts. *See also* charts
 converting to a chart sheet, 415
 copying, 416

- embedded charts (*continued*)
 deleting, 416
 modifying type, 973–974
 moving, 414–415
 overview, 403–404
 printing, 418–419
 renaming, 418
 resizing, 414–415
- empty strings, filling in, 808–809
- Encrypt Document dialog box, 714
- Encrypt with Password command, 714
- Enter mode
 changing to Edit mode, 36
 overview, 31
- EOMONTH function
 returning fiscal months, 321
 returning last day of month, 317–318
- equal sign (=)
 in condition arguments, 330
 as formula operator, 213
 precedence, 214
- Equation Editor, 516–517
- Equation Tools, Design, 516–517
- Error Alert tab, Data Validation dialog box, 555, 565
- error bars, adding, 468
- error values, 29
- errors
 formula errors
 AutoCorrect, 396–398
 blank cells not blank, 381–382
 categories, 379–380
 cells filled with hash marks, 381
 circular references, 391
 decimal precision problems, 387–388
 #DIV/0! errors, 237, 383
 extra space characters, 382
 forcing recalculation, 387
 mismatched parentheses, 380–381
 #NAME? errors, 237, 384
 #N/A errors, 237, 383
 #NULL! errors, 237, 384
 operator precedence, 386–387
 phantom links, 388
 #REF! errors, 237, 385
 ripple effect, 237
 tracing cell relationships, 389–390
 tracing error values, 383, 391
 #VALUE! errors, 237, 386
 printing error values, 382
 returned by lookup functions, hiding, 353–354
- ETL (extraction, transformation, loading), 777
- Evaluate Formula dialog box, 392–393
- events
 application events, 949, 959–962
 workbook events, 949, 951–955
 worksheet events, 949, 955–959
- Excel
 new chart types, 431–439
 uses for, 3
- Excel Options dialog box
 adding icons to Quick Access toolbar, 16–17
 Advanced tab, 34
 Customize Ribbon tab, 205–207
 Quick Access Toolbar tab, 202–203
- exclamation point (!)
 in link formula syntax, 578
 when referencing cells in other workbooks, 228
 when referencing cells in other worksheets, 228
- exponential notation, 363
- exponential smoothing, 701
- exponential smoothing tool, Analysis ToolPak, 701
- exponentiation (^) operator, 213, 214, 342
- exporting
 data, 550–551
 UI customizations, 204
- extracting
 characters from text string, 290–291
 filenames from URLs, 541–542
- F**
- fast Fourier transformations, 702
- File
 Account, About Excel, 863
 Export, Create PDF/XPS Document, 717
- Info
 Check for Issues, 167, 718
 Compatibility Mode, 168
 Edit Links to Files, 388
 Manage Workbook, 165
 Protect Workbook, 166–167
- Open
 Open as Copy, 174
 overview, 158, 522, 523
 Recent, 158, 159

- Options, Trust Center, 160, 161, 172
- Print
- overview, 177, 418
 - Printer, 184
 - Quick Print, 177–178
 - Settings, 183–184
- Save, 162–163
- Save As, 162–163
- File button, Excel window, 5
- files
- file extensions
 - filter options, 522
 - workbooks, 4
- file formats
- database, 522–523
 - HTML, 523–524
 - spreadsheet, 522
 - text, 523
 - XML, 524
- filenames
- extracting from URLs, 541–542
 - filtering, 160–161
 - rules, 163
- importing
- vs. opening files, 524
 - text files, 525–528
- opening, importing files *vs.*, 524
- Fill Color, 131–132
- fill handle, 22, 37
- fill text alignment, 128
- Filter button
- filtering tables, 114–116
 - filtering workbook names, 99
 - overview, 111–112
 - sorting tables, 111–112
 - toggling header row drop-down buttons, 119
- Filter buttons, 109
- FILTER function, 262–265
- filtering
- filenames, 160–161
 - PivotTables
 - with slicers, 635–637
 - with a timeline, 637–638
 - records, 264–265
 - table data, 111–117
- finalizing workbooks, 717–718
- Find and Replace dialog box
- dates, finding, 303
 - finding formatting, 395
 - overview, 81–83
 - replacing information, 394–395
 - replacing text in cells, 547–548
- FIND function
- finding a particular character, 291–292
 - finding second instance of a character, 292–293
- fiscal month, returning, 320–321
- fiscal quarter, calculating, 319–320
- fixed-value error bars, 468
- Flash Fill command
- creating data, 536
 - splitting data, 533–536
- flat files, loading data into Power Pivot data model, 738–742
- FLOOR function, 280
- FollowHyperlink worksheet event, 955
- Font Color, 131–132
- Font tab, Format Cells dialog box, 125
- fonts
- color, 131–132
 - document themes, customizing, 153–155
 - formatting worksheets, 124–126
- footers, adding to reports, 192–194
- Format Axis task pane, 456–460
- Format Cells dialog box
- Alignment tab, 127–131
 - Border tab, 132–134
 - custom number formatting, 497–498
 - displaying, 124
 - Font tab, 125
 - formatting worksheets, 124
 - number formatting, 45–47
- Format Cells option, Protect Sheet dialog box, 712
- Format Columns option, Protect Sheet dialog box, 712
- Format Comment dialog box, 106
- Format Comment option, 106
- Format Painter button, 134
- Format Rows option, Protect Sheet dialog box, 712
- Format task pane, 417, 444–445
- formatting
- avoiding pasting, 33
 - calculated columns, 759–760
 - charts
 - copying formatting, 417–418, 451
 - overview, 413–414, 416–417
 - conditional formatting
 - absolute references, 141–142

- formatting (*continued*)
 applying, 135
 color scales, 137–138
 copying cells containing, 145–146
 data bars, 135–137
 deleting, 146
 examples, 142–144
 formula-based rules, 139–142
 graphical formats, 135–139
 icon sets, 138–139
 locating cells that contain, 146
 managing rules, 145
 overview, 135
 relative references, 141–142
 selecting cells containing, 146
finding and replacing, 395
icons, 506–507
numbers
 automatic, 43
 custom number formatting, 47, 492–498
 Format Cells dialog box, 45–47, 491–492
 hiding zeros, 495–496
 number format strings, 492–498
 numbers in millions, 495
 numbers in thousands, 494–495
 overview, 4
 shortcut keys for, 43–44, 490–491
 suppressing zeros, 495–496
 in text strings, 297–299
 using Ribbon buttons, 43, 489–490
PivotTables, 607–608
shapes, 506–507
worksheet formatting
 borders and lines, 132–134
 colors and shading, 131–132
 document themes, 150–155
 fonts, 124–126
 with the Format Cells dialog box, 124
 with the Home tab, 122
 with the Mini toolbar, 123
 named styles, 146–150
 overview, 121–122
 text alignment, 127–131
Forms controls, 936
Formula bar
 cell edit mode, entering, 224
 function entry tips, 223–224
 overview, 5
 using as calculator, 244
formula errors
 AutoCorrect, 396–398
 blank cells not blank, 381–382
 categories, 379–380
 cells filled with hash marks, 381
 circular references, 391
 decimal precision problems, 387–388
 #DIV/0! errors, 237, 383
 extra space characters, 382
 forcing recalculation, 387
 mismatched parentheses, 380–381
 #NAME? errors, 237, 384
 #N/A errors, 237, 383
 #Null! errors, 237, 384
 #NUM! errors, 237, 384–385
 operator precedence, 386–387
 phantom links, 388
 #REF! errors, 237, 385
 ripple effect, 237
 tracing cell relationships, 389–390
 tracing error values, 383, 391
 #VALUE! errors, 237, 386
Formula Evaluator, 392–393
formulas
 @ symbol, 255
 cell references
 to cells outside the worksheet, 228–229
 changing types, 227–228
 overview, 225–229
 converting
 to text, 245
 to values, 245
 copying, 244–245
 for data validation
 accepting dates by day of week, 562
 accepting larger value than previous cell, 561
 accepting nonduplicate entries only, 561
 accepting text beginning with specific
 character, 561–562
 accepting text only, 561
 accepting value not exceeding a total, 563
 creating dependent lists, 563–564
 overview, 558–559

- editing, 224–225
 entering
 manually, 218–219
 overview, 218–219
 by pointing, 220
 using Formula AutoComplete, 219–220
 entering on tablets, 49
 formula-based rules, 139–142
 functions in, 216–218, 221–223
 hard-coding values, 244
 line breaks, adding, 295–296
 for lookups (*See* lookup formulas)
 named formulas, 240–241
 naming techniques, 239–243
 operators in, 213–216
 overview, 29, 211–212
 range names, pasting, 220–221
 selecting cells containing, 80
 in tables, 232–236
 variables, 229–232
 viewing, 389
- Formulas**
- Defined Names
 - Create from Selection, 242
 - Define Name, 240
 - Name Manager, 388
 - Formula Auditing
 - 391–392, 435–436
 - Error Checking, 383
 - Evaluate Formula, 392–393
 - Remove Arrows, 383, 390
 - Show Formulas, 389
 - Trace Dependents, 390
 - Trace Precedents, 390
 - Watch Window, 65–66
 - Function Library
 - Date & Time, 307
 - Insert, 221
 - Insert Function, 901
 - Formulas option, Go To Special dialog box, 80
 - Formulas tab, 9, 49
 - For-Next loops, 893–894
 - Fourier analysis tool, 702
 - Fraction number format, 47
 - free-floating text, adding to charts, 449
 - Freeze Panes command, 63–65
 - freezing panes, 63–65
 - frequency distributions, creating with PivotTables, 626–627
 - Frontline Systems, 695
 - F-test tool, Analysis ToolPak, 701–702
 - full outer joins, 846
 - full-screen mode, 199
 - Function Arguments dialog box, 221, 222–223
 - Function Library, 221
 - functions
 - arguments, 217–218
 - conversion functions, 283–284
 - creating
 - with no arguments, 905–906
 - with one arguments, 906–907
 - with a range argument, 908–909
 - with two arguments, 907–908
 - as wrapper for built-in functions, 909–910
 - Cube functions, 773–775
 - custom functions
 - calling from procedures, 904
 - creating, 900
 - debugging, 910
 - inserting, 910–912
 - using in worksheet formulas, 904–905
 - date and time functions, 307–327
 - dynamic array
 - FILTER, 262–265
 - overview, 255
 - RANDARRAY, 259–260
 - SEQUENCE, 260–261
 - SORT, 256–257
 - SORTBY, 257–258
 - UNIQUE, 258–259
 - XLOOKUP, 265–269
 - in formulas, 216–218
 - inserting into formulas, 221–223
 - lookup functions, 348–363
 - text functions, 286–299
 - VBA functions, 874–875
 - funnel charts, 437–438
 - fuzzy matching, 851–853

G

- General number format, 46
 general text alignment, 127
 Generate GetPivotData command, 640

- GETPIVOTDATA function, 639–640
Go to Special dialog box, 79–81, 388–389
Goal Seek, 675–678
goals
 calculating percent of, 271–272
 Goal Seek, 675–678
graphical conditional formats
 color scales, 137–138
 data bars, 135–137
 icon sets, 138–139
graphics
 Artistic Effects, 516
 background images, adding, 516
 color scale conditional formatting, 137–138
 data bar conditional formats, 135–137
 Equation Editor and, 516–517
 equations, inserting as graphics objects, 516–517
 file formats, 515–516
 icon set conditional formats, 138–139
 linked pictures, creating, 510–512
 screenshots, inserting, 516
 shapes (*See* shapes)
 SmartArt, 513–514
 Sparklines (*See* Sparklines)
 SVG (Scalable Vector Graphics) icons, inserting, 504
 Unicode characters, 498–502
 WordArt, 514–515
Graphics Tools, Format
 Accessibility, 506
 Change Graphic, 506
 Graphics Styles, 506
 Play 3D, 506
 Size, 506
graphs. *See* charts
greater than (>) operator
 in condition arguments, 330
 as formula operator, 213, 214
 in workbook names, 163
greater than or equal to (>=) operator, 213
gridlines
 adding, 452
 modifying, 452
 overview, 406
 printing, 189–190
 removing, 452
 simulating using borders, 189–190
 toggling, 133
Group By command, Power Query, 831–832, 839
Group mode, 78–79
grouping PivotTable items
 automatically, 621–625
 by date, 621–624
 manually, 619–621
 by time, 625
GroupName property, OptionButton control, 945
- ## H
- hard-coding values in formulas, 244
hash (#) marks
 cells filled with, 47, 69, 381
 in formula errors, 236
header row, 108, 109
headers
 adding to reports, 192–194
 column headers, printing, 190
 row headers
 overview, 108, 109
 printing, 190
 toggling drop-down buttons, 118
Help system, VBA, 963
Help tab, 10
HEX2DEC function, 538
hexadecimal values, converting to decimal, 538
Hidden and Empty Cell Settings dialog box, 468
hiding
 calculated columns, 760–762
 columns, 70–71
 data series, 461–462
 Mini toolbar, 123
 Ribbon, 54
 rows, 70–71
 shape objects, 505
 status bar, 54
 worksheets, 60
 zeros, 495–496
Highlight Cell Rules, 139
histogram charts, 431–433
histogram tool, Analysis ToolPak, 703–704
HLOOKUP function
 defined, 348
 looking up values horizontally, 352–353
 looking up values using closest match, 354–357
Home tab

- Alignment
 - Merge & Center, 130–131
 - Orientation, 131
- Cells, Insert, 111
- Conditional Formatting
 - Manage Rules, 367
 - New Rule, 367–368, 370, 371, 373, 375
- Editing
 - Clear, 33, 146
 - Fill, 88
 - Find & Select, 74, 79, 81
 - Sort & Filter, 113, 116
- Font
 - Borders, 132
 - Fill Color, 131–132
 - Font Color, 131–132
 - overview, 125
- Font Size, 125
- formatting tools, 122
- Number, 43, 45
- overview, 9
- Styles
 - Cell Styles, 146–150
 - Conditional Formatting, 135–146
- horizontal alignment options, 127–128
- horizontal axis, 405, 406
- horizontal scrollbar, 5, 6
- hot keys, 19
- HOUR function, 323–324
- hour(s)
 - adding to a time, 326–327
 - decimal hours, converting, 326
 - returning, 323–324
 - rounding time values, 325
- HTML files, 523–524, 551
- hyperlinks, deleting from cells, 33

- I
- icons
 - assigning icon sets, 138–139
 - formatting, 506–507
 - Quick Access toolbar, 200–201
 - attaching macros to, 203
 - rearranging, 203
 - removing, 204
 - SVG (Scalable Vector Graphics) icons, inserting, 504
- IF function
 - avoiding divide-by-zero errors, 278
 - calculating elapsed time, 324–325
 - checking multiple conditions, 330–331, 333–336
 - checking simple conditions, 329–330
 - defined, 348
 - overview, 229–230
 - [if_empty] argument, 262
- IFERROR function
 - checking for any error condition, 383
 - defined, 348
- If-Then construct, 893
- Ignore Integer Constraints option, Solver Options dialog box, 686
- Image ActiveX control, 944
- importing
 - data
 - database files, 522–523
 - filling gaps, 545–547
 - HTML files, 523–524
 - vs. opening files, 524
 - overview, 521
 - splitting text, 531–536
 - spreadsheet files, 522
 - text files, 523, 525–528
 - XML files, 524
 - UI customizations, 204
- incorrect reference errors, 379
- Increase Decimal button, 490
- INDEX function
 - defined, 348
 - looking up exact values, 351–352
 - looking up last value in column, 362–363
 - looking up value using closest match, 356–357
 - looking up values based on two-way matrix, 359–361
- INDIRECT function, validating conditional data, 331–332
- indirect precedents, 389
- infographic widgets, creating, 509–510
- inner joins, 846
- Input Message tab, Data Validation dialog box, 555
- input messages, displaying, 564
- Input ranges, histograms and, 703–704
- Input X Range option, Regression dialog box, 707
- Input Y Range option, Regression dialog box, 707
- InputBox function, 914–915, 968–969

- Insert
 Charts
 PivotChart, 640
 Recommended Charts, 419
 Filters
 Slicer, 636
 Timeline, 637
 Icons, 504
 Illustrations
 Online Pictures, 515
 Pictures, 515
 Screenshot, 516
 SmartArt, 513
 Module, 922
 Shapes, 502
 Symbols, Equation, 516–517
Tables
 Recommended PivotTables, 600–602
 Table, 110
Text
 Signature Line, 719
 Text Box, 285
 WordArt, 514
UserForm, 919
Insert Cells command, 67
Insert Columns option, Protect Sheet dialog box, 712
Insert dialog box, 67–68
Insert Function command, 221, 901
Insert Function dialog box, 221–223, 901
Insert Hyperlinks option, Protect Sheet dialog box, 712
Insert Rows option, Protect Sheet dialog box, 712
Insert Sheet Columns command, 67
Insert Sheet Rows command, 67
Insert Slicer command, 116–117
Insert Slicers dialog box, 636
Insert Statistic Chart command, 433
Insert tab, 9
insertion points, when editing cells, 34–35
Inspect Document command, 718
inspecting workbooks, 718–719
installing
 add-ins, 983
 Analysis ToolPak, 698
 Solver, 679
Integer data type, 977
IntegralHeight property
 ListBox control, 944
 TextBox control, 947
interface, tablets, 48–49
internal data model
 feeding, 747–754
 managing queries and connections, 755
 managing relationships, 754–755
intersection operator, 241–243
invisible digital signatures, 720
ISERR function, 354
ISERROR function, 354
ISNA function, 354
ISSOD function, 559–560
ISTEXT function, 561
italic, applying, 126
Iterations option, Solver Options dialog box, 686
- J**
- joining
 columns, 540–541
 Power Query queries, 845–853
 text strings, 286–287
Julian date format, converting to, 315–316
justified text, 128, 129
- K**
- keyboard
 accessing the Ribbon, 12–14
 navigating dialog boxes, 18–1920
 navigating worksheets, 7–8
 number formatting, 43–44
 selecting chart elements, 443
 using within task panes, 20
- L**
- Label ActiveX control, 944
labels
 adding to data series, 465–467
 axes labels, positioning, 456
 category labels, orienting, 459–460
 dynamic labels, creating, 510
landscape mode, printing in, 185
LargeChange property, ScrollBar control, 946
Last Cell option, Go To Special dialog box, 80
leap year bug, 306

- Left and Top property, 940
 left anti joins, 846
LEFT function
 converting text to sentence case, 288–289
 extracting parts of text string, 290–291, 814
 overview, 561
 padding numbers with zeros, 297
 splitting text, 531
left outer joins, 846
left text alignment, 127
legacy array formulas
 downside of, 250
 editing, 249
 overview, 247–248
legends
 adding to charts, 416, 450
 chart legends, 450–451
 moving, 450
 overview, 406
 selecting individual elements within, 443, 450
 using shapes in lieu of, 451
LEN function, counting specific characters
 in a cell, 294–295
less than or equal to (\leq), 213
less than ($<$) symbol
 as formula operator, 213, 214
 in workbook names, 163
LET function, 230
line breaks
 adding to charts, 295–296
 removing, 38
line charts, 424–426
Line Sparkline, 476
linear equations, solving with Solver, 686–688
lines, adding, 132–134
linked pictures, creating, 510–512
LinkedCell property
 CheckBox control, 942
 ComboBox control, 942–943
 ListBox control, 944
 OptionButton control, 945
 overview, 940
 ScrollBar control, 946
 SpinButton control, 946
linking tables to Power Pivot, 725–733
linking workbooks
 changing link source, 583
 changing startup prompt, 582
 creating external reference formulas, 578–580
 creating links to unsaved workbooks, 580–581
 intermediary links issue, 585
 modifying source workbook, 584–585
 moving source workbook, 584
 opening workbook with links, 581–582
 overview, 577–578
 phantom links, 388
 renaming source workbook, 584
 Save As command issue, 584
 severing links, 583
 updating links, 582–583
ListBox ActiveX control, 944
ListFillRange property, 940
 ComboBox control, 942–943
 ListBox control, 944
ListRows property, ComboBox control, 943
ListStyle property
 ComboBox control, 943
 ListBox control, 944
Locked attribute, 710–712
locking formula cells in template files, 175
logical arguments, 334–335
logical errors, 379
logical operators, 213
Long data type, 977
lookup formulas
 for exact value based on any lookup
 column, 351–352
 for exact value based on left lookup
 column, 348–350
 hiding errors, 353–354
 for last value in a column, 362–363
 for values based on multiple criteria, 361–362
 for values based on two-way matrix, 359–361
 for values by closest match, 354–357
 for values from multiple tables, 357–359
 for values horizontally, 352–353
LOOKUP function
 defined, 348
 returning last value, 363
lookup tables, classifying values, 538–540
LOWER function
 converting text to lowercase, 536
 converting text to sentence case, 288–289
lowercase, converting text to, 288–289, 536

M

- M language, 785–786
m time unit code, DateDIF function, 309
macro recorder indicator, 5, 6
macros
 adding to Quick Access toolbar, 203, 886–887, 931–932
 assigning to buttons, 885–886, 931
 assigning to shortcut keys, 885
 creating
 for ActiveX controls, 941–942
 by recording macros, 876–887
 by writing VBA code, 887–896
Developer tab, displaying, 870
Function procedures, 874–875
listing, 203
overview, 869
saving workbooks containing, 872
security of, 871–872
speeding, 975–978
storing in Personal Macro Workbook, 884–885
Sub procedures, 873–874
uses, 870
major tick marks, 456
Manage Relationships dialog box, 754–755
Manage Rules command, 145
manual what-if analysis, 653
map charts, 438–439
margins
 adjusting, 186–187
 aligning headers/footers, 194
 default settings, 17
 when using multiline headers/footers, 194
whitespace, hiding, 180
Mark as Final command, 718
marking workbooks as final, 717–718
MATCH function
 defined, 348
 looking up exact value based on any lookup column, 351–352
 looking up value using closest match, 356–357
 lookup up values based on two-way matrix, 359–361
matching text in a list, 542–543
Max Feasible Solutions option, Solver Options dialog box, 686
Max property
 ScrollBar control, 946
 SpinButton control, 946
Max Subproblems option, Solver Options dialog box, 686
Max Time option, Solver Options dialog box, 686
maximizing windows, 54
MaxLength priority, TextBox control, 947
md time unit code, DateDIF function, 309
.mdb file extension, 522
means
 covariance, 701
 F-tests, 701, 702
 T-tests, 707–708
 two-sample tests for, 707, 708
 Z-tests, 708
Merge Across option, Merge & Center control, 130
Merge Cells option, Merge & Center control, 130
Merge & Center command, 12
Merge & Center control, 130–131
Merge Columns command, 809–811
merging
 cells, 129–130
 Power Query queries, 845–853
 scenarios, 666
metadata, assigning to workbooks, 166
methods (of objects), 891
Microsoft Office Signature Line command, 719–70
MID function
 extracting parts of text string, 290–291, 814
 splitting text, 531
Min property
 ScrollBar control, 946
 SpinButton control, 946
Mini toolbar
 customizing chart elements, 446
 formatting controls, 123
 hiding, 123
 overview, 16
minimizing windows, 54
minor tick marks, 456
minus (-), 574
MINUTE function, 323–324
minute(s)
 adding to a time, 326–327
 decimal minutes, converting, 326

- returning, 323–324
 rounding time values, 325
 mismatched parentheses, 380–381
 mixed references, 225–227
MOD function
 alternate-row shading, displaying, 143–144
 checkerboard shading, creating, 144
 shading groups of rows, 144
modal dialog boxes, 18
modeless dialog boxes, 18
Monospaced option, 800
MONTH function, 313–314
month(s)
 between dates, calculating, 314–315
 last date of, returning, 317–318
 last weekday of, calculating, 322–323
 returning, 313–314
mouse
 scrolling the wheel, 8–9
 selecting chart elements, 442–443
 zooming worksheets, 9
Move or Copy dialog box, 58–59
moving
 charts, 414–415
 legends, 450
 Plot Area, 448
 Quick Access toolbar, 200
 ranges, 83–84, 967
 windows, 54
 worksheets, 58–59
moving average tool, Analysis ToolPak, 704–705
moving averages, 704–705
MsgBox function, 915–918
MultiLine property, TextBox control, 947
multiple selection, 74
multiplication (*) operator, 213, 214, 918
Multiply option, Paste Special dialog box, 94
MultiSelect property, ListBox control, 944
- ## N
- Name box, 96
name box, Excel window, 5, 6
Name Manager, 97–98
Name property, 940
named styles
 applying, 147
 controlling with templates, 150
 creating, 149
 merging from other workbooks, 150
 modifying, 147–149
 overview, 146–147
named views, creating, 197–198
#NAME? errors, 237, 384
names
 constants, 240
 formulas, 240–241
 range names
 deleting, problems with, 384
 pasting into formulas, 220–221
 sheet tabs, 57–58
 worksheets, 57–58
#N/A errors, 237, 383
negation (-) operator, 386, 387
negative values, fixing trailing minus signs, 549
NETWORKDAYS function, 310
NETWORKDAYS.INTL function, 310–311
New Formatting Rule dialog box
 creating formula-based rules, 139
 customizing colors, 138
 customizing icon sets, 139
 highlighting values based on dates, 372–377
 highlighting values in one list and another, 371–372
 highlighting values in one list but not
 another, 369–371
 overview, 136–137
New Name dialog box, 96–97
New Sheet button, 5, 6
New Sheet control, adding worksheets to
 workbooks, 57
NewSheet workbook event, 952, 954
non-numerics data, working with, 617–619
Normal Probability option, Regression dialog box,
 707
Normal view, printing in, 179–180
not equal to (<>) operator, 213
notes
 adding to cells, 104–108
 alternatives to, 107
 deleting, 108
 editing, 108
 formatting, 106–107
 showing, 105–106

NOW function, 308

#NULL! errors, 237, 384

Num Lock key, 8

number bases, converting between, 538

Number number format, 46

numbers

 formatting

 automatic, 43

 custom number formatting, 47, 492–498

 Format Cells dialog box, 45–47, 491–492

 hiding zeros, 495–496

 number format strings, 492–498

 numbers in millions, 495

 numbers in thousands, 494–495

 overview, 4, 489

 shortcut keys for, 43–44, 490–491

 suppressing zeros, 495–496

 in text strings, 297–299

 using Ribbon buttons, 43, 489–490

identified as text, 286

padding with zeros, 297

quarter numbers, calculating, 318–319

random number generation tool, 705–706

rounding

 to nearest penny, 279–280

 overview, 278

 to significant digits, 280–282

 using formulas, 279

#NUM! errors, 237, 384–385

0

Object data type, 977

objects

 Application, 888–889

 collections, 889

 Excel object model, 888–889

 methods, 891

 properties, 889–891

 qualified object references, 964

 Range, 892

 selecting, 80

 simplifying references, 976

Objects option, Go To Special dialog box, 80

OCT2DEC function, 538

octal values, converting to decimal, 538

ODBC connections, connecting to Power Query, 797

.ods file extension, 551

OFFSET function, 348

 calculating dynamic range, 487

 converting vertical data to horizontal, 544–545

OLAP databases, connecting to Power Query, 796

one-click printing, 177

OnKey application event, 959, 961–962

Online 3D Models dialog box, 504–505

OnTime application event, 959, 960–961

Open As Copy command, 174

Open dialog box, 522

Open workbook event, 952–953

OpenDocument Spreadsheet, exporting to, 551

opening

 add-ins, 980–981

 files, importing files vs., 524

 workbooks

 automatically, 161

 with external reference formulas, 581–582

Open event, 952–953

 overview, 158–161

 requiring password, 714–715

 Workbook_Open procedure, 952–953

operators

 defined, 212

 in formulas, 213–214

 logical, 213–214

 precedence, 386–387

 reference operators, 241–243

OptionButton ActiveX control, 945

OR function, 336

order argument, SERIES function, 465

outlines

 adding data, 575

 creating automatically, 572

 creating manually, 572–574

 displaying levels, 574

 overview, 567–570

 vs. PivotTables, 570

 preparing data for, 571

 removing, 575

P

padding data, 297

Page Break Preview view, 181–183

page breaks

- inserting, 187–188
- Normal view, 179–180
- Page Break Preview view, 181–183
- removing, 188
 - when changing page orientation, 185
- Page Layout tab, 9
 - Page Setup
 - Background, 190, 516
 - Breaks, 187–188
 - Margins, 186–187
 - Orientation, 185
 - Print Area, 185
 - Print Titles, 188
 - Size, 185
 - Scale to Fit, 189
 - Sheet Options
 - Gridlines, 189–190
 - Headings, 190
 - Themes
 - Colors, 154–155
 - Fonts, 153–155
 - Save Current Theme, 155
 - Themes, 152–153
- Page Layout view, 180–181
- Page Setup dialog box
 - Margins tab, 187
 - Preview box, 187
 - print settings, adjusting, 183–191
 - Sheet tab, 189
- page view buttons, 5
- parentheses (()), 273, 380–381
- Pareto charts, 433–434
- passwords
 - adding to workbooks, 165–166, 714, 715
 - limitations, 709
 - for VBA projects, 716
 - when assigning user permissions, 713
- Paste Special command, 92–95, 529, 587–588
- .pdf file extension, 551
- PDF files
 - creating, 198
 - exporting data to Power Query, 795–796
 - saving worksheets as, 717
- peekaboo tabs, 507–508
- penny, rounding to nearest, 279–280
- Percent Style button, 43, 490
- percentage error bars, 468
- Percentage number format, 47
- percentages
 - applying percent increase/decrease to values, 276–277
- calculating
 - overview, 271
 - percent distribution, 274–275
 - percent of goal, 271–272
 - percent variance, 272–273
 - percent variance with negative values, 273–274
 - running total, 275–276
 - divide-by-zero errors, 277–278
- permissions, assigning to worksheets, 713
- Personal Macro Workbook, 884–885
- PgDn key, 7, 8
- PgUp key, 7, 8
- phantom link errors, 388
- Picture property, 940, 944
- Picture Tools, Format, Adjust, 516
- PictureSizeMode property, Image control, 944
- pie charts, 426–427
- Pivot Column command, 821–823
- PivotChart Tools Analyze
 - Actions, Move Chart, 643
 - Show/Hide, 642
- PivotCharts
 - creating, 640–644
 - tips for using, 643–644
- pivoting columns, 821–823
- PivotTable Options dialog box, 608
- PivotTable Tools Analyze
 - Calculations
 - Fields, & Items, 628–634
 - OLAP Tools, 650
 - Data, Change Data Source, 624
 - Design, PivotTable Styles, 607–608
 - Filter, Insert Slicer, 649
 - Group
 - Group Field, 619
 - Group Selection, 619
 - Tools, PivotChart, 640, 643
- PivotTable Tools Design, Layout, Grand Totals, 633
- PivotTables
 - adding fields, 610
 - converting into Cube functions, 773–775

- PivotTables (*continued*)
 copying contents, 611
 creating automatically, 600–602
 creating based on data model, 644–650
 creating manually
 formatting options, 607–608
 number format, changing, 607–608
 setting up layout, 603–605
 specifying data, 602–603
 specifying location for, 603
 data appropriate for, 598–600
 examples, 596–598, 611–616
 filtering data
 with slicers, 635–637
 with a timeline, 637–638
 frequency distributions, creating, 626–627
 grouping items
 automatically, 621–625
 by date, 621–624
 manually, 619–621
 by time, 625
 modifying, 609–610
 vs. outlines, 570
 overview, 595
 referencing cells, 638–640
 removing fields, 610
 reverse, 635
 summary techniques, 608–609
 terminology, 606–607
 working with non-nuemrics data, 617–619
- plot area, 407, 448
- plus sign (+), in worksheet outlines, 574
- portrait mode, printing in, 185
- Power Pivot
 calculated measures
 creating, 770–772
 deleting, 773
 editing, 773
 overview, 770–772
 compatibility, 724
 data model
 adding tables, 727–728
 feeding, 747–754
 loading data from flat files, 738–742
 managing queries and connections, 755
 managing relationships, 754–755
 preparing tables, 726
 referencing fields from other tables, 768–770
 refreshing, 742–745
 using Cube functions in, 773–775
- DAX (Data Analysis Expression) functions
 building DAX-driven calculated
 columns, 762–770
 nesting, 770
 referencing fields from other tables, 768–770
 editing data connections, 745–746
 linking Excel tables, 725–733
 loading data
 from flat files, 738–742
 from relational databases, 733–738
 Manage, 733
 refreshing data, 742–745
 Ribbon interface, activating, 724
- Power Query
 aggregating data, 830–832
 appending data, 841–845
 basics of, 777–792
 blank fields, filling in, 808
 columns
 column-level actions, 788–790
 concatenating, 809–811
 creating custom columns, 823–830
 pivoting, 821–823
 selecting, 857
 splitting, 816–819
 unpivoting, 819–821
 copying queries, 859–860
 creating reference tables, 859
 data profiling with, 800–803
 duplicate records, removing, 805–807
 empty strings, filling in, 808–809
 getting query information, 855–856
 grouping data, 830–832
 importing data
 from database systems, 796–797
 from files, 793–796
 from other data systems, 797–798
 overview, 793–796
 M language, 785–786
 managing data source settings, 798–799
 managing queries, 787–788
 merging queries, 845–853

- organizing queries, 856–857
- performance issues, 862–865
- preventing automatic data type changes, 861–862
- productivity tips, 855–862
- query steps
 - managing, 785
 - overview, 784–785
 - renaming, 857–858
 - reusing, 837–841
 - viewing, 784
- refreshing data, 786–787
- saving queries, 782
- setting default load behavior, 860–861
- table-level actions, 790–792
- text
 - changing case, 811
 - cleaning and trimming, 812–813
 - extracting first and last characters, 815
 - extracting left, right, middle values, 814–815
 - extracting middle characters, 816
 - finding and replacing, 811–812
- viewing query dependencies, 860
- Power Query Editor
 - column-level actions, 788–790
 - M code, viewing and editing, 785–786
 - query steps, viewing, 784
 - table-level actions, 790–792
- precedents, selecting, 80
- Precedents option, Go To Special dialog box, 80
- primary key, 754
- Print Preview, 17
- printing
 - adding headers/footers to reports, 192–194
 - background images, 190
 - cell gridlines, 189–190
 - charts, 418–419
 - choosing printers, 184
 - column headers, 190
 - column titles, 188–189
 - copying page settings, 195
 - custom views, creating, 197–198
 - default settings, 17
 - error values, 382
 - named views, creating, 197–198
 - Normal view, 179–180
 - number of copies, specifying, 186
 - one-click printing, 177
 - Page Break Preview view, 181–183
 - page breaks (*See* page breaks)
 - Page Layout view, 180–181
 - page margins, adjusting, 186–187
 - page orientation, changing, 185
 - page settings, adjusting, 183–191
 - page setup settings, 183–191
 - paper size, specifying, 185
 - PDF files, creating, 198
 - preventing cells, 195–196
 - preventing objects, 196–197
 - previewing, 178
 - row headers, 190
 - row titles, 188–189
 - scaling output, 189
 - specifying what to print, 184–185
 - watermarks, inserting, 190–191
 - worksheets, 25–26
 - XPS files, 198
- PRN files
 - exporting to, 551
 - overview, 523
- Project Properties dialog box, 716–717
- PROPER function, 288, 536
- properties (of objects)
 - Application object, 890
 - overview, 889–891
 - Range object, 892
- properties, ActiveX controls
 - adjusting, 938–939
 - CheckBox, 942
 - ComboBox, 942–943
 - CommandButton, 943
 - ListBox control, 944
 - OptionButton, 945
 - ScrollBar, 945–946
 - shared, 939–940
 - SpinButton, 946
 - TextBox, 946–947
- Protect Sheet command, 710–711
- Protect Sheet dialog box, 710, 712
- Protect Structure and Windows dialog box, 716
- Protect Workbook command
 - Add a Digital Signature, 720
 - Encrypt with Password, 714

Protect Workbook command (*continued*)

 Mark as Final, 718

 overview, 716

Protected View, 160

protecting

 VBA Projects, 716–717

 workbooks

- adding digital signatures, 719–720
- inspecting workbooks, 718–719
- marking as final, 717–718
- overview, 60, 714–716
- worksheets, 60, 709, 710–713

Publish as PDF or XPS dialog box, 198, 717

Q

qualified object references, 964

quarter(s), calculating

- calendar quarter, 318–319
- fiscal quarter, 319–320

queries. *See* Power Query

Queries & Connections

- grouping queries, 856–857
- managing queries, 787–788
- refreshing data, 786–787

question mark (?)

- in condition arguments, 339
- searching for, 82
- as a wildcard, 268
- in workbook names, 163
- in worksheet names, 58

? (question mark) wildcard, 268

Quick Access toolbar

 customizing

- adding commands, 16–17, 201–202
- overview, 16–17, 199–204
- rearranging icons, 203
- removing icons, 204
- resetting to default state, 204
- sharing customizations, 204
- for specific workbooks, 203

 default tools, 16–17, 200

 displaying below the Ribbon, 16

 full-screen mode, 199

 macros, adding, 886–887, 931

 moving, 200

 overview, 5, 6

Redo button, 16–17

 tools included, 200

 Undo command, 16–17

Quick Access Toolbar tab, Excel Options dialog box, 202–203

Quick Print, 177–178

quote ("), in workbook filenames, 163

R

radar charts, 429–430

RAND function, 541

RANDARRAY function, 259–260

Random Number Generation dialog box, 705–706

random number generation tool, Analysis ToolPak, 705–706

Range object, 892

ranges. *See also* cells

 copying

 to adjacent cells, 88

 to the Office Clipboard, 89–91

 to other worksheets, 89

 overview, 83–84, 964–965

 paste options, 91–95

 using drag-and-drop, 87

 using Paste Options buttons, 86–87

 using Ribbon commands, 84

 using shortcut keys, 85–86

 using shortcut menu commands, 85

 using the Paste Special dialog box, 92–95

 using VBA, 964–965

 counting

 selected cells, 971

 values in, 282–283

 defined, 73

 determining type of selection, 970

 icon sets, assigning, 138–139

 looping through, 967–968

 moving, 83–84, 967

 named styles, 146–150

 names

 advantages, 95

 creating, 97–98

 deleting, problems with, 384

 managing, 97–98

 overview, 95

 pasting into formulas, 220–221

- prompting for cell values, 968–969
 range intersections, 241–243
 reference operators, 241–243
 selecting
 - complete rows and columns, 75
 - multisheet ranges, 76–79
 - noncontiguous ranges, 75–76
 - overview, 74
 - by searching, 81–83
 - special types of cells, 79–81
 selecting row/column, 966
 selecting to end of row/column, 966
 rank and percentile tool, Analysis ToolPak, 706
 rearranging
 - columns, 541
 - Ribbon tabs, 207
 Recent command, 158, 159
 Recommended Charts command, 419
 Recommended PivotTables command, 600–602
 recording macros
 - absolute vs. relative recording, 880–881
 - adding macros to Quick Access toolbar, 886–887
 - assigning macros to buttons, 885–886
 - assigning macros to shortcut keys, 885
 - editing macros, 879–880, 883
 - example, 877–878, 881
 - overview, 876–887
 - recording actions, 876
 - running macros, 881
 - storing macros in Personal Macro Workbook, 884–885
 - testing macros, 879, 883–884
 - viewing code, 878–879, 882, 884
 records, filtering, 264–265
 Redo button, 16–17
 reference tables, creating, 859
 referencing spill ranges, 254–255
 #REF! errors, 237, 385
 refreshing
 - Power Pivot data, 742–745
 - Power Query data, 786–787
 - worksheets, 592–593
 regional settings, 32
 Regression dialog box, 706–707
 regression tool, Analysis ToolPak, 706–707
 relational databases
 - connecting to Power Query, 796
 - loading data into Power Pivot data model, 733–738
 relative references, 141–142
 relative references, naming formulas, 225–227, 241
 Remove Duplicates command, 529, 805–807
 renaming
 - charts, 418
 - sheet tabs, 57–58
 Replace Values command, 789, 808, 809, 811
 replacing
 - cell contents, 33
 - formatting, 395
 - text
 - in cells, 547–548
 - overview, 394–395
 reports
 - filling gaps, 545–547
 - scenario summary reports, creating, 666–667
 Reset to Match Style command, 447
 Residuals option, Regression dialog box, 707
 resizing
 - charts, 414–415
 - Plot Area, 448
 - tables, 110
 - windows, 54
 resizing handle, 110
 restoring
 - Ribbon, 208
 - windows, 54
 return_type argument
 - WEEKDAY function, 314
 - WEEKNUM function, 314
 reverse PivotTables, 635
 Review tab
 - overview, 9
 - Proofing, Spelling, 396, 547
 Protect
 - Allow Edit Ranges, 713
 - Protect Sheet, 710
 - Protect Workbook, 716
 - Unprotect Sheet, 710
 Ribbon
 - accessing with keyboard, 12–14
 - chart elements, customizing, 446
 - commands
 - adding to groups, 207
 - adding to Quick Access toolbar, 16–17
 - copying cells/ranges, 84

- Ribbon (*continued*)
customizing
 adding commands to groups, 207
 creating new groups, 206
 creating new tabs, 206
 Customize Ribbon tab, Excel Options dialog box, 205–207
 reasons for, 205
 restoring to default state, 208
 sharing customizations, 204
 what can be customized, 205
hiding, 54
overview, 5, 6, 9
Power Pivot Ribbon interface, activating, 724
tabs
 contextual, 11
 creating, 206
 Developer tab, displaying, 870, 936
 overview, 9–11
 rearranging, 207
Ribbon (Ribbon), 49–51
Ribbon Display Options control, 6
right anti joins, 846
RIGHT function
 converting dates to Julian format, 315–316
 converting text to sentence case, 288–289
 extracting parts of text string, 290–291, 814
 splitting text, 531
right outer joins, 846
right text alignment, 128
ROUND function
 and decimal precision problems, 387
 rounding to significant digits, 280–282
 rounding to specified number of digits, 279
ROUNDDOWN function, 325
rounding
 numbers
 to the nearest penny, 279–280
 overview, 278
 to significant digits, 280–282
 using formulas, 279
 time values, 325
ROUNDUP function
 calculating calendar quarters, 319
 rounding time values, 325
Row Differences option, Go To Special dialog box, 80
ROW function, displaying alternate-row shading, 143–144
Row Height dialog box, 70
rows
 banded, 108, 109, 119
 data body, 109
 deleting, 68
 duplicate
 identifying, 530–531
 removing, 529–530
 header row, 108, 109
 headers, printing, 190
 height, changing, 68, 69–70
 hiding, 70–71
 highlighting based on a value, 142–143
 inserting, 66–68, 111
 orientation, switching, 410
 randomizing, 541
 row numbers, 5, 6
 selecting, 66
 titles, printing, 188–189
 total row, 108, 109–110
 unhiding, 70–71
rules, formula-based, 139–142
running totals, calculating, 275–276

S

- sampling tool, Analysis ToolPak, 707
Save As command, 162–163, 954
Save command, 162–163
saving
 Power Query queries, 782
workbooks
 as add-in files, 983, 987
 AutoRecover, 163–165
 BeforeSave event, 954
 containing macros, 872
 containing volatile functions, 56
 overview, 26, 161–163
 as template, 175–176
 using VBA, 972
worksheets
 in PDF format, 198, 717
 in XPS format, 198, 717
scaling printed output, 189

scatter grams, 427–428
 scatter plots, 427–428
Scenario Manager
 adding scenarios, 663–664
 defining scenarios, 662–664
 displaying scenarios, 664–665
 generating summary reports, 666–667
 merging scenarios, 666
 modifying scenarios, 666
Scenario Values dialog box, 664
Scientific number format, 47
 screen updating, disabling, 975
 screenshots, inserting, 516
ScreenTips, 201
Scroll Lock key, 7
ScrollBar ActiveX control, 945–946
ScrollBars property, *TextBox control*, 945–946, 947
 scrolling sheet tabs, 56
Search box, 5, 6, 14
SECOND function, 323–324
second(s)
 adding to a time, 326–327
 decimal seconds, converting, 326
 returning, 323–324
security
 macro security, 871–872
 Protected View, 160
Select Case construct, 895
Select Locked Cells option, *Protect Sheet dialog box*, 712
Select Unlocked Cells option, *Protect Sheet dialog box*, 712
selected cell indicator, 5, 6
SelectionChange worksheet event, 955, 958–959
semantic errors, 379
semicolon (;)
 in custom number formats, 47
 as list separator, 218
 in number formatting strings, 492
 when suppressing zeros, 496
sentence case, converting text to, 288–289
SEQUENCE function, 260–261
serial numbers
 dates, 301–302
 times, 303–304
SERIES formula, editing, 464–465
series_name argument, *SERIES function*, 464
Set Precision as Displayed option, 46
shading
 checkerboard shading, creating, 144
 displaying alternate-row shading, 143–144
 formatting worksheets, 131–132
 groups of rows, 144
shapes
 formatting, 506–507
 hiding, 505
 infographic widgets, creating, 509–510
 inserting, 502–503
 layering, 509
 linked pictures, creating, 510–512
 peekaboo tabs, 507–508
 preventing from printing, 196
 selecting, 505
Shapes gallery, 502–503
sheet tab buttons, 5, 6
sheet tabs
 color, changing, 58
 overview, 5, 6
 renaming, 57–58
 scrolling, 56
SheetActivate workbook event, 952, 953
SheetBeforeRightClick workbook event, 952
SheetBefpreDoubleClick workbook event, 952
SheetChange workbook event, 952
SheetDeactivate workbook event, 952
SheetSelectionChange workbook event, 952
shortcut keys
 assigning macros to, 885
 copying cells/ranges, 85–86
 formatting fonts, 126
 for number formatting, 490–491
 Quick Access toolbar commands, 204
 for VBA macros, 876–877
shortcut menus
 BeforeRightClick event, 959
 copying cells/ranges, 85, 86
 overview, 15–16
Show Formulas, 389
Show Iteration Results option, *Solver Options dialog box*, 686
Show Whitespace option, 800

- Shrink to Fit, 129
- shrinking text, 129
- Signature Setup dialog box, 719–720
- simple buttons, 12
- Single data type, 977
- single-cell goal seeking, 675–678
- single-factor analysis of variance, 698, 699–700
- sizes argument, SERIES function, 465
- Skip Blanks option, Paste Special dialog box, 94
- slash (/)
 - in workbook names, 163
 - in worksheet names, 58
- slicers
 - filtering PivotTables, 635–637
 - filtering tables, 116–117
- SmallChange property
 - ScrollBar control, 946
 - SpinButton control, 946
- SmartArt
 - overview, 513–514
 - preventing from printing, 196
- SmartArt Tools
 - Design
 - Create Graphic, 514
 - Layouts, 514
 - SmartArt Styles, 514
 - Format, 514
- smoothing data
 - exponential smoothing, 701
 - using moving averages, 704–705
- Solver
 - examples
 - allocating resources, 691–693
 - maximizing profits, 679–684
 - minimizing shipping costs, 688–691
 - optimizing investment portfolio, 693–695
 - solving linear equations, 686–688
 - installing, 679
 - options, 685–686
 - overview, 675, 678–679
 - problems appropriate for, 679
 - starting, 680
 - uses, 679
- Solver Options dialog box, 685–686
- Solver Parameters dialog box, 680–682
- Solver Results dialog box, 683
- SORT function, 256–257
- Sort option, Protect Sheet dialog box, 712
- SORTBY function, 257–258
- sorting
 - table data, 111–117
 - tables, 111–117
- source workbooks. *See also* linking workbooks
 - changing link source, 583
 - determining for a particular workbook, 584
 - Edit Links dialog box listing of, 581–582
 - links to nonexistent documents, 580
 - links to unsaved documents, 580–581
 - modifying, 584–585
 - moving, 584
 - renaming, 584
 - Save As command issue, 584
 - updating linked formulas, 582–583
 - updating links when source workbook is unavailable, 581
- spacebar, effect of pressing in cells, 381
- spaces
 - extra spaces in formulas, 382
 - identifying values containing, 382
 - removing from text, 289–290
- Sparklines
 - auto-updating, 486
 - changing axis scaling, 482–483
 - changing colors, 481
 - changing line width, 481
 - changing type, 481
 - creating, 477–479
 - displaying for dynamic range, 486–487
 - faking reference lines, 483–484
 - handling hidden or missing data, 480
 - highlighting data points, 481–482
 - sizing cells, 480
 - specifying data axis, 484–486
 - types of, 475–476
- Special number format, 47
- spell-checking, 396, 547
- Spelling dialog box, 396, 547
- spill ranges
 - overview, 252–253
 - referencing, 254–255
- #SPILL! error, 380, 385
- spin buttons, 12

- SpinBox ActiveX control, 946
 Split again command, 63
 split button controls, 12
 split buttons, 12
 Split command, 63
 splitting text
 Flash Fill command, 533–536
 overview, 531–536
 Text to Columns command, 532–533
 splitting worksheets into panes, 63
 spreadsheet file formats, 522
 SQL Server, loading data into Power Pivot data model, 733–738
 square brackets ([])
 in column names, 233, 236
 in link formula syntax, 578
 when applying custom format colors, 496
 in workbook names, 228
 in worksheet names, 58
 standard deviation error bars, 468
 standard error error bars, 468
 Start_Num argument, FIND function, 292–293
 statistics
 analysis of variance, 699–700
 correlation, 700–701
 descriptive, 701, 702
 F-tests, 701, 702
 status bar
 hiding, 54
 overview, 5, 6
 stock charts, 431
 StrConv function, 913
 strikethrough, applying, 126
 String data type, 977
 Style property, ComboBox control, 943
 Sub procedures, 873–874
 SUBSTITUTE function
 counting characters in text string, 294–295
 replacing text in cells, 548
 substituting text strings, 293–294
 substituting text, 293–294
 subtraction (-) operator, 213, 214, 387
 SUM function
 calculating percent distribution, 274–275
 overview, 563
 running total, calculating, 275–276
 summing all values meeting condition, 336–339
 SUMIF function, 336–341
 SUMIFS function, 339–340, 341, 342–344
 SUMPRODUCT function, 362
 sunburst charts, 436
 surface charts, 430
 SVG (Scalable Vector Graphics) icons, inserting, 504
 SYLK (Symbolic Link) files, 523, 551
 symbols, in worksheet outlines, 575–576
 syntax errors
 AutoCorrect and, 381
 mismatched parentheses, 380
 overview, 379, 380

T

- tab list, 5, 6
 tabbed dialog boxes, 19–20
 Table Import Wizard, loading data
 from external Excel files, 739–741
 from relational databases, 733–738
 from SQL Server, 733–738
 from text files, 741–742
 Table Styles command, 117–118
 Table Tools Design
 Table Style Options, 118–119
 Table Styles, 117–118
 Tools, Insert Slicer, 116–117
 tables
 adding data, 111
 components, 108–110
 creating, 110–111
 filtering data, 111–117, 114–117
 formulas within, 232–236
 linking to Power Pivot, 725–733
 overview, 108
 referencing data, 235–236
 resizing, 110
 sorting data, 111–117
 style, changing, 117–120
 summarizing data, 232–234
 tablets
 entering formulas, 49
 interface, 48–49
 overview, 47–48
 Ribbon, Draw tab, 49–51
 tabs
 color, changing, 58

- tabs (*continued*)
 contextual, 11
 creating, 206
 in dialog boxes, 19–20
 rearranging, 207
- task panes, 20
- templates
 chart templates, creating, 473–474
 custom, 174–176
 default, 172–174
 modifying, 172
 viewing, 169–170
- TETXJOIN function
 joining columns, 540–541
 joining text strings, 287
- text
 adding to cells, 548–549
 alignment, changing, 127–131
 case, changing, 288–289, 536, 811
 cleaning, 296–297, 812–813
 converting formulas to, 245
 counting specific characters, 294–295
 as data type, 28
 displaying at angle, 131
 editing in comments, 102–103
 extracting characters, 290–291, 815–816
 finding second instance of a character, 292–293
 finding specific characters, 291–292, 811–812
 formatting numbers in, 297–299
 free-floating, adding to charts, 449
 joining strings, 286–287
 line breaks, adding, 295–296
 matching in a list, 542–543
 numbers identified as, 286
 padding with zeros, 297
 removing from cells, 547–548
 removing spaces, 289–290
 replacing, 394–395
 in cells, 547–548
 shrinking, 129
 splitting, 531–536
 substituting, 293–294
 trimming, 812–813
 Unicode characters, 498–502
 wrapping, 129
- text boxes, creating, 285
- text files
 exporting data to Power Query, 795
 exporting to, 550–551
 formats, 523
 loading data into Power Pivot data model, 741–742
- TEXT function
 formatting numbers in text strings, 297–299
 returning fiscal months, 321
- Text Import Wizard, 523, 525–528
- Text number format, 47
- text strings, 212
- Text to Columns command, 532–533
- TextAlign property, TextBox control, 947
- TextBox ActiveX control, 946–947
- ThisWorkbook module, 950
- 3D models, inserting, 504–505
- tick marks, 456
- tilde (`) wildcard, 268
- TIME function, 326–327
- Time number format, 47
- time-based category axis, 459
- timelines, filtering PivotTables, 637–638
- time(s)
 current, returning, 308
 custom number formatting, 497–498
 elapsed, calculating, 324–325
 entering, 304–305
 formatting, 305–306
 hour, returning, 323–324
 minute, returning, 323–324
 rounding time values, 325
 second, returning, 323–324
- time-stamping worksheets, 41
- title bar, 5, 6
- title case, converting text to, 288–289, 536
- titles
 chart titles, 449
 column titles, 188–189
 row titles, 188–189
- TODAY function, 308
- toggle buttons, 12
- ToggleButton ActiveX control, 947
- Tools, References, 904
- top text alignment, 129
- Top/Bottom Rules, 139
- total row, 108, 109–110

- touch-related commands, 7
 Trace Error command, 383
 tracing
 cell relationships, 389–390
 error values, 391
 formula error values, 383
 trailing minus signs, fixing, 549
 Transform tab (Power Query)
 Capitalize Each Word, 811
 Clean, 813
 Data Type, 826
 Extract, 814–815
 Group By, 831–832
 Pivot Column, 821–823
 Trim, 813
 Transpose option, Paste Special dialog box, 94
 treemap charts, 437
 trendlines, adding, 468–470
TRIM function
 identifying values containing spaces, 382
 removing spaces from text, 289–290, 537
 trimming text, 812–813
 Trust Center, 160, 161, 172
 T-test tool, Analysis ToolPak, 707–708
 Tufte, Edward, 477
 two-factor analysis of variance, 699–700
 two-sample tests
 for means, 707, 708
 for variance, 701–702, 707
 TXT files
 exporting to, 550–551
 overview, 523
- U**
- UI customizations
 importing/exporting, 204
 Quick Access toolbar, 204
 Ribbon, 204–208
 underline, applying, 126
 Undo command, 16–17
 Unhide dialog box, 60
 unhiding
 columns, 70–71
 rows, 70–71
 worksheets, 60
 Unicode characters, 498–502
- UNIQUE function, 258–259
 unlocking cells, 710–712
 Unmerge Cells option, Merge & Center control, 130
 Unpivot Columns command, 819–820
 Unpivot Other Columns command, 821
 unpivoting columns, 819–821
 Unprotect Sheet command, 710
UPPER function, 288–289, 536–537
 uppercase, converting text to, 288–289, 536
 URLs, extracting filenames, 541–542
 Use AutoFilter option, Protect Sheet dialog box, 712
 Use Automatic Scaling option, Solver Options dialog box, 686
 Use PivotTable & PivotChart option, Protect Sheet dialog box, 712
 user interface, customizing
 Quick Access toolbar, 204
 Ribbon, 204–208
 sharing customizations, 204
 user-defined data type, 977
UserForms
 accelerator keys, adding, 932–933
 ActiveX controls
 adding, 919–920
 changing properties of, 920–921
 tab order, 933
 add-ins example, 983–988
 alternatives, 914–918
 closing, 924
 creating, 918–923, 923–924, 926–928
 displaying, 922–923, 930
 event handling, 921–922, 925–926, 928–930
 inserting in VBE, 919
 reasons for creating, 913–914
 testing, 924–925, 931
 user-level permissions, assigning to worksheets, 713
- V**
- Validation option, Paste Special dialog box, 93
 value axes
 in bubble charts, 430
 in combination charts, 470–471
 defined, 406
 modifying, 452–456
 in stock charts, 431
 Value property, 940

- values
 applying a percent increase or decrease to, 276–277
 calculating percent variance with negative values, 273–274
 classifying, 538–540
 converting, 538
 converting formulas to, 245
 counting in a range, 282–283
 entering into worksheets, 30–32
 error values, 237
 graphical conditional format, 135–139
 hard-coding, 244
 number formatting, 42–47
 numeric, 27–28
 summing, 24
values argument, SERIES function, 464–465
#VALUE! errors, 237, 386
variables
 declaring data types, 976–978
 formulas, 229–232
 overview, 892
variances
 covariance tool, 701
 F-tests, 701, 702
 percent variance, 272–273
 T-tests, 707–708
 two-sample tests for, 701–702, 707
 Z-tests, 708
Variant data type, 977
VBA
 chart-related tasks, 973–975
 comments, 888
 data types, 976–978
 Excel object model, 888–889
 functions (*See* functions)
 Help system, 963
 inserting modules, 888
 macros (*See* macros)
 objects
 Application, 888–889, 890
 collections, 889
 Excel object model, 888–889
 methods, 891
 properties, 889–891
 qualified object references, 964
 Range, 892
 simplifying references, 976
 programming constructs, 893–896
 protecting VBA Projects, 716–717
range-related tasks
 copying ranges, 964–965
 counting selected cells, 971
 determining type of selection, 970
 looping through ranges, 967–968
 moving ranges, 967
 overview, 963–971
 prompting for cell values, 968–969
 selecting rows/columns, 966
 selecting to end of row/column, 966
variables, 892
workbook-related tasks, 972–973
VBE (Visual Basic Editor)
 overview, 941–942
 protecting VBA projects, 716–717
vector-based images, 515
vertical alignment options, 129
vertical axes. *See* value axes
vertical bar (|)
 as insertion point, 34
 in workbook names, 163
vertical data, converting to horizontal, 543–545
vertical scrollbar, 5, 6
View Side by Side feature, 63
View tab
 overview, 9
 Properties Window, 920
 Show, Gridlines, 133
 Toolbox, 919
 Window
 Freeze Panes, 63–65
 Reset Window Position, 63
 Split, 63
 Synchronous Scrolling, 63
 View Side by Side, 63
Visible Cells Only option, Go To Special dialog box, 81
visible digital signatures, signing workbooks, 719–720
Visible property, 940
visualizations. *See* graphics
VLOOKUP function
 defined, 348
 looking up exact values, 348–350
 looking up value using closest match, 354–357
volatile functions, saving workbooks with, 56

W

Watch Window, 65–66
 waterfall charts, 434, 435
 watermarks, inserting, 190–191
 week, number in year, returning, 313–314
 WEEKDAY function, 313–314, 562
 weekday(s)
 last of month, calculating, 322–323
 Nth weekday of month, calculating, 321–322
 number from given date, extracting, 313–314
 weekend days, identifying, 142
 WEEKNUM function, 313–314
 what-if analysis
 data tables
 limitations, 661
 one-input data tables, 654–657
 overview, 653–654
 two-input data tables, 657–661
 example, 651–652
 manual approach, 653
 Scenario Manager
 adding scenarios, 663–664
 defining scenarios, 662–664
 displaying scenarios, 664–665
 generating summary reports, 666–667
 merging scenarios, 666
 modifying scenarios, 666
 overview, 661
 Width and Height property, 940
 wildcards, XLOOKUP function with, 268–269
 window controls, 5, 6
 WindowActivate workbook event, 952
 WindowDeactivate workbook event, 952
 windows
 closing, 55–56
 maximizing, 54
 minimizing, 54
 moving, 54
 resizing, 54
 restoring, 54
 states, 54
 switching among, 55
 Windows Clipboard, 85
 Win/Loss Sparkline, 476
 With-End With construct, 894–895
 WordArt, 514–515
 WordWrap property, TextBox control, 947

workbook events
 BeforeClose, 954–955
 BeforeSave, 954
 NewSheet, 954
 Open, 952–953
 overview, 949, 951–955
 SheetActivate, 953
 Workbook_BeforeClose procedure, 954–955
 Workbook_BeforeSave procedure, 954
 Workbook_NewSheet procedure, 954
 Workbook_Open procedure, 952–953
 workbooks
 activating sheets, 953
 adding worksheets, 56–57, 954
 closing
 BeforeClose event, 954–955
 overview, 168
 using VBA, 972
 Workbook_BeforeClose procedure, 954–955
 converting to add-ins, 983
 creating
 introductory example, 20–26
 keyboard shortcut for, 158
 from Start screen, 157–158
 from a template, 170–171
 using Blank Workbook command, 20, 157, 158
 exporting data to Power Query, 794–795
 file extension, 4
 finalizing, 717–718
 inspecting, 718–719
 linking
 changing link source, 583
 changing startup prompt, 582
 creating external reference formulas, 578–580
 creating links to unsaved workbooks, 580–581
 intermediary links issue, 585
 modifying source workbook, 584–585
 moving source workbook, 584
 opening workbook with links, 581–582
 overview, 577–578
 phantom links, 388
 renaming source workbook, 584
 Save As command issue, 584
 severing links, 583
 updating links, 582–583
 marking as final, 717–718
 opening

- workbooks (*continued*)
 - automatically, 161
 - with external reference formulas, 581–582
- Open event, 952–953
- overview, 158–161
- requiring password, 714–715
- Workbook_Open procedure, 952–953
- overview, 4
- protecting
 - adding digital signatures, 719–720
 - overview, 60, 709, 714–716
- saving
 - as add-in files, 983, 987
 - AutoRecover, 163–165
 - BeforeSave event, 954
 - containing volatile functions, 56
 - overview, 26, 161–163
 - as template, 175–176
 - using VBA, 972
 - workbooks containing macros, 872
- WORKDAY.INTL function, 311–313
- worksheet events
 - BeforeRightClick, 959
 - Change, 955, 956
 - overview, 949, 955–959
 - SelectionChange, 958–959
- worksheet functions, 212
- worksheet outlines
 - adding data, 575
 - adjusting symbols, 575
 - creating automatically, 572
 - creating manually, 572–574
 - displaying levels, 574
 - hiding symbols, 575–576
 - overview, 567–570
 - vs. PivotTables, 570
 - preparing data for, 571
 - removing, 575
- Worksheet_BeforeRightClick procedure, 959
- Worksheet_Change procedure, 955, 956
- worksheets
 - activating, 56
 - adding to workbooks, 56–57
 - comparing, 62–63
 - consolidating
 - example, 590–592
- overview, 585–594
- refreshing consolidation, 592–593
- sources, 593–594
- using external reference formulas, 587
- using the Consolidate dialog box, 588–589
- using the Paste Special dialog box, 587–588
- copying, 58–59
- data entry
 - AutoComplete, 37–38
 - AutoCorrect, 39
 - AutoFill, 37
 - data forms, 40–41
 - dates, 31–32
 - Enter mode, 31
 - numbers, 36, 40
 - numeric values, 30
 - text, 30–31
 - times, 32
- date-stamping, 41
- deleting, 57
- filling gaps, 545–547
- formatting
 - borders and lines, 132–134
 - colors and shading, 131–132
 - conditional formatting, 135–146
 - document themes, 150–155
 - fonts, 124–126
 - with the Format Cells dialog box, 124
 - with the Home tab, 122
 - with the Mini toolbar, 123
 - named styles, 146–150
 - overview, 121–122
 - text alignment, 127–131
- hiding, 60
- moving, 58–59
- naming, 57–58
- navigating, 4–9
- navigating cells, 183–191
- number formatting
 - automatic, 43
 - custom formats, 47
 - Format Cells dialog box, 45–47
 - overview, 42–47
 - Ribbon controls, 43
 - shortcut keys, 43–44
- overview, 4

printing, 25–26
protecting, 60, 709, 710–713
saving

 in PDF format, 198, 717
 in XPS format, 198, 717

tab color, changing, 58

time-stamping, 41

unhiding, 60

views

- comparing worksheets, 62–63
- creating, 61
- freezing panes, 63–65
- monitoring cells with Watch Window, 65–66
- overview, 60–61
- splitting worksheets into panes, 63
- using multiple views, 61–62

zooming

- with the mouse, 9
- overview, 60–61
- Zoom control, 6

`Worksheet_SelectionChange`
procedure, 958–959

Wrap Text, 129

X

XLA files, 522
XLM files, 522
`XLOOKUP` function, 265–269
XLS files, 522
.xlsm file extension, 873
XLSTART folder, workbooks stored in, 161
.xlsx file extension, 4

XLT files, 522
XML files, 524
XML Paper Specification Document (XPS) files
 overview, 551
 saving worksheets as, 198, 717
xVelocity analytics engine, 72
XY charts, 427–428

Y

y time unit code, `DATEDIF` function, 309
yd time unit code, `DATEDIF` function, 309
`YEAR` function, 313–314
`YEARFRAC` function, 316–317
year(s)
 between dates, calculating, 314–315
 extracting year from given date, 313–314
 fiscal quarters, calculating, 319–320
 percent completed/remaining, calculating, 316–317
 returning, 313–314
`ym` unit time code, `DATEDIF` function, 309

Z

zeros
 divide-by-zero errors, 277–278
 hiding, 495–496
 padding numbers with, 297
 suppressing, 495–496
Zoom control, 5, 6
zooming worksheets
 with mouse, 9
 Zoom control, 6
`Z-test` tool, Analysis ToolPak, 708

WILEY END USER LICENSE AGREEMENT

Go to www.wiley.com/go/eula to access Wiley's ebook EULA.